The

Intel Microprocessors

8086/8088, 80186/80188, 80286, 80386, 80486 Pentium, Pentium Pro
Processor, Pentium II, Pentium 4, and Core2 with 64-bit Extensions

Architecture, Programming, and Interfacing

EIGHTH EDITION

IJP–LC300–01
9950D

Barry B. Brey

THE INTEL MICROPROCESSORS

8086/8088, 80186/80188, 80286, 80386, 80486, Pentium, Pentium Pro Processor, Pentium II, Pentium III, Pentium 4, and Core2 with 64-Bit Extensions

Architecture, Programming, and Interfacing

Eighth Edition

BARRY B. BREY

PEARSON

Prentice
Hall

Editor in Chief: Vernon Anthony
Acquisitions Editor: Wyatt Morris
Editorial Assistant: Christopher Reed
Production Coordination: GGS Book Services
Project Manager: Jessica Sykes
Operations Specialist: Laura Weaver
Design Coordinator: Mike Fruhbeis
Cover Designer: Ilze Lemesis
Cover image: iStockphoto
Director of Marketing: David Gesell
Marketing Manager: Jimmy Stephens
Marketing Assistant: Les Roberts

This book was set in Times by GGS Book Services. It was printed and bound by Hamilton
Printing. The cover was printed by Phoenix Color Corp.

Pearson Prentice Hall™ is a trademark of Pearson Education, Inc.
Pearson® is a registered trademark of Pearson plc
Prentice Hall® is a registered trademark of Pearson Education, Inc.

Pearson Education Ltd., London   Pearson Education Australia Pty. Limited
Pearson Education Singapore Pte. Ltd.   Pearson Education North Asia Ltd., Hong Kong
Pearson Education Canada, Inc.   Pearson Educación de Mexico, S.A. de C.V.
Pearson Education—Japan   Pearson Education Malaysia Pte. Ltd.

PEARSON
Prentice
Hall

This text is dedicated to my progenies, Brenda (the programmer) and Gary (the veterinarian technician), and to my constant four-legged companions: Romy, Sassy, Sir Elton, Eye Envy, and Baby Hooter.

This page intentionally left blank

PREFACE

This practical reference text is written for students who require a thorough knowledge of programming and interfacing of the Intel family of microprocessors. Today, anyone functioning or striving to function in a field of study that uses computers must understand assembly language programming, a version of C language, and interfacing. Intel microprocessors have gained wide, and at times exclusive, application in many areas of electronics, communications, and control systems, particularly in desktop computer systems. A major addition to this eighth edition explains how to interface C/C++ using Visual C++ Express, which is a free download from Microsoft, with assembly language for both the older DOS and the Windows environments. Many applications include Visual C++ as a basis for learning assembly language using the inline assembler. Updated sections that detail new events in the fields of microprocessors and microprocessor interfacing have been added.

ORGANIZATION AND COVERAGE

To cultivate a comprehensive approach to learning, each chapter begins with a set of objectives that briefly define its content. Chapters contain many programming applications and examples that illustrate the main topics. Each chapter ends with a numerical summary, which doubles as a study guide, and reviews the information just presented. Questions and problems are provided for reinforcement and practice, including research paper suggestions.

This text contains many example programs using the Microsoft Macro Assembler program and the inline assembler in the Visual C++ environment, which provide a learning opportunity to program the Intel family of microprocessors. Operation of the programming environment includes the linker, library, macros, DOS function, BIOS functions, and Visual C/C++ program development. The inline assembler (C/C++) is illustrated for both the 16- and 32-bit programming environments of various versions of Visual C++. The text is written to use Visual C++ Express 2005 or 2008 as a development environment, but any version of Visual Studio can also be used with almost no change.

This text also provides a thorough description of family members, memory systems, and various I/O systems that include disk memory, ADC and DAC, 16550 UART, PIAs, timers, keyboard/display controllers, arithmetic coprocessors, and video display systems. Also discussed are

the personal computer system buses (AGP, ISA, PCI, PCI Express, USB, serial ports, and parallel port). Through these systems, a practical approach to microprocessor interfacing can be learned.

APPROACH

Because the Intel family of microprocessors is quite diverse, this text initially concentrates on real mode programming, which is compatible with all versions of the Intel family of micro-processors. Instructions for each family member, which include the 80386, 80486, Pentium, Pentium Pro, Pentium II, Pentium III, and Pentium 4 processors, are compared and contrasted with those for the 8086/8088 microprocessors. This entire series of microprocessors is very sim-ilar, which allows more advanced versions and their instructions to be learned with the basic 8086/8088. Please note that the 8086/8088 are still used in embedded systems along with their updated counterparts, the 80186/80188 and 80386EX embedded microprocessor.

This text also explains the programming and operation of the numeric coprocessor, MMX extension, and the SIMD extension, which function in a system to provide access to floating-point calculations that are important in control systems, video graphics, and computer-aided design (CAD) applications. The numeric coprocessor allows a program to access complex arithmetic operations that are otherwise difficult to achieve with normal microprocessor pro-gramming. The MMX and SIMD instructions allow both integer and floating-point data to be manipulated in parallel at very high speed.

This text also describes the pin-outs and function of the 8086–80486 and all versions of the Pentium microprocessor. First, interfacing is explained using the 8086/8088 with some of the more common peripheral components. After explaining the basics, a more advanced emphasis is placed on the 80186/80188, 80386, 80486, and Pentium through Pentium 4 microprocessors. Coverage of the 80286, because of its similarity to the 8086 and 80386, is minimized so the 80386, 80486, and Pentium versions can be covered in complete detail.

Through this approach, the operation of the microprocessor and programming with the advanced family members, along with interfacing all family members, provides a working and practical background of the Intel family of microprocessors. Upon completing a course using this text, you will be able to:

1. Develop software to control an application interface microprocessor. Generally, the software developed will also function on all versions of the microprocessor. This software also includes DOS-based and Windows-based applications. The main emphasis is on developing inline assembly and C++ mixed language programs in the Windows environment.
2. Program using MFC controls, handlers, and functions to use the keyboard, video display system, and disk memory in assembly language and C++.
3. Develop software that uses macro sequences, procedures, conditional assembly, and flow control assembler directives that are linked to a Visual C++ program.
4. Develop software for code conversions using lookup tables and algorithms.
5. Program the numeric coprocessor to solve complex equations.
6. Develop software for the MMX and SIMD extensions.
7. Explain the differences between the family members and highlight the features of each member.
8. Describe and use real and protected mode operation of the microprocessor.
9. Interface memory and I/O systems to the microprocessor.
10. Provide a detailed and comprehensive comparison of all family members and their software and hardware interfaces.
11. Explain the function of the real-time operating system in an embedded application.
12. Explain the operation of disk and video systems.
13. Interface small systems to the ISA, PCI, serial ports, parallel port, and USB bus in a personal computer system.

CONTENT OVERVIEW

Chapter 1 introduces the Intel family of microprocessors with an emphasis on the microprocessor-based computer system: its history, operation, and the methods used to store data in a microprocessor-based system. Number systems and conversions are also included. Chapter 2 explores the programming model of the microprocessor and system architecture. Both real and protected mode operations are explained.

Once an understanding of the basic machine is grasped, Chapters 3 through 6 explain how each instruction functions with the Intel family of microprocessors. As instructions are explained, simple applications are presented to illustrate the operation of the instructions and develop basic programming concepts.

Chapter 7 introduces the use of Visual C/C++ Express with the inline assembler and separate assembly language programming modules. It also explains how to configure Visual C++ Express for use with assembly language applications.

After the basis for programming is developed, Chapter 8 provides applications using the Visual C++ Express with the inline assembler program. These applications include programming using the keyboard and mouse through message handlers in the Windows environment. Disk files are explained using the File class, as well as keyboard and video operations on a personal computer system through Windows. This chapter provides the tools required to develop virtually any program on a personal computer system through the Windows environment.

Chapter 9 introduces the 8086/8088 family as a basis for learning basic memory and I/O interfacing, which follow in later chapters. This chapter shows the buffered system as well as the system timing.

Chapter 10 explains memory interface using both integrated decoders and programmable logic devices using VHDL. The 8-, 16-, 32-, and 64-bit memory systems are provided so the 8086–80486 and the Pentium through Pentium 4 microprocessors can be interfaced to memory.

Chapter 11 provides a detailed look at basic I/O interfacing, including PIAs, timers, the 16550 UART, and ADC/DAC. It also describes the interface of both DC and stepper motors.

Once these basic I/O components and their interface to the microprocessor are understood, Chapters 12 and 13 provide detail on advanced I/O techniques that include interrupts and direct memory access (DMA). Applications include a printer interface, real-time clock, disk memory, and video systems.

Chapter 14 details the operation and programming for the 8087–Pentium 4 family of arithmetic coprocessors, as well as MMX and SIMD instructions. Today few applications function efficiently without the power of the arithmetic coprocessor. Remember that all Intel microprocessors since the 80486 contain a coprocessor; since the Pentium, an MMX unit; and since the Pentium II, an SIMD unit.

Chapter 15 shows how to interface small systems to the personal computer through the use of the parallel port, serial ports, and the ISA, and PCI bus interfaces.

Chapters 16 and 17 cover the advanced 80186/80188–80486 microprocessors and explore their differences with the 8086/8088, as well as their enhancements and features. Cache memory, interleaved memory, and burst memory are described with the 80386 and 80486 microprocessors. Chapter 16 also covers real-time operating systems (RTOS), and Chapter 17 also describes memory management and memory paging.

Chapter 18 details the Pentium and Pentium Pro microprocessors. These microprocessors are based upon the original 8086/8088.

Chapter 19 introduces the Pentium II, Pentium III, Pentium 4, and Core2 microprocessors. It covers some of the new features, package styles, and the instructions that are added to the original instruction set.

Appendices are included to enhance the text. Appendix A provides an abbreviated listing of the DOS INT 21H function calls because the use of DOS has waned. It also details the use of

the assembler program and the Windows Visual C++ interface. A complete listing of all 8086–Pentium 4 and Core2 instructions, including many example instructions and machine coding in hexadecimal as well as clock timing information, is found in Appendix B. Appendix C provides a compact list of all the instructions that change the flag bits. Answers for the even-numbered questions and problems are provided in Appendix D.

To access supplementary materials online, instructors need to request an instructor access code. Go to **www.pearsonhighered.com/irc**, where you can register for an instructor access code. Within 48 hours after registering, you will receive a confirming e-mail, including an instructor access code. Once you have received your code, go to the site and log on for full instructions on downloading the materials you wish to use.

Acknowledgments

I greatly appreciate the feedback from the following reviewers:
   James K. Archibald, Brigham Young University
   William H. Murray III, Broome Community College.

STAY IN TOUCH

We can stay in touch through the Internet. My Internet site contains information about all of my textbooks and many important links that are specific to the personal computer, microprocessors, hardware, and software. Also available is a weekly lesson that details many of the aspects of the personal computer. Of particular interest is the "Technical Section," which presents many notes on topics that are not covered in this text. Please feel free to contact me at bbrey@ee.net if you need any type of assistance. I usually answer all of my e-mail within 24 hours.

My website is **http://members.ee.net/brey**

BRIEF CONTENTS

CONTENTS

CHAPTER 1

Introduction to the Microprocessor and Computer

INTRODUCTION

This chapter provides an overview of the Intel family of microprocessors. Included is a discussion of the history of computers and the function of the microprocessor in the microprocessor-based computer system. Also introduced are terms and jargon used in the computer field, so that **computerese** is understood and applied when discussing microprocessors and computers.

The block diagram and a description of the function of each block detail the operation of a computer system. Blocks, in the block diagram, show how the memory and input/output (I/O) system of the personal computer interconnect. Detailed is the way data are stored in the memory so each data type can be used as software is developed. Numeric data are stored as integers, floating-point, and binary-coded decimal (BCD); alphanumeric data are stored by using the ASCII (**American Standard Code for Information Interchange**) code and the Unicode.

CHAPTER OBJECTIVES

Upon completion of this chapter, you will be able to:

1. Converse by using appropriate computer terminology such as bit, byte, data, real memory system, protected mode memory system, Windows, DOS, I/O, and so forth.
2. Briefly detail the history of the computer and list applications performed by computer systems.
3. Provide an overview of the various 80X86 and Pentium family members.
4. Draw the block diagram of a computer system and explain the purpose of each block.
5. Describe the function of the microprocessor and detail its basic operation.
6. Define the contents of the memory system in the personal computer.
7. Convert between binary, decimal, and hexadecimal numbers.
8. Differentiate and represent numeric and alphabetic information as integers, floating-point, BCD, and ASCII data.

1–1   A HISTORICAL BACKGROUND

This first section outlines the historical events leading to the development of the microprocessor and, specifically, the extremely powerful and current 80X86,[1] Pentium, Pentium Pro, Pentium III, Pentium 4,[2] and Core2 microprocessors. Although a study of history is not essential to understand the microprocessor, it furnishes interesting reading and provides a historical perspective of the fast-paced evolution of the computer.

The Mechanical Age

The idea of a computing system is not new—it has been around long before modem electrical and electronic devices were developed. The idea of calculating with a machine dates to 500 BC when the Babylonians, the ancestors of the present-day Iraqis, invented the **abacus**, the first mechanical calculator. The abacus, which uses strings of beads to perform calculations, was used by the ancient Babylonian priests to keep track of their vast storehouses of grain. The abacus, which was used extensively and is still in use today, was not improved until 1642, when mathematician Blaise Pascal invented a calculator that was constructed of gears and wheels. Each gear contained 10 teeth that, when moved one complete revolution, advanced a second gear one place. This is the same principle that is used in the automobile's odometer mechanism and is the basis of all mechanical calculators. Incidentally, the PASCAL programming language is named in honor of Blaise Pascal for his pioneering work in mathematics and with the mechanical calculator.

The arrival of the first practical geared mechanical machines used to automatically compute information dates to the early 1800s. This is before humans invented the lightbulb or before much was known about electricity. In this dawn of the computer age, humans dreamed of mechanical machines that could compute numerical facts with a program—not merely calculating facts, as with a calculator.

In 1937 it was discovered through plans and journals that one early pioneer of mechanical computing machinery was Charles Babbage, aided by Augusta Ada Byron, the Countess of Lovelace. Babbage was commissioned in 1823 by the Royal Astronomical Society of Great Britain to produce a programmable calculating machine. This machine was to generate navigational tables for the Royal Navy. He accepted the challenge and began to create what he called his **Analytical Engine**. This engine was a steam-powered mechanical computer that stored a thousand 20-digit decimal numbers and a variable program that could modify the function of the machine to perform various calculating tasks. Input to his engine was through punched cards, much as computers in the 1950s and 1960s used punched cards. It is assumed that he obtained the idea of using punched cards from Joseph Jacquard, a Frenchman who used punched cards as input to a weaving machine he invented in 1801, which is today called *Jacquard's loom*. Jacquard's loom used punched cards to select intricate weaving patterns in the cloth that it produced. The punched cards programmed the loom.

After many years of work, Babbage's dream began to fade when he realized that the machinists of his day were unable to create the mechanical parts needed to complete his work. The Analytical Engine required more than 50,000 machined parts, which could not be made with enough precision to allow his engine to function reliably.

The Electrical Age

The 1800s saw the advent of the electric motor (conceived by Michael Faraday); with it came a multitude of motor-driven adding machines, all based on the mechanical calculator developed by Blaise Pascal. These electrically driven mechanical calculators were common pieces of office

---

[1]80X86 is an accepted acronym for 8086, 8088, 80186, 80188, 80286, 80386, and 80486 microprocessors and also include the Pentium series.

[2]Pentium, Pentium Pro, Pentium II, Pentium III, Pentium 4, and Core2 are registered trademarks of Intel Corporation.

equipment until well into the early 1970s, when the small handheld electronic calculator, first introduced by Bomar Corporation and called the **Bomar Brain**, appeared. Monroe was also a leading pioneer of electronic calculators, but its machines were desktop, four-function models the size of cash registers.

In 1889, Herman Hollerith developed the punched card for storing data. Like Babbage, he too apparently borrowed the idea of a punched card from Jacquard. He also developed a mechanical machine—driven by one of the new electric motors—that counted, sorted, and collated information stored on punched cards. The idea of calculating by machinery intrigued the United States government so much that Hollerith was commissioned to use his punched-card system to store and tabulate information for the 1890 census.

In 1896, Hollerith formed a company called the Tabulating Machine Company, which developed a line of machines that used punched cards for tabulation. After a number of mergers, the Tabulating Machine Company was formed into the International Business Machines Corporation, now referred to more commonly as IBM, Inc. The punched cards used in early computer systems are often called **Hollerith cards**, in honor of Herman Hollerith. The 12-bit code used on a punched card is called the **Hollerith code**.

Mechanical machines driven by electric motors continued to dominate the information processing world until the construction of the first electronic calculating machine in 1941. A German inventor named Konrad Zuse, who worked as an engineer for the Henschel Aircraft Company in Berlin, invented the first modern electromechanical computer. His Z3 calculating computer, as pictured in Figure 1–1, was probably invented for use in aircraft and missile design during World War II for the German war effort. The Z3 was a relay logic machine that was clocked at 5.33 Hz (far slower than the latest multiple GHz microprocessors). Had Zuse been given adequate funding by the German government, he most likely would have developed a
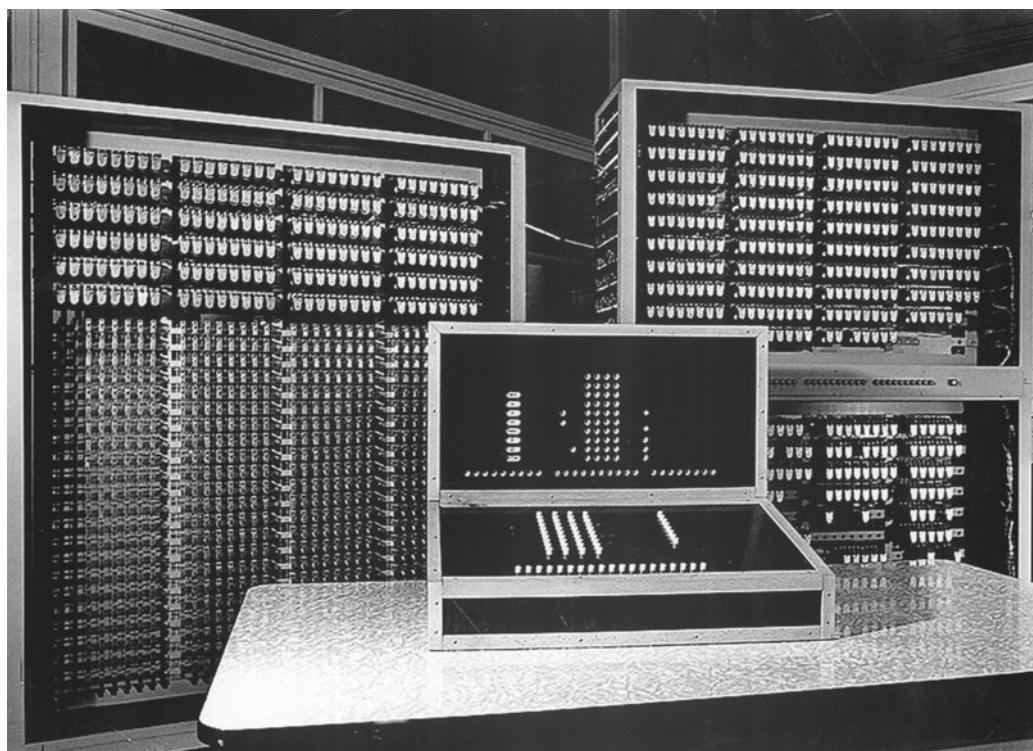
FIGURE 1–1   The Z3 computer developed by Konrad Zuse uses a 5.33 hertz clocking frequency. (Photo courtesy of Horst Zuse, the son of Konrad.)

much more powerful computer system. Zuse is today finally receiving some belated honor for his pioneering work in the area of digital electronics, which began in the 1930s, and for his Z3 computer system. In 1936 Zuse constructed a mechanical version of his system and later in 1939 Zuse constructed his first electromechanical computer system, called the Z2.

It has recently been discovered (through the declassification of British military documents) that the first electronic computer was placed into operation in 1943 to break secret German military codes. This first electronic computing system, which used vacuum tubes, was invented by Alan Turing. Turing called his machine **Colossus**, probably because of its size. A problem with Colossus was that although its design allowed it to break secret German military codes generated by the mechanical **Enigma machine**, it could not solve other problems. Colossus was not programmable—it was a fixed-program computer system, which today is often called a **special-purpose computer**.

The first general-purpose, programmable electronic computer system was developed in 1946 at the University of Pennsylvania. This first modem computer was called the **ENIAC** (**Electronic Numerical Integrator and Calculator**). The ENIAC was a huge machine, containing over 17,000 vacuum tubes and over 500 miles of wires. This massive machine weighed over 30 tons, yet performed only about 100,000 operations per second. The ENIAC thrust the world into the age of electronic computers. The ENIAC was programmed by rewiring its circuits—a process that took many workers several days to accomplish. The workers changed the electrical connections on plug-boards that looked like early telephone switchboards. Another problem with the ENIAC was the life of the vacuum tube components, which required frequent maintenance.

Breakthroughs that followed were the development of the transistor on December 23, 1947 at Bell Labs by John Bardeen, William Shockley, and Walter Brattain. This was followed by the 1958 invention of the integrated circuit by Jack Kilby of Texas Instruments. The integrated circuit led to the development of digital integrated circuits (RTL, or resistor-to-transistor logic) in the 1960s and the first microprocessor at Intel Corporation in 1971. At that time, Intel engineers Federico Faggin, Ted Hoff, and Stan Mazor developed the 4004 microprocessor (U.S. Patent 3,821,715)—the device that started the microprocessor revolution that continues today at an ever-accelerating pace.

Programming Advancements

Now that programmable machines were developed, programs and programming languages began to appear. As mentioned earlier, the first programmable electronic computer system was programmed by rewiring its circuits. Because this proved too cumbersome for practical application, early in the evolution of computer systems, computer languages began to appear in order to control the computer. The first such language, **machine language**, was constructed of ones and zeros using binary codes that were stored in the computer memory system as groups of instructions called a program. This was more efficient than rewiring a machine to program it, but it was still extremely time-consuming to develop a program because of the sheer number of program codes that were required. Mathematician John von Neumann was the first modern person to develop a system that accepted instructions and stored them in memory. Computers are often called **von Neumann machines** in honor of John von Neumann. (Recall that Babbage also had developed the concept long before von Neumann.)

Once computer systems such as the UNIVAC became available in the early 1950s, **assembly language** was used to simplify the chore of entering binary code into a computer as its instructions. The assembler allows the programmer to use mnemonic codes, such as ADD for addition, in place of a binary number such as 0100 0111. Although assembly language was an aid to programming, it wasn't until 1957, when Grace Hopper developed the first high-level programming language called **FLOWMATIC**, that computers became easier to program. In the

same year, IBM developed FORTRAN (**FORmula TRANslator**) for its computer systems. The FORTRAN language allowed programmers to develop programs that used formulas to solve mathematical problems. Note that FORTRAN is still used by some scientists for computer programming. Another similar language, introduced about a year after FORTRAN, was ALGOL (**ALGOrithmic Language**).

The first truly successful and widespread programming language for business applications was COBOL (**COmputer Business Oriented Language**). Although COBOL usage has diminished considerably in recent years, it is still a player in some large business and government systems. Another once-popular business language is RPG (**Report Program Generator**), which allows programming by specifying the form of the input, output, and calculations.

Since these early days of programming, additional languages have appeared. Some of the more common modern programming languages are BASIC, C#, C/C++, Java, PASCAL, and ADA. The BASIC and PASCAL languages were both designed as teaching languages, but have escaped the classroom. The BASIC language is used in many computer systems and may be one of the most common programming languages today. The BASIC language is probably the easiest of all to learn. Some estimates indicate that the BASIC language is used in the personal computer for 80% of the programs written by users. In the past decade, a new version of BASIC, Visual BASIC, has made programming in the Windows environment easier. The Visual BASIC language may eventually supplant C/C++ and PASCAL as a scientific language, but it is doubtful. It is more apparent that the C# language is gaining headway and *may* actually replace C/C++ and most other languages including Java and *may* eventually replace BASIC. This of course is conjecture and only the future will show which language eventually becomes dominant.

In the scientific community, primarily C/C++ and occasionally PASCAL and FORTRAN appear as control programs. One recent survey of embedded system developers showed that C was used by 60% and that 30% used assembly language. The remainder used BASIC and JAVA. These languages, especially C/C++, allow the programmer almost complete control over the programming environment and computer system. In many cases, C/C++ is replacing some of the low-level machine control software or drivers normally reserved for assembly language. Even so, assembly language still plays an important role in programming. Many video games written for the personal computer are written almost exclusively in assembly language. Assembly language is also interspersed with C/C++ to perform machine control functions efficiently. Some of the newer parallel instructions found on the newest Pentium and Core2 microprocessors are only programmable in assembly language.

The ADA language is used heavily by the Department of Defense. The ADA language was named in honor of Augusta Ada Byron, Countess of Lovelace. The Countess worked with Charles Babbage in the early 1800s in the development of software for his Analytical Engine.

The Microprocessor Age

The world's first microprocessor, the Intel 4004, was a 4-bit microprocessor–programmable controller on a chip. It addressed a mere 4096, 4-bit-wide memory locations. (A **bit** is a binary digit with a value of one or zero. A 4-bit-wide memory location is often called a **nibble**.) The 4004 instruction set contained only 45 instructions. It was fabricated with the then-current state-of-the-art P-channel MOSFET technology that only allowed it to execute instructions at the slow rate of 50 KIPs (**kilo-instructions per second**). This was slow when compared to the 100,000 instructions executed per second by the 30-ton ENIAC computer in 1946. The main difference was that the 4004 weighed much less than an ounce.

At first, applications abounded for this device. The 4-bit microprocessor debuted in early video game systems and small microprocessor-based control systems. One such early video game, a shuffleboard game, was produced by Bailey. The main problems with this early microprocessor were its speed, word width, and memory size. The evolution of the 4-bit microprocessor ended

TABLE 1–1 Early 8-bit microprocessors.

Manufacturer	Part Number
Fairchild	F-8
Intel	8080
MOS Technology	6502
Motorola	MC6800
National Semiconductor	IMP-8
Rockwell International	PPS-8
Zilog	Z-8

when Intel released the 4040, an updated version of the earlier 4004. The 4040 operated at a higher speed, although it lacked improvements in word width and memory size. Other companies, particularly Texas Instruments (TMS-1000), also produced 4-bit microprocessors. The 4-bit microprocessor still survives in low-end applications such as microwave ovens and small control systems and is still available from some microprocessor manufacturers. Most calculators are still based on 4-bit microprocessors that process 4-bit BCD (**binary-coded decimal**) codes.

Later in 1971, realizing that the microprocessor was a commercially viable product, Intel Corporation released the 8008—an extended 8-bit version of the 4004 microprocessor. The 8008 addressed an expanded memory size (16K bytes) and contained additional instructions (a total of 48) that provided an opportunity for its application in more advanced systems. (A **byte** is generally an 8-bit-wide binary number and a **K** is 1024. Often, memory size is specified in K bytes.)

As engineers developed more demanding uses for the 8008 microprocessor, they discovered that its somewhat small memory size, slow speed, and instruction set limited its usefulness. Intel recognized these limitations and introduced the 8080 microprocessor in 1973—the first of the modern 8-bit microprocessors. About six months after Intel released the 8080 microprocessor, Motorola Corporation introduced its MC6800 microprocessor. The floodgates opened and the 8080—and, to a lesser degree, the MC6800—ushered in the age of the microprocessor. Soon, other companies began to introduce their own versions of the 8-bit microprocessor. Table 1–1 lists several of these early microprocessors and their manufacturers. Of these early microprocessor producers, only Intel and Motorola (IBM also produces Motorola-style microprocessors) continue successfully to create newer and improved versions of the microprocessor. Motorola has sold its microprocessor division, and that company is now called Freescale Semiconductors, Inc. Zilog still manufactures microprocessors, but remains in the background, concentrating on microcontrollers and embedded controllers instead of general-purpose microprocessors. Rockwell has all but abandoned microprocessor development in favor of modem circuitry. Motorola has declined from having nearly 50% share of the microprocessor market to a much smaller share. Intel today has nearly 100% of the desktop and notebook market.

What Was Special about the 8080?   Not only could the 8080 address more memory and execute additional instructions, but it executed them 10 times faster than the 8008. An addition that took 20 μs (50,000 instructions per second) on an 8008-based system required only 2.0 μs (500,000 instructions per second) on an 8080-based system. Also, the 8080 was compatible with TTL (transistor-transistor logic), whereas the 8008 was not directly compatible. This made interfacing much easier and less expensive. The 8080 also addressed four times more memory (64K bytes) than the 8008 (l6K bytes). These improvements are responsible for ushering in the era of the 8080 and the continuing saga of the microprocessor. Incidentally, the first personal computer, the MITS Altair 8800, was released in 1974. (Note that the number 8800 was probably chosen to avoid copyright violations with Intel.) The BASIC language interpreter, written for the Altair 8800 computer, was developed in 1975 by Bill Gates and Paul Allen, the founders of

Microsoft Corporation. The assembler program for the Altair 8800 was written by Digital Research Corporation, which once produced DR-DOS for the personal computer.

The 8085 Microprocessor.   In 1977, Intel Corporation introduced an updated version of the 8080—the 8085. The 8085 was to be the last 8-bit, general-purpose microprocessor developed by Intel. Although only slightly more advanced than an 8080 microprocessor, the 8085 executed software at an even higher speed. An addition that took 2.0 μs (500,000 instructions per second on the 8080) required only 1.3 μs (769,230 instructions per second) on the 8085. The main advantages of the 8085 were its internal clock generator, internal system controller, and higher clock frequency. This higher level of component integration reduced the 8085's cost and increased its usefulness. Intel has managed to sell well over 100 million copies of the 8085 microprocessor, its most successful 8-bit, general-purpose microprocessor. Because the 8085 is also manufactured (second-sourced) by many other companies, there are over 200 million of these microprocessors in existence. Applications that contain the 8085 will likely continue to be popular. Another company that sold 500 million 8-bit microprocessors is Zilog Corporation, which produced the Z-80 microprocessor. The Z-80 is machine language–compatible with the 8085, which means that there are over 700 million microprocessors that execute 8085/Z-80 compatible code!

The Modern Microprocessor

In 1978, Intel released the 8086 microprocessor; a year or so later, it released the 8088. Both devices are 16-bit microprocessors, which executed instructions in as little as 400 ns (2.5 MIPs, or 2.5 **millions of instructions per second**). This represented a major improvement over the execution speed of the 8085. In addition, the 8086 and 8088 addressed 1M byte of memory, which was 16 times more memory than the 8085. (A **1M-byte memory** contains 1024K byte-sized memory locations or 1,048,576 bytes.) This higher execution speed and larger memory size allowed the 8086 and 8088 to replace smaller minicomputers in many applications. One other feature found in the 8086/8088 was a small 4- or 6-byte instruction cache or queue that prefetched a few instructions before they were executed. The queue sped the operation of many sequences of instructions and proved to be the basis for the much larger instruction caches found in modem microprocessors.

The increased memory size and additional instructions in the 8086 and 8088 have led to many sophisticated applications for microprocessors. Improvements to the instruction set included multiply and divide instructions, which were missing on earlier microprocessors. In addition, the number of instructions increased from 45 on the 4004, to 246 on the 8085, to well over 20,000 variations on the 8086 and 8088 microprocessors. Note that these microprocessors are called CISC (**complex instruction set computers**) because of the number and complexity of instructions. The additional instructions eased the task of developing efficient and sophisticated applications, even though the number of instructions are at first overwhelming and time-consuming to learn. The 16-bit microprocessor also provided more internal register storage space than the 8-bit microprocessor. The additional registers allowed software to be written more efficiently.

The 16-bit microprocessor evolved mainly because of the need for larger memory systems. The popularity of the Intel family was ensured in 1981, when IBM Corporation decided to use the 8088 microprocessor in its personal computer. Applications such as spreadsheets, word processors, spelling checkers, and computer-based thesauruses were memory-intensive and required more than the 64K bytes of memory found in 8-bit microprocessors to execute efficiently. The 16-bit 8086 and 8088 provided 1M byte of memory for these applications. Soon, even the 1M-byte memory system proved limiting for large databases and other applications. This led Intel to introduce the 80286 microprocessor, an updated 8086, in 1983.

The 80286 Microprocessor.   The 80286 microprocessor (also a 16-bit architecture microprocessor) was almost identical to the 8086 and 8088, except it addressed a 16M-byte memory system instead of a 1M-byte system. The instruction set of the 80286 was almost identical to the 8086 and 8088, except for a few additional instructions that managed the extra 15M bytes of memory. The clock speed of the 80286 was increased, so it executed some instructions in as little as 250 ns (4.0 MIPs) with the original release 8.0 MHz version. Some changes also occurred to the internal execution of the instructions, which led to an eightfold increase in speed for many instructions when compared to 8086/8088 instructions.

The 32-Bit Microprocessor.   Applications began to demand faster microprocessor speeds, more memory, and wider data paths. This led to the arrival of the 80386 in 1986 by Intel Corporation. The 80386 represented a major overhaul of the 16-bit 8086–80286 architecture. The 80386 was Intel's first practical 32-bit microprocessor that contained a 32-bit data bus and a 32-bit memory address. (Note that Intel produced an earlier, although unsuccessful, 32-bit microprocessor called the iapx-432.) Through these 32-bit buses, the 80386 addressed up to 4G bytes of memory. (**1G** of memory contains 1024M, or 1,073,741,824 locations.) A 4G-byte memory can store an astounding 1,000,000 typewritten, double-spaced pages of ASCII text data. The 80386 was available in a few modified versions such as the 80386SX, which addressed 16M bytes of memory through a 16-bit data and 24-bit address bus, and the 80386SL/80386SLC, which addressed 32M bytes of memory through a 16-bit data and 25-bit address bus. An 80386SLC version contained an internal cache memory that allowed it to process data at even higher rates. In 1995, Intel released the 80386EX microprocessor. The 80386EX microprocessor is called an **embedded PC** because it contains all the components of the AT class personal computer on a single integrated circuit. The 80386EX also contains 24 lines for input/output data, a 26-bit address bus, a 16-bit data bus, a DRAM refresh controller, and programmable chip selection logic.

Applications that require higher microprocessor speeds and large memory systems include software systems that use a GUI, or **graphical user interface**. Modem graphical displays often contain 256,000 or more picture elements (pixels, or pels). The least sophisticated VGA (**variable graphics array**) video display has a resolution of 640 pixels per scanning line with 480 scanning lines (this is the resolution used when the computer boots and display the boot screen). To display one screen of information, each picture element must be changed, which requires a high-speed microprocessor. Virtually all new software packages use this type of video interface. These GUI-based packages require high microprocessor speeds and accelerated video adapters for quick and efficient manipulation of video text and graphical data. The most striking system, which requires high-speed computing for its graphical display interface, is Microsoft Corporation's Windows.[3] We often call a GUI a WYSIWYG (**what you see is what you get**) display.

The 32-bit microprocessor is needed because of the size of its data bus, which transfers real (single-precision floating-point) numbers that require 32-bit-wide memory. In order to efficiently process 32-bit real numbers, the microprocessor must efficiently pass them between itself and memory. If the numbers pass through an 8-bit data bus, it takes four read or write cycles; when passed through a 32-bit data bus, however, only one read or write cycle is required. This significantly increases the speed of any program that manipulates real numbers. Most high-level languages, spreadsheets, and database management systems use real numbers for data storage. Real numbers are also used in graphical design packages that use vectors to plot images on the video screen. These include such CAD (**computer-aided drafting/design**) systems as AUTOCAD, ORCAD, and so forth.

---

[3]Windows is a registered trademark of Microsoft Corporation and is currently available as Windows 98, Windows 2000, Windows ME, and Windows XP.

Besides providing higher clocking speeds, the 80386 included a memory management unit that allowed memory resources to be allocated and managed by the operating system. Earlier microprocessors left memory management completely to the software. The 80386 included hardware circuitry for memory management and memory assignment, which improved its efficiency and reduced software overhead.

The instruction set of the 80386 microprocessor was upward-compatible with the earlier 8086, 8088, and 80286 microprocessors. Additional instructions referenced the 32-bit registers and managed the memory system. Note that memory management instructions and techniques used by the 80286 are also compatible with the 80386 microprocessor. These features allowed older, 16-bit software to operate on the 80386 microprocessor.

The 80486 Microprocessor. In 1989, Intel released the 80486 microprocessor, which incorporated an 80386-like microprocessor, an 80387-like numeric coprocessor, and an 8K-byte cache memory system into one integrated package. Although the 80486 microprocessor was not radically different from the 80386, it did include one substantial change. The internal structure of the 80486 was modified from the 80386 so that about half of its instructions executed in one clock instead of two clocks. Because the 80486 was available in a 50 MHz version, about half of the instructions executed in 25 ns (50 MIPs). The average speed improvement for a typical mix of instructions was about 50% over the 80386 that operated at the same clock speed. Later versions of the 80486 executed instructions at even higher speeds with a 66 MHz double-clocked version (80486DX2). The double-clocked 66 MHz version executed instructions at the rate of 66 MHz, with memory transfers executing at the rate of 33 MHz. (This is why it was called a double-clocked microprocessor.) A triple-clocked version from Intel, the 80486DX4, improved the internal execution speed to 100 MHz with memory transfers at 33 MHz. Note that the 80486DX4 microprocessor executed instructions at about the same speed as the 60 MHz Pentium. It also contained an expanded 16K-byte cache in place of the standard 8K-byte cache found on earlier 80486 microprocessors. Advanced Micro Devices (AMD) has produced a triple-clocked version that runs with a bus speed of 40 MHz and a clock speed of 120 MHz. The future promises to bring microprocessors that internally execute instructions at rates of up to 10 GHz or higher.

Other versions of the 80486 were called OverDrive[4] processors. The OverDrive processor was actually a double-clocked version of the 80486DX that replaced an 80486SX or slower-speed 80486DX. When the OverDrive processor was plugged into its socket, it disabled or replaced the 80486SX or 80486DX, and functioned as a doubled-clocked version of the microprocessor. For example, if an 80486SX, operating at 25 MHz, was replaced with an OverDrive microprocessor, it functioned as an 80486DX2 50 MHz microprocessor using a memory transfer rate of 25 MHz.

Table 1–2 lists many microprocessors produced by Intel and Motorola with information about their word and memory sizes. Other companies produce microprocessors, but none have attained the success of Intel and, to a lesser degree, Motorola.

The Pentium Microprocessor. The Pentium, introduced in 1993, was similar to the 80386 and 80486 microprocessors. This microprocessor was originally labeled the P5 or 80586, but Intel decided not to use a number because it appeared to be impossible to copyright a number. The two introductory versions of the Pentium operated with a clocking frequency of 60 MHz and 66 MHz, and a speed of 110 MIPs, with a higher-frequency 100 MHz one and one-half clocked version that operated at 150 MIPs. The double-clocked Pentium, operating at 120 MHz and 133 MHz, was also available, as were higher-speed versions. (The fastest version produced by Intel is the 233 MHz Pentium, which is a three and one-half clocked version.) Another difference was that the cache size was increased to 16K bytes from the 8K cache found in the basic version

TABLE 1–2   Many modern Intel and Motorola microprocessors.

Manufacturer	Part Number	Data Bus Width	Memory Size
Intel	8048	8	2K internal
	8051	8	8K internal
	8085A	8	64K
	8086	16	1M
	8088	8	1M
	8096	16	8K internal
	80186	16	1M
	80188	8	1M
	80251	8	16K internal
	80286	16	16M
	80386EX	16	64M
	80386DX	32	4G
	80386SL	16	32M
	80386SLC	16	32M + 8K cache
	80386SX	16	16M
	80486DX/DX2	32	4G + 8K cache
	80486SX	32	4G + 8K cache
	80486DX4	32	4G + 16 cache
	Pentium	64	4G + 16K cache
	Pentium OverDrive	32	4G + 16K cache
	Pentium Pro	64	64G + 16K L1 cache + 256K L2 cache
	Pentium II	64	64G + 32K L1 cache + 256K L2 cache
	Pentium III	64	64G + 32K L1 cache + 256K L2 cache
	Pentium 4	64	64G+32K L1 cache+ 512K L2 cache (or larger) (1T for 64-bit extensions)
	Pentium4 D (Dual Core)	64	1T + 32K L1 cache + 2 or 4 M L2 cache
	Core2	64	1T + 32K L1 cache + a shared 2 or 4 M L2 cache
	Itanium (Dual Core)	128	1T + 2.5 M L1 and L2 cache + 24 M L3 cache
Motorola	6800	8	64K
	6805	8	2K
	6809	8	64K
	68000	16	16M
	68008D	8	4M
	68008Q	8	1M
	68010	16	16M
	68020	32	4G
	68030	32	4G + 256 cache
	68040	32	4G + 8K cache
	68050	32	Proposed, but never released
	68060	64	4G + 16K cache
	PowerPC	64	4G + 32K cache

of the 80486. The Pentium contained an 8K-byte instruction cache and an 8K-byte data cache, which allowed a program that transfers a large amount of memory data to still benefit from a cache. The memory system contained up to 4G bytes, with the data bus width increased from the 32 bits found in the 80386 and 80486 to a full 64 bits. The data bus transfer speed was either 60 MHz or 66 MHz, depending on the version of the Pentium. (Recall that the bus speed of the 80486 was 33 MHz.) This wider data bus width accommodated double-precision floating-point numbers used for modem high-speed, vector-generated graphical displays. These higher bus speeds should allow virtual reality software and video to operate at more realistic rates on current and future Pentium-based platforms. The widened data bus and higher execution speed of the Pentium allow full-frame video displays to operate at scan rates of 30 Hz or higher—comparable to commercial television. Recent versions of the Pentium also included additional instructions, called multimedia extensions, or MMX instructions. Although Intel hoped that the MMX instructions would be widely used, it appears that few software companies have used them. The main reason is there is no high-level language support for these instructions.

Intel had also released the long-awaited Pentium OverDrive (P24T) for older 80486 systems that operate at either 63 MHz or 83 MHz clock. The 63 MHz version upgrades older 80486DX2 50 MHz systems; the 83 MHz version upgrades the 80486DX2 66 MHz systems. The upgraded 83 MHz system performs at a rate somewhere between a 66 MHz Pentium and a 75 MHz Pentium. If older VESA local bus video and disk-caching controllers seem too expensive to toss out, the Pentium OverDrive represents an ideal upgrade path from the 80486 to the Pentium.

Probably the most ingenious feature of the Pentium is its dual integer processors. The Pentium executes two instructions, which are not dependent on each other, simultaneously because it contains two independent internal integer processors called superscaler technology. This allows the Pentium to often execute two instructions per clocking period. Another feature that enhances performance is a jump prediction technology that speeds the execution of program loops. As with the 80486, the Pentium also employs an internal floating-point coprocessor to handle floating-point data, albeit at a five times speed improvement. These features portend continued success for the Intel family of microprocessors. Intel also may allow the Pentium to replace some of the RISC (**reduced instruction set computer**) machines that currently execute one instruction per clock. Note that some newer RISC processors execute more than one instruction per clock through the introduction of superscaler technology. Motorola, Apple, and IBM produce the PowerPC, a RISC microprocessor that has two integer units and a floating-point unit. The PowerPC certainly boosts the performance of the Apple Macintosh, but at present is slow to efficiently emulate the Intel family of microprocessors. Tests indicate that the current emulation software executes DOS and Windows applications at speeds slower than the 80486DX 25 MHz microprocessor. Because of this, the Intel family should survive for many years in personal computer systems. Note that there are currently 6 million Apple Macintosh[5] systems and well over 260 million personal computers based on Intel microprocessors. In 1998, reports showed that 96% of all PCs were shipped with the Windows operating system.
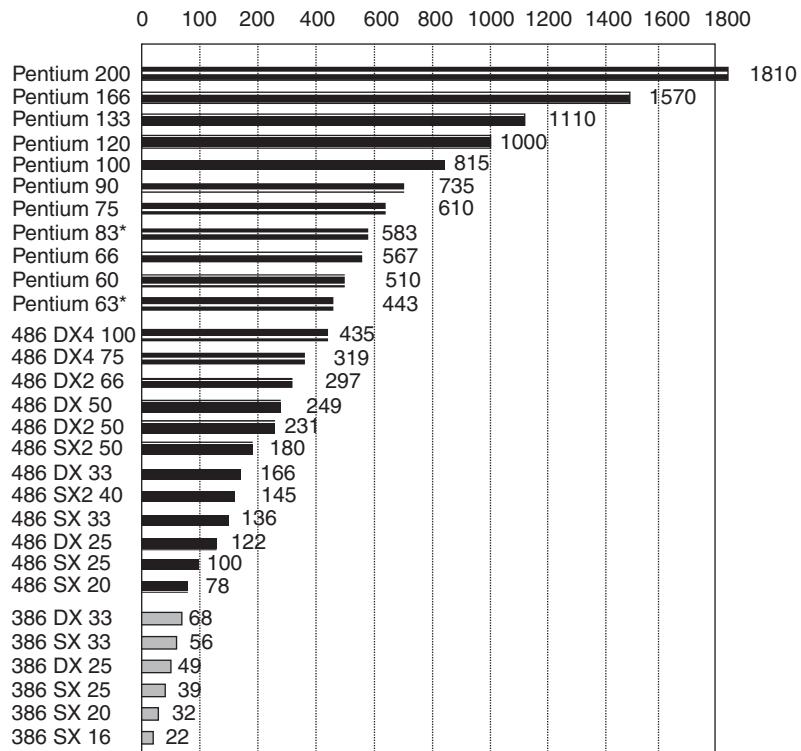
Recently Apple computer replaced the PowerPC with the Intel Pentium in most of its computer systems. It appears that the PowerPC could not keep pace with the Pentium line from Intel.

In order to compare the speeds of various microprocessors, Intel devised the iCOMP-rating index. This index is a composite of SPEC92, ZD Bench, and Power Meter. The iCOMP1 rating index is used to rate the speed of all Intel microprocessors through the Pentium. Figure 1–2 shows relative speeds of the 80386DX 25 MHz version at the low end to the Pentium 233 MHz version at the high end of the spectrum.

Since the release of the Pentium Pro and Pentium II, Intel has switched to the iCOMP2-rating index, which is scaled by a factor of 10 from the iCOMP1 index. A microprocessor with an index of 1000 using iCOMP1 is rated as 100 using iCOMP2. Another difference is the benchmarks used for

---

[5]Macintosh is a registered trademark of Apple Computer Corporation.

FIGURE 1–2   The Intel
iCOMP-rating index.

	0	100	200	400	600	800	1000	1200	1400	1600	1800

Pentium 200 — 1810
Pentium 166 — 1570
Pentium 133 — 1110
Pentium 120 — 1000
Pentium 100 — 815
Pentium 90 — 735
Pentium 75 — 610
Pentium 83* — 583
Pentium 66 — 567
Pentium 60 — 510
Pentium 63* — 443
486 DX4 100 — 435
486 DX4 75 — 319
486 DX2 66 — 297
486 DX 50 — 249
486 DX2 50 — 231
486 SX2 50 — 180
486 DX 33 — 166
486 SX2 40 — 145
486 SX 33 — 136
486 DX 25 — 122
486 SX 25 — 100
486 SX 20 — 78
386 DX 33 — 68
386 SX 33 — 56
386 DX 25 — 49
386 SX 25 — 39
386 SX 20 — 32
386 SX 16 — 22

Note: *Pentium OverDrive, the first part of
the scale is not linear, and the 166 MHz
and 200 MHz are MMX technology.

the scores. Figure 1–3 shows the iCOMP2 index listing the Pentium III at speeds up to 1000 MHz. Figure 1–4 shows SYSmark 2002 for the Pentium III and Pentium 4. Unfortunately Intel has not released any benchmarks that compare versions of the microprocessor since the SYSmark 2002. Newer benchmarks are available, but they do not compare one version with another.

Pentium Pro Processor.   A recent entry from Intel is the Pentium Pro processor, formerly named the P6 microprocessor. The Pentium Pro processor contains 21 million transistors, integer units, as well as a floating-point unit to increase the performance of most software. The basic clock frequency was 150 MHz and 166 MHz in the initial offering made available in late 1995. In addition to the internal 16K level-one (L1) cache (8K for data and 8K for instructions) the Pentium Pro processor also contains a 256K level-two (L2) cache. One other significant change is that the Pentium Pro processor uses three execution engines, so it can execute up to three instructions at a time, which can conflict and still execute in parallel. This represents a change from the Pentium, which executes two instructions simultaneously as long as they do not conflict. The Pentium Pro microprocessor has been optimized to efficiently execute 32-bit code; for this reason, it was often bundled with Windows NT rather than with normal versions of Windows 95. Intel launched the Pentium Pro processor for the server market. Still another change is that the Pentium Pro can address either a 4G-byte memory system or a 64G-byte memory system. The Pentium Pro has a 36-bit address bus if configured for a 64G memory system.

Pentium II and Pentium Xeon Microprocessors.   The Pentium II microprocessor (released in 1997) represents a new direction for Intel. Instead of being an integrated circuit as with prior versions of the microprocessor, Intel has placed the Pentium II on a small circuit board. The main reason for the change is that the L2 cache found on the main circuit board of the Pentium was not

FIGURE 1–3 The Intel iCOMP2-rating index.

	iCOMP2
Pentium III 1000 MHz	1277
Pentium III 933 MHz	1207
Pentium III 866 MHz	1125
Pentium III 800 MHz	1048
Pentium III 750 MHz	989
Pentium III 700 MHz	942
Pentium III 650 MHz	884
Pentium III 600 MHz	753
Pentium III 550 MHz	693
Pentium III 500 MHz	642
Pentium II 450 MHz	483
Pentium II 400 MHz	440
Pentium II 350 MHz	386
Pentium II 333 MHz	366
Pentium II 300 MHz	332
Pentium II 266 MHz	303
Pentium II 233 MHz	267
Pentium II* 266 MHz	213
Pentium 233 MHz	203

Note: *Pentium II Celeron, no cache.
iCOMP2 numbers are shown above. To convert to iCOMP3, multiply by 2.568.

fast enough to function properly with the Pentium II. On the Pentium system, the L2 cache operates at the system bus speed of 60 MHz or 66 MHz. The L2 cache and microprocessor are on a circuit board called the Pentium II module. This onboard L2 cache operates at a speed of 133 MHz and stores 512K bytes of information. The microprocessor on the Pentium II module is actually Pentium Pro with MMX extensions.

In 1998, Intel changed the bus speed of the Pentium II. Because the 266 MHz through the 333 MHz Pentium II microprocessors used an external bus speed of 66 MHz, there was a bottleneck, so the newer Pentium II microprocessors use a 100 MHz bus speed. The Pentium II microprocessors rated at 350 MHz, 400 MHz, and 450 MHz all use this higher 100 MHz memory bus speed. The higher speed memory bus requires the use of 8 ns SDRAM in place of the 10 ns SDRAM found in use with the 66 MHz bus speed.

In mid-1998 Intel announced a new version of the Pentium II called the Xeon,[6] which was specifically designed for high-end workstation and server applications. The main difference between the Pentium II and the Pentium II Xeon is that the Xeon is available with a L1 cache size of 32K bytes and a L2 cache size of either 512K, 1M, or 2M bytes. The Xeon functions with the 440GX chip set. The Xeon is also designed to function with four Xeons in the same system, which is similar to the Pentium Pro. This newer product represents a change in Intel's strategy: Intel now produces a professional version and a home/business version of the Pentium II microprocessor.

Pentium III Microprocessor.   The Pentium III microprocessor uses a faster core than the Pentium II, but it is still a P6 or Pentium Pro processor. It is also available in the slot 1 version mounted on a plastic cartridge and a socket 370 version called a flip-chip, which looks like the older Pentium package. Intel claims the flip-chip version costs less. Another difference is that the Pentium III is available with clock frequencies of up to 1 GHz. The slot 1 version contains a 512K cache and the flip-chip version contains a 256K cache. The speeds are comparable because the cache in the slot 1 version runs at one-half the clock speed, while the cache in the flip-chip version runs at the clock speed. Both versions use a memory bus speed of 100 MHz, while the Celeron[7] uses memory bus clock speed of 66 MHz.

The speed of the front side bus, the connection from the microprocessor to the memory controller, PCI controller, and AGP controller, is now either 100 MHz or 133 MHz. Although the memory still runs at 100 MHz, this change has improved performance.

Pentium 4 and Core2 Microprocessors.   The Pentium 4 microprocessor was first made available in late 2000. The most recent version of the Pentium is called the Core2 by Intel. The Pentium 4 and Core2, like the Pentium Pro through the Pentium III, use the Intel P6 architecture. The main difference is that the Pentium 4 is available in speeds to 3.2 GHz and faster and the chip sets that support the Pentium 4 use the RAMBUS or DDR memory technologies in place of once standard SDRAM technology. The Core2 is available at speeds of up to 3 GHz. These higher microprocessor speeds are made available by an improvement in the size of the internal

[6]Xeon is a registered trademark of Intel Corporation.

[7]Celeron is a trademark of Intel Corporation.

TABLE 1–3 Intel microprocessor core (P) versions.

Core (P) Version	Microprocessor
P1	8086 and 8088 (80186 and 80188)
P2	80286
P3	80386
P4	80486
P5	Pentium
P6	Pentium Pro, Pentium II, Pentium III, Pentium 4, and Core2
P7	Itanium

integration, which at present is the 0.045 micron or 45 nm technology. It is also interesting to note that Intel has changed the level 1 cache size from 32K to 8K bytes and most recently to 64K. Research must have shown that this size is large enough for the initial release version of the microprocessor, with future versions possibly containing a 64K L1 cache. The level 2 cache remains at 256K bytes as in the Pentium coppermine version with the latest versions containing a 512K cache. The Pentium 4 Extreme Edition contains a 2M L2 cache and the Pentium 4e contains a 1M level 2 cache, whereas the Core2 contains either a 2M or 4M L2 cache.

Another change likely to occur is a shift from aluminum to copper interconnections inside the microprocessor. Because copper is a better conductor, it should allow increased clock frequencies for the microprocessor in the future. This is especially true now that a method for using copper has surfaced at IBM Corporation. Another event to look for is a change in the speed of the front side bus, which will likely increase beyond the current maximum 1033 MHz.

Table 1–3 shows the various Intel P numbers and the microprocessors that belong to each class. The P versions show what internal core microprocessor is found in each of the Intel microprocessors. Notice that all of the microprocessors since the Pentium Pro use the same basic microprocessor core.

Pentium 4 and Core2, 64-bit and Multiple Core Microprocessors. Recently Intel has included new modifications to the Pentium 4 and Core2 that include a 64-bit core and multiple cores. The 64-bit modification allows the microprocessor to address more than 4G bytes of memory through a wider 64-bit address. Currently, 40 address pins in these newer versions allow up to 1T (terabytes) of memory to be accessed. The 64-bit machine also allows 64-bit integer arithmetic, but this is much less important than the ability to address more memory.

The biggest advancement in the technology is not the 64-bit operation, but the inclusion of multiple cores. Each core executes a separate task in a program, which increases the speed of execution if a program is written to take advantage of the multiple cores. Programs that do this are called **multithreaded** applications. Currently, Intel manufactures dual and quad core versions, but in the future the number of cores will likely increase to eight or even sixteen. The problem faced by Intel is that the clock speed cannot be increased to a much higher rate, so multiple cores are the current solution to providing faster microprocessors. Does this mean that higher clock speeds are not possible? Only the future portends whether they are or are not.

Intel recently demonstrated a version of the Core2 that contains 80 cores that uses the 45 nm fabrication technology. Intel expects to release an 80-core version some time in the next 5 years. The fabrication technology will become slightly smaller with 35 nm and possibly 25 nm technology.

The Future of Microprocessors. No one can really make accurate predictions, but the success of the Intel family should continue for quite a few years. What may occur is a change to RISC technology, but more likely are improvements to a new technology jointly by Intel and Hewlett-Packard
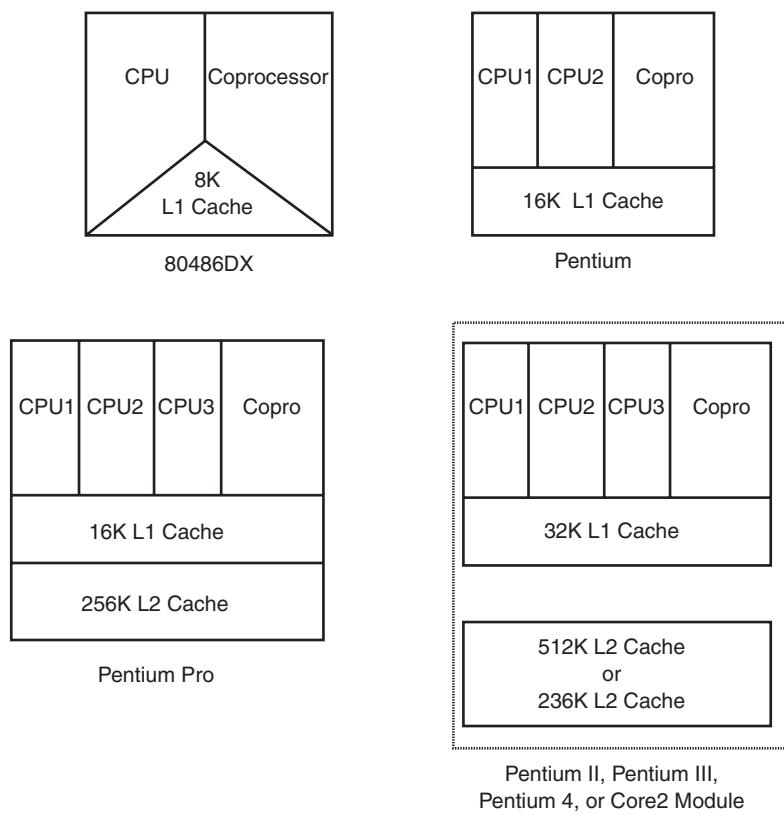
called hyper-threading technology. Even this new technology embodies the CISC instruction set of the 80X86 family of microprocessors, so that software for the system will survive. The basic premise behind this technology is that many microprocessors communicate directly with each other, allowing parallel processing without any change to the instruction set or program. Currently, the superscaler technology uses many microprocessors, but they all share the same register set. This new technology contains many microprocessors, each containing its own register set that is linked with the other microprocessors' registers. This technology offers true parallel processing without writing any special program.

The hyper-threading technology should continue into the future, bringing even more parallel processors (at present two processors). There are suggestions that Intel may also incorporate the chip set into the microprocessor package.

In 2002, Intel released a new microprocessor architecture that is 64 bits in width and has a 128-bit data bus. This new architecture, named the Itanium,[8] is a joint venture called EPIC (Explicitly Parallel Instruction Computing) of Intel and Hewlett-Packard. The Itanium architecture allows greater parallelism than traditional architectures, such as the Pentium III or Pentium 4. These changes include 128 general-purpose integer registers, 128 floating-point registers, 64 predicate registers, and many execution units to ensure enough hardware resources for software. The Itanium is designed for the server market and may or may not trickle down to the home/business market in the future.

Figure 1–5 is a conceptual view, comparing the 80486 through Pentium 4 microprocessors. Each view shows the internal structure of these microprocessors: the CPU, coprocessor, and

FIGURE 1–5 Conceptual views of the 80486, Pentium Pro, Pentium II, Pentium III, Pentium 4, and Core2 microprocessors.

8K L1 Cache

80486DX

16K L1 Cache

Pentium

16K L1 Cache
256K L2 Cache

Pentium Pro

32K L1 Cache
512K L2 Cache or 236K L2 Cache

Pentium II, Pentium III, Pentium 4, or Core2 Module

[8]Itanium is a trademark of Intel Corporation.

cache memory. This illustration shows the complexity and level of integration in each version of the microprocessor.

Because clock frequencies seemed to have peaked and the surge to multiple cores has begun, about the only major change to the Pentium will probably be a wider memory path (128 bits). Another consideration is the memory speed. Today, dynamic RAMs are the mainstay, but the speed of dynamic RAM memory has not changed for many years. A push to static RAM memory will eventually appear and will increase the performance of the PC. The main problem today with large static RAM is heat. Static RAM operates 50 times faster than dynamic RAM. Imagine a computer that contains a memory composed of static RAM.

Another problem is the speed of the mass storage connected to a computer. The transfer speed of hard disk drives has changed little in the past few years. A new technology is needed for mass storage. Flash memory could be a solution, because its write speed is comparable to hard disk memory. One change that would increase the speed of the computer system is the placement of possibly 4G bytes of flash memory to store the operation system for common applications. This would allow the operating system to load in a second or two instead of the many seconds required to boot a modern computer system.

1–2   THE MICROPROCESSOR-BASED PERSONAL COMPUTER SYSTEM

Computer systems have undergone many changes recently. Machines that once filled large areas have been reduced to small desktop computer systems because of the microprocessor. Although these desktop computers are compact, they possess computing power that was only dreamed of a few years ago. Million-dollar mainframe computer systems, developed in the early 1980s, are not as powerful as the Pentium Core2-based computers of today. In fact, many smaller companies have replaced their mainframe computers with microprocessor-based systems. Companies such as DEC (Digital Equipment Corporation, now owned by Hewlett-Packard Company) have stopped producing mainframe computer systems in order to concentrate their resources on microprocessor-based computer systems.

This section shows the structure of the microprocessor-based personal computer system. This structure includes information about the memory and operating system used in many microprocessor-based computer systems.

See Figure 1–6 for the block diagram of the personal computer. This diagram also applies to any computer system, from the early mainframe computers to the latest microprocessor-based systems. The block diagram is composed of three blocks that are interconnected by buses. (A **bus** is the set of common connections that carry the same type of information. For example, the address bus, which contains 20 or more connections, conveys the memory address to the memory.) These blocks and their function in a personal computer are outlined in this section of the text.

The Memory and I/O System

The memory structure of all Intel-based personal computers is similar. This includes the first personal computers based upon the 8088, introduced in 1981 by IBM, to the most powerful high-speed versions of today, based on the Pentium 4 or Core2. Figure 1–7 illustrates the memory map of a personal computer system. This map applies to any IBM personal computer or to any of the many IBM-compatible clones that are in existence.

The memory system is divided into three main parts: TPA (transient program area), system area, and XMS (extended memory system). The type of microprocessor in your computer determines whether an extended memory system exists. If the computer is based upon a really old 8086 or 8088 (a PC or XT), the TPA and systems area exist, but there is no extended memory

Buses

Memory system		Microprocessor		I/O system

Dynamic RAM (DRAM)
Static RAM (SRAM)
Cache
Read-only (ROM)
Flash memory
EEPROM
SDRAM
RAMBUS
DDR DRAM

8086
8088
80186
80188
80286
80386
80486
Pentium
Pentium Pro
Pentium II
Pentium III
Pentium 4
Core2

Printer
Serial communications
Floppy disk drive
Hard disk drive
Mouse
CD-ROM drive
Plotter
Keyboard
Monitor
Tape backup
Scanner
DVD

FIGURE 1–6   The block diagram of a microprocessor-based computer system.

area. The PC and XT computers contain 640K bytes of TPA and 384K bytes of system memory, for a total memory size of 1M bytes. We often call the first 1M byte of memory the real or conventional memory system because each Intel microprocessor is designed to function in this area by using its real mode of operation.

Computer systems based on the 80286 through the Core2 not only contain the TPA (640K bytes) and system area (384K bytes), they also contain extended memory. These machines are

FIGURE 1–7   The memory map of a personal computer.

Extended memory

15M bytes in the 80286 or 80386SX
31M bytes in the 80386SL/SLC
63M bytes in the 80386EX
4095M bytes in the 80386DX, 80486, and Pentium
64G bytes in the Pentium Pro, Pentium II, Pentium III, Pentium 4, and Core2

System area
384K bytes

1M bytes of real (conventional) memory

TPA
640K bytes

often called AT class machines. The PS/l and PS/2, produced by IBM, are other versions of the same basic memory design. Sometimes, these machines are also referred to as ISA (industry standard architecture) or EISA (extended ISA) machines. The PS/2 is referred to as a micro-channel architecture system, or ISA system, depending on the model number.

A change beginning with the introduction of the Pentium microprocessor and the ATX class machine is the addition of a bus called the PCI (**peripheral component interconnect**) bus, now being used in all Pentium through Core2 systems. Extended memory contains up to 15M bytes in the 80286 and 80386SX-based computers, and up to 4095M bytes in the 80386DX, 80486, and Pentium microprocessors, in addition to the first 1M byte of real or conventional memory. The Pentium Pro through Core2 computer systems have up to 1M less than 4G or 1 M less than 64G of extended memory. Servers tend to use the larger 64G memory map, while home/business computers use the 4G-byte memory map. The ISA machine contains an 8-bit peripheral bus that is used to interface 8-bit devices to the computer in the 8086/8088-based PC or XT computer system. The AT class machine, also called an ISA machine, uses a l6-bit peripheral bus for interface and may contain the 80286 or above microprocessor. The EISA bus is a 32-bit peripheral interface bus found in a few older 80386DX- and 80486-based systems. Note that each of these buses is compatible with the earlier versions. That is, the 8-bit interface card functions in the 8-bit ISA, l6-bit ISA, or 32-bit EISA bus system. Likewise, a l6-bit inter-face card functions in the l6-bit ISA or 32-bit EISA system.

Another bus type found in many 80486-based personal computers is called the **VESA** local bus, or VL bus. The local bus interfaces disk and video to the microprocessor at the local bus level, which allows 32-bit interfaces to function at the same clocking speed as the microproces-sor. A recent modification to the VESA local bus supports the 64-bit data bus of the Pentium microprocessor and competes directly with the PCI bus, although it has generated little, if any, interest. The ISA and EISA standards function at only 8 MHz, which reduces the performance of the disk and video interfaces using these standards. The PCI bus is either a 32- or 64-bit bus that is specifically designed to function with the Pentium through Core2 microprocessors at a bus speed of 33 MHz.
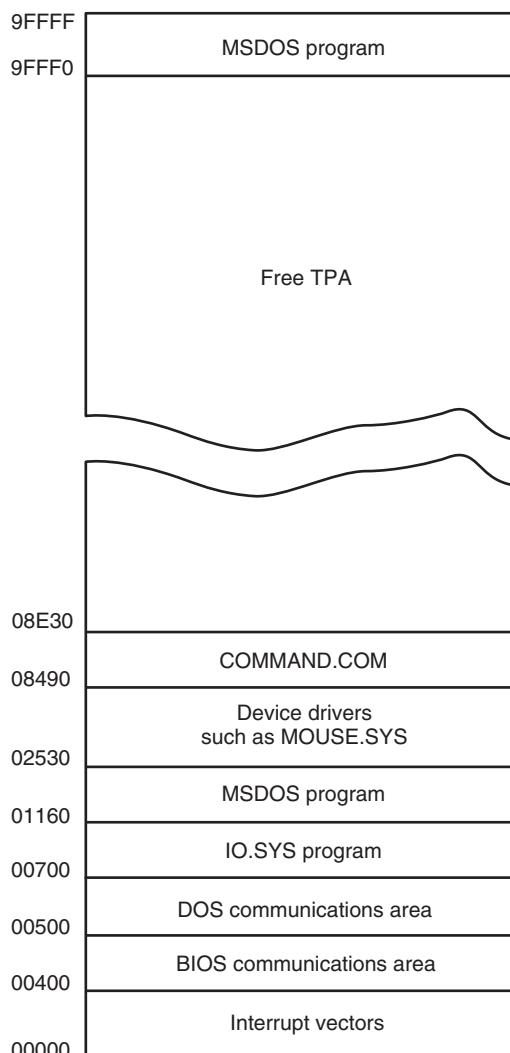
Three newer buses have appeared in ATX class systems. The first to appear was the USB (**universal serial bus**). The universal serial bus is intended to connect peripheral devices such as keyboards, a mouse, modems, and sound cards to the microprocessor through a serial data path and a twisted pair of wires. The main idea is to reduce system cost by reducing the number of wires. Another advantage is that the sound system can have a separate power supply from the PC, which means much less noise. The data transfer rates through the USB are 10 Mbps at pre-sent for USB1; they increase to 480 Mbps in USB2.

The second newer bus is the AGP (**advanced graphics port**) for video cards. The advanced graphics port transfers data between the video card and the microprocessor at higher speeds (66 MHz, with a 64-bit data path, or 533M bytes per second) than were possible through any other bus or connection. The latest AGP speed is 8X or 2G bytes per second. This video sub-system change has been made to accommodate the new DVD players for the PC.

The latest new buses to appear are the serial ATA interface (**SATA**) for hard disk drives and the PCI Express bus for the video card. The SATA bus transfers data from the PC to the hard disk drive at rates of 150M bytes per second or 300M bytes for SATA-2. The serial ATA standard will eventually reach speeds of 450M bytes per second. Today PCI Express bus video cards operate at 16X speeds.

The TPA. The transient program area (TPA) holds the DOS (**disk operating system**) operating system and other programs that control the computer system. The TPA is a DOS con-cept and not really applicable in Windows. The TPA also stores any currently active or inactive DOS application programs. The length of the TPA is 640K bytes. As mentioned, this area of memory holds the DOS operating system, which requires a portion of the TPA to function.

FIGURE 1–8 The memory map of the TPA in a personal computer. (Note that this map will vary between systems.)

In practice, the amount of memory remaining for application software is about 628K bytes if MSDOS[9] version 7.x is used as an operating system. Earlier versions of DOS required more of the TPA area and often left only 530K bytes or less for application programs. Figure 1–8 shows the organization of the TPA in a computer system running DOS.

The DOS memory map shows how the many areas of the TPA are used for system programs, data, and drivers. It also shows a large area of memory available for application programs. To the left of each area is a hexadecimal number that represents the memory addresses that begin and end each data area. Hexadecimal memory addresses or memory locations are used to number each byte of the memory system. (A hexadecimal number is a number represented in radix 16 or base 16, with each digit representing a value from 0 to 9 and A to F. We often end a hexadecimal number with an H to indicate that it is a hexadecimal value. For example, 1234H is 1234 hexadecimal. We also represent hexadecimal data as 0xl234 for a 1234 hexadecimal.)

---

[9]MSDOS is a trademark of Microsoft Corporation and version 7.x is supplied with Windows XP.

The Interrupt vectors access various features of the DOS, BIOS (basic I/O system), and applications. The system BIOS is a collection of programs stored in either a read-only (ROM) or flash memory that operates many of the I/O devices connected to your computer system. The system BIOS and DOS communications areas contain transient data used by programs to access I/O devices and the internal features of the computer system. These are stored in the TPA so they can be changed as the DOS operates.

The IO.SYS is a program that loads into the TPA from the disk whenever an MSDOS system is started. The IO.SYS contains programs that allow DOS to use the keyboard, video display, printer, and other I/O devices often found in the computer system. The IO.SYS program links DOS to the programs stored on the system BIOS ROM.

The size of the driver area and number of drivers changes from one computer to another. **Drivers** are programs that control installable I/O devices such as a mouse, disk cache, hand scanner, CD-ROM memory (**Compact Disk Read-Only Memory**), DVD (**Digital Versatile Disk**), or installable devices, as well as programs. Installable drivers are programs that control or drive devices or programs that are added to the computer system. DOS drivers are normally files that have an extension of .SYS, such as MOUSE.SYS; in DOS version 3.2 and later, the files have an extension of .EXE, such as EMM386.EXE. Note that even though these files are not used by Windows, they are still used to execute DOS applications, even with Windows XP. Windows uses a file called SYSTEM.INI to load drivers used by Windows. In newer versions of Windows such as Windows XP, a registry is added to contain information about the system and the drivers used by the system. You can view the registry with the REGEDIT program.

The COMMAND.COM program (**command processor**) controls the operation of the computer from the keyboard when operated in the DOS mode. The COMMAND.COM program processes the DOS commands as they are typed from the keyboard. For example, if DIR is typed, the COMMAND.COM program displays a directory of the disk files in the current disk directory. If the COMMAND.COM program is erased, the computer cannot be used from the keyboard in DOS mode. Never erase COMMAND.COM, IO.SYS, or MSDOS.SYS to make room for other software, or your computer will not function.

The System Area. The DOS system area, although smaller than the TPA, is just as important. The **system area** contains programs on either a read-only memory (ROM) or flash memory, and areas of read/write (RAM) memory for data storage. Figure 1–9 shows the system area of a typical personal computer system. As with the map of the TPA, this map also includes the hexadecimal memory addresses of the various areas.
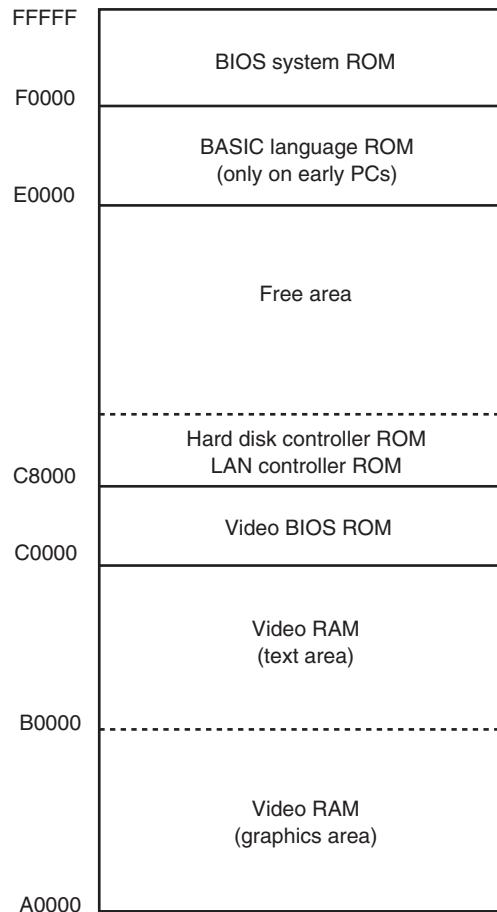
The first area of the system space contains video display RAM and video control programs on ROM or flash memory. This area starts at location A0000H and extends to location C7FFFH. The size and amount of memory used depends on the type of video display adapter attached to the system. Display adapters generally have their video RAM located at A0000H–AFFFFH, which stores graphical or bit-mapped data, and the memory at B0000H–BFFFFH stores text data. The video BIOS, located on a ROM or flash memory, is at locations C0000H–C7FFFH and contains programs that control the DOS video display.

The area at locations C8000H–DFFFFH is often open or free. This area is used for the expanded memory system (EMS) in a PC or XT system, or for the upper memory system in an AT system. Its use depends on the system and its configuration. The expanded memory system allows a 64K-byte page frame of memory to be used by application programs. This 64K-byte page frame (usually locations D0000H through DFFFFH) is used to expand the memory system by switching in pages of memory from the EMS into this range of memory addresses.

Memory locations E0000H–EFFFFH contain the cassette BASIC language on ROM found in early IBM personal computer systems. This area is often open or free in newer computer systems.

Finally, the system BIOS ROM is located in the top 64K bytes of the system area (F0000H–FFFFFH). This ROM controls the operation of the basic I/O devices connected to the
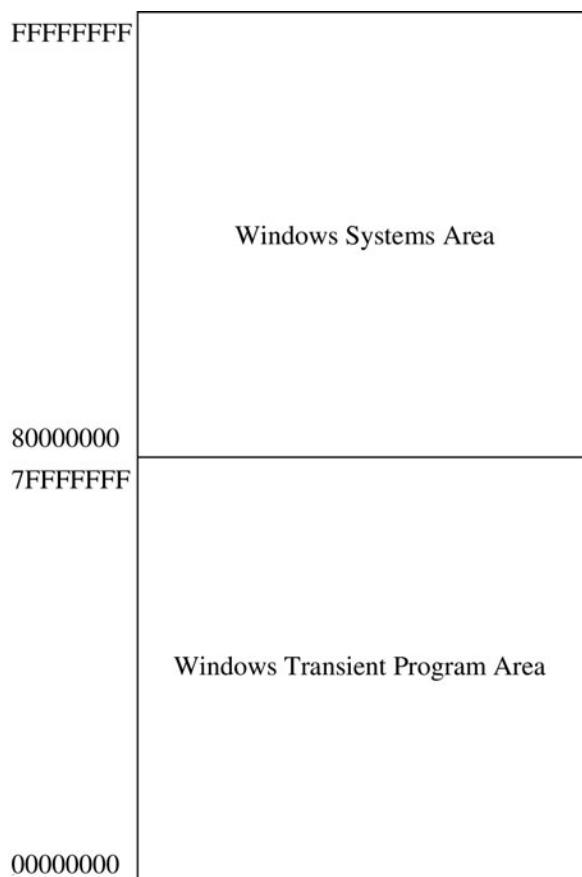
FIGURE 1–9   The system area of a typical personal computer.

computer system. It does not control the operation of the video system, which has its own BIOS ROM at location C0000H. The first part of the system BIOS (F0000H–F7FFFH) often contains programs that set up the computer; the second part contains procedures that control the basic I/O system.

Windows Systems.   Modern computers use a different memory map with Windows than the DOS memory maps of Figures 1–8 and 1–9. The Windows memory map appears in Figure 1–10 and has two main areas, a TPA and a system area. The difference between it and the DOS memory map are the sizes and locations of these areas.

The Windows TPA is the first 2G bytes of the memory system from locations 00000000H to 7FFFFFFFH. The Windows system area is the last 2G bytes of memory from locations 80000000H to FFFFFFFFH. It appears that the same idea used to construct the DOS memory map was also used in a modern Windows-based system. The system area is where the system BIOS is located and also the video memory. Also located in the system area is the actual Windows program and drivers. Every program that is written for Windows can use up to 2G bytes of memory located at linear addresses 00000000H through 7FFFFFFFH. This is even true in a 64-bit system, which does allow access to more memory, but not as a direct part of Windows. Information that is larger than 2G must be swapped into the Windows TPA area from another area of memory. In future versions of Windows and the Pentium, this will most likely be changed. The current version of Windows 64 (which is now a part of Windows Vista) supports up to 8G bytes of Windows memory.

FIGURE 1–10   The memory map used by Windows XP.

FFFFFFFF

Windows Systems Area

80000000
7FFFFFFF

Windows Transient Program Area

00000000

Does this mean that any program written for Windows will begin at physical address 00000000H? No, the memory system physical map is much different for the linear programming model shown in Figure 1–10. Every process in a Windows Vista, Windows XP, or Windows 2000 system has its own set of page tables, which define where in the physical memory each 4K-byte page of the process is located. This means that the process can be located anywhere in the memory, even in noncontiguous pages. Page tables and the paging structure of the microprocessor are discussed later in this chapter and are beyond the scope of the text at this point. As far as an application is concerned, you will always have 2G bytes of memory even if the computer has less memory. The operating system (Windows) handles assigning physical memory to the application and if not enough physical memory exists, it uses the hard disk drive for any that is not available.

I/O Space.   The I/O (input/output) space in a computer system extends from I/O port 0000H to port FFFFH. (An **I/O port address** is similar to a memory address, except that instead of addressing memory, it addresses an I/O device.) The I/O devices allow the microprocessor to communicate between itself and the outside world. The I/O space allows the computer to access up to 64K different 8-bit I/O devices, 32K different 16-bit devices, or 16K different 32-bit devices. The 64-bit extensions support the same I/O space and I/O sizes as the 32-bit version and does not add 64-bit I/O devices to the system. A great number of these locations are available for expansion in most computer systems. Figure 1–11 shows the I/O map found in many personal computer systems. To view the map on your computer in Windows, go to the Control Panel, Performance and Maintenance, System, Hardware tab, Device Manager, View tab, then select resources by type and click on the plus next to Input/Output (I/O).

FIGURE 1–11 Some I/O locations in a typical personal computer.

```
□··■■■■ Input/output (IO)
    │··· 🔲 [00000000 - 0000000F] Direct memory access controller
    │··· 🔲 [00000000 - 00000CF7] PCI bus
    │··· 🔲 [00000010 - 0000001F] Motherboard resources
    │··· 🔲 [00000020 - 00000021] Programmable interrupt controller
    │··· 🔲 [00000022 - 0000002D] Motherboard resources
    │··· 🔲 [0000002E - 0000002F] Motherboard resources
    │··· 🔲 [00000030 - 0000003F] Motherboard resources
    │··· 🔲 [00000040 - 00000043] System timer
    │··· 🔲 [00000044 - 0000005F] Motherboard resources
    │··· ⌨ [00000060 - 00000060] Easy Internet Keyboard
    │··· 🔲 [00000061 - 00000061] System speaker
    │··· 🔲 [00000062 - 00000063] Motherboard resources
    │··· ⌨ [00000064 - 00000064] Easy Internet Keyboard
    │··· 🔲 [00000065 - 0000006F] Motherboard resources
    │··· 🔲 [00000070 - 00000073] System CMOS/real time clock
    │··· 🔲 [00000074 - 0000007F] Motherboard resources
    │··· 🔲 [00000080 - 00000090] Direct memory access controller
    │··· 🔲 [00000091 - 00000093] Motherboard resources
    │··· 🔲 [00000094 - 0000009F] Direct memory access controller
    │··· 🔲 [000000A0 - 000000A1] Programmable interrupt controller
    │··· 🔲 [000000A2 - 000000BF] Motherboard resources
    │··· 🔲 [000000C0 - 000000DF] Direct memory access controller
    │··· 🔲 [000000E0 - 000000EF] Motherboard resources
    │··· 🔲 [000000F0 - 000000FF] Numeric data processor
    │··· 💾 [00000170 - 00000177] Secondary IDE Channel
    │··· 💾 [000001F0 - 000001F7] Primary IDE Channel
    │··· 🎮 [00000200 - 00000207] Standard Game Port
    │··· 🔲 [00000274 - 00000277] ISAPNP Read Data Port
    │··· 🔲 [00000279 - 00000279] ISAPNP Read Data Port
    │··· 🔌 [000002F8 - 000002FF] Communications Port (COM2)
    │··· 💾 [00000376 - 00000376] Secondary IDE Channel
    │··· 🔌 [00000378 - 0000037F] Printer Port (LPT1)
    │··· 📺 [000003B0 - 000003BB] ALL-IN-WONDER 9700 SERIES
    │··· 📺 [000003B0 - 000003BB] Intel(R) 82845G/GL/GE/PE/GV Processor to AGP Controller - 2561
    │··· 📺 [000003C0 - 000003DF] ALL-IN-WONDER 9700 SERIES
    │··· 📺 [000003C0 - 000003DF] Intel(R) 82845G/GL/GE/PE/GV Processor to AGP Controller - 2561
    │··· 🔲 [000003F0 - 000003F1] Motherboard resources
    │··· 💾 [000003F2 - 000003F5] Standard floppy disk controller
    │··· 💾 [000003F6 - 000003F6] Primary IDE Channel
    │··· 💾 [000003F7 - 000003F7] Standard floppy disk controller
    │··· 🔌 [000003F8 - 000003FF] Communications Port (COM1)
    │··· 🔲 [000004D0 - 000004D1] Motherboard resources
    │··· 🔲 [000004D6 - 000004D6] Motherboard resources
    │··· 🔲 [00000A79 - 00000A79] ISAPNP Read Data Port
    │··· 🔲 [00000D00 - 0000FFFF] PCI bus
    │··· 🎮 [00008400 - 0000843F] SoundMAX Integrated Digital Audio
    └··· 🎮 [00008800 - 000088FF] SoundMAX Integrated Digital Audio
```

The I/O area contains two major sections. The area below I/O location 0400H is considered reserved for system devices; many are depicted in Figure 1–11. The remaining area is available I/O space for expansion that extends from I/O port 0400H through FFFFH. Generally, I/O addresses between 0000H and 00FFH address components on the main board of the computer, while addresses between 0100H and 03FFH address devices located on plug-in cards (or on the main board). Note that the limitation of I/O addresses between 0000 and 03FFH comes from the original PC standard, as specified by IBM. When using the ISA bus, you must only use addresses between 0000H and 03FFH. The PCI bus uses I/O address between 0400H and FFFFH.

Various I/O devices that control the operation of the system are usually not directly addressed. Instead, the system BIOS ROM addresses these basic devices, which can vary slightly in location and function from one computer to the next. Access to most I/O devices should always be made

through Windows, DOS, or BIOS function calls to maintain compatibility from one computer system to another. The map shown in Figure 1–11 is provided as a guide to illustrate the I/O space in the system.

The Microprocessor

At the heart of the microprocessor-based computer system is the microprocessor integrated circuit. The microprocessor, sometimes referred to as the CPU (**central processing unit**), is the controlling element in a computer system. The microprocessor controls memory and I/O through a series of connections called buses. The buses select an I/O or memory device, transfer data between an I/O device or memory and the microprocessor, and control the I/O and memory system. Memory and I/O are controlled through instructions that are stored in the memory and executed by the microprocessor.

The microprocessor performs three main tasks for the computer system: (1) data transfer between itself and the memory or I/O systems, (2) simple arithmetic and logic operations, and (3) program flow via simple decisions. Although these are simple tasks, it is through them that the microprocessor performs virtually any series of operations or tasks.

The power of the microprocessor is in its capability to execute billions of millions of instructions per second from a program or software (**group of instructions**) stored in the memory system. This stored program concept has made the microprocessor and computer system very powerful devices. (Recall that Babbage also wanted to use the stored program concept in his Analytical Engine.)

Table 1–4 shows the arithmetic and logic operations executed by the Intel family of microprocessors. These operations are very basic, but through them, very complex problems are solved. Data are operated upon from the memory system or internal registers. Data widths are variable and include a **byte** (8 bits), **word** (16 bits), and **doubleword** (32 bits). Note that only the 80386 through the Core2 directly manipulate 8-, 16-, and 32-bit numbers. The earlier 8086–80286 directly manipulated 8- and 16-bit numbers, but not 32-bit numbers. Beginning with the 80486, the microprocessor contained a numeric coprocessor that allowed it to perform complex arithmetic using floating-point arithmetic. The numeric coprocessor, which is similar to a calculator chip, was an additional component in the 8086- through the 80386-based personal computer. The numeric coprocessor is also capable of performing integer operations on **quadwords** (64 bits). The MMX and SIMD units inside the Pentium through Core2 function with integers and floating-point number in parallel. The SIMD unit requires numbers stored as **octalwords** (128 bits).

Another feature that makes the microprocessor powerful is its ability to make simple decisions based upon numerical facts. For example, a microprocessor can decide if a number is zero, if it is positive, and so forth. These simple decisions allow the microprocessor to modify the

TABLE 1–4 Simple arithmetic and logic operations.

Operation	Comment
Addition	
Subtraction	
Multiplication	
Division	
AND	Logical multiplication
OR	Logic addition
NOT	Logical inversion
NEG	Arithmetic inversion
Shift	
Rotate	

TABLE 1–5 Decisions found in the 8086 through Core2 microprocessors.

Decision	Comment
Zero	Test a number for zero or not-zero
Sign	Test a number for positive or negative
Carry	Test for a carry after addition or a borrow after subtraction
Parity	Test a number for an even or an odd number of ones
Overflow	Test for an overflow that indicates an invalid result after a signed addition or a signed subtraction

program flow, so that programs appear to think through these simple decisions. Table 1–5 lists the decision-making capabilities of the Intel family of microprocessors.

Buses.   A bus is a common group of wires that interconnect components in a computer system. The buses that interconnect the sections of a computer system transfer address, data, and control information between the microprocessor and its memory and I/O systems. In the microprocessor-based computer system, three buses exist for this transfer of information: address, data, and control. Figure 1–12 shows how these buses interconnect various system components such as the microprocessor, read/write memory (RAM), read-only memory (ROM or flash), and a few I/O devices.

The address bus requests a memory location from the memory or an I/O location from the I/O devices. If I/O is addressed, the address bus contains a 16-bit I/O address from 0000H through FFFFH. The 16-bit I/O address, or port number, selects one of 64K different I/O devices. If memory is addressed, the address bus contains a memory address, which varies in width with the different versions of the microprocessor. The 8086 and 8088 address 1M byte of memory, using a 20-bit address that selects locations 00000H–FFFFFH. The 80286 and 80386SX address 16M bytes of memory using a 24-bit address that selects locations 000000H–FFFFFFH. The 80386SL, 80386SLC, and 80386EX address 32M bytes of memory, using 25-bit address that selects locations 0000000H–1FFFFFFH. The 80386DX, 80486SX, and 80486DX address
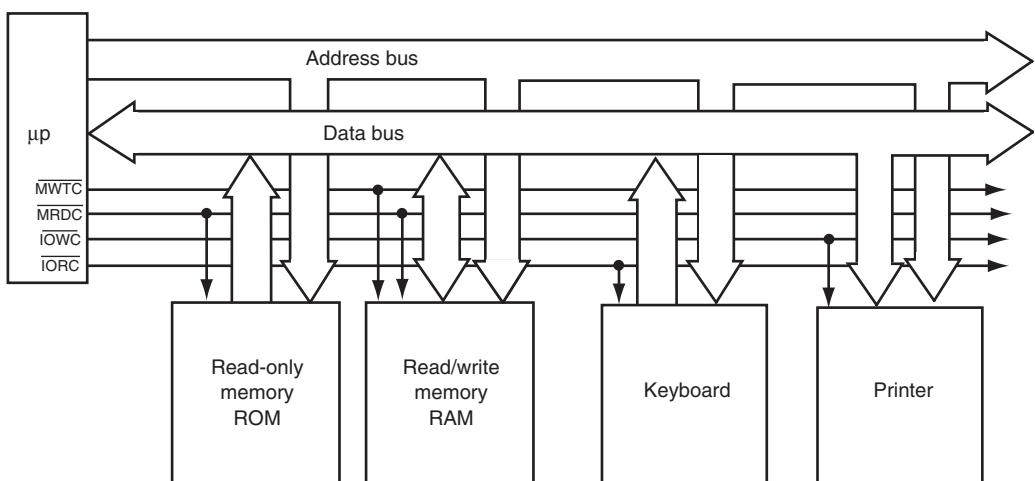
FIGURE 1–12   The block diagram of a computer system showing the address, data, and control bus structure.

TABLE 1–6 The Intel family of microprocessor bus and memory sizes.

Microprocessor	Data Bus Width	Address Bus Width	Memory Size
8086	16	20	1M
8088	8	20	1M
80186	16	20	1M
80188	8	20	1M
80286	16	24	16M
80386SX	16	24	16M
80386DX	32	32	4G
80386EX	16	26	64M
80486	32	32	4G
Pentium	64	32	4G
Pentium Pro–Core2	64	32	4G
Pentium Pro–Core2 (if extended addressing is enabled)	64	36	64G
Pentium 4 and Core2 with 64-bit extensions enabled	64	40	1T
Itanium	128	40	1T

4G bytes of memory, using a 32-bit address that selects locations 00000000H–FFFFFFFFH. The Pentium also addresses 4G bytes of memory, but it uses a 64-bit data bus to access up to 8 bytes of memory at a time. The Pentium Pro through Core2 microprocessors have a 64-bit data bus and a 32-bit address bus that address 4G of memory from location 00000000H–FFFFFFFFH, or a 36-bit address bus that addresses 64G of memory at locations 000000000H–FFFFFFFFFH, depending on their configuration. Refer to Table 1–6 for complete listing of bus and memory sizes of the Intel family of microprocessors.

The 64-bit extensions to the Pentium family provide 40 address pins in its current version that allow up to 1T byte of memory to be accessed through its 10 digit hexadecimal address. Note that 2^{40} is 1 terra. In future editions of the 64-bit microprocessors Intel plans to expand the number of address bits to 52, and ultimately to 64 bits. A 52-bit address bus allows 4P (Peta) bytes of memory to be accessed and a 64-bit address bus allows 16E (Exa) bytes of memory.

The data bus transfers information between the microprocessor and its memory and I/O address space. Data transfers vary in size, from 8 bits wide to 64 bits wide in various members of the Intel microprocessor family. For example, the 8088 has an 8-bit data bus that transfers 8 bits of data at a time. The 8086, 80286, 80386SL, 80386SX, and 80386EX transfer 16 bits of data through their data buses; the 80386DX, 80486SX, and 80486DX transfer 32 bits of data; and the Pentium through Core2 microprocessors transfer 64 bits of data. The advantage of a wider data bus is speed in applications that use wide data. For example, if a 32-bit number is stored in memory, it takes the 8088 microprocessor four transfer operations to complete because its data bus is only 8 bits wide. The 80486 accomplishes the same task with one transfer because its data bus is 32 bits wide. Figure 1–13 shows the memory widths and sizes of the 8086–80486 and Pentium through Core2 microprocessors. Notice how the memory sizes and organizations differ between various members of the Intel microprocessor family. In all family members, the memory is numbered by byte. Notice that the Pentium through Core2 microprocessors all contain a 64-bit-wide data bus.

The control bus contains lines that select the memory or I/O and cause them to perform a read or write operation. In most computer systems, there are four control bus connections: $\overline{MRDC}$ (**memory read control**), $\overline{MWTC}$ (**memory write control**), $\overline{IORC}$ (**I/O read control**), and $\overline{IOWC}$ (**I/O write control**). Note that the overbar indicates that the control signal is active-low; that is,

FIGURE 1–13 *(continued)*

1–3 NUMBER SYSTEMS

The use of the microprocessor requires a working knowledge of binary, decimal, and hexadecimal numbering systems. This section of the text provides a background for those who are unfamiliar with these numbering systems. Conversions between decimal and binary, decimal and hexadecimal, and binary and hexadecimal are described.

Digits

Before numbers are converted from one number base to another, the digits of a number system must be understood. Early in our education, we learned that a decimal (base 10) number is constructed with 10 digits: 0 through 9. The first digit in any numbering system is always zero. For example, a base 8 (**octal**) number contains 8 digits: 0 through 7; a base 2 (**binary**) number contains 2 digits: 0 and 1. If the base of a number exceeds 10, the additional digits use the

letters of the alphabet, beginning with an A. For example, a base 12 number contains 10 digits: 0 through 9, followed by A for 10 and B for 11. Note that a base 10 number does contain a *10* digit, just as a base 8 number does not contain an *8* digit. The most common numbering systems used with computers are decimal, binary, and hexadecimal (base 16). (Many years ago octal numbers were popular.) Each of these number systems are described and used in this section the chapter.

Positional Notation

Once the digits of a number system are understood, larger numbers are constructed by using positional notation. In grade school, we learned that the position to the left of the units position is the tens position, the position to the left of the tens position is the hundreds position, and so forth. (An example is the decimal number 132: This number has 1 hundred, 3 tens, and 2 units.) What probably was not learned was the exponential value of each position: The units position has a weight of 10^0, or 1; the tens position has weight of 10^1, or 10; and the hundreds position has a weight of 10^2, or 100. The exponential powers of the positions are critical for understanding numbers in other numbering systems. The position to the left of the radix (**number base**) point, called a decimal point only in the decimal system, is always the units position in any number system. For example, the position to the left of the binary point is always 2^0, or 1; the position to the left of the octal point is 8^0, or 1. In any case, any number raised to its zero power is always 1, or the units position.

The position to the left of the units position is always the number base raised to the first power; in a decimal system, this is 10^1, or 10. In a binary system, it is 2^1, or 2; and in an octal system, it is 8^1, or 8. Therefore, an 11 decimal has a different value from an 11 binary. The decimal number is composed of 1 ten plus 1 unit, and has a value of 11 units; while the binary number 11 is composed of 1 two plus 1 unit, for a value of 3 decimal units. The 11 octal has a value of 9 decimal units.

In the decimal system, positions to the right of the decimal point have negative powers. The first digit to the right of the decimal point has a value of 10^{-1}, or 0.1. In the binary system the first digit to the right of the binary point has a value of 2^{-1}, or 0.5. In general, the principles that apply to decimal numbers also apply to numbers in any other number system.

Example 1–1 shows 110.101 in binary (often written as 110.101_2). It also shows the power and weight or value of each digit position. To convert a binary number to decimal, add weights of each digit to form its decimal equivalent. The 110.101_2 is equivalent to a 6.625 in decimal (4 + 2 + 0.5 + 0.125). Notice that this is the sum of 2^2 (or 4) plus 2^1 (or 2), but 2^0 (or 1) is not added because there are no digits under this position. The fraction part is composed of 2^{-1} (.5) plus 2^{-3} (or .125), but there is no digit under the 2^{-2} (or .25) so .25 is not added.

EXAMPLE 1–1

```
Power               2²      2¹     2⁰     2⁻¹     2⁻²      2⁻³
Weight              4       2      1      .5      .25      .125
Number              1       1      0   .  1       0        1
Numeric Value       4   +   2  +   0  +  .5   +   0    +   .125   =   6.625
```

Suppose that the conversion technique is applied to a base 6 number, such as 25.2_6. Example 1–2 shows this number placed under the powers and weights of each position. In the example, there is a 2 under 6^1, which has a value of 12_{10} (2 × 6), and a 5 under 6^0, which has a value of 5 (5 × 1). The whole number portion has a decimal value of 12 + 5, or 17. The number to the right of the hex point is a 2 under 6^{-1}, which has a value of .333 (2 × .167). The number 25.2_6, therefore, has a value of 17.333.

EXAMPLE 1–2

```
Power                 6¹    6⁰   6⁻¹
Weight                6     1    .167
Number                2     5    .2
Numeric Value         12 + 5 + .333 = 17.333
```

Conversion to Decimal

The prior examples have shown that to convert from any number base to decimal, determine the weights or values of each position of the number, and then sum the weights to form the decimal equivalent. Suppose that a 125.7_8 octal is converted to decimal. To accomplish this conversion, first write down the weights of each position of the number. This appears in Example 1–3. The value of 125.7_8 is 85.875 decimal, or 1×64 plus 2×8 plus 5×1 plus $7 \times .125$.

EXAMPLE 1–3

```
Power                 8²    8¹   8⁰   8⁻¹
Weight                64    8    1    .125
Number                1     2    5    .7
Numeric Value         64 + 16 + 5 + .875 = 85.875
```

Notice that the weight of the position to the left of the units position is 8. This is 8 times 1. Then notice that the weight of the next position is 64, or 8 times 8. If another position existed, it would be 64 times 8, or 512. To find the weight of the next higher-order position, multiply the weight of the current position by the number base (or 8, in this example). To calculate the weights of position to the right of the radix point, divide by the number base. In the octal system, the position immediately to the right of the octal point is $1/8$, or .125. The next position is $^{.125}/_8$, or .015625, which can also be written as $^1/_{64}$. Also note that the number in Example 1–3 can also be written as the decimal number $85^7/_8$.

Example 1–4 shows the binary number 11011.0111 written with the weights and powers of each position. If these weights are summed, the value of the binary number converted to decimal is 27.4375.

EXAMPLE 1–4

```
Power          2⁴    2³   2²   2¹   2⁰   2⁻¹  2⁻²   2⁻³    2⁻⁴
Weight         16    8    4    2    1    .5   .25   .125   .0625
Number         1     1    0    1    1 .  0    1     1      1
Numeric Value  16 + 8 + 0 + 2 + 1 + 0 + .25 + .125 + .0625 = 27.4375
```

It is interesting to note that 2^{-1} is also $1/2$, 2^{-2} is $1/4$, and so forth. It is also interesting to note that 2^{-4} is $1/16$, or .0625. The fractional part of this number is $7/16$ or .4375 decimal. Notice that 0111 is a 7 in binary code for the numerator and the rightmost one is in the $1/16$ position for the denominator. Other examples: The binary fraction of .101 is $5/8$ and the binary fraction of .001101 is $^{13}/_{64}$.

Hexadecimal numbers are often used with computers. A 6A.CH (**H** for hexadecimal) is illustrated with its weights in Example 1–5. The sum of its digits is 106.75, or 106¾. The whole number part is represented with 6×16 plus 10 (A) $\times 1$. The fraction part is 12 (C) as a numerator and 16 (16^{-1}) as the denominator, or $^{12}/_{16}$, which is reduced to $^3/_4$.

EXAMPLE 1–5

```
Power             16¹  16⁰  16⁻¹
Weight            16   1    .0625
Number            6    A  . C
Number Value      96 + 10 + .75 = 106.75
```

Conversion from Decimal

Conversions from decimal to other number systems are more difficult to accomplish than conversion to decimal. To convert the whole number portion of a number to decimal, divide by 1 radix. To convert the fractional portion, multiply by the radix.

Whole Number Conversion from Decimal.   To convert a decimal whole number to another number system, divide by the radix and save the remainders as significant digits of the result. An algorithm for this conversion as is follows:

1. Divide the decimal number by the radix (number base).
2. Save the remainder (first remainder is the least significant digit).
3. Repeat steps 1 and 2 until the quotient is zero.

For example, to convert a 10 decimal to binary, divide it by 2. The result is 5, with a remainder of 0. The first remainder is the units position of the result (in this example, a 0). Next divide the 5 by 2. The result is 2, with a remainder of 1. The 1 is the value of the twos (2^1) position. Continue the division until the quotient is a zero. Example 1–6 shows this conversion process. The result is written as 1010_2 from the bottom to the top.

EXAMPLE 1–6

```
2) 10       remainder = 0
 2)  5      remainder = 1
  2)  2     remainder = 0
   2)  1    remainder = 1         result = 1010
       0
```

To convert a 10 decimal into base 8, divide by 8, as shown in Example 1–7. A 10 decimal is a 12 octal.

EXAMPLE 1–7

```
8) 10       remainder = 2
 8)  1      remainder = 1         result = 12
     0
```

Conversion from decimal to hexadecimal is accomplished by dividing by 16. The remainders will range in value from 0 through 15. Any remainder of 10 through 15 is then converted to the letters A through F for the hexadecimal number. Example 1–8 shows the decimal number 109 converted to a 6DH.

EXAMPLE 1–8

```
16)  109    remainder = 13 (D)
 16)   6    remainder = 6        result = 6D
       0
```

Converting from a Decimal Fraction.   Conversion from a decimal fraction to another number base is accomplished with multiplication by the radix. For example, to convert a decimal fraction into binary, multiply by 2. After the multiplication, the whole number portion of the result is

saved as a significant digit of the result, and the fractional remainder is again multiplied by the radix. When the fraction remainder is zero, multiplication ends. Note that some numbers are never-ending (repetend). That is, a zero is never a remainder. An algorithm for conversion from a decimal fraction is as follows:

1. Multiply the decimal fraction by the radix (number base).
2. Save the whole number portion of the result (even if zero) as a digit. Note that the first result is written immediately to the right of the radix point.
3. Repeat steps 1 and 2, using the fractional part of step 2 until the fractional part of step 2 is zero.

Suppose that a .125 decimal is converted to binary. This is accomplished with multiplications by 2, as illustrated in Example 1–9. Notice that the multiplication continues until the fractional remainder is zero. The whole number portions are written as the binary fraction (0.001) in this example.

EXAMPLE 1–9

```
    .125
x      2
    0.25     digit is 0

    .25
x      2
    0.5      digit is 0

     .5
x      2
    1.0      digit is 1     result = 0.001₂
```
$$\text{result} = 0.001_2$$

This same technique is used to convert a decimal fraction into any number base. Example 1–10 shows the same decimal fraction of .125 from Example 1–9 converted to octal by multiplying by 8.

EXAMPLE 1–10

```
    .125
x      8
    1.0      digit is 1     result = 0.1₈
```
$$\text{result} = 0.1_8$$

Conversion to a hexadecimal fraction appears in Example 1–11. Here, the decimal .046875 is converted to hexadecimal by multiplying by 16. Note that .046875 is 0.0CH.

EXAMPLE 1–11

```
    .046875
x        16
      0.75     digit is 0

   .75
x    16
   12.0        digit is 12(C)   result = 0.0C₁₆
```
$$\text{result} = 0.0C_{16}$$

Binary-Coded Hexadecimal

Binary-coded hexadecimal (BCH) is used to represent hexadecimal data in binary code. A binary-coded hexadecimal number is a hexadecimal number written so that each digit is represented by a 4-bit binary number. The values for the BCH digits appear in Table 1–7.

Hexadecimal numbers are represented in BCH code by converting each digit to BCH code with a space between each coded digit. Example 1–12 shows a 2AC converted to BCH code. Note that each BCH digit is separated by a space.

TABLE 1–7 Binary-coded hexadecimal (BCH) code.

Hexadecimal Digit	BCH Code
0	0000
1	0001
2	0010
3	0011
4	0100
5	0101
6	0110
7	0111
8	1000
9	1001
A	1010
B	1011
C	1100
D	1101
E	1110
F	1111

EXAMPLE 1–12

```
2AC = 0010 1010 1100
```

The purpose of BCH code is to allow a binary version of a hexadecimal number to be written in a form that can easily be converted between BCH and hexadecimal. Example 1–13 shows a BCH coded number converted back to hexadecimal code.

EXAMPLE 1–13

```
1000 0011 1101 . 1110 = 83D.E
```

Complements

At times, data are stored in complement form to represent negative numbers. There are two systems that are used to represent negative data: **radix** and **radix − 1** complements. The earliest system was the radix −1 complement, in which each digit of the number is subtracted from the radix −1 to generate the radix −1 complement to represent a negative number.

Example 1–14 shows how the 8-bit binary number 01001100 is one's (radix −1) complemented to represent it as a negative value. Notice that each digit of the number is subtracted from one to generate the radix −1 (one's) complement. In this example, the negative of 01001100 is 10110011. The same technique can be applied to any number system, as illustrated in Example 1–15, in which the fifteen's (radix −1) complement of a 5CD hexadecimal is computed by subtracting each digit from a fifteen.

EXAMPLE 1–14

```
    1111 1111
 −  0100 1100
    1011 0011
```

EXAMPLE 1–15

```
   15 15 15
 −  5  C  D
    A  3  2
```

Today, the radix −1 complement is not used by itself; it is used as a step for finding the radix complement. The radix complement is used to represent negative numbers in modem computer systems. (The radix −1 complement was used in the early days of computer technology.) The main problem with the radix −1 complement is that a negative or a positive zero exists; in the radix complement system, only a positive zero can exist.

To form the radix complement, first find the radix −1 complement, and then add a one to the result. Example 1–16 shows how the number 0100 1000 is converted to a negative value by two's (radix) complementing it.

EXAMPLE 1–16

```
  1111 1111
− 0100 1000
  1011 0111   (one's complement)
+         1
  1011 1000   (two's complement)
```

To prove that a 0100 1000 is the inverse (negative) of a 1011 1000, add the two together to form an 8-digit result. The ninth digit is dropped and the result is zero because a 0100 1000 is a positive 72, while a 1011 1000 is a negative 72. The same technique applies to any number system. Example 1–17 shows how the inverse of a 345 hexadecimal is found by first fifteen's complementing the number, and then by adding one to the result to form the sixteen's complement. As before, if the original 3-digit number 345 is added to the inverse of CBB, the result is a 3-digit 000. As before, the fourth bit (carry) is dropped. This proves that 345 is the inverse of CBB. Additional information about one's and two's complements is presented with signed numbers in the next section of the text.

EXAMPLE 1–17

```
  15 15 15
−  3  4  5
   C  B  A   (fifteen's complement)
+         1
   C  B  B   (sixteen's complement)
```

1–4   COMPUTER DATA FORMATS

Successful programming requires a precise understanding of data formats. In this section, many common computer data formats are described as they are used with the Intel family of microprocessors. Commonly, data appear as ASCII, Unicode, BCD, signed and unsigned integers, and floating-point numbers (real numbers). Other forms are available, but are not presented here because they are not commonly found.

ASCII and Unicode Data

ASCII (**American Standard Code for Information Interchange**) data represent alphanumeric characters in the memory of a computer system (see Table 1–8). The standard ASCII code is a 7-bit code, with the eighth and most significant bit used to hold parity in some antiquated systems. If ASCII data are used with a printer, the most significant bits are a 0 for alphanumeric printing and 1 for graphics printing. In the personal computer, an extended ASCII character set is selected by placing a 1 in the leftmost bit. Table 1–9 shows the extended ASCII character set, using code 80H–FFH. The extended ASCII characters store some foreign letters and punctuation, Greek

TABLE 1–8   ASCII code.

							Second									
First	X0	X1	X2	X3	X4	X5	X6	X7	X8	X9	XA	XB	XC	XD	XE	XF
0X	NUL	SOH	STX	ETX	EOT	ENQ	ACK	BEL	BS	HT	LF	VT	FF	CR	SO	SI
1X	DLE	DC1	DC2	DC3	DC4	NAK	SYN	ETB	CAN	EMS	SUB	ESC	FS	GS	RS	US
2X	SP	!	"	#	$	%	&	'	(	)	*	+	,	–	.	/
3X	0	1	2	3	4	5	6	7	8	9	:	;	<	=	>	?
4X	@	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
5X	P	Q	R	S	T	U	V	W	X	Y	Z	[	\	]	^	_
6X	`	a	b	c	d	e	f	g	h	i	j	k	l	m	n	o
7X	p	q	r	s	t	u	v	w	x	y	z	{	|	}	~	⠿

TABLE 1–9   Extended ASCII code, as printed by the IBM ProPrinter.

First	X0	X1	X2	X3	X4	X5	X6	X7	X8	X9	XA	XB	XC	XD	XE	XF
0X		☺	☻	♥	♦	♣	♠	●	◘	○	◙	♂	♀	♪	♫	☼
1X	►	◄	↕	‼	¶	§	▬	↨	↑	↓	→	←	∟	↔	▲	▼
8X	Ç	ü	é	â	ä	à	å	ç	ê	ë	è	ï	î	ì	Ä	Å
9X	É	æ	Æ	ô	ö	ò	û	ù	ÿ	Ö	Ü	¢	£	¥	₧	ƒ
AX	á	í	ó	ú	ñ	Ñ	ª	º	¿	⌐	¬	½	¼	¡	«	»
BX	░	▒	▓	│	┤	╡	╢	╖	╕	╣	║	╗	╝	╜	╛	┐
CX	└	┴	┬	├	─	┼	╞	╟	╚	╔	╩	╦	╠	═	╬	╧
DX	╨	╤	╥	╙	╘	╒	╓	╫	╪	┘	┌	█	▄	▌	▐	▀
EX	α	β	Γ	π	Σ	σ	µ	γ	Φ	Θ	Ω	δ	∞	φ	∈	∩
FX	≡	±	≥	≤	⌠	⌡	÷	≈	°	∙	·	√	ⁿ	²	■	

characters, mathematical characters, box-drawing characters, and other special characters. Note that extended characters can vary from one printer to another. The list provided is designed to be used with the IBM ProPrinter, which also matches the special character set found with most word processors.

The ASCII control characters, also listed in Table 1–8, perform control functions in a computer system, including clear screen, backspace, line feed, and so on. To enter the control codes through the computer keyboard, hold down the Control key while typing a letter. To obtain the control code 01H, type a Control-A; a 02H is obtained by a Control-B, and so on. Note that the control codes appear on the screen, from the DOS prompt, as ^A for Control-A, ^B for Control-B, and so forth. Also note that the carriage return code (CR) is the Enter key on most modem keyboards. The purpose of CR is to return the cursor or print head to the left margin. Another code that appears in many programs is the line feed code (LF), which moves the cursor down one line.

To use Table 1–8 or 1–9 for converting alphanumeric or control characters into ASCII characters, first locate the alphanumeric code for conversion. Next, find the first digit of the hexadecimal ASCII code. Then find the second digit. For example, the capital letter "A" is ASCII code 41H, and the lowercase letter "a" is ASCII code 61H. Many Windows-based applications, since Windows 95, use the **Unicode** system to store alphanumeric data. This system

stores each character as 16-bit data. The codes 0000H–00FFH are the same as standard ASCII code. The remaining codes, 0100H–FFFFH, are used to store all special characters from many worldwide character sets. This allows software written for the Windows environment to be used in many countries around the world. For complete information on Unicode, visit *http://www.unicode.org*.

ASCII data are most often stored in memory by using a special directive to the assembler program called define byte(s), or DB. (The assembler is a program that is used to program a computer in its native binary machine language.) An alternative to DB is the word BYTE. The DB and BYTE directives, and several examples of their usage with ASCII-coded character strings, are listed in Example 1–18. Notice how each character string is surrounded by apostrophes (')—never use the quote ("). Also notice that the assembler lists the ASCII-coded value for each character to the left of the character string. To the far left is the hexadecimal memory location where the character string is first stored in the memory system. For example, the character string WHAT is stored beginning at memory address 001D, and the first letter is stored as 57 (W), followed by 68 (H), and so forth. Example 1–19 shows the same three strings defined as String^ character strings for use with Visual C++ Express 2005 and 2008. Note that Visual C++ uses quotes to surround strings. If an earlier version of C++ is used, then the string is defined with a CString for Microsoft Visual C++ instead of a String^. The ^ symbol indicates that String is a member of the garbage collection heap for managing the storage. A garbage collector cleans off the memory system (frees unused memory) when the object falls from visibility or scope in a C++ program and it also prevents memory leaks.

EXAMPLE 1–18

```
0000    42 61 72 72 79    NAMES   DB   'Barry B. Brey'
        20 42 2E 20 42
        72 65 79
OOOD    57 68 65 20 63    MESS    DB   'Where can it be?'
        20 63 61 6E 20
        69 74 20 62 65
        3F
001D    57 69 20 74 20    WHAT    DB   'What is on first.'
        69 73 20 6F 6E
        20 66 69 72 73
        74 2E
```

EXAMPLE 1–19

```
String^ NAMES = "Barry B. Brey"   // C++ Express version

String^ MESS = "Where can it be?"

String^ WHAT = "What is on first."
```

BCD (Binary-Coded Decimal) Data

Binary-coded decimal (BCD) information is stored in either packed or unpacked forms. Packed BCD data are stored as two digits per byte and unpacked BCD data are stored as one digit per byte. The range of a BCD digit extends from 0000_2 to 1001_2, or 0–9 decimal. Unpacked BCD data are returned from a keypad or keyboard. Packed BCD data are used for some of the instructions included for BCD addition and subtraction in the instruction set of the microprocessor.

Table 1–10 shows some decimal numbers converted to both the packed and unpacked BCD forms. Applications that require BCD data are point-of-sales terminals and almost any device that performs a minimal amount of simple arithmetic. If a system requires complex arithmetic, BCD data are seldom used because there is no simple and efficient method of performing complex BCD arithmetic.

TABLE 1–10   Packed and unpacked BCD data.

Decimal	Packed		Unpacked		
12	0001 0010		0000 0001	0000 0010	
623	0000 0110	0010 0011	0000 0110	0000 0010	0000 0011
910	0000 1001	0001 0000	0000 1001	0000 0001	0000 0000

Example 1–20 shows how to use the assembler to define both packed and unpacked BCD data. Example 1–21 shows how to do this using Visual C++ and char or bytes. In all cases, the convention of storing the least-significant data first is followed. This means that to store 83 into memory, the 3 is stored first, and then followed by the 8. Also note that with packed BCD data, the letter H (hexadecimal) follows the number to ensure that the assembler stores the BCD value rather than a decimal value for packed BCD data. Notice how the numbers are stored in memory as unpacked, one digit per byte; or packed, as two digits per byte.

EXAMPLE 1–20

```
                    ;Unpacked BCD data (least-significant data first)
                    ;
0000 03 04 05       NUMB1 DB    3,4,5               ;defines number 543
0003 07 08          NUMB2 DB    7,8                 ;defines number 87
                    ;
                    ;Packed BCD data (least-significant data first)
                    ;
0005 37 34          NUMB3 DB    37H,34H             ;defines number 3437
0007 03 45          NUMB4 DB    3,45H               ;defines number 4503
```

EXAMPLE 1–21

```
                    //Unpacked BCD data (least-significant data first)
                    //
                    char Numb1 = 3,4,5;         ;defines number 543
                    char Numb2 = 7,8            ;defines number 87
                    //
                    //Packed BCD data (least-significant data first)
                    //
                    char Numb3 = 0x37,0x34      ;defines number 3437
                    char Numb4 = 3,0x45         ;defines number 4503
```

Byte-Sized Data

Byte-sized data are stored as *unsigned* and *signed* integers. Figure 1–14 illustrates both the unsigned and signed forms of the byte-sized integer. The difference in these forms is the weight of the leftmost bit position. Its value is 128 for the unsigned integer and minus 128 for the signed integer. In the signed integer format, the leftmost bit represents the sign bit of the number, as well as a weight of minus 128. For example, 80H represents a value of 128 as an unsigned number; as a signed number, it represents a value of minus 128. Unsigned integers range in value from 00H to FFH (0–255). Signed integers range in value from −128 to 0 to + 127.

Although negative signed numbers are represented in this way, they are stored in the two's complement form. The method of evaluating a signed number by using the weights of each bit position is much easier than the act of two's complementing a number to find its value. This is especially true in the world of calculators designed for programmers.
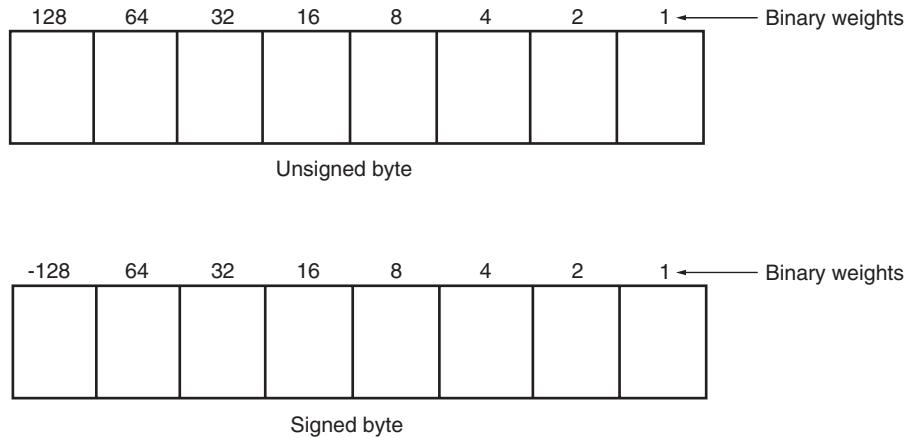
FIGURE 1–14   The unsigned and signed bytes illustrating the weights of each binary-bit position.

Whenever a number is two's complemented, its sign changes from negative to positive or positive to negative. For example, the number 00001000 is a +8. Its negative value (−8) is found by two's complementing the +8. To form a two's complement, first one's complement the number. To one's complement a number, invert each bit of a number from zero to one or from one to zero. Once the one's complement is formed, the two's complement is found by adding a one to the one's complement. Example 1–22 shows how numbers are two's complemented using this technique.

EXAMPLE 1–22

```
+ 8 = 00001000
      11110111   (one's complement)
+            1
− 8 = 11111000   (two's complement)
```

Another, and probably simpler, technique for two's complementing a number starts with the rightmost digit. Start by writing down the number from right to left. Write the number exactly as it appears until the first one. Write down the first one, and then invert all bits to its left. Example 1–23 shows this technique with the same number as in Example 1–22.

EXAMPLE 1–23

```
+8 = 00001000
        1000 (write number to first 1)
     1111     (invert the remaining bits)
-8 = 11111000
```

To store 8-bit data in memory using the assembler program, use the DB directive as in prior examples or char as in Visual C++ examples. Example 1–24 lists many forms of 8-bit numbers stored in memory using the assembler program. Notice in the example that a hexadecimal number is defined with the letter H following the number, and that a decimal number is written as is, without anything special. Example 1–25 shows the same byte data defined for use with a Visual C++ program. In C/C++ the hexadecimal value is preceded by a 0x to indicate a hexadecimal value.

EXAMPLE 1–24

```
            ;Unsigned byte-sized data
            ;
0000  FE    DATA1 DB   254   ;define 254 decimal
0001  87    DATA2 DB   87H   ;define 87 hexadecimal
0002  47    DATA3 DB   71    ;define 71 decimal
            ;
            ;Signed byte-sized data
            ;
0003  9C    DATA4 DB   -100  ;define -100 decimal
0004  64    DARA5 DB   +100  ;define +100 decimal
0005  FF    DATA6 DB   -1    ;define -1 decimal
0006  38    DATA7 DB   56    ;define 56 decimal
```

EXAMPLE 1–25

```
//Unsigned byte-sized data
//
unsigned char Data1 = 254;      //define 254 decimal
unsigned char Data2 = 0x87;     //define 87 hexadecimal
unsigned char Data3 = 71        //define 71 decimal
//
//Signed byte-sized data
//
char Data4 = -100;              //define -100 decimal
char Data5 = +100;              //define +100 decimal
char Data6 = -1;                //define -1 decimal
char Data7 = 56;                //define 56 decimal
```

Word-Sized Data

A word (16-bits) is formed with two bytes of data. The least significant byte is always stored in the lowest-numbered memory location, and the most significant byte is stored in the highest. This method of storing a number is called the **little endian** format. An alternate method, not used with the Intel family of microprocessors, is called the **big endian** format. In the big endian format, numbers are stored with the lowest location containing the most significant data. The big endian format is used with the Motorola family of microprocessors. Figure 1–15 (a) shows the weights of each bit position in a word of data, and Figure 1–15 (b) shows how the number 1234H appears when stored in the memory locations 3000H and 3001H. The only difference between a signed and an unsigned word in the leftmost bit is position. In the unsigned form, the leftmost bit is unsigned and has a weight of 32,768; in the signed form, its weight is −32,768. As with byte-sized signed data, the signed word is in two's complement form when representing a negative number. Also, notice that the low-order byte is stored in the lowest-numbered memory location (3000H) and the high-order byte is stored in the highest-numbered location (3001H).

Example 1–26 shows several signed and unsigned word-sized data stored in memory using the assembler program. Example 1–27 shows how to store the same numbers in a Visual C++

EXAMPLE 1–26

```
             ;Unsigned word-sized data
             ;
0000  09F0   DATA1  DW    2544     ;define 2544 decimal
0002  87AC   DATA2  DW    87ACH    ;define 87AC hexadecimal
0004  02C6   DATA3  DW    710      ;define 710 decimal
             ;
             ;Signed word-sized data
             ;
0006  CBA8   DATA4  DW    -13400   ;define -13400 decimal
0008  00C6   DATA5  DW    +198     ;define +198 decimal
000A  FFFF   DATA6  DW    -1       ;define -1 decimal
```

(a) Unsigned word

(b) The contents of memory location 3000H and 3001H are the word 1234H.
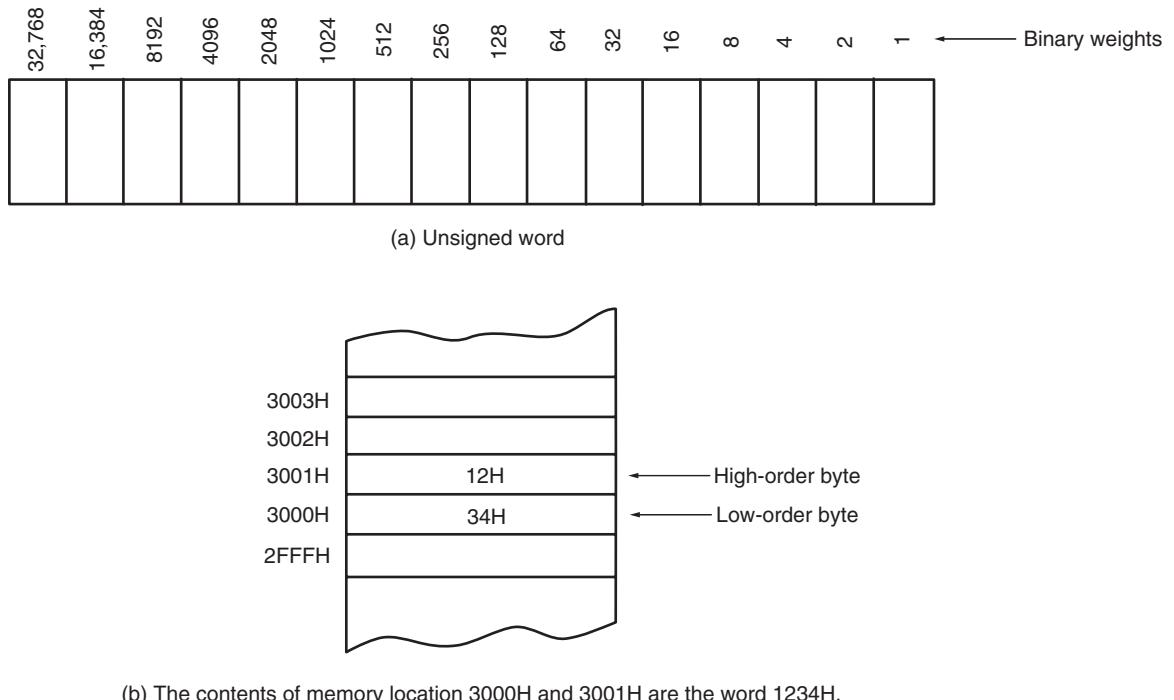
FIGURE 1–15　The storage format for a 16-bit word in (a) a register and (b) two bytes of memory.

EXAMPLE 1–27

```
//Unsigned word-sized data
//
unsigned short Data1 = 2544;    //define 2544 decimal
unsigned short Data2 = 0x87AC   //define 87AC hexadecimal
unsigned short Data3 = 710;     //define 710 decimal
//
//Signed word-sized data
//
short Data4 = -13400;           //define -13400 decimal
short Data5 = +198;             //define +198 decimal
short Data6 = -1;               //define -1 decimal
```

program (assuming version 5.0 or newer), which uses the **short** directive to store a 16-bit integer. Notice that the **define word(s)** directive, or **DW**, causes the assembler to store words in the memory instead of bytes, as in prior examples. The WORD directive can also be used to define a word. Notice that the word *data* is displayed by the assembler in the same form as entered. For example, a l000H is displayed by the assembler as a 1000. This is for our convenience because the number is actually stored in the memory as 00 l0 in two consecutive memory bytes.

Doubleword-Sized Data

Doubleword-sized data requires four bytes of memory because it is a 32-bit number. Doubleword data appear as a product after a multiplication and also as a dividend before a division. In the 80386 through the Core2, memory and registers are also 32 bits in width. Figure 1–16 shows the form used to store doublewords in the memory and the binary weights of each bit position.

When a doubleword is stored in memory, its least significant byte is stored in the lowest numbered memory location, and its most significant byte is stored in the highest-numbered

Binary weights: 2,147,438,648 1,073,741,824 536,870,912 268,435,456 134,217,728 67,108,864 33,554,432 16,777,216 8,388,608 4,194,304 2,097,152 1,048,576 524,288 262,144 131,072 65,536 32,768 16,384 8192 4096 2048 1024 512 256 128 64 32 16 8 4 2 1

(a) Unsigned doubleword

Address	Contents	
00103H	12H	← High-order byte
00102H	34H	
00101H	56H	
00100H	78H	← Low-order byte
000FFH		

(b) The contents of memory location 00100H–00103H are the doubleword 12345678H.

FIGURE 1–16   The storage format for a 32-bit word in (a) a register and (b) 4 bytes of memory.

memory location using the little endian format. Recall that this is also true for word-sized data. For example, 12345678H that is stored in memory locations 00100H–00103H is stored with the 78H in memory location 00100H, the 56H in location 00101H, the 34H in location 00102H, and the 12H in location 00103H.

To define doubleword-sized data, use the assembler directive **define doubleword(s)**, or **DD**. (You can also use the DWORD directive in place of DD.) Example 1–28 shows both signed and unsigned numbers stored in memory using the DD directive. Example 1–29 shows how to define the same doublewords in Visual C++ using the **int** directive.

EXAMPLE 1–28

```
                     ;Unsigned doubleword-sized data
                     ;
0000  0003E1C0       DATA1 DD    254400            ;define 254400 decimal
0004  87AC1234       DATA2 DD    87AC1234H         ;define 87AC1234 hexadecimal
0008  00000046       DATA3 DD    70                ;define 70 decimal
                     ;
                     ;Signed doubleword-sized data
                     ;
000C  FFEB8058       DATA4 DD    -1343400          ;define -1343400 decimal
0010  000000C6       DATA5 DD    +198              ;define +198 decimal
0014  FFFFFFFF       DATA6 DD    -1                ;define -1 decimal
```

EXAMPLE 1–29

```
//Unsigned doubleword-sized data
//
unsigned int Data1 = 254400;            //define 254400 decimal
unsigned int Data2 = 0x87AC1234;        //define 87AC1234 hexadecimal
unsigned int Data3 = 70;                //define 70 decimal
//
//Signed doubleword-sized data
//
int Data4 = -1343400;                   //define -1342400 decimal
int Data5 = +198;                       //define +198 decimal
int Data6 = -1;                         //define -1 decimal
```

Integers may also be stored in memory that is of any width. The forms listed here are standard forms, but that doesn't mean that a 256-byte wide integer can't be stored in the memory. The microprocessor is flexible enough to allow any size of data in assembly language. When non-standard-width numbers are stored in memory, the DB directive is normally used to store them. For example, the 24-bit number 123456H is stored using a DB 56H, 34H, 12H directive. Note that this conforms to the little endian format. This could also be done in Visual C++ using the char directive.

Real Numbers

Because many high-level languages use the Intel family of microprocessors, real numbers are often encountered. A real number, or a **floating-point number**, as it is often called, contains two parts: a mantissa, significand, or fraction; and an exponent. Figure 1–17 depicts both the 4- and 8-byte forms of real numbers as they are stored in any Intel system. Note that the 4-byte number is called **single-precision** and the 8-byte form is called **double-precision**. The form presented here is the same form specified by the IEEE[10] standard, IEEE-754, version 10.0. The standard has been adopted as the standard form of real numbers with virtually all programming languages and many applications packages. The standard also applies the data manipulated by the numeric coprocessor in the personal computer. Figure 1–17 (a) shows the single-precision form that contains a sign-bit, an 8-bit exponent, and a 24-bit fraction (mantissa). Note that because applications often require double-precision floating-point numbers [see Figure 1–17 (b)], the Pentium–Core2 with their 64-bit data bus perform memory transfers at twice the speed of the 80386/80486 microprocessors.

Simple arithmetic indicates that it should take 33 bits to store all three pieces of data. Not true—the 24-bit mantissa contains an **implied** (hidden) one-bit that allows the mantissa to represent 24 bits while being stored in only 23 bits. The hidden bit is the first bit of the normalized real number. When normalizing a number, it is adjusted so that its value is at least 1, but less than 2. For example, if 12 is converted to binary (1100_2), it is normalized and the result is 1.1×2^3. The whole number 1 is not stored in the 23-bit mantissa portion of the number; the 1 is the hidden one-bit. Table 1–11 shows the single-precision form of this number and others.

The exponent is stored as a biased exponent. With the single-precision form of the real number, the bias is 127 (7FH) and with the double-precision form, it is 1023 (3FFH). The bias

(a)

(b)
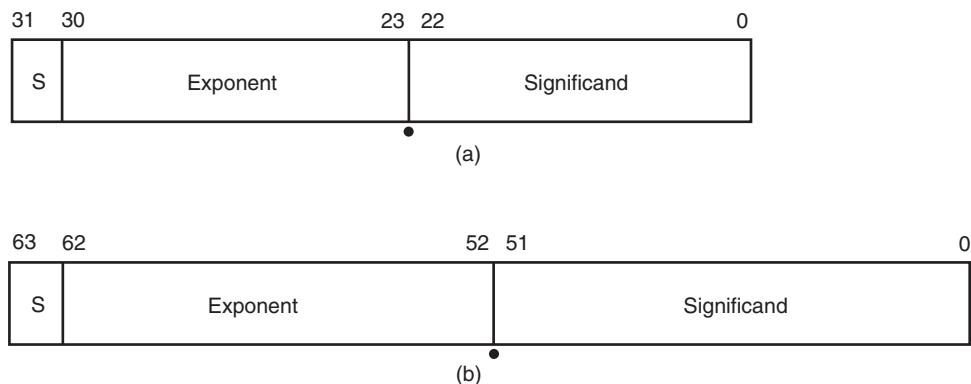
FIGURE 1–17 The floating-point numbers in (a) single-precision using a bias of 7FH and (b) double-precision using a bias of 3FFH.

---

[10]IEEE is the Institute of Electrical and Electronic Engineers.

TABLE 1–11   Single-precision real numbers.

Decimal	Binary	Normalized	Sign	Biased Exponent	Mantissa
+12	1100	1.1×2^3	0	10000010	10000000 00000000 00000000
−12	1100	1.1×2^3	1	10000010	10000000 00000000 00000000
+100	1100100	1.1001×2^6	0	10000101	10010000 00000000 00000000
−1.75	1.11	1.11×2^0	1	01111111	11000000 00000000 00000000
+0.25	0.01	1.0×2^{-2}	0	01111101	00000000 00000000 00000000
+0.0	0	0	0	00000000	00000000 00000000 00000000

and exponent are added before being stored in the exponent portion of the floating-point number. In the previous example, there is an exponent of 2^3, represented as a biased exponent of $127 + 3$ or 130 (82H) in the single-precision form, or as 1026 (402H) in the double-precision form.

There are two exceptions to the rules for floating-point numbers. The number 0.0 is stored as all zeros. The number infinity is stored as all ones in the exponent and all zeros in the mantissa. The sign-bit indicates either a positive or a negative infinity.

As with other data types, the assembler can be used to define real numbers in both single- and double-precision forms. Because single-precision numbers are 32-bit numbers, use the DD directive or use the **define quadword(s)**, or DQ, directive to define 64-bit double-precision real numbers. Optional directives for real numbers are REAL4, REAL8, and REAL10 for defining single-, double-, and extended precision real numbers. Example 1–30 shows numbers defined in real number format for the assembler. If using the inline assembler in Visual C++ single-precision numbers are defined as **float** and double-precision numbers are defined as **double** as shown in Example 1–31. There is no way to define the extended-precision floating-point number for use in Visual C++.

EXAMPLE 1–30

```
                                ;single-precision real numbers
                                ;
0000   3F9DF3B6                 NUMB1    DD     1.234      ;define 1.234
0004   C1BB3333                 NUMB2    DD     -23.4      ;define -23.4
0008   43D20000                 NUMB3    REAL4 4.2E2       ;define 420
                                ;
                                ;double-precision real numbers
                                ;
000C   405ED9999999999A         NUMB4    DQ     123.4      ;define 123.4
0014   C1BB333333333333         NUMB5    REAL8 -23.4       ;define -23.4
                                ;
                                ;Extended-precision real numbers
                                ;
001C   4005F6CCCCCCCCCCCCCD     NUMB6    REAL10 123.4      ;define 123.4
```

EXAMPLE 1–31

```
//Single-precision real numbers
//
float Numb1 = 1.234;
float Numb2 = -23.4;
float Numb3 = 4.3e2;
//
//Double-precision real numbers
//
double Numb4 = 123.4;
double Numb5 = -23.4;
```

1–5   SUMMARY

1. The mechanical computer age began with the advent of the abacus in 500 B.C. This first mechanical calculator remained unchanged until 1642, when Blaise Pascal improved it. An early mechanical computer system was the Analytical Engine developed by Charles Babbage in 1823. Unfortunately, this machine never functioned because of the inability to create the necessary machine parts.

2. The first electronic calculating machine was developed during World War II by Konrad Zuse, an early pioneer of digital electronics. His computer, the Z3, was used in aircraft and missile design for the German war effort.

3. The first electronic computer, which used vacuum tubes, was placed into operation in 1943 to break secret German military codes. This first electronic computer system, the Colossus, was invented by Alan Turing. Its only problem was that the program was fixed and could not be changed.

4. The first general-purpose, programmable electronic computer system was developed in 1946 at the University of Pennsylvania. This first modern computer was called the ENIAC (Electronics Numerical Integrator and Calculator).

5. The first high-level programming language, called FLOWMATIC, was developed for the UNIVAC I computer by Grace Hopper in the early 1950s. This led to FORTRAN and other early programming languages such as COBOL.

6. The world's first microprocessor, the Intel 4004, was a 4-bit microprocessor—a programmable controller on a chip—that was meager by today's standards. It addressed a mere 4096 4-bit memory locations. Its instruction set contained only 45 different instructions.

7. Microprocessors that are common today include the 8086/8088, which were the first 16-bit microprocessors. Following these early 16-bit machines were the 80286, 80386, 80486, Pentium, Pentium Pro, Pentium II, Pentium III, Pentium 4, and Core2 processors. The architecture has changed from 16 bits to 32 bits and, with the Itanium, to 64 bits. With each newer version, improvements followed that increased the processor's speed and performance. From all indications, this process of speed and performance improvement will continue, although the performance increases may not always come from an increased clock frequency.

8. The DOS-based personal computers contain memory systems that include three main areas: TPA (transient program area), system area, and extended memory. The TPA hold: application programs, the operating system, and drivers. The system area contains memory used for video display cards, disk drives, and the BIOS ROM. The extended memory area is only available to the 80286 through the Core2 microprocessor in an AT-style or ATX-style personal computer system. The Windows-based personal computers contain memory systems that include two main areas: TPA and systems area.

9. The 8086/8088 address 1M byte of memory from locations 00000H–FFFFFH. The 80286 and 80386SX address 16M bytes of memory from locations 000000H–FFFFFFH. The 80386SL addresses 32M bytes of memory from locations 0000000H–1FFFFFFH. The 80386DX through the Core2 address 4G bytes of memory from locations 00000000H–FFFFFFFFH. In addition, the Pentium Pro through the Core2 can operate with a 36-bit address and access up to 64G bytes of memory from locations 000000000H–FFFFFFFFFH. A Pentium 4 or Core2 operating with 64-bit extensions addresses memory from locations 0000000000H–FFFFFFFFFFH for 1T byte of memory.

10. All versions of the 8086 through the Core2 microprocessors address 64K bytes of I/O address space. These I/O ports are numbered from 0000H to FFFFH with I/O ports 0000H–03FFH reserved for use by the personal computer system. The PCI bus allows ports 0400H–FFFFH.

11. The operating system in early personal computers was either MSDOS (Microsoft disk operating system) or PCDOS (personal computer disk operating system from IBM). The operating

system performs the task of operating or controlling the computer system, along with its I/O devices. Modern computers use Microsoft Windows in place of DOS as an operating system.

12. The microprocessor is the controlling element in a computer system. The microprocessor performs data transfers, does simple arithmetic and logic operations, and makes simple decisions. The microprocessor executes programs stored in the memory system to perform complex operations in short periods of time.

13. All computer systems contain three buses to control memory and I/O. The address bus is used to request a memory location or I/O device. The data bus transfers data between the microprocessor and its memory and I/O spaces. The control bus controls the memory and I/O, and requests reading or writing of data. Control is accomplished with $\overline{IORC}$ (I/O read control), $\overline{IOWC}$ (I/O write control), $\overline{MRDC}$ (memory read control), and $\overline{MWTC}$ (memory write control).

14. Numbers are converted from any number base to decimal by noting the weights of each position. The weight of the position to the left of the radix point is always the units position in any number system. The position to the left of the units position is always the radix times one. Succeeding positions are determined by multiplying by the radix. The weight of the position to the right of the radix point is always determined by dividing by the radix.

15. Conversion from a whole decimal number to any other base is accomplished by dividing by the radix. Conversion from a fractional decimal number is accomplished by multiplying by the radix.

16. Hexadecimal data are represented in hexadecimal form or in a code called binary-coded hexadecimal (BCH). A binary-coded hexadecimal number is one that is written with a 4-bit binary number that represents each hexadecimal digit.

17. The ASCII code is used to store alphabetic or numeric data. The ASCII code is a 7-bit code; it can have an eighth bit that is used to extend the character set from 128 codes to 256 codes. The carriage return (Enter) code returns the print head or cursor to the left margin. The line feed code moves the cursor or print head down one line. Most modern applications use Unicode, which contains ASCII at codes 0000H–00FFH.

18. Binary-coded decimal (BCD) data are sometimes used in a computer system to store decimal data. These data are stored either in packed (two digits per byte) or unpacked (one digit per byte) form.

19. Binary data are stored as a byte (8 bits), word (16 bits), or doubleword (32 bits) in a computer system. These data may be unsigned or signed. Signed negative data are always stored in the two's complement form. Data that are wider than 8 bits are always stored using the little endian format. In 32-bit Visual C++ these data are represented with char (8 bits), short (16 bits) and int (32 bits).

20. Floating-point data are used in computer systems to store whole, mixed, and fractional numbers. A floating-point number is composed of a sign, a mantissa, and an exponent.

21. The assembler directives DB or BYTE define bytes, DW or WORD define words, DD or DWORD define doublewords, and DQ or QWORD define quadwords.

1–6　QUESTIONS AND PROBLEMS

1. Who developed the Analytical Engine?
2. The 1890 census used a new device called a punched card. Who developed the punched card?
3. Who was the founder of IBM Corporation?
4. Who developed the first electronic calculator?
5. The first electronic computer system was developed for what purpose?
6. The first general-purpose, programmable computer was called the _____.
7. The world's first microprocessor was developed in 1971 by _____.

8. Who was the Countess of Lovelace?
9. Who developed the first high-level programming language called FLOWMATIC?
10. What is a von Neumann machine?
11. Which 8-bit microprocessor ushered in the age of the microprocessor?
12. The 8085 microprocessor, introduced in 1977, has sold _____ copies.
13. Which Intel microprocessor was the first to address 1M bytes of memory?
14. The 80286 addresses _____ bytes of memory.
15. How much memory is available to the 80486 microprocessor?
16. When did Intel introduce the Pentium microprocessor?
17. When did Intel introduce the Pentium Pro processor?
18. When did Intel introduce the Pentium 4 microprocessor?
19. Which Intel microprocessor addresses 1T of memory?
20. What is the acronym MIPs?
21. What is the acronym CISC?
22. A binary bit stores a(n) _____or a(n) _____.
23. A computer K (pronounced kay) is equal to _____ bytes.
24. A computer M (pronounced meg) is equal to _____ K bytes.
25. A computer G (pronounced gig) is equal to _____ M bytes.
26. A computer P (pronounced peta) is equal to _____ T bytes.
27. How many typewritten pages of information are stored in a 4G-byte memory?
28. The first 1M byte of memory in a DOS-based computer system contains a(n) _____ and a(n) _____area.
29. How large is the Windows application programming area?
30. How much memory is found in the DOS transient program area?
31. How much memory is found in the Windows systems area?
32. The 8086 microprocessor addresses _____ bytes of memory.
33. The Core2 microprocessor addresses _____ bytes of memory.
34. Which microprocessors address 4G bytes of memory?
35. Memory above the first 1M byte is called _____ memory.
36. What is the system BIOS?
37. What is DOS?
38. What is the difference between an XT and an AT computer system?
39. What is the VESA local bus?
40. The ISA bus holds _____-bit interface cards.
41. What is the USB?
42. What is the AGP?
43. What is the XMS?
44. What is the SATA interface and where is it used in a system?
45. A driver is stored in the _____ area.
46. The personal computer system addresses _____ bytes of I/O space.
47. What is the purpose of the BIOS?
48. Draw the block diagram of a computer system.
49. What is the purpose of the microprocessor in a microprocessor-based computer?
50. List the three buses found in all computer systems.
51. Which bus transfers the memory address to the I/O device or to the memory?
52. Which control signal causes the memory to perform a read operation?
53. What is the purpose of the $\overline{IORC}$ signal?
54. If the $\overline{MRDC}$ signal is a logic 0, which operation is performed by the microprocessor?
55. Define the purpose of the following assembler directives:
    (a) DB
    (b) DQ

    (c) DW
    (d) DD
56. Define the purpose of the following 32-bit Visual C++ directives:
    (a) char
    (b) short
    (c) int
    (d) float
    (e) double
57. Convert the following binary numbers into decimal:
    (a) 1101.01
    (b) 111001.0011
    (c) 101011.0101
    (d) 111.0001
58. Convert the following octal numbers into decimal:
    (a) 234.5
    (b) 12.3
    (c) 7767.07
    (d) 123.45
    (e) 72.72
59. Convert the following hexadecimal numbers into decimal:
    (a) A3.3
    (b) 129.C
    (c) AC.DC
    (d) FAB.3
    (e) BB8.0D
60. Convert the following decimal integers into binary, octal, and hexadecimal:
    (a) 23
    (b) 107
    (c) 1238
    (d) 92
    (e) 173
61. Convert the following decimal numbers into binary, octal, and hexadecimal:
    (a) 0.625
    (b) .00390625
    (c) .62890625
    (d) 0.75
    (e) .9375
62. Convert the following hexadecimal numbers into binary-coded hexadecimal code (BCH):
    (a) 23
    (b) AD4
    (c) 34.AD
    (d) BD32
    (e) 234.3
63. Convert the following binary-coded hexadecimal numbers into hexadecimal:
    (a) 1100 0010
    (b) 0001 0000 1111 1101
    (c) 1011 1100
    (d) 0001 0000
    (e) 1000 1011 1010
64. Convert the following binary numbers to the one's complement form:
    (a) 1000 1000
    (b) 0101 1010

    (c) 0111 0111

    (d) 1000 0000

65. Convert the following binary numbers to the two's complement form:

    (a) 1000 0001

    (b) 1010 1100

    (c) 1010 1111

    (d) 1000 0000

66. Define byte, word, and doubleword.

67. Convert the following words into ASCII-coded character strings:

    (a) FROG

    (b) Arc

    (c) Water

    (d) Well

68. What is the ASCII code for the Enter key and what is its purpose?

69. What is the Unicode?

70. Use an assembler directive to store the ASCII-character string 'What time is it?' in the memory.

71. Convert the following decimal numbers into 8-bit signed binary numbers:

    (a) +32

    (b) −12

    (c) +100

    (d) −92

72. Convert the following decimal numbers into signed binary words:

    (a) +1000

    (b) −120

    (c) +800

    (d) −3212

73. Use an assembler directive to store byte-sized −34 into the memory.

74. Create a byte-sized variable called Fred1 and store a −34 in it in Visual C++.

75. Show how the following 16-bit hexadecimal numbers are stored in the memory system (use the standard Intel little endian format):

    (a) 1234H

    (b) A122H

    (c) B100H

76. What is the difference between the big endian and little endian formats for storing numbers that are larger than 8 bits in width?

77. Use an assembler directive to store a 123A hexadecimal into memory.

78. Convert the following decimal numbers into both packed and unpacked BCD forms:

    (a) 102

    (b) 44

    (c) 301

    (d) 1000

79. Convert the following binary numbers into signed decimal numbers:

    (a) 10000000

    (b) 00110011

    (c) 10010010

    (d) 10001001

80. Convert the following BCD numbers (assume that these are packed numbers) to decimal numbers:

    (a) 10001001

    (b) 00001001

    (c) 00110010

    (d) 00000001

81. Convert the following decimal numbers into single-precision floating-point numbers:
    (a) +1.5
    (b) −10.625
    (c) +100.25
    (d) −1200
82. Convert the following single-precision floating-point numbers into decimal numbers:
    (a) 0 10000000 11000000000000000000000
    (b) 1 01111111 00000000000000000000000
    (c) 0 10000010 10010000000000000000000
83. Use the Internet to write a short report about any one of the following computer pioneers:
    (a) Charles Babbage
    (b) Konrad Zuse
    (c) Joseph Jacquard
    (d) Herman Hollerith
84. Use the Internet to write a short report about any one of the following computer languages:
    (a) COBOL
    (b) ALGOL
    (c) FORTRAN
    (d) PASCAL
85. Use the Internet to write a short report detailing the features of the Itanium 2 microprocessor.
86. Use the Internet to detail the Intel 45 nm (nanometer) fabrication technology.

CHAPTER 2

The Microprocessor and Its Architecture

INTRODUCTION

This chapter presents the microprocessor as a programmable device by first looking at its internal programming model and then how its memory space is addressed. The architecture of the family of Intel microprocessors is presented simultaneously, as are the ways that the family members address the memory system.

The addressing modes for this powerful family of microprocessors are described for the real, protected, and flat modes of operation. Real mode memory (DOS memory) exists at locations 00000H–FFFFFH, the first 1M byte of the memory system, and is present on all versions of the microprocessor. Protected mode memory (Windows memory) exists at any location in the entire protected memory system, but is available only to the 80286–Core2, not to the earlier 8086 or 8088 microprocessors. Protected mode memory for the 80286 contains 16M bytes; for the 80386–Pentium, 4G bytes; and for the Pentium Pro through the Core2, either 4G or 64G bytes. With the 64-bit extensions enabled, the Pentium 4 and Core2 address 1T byte of memory in a flat memory model. Windows Vista or Windows 64 is needed to operate the Pentium 4 or Core2 in 64-bit mode using the flat mode memory to access the entire 1T byte of memory.

CHAPTER OBJECTIVES

Upon completion of this chapter, you will be able to:

1. Describe the function and purpose of each program-visible register in the 8086–Core2 microprocessors, including the 64-bit extensions.
2. Detail the flag register and the purpose of each flag bit.
3. Describe how memory is accessed using real mode memory-addressing techniques.
4. Describe how memory is accessed using protected mode memory-addressing techniques.
5. Describe how memory is accessed using the 64-bit flat memory model.
6. Describe the program-invisible registers found within the 80286 through Core2 microprocessors.
7. Detail the operation of the memory-paging mechanism.

2–1   INTERNAL MICROPROCESSOR ARCHITECTURE

Before a program is written or any instruction investigated, the internal configuration of the microprocessor must be known. This section of the chapter details the program-visible internal architecture of the 8086–Core2 microprocessors. Also detailed are the function and purpose of each of these
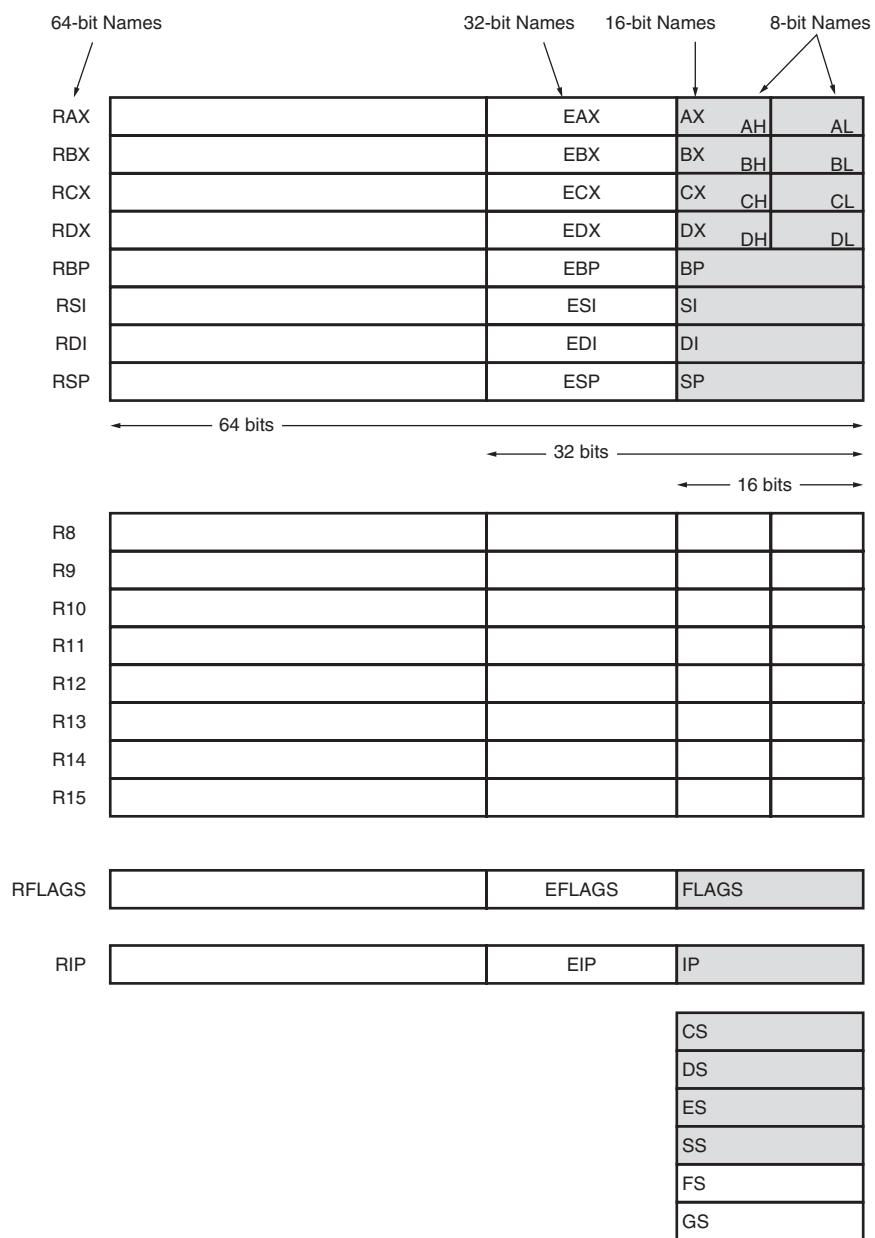
internal registers. Note that in a multiple core microprocessor each core contains the same programming model. The only difference is that each core runs a separate task or thread simultaneously.

The Programming Model

The programming model of the 8086 through the Core2 is considered to be **program visible** because its registers are used during application programming and are specified by the instructions. Other registers, detailed later in this chapter, are considered to be **program invisible** because they are not addressable directly during applications programming, but may be used indirectly during system programming. Only the 80286 and above contain the program-invisible registers used to control and operate the protected memory system and other features of the microprocessor.

Figure 2–1 illustrates the programming model of the 8086 through the Core2 microprocessor including the 64-bit extensions. The earlier 8086, 8088, and 80286 contain 16-bit internal

FIGURE 2–1 The programming model of the 8086 through the Core2 microprocessor including the 64-bit extensions.

architectures, a subset of the registers shown in Figure 2–1. The 80386 through the Core2 microprocessors contain full 32-bit internal architectures. The architectures of the earlier 8086 through the 80286 are fully upward-compatible to the 80386 through the Core2. The shaded areas in this illustration represent registers that are found in early versions of the 8086, 8088, or 80286 microprocessors and are provided on the 80386–Core2 microprocessors for compatibility to the early versions

The programming model contains 8-, 16-, and 32-bit registers. The Pentium 4 and Core2 also contain 64-bit registers when operated in the 64-bit mode as illustrated in the programming model. The 8-bit registers are AH, AL, BH, BL, CH, CL, DH, and DL and are referred to when an instruction is formed using these two-letter designations. For example, an ADD AL,AH instruction adds the 8-bit contents of AH to AL. (Only AL changes due to this instruction.) The 16-bit registers are AX, BX, CX, DX, SP, BP, DI, SI, IP, FLAGS, CS, DS, ES, SS, FS, and GS. Note that the first 4 16 registers contain a pair of 8-bit registers. An example is AX, which contains AH and AL. The 16-bit registers are referenced with the two-letter designations such as AX. For example, an ADD DX, CX instruction adds the 16-bit contents of CX to DX. (Only DX changes due to this instruction.) The extended 32-bit registers are EAX, EBX, ECX, EDX, ESP, EBP, EDI, ESI, EIP, and EFLAGS. These 32-bit extended registers, and 16-bit registers FS and GS, are available only in the 80386 and above. The 16-bit registers are referenced by the designations FS or GS for the two new 16-bit registers, and by a three-letter designation for the 32-bit registers. For example, an ADD ECX, EBX instruction adds the 32-bit contents of EBX to ECX. (Only ECX changes due to this instruction.)

Some registers are general-purpose or multipurpose registers, while some have special purposes. The multipurpose registers include EAX, EBX, ECX, EDX, EBP, EDI, and ESI. These registers hold various data sizes (bytes, words, or doublewords) and are used for almost any purpose, as dictated by a program.

The 64-bit registers are designated as RAX, RBX, and so forth. In addition to the renaming of the registers for 64-bit widths, there are also additional 64-bit registers that are called R8 through R15. The 64-bit extensions have multiplied the available register space by more than 8 times in the Pentium 4 and the Core2 when compared to the original microprocessor architecture as indicated in the shaded area in Figure 2–1. An example 64-bit instruction is ADD RCX, RBX, instruction, which adds the 64-bit contents of RBX to RCX. (Only RCX changes due to this instruction.) One difference exists: these additional 64-bit registers (R8 through R15) are addressed as a byte, word, doubleword, or **quadword**, but only the rightmost 8 bits is a byte. R8 through R15 have no provision for directly addressing bits 8 through 15 as a byte. In the 64-bit mode, a legacy high byte register (AH, BH, CH, or DH) cannot be addressed in the same instruction with an R8 through R15 byte. Because legacy software does not access R8 through R15, this causes no problems with existing 32-bit programs, which function without modification.

Table 2–1 shows the overrides used to access portions of a 64-bit register. To access the low-order byte of the R8 register, use R8B (where B is the low-order byte). Likewise, to access the low-order word of a numbered register, such as R10, use R10W in the instruction. The letter D is used to access a doubleword. An example instruction that copies the low-order doubleword from R8 to R11 is MOV R11D, R8D. There is no special letter for the entire 64-bit register.

TABLE 2–1 Flat mode 64-bit access to numbered registers.

Register Size	Override	Bits Accessed	Example
8 bits	B	7–0	MOV R9B, R10B
16 bits	W	15–0	MOV R10W, AX
32 bits	D	31–0	MOV R14D, R15D
64 bits	—	63–0	MOV R13, R12

Multipurpose Registers

RAX **(accumulator)**	RAX is referenced as a 64-bit register (RAX), a 32-bit register (EAX), a 16-bit register (AX), or as either of two 8-bit registers (AH and AL). Note that if an 8- or 16-bit register is addressed, only that portion of the 32-bit register changes without affecting the remaining bits. The accumulator is used for instructions such as multiplication, division, and some of the adjustment instructions. For these instructions, the accumulator has a special purpose, but is generally considered to be a multipurpose register. In the 80386 and above, the EAX register may also hold the offset address of a location in the memory system. In the 64-bit Pentium 4 and Core2, RAX holds a 64-bit offset address, which allows 1T (terra) byte of memory to be accessed through a 40-bit address bus. In the future, Intel plans to expand the address bus to 52 bits to address 4P (peta) bytes of memory.
RBX **(base index)**	RBX is addressable as RBX, EBX, BX, BH, or BL. The BX register sometimes holds the offset address of a location in the memory system in all versions of the microprocessor. In the 80386 and above, EBX also can address memory data. In the 64-bit Pentium 4 and Core2, RBX can also address memory data.
RCX **(count)**	RCX, which is addressable as RCX, ECX, CX, CH, or CL, is a general-purpose register that also holds the count for various instructions. In the 80386 and above, the ECX register also can hold the offset address of memory data. In the 64-bit Pentium 4, RCX can also address memory data. Instructions that use a count are the repeated string instructions (REP/REPE/REPNE); and shift, rotate, and LOOP/LOOPD instructions. The shift and rotate instructions use CL as the count, the repeated string instructions use CX, and the LOOP/LOOPD instructions use either CX or ECX. If operated in the 64-bit mode, LOOP uses the 64-bit RCX register for the loop counter.
RDX **(data)**	RDX, which is addressable as RDX, EDX, DX, DH, or DL, is a general-purpose register that holds a part of the result from a multiplication or part of the dividend before a division. In the 80386 and above, this register can also address memory data.
RBP **(base pointer)**	RBP, which is addressable as RBP, EBP, or BP, points to a memory location in all versions of the microprocessor for memory data transfers.
RDI **(destination index)**	RDI, which is addressable as RDI, EDI, or DI, often addresses string destination data for the string instructions.
RSI **(source index)**	RSI is used as RSI, ESI, or SI. The source index register often addresses source string data for the string instructions. Like RDI, RSI also functions as a general-purpose register. As a 16-bit register, it is addressed as SI; as a 32-bit register, it is addressed as ESI; and as a 64-bit register, it is addressed as RSI.
R8 through R15	These registers are only found in the Pentium 4 and Core2 if 64-bit extensions are enabled. As mentioned, data in these registers are addressed as 64-, 32-, 16-, or 8-bit sizes and are of general purpose. Most applications will not use these registers until 64-bit processors are common. Please note that the 8-bit portion is the rightmost 8-bit only; bits 8 to 15 are not directly addressable as a byte.
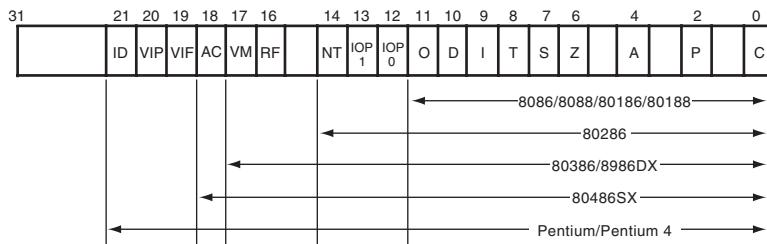
FIGURE 2–2 The EFLAG and FLAG register counts for the entire 8086 and Pentium microprocessor family.

Special-Purpose Registers. The special-purpose registers include RIP, RSP, and RFLAGS; and the segment registers include CS, DS, ES, SS, FS, and GS.

RIP (instruction pointer)
RIP addresses the next instruction in a section of memory defined as a code segment. This register is IP (16 bits) when the microprocessor operates in the real mode and EIP (32 bits) when the 80386 and above operate in the protected mode. Note that the 8086, 8088, and 80286 do not contain an EIP register and only the 80286 and above operate in the protected mode. The instruction pointer, which points to the next instruction in a program, is used by the microprocessor to find the next sequential instruction in a program located within the code segment. The instruction pointer can be modified with a jump or a call instruction. In the 64-bit mode, RIP contains a 40-bit address at present to address a 1T flat address space.

RSP (stack pointer)
RSP addresses an area of memory called the stack. The stack memory stores data through this pointer and is explained later in the text with the instructions that address stack data. This register is referred to as SP if used as a 16-bit register and ESP if referred to as a 32-bit register.

RFLAGS
RFLAGS indicate the condition of the microprocessor and control its operation. Figure 2–2 shows the flag registers of all versions of the microprocessor. (Note the flags are upward-compatible from the 8086/8088 through the Core2 microprocessors.) The 8086–80286 contain a FLAG register (16 bits) and the 80386 and above contain an EFLAG register (32-bit extended flag register). The 64-bit RFLAGS contain the EFLAG register, which is unchanged in the 64-bit version.

The rightmost five flag bits and the overflow flag change after many arithmetic and logic instructions execute. The flags never change for any data transfer or program control operation. Some of the flags are also used to control features found in the microprocessor. Following is a list of each flag bit, with a brief description of their function. As instructions are introduced in subsequent chapters, additional detail on the flag bits is provided. The rightmost five flags and the overflow flag are changed by most arithmetic and logic operations, although data transfers do not affect them.

C (carry)
Carry holds the carry after addition or the borrow after subtraction. The carry flag also indicates error conditions, as dictated by some programs and procedures. This is especially true of the DOS function calls.

P (parity)
Parity is a logic 0 for odd parity and a logic 1 for even parity. Parity is the count of ones in a number expressed as even or odd. For example, if a number contains three binary one bits, it has odd parity. If a number contains no one bits, it has even parity. The parity flag finds little application in modern programming and was implemented in early Intel microprocessors for checking data in data communications environments. Today parity checking is often accomplished by the data communications equipment instead of the microprocessor.

A (auxiliary carry)	The auxiliary carry holds the carry (half-carry) after addition or the borrow after subtraction between bit positions 3 and 4 of the result. This highly specialized flag bit is tested by the DAA and DAS instructions to adjust the value of AL after a BCD addition or subtraction. Otherwise, the A flag bit is not used by the microprocessor or any other instructions.
Z (zero)	The zero flag shows that the result of an arithmetic or logic operation is zero. If Z = 1, the result is zero; if Z = 0, the result is not zero. This may be confusing, but that is how Intel decided to name this flag.
S (sign)	The sign flag holds the arithmetic sign of the result after an arithmetic or logic instruction executes. If S = 1, the sign bit (leftmost bit of a number) is set or negative; if S = 0, the sign bit is cleared or positive.
T (trap)	The trap flag enables trapping through an on-chip debugging feature. (A program is debugged to find an error or bug.) If the T flag is enabled (1), the microprocessor interrupts the flow of the program on conditions as indicated by the debug registers and control registers. If the T flag is a logic 0, the trapping (debugging) feature is disabled. The Visual C++ debugging tool uses the trap feature and debug registers to debug faulty software.
I (interrupt)	The interrupt flag controls the operation of the INTR (interrupt request) input pin. If I = 1, the INTR pin is enabled; if I = 0, the INTR pin is disabled. The state of the I flag bit is controlled by the STI (set I flag) and CLI (clear I flag) instructions.
D (direction)	The direction flag selects either the increment or decrement mode for the DI and/or SI registers during string instructions. If D = 1, the registers are automatically decremented; if D = 0, the registers are automatically incremented. The D flag is set with the STD (set direction) and cleared with the CLD (clear direction) instructions.
O (overflow)	Overflows occur when signed numbers are added or subtracted. An overflow indicates that the result has exceeded the capacity of the machine. For example, if 7FH (+127) is added—using an 8-bit addition—to 01H (+1), the result is 80H (–128). This result represents an overflow condition indicated by the overflow flag for signed addition. For unsigned operations, the overflow flag is ignored.
IOPL (I/O privilege level)	IOPL is used in protected mode operation to select the privilege level for I/O devices. If the current privilege level is higher or more trusted than the IOPL, I/O executes without hindrance. If the IOPL is lower than the current privilege level, an interrupt occurs, causing execution to suspend. Note that an IOPL of 00 is the highest or most trusted and an IOPL of 11 is the lowest or least trusted.
NT (nested task)	The nested task flag indicates that the current task is nested within another task in protected mode operation. This flag is set when the task is nested by software.
RF (resume)	The resume flag is used with debugging to control the resumption of execution after the next instruction.
VM (virtual mode)	The VM flag bit selects virtual mode operation in a protected mode system. A virtual mode system allows multiple DOS memory partitions that are 1M byte in length to coexist in the memory

system. Essentially, this allows the system program to execute multiple DOS programs. VM is used to simulate DOS in the modern Windows environment.

AC **(alignment check)**	The alignment check flag bit activates if a word or doubleword is addressed on a non-word or non-doubleword boundary. Only the 80486SX microprocessor contains the alignment check bit that is primarily used by its companion numeric coprocessor, the 80487SX, for synchronization.
VIF **(virtual interrupt)**	The VIF is a copy of the interrupt flag bit available to the Pentium–Pentium 4 microprocessors.
VIP (virtual **interrupt pending)**	VIP provides information about a virtual mode interrupt for the Pentium–Pentium 4 microprocessors. This is used in multitasking environments to provide the operating system with virtual interrupt flags and interrupt pending information.
ID (identification)	The ID flag indicates that the Pentium–Pentium 4 microprocessors support the CPUID instruction. The CPUID instruction provides the system with information about the Pentium microprocessor, such as its version number and manufacturer.

Segment Registers. Additional registers, called segment registers, generate memory addresses when combined with other registers in the microprocessor. There are either four or six segment registers in various versions of the microprocessor. A segment register functions differently in the real mode when compared to the protected mode operation of the microprocessor. Details on their function in real and protected mode are provided later in this chapter. In the 64-bit flat model, segment registers have little use in a program except for the code segment register. Following is a list of each segment register, along with its function in the system:

CS (code)	The code segment is a section of memory that holds the code (programs and procedures) used by the microprocessor. The code segment register defines the starting address of the section of memory holding code. In real mode operation, it defines the start of a 64K-byte section of memory; in protected mode, it selects a descriptor that describes the starting address and length of a section of memory holding code. The code segment is limited to 64K bytes in the 8088–80286, and 4G bytes in the 80386 and above when these microprocessors operate in the protected mode. In the 64-bit mode, the code segment register is still used in the flat model, but its use differs from other programming modes as explained in Section 2-5.
DS (data)	The data segment is a section of memory that contains most data used by a program. Data are accessed in the data segment by an offset address or the contents of other registers that hold the offset address. As with the code segment and other segments, the length is limited to 64K bytes in the 8086–80286, and 4G bytes in the 80386 and above.
ES (extra)	The extra segment is an additional data segment that is used by some of the string instructions to hold destination data.
SS (stack)	The stack segment defines the area of memory used for the stack. The stack entry point is determined by the stack segment and stack pointer registers. The BP register also addresses data within the stack segment.

FS and GS   The FS and GS segments are supplemental segment registers available in the 80386–Core2 microprocessors to allow two additional memory segments for access by programs. Windows uses these segments for internal operations, but no definition of their usage is available.

2–2   REAL MODE MEMORY ADDRESSING

The 80286 and above operate in either the real or protected mode. Only the 8086 and 8088 operate exclusively in the real mode. In the 64-bit operation mode of the Pentium 4 and Core2, there is no real mode operation. This section of the text details the operation of the microprocessor in the real mode. **Real mode operation** allows the microprocessor to address only the first 1M byte of memory space—even if it is the Pentium 4 or Core2 microprocessor. Note that the first 1M byte of memory is called the **real memory**, **conventional memory**, or **DOS memory** system. The DOS operating system requires that the microprocessor operates in the real mode. Windows does not use the real mode. Real mode operation allows application software written for the 8086/8088, which only contains 1M byte of memory, to function in the 80286 and above without changing the software. The upward compatibility of software is partially responsible for the continuing success of the Intel family of microprocessors. In all cases, each of these microprocessors begins operation in the real mode by default whenever power is applied or the microprocessor is reset. Note that if the Pentium 4 or Core2 operate in the 64-bit mode, it cannot execute real mode applications; hence, DOS applications will not execute in the 64-bit mode unless a program that emulates DOS is written for the 64-bit mode.

Segments and Offsets
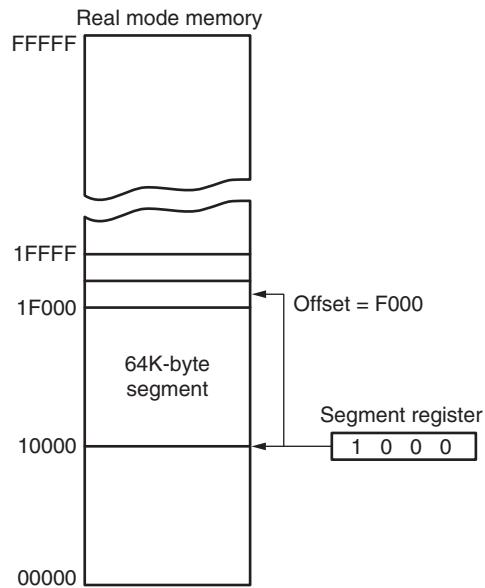
A combination of a segment address and an offset address accesses a memory location in the real mode. All real mode memory addresses must consist of a segment address plus an offset address. The **segment address**, located within one of the segment registers, defines the beginning address of any 64K-byte memory segment. The **offset address** selects any location within the 64K byte memory segment. Segments in the real mode always have a length of 64K bytes. Figure 2–3 shows how the **segment plus offset** addressing scheme selects a memory location. This illustration shows a memory segment that begins at location 10000H and ends at location IFFFFH—64K bytes in length. It also shows how an offset address, sometimes called a **displacement**, of F000H selects location 1F000H in the memory system. Note that the offset or displacement is the distance above the start of the segment, as shown in Figure 2–3.

The segment register in Figure 2–3 contains 1000H, yet it addresses a starting segment at location 10000H. In the real mode, each segment register is internally appended with a **0H** on its rightmost end. This forms a 20-bit memory address, allowing it to access the start of a segment. The microprocessor must generate a 20-bit memory address to access a location within the first 1M of memory. For example, when a segment register contains 1200H, it addresses a 64K-byte memory segment beginning at location 12000H. Likewise, if a segment register contains 1201H, it addresses a memory segment beginning at location 12010H. Because of the internally appended 0H, real mode segments can begin only at a l6-byte boundary in the memory system. This l6-byte boundary is often called a **paragraph**.

Because a real mode segment of memory is 64K in length, once the beginning address is known, the **ending address** is found by adding FFFFH. For example, if a segment register contains 3000H, the first address of the segment is 30000H, and the last address is 30000H + FFFFH or 3FFFFH. Table 2–2 shows several examples of segment register contents and the starting and ending addresses of the memory segments selected by each segment address.

The offset address, which is a part of the address, is added to the start of the segment to address a memory location within the memory segment. For example, if the segment address is

THE MICROPROCESSOR AND ITS ARCHITECTURE

FIGURE 2–3 The real
mode memory-addressing
scheme, using a segment
address plus an offset.

1000H and the offset address is 2000H, the microprocessor addresses memory location 12000H.
The offset address is always added to the starting address of the segment to locate the data. The
segment and offset address is sometimes written as 1000:2000 for a segment address of 1000H
with an offset of 2000H.

In the 80286 (with special external circuitry) and the 80386 through the Pentium 4, an extra
64K minus 16 bytes of memory is addressable when the segment address is FFFFH and the
HIMEM.SYS driver for DOS is installed in the system. This area of memory (0FFFF0H–
10FFEFH) is referred to as **high memory**. When an address is generated using a segment address
of FFFFH, the A20 address pin is enabled (if supported in older systems) when an offset is added.
For example, if the segment address is FFFFH and the offset address is 4000H, the machine
addresses memory location FFFF0H + 4000H or 103FF0H. Notice that the A20 address line is
the one in address 103FF0H. If A20 is not supported, the address is generated as 03FF0H because
A20 remains a logic zero.

Some addressing modes combine more than one register and an offset value to form an
offset address. When this occurs, the sum of these values may exceed FFFFH. For example, the
address accessed in a segment whose segment address is 4000H and whose offset address is
specified as the sum of F000H plus 3000H will access memory location 42000H instead of loca-
tion 52000H. When the F000H and 3000H are added, they form a l6-bit (**modulo 16**) sum of
2000H used as the offset address; not 12000H, the true sum. Note that the carry of 1
(F000H + 3000H = 12000H) is dropped for this addition to form the offset address of 2000H.
The address is generated as 4000:2000 or 42000H.

TABLE 2–2 Example
of real mode segment
addresses.

Segment Register	Starting Address	Ending Address
2000H	20000H	2FFFFH
2001H	20010H	3000FH
2100H	21000H	30FFFH
AB00H	AB000H	BAFFFH
1234H	12340H	2233FH

Default Segment and Offset Registers

The microprocessor has a set of rules that apply to segments whenever memory is addressed. These rules, which apply in the real and protected mode, define the segment register and offset register combination. For example, the code segment register is always used with the instruction pointer to address the next instruction in a program. This combination is **CS:IP** or **CS:EIP**, depending upon the microprocessor's mode of operation. The **code segment** register defines the start of the code segment and the **instruction pointer** locates the next instruction within the code segment. This combination (CS:IP or CS:EIP) locates the next instruction executed by the microprocessor. For example, if CS = 1400H and IP/EIP = 1200H, the microprocessor fetches its next instruction from memory location 14000H + 1200H or 15200H.

Another of the default combinations is the **stack**. Stack data are referenced through the stack segment at the memory location addressed by either the stack pointer (SP/ESP) or the pointer (BP/EBP). These combinations are referred to as SS:SP (SS:ESP), or SS:BP (SS:EBP). For example, if SS = 2000H and BP = 3000H, the microprocessor addresses memory location 23000H for the stack segment memory location. Note that in real mode, only the rightmost 16 bits of the extended register address a location within the memory segment. In the 80386–Pentium 4, never place a number larger than FFFFH into an offset register if the microprocessor is operated in the real mode. This causes the system to halt and indicate an addressing error.

Other defaults are shown in Table 2–3 for addressing memory using any Intel microprocessor with 16-bit registers. Table 2–4 shows the defaults assumed in the 80386 and above using 32-bit registers. Note that the 80386 and above have a far greater selection of segment/offset address combinations than do the 8086 through the 80286 microprocessors.

The 8086–80286 microprocessors allow four memory segments and the 80386–Core2 microprocessors allow six memory segments. Figure 2–4 shows a system that contains four memory segments. Note that a memory segment can touch or even overlap if 64K bytes of memory are not required for a segment. Think of segments as windows that can be moved over any area of memory to access data or code. Also note that a program can have more than four or six segments, but only access four or six segments at a time.

Suppose that an application program requires 1000H bytes of memory for its code, 190H bytes of memory for its data, and 200H bytes of memory for its stack. This application does not require an extra segment. When this program is placed in the memory system by DOS, it is loaded in the TPA at the first available area of memory above the drivers and other TPA program. This area is indicated by a **free-pointer** that is maintained by DOS. Program loading is handled automatically by the **program loader** located within DOS. Figure 2–5 shows how an application is stored in the memory system. The segments show an overlap because the amount of data in them does not require 64K bytes of memory. The side view of the segments clearly shows the overlap. It also shows how segments can be moved over any area of memory by changing the segment starting address. Fortunately, the DOS program loader calculates and assigns segment starting addresses.

Segment and Offset Addressing Scheme Allows Relocation

The segment and offset addressing scheme seems unduly complicated. It is complicated, but it also affords an advantage to the system. This complicated scheme of segment plus offset addressing

TABLE 2–3 Default 16-bit segment and offset combinations.

Segment	Offset	Special Purpose
CS	IP	Instruction address
SS	SP or BP	Stack address
DS	BX, DI, SI, an 8- or 16-bit number	Data address
ES	DI for string instructions	String destination address

TABLE 2–4 Default 32-bit segment and offset combinations.

Segment	Offset	Special Purpose
CS	EIP	Instruction address
SS	ESP or EBP	Stack address
DS	EAX, EBX, ECX, EDX, ESI, EDI, an 8- or 32-bit number	Data address
ES	EDI for string instructions	String destination address
FS	No default	General address
GS	No default	General address

allows DOS programs to be relocated in the memory system. It also allows programs written to function in the real mode to operate in a protected mode system. A **relocatable program** is one that can be placed into any area of memory and executed without change. **Relocatable data** are data that can be placed in any area of memory and used without any change to

FIGURE 2–4 A memory system showing the placement of four memory segments.
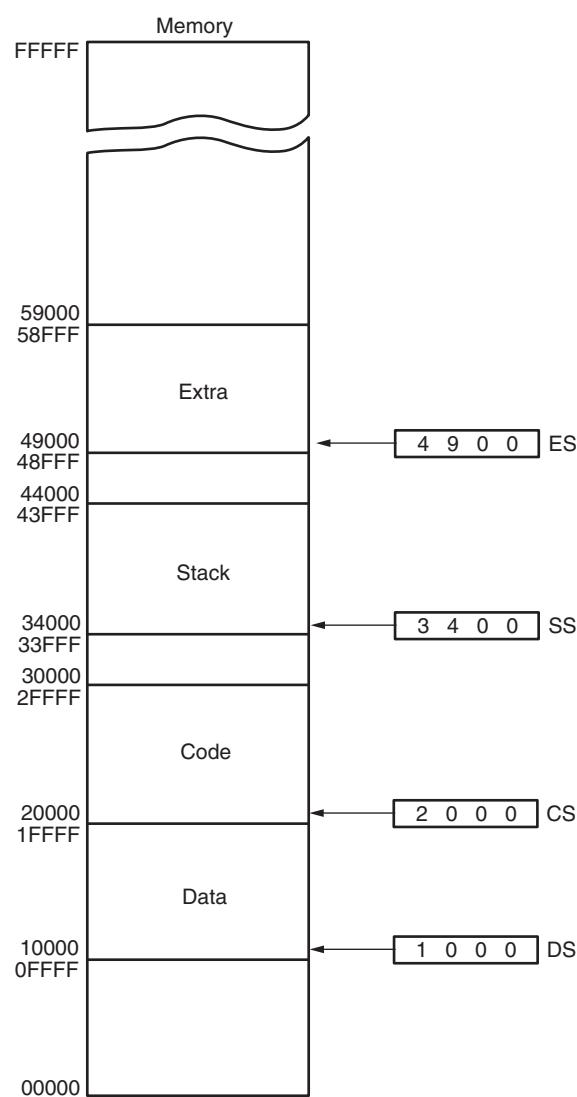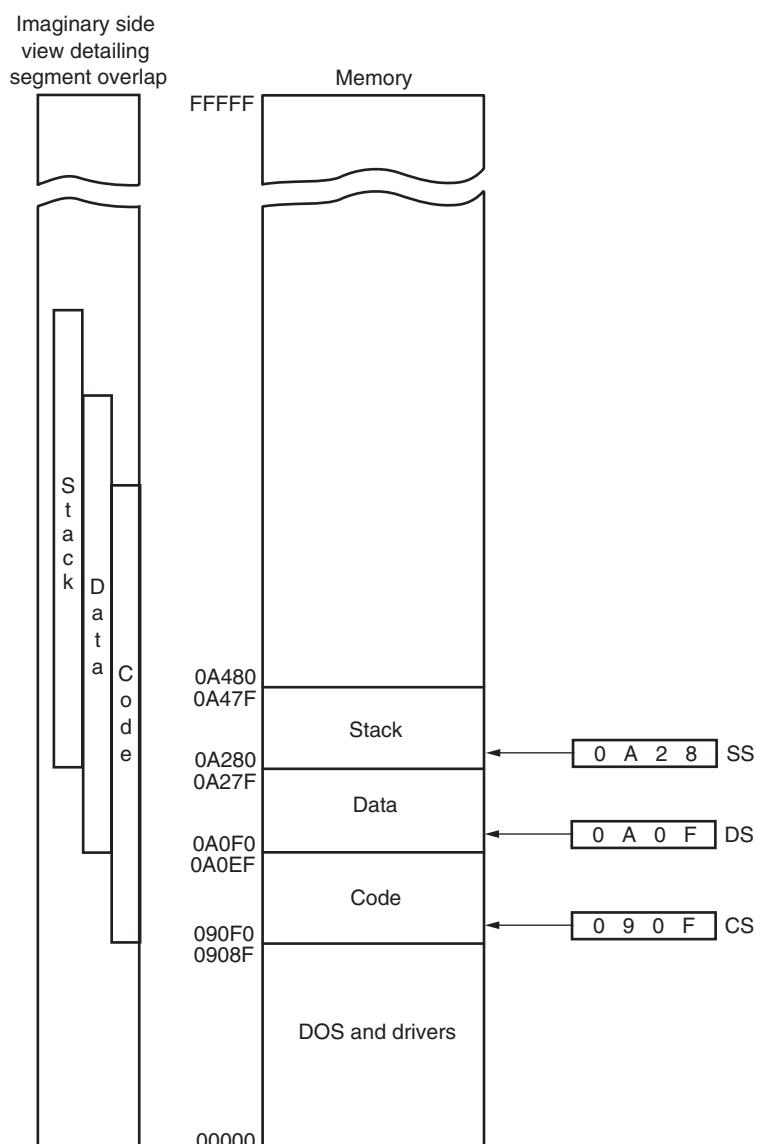
FIGURE 2–5 An application program containing a code, data, and stack segment loaded into a DOS system memory.

the program. The segment and offset addressing scheme allows both programs and data to be relocated without changing a thing in a program or data. This is ideal for use in a general-purpose computer system in which not all machines contain the same memory areas. The personal computer memory structure is different from machine to machine, requiring relocatable software and data.

Because memory is addressed within a segment by an offset address, the memory segment can be moved to any place in the memory system without changing any of the offset addresses. This is accomplished by moving the entire program, as a block, to a new area and then changing only the contents of the segment registers. If an instruction is 4 bytes above the start of the segment, its offset address is 4. If the entire program is moved to a new area of memory, this offset address of 4 still points to 4 bytes above the start of the segment. Only the contents of the segment register must be changed to address the program in the new area of memory. Without this feature, a program would have to be extensively rewritten or altered before it is moved. This would require additional time or many versions of a program for the many different configurations of computer systems. This concept also applies to programs written to

execute in the protected mode for Windows. In the Windows environment all programs are written assuming that the first 2G of memory are available for code and data. When the program is loaded, it is placed in the actual memory, which may be anywhere and a portion may be located on the disk in the form of a swap file.

<div style="display:flex">2–3</div>

INTRODUCTION TO PROTECTED MODE MEMORY ADDRESSING

Protected mode memory addressing (80286 and above) allows access to data and programs located above the first 1M byte of memory, as well as within the first 1M byte of memory. **Protected mode** is where Windows operates. Addressing this extended section of the memory system requires a change to the segment plus an offset addressing scheme used with real mode memory addressing. When data and programs are addressed in extended memory, the offset address is still used to access information located within the memory segment. One difference is that the segment address, as discussed with real mode memory addressing, is no longer present in the protected mode. In place of the segment address, the segment register contains a **selector** that selects a descriptor from a descriptor table. The **descriptor** describes the memory segment's location, length, and access rights. Because the segment register and offset address still access memory, protected mode instructions are identical to real mode instructions. In fact, most programs written to function in the real mode will function without change in the protected mode. The difference between modes is in the way that the segment register is interpreted by the microprocessor to access the memory segment. Another difference, in the 80386 and above, is that the offset address can be a 32-bit number instead of a 16-bit number in the protected mode. A 32-bit offset address allows the microprocessor to access data within a segment that can be up to 4G bytes in length. Programs that are written for the 32-bit protected mode execute in the 64-bit mode of the Pentium 4.

Selectors and Descriptors

The selector, located in the segment register, selects one of 8192 descriptors from one of two tables of descriptors. The **descriptor** describes the location, length, and access rights of the segment of memory. Indirectly, the segment register still selects a memory segment, but not directly as in the real mode. For example, in the real mode, if CS = 0008H, the code segment begins at location 00080H. In the protected mode, this segment number can address any memory location in the entire system for the code segment, as explained shortly.

There are two descriptor tables used with the segment registers: one contains global descriptors and the other contains local descriptors. The **global descriptors** contain segment definitions that apply to all programs, whereas the **local descriptors** are usually unique to an application. You might call a global descriptor a **system descriptor** and call a local descriptor an **application descriptor**. Each descriptor table contains 8192 descriptors, so a total of 16,384 total descriptors are available to an application at any time. Because the descriptor describes a memory segment, this allows up to 16,384 memory segments to be described for each application. Since a memory segment can be up to 4G bytes in length, this means that an application could have access to 4G × 16,384 bytes of memory or 64T bytes.

Figure 2–6 shows the format of a descriptor for the 80286 through the Core2. Note that each descriptor is 8 bytes in length, so the global and local descriptor tables are each a maximum of 64K bytes in length. Descriptors for the 80286 and the 80386–Core2 differ slightly, but the 80286 descriptor is upward-compatible.

The **base address** portion of the descriptor indicates the starting location of the memory segment. For the 80286 microprocessor, the base address is a 24-bit address, so segments begin at any location in its 16M bytes of memory. Note that the paragraph boundary limitation is

80286

31				0	
0000 0000 0000 0000		Access Rights	Base B23 B16		4
Base B15 B0		Limit L15 L0			0

Offset

80386–P4

31							0	
Base B31 B24	G	D	0	A V	Limit L19 L16	Access Rights	Base B23 B16	4
Base B15 B0					Limit L15 L0			0

Offset

64–bit P4

31							0	
0000 0000	G	D	L	A V	0000	Access Rights	0000 0000	4
0000 0000 0000 0000					0000 0000 0000 0000			0

Offset
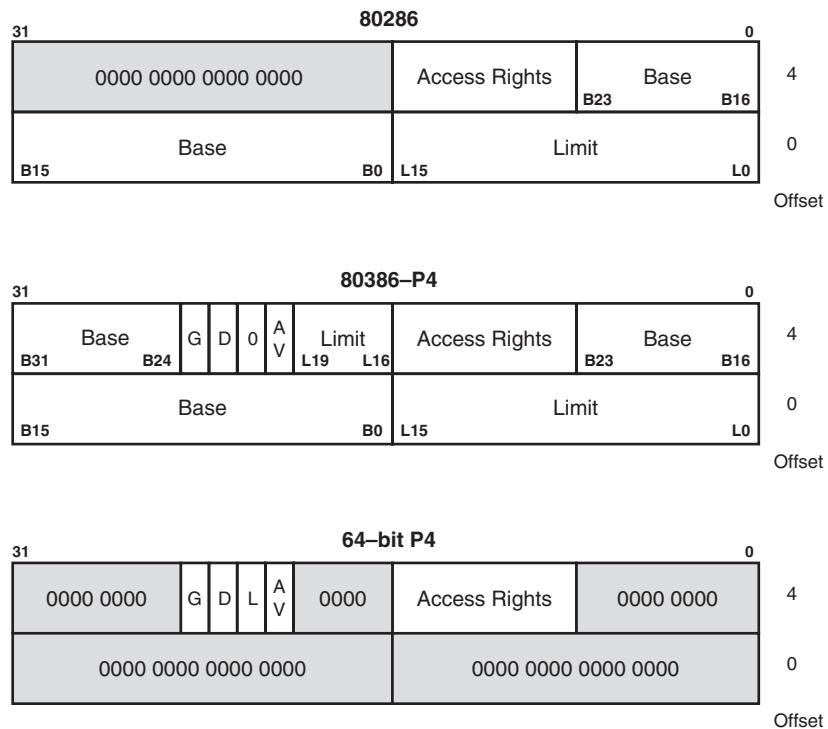
FIGURE 2–6   The 80286 through Core2 64-bit descriptors.

removed in these microprocessors when operated in the protected mode so segments may begin at any address. The 80386 and above use a 32-bit base address that allows segments to begin at any location in its 4G bytes of memory. Notice how the 80286 descriptor's base address is upward-compatible to the 80386 through the Pentium 4 descriptor because its most-significant 16 bits are 0000H. Refer to Chapters 18 and 19 for additional detail on the 64G memory space provided by the Pentium Pro through the Core2.

The segment limit contains the last offset address found in a segment. For example, if a segment begins at memory location F00000H and ends at location F000FFH, the base address is F00000H and the limit is FFH. For the 80286 microprocessor, the base address is F00000H and the limit is 00FFH. For the 80386 and above, the base address is 00F00000H and the limit is 000FFH. Notice that the 80286 has a 16-bit limit and the 80386 through the Pentium 4 have a 20-bit limit. An 80286 can access memory segments that are between 1 and 64K bytes in length. The 80386 and above access memory segments that are between 1 and 1M byte, or 4K and 4G bytes in length.

There is another feature found in the 80386 through the Pentium 4 descriptor that is not found in the 80286 descriptor: the G bit, or **granularity bit**. If G = 0, the limit specifies a segment limit of 00000H to FFFFFH. If G = 1, the value of the limit is multiplied by 4K bytes (appended with FFFH). The limit is then 00000FFFFH to FFFFFFFFH, if G = 1. This allows a segment length of 4K to 4G bytes in steps of 4K bytes. The reason that the segment length is 64K bytes in the 80286 is that the offset address is always 16 bits because of its 16-bit internal architecture. The 80386 and above use a 32-bit architecture that allows an offset address, in the protected mode operation, of the 32 bits. This 32-bit offset address allows segment lengths of 4G bytes and the 16-bit offset address allows segment lengths of 64K bytes. Operating systems operate in a 16- or 32-bit environment. For example, DOS uses a 16-bit environment, while most Windows applications use a 32-bit environment called **WIN32**.

In the 64-bit descriptor, the L bit (probably means large, but Intel calls it the 64-bit) selects 64-bit addresses in a Pentium 4 or Core2 with 64-bit extensions when L = 1 and 32-bit

compatibility mode when L = 0. In 64-bit protected operation, the code segment register is still used to select a section of code from the memory. Notice that the 64-bit descriptor has no limit or base address. It only contains an access rights byte and the control bits. In the 64-bit mode, there is no segment or limit in the descriptor and the base address of the segment, although not placed in the descriptor, is 00 0000 0000H. This means that all code segments start at address zero for 64-bit operation. There are no limit checks for a 64-bit code segment.

Example 2-1 shows the segment start and end if the base address is 10000000H, the limit is 001FFH, and the G bit = 0.

EXAMPLE 2–1

```
Base = Start = 10000000H
G = 0
End = Base + Limit = 10000000H + 001FFH = 100001FFH
```

Example 2-2 uses the same data as Example 2-1, except that the G bit = 1. Notice that the limit is appended with FFFH to determine the ending segment address.
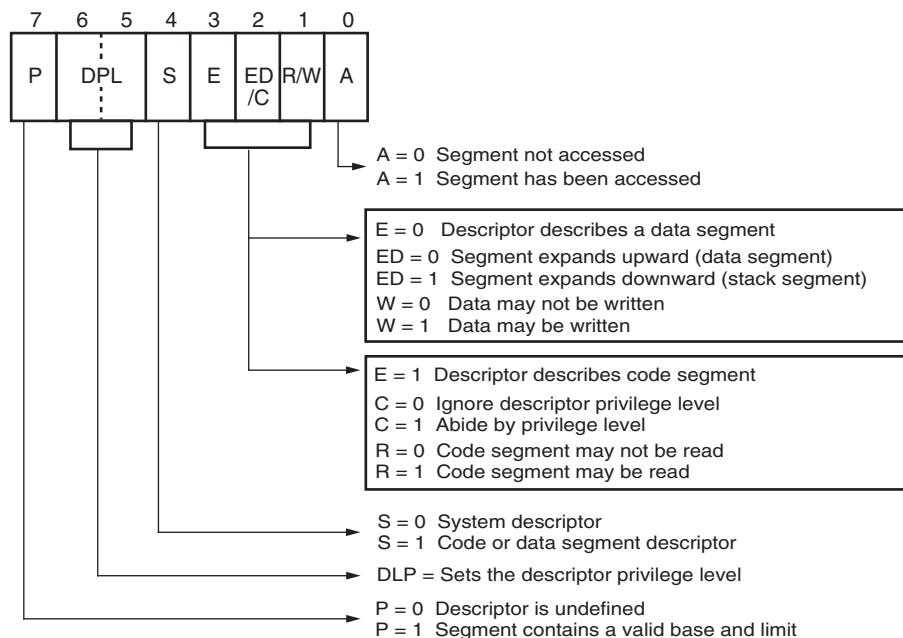
EXAMPLE 2–2

```
Base = Start = 10000000H
G = 1
End = Base + Limit = 10000000H + 001FFFFFH = 101FFFFFH
```

The AV bit, in the 80386 and above descriptor, is used by some operating systems to indicate that the segment is available (AV = 1) or not available (AV = 0). The D bit indicates how the 80386 through the Core2 instructions access register and memory data in the protected or real mode. If D = 0, the instructions are 16-bit instructions, compatible with the 8086–80286 microprocessors. This means that the instructions use 16-bit offset addresses and 16-bit register by default. This mode is often called the 16-bit instruction mode or DOS mode. If D = 1, the instructions are 32-bit instructions. By default, the 32-bit instruction mode assumes that all offset addresses and all registers are 32 bits. Note that the default for register size and offset address is overridden in both the 16- and 32-bit instruction modes. Both the MSDOS and PCDOS operating systems require that the instructions are always used in the 16-bit instruction mode. Windows 3.1, and any application that was written for it, also requires that the 16-bit instruction mode is selected. Note that the instruction mode is accessible only in a protected mode system such as Windows Vista. More detail on these modes and their application to the instruction set appears in Chapters 3 and 4.

The **access rights byte** (see Figure 2–7) controls access to the protected mode segment. This byte describes how the segment functions in the system. The access rights byte allows complete control over the segment. If the segment is a data segment, the direction of growth is specified. If the segment grows beyond its limit, the microprocessor's operating system program is interrupted, indicating a general protection fault. You can even specify whether a data segment can be written or is write-protected. The code segment is also controlled in a similar fashion and can have reading inhibited to protect software. Again, note that in 64-bit mode there is only a code segment and no other segment descriptor types. A 64-bit flat model program contains its data and stacks in the code segment.

Descriptors are chosen from the descriptor table by the segment register. Figure 2–8 shows how the segment register functions in the protected mode system. The segment register contains a 13-bit selector field, a table selector bit, and a requested privilege level field. The 13-bit **selector** chooses one of the 8192 descriptors from the descriptor table. The **TI bit** selects either the global descriptor table (TI = 0) or the local descriptor table (TI = 1). The **requested privilege level** (RPL) requests the access privilege level of a memory segment. The highest privilege level is 00 and the lowest is 11. If the requested privilege level matches or is higher in priority than the privilege level set by the access rights byte, access is granted. For example, if the

Note: Some of the letters used to describe the bits in the access rights bytes vary in Intel documentation.

FIGURE 2–7   The access rights byte for the 80286 through Core2 descriptor.

requested privilege level is 10 and the access rights byte sets the segment privilege level at 11, access is granted because 10 is higher in priority than privilege level 11. Privilege levels are used in multiuser environments. Windows uses privilege level 00 (**ring 0**) for the kernel and driver programs and level 11 (**ring 3**) for applications. Windows does not use levels 01 or 10. If privilege levels are violated, the system normally indicates an application or privilege level violation.

Figure 2–9 shows how the segment register, containing a selector, chooses a descriptor from the global descriptor table. The entry in the global descriptor table selects a segment in the memory system. In this illustration, DS contains 0008H, which accesses the descriptor number 1 from the global descriptor table using a requested privilege level of 00. Descriptor number 1 contains a descriptor that defines the base address as 00100000H with a segment limit of 000FFH. This means that a value of 0008H loaded into DS causes the microprocessor to use memory locations 00100000H–001000FFH for the data segment with this example descriptor table. Note that descriptor zero is called the null descriptor, must contain all zeros, and may not be used for accessing memory.
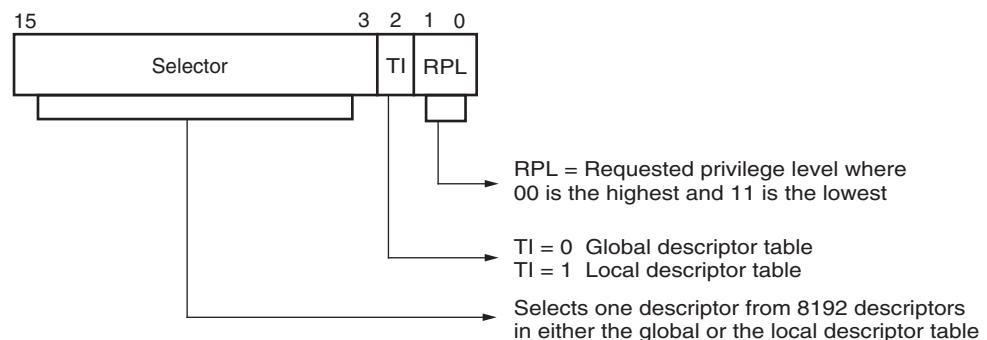
FIGURE 2–8   The contents of a segment register during protected mode operation of the 80286 through Core2 microprocessors.
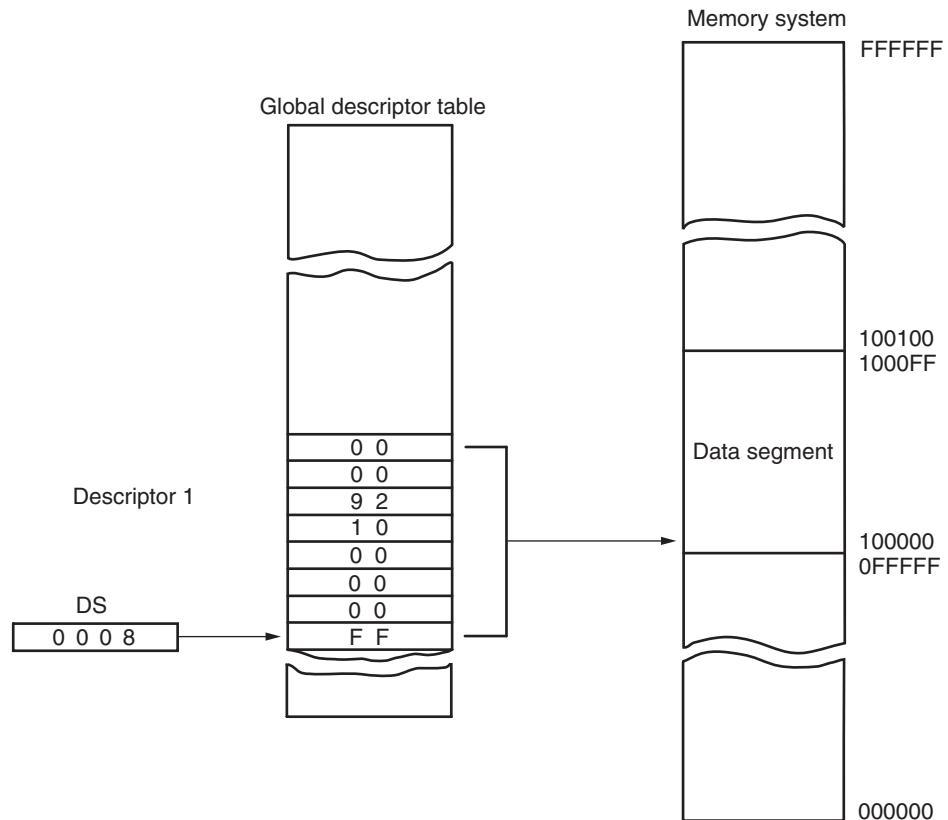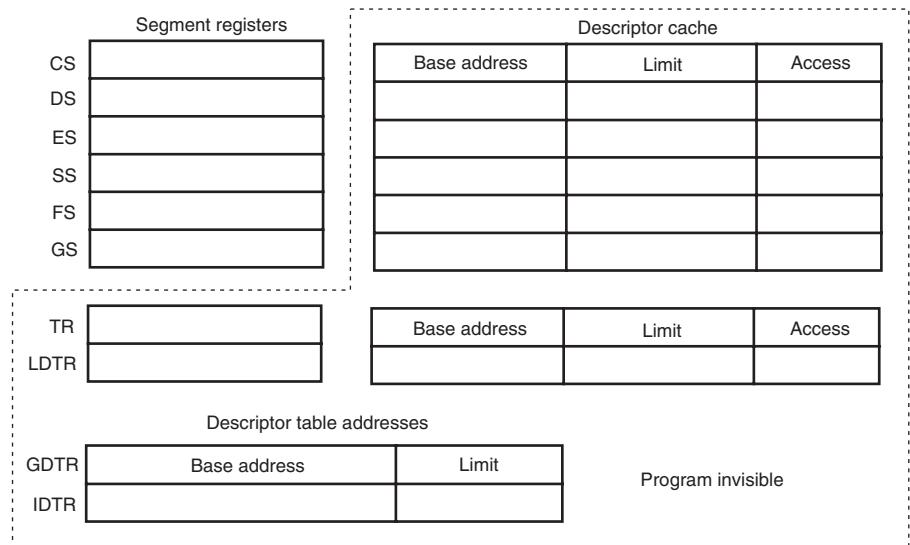
FIGURE 2–9   Using the DS register to select a description from the global descriptor table. In this example, the DS register accesses memory locations 00100000H–001000FFH as a data segment.

Program-Invisible Registers

The global and local descriptor tables are found in the memory system. In order to access and specify the address of these tables, the 80286–Core2 contain program-invisible registers. The program-invisible registers are not directly addressed by software so they are given this name (although some of these registers are accessed by the system software). Figure 2–10 illustrates the program-invisible registers as they appear in the 80286 through the Core2. These registers control the microprocessor when operated in protected mode.

Each of the segment registers contains a program-invisible portion used in the protected mode. The program-invisible portion of these registers is often called cache memory because cache is any memory that stores information. This cache is not to be confused with the level 1 or level 2 caches found with the microprocessor. The program-invisible portion of the segment register is loaded with the base address, limit, and access rights each time the number segment register is changed. When a new segment number is placed in a segment register, the microprocessor accesses a descriptor table and loads the descriptor into the program-invisible portion of the segment register. It is held there and used to access the memory segment until the segment number is again changed. This allows the microprocessor to repeatedly access a memory segment without referring to the descriptor table (hence the term *cache*).

The GDTR (**global descriptor table register**) and IDTR (**interrupt descriptor table register**) contain the base address of the descriptor table and its limit. The limit of each descriptor

Notes:
1. The 80286 does not contain FS and GS nor the program-invisible portions of these registers.
2. The 80286 contains a base address that is 24-bits and a limit that is 16-bits.
3. The 80386/80486/Pentium/Pentium Pro contain a base address that is 32-bits and a limit that is 20-bits.
4. The access rights are 8-bits in the 80286 and 12-bits in the 80386/80486/Pentium–Core2.

FIGURE 2–10   The program-invisible register within the 80286–Core2 microprocessors.

table is 16 bits because the maximum table length is 64K bytes. When the protected mode operation is desired, the address of the global descriptor table and its limit are loaded into the GDTR.

Before using the protected mode, the interrupt descriptor table and the IDTR must also be initialized. More detail is provided on protected mode operation later in the text. At this point, programming and additional description of these registers are impossible.

The location of the local descriptor table is selected from the global descriptor table. One of the global descriptors is set up to address the local descriptor table. To access the local descriptor table, the LDTR (**local descriptor table register**) is loaded with a selector, just as a segment register is loaded with a selector. This selector accesses the global descriptor table and loads the address, limit, and access rights of the local descriptor table into the cache portion of the LDTR.

The TR (task register) holds a selector, which accesses a descriptor that defines a task. A task is most often a procedure or application program. The descriptor for the procedure or application program is stored in the global descriptor table, so access can be controlled through the privilege levels. The task register allows a context or task switch in about 17 μs. Task switching allows the microprocessor to switch between tasks in a fairly short amount of time. The task switch allows multitasking systems to switch from one task to another in a simple and orderly fashion.

2–4   MEMORY PAGING

The **memory paging mechanism** located within the 80386 and above allows any physical memory location to be assigned to any linear address. The **linear address** is defined as the address generated by a program. The **physical address** is the actual memory location accessed by a program. With the memory paging unit, the linear address is invisibly translated to any physical address, which allows an application written to function at a specific address to be relocated through the paging mechanism. It also allows memory to be placed into areas where no memory exists. An example is the upper memory blocks provided by EMM386.EXE in a DOS system.

The EMM386.EXE program reassigns extended memory, in 4K blocks, to the system memory between the video BIOS and the system BIOS ROMS for upper memory blocks. Without the paging mechanism, the use of this area of memory is impossible.

In Windows, each application is allowed a 2G linear address space from location 00000000H–7FFFFFFFH even though there may not be enough memory or memory available at these addresses. Through paging to the hard disk drive and paging to the memory through the memory paging unit, any Windows application can be executed.

Paging Registers

The paging unit is controlled by the contents of the microprocessor's control registers. See Figure 2–11 for the contents of control registers CR0 through CR4. Note that these registers are available to the 80386 through the Core2 microprocessors. Beginning with the Pentium, an additional control register labeled CR4 controls extensions to the basic architecture provided in the Pentium or newer microprocessor. One of these features is a 2M- or a 4M-byte page that is enabled by controlling CR4.

The registers important to the paging unit are CR0 and CR3. The leftmost bit (PG) position of CR0 selects paging when placed at a logic 1 level. If the PG bit is cleared (0), the linear address generated by the program becomes the physical address used to access memory. If the PG bit is set (1), the linear address is converted to a physical address through the paging mechanism. The paging mechanism functions in both the real and protected modes.

CR3 contains the page directory base or root address, and the PCD and PWT bits. The PCD and PWT bits control the operation of the PCD and PWT pins on the microprocessor. If PCD is set (1), the PCD pin becomes a logic one during bus cycles that are not paged. This allows the external hardware to control the level 2 cache memory. (Note that the level 2 cache memory is an internal [on modern versions of the Pentium] high-speed memory that functions as a buffer between the microprocessor and the main DRAM memory system.) The PWT bit also appears on the PWT pin during bus cycles that are not paged to control the write-through cache in the system. The page directory base address locates the directory for the page translation unit. Note that this address locates the page directory at any 4K boundary in the memory system because it is appended internally with 000H. The page directory contains 1024 directory entries of 4 bytes each. Each page directory entry addresses a page table that contains 1024 entries.
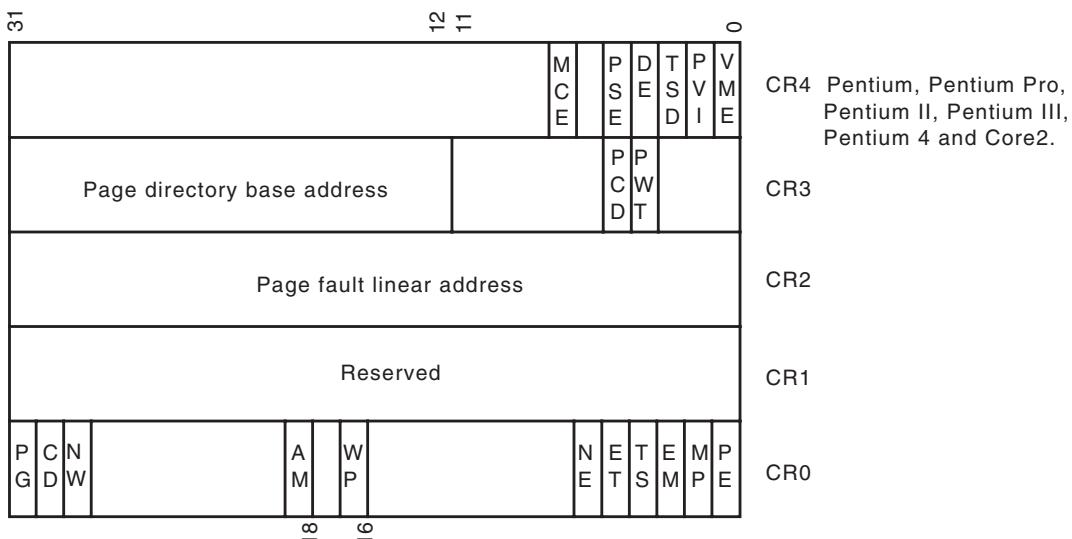
FIGURE 2–11   The control register structure of the microprocessor.

(a)

(b)
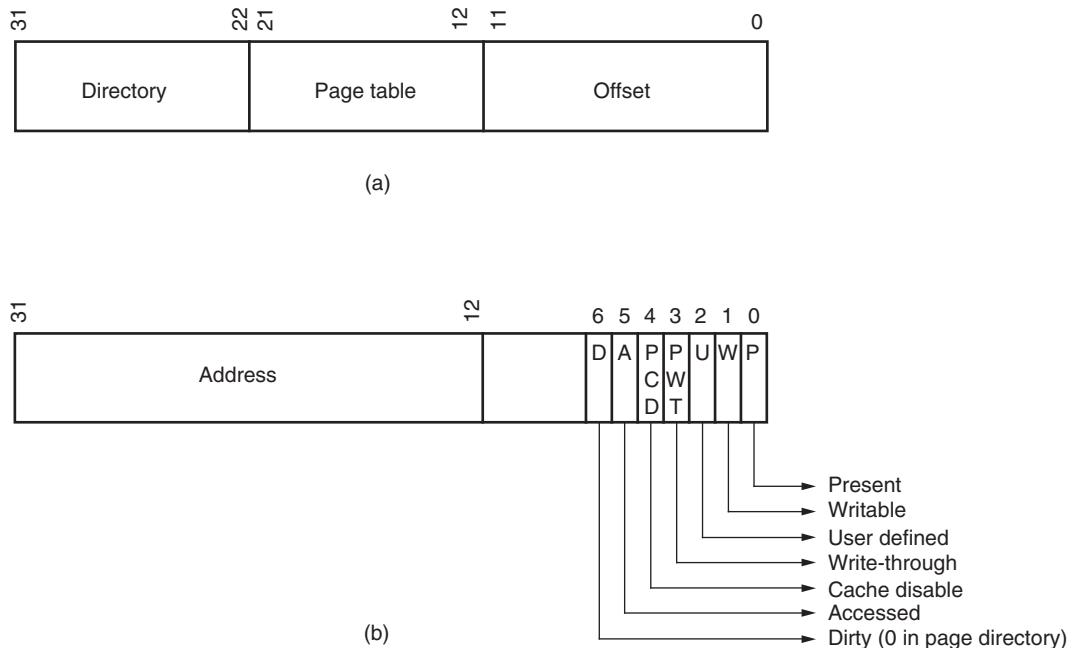
FIGURE 2–12   The format for the linear address (a) and a page directory or page table entry (b).

The linear address, as it is generated by the software, is broken into three sections that are used to access the **page directory entry, page table entry**, and **memory page offset address**. Figure 2–12 shows the linear address and its makeup for paging. Notice how the leftmost 10 bits address an entry in the page directory. For linear address 00000000H–003FFFFFH, the first page directory is accessed. Each page directory entry represents or repages a 4M section of the memory system. The contents of the page directory select a page table that is indexed by the next 10 bits of the linear address (bit positions 12–21). This means that address 00000000H–00000FFFH selects page directory entry of 0 and page table entry of 0. Notice this is a 4K-byte address range. The offset part of the linear address (bit positions 0–11) next selects a byte in the 4K-byte memory page. In Figure 2–12, if the page table entry 0 contains address 00100000H, then the physical address is 00100000H-00100FFFH for linear address 00000000H–00000FFFH. This means that when the program accesses a location between 00000000H and 00000FFFH, the microprocessor physically addresses location 00100000H–00100FFFH.

Because the act of repaging a 4K-byte section of memory requires access to the page directory and a page table, which are both located in memory, Intel has incorporated a special type of cache called the TLB (**translation look-aside buffer**). In the 80486 microprocessor, the cache holds the 32 most recent page translation addresses. This means that the last 32 page table translations are stored in the TLB, so if the same area of memory is accessed, the address is already present in the TLB, and access to the page directory and page tables is not required. This speeds program execution. If a translation is not in the TLB, the page directory and page table must be accessed, which requires additional execution time. The Pentium–Pentium 4 microprocessors contain separate TLBs for each of their instruction and data caches.

The Page Directory and Page Table

Figure 2–13 shows the page directory, a few page tables, and some memory pages. There is only one page directory in the system. The page directory contains 1024 doubleword addresses that locate up to 1024 page tables. The page directory and each page table are 4K bytes in length. If
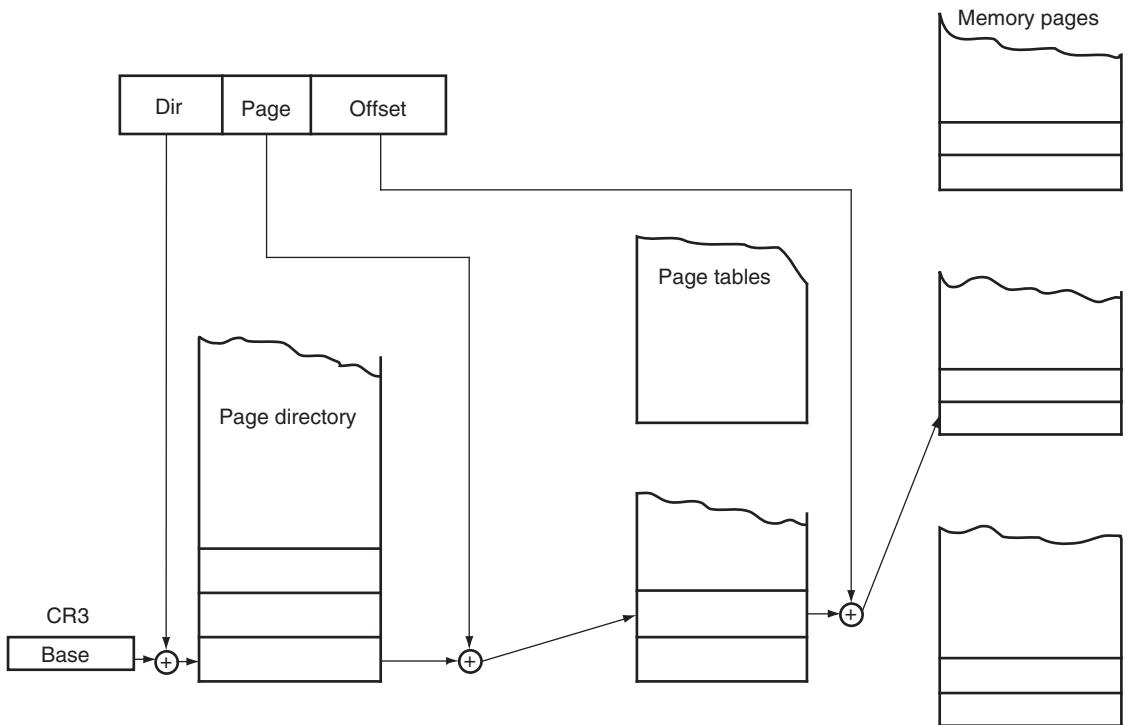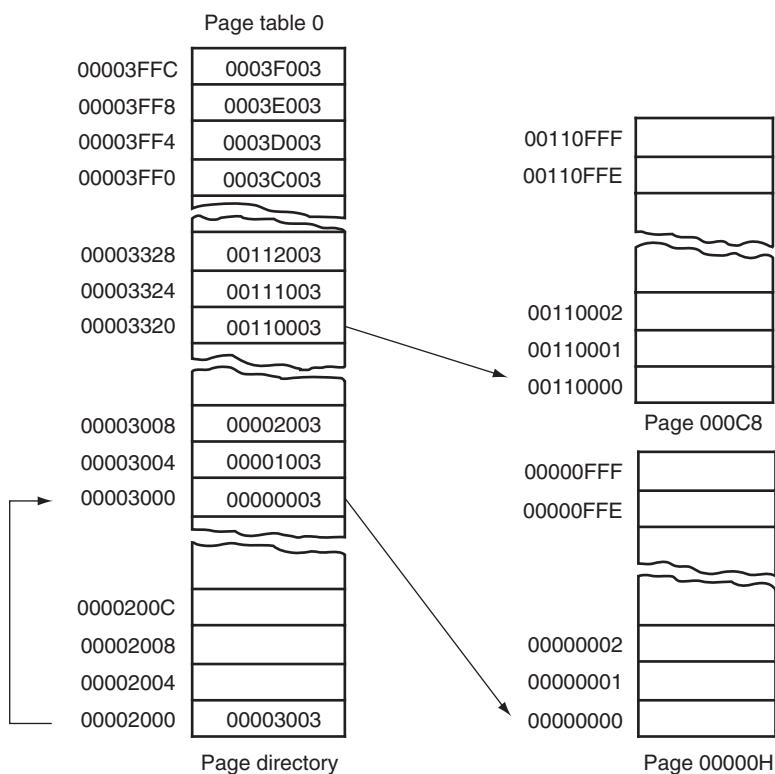
FIGURE 2–13 The paging mechanism in the 80386 through Core2 microprocessors.

FIGURE 2–14 The page directory, page table 0, and two memory pages. Note how the address of page 000C8000–000C9000 has been moved to 00110000–00110FFF.

the entire 4G byte of memory is paged, the system must allocate 4K bytes of memory for the page directory, and 4K times 1024 or 4M bytes for the 1024 page tables. This represents a considerable investment in memory resources.

The DOS system and EMM386.EXE use page tables to redefine the area of memory between locations C8000H–EFFFFH as upper memory blocks. This is done by repaging extended memory to backfill this part of the conventional memory system to allow DOS access to additional memory. Suppose that the EMM386.EXE program allows access to 16M bytes of extended and conventional memory through paging and locations C8000H–EFFFFH must be repaged to locations 110000–138000H, with all other areas of memory paged to their normal locations. Such a scheme is depicted in Figure 2–14.

Here, the page directory contains four entries. Recall that each entry in the page directory corresponds to 4M bytes of physical memory. The system also contains four page tables with 1024 entries each. Recall that each entry in the page table repages 4K bytes of physical memory. This scheme requires a total of 16K of memory for the four page tables and 16 bytes of memory for the page directory.

As with DOS, the Windows program also repages the memory system. At present, Windows version 3.11 supports paging for only 16M bytes of memory because of the amount of memory required to store the page tables. Newer versions of Windows repage the entire memory system. On the Pentium–Core2 microprocessors, pages can be 4K, 2M, or 4M bytes in length. In the 2M and 4M variations, there is only a page directory and a memory page, but no page table.

2–5   FLAT MODE MEMORY
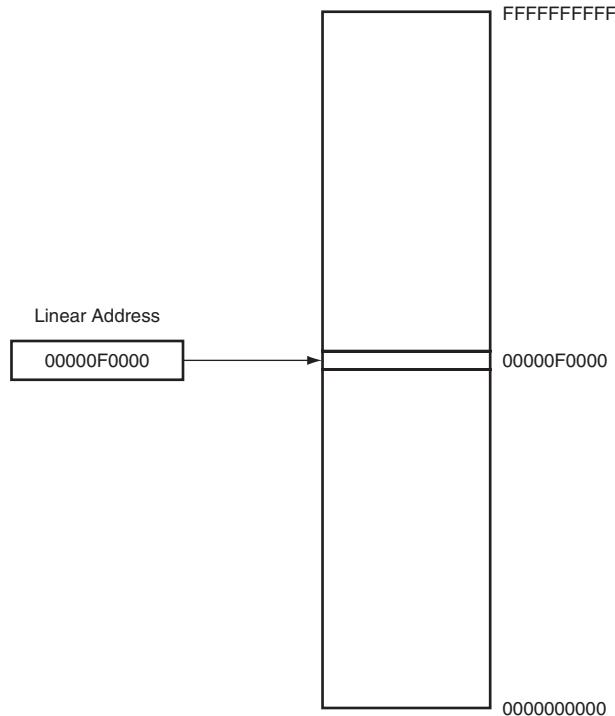
The memory system in a Pentium-based computer (Pentium 4 or Core2) that uses the 64-bit extensions uses a flat mode memory system. A *flat mode memory* system is one in which there is no segmentation. The address of the first byte in the memory is at 00 0000 0000H and the last location is at FF FFFF FFFFH (address is 40-bits). The flat model does not use a segment register to address a location in the memory. The CS segment register is used to select a descriptor from the descriptor table that defines the access rights of only a code segment. The segment register still selects the privilege level of the software. The flat model does not select the memory address of a segment using the base and limit in the descriptor (see Figure 2–6). In 64-bit mode the actual address is not modified by the descriptor as in 32-bit protected mode. The offset address is the actual physical address in 64-bit mode. Refer to Figure 2–15 for the flat mode memory model.

This form of addressing is much easier to understand, but offers little protection to the system, through the hardware, as did the protected mode system discussed in Section 2.3. The real mode system is *not* available if the processor operates in the 64-bit mode. Protection and paging are allowed in the 64-bit mode. The CS register is still used in the protected mode operation in the 64-bit mode.

In the 64-bit mode if set to IA32 compatibility (when the L bit −0 is in the descriptor), an address is 64-bits, but since only 40 bits of the address are brought out to the address pins, any address above 40 bits is truncated. Instructions that use a displacement address can only use a 32-bit displacement, which allows a range of ±2G from the current instruction. This addressing mode is called RIP relative addressing, and is explained in Chapter 3. The move immediate instruction allows a full 64-bit address and access to any flat mode memory location. Other instructions do not allow access to a location above 4G because the offset address is still 32-bits.

If the Pentium is operated in the full 64-bit mode (where the L = 1 in the descriptor), the address may be 64-bits or 32-bits. This is shown in examples in the next chapter with addressing modes and in more detail in Chapter 4. Most programs today are operated in the IA32 compatible mode so current versions of Windows software operates properly, but this will change in a

FIGURE 2–15   The 64-bit
flat mode memory model.

FFFFFFFFFF

Linear Address

| 00000F0000 |

00000F0000

0000000000

few years as memory becomes larger and most people have 64-bit computers. This is another
example of how the industry makes the software obsolete as the hardware changes.

2–6   SUMMARY

1. The programming model of the 8086 through 80286 contains 8- and 16-bit registers. The
   programming model of the 80386 and above contains 8-, 16-, and 32-bit extended registers
   as well as two additional 16-bit segment registers: FS and GS.
2. The 8-bit registers are AH, AL, BH, BL, CH, CL, DH, and DL. The 16-bit registers are AX,
   BX, CX, DX, SP, BP, DI, and SI. The segment registers are CS, DS, ES, SS, FS, and GS. The
   32-bit extended registers are EAX, EBX, ECX, EDX, ESP, EBP, EDI, and ESI. The 64-bit
   registers in a Pentium 4 with 64-bit extensions are RAX, RBX, RCX, RDX, RSP, RBP, RDI,
   RSI, and R8 through R15. In addition, the microprocessor contains an instruction pointer
   (IP/EIP/RIP) and flag register (FLAGS, EFLAGS, or RFLAGS).
3. All real mode memory addresses are a combination of a segment address plus an offset
   address. The starting location of a segment is defined by the 16-bit number in the segment
   register that is appended with a hexadecimal zero at its rightmost end. The offset address is
   a 16-bit number added to the 20-bit segment address to form the real mode memory address.
4. All instructions (code) are accessed by the combination of CS (segment address) plus IP or
   EIP (offset address).
5. Data are normally referenced through a combination of the DS (data segment) and either an
   offset address or the contents of a register that contains the offset address. The 8086–Core2
   use BX, DI, and SI as default offset registers for data if 16-bit registers are selected. The
   80386 and above can use the 32-bit registers EAX, EBX, ECX, EDX, EDI, and ESI as
   default offset registers for data.

6. Protected mode operation allows memory above the first 1M byte to be accessed by the 80286 through the Core2 microprocessors. This extended memory system (XMS) is accessed via a segment address plus an offset address, just as in the real mode. The difference is that the segment address is not held in the segment register. In the protected mode, the segment starting address is stored in a descriptor that is selected by the segment register.

7. A protected mode descriptor contains a base address, limit, and access rights byte. The base address locates the starting address of the memory segment; the limit defines the last location of the segment. The access rights byte defines how the memory segment is accessed via a program. The 80286 microprocessor allows a memory segment to start at any of its 16M bytes of memory using a 24-bit base address. The 80386 and above allow a memory segment to begin at any of its 4G bytes of memory using a 32-bit base address. The limit is a 16-bit number in the 80286 and a 20-bit number in the 80386 and above. This allows an 80286 memory segment limit of 64K bytes, and an 80386 and above memory segment limit of either 1M bytes ($G = 0$) or 4G bytes ($G = 1$). The L bit selects 64-bit address operation in the code descriptor.

8. The segment register contains three fields of information in the protected mode. The leftmost 13 bits of the segment register address one of 8192 descriptors from a descriptor table. The TI bit accesses either the global descriptor table ($TI = 0$) or the local descriptor table ($TI = 1$). The rightmost 2 bits of the segment register select the requested priority level for the memory segment access.

9. The program-invisible registers are used by the 80286 and above to access the descriptor tables. Each segment register contains a cache portion that is used in protected mode to hold the base address, limit, and access rights acquired from a descriptor. The cache allows the microprocessor to access the memory segment without again referring to the descriptor table until the segment register's contents are changed.

10. A memory page is 4K bytes in length. The linear address, as generated by a program, can be mapped to any physical address through the paging mechanism found within the 80386 through the Pentium 4 microprocessor.

11. Memory paging is accomplished through control registers CR0 and CR3. The PG bit of CR0 enables paging, and the contents of CR3 addresses the page directory. The page directory contains up to 1024 page table addresses that are used to access paging tables. The page table contains 1024 entries that locate the physical address of a 4K-byte memory page.

12. The TLB (translation look-aside buffer) caches the 32 most recent page table translations. This precludes page table translation if the translation resides in the TLB, speeding the execution of the software.

13. The flat mode memory contains 1T byte of memory using a 40-bit address. In the future, Intel plans to increase the address width to 52 bits to access 4P bytes of memory. The flat mode is only available in the Pentium 4 and Core2 that have their 64-bit extensions enabled.

2–7   QUESTIONS AND PROBLEMS

1. What are program-visible registers?
2. The 80286 addresses registers that are 8 and _____ bits wide.
3. The extended registers are addressable by which microprocessors?
4. The extended BX register is addressed as _____.
5. Which register holds a count for some instructions?
6. What is the purpose of the IP/EIP register?
7. The carry flag bit is not modified by which arithmetic operations?
8. Will an overflow occur if a signed FFH is added to a signed 01H?
9. A number that contains 3 one bits is said to have _____ parity.

10. Which flag bit controls the INTR pin on the microprocessor?
11. Which microprocessors contain an FS segment register?
12. What is the purpose of a segment register in the real mode operation of the microprocessor?
13. In the real mode, show the starting and ending addresses of each segment located by the following segment register values:
    (a) 1000H
    (b) 1234H
    (c) 2300H
    (d) E000H
    (e) AB00H
14. Find the memory address of the next instruction executed by the microprocessor, when operated in the real mode, for the following CS:IP combinations:
    (a) CS = 1000H and IP = 2000H
    (b) CS = 2000H and IP = 1000H
    (c) CS = 2300H and IP = 1A00H
    (d) CS = 1A00H and IP = B000H
    (e) CS = 3456H and IP = ABCDH
15. Real mode memory addresses allow access to memory below which memory address?
16. Which register or registers are used as an offset address for the string instruction destination in the microprocessor?
17. Which 32-bit register or registers are used to hold an offset address for data segment data in the Pentium 4 microprocessor?
18. The stack memory is addressed by a combination of the _____ segment plus _____ offset.
19. If the base pointer (BP) addresses memory, the _____ segment contains the data.
20. Determine the memory location addressed by the following real mode 80286 register combinations:
    (a) DS = 1000H and DI = 2000H
    (b) DS = 2000H and SI = 1002H
    (c) SS = 2300H and BP = 3200H
    (d) DS = A000H and BX = 1000H
    (e) SS = 2900H and SP = 3A00H
21. Determine the memory location addressed by the following real mode Core2 register combinations:
    (a) DS = 2000H and EAX = 00003000H
    (b) DS = 1A00H and ECX = 00002000H
    (c) DS = C000H and ESI = 0000A000H
    (d) SS = 8000H and ESP = 00009000H
    (e) DS = 1239H and EDX = 0000A900H
22. Protected mode memory addressing allows access to which area of the memory in the 80286 microprocessor?
23. Protected mode memory addressing allows access to which area of the memory in the Pentium 4 microprocessor?
24. What is the purpose of the segment register in protected mode memory addressing?
25. How many descriptors are accessible in the global descriptor table in the protected mode?
26. For an 80286 descriptor that contains a base address of A00000H and a limit of 1000H, what starting and ending locations are addressed by this descriptor?
27. For a Core2 descriptor that contains a base address of 01000000H, a limit of 0FFFFH, and G = 0, what starting and ending locations are addressed by this descriptor?
28. For a Core2 descriptor that contains a base address of 00280000H, a limit of 00010H, and G = 1, what starting and ending locations are addressed by this descriptor?

29. If the DS register contains 0020H in a protected mode system, which global descriptor table entry is accessed?
30. If DS = 0103H in a protected mode system, the requested privilege level is _____.
31. If DS = 0105H in a protected mode system, which entry, table, and requested privilege level are selected?
32. What is the maximum length of the global descriptor table in the Pentium 4 microprocessor?
33. Code a descriptor that describes a memory segment that begins at location 210000H and ends at location 21001FH. This memory segment is a code segment that can be read. The descriptor is for an 80286 microprocessor.
34. Code a descriptor that describes a memory segment that begins at location 03000000H and ends at location 05FFFFFFH. This memory segment is a data segment that grows upward in the memory system and can be written. The descriptor is for a Pentium 4 microprocessor.
35. Which register locates the global descriptor table?
36. How is the local descriptor table addressed in the memory system?
37. Describe what happens when a new number is loaded into a segment register when the microprocessor is operated in the protected mode.
38. What are the program-invisible registers?
39. What is the purpose of the GDTR?
40. How many bytes are found in a memory page?
41. What register is used to enable the paging mechanism in the 80386, 80486, Pentium, Pentium Pro, Pentium 4, and Core2 microprocessors?
42. How many 32-bit addresses are stored in the page directory?
43. Each entry in the page directory translates how much linear memory into physical memory?
44. If the microprocessor sends linear address 00200000H to the paging mechanism, which paging directory entry is accessed, and which page table entry is accessed?
45. What value is placed in the page table to redirect linear address 20000000H to physical address 30000000H?
46. What is the purpose of the TLB located within the Pentium class microprocessor?
47. Using the Internet, write a short report that details the TLB. *Hint:* You might want to go to the Intel Web site and search for information.
48. Locate articles about paging on the Internet and write a report detailing how paging is used in a variety of systems.
49. What is the flat mode memory system?
50. A flat mode memory system in the current version of the 64-bit Pentium 4 and Core2 allow these microprocessors to access _____ bytes of memory.

CHAPTER 3

Addressing Modes

INTRODUCTION

Efficient software development for the microprocessor requires a complete familiarity with the addressing modes employed by each instruction. In this chapter, the MOV (move data) instruction is used to describe the data-addressing modes. The MOV instruction transfers bytes or words of data between two registers or between registers and memory in the 8086 through the 80286. Bytes, words, or doublewords are transferred in the 80386 and above by a MOV. In describing the program memory-addressing modes, the CALL and JUMP instructions show how to modify the flow of the program.

The data-addressing modes include register, immediate, direct, register indirect, base-plus index, register-relative, and base relative-plus-index in the 8086 through the 80286 microprocessor. The 80386 and above also include a scaled-index mode of addressing memory data. The program memory-addressing modes include program relative, direct, and indirect. This chapter explains the operation of the stack memory so that the PUSH and POP instructions and other stack operations will be understood.

CHAPTER OBJECTIVES

Upon completion of this chapter, you will be able to:

1. Explain the operation of each data-addressing mode.
2. Use the data-addressing modes to form assembly language statements.
3. Explain the operation of each program memory-addressing mode.
4. Use the program memory-addressing modes to form assembly and machine language statements.
5. Select the appropriate addressing mode to accomplish a given task.
6. Detail the difference between addressing memory data using real mode and protected mode operation.
7. Describe the sequence of events that place data onto the stack or remove data from the stack.
8. Explain how a data structure is placed in memory and used with software.

3–1   DATA-ADDRESSING MODES
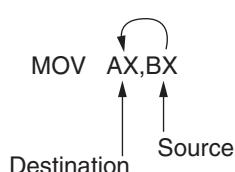
Because the MOV instruction is a very common and flexible instruction, it provides a basis for the explanation of the data-addressing modes. Figure 3–1 illustrates the MOV instruction and defines the direction of data flow. The **source** is to the right and the **destination** is to the left, next to the opcode MOV. (An **opcode**, or operation code, tells the microprocessor which operation to perform.) This direction of flow, which is applied to all instructions, is awkward at first. We naturally assume that things move from left to right, whereas here they move from right to left. Notice that a comma always separates the destination from the source in an instruction. Also, note that memory-to-memory transfers are *not* allowed by any instruction except for the MOVS instruction.

In Figure 3–1, the MOV AX, BX instruction transfers the word contents of the source register (BX) into the destination register (AX). The source never changes, but the destination always changes.[1] It is crucial to remember that a MOV instruction always *copies* the source data into the destination. The MOV never actually picks up the data and moves it. Also, note the flag register remains unaffected by most data transfer instructions. The source and destination are often called **operands**.
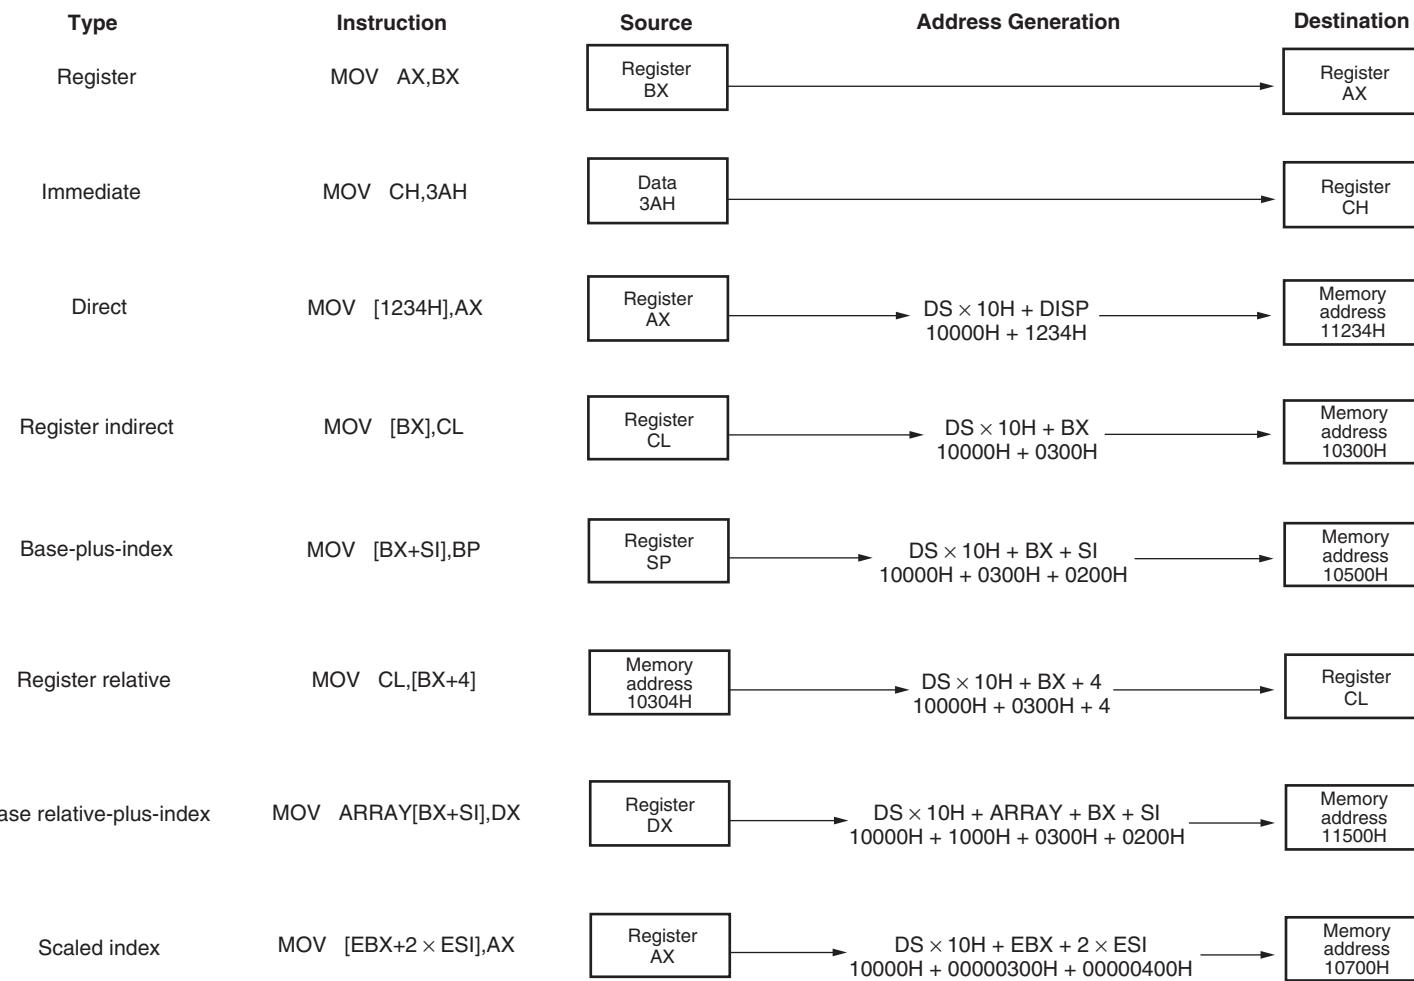
Figure 3–2 shows all possible variations of the data-addressing modes using the MOV instruction. This illustration helps to show how each data-addressing mode is formulated with the MOV instruction and also serves as a reference on data-addressing modes. Note that these are the same data-addressing modes found with all versions of the Intel microprocessor, except for the scaled-index-addressing mode, which is found only in the 80386 through the Core2. The RIP relative addressing mode is not illustrated and is only available on the Pentium 4 and the Core2 when operated in the 64-bit mode. The data-addressing modes are as follows:

Register addressing   Register addressing transfers a copy of a byte or word from the source register or contents of a memory location to the destination register or memory location. (Example: The MOV CX, DX instruction copies the word-sized contents of register DX into register CX.) In the 80386 and above, a doubleword can be transferred from the source register or memory location to the destination register or memory location. (Example: The MOV ECX, EDX instruction copies the doubleword-sized contents of register EDX into register ECX.) In the Pentium 4 operated in the 64-bit mode, any 64-bit register is also allowed. An example is the MOV RDX, RCX instruction that transfers a copy of the quadword contents of register RCX into register RDX.

Immediate addressing   Immediate addressing transfers the source, an immediate byte, word, doubleword, or quadword of data, into the destination register or memory location. (Example: The MOV AL, 22H instruction copies a byte-sized 22H into register AL.) In the 80386 and above, a doubleword of immediate data can be transferred into a register or

FIGURE 3–1   The MOV instruction showing the source, destination, and direction of data flow.

---

[1]The exceptions are the CMP and TEST instructions, which never change the destination. These instructions are described in later chapters.

Type	Instruction	Source	Address Generation	Destination
Register	MOV AX,BX	Register BX		Register AX
Immediate	MOV CH,3AH	Data 3AH		Register CH
Direct	MOV [1234H],AX	Register AX	DS × 10H + DISP 10000H + 1234H	Memory address 11234H
Register indirect	MOV [BX],CL	Register CL	DS × 10H + BX 10000H + 0300H	Memory address 10300H
Base-plus-index	MOV [BX+SI],BP	Register SP	DS × 10H + BX + SI 10000H + 0300H + 0200H	Memory address 10500H
Register relative	MOV CL,[BX+4]	Memory address 10304H	DS × 10H + BX + 4 10000H + 0300H + 4	Register CL
Base relative-plus-index	MOV ARRAY[BX+SI],DX	Register DX	DS × 10H + ARRAY + BX + SI 10000H + 1000H + 0300H + 0200H	Memory address 11500H
Scaled index	MOV [EBX+2 × ESI],AX	Register AX	DS × 10H + EBX + 2 × ESI 10000H + 00000300H + 00000400H	Memory address 10700H

Notes: EBX = 00000300H, ESI = 00000200H, ARRAY = 1000H, and DS = 1000H

FIGURE 3–2   8086–Core2 data-addressing modes.

memory location. (Example: The MOV EBX, 12345678H instruction copies a doubleword-sized l2345678H into the 32-bit-wide EBX register.) In 64-bit operation of the Pentium 4 or Core2, only a MOV immediate instruction allows access to any location in the memory using a 64-bit linear address.

Direct addressing	Direct addressing moves a byte or word between a memory location and a register. The instruction set does not support a memory-to-memory transfer, except with the MOVS instruction. (Example: The MOV CX, LIST instruction copies the word-sized contents of memory location LIST into register CX.) In the 80386 and above, a doubleword-sized memory location can also be addressed. (Example: The MOV ESI, LIST instruction copies a 32-bit number, stored in four consecutive bytes of memory, from location LIST into register ESI.) The direct memory instructions in the 64-bit mode use a full 64-bit linear address.
Register indirect addressing	Register indirect addressing transfers a byte or word between a register and a memory location addressed by an index or base register. The index and base registers are BP, BX, DI, and S1. (Example: The MOV AX, [BX] instruction copies the word-sized data from the data segment offset address indexed by BX into register AX.) In the 80386 and above, a byte, word, or doubleword is transferred between a register and a memory location addressed by any register: EAX, EBX, ECX, EDX, EBP, EDI, or ESI. (Example: The MOV AL, [ECX] instruction loads AL from the data segment offset address selected by the contents of ECX.) In 64-bit mode, the indirect address remains 32 bits in size, which means this form of addressing at present only allows access to 4G bytes of address space if the program operates in the 32-bit compatible mode. In the full 64-bit mode, any address is accessed using either a 64-bit address or the address contained in a register.
Base-plus-index addressing	Base-plus-index addressing transfers a byte or word between a register and the memory location addressed by a base register (BP or BX) plus an index register (DI or SI). (Example: The MOV [BX+DI], CL instruction copies the byte-sized contents of register CL into the data segment memory location addressed by BX plus DI.) In the 80386 and above, any two registers (EAX, EBX, ECX, EDX, EBP, EDI, or ESI) may be combined to generate the memory address. (Example: The MOV [EAX+EBX], CL instruction copies the byte-sized contents of register CL into the data segment memory location addressed by EAX plus EBX.)
Register relative addressing	Register relative addressing moves a byte or word between a register and the memory location addressed by an index or base register plus a displacement. (Example: MOV AX,[BX+4] or MOV AX,ARRAY[BX]. The first instruction loads AX from the data segment address formed by BX plus 4. The second instruction loads AX from the data segment memory location in ARRAY plus the contents of BX.) The 80386 and above use any 32-bit register except ESP to address memory. (Example: MOV AX,[ECX+4] or MOV AX,ARRAY[EBX]. The first instruction loads AX from the data segment address formed by ECX plus 4. The second instruction loads AX from the data segment memory location ARRAY plus the contents of EBX.)

Base relative-plus-index addressing	Base relative-plus-index addressing transfers a byte or word between a register and the memory location addressed by a base and an index register plus a displacement. (Example: MOV AX, ARRAY[BX+DI] or MOV AX, [BX+DI+4]. These instructions load AX from a data segment memory location. The first instruction uses an address formed by adding ARRAY, BX, and DI and the second by adding BX, DI, and 4.) In the 80386 and above, MOV EAX, ARRAY[EBX+ECX] loads EAX from the data segment memory location accessed by the sum of ARRAY, EBX, and ECX.
Scaled-index addressing	Scaled-index addressing is available only in the 80386 through the Pentium 4 microprocessor. The second register of a pair of registers is modified by the scale factor of 2×, 4×, or 8× to generate the operand memory address. (Example: A MOV EDX, [EAX+4*EBX] instruction loads EDX from the data segment memory location addressed by EAX plus four times EBX.) Scaling allows access to word (2×), doubleword (4×), or quadword (8×) memory array data. Note that a scaling factor of 1× also exists, but it is normally implied and does not appear explicitly in the instruction. The MOV AL, [EBX+ECX] is an example in which the scaling factor is a one. Alternately, the instruction can be rewritten as MOV AL, [EBX+1*ECX]. Another example is a MOV AL, [2*EBX] instruction, which uses only one scaled register to address memory.
RIP relative addressing	This addressing mode is only available to the 64-bit extensions on the Pentium 4 or Core2. This mode allows access to any location in the memory system by adding a 32-bit displacement to the 64-bit contents of the 64-bit instruction pointer. For example, if RIP = 1000000000H and a 32-bit displacement is 300H, the location accessed is 1000000300H. The displacement is signed so data located within ±2G from the instruction is accessible by this addressing mode.

Register Addressing

Register addressing is the most common form of data addressing and, once the register names are learned, is the easiest to apply. The microprocessor contains the following 8-bit register names used with register addressing: AH, AL, BH, BL, CH, CL, DH, and DL. Also present are the following 16-bit register names: AX, BX, CX, DX, SP, BP, SI, and DI. In the 80386 and above, the extended 32-bit register names are: EAX, EBX, ECX, EDX, ESP, EBP, EDI, and ESI. In the 64-bit mode of the Pentium 4, the register names are: RAX, RBX, RCX, RDX, RSP, RBP, RDI, RSI, and R8 through R15. With register addressing, some MOV instructions and the PUSH and POP instructions also use the 16-bit segment register names (CS, ES, DS, SS, FS, and GS). It is important for instructions to use registers that are the same size. *Never* mix an 8-bit register with a 16-bit register, an 8-bit register with a 32-bit register, or a l6-bit register with a 32-bit register because this is not allowed by the microprocessor and results in an error when assembled. Likewise never mix 64-bit registers with any other size register. This is even true when a MOV AX, AL (MOV EAX, AL) instruction may seem to make sense. Of course, the MOV AX, AL or MOV EAX, AL instructions are *not* allowed because the registers are of different sizes. Note that a few instructions, such as SHL DX, CL, are exceptions to this rule, as indicated in later chapters. It is also important to note that *none* of the MOV instructions affect the flag bits. The flag bits are normally modified by arithmetic or logic instructions.

Table 3–1 shows many variations of register move instructions. It is impossible to show all combinations because there are too many. For example, just the 8-bit subset of the MOV instruction
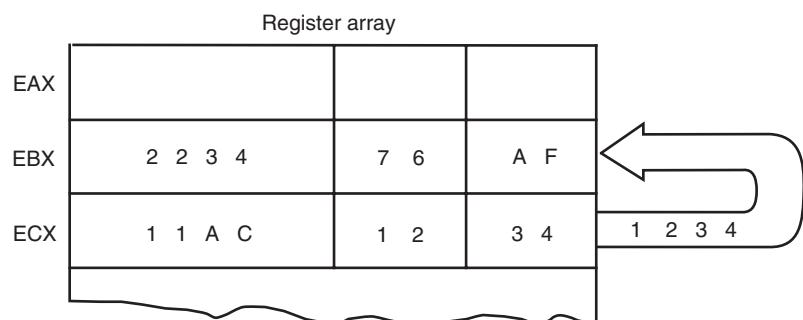
TABLE 3–1 Examples of register-addressed instructions.

Assembly Language	Size	Operation
MOV AL,BL	8 bits	Copies BL into AL
MOV CH,CL	8 bits	Copies CL into CH
MOV R8B,CL	8 bits	Copies CL to the byte portion of R8 (64-bit mode)
MOV R8B,CH	8 bits	Not allowed
MOV AX,CX	16 bits	Copies CX into AX
MOV SP,BP	16 bits	Copies BP into SP
MOV DS,AX	16 bits	Copies AX into DS
MOV BP,R10W	16 bits	Copies R10 into BP (64-bit mode)
MOV SI,DI	16 bits	Copies DI into SI
MOV BX,ES	16 bits	Copies ES into BX
MOV ECX,EBX	32 bits	Copies EBX into ECX
MOV ESP,EDX	32 bits	Copies EDX into ESP
MOV EDX,R9D	32 bits	Copies R9 into EDX (64-bit mode)
MOV RAX,RDX	64 bits	Copies RDX into RAX
MOV DS,CX	16 bits	Copies CX into DS
MOV ES,DS	—	Not allowed (segment-to-segment)
MOV BL,DX	—	Not allowed (mixed sizes)
MOV CS,AX	—	Not allowed (the code segment register may not be the destination register)

has 64 different variations. A segment-to-segment register MOV instruction is about the only type of register MOV instruction *not* allowed. Note that the code segment register is *not* normally changed by a MOV instruction because the address of the next instruction is found by both IP/EIP and CS. If only CS were changed, the address of the next instruction would be unpredictable. Therefore, changing the CS register with a MOV instruction is not allowed.

Figure 3–3 shows the operation of the MOV BX, CX instruction. Note that the source register's contents do not change, but the destination register's contents do change. This instruction moves *(copies)* a l234H from register CX into register BX. This *erases* the old contents (76AFH) of register BX, but the contents of CX remain unchanged. The contents of the destination register or destination memory location change for all instructions except the CMP and TEST instructions. Note that the MOV BX, CX instruction does not affect the leftmost 16 bits of register EBX.

FIGURE 3–3 The effect of executing the MOV BX, CX instruction at the point just before the BX register changes. Note that only the rightmost 16 bits of register EBX change.

Example 3–1 shows a sequence of assembled instructions that copy various data between 8-, 16-, and 32-bit registers. As mentioned, the act of moving data from one register to another changes only the destination register, never the source. The last instruction in this example (MOV CS,AX) assembles without error, but causes problems if executed. If only the contents of CS change without changing IP, the next step in the program is unknown and therefore causes the program to go awry.

EXAMPLE 3–1

```
0000 8B C3          MOV AX,BX    ;copy contents of BX into AX
0002 8A CE          MOV CL,DH    ;copy contents of DH into CL
0004 8A CD          MOV CL,CH    ;copy contents of CH into CL
0006 66|8B C3       MOV EAX,EBX  ;copy contents of EBX into EAX
0009 66|8B D8       MOV EBX,EAX  ;copy contents of EAX into EBX
000C 66|8B C8       MOV ECX,EAX  ;copy contents of EAX into ECX
000F 66|8B D0       MOV EDX,EAX  ;copy contents of EAX into EDX
0012 8C C8          MOV AX,CS    ;copy CS into DS (two steps)
0014 8E D8          MOV DS,AX
0016 8E C8          MOV CS,AX    ;copy AX into CS (causes problems)
```

Immediate Addressing
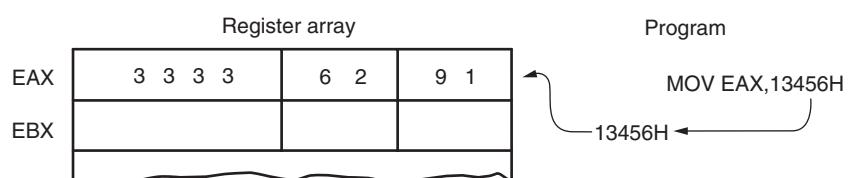
Another data-addressing mode is immediate addressing. The term *immediate* implies that the data immediately follow the hexadecimal opcode in the memory. Also note that immediate data are constant data, whereas the data transferred from a register or memory location are variable data. Immediate addressing operates upon a byte or word of data. In the 80386 through the Core2 microprocessors, immediate addressing also operates on doubleword data. The MOV immediate instruction transfers a copy of the immediate data into a register or a memory location. Figure 3–4 shows the operation of a MOV EAX,13456H instruction. This instruction copies the 13456H from the instruction, located in the memory immediately following the hexadecimal opcode, into register EAX. As with the MOV instruction illustrated in Figure 3–3, the source data overwrites the destination data.

In symbolic assembly language, the symbol # precedes immediate data in some assemblers. The MOV AX,#3456H instruction is an example. Most assemblers do not use the # symbol, but represent immediate data as in the MOV AX,3456H instruction. In this text, the # symbol is not used for immediate data. The most common assemblers—Intel ASM, Microsoft MASM,[2] and Borland TASM[3]—do not use the # symbol for immediate data, but an older assembler used with some Hewlett-Packard logic development system does, as may others.

As mentioned, the MOV immediate instruction under 64-bit operation can include a 64-bit immediate number. An instruction such as MOV RAX,123456780A311200H is allowed in the 64-bit mode.

The symbolic assembler portrays immediate data in many ways. The letter **H** appends hexadecimal data. If hexadecimal data begin with a letter, the assembler requires that the data

FIGURE 3–4 The operation of the MOV EAX,3456H instruction. This instruction copies the immediate data (13456H) into EAX.

---

[2]MASM (MACRO assembler) is a trademark of Microsoft Corporation.

[3]TASM (Turbo assembler) is a trademark of Borland Corporation.

TABLE 3–2   Examples of immediate addressing using the MOV instruction.

Assembly Language	Size	Operation
MOV BL,44	8 bits	Copies 44 decimal (2CH) into BL
MOV AX,44H	16 bits	Copies 0044H into AX
MOV SI,0	16 bits	Copies 0000H into SI
MOV CH,100	8 bits	Copies 100 decimal (64H) into CH
MOV AL,'A'	8 bits	Copies ASCII A into AL
MOV AH,1	8 bits	Not allowed in 64-bit mode, but allowed in 32- or 16-bit modes
MOV AX,'AB'	16 bits	Copies ASCII BA* into AX
MOV CL,11001110B	8 bits	Copies 11001110 binary into CL
MOV EBX,12340000H	32 bits	Copies 12340000H into EBX
MOV ESI,12	32 bits	Copies 12 decimal into ESI
MOV EAX,100B	32 bits	Copies 100 binary into EAX
MOV RCX,100H	64 bits	Copies 100H into RCX

*Note: This is not an error. The ASCII characters are stored as BA, so exercise care when using word-sized pairs of ASCII characters.

start with a **0**. For example, to represent a hexadecimal F2, 0F2H is used in assembly language. In some assemblers (though not in MASM, TASM, or this text), hexadecimal data are represented with an 'h, as in MOV AX,#'h1234. Decimal data are represented as is and require no special codes or adjustments. (An example is the 100 decimal in the MOV AL,100 instruction.) An ASCII-coded character or characters may be depicted in the immediate form if the ASCII data are enclosed in apostrophes. (An example is the MOV BH, 'A' instruction, which moves an ASCII-coded letter A [41H] into register BH.) Be careful to use the apostrophe (') for ASCII data and not the single quotation mark ('). Binary data are represented if the binary number is followed by the letter B, or, in some assemblers, the letter Y. Table 3–2 shows many different variations of MOV instructions that apply immediate data.

Example 3–2 shows various immediate instructions in a short assembly language program that places 0000H into the 16-bit registers AX, BX, and CX. This is followed by instructions that use register addressing to copy the contents of AX into registers SI, DI, and BP. This is a complete program that uses programming models for assembly and execution with MASM. The .MODEL TINY statement directs the assembler to assemble the program into a single code segment. The .CODE statement or directive indicates the start of the code segment; the .STARTUP statement indicates the starting instruction in the program; and the .EXIT statement causes the program to exit to DOS. The END statement indicates the end of the program file. This program is assembled with MASM and executed with CodeView[4] (CV) to view its execution. Note that the most recent version of TASM will also accept MASM code without any changes. To store the program into the system use the DOS EDIT program, Windows NotePad,[5] or Programmer's WorkBench[6] (PWB). Note that a TINY program always assembles as a command (.COM) program.

---

[4]CodeView is a registered trademark of Microsoft Corporation.

[5]NotePad is a registered trademark of Microsoft Corporation.

[6]Programmer's WorkBench is a registered trademark of Microsoft Corporation.

EXAMPLE 3–2

```
                    .MODEL TINY             ;choose single segment model
0000                .CODE                   ;start of code segment
                    .STARTUP                ;start of program

0100 B8 0000        MOV AX,0                ;place 0000H into AX
0103 BB 0000        MOV BX,0                ;place 0000H into BX
0106 B9 0000        MOV CX,0                ;place 0000H into CX

0109 8B F0          MOV SI,AX               ;copy AX into SI
010B 8B F8          MOV DI,AX               ;copy AX into DI
010D 8B E8          MOV BP,AX               ;copy AX into BP

                    .EXIT                   ;exit to DOS
                    END                     ;end of program
```

Each statement in an assembly language program consists of four parts or fields, as illustrated in Example 3–3. The leftmost field is called the *label* and it is used to store a symbolic name for the memory location that it represents. All labels must begin with a letter or one of the following special characters: @, $, -, or ? A label may be of any length from 1 to 35 characters. The label appears in a program to identify the name of a memory location for storing data and for other purposes that are explained as they appear. The next field to the right is called the *opcode field;* it is designed to hold the instruction, or opcode. The MOV part of the move data instruction is an example of an opcode. To the right of the opcode field is the *operand field,* which contains information used by the opcode. For example, the MOV AL,BL instruction has the opcode MOV and operands AL and BL. Note that some instructions contain between zero and three operands. The final field, the *comment field,* contains a comment about an instruction or a group of instructions. A comment always begins with a semicolon (;).

EXAMPLE 3–3

```
Label       Opcode  Operand    Comment

DATA1       DB      23H        ;define DATA1 as a byte of 23H
DATA2       DW      1000H      ;define DATA2 as a word of 1000H

START:      MOV     AL,BL      ;copy BL into AL
            MOV     BH,AL      ;copy AL into BH
            MOV     CX,200     ;copy 200 into CX
```

When the program is assembled and the list (.LST) file is viewed, it appears as the program listed in Example 3–2. The hexadecimal number at the far left is the offset address of the instruction or data. This number is generated by the assembler. The number or numbers to the right of the offset address are the machine-coded instructions or data that are also generated by the assembler. For example, if the instruction MOV AX,0 appears in a file and it is assembled, it appears in offset memory location 0100 in Example 3–2. Its hexadecimal machine language form is B8 0000. The B8 is the opcode in machine language and the 0000 is the 16-bit-wide data with a value of zero. When the program was written, only the MOV AX,0 was typed into the editor; the assembler generated the machine code and addresses, and stored the program in a file with the extension .LST. Note that all programs shown in this text are in the form generated by the assembler.

EXAMPLE 3–4

```
int MyFunction(int temp)
{
      _asm
      {
          mov   eax,temp
          add   eax,20h
          mov   temp,eax
      }
      return temp;      // return a 32-bit integer
}
```

Programs are also written using the inline assembler in some Visual C++ programs. Example 3–4 shows a function in a Visual C++ program that includes some code written with the inline assembler. This function adds 20H to the number returned by the function. Notice that the assembly code accesses C++ variable temp and all of the assembly code is placed in an **_asm** code block. Many examples in this text are written using the inline assembler within a C++ program.

Direct Data Addressing

Most instructions can use the direct data-addressing mode. In fact, direct data addressing is applied to many instructions in a typical program. There are two basic forms of direct data addressing: (1) direct addressing, which applies to a MOV between a memory location and AL, AX, or EAX, and (2) displacement addressing, which applies to almost any instruction in the instruction set. In either case, the address is formed by adding the displacement to the default data segment address or an alternate segment address. In 64-bit operation, the direct-addressing instructions are also used with a 64-bit linear address, which allows access to any memory location.

Direct Addressing.   Direct addressing with a MOV instruction transfers data between a memory location, located within the data segment, and the AL (8-bit), AX (l6-bit), or EAX (32-bit) register. A MOV instruction using this type of addressing is usually a 3-byte long instruction. (In the 80386 and above, a register size prefix may appear before the instruction, causing it to exceed 3 bytes in length.)

The MOV AL,DATA instruction, as represented by most assemblers, loads AL from the data segment memory location DATA (1234H). Memory location DATA is a symbolic memory location, while the 1234H is the actual hexadecimal location. With many assemblers, this instruction is represented as a MOV AL,[1234H][7] instruction. The [1234H] is an absolute memory location that is not allowed by all assembler programs. Note that this may need to be formed as MOV AL, DS:[1234H] with some assemblers, to show that the address is in the data segment. Figure 3–5 shows how this instruction transfers a copy of the byte-sized contents of memory location 11234H into AL. The effective address is formed by adding 1234H (the offset address) and 10000H (the data segment address of 1000H times 10H) in a system operating in the real mode.

Table 3–3 lists the direct-addressed instructions. These instructions often appear in programs, so Intel decided to make them special 3-byte-long instructions to reduce the length of programs. All other instructions that move data from a memory location to a register, called displacement-addressed instructions, require 4 or more bytes of memory for storage in a program.

Displacement Addressing.   Displacement addressing is almost identical to direct addressing, except that the instruction is 4 bytes wide instead of 3. In the 80386 through the Pentium 4,
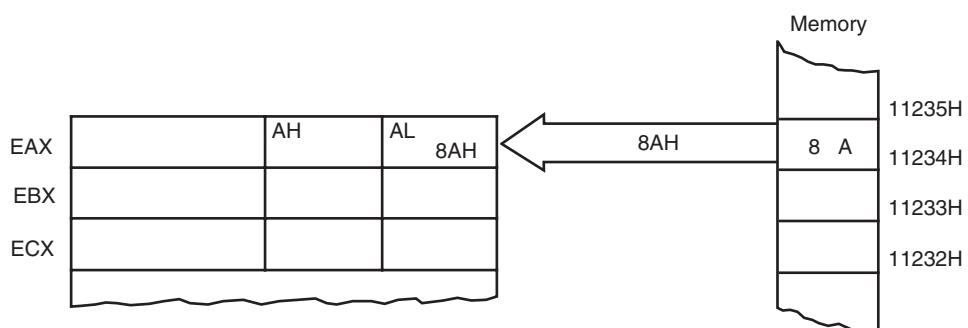
FIGURE 3–5   The operation of the MOV AL,[1234H] instruction when DS = 1000H.

---

[7]This form may be typed into a MASM program, but it most often appears when the debugging tool is executed.

TABLE 3–3 Direct addressed instructions using EAX, AX, and AL and RAX in 64-bit mode.

Assembly Language	Size	Operation
MOV AL,NUMBER	8 bits	Copies the byte contents of data segment memory location NUMBER into AL
MOV AX,COW	16 bits	Copies the word contents of data segment memory location COW into AX
MOV EAX,WATER*	32 bits	Copies the doubleword contents of data segment location WATER into EAX
MOV NEWS,AL	8 bits	Copies AL into byte memory location NEWS
MOV THERE,AX	16 bits	Copies AX into word memory location THERE
MOV HOME,EAX*	32 bits	Copies EAX into doubleword memory location HOME
MOV ES:[2000H],AL	8 bits	Copies AL into extra segment memory at offset address 2000H
MOV AL,MOUSE	8 bits	Copies the contents of location MOUSE into AL; in 64-bit mode MOUSE can be any address
MOV RAX,WHISKEY	64 bits	Copies 8 bytes from memory location WHISKEY into RAX

*Note: The 80386–Pentium 4 at times use more than 3 bytes of memory for 32-bit instructions.

this instruction can be up to 7 bytes wide if both a 32-bit register and a 32-bit displacement are specified. This type of direct data addressing is much more flexible because most instructions use it.

If the operation of the MOV CL,DS:[1234H] instruction is compared to that of the MOV AL,DS:[1234H] instruction of Figure 3–5, we see that both basically perform the same operation except for the destination register (CL versus AL). Another difference only becomes apparent upon examining the assembled versions of these two instructions. The MOV AL,DS:[1234H] instruction is 3 bytes long and the MOV CL,DS:[1234H] instruction is 4 bytes long, as illustrated in Example 3–5. This example shows how the assembler converts these two instructions into hexadecimal machine language. You must include the segment register DS: in this example, before the [offset] part of the instruction. You may use any segment register, but in most cases, data are stored in the data segment, so this example uses DS:[1234H].

EXAMPLE 3–5

```
0000 A0 1234 R          MOV AL,DS:[1234H]
0003 BA 0E 1234 R       MOV CL,DS:[1234H]
```

Table 3–4 lists some MOV instructions using the displacement form of direct addressing. Not all variations are listed because there are many MOV instructions of this type. The segment registers can be stored or loaded from memory.

Example 3–6 shows a short program using models that address information in the data segment. Note that the data segment begins with a .DATA statement to inform the assembler where the data segment begins. The model size is adjusted from TINY, as shown in Example 3–3, to SMALL so that a data segment can be included. The SMALL model allows one data segment and one code segment. The SMALL model is often used whenever memory data are required for a program. A SMALL model program assembles as an execute (.EXE) program file. Notice how this example allocates memory locations in the data segment by using the DB and DW directives. Here the .STARTUP statement not only indicates the start of the code, but it also loads the data segment register with the

TABLE 3–4 Examples of direct data addressing using a displacement.

Assembly Language	Size	Operation
MOV CH,DOG	8 bits	Copies the byte contents of data segment memory location DOG into CH
MOV CH,DS:[1000H]*	8 bits	Copies the byte contents of data segment memory offset address 1000H into CH
MOV ES,DATA6	16 bits	Copies the word contents of data segment memory location DATA6 into ES
MOV DATA7,BP	16 bits	Copies BP into data segment memory location DATA7
MOV NUMBER,SP	16 bits	Copies SP into data segment memory location NUMBER
MOV DATA1,EAX	32 bits	Copies EAX into data segment memory location DATA1
MOV EDI,SUM1	32 bits	Copies the doubleword contents of data segment memory location SUM1 into EDI

*This form of addressing is seldom used with most assemblers because an actual numeric offset address is rarely accessed.

segment address of the data segment. If this program is assembled and executed with CodeView, the instructions can be viewed as they execute and change registers and memory locations.

EXAMPLE 3–6

```
                        .MODEL SMALL       ;choose small model
0000                    .DATA              ;start data segment

0000 10         DATA1 DB 10H               ;place 10H into DATA1
0001 00         DATA2 DB 0                 ;place 00H into DATA2
0002 0000       DATA3 DW 0                 ;place 0000H into DATA3
0004 AAAA       DATA4 DW 0AAAAH            ;place AAAAH into DATA4

0000                    .CODE              ;start code segment
                        .STARTUP           ;start program

0017 A00000 R           MOV AL,DATA1       ;copy DATA1 into AL
001A 8A 26 0001 R       MOV AH,DATA2       ;copy DATA2 into AH
001E A3 0002 R          MOV DATA3,AX       ;copy AX into DATA3
0021 8B 1E 0004 R       MOV BX,DATA4       ;copy DATA4 into BX

                        .EXIT;             exit to DOS
                        END;               end program listing
```
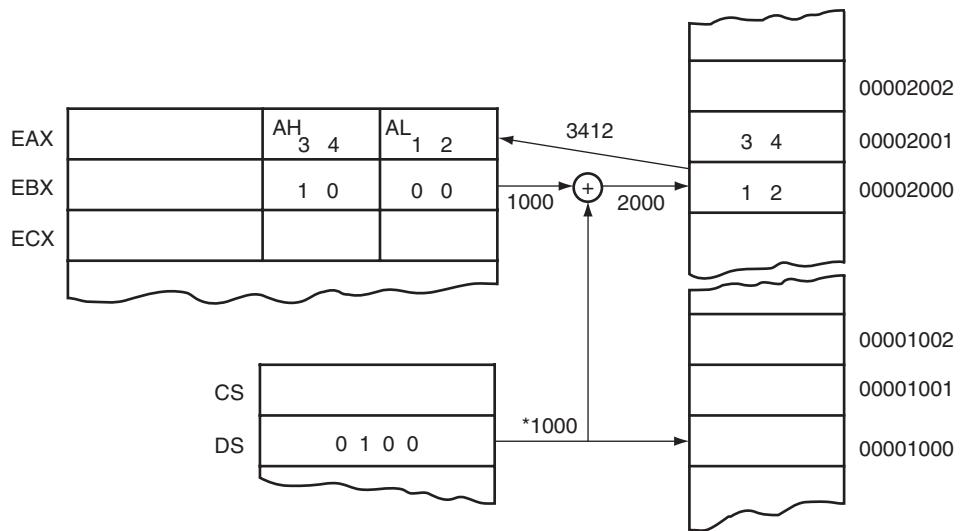
Register Indirect Addressing

Register indirect addressing allows data to be addressed at any memory location through an offset address held in any of the following registers: BP, BX, DI, and SI. For example, if register BX contains 1000H and the MOV AX,[BX] instruction executes, the word contents of data segment offset address 1000H are copied into register AX. If the microprocessor is operated in the real mode and DS = 0100H, this instruction addresses a word stored at memory bytes 2000H and 2001H, and transfers it into register AX (see Figure 3–6). Note that the contents of 2000H are moved into AL and the contents of 2001H are moved into AH. The [ ] symbols denote indirect addressing in assembly language. In addition to using the BP, BX, DI, and SI registers to indirectly address memory, the 80386 and above allow register indirect addressing with any extended register except ESP. Some typical instructions using indirect addressing appear in Table 3–5. If a Pentium 4 or Core2 is available that operates in the 64-bit mode, any 64-bit register is used to hold a 64-bit linear address. In the 64-bit mode, the segment registers serve no purpose in addressing a location in the flat model.

*After DS is appended with a 0.

FIGURE 3–6   The operation of the MOV AX,[BX] instruction when BX = 1000H and DS = 0100H. Note that this instruction is shown after the contents of memory are transferred to AX.

The **data segment** is used by default with register indirect addressing or any other addressing mode that uses BX, DI, or SI to address memory. If the BP register addresses memory, the **stack segment** is used by default. These settings are considered the default for these four index and base registers. For the 80386 and above, EBP addresses memory in the stack segment by default; EAX, EBX, ECX, EDX, EDI, and ESI address memory in the data segment by fault. When using a 32-bit register to address memory in the real mode, the contents of the 32-bit register must never

TABLE 3–5   Examples of register indirect addressing.

Assembly Language	Size	Operation
MOV CX,[BX]	16 bits	Copies the word contents of the data segment memory location addressed by BX into CX
MOV [BP],DL*	8 bits	Copies DL into the stack segment memory location addressed by BP
MOV [DI],BH	8 bits	Copies BH into the data segment memory location addressed by DI
MOV [DI],[BX]	—	Memory-to-memory transfers are not allowed except with string instructions
MOV AL,[EDX]	8 bits	Copies the byte contents of the data segment memory location addressed by EDX into AL
MOV ECX,[EBX]	32 bits	Copies the doubleword contents of the data segment memory location addressed by EBX into ECX
MOV RAX,[RDX]	64 bits	Copies the quadword contents of the memory location address by the linear address located in RDX into RAX (64-bit mode)

*Note: Data addressed by BP or EBP are by default in the stack segment, while other indirect addressed instructions use the data segment by default.
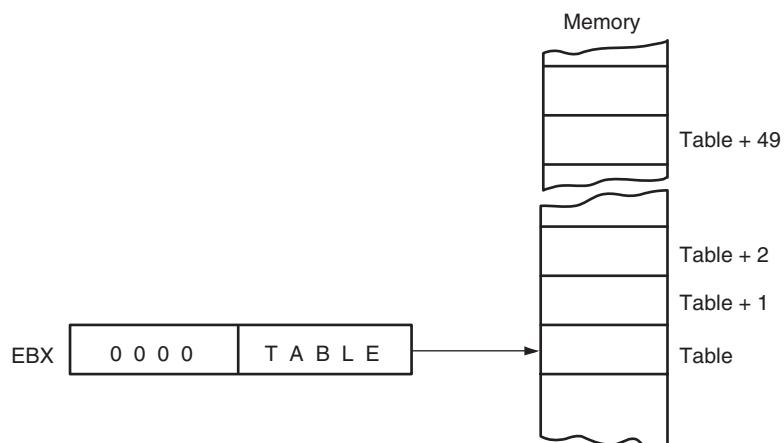
exceed 0000FFFFH. In the protected mode, any value can be used in a 32-bit register that is used to indirectly address memory, as long as it does not access a location outside of the segment, as dictated by the access rights byte. An example 80386–Pentium 4 instruction is MOV EAX,[EBX]. This instruction loads EAX with the doubleword-sized number stored at the data segment offset address indexed by EBX. In the 64-bit mode, the segment registers are not used in the address calculation because the register contains the actual linear memory address.

In some cases, indirect addressing requires specifying the size of the data. The size is specified by the **special assembler directive** BYTE PTR, WORD PTR, DWORD PTR, or QWORD PTR. These directives indicate the size of the memory data addressed by the memory **pointer** (PTR). For example, the MOV AL,[DI] instruction is clearly a byte-sized move instruction, but the MOV [DI],10H instruction is ambiguous. Does the MOV [DI],10H instruction address a byte-, word-, doubleword-, or quadword-sized memory location? The assembler can't determine the size of the 10H. The instruction MOV BYTE PTR [DI],10H clearly designates the location addressed by DI as a byte-sized memory location. Likewise, the MOV DWORD PTR [DI],10H clearly identifies the memory location as doubleword-sized. The BYTE PTR, WORD PTR, DWORD PTR, and QWORD PTR directives are used only with instructions that address a memory location through a pointer or index register with immediate data, and for a few other instructions that are described in subsequent chapters. Another directive that is occasionally used is the QWORD PTR, where a QWORD is a quadword (64-bits mode). If programs are using the SIMD instructions, the OWORD PTR, an octal word, is also used to represent a 128-bit-wide number.

Indirect addressing often allows a program to refer to tabular data located in the memory system. For example, suppose that you must create a table of information that contains 50 samples taken from memory location 0000:046C. Location 0000:046C contains a counter in DOS that is maintained by the personal computer's real-time clock. Figure 3–7 shows the table and the BX register used to sequentially address each location in the table. To accomplish this task, load the starting location of the table into the BX register with a MOV immediate instruction. After initializing the starting address of the table, use register indirect addressing to store the 50 samples sequentially.

The sequence shown in Example 3–7 loads register BX with the starting address of the table and it initializes the count, located in register CX, to 50. The **OFFSET** directive tells the assembler to load BX with the offset address of memory location TABLE, not the contents of TABLE. For example, the MOV BX,DATAS instruction copies the contents of memory location DATAS into BX, while the MOV BX,OFFSET DATAS instruction copies the offset address DATAS into BX. When the OFFSET directive is used with the MOV instruction, the assembler calculates the offset address and then uses a MOV immediate instruction to load the address in the specified 16-bit register.

FIGURE 3–7   An array (TABLE) containing 50 bytes that are indirectly addressed through register BX.

EXAMPLE 3–7

```
                              .MODEL SMALL          ;select small model
0000                          .DATA                 ;start data segment

0000 0032 [       DATAS  DW  50 DUP(?) ;setup array of 50 words
        0000
             ]
0000                          .CODE                 ;start code segment
                              .STARTUP              ;start program
0017 B8 0000                  MOV  AX,0
001A 8E C0                    MOV  ES,AX             ;address segment 0000 with ES
001C B8 0000 R                MOV  BX,OFFSET DATAS   ;address DATAS array with BX
001F B9 0032                  MOV  CX,50            ;load counter with 50
0022                AGAIN:
0022 26:A1 046C               MOV  AX,ES:[046CH]    ;get clock value
0026 89 07                    MOV  [BX],AX          ;save clock value in DATAS
0028 43                       INC  BX               ;increment BX to next element
0029 43                       INC BX
002A E2 F6                    LOOP AGAIN            ;repeat 50 times

                              .EXIT                 ;exit to DOS
                              END                   ;end program listing
```

Once the counter and pointer are initialized, a repeat-until CX = 0 loop executes. Here data are read from extra segment memory location 46CH with the MOV AX,ES:[046CH] instruction and stored in memory that is indirectly addressed by the offset address located in register BX. Next, BX is incremented (1 is added to BX) twice to address the next word in the table. Finally, the LOOP instruction repeats the LOOP 50 times. The LOOP instruction decrements (subtracts 1 from) the counter (CX); if CX is not zero, LOOP causes a jump to memory location AGAIN. If CX becomes zero, no jump occurs and this sequence of instructions ends. This example copies the most recent 50 values from the clock into the memory array DATAS. This program will often show the same data in each location because the contents of the clock are changed only 18.2 times per second. To view the program and its execution, use the CodeView program. To use CodeView, type CV XXXX.EXE, where XXXX.EXE is the name of the program that is being debugged. You can also access it as DEBUG from the Programmer's WorkBench program under the RUN menu. Note that CodeView functions only with .EXE or .COM files. Some useful CodeView switches are /50 for a 50-line display and /S for use of high-resolution video displays in an application. To debug the file TEST.COM with 50 lines, type CV /50 /S TEST.COM at the DOS prompt.

Base-Plus-Index Addressing

Base-plus-index addressing is similar to indirect addressing because it indirectly addresses memory data. In the 8086 through the 80286, this type of addressing uses one base register (BP or BX) and one index register (DI or SI) to indirectly address memory. The base register often holds the beginning location of a memory array, whereas the index register holds the relative position of an element in the array. Remember that whenever BP addresses memory data, both the stack segment register and BP generate the effective address.

In the 80386 and above, this type of addressing allows the combination of any two 32-bit extended registers except ESP. For example, the MOV DL,[EAX+EBX] instruction is an example using EAX (as the base) plus EBX (as the index). If the EBP register is used, the data are located in the stack segment instead of in the data segment.

Locating Data with Base-Plus-Index Addressing. Figure 3–8 shows how data are addressed by the MOV DX,[BX+DI] instruction when the microprocessor operates in the real mode. In this example, BX = 1000H, DI = 0010H, and DS = 0100H, which translate into memory address 02010H. This instruction transfers a copy of the word from location 02010H into the DX register.
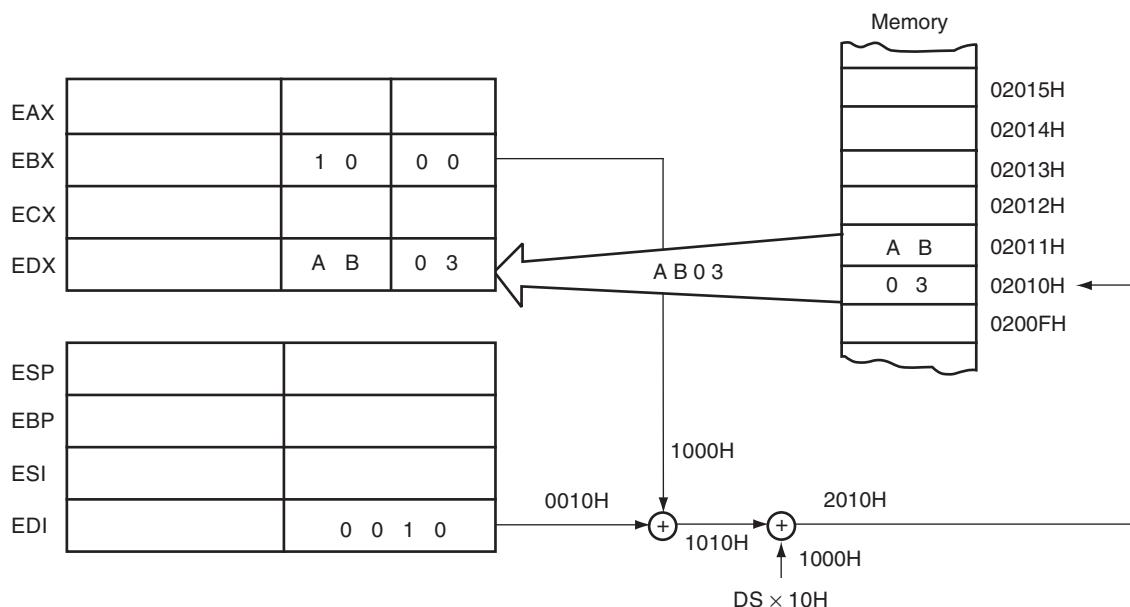
Memory

FIGURE 3–8   An example showing how the base-plus-index addressing mode functions for the MOV DX,[BX+DI] instruction. Notice that memory address 02010H is accessed because DS = 0100H, BX = 100H, and DI = 0010H.

Table 3–6 lists some instructions used for base-plus-index addressing. Note that the Intel assembler requires that this addressing mode appear as [BX][DI] instead of [BX+DI]. The MOV DX,[BX+DI] instruction is MOV DX,[BX][DI] for a program written for the Intel ASM assembler. This text uses the first form in all example programs, but the second form can be used in many assemblers, including MASM from Microsoft. Instructions like MOV DI,[BX+DI] will assemble, but will not execute correctly.

Locating Array Data Using Base-Plus-Index Addressing.   A major use of the base-plus-index addressing mode is to address elements in a memory array. Suppose that the elements in an array

TABLE 3–6   Examples of base-plus-index addressing.

Assembly Language	Size	Operation
MOV CX,[BX+DI]	16 bits	Copies the word contents of the data segment memory location addressed by BX plus DI into CX
MOV CH,[BP+SI]	8 bits	Copies the byte contents of the stack segment memory location addressed by BP plus SI into CH
MOV [BX+SI],SP	16 bits	Copies SP into the data segment memory location addressed by BX plus SI
MOV [BP+DI],AH	8 bits	Copies AH into the stack segment memory location addressed by BP plus DI
MOV CL,[EDX+EDI]	8 bits	Copies the byte contents of the data segment memory location addressed by EDX plus EDI into CL
MOV [EAX+EBX],ECX	32 bits	Copies ECX into the data segment memory location addressed by EAX plus EBX
MOV [RSI+RBX],RAX	64 bit	Copies RAX into the linear memory location addressed by RSI plus RBX (64-bit mode)
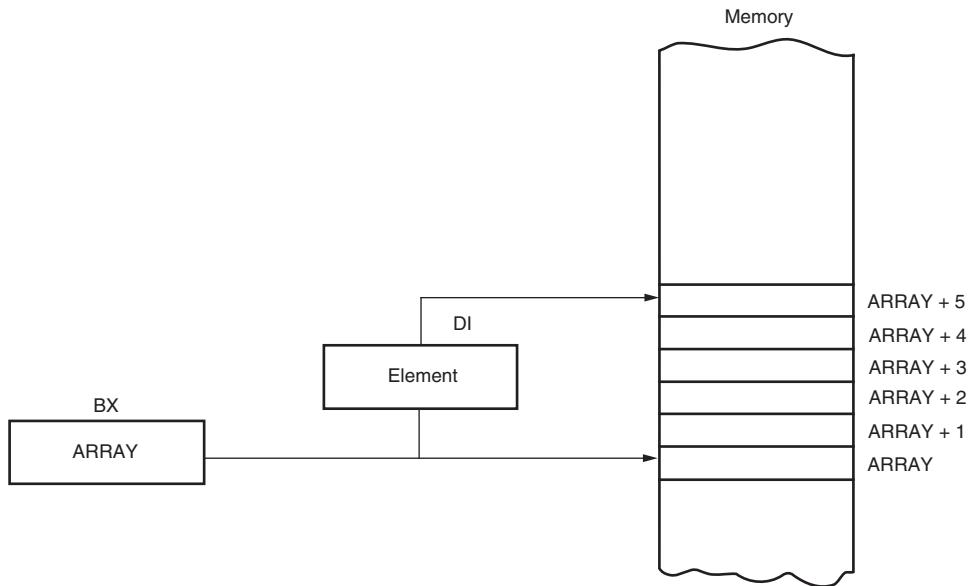
FIGURE 3–9 An example of the base-plus-index addressing mode. Here an element (DI) of an ARRAY (BX) is addressed.

located in the data segment at memory location ARRAY must be accessed. To accomplish this, load the BX register (base) with the beginning address of the array and the DI register (index) with the element number to be accessed. Figure 3–9 shows the use of BX and DI to access an element in an array of data.

A short program, listed in Example 3–8, moves array element 10H into array element 20H. Notice that the array element number, loaded into the DI register, addresses the array element. Also notice how the contents of the ARRAY have been initialized so that element 10H contains 29H.

EXAMPLE 3–8

```
                              .MODEL SMALL          ;select small model
0000                          .DATA                 ;start data segment
0000 0010 [         ARRAY DB  16 DUP(?)             ;setup array of 16 bytes
       00
          ]
0010 29                       DB 29H                ;element 10H
0011 001E [                   DB 20 dup(?)
       00
          ]
0000                          .CODE                 ;start code segment
                              .STARTUP

0017 B8 0000 R                MOV  BX,OFFSET ARRAY  ;address ARRAY
001A BF 0010                  MOV  DI,10H           ;address element 10H
001D 8A 01                    MOV  AL,[BX+DI]       ;get element 10H
001F BF 0020                  MOV  DI,20H           ;address element 20H
0022 88 01                    MOV  [BX+DI],AL       ;save in element 20H

                              .EXIT                 ;exit to DOS
                              END                   ;end program
```
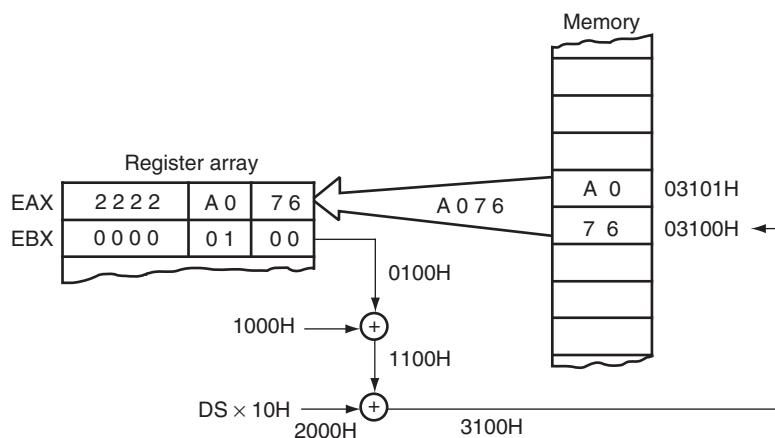
Register Relative Addressing

Register relative addressing is similar to base-plus-index addressing and displacement addressing. In register relative addressing, the data in a segment of memory are addressed by

FIGURE 3–10 The operation of the MOV AX, [BX+1000H] instructon, when BX = 0100H and DS = 0200H.

adding the displacement to the contents of a base or an index register (BP, BX, DI, or SI). Figure 3–10 shows the operation of the MOV AX,[BX+1000H] instruction. In this example, BX = 0100H and DS = 0200H, so the address generated is the sum of DS × 0H, BX, and the displacement of 1000H, which addresses location 03100H. Remember that BX, DI, or SI addresses the data segment and BP addresses the stack segment. In the 80386 and above, the displacement can be a 32-bit number and the register can be any 32-bit register except the ESP register. Remember that the size of a real mode segment is 64K bytes long. Table 3–7 lists a few instructions that use register relative addressing

The displacement is a number added to the register within the [ ], as in the MOV AL,[DI+2] instruction, or it can be a displacement is subtracted from the register, as in MOV AL,[SI–l]. A displacement also can be an offset address appended to the front of the [ ], as in MOV AL,DATA[DI]. Both forms of displacements also can appear simultaneously, as in the MOV AL,DATA[DI+3] instruction. Both forms of the displacement add to the base or base plus index register within the [ ] symbols. In the 8086–80286 microprocessors, the value of the displacement is limited to a 16-bit signed number with a value ranging between +32,767 (7FFFH)

TABLE 3–7 Examples of register relative addressing.

Assembly Language	Size	Operation
MOV AX,[DI+100H]	16 bits	Copies the word contents of the data segment memory location addressed by DI plus 100H into AX
MOV ARRAY[SI],BL	8 bits	Copies BL into the data segment memory location addressed by ARRAY plus SI
MOV LIST[SI+2],CL	8 bits	Copies CL into the data segment memory location addressed by the sum of LIST, SI, and 2
MOV DI,SET_IT[BX]	16 bits	Copies the word contents of the data segment memory location addressed by SET_IT plus BX into DI
MOV DI,[EAX+10H]	16 bits	Copies the word contents of the data segment location addressed by EAX plus 10H into DI
MOV ARRAY[EBX],EAX	32 bits	Copies EAX into the data segment memory location addressed by ARRAY plus EBX
MOV ARRAY[RBX],AL	8 bits	Copies AL into the memory location ARRAY plus RBX (64-bit mode)
MOV ARRAY[RCX],EAX	32 bits	Copies EAX into memory location ARRAY plus RCX (64-bit mode)
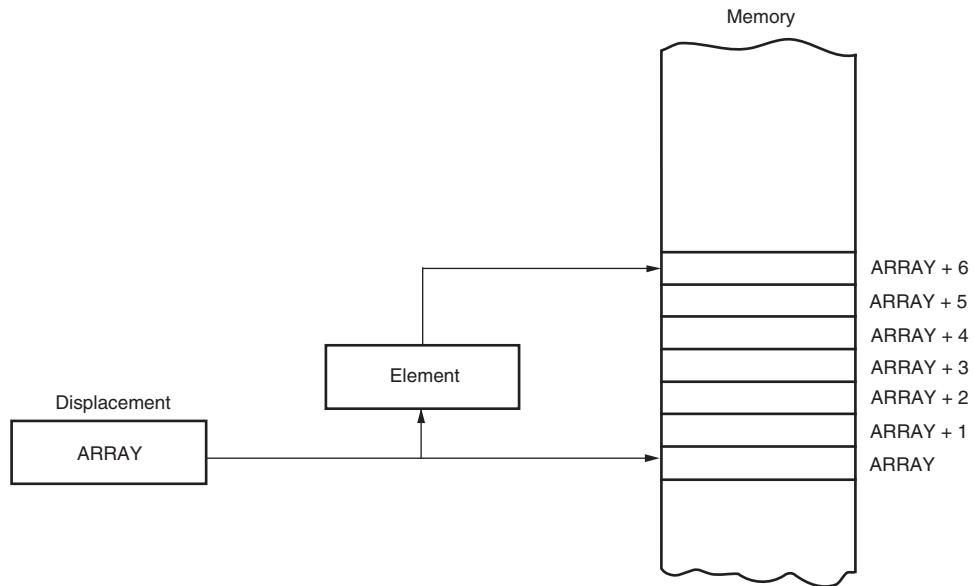
FIGURE 3–11 Register relative addressing used to address an element of ARRAY. The displacement addresses the start of ARRAY, and DI accesses an element.

and –32,768 (8000H); in the 80386 and above, a 32-bit displacement is allowed with a value ranging between +2,147,483,647 (7FFFFFFFH) and −2,147,483,648 (80000000H).

Addressing Array Data with Register Relative. It is possible to address array data with register relative addressing, such as one does with base-plus-index addressing. In Figure 3–11, register relative addressing is illustrated with the same example as for base-plus-index addressing. This shows how the displacement ARRAY adds to index register DI to generate a reference to an array element.

Example 3–9 shows how this new addressing mode can transfer the contents of array element 10H into array element 20H. Notice the similarity between this example and Example 3–8. The main difference is that, in Example 3–9, register BX is not used to address memory ARRAY; instead, ARRAY is used as a displacement to accomplish the same task.

EXAMPLE 3–9
```
                        .MODEL SMALL            ;select small model
0000                    .DATA                   ;start data segment
0000 0010 [     ARRAY   DB  16 dup(?)           ;setup ARRAY
        00
          ]
0010 29                 DB  29                  ;element 10H
0011 001E [             DB  30 dup(?)
        00
          ]
0000                    .CODE                   ;start code segment
                        .STARTUP                ;start program
0017 BF 0010            MOV  DI,10H             ;address element 10H
001A 8A 85 0000 R       MOV  AL,ARRAY[DI]       ;get ARRAY element 10H
001E BF 0020            MOV  DI,20H             ;address element 20H
0021 88 85 0000 R       MOV  ARRAY[DI],AL       ;save it in element 20H
                        .EXIT                   ;exit to DOS
                        END                     ;end of program
```
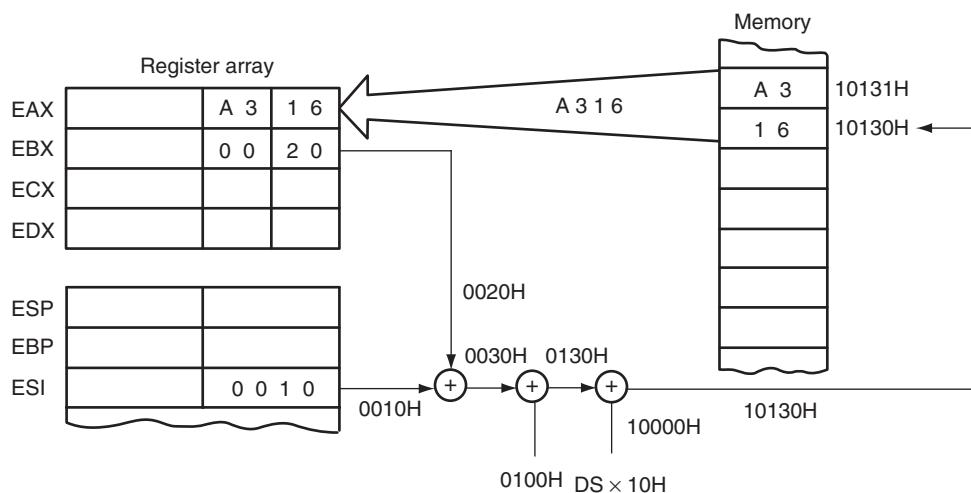
FIGURE 3–12 An example of base relative-plus-index addressing using a MOV AX,[BX+SI+100H] instruction. Note: DS = 1000H.

Base Relative-Plus-Index Addressing
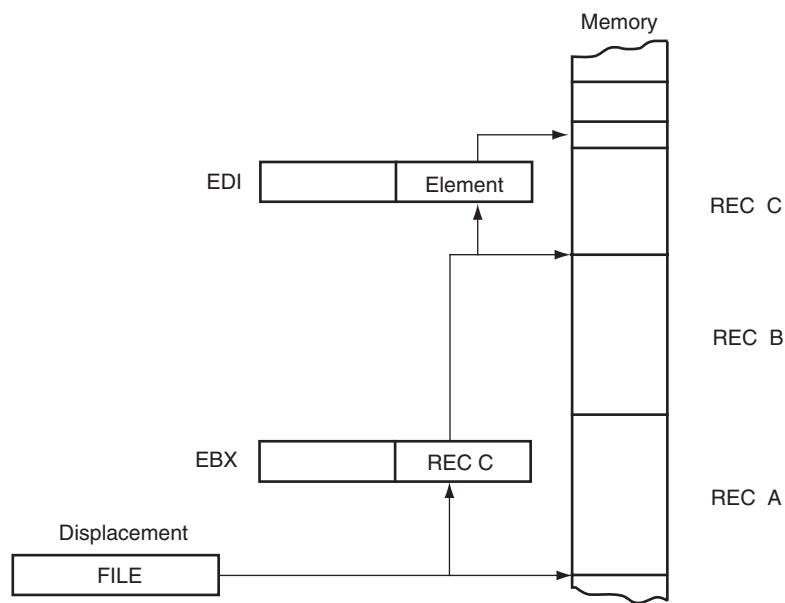
The base relative-plus-index addressing mode is similar to base-plus-index addressing, but it adds a displacement, besides using a base register and an index register, to form the memory address. This type of addressing mode often addresses a two-dimensional array of memory data.

Addressing Data with Base Relative-Plus-Index.   Base relative-plus-index addressing is the least-used addressing mode. Figure 3–12 shows how data are referenced if the instruction executed by the microprocessor is MOV AX,[BX+SI+100H]. The displacement of 100H adds to BX and SI to form the offset address within the data segment. Registers BX = 0020H, SI = 0100H, and DS = 1000H, so the effective address for this instruction is 10130H—the sum of these registers plus a displacement of 100H. This addressing mode is too complex for frequent use in programming. Some typical instructions using base relative-plus-index addressing appear in Table 3–8. Note that with the 80386 and above, the effective address is generated by the sum of two 32-bit registers plus a 32-bit displacement.

TABLE 3–8 Example base relative-plus-index instructions.

Assembly Language	Size	Operation
MOV DH,[BX+DI+20H]	8 bits	Copies the byte contents of the data segment memory location addressed by the sum of BX, DI and 20H into DH
MOV AX,FILE[BX+DI]	16 bits	Copies the word contents of the data segment memory location addressed by the sum of FILE, BX and DI into AX
MOV LIST[BP+DI],CL	8 bits	Copies CL into the stack segment memory location addressed by the sum of LIST, BP, and DI
MOV LIST[BP+SI+4],DH	8 bits	Copies DH into the stack segment memory location addressed by the sum of LIST, BP, SI, and 4
MOV EAX,FILE[EBX+ECX+2]	32 bits	Copies the doubleword contents of the memory location addressed by the sum of FILE, EBX, ECX, and 2 into EAX

FIGURE 3–13 Base relative-plus-index addressing used to access a FILE that contains multiple records (REC).

Addressing Arrays with Base Relative-Plus-Index.   Suppose that a file of many records exists in memory and each record contains many elements. The displacement addresses the file, the base register addresses a record, and the index register addresses an element of a record. Figure 3–13 illustrates this very complex form of addressing.

Example 3–10 provides a program that copies element 0 of record A into element 2 of record C by using the base relative-plus-index mode of addressing. This example FILE contains four records and each record contains 10 elements. Notice how the THIS BYTE statement is used to define the label FILE and RECA as the same memory location.

EXAMPLE 3–10

```
                              .MODEL SMALL          ;select small model
0000                          .DATA                 ;start data segment
0000 = 0000        FILE   EQU  THIS BYTE            ;assign FILE to this byte
0000 000A [        RECA   DB  10 dup(?)             ;10 bytes for record A
         00
       ]
000A 000A [        RECB   DB  10 dup(?)             ;10 bytes for record B
         00
       ]
0014 000A [        RECC   DB  10 dup(?)             ;10 bytes for record C
         00
       ]
001E 000A [        RECD   DB  10 dup(?)             ;10 bytes for record D
         00
       ]
0000                          .CODE                 ;start code segment
                              .STARTUP              ;start program
0017 BB 0000 R                MOV  BX,OFFSET RECA   ;address record A
001A BF 0000                  MOV  DI,0             ;address element 0
001D 8A 81 0000 R             MOV  AL,FILE[BX+DI]   ;get data
0021 BB 0014 R                MOV  BX,OFFSET RECC   ;address record C
0024 BF 0002                  MOV  DI,2             ;address element 2
0027 88 81 0000 R             MOV  FILE[BX+DI],AL   ;save data
                              .exit                 ;exit to DOS
                              end                   ;end of program
```

TABLE 3–9   Examples of scaled-index addressing.

Assembly Language	Size	Operation
MOV EAX,[EBX+4*ECX]	32 bits	Copies the doubleword contents of the data segment memory location addressed by the sum of 4 times ECX plus EBX into EAX
MOV [EAX+2*EDI+100H],CX	16 bits	Copies CX into the data segment memory location addressed by the sum of EAX, 100H, and 2 times EDI
MOV AL,[EBP+2*EDI+2]	8 bits	Copies the byte contents of the stack segment memory location addressed by the sum of EBP, 2, and 2 times EDI into AL
MOV EAX,ARRAY[4*ECX]	32 bits	Copies the doubleword contents of the data segment memory location addressed by the sum of ARRAY and 4 times ECX into EAX

Scaled-Index Addressing

Scaled-index addressing is the last type of data-addressing mode discussed. This data-addressing mode is unique to the 80386 through the Core2 microprocessors. Scaled-index addressing uses two 32-bit registers (a base register and an index register) to access the memory. The second register (index) is multiplied by a scaling factor. The scaling factor can be $1\times$, $2\times$, $4\times$, or $8\times$. A scaling factor of $1\times$ is implied and need not be included in the assembly language instruction (MOV AL,[EBX+ECX]). A scaling factor of $2\times$ is used to address word-sized memory arrays, a scaling factor of $4\times$ is used with doubleword-sized memory arrays, and a scaling factor of $8\times$ is used with quadword-sized memory arrays.

An example instruction is MOV AX,[EDI+2*ECX]. This instruction uses a scaling factor of $2\times$, which multiplies the contents of ECX by 2 before adding it to the EDI register to form the memory address. If ECX contains a 00000000H, word-sized memory element 0 is addressed; if ECX contains a 00000001H, word-sized memory element 1 is accessed, and so forth. This scales the index (ECX) by a factor of 2 for a word-sized memory array. Refer to Table 3–9 for some examples of scaled-index addressing. As you can imagine, there are an extremely large number of the scaled-index addressed register combinations. Scaling is also applied to instructions that use a single indirect register to access memory. The MOV EAX,[4*EDI] is a scaled-index instruction that uses one register to indirectly address memory. In the 64-bit mode, an instruction such as MOV RAX,[8*RDI] might appear in a program.

Example 3–11 shows a sequence of instructions that uses scaled-index addressing to access a word-sized array of data called LIST. Note that the offset address of LIST is loaded into register EBX with the MOV EBX,OFFSET LIST instruction. Once EBX addresses array LIST, the elements (located in ECX) of 2, 4, and 7 of this word-wide array are added, using a scaling factor of 2 to access the elements. This program stores the 2 at element 2 into elements 4 and 7. Also notice the .386 directive to select the 80386 microprocessor. This directive must follow the .MODEL statement for the assembler to process 80386 instructions for DOS. If the 80486 is in use, the .486 directive appears after the .MODEL statement; if the Pentium is in use, then use .586; and if the Pentium Pro, Pentium II, Pentium III, Pentium 4, or Core2 is in use, then use the .686 directive. If the microprocessor selection directive appears before the .MODEL statement, the microprocessor executes instructions in the 32-bit protected mode, which must execute in Windows.

EXAMPLE 3–11

```
                         .MODEL SMALL           ;select small model
                         .386                   ;select 80386 microprocessor
0000                     .DATA                  ;start data segment
0000 0000 0001 0002 LIST DW  0,1,2,3,4          ;define array LIST
     0003 0004
```

```
000A 0005 0006 0007         DW  5,6,7,8,9
     0008 0009
0000                        .CODE                 ;start code segment
0010 66|BB 00000000 R       MOV EBX,OFFSET LIST    ;address array LIST
0016 66|B9 00000002         MOV ECX,2              ;address element 2
001C 67&8B 04 4B            MOV AX,EBX+2*ECX]      ;get element 2
0020 66|B9 00000004         MOV ECX,4              ;address element 4
0026 67&89 04 4B            MOV [EBX+2*ECX],AX     ;store in element 4
002A 66|B9 00000007         MOV ECX,7              ;address element 7
0030 67&89 04 4B            MOV [EBX+2*ECX],AX     ;store in element 7
                            .exit                  ;exit to DOS
                            end
```

RIP Relative Addressing

This form of addressing uses the 64-bit instruction pointer register in the 64-bit mode to address a linear location in the flat memory model. The inline assembler program available to Visual C++ does not contain any way of using this addressing mode or any other 64-bit addressing mode. The Microsoft Visual C++ does not at present support developing 64-bit assembly code. The instruction pointer is normally addressed using a * as in *+34, which is 34 bytes ahead in a program. When Microsoft finally places an inline assembler into Visual C++ for the 64-bit mode, this most likely will be the way that RIP relative addressing will appear.

One source is Intel, which does produce a compiler with an inline assembler for 64-bit code (http://www.intel.com/cd/software/products/asmo-na/eng/compilers/cwin/279582.htm).

Data Structures

A data structure is used to specify how information is stored in a memory array and can be quite useful with applications that use arrays. It is best to think of a data structure as a template for data. The start of a structure is identified with the STRUC assembly language directive and the end with the ENDS statement. A typical data structure is defined and used three times in Example 3–12. Notice that the name of the structure appears with the STRUC and with ENDS statement. The example shows the data structure as it was typed without the assembled version.

EXAMPLE 3–12

```
;define the INFO data structure
;
INFO       STRUC

NAMES   DB  32 dup(?)   ;reserve 32 bytes for a name
STREET  DB  32 dup(?)   ;reserve 32 bytes for the street address
CITY    DB  16 dup(?)   ;reserve 16 bytes for the city
STATE   DB   2 dup(?)   ;reserve 2 bytes for the state
ZIP     DB   5 dup(?)   ;reserve 5 bytes for the zipcode

INFO       ENDS

NAME1    INFO <'Bob Smith', '123 Main Street', 'Wanda', 'OH', '44444'>
NAME2    INFO <'Steve Doe', '222 Moose Lane', 'Miller', 'PA', '18100'>
NAME3    INFO <'Jim Dover', '303 Main Street', 'Orender', 'CA', '90000'>
```

The data structure in Example 3–12 defines five fields of information. The first is 32 bytes long and holds a name; the second is 32 bytes long and holds a street address; the third is 16 bytes long for the city; the fourth is 2 bytes long for the state; the fifth is 5 bytes long for the ZIP code. Once the structure is defined (INFO), it can be filled, as illustrated, with names and addresses. Three example uses for INFO are illustrated. Note that literals are surrounded with apostrophes and the entire field is surrounded with < > symbols when the data structure is used to define data.

When data are addressed in a structure, use the structure name and the field name to select a field from the structure. For example, to address the STREET in NAME2, use the operand NAME2.STREET, where the name of the structure is first followed by a period and then by the name of the field. Likewise, use NAME3.CITY to refer to the city in structure NAME3.

A short sequence of instructions appears in Example 3-13 that clears the name field in structure NAME1, the address field in structure NAME2, and the ZIP code field in structure NAME3. The function and operation of the instructions in this program are defined in later chapters in the text. You may wish to refer to this example once you learn these instructions.

EXAMPLE 3–13

```
                         ;clear NAMES in array NAME1
                         ;
0000 B9 0020             MOV  CX,32
0003 B0 00               MOV  AL,0
0005 BE 0000 R           MOV  DI,OFFSET NAME1.NAMES
0008 F3/AA               REP  STOSB
                         ;
                         ;clear STREET in array NAME2
                         ;
000A B9 0020             MOV  CX,32
000D B0 00               MOV  AL,0
000F BE 0077 R           MOV  DI,OFFSET NAME2.STREET
0012 F3/AA               REP  STOSB
                         ;
                         ;clear ZIP in NAME3
                         ;
0014 B9 0005             MOV  CX,5
0017 B0 00               MOV  AL,0
0019 BE 0100 R           MOV  DI,OFFSET NAME3.ZIP
001C F3/AA               REP STOSB
```

3–2   PROGRAM MEMORY-ADDRESSING MODES

Program memory-addressing modes, used with the JMP (jump) and CALL instructions, consist of three distinct forms: direct, relative, and indirect. This section introduces these three addressing forms, using the JMP instruction to illustrate their operation.

Direct Program Memory Addressing

Direct program memory addressing is what many early microprocessors used for all jumps and calls. Direct program memory addressing is also used in high-level languages, such as the BASIC language GOTO and GOSUB instructions. The microprocessor uses this form of addressing, but not as often as relative and indirect program memory addressing are used.

The instructions for direct program memory addressing store the address with the opcode. For example, if a program jumps to memory location 10000H for the next instruction, the address (10000H) is stored following the opcode in the memory. Figure 3–14 shows the direct intersegment JMP instruction and the 4 bytes required to store the address 10000H. This JMP instruction loads CS with 1000H and IP with 0000H to jump to memory location 10000H for the next instruction. (An **intersegment jump** is a jump to any memory location within the entire memory system.) The direct jump is often called a *far jump* because it can jump to any memory location for the next

FIGURE 3–14 The 5-byte machine language version of a JMP [10000H] instruction.

Opcode	Offset (low)	Offset (high)	Segment (low)	Segment (high)
E A	0 0	0 0	0 0	1 0

FIGURE 3–15   A JMP [2] instruction. This instruction skips over the 2 bytes of memory that follow the JMP instruction.

```
10000   EB  ⎤
10001   02  ⎦ ─ JMP [2] ┐
10002   —                │
10003   —                │
10004   ◄────────────────┘
```

instruction. In the real mode, a far jump accesses any location within the first 1M byte of memory by changing both CS and IP. In protected mode operation, the far jump accesses a new code segment descriptor from the descriptor table, allowing it to jump to any memory location in the entire 4G-byte address range in the 80386 through Core2 microprocessors.

In the 64-bit mode for the Pentium 4 and Core2, a jump or a call can be to any memory location in the system. The CS segment is still used, but not for the address of the jump or the call. The CS register contains a pointer to a descriptor that describes the access rights and privilege level of the code segment, but not the address of the jump or call.

The only other instruction that uses direct program addressing is the intersegment or far CALL instruction. Usually, the name of a memory address, called a *label*, refers to the location that is called or jumped to instead of the actual numeric address. When using a label with the CALL or JMP instruction, most assemblers select the best form of program addressing.

Relative Program Memory Addressing

Relative program memory addressing is not available in all early microprocessors, but it is available to this family of microprocessors. The term *relative* means "relative to the instruction pointer (IP)." For example, if a JMP instruction skips the next 2 bytes of memory, the address in relation to the instruction pointer is a 2 that adds to the instruction pointer. This develops the address of the next program instruction. An example of the relative JMP instruction is shown in Figure 3–15. Notice that the JMP instruction is a 1-byte instruction, with a 1-byte or a 2-byte displacement that adds to the instruction pointer. A 1-byte displacement is used in **short** jumps, and a 2-byte displacement is used with **near** jumps and calls. Both types are considered to be intrasegment jumps. (An **intrasegment jump** is a jump anywhere within the current code segment.) In the 80386 and above, the displacement can also be a 32-bit value, allowing them to use relative addressing to any location within their 4G-byte code segments.

Relative JMP and CALL instructions contain either an 8-bit or a 16-bit signed displacement that allows a forward memory reference or a reverse memory reference. (The 80386 and above can have an 8-bit or 32-bit displacement.) All assemblers automatically calculate the distance for the displacement and select the proper 1-, 2- or 4-byte form. If the distance is too far for a 2-byte displacement in an 8086 through an 80286 microprocessor, some assemblers use the direct jump. An 8-bit displacement *(short)* has a jump range of between +127 and −128 bytes from the next instruction; a 16-bit displacement *(near)* has a range of ±32K bytes. In the 80386 and above, a 32-bit displacement allows a range of ±2G bytes. The 32-bit displacement can only be used in the protected mode.

Indirect Program Memory Addressing

The microprocessor allows several forms of program indirect memory addressing for the JMP and CALL instructions. Table 3–10 lists some acceptable program indirect jump instructions, which can use any 16-bit register (AX, BX, CX, DX, SP, BP, DI, or SI); any relative register ([BP], [BX], [DI], or [SI]); and any relative register with a displacement. In the 80386 and above, an extended register can also be used to hold the address or indirect address of a relative JMP or CALL. For example, the JMP EAX jumps to the location address by register EAX.

If a 16-bit register holds the address of a JMP instruction, the jump is near. For example, if the BX register contains 1000H and a JMP BX instruction executes, the microprocessor jumps to offset address 1000H in the current code segment.

TABLE 3–10 Examples of indirect program memory addressing.

Assembly Language	Operation
JMP AX	Jumps to the current code segment location addressed by the contents of AX
JMP CX	Jumps to the current code segment location addressed by the contents of CX
JMP NEAR PTR[BX]	Jumps to the current code segment location addressed by the contents of the data segment location addressed by BX
JMP NEAR PTR[DI+2]	Jumps to the current code segment location addressed by the contents of the data segment memory location addressed by DI plus 2
JMP TABLE[BX]	Jumps to the current code segment location addressed by the contents of the data segment memory location address by TABLE plus BX
JMP ECX	Jumps to the current code segment location addressed by the contents of ECX
JMP RDI	Jumps to the linear address contained in the RDI register (64-bit mode)

FIGURE 3–16 A jump table that stores addresses of various programs. The exact address chosen from the TABLE is determined by an index stored with the jump instruction.

```
TABLE   DW   LOC0
        DW   LOC1
        DW   LOC2
        DW   LOC3
```

If a relative register holds the address, the jump is also considered to be an indirect jump. For example, JMP [BX] refers to the memory location within the data segment at the offset address contained in BX. At this offset address is a l6-bit number that is used as the offset address in the intrasegment jump. This type of jump is sometimes called an *indirect-indirect* or *double-indirect jump*.

Figure 3–16 shows a jump table that is stored, beginning at memory location TABLE. This jump table is referenced by the short program of Example 3–14. In this example, the BX register is loaded with a 4 so, when it combines in the JMP TABLE[BX] instruction with TABLE, the effective address is the contents of the second entry in the 16-bit-wide jump table.

EXAMPLE 3–14
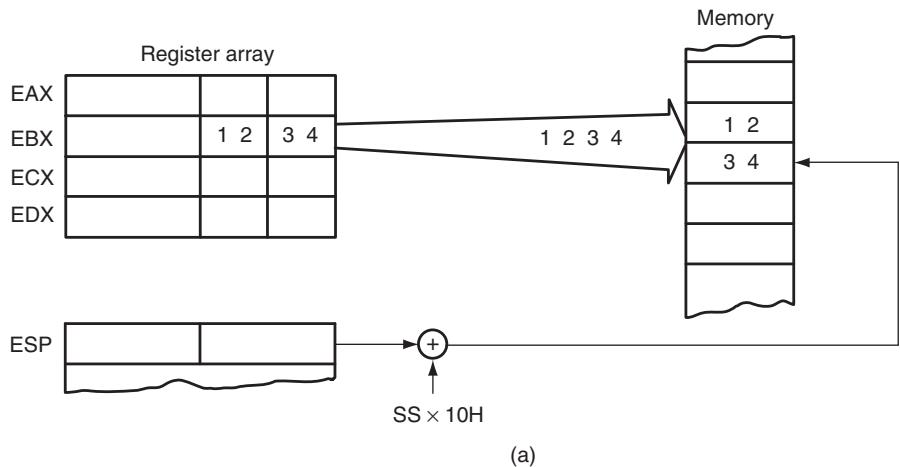
```
                    ;Using indirect addressing for a jump
                    ;
0000 BB 0004        MOV BX,4         ;address LOC2
0003 FF A7 23A1 R   JMP TABLE[BX]    ;jump to LOC2
```
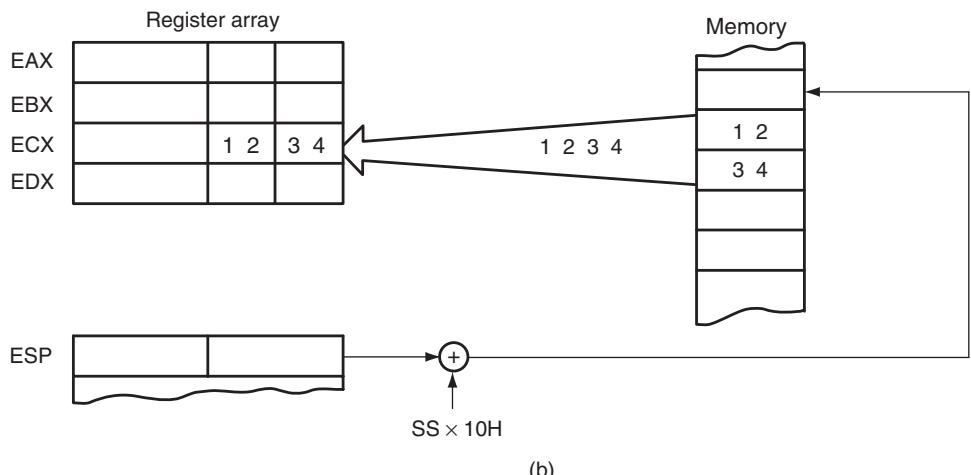
3–3 STACK MEMORY-ADDRESSING MODES

The stack plays an important role in all microprocessors. It holds data temporarily and stores the return addresses used by procedures. The stack memory is an LIFO (**last-in, first-out**) memory, which describes the way that data are stored and removed from the stack. Data are placed onto the stack with a **PUSH instruction** and removed with a **POP instruction**. The CALL instruction also uses the stack to hold the return address for procedures and a RET (return) instruction to remove the return address from the stack.

The stack memory is maintained by two registers: the stack pointer (SP or ESP) and the stack segment register (SS). Whenever a word of data is pushed onto the stack [see Figure 3–17(a)], the high-order 8 bits are placed in the location addressed by SP – 1. The low-order 8 bits are placed in the location addressed by SP – 2. The SP is then decremented by 2 so that the next word of data

FIGURE 3–17 The PUSH and POP instructions: (a) PUSH BX places the contents of BX onto the stack; (b) POP CX removes data from the stack and places them into CX. Both instructions are shown after execution.

is stored in the next available stack memory location. The SP/ESP register always points to an area of memory located within the stack segment. The SP/ESP register adds to SS × 10H to form the stack memory address in the real mode. In protected mode operation, the SS register holds a selector that accesses a descriptor for the base address of the stack segment.

Whenever data are popped from the stack [see Figure 3–17(b)], the low-order 8 bits are removed from the location addressed by SP. The high-order 8 bits are removed from the location addressed by SP + 1. The SP register is then incremented by 2. Table 3–11 lists some of the PUSH and POP instructions available to the microprocessor. Note that PUSH and POP store or retrieve words of data—never bytes—in the 8086 through the 80286 microprocessors. The 80386 and above allow words or doublewords to be transferred to and from the stack. Data may be pushed onto the stack from any 16-bit register or segment register; in the 80386 and above, from any 32-bit extended register. Data may be popped off the stack into any register or any segment register except CS. The reason that data may not be popped from the stack into CS is that this only

TABLE 3–11   Example PUSH and POP instructions.

Assembly Language	Operation
POPF	Removes a word from the stack and places it into the flag register
POPFD	Removes a doubleword from the stack and places it into the EFLAG register
PUSHF	Copies the flag register to the stack
PUSHFD	Copies the EFLAG register to the stack
PUSH AX	Copies the AX register to the stack
POP BX	Removes a word from the stack and places it into the BX register
PUSH DS	Copies the DS register to the stack
PUSH 1234H	Copies a word-sized 1234H to the stack
POP CS	This instruction is illegal
PUSH WORD PTR[BX]	Copies the word contents of the data segment memory location addressed by BX onto the stack
PUSHA	Copies AX, CX, DX, BX, SP, BP, DI, and SI to the stack
POPA	Removes the word contents for the following registers from the stack: SI, DI, BP, SP, BX, DX, CX, and AX
PUSHAD	Copies EAX, ECX, EDX, EBX, ESP, EBP, EDI, and ESI to the stack
POPAD	Removes the doubleword contents for the following registers from the stack: ESI, EDI, EBP, ESP, EBX, EDX, ECX, and EAX
POP EAX	Removes a doubleword from the stack and places it into the EAX register
POP RAX	Removes a quadword from the stack and places it into the RAC register (64-bit mode)
PUSH EDI	Copies EDI to the stack
PUSH RSI	Copies RSI into the stack (64-bit mode)
PUSH QWORD PTR[RDX]	Copies the quadword contents of the memory location addressed by RDX onto the stack

changes part of the address of the next instruction. In the Pentium 4 or Core2 operated in 64-bit mode, the 64-bit registers can be pushed or popped from the stack, but they are 8 bytes in length.

The PUSHA and POPA instructions either push or pop all of the registers, except segment registers, onto the stack. These instructions are not available on the early 8086/8088 processors. The push immediate instruction is also new to the 80286 through the Core2 microprocessors. Note the examples in Table 3–11, which show the order of the registers transferred by the PUSHA and POPA instructions. The 80386 and above also allow extended registers to be pushed or popped. The 64-bit mode for the Pentium 4 and Core2 does not contain a PUSHA or POPA instruction.

Example 3–15 lists a short program that pushes the contents of AX, BX, and CX onto the stack. The first POP retrieves the value that was pushed onto the stack from CX and places it into AX. The second POP places the original value of BX into CX. The last POP places the value of AX into BX.

EXAMPLE 3–15

```
                    .MODEL TINY        ;select tiny model
0000                .CODE              ;start code segment
                    .STARTUP           ;start program
0100 B8 1000        MOV   AX,1000H      ;load test data
0103 BB 2000        MOV   BX,2000H
```

```
0106 B9 3000            MOV   CX,3000H

0109 50                 PUSH  AX              ;1000H to stack
010A 53                 PUSH  BX              ;2000H to stack
010B 51                 PUSH  CX              ;3000H to stack

010C 58                 POP   AX              ;3000H to AX
010D 59                 POP   CX              ;2000H to CBX
010E 5B                 POP   BX              ;1000H to BX
                        .exit                 ;exit to DOS
                        end                   ;end program
```

3–4   SUMMARY

1. The data-addressing modes include register, immediate, direct, register indirect, base-plus-index, register relative, and base relative-plus-index addressing. The 80386 through the Pentium 4 microprocessors have an additional addressing mode called scaled-index addressing.

2. The program memory-addressing modes include direct, relative, and indirect addressing.

3. Table 3–12 lists all real mode data-addressing modes available to the 8086 through the 80286 microprocessors. Note that the 80386 and above use these modes, plus the many defined through this chapter. In the protected mode, the function of the segment register is to address a descriptor that contains the base address of the memory segment.

4. The 80386 through Core2 microprocessors have additional addressing modes that allow the extended registers EAX, EBX, ECX, EDX, EBP, EDI, and ESI to address memory. Although these addressing modes are too numerous to list in tabular form, in general, any of these registers function in the same way as those listed in Table 3–12. For example, the MOV AL,TABLE[EBX+2*ECX+10H] is a valid addressing mode for the 80386–Core2 microprocessors.

5. The 64-bit mode for the Pentium 4 and Core2 microprocessors use the same addressing modes as the Pentium 4 or Core2 in 32-bit mode, except the registers contain a linear address and they are 64 bits in width. An additional addressing mode called RIP relative exists for the 64-bit mode that addresses data relative to the address in the instruction pointer.

6. The MOV instruction copies the contents of the source operand into the destination operand. The source never changes for any instruction.

7. Register addressing specifies any 8-bit register (AH, AL, BH, BL, CH, CL, DH, or DL) or any 16-bit register (AX, BX, CX, DX, SP, BP, SI, or DI). The segment registers (CS, DS, ES, or SS) are also addressable for moving data between a segment register and a 16-bit register/memory location or for PUSH and POP. In the 80386 through the Core2 microprocessors, the extended registers also are used for register addressing; they consist of EAX, EBX, ECX, EDX, ESP, EBP, EDI, and ESI. Also available to the 80386 and above are the FS and GS segment registers. In the 64-bit mode, the registers are RAX, RBX, RCX, RDX, RSP, RBP, RDI, RSI, and R8 through R15.

8. The MOV immediate instruction transfers the byte or word that immediately follows the opcode into a register or a memory location. Immediate addressing manipulates constant data in a program. In the 80386 and above, doubleword immediate data may also be loaded into a 32-bit register or memory location.

9. The .MODEL statement is used with assembly language to identify the start of a file and the type of memory model used with the file. If the size is TINY, the program exists in one segment, the code segment, and is assembled as a command (.COM) program. If the SMALL

TABLE 3–12 Example real mode data-addressing modes.

Assembly Language	Address Generation
MOV AL,BL	8-bit register addressing
MOV AX,BX	16-bit register addressing
MOV EAX,ECX	32-bit register addressing
MOV DS,DX	Segment register addressing
MOV AL,LIST	(DS x 10H) + LIST
MOV CH,DATA1	(DS x 10H) + DATA1
MOV ES,DATA2	(DS x 10H) + DATA2
MOV AL,12	Immediate data of 12
MOV AL,[BP]	(SS x 10H) + BP
MOV AL,[BX]	(DS x 10H) + BX
MOV AL,[DI]	(DS x 10H) + DI
MOV AL,[SI]	(DS x 10H) + SI
MOV AL,[BP+2]	(SS x 10H) + BP + 2
MOV AL,[BX–4]	(DS x 10H) + BX − 4
MOV AL,[DI+1000H]	(DS x 10H) + DI + 1000H
MOV AL,[SI+300H]	(DS x 10H) + SI + 300H
MOV AL,LIST[BP]	(SS x 10H) + LIST + BP
MOV AL,LIST[BX]	(DS x 10H) + LIST + BX
MOV AL,LIST[DI]	(DS x 10H) + LIST + DI
MOV AL,LIST[SI]	(DS x 10H) + LIST + SI
MOV AL,LIST[BP+2]	(SS x 10H) + LIST + BP + 2
MOV AL,LIST[BX–6]	(DS x 10H) + LIST + BX − 6
MOV AL,LIST[DI+100H]	(DS x 10H) + LIST + DI + 100H
MOV AL,LIST[SI+200H]	(DS x 10H) + LIST + SI + 200H
MOV AL,[BP+DI]	(SS x 10H) + BP + DI
MOV AL,[BP+SI]	(SS x 10H) + BP + SI
MOV AL,[BX+DI]	(DS x 10H) + BX + DI
MOV AL,[BX+SI]	(DS x 10H) + BX + SI
MOV AL,[BP+DI+8]	(SS x 10H) + BP + DI + 8
MOV AL,[BP+SI–8]	(SS x 10H) + BP + SI − 8
MOV AL,[BX+DI+10H]	(DS x 10H) + BX + DI + 10H
MOV AL,[BX+SI–10H]	(DS x 10H) + BX + SI − 10H
MOV AL,LIST[BP+DI]	(SS x 10H) + LIST + BP + DI
MOV AL,LIST[BP+SI]	(SS x 10H) + LIST + BP + SI
MOV AL,LIST[BX+DI]	(DS x 10H) + LIST + BX + DI
MOV AL,LIST[BX+SI]	(DS x 10H) + LIST + BX + SI
MOV AL,LIST[BP+DI+2]	(SS x 10H) + LIST + BP + DI + 2
MOV AL,LIST[BP+SI–7]	(SS x 10H) + LIST + BP + SI − 7
MOV AL,LIST[BX+DI+3]	(DS x 10H) + LIST + BX + DI + 3
MOV AL,LIST[BX+SI–2]	(DS x 10H) + LIST + BX + SI − 2

model is used, the program uses a code and data segment and assembles as an execute (.EXE) program. Other model sizes and their attributes are listed in Appendix A.

10. Direct addressing occurs in two forms in the microprocessor: (1) direct addressing and (2) displacement addressing. Both forms of addressing are identical except that direct addressing is used to transfer data between EAX, AX, or AL and memory; displacement addressing is used with any register-memory transfer. Direct addressing requires 3 bytes of memory, whereas displacement addressing requires 4 bytes. Note that some of these instructions in the 80386 and above may require additional bytes in the form of prefixes for register and operand sizes.

11. Register indirect addressing allows data to be addressed at the memory location pointed to by either a base (BP and BX) or index register (DI and SI). In the 80386 and above, extended registers EAX, EBX, ECX, EDX, EBP, EDI, and ESI are used to address memory data.

12. Base-plus-index addressing often addresses data in an array. The memory address for this mode is formed by adding a base register, index register, and the contents of a segment register times 10H. In the 80386 and above, the base and index registers may be any 32-bit register except EIP and ESP.

13. Register relative addressing uses a base or index register, plus a displacement to access memory data.

14. Base relative-plus-index addressing is useful for addressing a two-dimensional memory array. The address is formed by adding a base register, an index register, displacement, and the contents of a segment register times 10H.

15. Scaled-index addressing is unique to the 80386 through the Core2. The second of two registers (index) is scaled by a factor of 2×, 4×, or 8× to access words, doublewords, or quadwords in memory arrays. The MOV AX,[EBX+2*ECX] and the MOV [4*ECX],EDX are examples of scaled-index instructions.

16. Data structures are templates for storing arrays of data and are addressed by array name and field. For example, array NUMBER and field TEN of array NUMBER is addressed as NUMBER.TEN.

17. Direct program memory addressing is allowed with the JMP and CALL instructions to any location in the memory system. With this addressing mode, the offset address and segment address are stored with the instruction.

18. Relative program addressing allows a JMP or CALL instruction to branch forward or backward in the current code segment by ±32K bytes. In the 80386 and above, the 32-bit displacement allows a branch to any location in the current code segment by using a displacement value of ±2G bytes. The 32-bit displacement can be used only in protected mode.

19. Indirect program addressing allows the JMP or CALL instructions to address another portion of the program or subroutine indirectly through a register or memory location.

20. The PUSH and POP instructions transfer a word between the stack and a register or memory location. A PUSH immediate instruction is available to place immediate data on the stack. The PUSHA and POPA instructions transfer AX, CX, DX, BX, BP, SP, SI, and DI between the stack and these registers. In the 80386 and above, the extended register and extended flags can also be transferred between registers and the stack. A PUSHFD stores the EFLAGS, whereas a PUSHF stores the FLAGS. POPA and PUSHA are not available in the 64-bit mode.

3–5   QUESTIONS AND PROBLEMS

1. What do the following MOV instructions accomplish?
   (a) MOV AX,BX
   (b) MOV BX,AX
   (c) MOV BL,CH
   (d) MOV ESP,EBP
   (e) MOV RAX,RCX
2. List the 8-bit registers that are used for register addressing.

3. List the 16-bit registers that are used for register addressing.
4. List the 32-bit registers that are used for register addressing in the 80386 through the Core2 microprocessors.
5. List the 64-bit registers available to the 64-bit mode of the Pentium 4 and Core2.
6. List the 16-bit segment registers used with register addressing by MOV, PUSH, and POP.
7. What is wrong with the MOV BL,CX instruction?
8. What is wrong with the MOV DS,SS instruction?
9. Select an instruction for each of the following tasks:
    (a) copy EBX into EDX
    (b) copy BL into CL
    (c) copy SI into BX
    (d) copy DS into AX
    (e) copy AL into AH
    (f) copy R8 into R10
10. Select an instruction for each of the following tasks:
    (a) move 12H into AL
    (b) move 123AH into AX
    (c) move 0CDH into CL
    (d) move 1000H into RAX
    (e) move 1200A2H into EBX
11. What special symbol is sometimes used to denote immediate data?
12. What is the purpose of the .MODEL TINY statement?
13. What assembly language directive indicates the start of the CODE segment?
14. What is a label?
15. The MOV instruction is placed in what field of a statement?
16. A label may begin with what characters?
17. What is the purpose of the .EXIT directive?
18. Does the .MODEL TINY statement cause a program to assemble as an execute (.EXE) program?
19. What tasks does the .STARTUP directive accomplish in the small memory model?
20. What is a displacement? How does it determine the memory address in a MOV DS:[2000H],AL instruction?
21. What do the symbols [ ] indicate?
22. Suppose that DS = 0200H, BX = 0300H, and DI = 400H. Determine the memory address accessed by each of the following instructions, assuming real mode operation:
    (a) MOV AL,[1234H]
    (b) MOV EAX,[BX]
    (c) MOV [DI],AL
23. What is wrong with a MOV [BX],[DI] instruction?
24. Choose an instruction that requires BYTE PTR.
25. Choose an instruction that requires WORD PTR.
26. Choose an instruction that requires DWORD PTR.
27. Select an instruction that requires QWORD PTR.
28. Explain the difference between the MOV BX,DATA instruction and the MOV BX,OFFSET DATA instruction.
29. Suppose that DS = 1000H, SS = 2000H, BP = 1000H, and DI = 0100H. Determine the memory address accessed by each of the following instructions, assuming real mode operation:
    (a) MOV AL,[BP+DI]

    (b) MOV CX,[DI]

    (c) MOV EDX,[BP]

30. What, if anything, is wrong with a MOV AL,[BX][SI] instruction?

31. Suppose that DS = 1200H, BX = 0100H, and SI = 0250H. Determine the address accessed by each of the following instructions, assuming real mode operation:

    (a) MOV [100H],DL

    (b) MOV [SI+100H],EAX

    (c) MOV DL,[BX+100H]

32. Suppose that DS = 1100H, BX = 0200H, LIST = 0250H, and SI = 0500H. Determine the address accessed by each of the following instructions, assuming real mode operation:

    (a) MOV LIST[SI],EDX

    (b) MOV CL,LIST[BX+SI]

    (c) MOV CH,[BX+SI]

33. Suppose that DS = 1300H, SS = 1400H, BP = 1500H, and SI = 0100H. Determine the address accessed by each of the following instructions, assuming real mode operation:

    (a) MOV EAX,[BP+200H]

    (b) MOV AL,[BP+SI−200H]

    (c) MOV AL,[SI−0100H]

34. Which base register addresses data in the stack segment?

35. Suppose that EAX = 00001000H, EBX = 00002000H, and DS = 0010H. Determine the addresses accessed by the following instructions, assuming real mode operation:

    (a) MOV ECX,[EAX+EBX]

    (b) MOV [EAX+2*EBX],CL

    (c) MOV DH,[EBX+4*EAX+1000H]

36. Develop a data structure that has five fields of one word each named Fl, F2, F3, F4, and F5 with a structure name of FIELDS.

37. Show how field F3 of the data structure constructed in question 36 is addressed in a program.

38. What are the three program memory-addressing modes?

39. How many bytes of memory store a far direct jump instruction? What is stored in each of the bytes?

40. What is the difference between an intersegment and intrasegment jump?

41. If a near jump uses a signed 16-bit displacement, how can it jump to any memory location within the current code segment?

42. The 80386 and above use a _____-bit displacement to jump to any location within the 4G-byte code segment.

43. What is a far jump?

44. If a JMP instruction is stored at memory location 100H within the current code segment, it cannot be a _____ jump if it is jumping to memory location 200H within the current code segment.

45. Show which JMP instruction assembles (short, near, or far) if the JMP THERE instruction is stored at memory address 10000H and the address of THERE is:

    (a) l0020H

    (b) 11000H

    (c) 0FFFEH

    (d) 30000H

46. Form a JMP instruction that jumps to the address pointed to by the BX register.

47. Select a JMP instruction that jumps to the location stored in memory at the location TABLE. Assume that it is a near JMP.

48. How many bytes are stored on the stack by a PUSH AX?

49. Explain how the PUSH [DI] instruction functions.
50. What registers are placed on the stack by the PUSHA instruction? In what order?
51. What does the PUSHAD instruction accomplish?
52. Which instruction places the EFLAGS on the stack in the Pentium 4 microprocessor?
53. Is a PUSHA available in the 64-bit mode of the Pentium 4 or the Core2?

CHAPTER 4

Data Movement Instructions

INTRODUCTION

This chapter concentrates on the data movement instructions. The data movement instructions include MOV, MOVSX, MOVZX, PUSH, POP, BSWAP, XCHG, XLAT, IN, OUT, LEA, LDS, LES, LFS, LGS, LSS, LAHF, SAHF, and the string instructions MOVS, LODS, STOS, INS, and OUTS. The latest data transfer instruction implemented on the Pentium Pro and above is the CMOV (conditional move) instruction. The data movement instructions are presented first because they are more commonly used in programs and are easy to understand.

The microprocessor requires an assembler program, which generates machine language, because machine language instructions are too complex to efficiently generate by hand. This chapter describes the assembly language syntax and some of its directives. [This text assumes that the user is developing software on an IBM personal computer or clone. It is recommended that the Microsoft MACRO assembler (MASM) be used as the development tool, but the Intel Assembler (ASM), Borland Turbo assembler (TASM), or similar software function equally as well. The most recent version of TASM completely emulates the MASM program. This text presents information that functions with the Microsoft MASM assembler, but most programs assemble without modification with other assemblers. Appendix A explains the Microsoft assembler and provides detail on the linker program.] As a more modern alternative, the Visual C++ Express compiler and its inline assembler program may also be used as a development system. Both are explained in detail in the text.

CHAPTER OBJECTIVES

Upon completion of this chapter, you will be able to:

1. Explain the operation of each data movement instruction with applicable addressing modes.
2. Explain the purposes of the assembly language pseudo-operations and key words such as ALIGN, ASSUME, DB, DD, DW, END, ENDS, ENDP, EQU, .MODEL, OFFSET, ORG, PROC, PTR, SEGMENT, USEI6, USE32, and USES.
3. Select the appropriate assembly language instruction to accomplish a specific data movement task.
4. Determine the symbolic opcode, source, destination, and addressing mode for a hexadecimal machine language instruction.
5. Use the assembler to set up a data segment, stack segment, and code segment.

6. Show how to set up a procedure using PROC and ENDP.
7. Explain the difference between memory models and full-segment definitions for the MASM assembler.
8. Use the Visual C++ online assembler to perform data movement tasks.

4–1   MOV REVISITED

The MOV instruction, introduced in Chapter 3, explains the diversity of 8086–Core2 addressing modes. In this chapter, the MOV instruction introduces the machine language instructions available with various addressing modes and instructions. Machine code is introduced because it may occasionally be necessary to interpret machine language programs generated by an assembler or inline assembler of Visual C++. Interpretation of the machine's native language (**machine language**) allows debugging or modification at the machine language level. Occasionally, machine language patches are made by using the DEBUG program available with DOS and also in Visual C++ for Windows, which requires some knowledge of machine language. Conversion between machine and assembly language instructions is illustrated in Appendix B.

Machine Language

Machine language is the native binary code that the microprocessor understands and uses as its instructions to control its operation. Machine language instructions for the 8086 through the Core2 vary in length from 1 to as many as 13 bytes. Although machine language appears complex, there is order to this microprocessor's machine language. There are well over 100,000 variations of machine language instructions, meaning that there is no complete list of these variations. Because of this, some binary bits in a machine language instruction are given, and the remaining bits are determined for each variation of the instruction.

Instructions for the 8086 through the 80286 are 16-bit mode instructions that take the form found in Figure 4–1(a). The 16-bit mode instructions are compatible with the 80386 and above if they are programmed to operate in the 16-bit instruction mode, but they may be prefixed, as shown in Figure 4–1(b). The 80386 and above assume that all instructions are 16-bit mode instructions when the machine is operated in the *real mode (DOS)*. In the *protected mode (Windows)*, the upper byte of the descriptor contains the D-bit that selects either the 16- or 32-bit instruction mode. At present, only Windows 95 through Windows XP and Linux operate in the 32-bit instruction mode. The 32-bit mode instructions are in the form shown in Figure 4–1(b).
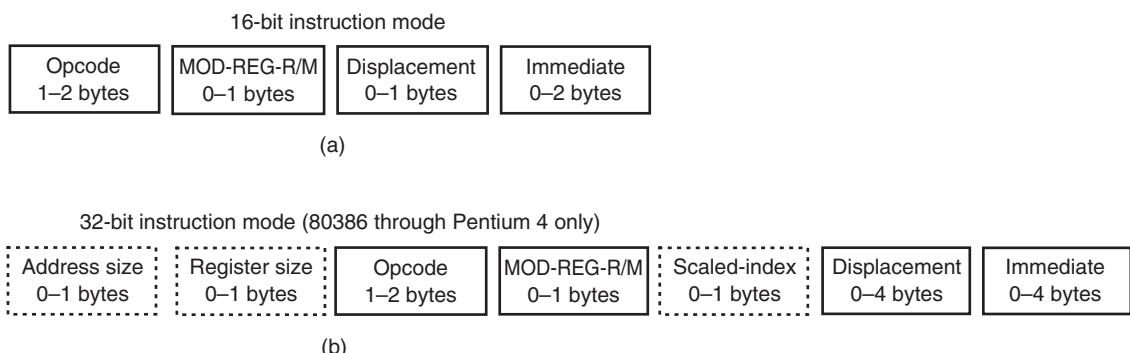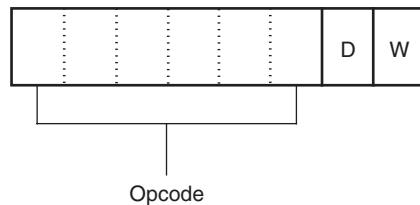
FIGURE 4–1   The formats of the 8086–Core2 instructions. (a) The 16-bit form and (b) the 32-bit form.

FIGURE 4–2 Byte 1 of many machine language instructions, showing the position of the D- and W-bits.

Opcode

These instructions occur in the 16-bit instruction mode by the use of prefixes, which are explained later in this chapter.

The first 2 bytes of the 32-bit instruction mode format are called override prefixes because they are not always present. The first modifies the size of the operand address used by the instruction and the second modifies the register size. If the 80386 through the Pentium 4 operate as 16-bit instruction mode machines (real or protected mode) and a 32-bit register is used, the **register-size prefix** (66H) is appended to the front of the instruction. If operated in the 32-bit instruction mode (protected mode only) and a 32-bit register is used, the register-size prefix is absent. If a 16-bit register appears in an instruction in the 32-bit instruction mode, the register-size 16-bit instruction mode, the register-size prefix is present to select a l6-bit register. The **address size-prefix** (67H) is used in a similar fashion, as explained later in this chapter. The prefixes toggle the size of the register and operand address from l6-bit to 32-bit or from 32-bit to l6-bit for the prefixed instruction. Note that the l6-bit instruction mode uses 8- and l6-bit registers and addressing modes, while the 32-bit instruction mode uses 8- and 32-bit registers and addressing modes by default. The prefixes override these defaults so that a 32-bit register can be used in the l6-bit mode or a l6-bit register can be used in the 32-bit mode. The **mode of operation** (16 or 32 bits) should be selected to function with the current application. If 8- and 32-bit data pervade the application, the 32-bit mode should be selected; likewise, if 8- and l6-bit data pervade, the l6-bit mode should be selected. Normally, mode selection is a function of the operating system. (Remember that DOS can operate only in the l6-bit mode, where Windows can operate in both modes.)

The Opcode.   The **opcode** selects the operation (addition, subtraction, move, and so on) that is performed by the microprocessor. The opcode is either 1 or 2 bytes long for most machine language instructions. Figure 4–2 illustrates the general form of the first opcode byte of many, but *not* all, machine language instructions. Here, the first 6 bits of the first byte are the binary opcode. The remaining 2 bits indicate the **direction** (D)—not to be confused with the instruction mode bit (16/32) or direction flag bit (used with string instructions)—of the data flow, and indicate whether the data are a byte or a word (W). In the 80386 and above, words and doublewords are both specified when W = 1. The instruction mode and register-size prefix (66H) determine whether W represents a word or a doubleword.

If the direction bit (D) = 1, data flow *to* the register REG field from the R/M field located in the second byte of an instruction. If the D-bit = 0 in the opcode, data flow to the R/M field *from* the REG field. If the W-bit = 1, the data size is a *word* or *doubleword;* if the W-bit = 0, the data size is always a *byte.* The W-bit appears in most instructions, while the D-bit appears mainly with the MOV and some other instructions. Refer to Figure 4–3 for the binary bit pattern of the second opcode byte (reg-mod-r/m) of many instructions. Figure 4–3 shows the location of the MOD (mode), REG (register), and R/M (register/memory) fields.

FIGURE 4–3 Byte 2 of many machine language instructions, showing the position of the MOD, REG, and R/M fields.

MOD   REG   R/M
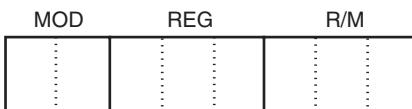
TABLE 4–1 MOD field for the 16-bit instruction mode.

MOD	Function
00	No displacement
01	8-bit sign-extended displacement
10	16-bit signed displacement
11	R/M is a register

MOD Field.   The MOD field specifies the addressing mode (MOD) for the selected instruction. The MOD field selects the type of addressing and whether a displacement is present with the selected type. Table 4–1 lists the operand forms available to the MOD field for l6-bit instruction mode, unless the operand address-size override prefix (67H) appears. If the MOD field contains an 11, it selects the register-addressing mode. Register addressing uses the R/M field to specify a register instead of a memory location. If the MOD field contains a 00, 01, or 10, the R/M field selects one of the data memory-addressing modes. When MOD selects a data memory addressing mode, it indicates that the addressing mode contains no displacement (00), an 8-bit sign-extended displacement (01), or a l6-bit displacement (10). The MOV AL,[DI] instruction is an example that contains no displacement, a MOV AL,[DI+2] instruction uses an 8-bit displacement (+2), and a MOV AL,[DI+1000H] instruction uses a 16-bit displacement (+1000H).

All 8-bit displacements are sign-extended into 16-bit displacements when the microprocessor executes the instruction. If the 8-bit displacement is 00H–7FH (positive), it is sign-extended to 0000H–007FH before adding to the offset address. If the 8-bit displacement is 80H–FFH (negative), it is sign-extended to FF80H–FFFFH. To sign-extend a number, its sign-bit is copied to the next higher-order byte, which generates either a 00H or an FFH in the next higher-order byte. Some assembler programs do not use the 8-bit displacements and in place default to all 16-bit displacements.

In the 80386 through the Core2 microprocessors, the MOD field may be the same as shown in Table 4–1 for 16-bit instruction mode; if the instruction mode is 32 bits, the MOD field is as it appears in Table 4–2. The MOD field is interpreted as selected by the address-size override prefix or the operating mode of the microprocessor. This change in the interpretation of the MOD field and instruction supports many of the numerous additional addressing modes allowed in the 80386 through the Core2. The main difference is that when the MOD field is a 10, this causes the 16-bit displacement to become a 32-bit displacement, to allow any protected mode memory location (4G bytes) to be accessed. The 80386 and above only allow an 8- or 32-bit displacement when operated in the 32-bit instruction mode, unless the address-size override prefix appears. Note that if an 8-bit displacement is selected, it is sign-extended into a 32-bit displacement by the microprocessor.

Register Assignments.   Table 4–3 lists the register assignments for the REG field and the R/M field (MOD = 11). This table contains three lists of register assignments: one is used when the W bit = 0 (bytes), and the other two are used when the W bit = 1 (words or doublewords). Note that doubleword registers are only available to the 80386 through the Core2.

TABLE 4–2 MOD field for the 32-bit instruction mode (80386–Core2 only).

MOD	Function
00	No displacement
01	8-bit sign-extended displacement
10	32-bit signed displacement
11	R/M is a register

TABLE 4–3 REG and R/M (when) MOD = 11 assignments.

Code	W = 0 (Byte)	W = 1 (Word)	W = 1 (Doubleword)
000	AL	AX	EAX
001	CL	CX	ECX
010	DL	DX	EDX
011	BL	BX	EBX
100	AH	SP	ESP
101	CH	BP	EBP
110	DH	SI	ESI
111	BH	DI	EDI

Suppose that a 2-byte instruction, 8BECH, appears in a machine language program. Because neither a 67H (operand address-size override prefix) nor a 66H (register-size override prefix) appears as the first byte, the first byte is the opcode. If the microprocessor is operated in the 16-bit instruction mode, this instruction is converted to binary and placed in the instruction format of bytes 1 and 2, as illustrated in Figure 4–4. The opcode is 100010. If you refer to Appendix B, which lists the machine language instructions, you will find that this is the opcode for a MOV instruction. Notice that both the D and W bits are a logic 1, which means that a word moves into the destination register specified in the REG field. The REG field contains a 101, indicating register BP, so the MOV instruction moves data into register BP. Because the MOD field contains a 11, the R/M field also indicates a register. Here, R/M = 100 (SP); therefore, this instruction moves data from SP into BP and is written in symbolic form as a MOV BP,SP instruction.

Suppose that a 668BE8H instruction appears in an 80386 or above, operated in the 16-bit instruction mode. The first byte (66H) is the register-size override prefix that selects 32-bit register operands for the 16-bit instruction mode. The remainder of the instruction indicates that the opcode is a MOV with a source operand of EAX and a destination operand of EBP. This instruction is a MOV EBP,EAX. The same instruction becomes a MOV BP,AX instruction in the 80386 and above if it is operated in the 32-bit instruction mode, because the register-size override prefix selects a 16-bit register. Luckily, the assembler program keeps track of the register- and address-size prefixes and the mode of operation. Recall that if the .386 switch is placed before the .MODEL statement, the 32-bit mode is selected; if it is placed after the .MODEL statement, the 16-bit mode is selected. All programs written using the inline assembler in Visual C++ are always in the 32-bit mode.

R/M Memory Addressing. If the MOD field contains a 00, 01, or 10, the R/M field takes on a new meaning. Table 4–4 lists the memory-addressing modes for the R/M field when MOD is a 00, 01, or 10 for the 16-bit instruction mode.

```
Opcode       D  W        MOD   REG   R/M
1 0 0 0 1 0  1  1        1 1   1 0 1  1 0 0
```

Opcode = MOV
D = Transfer to register (REG)
W = Word
MOD = R/M is a register
REG = BP
R/M = SP

FIGURE 4–4 The 8BEC instruction placed into bytes 1 and 2 formats from Figures 4–2 and 4–3. This instruction is a MOV BP,SP.

TABLE 4–4 16-bit R/M memory-addressing modes.

R/M Code	Addressing Mode
000	DS:[BX+SI]
001	DS:[BX+DI]
010	SS:[BP+SI]
011	SS:[BP+DI]
100	DS:[SI]
101	DS:[DI]
110	SS:[BP]*
111	DS:[BX]

*Note: See text section, Special Addressing Mode.

All of the 16-bit addressing modes presented in Chapter 3 appear in Table 4–4. The displacement, discussed in Chapter 3, is defined by the MOD field. If MOD = 00 and R/M = 101, the addressing mode is [DI]. If MOD = 01 or 10, the addressing mode is [DI + 33H], or LIST [DI + 22H] for the 16-bit instruction mode. This example uses LIST, 33H, and 22H as arbitrary values for the displacement.

Figure 4–5 illustrates the machine language version of the 16-bit instruction MOV DL,[DI] or instruction (8AI5H). This instruction is 2 bytes long and has an opcode 100010, D = 1 (to REG from R/M), W = 0 (byte), MOD = 00 (no displacement), REG = 010 (DL), and R/M = 101 ([DI]). If the instruction changes to MOV DL,[DI + 1], the MOD field changes to 01 for an 8-bit displacement, but the first 2 bytes of the instruction otherwise remain the same. The instruction now becomes 8A5501H instead of 8A15H. Notice that the 8-bit displacement appends to the first 2 bytes of the instruction to form a 3-byte instruction instead of 2 bytes. If the instruction is again changed to a MOV DL,[DI + 1000H], the machine language form becomes 8A750010H. Here, the 16-bit displacement of 1000H (coded as 0010H) appends the opcode.

Special Addressing Mode.   There is a special addressing mode that does not appear in Tables 4–2, 4–3, or 4–4. It occurs whenever memory data are referenced by only the displacement mode of addressing for 16-bit instructions. Examples are the MOV [1000H],DL and MOV NUMB,DL instructions. The first instruction moves the contents of register DL into data segment memory location 1000H. The second instruction moves register DL into symbolic data segment memory location NUMB.

Whenever an instruction has only a displacement, the MOD field is always a 00 and the R/M field is always 110. As shown in the tables, the instruction contains no displacement and uses addressing mode [BP]. You cannot actually *use* addressing mode [BP] without a displacement in machine language. The assembler takes care of this by using an 8-bit displacement (MOD = 01)

Opcode = MOV
D = Transfer to register (REG)
W = Byte
MOD = No displacement
REG = DL
R/M = DS:[DI]
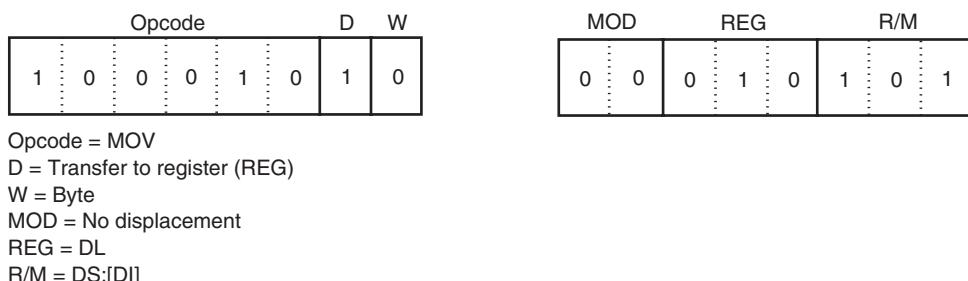
FIGURE 4–5 A MOV DL,[DI] instruction converted to its machine language form.

Byte 1

Opcode						D	W
1	0	0	0	1	0	0	0

Byte 2

MOD		REG			R/M		
0	0	0	1	0	1	1	0

Byte 3 — Displacement—low

0	0	0	0	0	0	0	0

Byte 4 — Displacement—high

0	0	0	1	0	0	0	0

Opcode = MOV
D = Transfer from register (REG)
W = Byte
MOD = because R/M is [BP] (special addressing)
REG = DL
R/M = DS:[BP]
Displacement = 1000H
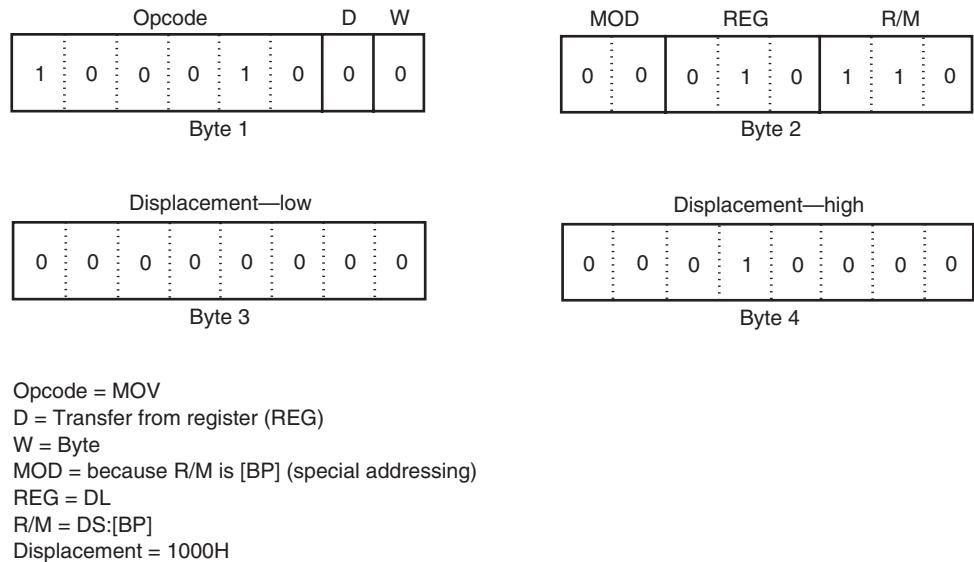
FIGURE 4–6   The MOV [1000H],DI instruction uses the special addressing mode.

of 00H whenever the [BP] addressing mode appears in an instruction. This means that the [BP] addressing mode assembles as a [BP + 0], even though a [BP] is used in the instruction. The same special addressing mode is also available for the 32-bit mode.

Figure 4–6 shows the binary bit pattern required to encode the MOV [1000H],DL instruction in machine language. If the individual translating this symbolic instruction into machine language does not know about the special addressing mode, the instruction would incorrectly translate to a MOV [BP],DL instruction. Figure 4–7 shows the actual form of the MOV [BP],DL instruction. Notice that this is a 3-byte instruction with a displacement of 00H.

Byte 1

Opcode						D	W
1	0	0	0	1	0	0	0

Byte 2

MOD		REG			R/M		
0	1	0	1	0	1	1	0

Byte 3 — 8-bit displacement

0	0	0	0	0	0	0	0

Opcode = MOV
D = Transfer from register (REG)
W = Byte
MOD = because R/M is [BP] (special addressing)
REG = DL
R/M = DS:[BP]
Displacement = 00H
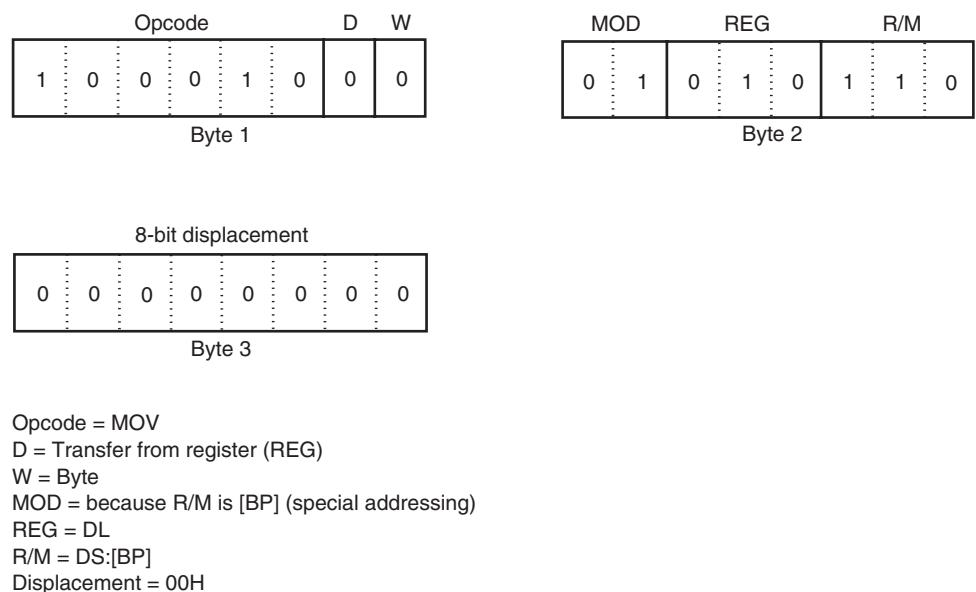
FIGURE 4–7   The MOV [BP],DL instruction converted to binary machine language.

TABLE 4–5 32-bit address-
ing modes selected by R/M.

R/M Code	Function
000	DS:[EAX]
001	DS:[ECX]
010	DS:[EDX]
011	DS:[EBX]
100	Uses scaled-index byte
101	SS:[EBP]*
110	DS:[ESI]
111	DS:[EDI]

*Note: See text section, Special Addressing Mode.
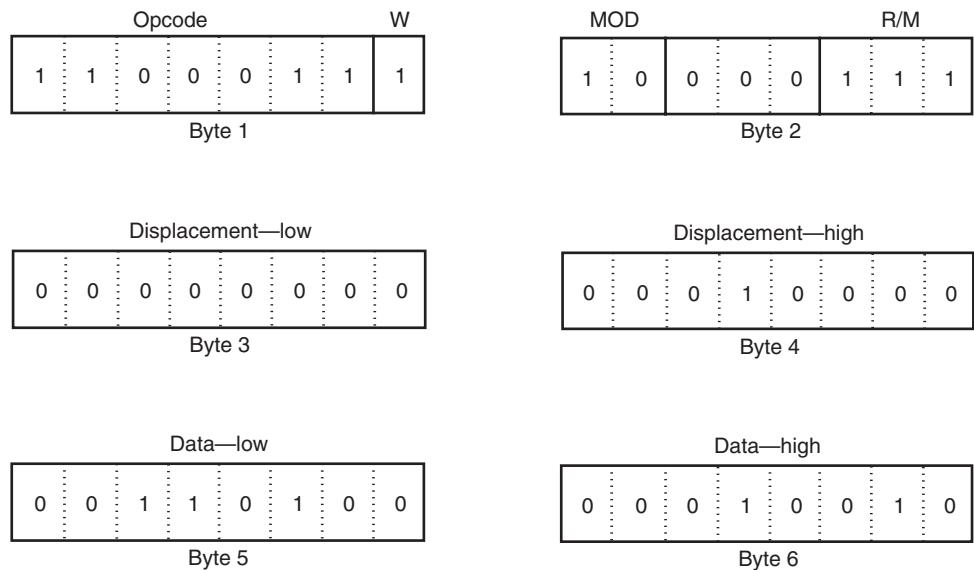
32-Bit Addressing Modes.   The 32-bit addressing modes found in the 80386 and above are obtained by either running these machines in the 32-bit instruction mode or in the 16-bit instruction mode by using the address-size prefix 67H. Table 4–5 shows the coding for R/M used to specify the 32-bit addressing modes. Notice that when R/M = 100, an additional byte called a scaled-index byte appears in the instruction. The scaled-index byte indicates the additional forms of scaled-index addressing that do not appear in Table 4–5. The scaled-index byte is mainly used when two registers are added to specify the memory address in an instruction. Because the scaled-index byte is added to the instruction, there are 7 bits in the opcode and 8 bits in the scaled-index byte to define. This means that a scaled-index instruction has 2^{15} (32K) possible combinations. There are over 32,000 different variations of the MOV instruction alone in the 80386 through the Core2 microprocessors.

     Figure 4–8 shows the format of the scaled-index byte as selected by a value of 100 in the R/M field of an instruction when the 80386 and above use a 32-bit address. The leftmost 2 bits select a scaling factor (multiplier) of 1×, 2×, 4×, or 8×. Note that a scaling factor of 1× is implicit if none is used in an instruction that contains two 32-bit indirect address registers. The index and base fields both contain register numbers, as indicated in Table 4–3 for 32-bit registers.

     The instruction MOV EAX,[EBX+4*ECX] is encoded as 67668B048BH. Notice that both the *address size* (67H) and *register size* (66H) override prefixes appear in the instruction. This coding (67668B048BH) is used when the 80386 and above microprocessors are operated in the 16-bit instruction mode for this instruction. If the microprocessor operates in the 32-bit instruction mode, both prefixes disappear and the instruction becomes an 8B048BH instruction. The use of the prefixes depends on the mode of operation of the microprocessor. Scaled-index addressing can also use a single register multiplied by a scaling factor. An example is the MOV AL,[2*ECX] instruction. The contents of the data segment location addressed by two times ECX are copied into AL.

An Immediate Instruction.   Suppose that the MOV WORD PTR [BX+1000H],1234H instruction is chosen as an example of a 16-bit instruction using immediate addressing. This instruction moves a 1234H into the word-sized memory location addressed by the sum of 1000H, BX, and

FIGURE 4–8 The scaled-index byte.

ss
00 = ×1
01 = ×2
10 = ×4
11 = ×8

Opcode | W

| 1 | 1 | 0 | 0 | 0 | 1 | 1 | 1 |

Byte 1

MOD | R/M

| 1 | 0 | 0 | 0 | 0 | 1 | 1 | 1 |

Byte 2

Displacement—low

| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

Byte 3

Displacement—high

| 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 |

Byte 4

Data—low

| 0 | 0 | 1 | 1 | 0 | 1 | 0 | 0 |

Byte 5

Data—high

| 0 | 0 | 0 | 1 | 0 | 0 | 1 | 0 |

Byte 6

Opcode = MOV (immediate)
W = Word
MOD = 16-bit displacement
REG = 000 (not used in immediate addressing)
R/M = DS:[BX]
Displacement = 1000H
Data = 1234H

FIGURE 4–9 A MOV WORD PTR [BX + 1000H], 1234H instruction converted to binary machine language.

DS × 10H. This 6-byte instruction uses 2 bytes for the opcode, W, MOD, and R/M fields. Two of the 6 bytes are the data of 1234H; 2 of the 6 bytes are the displacement of 1000H. Figure 4–9 shows the binary bit pattern for each byte of this instruction.

This instruction, in symbolic form, includes WORD PTR. The WORD PTR directive indicates to the assembler that the instruction uses a word-sized memory pointer. If the instruction moves a byte of immediate data, BYTE PTR replaces WORD PTR in the instruction. Likewise, if the instruction uses a doubleword of immediate data, the DWORD PTR directive replaces BYTE PTR. Most instructions that refer to memory through a pointer do not need the BYTE PTR, WORD PTR, or DWORD PTR directives. These directives are necessary only when it is not clear whether the operation is a byte, word, or doubleword. The MOV [BX],AL instruction is clearly a byte move; the MOV [BX],9 instruction is not exact, and could therefore be a byte-, word-, or doubleword-sized move. Here, the instruction must be coded as MOV BYTE PTR [BX],9, MOV WORD PTR [BX],9, or MOV DWORD PTR [BX],9. If not, the assembler flags it as an error because it cannot determine the intent of the instruction.

Segment MOV Instructions. If the contents of a segment register are moved by the MOV, PUSH, or POP instructions, a special set of register bits (REG field) selects the segment register (see Table 4–6).

Figure 4–10 shows a MOV BX,CS instruction converted to binary. The opcode for this type of MOV instruction is different for the prior MOV instructions. Segment registers can be moved between any 16-bit register or 16-bit memory location. For example, the MOV [DI],DS instruction stores the contents of DS into the memory location addressed by DI in the data

TABLE 4–6   Segment register selection.

Code	Segment Register
000	ES
001	CS*
010	SS
011	DS
100	FS
101	GS

*Note: MOV CS,R/M and POP CS are not allowed.

Opcode									MOD	REG	R/M			
1	0	0	0	1	1	0	0		1	1	0	0	1	0	1	1

Opcode = MOV
MOD = R/M is a register
REG = CS
R/M = BX

FIGURE 4–10   A MOV BX,CS instruction converted to binary machine language.

segment. An immediate segment register MOV is not available in the instruction set. To load a segment register with immediate data, first load another register with the data and then move it to a segment register.

Although this discussion has not been a complete coverage of machine language coding, it provides enough information for machine language programming. Remember that a program written in symbolic assembly language *(assembly language)* is rarely assembled by hand into binary machine language. An assembler program converts symbolic assembly language into machine language. With the microprocessor and its over 100,000 instruction variations, let us hope that an assembler is available for the conversion, because the process is very time-consuming, although not impossible.

The 64-Bit Mode for the Pentium 4 and Core2

None of the information presented thus far addresses the issue of 64-bit operation of the Pentium 4 or Core2. In the 64-bit mode, an additional prefix called REX (*register extension*) is added. The REX prefix, which is encoded as a 40H–4FH, follows other prefixes and is placed immediately before the opcode to modify it for 64-bit operation. The purpose of the REX prefix is to modify the reg and r/m fields in the second byte of the instruction. REX is needed to be able to address registers R8 through R15. Figure 4–11 illustrates the structure of REX and also its application to the second byte of the opcode.

The register and memory address assignments for the rrrr and mmmm fields are shown in Table 4–7 for 64-bit operations. The reg field can only contain register assignments as in other modes of operation and the r/m field contains either a register or memory assignment.

Figure 4–12 shows the scaled-index byte with the REX prefix for more complex addressing modes and also for using a scaling factor in the 64-bit mode of operation. As with 32-bit instructions, the modes allowed by the scaled-index byte are fairly all conclusive allowing pairs of registers to address memory and also an index factor of 2×, 4×, or 8×. An example is the instruction MOV RAXW,[RDX+RCX−12], which requires the scaled-index byte with an index of 1, which is understood but never entered into the instruction.

FIGURE 4–11 The application
of REX without scaled index.

W = 1 (64 bits)
W = 0 (CS descriptor)
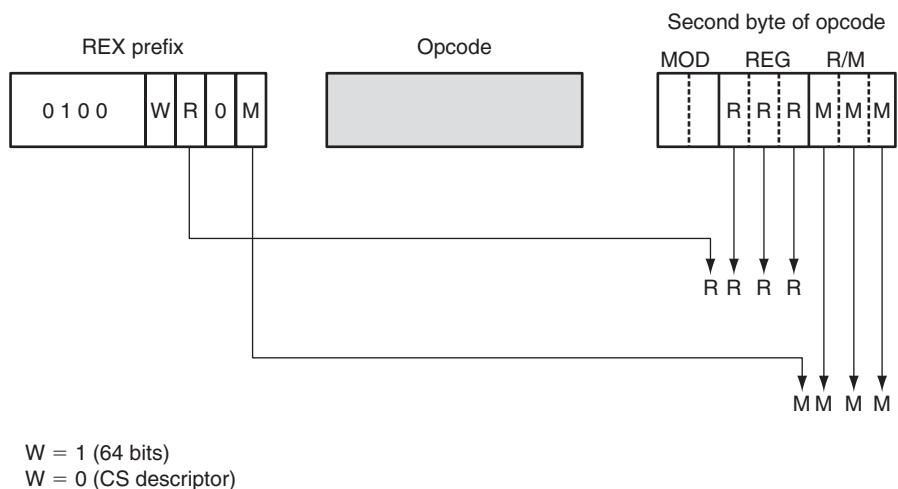
TABLE 4–7 The 64-bit reg-
ister and memory designators
for rrrr and mmmm.

Code	Register	Memory
0000	RAX	[RAX]
0001	RCX	[RCX]
0010	RDX	[RDX]
0011	RBX	[RBX]
0100	RSP	See note
0101	RBP	[RBP]
0110	RSI	[RSI]
0111	RDI	[RDI]
1000	R8	[R8]
1001	R9	[R9]
1010	R10	[R10]
1011	R11	[R11]
1100	R12	[R12]
1101	R13	[R13]
1110	R14	[R14]
1111	R15	[R15]

Note: This addressing mode specifies the inclusion
of the scaled-index byte.

W = 1 (64 bits)
W = 0 (CS descriptor)
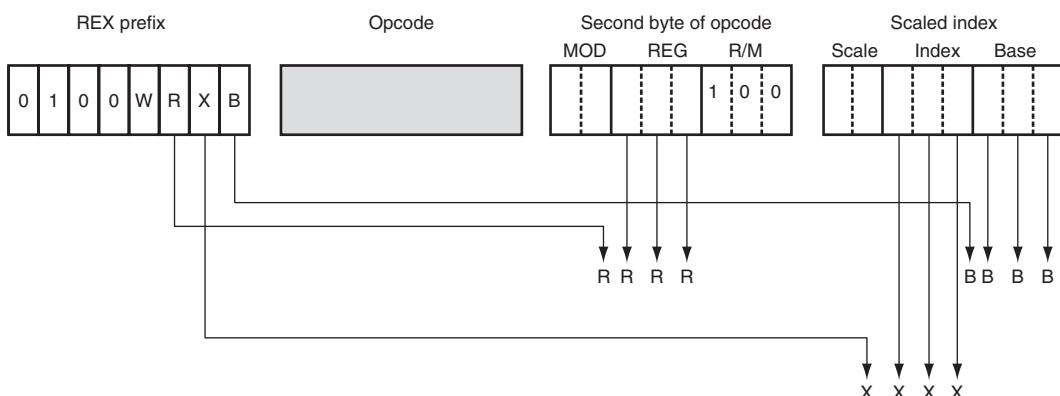
FIGURE 4–12 The scaled-index byte and REX prefix for 64-bit operations.

4–2   **PUSH/POP**

The PUSH and POP instructions are important instructions that *store* and *retrieve* data from the LIFO (last-in, first-out) stack memory. The microprocessor has six forms of the PUSH and POP instructions: register, memory, immediate, segment register, flags, and all registers. The PUSH and POP immediate and the PUSHA and POPA (all registers) forms are not available in the earlier 8086/8088 microprocessors, but are available to the 80286 through the Core2.

Register addressing allows the contents of any 16-bit register to be transferred to or from the stack. In the 80386 and above, the 32-bit extended registers and flags (EFLAGS) can also be pushed or popped from the stack. Memory-addressing PUSH and POP instructions store the contents of a 16-bit memory location (or 32 bits in the 80386 and above) on the stack or stack data into a memory location. Immediate addressing allows immediate data to be pushed onto the stack, but not popped off the stack. Segment register addressing allows the contents of any segment register to be pushed onto the stack or removed from the stack (ES may be pushed, but data from the stack may never be popped into ES). The flags may be pushed or popped from that stack, and the contents of all the registers may be pushed or popped.

PUSH

The 8086–80286 PUSH instruction always transfers 2 bytes of data to the stack; the 80386 and above transfer 2 or 4 bytes, depending on the register or size of the memory location. The source of the data may be any internal 16- or 32-bit register, immediate data, any segment register, or any 2 bytes of memory data. There is also a PUSHA instruction that copies the contents of the internal register set, except the segment registers, to the stack. The PUSHA (**push all**) instruction copies the registers to the stack in the following order: AX, CX, DX, BX, SP, BP, SI, and DI. The value for SP that is pushed onto the stack is whatever it was before the PUSHA instruction executed. The PUSHF (**push flags**) instruction copies the contents of the flag register to the stack. The PUSHAD and POPAD instructions push and pop the contents of the 32-bit register set found in the 80386 through the Pentium 4. The PUSHA and POPA instructions do not function in the 64-bit mode of operation for the Pentium 4.

Whenever data are pushed onto the stack, the first (most-significant) data byte moves to the stack segment memory location addressed by $SP - 1$. The second (least-significant) data byte moves into the stack segment memory location addressed by $SP - 2$. After the data are stored by a PUSH, the contents of the SP register decrement by 2. The same is true for a doubleword push, except that 4 bytes are moved to the stack memory (most-significant byte first), after which the stack pointer decrements by 4. Figure 4–13 shows the operation of the PUSH AX instruction. This instruction copies the contents of AX onto the stack where address $SS:[SP - 1] = AH$, $SS:[SP - 2] = AL$, and afterwards $SP = SP - 2$. In 64-bit mode, 8 bytes of the stack are used to store the number pushed onto the stack.

The PUSHA instruction pushes all the internal 16-bit registers onto the stack, as illustrated in Figure 4–14. This instruction requires 16 bytes of stack memory space to store all eight 16-bit registers. After all registers are pushed, the contents of the SP register are decremented by 16. The PUSHA instruction is very useful when the entire register set (microprocessor environment) of the 80286 and above must be saved during a task. The PUSHAD instruction places the 32-bit register set on the stack in the 80386 through the Core2. PUSHAD requires 32 bytes of stack storage space.
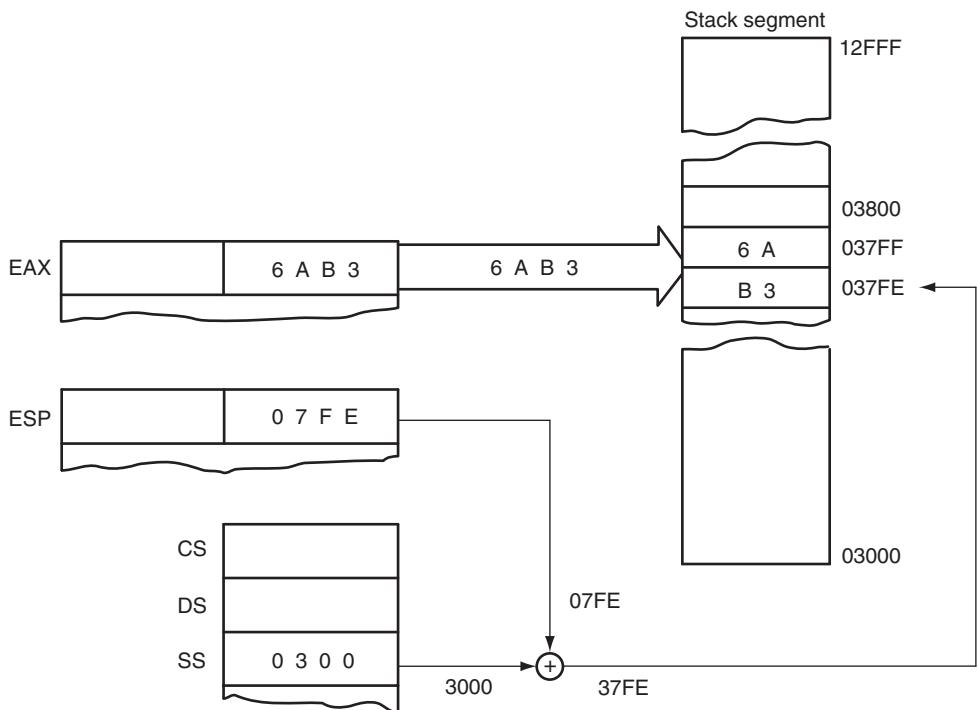
FIGURE 4–13 The effect of the PUSH AX instruction on ESP and stack memory locations 37FFH and 37FEH. This instruction is shown at the point after execution.

The PUSH immediate data instruction has two different opcodes, but in both cases, a 16-bit immediate number moves onto the stack; if PUSHD is used, a 32-bit immediate datum is pushed. If the values of the immediate data are 00H–FFH, the opcode is a 6AH; if the data are 0100H–FFFFH, the opcode is 68H. The PUSH 8 instruction, which pushes 0008H onto the

FIGURE 4–14 The operation of the PUSHA instruction, showing the location and order of stack data.
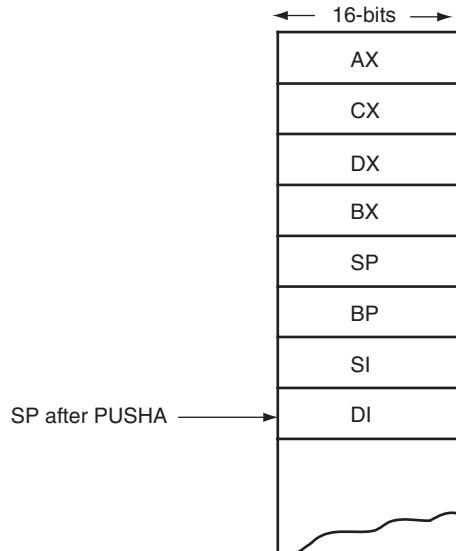
TABLE 4–8 The PUSH
instruction.

Symbolic	Example	Note
PUSH reg16	PUSH BX	16-bit register
PUSH reg32	PUSH EDX	32-bit register
PUSH mem16	PUSH WORD PTR[BX]	16-bit pointer
PUSH mem32	PUSH DWORD PTR[EBX]	32-bit pointer
PUSH mem64	PUSH QWORD PTR[RBX]	64-bit pointer (64-bit mode)
PUSH seg	PUSH DS	Segment register
PUSH imm8	PUSH 'R'	8-bit immediate
PUSH imm16	PUSH 1000H	16-bit immediate
PUSHD imm32	PUSHD 20	32-bit immediate
PUSHA	PUSHA	Save all 16-bit registers
PUSHAD	PUSHAD	Save all 32-bit registers
PUSHF	PUSHF	Save flags
PUSHFD	PUSHFD	Save EFLAGS

stack, assembles as 6A08H. The PUSH 1000H instruction assembles as 680010H. Another example of PUSH immediate is the PUSH 'A' instruction, which pushes a 0041H onto the stack. Here, the 41H is the ASCII code for the letter A.

Table 4–8 lists the forms of the PUSH instruction that include PUSHA and PUSHF. Notice how the instruction set is used to specify different data sizes with the assembler.

POP

The POP instruction performs the inverse operation of a PUSH instruction. The POP instruction removes data from the stack and places it into the target 16-bit register, segment register, or a 16-bit memory location. In the 80386 and above, a POP can also remove 32-bit data from the stack and use a 32-bit address. The POP instruction is not available as an immediate POP. The POPF (pop flags) instruction removes a 16-bit number from the stack and places it into the flag register; the POPFD removes a 32-bit number from the stack and places it into the extended flag register. The POPA (pop all) instruction removes 16 bytes of data from the stack and places them into the following registers, in the order shown: DI, SI, BP, SP, BX, DX, CX, and AX. This is the reverse order from the way they were placed on the stack by the PUSHA instruction, causing the same data to return to the same registers. In the 80386 and above, a POPAD instruction reloads the 32-bit registers from the stack.

Suppose that a POP BX instruction executes. The first byte of data removed from the stack (the memory location addressed by SP in the stack segment) moves into register BL. The second byte is removed from stack segment memory location SP + 1 and is placed into register BH. After both bytes are removed from the stack, the SP register is incremented by 2. Figure 4–15 shows how the POP BX instruction removes data from the stack and places them into register BX.

The opcodes used for the POP instruction and all of its variations appear in Table 4–9. Note that a POP CS instruction is not a valid instruction in the instruction set. If a POP CS instruction executes, only a portion of the address (CS) of the next instruction changes. This makes the POP CS instruction unpredictable and therefore not allowed.

Initializing the Stack

When the stack area is initialized, load both the stack segment (SS) register and the stack pointer (SP) register. It is normal to designate an area of memory as the stack segment by loading SS with the bottom location of the stack segment.
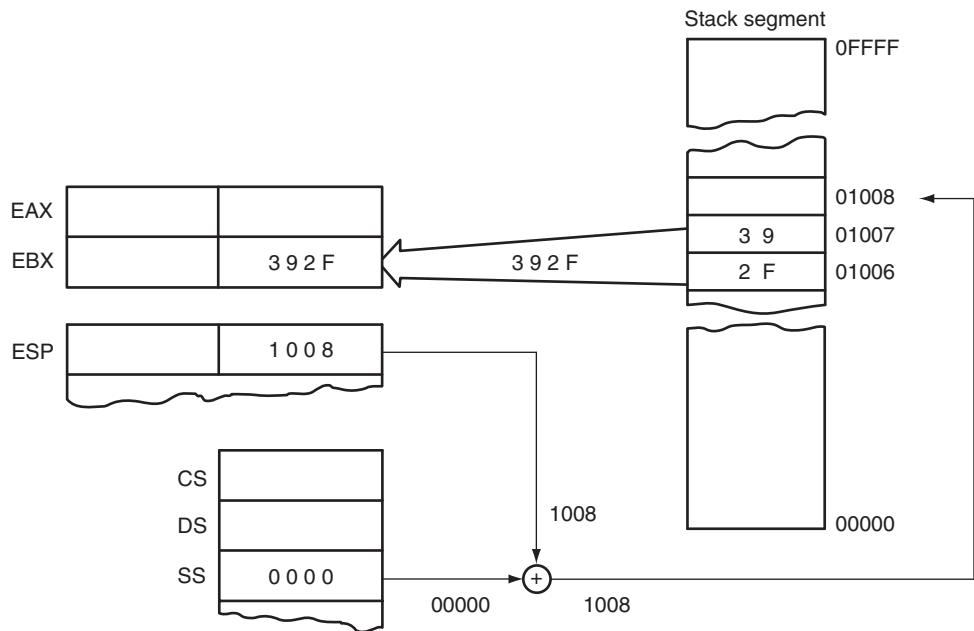
FIGURE 4–15 The POP BX instruction, showing how data are removed from the stack. This instruction is shown after execution.

TABLE 4–9 The POP instructions.

Symbolic	Example	Note
POP reg16	POP CX	16-bit register
POP reg32	POP EBP	32-bit register
POP mem16	POP WORD PTR[BX+1]	16-bit pointer
POP mem32	POP DATA3	32-bit memory address
POP mem64	POP FROG	64-bit memory address (64-bit mode)
POP seg	POP FS	Segment register
POPA	POPA	Pops all 16-bit registers
POPAD	POPAD	Pops all 32-bit registers
POPF	POPF	Pops flags
POPFD	POPFD	Pops EFLAGS

For example, if the stack segment is to reside in memory locations 10000H–1FFFFH, load SS with a 1000H. (Recall that the rightmost end of the stack segment register is appended with a 0H for real mode addressing.) To start the stack at the top of this 64K-byte stack segment, the stack pointer (SP) is loaded with a 0000H. Likewise, to address the top of the stack at location 10FFFH, use a value of 1000H in SP. Figure 4–16 shows how this value causes data to be pushed onto the top of the stack segment with a PUSH CX instruction. Remember that all segments are cyclic in nature—that is, the top location of a segment is contiguous with the bottom location of the segment.

In assembly language, a stack segment is set up as illustrated in Example 4–1. The first statement identifies the start of the stack segment and the last statement identifies the end of the stack segment. The assembler and linker programs place the correct stack segment address in SS and the length of the segment (top of the stack) into SP. There is no need to load these registers in your program unless you wish to change the initial values for some reason.
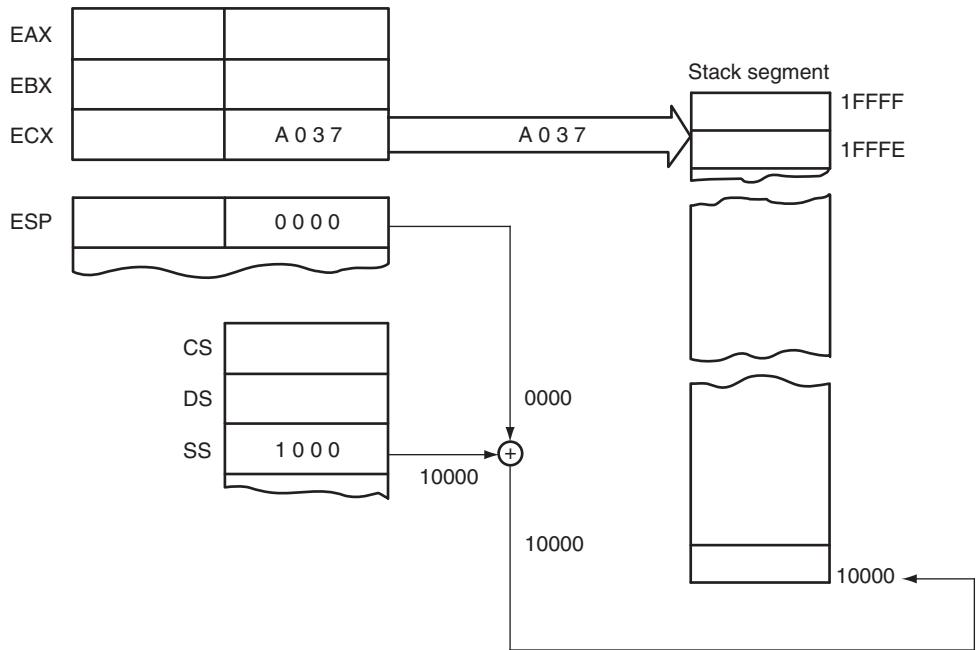
FIGURE 4–16 The PUSH CX instruction, showing the cyclical nature of the stack segment. This instruction is shown just before execution, to illustrate that the stack bottom is contiguous to the top.

EXAMPLE 4–1

```
0000                      STACK_SEG    SEGMENT STACK
0000 0100[                             DW    100H DUP(?)
        ????
          ]
0200                      STACK_SEG    ENDS
```

An alternative method for defining the stack segment is used with one of the memory models for the MASM assembler only (refer to Appendix A). Other assemblers do not use models; if they do, the models are not exactly the same as with MASM. Here, the .STACK statement, followed by the number of bytes allocated to the stack, defines the stack area (see Example 4–2). The function is identical to Example 4–1. The .STACK statement also initializes both SS and SP. Note that this text uses memory models that are designed for the Microsoft Macro Assembler program MASM.

EXAMPLE 4–2

```
            .MODEL SMALL
            .STACK 200H  ;set stack size
```

If the stack is not specified by using either method, a warning will appear when the program is linked. The warning may be ignored if the stack size is 128 bytes or fewer. The system automatically assigns (through DOS) at least 128 bytes of memory to the stack. This memory section is located in the **program segment prefix** (PSP), which is appended to the beginning of each program file. If you use more memory for the stack, you will erase information in the PSP that is critical to the operation of your program and the computer. This error often causes the computer program to crash. If the TINY memory model is used, the stack is automatically located at the very end of the segment, which allows for a larger stack area.

4–3   LOAD-EFFECTIVE ADDRESS

There are several load-effective address instructions in the microprocessor instruction set. The LEA instruction loads any 16-bit register with the offset address, as determined by the addressing mode selected for the instruction. The LDS and LES variations load any 16-bit register with the offset address retrieved from a memory location, and then load either DS or ES with a segment address retrieved from memory. In the 80386 and above, LFS, LGS, and LSS are added to the instruction set, and a 32-bit register can be selected to receive a 32-bit offset from memory. In the 64-bit mode for the Pentium 4, the LDS and LES instructions are *invalid* and not used because the segments have no function in the flat memory model. Table 4–10 lists the load-effective address instructions.

LEA

The LEA instruction loads a 16- or 32-bit register with the offset address of the data specified by the operand. As the first example in Table 4–9 shows, the operand address NUMB is loaded into register AX, not the contents of address NUMB.

By comparing LEA with MOV, we observe that LEA BX,[DI] loads the offset address specified by [DI] (contents of DI) into the BX register; MOV BX,[DI] loads the data stored at the memory location addressed by [DI] into register BX.

Earlier in the text, several examples were presented by using the OFFSET directive. The OFFSET directive performs the same function as an LEA instruction if the operand is a displacement. For example, the MOV BX,OFFSET LIST performs the same function as LEA BX,LIST. Both instructions load the offset address of memory location LIST into the BX register. See Example 4–3 for a short program that loads SI with the address of DATA1 and DI with the address of DATA2. It then exchanges the contents of these memory locations. Note that the LEA and MOV with OFFSET instructions are both the same length (3 bytes).

EXAMPLE 4–3

```
                          .MODEL  SMALL          ;select small model
0000                      .DATA                  ;start data segment
0000 2000      DATA1      DW      2000H          ;define DATA1
0002 3000      DATA2      DW      3000H          ;define DATA2
0000                      .CODE                  ;start code segment
                          .STARTUP               ;start program
0017 BE 0000 R            LEA SI,DATA1           ;address DATA1 with SI
001A BF 0002 R            MOV DI,OFFSET DATA2    ;address DATA2 with DI
001D 8B 1C                MOV BX,[SI]            ;exchange DAT1 with DATA2
001F 8B 0D                MOV CX,[DI]
0021 89 0C                MOV [SI],CX
0023 89 1D                MOV [DI],BX
                          .EXIT
                          END
```

TABLE 4–10   Load-effective address instructions.

Assembly Language	Operation
LEA AX,NUMB	Loads AX with the offset address of NUMB
LEA EAX,NUMB	Loads EAX with the offset address of NUMB
LDS DI,LIST	Loads DS and DI with the 32-bit contents of data segment memory location LIST
LDS EDI,LIST1	Loads the DS and EDI with the 48-bit contents of data segment memory location LIST1
LES BX,CAT	Loads ES and BX with the 32-bit contents of data segment memory location CAT
LFS DI,DATA1	Loads FS and DI with the 32-bit contents of data segment memory location DATA1
LGS SI,DATA5	Loads GS and SI with the 32-bit contents of data segment memory location DATA5
LSS SP,MEM	Loads SS and SP with the 32-bit contents of data segment memory location MEM

But why is the LEA instruction available if the OFFSET directive accomplishes the same task? First, OFFSET only functions with simple operands such as LIST. It may not be used for an operand such as [DI], LIST [SI], and so on. The OFFSET directive is more efficient than the LEA instruction for simple operands. It takes the microprocessor longer to execute the LEA BX,LIST instruction than the MOV BX,OFFSET LIST. The 80486 microprocessor, for example, requires two clocks to execute the LEA BX,LIST instruction and only one clock to execute MOV BX,OFFSET LIST. The reason that the MOV BX,OFFSET LIST instruction executes faster is because the assembler calculates the offset address of LIST, whereas the microprocessor calculates the address for the LEA instruction. The MOV BX,OFFSET LIST instruction is actually assembled as a move immediate instruction and is more efficient.

Suppose that the microprocessor executes an LEA BX,[DI] instruction and DI contains a 1000H. Because DI contains the offset address, the microprocessor transfers a copy of DI into BX. A MOV BX,DI instruction performs this task in less time and is often preferred to the LEA BX,[DI] instruction.

Another example is LEA SI,[BX + DI]. This instruction adds BX to DI and stores the sum in the SI register. The sum generated by this instruction is a modulo-64K sum. (A modulo-64K sum drops any carry out of the 16-bit result.) If BX = 1000H and DI = 2000H, the offset address moved into SI is 3000H. If BX = 1000H and DI = FF00H, the offset address is 0F00H instead of 10F00H. Notice that the second result is a modulo-64K sum of 0F00H.

LDS, LES, LFS, LGS, and LSS

The LDS, LES, LFS, LGS, and LSS instructions load any 16-bit or 32-bit register with an offset address, and the DS, ES, FS, GS, or SS segment register with a segment address. These instructions use any of the memory-addressing modes to access a 32-bit or 48-bit section of memory that contains both the segment and offset address. The 32-bit section of memory contains a 16-bit offset and segment address, while the 48-bit section contains a 32-bit offset and a segment address. These instructions may not use the register addressing mode (MOD = 11). Note that the LFS, LGS, and LSS instructions are only available on 80386 and above, as are the 32-bit registers.

Figure 4–17 illustrates an example LDS BX,[DI] instruction. This instruction transfers the 32-bit number, addressed by DI in the data segment, into the BX and DS registers. The LDS, LES, LFS, LGS, and LSS instructions obtain a new far address from memory. The offset address appears first, followed by the segment address. This format is used for storing all 32-bit memory addresses.

A far address can be stored in memory by the assembler. For example, the ADDR DD FAR PTR FROG instruction stores the offset and segment address (far address) of FROG in 32 bits of memory at location ADDR. The DD directive tells the assembler to store a doubleword (32-bit number) in memory address ADDR.

In the 80386 and above, an LDS EBX,[DI] instruction loads EBX from the 4-byte section of memory addressed by DI in the data segment. Following this 4-byte offset is a word that is loaded to the DS register. Notice that instead of addressing a 32-bit section of memory, the 80386 and above address a 48-bit section of the memory whenever a 32-bit offset address is loaded to a 32-bit register. The first 4 bytes contain the offset value loaded to the 32-bit register and the last 2 bytes contain the segment address.

The most useful of the load instructions is the LSS instruction. Example 4–4 shows a short program that creates a new stack area after saving the address of the old stack area. After executing some dummy instructions, the old stack area is reactivated by loading both SS and SP with the LSS instruction. Note that the CLI (**disable interrupt**) and STI (**enable interrupt**) instructions must be included to disable interrupts. (This topic is discussed near the end of this chapter.) Because the LSS instruction functions in the 80386 or above, the .386 statement appears after
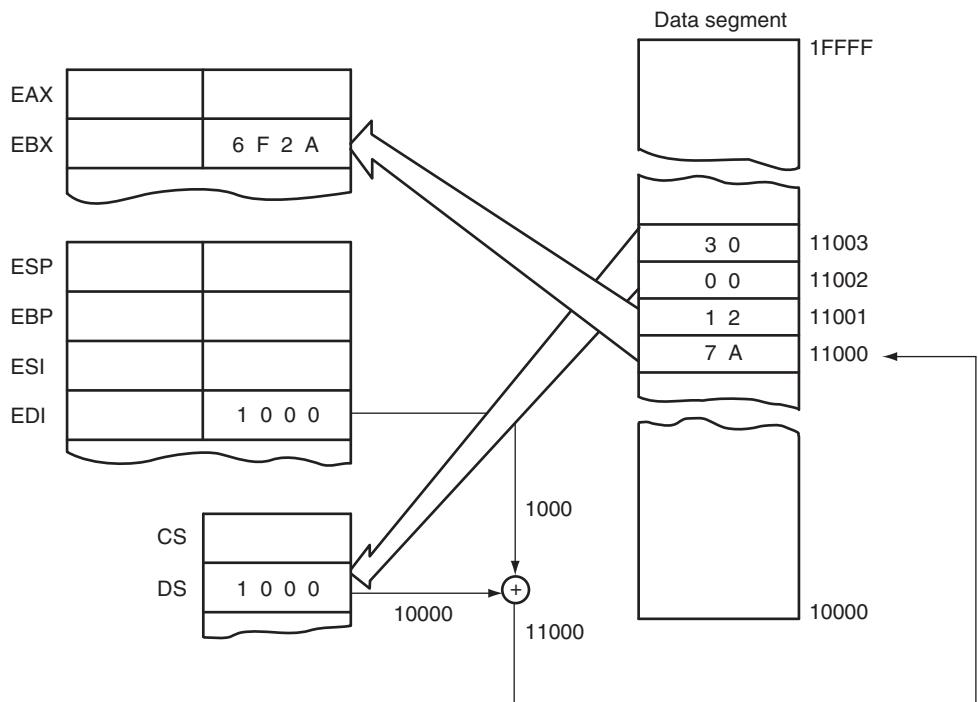
FIGURE 4–17 The LDS BX,[DI] instruction loads register BX from addresses 11000H and 11001H and register DS from locations 11002H and 11003H. This instruction is shown at the point just before DS changes to 3000H and BX changes to 127AH.

the .MODEL statement to select the 80386 microprocessor. Notice how the WORD PTR directive is used to override the doubleword (DD) definition for the old stack memory location. If an 80386 or newer microprocessor is in use, it is suggested that the .386 switch be used to develop software for the 80386 microprocessor. This is true even if the microprocessor is a Pentium, Pentium Pro, Pentium II, Pentium III, Pentium 4, or Core2. The reason is that the 80486–Core2 microprocessors add only a few additional instructions to the 80386 instruction set, which are seldom used in software development. If the need arises to use any of the CMPXCHG, CMPX-CHG8 (new to the Pentium), XADD or BSWAP instructions, select either the .486 switch for the 80486 microprocessor of the .586 switch for the Pentium. You can even specify the Pentium II –Core2 using the .686 switch.

EXAMPLE 4–4

```
                        .MODEL SMALL          ;select small model
                        .386                  ;select 80386
0000                    .DATA                 ;start data segment
0000 00000000   SADDR   DD      ?             ;old stack address
0004 1000 [     SAREA   DW      1000H DUP(?)  ;new stack area
        ????
            ]
2004 = 2004     STOP    EQU THIS WORD         ;define top of new stack
0000                    .CODE                 ;start code segment
                        .STARTUP              ;start program
0010 FA                 CLI                   ;disable interrupts
0011 8B C4              MOV AX,SP             ;save old SP
0013 A3 0000 R          MOV WORD PTR SADDR,AX
0016 8C D0              MOV AX,SS             ;save old SS
0018 A3 0002 R          MOV WORD PTR SADDR+2,AX
```

```
001B 8C D8                  MOV AX,DS              ;load new SS
001D 8E D0                  MOV SS,AX
001F B8 2004 R              MOV AX,OFFSET STOP     ;load new SP
0022 8B E0                  MOV SP,AX
0024 FB                     STI                    ;enable interrupts

0025 8B C0                  MOV AX,AX              ;do some dummy instructions
0027 8B C0                  MOV AX,AX
0029 9F B2 26 0000 R        LSS SP,SADDR           ;get old stack
                            .EXIT                  ;exit to DOS
                            END                    ;end program listing
```

4–4   STRING DATA TRANSFERS

There are five string data transfer instructions: LODS, STOS, MOVS, INS, and OUTS. Each string instruction allows data transfers that are either a single byte, word, or doubleword (or if repeated, a block of bytes, words, or doublewords). Before the string instructions are presented, the operation of the D flag-bit (direction), DI, and SI must be understood as they apply to the string instructions. In the 64-bit mode of the Pentium 4 and Core2, quadwords are also used with the string instructions such as LODSQ.

The Direction Flag

The direction flag (D, located in the flag register) selects the auto-increment ($D = 0$) or the auto-decrement ($D = 1$) operation for the DI and SI registers during string operations. The direction flag is used only with the string instructions. The CLD instruction clears the D flag ($D = 0$) and the STD instruction sets it ($D = 1$). Therefore, the CLD instruction selects the auto-increment mode ($D = 0$) and STD selects the auto-decrement mode ($D = 1$).

Whenever a string instruction transfers a byte, the contents of DI and/or SI are incremented or decremented by 1. If a word is transferred, the contents of DI and/or SI are incremented or decremented by 2. Doubleword transfers cause DI and/or SI to increment or decrement by 4. Only the actual registers used by the string instruction are incremented or decremented. For example, the STOSB instruction uses the DI register to address a memory location. When STOSB executes, only the DI register is incremented or decremented without affecting SI. The same is true of the LODSB instruction, which uses the SI register to address memory data. A LODSB instruction will only increment or decrement SI without affecting DI.

DI and SI

During the execution of a string instruction, memory accesses occur through either or both of the DI and SI registers. The DI offset address accesses data in the extra segment for all string instructions that use it. The SI offset address accesses data, by default, in the data segment. The segment assignment of SI may be changed with a segment override prefix, as described later in this chapter. The DI segment assignment is always in the extra segment when a string instruction executes. This assignment cannot be changed. The reason that one pointer addresses data in the extra segment and the other in the data segment is so that the MOVS instruction can move 64K bytes of data from one segment of memory to another.

When operating in the 32-bit mode in the 80386 microprocessor or above, the EDI and ESI registers are used in place of DI and SI. This allows string using any memory location in the entire 4G-byte protected mode address space of the microprocessor.

LODS

The LODS instruction loads AL, AX, or EAX with data stored at the data segment offset address indexed by the SI register. (Note that only the 80386 and above use EAX.) After loading AL with

TABLE 4–11 Forms of the LODS instruction.

Assembly Language	Operation
LODSB	AL = DS:[SI]; SI = SI ± 1
LODSW	AX = DS:[SI]; SI = SI ± 2
LODSD	EAX = DS:[SI]; SI = SI ± 4
LODSQ	RAX = [RSI]; RSI = RSI ± 8 (64-bit mode)
LODS LIST	AL = DS:[SI]; SI = SI ± 1 (if LIST is a byte)
LODS DATA1	AX = DS:[SI]; SI = SI ± 2 (if DATA1 is a word)
LODS FROG	EAX = DS:[SI]; SI = SI ± 4 (if FROG is a doubleword)

Note: The segment register can be overridden with a segment override prefix as in LODS ES:DATA4.

a byte, AX with a word, or EAX with a doubleword, the contents of SI increment, if D = 0 or decrement, if D = 1. A 1 is added to or subtracted from SI for a byte-sized LODS, a 2 is added or subtracted for a word-sized LODS, and a 4 is added or subtracted for a doubleword-sized LODS.

Table 4–11 lists the permissible forms of the LODS instruction. The LODSB (loads a byte) instruction causes a byte to be loaded into AL, the LODSW (loads a word) instruction causes a word to be loaded into AX, and the LODSD (loads a doubleword) instruction causes a doubleword to be loaded into EAX. Although rare, as an alternative to LODSB, LODSW, LODSD, and LODSQ, the LODS instruction may be followed by a byte-, word- or doubleword-sized operand to select a byte, word, or doubleword transfer. Operands are often defined as bytes with DB, as words with DW, and as doublewords with DD. The DB pseudo-operation defines byte(s), the DW pseudo-operation defines word(s), and the DD pseudo-operations define doubleword(s).

Figure 4–18 shows the effect of executing the LODSW instruction if the D flag = 0, SI = 1000H, and DS = 1000H. Here, a 16-bit number stored at memory locations 11000H and 1l001H moves into AX. Because D = 0 and this is a word transfer, the contents of SI increment by 2 after AX loads with memory data.

STOS

The STOS instruction stores AL, AX, or EAX at the extra segment memory location addressed by the DI register. (Note that only the 80386–Core2 use EAX and doublewords.) Table 4–12 lists all forms of the STOS instruction. As with LODS, a STOS instruction may be appended with a B, W, or D for byte, word, or doubleword transfers. The STOSB (**stores a byte**) instruction stores the byte in AL at the extra segment memory location addressed by DI. The STOSW (**stores a word**) instruction stores AX in the extra segment memory location addressed by DI. A doubleword is stored in the extra segment location addressed by DI with the STOSD (**stores a doubleword**) instruction. After the byte (AL), word (AX), or doubleword (EAX) is stored, the contents of DI increment or decrement.

STOS with a REP.   The **repeat prefix** (REP) is added to any string data transfer instruction, except the LODS instruction. It doesn't make any sense to perform a repeated LODS operation. The REP prefix causes CX to decrement by 1 each time the string instruction executes. After CX decrements, the string instruction repeats. If CX reaches a value of 0, the instruction terminates and the program continues with the next sequential instruction. Thus, if CX is loaded with 100 and a REP STOSB instruction executes, the microprocessor automatically repeats the STOSB instruction 100 times. Because the DI register is automatically incremented or decremented after each datum is stored, this instruction stores the contents of AL in a block of memory instead of a single byte of memory. In the Pentium 4 operated in 64-bit mode, the RCX register is used with the REP prefix.
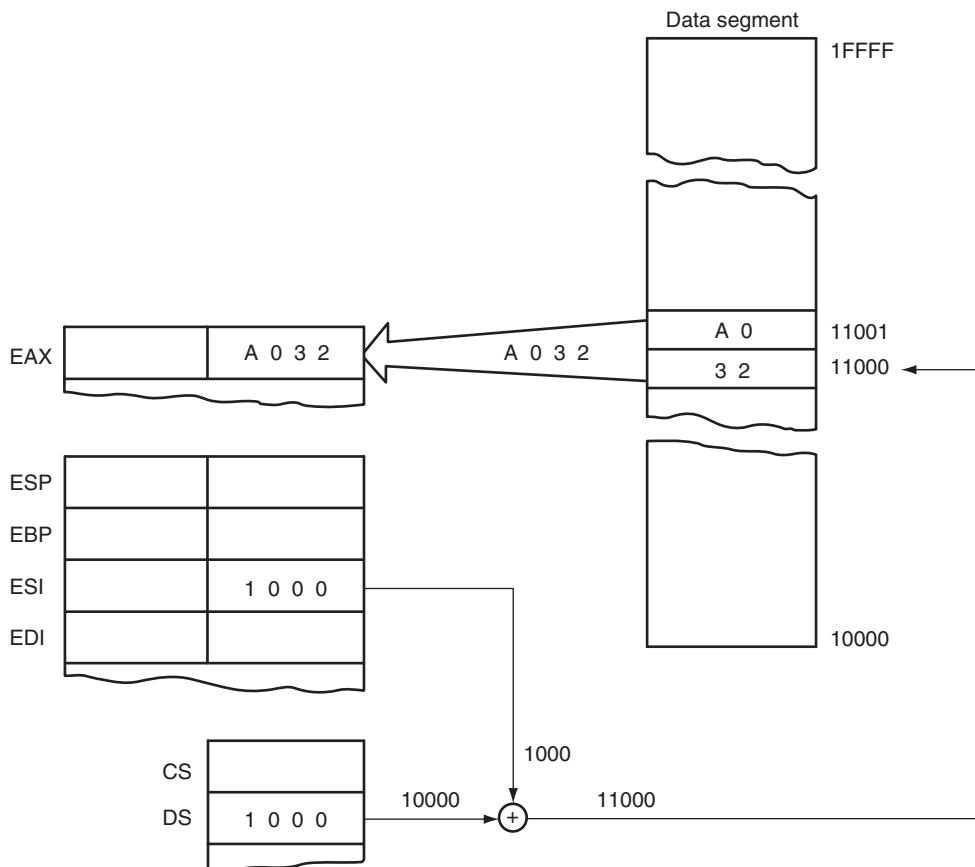
FIGURE 4–18 The operation of the LODSW instruction if DS = 1000H, D = 0, 11000H = 32, and 11001H = A0. This instruction is shown after AX is loaded from memory, but before SI increments by 2.

Suppose that the STOSW instruction is used to clear an area of memory called Buffer using a count called Count and the program is to function call ClearBuffer in the C++ environment using the inline assembler. (See Example 4–5.) Note that both the Count and Buffer address are transferred to the function. The REP STOSW instruction clears the memory buffer called Buffer. Notice that Buffer is a pointer to the actual buffer that is cleared by this function.

TABLE 4–12 Forms of the STOS instruction.

Assembly Language	Operation
STOSB	ES:[DI] = AL; DI = DI ± 1
STOSW	ES:[DI] = AX; DI = DI ± 2
STOSD	ES:[DI] = EAX; DI = DI ± 4
STOSQ	[RDI] = RAX; RDI = RDI ± 8 (64-bit mode)
STOS LIST	ES:[DI] = AL; DI = DI ± 1 (if LIST is a byte)
STOS DATA3	ES:[DI] = AX; DI = DI ± 2 (if DATA3 is a word)
STOS DATA4	ES:[DI] = EAX; DI = DI ± 4 (if DATA4 is a doubleword)

EXAMPLE 4–5

```
void ClearBuffer (int count, short* buffer)
{
   _asm{
          push edi              ;save registers
          push es
          push ds
          mov  ax,0
          mov  ecx, count
          mov  edi, buffer
          pop  es               ;load ES with DS
          rep  stosw            ;clear Buffer
          pop  es               ;restore registers
          pop  edi
   }
}
```

The operands in a program can be modified by using arithmetic or logic operators such as multiplication (*). Other operators appear in Table 4–13.

MOVS

One of the more useful string data transfer instructions is MOVS, because it transfers data from one memory location to another. This is the only memory-to-memory transfer allowed in the 8086–Pentium 4 microprocessors. The MOVS instruction transfers a byte, word, or doubleword from the data segment location addressed by SI to the extra segment location addressed by SI. As with the other string instructions, the pointers then are incremented or decremented, as dictated by the direction flag. Table 4–14 lists all the permissible forms of the MOVS instruction. Note that only the source operand (SI), located in the data segment, may be overridden so that another segment may be used. The destination operand (DI) must always be located in the extra segment.

It is often necessary to transfer the contents of one area of memory to another. Suppose that we have two blocks of doubleword memory, blockA and blockB, and we need to copy blockA into blockB. This can be accomplished using the MOVSD instruction as illustrated in Example 4–6, which is a C++ language function written using the inline assembler. The function receives three pieces of information from the caller: blockSize and the addresses of blockA and blockB. Note that all data are in the data segment in a Visual C++ program so we need to copy DS into ES, which is done using a PUSH DS followed by a POP ES. We also need to save all registers that we changed except for EAX, EBX, ECX, and EDX.

Example 4–7 shows the same function written in C++ exclusively, so the two methods can be compared and contrasted. Example 4–8 shows the assembly language version of

TABLE 4–13   Common operand modifiers.

Operator	Example	Comment
+	MOV AL,6+3	Copies 9 into AL
−	MOV AL,6−3	Copies 3 into AL
*	MOV AL,4*3	Copies 12 into AL
/	MOV AX,12/5	Copies 2 into AX (remainder is lost)
MOD	MOV AX,12 MOD 7	Copies 5 into AX (quotient is lost)
AND	MOV AX,12 AND 4	Copies 4 into AX (1100 AND 0100 = 0100)
OR	MOV EAX,12 OR 1	Copies 13 into EAX (1100 OR 0001 = 1101)
NOT	MOV AL,NOT 1	Copies 254 into AL (NOT 0000 0001 = 1111 1110 or 254)

TABLE 4–14 Forms of the MOVS instruction.

Assembly Language	Operation
MOVSB	ES:[DI] = DS:[SI]; DI = DI ± 1; SI = SI ± 1 (byte transferred)
MOVSW	ES:[DI] = DS:[SI]; DI = DI ± 2; SI = SI ± 2 (word transferred)
MOVSD	ES:[DI] = DS:[SI]; DI = DI ± 4; SI = SI ± 4 (doubleword transferred)
MOVSQ	[RDI] = [RSI]; RDI = RDI ± 8; RSI = RSI ± 8 (64-bit mode)
MOVS BYTE1, BYTE2	ES:[DI] = DS:[SI]; DI = DI ± 1; SI = SI ± 1 (byte transferred if BYTE1 and BYTE2 are bytes)
MOVS WORD1,WORD2	ES:[DI] = DS:[SI]; DI = DI ± 2; SI = SI ± 2 (word transferred if WORD1 and WORD2 are words)
MOVS TED,FRED	ES:[DI] = DS:[SI]; DI = DI ± 4; SI = SI ± 4 (doubleword transferred if TED and FRED are doublewords)

Example 4–7 for comparison to Example 4–6. Notice how much shorter the assembly language version is compared to the C++ version generated in Example 4–8. Admittedly the C++ version is a little easier to type, but if execution speed is important, Example 4–6 will run much faster than Example 4–7.

EXAMPLE 4–6

```
//Function that copies blockA into blockB using the inline assembler
//
void TransferBlocks (int blockSize, int* blockA, int* blockB)
{
      _asm{
              push es                   ;save registers
              push edi
              push esi
              push ds                   ;copy DS into ES
              pop  es
              mov  esi, blockA          ;address blockA
              mov  edi, blockB          ;address blockB
              mov  ecx, blockSize       ;load count
              rep  movsd                ;move data
              pop  esi
              pop  edi
              pop  es                   ;restore registers
      }
}
```

EXAMPLE 4–7

```
//C++ version of Example 4-6
//
void TransferBlocks (int blockSize, int* blockA, int* blockB)
{
      for (int a = 0; a < blockSize; a++)
      {
              blockA = blockB++;
              blockA++;
      }
}
```

EXAMPLE 4–8

```
void TransferBlocks(int blockSize, int* blockA, int* blockB)
{
004136A0  push          ebp
004136A1  mov           ebp,esp
004136A3  sub           esp,0D8h
004136A9  push          ebx
004136AA  push          esi
004136AB  push          edi
004136AC  push          ecx
004136AD  lea           edi,[ebp-0D8h]
004136B3  mov           ecx,36h
004136B8  mov           eax,0CCCCCCCCh
004136BD  rep stos      dword ptr [edi]
004136BF  pop           ecx
004136C0  mov           dword ptr [ebp-8],ecx
          for( int a = 0; a < blockSize; a++ )
004136C3  mov           dword ptr [a],0
004136CA  jmp           TransferBlocks+35h (4136D5h)
004136CC  mov           eax,dword ptr [a]
004136CF  add           eax, 1
004136D2  mov           dword ptr [a],eax
004136D5  mov           eax,dword ptr [a]
004136D8  cmp           eax,dword ptr [blockSize]
004136DB  jge           TransferBlocks+57h (4136F7h)
          {
              blockA = blockB++;
004136DD  mov           eax,dword ptr [blockB]
004136E0  mov           dword ptr [blockA],eax
004136E3  mov           ecx,dword ptr [blockB]
004136E6  add           ecx,4
004136E9  mov           dword ptr [blockB],ecx
          blockA++;
004136EC  mov           eax,dword ptr [blockA]
004136EF  add           eax,4
004136F2  mov           dword ptr [blockA],eax
          }
004136F5  jmp           TransferBlocks+2Ch (4136CCh)
 }
004136F7  pop           edi
004136F8  pop           esi
004136F9  pop           ebx
004136FA  mov           esp,ebp
004136FC  pop           ebp
004136FD  ret           0Ch
```

INS

The INS (input string) instruction (not available on the 8086/8088 microprocessors) transfers a byte, word, or doubleword of data from an I/O device into the extra segment memory location addressed by the DI register. The I/O address is contained in the DX register. This instruction is useful for inputting a block of data from an external I/O device directly into the memory. One application transfers data from a disk drive to memory. Disk drives are often considered and interfaced as I/O devices in a computer system.

As with the prior string instructions, there are three basic forms of the INS. The INSB instruction inputs data from an 8-bit I/O device and stores it in the byte-sized memory location indexed by SI. The INSW instruction inputs 16-bit I/O data and stores it in a word-sized memory location. The INSD instruction inputs a doubleword. These instructions can be repeated using the REP prefix, which allows an entire block of input data to be stored in the memory from an I/O device. Table 4–15 lists the various forms of the INS instruction. Note that in the 64-bit mode there is no 64-bit input, but the memory address is 64 bits and located in RDI for the INS instructions.

TABLE 4–15 Forms of the INS instruction.

Assembly Language	Operation
INSB	ES:[DI] = [DX]; DI = DI ± 1 (byte transferred)
INSW	ES:[DI] = [DX]; DI = DI ± 2 (word transferred)
INSD	ES:[DI] = [DX]; DI = DI ± 4 (doubleword transferred)
INS LIST	ES:[DI] = [DX]; DI = DI ± 1 (if LIST is a byte)
INS DATA4	ES:[DI] = [DX]; DI = DI ± 2 (if DATA4 is a word)
INS DATA5	ES:[DI] = [DX]; DI = DI ± 4 (if DATA5 is a doubleword)

Note: [DX] indicates that DX is the I/O device address. These instructions are not available on the 8086 and 8088 microprocessors.

Example 4–9 shows a sequence of instructions that inputs 50 bytes of data from an I/O device whose address is 03ACH and stores the data in extra segment memory array LISTS. This software assumes that data are available from the I/O device at all times. Otherwise, the software must check to see if the I/O device is ready to transfer data precluding the use of a REP prefix.

EXAMPLE 4–9

```
                    ;Using the REP INSB to input data to a memory array
                    ;
0000 BF 0000 R          MOV DI,OFFSET LISTS    ;address array
0003 BA 03AC            MOV DX,3ACH            ;address I/O
0006 FC                 CLD                    ;auto-increment
0007 B9 0032            MOV CX,50              ;load counter
000A F3/6C             REP INSB               ;input data
```

OUTS

The OUTS (output string) instruction (not available on the 8086/8088 microprocessors) transfers a byte, word, or doubleword of data from the data segment memory location address by SI to an I/O device. The I/O device is addressed by the DX register as it is with the INS instruction. Table 4–16 shows the variations available for the OUTS instruction. In the 64-bit mode for the Pentium 4 and Core2, there is no 64-bit output, but the address in RSI is 64 bits wide.

Example 4–10 shows a short sequence of instructions that transfer data from a data segment memory array (ARRAY) to an I/O device at I/O address 3ACH. This software assumes that the I/O device is always ready for data.

EXAMPLE 4–10

```
                    ;Using the REP OUTSB to output data from a memory array
                    ;
0000 BE 0064 R          MOV SI,OFFSET ARRAY    ;address array
0003 BA 03AC            MOV DX,3ACH            ;address I/O
0006 FC                 CLD                    ;auto-increment
0007 B9 0064            MOV CX,100             ;load counter
000A F3/6E             REP OUTSB              ;output data
```

TABLE 4–16 Forms of the OUTS instruction.

Assembly Language	Operation
OUTSB	[DX] = DS:[SI]; SI = SI ± 1 (byte transferred)
OUTSW	[DX] = DS:[SI]; SI = SI ± 2 (word transferred)
OUTSD	[DX] = DS:[SI]; SI = SI ± 4 (doubleword transferred)
OUTS DATA7	[DX] = DS:[SI]; SI = SI ± 1 (if DATA7 is a byte)
OUTS DATA8	[DX] = DS:[SI]; SI = SI ± 2 (if DATA8 is a word)
OUTS DATA9	[DX] = DS:[SI]; SI = SI ± 4 (if DATA9 is a doubleword)

Note: [DX] indicates that DX is the I/O device address. These instructions are not available on the 8086 and 8088 microprocessors.

4–5   **MISCELLANEOUS DATA TRANSFER INSTRUCTIONS**

Don't be fooled by the term *miscellaneous;* these instructions are used in programs. The data transfer instructions detailed in this section are XCHG, LAHF, SAHF, XLAT, IN, OUT, BSWAP, MOVSX, MOVZX, and CMOV. Because the miscellaneous instructions are not used as often as a MOV instruction, they have been grouped together and presented in this section.

XCHG

The XCHG (exchange) instruction exchanges the contents of a register with the contents of any other register or memory location. The XCHG instruction cannot exchange segment registers or memory-to-memory data. Exchanges are byte-, word-, or doubleword-sized (80386 and above), and they use any addressing mode discussed in Chapter 3, except immediate addressing. Table 4–17 shows some examples of the XCHG instruction. In the 64-bit mode, data sizes may also be 64 bits for the exchange instruction.

The XCHG instruction, using the 16-bit AX register with another 16-bit register, is the most efficient exchange. This instruction occupies 1 byte of memory. Other XCHG instructions require 2 or more bytes of memory, depending on the addressing mode selected.

When using a memory-addressing mode and the assembler, it doesn't matter which operand addresses memory. The XCHG AL,[DI] instruction is identical to the XCHG [DI],AL instruction, as far as the assembler is concerned.

If the 80386 through the Core2 microprocessor is available, the XCHG instruction can exchange doubleword data. For example, the XCHG EAX,EBX instruction exchanges the contents of the EAX register with the EBX register.

LAHF and SAHF

The LAHF and SAHF instructions are seldom used because they were designed as bridge instructions. These instructions allowed 8085 (an early 8-bit microprocessor) software to be translated into 8086 software by a translation program. Because any software that required translation was completed many years ago, these instructions have little application today. The LAHF instruction transfers the rightmost 8 bits of the flag register into the AH register. The SAHF instruction transfers the AH register into the rightmost 8 bits of the flag register.

At times, the SAHF instruction may find some application with the numeric coprocessor. The numeric coprocessor contains a status register that is copied into the AX register with the FSTSW AX instruction. The SAHF instruction is then used to copy from AH into the flag register. The flags are then tested for some of the conditions of the numeric coprocessor. This is detailed in Chapter 14, which explains the operation and programming of the numeric coprocessor. Because LAHF and LAFH are legacy instructions, they do not function in the 64-bit mode and are invalid instructions.

TABLE 4–17   Forms of the XCHG instruction.

Assembly Language	Operation
XCHG AL,CL	Exchanges the contents of AL with CL
XCHG CX,BP	Exchanges the contents of CX with BP
XCHG EDX,ESI	Exchanges the contents of EDX with ESI
XCHG AL,DATA2	Exchanges the contents of AL with data segment memory location DATA2
XCHG RBX,RCX	Exchange the contents of RBX with RCX (64-bit mode)
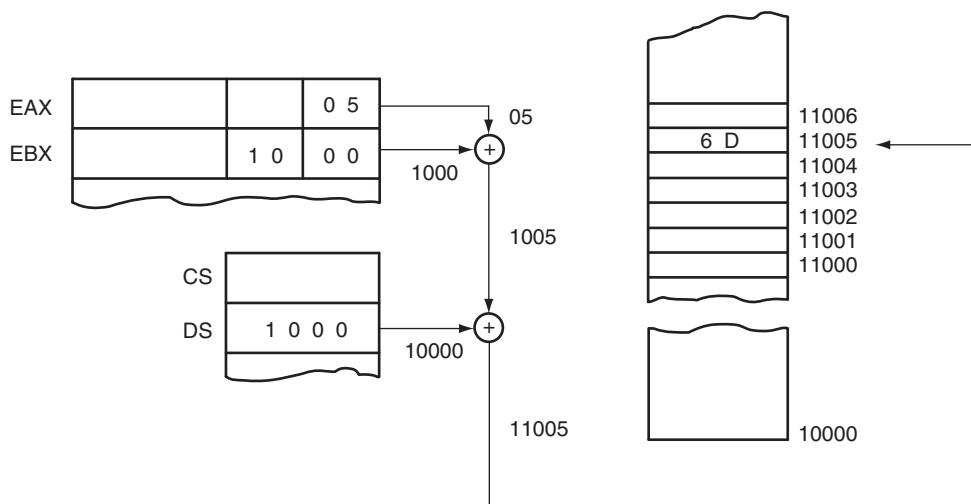
FIGURE 4–19   The operation of the XLAT instruction at the point just before 6DH is loaded into AL.

XLAT

The XLAT (translate) instruction converts the contents of the AL register into a number stored in a memory table. This instruction performs the direct table lookup technique often used to convert one code to another. An XLAT instruction first adds the contents of AL to BX to form a memory address within the data segment. It then copies the contents of this address into AL. This is the only instruction that adds an 8-bit number to a l6-bit number.

Suppose that a 7-segment LED display lookup table is stored in memory at address TABLE. The XLAT instruction then uses the lookup table to translate the BCD number in AL to a 7-segment code in AL. Example 4–11 provides a sequence of instructions that converts from a BCD code to a 7-segment code. Figure 4–19 shows the operation of this example program if TABLE = 1000H, DS = 1000H, and the initial value of AL = 05H (5 BCD). After the translation, AL = 6DH.

EXAMPLE 4–11

```
                    TABLE  DB   3FH, 06H, 5BH, 4FH    ;lookup table
                           DB   66H, 6DH, 7DH, 27H
                           DB   7FH, 6FH

0017 B0 05          LOOK:  MOV AL,5                ;load AL with 5 (a test number)
0019 BB 1000 R             MOV BX,OFFSET TABLE     ;address lookup table
001C D7                    XLAT                    ;convert
```

IN and OUT

Table 4–18 lists the forms of the IN and OUT instructions, which perform I/O operations. Notice that the contents of AL, AX, or EAX are transferred only between the I/O device and the microprocessor. An IN instruction transfers data from an external I/O device into AL, AX, or EAX; an OUT transfers data from AL, AX, or EAX to an external I/O device. (Note that only the 80386 and above contain EAX.)

Two forms of I/O device (port) addressing exist for IN and OUT: fixed port and variable port. *Fixed-port addressing* allows data transfer between AL, AX, or EAX using an 8-bit I/O port

TABLE 4–18   In and OUT
instructions.

Assembly Language	Operation
IN AL,p8	8 bits are input to AL from I/O port p8
IN AX,p8	16 bits are input to AX from I/O port p8
IN EAX,p8	32 bits are input to EAX from I/O port p8
IN AL,DX	8 bits are input to AL from I/O port DX
IN AX,DX	16 bits are input to AX from I/O port DX
IN EAX,DX	32 bits are input to EAX from I/O port DX
OUT p8,AL	8 bits are output to I/O port p8 from AL
OUT p8,AX	16 bits are output to I/O port p8 from AX
OUT p8,EAX	32 bits are output to I/O port p8 from EAX
OUT DX,AL	8 bits are output to I/O port DX from AL
OUT DX,AX	16 bits are output to I/O port DX from AX
OUT DX,EAX	32 bits are output to I/O port DX from EAX

Note: p8 = an 8-bit I/O port number (0000H to 00FFH) and DX = the 16-bit I/O
port number (0000H to FFFFH) held in register DX.

address. It is called *fixed-port addressing* because the port number follows the instruction's
opcode, just as it did with immediate addressing. Often, instructions are stored in ROM. A fixed-
port instruction stored in ROM has its port number permanently fixed because of the nature of
read-only memory. A fixed-port address stored in RAM can be modified, but such a modification
does not conform to good programming practices.

The port address appears on the address bus during an I/O operation. For the 8-bit
fixed-port I/O instructions, the 8-bit port address is zero-extended into a 16-bit address. For
example, if the IN AL,6AH instruction executes, data from I/O address 6AH are input to AL.
The address appears as a 16-bit 006AH on pins A0–A15 of the address bus. Address bus
bits A16–A19 (8086/8088), A16–A23 (80286/80386SX), A16–A24 (80386SL/80386SLC/
80386EX), or A16–A31 (80386–Core2) are undefined for an IN or OUT instruction. Note
that Intel reserves the last 16 I/O ports (FFF0H–FFFFH) for use with some of its peripheral
components.

Variable-port addressing allows data transfers between AL, AX, or EAX and a 16-bit port
address. It is called *variable-port addressing* because the I/O port number is stored in register DX,
which can be changed (varied) during the execution of a program. The 16-bit I/O port address
appears on the address bus pin connections A0–A15. The IBM PC uses a 16-bit port address to
access its I/O space. The ISA bus I/O space for a PC is located at I/O port 0000H–03FFH. Note
that PCI bus cards may use I/O addresses above 03FFH.

Figure 4–20 illustrates the execution of the OUT 19H,AX instruction, which transfers
the contents of AX to I/O port 19H. Notice that the I/O port number appears as a 0019H on
the 16-bit address bus and that the data from AX appears on the data bus of the microproces-
sor. The system control signal $\overline{\text{IOWC}}$ (I/O write control) is a logic zero to enable the I/O
device.

A short program that clicks the speaker in the personal computer appears in Example
4–12. The speaker (in DOS only) is controlled by accessing I/O port 61H. If the rightmost
2 bits of this port are set (11) and then cleared (00), a click is heard on the speaker. Note that
this program uses a logical OR instruction to set these 2 bits and a logical AND instruction to
clear them. These logic operation instructions are described in Chapter 5. The MOV
CX,8000H instruction, followed by the LOOP L1 instruction, is used as a time delay. If the
count is increased, the click will become longer; if shortened, the click will become shorter.
To obtain a series of clicks that can be heard, the program must be modified to repeat many
times.

Microprocessor-based system
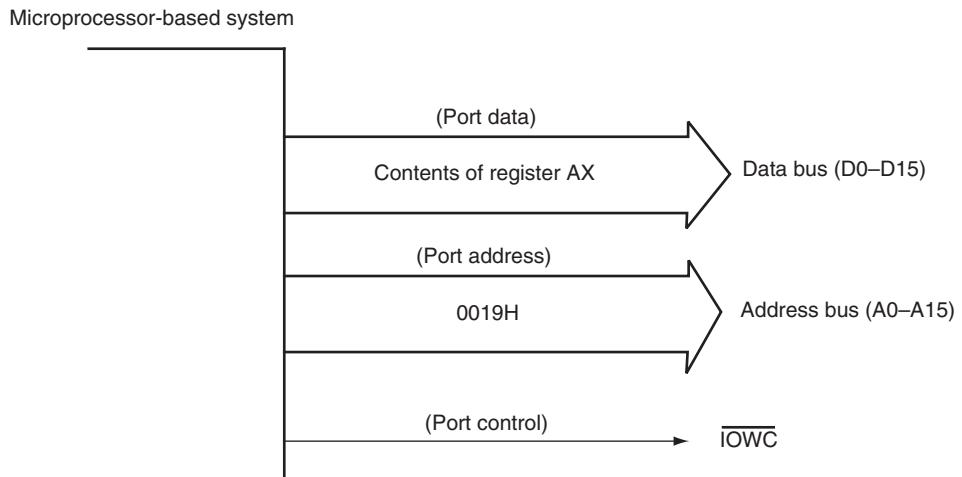
FIGURE 4–20   The signals found in the microprocessor-based system for an OUT 19H,AX instruction.

EXAMPLE 4–12

```
                    .MODEL TINY        ;select tiny model
0000                .CODE              ;start code segment
                    .STARTUP           ;start program
0100 E4 61              IN AL,61H      ;read I/O port 61H
0102 0C 03              OR AL,3        ;set rightmost two bits
0104 E6 61              OUT 61H,AL     ;speaker on
0106 B9 8000            MOV CX,8000H   ;load delay count
0109                L1:
0109 E2 FE              LOOP L1        ;time delay
010B E4 61              IN AL,61H      ;speaker off
010D 24 FC              AND AL,0FCH
010F E6 61              OUT 61H,AL
                    .EXIT
                    END
```

MOVSX and MOVZX

The MOVSX (**move and sign-extend**) and MOVZX (**move and zero-extend**) instructions are found in the 80386–Pentium 4 instruction sets. These instructions move data, and at the same time either sign- or zero-extend it. Table 4–19 illustrates these instructions with several examples of each.

When a number is zero-extended, the most significant part fills with zeros. For example, if an 8-bit 34H is zero-extended into a 16-bit number, it becomes 0034H. Zero-extension is often used to convert unsigned 8- or 16-bit numbers into unsigned 16- or 32-bit numbers by using the MOVZX instruction.

A number is sign-extended when its sign-bit is copied into the most significant part. For example, if an 8-bit 84H is sign-extended into a 16-bit number, it becomes FF84H. The sign-bit of an 84H is a 1, which is copied into the most significant part of the sign-extended result. Sign-extension is most often used to convert 8- or 16-bit signed numbers into 16- or 32-bit signed numbers by using the MOVSX instruction.

BSWAP

The BSWAP (**byte swap**) instruction is available only in the 80486–Pentium 4 microprocessors. This instruction takes the contents of any 32-bit register and swaps the first byte with the fourth, and the second with the third. For example, the BSWAP EAX instruction with

TABLE 4–19 The MOVSX and MOVZX instructions.

Assembly Language	Operation
MOVSX CX,BL	Sign-extends BL into CX
MOVSX ECX,AX	Sign-extends AX into ECX
MOVSX BX,DATA1	Sign-extends the byte at DATA1 into BX
MOVSX EAX,[EDI]	Sign-extends the word at the data segment memory location addressed by EDI into EAX
MOVSX RAX,[RDI]	Sign-extends the doubleword at address RDI into RAX (64-bit mode)
MOVZX DX,AL	Zero-extends AL into DX
MOVZX EBP,DI	Zero-extends DI into EBP
MOVZX DX,DATA2	Zero-extends the byte at DATA2 into DX
MOVZX EAX,DATA3	Zero-extends the word at DATA3 into EAX
MOVZX RBX,ECX	Zero-extends ECX into RBX

EAX = 00112233H swaps bytes in EAX, resulting in EAX = 33221100H. Notice that the order of all 4 bytes is reversed by this instruction. This instruction is used to convert data between the big and little endian forms. In 64-bit operation for the Pentium 4, all 8 bytes in the selected operand are swapped.

CMOV

The CMOV (**conditional move**) class of instruction is new to the Pentium Pro–Core2 instruction sets. There are many variations of the CMOV instruction. Table 4–20 lists these variations

TABLE 4–20 The conditional move instructions.

Assembly Language	Flag(s) Tested	Operation
CMOVB	C = 1	Move if below
CMOVAE	C = 0	Move if above or equal
CMOVBE	Z = 1 or C = 1	Move if below or equal
CMOVA	Z = 0 and C = 0	Move of above
CMOVE or CMOVZ	Z = 1	Move if equal or move if zero
CMOVNE or CMOVNZ	Z = 0	Move if not equal or move if not zero
CMOVL	S ! = O	Move if less than
CMOVLE	Z = 1 or S ! = O	Move if less than or equal
CMOVG	Z = 0 and S = O	Move if greater than
CMOVGE	S = O	Move if greater than or equal
CMOVS	S = 1	Move if sign (negative)
CMOVNS	S = 0	Move if no sign (positive)
CMOVC	C = 1	Move if carry
CMOVNC	C = 0	Move if no carry
CMOVO	O = 1	Move if overflow
CMOVNO	O = 0	Move if no overflow
CMOVP or CMOVPE	P = 1	Move if parity or move if parity even
CMOVNP or CMOVPO	P = 0	Move if no parity or move if parity odd

TABLE 4–21 Instructions that include segments override prefixes.

Assembly Language	Segment Accessed	Default Segment
MOV AX,DS:[BP]	Data	Stack
MOV AX,ES:[BP]	Extra	Stack
MOV AX,SS:[DI]	Stack	Data
MOV AX,CS:LIST	Code	Data
MOV ES:[SI],AX	Extra	Data
LODS ES:DATA1	Extra	Data
MOV EAX,FS:DATA2	FS	Data
MOV GS:[ECX],BL	GS	Data

of CMOV. These instructions move the data only if the condition is true. For example, the CMOVZ instruction moves data only if the result from some prior instruction was a zero. The destination is limited to only a 16- or 32-bit register, but the source can be a 16- or 32-bit register or memory location.

Because this is a new instruction, you cannot use it with the assembler unless the .686 switch is added to the program.

4–6   SEGMENT OVERRIDE PREFIX

The **segment override prefix**, which may be added to almost any instruction in any memory-addressing mode, allows the programmer to deviate from the default segment. The segment override prefix is an additional byte that appends the front of an instruction to select an alternate segment register. About the only instructions that cannot be prefixed are the jump and call instructions that must use the code segment register for address generation. The segment override is also used to select the FS and GS segments in the 80386 through the Core2 microprocessors.

For example, the MOV AX,[DI] instruction accesses data within the data segment by default. If required by a program, this can be changed by prefixing the instruction. Suppose that the data are in the extra segment instead of in the data segment. This instruction addresses the extra segment if changed to MOV AX,ES:[DI].

Table 4–21 shows some altered instructions that address different memory segments that are different from normal. Each time an instruction is prefixed with a segment override prefix, the instruction becomes 1 byte longer. Although this is not a serious change to the length of the instruction, it does add to the instruction's execution time. It is usually customary to limit the use of the segment override prefix and remain in the default segments so that shorter and more efficient software is written.

4–7   ASSEMBLER DETAIL

The assembler (MASM)[1] for the microprocessor can be used in two ways: (1) with models that are unique to a particular assembler, and (2) with full-segment definitions that allow complete control over the assembly process and are universal to all assemblers. This section of the text pre-

---

[1]The assembler used throughout this text is the Microsoft MACRO assembler called MASM, version 6.1X.

sents both methods and explains how to organize a program's memory space by using the assembler. It also explains the purpose and use of some of the more important directives used with this assembler. Appendix A provides additional detail about the assembler.

In most cases, the inline assembler found in Visual C++ is used for developing assembly code for use in a C++ program, but there are occasions that require separate assembly modules writing using the assembler. This section of the text contrasts, where possible, the inline assembler and the assembler.

Directives

Before the format of an assembly language program is discussed, some details about the directives (**pseudo-operations**) that control the assembly process must be learned. Some common assembly language directives appear in Table 4–22. **Directives** indicate how an operand or section of a program is to be processed by the assembler. Some directives generate and store information in the memory; others do not. The DB (define byte) directive stores bytes of data in the memory, whereas the BYTE PTR directive never stores data. The BYTE PTR directive indicates the size of the data referenced by a pointer or index register. Note that none of the directives function in the inline assembler program that is a part of Visual C++. If you are using the inline assembler exclusively, you can skip this part of the text. Be aware that complex sections of assembly code are still written using MASM.

Note that by default the assembler accepts only 8086/8088 instructions, unless a program is preceded by the .686 or .686P directive or one of the other microprocessor selection switches. The .686 directive tells the assembler to use the Pentium Pro instruction set in the real mode, and the .686P directive tells the assembler to use the Pentium Pro protected mode instruction set. Most modern software is written assuming that the microprocessor is a Pentium Pro or newer, so the .686 switch is often used. Windows 95 was the first major operating system to use a 32-bit architecture that conforms to the 80386. Windows XP requires a Pentium class machine (.586 switch) using at least a 233MHz microprocessor.

Storing Data in a Memory Segment.   The DB (**define byte**), DW (**define word**), and DD (**define doubleword**) directives, first presented in Chapter 1, are most often used with MASM to define and store memory data. If a numeric coprocessor executes software in the system, the DQ (**define quadword**) and DT (**define ten bytes**) directives are also common. These directives label a memory location with a symbolic name and indicate its size.

Example 4–13 shows a memory segment that contains various forms of data definition directives. It also shows the full-segment definition with the first SEGMENT statement to indicate the start of the segment and its symbolic name. Alternately, as in past examples in this and prior chapters, the SMALL model can be used with the .DATA statement. The last statement in this example contains the ENDS directive, which indicates the end of the segment. The name of the segment (LIST_SEG) can be anything that the programmer desires to call it. This allows a program to contain as many segments as required.

EXAMPLE 4–13

```
                ;Using the DB, DW, and DD directives
                ;
0000            LIST_SEG    SEGMENT

0000 01 02 03   DATA1 DB    1,2,3           ;define bytes
0003 45               DB    45H             ;hexadecimal
0004 41               DB    'A'             ;ASCII
0005 F0               DB    11110000B       ;binary
0006 000C 000D  DATA2 DW    12,13           ;define words
000A 0200             DW    LIST1           ;symbolic
```

TABLE 4–22   Common MASM directives.

Directive	Function
.286	Selects the 80286 instruction set
.286P	Selects the 80286 protected mode instruction set
.386	Selects the 80386 instruction set
.386P	Selects the 80386 protected mode instruction set
.486	Selects the 80486 instruction set
.486P	Selects the 80498 protected mode instruction set
.586	Selects the Pentium instruction set
.586P	Selects the Pentium protected mode instruction set
.686	Selects the Pentium Pro–Core2 instruction set
.686P	Selects the Pentium Pro–Core2 protected mode instruction set
.287	Selects the 80287 math coprocessor
.387	Selects the 80387 math coprocessor
.CODE	Indicates the start of the code segment (models only)
.DATA	Indicates the start of the data segment (models only)
.EXIT	Exits to DOS (models only)
.MODEL	Selects the programming model
.STACK	Selects the start of the stack segment (models only)
.STARTUP	Indicates the starting instruction in a program (models only)
ALIGN n	Align to boundary n (n = 2 for words, n = 4 for doublewords)
ASSUME	Informs the assembler to name each segment (full segments only)
BYTE	Indicates byte-sized as in BYTE PTR
DB	Defines byte(s) (8 bits)
DD	Defines doubleword(s) (32 bits)
DQ	Defines quadwords(s) (64 bits)
DT	Defines ten byte(s) (80 bits)
DUP	Generates duplicates
DW	Define word(s) (16 bits)
DWORD	Indicates doubleword-sized, as in DWORD PTR
END	Ends a program file
ENDM	Ends a MACRO sequence
ENDP	Ends a procedure
ENDS	Ends a segment or data structure
EQU	Equates data or a label to a label
FAR	Defines a far pointer, as in FAR PTR
MACRO	Designates the start of a MACRO sequence
NEAR	Defines a near pointer, as in NEAR PTR
OFFSET	Specifies an offset address
ORG	Sets the origin within a segment
OWORD	Indicates octalwords, as in OWORD PTR
PROC	Starts a procedure
PTR	Designates a pointer
QWORD	Indicates quadwords, as in QWORD PTR
SEGMENT	Starts a segment for full segments
STACK	Starts a stack segment for full segments
STRUC	Defines the start of a data structure
USES	Automatically pushes and pops registers
USE16	Uses 16-bit instruction mode
USE32	Uses 32-bit instruction mode
WORD	Indicates word-sized, as in WORD PTR

```
000C 2345                  DW   2345H          ;hexadecimal
000E 00000300      DATA3 DD   300H             ;define doubleword
0012 4007DF3B            DD   2.123            ;real
0016 544269E1            DD   3.34E+12         ;real
001A 00              LISTA DB   ?              ;reserve 1 byte
001B 000A[           LISTB DB   10 DUP(?)      ;reserve 10 bytes
        ??
          ]
0025 00                   ALIGN 2              ;set word boundary
0026 0100[
        0000
          ]     LISTC DW  100H DUP(0)     ;reserve 100H words
0226 0016[           LISTD DD  22 DUP(?)       ;reserve 22 doublewords
      ???????
              ]
027E 0064[           SIXES DB  100 DUP(6)      ;reserve 100 bytes
        06
          ]
02E2                 LIST_SEG    ENDS
```

Example 4–13 shows various forms of data storage for bytes at DATA1. More than 1 byte can be defined on a line in binary, hexadecimal, decimal, or ASCII code. The DATA2 label shows how to store various forms of word data. Doublewords are stored at DATA3; they include floating-point, single-precision real numbers.

Memory is reserved for use in the future by using a question mark (?) as an operand for a DB, DW, or DD directive. When a ? is used in place of a numeric or ASCII value, the assembler sets aside a location and does not initialize it to any specific value. (Actually, the assembler usually stores a zero into locations specified with a?.) The DUP (duplicate) directive creates an array, as shown in several ways in Example 4–12. A 10 DUP (?) reserves 10 locations of memory, but stores no specific value in any of the 10 locations. If a number appears within the ( ) part of the DUP statement, the assembler initializes the reserved section of memory with the data indicated. For example, the LIST2 DB 10 DUP (2) instruction reserves 10 bytes of memory for array LIST2 and initializes each location with a 02H.

The ALIGN directive, used in this example, makes sure that the memory arrays are stored on word boundaries. An ALIGN 2 places data on *word boundaries* and an ALIGN 4 places them on *doubleword boundaries*. In the Pentium–Pentium 4, quadword data for double-precision floating-point numbers should use ALIGN 8. It is important that word-sized data are placed at word boundaries and doubleword-sized data are placed at doubleword boundaries. If not, the microprocessor spends additional time accessing these data types. A word stored at an odd-numbered memory location takes twice as long to access as a word stored at an even-numbered memory location. Note that the ALIGN directive cannot be used with memory models because the size of the model determines the data alignment. If all doubleword data are defined first, followed by word-sized and then byte-sized data, the ALIGN statement is not necessary to align data correctly.

ASSUME, EQU, and ORG.   The equate directive (EQU) equates a numeric, ASCII, or label to another label. Equates make a program clearer and simplify debugging. Example 4–14 shows several equate statements and a few instructions that show how they function in a program.

EXAMPLE 4–14

```
          ;Using equate directive
          ;
= 000A      TEN    EQU 10
= 0009      NINE   EQU 9

0000 B0 0A          MOV AL,TEN
0002 04 09          ADD AL,NINE
```

The THIS directive always appears as THIS BYTE, THIS WORD, THIS DWORD, or THIS QWORD. In certain cases, data must be referred to as both a byte and a word. The assembler can only assign either a byte, word, or doubleword address to a label. To assign a byte label to a word, use the software listed in Example 4–15.

EXAMPLE 4–15

```
                        ;Using the THIS and ORG directives
                        ;
0000                    DATA_SEG    SEGMENT

0300                        ORG   300H

= 0300                  DATA1  EQU   THIS BYTE
0300                    DATA2  DW    ?
0302                    DATA_SEG    ENDS

0000                    CODE_SEG    SEGMENT 'CODE'
                            ASSUME CS:CODE_SEG, DS:DATA_SEG
0000 8A 1E 0300 R           MOV BL,DATA1
0004 A1 0300 R              MOV AX,DATA2
0007 8A 3E 0301 R           MOV BH,DATA1+1

000B                    CODE_SEG    ENDS
```

This example also illustrates how the ORG (origin) statement changes the starting offset address of the data in the data segment to location 300H. At times, the origin of data or the code must be assigned to an absolute offset address with the ORG statement. The ASSUME statement tells the assembler what names have been chosen for the code, data, extra, and stack segments. Without the ASSUME statement, the assembler assumes nothing and automatically uses a segment override prefix on all instructions that address memory data. The ASSUME statement is only used with full-segment definitions, as described later in this section of the text.

PROC and ENDP.   The PROC and ENDP directives indicate the start and end of a procedure (subroutine). These directives *force structure* because the procedure is clearly defined. Note that if structure is to be violated for whatever reason, use the CALLF, CALLN, RETF, and RETN instructions. Both the PROC and ENDP directives require a label to indicate the name of the procedure. The PROC directive, which indicates the start of a procedure, must also be followed with a NEAR or FAR. A NEAR procedure is one that resides in the same code segment as the program. A FAR procedure may reside at any location in the memory system. Often the call NEAR procedure is considered to be *local,* and the call FAR procedure is considered to be *global.* The term *global* denotes a procedure that can be used by any program; *local* defines a procedure that is only used by the current program. Any labels that are defined within the procedure block are also defined as either local (NEAR) or global (FAR).

Example 4–16 shows a procedure that adds BX, CX, and DX and stores the sum in register AX. Although this procedure is short and may not be particularly useful, it does illustrate how to use the PROC and ENDP directives to delineate the procedure. Note that information about the operation of the procedure should appear as a grouping of comments that show the registers changed by the procedure and the result of the procedure.

EXAMPLE 4–16

```
                        ;A procedure that adds BX, CX, and DX with the
                        ;sum stored in AX
                        ;
0000                    ADDEM PROC  FAR         ;start of procedure
```

```
0000 03 D9                    ADD    BX,CX
0002 03 DA                    ADD    BX,DX
0004 8B C3                    MOV    AX,BX
0006 CB                       RET

0007                ADDEM ENDP                   ;end of procedure
```

If version 6.x of the Microsoft MASM assembler program is available, the PROC directive specifies and automatically saves any registers used within the procedure. The USES statement indicates which registers are used by the procedure, so that the assembler can automatically save them before your procedure begins and restore them before the procedure ends with the RET instruction. For example, the ADDS PROC USES AX BX CX statement automatically pushes AX, BX, and CX on the stack before the procedure begins and pops them from the stack before the RET instruction executes at the end of the procedure. Example 4–17 illustrates a procedure written using MASM version 6.x that shows the USES statement. Note that the registers in the list are not separated by commas, but by spaces, and the PUSH and POP instructions are displayed in the procedure listing because it was assembled with the .LIST ALL directive. The instructions prefaced with an asterisk (*) are inserted by the assembler and were not typed in the source file. The USES statement appears elsewhere in this text, so if MASM version 5.10 is in use, the code will need to be modified.

EXAMPLE 4–17

```
                             ;A procedure that includes the USES directive to
                             ;save BX, CX, and DX on the stack and restore them
                             ;before the return instruction.
0000                ADDS    PROC    NEAR    USES BX CX DX

0000 53       *             push    bx
0001 51       *             push    cx
0002 52       *             push    dx
0003 03 D8                  ADD     BX,AX
0005 03 CB                  ADD     CX,BX
0007 03 D1                  ADD     DX,CX
0009 8B C2                  MOV     AX,DX
                            RET
000B 5A       *             pop     dx
000C 59       *             pop     cx
000D 5B       *             pop     bx
000E C3       *             ret     0000h

000F                ADDS    ENDP
```

Memory Organization

The assembler uses two basic formats for developing software: One method uses models and the other uses full-segment definitions. Memory models, as presented in this section and briefly in earlier chapters, are unique to the MASM assembler program. The TASM assembler also uses memory models, but they differ somewhat from the MASM models. The full-segment definitions are common to most assemblers, including the Intel assembler, and are often used for software development. The models are easier to use for simple tasks. The full-segment definitions offer better control over the assembly language task and are recommended for complex programs. The model was used in early chapters because it is easier to understand for the beginning programmer. Models are also used with assembly language procedures that are used by high-level languages such as C/C++. Although this text fully develops and uses the memory model definitions for its programming examples, realize that full-segment definitions offer some advantages over memory models, as discussed later in this section.

Models.   There are many models available to the MASM assembler, ranging from tiny to huge. Appendix A contains a table that lists all the models available for use with the assembler. To designate a model, use the .MODEL statement followed by the size of the memory system. The TINY model requires that all software and data fit into one 64K-byte memory segment; it is useful for many small programs. The SMALL model requires that only one data segment be used with one code segment for a total of 128K bytes of memory. Other models are available, up to the HUGE model.

Example 4–18 illustrates how the .MODEL statement defines the parameters of a short program that copies the contents of a 100-byte block of memory (LISTA) into a second 100-byte block of memory (LISTB). It also shows how to define the stack, data, and code segments. The .EXIT 0 directive returns to DOS with an error code of 0 (no error). If no parameter is added to .EXIT, it still returns to DOS, but the error code is not defined. Also note that special directives such as @DATA (see Appendix A) are used to identify various segments. If the .STARTUP directive is used (MASM version 6.x), the MOV AX,@DATA followed by MOV DS,AX statements can be eliminated. The .STARTUP directive also eliminates the need to store the starting address next to the END label. Models are important with both Microsoft Visual C++ and Borland C++ development systems if assembly language is included with C++ programs. Both development systems use inline assembly programming for adding assembly language instructions and require an understanding of programming models.

EXAMPLE 4–18

```
                          .MODEL SMALL          ;select small model
                          .STACK 100H           ;define stack
                          .DATA                 ;start data segment

0000 0064[        LISTA  DB    100 DUP(?)
        ??
          ]
0064 0064[        LISTB  DB    100 DUP(?)
        ??
          ]

                          .CODE                 ;start code segment

0000 B9 —— ?      HERE:   MOV   AX,@DATA         ;load ES and DS
0003 8E C0                MOV   ES,AX
0005 8E D8                MOV   DS,AX
0007 FC                   CLD                    ;move data
0008 BE 0000 R            MOV   SI,OFFSET LISTA
000B BF 0064 R            MOV   DI,OFFSET LISTB
000E B9 0064              MOV   CX,100
0011 F3/A4                REP   MOVSB

0013                      .EXIT 0                ;exit to DOS
                          END HERE
```

Full-Segment Definitions.   Example 4–19 illustrates the same program using full segment definitions. Full-segment definitions are also used with the Borland and Microsoft C/C++ environments for procedures developed in assembly language. The program in Example 4–19 appears longer than the one pictured in Example 4–18, but it is more structured than the model method of setting up a program. The first segment defined is the STACK_SEG, which is clearly delineated with the SEGMENT and ENDS directives. Within these directives, a DW 100 DUP (?) sets aside 100H words for the stack segment. Because the word STACK appears next to SEGMENT, the assembler and linker automatically load both the stack segment register (SS) and stack pointer (SP).

EXAMPLE 4–19

```
0000                    STACK_SEG     SEGMENT      'STACK'
0000 0064[              DW    100H DUP(?)
        ????
                ]
0200                    STACK_SEG     ENDS

0000                    DATA_SEG      SEGMENT      'DATA'
0000 0064[              LISTA   DB    100 DUP(?)
        ??
                ]
0064 0064[              LISTB   DB    100 DUP(?)
        ??
                ]
00CB                    DATA_SEG      ENDS

0000                    CODE_SEG      SEGMENT      'CODE'
                        ASSUME CS:CODE_SEG, DS:DATA_SEG
                        ASSUME SS:STACK_SEG
0000                    MAIN    PROC   FAR
0000 B8 —— R            MOV     AX,DATA_SEG          ;load DS and ES
0003 8E C0              MOV     ES,AX
0005 8E D8              MOV     DS,AX
0007 FC                 CLD                          ;save data
0008 BE 0000 R          MOV     SI,OFFSET LISTA
000B BF 0064 R          MOV     DI,OFFSET LISTB
000E B9 0064            MOV     CX,100
0011 F3/A4              REP     MOVSB
0013 B4 4C              MOV     AH,4CH               ;exit to DOS
0015 CD 21              INT     21H
0017                    MAIN    ENDP
0017                    CODE_SEG      ENDS
                        END     MAIN
```

Next, the data are defined in the DATA_SEG. Here, two arrays of data appear as LISTA and LISTB. Each array contains 100 bytes of space for the program. The names of the segments in this program can be changed to any name. Always include the group name 'DATA', so that the Microsoft program CodeView can be effectively used to symbolically debug this software. CodeView is a part of the MASM package used to debug software. To access CodeView, type CV, followed by the file name at the DOS command line; if operating from Programmer's WorkBench, select Debug under the Run menu. If the group name is not placed in a program, CodeView can still be used to debug a program, but the program will not be debugged in symbolic form. Other group names such as 'STACK', 'CODE', and so forth are listed in Appendix A. You must at least place the word 'CODE' next to the code segment SEGMENT statement if you want to view the program symbolically in CodeView.

The CODE_SEG is organized as a far procedure because most software is procedure-oriented. Before the program begins, the code segment contains the ASSUME statement. The ASSUME statement tells the assembler and linker that the name used for the code segment (CS) is CODE_SEG; it also tells the assembler and linker that the data segment is DATA_SEG and the stack segment is STACK_SEG. Notice that the group name 'CODE' is used for the code segment for use by CodeView. Other group names appear in Appendix A with the models.

After the program loads both the extra segment register and data segment register with the location of the data segment, it transfers 100 bytes from LISTA to LISTB. Following this is a sequence of two instructions that return control back to DOS (the disk operating system). Note that the program loader does not automatically initialize DS and ES. These registers must be loaded with the desired segment addresses in the program.

The last statement in the program is END MAIN. The END statement indicates the end of the program and the location of the first instruction executed. Here, we want the machine to execute the main procedure so the MAIN label follows the END directive.

In the 80386 and above, an additional directive is found attached to the code segment. The USE16 or USE32 directive tells the assembler to use either the 16- or 32-bit instruction modes for the microprocessor. Software developed for the DOS environment must use the USE16 directive for the 80386 through the Core2 program to function correctly because MASM assumes that all segments are 32 bits and all instruction modes are 32 bits by default.

A Sample Program

Example 4–20 provides a sample program, using full-segment definitions, that reads a character from the keyboard and displays it on the CRT screen. Although this program is trivial, it illustrates a complete workable program that functions on any personal computer using DOS, from the earliest 8088-based system to the latest Core2-based system. This program also illustrates the use of a few DOS function calls. (Appendix A lists the DOS function calls with their parameters.) The BIOS function calls allow the use of the keyboard, printer, disk drives, and everything else that is available in your computer system.

This example program uses only a code segment because there is no data. A stack segment should appear, but it has been left out because DOS automatically allocates a l28-byte stack for all programs. The only time that the stack is used in this example is for the INT 21H instructions that call a procedure in DOS. Note that when this program is linked, the linker signals that no stack segment is present. This warning may be ignored in this example because the stack is fewer than 128 bytes.

Notice that the entire program is placed into a far procedure called MAIN. It is good programming practice to write all software in procedural form, which allows the program to be used as a procedure at some future time if necessary. It is also fairly important to document register use and any parameters required for the program in the program header, which is a section of comments that appear at the start of the program.

The program uses DOS functions 06H and 4CH. The function number is placed in AH before the INT 21H instruction executes. The 06H function reads the keyboard if DL = 0FFH, or displays the ASCII contents of DL if it is not 0FFH. Upon close examination, the first section of the program moves 06H into AH and 0FFH into DL, so that a key is read from the keyboard. The INT 21H tests the keyboard; if no key is typed, it returns equal. The JE instruction tests the equal condition and jumps to MAIN if no key is typed.

When a key is typed, the program continues to the next step, which compares the contents of AL with an @ symbol. Upon return from the INT 21H, the ASCII character of the typed key is found in AL. In this program, if an @ symbol is typed, the program ends. If the @ symbol is not typed, the program continues by displaying the character typed on the keyboard with the next INT 21H instruction.

The second INT 21H instruction moves the ASCII character into DL so it can be displayed on the CRT screen. After displaying the character, a JMP executes. This causes the program to continue at MAIN, where it repeats reading a key.

If the @ symbol is typed, the program continues at MAIN1, where it executes the DOS function code number 4CH. This causes the program to return to the DOS prompt so that the computer can be used for other tasks.

More information about the assembler and its application appears in Appendix A and in the next several chapters. Appendix A provides a complete overview of the assembler, linker, and DOS functions. It also provides a list of the BIOS (basic I/O system) functions. The information provided in the following chapters clarifies how to use the assembler for certain tasks at different levels of the text.

EXAMPLE 4–20

```
                    ;An example DOS full-segment program that reads a key and
                    ;displays it. Note that an @ key ends the program.
                    ;
0000                CODE_SEG      SEGMENT 'CODE'
                         ASSUME CS:CODE_SEG

0000                MAIN     PROC   FAR

0000 B4 06                   MOV  AH,06H          ;read a key
0002 B2 FF                   MOV  DL,0FFH
0004 CD 21                   INT  21H
0006 74 F8                   JE   MAIN            ;if no key typed
0008 3C 40                   CMP  AL,'@'
000A 74 08                   JE   MAIN1           ;if an @ key
000C B4 06                   MOV  AH,06H          ;display key (echo)
000E 8A D0                   MOV  DL,AL
0010 CD 21                   INT  21H
0012 EB EC                   JMP  MAIN            ;repeat
0014                MAIN1:
0014 B4 4C                   MOV  AH,4CH          ;exit to DOS
0016 CD 21                   INT  21H

0018                MAIN  ENDP
0018                      END  MAIN
```

EXAMPLE 4–21

```
                    ;An example DOS model program that reads a key and displays
                    ;it. Note that an @ key ends the program.
                    ;
                    .MODEL TINY
0000                .CODE
                    .STARTUP

0100                MAIN:
0100 B4 06                   MOV  AH,6            ;read a key
0102 B2 FF                   MOV  DL,0FFH
0104 CD 21                   INT  21H
0106 74 F8                   JE   MAIN            ;if no key typed
0108 3C 40                   CMP  AL, '@'
010A 74 08                   JE   MAIN1           ;if an @ key
010C B4 06                   MOV  AH,06H          ;display key (echo)
010E 8A D0                   MOV  DL,AL
0110 CD 21                   INT  21H
0112 EB EC                   JMP  MAIN            ;repeat
0114                MAIN1:

                    .EXIT                         ;exit to DOS
                    END
```

Example 4–21 shows the program listed in Example 4–20, except models are used instead of full-segment descriptions. Please compare the two programs to determine the differences. Notice how much shorter and cleaner looking the models can make a program.

4–8   SUMMARY

1. Data movement instructions transfer data between registers, a register and memory, a register and the stack, memory and the stack, the accumulator and I/O, and the flags and the stack. Memory-to-memory transfers are allowed only with the MOVS instruction.

2. Data movement instructions include MOV, PUSH, POP, XCHG, XLAT, IN, OUT, LEA, LOS, LES, LSS, LGS, LFS, LAHF, SAHF, and the following string instructions: LODS, STOS, MOVS, INS, and OUTS.

3. The first byte of an instruction contains the opcode. The opcode specifies the operation performed by the microprocessor. The opcode may be preceded by one or more override prefixes in some forms of instructions.

4. The D-bit, located in many instructions, selects the direction of data flow. If D = 0, the data flow from the REG field to the R/M field of the instruction. If D = 1, the data flow from the R/M field to the REG field.

5. The W-bit, found in most instructions, selects the size of the data transfer. If W = 0, the data are byte-sized; if W = 1, the data are word-sized. In the 80386 and above, W = 1 specifies either a word or doubleword register.

6. MOD selects the addressing mode of operation for a machine language instruction's R/M field. If MOD = 00, there is no displacement; if MOD = 01, an 8-bit sign-extended displacement appears; if MOD = 10, a 16-bit displacement occurs; and if MOD = 11, a register is used instead of a memory location. In the 80386 and above, the MOD bits also specify a 32-bit displacement.

7. A 3-bit binary register code specifies the REG and R/M fields when the MOD = 11. The 8-bit registers are AH, AL, BH, BL, CH, CL, DH, and DL. The l6-bit registers are AX, BX, CX, DX, SP, BP, DI, and SI. The 32-bit registers are EAX, EBX, ECX, EDX, ESP, EBP, EDI, and ESI. To access the 64-bit registers, a new prefix is added called the REX prefix that contains a fourth bit for accessing registers R8 through R15.

8. When the R/M field depicts a memory mode, a 3-bit code selects one of the following modes: [BX+DI], [BX+SI], [BP+DI], [BP+SI], [BX], [BP], [DI], or [SI] for 16-bit instructions. In the 80386 and above, the R/M field specifies EAX, EBX, ECX, EDX, EBP, EDI, and ESI or one of the scaled-index modes of addressing memory data. If the scaled-index mode is selected (R/M = 100), an additional byte (scaled-index byte) is added to the instruction to specify the base register, index register, and the scaling factor.

9. By default, all memory-addressing modes address data in the data segment unless BP or EBP addresses memory. The BP or EBP register addresses data in the stack segment.

10. The segment registers are addressed only by the MOV, PUSH, or POP instructions. The instruction may transfer a segment register to a 16-bit register, or vice versa. MOV CS,reg or POP CS instructions are not allowed because they change only a portion of the address. The 80386 through the Pentium 4 include two additional segment registers, FS and GS.

11. Data are transferred between a register or a memory location and the stack by the PUSH and POP instructions. Variations of these instructions allow immediate data to be pushed onto the stack, the flags to be transferred between the stack, and all 16-bit registers can be transferred between the stack and the registers. When data are transferred to the stack, 2 bytes (8086–80286) always move. The most significant byte is placed at the location addressed by SP − 1, and the least significant byte is placed at the location addressed by SP − 2. After placing the data on the stack, SP is decremented by 2. In the 80386–Core2, 4 bytes of data from a memory location or register may also be transferred to the stack.

12. Opcodes that transfer data between the stack and the flags are PUSHF and POPF. Opcodes that transfer all the 16-bit registers between the stack and the registers are PUSHA and POPA. In the 80386 and above, PUSHFD and POPFD transfer the contents of the EFLAGS between the microprocessor and the stack, and PUSHAD and POPAD transfer all the 32-bit registers. The PUSHA and POPA instructions are invalid in the 64-bit mode.

13. LEA, LDS, and LES instructions load a register or registers with an effective address. The LEA instruction loads any 16-bit register with an effective address; LDS and LES load any 16-bit register and either DS or ES with the effective address. In the 80386 and above, additional instructions include LFS, LGS, and LSS, which load a 16-bit register and FS, GS, or SS.

14. String data transfer instructions use either or both DI and SI to address memory. The DI off-set address is located in the extra segment, and the SI offset address is located in the data segment. If the 80386–Core2 is operated in protected mode, ESI and EDI are used with the string instructions.

15. The direction flag (D) chooses the auto-increment or auto-decrement mode of operation for DI and SI for string instructions. To clear D to 0, use the CLD instruction to select the auto-increment mode; to set D to 1, use the STD instruction to select the auto-decrement mode. Either or both DI and SI increment/decrement by 1 for a byte operation, by 2 for a word operation, and by 4 for a doubleword operation.

16. LODS loads AL, AX, or EAX with data from the memory location addressed by SI; STOS stores AL, AX, or EAX in the memory location addressed by DI; and MOVS transfers a byte, a word, or a doubleword from the memory location addressed by SI into the location addressed by DI.

17. INS inputs data from an I/O device addressed by DX and stores it in the memory location addressed by DI. The OUTS instruction outputs the contents of the memory location addressed by SI and sends it to the I/O device addressed by DX.

18. The REP prefix may be attached to any string instruction to repeat it. The REP prefix repeats the string instruction the number of times found in register CX.

19. Arithmetic and logic operators can be used in assembly language. An example is MOV AX,34*3, which loads AX with 102.

20. Translate (XLAT) converts the data in AL into a number stored at the memory location addressed by BX plus AL.

21. IN and OUT transfer data between AL, AX, or EAX and an external I/O device. The address of the I/O device is either stored with the instruction (fixed-port addressing) or in register DX (variable-port addressing).

22. The Pentium Pro–Core2 contain a new instruction called CMOV, or conditional move. This instruction only performs the move if the condition is true.

23. The segment override prefix selects a different segment register for a memory location than the default segment. For example, the MOV AX,[BX] instruction uses the data segment, but the MOV AX,ES:[BX] instruction uses the extra segment because of the ES: override prefix. Using the segment override prefix is the only way to address the FS and GS segments in the 80386 through the Pentium 4.

24. The MOVZX (move and zero-extend) and MOVSX (move and sign-extend) instructions, found in the 80386 and above, increase the size of a byte to a word or a word to a double-word. The zero-extend version increases the size of the number by inserting leading zeros. The sign-extend version increases the size of the number by copying the sign-bit into the more significant bits of the number.

25. Assembler directives DB (define byte), DW (define word), DD (define doubleword), and DUP (duplicate) store data in the memory system.

26. The EQU (equate) directive allows data or labels to be equated to labels.

27. The SEGMENT directive identifies the start of a memory segment and ENDS identifies the end of a segment when full-segment definitions are in use.

28. The ASSUME directive tells the assembler what segment names you have assigned to CS, DS, ES, and SS when full-segment definitions are in effect. In the 80386 and above, ASSUME also indicates the segment name for FS and GS.

29. The PROC and ENDP directives indicate the start and end of a procedure. The USES directive (MASM version 6.x) automatically saves and restores any number of registers on the stack if they appear with the PROC directive.

30. The assembler assumes that software is being developed for the 8086/8088 microprocessor unless the .286, .386, .486, .586, or .686 directive is used to select one of these other micro-processors. This directive follows the .MODEL statement to use the 16-bit instruction mode and precedes it for the 32-bit instruction mode.

31. Memory models can be used to shorten the program slightly, but they can cause problems for larger programs. Also be aware that memory models are not compatible with all assembler programs.

4–9   QUESTIONS AND PROBLEMS

1. The first byte of an instruction is the _____, unless it contains one of the override prefixes.
2. Describe the purpose of the D- and W-bits found in some machine language instructions.
3. In a machine language instruction, what information is specified by the MOD field?
4. If the register field (REG) of an instruction contains 010 and W = 0, what register is selected, assuming that the instruction is a 16-bit mode instruction?
5. How are the 32-bit registers selected for the Pentium 4 microprocessor?
6. What memory-addressing mode is specified by R/M = 001 with MOD = 00 for a 16-bit instruction?
7. Identify the default segment registers assigned to the following:
   (a) SP
   (b) EBX
   (c) DI
   (d) EBP
   (e) SI
8. Convert an 8B07H from machine language to assembly language.
9. Convert an 8B9E004CH from machine language to assembly language.
10. If a MOV SI,[BX+2] instruction appears in a program, what is its machine language equivalent?
11. If a MOV ESI,[EAX] instruction appears in a program for the Core2 microprocessor operating in the 16-bit instruction mode, what is its machine language equivalent?
12. What is the purpose of REX?
13. What is wrong with a MOV CS,AX instruction?
14. Form a short sequence of instructions that load the data segment register with a 1000H.
15. The PUSH and POP instructions always transfer a(n) _____ -bit number between the stack and a register or memory location in the 80386–Core2 microprocessors when operated in the 32-bit mode.
16. Create an instruction that places RAX onto the stack in the 64-bit mode for the Pentium 4.
17. What segment register may not be popped from the stack?
18. Which registers move onto the stack with the PUSHA instruction?
19. Which registers move onto the stack for a PUSHAD instruction?
20. Describe the operation of each of the following instructions:
   (a) PUSH AX
   (b) POP ESI
   (c) PUSH [BX]
   (d) PUSHFD
   (e) POP DS
   (f) PUSHD 4
21. Explain what happens when the PUSH BX instruction executes. Make sure to show where BH and BL are stored. (Assume that SP = 0100H and SS = 0200H.)
22. Repeat question 21 for the PUSH EAX instruction.
23. The 16-bit POP instruction (except for POPA) increments SP by _____.
24. What values appear in SP and SS if the stack is addressed at memory location 02200H?

25. Compare the operation of a MOV DI,NUMB instruction with an LEA DI,NUMB instruction.
26. What is the difference between an LEA SI,NUMB instruction and a MOV SI,OFFSET NUMB instruction?
27. Which is more efficient, a MOV with an OFFSET or an LEA instruction?
28. Describe how the LDS BX,NUMB instruction operates.
29. What is the difference between the LDS and LSS instructions?
30. Develop a sequence of instructions that moves the contents of data segment memory locations NUMB and NUMB+1 into BX, DX, and SI.
31. What is the purpose of the direction flag?
32. Which instructions set and clear the direction flag?
33. Which string instruction(s) use both DI and SI to address memory data?
34. Explain the operation of the LODSB instruction.
35. Explain the operation of the LODSQ instruction for the 64-bit mode of the Pentium 4 or Core2.
36. Explain the operation of the OUTSB instruction.
37. Explain the operation of the STOSW instruction.
38. Develop a sequence of instructions that copy 12 bytes of data from an area of memory addressed by SOURCE into an area of memory addressed by DEST.
39. What does the REP prefix accomplish and what type of instruction is it used with?
40. Select an assembly language instruction that exchanges the contents of the EBX register with the ESI register.
41. Where is the I/O address (port number) stored for an INSB instruction?
42. Would the LAHF and SAHF instructions normally appear in software?
43. Write a short program that uses the XLAT instruction to convert the BCD numbers 0–9 into ASCII-coded numbers 30H–39H. Store the ASCII-coded data in a TABLE located within the data segment.
44. Explain how the XLAT instruction transforms the contents of the AL register.
45. Explain what the IN AL,12H instruction accomplishes.
46. Explain how the OUT DX,AX instruction operates.
47. What is a segment override prefix?
48. Select an instruction that moves a byte of data from the memory location addressed by the BX register in the extra segment into the AH register.
49. Develop a sequence of instructions that exchanges the contents of AX with BX, ECX with EDX, and SI with DI.
50. What is an assembly language directive?
51. What is accomplished by the CMOVNE CX,DX instruction in the Pentium 4 microprocessor?
52. Describe the purpose of the following assembly language directives: DB, DW, and DD.
53. Select an assembly language directive that reserves 30 bytes of memory for array LIST1.
54. Describe the purpose of the EQU directive.
55. What is the purpose of the .686 directive?
56. What is the purpose of the .MODEL directive?
57. If the start of a segment is identified with .DATA, what type of memory organization is in effect?
58. If the SEGMENT directive identifies the start of a segment, what type of memory organization is in effect?
59. What does the INT 21H accomplish if AH contains a 4CH?
60. What directives indicate the start and end of a procedure?
61. Explain the purpose of the USES statement as it applies to a procedure with version 6.x of MASM.
62. Develop a near procedure that stores AL in four consecutive memory locations within the data segment, as addressed by the DI register.
63. How is the Pentium 4 microprocessor instructed to use the 16-bit instruction mode?
64. Develop a far procedure that copies contents of the word-sized memory location CS:DATA4 into AX, BX, CX, DX, and SI.

CHAPTER 5

Arithmetic and Logic Instructions

INTRODUCTION

In this chapter, we examine the arithmetic and logic instructions. The arithmetic instructions include addition, subtraction, multiplication, division, comparison, negation, increment, and decrement. The logic instructions include AND, OR, Exclusive-OR, NOT, shifts, rotates, and the logical compare (TEST). This chapter also presents the 80386 through the Core2 instructions XADD, SHRD, SHLD, bit tests, and bit scans. The chapter concludes with a discussion of string comparison instructions, which are used for scanning tabular data and for comparing sections of memory data. Both comparison tasks are performed efficiently with the string scan (SCAS) and string compare (CMPS) instructions.

If you are familiar with an 8-bit microprocessor, you will recognize that the 8086 through the Core2 instruction set is superior to most 8-bit microprocessors because most of the instructions have two operands instead of one. Even if this is your first microprocessor, you will quickly learn that this microprocessor possesses a powerful and easy-to-use set of arithmetic and logic instructions.

CHAPTER OBJECTIVES

Upon completion of this chapter, you will be able to:

1. Use arithmetic and logic instructions to accomplish simple binary, BCD, and ASCII arithmetic.
2. Use AND, OR, and Exclusive-OR to accomplish binary bit manipulation.
3. Use the shift and rotate instructions.
4. Explain the operation of the 80386 through the Core2 exchange and add, compare and exchange, double-precision shift, bit test, and bit scan instructions.
5. Check the contents of a table for a match with the string instructions.

5-1   ADDITION, SUBTRACTION, AND COMPARISON

The bulk of the arithmetic instructions found in any microprocessor include addition, subtraction, and comparison. In this section, addition, subtraction, and comparison instructions are illustrated. Also shown are their uses in manipulating register and memory data.

Addition

Addition (ADD) appears in many forms in the microprocessor. This section details the use of the ADD instruction for 8-, 16-, and 32-bit binary addition. A second form of addition, called **add-with-carry**, is introduced with the ADC instruction. Finally, the increment instruction (INC) is presented. Increment is a special type of addition that adds 1 to a number. In Section 5–3, other forms of addition are examined, such as BCD and ASCII. Also described is the XADD instruction, found in the 80486 through the Pentium 4.

Table 5–1 illustrates the addressing modes available to the ADD instruction. (These addressing modes include almost all those mentioned in Chapter 3.) However, because there are more than 32,000 variations of the ADD instruction in the instruction set, it is impossible to list them all in this table. The only types of addition *not* allowed are memory-to-memory and segment register. The segment registers can only be moved, pushed, or popped. Note that, as with all other instructions, the 32-bit registers are available only with the 80386 through the Core2. In the 64-bit mode of the Pentium 4 and Core2, the 64-bit registers are also used for addition.

TABLE 5–1 Example addition instructions.

Assembly Language	Operation
ADD AL,BL	AL = AL + BL
ADD CX,DI	CX = CX + DI
ADD EBP,EAX	EBP = EBP + EAX
ADD CL,44H	CL = CL + 44H
ADD BX,245FH	BX = BX + 245FH
ADD EDX,12345H	EDX = EDX + 12345H
ADD [BX],AL	AL adds to the byte contents of the data segment memory location addressed by BX with the sum stored in the same memory location
ADD CL,[BP]	The byte contents of the stack segment memory location addressed by BP add to CL with the sum stored in CL
ADD AL,[EBX]	The byte contents of the data segment memory location addressed by EBX add to AL with the sum stored in AL
ADD BX,[SI+2]	The word contents of the data segment memory location addressed by SI + 2 add to BX with the sum stored in BX
ADD CL,TEMP	The byte contents of data segment memory location TEMP add to CL with the sum stored in CL
ADD BX,TEMP[DI]	The word contents of the data segment memory location addressed by TEMP + DI add to BX with the sum stored in BX
ADD [BX+D],DL	DL adds to the byte contents of the data segment memory location addressed by BX + DI with the sum stored in the same memory location
ADD BYTE PTR [DI],3	A 3 adds to the byte contents of the data segment memory location addressed by DI with the sum stored in the same location
ADD BX,[EAX+2*ECX]	The word contents of the data segment memory location addressed by EAX plus 2 times ECX add to BX with the sum stored in BX
ADD RAX,RBX	RBX adds to RAX with the sum stored in RAX (64-bit mode)
ADD EDX,[RAX+RCX]	The doubleword in EDX is added to the doubleword addressed by the sum of RAX and RCX and the sum is stored in EDX (64-bit mode)

Register Addition.   Example 5–1 shows a simple sequence of instructions that uses register addition to add the contents of several registers. In this example, the contents of AX, BX, CX, and DX are added to form a 16-bit result stored in the AX register.

EXAMPLE 5–1

```
0000 03 C3        ADD AX,BX
0002 03 C1        ADD AX,CX
0004 03 C2        ADD AX,DX
```

Whenever arithmetic and logic instructions execute, the contents of the flag register change. Note that the contents of the interrupt, trap, and other flags do not change due to arithmetic and logic instructions. Only the flags located in the rightmost 8 bits of the flag register and the overflow flag change. These rightmost flags denote the result of the arithmetic or a logic operation. Any ADD instruction modifies the contents of the sign, zero, carry, auxiliary carry, parity, and overflow flags. The flag bits never change for most of the data transfer instructions presented in Chapter 4.

Immediate Addition.   Immediate addition is employed whenever constant or known data are added. An 8-bit immediate addition appears in Example 5-2. In this example, DL is first loaded with 12H by using an immediate move instruction. Next, 33H is added to the 12H in DL by an immediate addition instruction. After the addition, the sum (45H) moves into register DL and the flags change, as follows:

$$Z = 0 \text{ (result not zero)}$$
$$C = 0 \text{ (no carry)}$$
$$A = 0 \text{ (no half-carry)}$$
$$S = 0 \text{ (result positive)}$$
$$P = 0 \text{ (odd parity)}$$
$$O = 0 \text{ (no overflow)}$$

EXAMPLE 5–2

```
0000 B2 12        MOV DL,12H
0002 80 C2 33     ADD DL,33H
```

Memory-to-Register Addition.   Suppose that an application requires memory data to be added to the AL register. Example 5–3 shows an example that adds two consecutive bytes of data, stored at the data segment offset locations NUMB and NUMB+1, to the AL register.

EXAMPLE 5–3

```
0000 BF 0000 R    MOV DI,OFFSET NUMB       ;address NUMB
0003 B0 00        MOV AL,0                 ;clear sum
0005 02 05        ADD AL,[DI]              ;add NUMB
0007 02 45 01     ADD AL,[DI+1]            ;add NUMB+1
```

The first instruction loads the destination index register (DI) with offset address NUMB. The DI register, used in this example, addresses data in the data segment beginning at memory location NUMB. After clearing the sum to zero, the ADD AL,[DI] instruction adds the contents of memory location NUMB to AL. Finally, the ADD AL,[DI+I] instruction adds the contents of memory location NUMB plus 1 byte to the AL register. After both ADD instructions execute, the result appears in the AL register as the sum of the contents of NUMB plus the contents of NUMB+1.

Array Addition.   Memory arrays are sequential lists of data. Suppose that an array of data (ARRAY) contains 10 bytes, numbered from element 0 through element 9. Example 5–4 shows how to add the contents of array elements 3, 5, and 7 together.

This example first clears AL to 0, so it can be used to accumulate the sum. Next, register SI is loaded with a 3 to initially address array element 3. The ADD AL,ARRAY[SI] instruction adds the contents of array element 3 to the sum in AL. The instructions that follow add array elements 5 and 7 to the sum in AL, using a 3 in SI plus a displacement of 2 to address element 5, and a displacement of 4 to address element 7.

EXAMPLE 5–4

```
0000 B0 00                  MOV AL,0           ;clear sum
0002 BE 0003                MOV SI,3           ;address element 3
0005 02 84 0000 R           ADD AL,ARRAY[SI]   ;add element 3
0009 02 84 0002 R           ADD AL,ARRAY[SI+2] ;add element 5
000D 02 84 0004 R           ADD AL,ARRAY[SI+4] ;add element 7
```

Suppose that an array of data contains 16-bit numbers used to form a 16-bit sum in register AX. Example 5–5 shows a sequence of instructions written for the 80386 and above, showing the scaled-index form of addressing to add elements 3, 5, and 7 of an area of memory called ARRAY. In this example, EBX is loaded with the address ARRAY, and ECX holds the array element number. Note how the scaling factor is used to multiply the contents of the ECX register by 2 to address words of data. (Recall that words are 2 bytes long.)

EXAMPLE 5–5

```
0000 66|BB 00000000 R       MOV EBX,OFFSET ARRAY   ;address ARRAY
0006 66|B9 00000003         MOV ECX,3              ;address element 3
000C 67&8B 04 4B            MOV AX,[EBX+2*ECX]     ;get element 3
0010 66|B9 00000005         MOV ECX,5              ;address element 5
0016 67&03 04 4B            ADD AX,[EBX+2*ECX]     ;add element 5
001A 66|B0 00000007         MOV ECX,7              ;address element 7
0020 67&03 04 4B            ADD AX,[EBX+2*ECX]     ;add element 7
```

Increment Addition. Increment addition (INC) adds 1 to a register or a memory location. The INC instruction adds 1 to any register or memory location, except a segment register. Table 5–2 illustrates some of the possible forms of the increment instructions available to the 8086–Core2 processors. As with other instructions presented thus far, it is impossible to show all variations of the INC instruction because of the large number available.

With indirect memory increments, the size of the data must be described by using the BYTE PTR, WORD PTR, DWORD PTR, or QWORD PTR directives. The reason is that the

TABLE 5–2 Example increment instructions.

Assembly Language	Operation
INC BL	BL = BL + 1
INC SP	SP = SP + 1
INC EAX	EAX = EAX + 1
INC BYTE PTR[BX]	Adds 1 to the byte contents of the data segment memory location addressed by BX
INC WORD PTR[SI]	Adds 1 to the word contents of the data segment memory location addressed by SI
INC DWORD PTR[ECX]	Adds 1 to the doubleword contents of the data segment memory location addressed by ECX
INC DATA1	Adds 1 to the contents of data segment memory location DATA1
INC RCX	Adds 1 to RCX (64-bit mode)

assembler program cannot determine if, for example, the INC [DI] instruction is a byte-, word-, or doubleword-sized increment. The INC BYTE PTR [DI] instruction clearly indicates byte-sized memory data; the INC WORD PTR [DI] instruction unquestionably indicates a word-sized memory data; and the INC DWORD PTR [DI] instruction indicates doubleword-sized data. In 64-bit mode operation of the Pentium 4 and Core2, the INC QWORD PTR [RSI] instruction indicates quadword-sized data.

Example 5–6 shows how to modify Example 5–3 to use the increment instruction for addressing NUMB and NUMB+1. Here, an INC DI instruction changes the contents of register DI from offset address NUMB to offset address NUMB+1. Both program sequences shown in Examples 5–3 and 5–6 add the contents of NUMB and NUMB+1. The difference between them is the way that the address is formed through the contents of the DI register using the increment instruction.

EXAMPLE 5–6

```
0000 BF 0000 R      MOV DI,OFFSET NUMB          ;address NUMB
0003 B0 00          MOV AL,0                    ;clear sum
0005 02 05          ADD AL,[DI]                 ;add NUMB
0007 47             INC DI                      ;increment DI
0008 02 05          ADD AL,[DI]                 ;add NUMB+1
```

Increment instructions affect the flag bits, as do most other arithmetic and logic operations. The difference is that increment instructions do not affect the carry flag bit. Carry doesn't change because we often use increments in programs that depend upon the contents of the carry flag. Note that increment is used to point to the next memory element in a byte-sized array of data only. If word-sized data are addressed, it is better to use an ADD DI,2 instruction to modify the DI pointer in place of two INC DI instructions. For doubleword arrays, use the ADD DI,4 instruction to modify the DI pointer. In some cases, the carry flag must be preserved, which may mean that two or four INC instructions might appear in a program to modify a pointer.
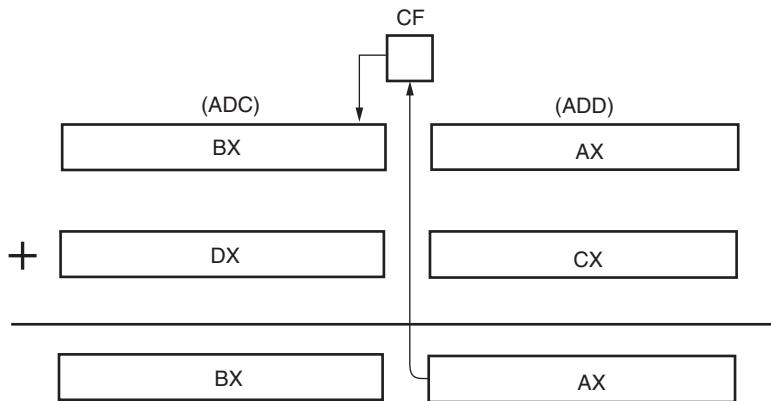
Addition-with-Carry.   An addition-with-carry instruction (ADC) adds the bit in the carry flag (C) to the operand data. This instruction mainly appears in software that adds numbers that are wider than 16 bits in the 8086–80286 or wider than 32 bits in the 80386–Core2.

Table 5–3 lists several add-with-carry instructions, with comments that explain their operation. Like the ADD instruction, ADC affects the flags after the addition.

TABLE 5–3   Example add-with-carry instructions.

Assembly Language	Operation
ADC AL,AH	AL = AL + AH + carry
ADC CX,BX	CX = CX + BX + carry
ADC EBX,EDX	EBX = EBX + EDX + carry
ADC RBX,0	RBX = RBX + 0 + carry (64-bit mode)
ADC DH,[BX]	The byte contents of the data segment memory location addressed by BX add to DH with the sum stored in DH
ADC BX,[BP+2]	The word contents of the stack segment memory location addressed by BP plus 2 add to BX with the sum stored in BX
ADC ECX,[EBX]	The doubleword contents of the data segment memory location addressed by EBX add to ECX with the sum stored in ECX

FIGURE 5–1 Addition-with-carry showing how the carry flag (C) links the two 16-bit additions into one 32-bit addition.

Suppose that a program is written for the 8086–80286 to add the 32-bit number in BX and AX to the 32-bit number in DX and CX. Figure 5–1 illustrates this addition so that the placement and function of the carry flag can be understood. This addition cannot be easily performed without adding the carry flag bit because the 8086–80286 only adds 8- or 16-bit numbers. Example 5–7 shows how the contents of registers AX and CX add to form the least significant 16 bits of the sum. This addition may or may not generate a carry. A carry appears in the carry flag if the sum is greater than FFFFH. Because it is impossible to predict a carry, the most significant 16 bits of this addition are added with the carry flag using the ADC instruction. The ADC instruction adds the 1 or the 0 in the carry flag to the most significant 16 bits of the result. This program adds BX–AX to DX–CX, with the sum appearing in BX–AX.

EXAMPLE 5–7

```
0000 03 C1          ADD AX,CX
0002 13 DA          ADC BX,DX
```

Suppose the same software is rewritten for the 80386 through the Core2, but modified to add two 64-bit numbers in the 32-bit mode. The changes required for this operation are the use of the extended registers to hold the data and modifications of the instructions for the 80386 and above. These changes are shown in Example 5–8, which adds two 64-bit numbers. In the 64-bit mode of the Pentium 4 and Core2, this addition is handled with a single ADD instruction if the location of the operands is changed to RAX and RBX as in the instruction ADD RAX,RBX, which adds RBX to RAX.

EXAMPLE 5–8

```
0000 66|03 C1       ADD EAX,ECX
0003 66|13 DA       ADC EBX,EDX
```

Exchange and Add for the 80486–Core2 Processors. A new type of addition called exchange and add (XADD) appears in the 80486 instruction set and continues through the Core2. The XADD instruction adds the source to the destination and stores the sum in the destination, as with any addition. The difference is that after the addition takes place, the original value of the destination is copied into the source operand. This is one of the few instructions that change the source.

For example, if BL = 12H and DL = 02H, and the XADD BL,DL instruction executes, the BL register contains the sum of 14H and DL becomes 12H. The sum of 14H is generated and the original destination of 12H replaces the source. This instruction functions with any register size and any memory operand, just as with the ADD instruction.

Subtraction

Many forms of subtraction (SUB) appear in the instruction set. These forms use any addressing mode with 8-, 16-, or 32-bit data. A special form of subtraction (decrement, or DEC) subtracts 1 from any register or memory location. Section 5–3 shows how BCD and ASCII data subtract. As with addition, numbers that are wider than 16 bits or 32 bits must occasionally be subtracted. The **subtract-with-borrow instruction** (SBB) performs this type of subtraction. In the 80486 through the Core2 processors, the instruction set also includes a compare and exchange instruction. In the 64-bit mode for the Pentium 4 and Core2, a 64-bit subtraction is also available.

Table 5–4 lists some of the many addressing modes allowed with the subtract instruction (SUB). There are well over 1000 possible subtraction instructions, far too many to list here. About the only types of subtraction not allowed are memory-to-memory and segment register subtractions. Like other arithmetic instructions, the subtract instruction affects the flag bits.

Register Subtraction.   Example 5–9 shows a sequence of instructions that perform register subtraction. This example subtracts the 16-bit contents of registers CX and DX from the contents of register BX. After each subtraction, the microprocessor modifies the contents of the flag register. The flags change for most arithmetic and logic operations.

EXAMPLE 5–9

```
0000 2B D9          SUB BX,CX
0002 2B DA          SUB BX,DX
```

Immediate Subtraction.   As with addition, the microprocessor also allows immediate operands for the subtraction of constant data. Example 5–10 presents a short sequence of instructions that subtract 44H from 22H. Here, we first load the 22H into CH using an immediate move

TABLE 5–4   Example subtraction instructions.

Assembly Language	Operation
SUB CL,BL	CL = CL – BL
SUB AX,SP	AX = AX – SP
SUB ECX,EBP	ECX = ECX – EBP
SUB RDX,R8	RDX = RDX – R8 (64-bit mode)
SUB DH,6FH	DH = DH – 6FH
SUB AX,0CCCCH	AX = AX – 0CCCCH
SUB ESI,2000300H	ESI = ESI – 2000300H
SUB [DI],CH	Subtracts CH from the byte contents of the data segment memory addressed by DI and stores the difference in the same memory location
SUB CH,[BP]	Subtracts the byte contents of the stack segment memory location addressed by BP from CH and stores the difference in CH
SUB AH,TEMP	Subtracts the byte contents of memory location TEMP from AH and stores the difference in AH
SUB DI,TEMP[ESI]	Subtracts the word contents of the data segment memory location addressed by TEMP plus ESI from DI and stores the difference in DI
SUB ECX,DATA1	Subtracts the doubleword contents of memory location DATA1 from ECX and stores the difference in ECX
SUB RCX,16	RCX = RCX – 18 (64-bit mode)

instruction. Next, the SUB instruction, using immediate data 44H, subtracts 44H from the 22H. After the subtraction, the difference (0DEH) moves into the CH register. The flags change as follows for this subtraction:

$$Z = 0 \text{ (result not zero)}$$
$$C = 1 \text{ (borrow)}$$
$$A = 1 \text{ (half-borrow)}$$
$$S = 1 \text{ (result negative)}$$
$$P = 1 \text{ (even parity)}$$
$$O = 0 \text{ (no overflow)}$$

EXAMPLE 5–10

```
0000 B5 22          MOV CH,22H
0002 80 ED 44       SUB CH,44H
```

Both carry flags (C and A) hold borrows after a subtraction instead of carries, as after an addition. Notice in this example that there is no overflow. This example subtracted 44H (+68) from 22H (+34), resulting in a 0DEH (−34). Because the correct 8-bit signed result is −34, there is no overflow in this example. An 8-bit overflow occurs only if the signed result is greater than +127 or less than −128.

Decrement Subtraction.   Decrement subtraction (DEC) subtracts 1 from a register or the contents of a memory location. Table 5–5 lists some decrement instructions that illustrate register and memory decrements.

The decrement indirect memory data instructions require BYTE PTR, WORD PTR, DWORD PTR, or QWORD PTR because the assembler cannot distinguish a byte from a word or doubleword when an index register addresses memory. For example, DEC [SI] is vague because the assembler cannot determine whether the location addressed by SI is a byte, word, or doubleword. Using DEC BYTE PTR[SI], DEC WORD PTR[DI], or DEC DWORD PTR[SI] reveals

TABLE 5–5   Example decrement instructions.

Assembly Language	Operation
DEC BH	BH = BH − 1
DEC CX	CX = CX − 1
DEC EDX	EDX = EDX − 1
DEC R14	R14 = R14 − 1 (64-bit mode)
DEC BYTE PTR[DI]	Subtracts 1 from the byte contents of the data segment memory location addressed by DI
DEC WORD PTR[BP]	Subtracts 1 from the word contents of the stack segment memory location addressed by BP
DEC DWORD PTR[EBX]	Subtracts 1 from the doubleword contents of the data segment memory location addressed by EBX
DEC QWORD PTR[RSI]	Subtracts 1 from the quadword contents of the memory location addressed by RSI (64-bit mode)
DEC NUMB	Subtracts 1 from the contents of data segment memory location NUMB

TABLE 5–6 Example subtraction-with-borrow instructions.

Assembly Language	Operation
SBB AH,AL	AH = AH – AL – carry
SBB AX,BX	AX = AX – BX – carry
SBB EAX,ECX	EAX = EAX – ECX – carry
SBB CL,2	CL = CL – 2 – carry
SBB RBP,8	RBP = RBP– 2 – carry (64-bit mode)
SBB BYTE PTR[DI],3	Both 3 and carry subtract from the data segment memory location addressed by DI
SBB [DI],AL	Both AL and carry subtract from the data segment memory location addressed by DI
SBB DI,[BP+2]	Both carry and the word contents of the stack segment memory location addressed by BP plus 2 subtract from DI
SBB AL,[EBX+ECX]	Both carry and the byte contents of the data segment memory location addressed by EBX plus ECX subtract from AL
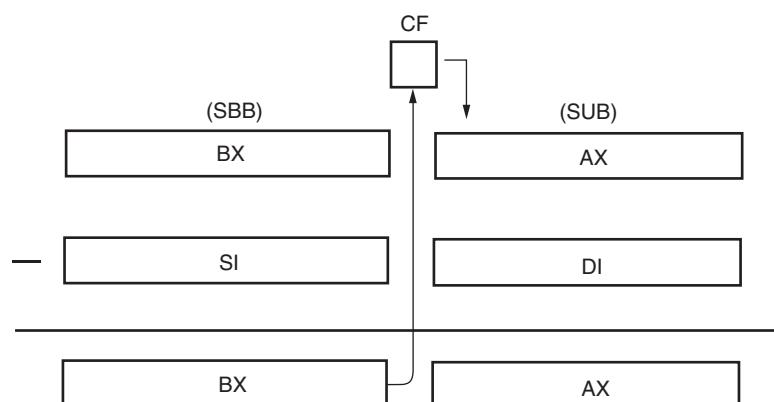
the size of the data to the assembler. In the 64-bit mode, a DEC QWORD PTR[RSI] decrement the 64-bit number stored at the address pointed to by the RSI register.

Subtraction-with-Borrow. A subtraction-with-borrow (SBB) instruction functions as a regular subtraction, except that the carry flag (C), which holds the borrow, also subtracts from the difference. The most common use for this instruction is for subtractions that are wider than 16 bits in the 8086–80286 microprocessors or wider than 32 bits in the 80386–Core2. Wide subtractions require that borrows propagate through the subtraction, just as wide additions propagate the carry.

Table 5–6 lists several SBB instructions with comments that define their operations. Like the SUB instruction, SBB affects the flags. Notice that the immediate subtract from memory instruction in this table requires a BYTE PTR, WORD PTR, DWORD PTR, or QWORD PTR directive.

When the 32-bit number held in BX and AX is subtracted from the 32-bit number held in SI and DI, the carry flag propagates the borrow between the two 16-bit subtractions. The carry flag holds the borrow for subtraction. Figure 5–2 shows how the borrow propagates through the carry flag (C) for this task. Example 5–11 shows how this subtraction is performed by a program. With wide subtraction, the least significant 16- or 32-bit data are subtracted with the SUB

FIGURE 5–2 Subtraction-with-borrow showing how the carry flag (C) propagates the borrow.

instruction. All subsequent and more significant data are subtracted by using the SBB instruction. The example uses the SUB instruction to subtract DI from AX, then uses SBB to subtract-with-borrow SI from BX.

EXAMPLE 5–11

```
0000 2B C7        SUB AX,DI
0002 1B DE        SBB BX,SI
```

Comparison

The comparison instruction (CMP) is a subtraction that changes only the flag bits; the destination operand never changes. A comparison is useful for checking the entire contents of a register or a memory location against another value. A CMP is normally followed by a conditional jump instruction, which tests the condition of the flag bits.

Table 5–7 lists a variety of comparison instructions that use the same addressing modes as the addition and subtraction instructions already presented. Similarly, the only disallowed forms of compare are memory-to-memory and segment register compares.

Example 5–12 shows a comparison followed by a conditional jump instruction. In this example, the contents of AL are compared with 10H. Conditional jump instructions that often follow the comparison are JA (jump above) or JB (jump below). If the JA follows the comparison, the jump occurs if the value in AL is above 10H. If the JB follows the comparison, the jump occurs if the value in AL is below 10H. In this example, the JAE instruction follows the comparison. This instruction causes the program to continue at memory location SUBER if the value in AL is 10H or above. There is also a JBE (jump below or equal) instruction that could follow the comparison to jump if the outcome is below or equal to 10H. Later chapters provide additional detail on the comparison and conditional jump instructions.

TABLE 5–7 Example comparison instructions.

Assembly Language	Operation
CMP CL,BL	CL – BL
CMP AX,SP	AX – SP
CMP EBP,ESI	EBP – ESI
CMP RDI,RSI	RDI – RSI (64-bit mode)
CMP AX,2000H	AX – 2000H
CMP R10W,12H	R10 (word portion) – 12H (64-bit mode)
CMP [DI],CH	CH subtracts from the byte contents of the data segment memory location addressed by DI
CMP CL,[BP]	The byte contents of the stack segment memory location addressed by BP subtracts from CL
CMP AH,TEMP	The byte contents of data segment memory location TEMP subtracts from AH
CMP DI,TEMP[BX]	The word contents of the data segment memory location addressed by TEMP plus BX subtracts from DI
CMP AL,[EDI+ESI]	The byte contents of the data segment memory location addressed by EDI plus ESI subtracts from AL

EXAMPLE 5–12

```
0000 3C 10         CMP AL,10H          ;compare AL against 10H
0002 73 1C         JAE SUBER           ;if AL is 10H or above
```

Compare and Exchange (80486–Core2 Processors Only).   The compare and exchange instruction (CMPXCHG), found only in the 80486 through the Core2 instruction sets, compares the destination operand with the accumulator. If they are equal, the source operand is copied into the destination; if they are not equal, the destination operand is copied into the accumulator. This instruction functions with 8-, 16-, or 32-bit data.

The CMPXCHG CX,DX instruction is an example of the compare and exchange instruction. This instruction first compares the contents of CX with AX. If CX equals AX, DX is copied into AX; if CX is not equal to AX, CX is copied into AX. This instruction also compares AL with 8-bit data and EAX with 32-bit data if the operands are either 8- or 32-bit.

In the Pentium–Core2 processors, a CMPXCHG8B instruction is available that compares two quadwords. This is the only new data manipulation instruction provided in the Pentium–Core2 when they are compared with prior versions of the microprocessor. The compare-and-exchange-8-bytes instruction compares the 64-bit value located in EDX:EAX with a 64-bit number located in memory. An example is CMPXCHG8B TEMP. If TEMP equals EDX:EAX, TEMP is replaced with the value found in ECX:EBX; if TEMP does not equal EDX:EAX, the number found in TEMP is loaded into EDX:EAX. The Z (zero) flag bit indicates that the values are equal after the comparison.

This instruction has a bug that will cause the operating system to crash. More information about this flaw can be obtained at www.intel.com. There is also a CMPXCHG16B instruction available to the Pentium 4 when operated in 64-bit mode.

5-2   MULTIPLICATION AND DIVISION

Only modern microprocessors contain multiplication and division instructions. Earlier 8-bit microprocessors could not multiply or divide without the use of a program that multiplied or divided by using a series of shifts and additions or subtractions. Because microprocessor manufacturers were aware of this inadequacy, they incorporated multiplication and division instructions into the instruction sets of the newer microprocessors. The Pentium–Core2 processors contain special circuitry that performs a multiplication in as little as one clocking period, whereas it took over 40 clocking periods to perform the same multiplication in earlier Intel microprocessors.

Multiplication

Multiplication is performed on bytes, words, or doublewords, and can be signed integer (IMUL) or unsigned integer (MUL). Note that only the 80386 through the Core2 processors multiply 32-bit doublewords. The product after a multiplication is always a double-width product. If two 8-bit numbers are multiplied, they generate a 16-bit product; if two 16-bit numbers are multiplied, they generate a 32-bit product; and if two 32-bit numbers are multiplied, a 64-bit product is generated. In the 64-bit mode of the Pentium 4, two 64-bit numbers are multiplied to generate a 128-bit product.

Some flag bits (overflow and carry) change when the multiply instruction executes and produce predictable outcomes. The other flags also change, but their results are unpredictable and therefore are unused. In an 8-bit multiplication, if the most significant 8 bits of the result are zero, both C and O flag bits equal zero. These flag bits show that the result is 8 bits wide (C = 0) or 16 bits wide (C = 1). In a 16-bit multiplication, if the most significant 16-bits part of

TABLE 5–8 Example 8-bit multiplication instructions.

Assembly Language	Operation
MUL CL	AL is multiplied by CL; the unsigned product is in AX
IMUL DH	AL is multiplied by DH; the signed product is in AX
IMUL BYTE PTR[BX]	AL is multiplied by the byte contents of the data segment memory location addressed by BX; the signed product is in AX
MUL TEMP	AL is multiplied by the byte contents of data segment memory location TEMP; the unsigned product is in AX

the product is 0, both C and O clear to zero. In a 32-bit multiplication, both C and O indicate that the most significant 32 bits of the product are zero.

8-Bit Multiplication. With 8-bit multiplication, the multiplicand is always in the AL register, whether signed or unsigned. The multiplier can be any 8-bit register or any memory location. Immediate multiplication is not allowed unless the special signed immediate multiplication instruction, discussed later in this section, appears in a program. The multiplication instruction contains one operand because it always multiplies the operand times the contents of register AL. An example is the MUL BL instruction, which multiplies the unsigned contents of AL by the unsigned contents of BL. After the multiplication, the unsigned product is placed in AX—a double-width product. Table 5–8 illustrates some 8-bit multiplication instructions.

Suppose that BL and CL each contain two 8-bit unsigned numbers, and these numbers must be multiplied to form a 16-bit product stored in DX. This procedure cannot be accomplished by a single instruction because we can only multiply a number times the AL register for an 8-bit multiplication. Example 5–13 shows a short program that generates DX = BL × CL. This example loads register BL and CL with example data 5 and 10. The product, a 50, moves into DX from AX after the multiplication by using the MOV DX,AX instruction.

EXAMPLE 5–13

```
0000 B3 05        MOV   BL,5          ;load data
0002 B1 0A        MOV   CL,10
0004 8A C1        MOV   AL,CL         ;position data
0006 F6 E3        MUL   BL            ;multiply
0008 8B D0        MOV   DX,AX         ;position product
```

For signed multiplication, the product is in binary form, if positive, and in two's complement form, if negative. These are the same forms used to store all positive and negative signed numbers used by the microprocessor. If the program of Example 5–13 multiplies two signed numbers, only the MUL instruction is changed to IMUL.

16-Bit Multiplication. Word multiplication is very similar to byte multiplication. The difference is that AX contains the multiplicand instead of AL, and the 32-bit product appears in DX–AX instead of AX. The DX register always contains the most significant 16 bits of the product, and AX contains the least significant 16 bits. As with 8-bit multiplication, the choice of the multiplier is up to the programmer. Table 5–9 shows several different 16-bit multiplication instructions.

A Special Immediate 16-Bit Multiplication. The 8086/8088 microprocessors could not perform immediate multiplication; the 80186 through the Core2 processors can do so by using a special version of the multiply instruction. Immediate multiplication must be signed multiplication, and the instruction format is different because it contains three operands. The first operand is the 16-bit destination register; the second operand is a register or memory location

TABLE 5–9   Example 16-bit multiplication instructions.

Assembly Language	Operation
MUL CX	AX is multiplied by CX; the unsigned product is in DX–AX
IMUL DI	AX is multiplied by DI; the signed product is in DX–AX
MUL WORD PTR[SI]	AX is multiplied by the word contents of the data segment memory location addressed by SI; the unsigned product is in DX–AX

that contains the 16-bit multiplicand; and the third operand is either 8-bit or 16-bit immediate data used as the multiplier.

The IMUL CX,DX,12H instruction multiplies 12H times DX and leaves a 16-bit signed product in CX. If the immediate data are 8 bits, they sign-extend into a 16-bit number before the multiplication occurs. Another example is IMUL BX,NUMBER,1000H, which multiplies NUM-BER times 1000H and leaves the product in BX. Both the destination and multiplicand must be 16-bit numbers. Although this is immediate multiplication, the restrictions placed upon it limit its utility, especially the fact that it is a signed multiplication and the product is 16 bits wide.

32-Bit Multiplication.   In the 80386 and above, 32-bit multiplication is allowed because these microprocessors contain 32-bit registers. As with 8- and 16-bit multiplication, 32-bit multiplication can be signed or unsigned by using the IMUL and MUL instructions. With 32-bit multiplication, the contents of EAX are multiplied by the operand specified with the instruction. The product (64 bits wide) is found in EDX–EAX, where EAX contains the least significant 32 bits of the product. Table 5–10 lists some of the 32-bit multiplication instructions found in the 80386 and above instruction set.

64-Bit Multiplication.   The result of a 64-bit multiplication in the Pentium 4 appears in the RDX:RAX register pair as a 128-bit product. Although multiplication of this size is relatively rare, the Pentium 4 and Core2 can perform it on both signed and unsigned numbers. Table 5–11 shows a few examples of this high precision multiplication.

TABLE 5–10   Example 32-bit multiplication instructions.

Assembly Language	Operation
MUL ECX	EAX is multiplied by ECX; the unsigned product is in EDX–EAX
IMUL EDI	EAX is multiplied by EDI; the signed product is in EDX–EAX
MUL DWORD PTR[ESI]	EAX is multiplied by the doubleword contents of the data segment memory location address by ESI; the unsigned product is in EDX–EAX

TABLE 5–11   Example 64-bit multiplication instructions.

Assembly Language	Operation
MUL RCX	RAX is multiplied by RCX; the unsigned product is in RDX–RAX
IMUL RDI	RAX is multiplied by RDI; the signed product is in RDX–RAX
MUL QWORD PTR[RSI]	RAX is multiplied by the quadword contents of the memory location address by RSI; the unsigned product is in RDX–RAX

Division

As with multiplication, division occurs on 8- or 16-bit numbers in the 8086–80286 microprocessors, and on 32-bit numbers in the 80386 and above microprocessor. These numbers are signed (IDIV) or unsigned (DIV) integers. The dividend is always a double-width dividend that is divided by the operand. This means that an 8-bit division divides a 16-bit number by an 8-bit number; a 16-bit division divides a 32-bit number by a 16-bit number; and a 32-bit division divides a 64-bit number by a 32-bit number. There is no immediate division instruction available to any microprocessor. In the 64-bit mode of the Pentium 4 and Core2, a 64-bit division divides a 128-bit number by a 64-bit number.

None of the flag bits change predictably for a division. A division can result in two different types of errors; one is an attempt to divide by zero and the other is a divide overflow. A divide overflow occurs when a small number divides into a large number. For example, suppose that AX = 3000 and that it is divided by 2. Because the quotient for an 8-bit division appears in AL, the result of 1500 causes a divide overflow because the 1500 does not fit into AL. In either case, the microprocessor generates an interrupt if a divide error occurs. In most systems, a divide error interrupt displays an error message on the video screen. The divide error interrupt and all other interrupts for the microprocessor are explained in Chapter 6.

8-Bit Division. An 8-bit division uses the AX register to store the dividend that is divided by the contents of any 8-bit register or memory location. The quotient moves into AL after the division with AH containing a whole number remainder. For a signed division, the quotient is positive or negative; the remainder always assumes the sign of the dividend and is always an integer. For example, if AX = 0010H (+16) and BL = 0FDH(−3) and the IDIV BL instruction executes, AX = 01FBH. This represents a quotient of −5 (AL) with a remainder of 1 (AH). If, on the other hand, a −16 is divided by +3, the result will be a quotient of −5 (AL) with a remainder of −1 (AH). Table 5–12 lists some of the 8-bit division instructions.

With 8-bit division, the numbers are usually 8 bits wide. This means that one of them, the dividend, must be converted to a 16-bit wide number in AX. This is accomplished differently for signed and unsigned numbers. For the unsigned number, the most significant 8 bits must be cleared to zero (zero-extended). The MOVZX instruction described in Chapter 4 can be used to zero-extend a number in the 80386 through the Core2 processors. For signed numbers, the least significant 8 bits are sign-extended into the most significant 8 bits. In the microprocessor, a special instruction sign-extends AL into AH, or converts an 8-bit signed number in AL into a 16-bit signed number in AX. The CBW (**convert byte to word**) instruction performs this conversion. In the 80386 through the Core2, a MOVSX instruction (see Chapter 4) sign-extends a number.

TABLE 5–12 Example 8-bit division instructions.

Assembly Language	Operation
DIV CL	AX is divided by CL; the unsigned quotient is in AL and the unsigned remainder is in AH
IDIV BL	AX is divided by BL; the signed quotient is in AL and the signed remainder is in AH
DIV BYTE PTR[BP]	AX is divided by the byte contents of the stack segment memory location addressed by BP; the unsigned quotient is in AL and the unsigned remainder is in AH

EXAMPLE 5–14

```
0000 A0 0000 R          MOV  AL,NUMB        ;get NUMB
0003 B4 00              MOV  AH,0           ;zero-extend
0005 F6 36 0002 R       DIV  NUMB1          ;divide by NUMB1
0009 A2 0003 R          MOV  ANSQ,AL        ;save quotient
000C 88 26 0004 R       MOV  ANSR,AH        ;save remainder
```

Example 5–14 illustrates a short program that divides the unsigned byte contents of memory location NUMB by the unsigned contents of memory location NUMB1. Here, the quotient is stored in location ANSQ and the remainder is stored in location ANSR. Notice how the contents of location NUMB are retrieved from memory and then zero-extended to form a 16-bit unsigned number for the dividend.

16-Bit Division.   Sixteen-bit division is similar to 8-bit division, except that instead of dividing into AX, the 16-bit number is divided into DX–AX, a 32-bit dividend. The quotient appears in AX and the remainder appears in DX after a 16-bit division. Table 5–13 lists some of the 16-bit division instructions.

As with 8-bit division, numbers must often be converted to the proper form for the dividend. If a 16-bit unsigned number is placed in AX, DX must be cleared to zero. In the 80386 and above, the number is zero-extended by using the MOVZX instruction. If AX is a 16-bit signed number, the CWD (**convert word to doubleword**) instruction sign-extends it into a signed 32-bit number. If the 80386 and above is available, the MOVSX instruction can also be used to sign-extend a number.

EXAMPLE 5–15

```
0000 B8 FF9C            MOV  AX,–100        ;load a –100
0003 B9 0009            MOV  CX,9           ;load +9
0006 99                 CWD                 ;sign-extend
0007 F7 F9              IDIV CX
```

Example 5–15 shows the division of two 16-bit signed numbers. Here, –100 in AX is divided by +9 in CX. The CWD instruction converts the –100 in AX to –100 in DX–AX before the division. After the division, the results appear in DX–AX as a quotient of –11 in AX and a remainder of –1 in DX.

32-Bit Division.   The 80386 through the Pentium 4 processors perform 32-bit division on signed or unsigned numbers. The 64-bit contents of EDX–EAX are divided by the operand specified by the instruction, leaving a 32-bit quotient in EAX and a 32-bit remainder in EDX. Other than the size of the registers, this instruction functions in the same manner as the 8- and 16-bit divisions. Table 5–14 shows some 32-bit division instructions. The CDQ (**convert doubleword to quadword**) instruction is used before a signed division to convert the 32-bit contents of EAX into a 64-bit signed number in EDX–EAX.

TABLE 5–13   Example 16-bit division instructions.

Assembly Language	Operation
DIV CX	DX–AX is divided by CX; the unsigned quotient is AX and the unsigned remainder is in DX
IDIV SI	DX–AX is divided by SI; the signed quotient is in AX and the signed remainder is in DX
DIV NUMB	DX–AX is divided by the word contents of data segment memory NUMB; the unsigned quotient is in AX and the unsigned remainder is in DX

TABLE 5–14   Example 32-bit division instructions.

Assembly Language	Operation
DIV ECX	EDX–EAX is divided by ECX; the unsigned quotient is in EAX and the unsigned remainder is in EDX
IDIV DATA4	EDX–EAX is divided by the doubleword contents in data segment memory location DATA4; the signed quotient is in EAX and the signed remainder is in EDX
DIV DWORD PTR[EDI]	EDX–EAX is divided by the doubleword contents of the data segment memory location addressed by EDI; the unsigned quotient is in EAX and the unsigned remainder is in EDX

The Remainder.   What is done with the remainder after a division? There are a few possible choices. The remainder could be used to round the quotient or just dropped to truncate the quotient. If the division is unsigned, rounding requires that the remainder be compared with half the divisor to decide whether to round up the quotient. The remainder could also be converted to a fractional remainder.

EXAMPLE 5–16

```
0000 F6 F3                      DIV   BL            ;divide
0002 02 E4                      ADD   AH,AH         ;double remainder
0004 3A E3                      CMP   AH,BL         ;test for rounding
0006 72 02                      JB    NEXT          ;if OK
0008 FE C0                      INC   AL            ;round
000A                    NEXT:
```

Example 5–16 shows a sequence of instructions that divide AX by BL and round the unsigned result. This program doubles the remainder before comparing it with BL to decide whether to round the quotient. Here, an INC instruction rounds the contents of AL after the comparison.

Suppose that a fractional remainder is required instead of an integer remainder. A fractional remainder is obtained by saving the quotient. Next, the AL register is cleared to zero. The number remaining in AX is now divided by the original operand to generate a fractional remainder.

EXAMPLE 5–17

```
0000 B8 000D            MOV   AX,13           ;load 13
0003 B3 02              MOV   BL,2            ;load 2
0005 F6 F3              DIV   BL              ;13/2
0007 A2 0003 R          MOV   ANSQ,AL         ;save quotient
000A B0 00              MOV   AL,0            ;clear AL
000C F6 F3              DIV   BL              ;generate remainder
000E A2 0004 R          MOV   ANSR,AL         ;save remainder
```

Example 5–17 shows how 13 is divided by 2. The 8-bit quotient is saved in memory location ANSQ, and then AL is cleared. Next, the contents of AX are again divided by 2 to generate a fractional remainder. After the division, the AL register equals 80H. This is 10000000_2. If the binary point (radix) is placed before the leftmost bit of AL, the fractional remainder in AL is 0.10000000_2 or 0.5 decimal. The remainder is saved in memory location ANSR in this example.

64-Bit Division.   The Pentium 4 processor operated in 64-bit mode performs 64-bit division on signed or unsigned numbers. The 64-bit division uses the RDX:RAX register pair to hold the dividend and the quotient is found in RAX and the remainder is in RDX after the division. Table 5–15 illustrates a few 64-bit division instructions.

TABLE 5–15   Example 64-bit division instructions.

Assembly Language	Operation
DIV RCX	RDX–RAX is divided by RCX; the unsigned quotient is in RAX and the unsigned remainder is in RDX
IDIV DATA4	RDX–RAX is divided by the quadword contents in memory location DATA4; the signed quotient is in RAX and the signed remainder is in RDX
DIV QWORD PTR[RDI]	RDX–RAX is divided by the quadword contents of the memory location addressed by RDI; the unsigned quotient is in RAX and the unsigned remainder is in RDX

5-3   BCD AND ASCII ARITHMETIC

The microprocessor allows arithmetic manipulation of both BCD (**binary-coded decimal**) and ASCII (**American Standard Code for Information Interchange**) data. This is accomplished by instructions that adjust the numbers for BCD and ASCII arithmetic.

The BCD operations occur in systems such as point-of-sales terminals (e.g., cash registers) and others that seldom require complex arithmetic. The ASCII operations are performed on ASCII data used by many programs. In many cases, BCD or ASCII arithmetic is rarely used today, but some of the operations can be used for other purposes.

None of the instructions detailed in this section of the chapter function in the 64-bit mode of the Pentium 4 or Core2. In the future it appears that the BCD and ASCII instruction will become obsolete.

BCD Arithmetic

Two arithmetic techniques operate with BCD data: addition and subtraction. The instruction set provides two instructions that correct the result of a BCD addition and a BCD subtraction. The DAA (**decimal adjust after addition**) instruction follows BCD addition, and the DAS (**decimal adjust after subtraction**) follows BCD subtraction. Both instructions correct the result of the addition or subtraction so that it is a BCD number.

For BCD data, the numbers always appear in the packed BCD form and are stored as two BCD digits per byte. The adjustment instructions function only with the AL register after BCD addition and subtraction.

DAA Instruction.   The DAA instruction follows the ADD or ADC instruction to adjust the result into a BCD result. Suppose that DX and BX each contain 4-digit packed BCD numbers. Example 5–18 provides a short sample program that adds the BCD numbers in DX and BX, and stores the result in CX.

EXAMPLE 5–18

```
0000 BA 1234             MOV  DX,1234H       ;load 1234 BCD
0003 BB 3099             MOV  BX,3099H       ;load 3099 BCD
0006 8A C3               MOV  AL,BL          ;sum BL and DL
0008 02 C2               ADD  AL,DL
000A 27                  DAA
000B 8A C8               MOV  CL,AL          ;answer to CL
000D 9A C7               MOV  AL,BH          ;sum BH, DH and carry
000F 12 C6               ADC  AL,DH
0011 27                  DAA
0012 8A E8               MOV  CH,AL          ;answer to CH
```

Because the DAA instruction functions only with the AL register, this addition must occur 8 bits at a time. After adding the BL and DL registers, the result is adjusted with a DAA instruction before being stored in CL. Next, add BH and DH registers with carry; the result is then adjusted with DAA before being stored in CH. In this example, a 1234 is added to 3099 to generate a sum of 4333, which moves into CX after the addition. Note that 1234 BCD is the same as 1234H.

DAS Instruction. The DAS instruction functions as does the DAA instruction, except that it follows a subtraction instead of an addition. Example 5-19 is the same as Example 5–18, except that it subtracts instead of adds DX and BX. The main difference in these programs is that the DAA instructions change to DAS, and the ADD and ADC instructions change to SUB and SBB instructions.

EXAMPLE 5–19

```
0000 BA 1234          MOV  DX,1234H      ;load 1234 BCD
0003 BB 3099          MOV  BX,3099H      ;load 3099 BCD
0006 8A C3            MOV  AL,BL         ;subtract DL from BL
0008 2A C2            SUB  AL,DL
000A 2F               DAS
000B 8A C8            MOV  CL,AL         ;answer to CL
000D 9A C7            MOV  AL,BH         ;subtract DH
000F 1A C6            SBB  AL,DH
0011 2F               DAS
0012 8A E8            MOV  CH,AL         ;answer to CH
```

ASCII Arithmetic

The ASCII arithmetic instructions function with ASCII-coded numbers. These numbers range in value from 30H to 39H for the numbers 0–9. There are four instructions used with ASCII arithmetic operations: AAA (**ASCII adjust after addition**), AAD (**ASCII adjust before division**), AAM (**ASCII adjust after multiplication**), and AAS (**ASCII adjust after subtraction**). These instructions use register AX as the source and as the destination.

AAA Instruction. The addition of two one-digit ASCII-coded numbers will not result in any useful data. For example, if 31H and 39H are added, the result is 6AH. This ASCII addition (1 + 9) should produce a two-digit ASCII result equivalent to a 10 decimal, which is a 31H and a 30H in ASCII code. If the AAA instruction is executed after this addition, the AX register will contain a 0100H. Although this is not ASCII code, it can be converted to ASCII code by adding 3030H to AX which generates 3130H. The AAA instruction clears AH if the result is less than 10, and adds 1 to AH if the result is greater than 10.

EXAMPLE 5–20

```
0000 B8 0031          MOV  AX,31H        ;load ASCII 1
0003 04 39            ADD  AL,39H        ;add ASCII 9
0005 37               AAA                ;adjust sum
0006 05 3030          ADD  AX,3030H      ;answer to ASCII
```

Example 5–20 shows the way ASCII addition functions in the microprocessor. Please note that AH is cleared to zero before the addition by using the MOV AX,31H instruction. The operand of 0031H places 00H in AH and 31H into AL.

AAD Instruction. Unlike all other adjustment instructions, the AAD instruction appears before a division. The AAD instruction requires that the AX register contain a two-digit unpacked BCD number (not ASCII) before executing. After adjusting the AX register with AAD, it is divided by an unpacked BCD number to generate a single-digit result in AL with any remainder in AH.

Example 5–21 illustrates how 72 in unpacked BCD is divided by 9 to produce a quotient of 8. The 0702H loaded into the AX register is adjusted by the AAD instruction to 0048H. Notice that this converts a two-digit unpacked BCD number into a binary number so it can be divided with the binary division instruction (DIV). The AAD instruction converts the unpacked BCD numbers between 00 and 99 into binary.

EXAMPLE 5–21

```
0000 B3 09              MOV   BL,9           ;load divisor
0002 B8 0702            MOV   AX,702H        ;load dividend
0005 D5 0A              AAD                  ;adjust
0007 F6 F3              DIV   BL             ;divide
```

AAM Instruction.   The AAM instruction follows the multiplication instruction after multiplying two one-digit unpacked BCD numbers. Example 5–22 shows a short program that multiplies 5 times 5. The result after the multiplication is 0019H in the AX register. After adjusting the result with the AAM instruction, AX contains 0205H. This is an unpacked BCD result of 25. If 3030H is added to 0205H, it has an ASCII result of 3235H.

EXAMPLE 5–22

```
0000 B0 05              MOV   AL,5           ;load multiplicand
0002 B1 03              MOV   CL,3           ;load multiplier
0004 F6 E1              MUL   CL
0006 DA 0A              AAM                  ;adjust
```

The AAM instruction accomplishes this conversion by dividing AX by 10. The remainder is found in AL, and the quotient is in AH. Note that the second byte of the instruction contains 0AH. If the 0AH is changed to another value, AAM divides by the new value. For example, if the second byte is changed to 0BH, the AAM instruction divides by 11. This is accomplished with DB 0D4H, 0BH in place of AAM, which forces the AMM instruction to multiply by 11.

One side benefit of the AAM instruction is that AAM converts from binary to unpacked BCD. If a binary number between 0000H and 0063H appears in the AX register, the AAM instruction converts it to BCD. For example, if AX contains a 0060H before AAM, it will contain 0906H after AAM executes. This is the unpacked BCD equivalent of 96 decimal. If 3030H is added to 0906H, the result changes to ASCII code.

Example 5–23 shows how the l6-bit binary content of AX is converted to a four-digit ASCII character string by using division and the AAM instruction. Note that this works for numbers between 0 and 9999. First DX is cleared and then DX–AX is divided by 100. For example, if AX $= 245_{10}$, AX $= 2$ and DX $= 45$ after the division. These separate halves are converted to BCD using AAM, and then 3030H is added to convert to ASCII code.

EXAMPLE 5–23

```
0000 33 D2              XOR   DX,DX          ;clear DX
0002 B9 0064            MOV   CX,100         ;divide DX-AX by 100
0005 F7 F1              DIV   CX
0007 D4 0A              AAM                  ;convert to BCD
0009 05 3030            ADD   AX,3030H       ;convert to ASCII
000C 92                 XCHG  AX,DX          ;repeat for remainder
000D D4 0A              AAM
000F 05 3030            ADD   AX,3030H
```

Example 5–24 uses the DOS 21H function AH $=$ 02H to display a sample number in decimal on the video display using the AAM instruction. Notice how AAM is used to convert AL into BCD. Next, ADD AX,3030H converts the BCD code in AX into ASCII for display with DOS INT 21H. Once the data are converted to ASCII code, they are displayed by loading DL

with the most significant digit from AH. Next, the least significant digit is displayed from AL. Note that the DOS INT 2lH function calls change AL.

EXAMPLE 5–24

```
                           ;A program that displays the number in AL, loaded
                           ;with the first instruction (48H).
                           ;
                           .MODEL TINY        ;select tiny model
0000                       .CODE              ;start code segment
                           .STARTUP           ;start program
0100 B0 48                     MOV  AL,48H     ;load test data
0102 B4 00                     MOV  AH,0       ;clear AH
0104 D4 0A                     AAM             ;convert to BCD
0106 05 3030                   ADD  AX,3030H   ;convert to ASCII
0109 8A D4                     MOV  DL,AH      ;display most-significant digit
010B B4 02                     MOV  AH,2
010D 50                        PUSH AX
010E CD 21                     INT  21H
0110 58                        POP  AX
0111 8A D0                     MOV  DL,AL      ;display least-significant digit
0113 CD 21                     INT  21H
                           .EXIT              ;exit to DOS
                           END
```

AAS Instruction.   Like other ASCII adjust instructions, AAS adjusts the AX register after an ASCII subtraction. For example, suppose that 35H subtracts from 39H. The result will be 04H, which requires no correction. Here, AAS will modify neither AH nor AL. On the other hand, if 38H is subtracted from 37H, then AL will equal 09H and the number in AH will decrement by 1. This decrement allows multiple-digit ASCII numbers to be subtracted from each other.

5-4   BASIC LOGIC INSTRUCTIONS

The basic logic instructions include AND, OR, Exclusive-OR, and NOT. Another logic instruction is TEST, which is explained in this section of the text because the operation of the TEST instruction is a special form of the AND instruction. Also explained is the NEG instruction, which is similar to the NOT instruction.

Logic operations provide binary bit control in low-level software. The logic instructions allow bits to be set, cleared, or complemented. Low-level software appears in machine language or assembly language form and often controls the I/O devices in a system. All logic instructions affect the flag bits. Logic operations always clear the carry and overflow flags, while the other flags change to reflect the condition of the result.

When binary data are manipulated in a register or a memory location, the rightmost bit position is always numbered bit 0. Bit position numbers increase from bit 0 toward the left, to bit 7 for a byte, and to bit 15 for a word. A doubleword (32 bits) uses bit position 31 as its leftmost bit and a quadword (64-bits) uses bit position 63 as it leftmost bit.

AND

The AND operation performs logical multiplication, as illustrated by the truth table in Figure 5–3. Here, two bits, A and B, are ANDed to produce the result X. As indicated by the truth table, X is a logic 1 only when both A and B are logic 1s. For all other input combinations of A and B, X is a logic 0. It is important to remember that 0 AND anything is always 0, and 1 AND 1 is always 1.

FIGURE 5–3   (a) The truth table for the AND operation and (b) the logic symbol of an AND gate.

A	B	T
0	0	0
0	1	0
1	0	0
1	1	1

(a)

(b)

The AND instruction can replace discrete AND gates if the speed required is not too great, although this is normally reserved for embedded control applications. (Note that Intel has released the 80386EX embedded controller, which embodies the basic structure of the personal computer system.) With the 8086 microprocessor, the AND instruction often executes in about a microsecond. With newer versions, the execution speed is greatly increased. Take the 3.0 GHz Pentium with its clock time of 1/3 ns that executes up to three instruction per clock (1/9 ns per AND operation). If the circuit that the AND instruction replaces operates at a much slower speed than the microprocessor, the AND instruction is a logical replacement. This replacement can save a considerable amount of money. A single AND gate integrated circuit (74HCT08) costs approximately 40¢, while it costs less than 1/100¢ to store the AND instruction in read-only memory. Note that a logic circuit replacement such as this only appears in control systems based on microprocessors and does not generally find application in the personal computer.

The AND operation clears bits of a binary number. The task of clearing a bit in a binary number is called **masking**. Figure 5–4 illustrates the process of masking. Notice that the leftmost 4 bits clear to 0 because 0 AND anything is 0. The bit positions that AND with 1s do not change. This occurs because if a 1 ANDs with a 1, a 1 results; if a 1 ANDs with a 0, a 0 results.

The AND instruction uses any addressing mode except memory-to-memory and segment register addressing. Table 5–16 lists some AND instructions and comments about their operations.

An ASCII-coded number can be converted to BCD by using the AND instruction to mask off the leftmost four binary bit positions. This converts the ASCII 30H to 39H to 0–9. Example 5–25 shows a short program that converts the ASCII contents of BX into BCD. The AND instruction in this example converts two digits from ASCII to BCD simultaneously.

EXAMPLE 5–25

```
0000 BB 3135            MOV BX,3135H       ;load ASCII
0003 81 E3 0F0F         AND BX,0F0FH       ;mask BX
```

OR

The **OR operation** performs logical addition and is often called the *Inclusive-OR* function. The OR function generates a logic 1 output if any inputs are 1. A 0 appears at the output only when all inputs are 0. The truth table for the OR function appears in Figure 5–5. Here, the inputs A and

FIGURE 5–4   The operation of the AND function showing how bits of a number are cleared to zero.

```
   x x x x  x x x x   Unknown number
•  0 0 0 0  1 1 1 1   Mask
   ─────────────────
   0 0 0 0  x x x x   Result
```

TABLE 5–16 Example AND instructions.

Assembly Language	Operation
AND AL,BL	AL = AL and BL
AND CX,DX	CX = CX and DX
AND ECX,EDI	ECX = ECX and EDI
AND RDX,RBP	RDX = RDX and RBP (64-bit mode)
AND CL,33H	CL = CL and 33H
AND DI,4FFFH	DI = DI and 4FFFH
AND ESI,34H	ESI = ESI and 34H
AND RAX,1	RAX = RAX and 1 (64-bit mode)
AND AX,[DI]	The word contents of the data segment memory location addressed by DI are ANDed with AX
AND ARRAY[SI],AL	The byte contents of the data segment memory location addressed by ARRAY plus SI are ANDed with AL
AND [EAX],CL	CL is ANDed with the byte contents of the data segment memory location addressed by ECX

FIGURE 5–5 (a) The truth table for the OR operation and (b) the logic symbol of an OR gate.

A	B	T
0	0	0
0	1	1
1	0	1
1	1	1

(a)

(b)

B OR together to produce the X output. It is important to remember that 1 ORed with anything yields a 1.

In embedded controller applications, the OR instruction can also replace discrete OR gates. This results in considerable savings because a quad, two-input OR gate (74HCT32) costs about 40¢, while the OR instruction costs less than 1/100¢ to store in a read-only memory.

Figure 5–6 shows how the OR gate sets (1) any bit of a binary number. Here, an unknown number (XXXX XXXX) ORs with a 0000 1111 to produce a result of XXXX 1111. The rightmost 4 bits set, while the leftmost 4 bits remain unchanged. The OR operation sets any bit; the AND operation clears any bit.

The OR instruction uses any of the addressing modes allowed to any other instruction except segment register addressing. Table 5–17 illustrates several example OR instructions with comments about their operation.

FIGURE 5–6 The operation of the OR function showing how bits of a number are set to one.

```
  x x x x  x x x x   Unknown number
+ 0 0 0 0  1 1 1 1   Mask
─────────────────
  x x x x  1 1 1 1   Result
```

TABLE 5–17   Example OR instructions.

Assembly Language	Operation
OR AH,BL	AL = AL or BL
OR SI,DX	SI = SI or DX
OR EAX,EBX	EAX = EAX or EBX
OR R9,R10	R9 = R9 or R10 (64-bit mode)
OR DH,0A3H	DH = DH or 0A3H
OR SP,990DH	SP = SP or 990DH
OR EBP,10	EBP = EBP or 10
OR RBP,1000H	RBP = RBP or 1000H (64-bit mode)
OR DX,[BX]	DX is ORed with the word contents of data segment memory location addressed by BX
OR DATES[DI + 2],AL	The byte contents of the data segment memory location addressed by DI plus 2 are ORed with AL

Suppose that two BCD numbers are multiplied and adjusted with the AAM instruction. The result appears in AX as a two-digit unpacked BCD number. Example 5–26 illustrates this multiplication and shows how to change the result into a two-digit ASCII-coded number using the OR instruction. Here, OR AX,3030H converts the 0305H found in AX to 3335H. The OR operation can be replaced with an ADD AX,3030H to obtain the same results.

EXAMPLE 5–26

```
0000 B0 05              MOV   AL,5            ;load data
0002 B3 07              MOV   BL,7
0004 F6 E3              MUL   BL
0006 D4 0A              AAM                   ;adjust
0008 0D 3030            OR    AX,3030H        ;convert to ASCII
```

Exclusive-OR

The **Exclusive-OR** instruction (XOR) differs from Inclusive-OR (OR). The difference is that a 1,1 condition of the OR function produces a 1; the 1,1 condition of the Exclusive-OR operation produces a 0. The Exclusive-OR operation excludes this condition; the Inclusive-OR includes it.

Figure 5–7 shows the truth table of the Exclusive-OR function. (Compare this with Figure 5–5 to appreciate the difference between these two OR functions.) If the inputs of the

FIGURE 5–7   (a) The truth table for the Exclusive-OR operation and (b) the logic symbol of an Exclusive-OR gate.

A	B	T
0	0	0
0	1	1
1	0	1
1	1	0

(a)

(b)

TABLE 5–18 Example Exclusive-OR instructions.

Assembly Language	Operation
XOR CH,DL	CH = CH xor DL
XOR SI,BX	SI = SI xor BX
XOR EBX,EDI	EBX = EBX xor EDI
XOR RAX,RBX	RAX = RAX xor RBX (64-bit mode)
XOR AH,0EEH	AH = AH xor 0EEH
XOR DI,00DDH	DI = DI xor 00DDH
XOR ESI,100	ESI = ESI xor 100
XOR R12,20	R12 = R12 xor 20 (64-bit mode)
XOR DX,[SI]	DX is Exclusive-ORed with the word contents of the data segment memory location addressed by SI
XOR DEAL[BP+2],AH	AH is Exclusive-ORed with the byte contents of the stack segment memory location addressed by BP plus 2

Exclusive-OR function are both 0 or both 1, the output is 0. If the inputs are different, the output is 1. Because of this, the Exclusive-OR is sometimes called a comparator.

The XOR instruction uses any addressing mode except segment register addressing. Table 5–18 lists several Exclusive-OR instructions and their operations.

As with the AND and OR functions, Exclusive-OR can replace discrete logic circuitry in embedded applications. The 74HCT86 quad, two-input Exclusive-OR gate is replaced by one XOR instruction. The 74HCT86 costs about 40¢, whereas the instruction costs less than 1/100¢ to store in the memory. Replacing just one 74HCT86 saves a considerable amount of money, especially if many systems are built.

The Exclusive-OR instruction is useful if some bits of a register or memory location must be inverted. This instruction allows part of a number to be inverted or complemented. Figure 5–8 shows how just part of an unknown quantity can be inverted by XOR. Notice that when a 1 Exclusive-ORs with X, the result is X. If a 0 Exclusive-ORs with X, the result is X.

Suppose that the leftmost 10 bits of the BX register must be inverted without changing the rightmost 6 bits. The XOR BX,0FFC0H instruction accomplishes this task. The AND instruction clears (0) bits, the OR instruction sets (1) bits, and now the Exclusive-OR instruction inverts bits. These three instructions allow a program to gain complete control over any bit stored in any register or memory location. This is ideal for control system applications in which equipment must be turned on (1), turned off (0), and toggled from on to off or off to on.

A common use for the Exclusive-OR instruction is to clear a register to zero. For example, the XOR CH,CH instruction clears register CH to 00H and requires 2 bytes of memory to store the instruction. Likewise, the MOV CH, 00H instruction also clears CH to 00H, but requires 3 bytes of memory. Because of this saving, the XOR instruction is often used to clear a register in place of a move immediate.

Example 5–27 shows a short sequence of instructions that clears bits 0 and 1 of CX, sets bits 9 and 10 of CX, and inverts bit 12 of CX. The OR instruction is used to set bits, the AND instruction is used to clear bits, and the XOR instruction inverts bits.

FIGURE 5–8 The operation of the Exclusive-OR function showing how bits of a number are inverted.

```
  x x x x  x x x x   Unknown number
⊕ 0 0 0 0  1 1 1 1   Mask
  ─────────────────
  x x x x  x̄ x̄ x̄ x̄   Result
```

EXAMPLE 5–27

```
0000 81 C9 0600          OR  CX,0600H        ;set bits 9 and 10
0004 83 E1 FC            AND CX,0FFFCH        ;clear bits 0 and 1
0007 81 F1 1000          XOR CX,1000H         ;invert bit 12
```

Test and Bit Test Instructions

The **TEST instruction** performs the AND operation. The difference is that the AND instruction changes the destination operand, whereas the TEST instruction does not. A TEST only affects the condition of the flag register, which indicates the result of the test. The TEST instruction uses the same addressing modes as the AND instruction. Table 5–19 lists some TEST instructions and their operations.

The TEST instruction functions in the same manner as a CMP instruction. The difference is that the TEST instruction normally tests a single bit (or occasionally multiple bits), whereas the CMP instruction tests the entire byte, word, or doubleword. The zero flag (Z) is a logic 1 (indicating a zero result) if the bit under test is a zero, and Z = 0 (indicating a nonzero result) if the bit under test is not zero.

Usually the TEST instruction is followed by either the JZ (jump if zero) or JNZ (jump if not zero) instruction. The destination operand is normally tested against immediate data. The value of immediate data is 1 to test the rightmost bit position, 2 to test the next bit, 4 for the next, and so on.

Example 5–28 lists a short program that tests the rightmost and leftmost bit positions of the AL register. Here, 1 selects the rightmost bit and 128 selects the leftmost bit. (Note: A 128 is an 80H.) The JNZ instruction follows each test to jump to different memory locations, depending on the outcome of the tests. The JNZ instruction jumps to the operand address (RIGHT or LEFT in the example) if the bit under test is not zero.

EXAMPLE 5–28

```
0000 A8 01               TEST AL,1           ;test right bit
0002 75 1C               JNZ  RIGHT          ;if set
0004 A8 80               TEST AL,128         ;test left bit
0006 75 38               JNZ  LEFT           ;if set
```

The 80386 through the Pentium 4 processors contain additional test instructions that test single bit positions. Table 5–20 lists the four different bit test instructions available to these microprocessors.

All four forms of the bit test instruction test the bit position in the destination operand selected by the source operand. For example, the BT AX,4 instruction tests bit position 4 in AX. The result of the test is located in the carry flag bit. If bit position 4 is a 1, carry is set; if bit position 4 is a 0, carry is cleared.

TABLE 5–19   Example TEST instructions.

Assembly Language	Operation
TEST DL,DH	DL is ANDed with DH
TEST CX,BX	CX is ANDed with BX
TEST EDX,ECX	EDX is ANDed with ECX
TEST RDX,R15	RDX is ANDed with R15 (64-bit mode)
TEST AH,4	AH is ANDed with 4
TEST EAX,256	EAX is ANDed with 256

TABLE 5–20 Bit test instructions.

Assembly Language	Operation
BT	Tests a bit in the destination operand specified by the source operand
BTC	Tests and complements a bit in the destination operand specified by the source operand
BTR	Tests and resets a bit in the destination operand specified by the source operand
BTS	Tests and sets a bit in the destination operand specified by the source operand

The remaining 3-bit test instructions also place the bit under test into the carry flag and change the bit under test afterward. The BTC AX,4 instruction complements bit position 4 after testing it, the BTR AX,4 instruction clears it (0) after the test, and the BTS AX,4 instruction sets it (1) after the test.

Example 5–29 repeats the sequence of instructions listed in Example 5–27. Here, the BTR instruction clears bits in CX, BTS sets bits in CX, and BTC inverts bits in CX.

EXAMPLE 5–29

```
0000 0F BA E9 09          BTS   CX,9            ;set bit 9
0004 0F BA E9 0A          BTS   CX,10           ;set bit 10
0008 0F BA F1 00          BTR   CX,0            ;clear bit 0
000C 0F BA F1 01          BTR   CX,1            ;clear bit 1
0010 0F BA F9 0C          BTC   CX,12           ;complement bit 12
```

NOT and NEG

Logical inversion, or the one's complement (NOT), and arithmetic sign inversion, or the two's complement (NEG), are the last two logic functions presented (except for shift and rotate in the next section of the text). These are two of a few instructions that contain only one operand. Table 5–21 lists some variations of the NOT and NEG instructions. As with most other instructions, NOT and NEG can use any addressing mode except segment register addressing.

TABLE 5–21 Example NOT and NEG instructions.

Assembly Language	Operation
NOT CH	CH is one's complemented
NEG CH	CH is two's complemented
NEG AX	AX is two's complemented
NOT EBX	EBX is one's complemented
NEG ECX	ECX is two's complemented
NOT RAX	RAX is one's complemented (64-bit mode)
NOT TEMP	The contents of data segment memory location TEMP is one's complemented
NOT BYTE PTR[BX]	The byte contents of the data segment memory location addressed by BX are one's complemented

The NOT instruction inverts all bits of a byte, word, or doubleword. The NEG instruction two's complements a number, which means that the arithmetic sign of a signed number changes from positive to negative or from negative to positive. The NOT function is considered logical, and the NEG function is considered an arithmetic operation.

5-5   SHIFT AND ROTATE

Shift and rotate instructions manipulate binary numbers at the binary bit level, as did the AND, OR, Exclusive-OR, and NOT instructions. Shifts and rotates find their most common applications in low-level software used to control I/O devices. The microprocessor contains a complete complement of shift and rotate instructions that are used to shift or rotate any memory data or register.

Shift

Shift instructions position or move numbers to the left or right within a register or memory location. They also perform simple arithmetic such as multiplication by powers of 2^{+n} (left shift) and division by powers of 2^{-n} (right shift). The microprocessor's instruction set contains four different shift instructions: Two are logical shifts and two are arithmetic shifts. All four shift operations appear in Figure 5–9.

Notice in Figure 5–9 that there are two right shifts and two left shifts. The logical shifts move a 0 into the rightmost bit position for a logical left shift and a 0 into the leftmost bit position for a logical right shift. There are also two arithmetic shifts. The arithmetic shift left and logical left shift are identical. The arithmetic right shift and logical right shift are different because the arithmetic right shift copies the sign-bit through the number, whereas the logical right shift copies a 0 through the number.

FIGURE 5–9   The shift instructions showing the operation and direction of the shift.
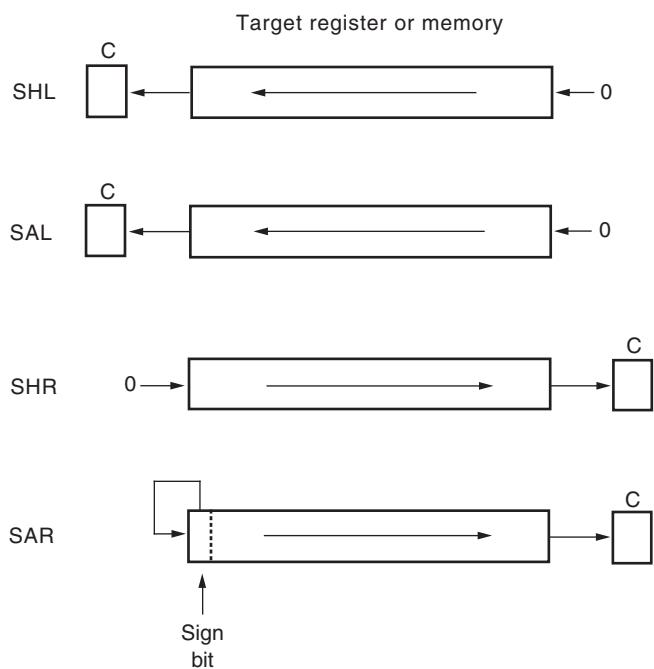
TABLE 5–22   Example shift instructions.

Assembly Language	Operation
SHL AX,1	AX is logically shifted left 1 place
SHR BX,12	BX is logically shifted right 12 places
SHR ECX,10	ECX is logically shifted right 10 places
SHL RAX,50	RAX is logically shifted left 50 places (64-bit mode)
SAL DATA1,CL	The contents of data segment memory location DATA1 are arithmetically shifted left the number of spaces specified by CL
SHR RAX,CL	RAX is logically shifted right the number of spaces specified by CL (64-bit mode)
SAR SI,2	SI is arithmetically shifted right 2 places
SAR EDX,14	EDX is arithmetically shifted right 14 places

Logical shift operations function with unsigned numbers, and arithmetic shifts function with signed numbers. Logical shifts multiply or divide unsigned data, and arithmetic shifts multiply or divide signed data. A shift left always multiplies by 2 for each bit position shifted, and a shift right always divides by 2 for each bit position shifted. Shifting a number two places, to the left or right, multiplies or divides by 4.

Table 5–22 illustrates some addressing modes allowed for the various shift instructions. There are two different forms of shifts that allow any register (except the segment register) or memory location to be shifted. One mode uses an immediate shift count, and the other uses register CL to hold the shift count. Note that CL must hold the shift count. When CL is the shift count, it does not change when the shift instruction executes. Note that the shift count is a modulo-32 count, which means that a shift count of 33 will shift the data one place (33/32 = remainder of 1). The same applies to a 64-bit number, but the shift count is modulo-64.

Example 5–30 shows how to shift the DX register left 14 places in two different ways. The first method uses an immediate shift count of 14. The second method loads 14 into CL and then uses CL as the shift count. Both instructions shift the contents of the DX register logically to the left 14 binary bit positions or places.

EXAMPLE 5–30

```
0000 C1 E2 0E            SHL DX,14

                 or

0003 B1 0E               MOV CL,14
0005 D3 E2               SHL DX,CL
```

Suppose that the contents of AX must be multiplied by 10, as shown in Example 5–31. This can be done in two ways: by the MUL instruction or by shifts and additions. A number is doubled when it shifts left one place. When a number is doubled, and then added to the number times 8, the result is 10 times the number. The number 10 decimal is 1010 in binary. A logic 1 appears in both the 2's and 8's positions. If 2 times the number is added to 8 times the number, the result is 10 times the number. Using this technique, a program can be written to multiply by any constant. This technique often executes faster than the multiply instruction found in earlier versions of the Intel microprocessor.

EXAMPLE 5–31

```
                    ;Multiply AX by 10 (1010)
                    ;
0000 D1 E0              SHL  AX,1               ;AX times 2
0002 8B D8              MOV  BX,AX
0004 C1 E0 02           SHL  AX,2               ;AX times 8
0007 03 C3              ADD  AX,BX              ;AX times 10
                    ;
                    ;Multiply AX by 18 (10010)
                    ;
0009 D1 E0              SHL  AX,1               ;AX times 2
000B 8B D8              MOV  BX,AX
000D C1 E0 03           SHL  AX,3               ;AX times 16
0010 03 C3              ADD  AX,BX              ;AX times 18
                    ;
                    ;Multiply AX by 5 (101)
                    ;
0012 8B D8              MOV  BX,AX
0014 C1 E0 02           SHL  AX,2               ;AX times 4
0017 03 C3              ADD  AX,BX              ;AX times 5
```

Double-Precision Shifts (80386–Core2 Only). The 80386 and above contain two double preci-
sion shifts: SHLD (shift left) and SHRD (shift right). Each instruction contains three operands,
instead of the two found with the other shift instructions. Both instructions function with two
16-or 32-bit registers, or with one 16- or 32-bit memory location and a register.

The SHRD AX,BX,12 instruction is an example of the double-precision shift right instruc-
tion. This instruction logically shifts AX right by 12 bit positions. The rightmost 12 bits of BX
shift into the leftmost 12 bits of AX. The contents of BX remain unchanged by this instruction.
The shift count can be an immediate count, as in this example, or it can be found in register CL,
as with other shift instructions.

The SHLD EBX,ECX,16 instruction shifts EBX left. The leftmost 16 bits of ECX fill the
rightmost 16 bits of EBX after the shift. As before, the contents of ECX, the second operand,
remain unchanged. This instruction, as well as SHRD, affects the flag bits.

Rotate

Rotate instructions position binary data by rotating the information in a register or memory loca-
tion, either from one end to another or through the carry flag. They are often used to shift or posi-
tion numbers that are wider than 16 bits in the 8086–80286 microprocessors or wider than 32 bits
in the 80386 through the Core2. The four available rotate instructions appear in Figure 5–10.

Numbers rotate through a register or memory location, through the C flag (carry), or through
a register or memory location only. With either type of rotate instruction, the programmer can select
either a left or a right rotate. Addressing modes used with rotate are the same as those used with
shifts. A rotate count can be immediate or located in register CL. Table 5–23 lists some of the pos-
sible rotate instructions. If CL is used for a rotate count, it does not change. As with shifts, the count
in CL is a modulo-32 count for a 32-bit operation and modulo-64 for a 64-bit operation.

Rotate instructions are often used to shift wide numbers to the left or right. The program
listed in Example 5–32 shifts the 48-bit number in registers DX, BX, and AX left one binary
place. Notice that the least significant 16 bits (AX) shift left first. This moves the leftmost bit of
AX into the carry flag bit. Next, the rotate BX instruction rotates carry into BX, and its leftmost
bit moves into carry. The last instruction rotates carry into DX, and the shift is complete.

EXAMPLE 5–32

```
0000 D1 E0       SHL AX,1
0002 D1 D3       RCL BX,1
0004 D1 D2       RCL DX,1
```

FIGURE 5–10 The rotate instructions showing the direction and operation of each rotate.
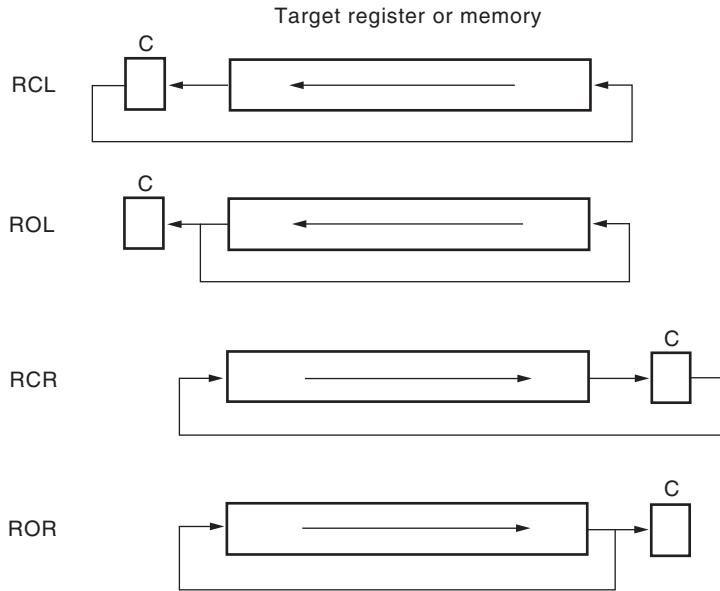
TABLE 5–23 Example rotate instructions.

Assembly Language	Operation
ROL SI,14	SI rotates left 14 places
RCL BL,6	BL rotates left through carry 6 places
ROL ECX,18	ECX rotates left 18 places
ROL RDX,40	RDX rotates left 40 places
RCR AH,CL	AH rotates right through carry the number of places specified by CL
ROR WORD PTR[BP],2	The word contents of the stack segment memory location addressed by BP rotate right 2 places

Bit Scan Instructions

Although the bit scan instructions don't shift or rotate numbers, they do scan through a number searching for a 1-bit. Because this is accomplished within the microprocessor by shifting the number, bit scan instructions are included in this section of the text.

The bit scan instructions BSF (bit scan forward) and BSR (bit scan reverse) are available only in the 80386–Pentium 4 processors. Both forms scan through the source number, searching for the first 1-bit. The BSF instruction scans the number from the leftmost bit toward the right, and BSR scans the number from the rightmost bit toward the left. If a 1-bit is encountered, the zero flag is set and the bit position number of the 1-bit is placed into the destination operand. If no 1-bit is encountered (i.e., the number contains all zeros), the zero flag is cleared. Thus, the result is not-zero if no 1-bit is encountered.

For example, if EAX = 60000000H and the BSF EBX,EAX instruction executes, the number is scanned from the leftmost bit toward the right. The first 1-bit encountered is at bit position 30, which is placed into EBX and the zero flag bit is set. If the same value for EAX is used for the BSR instruction, the EBX register is loaded with 29 and the zero flag bit is set.

5-6   ## STRING COMPARISONS

As illustrated in Chapter 4, the string instructions are very powerful because they allow the programmer to manipulate large blocks of data with relative ease. Block data manipulation occurs with the string instructions MOVS, LODS, STOS, INS, and OUTS. In this section, additional string instructions that allow a section of memory to be tested against a constant or against another section of memory are discussed. To accomplish these tasks, use the SCAS (**string scan**) or CMPS (**string compare**) instructions.

SCAS

The SCAS (string scan instruction) compares the AL register with a byte block of memory, the AX register with a word block of memory, or the EAX register (80386-Core2) with a double-word block of memory. The SCAS instruction subtracts memory from AL, AX, or EAX without affecting either the register or the memory location. The opcode used for byte comparison is SCASB, the opcode used for the word comparison is SCASW, and the opcode used for a doubleword comparison is SCASD. In all cases, the contents of the extra segment memory location addressed by DI is compared with AL, AX, or EAX. Recall that this default segment (ES) cannot be changed with a segment override prefix.

Like the other string instructions, SCAS instructions use the direction flag (D) to select either auto-increment or auto-decrement operation for DI. They also repeat if prefixed by a conditional repeat prefix.

Suppose that a section of memory is 100 bytes long and begins at location BLOCK. This section of memory must be tested to see whether any location contains 00H. The program in Example 5–33 shows how to search this part of memory for 00H using the SCASB instruction. In this example, the SCASB instruction has an REPNE (**repeat while not equal**) prefix. The REPNE prefix causes the SCASB instruction to repeat until either the CX register reaches 0, or until an equal condition exists as the outcome of the SCASB instruction's comparison. Another conditional repeat prefix is REPE (**repeat while equal**). With either repeat prefix, the contents of CX decrements without affecting the flag bits. The SCASB instruction and the comparison it makes change the flags.

EXAMPLE 5–33

```
0000 BF 0011 R           MOV   DI,OFFSET BLOCK    ;address data
0003 FC                  CLD                      ;auto-increment
0004 B9 0064             MOV   CX,100             ;load counter
0007 32 C0               XOR   AL,AL              ;clear AL
0009 F2/AE               REPNE SCASB
```

Suppose that you must develop a program that skips ASCII-coded spaces in a memory array. (This task appears in the procedure listed in Example 5–34.) This procedure assumes that the DI register already addresses the ASCII-coded character string and that the length of the string is 256 bytes or fewer. Because this program is to skip spaces (20H), the REPE prefix is used with a SCASB instruction. The SCASB instruction repeats the comparison, searching for a 20H, as long as an equal condition exists.

EXAMPLE 5–34

```
0000 FC                  CLD               ;auto-increment
0001 B9 0100             MOV   CX,256      ;load counter
0004 B0 20               MOV   AL,20H      ;get space
0006 F3/AE               REPE SCASB
```

CMPS

The CMPS (compare strings instruction) always compares two sections of memory data as bytes (CMPSB), words (CMPSW), or doublewords (CMPSD). Note that only the 80386 through Core2 can use doublewords. In the Pentium 4 or Core2 operated in 64-bit mode, a CMPSQ instruction uses quadwords. The contents of the data segment memory location addressed by SI are compared with the contents of the extra segment memory location addressed by DI. The CMPS instruction increments or decrements both SI and DI. The CMPS instruction is normally used with either the REPE or REPNE prefix. Alternates to these prefixes are REPZ (repeat while zero) and REPNZ (repeat while not zero), but usually the REPE or REPNE prefixes are used in programming.

Example 5–35 illustrates a short procedure that compares two sections of memory searching for a match. The CMPSB instruction is prefixed with REPE. This causes the search to continue as long as an equal condition exists. When the CX register becomes 0 or an unequal condition exists, the CMPSB instruction stops execution. After the CMPSB instruction ends, the CX register is 0 or the flags indicate an equal condition when the two strings match. If CX is not 0 or the flags indicate a not-equal condition, the strings do not match.

EXAMPLE 5–35

```
0000 BE 0075 R              MOV  SI,OFFSET LINE   ;address LINE
0003 BF 007F R              MOV  DI,OFFSET TABLE  ;address TABLE
0006 FC                     CLD                   ;auto-increment
0007 B9 000A                MOV  CX,10            ;load counter
000A F3/A6                  REPE CMPSB            ;search
```

5–7   SUMMARY

1. Addition (ADD) can be 8, 16, 32, or 64 bits. The ADD instruction allows any addressing mode except segment register addressing. Most flags (C, A, S, Z, P, and O) change when the ADD instruction executes. A different type of addition, add-with-carry (ADC), adds two operands and the contents of the carry flag (C). The 80486 through the Core2 processors have an additional instruction (XADD) that combines an addition with an exchange.
2. The increment instruction (INC) adds 1 to the byte, word, or doubleword contents of a register or memory location. The INC instruction affects the same flag bits as ADD except the carry flag. The BYTE PTR, WORD PTR, DWORD PTR, or QWORD PTR directives appear with the INC instruction when the contents of a memory location are addressed by a pointer.
3. Subtraction (SUB) is a byte, word, doubleword, or quadword and is performed on a register or a memory location. The only form of addressing not allowed by the SUB instruction is segment register addressing. The subtract instruction affects the same flags as ADD and subtracts carry if the SBB form is used.
4. The decrement (DEC) instruction subtracts 1 from the contents of a register or a memory location. The only addressing modes not allowed with DEC are immediate or segment register addressing. The DEC instruction does not affect the carry flag and is often used with BYTE PTR, WORD PTR, DWORD PTR, or QWORD PTR.
5. The comparison (CMP) instruction is a special form of subtraction that does not store the difference; instead, the flags change to reflect the difference. Comparison is used to compare an entire byte or word located in any register (except segment) or memory location.

An additional comparison instruction (CMPXCHG), which is a combination of comparison and exchange instructions, is found in the 80486–Core2 processors. In the Pentium–Core2 processors, the CMPXCHG8B instruction compares and exchanges quadword data. In the 64-bit Pentium 4 and Core2, a COMPXCHG16B instruction is available.

6. Multiplication is byte, word, or doubleword, and it can be signed (IMUL) or unsigned (MUL). The 8-bit multiplication always multiplies register AL by an operand with the product found in AX. The 16-bit multiplication always multiplies register AX by an operand with the product found in DX–AX. The 32-bit multiply always multiplies register EAX by an operand with the product found in EDX–EAX. A special IMUL immediate instruction exists on the 80186–Core2 processors that contains three operands. For example, the IMUL BX,CX,3 instruction multiplies CX by 3 and leaves the product in BX. In the Pentium 4 and Core2 with 64-bit mode enabled, multiplication is 64 bits.

7. Division is byte, word, or doubleword, and it can be signed (IDIV) or unsigned (DIV). For an 8-bit division, the AX register divides by the operand, after which the quotient appears in AL and the remainder appears in AH. In the 16-bit division, the DX–AX register divides by the operand, after which the AX register contains the quotient and DX contains the remainder. In the 32-bit division, the EDX–EAX register is divided by the operand, after which the EAX register contains the quotient and the EDX register contains the remainder. Note that the remainder after a signed division always assumes the sign of the dividend.

8. BCD data add or subtract in packed form by adjusting the result of the addition with DAA or the subtraction with DAS. ASCII data are added, subtracted, multiplied, or divided when the operations are adjusted with AAA, AAS, AAM, and AAD. These instructions do not function in the 64-bit mode.

9. The AAM instruction has an interesting added feature that allows it to convert a binary number into unpacked BCD. This instruction converts a binary number between 00H–63H into unpacked BCD in AX. The AAM instruction divides AX by 10, and leaves the remainder in AL and quotient in AH. These instructions do not function in the 64-bit mode.

10. The AND, OR, and Exclusive-OR instructions perform logic functions on a byte, word, or doubleword stored in a register or memory location. All flags change with these instructions, with carry (C) and overflow (O) cleared.

11. The TEST instruction performs the AND operation, but the logical product is lost. This instruction changes the flag bits to indicate the outcome of the test.

12. The NOT and NEG instructions perform logical inversion and arithmetic inversion. The NOT instruction one's complements an operand, and the NEG instruction two's complements an operand.

13. There are eight different shift and rotate instructions. Each of these instructions shifts or rotates a byte, word, or doubleword register or memory data. These instructions have two operands: The first is the location of the data shifted or rotated, and the second is an immediate shift or rotate count or CL. If the second operand is CL, the CL register holds the shift or rotate count. In the 80386 through the Core2 processors, two additional double-precision shifts (SHRD and SHLD) exist.

14. The scan string (SCAS) instruction compares AL, AX, or EAX with the contents of the extra segment memory location addressed by DI.

15. The string compare (CMPS) instruction compares the byte, word, or doubleword contents of two sections of memory. One section is addressed by DI in the extra segment, and the other is addressed by SI in the data segment.

16. The SCAS and CMPS instructions repeat with the REPE or REPNE prefixes. The REPE prefix repeats the string instruction while an equal condition exists, and the REPNE repeats the string instruction while a not-equal condition exists.

5–8   **QUESTIONS AND PROBLEMS**

1. Select an ADD instruction that will:
   (a) add BX to AX
   (b) add 12H to AL
   (c) add EDI and EBP
   (d) add 22H to CX
   (e) add the data addressed by SI to AL
   (f) add CX to the data stored at memory location FROG
   (g) add 234H to RCX
2. What is wrong with the ADD RCX,AX instruction?
3. Is it possible to add CX to DS with the ADD instruction?
4. If AX = 1001H and DX = 20FFH, list the sum and the contents of each flag register bit (C, A, S, Z, and O) after the ADD AX,DX instruction executes.
5. Develop a short sequence of instructions that adds AL, BL, CL, DL, and AH. Save the sum in the DH register.
6. Develop a short sequence of instructions that adds AX, BX, CX, DX, and SP. Save the sum in the DI register.
7. Develop a short sequence of instructions that adds ECX, EDX, and ESI. Save the sum in the EDI register.
8. Develop a short sequence of instructions that adds RCX, RDX, and RSI. Save the sum in the R12 register.
9. Select an instruction that adds BX to DX, and also adds the contents of the carry flag (C) to the result.
10. Choose an instruction that adds 1 to the contents of the SP register.
11. What is wrong with the INC [BX] instruction?
12. Select a SUB instruction that will:
    (a) subtract BX from CX
    (b) subtract 0EEH from DH
    (c) subtract DI from SI
    (d) subtract 3322H from EBP
    (e) subtract the data address by SI from CH
    (f) subtract the data stored 10 words after the location addressed by SI from DX
    (g) subtract AL from memory location FROG
    (h) subtract R9 from R10
13. If DL = 0F3H and BH = 72H, list the difference after BH is subtracted from DL and show the contents of the flag register bits.
14. Write a short sequence of instructions that subtracts the numbers in DI, SI, and BP from the AX register. Store the difference in register BX.
15. Choose an instruction that subtracts 1 from register EBX.
16. Explain what the SBB [DI–4],DX instruction accomplishes.
17. Explain the difference between the SUB and CMP instruction.
18. When two 8-bit numbers are multiplied, where is the product found?
19. When two 16-bit numbers are multiplied, what two registers hold the product? Show the registers that contain the most and least significant portions of the product.
20. When two numbers multiply, what happens to the O and C flag bits?
21. Where is the product stored for the MUL EDI instruction?
22. Write a sequence of instructions that cube the 8-bit number found in DL. Load DL with a 5 initially, and make sure that your result is a l6-bit number.

23. What is the difference between the IMUL and MUL instructions?
24. Describe the operation of the IMUL BX,DX,100H instruction.
25. When 8-bit numbers are divided, in which register is the dividend found?
26. When l6-bit numbers are divided, in which register is the quotient found?
27. When 64-bit numbers are divided, in which register is the quotient found?
28. What errors are detected during a division?
29. Explain the difference between the IDIV and DIV instructions.
30. Where is the remainder found after an 8-bit division?
31. Where is the quotient found after a 64-bit division?
32. Write a short sequence of instructions that divides the number in BL by the number in CL and then multiplies the result by 2. The final answer must be a 16-bit number stored in the DX register.
33. Which instructions are used with BCD arithmetic operations?
34. Explain how the AAM instruction converts from binary to BCD.
35. Which instructions are used with ASCII arithmetic operations?
36. Develop a sequence of instructions that converts the unsigned number in AX (values of 0–65535) into a 5-digit BCD number stored in memory, beginning at the location addressed by the BX register in the data segment. Note that the most significant character is stored first and no attempt is made to blank leading zeros.
37. Develop a sequence of instructions that adds the 8-digit BCD number in AX and BX to the 8-digit BCD number in CX and DX. (AX and CX are the most significant registers. The result must be found in CX and DX after the addition.)
38. Does the AAM instruction function in the 64-bit mode?
39. Select an AND instruction that will:
    (a) AND BX with DX and save the result in BX
    (b) AND 0EAH with DH
    (c) AND DI with BP and save the result in DI
    (d) AND 1122H with EAX
    (e) AND the data addressed by BP with CX and save the result in memory
    (f) AND the data stored in four words before the location addressed by SI with DX and save the result in DX
    (g) AND AL with memory location WHAT and save the result at location WHAT
40. Develop a short sequence of instructions that clears (0) the three leftmost bits of DH without changing the remainder of DH and stores the result in BH.
41. Select an OR instruction that will:
    (a) OR BL with AH and save the result in AH
    (b) OR 88H with ECX
    (c) OR DX with SI and save the result in SI
    (d) OR 1122H with BP
    (e) OR the data addressed by RBX with RCX and save the result in memory
    (f) OR the data stored 40 bytes after the location addressed by BP with AL and save the result in AL
    (g) OR AH with memory location WHEN and save the result in WHEN
42. Develop a short sequence of instructions that sets (1) the rightmost 5 bits of DI without changing the remaining bits of DI. Save the results in SI.
43. Select the XOR instruction that will:
    (a) XOR BH with AH and save the result in AH
    (b) XOR 99H with CL
    (c) XOR DX with DI and save the result in DX
    (d) XOR lA23H with RSP

    (e) XOR the data addressed by EBX with DX and save the result in memory

    (f) XOR the data stored 30 words after the location addressed by BP with DI and save the result in DI

    (g) XOR DI with memory location WELL and save the result in DI

44. Develop a sequence of instructions that sets (1) the rightmost 4 bits of AX; clears (0) the leftmost 3 bits of AX; and inverts bits 7, 8, and 9 of AX.

45. Describe the difference between the AND and TEST instructions.

46. Select an instruction that tests bit position 2 of register CH.

47. What is the difference between the NOT and the NEG instruction?

48. Select the correct instruction to perform each of the following tasks:

    (a) shift DI right three places, with zeros moved into the leftmost bit

    (b) move all bits in AL left one place, making sure that a 0 moves into the rightmost bit position

    (c) rotate all the bits of AL left three places

    (d) rotate carry right one place through EDX

    (e) move the DH register right one place, making sure that the sign of the result is the same as the sign of the original number

49. What does the SCASB instruction accomplish?

50. For string instructions, DI always addresses data in the _____ segment.

51. What is the purpose of the D flag bit?

52. Explain what the REPE prefix does when coupled with the SCASB instruction.

53. What condition or conditions will terminate the repeated string instruction REPNE SCASB?

54. Describe what the CMPSB instruction accomplishes.

55. Develop a sequence of instructions that scans through a 300H-byte section of memory called LIST, located in the data segment, searching for a 66H.

56. What happens if AH = 02H and DL = 43H when the INT 21H instruction is executed?

CHAPTER 6

Program Control Instructions

INTRODUCTION

The program control instructions direct the flow of a program and allow the flow to change. A change in flow often occurs after a decision made with the CMP or TEST instruction is followed by a conditional jump instruction. This chapter explains the program control instructions, including the jumps, calls, returns, interrupts, and machine control instructions.

This chapter also presents the relational assembly language statements (.IF, .ELSE, .ELSEIF, .ENDIF, .WHILE, .ENDW, .REPEAT, and .UNTIL) that are available in version 6.xx and above of MASM or TASM, with version 5.xx set for MASM compatibility. These relational assembly language commands allow the programmer to develop control flow portions of the program with C/C++ language efficiency.

CHAPTER OBJECTIVES

Upon completion of this chapter, you will be able to:

1. Use both conditional and unconditional jump instructions to control the flow of a program.
2. Use the relational assembly language statements .IF, .REPEAT, .WHILE, and so forth in programs.
3. Use the call and return instructions to include procedures in the program structure.
4. Explain the operation of the interrupts and interrupt control instructions.
5. Use machine control instructions to modify the flag bits.
6. Use ENTER and LEAVE to enter and leave programming structures.

6–1   THE JUMP GROUP

The main program control instruction, **jump** (JMP), allows the programmer to skip sections of a program and branch to any part of the memory for the next instruction. A conditional jump instruction allows the programmer to make decisions based upon numerical tests. The results of numerical tests are held in the flag bits, which are then tested by conditional jump instructions. Another instruction similar to the conditional jump, the conditional set, is explained with the conditional jump instructions in this section.

In this section of the text, all jump instructions are illustrated with their uses in sample programs. Also revisited are the LOOP and conditional LOOP instructions, first presented in Chapter 3, because they are also forms of the jump instruction.

Unconditional Jump (JMP)

Three types of unconditional jump instructions (see Figure 6–1) are available to the microprocessor: short jump, near jump, and far jump. The **short jump** is a 2-byte instruction that allows jumps or branches to memory locations within +127 and –128 bytes from the address following the jump. The 3-byte **near jump** allows a branch or jump within ±32K bytes (or anywhere in the current code segment) from the instruction in the current code segment. Remember that segments are cyclic in nature, which means that one location above offset address FFFFH is offset address 0000H. For this reason, if you jump 2 bytes ahead in memory and the instruction pointer addresses offset address FFFFH, the flow continues at offset address 0001H. Thus, a displacement of ±32K bytes allows a jump to any location within the current code segment. Finally, the 5-byte **far jump** allows a jump to any memory location within the real memory system. The short and near jumps are often called **intrasegment jumps**, and the far jumps are often called **intersegment jumps**.

In the 80386 through the Core2 processors, the near jump is within ±2G if the machine is operated in the protected mode, with a code segment that is 4G bytes long. If operated in the real mode, the near jump is within ±32K bytes. In the protected mode, the 80386 and above use a 32-bit displacement that is not shown in Figure 6–1. If the Pentium 4 is operated in the 64-bit mode, a jump can be to any address in its 1T memory space.

Short Jump. Short jumps are called **relative jumps** because they can be moved, along with their related software, to any location in the current code segment without a change. This is because the jump address is not stored with the opcode. Instead of a jump address, a **distance**, or displacement, follows the opcode. The short jump displacement is a distance represented by a 1-byte signed number whose value ranges between +127 and −128. The short jump instruction appears in Figure 6–2. When the microprocessor executes a short jump, the displacement is sign-extended and added to the instruction pointer (IP/EIP) to generate the jump address within the current code segment. The short jump instruction branches to this new address for the next instruction in the program.

Example 6–1 shows how short jump instructions pass control from one part of the program to another. It also illustrates the use of a **label** (a symbolic name for a memory address) with the jump instruction. Notice how one jump (JMP SHORT NEXT) uses the SHORT directive to force

FIGURE 6–1 The three main forms of the JMP instruction. Note that Disp is either an 8- or 16-bit signed displacement or distance.

194   CHAPTER 6
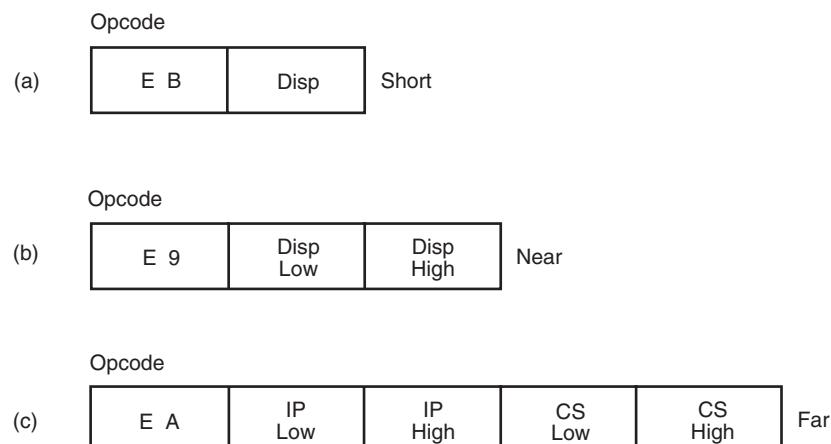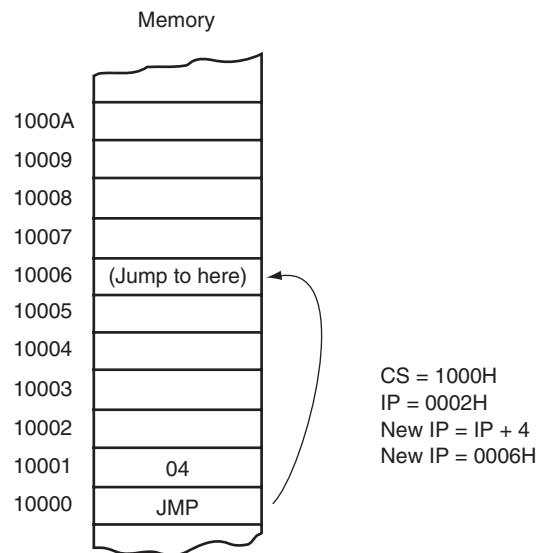
FIGURE 6–2 A short jump
to four memory locations
beyond the address of the
next instruction.

a short jump, while the other does not. Most assembler programs choose the best form of the
jump instruction so the second jump instruction (JMP START) also assembles as a short jump. If
the address of the next instruction (0009H) is added to the sign-extended displacement (0017H)
of the first jump, the address of NEXT is at location 0017H + 0009H or 0020H.

EXAMPLE 6–1

```
0000 33 DB                XOR  BX,BX
0002 B8 0001      START:   MOV  AX,1
0005 03 C3                 ADD  AX,BX
0007 EB 17                 JMP  SHORT NEXT


                   <skipped memory locations>


0020 8B D8        NEXT:    MOV  BX,AX
0022 EB DE                 JMP  START
```
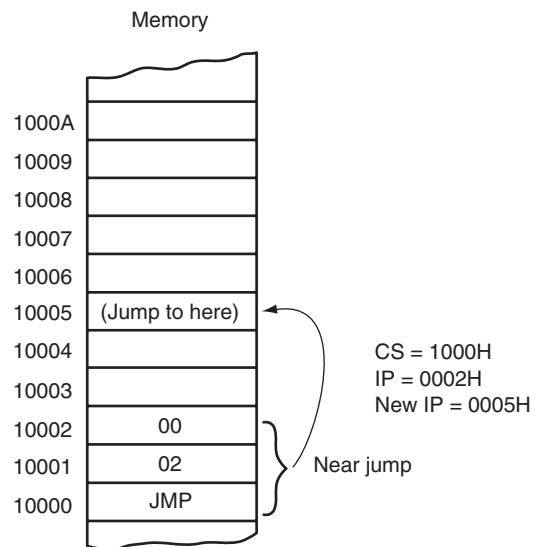
Whenever a jump instruction references an address, a label normally identifies the address.
The JMP NEXT instruction is an example; it jumps to label NEXT for the next instruction. It is
very rare to use an actual hexadecimal address with any jump instruction, but the assembler sup-
ports addressing in relation to the instruction pointer by using the $+a displacement. For exam-
ple, the JMP $+2 instruction jumps over the next two memory locations (bytes) following the
JMP instruction. The label NEXT must be followed by a colon (NEXT:) to allow an instruction
to reference it for a jump. If a colon does not follow a label, you cannot jump to it. Note that the
only time a colon is used after a label is when the label is used with a jump or call instruction.
This is also true in Visual C++.

Near Jump.   The near jump is similar to the short jump, except that the distance is farther. A
near jump passes control to an instruction in the current code segment located within ±32K bytes
from the near jump instruction. The distance is ±2G in the 80386 and above when operated in
protected mode. The near jump is a 3-byte instruction that contains an opcode followed by a
signed 16-bit displacement. In the 80386 through the Pentium 4 processors, the displacement is
32 bits and the near jump is 5 bytes long. The signed displacement adds to the instruction pointer
(IP) to generate the jump address. Because the signed displacement is in the range of ±32K, a

FIGURE 6–3 A near jump that adds the displacement (0002H) to the contents of IP.

near jump can jump to any memory location within the current real mode code segment. The protected mode code segment in the 80386 and above can be 4G bytes long, so the 32-bit displacement allows a near jump to any location within ±2G bytes. Figure 6–3 illustrates the operation of the real mode near jump instruction.

The near jump is also relocatable (as was the short jump) because it is also a relative jump. If the code segment moves to a new location in the memory, the distance between the jump instruction and the operand address remains the same. This allows a code segment to be relocated by simply moving it. This feature, along with the relocatable data segments, makes the Intel family of microprocessors ideal for use in a general-purpose computer system. Software can be written and loaded anywhere in the memory and function without modification because of the relative jumps and relocatable data segments.

Example 6–2 shows the same basic program that appeared in Example 6–1, except that the jump distance is greater. The first jump (JMP NEXT) passes control to the instruction at offset memory location 0200H within the code segment. Notice that the instruction assembles as E9 0200 R. The letter R denotes a **relocatable jump address** of 0200H. The relocatable address of 0200H is for the assembler program's internal use only. The actual machine language instruction assembles as E9 F6 01, which does not appear in the assembler listing. The actual displacement is 01F6H for this jump instruction. The assembler lists the jump address as 0200 R, so the address is easier to interpret as software is developed. If the linked execution file (.EXE) or command file (.COM) is displayed in hexadecimal code, the jump instruction appears as E9 F6 01.

EXAMPLE 6–2
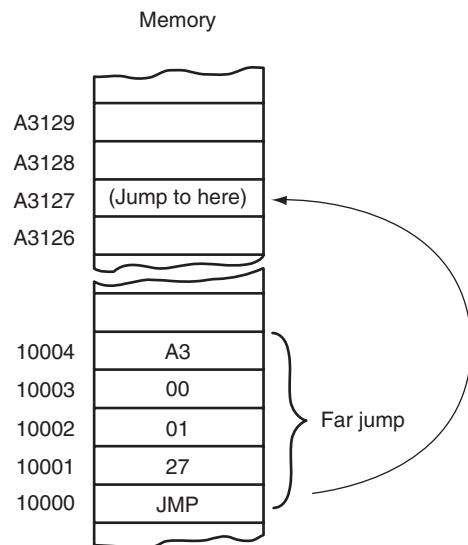
```
0000 33DB                 XOR  BX,BX
0002 B8 0001      START:  MOV  AX,1
0005 03 C3                ADD  AX,BX
0007 E9 0200 R            JMP  NEXT


                  <skipped memory locations>


0200 8B D8        NEXT:   MOV  BX,AX
0202 E9 0002 R            JMP  START
```

Far Jump. A far jump instruction (see Figure 6–4) obtains a new segment and offset address to accomplish the jump. Bytes 2 and 3 of this 5-byte instruction contain the new offset address; bytes

FIGURE 6–4   A far jump instruction replaces the contents of both CS and IP with 4 bytes following the opcode.

4 and 5 contain the new segment address. If the microprocessor (80286 through the Core2) is operated in the protected mode, the segment address accesses a descriptor that contains the base address of the far jump segment. The offset address, which is either 16 or 32 bits, contains the offset address within the new code segment.

Example 6–3 lists a short program that uses a far jump instruction. The far jump instruction sometimes appears with the FAR PTR directive, as illustrated. Another way to obtain a far jump is to define a label as a **far label.** A label is far only if it is external to the current code segment or procedure. The JMP UP instruction in the example references a far label. The label UP is defined as a far label by the EXTRN UP:FAR directive. **External labels** appear in programs that contain more than one program file. Another way of defining a label as global is to use a *double colon* (LABEL::) following the label in place of the single colon. This is required inside procedure blocks that are defined as near if the label is accessed from outside the procedure block.

When the program files are joined, the linker inserts the address for the UP label into the JMP UP instruction. It also inserts the segment address in the JMP START instruction. The segment address in JMP FAR PTR START is listed as - - - - R for relocatable; the segment address in JMP UP is listed as - - - - E for external. In both cases, the - - - - is filled in by the linker when it links or joins the program files.

EXAMPLE 6–3

```
                         EXTRN   UP:FAR

0000 33 DB               XOR BX,BX
0002 B8 0001     START:  ADD AX,1
0005 E9 0200 R           JMP NEXT

                 ;<skipped memory locations>

0200 8B D8       NEXT:   MOV BX,AX
0202 EA 0002 ── R        JMP FAR PTR START

0207 EA 0000 ── R        JMP UP
```

Jumps with Register Operands.   The jump instruction can also use a 16- or 32-bit register as an operand. This automatically sets up the instruction as an **indirect jump.** The address of the jump is in the register specified by the jump instruction. Unlike the displacement associated

with the near jump, the contents of the register are transferred directly into the instruction pointer. An indirect jump does not add to the instruction pointer, as with short and near jumps. The JMP AX instruction, for example, copies the contents of the AX register into the IP when the jump occurs. This allows a jump to any location within the current code segment. In the 80386 and above, a JMP EAX instruction also jumps to any location within the current code segment; the difference is that in protected mode the code segment can be 4G bytes long, so a 32-bit offset address is needed.

Example 6–4 shows how the JMP AX instruction accesses a jump table in the code segment. This DOS program reads a key from the keyboard and then modifies the ASCII code to 00H in AL for a '1', 01H for a '2', and 02H for a '3'. If a '1', '2', or '3' is typed, AH is cleared to 00H. Because the jump table contains 16-bit offset addresses, the contents of AX are doubled to 0, 2, or 4, so a 16-bit entry in the table can be accessed. Next, the offset address of the start of the jump table is loaded to SI, and AX is added to form the reference to the jump address. The MOV AX,[SI] instruction then fetches an address from the jump table, so the JMP AX instruction jumps to the addresses (ONE, TWO, or THREE) stored in the jump table.

EXAMPLE 6–4

```
              ;Instructions that read 1, 2, or 3 from the keyboard.
              ;The number is displayed as 1, 2, or 3 using a jump table
              ;
              .MODEL SMALL              ;select SMALL model
0000          .DATA                    ;start data segment
0000 0030 R   TABLE: DW ONE            ;jump table
0002 0034 R          DW TWO
0004 0038 R          DW THREE
0000          .CODE                    ;start code segment
              .STARTUP                 ;start program
0017 B4 01    TOP:   MOV AH,1          ;read key into AL
0019 CD 21           INT 21H
001B 2C 31           SUB AL,31         ;convert to BCD
001D 72 F9           JB TOP            ;if key < 1
001F 32 02           CMP AL,2
0021 77 F4           JA TOP            ;if key > 3
0023 B4 00           MOV AH,0          ;double key code
0025 03 C0           ADD AX,AX
0027 BE 0000 R       MOV SI,OFFSET TABLE ;address TABLE
002A 03 F0           ADD SI,AX         ;form lookup address
002C 8B 04           MOV AX,[SI]       ;get ONE, TWO or THREE
002E FF E0           JMP AX            ;jump to ONE, TWO or THREE
0030 B2 31    ONE:   MOV DL,'1'        ;get ASCII 1
0032 EB 06           JMP BOT
0034 B2 32    TWO:   MOV DL,'2'        ;get ASCII 2
0036 EB 02           JMP BOT
0038 B2 33    THREE: MOV DL,'3'        ;get ASCII 3
003A B4 02    BOT:   MOV AH,2          ;display number
003C CD 21           INT 21H
              .EXIT
              END
```

Indirect Jumps Using an Index. The jump instruction may also use the [ ] form of addressing to directly access the jump table. The jump table can contain offset addresses for near indirect jumps, or segment and offset addresses for far indirect jumps. (This type of jump is also known as a *double-indirect jump* if the register jump is called an *indirect jump*.) The assembler assumes that the jump is near unless the FAR PTR directive indicates a far jump instruction. Here Example 6–5 repeats Example 6–4 by using the JMP TABLE [SI] instead of JMP AX. This reduces the length of the program.

EXAMPLE 6–5

```
                    ;Instructions that read 1, 2, or 3 from the keyboard.
                    ;The number is displayed as 1, 2, or 3 using a jump table
                    ;
                    .MODEL  SMALL              ;select SMALL model
0000                .DATA                      ;start data segment
0000 002D R         TABLE: DW    ONE           ;jump table
0002 0031 R                DW    TWO
0004 0035 R                DW    THREE
0000                .CODE                      ;start code segment
                    .STARTUP                   ;start program
0017 B4 01          TOP:    MOV   AH,1         ;read key into AL
0019 CD 21                  INT   21H
001B 2C 31                  SUB   AL,31        ;convert to BCD
001D 72 F9                  JB    TOP          ;if key < 1
001F 32 02                  CMP   AL,2
0021 77 F4                  JA    TOP          ;if key > 3
0023 B4 00                  MOV   AH,0         ;double key code
0025 03 C0                  ADD   AX,AX
0027 B5 F0                  MOV   SI,AX        ;form lookup address
0029 FF A4 0000 R           JMP   TABLE[SI]    ;jump to ONE, TWO or THREE
002D B2 31          ONE:    MOV   DL,'1'       ;get ASCII 1
002F EB 06                  JMP   BOT
0031 B2 32          TWO:    MOV   DL,'2'       ;get ASCII 2
0033 EB 02                  JMP   BOT
0035 B2 33          THREE:  MOV   DL,'3'       ;get ASCII 3
0037 B4 02          BOT:    MOV   AH,2         ;display number
0039 CD 21                  INT   21H
                    .EXIT
                    END
```

The mechanism used to access the jump table is identical with a normal memory reference. The JMP TABLE [SI] instruction points to a jump address stored at the code segment offset location addressed by SI. It jumps to the address stored in the memory at this location. Both the register and indirect indexed jump instructions usually address a 16-bit offset. This means that both types of jumps are near jumps. If a JMP FAR PTR [SI] or JMP TABLE [SI], with TABLE data defined with the DD directive appears in a program, the microprocessor assumes that the jump table contains doubleword, 32-bit addresses (IP and CS).

Conditional Jumps and Conditional Sets

Conditional jump instructions are always short jumps in the 8086 through the 80286 microprocessors. This limits the range of the jump to within +127 bytes and −128 bytes from the location following the conditional jump. In the 80386 and above, conditional jumps are either short or near jumps (±32K). In the 64-bit mode of the Pentium 4, the near jump distance is ±2G for the conditional jumps. This allows these microprocessors to use a conditional jump to any location within the current code segment. Table 6–1 lists all the conditional jump instructions with their test conditions. Note that the Microsoft MASM version 6.x assembler automatically adjusts conditional jumps if the distance is too great.

The conditional jump instructions test the following flag bits: sign (S), zero (Z), carry (C), parity (P), and overflow (0). If the condition under test is true, a branch to the label associated with the jump instruction occurs. If the condition is false, the next sequential step in the program executes. For example, a JC will jump if the carry bit is set.

The operation of most conditional jump instructions is straightforward because they often test just one flag bit, although some test more than one. Relative magnitude comparisons require more complicated conditional jump instructions that test more than one flag bit.

TABLE 6–1 Conditional jump instructions.

Assembly Language	Tested Condition	Operation
JA	Z = 0 and C = 0	Jump if above
JAE	C = 0	Jump if above or equal
JB	C = 1	Jump if below
JBE	Z = 1 or C = 1	Jump if below or equal
JC	C = 1	Jump if carry
JE or JZ	Z = 1	Jump if equal or jump if zero
JG	Z = 0 and S = 0	Jump if greater than
JGE	S = 0	Jump if greater than or equal
JL	S != O	Jump if less than
JLE	Z = 1 or S != O	Jump if less than or equal
JNC	C = 0	Jump if no carry
JNE or JNZ	Z = 0	Jump if not equal or jump if not zero
JNO	O = 0	Jump if no overflow
JNS	S = 0	Jump if no sign (positive)
JNP or JPO	P = 0	Jump if no parity or jump if parity odd
JO	O = 1	Jump if overflow
JP or JPE	P = 1	Jump if parity or jump if parity even
JS	S = 1	Jump if sign (negative)
JCXZ	CX = 0	Jump if CX is zero
JECXZ	ECX = 0	Jump if ECX equals zero
JRCXZ	RCX = 0	Jump if RCX equals zero (64-bit mode)

Because both signed and unsigned numbers are used in programming, and because the order of these numbers is different, there are two sets of conditional jump instructions for magnitude comparisons. Figure 6–5 shows the order of both signed and unsigned 8-bit

FIGURE 6–5 Signed and unsigned numbers follow different orders.

numbers. The 16- and 32-bit numbers follow the same order as the 8-bit numbers, except that they are larger. Notice that an FFH (255) is above the 00H in the set of unsigned numbers, but an FFH (−1) is less than 00H for signed numbers. Therefore, an unsigned FFH is above 00H, but a signed FFH is less than 00H.

When signed numbers are compared, use the JG, JL, JGE, JLE, JE, and JNE instructions. The terms *greater than* and *less than* refer to signed numbers. When unsigned numbers are compared, use the JA, JB, JAB, JBE, JE, and JNE instructions. The terms *above* and *below* refer to unsigned numbers.

The remaining conditional jumps test individual flag bits, such as overflow and parity. Notice that JE has an alternative opcode JZ. All instructions have alternates, but many aren't used in programming because they don't usually fit the condition under test. (The alternates appear in Appendix B with the instruction set listing.) For example, the JA instruction (jump if above) has the alternative JNBE (jump if not below or equal). A JA functions exactly as a JNBE, but a JNBE is awkward in many cases when compared to a JA.

The conditional jump instructions all test flag bits except for JCXZ (jump if CX = 0) and JECXZ (jump if ECX = 0). Instead of testing flag bits, JCXZ directly tests the contents of the CX register without affecting the flag bits, and JECXZ tests the contents of the ECX register. For the JCXZ instruction, if CX = 0, a jump occurs, and if CX != 0, no jump occurs. Likewise for the JECXZ instruction, if ECX = 0, a jump occurs; if ECX != 0, no jump occurs. In the Pentium 4 or Core2 operated in the 64-bit mode, the JRCXZ instruction jumps is RCX = 0.

A program that uses JCXZ appears in Example 6–6. Here, the SCASB instruction searches a table for 0AH. Following the search, a JCXZ instruction tests CX to see if the count has reached zero. If the count is zero, the 0AH is not found in the table. The carry flag is used in this example to pass the not found condition back to the calling program. Another method used to test to see if the data are found is the JNE instruction. If JNE replaces JCXZ, it performs the same function. After the SCASB instruction executes, the flags indicate a not-equal condition if the data were not found in the table.

EXAMPLE 6–6

```
                    ;Instructions that search a table of 100H bytes for 0AH
                    ;The offset address of TABLE is assumed to be in SI
                    ;
0017 B9 0064            MOV CX,100          ;load counter
001A B0 0A              MOV AL,0AH          ;load AL with 0AH
001C FC                 CLD                 ;auto-increment
001D F2/AE              REPNE SCASB         ;search for 0AH
001F F9                 STC                 ;set carry if found
0020 E3 01              JCXZ NOT_FOUND      ;if not found
0022              NOT_FOUND
```

The Conditional Set Instructions. In addition to the conditional jump instructions, the 80386 through the Core2 processors also contain conditional set instructions. The conditions tested by conditional jumps are put to work with the conditional set instructions. The conditional set instructions set a byte to either 01H or clear a byte to 00H, depending on the outcome of the condition under test. Table 6–2 lists the available forms of the conditional set instructions.

These instructions are useful where a condition must be tested at a point much later in the program. For example, a byte can be set to indicate that the carry is cleared at some point in the program by using the SETNC MEM instruction. This instruction places 01H into memory location MEM if carry is cleared, and 00H into MEM if carry is set. The contents of MEM can be

TABLE 6–2 Conditional set instructions.

Assembly Language	Tested Condition	Operation
SETA	Z = 0 and C = 0	Set if above
SETAE	C = 0	Set if above or equal
SETB	C = 1	Set if below
SETBE	Z = 1 or C = 1	Set if below or equal
SETC	C = 1	Set if carry
SETE or SETZ	Z = 1	Set if equal or set if zero
SETG	Z = 0 and S = 0	Set if greater than
SETGE	S = 0	Set if greater than or equal
SETL	S != O	Set if less than
SETLE	Z = 1 or S != O	Set if less than or equal
SETNC	C = 0	Set if no carry
SETNE or SETNZ	Z = 0	Set if not equal or set if not zero
SETNO	O = 0	Set if no overflow
SETNS	S = 0	Set if no sign (positive)
SETNP or SETPO	P = 0	Set if no parity or set if parity odd
SETO	O = 1	Set if overflow
SETP or SETPE	P = 1	Set if parity or set if parity even
SETS	S = 1	Set if sign (negative)

tested at a later point in the program to determine if carry is cleared at the point where the SETNC MEM instruction executed.

LOOP

The LOOP instruction is a combination of a decrement CX and the JNZ conditional jump. In the 8086 through the 80286 processors, LOOP decrements CX; if CX != 0, it jumps to the address indicated by the label. If CX becomes 0, the next sequential instruction executes. In the 80386 and above, LOOP decrements either CX or ECX, depending upon the instruction mode. If the 80386 and above operate in the l6-bit instruction mode, LOOP uses CX; if operated in the 32-bit instruction mode, LOOP uses ECX. This default is changed by the LOOPW (using CX) and LOOPD (using ECX) instructions in the 80386 through the Core2. In the 64-bit mode, the loop counter is in RCX and is 64 bits wide.

　　Example 6–7 shows how data in one block of memory (BLOCK1) add to data in a second block of memory (BLOCK2), using LOOP to control how many numbers add. The LODSW and STOSW instructions access the data in BLOCK1 and BLOCK2. The ADD AX, ES:[DI] instruction accesses the data in BLOCK2 located in the extra segment. The only reason that BLOCK2 is in the extra segment is that DI addresses extra segment data for the STOSW instruction. The .STARTUP directive only loads DS with the address of the data segment. In this example, the extra segment also addresses data in the data segment, so the contents of DS are copied to ES through the accumulator. Unfortunately, there is no direct move from segment register to segment register instruction.

EXAMPLE 6–7

```
                      ;A program that sums the contents of BLOCK1 and BLOCK2
                      ;and stores the results on top of the data in BLOCK2.
                      ;
                      .MODEL SMALL                ;select SMALL model
0000                  .DATA                       ;start data segment
0000 0064[            BLOCK1 DW 100 DUP(?)        ;100 words for BLOCK1
       0000
         ]
00C8 0064[            BLOCK2 DW 100 DUP(?)        ;100 words for BLOCK2
       0000
         ]
0000                  .CODE                       ;start code segment
                      .STARTUP                    ;start program
0017 8C D8                  MOV AX,DS             ;overlap DS and ES
0019 8E C0                  MOV ES,AX
001B FC                     CLD                   ;select auto-increment
001C B9 0064                MOV CX,100            ;load counter
001F BE 0000 R              MOV SI,OFFSET BLOCK1  ;address BLOCK1
0022 BF 00C8 R              MOV DI,OFFSET BLOCK2  ;address BLOCK2
0025 AD        L1:          LODSW                 ;load AX with BLOCK1
0026 26:03 05               ADD AX,ES:[DI]        ;add BLOCK2
0029 AB                     STOSW                 ;save answer
002A E2 F9                  LOOP L1               ;repeat 100 times
                      .EXIT
                      END
```

Conditional LOOPs.   As with REP, the LOOP instruction also has conditional forms: LOOPE and
LOOPNE. The LOOPE (**loop while equal**) instruction jumps if CX != 0 while an equal condition
exists. It will exit the loop if the condition is not equal or if the CX register decrements to 0. The
LOOPNE (**loop while not equal**) instruction jumps if CX != 0 while a not-equal condition exists.
It will exit the loop if the condition is equal or if the CX register decrements to 0. In the 80386
through the Core2 processors, the conditional LOOP instruction can use either CX or ECX as the
counter. The LOOPEW/LOOPED or LOOPNEW/LOOPNED instructions override the instruction
mode if needed. Under 64-bit operation, the loop counter uses RCX and is 64 bits in width.

As with the conditional repeat instructions, alternates exist for LOOPE and LOOPNE. The
LOOPE instruction is the same as LOOPZ, and the LOOPNE instruction is the same as
LOOPNZ. In most programs, only the LOOPE and LOOPNE apply.

6–2   CONTROLLING THE FLOW OF THE PROGRAM

It is much easier to use the assembly language statements .IF, .ELSE, .ELSEIF, and .ENDIF to
control the flow of the program than it is to use the correct conditional jump statement. These
statements always indicate a special assembly language command to MASM. Note that the con-
trol flow assembly language statements beginning with a period are only available to MASM
version 6.xx, and not to earlier versions of the assembler such as 5.10. Other statements devel-
oped in this chapter include the .REPEAT–.UNTIL and .WHILE–.ENDW statements. These
statements (**the dot commands**) do not function when using the Visual C++ inline assembler.

Example 6–8(a) shows how these statements are used to control the flow of a program by
testing AL for the ASCII letters A through F. If the contents of AL are A through F, 7 is sub-
tracted from AL.

Accomplishing the same task using the Visual C++ inline assembler is usually handled in
C++ rather than in assembly language. Example 6–8(b) shows the same task using the inline
assembler in Visual C++ and conditional jumps in assembly language. It also shows how to use

a label in an assembly block in Visual C++. This illustrates that it is more difficult to accomplish the same task without the dot commands. Never use uppercase for assembly language commands with the inline assembler because some of them are reserved by C++ and will cause problems.

EXAMPLE 6–8 (a)

```
.IF AL >= 'A' && AL <= 'F'
        SUB AL,7
.ENDIF
SUB AL,30H
```

EXAMPLE 6–8 (b)

```
char temp;
_asm{
        mov     al,temp
        cmp     al,41h
        jb      Later
        cmp     al,46h
        ja      Later
        sub     al,7
Later:
        sub     al,30h
        mov     temp,al

}
```

In Example 6–8(a) notice how the && symbol represents the AND function in the .IF statement. There is no .if in Example 6–8(b) because the same operation was performed by using a few compare (CMP) instructions to accomplish the same task. See Table 6–3 for a complete list of relational operators used with the .IF statement. Note that many of these conditions (such as &&) are also used by many high-level languages such as C/C++.

Example 6–9 shows another example of the conditional .IF directive that converts all ASCII-coded letters to uppercase. First, the keyboard is read without echo using DOS INT 21H function 06H, and then the .IF statement converts the character into uppercase, if needed. In this example, the logical AND function (&&) is used to determine if the character is in lowercase. If it is lowercase, 20H is subtracted, converting to uppercase. This program reads a key from the

TABLE 6–3 Relational operators used with the .IF statement in assembly language.

Operator	Function
==	Equal or the same as
!=	Not equal
>	Greater than
>=	Greater than or equal
<	Less than
<=	Less than or equal
&	Bit test
!	Logical inversion
&&	Logical AND
||	Logical OR
|	Or

keyboard and converts it to uppercase before displaying it. Notice also how the program terminates when the control C key (ASCII = 03H) is typed. The .LISTALL directive causes all assembler-generated statements to be listed, including the label @Startup generated by the .STARTUP directive. The .EXIT directive also is expanded by .LISTALL to show the use of the DOS INT 21H function 4CH, which returns control to DOS.

EXAMPLE 6–9

```
                ;A DOS program that reads the keyboard and converts all
                ;lowercase data to uppercase before displaying it.
                ;
                ;This program is terminated with a control-C
                ;
                .MODEL TINY                 ;select tiny model
                .LISTALL                    ;list all statements
0000            .CODE                       ;start code segment
                .STARTUP                    ;start program
0100        * @Startup
0100 B4 06     MAIN1: MOV  AH,6             ;read key without echo
0102 B2 FF            MOV  DL,0FFH
0104 CD 21            INT  21H
0106 74 F8            JE   MAIN1            ;if no key
0108 3C 03            CMP  AL,3             ;test for control-C
010A 74 10            JE   MAIN2            ;if control-C

                .IF AL >= 'a' && AL <= 'z'

010C 3C 61     *              cmp al,'a'
010E 72 06     *              jb  @C0001
0110 3C 7A     *              cmp al,'z'
0112 77 02     *              ja  @C0001

0114 2C 20            SUB AL,20H

                .ENDIF

0116        * @C0001:
0116 8A D0            MOV  DL,AL             ;echo character to display
0118 CD 21            INT  21H
011A EB E4            JMP  MAIN1             ;repeat
011C          MAIN2:
                .EXIT
011C B4 4C     *      MOV AH,4CH
011E CD 21     *      INT 21H
                END
```

In this program, a lowercase letter is converted to uppercase by the use of the .IF AL >= 'a' && AL <= 'z' statement. If AL contains a value that is greater than or equal to a lowercase a, and less than or equal to a lowercase z (a value of a through z), the statement between the .IF and .ENDIF executes. This statement (SUB AL,20H) subtracts 20H from the lowercase letter to change it to an uppercase letter. Notice how the assembler program implements the .IF statement (see lines that begin with *). The label @C0001 is an assembler-generated label used by the conditional jump statements placed in the program by the .IF statement.

Another example that uses the conditional .IF statement appears in Example 6–10. This program reads a key from the keyboard, and then converts it to hexadecimal code. This program is not listed in expanded form.

In this example, the .IF AL >='a' && AL<='f' statement causes the next instruction (SUB AL,57H) to execute if AL contains letters a through f, converting them to hexadecimal. If it is not between letters a and f, the next .ELSEIF statement tests it for the letters A through F. If it is the letters A through F, 37H is subtracted from AL. If neither condition is true, 30H

is subtracted from AL before AL is stored at data segment memory location TEMP. The same conversion can be performed in a C++ function as illustrated in the program snippet of Example 6–10(b).

EXAMPLE 6–10(a)

```
                 ;A DOS program that reads key and stores its hexadecimal
                 ;value in memory location TEMP
                 ;
                 .MODEL SMALL                  ;select small model
0000             .DATA                         ;start data segment
0000 00          TEMP    DB    ?               ;define TEMP
0000             .CODE                         ;start code segment
                 .STARTUP                      ;start program
0017 B4 01               MOV  AH,1             ;read keyboard
0019 CD 21               INT  21H
                 .IF AL >= 'a' && AL <= 'f'
0023 2C 57                   SUB  AL,57H                    ;if lowercase
                 .ELSEIF .IF AL >= 'A' && AL <= 'F'
002F 2C 37                   SUB  AL,37H                    ;if uppercase
                 .ELSE
0033 2C 30                   SUB  AL,30H                    ; otherwise
                 .ENDIF
0035 A2 0000 R           MOV  TEMP,AL          ;save it in TEMP
                 .EXIT
                 END
```

EXAMPLE 6–10(b)

```
char Convert(char temp)
{
      if ( temp >= 'a' && temp <= 'f' )
            temp -= 0x57;
      else if ( temp >= 'A' && temp <= 'F' )
            temp -= 0x37;
      else
            temp -= 0x30;
      return temp;
}
```

WHILE Loops

As with most high-level languages, the assembler also provides the WHILE loop construct, available to MASM version 6.x. The .WHILE statement is used with a condition to begin the loop, and the .ENDW statement ends the loop.

Example 6–11 shows how the .WHILE statement is used to read data from the keyboard and store it into an array called BOP until the enter key (0DH) is typed. This program assumes that BUF is stored in the extra segment because the STOSB instruction is used to store the keyboard data in memory. Note that the .WHILE loop portion of the program is shown in expanded form so that the statements inserted by the assembler (beginning with a *) can be studied. After the Enter key (0DH) is typed, the string is appended with a $ so it can be displayed with DOS INT 21H function number 9.

EXAMPLE 6–11

```
                 ;A DOS program that reads a character string from the
                 ;keyboard and then displays it again.
                 ;
                 .MODEL SMALL                  ;select small model
0000             .DATA                         ;start data segment
0000 0D 0A       MES    DB  13,10              ;return and line feed
```

```
0002 0100[           BUF    DB 256 DUP(?)           ;character string buffer
        00
          ]
0000                 .CODE                          ;start code segment
                     .STARTUP                       ;start program
0017 8C D8                 MOV  AX,DX                ;overlap DS with ES
0019 8C C0                 MOV  ES,AX
001B FC                    CLD                       ;select auto-increment
001C BF 0002 R             MOV  DI,OFFSET BUF        ;address buffer

                     .WHILE AL != 0DH               ;loop while not enter
001F EB 05      *          jmp  @C0001
0021            *    @C0002:
0021 B4 01                 MOV  AH,1                 ;read key
0023 CD 21                 INT  21H
0025 AA                    STOSB                     ;store key code
                     .ENDW
0026            *    @C0001:
0026 3C 0D      *          cmp  al,0dh
0028 75 F7      *          jne  @C0002
002A C6 45 FF 24           MOV  BYTE PTR[DI-1]'&'
002E BA 0000 R             MOV  DX,OFFSET MES
0031 B4 09                 MOV  AH,9
0033 CD 21                 INT  21H                  ;display MES
                 .EXIT
                 END
```

The program in Example 6–11 functions perfectly, as long as we arrive at the .WHILE statement with AL containing some other value except 0DH. This can be corrected by adding a MOV AL,0DH instruction before the .WHILE statement in Example 6–11. Although not shown in an example, the .BREAK and .CONTINUE statements are available for use with the while loop. The .BREAK statement is often followed by the .IF statement to select the break condition as in .BREAK .IF AL == 0DH. The .CONTINUE statement, which can be used to allow the DO–.WHILE loop to continue if a certain condition is met, can be used with .BREAK. For example, .CONTINUE .IF AL == 15 allows the loop to continue if AL equals 15. Note that the .BREAK and .CONTINUE commands function in the same manner in a C++ program.

REPEAT-UNTIL Loops

Also available to the assembler is the REPEAT–UNTIL construct. A series of instructions is repeated until some condition occurs. The .REPEAT statement defines the start of the loop; the end is defined with the .UNTIL statement, which contains a condition. Note that .REPEAT and .UNTIL are available to version 6.x of MASM.

If Example 6–11 is again reworked by using the REPEAT-UNTIL construct, this appears to be the best solution. See Example 6–12 for the program that reads keys from the keyboard and stores keyboard data into extra segment array BUF until the enter key is pressed. This program also fills the buffer with keyboard data until the Enter key (0DH) is typed. Once the Enter key is typed, the program displays the character string using DOS INT 2lH function number 9, after appending the buffer data with the required dollar sign. Notice how the .UNTIL AL == 0DH statement generates code (statements beginning with *) to test for the Enter key.

EXAMPLE 6–12

```
;A DOS program that reads a character string from the
;keyboard and then displays it again.
;
```

```
                        .MODEL SMALL                ;select small model
0000                    .DATA                       ;start data segment
0000 0D 0A      MES   DB  13,10                      ;return and line feed
0002 0100[      BUF   DB  256 DUP(?)                 ;character string buffer
        00
          ]
0000                    .CODE                       ;start code segment
                        .STARTUP                    ;start program
0017 8C D8              MOV  AX,DX                   ;overlap DS with ES
0019 8C C0              MOV  ES,AX
001B FC                 CLD                          ;select auto-increment
001C BF 0002 R          MOV  DI,OFFSET BUF           ;address buffer

                        .REPEAT                      ;repeat until enter

001F            *  @C0001:
001F B4 01              MOV  AH,1                    ;read key
0021 CD 21              INT  21H
0023 AA                 STOSB                        ;store key code

                        .UNTIL AL == 0DH

0025 3C 0D      *       cmp  al,0dh
0027 75 F7      *       jne  @C0001
0028 C6 45 FF 24        MOV  BYTE PTR[DI-1]'&'
002C BA 0000 R          MOV  DX,OFFSET MES
002E B4 09              MOV  AH,9
0031 CD 21              INT  21H                     ;display MES
                .EXIT
                END
```

There is also an .UNTILCXZ instruction available that uses the LOOP instruction to check CX for a repeat loop. The .UNTILCXZ instruction uses the CX register as a counter to repeat a loop a fixed number of times. Example 6–13 shows a sequence of instructions that uses the .UNTILCXZ instruction used to add the contents of byte-sized array ONE to byte-sized array TWO. The sums are stored in array THREE. Note that each array contains 100 bytes of data, so the loop is repeated 100 times. This example assumes that array THREE is in the extra segment, and that arrays ONE and TWO are in the data segment. Notice how the LOOP instruction is inserted for the .UNTILCXZ.

EXAMPLE 6–13

```
012C B9 0064            MOV  CX,100                  ;set count
012F BF 00C8 R          MOV  DI,OFFSET THREE         ;address arrays
0132 BE 0000 R          MOV  SI,OFFSET ONE
0135 BB 0064 R          MOV  BX,OFFSET TWO

                        .REPEAT

0138            *  @C0001:
0138 AC                 LODSB
0139 02 07              ADD AL,[BX]
013B AA                 STOSB
013C 43                 INC BX

                        .UNTILCXZ

013D E2 F9      *       LOOP @C0001
```

6–3   PROCEDURES

The procedure (subroutine, method, or **function**) is an important part of any computer system's architecture. A procedure is a group of instructions that usually performs one task. A procedure is a reusable section of the software that is stored in memory once, but used as often as necessary. This saves memory space and makes it easier to develop software. The only disadvantage of a procedure is that it takes the computer a small amount of time to link to the procedure and return from it. The CALL instruction links to the procedure, and the RET (**return**) instruction returns from the procedure.

The stack stores the return address whenever a procedure is called during the execution of a program. The CALL instruction pushes the address of the instruction following the CALL (**return address**) on the stack. The RET instruction removes an address from the stack so the program returns to the instruction following the CALL.

With the assembler, there are specific rules for storing procedures. A procedure begins with the PROC directive and ends with the ENDP directive. Each directive appears with the name of the procedure. This programming structure makes it easy to locate the procedure in a program listing. The PROC directive is followed by the type of procedure: NEAR or FAR. Example 6–16 shows how the assembler uses the definition of both a near (intrasegment) and far (intersegment) procedure. In MASM version 6.x, the NEAR or FAR type can be followed by the USES statement. The USES statement allows any number of registers to be automatically pushed to the stack and popped from the stack within the procedure. The USES statement is also illustrated in Example 6–14.
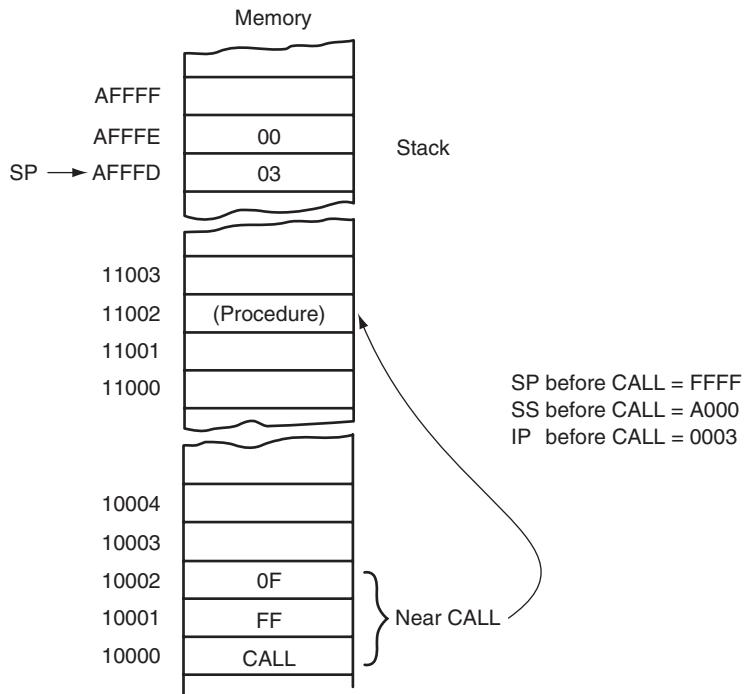
EXAMPLE 6–14

```
0000               SUMS   PROC   NEAR
0000 03 C3                ADD    AX,BX
0002 03 C1                ADD    AX,CX
0004 03 C2                ADD    AX,DX
0006 C3                   RET
0007               SUMS   ENDP

0007               SUMS1  PROC   FAR
0007 03 C3                ADD    AX,BX
0009 03 C1                ADD    AX,CX
000B 03 C2                ADD    AX,DX
000D CB                   RET
000E               SUMS1  ENDP

000E               SUMS3  PROC   NEAR USE BX CX DX
0011 03 C3                ADD    AX,BX
0013 03 C1                ADD    AX,CX
0015 03 C2                ADD    AX,DX
                          RET
001B               SUMS   ENDP
```

When these first two procedures are compared, the only difference is the opcode of the return instruction. The near return instruction uses opcode C3H and the far return uses opcode CBH. A near return removes a 16-bit number from the stack and places it into the instruction pointer to return from the procedure in the current code segment. A far return removes a 32-bit number from the stack and places it into both IP and CS to return from the procedure to any memory location.

Procedures that are to be used by all software (global) should be written as far procedures. Procedures that are used by a given task (local) are normally defined as near procedures. Most procedures are near procedures.

FIGURE 6–6 The effect of a near CALL on the stack and the instruction pointer.

CALL

The CALL instruction transfers the flow of the program to the procedure. The CALL instruction differs from the jump instruction because a CALL saves a return address on the stack. The return address returns control to the instruction that immediately follows the CALL in a program when a RET instruction executes.

Near CALL. The near CALL instruction is 3 bytes long; the first byte contains the opcode, and the second and third bytes contain the displacement, or distance of ±32K in the 8086 through the 80286 processors. This is identical to the form of the near jump instruction. The 80386 and above use a 32-bit displacement, when operating in the protected mode, that allows a distance of ±2G bytes. When the near CALL executes, it first pushes the offset address of the next instruction onto the stack. The offset address of the next instruction appears in the instruction pointer (IP or EIP). After saving this return address, it then adds the displacement from bytes 2 and 3 to the IP to transfer control to the procedure. There is no short CALL instruction. A variation on the opcode exists as CALLN, but this should be avoided in favor of using the PROC statement to define the CALL as near.

Why save the IP or EIP on the stack? The instruction pointer always points to the next instruction in the program. For the CALL instruction, the contents of IP/EIP are pushed onto the stack, so program control passes to the instruction following the CALL after a procedure ends. Figure 6–6 shows the return address (IP) stored on the stack and the call to the procedure.

Far CALL. The far CALL instruction is like a far jump because it can call a procedure stored in any memory location in the system. The far CALL is a 5-byte instruction that contains an opcode followed by the next value for the IP and CS registers. Bytes 2 and 3 contain the new contents of the IP, and bytes 4 and 5 contain the new contents for CS.

The far CALL instruction places the contents of both IP and CS on the stack before jumping to the address indicated by bytes 2 through 5 of the instruction. This allows the far CALL to call a procedure located anywhere in the memory and return from that procedure.

FIGURE 6–7   The effect of a
far CALL instruction.

Memory

```
        AFFFF  ┌──────────┐
        AFFFE  │    10    │
        AFFFD  │    00    │   Stack
        AFFFC  │    00    │
SP ──▶  AFFFB  │    05    │
               └──────────┘

        11003  ┌──────────┐
        11002  │(Procedure)│◀──┐
        11001  │          │    │
        11000  │          │    │
               └──────────┘    │
                               │   SP before CALL = FFFF
                               │   SS before CALL = A000
        10004  ┌──────────┐┐   │   IP  before CALL = 0005
               │    11    ││   │
        10003  │    00    ││   │
        10002  │    00    │├ Far CALL
        10001  │    02    ││   │
        10000  │   CALL   ││───┘
               └──────────┘┘
```
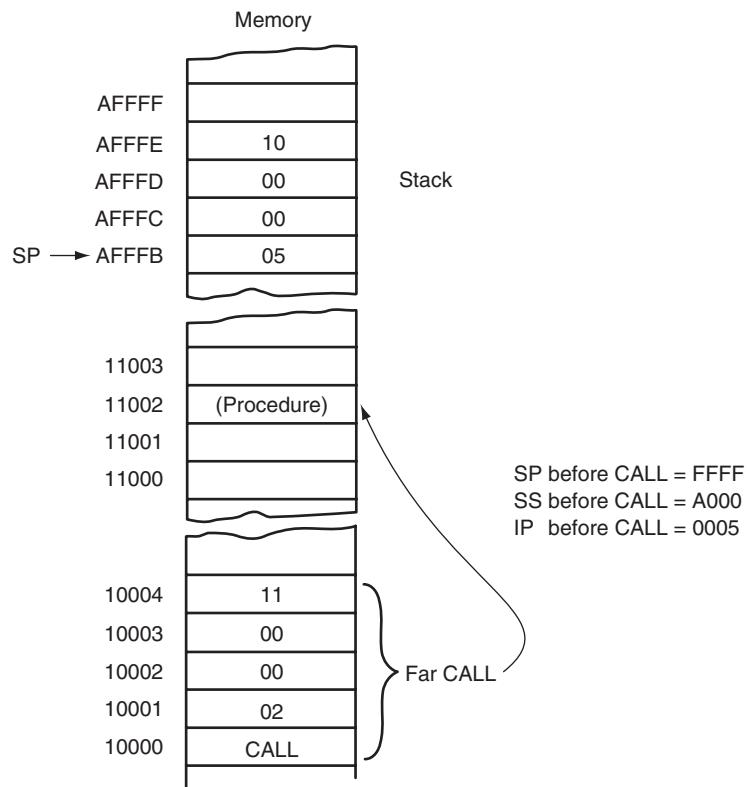
        Figure 6–7 shows how the far CALL instruction calls a far procedure. Here, the contents of
IP and CS are pushed onto the stack. Next, the program branches to the procedure. A variant of
the far call exists as CALLF, but this should be avoided in favor of defining the type of call
instruction with the PROC statement.
        In the 64-bit mode a far call is to any memory location and the information placed onto the
stack is an 8-byte number. Likewise, the far return instruction also retrieves an 8-byte return
address from the stack and places it into RIP.

CALLs with Register Operands.   Like jump instructions, call instructions also may contain a
register operand. An example is the CALL BX instruction, which pushes the contents of IP onto
the stack. It then jumps to the offset address, located in register BX, in the current code segment.
This type of CALL always uses a 16-bit offset address, stored in any 16-bit register except the
segment registers.
        Example 6–15 illustrates the use of the CALL register instruction to call a procedure that
begins at offset address DISP. (This call could also directly call the procedure by using the CALL
DISP instruction.) The OFFSET address DISP is placed into the BX register, and then the CALL
BX instruction calls the procedure beginning at address DISP. This program displays an "OK"
on the monitor screen.

EXAMPLE 6–15

```
                ;A DOS program that displays OK using the DISP procedure.
                ;
                .MODEL TINY                 ;select tiny model
0000            .CODE                       ;start code segment
                .STARTUP                    ;start program
```

```
0100 BB 0110 R              MOV  BX,OFFSET DISP   ;load BX with offset DISP
0103 B2 4F                  MOV  DL,'O'           ;display O
0105 FF D3                  CALL BX
0107 B2 4B                  MOV  DL,'K'           ;display K
0109 FF D3                  CALL BX
                 .EXIT
                 ;
                 ;Procedure that displays the ASCII character in DL
                 ;
0110             DISP  PROC  NEAR
0110 B4 02                  MOV  AH,2             ;select function 2
0112 CD 21                  INT  21H              ;execute DOS function 2
0114 C3                     RET
0115             DISP  ENDP
                 END
```

CALLs with Indirect Memory Addresses.   A CALL with an indirect memory address is particularly useful whenever different subroutines need to be chosen in a program. This selection process is often keyed with a number that addresses a CALL address in a lookup table. This is essentially the same as the indirect jump that used a lookup table for a jump address earlier in this chapter.

Example 6–16 shows how to access a table of addresses using an indirect CALL instruction. This table illustrated in the example contains three separate subroutine addresses referenced by the numbers 0, 1, and 2. This example uses the scaled-index addressing mode to multiply the number in EBX by 2 so it properly accesses the correct entry in the lookup table.

EXAMPLE 6–16

```
                 ;Instruction that calls procedure ZERO, ONE, or TWO
                 ;depending on the value in EBX
                 ;
        TABLE  DW   ZERO              ;address of procedure ZERO
               DW   ONE               ;address of procedure ONE
               DW   TWO               ;address of procedure TWO

               CALL  TABLE[2*EBX]
```

The CALL instruction also can reference far pointers if the instruction appears as CALL FAR PTR [4*EBX] or as CALL TABLE [4*EBX], if the data in the table are defined as double-word data with the DD directive. These instructions retrieve a 32-bit address (4 bytes long) from the data segment memory location addressed by EBX and use it as the address of a far procedure.
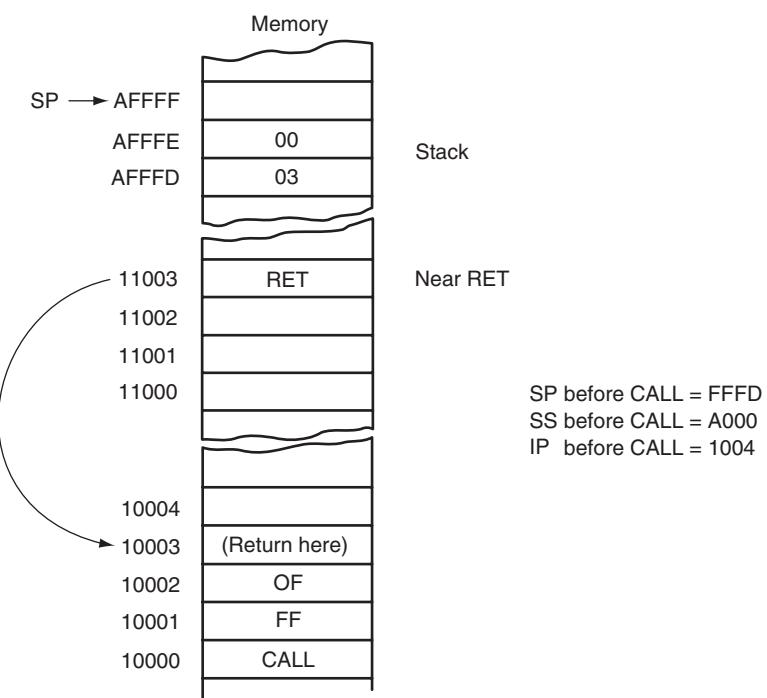
RET

The return instruction (RET) removes a 16-bit number (**near return**) from the stack and places it into IP, or removes a 32-bit number (**far return**) and places it into IP and CS. The near and far return instructions are both defined in the procedure's PROC directive, which automatically selects the proper return instruction. With the 80386 through the Pentium 4 processors operating in the protected mode, the far return removes 6 bytes from the stack. The first 4 bytes contain the new value for EIP and the last 2 contain the new value for CS. In the 80386 and above, a protected mode near return removes 4 bytes from the stack and places them into EIP.

When IP/EIP or IP/EIP and CS are changed, the address of the next instruction is at a new memory location. This new location is the address of the instruction that immediately follows the most recent CALL to a procedure. Figure 6–8 shows how the CALL instruction links to a procedure and how the RET instruction returns in the 8086–Core2 operating in the real mode.

There is one other form of the return instruction, which adds a number to the contents of the stack pointer (SP) after the return address is removed from the stack. A return that uses an immediate operand is ideal for use in a system that uses the C/C++ or PASCAL calling conventions. (This is true even though the C/C++ and PASCAL calling conventions require the caller to remove stack data for many functions.) These conventions push parameters on the stack before

FIGURE 6–8 The effect of a near return instruction on the stack and instruction pointer.

calling a procedure. If the parameters are to be discarded upon return, the return instruction contains a number that represents the number of bytes pushed to the stack as parameters.

Example 6–17 shows how this type of return erases the data placed on the stack by a few pushes. The RET 4 adds a 4 to SP after removing the return address from the stack. Because the PUSH AX and PUSH BX together place 4 bytes of data on the stack, this return effectively deletes AX and BX from the stack. This type of return rarely appears in assembly language programs, but it is used in high-level programs to clear stack data after a procedure. Notice how parameters are addressed on the stack by using the BP register, which by default addresses the stack segment. Parameter stacking is common in procedures written for C++ or PASCAL by using the C++ or PASCAL calling conventions.

EXAMPLE 6–17

```
0000 B8 001E                 MOV   AX,30
0003 BB 0028                 MOV   BX,40
0006 50                      PUSH  AX            ;stack parameter 1
0007 53                      PUSH  BX            ;stack parameter 2
0008 E8 0066                 CALL  ADDM          ;add stack parameters


0071             ADDM  PROC  NEAR
0071 55                      PUSH  BP            ;save BP
0072 8B EC                   MOV   BP,SP         ;address stack with BP
0074 8B 46 04                MOV   AX,[BP+4]     ;get parameter 1
0077 03 46 06                ADD   AX,[BP+6]     ;add parameter 2
007A 5D                      POP   BP            ;restore BP
007B C2 0004                 RET   4             ;return, dump parameters
007E             ADDM  ENDP
```

As with the CALLN and CALLF instructions, there are also variants of the return instruction: RETN and RETF. As with the CALLN and CALLF instructions, these variants should also be avoided in favor of using the PROC statement to define the type of call and return.

6–4   INTRODUCTION TO INTERRUPTS

An interrupt is either a **hardware-generated CALL** (externally derived from a hardware signal) or a **software-generated CALL** (internally derived from the execution of an instruction or by some other internal event). At times, an internal interrupt is called an *exception.* Either type interrupts the program by calling an **interrupt service procedure** (ISP) or interrupt handler.

This section explains software interrupts, which are special types of CALL instructions. This section describes the three types of software interrupt instructions (INT, INTO, and INT 3), provides a map of the interrupt vectors, and explains the purpose of the special interrupt return instruction (IRET).

Interrupt Vectors

An **interrupt vector** is a 4-byte number stored in the first 1024 bytes of the memory (00000H–003FFH) when the microprocessor operates in the real mode. In the protected mode, the vector table is replaced by an interrupt descriptor table that uses 8-byte descriptors to describe each of the interrupts. There are 256 different interrupt vectors, and each vector contains the address of an interrupt service procedure. Table 6–4 lists the interrupt vectors, with a brief description and the memory location of each vector for the real mode. Each vector contains a value for IP and CS that forms the address of the interrupt service procedure. The first 2 bytes contain the IP, and the last 2 bytes contain the CS.

TABLE 6–4   Interrupt vectors defined by Intel.

Number	Address	Microprocessor	Function
0	0H–3H	All	Divide error
1	4H–7H	All	Single-step
2	8–BH	All	NMI pin
3	CH–FH	All	Breakpoint
4	10H–13H	All	Interrupt on overflow
5	14H–17H	80186–Core2	Bound instruction
6	18H–1BH	80186–Core2	Invalid opcode
7	1CH–1FH	80186–Core2	Coprocessor emulation
8	20H–23H	80386–Core2	Double fault
9	24H–27H	80386	Coprocessor segment overrun
A	28H–2BH	80386–Core2	Invalid task state segment
B	2CH–2FH	80386–Core2	Segment not present
C	30H–33H	80386–Core2	Stack fault
D	34H–37H	80386–Core2	General protection fault (GPF)
E	38H–3BH	80386–Core2	Page fault
F	3CH–3FH	—	Reserved
10	40H–43H	80286–Core2	Floating-point error
11	44H–47H	80486SX	Alignment check interrupt
12	48H–4BH	Pentium–Core2	Machine check exception
13–1F	4CH–7FH	—	Reserved
20–FF	80H–3FFH	—	User interrupts

Intel reserves the first 32 interrupt vectors for the present and future microprocessor products. The remaining interrupt vectors (32–255) are available for the user. Some of the reserved vectors are for errors that occur during the execution of software, such as the divide error interrupt. Some vectors are reserved for the coprocessor. Still others occur for normal events in the system. In a personal computer, the reserved vectors are used for system functions, as detailed later in this section. Vectors 1–6, 7, 9, 16, and 17 function in the real mode and protected mode; the remaining vectors function only in the protected mode.

Interrupt Instructions

The microprocessor has three different interrupt instructions that are available to the programmer: INT, INTO, and INT 3. In the real mode, each of these instructions fetches a vector from the vector table, and then calls the procedure stored at the location addressed by the vector. In the protected mode, each of these instructions fetches an interrupt descriptor from the interrupt descriptor table. The descriptor specifies the address of the interrupt service procedure. The interrupt call is similar to a far CALL instruction because it places the return address (IP/EIP and CS) on the stack.

INTs.   There are 256 different software interrupt instructions (INTs) available to the programmer. Each INT instruction has a numeric operand whose range is 0 to 255 (00H–FFH). For example, the INT 100 uses interrupt vector 100, which appears at memory address 190H–193H. The address of the interrupt vector is determined by multiplying the interrupt type number by 4. For example, the INT 10H instruction calls the interrupt service procedure whose address is stored beginning at memory location 40H (10H × 4) in the real mode. In the protected mode, the interrupt descriptor is located by multiplying the type number by 8 instead of 4 because each descriptor is 8 bytes long.

Each INT instruction is 2 bytes long. The first byte contains the opcode, and the second byte contains the vector type number. The only exception to this is INT 3, a 1-byte special software interrupt used for breakpoints.

Whenever a software interrupt instruction executes, it (1) pushes the flags onto the stack, (2) clears the T and I flag bits, (3) pushes CS onto the stack, (4) fetches the new value for CS from the interrupt vector, (5) pushes IP/EIP onto the stack, (6) fetches the new value for IP/EIP from the vector, and (7) jumps to the new location addressed by CS and IP/EIP.

The INT instruction performs as a far CALL except that it not only pushes CS and IP onto the stack, but it also pushes the flags onto the stack. The INT instruction performs the operation of a PUSHF, followed by a far CALL instruction.

Notice that when the INT instruction executes, it clears the interrupt flag (I), which controls the external hardware interrupt input pin INTR (interrupt request). When I = 0, the microprocessor disables the INTR pin; when I = 1, the microprocessor enables the INTR pin.

Software interrupts are most commonly used to call system procedures because the address of the system function need not be known. The system procedures are common to all system and application software. The interrupts often control printers, video displays, and disk drives. Besides relieving the program from remembering the address of the system call, the INT instruction replaces a far CALL that would otherwise be used to call a system function. The INT instruction is 2 bytes long, whereas the far CALL is 5 bytes long. Each time that the INT instruction replaces a far CALL, it saves 3 bytes of memory in a program. This can amount to a sizable saving if the INT instruction often appears in a program, as it does for system calls.

IRET/IRETD.   The interrupt return instruction (IRET) is used only with software or hardware interrupt service procedures. Unlike a simple return instruction (RET), the IRET instruction will (1) pop stack data back into the IP, (2) pop stack data back into CS, and (3) pop stack data back into the flag register. The IRET instruction accomplishes the same tasks as the POPF, followed by a far RET instruction.

Whenever an IRET instruction executes, it restores the contents of I and T from the stack. This is important because it preserves the state of these flag bits. If interrupts were enabled before an interrupt service procedure, they are automatically re-enabled by the IRET instruction because it restores the flag register.

In the 80386 through the Core2 processors, the IRETD instruction is used to return from an interrupt service procedure that is called in the protected mode. It differs from the IRET because it pops a 32-bit instruction pointer (EIP) from the stack. The IRET is used in the real mode and the IRETD is used in the protected mode.

INT 3. An INT 3 instruction is a special software interrupt designed to function as a breakpoint. The difference between it and the other software interrupts is that INT 3 is a 1-byte instruction, while the others are 2-byte instructions.

It is common to insert an INT 3 instruction in software to interrupt or break the flow of the software. This function is called a breakpoint. A breakpoint occurs for any software interrupt, but because INT 3 is 1 byte long, it is easier to use for this function. Breakpoints help to debug faulty software.

INTO. Interrupt on overflow (INTO) is a conditional software interrupt that tests the overflow flag (O). If O = 0, the INTO instruction performs no operation; if O = 1 and an INTO instruction executes, an interrupt occurs via vector type number 4.

The INTO instruction appears in software that adds or subtracts signed binary numbers. With these operations, it is possible to have an overflow. Either the JO instruction or INTO instruction detects the overflow condition.

An Interrupt Service Procedure. Suppose that, in a particular system, a procedure is required to add the contents of DI, SI, BP, and BX and then save the sum in AX. Because this is a common task in this system, it may occasionally be worthwhile to develop the task as a software interrupt. Realize that interrupts are usually reserved for system events and this is merely an example showing how an interrupt service procedure appears. Example 6–18 shows this software interrupt. The main difference between this procedure and a normal far procedure is that it ends with the IRET instruction instead of the RET instruction, and the contents of the flag register are saved on the stack during its execution. It is also important to save all registers that are changed by the procedure using USES.
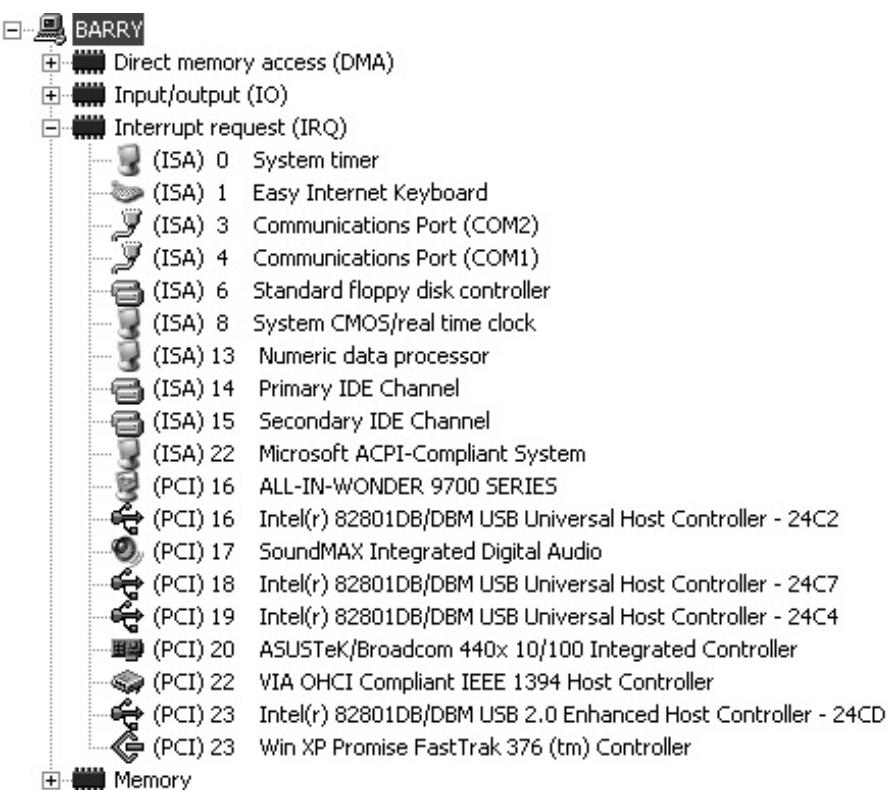
EXAMPLE 6–18

```
0000                    INTS   PROC   FAR USES AX
0000 03 C3                     ADD    AX,BX
0002 03 05                     ADD    AX,BP
0004 03 C7                     ADD    AX,DI
0006 03 C6                     ADD    AX,SI
0008 CF                        IRET
0009                    INTS   ENDP
```

Interrupt Control

Although this section does not explain hardware interrupts, two instructions are introduced that control the INTR pin. The **set interrupt flag** instruction (STI) places a 1 into the I flag bit, which enables the INTR pin. The **clear interrupt flag** instruction (CLI) places a 0 into the I flag bit, which disables the INTR pin. The STI instruction enables INTR and the CLI instruction disables INTR. In a software interrupt service procedure, hardware interrupts are enabled as one of the first steps. This is accomplished by the STI instruction. The reason interrupts are enabled early in an interrupt service procedure is that just about all of the I/O devices in the personal computer are interrupt-processed. If the interrupts are disabled too long, severe system problems result.

FIGURE 6–9   Interrupts in a typical personal computer.

Interrupts in the Personal Computer

The interrupts found in the personal computer differ somewhat from the ones presented in Table 6–4. The reason that they differ is that the original personal computers are 8086/8088-based systems. This meant that they only contained Intel-specified interrupts 0–4. This design has been carried forward so that newer systems are compatible with the early personal computers.

Access to the protected mode interrupt structure in use by Windows is accomplished through kernel functions Microsoft provides and cannot be directly addressed. Protected mode interrupts use an interrupt descriptor table, which is beyond the scope of the text at this point. Protected mode interrupts are discussed completely in later chapters.

Figure 6–9 illustrates the interrupts available in the author's computer. The interrupt assignments are viewable in the control panel of Windows under Performance and Maintenance by clicking on System and selecting Hardware and then Device Manager. Now click on View and select Device by Type and finally Interrupts.

64-Bit Mode Interrupts

The 64-bit system uses the IRETQ instruction to return from an interrupt service procedure. The main difference between IRET/IRETD and the IRETQ instruction is that IRETQ retrieves an 8-byte return address from the stack. The IRETQ instruction also retrieves the 32-bit EFLAG register from the stack and places it into the RFLAG register. It appears that Intel has no plans for using the leftmost 32 bits of the RFLAG register. Otherwise, 64-bit mode interrupts are the same as 32-bit mode interrupts.

6–5   ## MACHINE CONTROL AND MISCELLANEOUS INSTRUCTIONS

The last category of real mode instructions found in the microprocessor is the machine control and miscellaneous group. These instructions provide control of the carry bit, sample the BUSY/TEST pin, and perform various other functions. Because many of these instructions are used in hardware control, they need only be explained briefly at this point.

Controlling the Carry Flag Bit

The carry flag (C) propagates the carry or borrow in multiple-word/doubleword addition and subtraction. It also can indicate errors in assembly language procedures. Three instructions control the contents of the carry flag: STC (set carry), CLC (clear carry), and CMC (complement carry).

Because the carry flag is seldom used except with multiple-word addition and subtraction, it is available for other uses. The most common task for the carry flag is to indicate an error upon return from a procedure. Suppose that a procedure reads data from a disk memory file. This operation can be successful, or an error such as file-not-found can occur. Upon return from this procedure, if C = 1, an error has occurred; if C = 0, no error occurred. Most of the DOS and BIOS procedures use the carry flag to indicate error conditions. This flag is not available in Visual C/C++ for use with C++.

WAIT

The WAIT instruction monitors the hardware $\overline{BUSY}$ pin on the 80286 and 80386, and the $\overline{TEST}$ pin on the 8086/8088. The name of this pin was changed beginning with the 80286 microprocessor from $\overline{TEST}$ to $\overline{BUSY}$. If the WAIT instruction executes while the $\overline{BUSY}$ pin = 1, nothing happens and the next instruction executes. If the $\overline{BUSY}$ pin = 0 when the WAIT instruction executes, the microprocessor waits for the $\overline{BUSY}$ pin to return to a logic 1. This pin inputs a busy condition when at a logic 0 level.

The $\overline{BUSY}/\overline{TEST}$ pin of the microprocessor is usually connected to the $\overline{BUSY}$ pin of the 8087 through the 80387 numeric coprocessors. This connection allows the microprocessor to wait until the coprocessor finishes a task. Because the coprocessor is inside an 80486 through the Core2, the $\overline{BUSY}$ pin is not present in these microprocessors.

HLT

The halt instruction (HLT) stops the execution of software. There are three ways to exit a halt: by an interrupt, by a hardware reset, or during a DMA operation. This instruction normally appears in a program to wait for an interrupt. It often synchronizes external hardware interrupts with the software system. Note that DOS and Windows both use interrupts extensively, so HLT will not halt the computer when operated under these operating systems.

NOP

When the microprocessor encounters a no operation instruction (NOP), it takes a short time to execute. In early years, before software development tools were available, a NOP, which performs absolutely no operation, was often used to pad software with space for future machine language instructions. If you are developing machine language programs, which are extremely rare, it is recommended that you place 10 or so NOPS in your program at 50-byte intervals. This is done in case you need to add instructions at some future point. A NOP may also find application in time delays to waste time. Realize that a NOP used for timing is not very accurate because of the cache and pipelines in modem microprocessors.

LOCK Prefix

The LOCK prefix appends an instruction and causes the $\overline{LOCK}$ pin to become a logic 0. The $\overline{LOCK}$ pin often disables external bus masters or other system components. The LOCK prefix causes the $\overline{LOCK}$ pin to activate for only the duration of a locked instruction. If more than one sequential instruction is locked, the $\overline{LOCK}$ pin remains a logic 0 for the duration of the sequence of locked instructions. The LOCK:MOV AL,[SI] instruction is an example of a locked instruction.

ESC

The escape (ESC) instruction passes instructions to the floating-point coprocessor from the microprocessor. Whenever an ESC instruction executes, the microprocessor provides the memory address, if required, but otherwise performs a NOP. Six bits of the ESC instruction provide the opcode to the coprocessor and begin executing a coprocessor instruction.

The ESC opcode never appears in a program as ESC and in itself is considered obsolete as an opcode. In its place are a set of coprocessor instructions (FLD, FST, FMUL, etc.) that assemble as ESC instructions for the coprocessor. More detail is provided in Chapter 13, which details the 8087–Core2 numeric coprocessors.

BOUND

The BOUND instruction, first made available in the 80186 microprocessor, is a comparison instruction that may cause an interrupt (vector type number 5). This instruction compares the contents of any 16-bit or 32-bit register against the contents of two words or doublewords of memory: an upper and a lower boundary. If the value in the register compared with memory is not within the upper and lower boundary, a type 5 interrupt ensues. If it is within the boundary, the next instruction in the program executes.

For example, if the BOUND SI,DATA instruction executes, word-sized location DATA contains the lower boundary, and word-sized location DATA+2 bytes contains the upper boundary. If the number contained in SI is less than memory location DATA or greater than memory location DATA+2 bytes, a type 5 interrupt occurs. Note that when this interrupt occurs, the return address points to the BOUND instruction, not to the instruction following BOUND. This differs from a normal interrupt, where the return address points to the next instruction in the program.
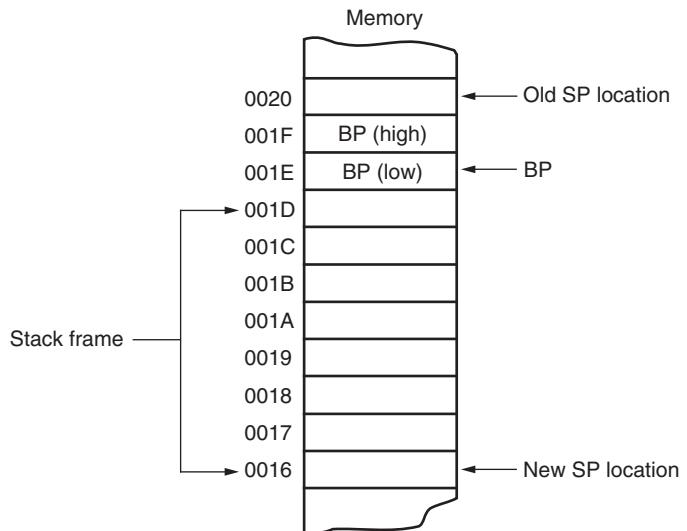
ENTER and LEAVE

The ENTER and LEAVE instructions, first made available to the 80186 microprocessor, are used with stack frames, which are mechanisms used to pass parameters to a procedure through the stack memory. The stack frame also holds local memory variables for the procedure. Stack frames provide dynamic areas of memory for procedures in multiuser environments.

The ENTER instruction creates a stack frame by pushing BP onto the stack and then loading BP with the uppermost address of the stack frame. This allows stack frame variables to be accessed through the BP register. The ENTER instruction contains two operands: The first operand specifies the number of bytes to reserve for variables on the stack frame, and the second specifies the level of the procedure.

Suppose that an ENTER 8,0 instruction executes. This instruction reserves 8 bytes of memory for the stack frame and the zero specifies level 0. Figure 6–10 shows the stack frame set up by this instruction. Note that this instruction stores BP onto the top of the stack. It then subtracts 8 from the stack pointer, leaving 8 bytes of memory space for temporary data storage. The uppermost location of this 8-byte temporary storage area is addressed by BP. The LEAVE instruction reverses this process by reloading both SP and BP with their prior values. The ENTER and LEAVE instructions were used to call C++ functions in Windows 3.1, but since then, CALL has been used in modern versions of Windows for C++ functions.

FIGURE 6–10 The stack frame created by the ENTER 8,0 instruction. Notice that BP is stored beginning at the top of the stack frame. This is followed by an 8-byte area called a stack frame.

6–6 SUMMARY

1. There are three types of unconditional jump instructions: short, near, and far. The short jump allows a branch to within +127 and −128 bytes. The near jump (using a displacement of ±32K) allows a jump to any location in the current code segment (intrasegment). The far jump allows a jump to any location in the memory system (intersegment). The near jump in an 80386 through a Core2 is within ±2G bytes because these microprocessors can use a 32-bit signed displacement.

2. Whenever a label appears with a JMP instruction or conditional jump, the label, located in the label field, must be followed by a colon (LABEL:). For example, the JMP DOGGY instruction jumps to memory location DOGGY:.

3. The displacement that follows a short or near jump is the distance from the next instruction to the jump location.

4. Indirect jumps are available in two forms: (1) jump to the location stored in a register and (2) jump to the location stored in a memory word (near indirect) or doubleword (far indirect).

5. Conditional jumps are all short jumps that test one or more of the flag bits: C, Z, O, P, or S. If the condition is true, a jump occurs; if the condition is false, the next sequential instruction executes. Note that the 80386 and above allow a 16-bit signed displacement for the conditional jump instructions. In 64-bit mode, the displacement is 32 bits allowing a range of ±2G.

6. A special conditional jump instruction (LOOP) decrements CX and jumps to the label when CX is not 0. Other forms of loop include LOOPE, LOOPNE, LOOPZ, and LOOPNZ. The LOOPE instruction jumps if CX is not 0 and if an equal condition exists. In the 80386 through the Core2, the LOOPD, LOOPED, and LOOPNED instructions also use the ECX register as a counter. In the 64-bit mode, these instructions use the RCX register as for iteration.

7. The 80386 through the Core2 contain conditional set instructions that either set a byte to 01H or clear it to 00H. If the condition under test is true, the operand byte is set to 01H; if the condition under test is false, the operand byte is cleared to 00H.

8. The .IF and .ENDIF statements are useful in assembly language for making decisions. The instructions cause the assembler to generate conditional jump statements that modify the flow of the program.

9. The .WHILE and .ENDW statements allow an assembly language program to use the WHILE construction, and the .REPEAT and .UNTIL statements allow an assembly language program to use the REPEAT-UNTIL construct.

10. Procedures are groups of instructions that perform one task and are used from any point in a program. The CALL instruction links to a procedure and the RET instruction returns from a procedure. In assembly language, the PROC directive defines the name and type of procedure. The ENDP directive declares the end of the procedure.

11. The CALL instruction is a combination of a PUSH and a JMP instruction. When CALL executes, it pushes the return address on the stack and then jumps to the procedure. A near CALL places the contents of IP on the stack, and a far CALL places both IP and CS on the stack.

12. The RET instruction returns from a procedure by removing the return address from the stack and placing it into IP (near return), or IP and CS (far return).

13. Interrupts are either software instructions similar to CALL or hardware signals used to call procedures. This process interrupts the current program and calls a procedure. After the procedure, a special IRET instruction returns control to the interrupted software.

14. Real mode interrupt vectors are 4 bytes long and contain the address (IP and CS) of the interrupt service procedure. The microprocessor contains 256 interrupt vectors in the first 1K bytes of memory. The first 32 are defined by Intel; the remaining 224 are user interrupts. In protected mode operation, the interrupt vector is 8 bytes long and the interrupt vector table may be relocated to any section of the memory system.

15. Whenever an interrupt is accepted by the microprocessor, the flags IP and CS are pushed onto the stack. Besides pushing the flags, the T and I flag bits are cleared to disable both the trace function and the INTR pin. The final event that occurs for the interrupt is that the interrupt vector is fetched from the vector table and a jump to the interrupt service procedure occurs.

16. Software interrupt instructions (INT) often replace system calls. Software interrupts save 3 bytes of memory each time they replace CALL instructions.

17. A special return instruction (IRET) must be used to return from an interrupt service procedure. The IRET instruction not only removes IP and CS from the stack, it also removes the flags from the stack.

18. Interrupt on an overflow (INTO) is a conditional interrupt that calls an interrupt service procedure if the overflow flag (O) = 1.

19. The interrupt enable flag (I) controls the INTR pin connection on the microprocessor. If the STI instruction executes, it sets I to enable the INTR pin. If the CLI instruction executes, it clears I to disable the INTR pin.

20. The carry flag bit (C) is clear, set, and complemented by the CLC, STC, and CMC instructions.

21. The WAIT instruction tests the condition of the $\overline{BUSY}$ or $\overline{TEST}$ pin on the microprocessor. If $\overline{BUSY}$ or $\overline{TEST}$ = 1, WAIT does not wait; but if $\overline{BUSY}$ or $\overline{TEST}$ = 0, WAIT continues testing the $\overline{BUSY}$ or $\overline{TEST}$ pin until it becomes a logic 1. Note that the 8086/8088 contains the $\overline{TEST}$ pin, while the 80286–80386 contain the $\overline{BUSY}$ pin. The 80486 through the Core2 do not contain a $\overline{BUSY}$ or $\overline{TEST}$ pin.

22. The LOCK prefix causes the $\overline{LOCK}$ pin to become a logic 0 for the duration of the locked instruction. The ESC instruction passes instruction to the numeric coprocessor.

23. The BOUND instruction compares the contents of any 16-bit register against the contents of two words of memory: an upper and a lower boundary. If the value in the register compared with memory is not within the upper and lower boundary, a type 5 interrupt ensues.

24. The ENTER and LEAVE instructions are used with stack frames. A stack frame is a mecha-
     nism used to pass parameters to a procedure through the stack memory. The stack frame also
     holds local memory variables for the procedure. The ENTER instruction creates the stack
     frame, and the LEAVE instruction removes the stack frame from the stack. The BP register
     addresses stack frame data.

6–7   QUESTIONS AND PROBLEMS

1. What is a short JMP?
2. Which type of JMP is used when jumping to any location within the current code
    segment?
3. Which JMP instruction allows the program to continue execution at any memory location in
    the system?
4. Which JMP instruction is 5 bytes long?
5. What is the range of a near jump in the 80386–Core2 microprocessors?
6. Which type of JMP instruction (short, near, or far) assembles for the following:
    (a) if the distance is 0210H bytes
    (b) if the distance is 0020H bytes
    (c) if the distance is 10000H bytes
7. What can be said about a label that is followed by a colon?
8. The near jump modifies the program address by changing which register or registers?
9. The far jump modifies the program address by changing which register or registers?
10. Explain what the JMP AX instruction accomplishes. Also identify it as a near or a far jump
    instruction.
11. Contrast the operation of a JMP DI with a JMP [DI].
12. Contrast the operation of a JMP [DI] with a JMP FAR PTR [DI].
13. List the five flag bits tested by the conditional jump instructions.
14. Describe how the JA instruction operates.
15. When will the JO instruction jump?
16. Which conditional jump instructions follow the comparison of signed numbers?
17. Which conditional jump instructions follow the comparison of unsigned numbers?
18. Which conditional jump instructions test both the Z and C flag bits?
19. When does the JCXZ instruction jump?
20. Which SET instruction is used to set AL if the flag bits indicate a zero condition?
21. The 8086 LOOP instruction decrements register _____ and tests it for a 0 to decide
    if a jump occurs.
22. The Pentium 4 LOOPD instruction decrements register _____ and tests it for a 0 to
    decide if a jump occurs.
23. The Core2 operated in 64-bit mode for a LOOP instruction decrements register
    _____ and tests it for a 0 to decide if a jump occurs.
24. Develop a short sequence of instructions that stores 00H into 150H bytes of memory, begin-
    ning at extra segment memory location DATAZ. You must use the LOOP instruction to help
    perform this task.
25. Explain how the LOOPE instruction operates.
26. Show the assembly language instructions are generated by the following sequence:

```
.IF AL==3
   ADD AL,2
.ENDIF
```

27. Develop a sequence of instructions that searches through a block of 100H bytes of memory. This program must count all the unsigned numbers that are above 42H and all that are below 42H. Byte-sized data segment memory location UP must contain the count of numbers above 42H, and data segment location DOWN must contain the count of numbers below 42H.

28. Develop a short sequence of instructions that uses the REPEAT-UNTIL construct to copy the contents of byte-sized memory BLOCKA into byte-sized memory BLOCKB until 00H is moved.

29. What happens if the .WHILE 1 instruction is placed in a program?

30. Using the WHILE construct, develop a sequence of instructions that add the byte-sized contents of BLOCKA to BLOCKB while the sum is not 12H.

31. What is the purpose of the .BREAK directive?

32. What is a procedure?

33. Explain how the near and far CALL instructions function.

34. The last executable instruction in a procedure must be a(n) _____.

35. How does the near RET instruction function?

36. How is a procedure identified as near or far?

37. Which directive identifies the start of a procedure?

38. Write a near procedure that cubes the contents of the CX register. This procedure may not affect any register except CX.

39. Explain what the RET 6 instruction accomplishes.

40. Write a procedure that multiplies DI by SI and then divides the result by 100H. Make sure that the result is left in AX upon returning from the procedure. This procedure may not change any register except AX.

41. Write a procedure that sums EAX, EBX, ECX, and EDX. If a carry occurs, place a logic 1 in EDI. If no carry occurs, place a 0 in EDI. The sum should be found in EAX after the execution of your procedure.

42. What is an interrupt?

43. Which software instructions call an interrupt service procedure?

44. How many different interrupt types are available in the microprocessor?

45. Illustrate the contents of an interrupt vector and explain the purpose of each part.

46. What is the purpose of interrupt vector type number 0?

47. How does the IRET instruction differ from the RET instruction?

48. What is the IRETD instruction?

49. What is the IRETQ instruction?

50. The INTO instruction only interrupts the program for what condition?

51. The interrupt vector for an INT 40H instruction is stored at which memory locations?

52. What instructions control the function of the INTR pin?

53. What instruction tests the $\overline{BUSY}$ pin?

54. When will the BOUND instruction interrupt a program?

55. An ENTER 16,0 instruction creates a stack frame that contains _____ bytes.

56. Which register moves to the stack when an ENTER instruction executes?

57. Which instruction passes opcodes to the numeric coprocessor?

CHAPTER 7

Using Assembly Language with C/C++

INTRODUCTION

Today, it is rare to develop a complete system using only assembly language. We often use C/C++ with some assembly language to develop a system. The assembly language portion usually solves tasks (difficult or inefficient to accomplish in C/C++) that often include control software for peripheral interfaces and driver programs that use interrupts. Another application of assembly language in C/C++ programs is the MMX and SEC instructions that are part of the Pentium class microprocessor and not supported in C/C++. Although C++ does have macros for these commands, they are more complicated to use than using assembly language. This chapter develops the idea of mixing C/C++ and assembly language. Many applications in later chapters also illustrate the use of both assembly language and C/C++ to accomplish tasks for the microprocessor.

This text uses Microsoft Visual C/C++ Express, but programs can often be adapted to any version of C/C++, as long as it is standard ANSI (**American National Standards Institute**) format C/C++. If you want, you can use C/C++ to enter and execute all the programming applications in this text. The 16-bit applications are written by using Microsoft Visual C/C++ version 1.52 or newer (available [**CL.EXE**] for no cost as a legacy application in the **Microsoft Windows Driver Development Kit** [DDK]); the 32-bit applications are written using Microsoft Visual C/C++ version 6 or newer and preferably Microsoft Visual C/C++ version .NET 2003 or Visual C++ Express. The examples in the text are written assuming that you have the latest version of **Visual C++ Express**, which is a free downloadable version of Visual C++. Please visit http://msdn.com to obtain the Visual C++ Express program.

CHAPTER OBJECTIVES

Upon completion of this chapter, you will be able to:

1. Use assembly language in _asm blocks within C/C+.
2. Learn the rules that apply to mixed language software development.
3. Use common C/C++ data and structures with assembly language.
4. Use both the 16-bit (DOS) interface and the 32-bit (Microsoft Windows) interface with assembly language code.
5. Use assembly language objects with C/C++ programs.

7–1   USING ASSEMBLY LANGUAGE WITH C++ FOR 16-BIT DOS APPLICATIONS

This section shows how to incorporate assembly language commands within a C/C++ program. This is important because the performance of a program often depends on the incorporation of assembly language sequences to speed its execution. As mentioned in the introduction to the chapter, assembly language is also used for I/O operations in embedded systems. This text assumes that you are using a version of the Microsoft C/C++ program, but any C/C++ program should function as shown, if it supports inline assembly commands. The only change might be setting up the C/C++ package to function with assembly language. This section of the text assumes that you are building l6-bit applications for DOS. Make sure that your software can build l6-bit applications before attempting any of the programs in this section. If you build a 32-bit application and attempt to use the DOS INT 21H function, the program will crash because DOS calls are not directly allowed. In fact, they are inefficient to use in a 32-bit application.

To build a 16-bit **DOS** application, you will need the legacy 16-bit compiler usually found in the C:\WINDDK\2600.1106\bin\win_me\bin16 directory of the Windows DDK. (The Windows driver development kit can be obtained for a small shipping charge from Microsoft Corporation.) The compiler is CL.EXE and the 16-bit linker program is LINK.EXE, both located in the directory or folder listed. Because the path in the computer that you are using probably points to the 32-bit linker program, it would be wise to work from this directory so the proper linker is used when linking the object files generated by the compiler. Compilation and linking must be performed at the command line because there is no visual interface or editor provided with the compiler and linker. Programs are generated using either Notepad or DOS Edit.

Basic Rules and Simple Programs

Before assembly language code can be placed in a C/C++ program, some rules must be learned. Example 7–1 shows how to place assembly code inside an assembly language block within a short C/C++ program. Note that all the assembly code in this example is placed in the **_asm** block. Labels are used as illustrated by the label **big:** in this example. It is also extremely important to use lowercase characters for any inline assembly code. If you use uppercase, you will find that some of the assembly language commands and registers are reserved or defined words in C/C++ language.

Example 7–1 uses no C/C++ commands except for the main procedure. Enter the program using either WordPad or Edit. This program (Example 7–1) reads one character from the console keyboard, and then filters it through assembly language so that only the numbers 0 through 9 are sent back to the video display. Although this programming example does not accomplish much, it does show how to set up and use some simple programming constructs in the C/C++ environment and also how to use the inline assembler.

EXAMPLE 7–1

```
//Accepts and displays one character of 1 through 9,
//all others are ignored.

void main(void)
{
      _asm
      {
            mov  ah,8           ;read key no echo
            int  21h
            cmp  al,'0'         ;filter key code
            jb   big
            cmp  al,'9'
            ja   big
            mov  dl,al          ;echo 0 - 9
```

```
            mov  ah,2
            int  21h
      big:
       }
}
```

The register AX was not saved in Example 7–1, but it was used by the program. It is very important to note that the AX, BX, CX, DX, and ES registers are never used by Microsoft C/C++. (The function of AX on a return from a procedure is explained later in this chapter.) These registers, which might be considered **scratchpad** registers, are available to use with assembly language. If you wish to use any of the other registers, make sure that you save them with a PUSH before they are used and restore them with a POP afterwards. If you fail to save the registers used by a program, the program may not function correctly and can crash the computer. If the 80386 or above microprocessor is used as a base for the program, the EAX, EBX, ECX, EDX, and ES registers do not need to be saved. If any other registers are used, they must be saved or the program will crash.

To compile the program, start the Command Prompt program located in the Start Menu under Accessories. Change the path to C:\WINDDK\2600.1106\bin\win_me\bin16 if that is where you have your Windows DDK. You will also need to go to the C:\WINDDK\2600.1106\ lib\win_me directory and copy slibce.lib to the C:\WINDDK\2600.1106\bin\win_me\bin16 directory. Make sure you saved the program in the same path and use the extension .c with the file name. If using Notepad, make sure you select All Files under File Type when saving. To compile the program, type CL /G3 filename.c>. This will generate the .exe file (/G3 is the 80386) for the program. (See Table 7–1 for a list of the /G compiler switches.) Any errors that appear are ignored by pressing the Enter key. These errors generate warnings that will not cause a problem when the program is executed. When the program is executed, you will only see a number echoed back to the DOS screen.

Example 7–2 shows how to use variables from C with a short assembly language program. In this example, the char variable type (a byte in C) is used to save space for a few 8-bit bytes of data. The program itself performs the operation $X + Y = Z$, where X and Y are two one-digit numbers, and Z is the result. As you might imagine, you could use the inline assembly in C to learn assembly language and write many of the programs in this textbook. The semicolon adds comments to the listing in the _asm block, just as with the normal assembler.

EXAMPLE 7–2

```
void main(void)
{
      char a, b;
      _asm
      {
            mov  ah,1             ;read first digit
            int  21h
            mov  a,al
            mov  ah,1             ;read a + sign
```

TABLE 7–1 Compiler (16-bit) G options.

Compiler Switch	Function
/G1	Selects the 8088/8086
/G2	Selects the 80188/80186/80286
/G3	Selects the 80386
/G4	Selects the 80486
/G5	Selects the Pentium
/G6	Selects the Pentium Pro–Pentium 4

Note: The 32-bit C++ compiler does not recognize /G1 or /G2.

```
            int  21h
            cmp  al,'+'
            jne  end1            ;if not plus
            mov  ah,1
            int  21h             ;read second number
            mov  b,al
            mov  ah,2            ;display =
            mov  dl,'='
            int  21h
            mov  ah,0
            mov  al,a            ;generate sum
            add  al,b
            aaa                  ;ASCII adjust for addition
            add  ax,3030h
            cmp  ah,'0'
            je   down
            push ax              ;display 10's position
            mov  dl,ah
            mov  ah,2
            int  21h
            pop  ax
    down:
            mov  dl,al           ;display units position
            mov  ah,2
            int  21h
    end1:
    }
}
```

What Cannot Be Used from MASM Inside an _asm Block

Although MASM contains some nice features, such as conditional commands (.IF, .WHILE, .REPEAT, etc.), the inline assembler does not include the conditional commands from MASM, nor does it include the MACRO feature found in the assembler. Data allocation with the inline assembler is handled by C instead of by using DB, DW, DD, etc. Just about all other features are supported by the inline assembler. These omissions from the inline assembler can cause some slight problems, as will be discussed in later sections of this chapter.

Using Character Strings

Example 7–3 illustrates a simple program that uses a character string defined with C and displays it so that each word is listed on a separate line. Notice the blend of both C statements and assembly language statements. The WHILE statement repeats the assembly language commands until the null (00H) is discovered at the end of the character string. If the null is not discovered, the assembly language instructions display a character from the string unless a space is located. For each space, the program displays a carriage return/line feed combination. This causes each word in the string to be displayed on a separate line.

EXAMPLE 7–3

```
// Example that displays showing one word per line

void main(void)
{
    char strings[] = "This is my first test application using _asm. \n";
    int sc = -1;
    while (strings[sc++] != 0)
    {
        _asm
        {
            push si
            mov  si,sc           ;get pointer
```

```
                        mov  dl,strings[si]    ;get character
                        cmp  dl,' '            ;if not space
                        jne  next
                        mov  ah,2              ;display new line
                        mov  dl,10
                        int  21h
                        mov  dl,13
                 next:  mov  ah,2              ;display character
                        int  21h
                        pop  si
                 }
           }
}
```

Suppose that you want to display more than one string in a program, but you still want to use assembly language to develop the software to display a string. Example 7–4 illustrates a program that creates a procedure displaying a character string. This procedure is called each time that a string is displayed in the program. Note that this program displays one string on each line, unlike Example 7–3.

EXAMPLE 7–4

```
// A program illustrating an assembly language procedure that
// displays C language character strings

char string1[] = "This is my first test program using _asm.";
char string2[] = "This is the second line in this program.";
char string3[] = "This is the third.";

void main(void)
{
      Str (string1);
      Str (string2);
      Str (string3);
}

Str (char *string_adr)
{
      _asm
      {
            mov  bx,string_adr    ;get address of string
            mov  ah,2
      top:
            mov  dl,[bx]
            inc  bx
            cmp  al,0             ;if null
            je   bot
            int  21h              ;display character
            jmp  top
      bot:
            mov  dl,13            ;display CR + LF
            int  21h
            mov  dl,10
            int  21h
      }
}
```

Using Data Structures

Data structures are an important part of most programs. This section shows how to interface a data structure created in C with an assembly language section that manipulates the data in the structure. Example 7–5 illustrates a short program that uses a data structure to store names, ages, and

salaries. The program then displays each of the entries by using a few assembly language procedures. Although the string procedure displays a character string, shown in Example 7–4, no carriage return/line feed combination is displayed—instead, a space is displayed. The Crlf procedure displays a carriage return/line feed combination. The Numb procedure displays the integer.

EXAMPLE 7–5

```
// Program illustrating an assembly language procedure that
// displays the contents of a C data structure.

// A simple data structure

typedef struct records
{
      char first_name[16];
      char last_name[16];
      int age;
      int salary;
} RECORD;

// Fill some records

RECORD record[4] =
{ {"Bill" ,"Boyd" , 56, 23000},
  {"Page", "Turner", 32, 34000},
  {"Bull", "Dozer", 39. 22000},
  {"Hy", "Society", 48, 62000}
};

// Program

void main(void)
{
      int pnt = -1;
      while (pnt++ < 3)
      {
            Str(record[pnt].last_name);
            Str(record[pnt].first_name);
            Numb(record[pnt].age);
            Numb(record[pnt].salary);
            Crlf();
      }
}

Str (char *string_adr[])
{
      _asm
      {
            mov   bx,string_adr
            mov   ah,2
      top:
            mov   dl,[bx]
            inc   bx
            cmp   al,0
            je    bot
            int   21h
            jmp   top
      bot:
            mov   al,20h
            int   21h
      }
}
```

```
Crlf()
{
      _asm
      {
              mov  ah,2
              mov  dl,13
              int  21h
              mov  dl,10
              int  21h
      }
}

Numb (int temp)
{
      _asm
      {
              mov  ax,temp
              mov  bx,10
              push bx
      L1:
              mov  dx,0
              div  bx
              push dx
              cmp  ax,0
              jne  L1
      L2:
              pop  dx
              cmp  dl,bl
              je   L3
              mov  ah,2
              add  dl,30h
              int  21h
              jmp  L2
      L3:
              mov  dl,20h
              int  21h
      }
}
```

An Example of a Mixed-Language Program

To see how this technique can be applied to any program, Example 7–6 shows how the program can do some operations in assembly language and some in C language. Here, the only assembly language portions of the program are the Dispn procedure that displays an integer and the Readnum procedure, which reads an integer. The program in Example 7–6 makes no attempt to detect or correct errors. Also, the program functions correctly only if the result is positive and less than 64K. Notice that this example uses assembly language to perform the I/O; the C portion performs all other operations to form the shell of the program.

EXAMPLE 7–6

```
/*

A program that functions as a simple calculator to perform addition,
subtraction, multiplication, and division. The format is X <oper> Y =.

*/

int temp;

void main(void)
{
```

```
              int temp1, oper;
              while (1)
              {
                    oper = Readnum();                    //get first number and operation
                    temp1 = temp;
                    if ( Readnum() == '=' )              //get second number
                    {
                          switch (oper)
                          {
                                case '+':
                                      temp += temp1;
                                      break;
                                case '-':
                                      temp = temp1 - temp;
                                      break;
                                case '/':
                                      temp = temp1 / temp;
                                      break;
                                case '*':
                                      temp *= temp1;
                                      break;
                          }
                          Dispn(temp);                   //display result
                    }
                    else
                          Break;
                    }
              }
        }

        int Readnum()
        {
              int a;
              temp = 0;
              _asm
              {
              Readnum1:
                    mov   ah,1
                    int   21h
                    cmp   al,30h
                    jb    Readnum2
                    cmp   al,39h
                    ja    Readnum2
                    sub   al,30h
                    shl   temp,1
                    mov   bx,temp
                    shl   temp,2
                    add   temp,bx
                    add   byte ptr temp,al
                    adc   byte ptr temp+1,0
                    jmp   Readnum1
              Readnum2:
                    Mov   ah,0
                    mov   a,ax
              }
              return a;
        }

        Dispn (int DispnTemp)
        {
              _asm
              {
                    mov   ax,DispnTemp
                    mov   bx,10
                    push bx
              Dispn1:
                    mov   dx,0
```

```
                div   bx
                push  dx
                cmp   ax,0
                jne   Dispn1
        Dispn2:
                pop   dx
                cmp   dl,bl
                je    Dispn3
                add   dl,30h
                mov   ah,2
                int   21h
                jmp   Dispn2
        Dispn3:
                mov   dl,13
                int   21h
                mov   dl,10
                int   21h
        }
}
```

7–2   USING ASSEMBLY LANGUAGE WITH VISUAL C/C++ FOR 32-BIT APPLICATIONS

A major difference exists between l6-bit and 32-bit applications. The 32-bit applications are written using Microsoft Visual C/C++ Express for Windows and the l6-bit applications are written using Microsoft C++ for DOS. The main difference is that Visual C/C++ Express for Windows is more common today, but Visual C/C++ Express cannot easily call DOS functions such as INT 2lH. It is suggested that embedded applications that do not require a visual interface be written in l6-bit C or C++, and applications that incorporate Microsoft Windows or Windows CE (available for use on a ROM or Flash[1] device for embedded applications) use 32-bit Visual C/C++ Express for Windows.

A 32-bit application is written by using any of the 32-bit registers, and the memory space is essentially limited to 2G bytes for Windows. The free version of Visual C++ Express does not support 64-bit applications written in assembly language at this time. The only difference is that you may not use the DOS function calls; instead use the console getch() or getche() and putch C/C++ language functions available for use with DOS console applications. Embedded applications use direct assembly language instructions to access I/O devices in an embedded system. In the Visual interface, all I/O is handled by the Windows operating system framework.

Console applications in WIN32 run in native mode, which allow assembly language to be included in the program without anything other than the _asm keyword. Windows forms applications are more challenging because they operate in the managed mode, which does not run in the native mode of the microprocessor. Managed applications operate in a pseudo mode that does not generate native code.

An Example that Uses Console I/O to Access the Keyboard and Display

Example 7–7 illustrates a simple console application that uses the console I/O commands to read and write data from the console. To enter this application (assuming Visual Studio .NET 2003 or Visual C++ Express is available), select a WIN32 console application in the new project option (see Figure 7–1). Notice that instead of using the customary stdio.h library, we use the conio.h library in this application. This example program displays any number between 0 and 1000 in all number bases between base 2 and base 16. Notice that the main program is not called main as it was in earlier versions of C/C++, but is called _tmain in the current version of Visual C/C++ Express when used with a console application. The argc is the argument count passed to the _tmain procedure from the command line, and the argv[] is an array that contains the command line argument strings.

_____

[1]Flash is a trademark of Intel Corporation.
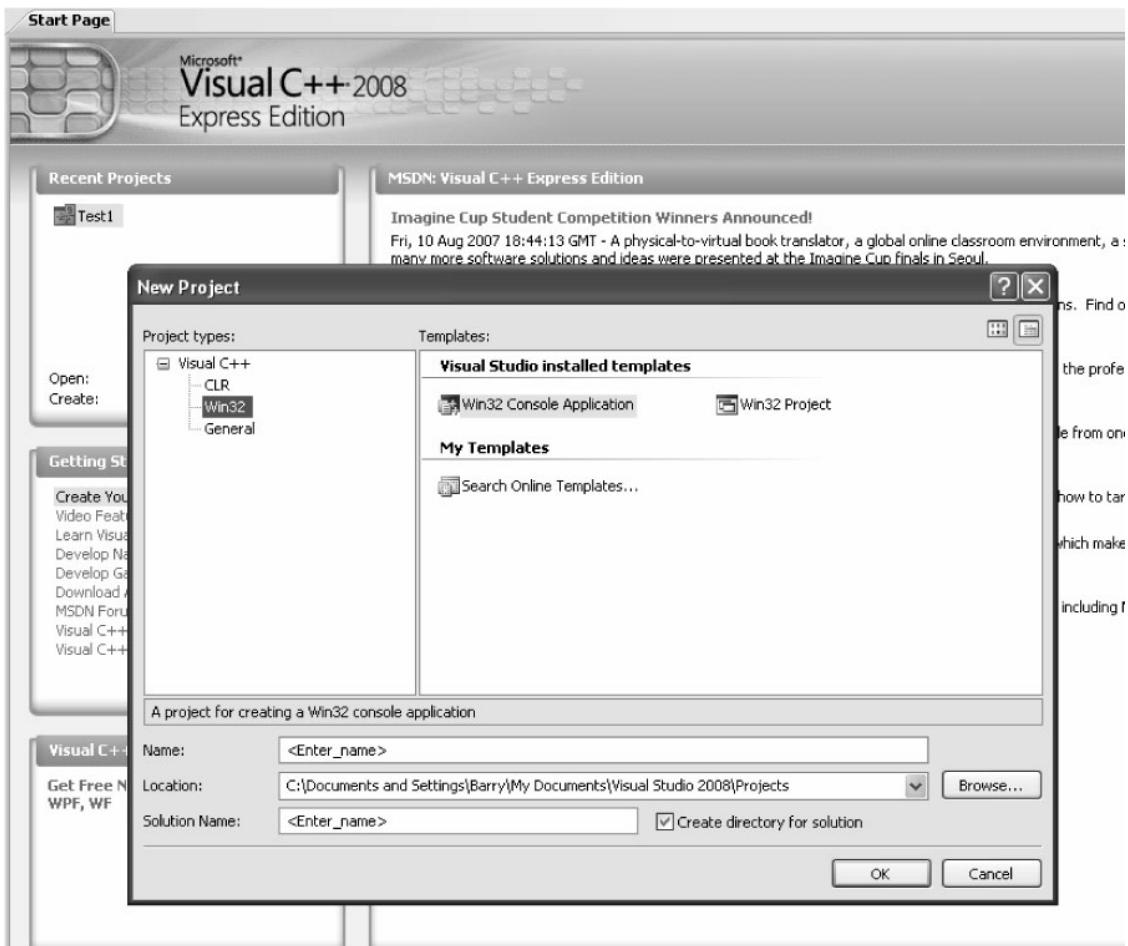
FIGURE 7–1 The new project screen selection of a WIN32 console application.

EXAMPLE 7–7

```
// Program that displays any number in all numbers bases
// between base 2 and base 16.

#include "stdafx.h"
#include <conio.h>

char *buffer = "Enter a number between 0 and 1000: ";
char *buffer1 = "Base: ";
int a, b = 0;

void disps(int base, int data);

int _tmain(int argc, _TCHAR* argv[])
{
     int i;
     _cputs(buffer);
     a = _getche();
     while ( a >= '0' && a <= '9' )
     {
          _asm sub a, 30h;
          b = b * 10 + a;
          a = _getche();
     }
```

```
            _putch(10);
            _putch(10);
            _putch(13);
            for ( i = 2; i < 17; i++ )
            {
                    _cputs(buffer1);
                    disps(10,i );
                    _putch(' ');
                    _putch('=');
                    _putch(' ');
                    disps(i, b);
                    _putch(10);
                    _putch(13);
            }
            getche();                                    //wait for any key
            return 0;
}

void disps(int base, int data)
{
            int temp;
            _asm
            {
                    mov  eax, data
                    mov  ebx, base
                    push ebx
            disps1:
                    mov  edx,0
                    div  ebx
                    push edx
                    cmp  eax,0
                    jne  disps1
            disps2:
                    pop  edx
                    cmp  ebx,edx
                    je   disps4
                    add  dl,30h
                    cmp  dl,39h
                    jbe  disps3
                    add  dl,7
            disps3:
                    mov  temp,edx
            }
            _putch(temp);
            _asm jmp disps2;
disps4:;
}
```

This example presents a mixture of assembly language and C/C++ language commands. The procedure disps (base,data) does most of the work for this program. It allows any integer (unsigned) to be displayed in any number base, which can be any value between base 2 and base 36. The upper limit occurs because letters of the alphabet only extend to the letter Z. If you need to convert larger number bases, a new scheme for bases over 36 must be developed. Perhaps the lowercase letters a through z can be used for base 37 to 52. Example 7–7 only displays the number that is entered in base 2 through base 16.

Directly Addressing I/O Ports

If a program is written that must access an actual port number, we can use console I/O commands such as the _inp(port) command to input byte data, and the _outp(port,byte_data) command to output byte data. When writing software for the personal computer, it is rare to directly address an I/O port, but when software is written for an embedded system, we often directly address an I/O port. An alternate to using the _inp and _outp commands is assembly language, which is more efficient

in most cases. Be aware that I/O ports may *not* be accessed in the Windows environment if you are using Windows NT, Windows 2000, Windows XP, or Windows Vista. The only way to access the I/O ports in these modern operating systems is to develop a kernel driver. At this point in the text it would not be practical to develop such a driver. If you are using Windows 98 or even Windows 95, you can use inp and outp instructions in C/C++ to access the I/O ports directly.

Developing a Visual C++ Application for Windows

This section of the text shows how to use Visual C++ Express to develop a dialog-based application for the Microsoft Foundation Classes library. The Microsoft Foundation Classes (MFC) is a collection of classes that allows us to use the Windows interface without a great deal of difficulty. The MFC has been renamed to the **common language runtime** (CLR) in Visual C++ Express. The easiest application to learn and develop is a program that uses a forms application as presented here. This basic application type is used to program and test all of the software examples in this textbook written in the Visual C++ Express programming environment.

To create a Visual C++ form-based application, start Visual C++ Express and click on Create Project near the upper left corner of the start screen. (If you do not have the Visual C++ Express program, it is available for free from Microsoft at http://msdn.com.) Download and install the latest version, even of it is a beta version. Figure 7–2 illustrates what is displayed
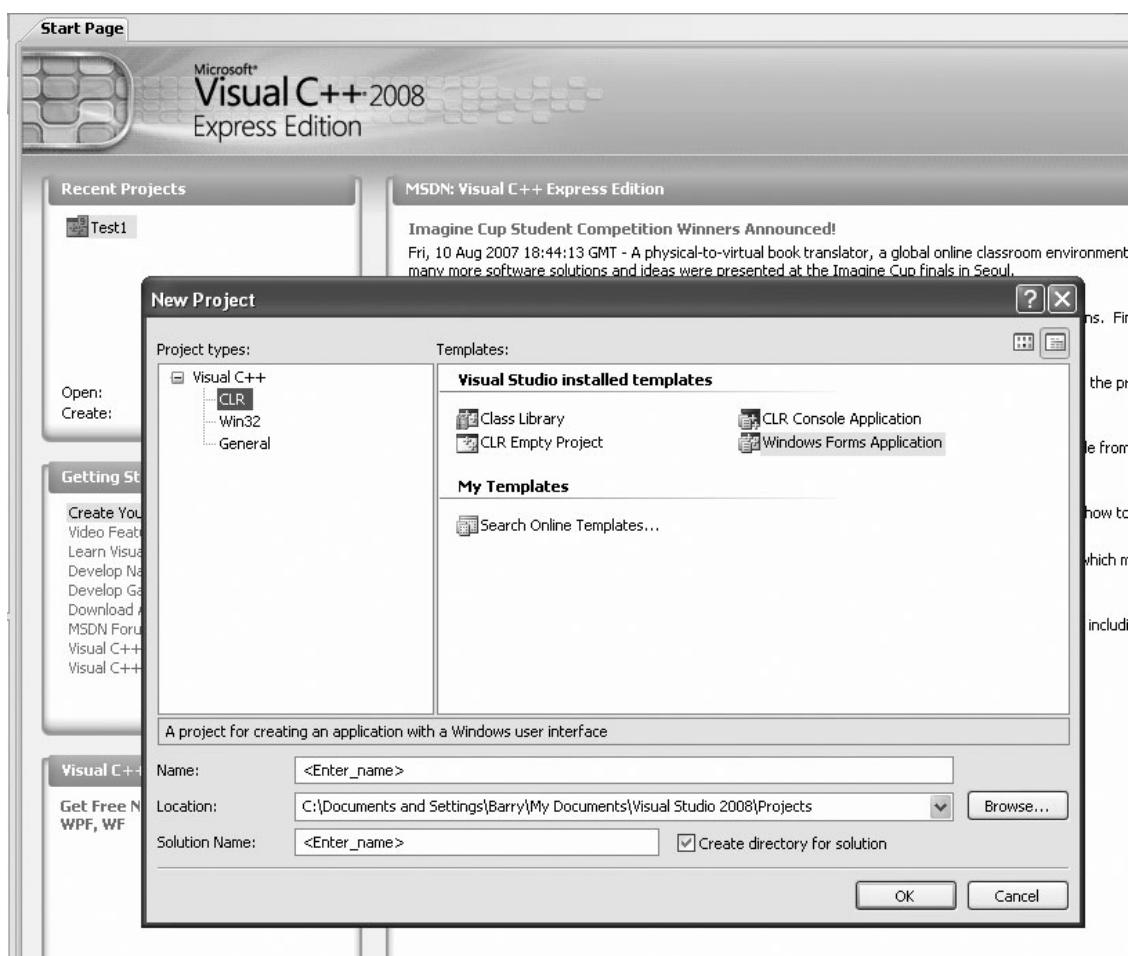
FIGURE 7–2   Starting a C++ program for Windows in Visual C++ Express.
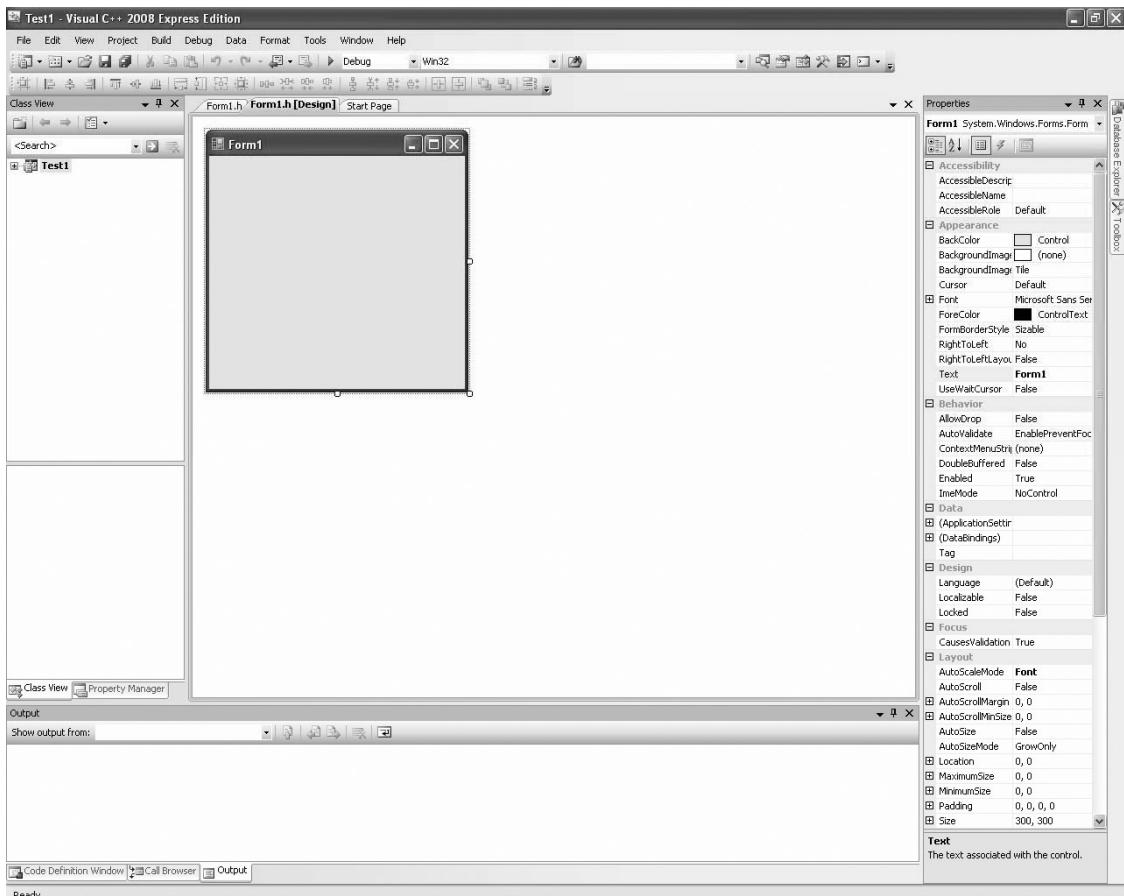
FIGURE 7–3   Design window screen shot.
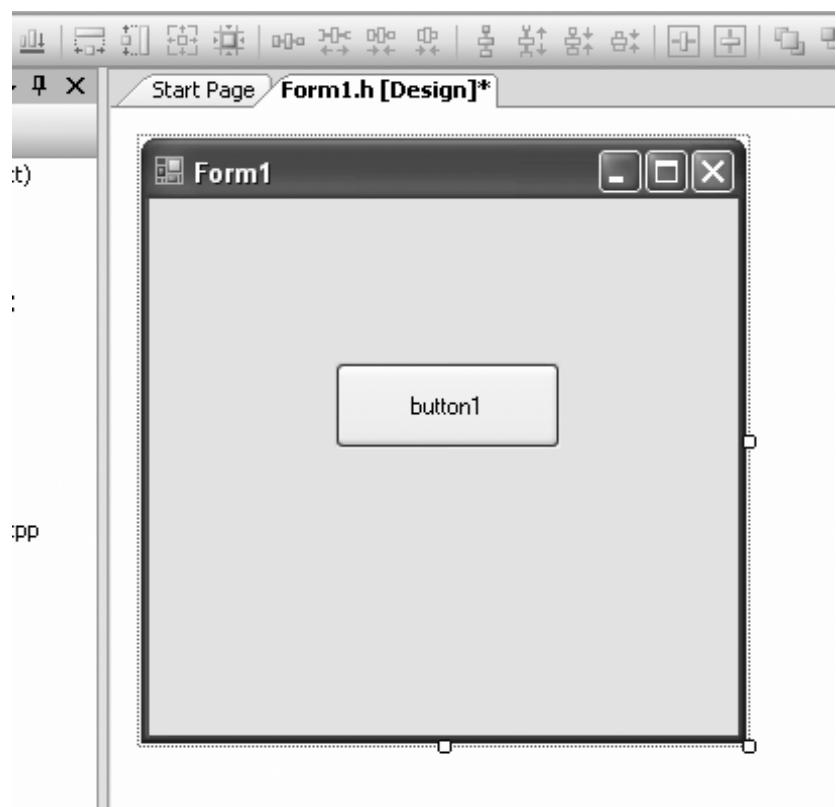
when the CLR Windows Forms application type is selected under Visual C++ Express Projects. Enter a name for the project and select an appropriate path for the project, then click on OK.

After a few moments the design screen should appear as in Figure 7–3. In the middle section is the form created by this application. To test the application, as it appears, just find the green arrow ▷ located somewhere above the form and below the Windows menu bar at the top of the screen and click on it to compile, link, and execute the dialog application. (Answer yes to "Would you like to build the application?"). Click on the X in the title bar to close the application. You have just created and tested your very first Visual C++ Express application.

When looking at the screen shot in Figure 7–3, several items are located in the image that are important to program creation and development. The right margin of the screen contains a Properties window, which contains the properties of the form. The left margin contains Solution Explorer. The tabs, located at the bottom of the Solution Explorer window, allow other views to be displayed such as a class view and so forth in this area. The tabs at the bottom of the Properties window allow the classes, properties, dynamic help, or output to be displayed in this window. Your screen may or may not appear as the one illustrated in Figure 7–3 because it can be modified and probably will be modified as you use the program.

To create a simple application, select the toolbox by clicking on Tools at the top of the screen or by opening the View dropdown menu and selecting Toolbox from the list. Windows is

an events-driven system so an object or a control is needed on the form to initiate an event. The control could be a button or almost any control object selected from the toolbox. Click on the button control near the top of the toolbox, which selects the button. Now move the mouse pointer (do not drag the button) over to the dialog application in the middle of the screen and draw, by left-clicking and resizing the button near the center (see Figure 7–4).

Once the button is placed on the screen, an event handler must be added to the application so that the act of pressing or clicking on the button can be handled. The event handlers are selected by going to the Properties window and clicking on the yellow lightning bolt ⚡. Make sure that the item selected for events is the button1 object. To switch back to the Properties window from the event window, click on the icon just to the left of the lightning bolt. Locate the Click event (should be the first event) and then double-click on the textbox to the right to install the event handler for Click. The view will now switch to the code view and change the location of the button click software.

The software currently in view is the button1_Click function, which is called when the user clicks on the button. This procedure is illustrated in Example 7–8. To test the button, change the software in Example 7–8 to the software in Example 7–9(a). Click on the green arrow to compile, link, and execute the dialog application and click on button1 when it is running. The label on button1 will change to "Wow, Hello" if the button has been made wide enough. This is the first working application, but it does not use any assembly code. Example 7–9(a) uses the Text member property of the button1 object to change the text displayed on button1. A variant that uses a character string object (String^) appears in Example 7–9(b) to display "Wow, Hello World."

EXAMPLE 7–8

```
private: System::Void button1_Click(System::Object^ sender,
      System::EventArgs^ e)
{

}
```

EXAMPLE 7–9

```
//Version (a)

private: System::Void button1_Click(System::Object^ sender,
      System::EventArgs^ e)
{
      button1->Text = "Wow, Hello";
}

//Version (b)

private: System::Void button1_Click(System::Object^ sender,
      System::EventArgs^ e)
{
      String^ str1 = "Wow, Hello World";
      button1->Text = str1;
}
```
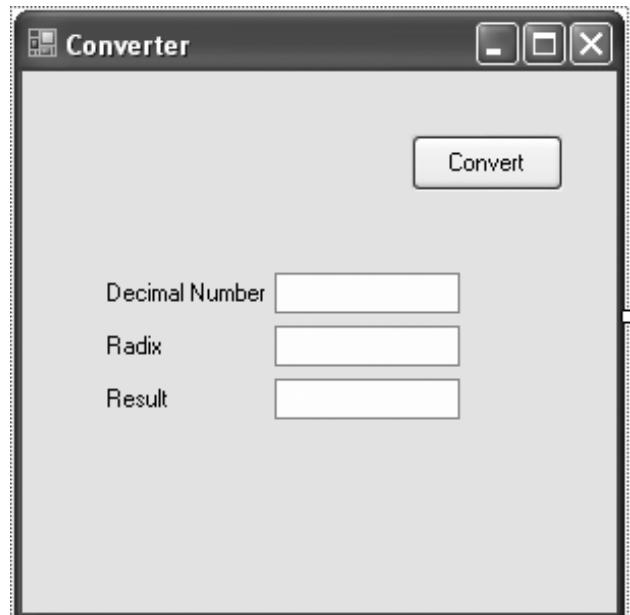
Now that a simple application has been written, we can modify it to illustrate a more compli-
cated application as shown in Figure 7–5. The caption on the button has been changed to the word
"Convert." To return to the design screen, select the tab at the top of the program window that is
labeled Form1.h[design]*. When in the Design window, change the caption on the button1 object
by clicking the button and then finding the Text property from the properties of button1 in the
Properties window. Change the Text property to "Convert." In Figure 7–5 notice that there are three
Label controls and three textbox controls in the illustration below and to the left of the Convert but-
ton. These controls are located in the toolbox. Draw them on the screen in approximately the same

FIGURE 7–5   The first
application.

place as in Figure 7–5. The labels are changed in properties for each label control. Change the text for each label as indicated.

Our goal in this example is to display any decimal number entered in the Decimal Number box as a number with any radix (number base) as selected by the number entered in the Radix box. The result appears in the Result box when the Convert button is clicked. To switch the view to the program view, click on the Form1.h tab at the top of the Design window.

To obtain the value from an edit control, use Text property to obtain a string version of the number. The problem is that in this case an integer is needed and not a string. The string must be converted to an integer. The Convert class provided in C++ performs conversion from most data types to most data types. In this case, the Convert class member function ToInt32 is used to transform the string into an integer. The difficult portion of this example is the conversion from base 10 to any number base. Example 7–10 shows how the Convert class is used to convert the string from the textbox into an integer. This will only function correctly if the number entered into textbox1 is an integer. If a letter or anything else is entered, the program will crash and display an error.

EXAMPLE 7–10

```
private: System::Void button1_Click(System::Object^ sender,
        System::EventArgs^ e)
{
        int number = Convert::ToInt32(textBox1->Text);
}
```

To handle input errors, a try-catch code block is used as illustrated in Example 7–11. The try portion tries the code and the catch statement catches any error and displays a message using the MessageBox class Show member function.

EXAMPLE 7–11

```
private: System::Void button1_Click(System::Object^ sender,
        System::EventArgs^ e)
{
        try
        {
                int number = Convert::ToInt32(textBox1->Text);
        }
        catch (...)                 // catch any error
        {
                MessageBox::Show("The input must be an integer!");
        }
}
```

The remainder of the application appears in the button1_Click function of Example 7–12. This program uses the **Horner's algorithm** to convert to any radix from 2 through 36. This conversion algorithm divides the number to be converted by the desired radix until the result is zero. After each division, the remainder is saved as a significant digit in the result and the quotient is divided again by the radix. Note that Windows does not use ASCII code, it uses Unicode so the Char (16-bit Unicode) is needed in place of the 8-bit char. Notice how the order of the remainders is placed into the result string by concatenating each digit to the left of the string. A 0x30 is added to each digit to convert to ASCII code in the example.

Horner's algorithm:

1. Divide the number by the desired radix.
2. Save the remainder and replace the number with the quotient.
3. Repeat steps 1 and 2 until the quotient is zero.

EXAMPLE 7–12

```
private: System::Void button1_Click(System::Object^ sender,
     System::EventArgs^ e)
{
     String^ result = "";
     int number;
     int radix;
     try
     {
          number = Convert::ToInt32(textBox1->Text);
          radix = Convert::ToInt32(textBox2->Text);
     }
     catch (...)                 // catch all Convert errors
     {
          MessageBox::Show("All inputs must be integers!");
     }
     if (radix < 2 || radix > 36)
     {
          MessageBox::Show("The radix must range between 2 and 36");
     }
     else
     {
          do    // conversion algorithm
          {
               char digit = number % radix;
               number /= radix;
               if (digit > 9)               // for letters
               {
                    digit += 7;             // add bias
               }
               digit += 0x30;               // convert to ASCII
               result = digit + result;
          }
          while (number != 0);
     }
     textBox3->Text = result;
}
```

Since this is an assembly language text, the Convert class is not going to be used for good reason; the function is quite large. To see just how large, you can put a breakpoint in the software to the left of a Convert function by clicking on the gray bar to the left of the line of code. A brown circle, a **breakpoint**, will appear. If you run the program, it will break (stop) at this point and enter the debugging mode so it can be viewed in assembly language form. To display the disassembled code, run the program until it breaks, and then go to the Debug menu and select Windows. In the Windows menu, near the bottom, find "Disassembly." The registers can also be displayed to step through a program in assembly language.

As you can see, if the program is debugged as described, this is a substantial amount of code that can be significantly reduced if it is rewritten using the inline assembler. Example 7–13 depicts the assembly language version of Convert::ToInt32 function. This function is considerably shorter (if debugged and viewed in the Disassemble window) and executes many times faster than the Convert in Example 7–12. This example points out the inefficiency of the code generated by a high-level language, which may not always be important, but many cases require tight and efficient code, and that can only be written in assembly language. My guess is that as a plateau is reached on processor speed, more things will be written in assembly language. In addition, the new instructions such as MMX and SSE are not available in high-level languages. They require a very good working knowledge of assembly code.

The main problem with using inline assembly code is that the code cannot be placed into a Windows-managed forms application in a managed class. In order to use the assembler, the
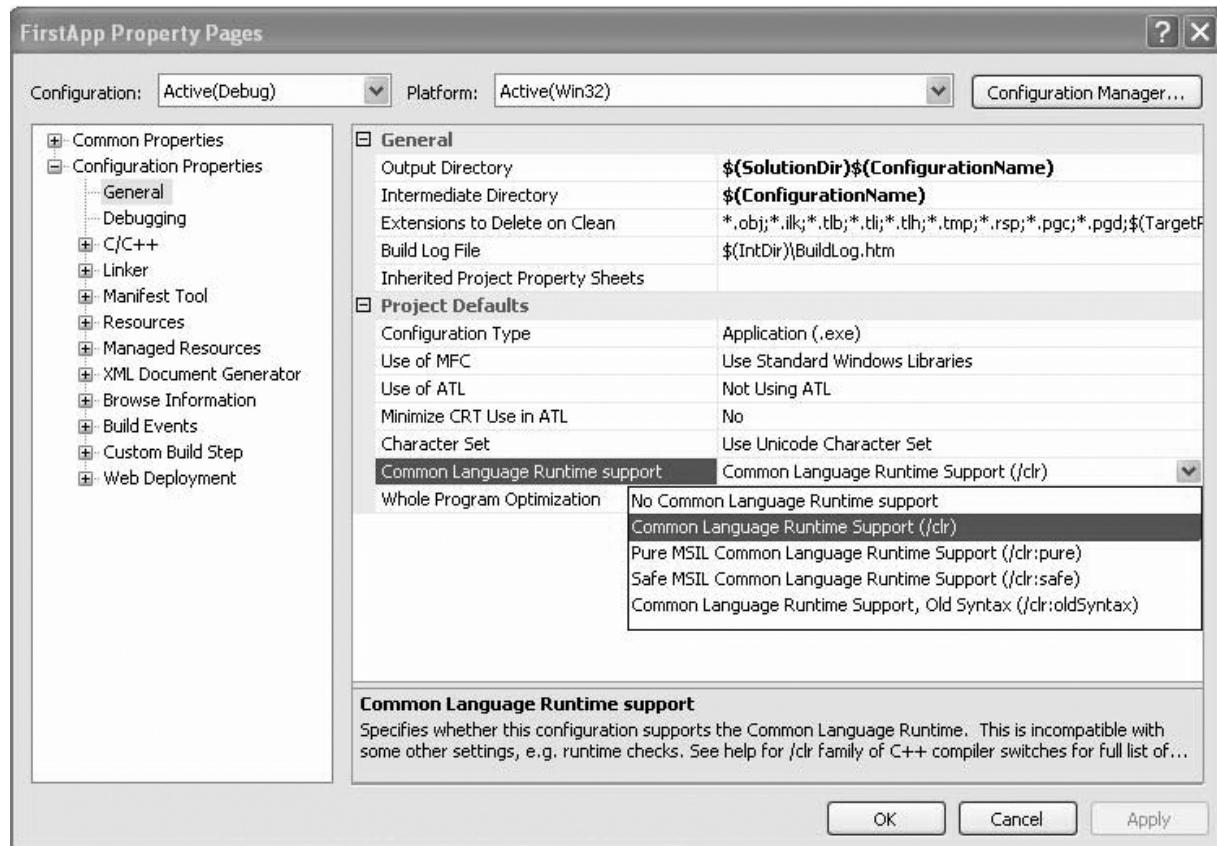
FIGURE 7–6   Changing to /clr for assembly language.

function must be placed before the managed class in order for it to compile. Therefore, in the project properties, Common Runtime Support must also be changed to /clr from the default setting of /clr:pure so it will compile successfully. (Refer to Figure 7–6 for a screen shot of how to change Common Language Runtime support to /clr.) A **managed** program runs under the virtual machine called .net and an **unmanaged** application operated in the native mode of the computer. The inline assembler generates native code for the microprocessor so it must be unmanaged and reside before the managed class in a program.

Example 7–13 illustrates how to replace part of the Horner's algorithm with assembly code in a function called Adjust. The adjust function tests the number for 9, and if it's greater than 9, it adds 0x07 and then 0x30 to convert it to ASCII, which is returned. Notice in the example that the assembly code is placed immediately following the using statements at the top of the program. This is where any assembly functions must be placed so a program can function correctly. The application starts in native mode and switches to managed mode when it encounters the managed class. By placing the assembly code before the managed class, it is available to the entire application and it executes in unmanaged or native mode.

At the end of Example 7–13 an alternative version of Adjust appears that is more efficient. The alternative version does not have a return instruction, so how can it function? What does not appear is that any assembly language function returns the value in AL for a byte, AX for a word or short, and EAX for an int. Note that the return value dictates the size of the value returned.

EXAMPLE 7–13

```
#pragma once

namespace FirstApp {

      using namespace System;
      using namespace System::ComponentModel;
      using namespace System::Collections;
      using namespace System::Windows::Forms;
      using namespace System::Data;
      using namespace System::Drawing;

             // short is for a 16-bit variable.

      short Adjust(short n)
      {
             _asm
             {
                  mov  ax,n
                  cmp  ax,9
                  jle  later
                  add  ax,7
             later:
                  add  ax,30h
                  mov  n,ax
             }
             return n;
      }

      /* as an alternative version

      short Adjust(short n)
      {
             _asm
             {
                  mov  ax,n
                  add  ax,30h
                  cmp  ax,39h
                  jle later
                  add  ax,7
             later:
             }
      }

      */
// managed class follows
```

Figure 7–14 shows the modification to the button1_click function so that Adjust is called in place of the code that appears in Example 7–12. The code used to set up the form application that appears between Examples 7–13 and 7–14 is not shown. Notice that the assembly function uses short in place of character. A short is a 16-bit number used in unmanaged mode and a Char is a 16-bit number used in managed mode to represent a Unicode character. Here a cast is used to convert to a Char because without it, the numeric values are displayed instead of ASCII code.

EXAMPLE 7–14

```
// The sole event handler in this application.

private: System::Void button1_Click(System::Object^ sender,
      System::EventArgs^ e)
{
      String^ result = "";
      try
```

```
        {
              unsigned int number = Convert::ToUInt32(textBox1->Text);
              unsigned int radix = Convert::ToUInt32(textBox2->Text);
              if (radix < 2 || radix > 36)
              {
                    MessageBox::Show("The radix must range between 2 and 36");
              }
              else
              {
                    do
                    {
                          result = (Char)Adjust(number % radix) + result;
                          number /= radix;
                    }
                    while (number != 0);
                    textBox3->Text = result;
              }
        }
        catch (...)
              // catch any error
        {
              MessageBox::Show
                    ("The decimal input must between 0 and 4294967295!");
        }
}
```

7–3   MIXED ASSEMBLY AND C++ OBJECTS

As mentioned in the prior sections, the inline assembler is limited because it cannot use MACRO sequences and the conditional program flow directives presented in Chapter 6. In some cases, it is better to develop assembly language modules that are then linked with C++ for more flexibility. This is especially true if the application is being developed by a team of programmers. This section of the chapter details the use of different objects that are linked to form a program using both assembly language and C++.

Linking Assembly Language with Visual C++

Example 7–15 illustrates a flat model procedure that will be linked to a C++ program. We denote that the assembly module is a C++ module by using the letter C after the word *flat* in the model statement. The linkage specified by the letter C is the same for the C or C++ languages. The flat model allows assembly language software to be any length up to 2G bytes. Note that the .586 switch appears before the model statement, which causes the assembler to generate code that functions in the protected 32-bit mode. The Reverse procedure, shown in Example 7–15, accepts a character string from a C++ program, reverses its order, and returns to the C++ program. Notice how this program uses conditional program flow instructions, which are not available with the inline assembler described in prior sections of this chapter. The assembly language module can have any name and it can contain more than one procedure, as long as each procedure contains a PUBLIC statement defining the name of the procedure as public. Any parameters that are transferred in the C++ program and the assembly language program are indicated with the backslash following the name of the procedure. This names the parameter for the assembly language program (it can be a different name in C++) and indicates the size of the parameter. The only thing that is not different in the C++ calling program and the assembly program is the order of the parameters. In this example, the parameter is a pointer to a character string and the result is returned as a replacement for the original string.

EXAMPLE 7–15

```
;
;External function that reverses the order of a string of characters
;
.586                                    ;select Pentium and 32-bit model
.model flat, C                          ;select flat model with C/C++ linkage
.stack 1024                             ;allocate stack space
.code                                   ;start code segment

public Reverse                          ;define Reverse as a public function

Reverse proc uses esi, \               ;define procedure
arraychar:ptr                           ;define external pointer

      mov esi,arraychar                 ;address string
      mov eax,0
      push eax                          ;indicate end of string

      .repeat                           ;push all the characters to the stack
            mov al,[esi]
            push eax
            inc esi
      .until byte ptr [esi] == 0

      mov esi,arraychar                 ;address string start

      .while eax != 0                   ;pop in reverse order
            pop  'eax
            mov  [esi],al
            inc  esi
      .endw
      Ret

Reverse     endp
End
```

Example 7–16 illustrates a C++ language program for DOS console applications that uses the Reverse assembly language procedure. The EXTERN statement is used to indicate that an external procedure called Reverse is to be used in the C++ program. The name of the procedure is case-sensitive, so make sure that it is spelled the same in both the assembly language module and the C++ language module. The EXTERN statement in Example 7–16 shows that the external assembly language procedure transfers a character string to the procedure and returns no data. If data are returned from the assembly language procedure, data are returned as a value in register EAX for bytes, words, or doublewords. If floating-point numbers are returned, they must be returned on the floating-point coprocessor stack. If a pointer is returned, it must be in EAX.

EXAMPLE 7–16

```
/* Program that reverses the order of a character string */

#include <stdio.h>
#include <conio.h>

extern "C" void Reverse(char *);

char chararray[17] = "So what is this?";

int main(int argc, char* argv[]
{
  printf ("%s \n", chararray);
  Reverse (char array);
  printf ("%s\n", chararray);
  getche();                             //wait to see result
  return 0;
}
```

Once both the C++ program and the assembly language program are written, the Visual C++ development system must be set up to link the two together. For linking and assembly, we will assume that the assembly language module is called Reverse.txt (you cannot add an .asm extension file to the file list for inclusion into a project, so just use the .txt extension and add a .txt file) and the C++ language module is called Main.cpp. Both modules are stored in the C:\PROJECT\MINE directory or some other directory of your choosing. After the modules are placed in the same project workspace, the Programmer's Workbench program is used to edit both assembly language and C++ language modules.

To set up the Visual C++ developer studio to compile, assemble, and link these files, follow these steps:

1. Start the developer studio and select New from the File menu.
   a. Choose New Project.
   b. When the Application Wizard appears, click on Visual C++ Projects.
   c. Select C++ Console Application, and name the project Mine.
   d. Then click on OK.
2. You will see the project in the Solution window at the left margin in the center. It will have a single file called Main.cpp, which is the C++ program file. Modify this to appear as in Example 7–16.
3. To add the assembly language module, right-click on the line Source Files and select Add from the menu. Choose Add New Item from the list. Scroll down the list of file types until you find Text Files and select it, then enter the file name as Reverse and click on Open. This creates the assembly module called Reverse.txt. You may enter the assembly code from Example 7–15 in this file.
4. Under the Source Files listing in the Solution Explorer, right-click on Reverse.txt and select Properties. Figure 7–7 shows what to enter in this wizard after you click on the Custom Build step. Make sure you enter the object file name (Reverse.obj) in the Outputs box and ml /c /Cx /coff Reverse.txt in the Command Line box. The Reverse assembly language file will assemble and be included in the project.
5. Assuming both Examples 7–15 and 7–16 have been entered and you have completed all steps, the program will function.

FIGURE 7–7   Using the assembler to assemble a module in Visual C++.
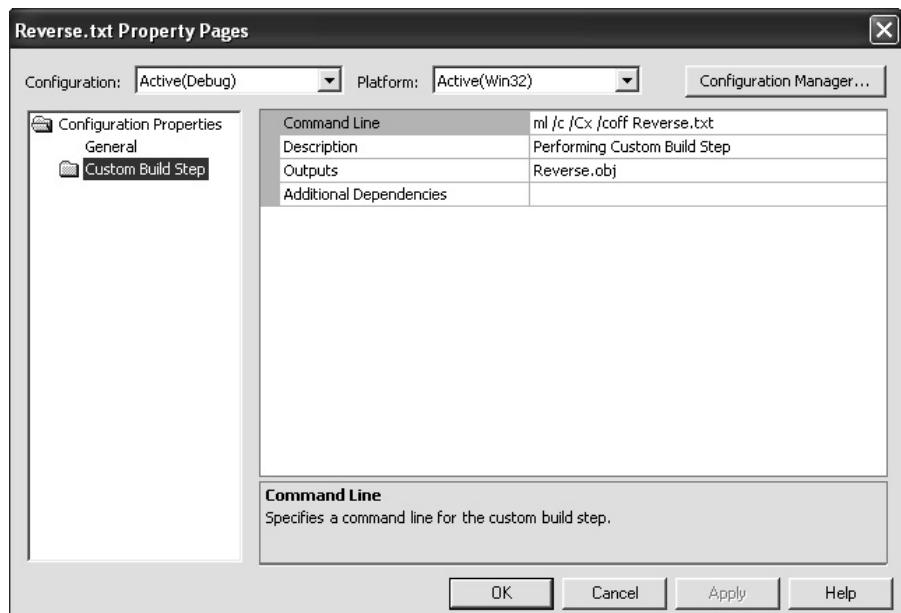
TABLE 7–2 Morse code.

A	. _	J	_ _ _	S	. . .
B	_ . . .	K	_ . _	T	_
C	_ . _ .	L	. _ . .	U	. . _
D	_ . .	M	_ _	V	. . . _
E	.	N	_ .	W	. _ _
F	. . _ .	O	_ _ _	X	_ . . _
G	_ _ .	P	. _ _ .	Y	_ . _ _
H	. . . .	Q	_ _ . _	Z	_ _ . .
I	. .	R	. _ .		

At last, you can execute the program. Click on the green arrow. You should see two lines of ASCII text data displayed. The first line is in correct forward order and the second is in reverse order. Although this is a trivial application, it does illustrate how to create and link C++ language with assembly language.

Now that we have a good understanding of interfacing assembly language with C++, we need a longer example that uses a few assembly language procedures with a C++ language program. Example 7–17 illustrates an assembly language package that includes a procedure (Scan) to test a character input against a lookup table and return a number that indicates the relative position in the table. A second procedure (Look) uses a number transferred to it and returns with a character string that represents Morse code. (The code is not important, but if you are interested, Table 7–2 lists Morse code.)

EXAMPLE 7–17

```
.586
.model flat, C
.data

table db  2,1,4,8,4,10,3,4      ;ABCD
      db  1,0,4,2,3,6,4,0        ;EFGH
      db  2,0,4,7,3,5,4,4        ;IJKL
      db  2,3,2,2,3,7,4,6        ;MNOP
      db  4,13,3,2,3,0,1,1       ;QRST
      db  3,1,4,1,3,3,4,9        ;UVWX
      db  4,11,4,12              ;YZ

.code

Public Scan
Public Look

Scan proc uses ebx,\
char:dword

      mov ebx,char
      .if bl >= 'a' && bl <= 'z'
            sub bl,20h
      .endif
      sub bl,41h
      add bl,bl
      add ebx,offset table
      mov ax,word ptr[ebx]
      ret

Scan   endp

Look   proc uses ebx ecx,\
numb:dword,\
pntr:ptr
```

```
       mov   ebx,pntr
       mov   eax,numb
       mov   ecx,0
       mov   cl,al
       .repeat
              shr ah,1
              .if carry?
                     mov byte ptr[ebx],'_'
              .else
                     mov byte ptr[ebx],'.'
              .endif
                     inc ebx
       .untilcxz
       mov byte ptr[ebx],0
       ret

Look   endp
end
```

The lookup table in Example 7–17 contains 2 bytes for each character between A and Z. For example, the code for A is a 2 for a Morse-coded character two of any combination of dashes or dots, and the 1 is the code for the letter a (. –), where the binary equivalent 01 (for two digits) is a dot followed by a dash. This lookup table is accessed by the Scan procedure to obtain the correct Morse code from the lookup table, which is returned in AX as a parameter to the C++ language call. The remaining assembly code is mundane.

Example 7–18 lists the C++ program, which calls the two procedures listed in Example 7–17. This software is simple to understand, so we do not explain it.

EXAMPLE 7–18

```cpp
// Moorse.cpp : Defines the entry point for the console application.

#include <iostream>
using namespace std;

extern "C" int Scan(int);
extern "C" void Look(int, char *);

int main(int argc, char* argv[])
{
       int a = 0;
       char chararray[] = "This, is the trick!\n";
       char chararray1[10];

       while ( chararray[a] != '\n' )
       {
              if ( chararray[a] < 'A' || chararray[a] > 'z' )
                     cout << chararray[a] << '\n';
              else
              {
                     Look ( Scan ( chararray[a] ), chararray1 );
                     cout << chararray[a] << " = " << chararray1 << '\n';
              }
              a++;
       }
       cout << "Type enter to quit!";
       cin.get();
       return 0;
}
```

Although the examples presented here are for console applications, the same method of instructing Visual Studio to assemble and link an assembly language module is also used for Visual applications for Windows. The main difference is that Windows applications do not use printf or cout. The next chapter explains how library files can also be used with Visual C++ and also gives many more programming examples.

Adding New Assembly Language Instructions to C/C++ Programs

From time to time, as new microprocessors are introduced by Intel, new assembly language instructions are also introduced. These new instructions cannot be used in C++ unless you develop a macro for C++ to include them in the program. An example is the CPUID assembly language instruction. This will function in an _asm block within C++ because the inline assembler does not recognize it. Another group of newer instructions includes the MMX and SEC instructions. These are also recognized, but in order to illustrate how a new instruction is added that is not in the assembler, we show the technique. To use any new instructions, first look up the machine language code from Appendix B or from Intel's website at www.intel.com. For example, the machine code for the CPUID instruction is 0F A2. This 2-byte instruction can be defined as a C++ macro, as illustrated in Example 7–19. To use the new macro in a C++ program, all we need to type is CPUID. The _emit macro stores the byte that follows it in the program.

EXAMPLE 7–19

```
#define CPUID _asm _emit 0x0f _asm _emit 0xa2
```

7–4	SUMMARY

1. The inline assembler is used to insert short, limited assembly language sequences into a C++ program. The main limitation of the inline assembler is that it cannot use macro sequences or conditional program flow instructions.
2. Two versions of C++ language are available. One is designed for 16-bit DOS console applications and the other for 32-bit Windows applications. The type chosen for an application depends on the environment, but in most cases programmers today use Windows and the 32-bit Visual Express version.
3. The 16-bit assembly language applications use the DOS INT 21H commands to access devices in the system. The 32-bit assembly language applications cannot efficiently or easily access the DOS INT 21H function calls even though many are available.
4. The most flexible and often-used method of interfacing assembly language in a C++ program is through separate assembly language modules. The only difference is that these separate assembly language modules must be defined by using the C directive following the .model statement to define the module linkage as C/C++ compatible.
5. The PUBLIC statement is used in an assembly language module to indicate that the procedure name is public and available to use with another module. External parameters are defined in an assembly language module by using the name of the procedure in the PROC statement. Parameters are returned through the EAX register to the calling C/C++ procedure from the assembly language procedure.
6. Assembly language modules are declared external to the C++ program by using the extern directive. If the extern directive is followed by the letter C, the directive is used in a C/C++ language program.
7. When using Visual Studio, we can instruct it to assemble an assembly language module by clicking on Properties for the module and adding the assembler language program (ml /c /Cx / coff Filename.txt) and output file as an object file (Filename.obj) in the Custom Build step for the module.
8. Assembly language modules can contain many procedures, but can never contain programs using the .startup directive.

7–5   QUESTIONS AND PROBLEMS

1. Does the inline assembler support assembly language macro sequences?
2. Can a byte be defined in the inline assembler by using the DB directive?
3. How are labels defined in the inline assembler?
4. Which registers can be used in assembly language (either inline or linked modules) without saving?
5. What register is used to return integer data from assembly language to the C++ language caller?
6. What register is used to return floating-point data from assembler language to the C++ language caller?
7. Is it possible to use the .if statement in the inline assembler?
8. In Example 7–3, explain how the mov dl,strings[si] instruction accesses strings data.
9. In Example 7–3, why was the SI register pushed and popped?
10. Notice in Example 7–5 that no C++ libraries (#include) are used. Do you think that compiled code for this program is smaller than a program to accomplish the same task in C++ language? Why?
11. What is the main difference between the 16-bit and 32-bit versions of C/C++ when using the inline assembler?
12. Can the INT 21H instruction, used to access DOS functions, be used in a program using the 32-bit version of the C/C++ compiler? Explain your answer.
13. What is the #include <conio.h> C/C++ library used for in a program?
14. Write a short C/C++ program that uses the _getche() function to read a key and the _putch() function to display the key. The program must end if an '@' is typed.
15. Would an embedded application that is not written for the PC ever use the conio.h library?
16. In Example 7–7, what is the purpose of the sequence of instructions _punch(10); followed by _punch(13);?
17. In Example 7–7, explain how a number is displayed in any number base.
18. Which is more flexible in its application, the inline assembler or assembly language modules that are linked to C++?
19. What is the purpose of a PUBLIC statement in an assembly code module?
20. How is an assembly code module prepared for use with C++ language?
21. In a C++ language program, the extern void GetIt(int); statement indicates what about function GetIt?
22. How is a 16-bit word of data defined in C++?
23. What is a control in a C++ Visual program and where is it obtained?
24. What is an event in a C++ Visual program and what is an event handler?
25. In Example 7–13, what size parameter is short?
26. Can the edit screen of C++ Visual Studio be used to enter and edit an assembly language programming module.
27. How are external procedures that are written in assembly language indicated to a C++ program?
28. Show how the RDTSC instruction (opcode is 0F 31) could be added to a C++ program using the _emit macro.
29. In Example 7–17, explain what data type is used by Scan.
30. Write a short assembly language module to be used with C++ that rotates a number three places to the left. Call your procedure RotateLeft3 and assume the number is an 8-bit char (byte in assembly).
31. Repeat question 30, but write the same function in C++ without the assembler.

32. Write a short assembly language module that receives a parameter (byte-sized) and returns a byte-sized result to a caller. Your procedure must take this byte and convert it into an upper-case letter. If an uppercase letter or anything else appears, the byte should not be modified.
33. How is a CLR Visual C++ Express application executed from Visual Studio?
34. What are properties in a Visual C++ application?
35. What is an ActiveX control or object?
36. Show how a single instruction assembly language instruction, such as inc ptr, is inserted into a Visual C++ program.

CHAPTER 8

Programming the Microprocessor

INTRODUCTION

This chapter develops programs and programming techniques using the inline assembler program from Visual C++ Express. The Visual C++ inline assembler has already been explained and demonstrated in prior chapters, but there are still more features to learn at this point.

Some programming techniques explained in this chapter include assembly language modules, keyboard and display manipulation, program modules, library files, using the mouse, using timers, and other important programming techniques. As an introduction to programming, this chapter provides a wealth of background on valuable programming techniques so that programs can be easily developed for the personal computer by using the inline assembler as a springboard for Visual C++ Express applications created for Windows.

CHAPTER OBJECTIVES

Upon completion of this chapter, you will be able to:

1. Use the MASM assembler and linker program to create programs that contain more than one module.
2. Explain the use of EXTRN and PUBLIC as they apply to modular programming.
3. Set up a library file that contains commonly used subroutines and learn how to use the DUMPBIN program.
4. Write and use MACRO and ENDM to develop macro sequences used with linear programming in modules that link to C++ code.
5. Show how both sequential and random access files are developed for use in a system.
6. Develop programs using event handlers to perform keyboard and display tasks.
7. Use conditional assembly language statements in programs.
8. Use the mouse in program examples.

8–1 MODULAR PROGRAMMING

Many programs are too large to be developed by one person. This means that programs are routinely developed by teams of programmers. The linker program is provided with Visual Studio so that programming modules can be linked together into a complete program. Linking is also available from the command prompt provided by Windows. This section of the text describes the linker, the linking task, library files, EXTRN, and PUBLIC as they apply to program modules and modular programming.

The Assembler and Linker

The **assembler program** converts a symbolic **source module** (file) into a hexadecimal **object file**. It is even a part of Visual Studio, located in the C:\Program Files\Microsoft Visual Studio .NET 2003\Vc7\bin folder. We have seen many examples of symbolic source files, written in assembly language, in prior chapters. Example 8–1 shows how the assembler dialog that appears as a source module named NEW.ASM is assembled. Note that this dialog is used with version 6.15 at the DOS command line. The version that comes with Visual C will not work for 16-bit DOS programs. If a 16-bit assembler and linker are needed, they can be obtained in the Windows Driver Development Kit (DDK). Whenever you create a source file, it should have the extension of ASM, but as we learned in the last chapter, that is not always possible. Source files are created by using NotePad or almost any other word processor or editor capable of generating an ASCII file.

EXAMPLE 8–1

```
C:\masm611\BIN>ml new.asm

Microsoft (R) Macro Assembler Version 6.11
Copyright (C) Microsoft Corp 1981-1993. All rights reserved.

 Assembling: new.asm

Microsoft (R) Segmented Executable Linker Version 5.60.220 Sep 9 1994
Copyright (C) Microsoft Corp 1984-1993. All rights reserved.

Object Modules [.obj]: new.obj
Run File [new.exe]: "new.exe"
List File [nul.map]: NUL
Libraries [.lib]:
Definitions File [nul.def]:
```

The assembler program (ML) requires the source file name following ML. In Example 8–1, the /Fl switch is used to create a listing file named NEW.LST. Although this is optional, it is recommended so that the output of the assembler can be viewed for troubleshooting problems. The source listing file (.LST) contains the assembled version of the source file and its hexadecimal machine language equivalent. The cross-reference file (.CRF), which is not generated in this example, lists all labels and pertinent information required for cross-referencing. An object file is also generated by ML as an input to the linker program. In many cases we only need to generate an object file, which is accomplished by using the /c switch.

The **linker program**, which executes as the second part of ML, reads the object files that are created by the assembler program and links them together into a single execution file. An **execution file** is created with the file name extension EXE. Execution files are selected by typing the file name at the DOS prompt (C:\). An example execution file is FROG.EXE, which is executed by typing FROG at the command prompt.

If a file is short enough (less than 64K bytes long), it can be converted from an execution file to a **command file** (.COM). The command file is slightly different from an execution file in

that the program must be originated at location 0100H before it can execute. This means that the program must be no larger than 64K–100H in length. The ML program generates a command file if the tiny model is used with a starting address of 100H. Command files are only used with DOS or if a true binary version (for a EPROM/FLASH burner) is needed. The main advantage of a command file is that it loads off the disk into the computer much more quickly than an execution file. It also requires less disk storage space than the equivalent execution file.

Example 8–2 shows the linker program protocol when it is used to link the files NEW, WHAT, and DONUT. The linker also links library files (LIBS) so procedures, located with LIBS, can be used with the linked execution file. To invoke the linker, type LINK at the command prompt, as illustrated in Example 8–2. Note that before files are linked, they must first be assembled and they must be error-free. ML not only links the files, but it also assembles them prior to linking.

EXAMPLE 8–2

```
C:\masm611\BIN>ml new.asm what.asm donut.asm
Microsoft (R) Macro Assembler Version 6.11
Copyright (C) Microsoft Corp 1981-1993. All rights reserved.

 Assembling: new.asm
 Assembling: what.asm
 Assembling: donut.asm

Microsoft (R) Segmented Executable Linker Version 5.60.220 Sep 9 1994
Copyright (C) Microsoft Corp 1984-1993. All rights reserved.

Object Modules [.obj]: new.obj+
Object Modules [.obj]: "what.obj"+
Object Modules [.obj]: "donut.obj"/t
Run File [new.com]: "new.com"
List File [nul.map]: NUL
Libraries [.lib]:
Definitions File [nul.def]:
```

In this example, after typing ML, the linker program asks for the "Object Modules," which are created by the assembler. In this example, we have three object modules: NEW, WHAT, and DONUT. If more than one object file exists, type the main program file first (NEW, in this example), followed by any other supporting modules.

Library files are entered after the file name and after the switch /LINK. In this example, library files were not entered. To use a library called NUMB.LIB while assembling a program called NEW.ASM, type ML NEW.ASM /LINK NUMB.LIB.

In the Windows environment you cannot link a program—you can only assemble a program. You must use Visual Studio to link the program files during the build. You can assemble a file or files and generate objects for use with Visual C++. Example 8–3 illustrates how a module is compiled, but not linked with ML. The /c switch (lowercase c) tells the assembler to compile and generate object files, **/Cx** preserves the case of all functions and variables, and /coff generates a **common object file format** output for the object files used in a 32-bit environment.

EXAMPLE 8–3

```
C:\Program Files\Microsoft Visual Studio .NET 2003\Vc7\bin>ml /c /Cx /coff new.asm
Microsoft (R) Macro Assembler Version 7.10.3077
Copyright (C) Microsoft Corporation. All rights reserved.

 Assembling: new.asm
```

PUBLIC and EXTRN

The PUBLIC and EXTRN directives are very important to modular programming because they allow communications between modules. We use PUBLIC to declare that labels of code, data, or entire segments are available to other program modules. EXTRN (external) declares that labels are external to a module. Without these statements, modules could not be linked together to create a program by using modular programming techniques. They might link, but one module would not be able to communicate to another.

The PUBLIC directive is placed in the opcode field of an assembly language statement to define a label as public, so that the label can be used (seen by) by other modules. The label declared as public can be a jump address, a data address, or an entire segment. Example 8–4 shows the PUBLIC statement used to define some labels and make them public to other modules in a program fragment. When segments are made public, they are combined with other public segments that contain data with the same segment name.

EXAMPLE 8–4

```
                        .model  flat,  c
                        .data

                             public Data1        ;declare Data1 and Data2 public
                             public Data2

                         Data1   db     100 dup(?)
0000 0064[
        00
          ]
0064 0064[               Data2   db     100 dup(?)
        00
          ]
                        .code
                        .startup

                             public Read         ;declare Read public

                        Read    proc   far
0006 B4 06                      mov ah,6
```

The EXTRN statement appears in both data and code segments to define labels as external to the segment. If data are defined as external, their sizes must be defined as BYTE, WORD, or DWORD. If a jump or call address is external, it must be defined as NEAR or FAR. Example 8–5 shows how the external statement is used to indicate that several labels are external to the program listed. Notice in this example that any external address or data is defined with the letter E in the hexadecimal assembled listing. It is assumed that Example 8–4 and Example 8–5 are linked together.

EXAMPLE 8–5

```
                        .model flat, c
                        .data
                             extrn Data1:byte
                             extrn Data2:byte
                             extrn Data3:word
                             extrn Data4:dword
                        .code
                             extrn Read:far
                        .startup
0005 Bf 0000 E               mov    dx,offset Data1
0008 B9 000A                 mov    cx,10
000B                    Start:
```

```
000B 9A 0000 ---- E        call    Read
0010 AA                    stosb
0011 E2 F8                 loop    Start
                        .exit
                        End
```

Libraries

Library files are collections of procedures that are used by many different programs. These procedures are assembled and compiled into a library file by the LIB program that accompanies the MASM assembler program. Libraries allow common procedures to be collected into one place so they can be used by many different applications. You may have noticed when setting up Visual C++ to build the assembly language modules in Chapter 7 that many library files were in the link list used by Visual C++. The library file (FILENAME.LIB) is invoked when a program is linked with the linker program.

Why bother with library files? A library file is a good place to store a collection of related procedures. When the library file is linked with a program, only the procedures required by that program are removed from the library file and added to the program. If any amount of assembly language programming is to be accomplished efficiently, a good set of library files is essential and saves many hours in recoding common functions.

Creating a Library File.   A library file is created with the LIB command, which executes the LIB.EXE program that is supplied with Visual Studio. A library file is a collection of assembled .OBJ files that contains procedures or tasks written in assembly language or any other language. Example 8–6 shows two separate functions (UpperCase and LowerCase) included in a module that is written for Windows, which will be used to structure a library file. Please notice that the name of the procedure must be declared PUBLIC in a library file and does not necessarily need to match the file name, although it does in this example. A variable is transferred to each file, so the EXTRN statement also appears in each procedure to gain access to an external variable. Example 8–7 shows the C++ protocols that are required to use the functions in this library file in a C++ program, provided the library is linked to the program.

EXAMPLE 8–6

```
.586
.model flat,c
.code
      public UpperCase
      public LowerCase
UpperCase proc ,\
      Data1:byte
      mov   al,Data1
      .if al >= 'a' && al <= 'z'
            sub al,20h
      .endif
      ret
UpperCase endp

LowerCase proc ,\
      Data2:byte
      mov   al,Data2
      .if al >= 'A' && al <= 'Z'
            add al,20h
      .endif
      ret
LowerCase endp
End
```

EXAMPLE 8–7

```
extern "C" char UpperCase(char);
extern "C" char LowerCase(char);
```

The LIB program begins with the copyright message from Microsoft, followed by the prompt *Library name.* The library name chosen is case for the CASE.LIB file. Because this is a new file, the library program must be prompted with the object file name. You must first assemble CASE.ASM with ML. The actual LIB command is listed in Example 8–8. Notice that the LIB program is invoked with the object name following it on the command line.

EXAMPLE 8–8

```
C:\Program Files\Microsoft Visual Studio .NET 2003\Vc7\bin>lib case.obj
Microsoft (R) Library Manager Version 7.10.3077
Copyright (C) Microsoft Corporation. All rights reserved.
```

A utility program called DUMPBIN.EXE is provided to display the contents of the library or any other file. Example 8–9 shows the outcome of a binary dump using the /all switch to show the library module CASE.LIB and all its components. Near the top of this listing are the public names for _UpperCase and _LowerCase. The Raw Data #1 section contains the actual hexadecimal-coded instructions for the two procedures.

EXAMPLE 8–9

```
C:\Program Files\Microsoft Visual Studio .NET 2003\Vc7\bin>dumpbin /all case.lib

Microsoft (R) COFF/PE Dumper Version 7.10.3077
Copyright (C) Microsoft Corporation.  All rights reserved.

Dump of file case.lib

File Type: LIBRARY

Archive member name at 8: /
401D4A83 time/date Sun Feb 01 13:50:43 2004
        uid
        gid
      0 mode
     22 size
correct header end

    2 public symbols

      C8 _LowerCase
      C8 _UpperCase

Archive member name at 66: /
401D4A83 time/date Sun Feb 01 13:50:43 2004
        uid
        gid
      0 mode
     26 size
correct header end

    1 offsets

        1       C8

    2 public symbols

        1 _LowerCase
        1 _UpperCase
```

```
Archive member name at C8: case.obj/
401D43A6 time/date Sun Feb 01 13:21:26 2004
         uid
         gid
  100666 mode
     228 size
correct header end
FILE HEADER VALUES
             14C machine (x86)
               3 number of sections
        401D43A6 time date stamp Sun Feb 01 13:21:26 2004
             124 file pointer to symbol table
               D number of symbols
               0 size of optional header
               0 characteristics

SECTION HEADER #1
   .text name
       0 physical address
       0 virtual address
      24 size of raw data
      8C file pointer to raw data (0000008C to 000000AF)
       0 file pointer to relocation table
       0 file pointer to line numbers
       0 number of relocations
       0 number of line numbers
60500020 flags
         Code
         16 byte align
         Execute Read

RAW DATA #1
  00000000: 55 8B EC 8A 45 08 3C 61 72 06 3C 7A 77 02 2C 20  U.ì.E.<ar.<zw.,
  00000010: C9 C3 55 8B EC 8A 45 08 3C 41 72 06 3C 5A 77 02  ÉAU.ì.E.<Ar.<Zw.
  00000020: 04 20 C9 C3                                      . ÉA

SECTION HEADER #2
   .data name
      24 physical address
       0 virtual address
       0 size of raw data
       0 file pointer to raw data
       0 file pointer to relocation table
       0 file pointer to line numbers
       0 number of relocations
       0 number of line numbers
C0500040 flags
         Initialized Data
         16 byte align
         Read Write

SECTION HEADER #3
.debug$S name
      24 physical address
       0 virtual address
      74 size of raw data
      B0 file pointer to raw data (000000B0 to 00000123)
       0 file pointer to relocation table
       0 file pointer to line numbers
       0 number of relocations
       0 number of line numbers
42100040 flags
         Initialized Data
         Discardable
         1 byte align
         Read Only
```

```
RAW DATA #3
  00000000: 04 00 00 00 F1 00 00 00 00 00 00 00 30 00 01 11  ....ñ.......0...
  00000010: 00 00 00 00 43 3A 5C 50 52 4F 47 52 41 7E 31 5C  ....C:\PROGRA~1\
  00000020: 4D 49 43 52 4F 53 7E 31 2E 4E 45 54 5C 56 63 37  MICROS~1.NET\Vc7
  00000030: 5C 62 69 6E 5C 63 61 73 65 2E 6F 62 6A 00 34 00  \bin\case.obj.4.
  00000040: 16 11 03 02 00 00 05 00 00 00 00 00 00 00 07 00  ................
  00000050: 0A 00 05 0C 4D 69 63 72 6F 73 6F 66 74 20 28 52  ....Microsoft (R
  00000060: 29 20 4D 61 63 72 6F 20 41 73 73 65 6D 62 6C 65  ) Macro Assemble
  00000070: 72 00 00 00                                      r...

COFF SYMBOL TABLE
000 00000000 DEBUG  notype        Filename    | .file
    C:\PROGRA~1\MICROS~1.NET\Vc7\bin\case.asm
004 000F0C05 ABS    notype        Static      | @comp.id
005 00000000 SECT1  notype        Static      | .text
    Section length   24, #relocs    0, #linenums    0, checksum      0
007 00000000 SECT2  notype        Static      | .data
    Section length    0, #relocs    0, #linenums    0, checksum      0
009 00000000 SECT3  notype        Static      | .debug$S
    Section length   74, #relocs    0, #linenums    0, checksum      0
00B 00000000 SECT1  notype ()     External    | _UpperCase
00C 00000012 SECT1  notype ()     External    | _LowerCase

String Table Size = 0x1A bytes

  Summary

        0 .data
       74 .debug$S
       24 .text
```

Once the library file is linked to your program file, only the library procedures actually used by your program are placed in the execution file. Don't forget to use the extern "C" statement in the C++ program to use a function from a library file.

In Visual C++ Express, a library is created by selecting the Class Library choice at the Create menu. This feature creates a DLL (**dynamic link library**) file that can be included in any C++ application. The DLL can contain C++ code or assembly code. To include the DLL in any program, under Project, select the "add reference" choice and browse to the DLL file. Once the DLL is added, place an #include at the start of the class where the DLL is to be used in a program.

Macros

A **macro** is a group of instructions that perform one task, just as a procedure performs one task. The difference is that a procedure is accessed via a CALL instruction, whereas a macro, and all the instructions defined in the macro, is inserted in the program at the point of usage. Creating a macro is very similar to creating a new opcode, which is actually a sequence of instructions, in this case, that can be used in the program. You type the name of the macro and any parameters associated with it, and the assembler then inserts them into the program. Macro sequences execute faster than procedures because there is no CALL or RET instruction to execute. The instructions of the macro are placed in your program by the assembler at the point where they are invoked. Be aware that macros will not function using the inline assembler; they only function in external assembly language modules.

The MACRO and ENDM directives delineate a macro sequence. The first statement of a macro is the MACRO instruction, which contains the name of the macro and any parameters associated with it. An example is MOVE MACRO A,B, which defines the macro name as MOVE. This new pseudo opcode uses two parameters: A and B. The last statement of a macro is the ENDM instruction, which is placed on a line by itself. Never place a label in front of the ENDM statement. If a label appears before ENDM, the macro will not assemble.

Example 8–10 shows how a macro is created and used in a program. The first six lines of code define the macro. This macro moves the word-sized contents of memory location B into word-sized memory location A. After the macro is defined in the example, it is used twice. The macro is expanded by the assembler in this example, so that you can see how it assembles to generate the moves. Any hexadecimal machine language statement followed by a number (1, in this example) is a macro expansion statement. The expansion statements are not typed in the source program; they are generated by the assembler (if .LISTALL is included in the program) to show that the assembler has inserted them into the program. Notice that the comment in the macro is preceded with ;; instead of ; as is customary. Macro sequences must always be defined before they are used in a program, so they generally appear at the top of the code segment.

EXAMPLE 8–10

```
                      MOVE    MACRO A,B
                              PUSH  AX
                              MOV   AX,B
                              MOV   A,AX
                              POP   AX
                              ENDM

                              MOVE VAR1,VAR2     ;;move VAR2 into VAR1

0000  50            1         PUSH  AX
0001  A1 0002 R     1         MOV   AX,VAR2
0004  A3 0000 R     1         MOV   VAR1,AX
0007  58            1         POP   AX

                              MOVE  VAR3,VAR4    ;;move VAR4 into VAR3
0008  50            1         PUSH  AX
0009  A1 0006 R     1         MOV   AX,VAR4
000C  A3 0004 R     1         MOV   VAR3,AX
000F  58            1         POP   AX
```

Local Variables in a Macro.   Sometimes, macros contain local variables. A **local variable** is one that appears in the macro, but is not available outside the macro. To define a local variable, we use the LOCAL directive. Example 8–11 shows how a local variable, used as a jump address, appears in a macro definition. If this jump address is not defined as local, the assembler will flag it with errors on the second and subsequent attempts to use the macro.

EXAMPLE 8–11

```
                      FILL    MACRO WHERE, HOW_MANY    ;;fill memory
                              LOCAL FILL1
                              PUSH  SI
                              PUSH  CX
                              MOV   SI,OFFSET WHERE
                              MOV   CX,HOW_MANY
                              MOV   AL,0
                      FILL1:  MOV   [SI],AL
                              INC   SI
                              LOOP  FILL1
                              POP   CX
                              POP   SI
                              ENDM

                              FILL  MES1,5

                      1       LOCAL FILL1
0014  56            1         PUSH  SI
0015  51            1         PUSH  CX
0016  BE 0000 R     1         MOV   SI,OFFSET MES1
0019  B9 0005       1         MOV   CX,5
001C  B0 00         1         MOV   AL,0
0029  88 04         1   ??0000:MOV   [SI],AL
```

```
002B   46                    1              INC    SI
002C   E2 FB                 1              LOOP   ??0000
002E   59                    1              POP    CX
002F   5E                    1              POP    SI


                                            FILL   MES2,10


                             1              LOCAL FILL1
0030   56                    1              PUSH   SI
0031   51                    1              PUSH   CX
0032   BE 0014 R             1              MOV    SI,OFFSET MES2
0035   B9 000A               1              MOV    CX,10
0038   B0 00                 1              MOV    AL,0
003A   88 04                 1     ??0001:MOV      [SI],AL
003C   46                    1              INC    SI
003D   E2 FB                 1              LOOP   ??0001
003F   59                    1              POP    CX
0040   5E                    1              POP    SI
                                   .EXIT
```

Example 8–11 shows a FILL macro that stores any number (parameter HOW_MANY) of 00H into the memory location addressed by parameter WHERE. Notice how the address FILL1 is treated when the macros are expanded. The assembler uses labels that start with ?? to designate them are assembler-generated labels.

The LOCAL directive must always be used on the line immediately following the MACRO statement or an error occurs. The LOCAL statement may have up to 35 labels, all separated with commas.

Placing MACRO Definitions in Their Own Module.   Macro definitions can be placed in the program file as shown, or they can be placed in their own macro module. A file can be created that contains only macros to be included with other program files. We use the INCLUDE directive to indicate that a program file will include a module that contains external macro definitions. Although this is not a library file, for all practical purposes it functions as a library of macro sequences.

When macro sequences are placed in a file (often with the extension INC or MAC), they do not contain PUBLIC statements as does a library. If a file called MACRO.MAC contains macro sequences, the INCLUDE statement is placed in the program file as INCLUDE C:\ASSM\MACRO.MAC. Notice that the macro file is on drive C, subdirectory ASSM in this example. The INCLUDE statement includes these macros, just as if you had typed them into the file. No EXTRN statement is needed to access the macro statements that have been included. Programs may contain both macro include files and library files.

8–2   USING THE KEYBOARD AND VIDEO DISPLAY

Today, there are few programs that don't use the keyboard and video display. This section of the text explains how to use the keyboard and video display connected to the IBM PC or any compatible computer running under Windows.

Reading the Keyboard

The keyboard of the personal computer is read by many different objects available to Visual C++. Data read from the keyboard are either in ASCII-coded or in extended ASCII-coded form. They are then either stored in 8-bit ASCII form or in 16-bit Unicode form. As mentioned in an earlier chapter, Unicode contains ASCII code in the codes 0000H–00FFH. The remaining codes are used for foreign language character sets. Do not use cin or getch to read keys in Visual C++ as we do in a DOS C++ console application; in place of cin or getch we use controls in Visual C++ that accomplish the same task.

TABLE 8–1   The keyboard scanning and extended ASCII codes as returned from the keyboard.
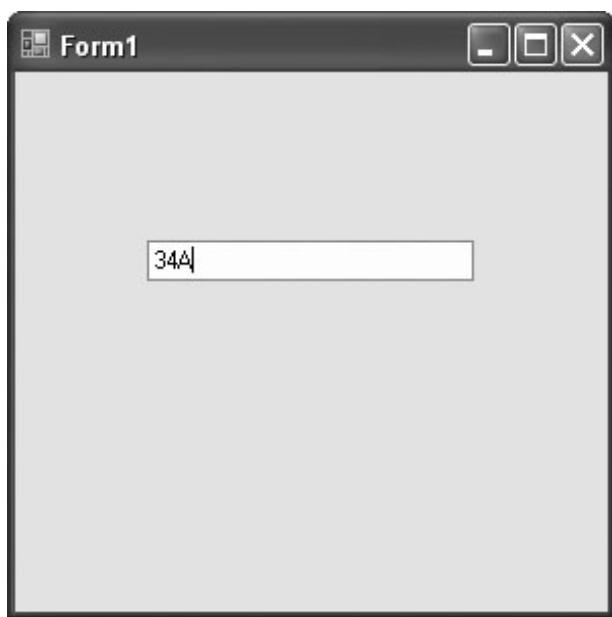
Key	Scan Code	Extended ASCII code with….			
		Nothing	Shift	Control	Alternate
Esc	01				01
1	02				78
2	03			03	79
3	04				7A
4	05				7B
5	06				7C
6	07				7D
7	08				7E
8	09				7F
9	0A				80
0	0B				81
−	0C				82
+	0D				83
Bksp	0E				0E
Tab	0F		0F	94	A5
Q	10				10
W	11				11
E	12				12
R	13				13
T	14				14
Y	15				15
U	16				16
I	17				17
O	18				18
P	19				19
[	1A				1A
]	1B				1B
Enter	1C				1C
Enter	1C				A6
Lctrl	1D				
Rctrl	1D				
A	1E				1E
S	1F				1F
D	20				20
F	21				21
G	22				22
H	23				23
J	24				24
K	25				25
L	26				26
;	27				27
'	28				28
,	29				29
Lshft	2A				
\	2B				

TABLE 8–1   *(continued)*

Key	Scan Code	Nothing	Shift	Control	Alternate
		Extended ASCII code with….			
Z	2C				2C
X	2D				2D
C	2E				2E
V	2F				2F
B	30				30
N	31				31
M	32				32
,	33				33
.	34				34
/	35				35
Gray/	35			95	A4
Rshft	36				
PrtSc	E0 2A E0 37				
L alt	38				
R alt	38				
Space	39				
Caps	3A				
F1	3B	3B	54	5E	68
F2	3C	3C	55	5F	69
F3	3D	3D	56	60	6A
F4	3E	3E	57	61	6B
F5	3F	3F	58	62	6C
F6	40	40	59	63	6D
F7	41	41	5A	64	6E
F8	42	42	5B	65	6F
F9	43	43	5C	66	70
F10	44	44	5D	67	71
F11	57	85	87	89	8B
F12	58	86	88	8A	8C
Num	45				
Scroll	46				
Home	E0 47	47	47	77	97
Up	48	48	48	8D	98
Pgup	E0 49	49	49	84	99
Gray-	4A				
Left	4B	4B	4B	73	9B
Center	4C				
Right	4D	4D	4D	74	9D
Gray +	4E				
End	E0 4F	4F	4F	75	9F
Down	E0 50	50	50	91	A0
Pgdn	E0 51	51	51	76	A1
Ins	E0 52	52	52	92	A2
Del	E0 53	53	53	93	A3
Pause	E0 10 45				

FIGURE 8–1 Using the textbox with filtering.

The ASCII-coded data appear as outlined in Table 1–8 in Section 1–4. The extended character set of Table 1–9 applies to printed or displayed data only, and not to keyboard data. Notice that the ASCII codes in Table 1–8 correspond to most of the keys on the keyboard. Also available through the keyboard are extended ASCII-coded keyboard data. Table 8–1 lists most of the extended ASCII codes obtained with various keys and key combinations. Notice that most keys on the keyboard have alternative key codes. Each function key has four sets of codes selected by the function key alone, the Shift-function key combination, the alternate-function key combination, and the Control-function key combination.

Creating a Visual C++ Express application that contains a simple textbox gives a better understanding of reading a key in Windows. Figure 8–1 shows such an application written as a forms-based application. Recall that to create a forms-based application:

1. Start Visual C++ Express.
2. Click on Create: Project.
3. Select a CLR Windows Forms Application, then give it a name and click on OK.

Once the new forms-based application is created, select the textbox control from the toolbox and draw it on the screen of the dialog box, as illustrated in Figure 8–1.

Setting Focus. The first thing that should be added to the application is a set focus to the textbox control. When focus is set, the cursor moves to the object, in this case the textbox. Focus is set to a control by using **textBox1->Focus()**, which in our case is because the textbox control is named textBox1. This statement is placed in the Form1_Load function, which must be installed by double-clicking on a blank area of the form. The Form1_Load function can also be installed by clicking on the yellow lightning bolt and selecting Load and then adding it by double-clicking on the blank textbox to its right. The application will now set focus to the textbox1 control when started. This means that the blinking cursor appears inside the textbox control.

When the application is executed and keys are typed into the textbox control, the program reads the keyboard and displays each character as it is typed. In some cases this may be undesirable and may require some filtering. One such case is if the program requires that the user enter only hexadecimal data. In order to intercept keystrokes as they are typed, the event handlers

KeyDown and KeyPress are used for the textbox. The KeyDown event handler is called when the key is pressed down, which is followed by a call to the KeyPress event handler. To insert these functions into the application for the textbox control, click on the textbox and then select the Properties window. Next find the yellow lightning bolt and click on it, and install KeyDown and KeyPress events handlers for the textbox1 control.

To illustrate filtering, this application uses the KeyDown function to look at each keyboard character that is typed before the program uses the keystroke. This allows the characters to be modified. Here the program only allows the numbers 0 through 9 and the letters A through F to be typed from the keyboard. If a lowercase letter is typed, it is converted into uppercase. If any other key is typed, it is ignored.

To accomplish filtering, use the **KeyEventArgs^** class argument e, which is passed to the KeyDown functon as illustrated in Example 8–12. In this example, C++ is used to accomplish the task of filtering the keyboard entry into the textbox control. The variable keyHandled is used to indicate whether or not the key is handled by the filtering. If keyHandled is set to false, the key has not been handled and it will appear in the textbox. Likewise, if keyhandled is set to true, the key has been handled and will *not* appear in the textbox. The condition of keyHandled is passed to Windows in the KeyPress event that also appears in Example 8–12. Note that Keys::D0 through Keys::D9 are the number keys on the QWERTY keyboard and Keys::NumPad0 through Keys::NumPad9 are on the numeric keypad. A D8 with the shift equal to false is the eight key and a D8 with shift equal to true is an asterisk key.

EXAMPLE 8–12

```cpp
private: System::Void Form1_Load(System::Object^ sender,
      System::EventArgs^ e)
{
      textBox1->Focus();           // set Focus to textbox1
}

bool keyHandled;

private: System::Void textBox1_KeyDown(System::Object^ sender,
      System::Windows::Forms::KeyEventArgs^ e)
{     // this is called first
      keyHandled = true;
      if (e->KeyCode >= Keys::NumPad0 && e->KeyCode <= Keys::NumPad9 ||
            e->KeyCode >= Keys::D0 && e->KeyCode <= Keys::D9 &&
            e->Shift == false ||
            e->KeyCode >= Keys::A && e->KeyCode <= Keys::F ||
            e->KeyCode == Keys::Back)
      {
            keyHandled = false;
      }
}

private: System::Void textBox1_KeyPress(System::Object^ sender,
      System::Windows::Forms::KeyPressEventArgs^ e)
{     // this is called second
      if (e->KeyChar >= 'a' && e->KeyChar <= 'f')
      {
            e->KeyChar -= 32;            // make uppercase
      }
      e->Handled = keyHandled;
}
```

The if statement in KeyPress event tests the e->KeyCode value for the letters a, b, c d, e, and f, which can be either uppercase or lowercase. The KeyDown event tests for the numbers 0–9 on both the keyboard and the number pad. Also the backspace key is tested. If any of these are typed, keyHandled is set to false to indicate that these keys are not handled by the KeyDown function. The KeyPress event determines if the letter a–f is typed and converts it into uppercase by subtracting 32. The 32 is the bias between uppercase and lowercase letters in the ASCII code.

Next, a return with e->Handled set into true or false occurs. A return true causes Windows to dispose of the keystroke. In this case, if a number is typed, or the letters a through f or A through F, the keystroke is passed to Windows framework by the normal return false at the end of the KeyPress function. This filters the keystrokes so only A–F or 0–9 appears in the edit box.

Example 8–12 is repeated using the inline assembler to accomplish the same task in Example 8–13. Here a function, called filter, returns a true or false that is passed to keyHandled in the KeyDown function. In this example, C++ seems to require less typing than assembly language, but it is important to be able to visualize both forms. Don't forget to change the project property, Common Language Runtime Support to /CLR, so this will function correctly (see Chapter 7). Notice that KeyValue is used with the assembly version to pass a char to the Filter function. Also note that an integer return value of 0 is false and 1 is true.

EXAMPLE 8–13

```
// Placed at the top of the program following the uses statements

int Filter(char key)
{
      int retval;          // 0 = false, 1 = true
      _asm
      {
            mov eax,1
            cmp key,8              ; backspace
            jb good
            cmp key,30h
            jb bad
            cmp key,39h
            jbe good
            cmp key,41h
            jb bad
            cmp key,46h
            jbe good
            cmp key,61h
            jb bad
            cmp key,66h
            jbe good
good:       dec eax
bad:        mov retval,eax
      }
      return retval;
}


private: System::Void Form1_Load(System::Object^ sender,
      System::EventArgs^ e)
{
      textBox1->Focus();
}

bool keyHandled;

// new version of textbox1_KeyDown

private: System::Void textBox1_KeyDown(System::Object^ sender,
      System::Windows::Forms::KeyEventArgs^ e)
{
      keyHandled = Filter(e->KeyValue);
}
private: System::Void textBox1_KeyPress(System::Object^ sender,
      System::Windows::Forms::KeyPressEventArgs^ e)
{
      if (e->KeyChar >= 'a' && e->KeyChar <= 'f')
      {
            e->KeyChar -= 32;
      }
      e->Handled = keyHandled;
}
```
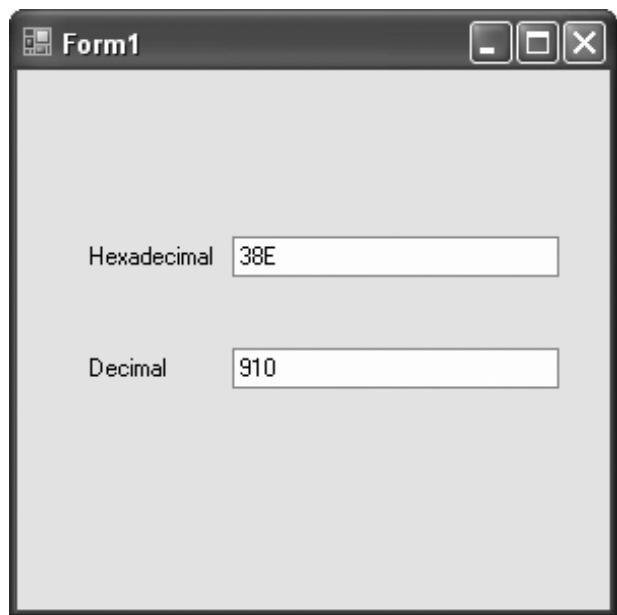
FIGURE 8–2 Hexadecimal to decimal conversion.

If this code is added to the application and executed, the only keys that will appear in the textbox control are 0–9 and A–F. Any amount of filtering can be done in a likewise manner in the Visual C++ Express environment. The properties of the textbox control include character casing, which could have been set to uppercase to shorten the filtering task, but here software accomplished the uppercase feature.

Using the Video Display

As with the keyboard, in Visual C++ objects are used to display information. The textbox control can be used to either read data or display data as can most objects. Modify the application presented in Figure 8–1 so it contains an additional textbox control as shown in Figure 8–2. Notice that a few label controls have been added to the form to identify the contents of the textbox controls. In this new application the keyboard data is still read into textbox control textBox1, but when the Enter key is typed, a decimal version of the data entered into textBox1 appears in textBox2—the second textbox control. Make sure that the second control is named **textBox2** to maintain compatibility with the software presented here.

To cause the program to react to the Enter key, ASCII code 13 (0DH or 0x0d), modify the KeyPress function of Example 8–13 as shown in Example 8–14. Notice how the Enter key is detected using an else if. Once the Enter key is detected, the contents of textBox1 are converted to decimal for display in textBox2, as shown in Example 8–15.

EXAMPLE 8–14

```
private: System::Void textBox1_KeyPress(System::Object^ sender,
      System::Windows::Forms::KeyPressEventArgs^ e)
{
      if (e->KeyChar >= 'a' && e->KeyChar <= 'f')
      {
            e->KeyChar -= 32;
      }
      else if (e->KeyChar == 13)
      {
```

```
                // software to display the decimal version in textBox2

                keyHandled = true;
        }
        e->Handled = keyHandled;
}
```

A slight problem arises with textbox data; the data entered into a textbox control is accessed as a string, but not as a hexadecimal number. In this example program (see Example 8–15), a function called Converts changes the hexadecimal character string into a number. The program now has two functions that contain assembly code.

EXAMPLE 8–15

```
// placed after the using statements at the top of the program

int Filter(char key)
{
        int retval;
        _asm
        {
                mov eax,1
                cmp key,8               ; backspace
                je  good
                cmp key,30h
                jb  bad
                cmp key,39h
                jbe good
                cmp key,41h
                jb  bad
                cmp key,46h
                jbe good
                cmp key,61h
                jb  bad
                cmp key,66h
                jbe good
good:           dec al
bad:            mov retval,eax
        }
        return retval;
}

int Converts(int number, short digit)
{
        _asm
        {
                mov eax,number
                shl eax,4
                mov dx,digit
                sub dx,30h
                cmp dx,9
                jbe later
                sub dx,7
later:          or  al,dl
                mov number,eax
        }
        return number;
}
private: System::Void Form1_Load(System::Object^ sender,
        System::EventArgs^ e)
{
        textBox1->Focus();
}

bool keyHandled;
```

```
private: System::Void textBox1_KeyDown(System::Object^ sender,
      System::Windows::Forms::KeyEventArgs^ e)
{
      keyHandled = Filter(e->KeyValue);
}
private: System::Void textBox1_KeyPress(System::Object^ sender,
      System::Windows::Forms::KeyPressEventArgs^ e)
{
      if (e->KeyChar >= 'a' && e->KeyChar <= 'f')
      {
            e->KeyChar -= 32;
      }
      else if (e->KeyChar == 13)
      {
            int number = 0;
            for (int a = 0; a < textBox1->Text->Length; a++)
            {
                  number = Converts(number, textBox1->Text[a]);
            }
            textBox2->Text = Convert::ToString(number);
            keyHandled = true;
      }
      e->Handled = keyHandled;
}
```

Example 8–15 shows the completed application. When the Enter key is pressed, the program obtains the character string from textBox1 and converts it, a character at a time, into an integer using the Converts function. Once an integer is created, it is changed into a character string using the Convert class and its member function ToString for display in textBox2.

The assembly language in the Converts function converts from ASCII to binary by subtracting 30h from each digit. This action converts ASCII numbers (0–9) 30H through 39H to the binary numbers 0 through 9. It does not convert the letters 41H through 46H (A through F) to binary because the result is 11H through 16H and not 0AH through 0FH. To adjust the values obtained for the letters, use a cmp (compare) instruction to detect 11H through 16H and then subtract an additional 7 to convert from 11H through 16H to 0AH to 0FH. Once the ASCII digit is in binary form, the integer at number is shifted left four binary places and the binary version of the ASCII digit is ORed to it to accumulate the converted digits in temp1.

Once the hexadecimal number from textBox1 is converted to binary in variable number, it is displayed in textBox2 using the ToString function in the Convert class. As before, a return true informs the Windows interface that the Enter key has been processed by the program. If a return false is used in KeyPress for the Enter key, the computer generates an error beep.

Using a Timer in a Program

Timers are important in programming. A timer is programmed to fire or trigger after a known amount of time, which is given in milliseconds. The allowable range for the timer is from 1 to 2G milliseconds. This allows a timer to be programmed for just about any time needed. If programmed with 1000, the timer fires in 1 second, and if programmed with 60000, the timer fires in 1 minute and so forth. A timer may be used a single time or multiple times and as many timers as needed (up to 2 billion) may appear in a program. The timer is found in the toolbox at the Design window.

To illustrate the timer and its use in a program, a design appears in Figure 8–3 that allows a user to demonstrate a shift or a rotate on a binary number. The shift or rotate is animated with the timer in this program. The design contains two label controls and two button controls plus the timer control. Add all five of these controls to the application. The timer does not appear on the form, but in an area near the bottom of the design screen. Notice that a few properties of the form are changed so the icon is not shown and the Text property of the form is changed to Shift/Rotate instead of Form1. Also the Text properties of the labels and buttons are changed as indicated.

FIGURE 8–3 The Shift/
Rotate application design
screen.

Once the form appears as shown in Figure 8–3, add event handler functions (*yellow light-ning bolt*) for the two command buttons (Click) and for the timer (Tick). To add an event handler, click on the button or timer, go to the Properties window at the right of the screen, and click on the yellow lightning bolt and select the event. The program contains three handlers, two for button clicks and one for a timer tick.

At the design screen go to the properties of the timer and set the interval to 500, which is ½ second. Do not enable the timer. The timer is enabled in the software when a button is clicked to either rotate a number or shift a number in ½-second animation steps.

Example 8–16 illustrates the software added to the application in the three handlers required to implement the application for Figure 8–3. The software for the two button click handlers is nearly identical except for the Boolean variable shift. The two statements in each place text onto the labels. If the shift button is pressed, "Shifted" is displayed, and if the rotate button is pressed, "Rotated" is displayed on label1. The second label has the test number 00011001 displayed. The Boolean variable shift is set to true for the shift button and false for the rotate button. In both button handlers, count is set to 8 to shift or rotate the number 8 places. Finally, the last statement in each button handler starts the timer. As an alternative, the enabled member could be set to true to start the timer. Once the timer is started or enabled, it fires in ½ second and calls the Tick handler for timer1 where all the work is accomplished in this program.

The first statement in the timer tick function sets the digit string to zero. This is the number that shifts into the right end of label2. For a shift, it will remain a zero and for a rotate it depends on the leftmost digit of the number in label2. If shift is false (rotate) and the leftmost digit of label2 is 1, the digit is changed to 1. After the if statement, the number is shifted or rotated and placed on label2. If the count reaches 0, after 8 iterations in 4 seconds, the timer is disabled by setting the enabled member to false to stop the animation.

EXAMPLE 8–16

```
bool shift;
int count;

private: System::Void button1_Click(System::Object^ sender,
      System::EventArgs^ e)
{
```

```
        label1->Text = "Shifted";
        label2->Text = "00011001";
        shift = true;
        count = 8;
        timer1->Start();
}
private: System::Void button2_Click(System::Object^ sender,
        System::EventArgs^ e)
{
        label1->Text = "Rotated";
        label2->Text = "00011001";
        shift = false;
        count = 8;
        timer1->Start();
}

private: System::Void timer1_Tick(System::Object^ sender,
        System::EventArgs^ e)
{
        String^ digit = "0";
        if (shift == false && label2->Text[0] == '1')
        {
                digit = "1";
        }
        label2->Text = label2->Text->Remove(0,1) + digit;
        if (−count == 0)
        {
                timer1->Enabled = false;
        }
}
```

This application is embellished by adding a pair of radio buttons to select right and left for the direction of the shift or rotate. Label2 can also be replaced with a textbox so the user could enter any binary number (with appropriate filtering) and shift or rotate it left or right.

The program does not use any assembly language, but if a breakpoint is inserted in the timer function, the assembly code can be viewed in Visual C++ Express. When the program Breaks, go to the Debug menu and select Windows. Next, select the Disassembly window to show the assembly language for the timer tick function.

The Mouse

The mouse pointing device, as well as a track ball, is accessed from within the framework of Visual C++ Express. Like many other devices under the control of Windows, the mouse can have message handlers installed in an application so that mouse movement and other functions can be used in a program. As we saw in prior examples, message handlers (*event handlers*) are installed in an application in the Properties section by clicking on the icon to the right of the lightning bolt. The mouse handlers are listed in Table 8–2.

For an illustration of how to use the mouse, refer to the example application presented in Figure 8–4. This example shows the mouse pointer coordinates as the pointer is moved within

TABLE 8–2 Mouse message handlers.

Handler	Trigger
MouseDown	Mouse button down
MouseEnter	Mouse pointer enters the control
MouseHover	Mouse not moved for awhile
MouseLeave	Mouse pointer leaves the control
MouseMove	Mouse moved
MouseUp	Mouse button released

FIGURE 8–4 Displaying
the mouse coordinates.

the dialog application. Although the software listing (see Example 8–17) does not use any
assembly language, it does illustrate how to obtain and display the position of the mouse pointer.
Notice how the MouseEventArgs^ are used to obtain the location of the mouse pointer using the
X and Y coordinates.

EXAMPLE 8–17

```
private: System::Void Form1_MouseMove(System::Object^ sender,
      System::Windows::Forms::MouseEventArgs^ e)
{
      label1->Text = "X-coordinate = " + Convert::ToString(e->Location.X);
      label2->Text = "Y-coordinate = " + Convert::ToString(e->Location.Y);
}
```

Example 8–17 illustrates the only part of the application that is modified to display the
mouse coordinates. The MouseMove function is installed when the MouseMove event handler is
installed in the program. This application uses two labels to display the mouse coordinates.
These two objects are named label1 and label2 for the application. The MouseMove function
returns the position of the mouse pointer in the Location data structure as members X and Y. This
example uses the Convert class to convert the integer returned as the mouse point X or Y into an
ASCII character string for placement on a label.

The mouse pointer does not track through the two labels in the application. In order to have
the application display the coordinates in the labels as well as the form, two additional
MouseMove handlers must be installed for the labels. Example 8–18 shows the addition of two

more MouseMove functions and biases added to the X and Y coordinates so the mouse tracks through the labels. Where are the bias numbers obtained? The biases are in the Location properties of each label where the label position is given as X, Y. The numbers in the application will depend on where the labels are placed on the form and could be obtained by using label1->Location.X and so forth.

EXAMPLE 8–18

```
private: System::Void Form1_MouseMove(System::Object^ sender,
      System::Windows::Forms::MouseEventArgs^ 'e)
{
      label1->Text = "X-coordinate = " + Convert::ToString(e->Location.X);
      label2->Text = "Y-coordinate = " + Convert::ToString(e->Location.Y);
}
private: System::Void label1_MouseMove(System::Object^ sender,
      System::Windows::Forms::MouseEventArgs^ e)
{
      label1->Text = "X-coordinate = " + Convert::ToString(e->Location.X+159);
      label2->Text = "Y-coordinate = " + Convert::ToString(e->Location.Y+232);
}
private: System::Void label2_MouseMove(System::Object^ sender,
      System::Windows::Forms::MouseEventArgs^ e)
{
      label1->Text = "X-coordinate = " + Convert::ToString(e->Location.X+159);
      label2->Text = "Y-coordinate = " + Convert::ToString(e->Location.Y+246);
}
```

Install a mouse handler for the MouseDown event. Modify the application by adding the MouseDown event handler as illustrated in Example 8–19. The function causes the left button to change the color of the labels to red when clicked and the right button changes the color of the labels to blue when pushed. The color codes used with most functions in Visual C++ are found in the Color class for most common colors. The application tests for the left and right mouse buttons using the Button member of the MouseEventArgs object as shown. (Microsoft chose the name mouses for the MouseButtons enumerator.)

EXAMPLE 8–19

```
private: System::Void Form1_MouseDown(System::Object^ sender,
      System::Windows::Forms::MouseEventArgs^ e)
{
      if (e->Button == ::mouses::MouseButtons::Left)
      {
            label1->ForeColor = Color::Red;
            label2->ForeColor = Color::Red;
      }
      else if (e->Button == ::mouses::MouseButtons::Right)
      {
            label1->ForeColor = Color::Blue;
            label2->ForeColor = Color::Blue;
      }
}
```

8–3   DATA CONVERSIONS

In computer systems, data are seldom in the correct form. One main task of the system is to convert data from one form to another. This section of the chapter describes conversions between binary and ASCII data. Binary data are removed from a register or memory and converted to ASCII for the video display. In many cases, ASCII data are converted to binary as they are typed on the keyboard. We also explain converting between ASCII and hexadecimal data.

Converting from Binary to ASCII

Conversion from binary to ASCII is accomplished in three ways: (1) by the AAM instruction if the number is less than 100 (provided the 64-bit extensions are not used for the conversion), (2) by a series of decimal divisions (divide by 10), or (3) by using the C++ Convert class function ToString. Techniques 1 and 2 are presented in this section.

The AAM instruction converts the value in AX into a two-digit unpacked BCD number in AX. If the number in AX is 0062H (98 decimal) before AAM executes, AX contains 0908H after AAM executes. This is not ASCII code, but it is converted to ASCII code by adding 3030H to AX. Example 8–20 illustrates a program that uses the procedure that processes the binary value in AL (0–99) and displays it on the video screen as a decimal number. The procedure blanks a leading zero, which occurs for the numbers 0–9, with an ASCII space code. This example program displays the number 74 (**testdata**) on the video screen. To implement this program, create a forms-based application in Visual C++ and place a single label called label1 on the form. The number 74 will appear if the assembly language function in Example 8–20 is placed at the top of the program after the last using statement and the project is changed to a /CLR program. The call to the assembly language function is placed in the Load event handler for the form.

EXAMPLE 8–20

```
// place at top of program
// will not function in 64-bit mode
void ConvertAam(char number, char* data)
{
        _asm
        {
                mov ebx,data           ;pointer to ebx
                mov al,number          ;get test data
                mov ah,0               ;clear AH
                aam                    ;convert to BCD
                add ah,20h
                cmp al,20h             ;test for leading zero
                je D1                  ;if leading zero
                add ah,10h             ;convert to ASCII
D1:
                mov [ebx], ah
                add al,30h             ;convert to ASCII
                mov [ebx+1], al
        }
}


private: System::Void Form1_Load(System::Object^ sender,
        System::EventArgs^ e)
{
        char temp[2];                          // place for result
        ConvertAam(74, temp);
        Char a = temp[0];
        Char b = temp[1];
        label1->Text = Convert::ToString(a) + Convert::ToString(b);
}
```

The reason that AAM converts any number between 0 and 99 to a two-digit unpacked BCD number is because it divides AX by 10. The result is left in AX so AH contains the quotient and AL the remainder. This same scheme of dividing by 10 can be expanded to convert any whole number of any number system (if the divide-by number is changed) from binary to an ASCII-coded character string that can be displayed on the video screen. For example, if AX is divided by 8 instead of 10, the number is displayed in octal.

The algorithm (called **Horner's algorithm**) for converting from binary to decimal ASCII code is:

1. Divide by 10, then save the remainder on the stack as a significant BCD digit.
2. Repeat step 1 until the quotient is a 0.
3. Retrieve each remainder and add 30H to convert to ASCII before displaying or printing.

Example 8–21 shows how the unsigned 32-bit number is converted to ASCII and displayed on the video screen. Here, we divide EAX by 10 (for decimal) and save the remainder on the stack after each division for later conversion to ASCII. After all the digits have been converted, the result is displayed on the video screen by removing the remainders from the stack and converting them to ASCII code. This program also blanks any leading zeros that occur. As mentioned, any number base can be used by changing the **radix** variable in this example. Again, to implement this example create a forms application with the /CLR option and a single label called label1. If the number base is greater than 10, letters are used for the representation of characters beyond 9. The software functions from base 2 to base 36.

EXAMPLE 8–21

```
void Converts(int number, int radix, char* data)
{
      _asm
      {
            mov    ebx,data         ;initialize pointer
            push   radix
            mov    eax, number      ;get test data
L1:
            mov    edx,0            ;clear edx
            div    radix           ;divide by base
            push   edx             ;save remainder
            cmp    eax,0
            jnz    L1              ;repeat until 0
L2:
            pop    edx             ;get remainder
            cmp    edx,radix
            je     L4              ;if finished
            add    dl,30h          ;convert to ASCII
            cmp    dl,39h
            jbe    L3
            add    dl,7
L3:
            mov    [ebx],dl        ;save digit
            inc    ebx             ;point to next
            jmp    l2              ;repeat until done
L4:
            mov    byte ptr[ebx],0 ;save null in string
      }
}


private: System::Void Form1_Load(System::Object^ sender,
      System::EventArgs^ e)
{
      char temp[32];                     // place for result
      Converts(7423, 10, temp);
      String^ a = "";
      int count = 0;
      while (temp[count] != 0)           // convert to string
      {
            Char b = temp[count++];
            a += b;
      }
      label1->Text = a;
}
```

Converting from ASCII to Binary

Conversions from ASCII to binary usually start with keyboard entry. If a single key is typed, the conversion occurs when 30H is subtracted from the number. If more than one key is typed, conversion from ASCII to binary still requires 30H to be subtracted, but there is one additional step. After subtracting 30H, the number is added to the result after the prior result is first multiplied by 10.

The algorithm for converting from ASCII to binary is:

1. Begin with a binary result of 0.
2. Subtract 30H from the character to convert it to BCD.
3. Multiply the result by 10, and then add the new BCD digit.
4. Repeat steps 2 and 3 for each character of the number.

Example 8–22 illustrates a program that implements this algorithm. Here, the binary number is displayed from variable temp on label1 using the Convert class to convert it to a string. Each time this program executes, it reads a number from the char variable array numb and converts it to binary for display on the label.

EXAMPLE 8–22

```
int ConvertAscii(char* data)
{
      int number = 0;
      _asm
      {
            mov  ebx,data    ;intialize pointer
            mov  ecx,0
B1:
            mov  cl,[ebx]    ;get digit
            inc  ebx         ;address next digit
            cmp  cl,0        ;if null found
            je   B2
            sub  cl,30h      ;convert from ASCII to BCD
            mov  eax,10      ;x10
            mul  number
            add  eax,ecx     ;add digit
            mov  number,eax  ;save result
            jmp  B1
B2:
      }
      return number;
}


private: System::Void Form1_Load(System::Object^ sender,
      System::EventArgs^ e)
{
      char temp[] = "2316";                    // string
      int number = ConvertAscii(temp);
      label1->Text = Convert::ToString(number);
}
```

Displaying and Reading Hexadecimal Data

Hexadecimal data are easier to read from the keyboard and display than decimal data. These types of data are not used at the application level, but at the system level. System-level data are often hexadecimal, and must either be displayed in hexadecimal form or read from the keyboard as hexadecimal data.

Reading Hexadecimal Data. Hexadecimal data appear as 0 to 9 and A to F. The ASCII codes obtained from the keyboard for hexadecimal data are 30H to 39H for the numbers 0 through 9, and 41H to 46H (A–F) or 61H to 66H (a–f) for the letters. To be useful, a program that reads hexadecimal data must be able to accept both lowercase and uppercase letters as well as numbers.

Example 8–23 shows two functions: One (Conv) converts the contents of an unsigned char from ASCII code to a single hexadecimal digit, and the other (Readh) converts a String with up to eight hexadecimal digits into a numeric value that is returned as a 32-bit unsigned integer. This example illustrates a balanced mixture of C++ and assembly language to perform the conversion.

EXAMPLE 8–23

```
unsigned char Conv(unsigned char temp)
{
      _asm
      {
            cmp   temp,'9'
            jbe   Conv2            ;if 0 - 9
            cmp   temp,'a'
            jb    Conv1            ;if A - F
            sub   temp,20h         ;to uppercase
Conv1:
            sub   temp,7
Conv2:
            sub   temp,30h
      }
      return temp;
}



private: System::UInt32 ReadH(String^ temp)
{
      unsigned int numb = 0;
      for ( int a = 0; a < temp->Length; a++ )
      {
            numb <= 4;
            numb += Conv(temp[a]);
      }
      return numb;
}



private: System::Void Form1_Load(System::Object^ sender,
      System::EventArgs^ e)
{
      unsigned int temp = ReadH("2AB4");
      label1->Text = Convert::ToString(temp);      // display in decimal
}
```

Displaying Hexadecimal Data. To display hexadecimal data, a number must be divided into 2-, 4-, or 8-bit sections that are converted into hexadecimal digits. Conversion is accomplished by adding 30H to the numbers 0 to 9 or 37H to the letters A to F for each section.

A function (Disph) stores a string of the contents of the unsigned integer parameter passed to the function. This function converts the unsigned into a two-, four-, or eight-digit character string as selected by parameter size. The function is listed in Example 8–24. Disph(number, 2) converts an unsigned integer number into a two-digit hexadecimal String, where Disph(number, 4) converts it to a four-digit hexadecimal string and Disph(number, 8) converts to an eight-digit hexadecimal character string.

EXAMPLE 8–24

```
void Disph(unsigned int number, unsigned int size, char* temp)
{
        int a;
        number <i<= ( 8 - size ) * 4;    //adjust position
        for (a = 0; a < size; a++)
        {
                char temp1;
                _asm
                {
                        rol number, 4;
                        mov al,byte ptr number
                        and al,0fh              ;make 0 - f
                        add al,30h              ;convert to ASCII
                        cmp al,39h
                        jbe Disph1
                        add al,7
Disph1:
                        mov temp1,al
                }
                temp[a] = temp1;               //add digit to string
        }
        temp[a] = 0;                           // null string end
}
private: System::Void Form1_Load(System::Object^ sender,
        System::EventArgs^ e)
{
        char temp[9];
        Disph(1000,4,temp);
        String^ a = "";
        int count = 0;
        while (temp[count] != 0)              // convert to string
        {
                Char b = temp[count++];
                a += b;
        }
        label1->Text = a;
}
```

Using Lookup Tables for Data Conversions

Lookup tables are often used to convert data from one form to another. A lookup table is formed in the memory as a list of data that is referenced by a procedure to perform conversions. In many lookup tables, the **XLAT instruction** is often used to look up data in a table, provided that the table contains 8-bit-wide data and its length is less than or equal to 256 bytes.

Converting from BCD to Seven-Segment Code.   One simple application that uses a lookup table is BCD to seven-segment code conversion. Example 8–25 illustrates a lookup table that contains the seven-segment codes for the numbers 0 to 9. These codes are used with the seven-segment display pictured in Figure 8–5. This seven-segment display uses active high (logic 1) inputs to light a segment. The lookup table code (array temp1) is arranged so that the **a** segment is in bit position 0 and the **g** segment is in bit position 6. Bit position 7 is 0 in this example, but it can be used for displaying a decimal point, if required.

EXAMPLE 8–25
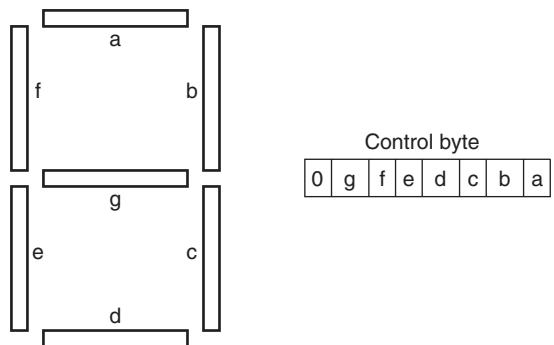
```
unsigned char LookUp(unsigned char temp)
{
        unsigned char temp1[] = {0x3f, 6, 0x5b, 0x4f, 0x66,
                                 0x6d, 0x7d, 7, 0x7f, 0x6f};
        _asm
        {
                lea ebx,temp1
                mov al,temp
```

FIGURE 8–5 The seven-segment display.


```
            xlat
            mov temp,al
      }
      return temp;
}
```

The LookUp function, which performs the conversion, contains only a few instructions and assumes that the temp parameter contains the BCD digit (0–9) to be converted to seven-segment code that is returned as an unsigned char. The first instruction addresses the lookup table by loading its address into EBX, and the others perform the conversion and return the seven-segment code as an unsigned char. Here the temp1 array is indexed by the BCD passed to the function in temp.

Using a Lookup Table to Access ASCII Data.   Some programming techniques require that numeric codes be converted to ASCII character strings. For example, suppose that you need to display the days of the week for a calendar program. Because the number of ASCII characters in each day is different, some type of lookup table must be used to reference the ASCII-coded days of the week.

The program in Example 8–26 shows a table, formed as an array, which references ASCII-coded character strings. Each character string contains an ASCII-coded day of the week. The table contains references to each day of the week. The function that accesses the day of the week uses the day parameter, with the numbers 0 to 6 to refer to Sunday through Saturday. If day contains a 2 when this function is invoked, the word *Tuesday* is displayed on the video screen. Please note that this function does not use any assembly code, since we are merely accessing an element in an array using the day of the week as an index. It is shown so additional uses for arrays can be presented, because they may have application in programs used with embedded microprocessors.
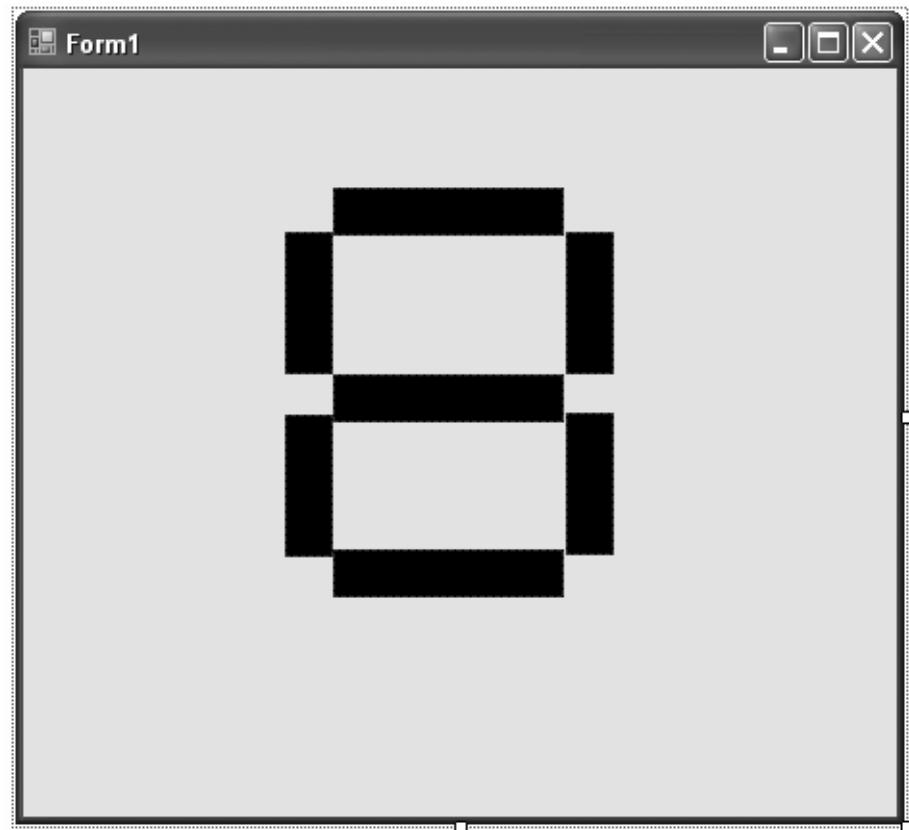
EXAMPLE 8–26

```
private: System::String^ GetDay(unsigned char day)
{
      array<String^>^ temp =
      {
            "Sunday",
            "Monday",
            "Tuesday",
            "Wednesday",
            "Thurday",
            "Friday",
            "Saturday",
      };
      return temp[day];
}
```

FIGURE 8–6 A seven-segment display.

An Example Program Using a Lookup Table

Figure 8–6 shows the screen of a dialog application called Display that displays the a seven-segment-style character on the screen for each numeric key typed on the keyboard. As we learned in prior examples, the keyboard can be intercepted in a Visual C++ program using the KeyDown and KeyPress handler functions, which is exactly what the program does to obtain the key from the keyboard. Next the code typed is filtered so only 0–9 are accepted and a lookup table is used to access the seven-segment code for display.

The display digit is drawn using panel control objects. The horizontal bars are drawn using dimensions 120×25 and the vertical bars are drawn using dimensions 25×75. The dimensions of an object appear in the extreme lower right corner of the resource screen in Visual Studio. Make sure that you add the panels in the same order as the display; that is, add label a first, followed by b, and so on, just as in the seven-segment display of Figure 8–5. Use panel1 through panel7 for the variable names of the panels in this application and don't forget to select a background color of black.

Add the function listed in Example 8–27 called Clear to the program to clear the display. This is used to clear the digit from the screen when the program first executes and also before a new digit is displayed. Notice that the Visible property of the panels is used to hide the digit. An alternate method changes the color of the panel.

EXAMPLE 8–27

```
private: System::Void Clear()
{
        panel1->Visible = false;
```

```
            panel2->Visible = false;
            panel3->Visible = false;
            panel4->Visible = false;
            panel5->Visible = false;
            panel6->Visible = false;
            panel7->Visible = false;
}
```

Once a key is typed, the KeyDown function (see Example 8–28) filters the keystroke and converts the keystroke into seven-segment code using the lookup table. After converting to seven-segment code, the ShowDigit function is called to show the digit on the screen. The ShowDigit function tests each bit of the seven-segment code and changes the visibility of each panel to display a digit. This program does not use any assembly code for its operation.

EXAMPLE 8–28

```
private: System::Void Clear()
{
        panel1->Visible = false;
        panel2->Visible = false;
        panel3->Visible = false;
        panel4->Visible = false;
        panel5->Visible = false;
        panel6->Visible = false;
        panel7->Visible = false;
}

private: System::Void Form1_KeyDown(System::Object^ sender,
        System::Windows::Forms::KeyEventArgs^ e)
{
        char lookup[] = {0x3f, 6, 0x5b, 0x4f, 0x66, 0x6d, 0x7d, 7, 0x7f, 0x6f};
        if (e->KeyCode >= Keys::D0 && e->KeyCode <= Keys::D9)
        {
                ShowDigit(lookup[e->KeyValue - 0x30]);  //display the digit
        }
}

private: System::Void ShowDigit(unsigned char code)
{
        Clear();
        if (( code & 1 ) == 1)            //test a segment
              panel1->Visible = true;
        if (( code & 2 ) == 2)            //test b segment
              panel4->Visible = true;
        if (( code & 4 ) == 4)            //test c segment
              panel5->Visible = true;
        if (( code & 8 ) == 8)            //test d segment
              panel3->Visible = true;
        if (( code & 16 ) == 16)          //test e segment
              panel6->Visible = true;
        if (( code & 32 ) == 32)          //test f segment
              panel7->Visible = true;
        if (( code & 64 ) == 64)          //test g segment
              panel2->Visible = true;
}

private: System::Void Form1_Load(System::Object^ sender,
        System::EventArgs^ e)
{
        Clear();
}
```

8–4   **DISK FILES**

Data are found stored on the disk in the form of files. The disk itself is organized in four main parts: the boot sector, the **file allocation table** (FAT), the root directory, and the data storage areas. The Windows NTFS (**New Technology File System**) contains a boot sector and a **master file table** (MFT). The first sector on the disk is the boot sector, which is used to load the disk operating system (DOS) from the disk into the memory when power is applied to the computer.

The FAT (or MFT) is where the names of files/subdirectories and their locations on the disk are stored by the operating system. All references to any disk file are handled through the FAT (or MFT). All other subdirectories and files are referenced through the **root directory** in the FAT system. The NTFS system does not have a root directory even though the file system may still appear to have a root directory. The disk files are all considered sequential access files, meaning that they are accessed a byte at a time, from the beginning of the file toward the end. Both the NTFS file system and the FAT file system are in use, with the hard disk drive on most modern Windows systems using NTFS and the floppy disk, CD-ROM, and DVD using the FAT system.

Disk Organization

Figure 8–7 illustrates the organization of sectors and tracks on the surface of the disk. This organization applies to both floppy and hard disk memory systems. The outer track is always track 0, and the inner track is 39 (double density) or 79 (high density) on floppy disks. The inner track on a hard disk is determined by the disk size, and could be 10,000 or higher for very large hard disks.

Figure 8–8 shows the organization of data on a disk. The length of the FAT is determined by the size of the disk. In the NTFS system, the length of the MFT is determined by the number
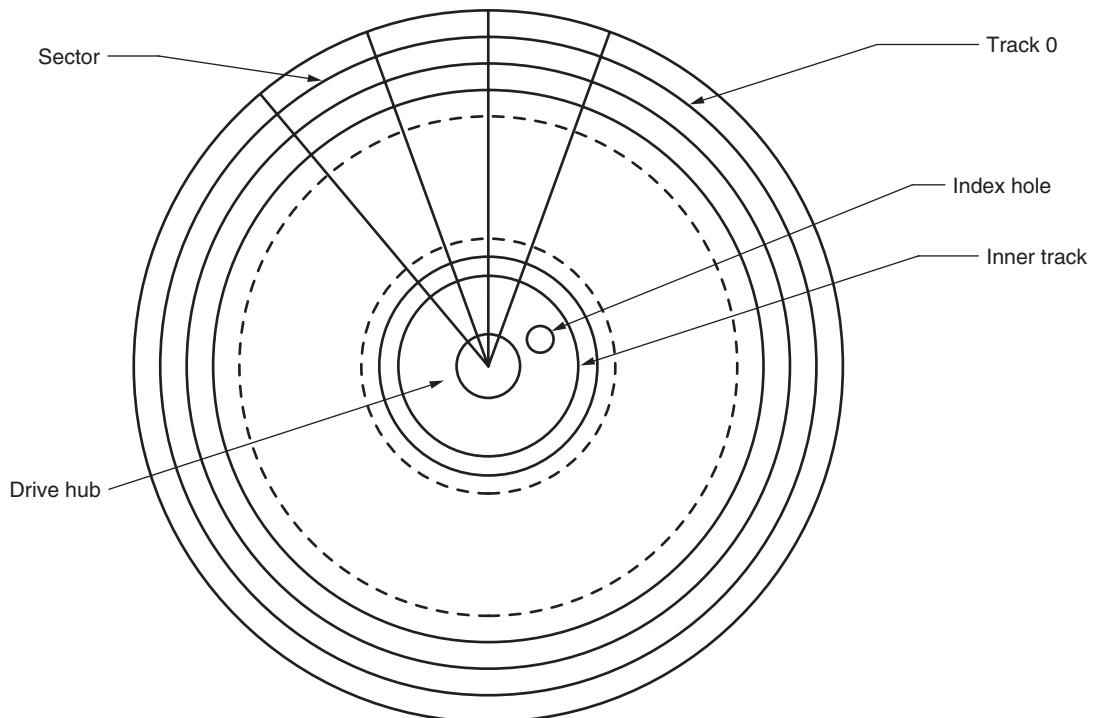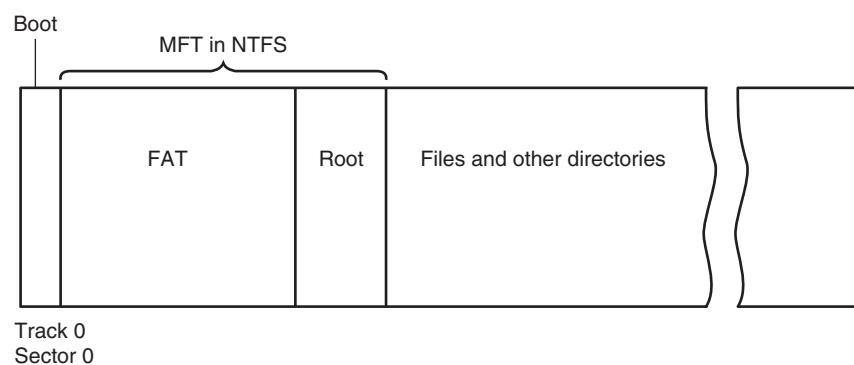
FIGURE 8–7   Structure of the disk.

FIGURE 8–8 Main data storage areas on a disk.

of files stored on the disk. Likewise, the length of the root directory, in a FAT volume, is deter-mined by the number of files and subdirectories located within it. The boot sector is always a sin-gle 512-byte-long sector located in the outer track at sector 0, the first sector.

The boot sector contains a **bootstrap loader** program that is read into RAM when the sys-tem is powered. The bootstrap loader then executes and loads the operating system into RAM. Next, the bootstrap loader passes control to the operating system program, allowing the com-puter to be under the control of and execute Windows, in most cases. This same sequence of events also occurs if the Linux operating system is found on the disk.
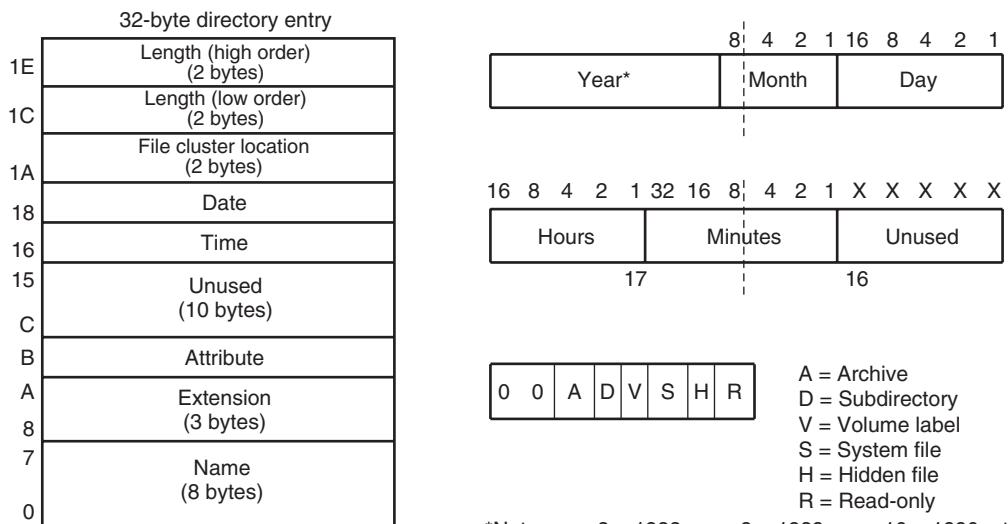
The FAT indicates which sectors are free, which are corrupted (unusable), and which con-tain data. The FAT table is referenced each time that the operating system writes data to the disk so that it can find a free sector. Each free cluster is indicated by 0000H in the FAT and each occu-pied sector is indicated by the cluster number. A **cluster** can be anything from one sector to any number of sectors in length. Many hard disk memory systems use four sectors per cluster, which means that the smallest file is 512 bytes $\times$ 4, or 2048 bytes long. In a system that uses NTFS, the cluster size is usually 4K bytes, which is eight sectors long.

Figure 8–9 shows the format of each directory entry in the root, or in any other directory or subdirectory. Each entry contains the name, extension, attribute, time, date, location, and length. The length of the file is stored as a 32-bit number. This means that a file can have a maximum length of 4G bytes. The location is the starting cluster number.

Windows NTFS uses a much larger directory entry or record (1,024 bytes) than that of the FAT system (32 bytes). The MFT record contains the file name, file date, attribute, and data. The data can be the entire contents of the file, or a pointer to where the data is stored on the disk called a file run. Generally files that are smaller than about 1500 bytes fit into the MFT record. Longer files fit into a file run or file runs. A **file run** is a series of contiguous clusters that store the file data. Figure 8–10 illustrates an MFT record in the Windows NTFS file system. The information attribute contains the create date, last modification date, create time, last modification time, and file attributes such as read-only, archive, and so forth. The security attribute stores all security information for the file for limiting access to the file in the Windows system. The header stores information about the record type, size, name (optional), and whether it is resident or not.

File Names

Files and programs are stored on a disk and referenced both by a file name and an extension to the file name. With the DOS operating system, the *file name* may only be from one to eight characters long. The file name contains just about any ASCII character, except for spaces or the " \ . / [ ] * , : < > I ; ? = characters. In addition to the file name, the file can have an optional

FIGURE 8–9   Format of any FAT directory or subdirectory entry.

one- to three-digit *extension* to the file name. Note that the name of a file and its extension are always separated by a period. If Windows 95 through Windows XP is in use, the file name can be of any length (up to 255 characters) and can even contain spaces. This is an improvement over the eight-character file name limitation of DOS. Also note that a Windows file can have more than one extension.

Directory and Subdirectory Names.   The DOS file management system arranges the data and programs on a disk into directories and subdirectories. In Windows directories and sub-directories are called file folders. The rules that apply to file names also apply to file folder names. The disk is structured so that it contains a root directory when first formatted. The root directory or folder for a hard disk used as drive C is C:\. Any other folder is placed in the root directory. For example, C:\DATA is folder DATA in the root directory. Each folder placed in the root directory can also have subdirectories or subfolders. Examples are the subfolders C:\DATA\AREA1 and C:\DATA\AREA2, in which the folder DATA contains two subfolders: AREA1 and AREA2. Subfolders can also have additional subfolders. For example, C:\DATA\AREA2\LIST depicts folder DATA, subfolder AREA, which contains a subfolder called LIST.

Sequential Access Files

All DOS files and Windows files are sequential files. A sequential file is stored and accessed from the beginning of the file toward the end, with the first byte and all bytes between it and the last accessed to read the last byte. Fortunately, files are read and written in C++ using the File class, which makes their access and manipulation easy. This section of the text describes how to

FIGURE 8–10   A record in the Master File Table in the NTFS system.

Header	Information Attribute	File Name Attribute	Data		Security Attribute

create, read, write, delete, and rename a sequential access file. To gain access to the File class, a new using must be added to the list of using statements at the top of the program. If file access is needed, add a **using namespace System::IO;** statement to the program.

File Creation. Before a file can be used, it must exist on the disk. A file is created by the File class using Create as an attribute that directs File to create a file. A file is created with create as illustrated in Example 8–29. Here the name of the file that is created by the program is stored in a String^ called FileName. Next, the File class is used to test and see if the file already exists before creating it. Finally, in the if statement the file is created.

In this example, if the file fails to open because the disk is full or the folder is not found, a Windows message box displays Cannot create file followed by the file name, and an exit from the program occurs when OK is clicked in the message box. To try this example, create a dialog application and place the code in the Load event handler. Choose a folder name that does not exist (test should probably work) and run the application. You should see the error message. If you change the FileName so it does not include the folder, you will not get the error message.

EXAMPLE 8–29

```
String^ fileName = "C:\\Test.txt";

if (File::Exists(fileName) == false)
{            // don't forget using namespace System::IO;
      try
      {
            File::Create(fileName);
      }
      catch (...)
      {
            MessageBox::Show("Cannot create " + fileName);
            Application::Exit();
      }
}

// Test.txt now exists with a length of 0 bytes
```

Writing to a File. Once a file exists, it can be written to. In fact, it would be highly unusual to create a file without writing something to it. Data are written to a file one byte at a time. The FileStream class is used to write a stream of data to the file. Data are always written starting at the very first byte in a file. Example 8–30 lists a program that creates a file in the root directory called Test1.txt and stores the letter A in each of its 256 bytes. If you execute this code and look at Test1.txt with NotePad, you will see a file filled with 256 letter As. Note that the file stream should be closed when finished using Close( ) function. Also notice in this example that an array of size byte is created using the garbage collection class in C++. It is important to use this class to create a managed array of data.

EXAMPLE 8–30

```
String^ fileName = "C:\\Test1.txt";
array<Byte>^ buffer = gcnew array<Byte>(256);

try
{
      FileStream^ fs = File::OpenWrite(fileName);
      for (int a = 0; a < 256; a++)
      {
      buffer[a] = 'A';
      }
```

```
        fs->Write(buffer, 0, buffer->Length);
        fs->Close();
}
catch (...)
{
        MessageBox::Show("Disk error");
        Application::Exit();
}
```

Suppose that a 32-bit integer must be written to a file. Because only bytes can be written, a method must be used to convert the four bytes of the integer into a form that can be written to a file. In C++ shifts are used to place the byte into the proper location to store in the array. Assembly language can also accomplish the same task in fewer bytes, as listed in Example 8–31. If you look at the assembly code for each method, you see that the assembly language method is much shorter and much faster. If speed and size are important, then the assembly code is by far the best choice, although in this case the code generated by C++ is fairly efficient.

EXAMPLE 8–31

```
int number = 0x20000;
array<Byte>^ buf = gcnew array<Byte>(4);
//C++ conversion

buf[0] = number;
buf[1] = number >> 8;
buf[2] = number >> 16;
buf[3] = number >> 24;

//Assembly language conversion
_asm
{
        mov     eax,number
        mov     buf[0],al
        mov     buf[1],ah
        bswap   eax             ;little endian to big endian
        mov     buf[2],ah
        mov     buf[3],al
}
```

Reading File Data. File data are read from the beginning of the file toward the end using the OpenRead member of File. Example 8–32 shows an example that reads the file written in Example 8–30 into a buffer called buffer1. The OpenRead function returns the number of bytes actually read from the file, but not used in this example. This works fine if the size of the file is known as it is here, but suppose that the length of the file is not known. The FileInfo class is used to find the length of a file as illustrated in Example 8–33.

EXAMPLE 8–32

```
String^ fileName = "C:\\Test1.txt";
array<Byte>^ buffer1 = gcnew array<Byte>(256);

try
{
        FileStream^ fs = File::OpenRead(fileName);
        fs->Read(buffer1, 0, 256);
        fs->Close();
}
catch (...)
{
        MessageBox::Show("Disk error");
        Application::Exit();
}
```

EXAMPLE 8–33

```
String^ fileName = "C:\\Test1.txt";
```

```
HexDump -- C:\mcc18\example\Exp8.c

00000610  61 73 65 20 27 48 27 3A 0D 0A 09 09 09 09 7B 0D   ase 'H':......{.
00000620  0A 09 09 09 09 09 53 65 74 43 68 61 6E 41 44 43   ......SetChanADC
00000630  28 20 41 44 43 5F 43 48 30 20 29 3B 09 2F 2F 20   ( ADC_CH0 );.//
00000640  66 6F 72 20 54 65 6D 70 0D 0A 09 09 09 09 09 43   for Temp.......C
00000650  6F 6E 76 65 72 74 41 44 43 28 29 3B 20 2F 2F 20   onvertADC(); //
00000660  53 74 61 72 74 20 63 6F 6E 76 65 72 73 69 6F 6E   Start conversion
00000670  0D 0A 09 09 09 09 09 77 68 69 6C 65 28 20 42 75   .......while( Bu
00000680  73 79 41 44 43 28 29 20 29 3B 20 2F 2F 20 57 61   syADC() ); // Wa
00000690  69 74 20 66 6F 72 20 63 6F 6D 70 6C 65 74 69 6F   it for completio
000006A0  6E 0D 0A 09 09 09 09 09 74 65 6D 70 20 3D 20 52   n.......temp = R
000006B0  65 61 64 41 44 43 28 29 20 2A 20 34 2E 38 39 3B   eadADC() * 4.89;
000006C0  0D 0A 09 09 09 09 09 69 74 6F 61 28 74 65 6D 70   .......itoa(temp
000006D0  2F 20 31 30 2C 20 62 75 66 66 65 72 29 3B 0D 0A   / 10, buffer);..
000006E0  09 09 09 09 09 70 75 74 73 55 53 41 52 54 28 62   .....putsUSART(b
000006F0  75 66 66 65 72 29 3B 0D 0A 09 09 09 09 09 62 72   uffer);.......br
00000700  65 61 6B 3B 0D 0A 09 09 09 09 7D 0D 0A 09 09 09   eak;......}.....
00000710  09 63 61 73 65 20 27 49 27 3A 0D 0A 09 09 09 09   .case 'I':......
00000720  7B 0D 0A 09 09 09 09 09 69 66 20 28 20 50 4F 52   {.......if ( POR
00000730  54 44 62 69 74 73 2E 52 44 33 20 3D 3D 20 30 20   TDbits.RD3 == 0
00000740  29 0D 0A 09 09 09 09 09 09 57 72 69 74 65 55 53   ).......WriteUS
00000750  41 52 54 28 27 31 27 29 3B 0D 0A 09 09 09 09 09   ART('1');.......
00000760  65 6C 73 65 0D 0A 09 09 09 09 09 09 57 72 69 74   else........Writ
00000770  65 55 53 41 52 54 28 27 30 27 29 3B 0D 0A 09 09   eUSART('0');....
00000780  09 09 09 62 72 65 61 6B 3B 0D 0A 09 09 09 09 7D   ...break;......}
00000790  0D 0A 09 09 09 7D 0D 0A 09 09 7D 0D 0A 09 7D 0D   .....}....}...}.
000007A0  0A 7D 0D 0A   .}..

                                                        Open
```

FIGURE 8–11 The HexDump program.

```
FileInfo^ fi = gcnew FileInfo(fileName);
int fileLength = fi->Length;
```

An Example Binary Dump Program. One tool not available with Windows is a program that displays the contents of a file in hexadecimal code. Although this may not be used by most programmers, it is used whenever software is developed so that the actual contents of a file can be viewed in hexadecimal format. Start a forms application in Windows and call it HexDump. Place a control called a Rich Textbox onto the form as illustrated in Figure 8–11. Under Properties for the Rich Textbox Control, make sure you change Locked to true and Scroll bars to Vertical. If you display a very large file, you will want to be able to scroll down through the code. Very large files take some time to load in this program.

This program uses the function (Disph) shown earlier in Example 8–24 to display the address as an eight-digit hexadecimal address and also to display the contents of the address in hexadecimal form as a two-digit number. Add Disph function to the program so it returns a String at the location addressed by char temp as the third parameter. The first two parameters contain two integers: one for the number and one for the number of digits called size, as shown in Example 8–34.

Example 8–34 shows the entire program required to perform a hexadecimal dump. Most of the program is generated by Visual C++, only the function at the top and a few at the end were entered to create the application. Note that to change the file for this program requires a change of the name of the file in the program. This can be modified by using an edit box to enter the file name, but it was not done in this example for sake of brevity. In this program 16 bytes are read at a time and formatted for display. This process continues until no bytes remain in the file. The

ASCII data that are displayed at the end of the hexadecimal listing are filtered so that any ASCII character under 32 (a space) are displayed as a period. This is important or control characters such as line feed, backspace, and the like will destroy the screen formatting of the ASCII text, and that is undesirable.

EXAMPLE 8–34

```
#pragma once
namespace HexDump1 {

      using namespace System;
      using namespace System::ComponentModel;
      using namespace System::Collections;
      using namespace System::Windows::Forms;
      using namespace System::Data;
      using namespace System::Drawing;
      using namespace System::IO;

// assembly code here compiled using the /CLR switch
void Disph(unsigned int number, unsigned int size, char* temp)
{
      int a;
      number <= ( 8 - size ) * 4;       //adjust position
      for (a = 0; a < size; a++)
      {
            char temp1;
            _asm
            {
                  rol number, 4;
                  mov al,byte ptr number
                  and al,0fh            ;make 0 - f
                  add al,30h            ;convert to ASCII
                  cmp al,39h
                  jbe Disph1
                  add al,7
Disph1:
                  mov temp1,al
            }
            temp[a] = temp1;           //add digit to string
      }
      temp[a] = 0;                      // null string end
}

      /// <summary>
      /// Summary for Form1
      ///
      /// WARNING: If you change the name of this class, you will need to
      change the
      ///          'Resource File Name' property for the managed resource
      compiler tool
      ///          associated with all .resx files this class depends on.
      Otherwise,
      ///          the designers will not be able to interact properly with
      localized
      ///          resources associated with this form.
      /// </summary>
      public ref class Form1 : public System::Windows::Forms::Form
      {
      public:
            Form1(void)
            {
                  InitializeComponent();
                  //
                  //TODO: Add the constructor code here
                  //
            }
```

```
        protected:
                /// <summary>
                /// Clean up any resources being used.
                /// </summary>
                ~Form1()
                {
                        if (components)
                        {
                                delete components;
                        }
                }
        private: System::Windows::Forms::RichTextBox^ richTextBox1;
        private: System::Windows::Forms::OpenFileDialog^ openFileDialog1;
        private: System::Windows::Forms::Button^ button1;
        protected:

        private:
                /// <summary>
                /// Required designer variable.
                /// </summary>
                System::ComponentModel::Container ^components;
#pragma region Windows Form Designer generated code
                /// <summary>
                /// Required method for Designer support - do not modify
                /// the contents of this method with the code editor.
                /// </summary>
                void InitializeComponent(void)
                {
                        this->richTextBox1 = (gcnew
System::Windows::Forms::RichTextBox());
                        this->openFileDialog1 = (gcnew
System::Windows::Forms::OpenFileDialog());
                        this->button1 = (gcnew System::Windows::Forms::Button());
                        this->SuspendLayout();
                        //
                        // richTextBox1
                        //
                        this->richTextBox1->Font = (gcnew
        System::Drawing::Font(L"Courier New", 9.75F,
        System::Drawing::FontStyle::Regular,
        System::Drawing::GraphicsUnit::Point,
                                static_cast<System::Byte>(0)));
                        this->richTextBox1->Location = System::Drawing::Point(12,
        12);
                        this->richTextBox1->Name = L"richTextBox1";
                        this->richTextBox1->ScrollBars =
        System::Windows::Forms::RichTextBoxScrollBars::Vertical;
                        this->richTextBox1->Size = System::Drawing::Size(657, 420);
                        this->richTextBox1->TabIndex = 0;
                        this->richTextBox1->Text = L"";
                        //
                        // openFileDialog1
                        //
                        this->openFileDialog1->FileName = L"openFileDialog1";
                        //
                        // button1
                        //
                        this->button1->Location = System::Drawing::Point(601, 438);
                        this->button1->Name = L"button1";
                        this->button1->Size = System::Drawing::Size(68, 25);
                        this->button1->TabIndex = 1;
                        this->button1->Text = L"Open";
                        this->button1->UseVisualStyleBackColor = true;
                        this->button1->Click += gcnew System::EventHandler(this,
        &Form1::button1_Click);
                        //
                        // Form1
                        //
                        this->AutoScaleDimensions = System::Drawing::SizeF(6, 13);
```

```
                        this->AutoScaleMode =
        System::Windows::Forms::AutoScaleMode::Font;
                        this->ClientSize = System::Drawing::Size(681, 468);
                        this->Controls->Add(this->button1);
                        this->Controls->Add(this->richTextBox1);
                        this->Name = L"Form1";
                        this->ShowIcon = false;
                        this->StartPosition =
        System::Windows::Forms::FormStartPosition::CenterScreen;
                        this->Text = L"HexDump";
                        this->ResumeLayout(false);
                }
#pragma endregion

private: System::String^ Disp(int number, int size)
{
        char temp[9];
        Disph(number,size,temp);
        String^ a = "";
        int count = 0;
        while (temp[count] != 0)            // convert to string
        {
                Char b = temp[count++];
                a += b;
        }
        return a;
}

private: System::Void button1_Click(System::Object^ sender,
        System::EventArgs^ e)
{
        array<Byte>^ buffer = gcnew array<Byte>(1000000);
        int fileLength;
        String^ line = "";
        if (openFileDialog1->ShowDialog() ==
                System::Windows::Forms::DialogResult::OK)
        {
                try
                {
                        FileStream^ fs = File::OpenRead(openFileDialog1->FileName);
                        fs->Read(buffer, 0, 1000000);
                        fs->Close();
                        FileInfo^ fi = gcnew FileInfo(openFileDialog1->FileName);
                        fileLength = fi->Length;
                        this->Text = "HexDump -- " + openFileDialog1->FileName;
                }
                catch (...)
                {
                        MessageBox::Show("Disk error");
                        Application::Exit();
                }
                for (int a = 0; a < fileLength; a++)
                {
                        if (a % 16 == 0)
                        {
                                if (a != 0)
                                {
                                        richTextBox1->Text += " " + line;
                                        line = "";
                                        richTextBox1->Text += "\n";
                                }
                                richTextBox1->Text += Disp(a, 8);
                        }
                        richTextBox1->Text += " " + Disp(buffer[a], 2);
                        if (buffer[a] >= 32 && buffer[a] < 128)
                                line += Convert::ToChar(buffer[a]);
                        else
                                line += ".";
```

```
                        }
                        richTextBox1->Text += " " + line;
                }
                else
                {
                        this->Text = "HexDump";
                }
        }
}

};
}
```

The File Pointer and Seek.   When a file is opened, written, or read, the file pointer addresses
the current location in the sequential file. When a file is opened, the file pointer always addresses
the first byte of the file. If a file is 1024 bytes long, and a read function reads 1023 bytes, the file
pointer addresses the last byte of the file, but not the end of the file.

The **file pointer** is a 32-bit number that addresses any byte in a file. The File Append
member function is used to add new information to the end of a file. The file pointer can be
moved from the start of the file or from the end of the file. Open moves the pointer to the start of
the file. In practice, both are used to access different parts of the file. The FileStream member
function Seek allows the file pointer to be moved to the start of a file (SeekOrigin::Begin), the
end of a file (SeekOrigin::End), or the current location in the file (SeekOrigin::Current). The first
number in the Seek function is the offset. If the third byte in the file is accessed, it is accessed
with a Seek(2, SeekOrigin::Begin) function. (The third byte is at offset 2.) Note that the second
number in the Write function is also an offset and can be used in the same manner as a Seek.

Suppose that a file exists on the disk and that you must append the file with 256 bytes of
new information. When the file is opened, the file pointer addresses the first byte of the file. If
you attempt to write without moving the file pointer to the end of the file, the new data will
overwrite the first 256 bytes of the file. Example 8–35 shows a sequence of instructions for
Appends, which adds 256 bytes of data to the end of the file, and then closes the file. This file is
appended with 256 new bytes of data from area Buffer.

EXAMPLE 8–35
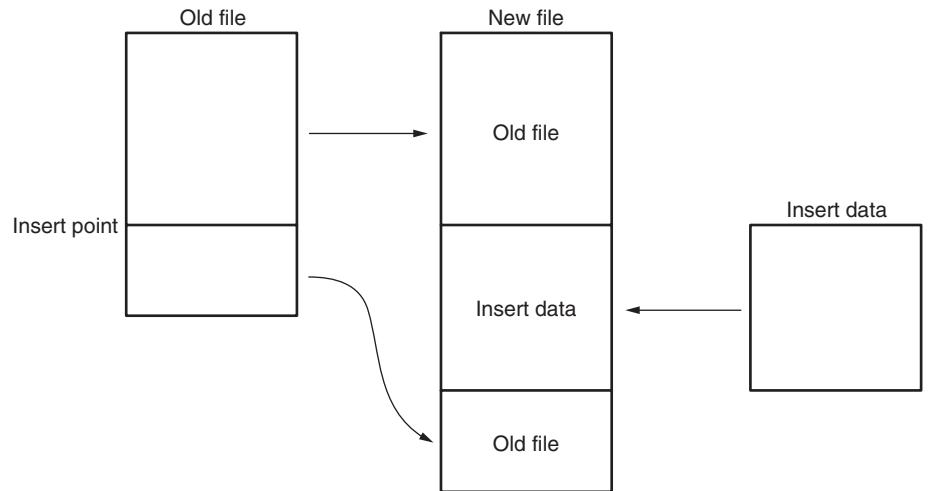
```
String^ fileName = "C:\\Test1.txt";
array<Byte>^ buffer = gcnew array<Byte>(256);

try
{
        FileStream^ fs = File::OpenWrite(fileName);
        for (int a = 0; a < 256; a++)
        {
                buffer[a] = 'S';
        }
        fs->Seek(0, SeekOrigin::End);
        fs->Write(buffer, 0, buffer->Length);
        fs->Close();
}
catch (...)
{
        MessageBox::Show("Disk error");
        Application::Exit();
}


// or the same operation is performed using the offset number
// in the Write function as follows:

String^ fileName = "C:\\Test1.txt";
array<Byte>^ buffer = gcnew array<Byte>(256);
```

FIGURE 8–12 Inserting
new data within an old file.


```
try
{
        FileStream^ fs = File::OpenWrite(fileName);
        for (int a = 0; a < 256; a++)
        {
                buffer[a] = 'S';
        }
        fs->Write(buffer, 256, buffer->Length);
        fs->Close();
}
catch (...)
{
        MessageBox::Show("Disk error");
        Application::Exit();
}
```

One of the more difficult file maneuvers is inserting new data in the middle of the file. Figure 8–12 shows how this is accomplished by creating a second file. Notice that the part of the file before the insertion point is copied into the new file. This is followed by the new information before the remainder of the file is appended after the insertion in the new file. Once the new file is complete, the old file is deleted and the new file is renamed to the old file name.

Example 8–36 shows a program that inserts new data into an old file. This program copies the Data.new file into the Data.old file at a point after the first 256 bytes of the Data.old file. The new data from buffer2 is next added to the file and then this is followed by the remainder of the old file. New File member functions are used to delete the old file and rename the new file to the old file name.

EXAMPLE 8–36

```
private: System::Void Form1_Load(System::Object^ sender,
        System::EventArgs^ e)
{
        String^ fileName1 = "C:\\Data.old";
        String^ fileName2 = "C:\\Data.new";
        int fileLength;
        array<Byte>^ buffer1 = gcnew array<Byte>(256);
        array<Byte>^ buffer2 = gcnew array<Byte>(6);

        try
        {
                FileStream^ fs1 = File::OpenWrite(fileName1);
                FileStream^ fs2 = File::OpenWrite(fileName2);
```

```
                    FileInfo^ fi = gcnew FileInfo(fileName1);
                    fileLength = fi->Length;
                    fs1->Read(buffer1, 0, 256);
                    fs2->Write(buffer1, 0, 256);
                    fs2->Write(buffer2, 0, 6);
                    fileLength -= 256;
                    while (fileLength > 0)
                    {
                            fs1->Read(buffer1, 0, 256);
                            fs2->Write(buffer1, 0, 256);
                            fileLength -= 256;
                    }
                    fs1->Close();
                    fs2->Close();
            }
            catch (...)
            {
                    MessageBox::Show("Disk error");
                    Application::Exit();
            }
            }
```

Random Access Files

Random access files are developed through software using sequential access files. A random access file is addressed by a record number rather than by going through the file searching for data. The Seek function becomes very important when random access files are created. Random access files are much easier to use for large volumes of data, which are often called databases.

Creating a Random Access File.   Planning is paramount when creating a random access file system. Suppose that a random access file is required for storing the names of customers. Each customer record requires 32 bytes for the last name, 32 bytes for the first name, and one byte for the middle initial. Each customer record contains two street address lines of 64 bytes each, a city line of 32 bytes, two bytes for the state code, and nine bytes for the Zip Code. The basic customer information alone requires 236 bytes; additional information expands the record to 512 bytes. Because the business is growing, provisions are made for 5000 customers. This means that the total random access file is 2,560,000 bytes long.

Example 8–37 illustrates a short program that creates a file called CUST.FIL and inserts 5000 blank records of 512 bytes each. A blank record contains 00H in each byte. This appears be a large file, but it fits on the smallest of hard disks.

EXAMPLE 8–37

```
private: System::Void Form1_Load(System::Object^ sender,
      System::EventArgs^ e)
{
      String^ fileName = "C:\\Cust.fil";
      array<Byte>^ buffer = gcnew array<Byte>(512);

      for ( int a = 0; a < 512; a++ )            //fill buffer
      {
            buffer[a] = 0;
      }

      try
      {
            FileStream^ fs = File::OpenWrite(fileName);
            for (int a = 0; a < 5000; a++)
            {
                  fs->Write(buffer, 0, 512);
            }
            fs->Close();
```

```
        }
        catch (...)
        {
                MessageBox::Show("Disk error");
                Application::Exit();
        }
}
```

Reading and Writing a Record.   Whenever a record must be read, the record number is found by using a Seek. Example 8–38 lists a function that is used to Seek to a record. This function assumes that a file has been opened as CustomerFile and that the CUST.FIL remains open at all times.

Notice how the record number is multiplied by 512 to obtain a count to move the file pointer using a Seek. In each case, the file pointer is moved from the start of the file to the desired record.

EXAMPLE 8–38

```
void CCusDatabaseDlg::FindRecord(unsigned int RecordNumber)
{
        File.Seek( RecordNumber * 512, CFile::begin );
}
```

Other functions (listed in Example 8–39) are needed to manage the customer database. These include WriteRecord, ReadRecord, FindLastNameRecord, FindBlankRecord, and so on. Some of these are listed in the example as well as the data structure that contains the information for each record.

EXAMPLE 8–39

```
// class placed before the form1 class for containing a record

public ref class Customer
{
        public: static array<Byte>^ FirstName = gcnew array<Byte>(32);
        public: static array<Byte>^ Mi = gcnew array<Byte>(1);
        public: static array<Byte>^ LastName = gcnew array<Byte>(32);
        public: static array<Byte>^ Street1 = gcnew array<Byte>(64);
        public: static array<Byte>^ Street2 = gcnew array<Byte>(64);
        public: static array<Byte>^ City = gcnew array<Byte>(32);
        public: static array<Byte>^ State = gcnew array<Byte>(2);
        public: static array<Byte>^ ZipCode = gcnew array<Byte>(9);
        public: static array<Byte>^ Other = gcnew array<Byte>(276);
};

// functions placed at the end of the form1 class

static array<Byte>^ buffer = gcnew array<Byte>(512);
static String^ fileName = "C:\\Cust.fil";
static FileStream^ fs;
static Customer Record;

private: System::Void Form1_Load(System::Object^ sender,
        System::EventArgs^ e)
{       // open the file when the application starts
        Customer Record;
        try
        {
                fs = File::OpenWrite(fileName);
                for (int a = 0; a < 5000; a++)
                {
                    fs->Write(buffer, 0, 512);
                }
                fs->Close();
        }
        catch (...)
        {
```

```
                    MessageBox::Show("Disk error");
                    Application::Exit();
        }
}

private: System::Void FindRecord(unsigned int RecordNumber)
{
        fs->Seek(RecordNumber * 512, SeekOrigin::Begin);
}

private: System::Void WriteRecord(unsigned int RecordNumber)
{
        FindRecord(RecordNumber);
        fs->Write(Record.FirstName, 0, 32);
        fs->Write(Record.Mi, 0, 1);
        fs->Write(Record.LastName, 0, 32);
        fs->Write(Record.Street1, 0, 64);
        fs->Write(Record.Street2, 0, 64);
        fs->Write(Record.City, 0, 32);
        fs->Write(Record.State, 0, 2);
        fs->Write(Record.ZipCode, 0, 9);
}

private: System::Void ReadRecord(unsigned int RecordNumber)
{
        FindRecord(RecordNumber);
        fs->Read(Record.FirstName, 0, 32);
        fs->Read(Record.Mi, 0, 1);
        fs->Read(Record.LastName, 0, 32);
        fs->Read(Record.Street1, 0, 64);
        fs->Read(Record.Street2, 0, 64);
        fs->Read(Record.City, 0, 32);
        fs->Read(Record.State, 0, 2);
        fs->Read(Record.ZipCode, 0, 9);
}


private: System::UInt32 FindFirstName(array<Byte>^ FirstName)
{
        for ( int a = 0; a < 5000; a++ )
        {
                ReadRecord(a);
                if (Record.FirstName == FirstName)
                {
                        return a;     //if found return record number
                }
        }
        return 5001;                 //if not found return 5001
}


private: System::UInt32 FindBlankRecord()
{
        for ( int a = 0; a < 5000; a++ )
        {
                ReadRecord(a);
                if (Record.LastName[0] == 0 )
                {
                        return a;
                }
        }
        return 0;
}
```

FIGURE 8–13   The DataTime application.

8–5   EXAMPLE PROGRAMS

Now that many of the basic programming building blocks have been discussed, we present some example application programs. Although these example programs may seem trivial, they show some additional programming techniques and illustrate programming styles for the microprocessor.

Time/Date Display Program
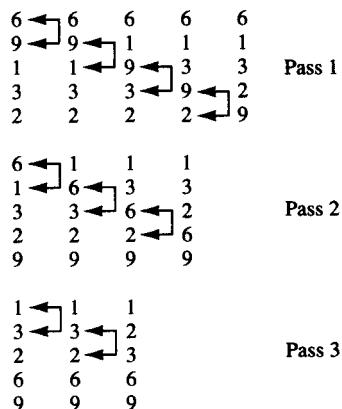
Although this program does not use assembly language, it does demonstrate how to obtain the date and time from the Windows API and how to format it for display. It also illustrates how to use a timer in Visual C++. Example 8–40 illustrates a program that uses a timer, set to interrupt the program once per second, to display the time and date. The time and date are obtained by using DateTime object to read the computer time and date into a variable called dt. The format member TimeDate is used to format the dt variable. Create a dialog application called DateTime and place two labels on it as shown in Figure 8–13.

EXAMPLE 8–40

```
private: System::Void Form1_Load(System::Object^ sender,
      System::EventArgs^ e)
{
      ShowDateTime();
}

private: System::Void ShowDateTime()
{
      DateTime dt = DateTime::Now;              // get current time
      label1->Text = dt.Hour.ToString() + ":" + dt.Minute.ToString();
      label2->Text = dt.Date.ToLongDateString();
}
```

FIGURE 8–14 A bubble sort showing data as they are sorted. Note: Sorting five numbers may require four passes.

private: System::Void timer1_Tick(System::Object^ sender,
```
        System::EventArgs^ e)
{
        ShowDateTime();
}
```
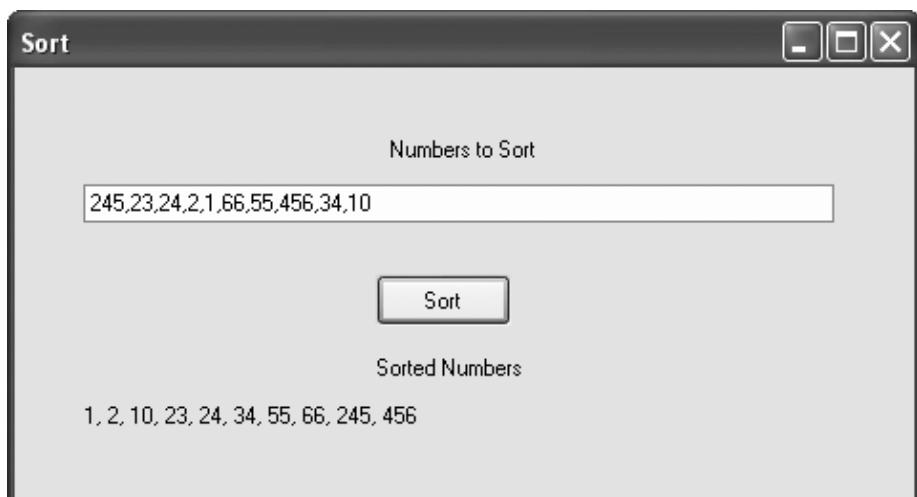
Numeric Sort Program

At times, numbers must be sorted into numeric order. This is often accomplished with a bubble sort. Figure 8–14 shows five numbers that are sorted with a bubble sort. Notice that the set of five numbers is tested four times with four passes. For each pass, two consecutive numbers are compared and sometimes exchanged. Also notice that during the first pass there are four comparisons, during the second three, and so forth.

Example 8–41 illustrates a program that accepts 10 numbers from the keyboard (32-bit integers). After these 32-bit numbers are accepted and stored in memory section numbers, they are sorted by using the bubble-sorting technique. This bubble sort uses a swap flag to determine whether any numbers were exchanged in a pass. If no numbers were exchanged, the numbers are in order and the sort terminates. This early termination normally increases the efficiency of the sort because numbers are rarely completely out of order.

Once the numbers are sorted, they are displayed in ascending order. Figure 8–15 shows how the application appears after it is executed.

FIGURE 8–15 The bubble sort.

EXAMPLE 8–41

```
void Sort(int* data)
{
        char flag;
        _asm
        {
                mov  ecx,9            ;9 for 10 numbers
L1:
                mov  flag,0           ;clear flag
                mov  edx,0
L2:
                mov  ebx,data
                mov  eax,[ebx+edx*4]
                cmp  eax,[ebx+edx*4+4]
                jbe  L3
                push eax              ;swap
                mov  eax,[ebx+edx*4+4]
                mov  [ebx+edx*4], eax
                pop  dword ptr [ebx+edx*4+4]
                mov  flag,1           ;set flag
L3:
                inc  edx
                cmp  edx,ecx
                jne  L2
                cmp  flag,0
                jz   L4               ;if no swaps
                loop L1
L4:
        }
}
bool isHandled;

private: System::Void button1_Click(System::Object^ sender,
        System::EventArgs^ e)
{
        int numbers[10];
        int count = 0;
        int digit = 0;
        int a;
        for(a = 0; a < 10; a++)
        {
                numbers[a] = 0;
                while (digit < textBox1->Text->Length && textBox1->Text[digit]
                        != ',')
                {
                        numbers[a] = numbers[a] * 10 +
                                (int)(textBox1->Text[digit] - 0x30);
                        digit++;
                }
                digit++;
                if (digit >= textBox1->Text->Length)
                {
                        break;
                }
        }
        if (a == 9)
        {
                Sort(numbers);
                label2->Text = "";
                for (int a = 0; a < 9; a++)
                {
                        label2->Text += numbers[a].ToString() + ", ";
                }
```

```
                    label2->Text += numbers[9].ToString();
        }
        else
        {
                MessageBox::Show(
                        "10 numbers must be entered separated by commas");
        }
}

private: System::Void textBox1_KeyDown(System::Object^ sender,
        System::Windows::Forms::KeyEventArgs^ e)
{
        isHandled = true;
        if (e->KeyCode >= Keys::D0 && e->KeyCode <= Keys::D9 ||
                e->KeyCode == Keys::Oemcomma || e->KeyCode == Keys::Back)
        {
                isHandled = false;
        }
}

private: System::Void textBox1_KeyPress(System::Object^ sender,
        System::Windows::Forms::KeyPressEventArgs^ e)
{
        e->Handled = isHandled;
}
```
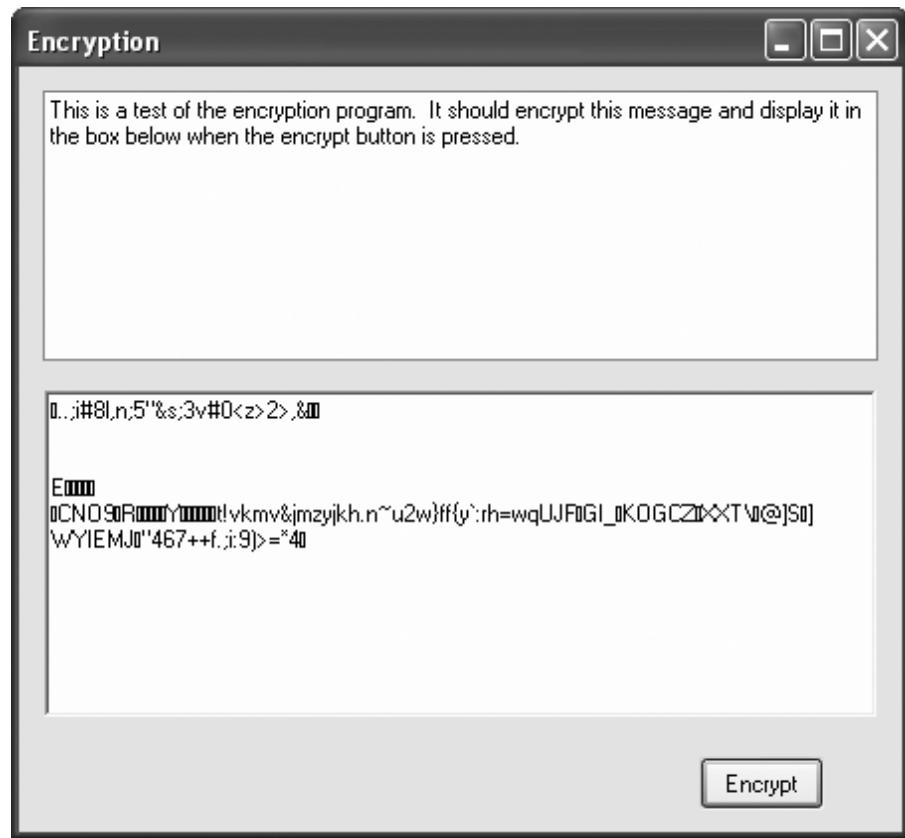
Data Encryption

Data encryption seems to be the vogue at this time because of the security aspect of many systems. To illustrate simple data encryption for a character string, suppose that each character in a string is exclusive-ORed with a number called an encryption key. This certainly changes the code of the character, but to make it a bit more random, suppose that the encryption key is changed after each character is encrypted. In this way patterns are much harder to detect in the encrypted message, making it harder to decipher.

To illustrate this simple scheme, Figure 8–16 shows a screen shot of the program to test the scheme, using a textbox control to accept a character string and a label to display the encrypted message. This example was generated using an initial encryption key of 0×45. If the initial value is changed, the encrypted message will change.

Example 8–42 lists the program used to generate the message in its encrypted form in a rich textbox control. The button event handler reads the contents of the textbox control, used for entering the character string to be encrypted, and uses a short assembly language function to encrypt the string. Notice how the program uses assembly language to Exclusive-OR each character of the string with the EncryptionKey and then how the EncryptionKey is modified for the next character. The technique used here increments the Encryption key and prevents the key from becoming larger than 7FH. This technique can be made more intricate to make it even more difficult to decipher. For example, suppose that the key is incremented on every other character and that is alternated with inverting the key, as shown in Example 8–43. Almost any combination of operations can be used to modify the key between passes to make it very difficult to decode. In practice we use a 128-bit key and the technique for modification is different, but nonetheless, this is basically how encryption is performed. Because Example 8–40 uses an 8-bit key, the encrypted message could be cracked by trying all 256 (2^8) possible keys, but if a 128-bit key is used, it requires far many more attempts (2^{128}) to crack—an almost impossible number of attempts.

FIGURE 8–16 Data encryption application.

EXAMPLE 8–42

```
char EncryptionKey = 0x45;
char Encrypt(char code)
{
        _asm
        {
                mov  al,code
                xor  al,EncryptionKey
                mov  code,al
                mov  al,EncryptionKey
                inc  al
                and  al,7fh
                mov  EncryptionKey,al
        }
        return code;
}

private: System::Void button1_Click(System::Object^ sender,
        System::EventArgs^ e)
{
        richTextBox1->Text = "";
        for (int a = 0; a < textBox1->Text->Length; a++)
        {
        richTextBox1->Text += Convert::ToChar(Encrypt(textBox1->Text[a]));
        }
}
```

EXAMPLE 8–43

```
//just the assembly language part of the program
```

```
char EncryptionKey = 0x45;
char everyOther = 0;

char Encrypt(char code)
{
        _asm
        {
                mov   al,code
                xor   al,EncryptionKey
                mov   code,al
                mov   al,everyOther
                inc   al
                and   al,1
                mov   everyOther,al
                mov   bl,EncryptionKey
                cmp   al,0
                jz    L1
                inc   bl
                jmp   L2
L1:
                not   bl
L2:
                and   bl,7fh
                mov   EncryptionKey,bl
        }
        return code;
}
```

8–6   SUMMARY

1. The assembler program (ML.EXE) assembles modules that contain PUBLIC variables and segments, plus EXTRN (external) variables. The linker program (LINK.EXE) links modules and library files to create a run-time program executed from the DOS command line. The run-time program usually has the extension EXE, but might contain the extension COM.

2. The MACRO and ENDM directives create a new opcode for use in programs. These macros are similar to procedures, except that there is no call or return. In place of them, the assembler inserts the code of the macro sequence into a program each time it is invoked. Macros can include variables that pass information and data to the macro sequence.

3. Setting focus to an object is accomplished by using the Focus( ) member variable found with most objects.

4. The Convert class in C++ is used to convert from one form to another in many cases, but not in all cases.

5. The mouse driver is accessed from Windows by installing handlers for various Windows events such as MouseMove, MouseDown, etc.

6. Conversion from binary to BCD is accomplished with the AAM instruction for numbers that are less than 100 or by repeated division by 10 for larger numbers. Once the number is converted to BCD, 30H is added to convert each digit to ASCII code for placement in a string.

7. When converting from an ASCII number to BCD, 30H is subtracted from each digit. To obtain the binary equivalent, multiply by 10 and then add each new digit.

8. Lookup tables are used for code conversion with the XLAT instruction if the code is an 8-bit code. If the code is wider than 8 bits, a short procedure that accesses a lookup table provides the conversion. Lookup tables are also used to hold addresses so that different parts of a program or different procedures can be selected.

9. Conditional assembly language statements allow portions of a program to be assembled if a condition is met. These are useful for tailoring software to an application. In Visual

C++ Express, a program that contains assembly code must be compiled with the /CLR switch.

10. The disk, memory system contains tracks that hold information stored in sectors. Many disk systems store 512 bytes of information per sector. Data on the disk are organized in a boot sector, file allocation table, root directory, and data storage area. The boot sector loads the DOS system from the disk into the computer memory system. The FAT or MFT indicates which sectors are present and whether they contain data. The root directory contains file names and subdirectories through which all disk files are accessed. The data storage area contains all subdirectories and data files.

11. Files are manipulated with the File object in Visual C++. To read a disk file, the file must be opened, read, and then closed. To write to a disk file, it must be opened, written, and then closed. When a file is opened, the file pointer addresses the first byte of the file. To access data at other locations, the file pointer is moved using a Seek before data are read or written.

12. A sequential access file is a file that is accessed sequentially from the beginning to the end. A random access file is a file that is accessed at any point. Although all disk files are sequential, they can be treated as random access files by using software.

8–7   QUESTIONS AND PROBLEMS

1. The assembler converts a source file to a(n) _____ file.
2. What files are generated from the source file TEST.ASM if it is processed by ML.EXE?
3. The linker program links object files and _____ files to create an execution file.
4. What does the PUBLIC directive indicate when placed in a program module?
5. What does the EXTRN directive indicate when placed in a program module?
6. What directive appears with labels defined as external?
7. Describe how a library file works when it is linked to other object files by the linker program.
8. What assembler language directives delineate a macro sequence?
9. What is a macro sequence?
10. How are parameters transferred to a macro sequence?
11. Develop a macro called ADD32 that adds the 32-bit contents of DX-CX to the 32-bit contents of BX-AX.
12. How is the LOCAL directive used within a macro sequence?
13. Develop a macro called ADDLIST PARA1,PARA2 that adds the contents of PARA1 to PARA2. Each of these parameters represents an area of memory. The number of bytes added are indicated by register CX before the macro is invoked.
14. Develop a macro that sums a list of byte-sized data invoked by the macro ADDM LIST,LENGTH. The label LIST is the starting address of the data block and LENGTH is the number of data added. The result must be a 16-bit sum found in AX at the end of the macro sequence.
15. What is the purpose of the INCLUDE directive?
16. Modify the function in Example 8–12 so that it filters the numbers 0 through 9 from only the keyboard and not the keypad and ignores all other characters.
17. Modify the function in Example 8–12 so that it generates a random 8-bit number in class variable char Random. (Hint: To accomplish this, increment Random each time that the KeyDown function is called.)
18. Modify the software you developed in question 17 so that it generates a random number between 9 and 62.
19. Modify the function listed in Example 8–15 so that the hexadecimal numbers use lowercase letters a through f instead of the uppercase letters.
20. Modify Example 8–16 so it will shift/rotate left or right. This is accomplished by adding a pair of radio buttons to select the direction.

21. What event handlers are used to access the mouse in the Visual C++ programming environment and what event causes each handler to be called?
22. How is the right mouse button detected in a program?
23. How is a double-click detected with the mouse?
24. Develop software that detects when both the right and left mouse buttons are pressed simultaneously.
25. How is a color selected in a program using Visual C++?
26. What is the purpose of the ForeColor property?
27. When a number is converted from binary to BCD, the _____ instruction accomplishes the conversion, provided the number is less than 100 decimal.
28. How is a large number (over 100 decimal) converted from binary to BCD?
29. How could a binary number be displayed as an octal number?
30. A BCD digit is converted to ASCII code by adding a(n) _____.
31. An ASCII-coded number is converted to BCD by subtracting _____.
32. Develop a function that reads an ASCII number from a textbox control as keys are typed (use KeyDown) on the keyboard and returns it as an unsigned int. The number in the textbox is an octal number that is converted to binary by the function.
33. Explain how a three-digit ASCII-coded number is converted to binary.
34. Develop a function that converts all lowercase ASCII-coded letters into uppercase ASCII-coded letters. Your procedure may not change any other character except the letters a–z and must return the converted character as a char.
35. Develop a lookup table that converts hexadecimal data 00H–0FH into the ASCII-coded characters that represent the hexadecimal digits. Make sure to show the lookup table and any software required for the conversion. It is suggested that a function is created to perform the conversion.
36. Explain the purpose of a boot sector, FAT, and root directory in the FAT system.
37. Explain the purpose of the MFT in the NTFS file system.
38. The surface of a disk is divided into tracks that are further subdivided into _____.
39. What is a bootstrap loader and where is it found?
40. What is a cluster?
41. The NTFS file system often uses cluster of _____ bytes in length.
42. What is the maximum length of a file?
43. What code is used to store the name of a file when long file names are in use?
44. DOS file names are at most _____ characters in length.
45. How many characters normally appear in an extension?
46. How many characters may appear in a long file name?
47. Develop a program that opens a file called TEST.LST, reads 512 bytes from the file into memory area Array, and closes the file.
48. Show how to rename file TEST.LST to TEST.LIS.
49. What is the purpose of the File Move member function?
50. What is a control?
51. Write a program that reads any decimal number between 0 and 2G and displays the 32-bit binary version on the video display.
52. Write a program that displays the binary powers of 2 (in decimal) on the video screen for the powers 0 through 7. Your display shows 2^n = value for each power of 2.
53. Using a timer to generate a random number, develop a program that displays random numbers between 1 and 47 (or whatever) for your state's lottery.
54. Modify the program in Example 8–28 so it also displays the letters A, b, C, d, E, and F for a hexadecimal seven-segment display.
55. Modify Example 8–42 to encrypt the message using an algorithm of your own design.
56. Develop a Decryption function (for a String) to accompany the encryption of question 55.

CHAPTER 9

8086/8088 Hardware Specifications

INTRODUCTION

In this chapter, the pin functions of both the 8086 and 8088 microprocessors are detailed and information is provided on the following hardware topics: clock generation, bus buffering, bus latching, timing, wait states, and minimum mode operation versus maximum mode operation. These simple microprocessors are explained first, because of their less intricate structures, as an introduction to the Intel microprocessor family.

Before it is possible to connect or interface anything to the microprocessor, it is necessary to understand the pin functions and timing. These rudimentary microprocessors contain the same basic pins as the latest Pentium 4 or Core2 microprocessor. Thus, the information in this chapter is essential to a complete understanding of memory and I/O interfacing, which we cover in the later chapters of the text.

CHAPTER OBJECTIVES

Upon completion of this chapter, you will be able to:

1. Describe the function of each 8086 and 8088 pin.
2. Understand the microprocessor's DC characteristics and indicate its fan-out to common logic families.
3. Use the clock generator chip (8284A) to provide the clock for the microprocessor.
4. Connect buffers and latches to the buses.
5. Interpret the timing diagrams.
6. Describe wait states and connect the circuitry required to cause various numbers of waits.
7. Explain the difference between minimum and maximum mode operation.

9–1　PIN-OUTS AND THE PIN FUNCTIONS

In this section, we explain the function and (in certain instances) the multiple functions of each of the microprocessor's pins. In addition, we discuss the DC characteristics to provide a basis for understanding the later sections on buffering and latching.

The Pin-Out

Figure 9–1 illustrates the pin-outs of the 8086 and 8088 microprocessors. As a close comparison reveals, there is virtually no difference between these two microprocessors—both are packaged in 40-pin **dual in-line** packages (DIPs).

As mentioned in Chapter 1, the 8086 is a 16-bit microprocessor with a 16-bit data bus and the 8088 is a 16-bit microprocessor with an 8-bit data bus. (As the pin-outs show, the 8086 has pin connections AD_0–AD_{15}, and the 8088 has pin connections AD_0–AD_7.) Data bus width therefore the only major difference between these microprocessors. This allows the 8086 to transfer 16-bit data more efficiently.

There is, however, a minor difference in one of the control signals. The 8086 has an $M/\overline{IO}$ pin, and the 8088 has an $IO/\overline{M}$ pin. The only other hardware difference appears on Pin 34 of both integrated circuits: on the 8088, it is an SS0 pin, whereas on the 8086, it is a $\overline{BHE}$ /S7 pin.

Power Supply Requirements

Both the 8086 and 8088 microprocessors require +5.0 V with a supply voltage tolerance of ±10 percent. The 8086 uses a maximum supply current of 360 mA, and the 8088 draws a maximum of 340 mA. Both microprocessors operate in ambient temperatures of between 32° F and 180° F. This range is not wide enough to be used outdoors in the winter or even in the summer, but extended temperature-range versions of the 8086 and 8088 microprocessors are available. There is also a CMOS version, which requires a very low supply current and has an extended temperature range. The 80C88 and 80C86 are CMOS versions that require only 10 mA of power supply current and function in temperature extremes of −40° F through +225° F.

DC Characteristics

It is impossible to connect anything to the pins of the microprocessor without knowing the input current requirement for an input pin and the output current drive capability for an output pin. This knowledge allows the hardware designer to select the proper interface components for use with the microprocessor without the fear of damaging anything.

FIGURE 9–1 (a) The pin-out of the 8086 in maximum mode; (b) the pin-out of the 8086 in minimum mode.
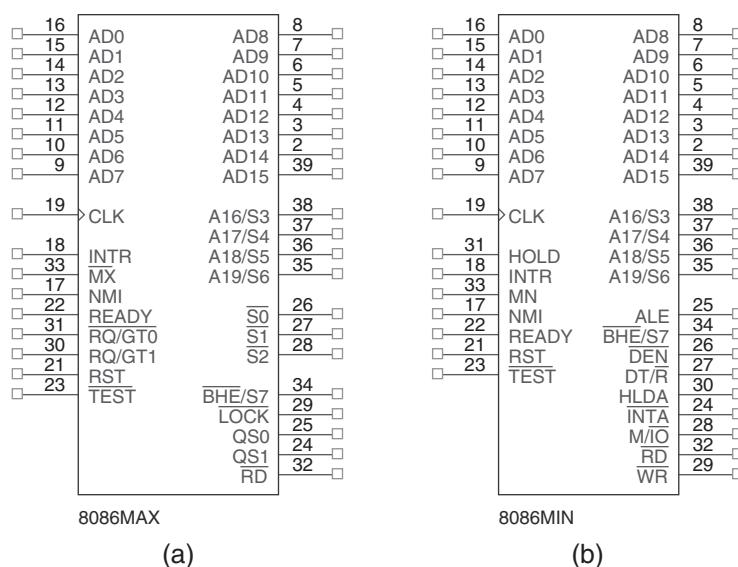
TABLE 9–1 Input characteristics of the 8086 and 8088 microprocessors.

Logic Level	Voltage	Current
0	0.8 V maximum	±10 μA maximum
1	2.0 V minimum	±10 μA maximum

Input Characteristics.   The input characteristics of these microprocessors are compatible with all the standard logic components available today. Table 9–1 depicts the input voltage levels and the input current requirements for any input pin on either microprocessor. The input current levels are very small because the inputs are the gate connections of MOSFETs and represent only leakage currents.

Output Characteristics.   Table 9–2 illustrates the output characteristics of all the output pins of these microprocessors. The logic 1 voltage level of the 8086/8088 is compatible with that of most standard logic families, but the logic 0 level is not. Standard logic circuits have a maximum logic 0 voltage of 0.4 V, and the 8086/8088 has a maximum of 0.45 V. Thus, there is a difference of 0.05 V.

This difference reduces the noise immunity from a standard level of 400 mV (0.8 V – 0.45 V) to 350 mV. (The noise immunity is the difference between the logic 0 output voltage and the logic 0 input voltage levels.) The reduction in noise immunity may result in problems with long wire connections or too many loads. It is therefore recommended that no more than 10 loads of any type or combination be connected to an output pin without buffering. If this loading factor is exceeded, noise will begin to take its toll in timing problems.

Table 9–3 lists some of the more common logic families and the recommended fan-out from the 8086/8088. The best choice of component types for the connection to an 8086/8088 output pin is an LS, 74ALS, or 74HC logic component. Note that some of the fan-out currents calculate to more than 10 unit loads. It is therefore recommended that if a fan-out of more than 10 unit loads is required, the system should be buffered.

Pin Connections

$AD_7–AD_0$ The 8088 **address/data bus** lines are the multiplexed address data bus of the 8088 and contain the rightmost 8 bits of the memory address or I/O port number whenever ALE is active (logic 1) or data whenever ALE is inactive (logic 0). These pins are at their high-impedance state during a hold acknowledge.

TABLE 9–2 Output characteristics of the 8086 and 8088 microprocessors.

Logic Level	Voltage	Current
0	0.45V maximum	2.0 mA maximum
1	2.4 V minimum	−400 μA maximum

TABLE 9–3 Recommended fan-out from any 8086/8088 pin connection.

Family	Sink Current	Source Current	Fan-out
TTL (74)	−1.6 mA	40 μA	1
TTL (74LS)	−0.4 mA	20 μA	5
TTL (74S)	−2.0 mA	50 μA	1
TTL (74ALS)	−0.1 mA	20 μA	10
TTL (74AS)	−0.5 mA	25 μA	10
TTL (74F)	−0.5 mA	25 μA	10
CMOS (74HC)	− 10 μA	10 μA	10
CMOS (CD)	− 10 μA	10 μA	10
NMOS	− 10 μA	10 μA	10

TABLE 9–4 Function of status bits S_3 and S_4.

S_4	S_3	Function
0	0	Extra segment
0	1	Stack segment
1	0	Code or no segment
1	1	Data segment

A_{15}–A_8 The 8088 **address bus** provides the upper-half memory address bits that are present throughout a bus cycle. These address connections go to their high-impedance state during a hold acknowledge.

AD_{15}–AD_8 The 8086 **address/data bus** lines compose the upper multiplexed address/data bus on the 8086. These lines contain address bits A_{15}–A_8 whenever ALE is a logic 1, and data bus connections D_{15}–D_8 when ALE is a logic 0. These pins enter a high-impedance state when a hold acknowledge occurs.

A_{19}/S_6–A_{16}/S_3 The **address/status bus** bits are multiplexed to provide address signals A_{19}–A_{16} and also status bits S_6–S_3. These pins also attain a high-impedance state during the hold acknowledge.

Status bit S_6 is always a logic 0, bit S_5 indicates the condition of the IF flag bit, and S_4 and S_3 show which segment is accessed during the current bus cycle. See Table 9–4 for the truth table of S_4 and S_3. These two status bits could be used to address four separate 1M byte memory banks by decoding them as A_{21} and A_{20}.

$\overline{RD}$ Whenever the **read signal** is a logic 0, the data bus is receptive to data from the memory or I/O devices connected to the system. This pin floats to its high-impedance state during a hold acknowledge.

READY   The **READY** input is controlled to insert wait states into the timing of the microprocessor. If the READY pin is placed at a logic 0 level, the microprocessor enters into wait states and remains idle. If the READY pin is placed at a logic 1 level, it has no effect on the operation of the microprocessor.

INTR   **Interrupt request** is used to request a hardware interrupt. If INTR is held high when IF = 1, the 8086/8088 enters an interrupt acknowledge cycle ($\overline{INTA}$ becomes active) after the current instruction has completed execution.

$\overline{TEST}$ The **Test** pin is an input that is tested by the WAIT instruction. If $\overline{TEST}$ is a logic 0, the WAIT instruction functions as an NOP and if $\overline{TEST}$ is a logic 1, the WAIT instruction waits for $\overline{TEST}$ to become a logic 0. The $\overline{TEST}$ pin is most often connected to the 8087 numeric coprocessor.

NMI   The **non-maskable interrupt** input is similar to INTR except that the NMI interrupt does not check to see whether the IF flag bit is a logic 1. If NMI is activated, this interrupt input uses interrupt vector 2.

RESET   The **reset** input causes the microprocessor to reset itself if this pin is held high for a minimum of four clocking periods. Whenever the 8086 or 8088 is reset, it begins executing instructions at memory location FFFF0H and disables future interrupts by clearing the IF flag bit.

CLK   The **clock** pin provides the basic timing signal to the microprocessor. The clock signal must have a duty cycle of 33 % (high for one third of the clocking period and low for two thirds) to provide proper internal timing for the 8086/8088.

VCC   This **power supply** input provides a +5.0 V, ±10 % signal to the microprocessor.

GND	The **ground** connection is the return for the power supply. Note that the 8086/8088 microprocessors have two pins labeled GND—both must be connected to ground for proper operation.
MN/$\overline{\text{MX}}$	The **minimum/maximum** mode pin selects either minimum mode or maximum mode operation for the microprocessor. If minimum mode is selected, the MN/$\overline{\text{MX}}$ pin must be connected directly to +5.0 V.
$\overline{\text{BHE}}$ S$_7$	The **bus high enable** pin is used in the 8086 to enable the most-significant data bus bits (D$_{15}$–D$_8$) during a read or a write operation. The state of S$_7$ is always a logic 1.

Minimum Mode Pins.   Minimum mode operation of the 8086/8088 is obtained by connecting the MN/$\overline{\text{MX}}$ pin directly to +5.0 V. Do not connect this pin to +5.0 V through a pull-up register, or it will not function correctly.

IO/$\overline{\text{M}}$ or M/$\overline{\text{IO}}$	The **IO/$\overline{\text{M}}$** (8088) or the **M/$\overline{\text{IO}}$** (8086) pin selects memory or I/O. This pin indicates that the microprocessor address bus contains either a memory address or an I/O port address. This pin is at its high-impedance state during a hold acknowledge.
$\overline{\text{WR}}$	The **write line** is a strobe that indicates that the 8086/8088 is outputting data to a memory or I/O device. During the time that the $\overline{\text{WR}}$ is a logic 0, the data bus contains valid data for memory or I/O. This pin floats to a high-impedance during a hold acknowledge.
$\overline{\text{INTA}}$	The **interrupt acknowledge** signal is a response to the INTR input pin. The $\overline{\text{INTA}}$ pin is normally used to gate the interrupt vector number onto the data bus in response to an interrupt request.
ALE	**Address latch enable** shows that the 8086/8088 address/data bus contains address information. This address can be a memory address or an I/O port number. Note that the ALE signal does not float during a hold acknowledge.
DT/$\overline{\text{R}}$	The **data transmit/receive** signal shows that the microprocessor data bus is transmitting (**DT/$\overline{\text{R}}$** = 1) or receiving (**DT/$\overline{\text{R}}$** = 0) data. This signal is used to enable external data bus buffers.
DEN	**Data bus enable** activates external data bus buffers.
HOLD	The **hold input** requests a direct memory access (DMA). If the HOLD signal is a logic 1, the microprocessor stops executing software and places its address, data, and control bus at the high-impedance state. If the HOLD pin is a logic 0, the microprocessor executes software normally.
HLDA	**Hold acknowledge** indicates that the 8086/8088 has entered the hold state.
$\overline{\text{SS0}}$	The $\overline{\text{SS0}}$ status line is equivalent to the S$_0$ pin in maximum mode operation of the microprocessor. This signal is combined with **IO/$\overline{\text{M}}$** and **DT/$\overline{\text{R}}$** to decode the function of the current bus cycle (see Table 9–5).

Maximum Mode Pins.   In order to achieve maximum mode for use with external coprocessors, connect the MN/$\overline{\text{MX}}$ pin to ground.

$\overline{\text{S2}}$, $\overline{\text{S1}}$, and $\overline{\text{S0}}$	The **status bits** indicate the function of the current bus cycle. These signals are normally decoded by the 8288 bus controller described later in this chapter. Table 9–6 shows the function of these three status bits in the maximum mode.
$\overline{\text{RQ}}$/$\overline{\text{GT1}}$ and $\overline{\text{RQ}}$/$\overline{\text{GT0}}$	The **request/grant** pins request direct memory accesses (DMA) during maximum mode operation. These lines are bidirectional and are used to both request and grant a DMA operation.

TABLE 9–5 Bus cycle status (8088) using $\overline{SS0}$.

$IO/\overline{M}$	$DT/\overline{R}$	$\overline{SS0}$	Function
0	0	0	Interrupt acknowledge
0	0	1	Memory read
0	1	0	Memory write
0	1	1	Halt
1	0	0	Opcode fetch
1	0	1	I/O read
1	1	0	I/O write
1	1	1	Passive

TABLE 9–6 Bus control function generated by the bus controller (8288).

$\overline{S2}$	$\overline{S1}$	$\overline{S0}$	Function
0	0	0	Interrupt acknowledge
0	0	1	I/O read
0	1	0	I/O write
0	1	1	Halt
1	0	0	Opcode fetch
1	0	1	Memory read
1	1	0	Memory write
1	1	1	Passive

TABLE 9–7 Queue status bits.

QS_1	QS_0	Function
0	0	Queue is idle
0	1	First byte of opcode
1	0	Queue is empty
1	1	Subsequent byte of opcode

$\overline{\textbf{LOCK}}$ The **lock** output is used to lock peripherals off the system. This pin is activated by using the LOCK: prefix on any instruction.

QS_1 **and** QS_0 The **queue status** bits show the status of the internal instruction queue. These pins are provided for access by the numeric coprocessor (8087). See Table 9–7 for the operation of the queue status bits.
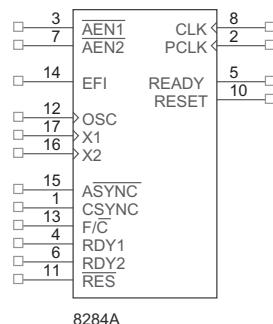
9–2   CLOCK GENERATOR (8284A)

This section describes the clock generator (8284A) and the RESET signal, and introduces the READY signal for the 8086/8088 microprocessors. (The READY signal and its associated circuitry are treated in detail in Section 9–5.)

The 8284A Clock Generator

The 8284A is an ancillary component to the 8086/8088 microprocessors. Without the clock generator, many additional circuits are required to generate the clock (CLK) in an 8086/8088-based system. The 8284A provides the following basic functions or signals: clock generation, RESET synchronization, READY synchronization, and a TTL-level peripheral clock signal. Figure 9–2 illustrates the pin-out of the 8284A clock generator.
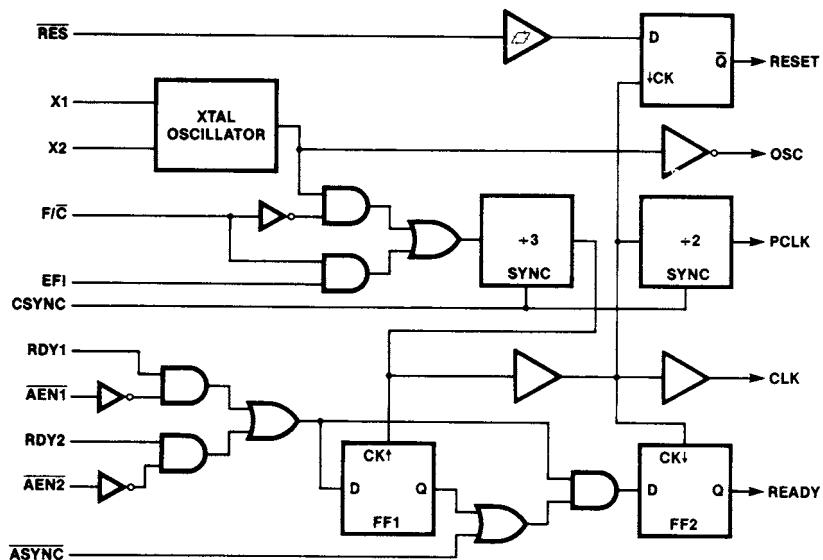
FIGURE 9–2   The pin-out of
the 8284A clock generator.

8284A

Pin Functions.   The 8284A is an 18-pin integrated circuit designed specifically for use with the 8086/8088 microprocessor. The following is a list of each pin and its function.

$\overline{\text{AEN1}}$ **and** $\overline{\text{AEN2}}$	The **address enable** pins are provided to qualify the bus ready signals, RDY1 and RDY2, respectively. Section 9–5 illustrates the use of these two pins, which are used to cause wait states, along with the RDY1 and RDY2 inputs. Wait states are generated by the READY pin of the 8086/8088 microprocessors, which is controlled by these two inputs.
RDY_1 **and** RDY_2	The **bus ready** inputs are provided, in conjunction with the $\overline{\text{AEN1}}$ and $\overline{\text{AEN2}}$ pins, to cause wait states in an 8086/8088-based system.
$\overline{\text{ASYNC}}$	The **ready synchronization** selection input selects either one or two stages of synchronization for the RDY_1 and RDY_2 inputs.
READY	**Ready** is an output pin that connects to the 8086/8088 READY input. This signal is synchronized with the RDY_1 and RDY_2 inputs.
X_1 **and** X_2	The **crystal oscillator** pins connect to an external crystal used as the timing source for the clock generator and all its functions.
$\text{F}/\overline{\text{C}}$	The **frequency/crystal** select input chooses the clocking source for the 8284A. If this pin is held high, an external clock is provided to the EFI input pin; if it is held low, the internal crystal oscillator provides the timing signal. The external frequency input is used when the $\text{F}/\overline{\text{C}}$ pin is pulled high. EFI supplies the timing whenever the $\text{F}/\overline{\text{C}}$ pin is high.
CLK	The **clock output** pin provides the CLK input signal to the 8086/8088 microprocessors and other components in the system. The CLK pin has an output signal that is one third of the crystal or EFI input frequency, and has a 33% duty cycle, which is required by the 8086/8088.
PCLK	The **peripheral clock** signal is one sixth the crystal or EFI input frequency, and has a 50% duty cycle. The PCLK output provides a clock signal to the peripheral equipment in the system.
OSC	The **oscillator output** is a TTL-level signal that is at the same frequency as the crystal or EFI input. The OSC output provides an EFI input to other 8284A clock generators in some multiple-processor systems.
$\overline{\text{RES}}$	The **reset input** is an active-low input to the 8284A. The $\overline{\text{RES}}$ pin is often connected to an RC network that provides power-on resetting.
RESET	The **reset output** is connected to the 8086/8088 RESET input pin.
CSYNC	The **clock synchronization** pin is used whenever the EFI input provides synchronization in systems with multiple processors. If the internal crystal oscillator is used, this pin must be grounded.
GND	The **ground pin** connects to ground.
VCC	This **power supply** pin connects to +5.0 V with a tolerance of ±10%.

FIGURE 9–3 The internal block diagram of the 8284A clock generator.

Operation of the 8284A

The 8284A is a relatively easy component to understand. Figure 9–3 illustrates the internal timing diagram of the 8284A clock generator.

Operation of the Clock Section. The top half of the logic diagram represents the clock and synchronization section of the 8284A clock generator. As the diagram shows, the crystal oscillator has two inputs: X_1 and X_2. If a crystal is attached to X_1 and X_2, the oscillator generates a square-wave signal at the same frequency as the crystal. The square-wave signal is fed to an AND gate and also to an inverting buffer that provides the OSC output signal. The OSC signal is sometimes used as an EFI input to other 8284A circuits in a system.

An inspection of the AND gate reveals that when $F/\overline{C}$ is a logic 0, the oscillator output is steered through to the divide-by-3 counter. If $F/\overline{C}$ is a logic 1, then EFI is steered through to the counter.

The output of the divide-by-3 counter generates the timing for ready synchronization, a signal for another counter (divide-by-2), and the CLK signal to the 8086/8088 microprocessor. The CLK signal is also buffered before it leaves the clock generator. Notice that the output of the first counter feeds the second. These two cascaded counters provide the divide-by-6 output at PCLK, the peripheral clock output.

Figure 9–4 shows how an 8284A is connected to the 8086/8088. Notice that $F/\overline{C}$ and CSYNC are grounded to select the crystal oscillator, and that a 15 MHz crystal provides the normal 5 MHz clock signal to the 8086/8088, as well as a 2.5 MHz peripheral clock signal.

Operation of the Reset Section. The reset section of the 8284A is very simple: It consists of a Schmitt trigger buffer and a single D-type flip-flop circuit. The D-type flip-flop ensures that the timing requirements of the 8086/8088 RESET input are met. This circuit applies the RESET signal to the microprocessor on the negative edge (1-to-0 transition) of each clock. The 8086/8088 microprocessors sample RESET at the positive edge (0-to-1 transition) of the clocks; therefore, this circuit meets the timing requirements of the 8086/8088.

Refer to Figure 9–4. Notice that an RC circuit provides a logic 0 to the $\overline{\text{RES}}$ input pin when power is first applied to the system. After a short time, the $\overline{\text{RES}}$ input becomes a logic 1 because the capacitor charges toward +5.0 V through the resistor. A pushbutton switch allows the microprocessor to be reset by the operator. Correct reset timing requires the RESET input to come a logic 1 no later than four clocks after system power is applied, and to be held high for at
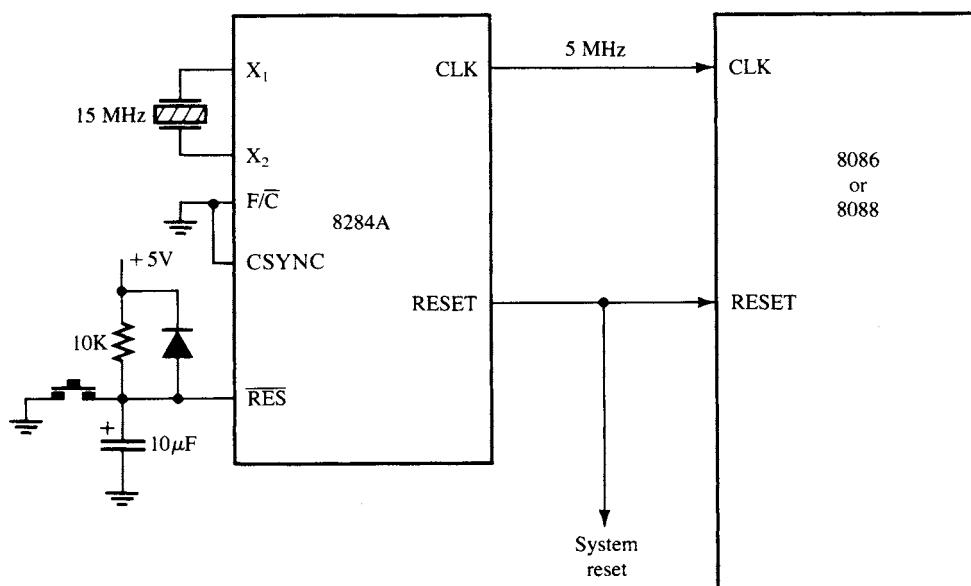
FIGURE 9–4 The clock generator (8284A) and the 8086 and 8088 microprocessors illustrating the connection for the clock and reset signals. A 15 MHz crystal provides the 5 MHz clock for the microprocessor.

least 50 μs. The flip-flop makes certain that RESET goes high in four clocks, and the RC time constant ensures that it stays high for at least 50 μs.

9–3   BUS BUFFERING AND LATCHING

Before the 8086/8088 microprocessors can be used with memory or I/O interfaces, their multiplexed buses must be demultiplexed. This section provides the detail required to demultiplex the buses and illustrates how the buses are buffered for very large systems. (Because the maximum fan-out is 10, the system must be buffered if it contains more than 10 other components.)

Demultiplexing the Buses

The address/data bus on the 8086/8088 is multiplexed (shared) to reduce the number of pins required for the 8086/8088 microprocessor integrated circuit. Unfortunately, this burdens the hardware designer with the task of extracting or demultiplexing information from these multiplexed pins.

Why not leave the buses multiplexed? Memory and I/O require that the address remains valid and stable throughout a read or write cycle. If the buses are multiplexed, the address changes at the memory and I/O, which causes them to read or write data in the wrong locations.

All computer systems have three buses: (1) an address bus that provides the memory and I/O with the memory address or the I/O port number, (2) a data bus that transfers data between the microprocessor and the memory and I/O in the system, and (3) a control bus that provides control signals to the memory and I/O. These buses must be present in order to interface to memory and I/O.

Demultiplexing the 8088.   Figure 9–5 illustrates the 8088 microprocessor and the components required to demultiplex its buses. In this case, two 74LS373 or 74LS573 transparent latches are used to demultiplex the address/data bus connections AD_7–AD_0 and the multiplexed address/status connections A_{19}/S_6–A_{16}/S_3.

These transparent latches, which are like wires whenever the address latch enable pin (ALE) becomes a logic 1, pass the inputs to the outputs. After a short time, ALE returns to its logic 0 condition, which causes the latches to remember the inputs at the time of the change to a
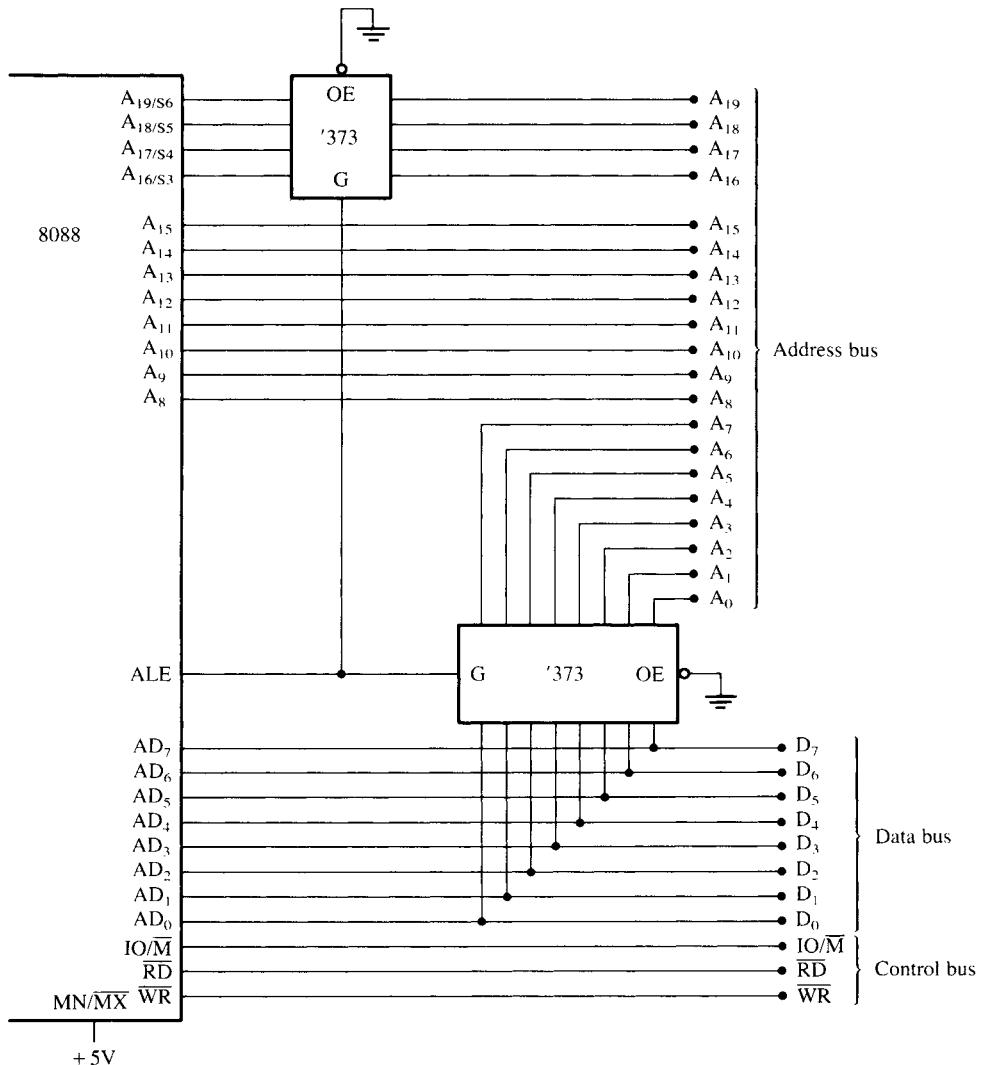
FIGURE 9–5   The 8088 microprocessor shown with a demultiplexed address bus. This is the model used to build many 8088-based systems.

logic 0. In this case, A_7–A_0 are stored in the bottom latch and A_{19}–A_{16} are stored in the top latch. This yields a separate address bus with connections A_{19}–A_0. These address connections allow the 8088 to address 1M byte of memory space. The fact that the data bus is separate allows it to be connected to any 8-bit peripheral device or memory component.

Demultiplexing the 8086.   Like the 8088, the 8086 system requires separate address, data, and control buses. It differs primarily in the number of multiplexed pins. In the 8088, only AD_7–AD_0 and A_{19}/S_6–A_{16}/S_3 are multiplexed. In the 8086, the multiplexed pins include AD_{15}–AD_0 A_{19}/S_6–A_{16}/S_3, and $\overline{BHE}/S_7$. All of these signals must be demultiplexed.

Figure 9–6 illustrates a demultiplexed 8086 with all three buses: address (A_{19}–A_0 and $\overline{BHE}$), data (D_{15}–D_0), and control ($M/\overline{IO}$, $\overline{RD}$, and $\overline{WR}$).

This circuit shown in Figure 9–6 is almost identical to the one pictured in Figure 9–5, except that an additional 74LS373 latch has been added to demultiplex the address/data bus pins AD_{15}–AD_8 and a $\overline{BHE}/S_7$ input has been added to the top 74LS373 to select the high-order memory bank in the 16-bit memory system of the 8086. Here, the memory and I/O system see the
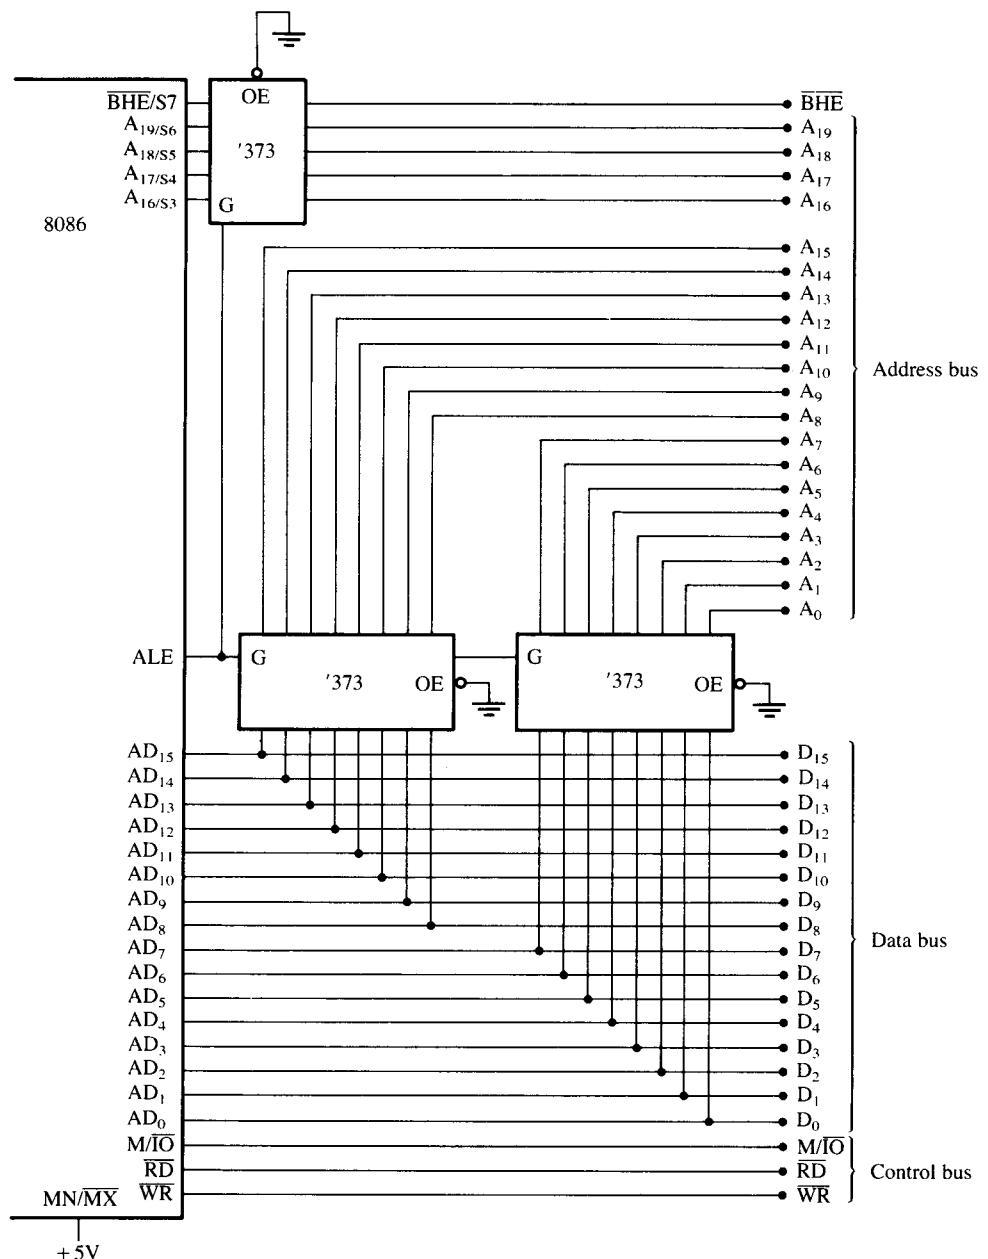
FIGURE 9–6 The 8086 microprocessor shown with a demultiplexed address bus. This is the model used to build many 8086-based systems.

8086 as a device with a 20-bit address bus (A_{19}–A_0), a 16-bit data bus (D_{15}–D_0), and a three-line control bus ($M/\overline{IO}$, $\overline{RD}$, and $\overline{WR}$).

The Buffered System

If more than 10 unit loads are attached to any bus pin, the entire 8086 or 8088 system must be buffered. The demultiplexed pins are already buffered by the 74LS373 or 74LS573 latches, which have been designed to drive the high-capacitance buses encountered in microcomputer

systems. The buffer's output currents have been increased so that more TTL unit loads may be driven: A logic 0 output provides up to 32 mA of sink current, and a logic 1 output provides up to 5.2 mA of source current.

A fully buffered signal will introduce a timing delay to the system. This causes no difficulty unless memory or I/O devices are used, which function at near the maximum speed of the bus. Section 9–4 discusses this problem and the time delays involved in more detail.

The Fully Buffered 8088. Figure 9–7 depicts a fully buffered 8088 microprocessor. Notice that the remaining eight address pins, A_{15}–A_8, use a 74LS244 octal buffer; the eight data bus pins, D_7–D_0, use a 74LS245 octal bidirectional bus buffer; and the control bus signals, $M/\overline{IO}$, $\overline{RD}$,
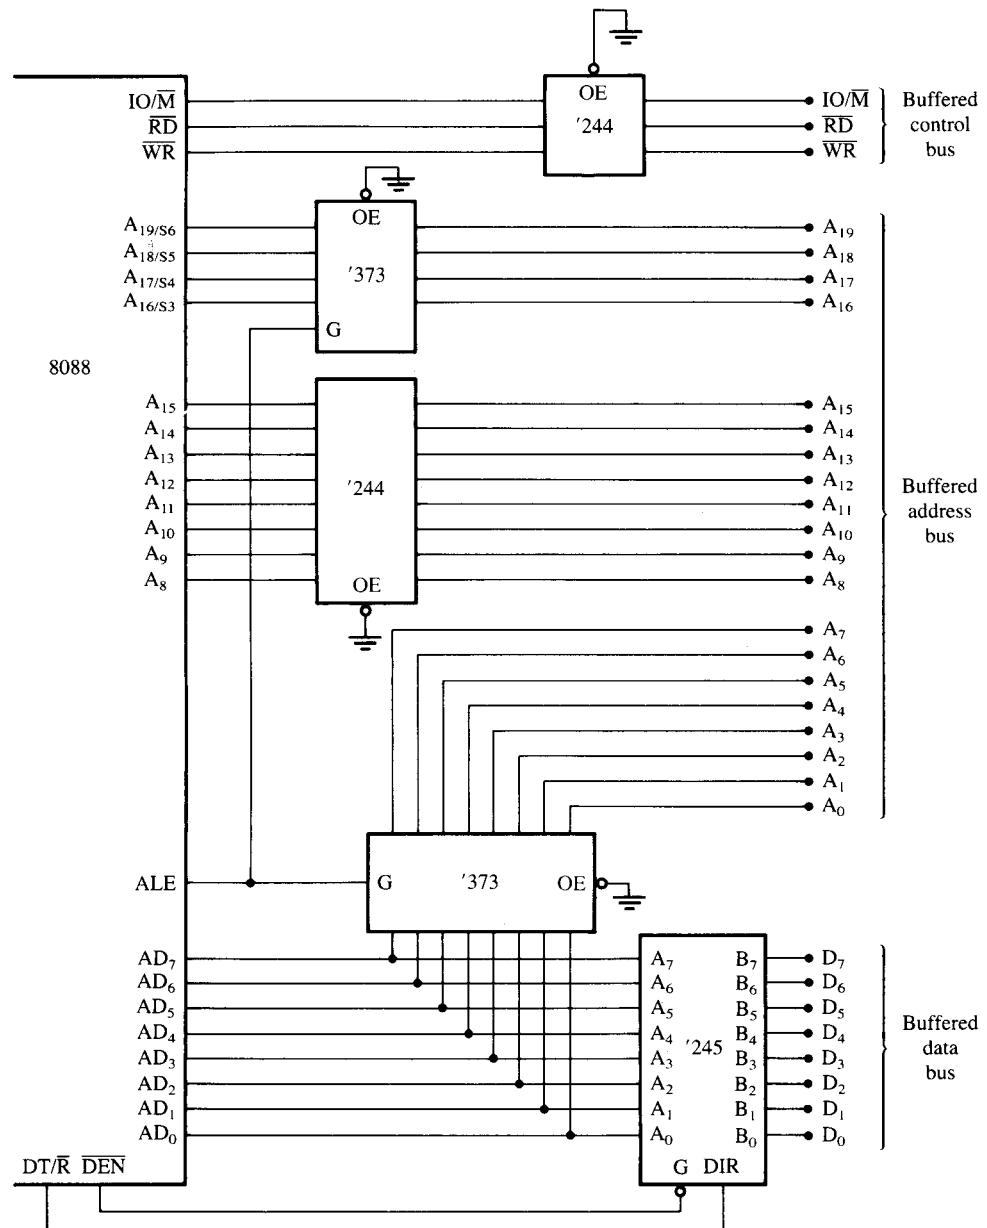
FIGURE 9–7 A fully buffered 8088 microprocessor.

and $\overline{\text{WR}}$, use a 74LS244 buffer. A fully buffered 8088 system requires two 74LS244s, one 74LS245, and two 74LS373s. The direction of the 74LS245 is controlled by the DT/$\overline{\text{R}}$ signal and is enabled and disabled by the $\overline{\text{DEN}}$ signal.

The Fully Buffered 8086.   Figure 9–8 illustrates a fully buffered 8086 microprocessor. Its address pins are already buffered by the 74LS373 address latches; its data bus employs two
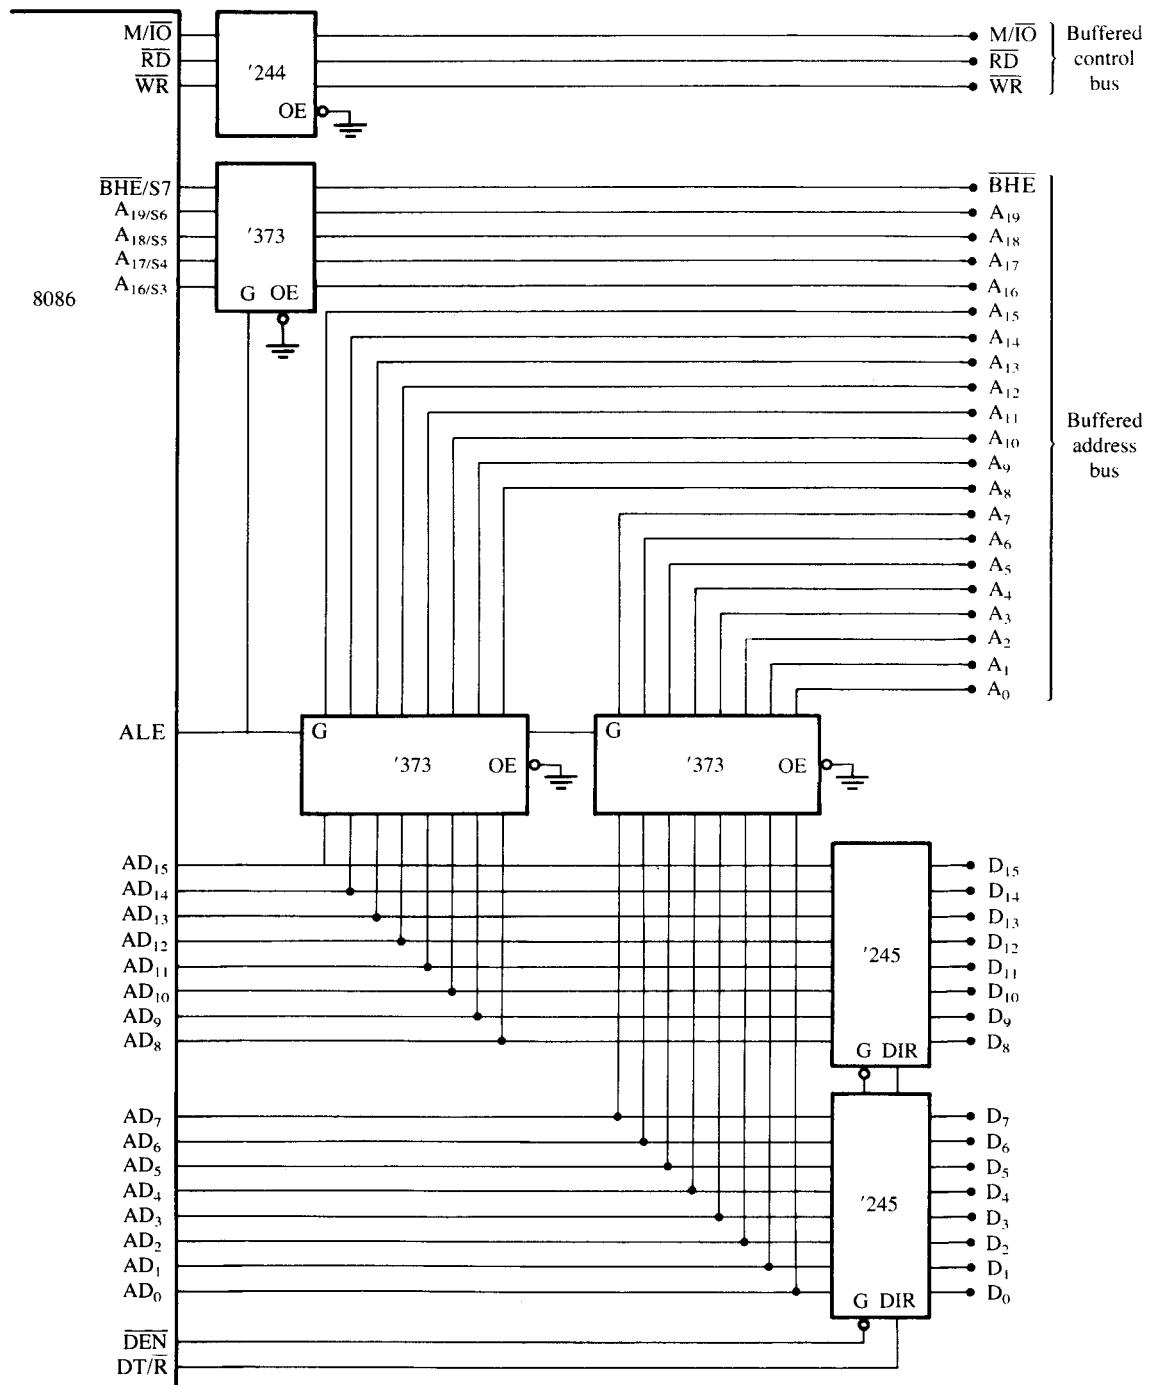
FIGURE 9–8   A fully buffered 8086 microprocessor.

74LS245 octal bidirectional bus buffers; and the control bus signals, M/$\overline{\text{IO}}$, $\overline{\text{RD}}$, and $\overline{\text{WR}}$ use a 74LS244 buffer. A fully buffered 8086 system requires one 74LS244, two 74LS245s, and three 74LS373s. The 8086 requires one more buffer than the 8088 because of the extra eight data bus connections, D_{15}–D_8. It also has a $\overline{\text{BHE}}$ signal that is buffered for memory-bank selection.

9–4   BUS TIMING

It is essential to understand system bus timing before choosing a memory or I/O device for inter-facing to the 8086 or 8088 microprocessors. This section provides insight into the operation of the bus signals and the basic read and write timing of the 8086/8088. It is important to note that we discuss only the times that affect memory and I/O interfacing in this section.

Basic Bus Operation

The three buses of the 8086 and 8088—address, data, and control—function exactly the same way as those of any other microprocessor. If data are written to the memory (see the simplified timing for write in Figure 9–9), the microprocessor outputs the memory address on the address bus, out-puts the data to be written into memory on the data bus, and issues a write ($\overline{\text{WR}}$) to memory and IO/$\overline{\text{M}}$ = 0 for the 8088 and M/$\overline{\text{IO}}$ = 1 for the 8086. If data are read from the memory (see the sim-plified timing for read in Figure 9–10), the microprocessor outputs the memory address on the address bus, issues a read memory signal ($\overline{\text{RD}}$), and accepts the data via the data bus.

Timing in General

The 8086/8088 microprocessors use the memory and I/O in periods called bus cycles. Each bus cycle equals four system-clocking periods (T states). Newer microprocessors divide the bus cycle into as few as two clocking periods. If the clock is operated at 5 MHz (the basic operating frequency for these two microprocessors), one 8086/8088 bus cycle is complete in 800 ns. This means that the microprocessor reads or writes data between itself and memory or I/O at a maxi-mum rate of 1.25 million times a second. (Because of the internal queue, the 8086/8088 can exe-cute 2.5 million instructions per second [MIPS] in bursts.) Other available versions of these microprocessors operate at much higher transfer rates due to higher clock frequencies.
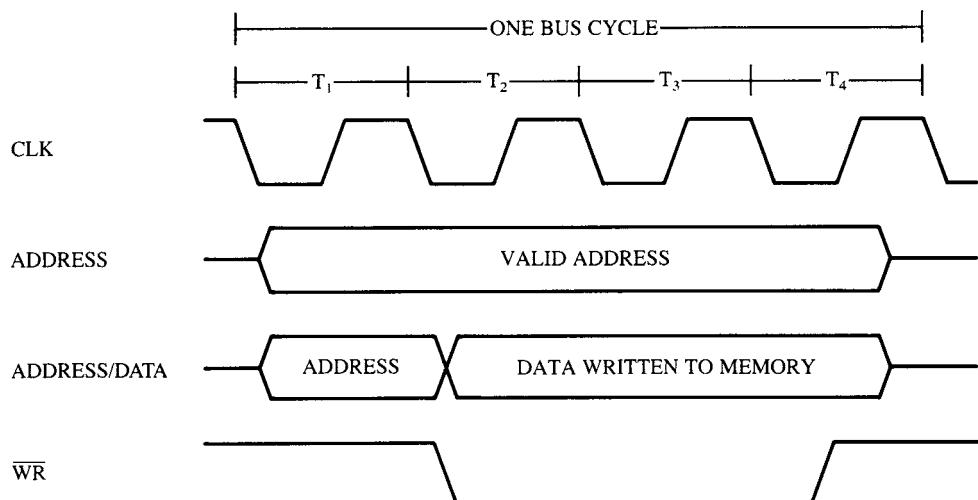
FIGURE 9–9   Simplified 8086/8088 write bus cycle.
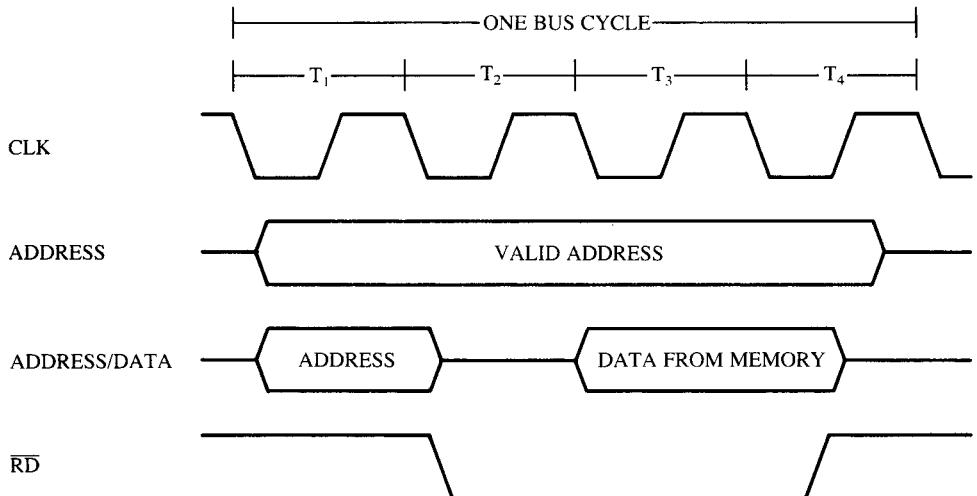
FIGURE 9–10   Simplified 8086/8088 read bus cycle.

During the first clocking period in a bus cycle, which is called T1, many things happen. The address of the memory or I/O location is sent out via the address bus and the address/data bus connections. (The address/data bus is multiplexed and sometimes contains memory-addressing information, sometimes data.) During TI, control signals ALE, DT/$\overline{R}$, and IO/$\overline{M}$ (8088) or M/$\overline{IO}$ (8086) are also output. The IO/$\overline{M}$ or M/$\overline{IO}$ signal indicates whether the address bus contains a memory address or an I/O device (port) number.

During T2, the 8086/8088 microprocessors issue the $\overline{RD}$ or $\overline{WR}$ signal, $\overline{DEN}$, and in the case of a write, the data to be written appear on the data bus. These events cause the memory or I/O device to begin to perform a read or a write. The $\overline{DEN}$ signal turns on the data bus buffers, if they are present in the system, so the memory or I/O can receive data to be written, or so the microprocessor can accept the data read from the memory or I/O for a read operation. If this happens to be a write bus cycle, the data are sent out to the memory or I/O through the data bus.

READY is sampled at the end of T_2, as illustrated in Figure 9–11. If READY is low at this time, T_3 becomes a wait state (T_w). (More detail is provided in Section 9–5.) This clocking period is provided to allow the memory time to access data. If the bus cycle happens to be a read bus cycle, the data bus is sampled at the end of T_3.

In T_4, all bus signals are deactivated in preparation for the next bus cycle. This is also the time when the 8086/8088 samples the data bus connections for data that are read from memory or I/O. In addition, at this point, the trailing edge of the $\overline{WR}$ signal transfers data to the memory or I/O, which activates and writes when the $\overline{WR}$ signal returns to a logic 1 level.

Read Timing

Figure 9–11 also depicts the read timing for the 8088 microprocessor. The 8086 read timing is identical except that the 8086 has 16 rather than eight data bus bits. A close look at this timing diagram should allow you to identify all the main events described for each T state.

The most important item contained in the read timing diagram is the amount of time allowed for the memory or I/O to read the data. Memory is chosen by its access time, which is the fixed amount of time that the microprocessor allows it to access data for the read operation. It is therefore extremely important that the memory chosen complies with the limitations of the system.

The microprocessor timing diagram does not provide a listing for memory access time. Instead, it is necessary to combine several times to arrive at the access time. To find memory
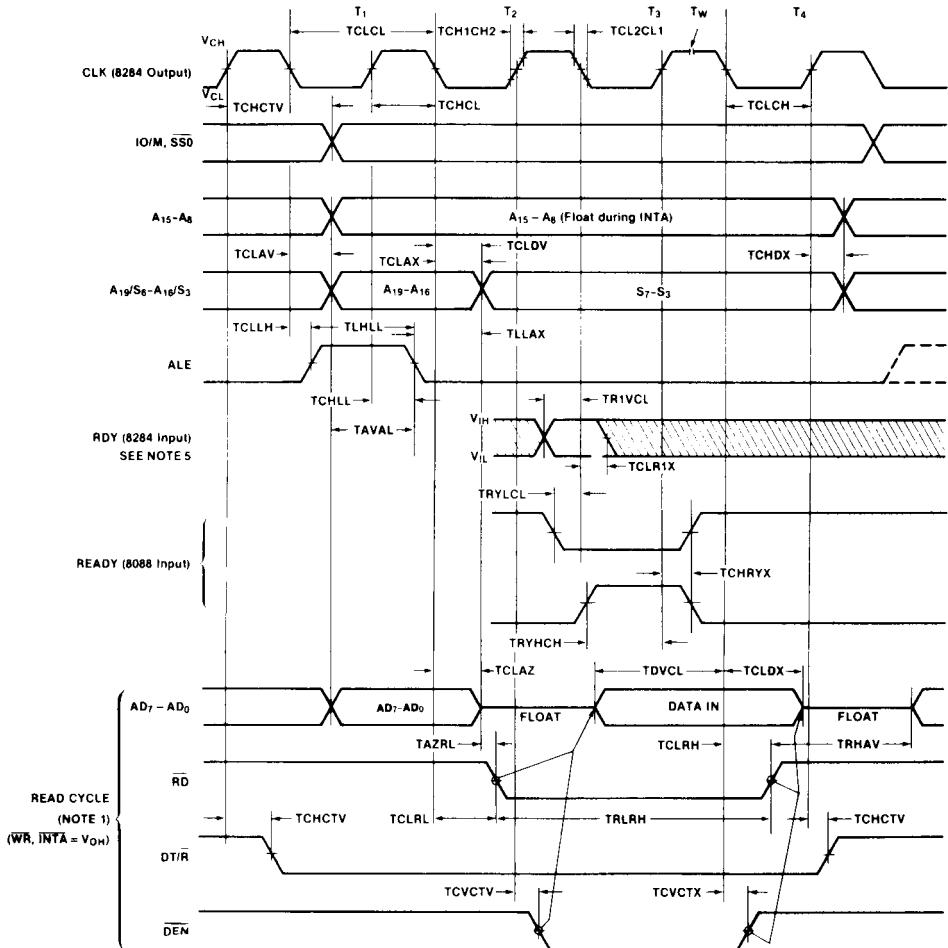
FIGURE 9–11 Minimum mode 8088 bus timing for a read operation.

access time in this diagram, first locate the point in T_3 when data are sampled. If you examine the timing diagram closely, you will notice a line that extends from the end of T_3 down to the data bus. At the end of T_3, the microprocessor samples the data bus.

Memory access time starts when the address appears on the memory address bus and continues until the microprocessor samples the memory data at T_3. Approximately three T states elapse between these times. (See Figure 9–12 for the following times.) The address does not appear until T_{CLAV} time (110 ns if the clock is 5 MHz) after the start of T_1. This means that T_{CLAV} time must be subtracted from the three clocking states (600 ns) that separate the appearance of the address (T_1) and the sampling of the data (T_3). One other time must also be subtracted: the data setup time (T_{DVCL}), which occurs before T_3. Memory access time is thus three clocking states minus the sum of T_{CLAV} and T_{DVCL}. Because T_{DVCL} is 30 ns with a 5 MHz clock, the allowed memory access time is only 460 ns (access time = 600 ns − 110 ns − 30 ns).

The memory devices chosen for connection to the 8086/8088 operating at 5 MHz must be able to access data in less than 460 ns, because of the time delay introduced by the address decoders and buffers in the system. At least a 30- or 40-ns margin should exist for the operation of these circuits. Therefore, the memory speed should be no slower than about 420 ns to operate correctly with the 8086/8088 microprocessors.

FIGURE 9–12 8088 AC characteristics.

A.C. CHARACTERISTICS (8088: T_A = 0°C to 70°C, V_{CC} = 5V ±10%)*
(8088-2: T_A = 0°C to 70°C, V_{CC} = 5V ±5%)

MINIMUM COMPLEXITY SYSTEM TIMING REQUIREMENTS

Symbol	Parameter	8088		8088-2		Units	Test Conditions
		Min.	Max.	Min.	Max.		
TCLCL	CLK Cycle Period	200	500	125	500	ns	
TCLCH	CLK Low Time	118		68		ns	
TCHCL	CLK High Time	69		44		ns	
TCH1CH2	CLK Rise Time		10		10	ns	From 1.0V to 3.5V
TCL2CL1	CLK Fall Time		10		10	ns	From 3.5V to 1.0V
TDVCL	Data in Setup Time	30		20		ns	
TCLDX	Data in Hold Time	10		10		ns	
TR1VCL	RDY Setup Time into 8284 (See Notes 1, 2)	35		35		ns	
TCLR1X	RDY Hold Time into 8284 (See Notes 1, 2)	0		0		ns	
TRYHCH	READY Setup Time into 8088	118		68		ns	
TCHRYX	READY Hold Time into 8088	30		20		ns	
TRYLCL	READY Inactive to CLK (See Note 3)	−8		−8		ns	
THVCH	HOLD Setup Time	35		20		ns	
TINVCH	INTR, NMI, TEST Setup Time (See Note 2)	30		15		ns	
TILIH	Input Rise Time (Except CLK)		20		20	ns	From 0.8V to 2.0V
TIHIL	Input Fall Time (Except CLK)		12		12	ns	From 2.0V to 0.8V

A.C. CHARACTERISTICS (Continued)

TIMING RESPONSES

Symbol	Parameter	8088		8088-2		Units	Test Conditions
		Min.	Max.	Min.	Max.		
TCLAV	Address Valid Delay	10	110	10	60	ns	
TCLAX	Address Hold Time	10		10		ns	
TCLAZ	Address Float Delay	TCLAX	80	TCLAX	50	ns	
TLHLL	ALE Width	TCLCH−20		TCLCH−10		ns	
TCLLH	ALE Active Delay		80		50	ns	
TCHLL	ALE Inactive Delay		85		55	ns	
TLLAX	Address Hold Time to ALE Inactive	TCHCL−10		TCHCL−10		ns	
TCLDV	Data Valid Delay	10	110	10	60	ns	C_L = 20-100 pF for all 8088 Outputs in addition to internal loads
TCHDX	Data Hold Time	10		10		ns	
TWHDX	Data Hold Time After $\overline{WR}$	TCLCH−30		TCLCH−30		ns	
TCVCTV	Control Active Delay 1	10	110	10	70	ns	
TCHCTV	Control Active Delay 2	10	110	10	60	ns	
TCVCTX	Control Inactive Delay	10	110	10	70	ns	
TAZRL	Address Float to READ Active	0		0		ns	
TCLRL	$\overline{RD}$ Active Delay	10	165	10	100	ns	
TCLRH	$\overline{RD}$ Inactive Delay	10	150	10	80	ns	
TRHAV	$\overline{RD}$ Inactive to Next Address Active	TCLCL−45		TCLCL−40		ns	
TCLHAV	HLDA Valid Delay	10	160	10	100	ns	
TRLRH	$\overline{RD}$ Width	2TCLCL−75		2TCLCL−50		ns	
TWLWH	$\overline{WR}$ Width	2TCLCL−60		2TCLCL−40		ns	
TAVAL	Address Valid to ALE Low	TCLCH−60		TCLCH−40		ns	
TOLOH	Output Rise Time		20		20	ns	From 0.8V to 2.0V
TOHOL	Output Fall Time		12		12	ns	From 2.0V to 0.8V

The only other timing factor that may affect memory operation is the width of the $\overline{RD}$ strobe. On the timing diagram, the read strobe is given as T_{RLRH}. The time for this strobe is 325 ns (5 MHz clock rate), which is wide enough for almost all memory devices manufactured with an access time of 400 ns or less.

Write Timing

Figure 9–13 illustrates the write-timing diagram for the 8088 microprocessor. (Again, the 8086 is nearly identical, so it need not be presented here in a separate timing diagram.)

The main differences between read and write timing are minimal. The $\overline{RD}$ strobe is replaced by the $\overline{WR}$ strobe, the data bus contains information for the memory rather than information from the memory, and $DT/\overline{R}$ remains a logic 1 instead of a logic 0 throughout the bus cycle.

When interfacing some memory devices, timing may be especially critical between the point at which $\overline{WR}$ becomes a logic 1 and the time when the data are removed from the data bus. This is the case because, as you will recall, memory data are written at the trailing edge of the $\overline{WR}$ strobe. According to the timing diagram, this critical period is T_{WHDX} or 88 ns when the 8088 is operated with a 5 MHz clock. Hold time is often much less than this; it is, in fact, often 0 ns for memory devices. The width of the $\overline{WR}$ strobe is T_{WLWH} or 340 ns at a 5 MHz clock rate. This rate is compatible with most memory devices that have an access time of 400 ns or less.
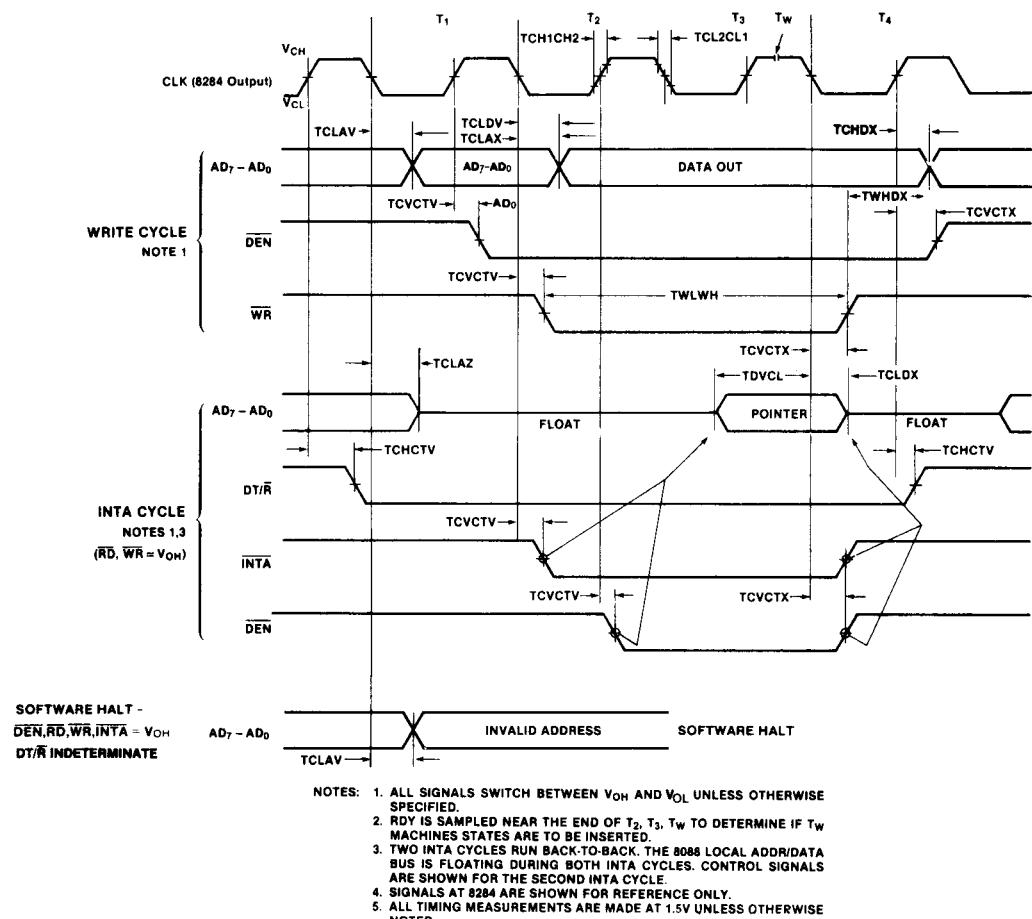
FIGURE 9–13   Minimum mode 8088 write bus timing.

9–5   READY AND THE WAIT STATE

As we mentioned earlier in this chapter, the READY input causes wait states for slower memory and I/O components. A wait state (T_w) is an extra clocking period, inserted between T_2 and T_3 to lengthen the bus cycle. If one wait state is inserted, then the memory access time, normally 460 ns with a 5 MHz clock, is lengthened by one clocking period (200 ns) to 660 ns.

In this section, we discuss the READY synchronization circuitry inside the 8284A clock generator, show how to insert one or more wait states selectively into the bus cycle, and examine the READY input and the synchronization times it requires.

The READY Input

The READY input is sampled at the end of T_2 and again, if applicable, in the middle of T_w. If READY is a logic 0 at the end of T_2, T_3 is delayed and T_w is inserted between T_2 and T_3. READY is next sampled at the middle of T_w to determine whether the next state is T_w or T_3. It is tested for a logic 0 on the 1-to-0 transition of the clock at the end of T_2, and for a 1 on the 0-to-1 transition of the clock in the middle of T_w.

The READY input to the 8086/8088 has some stringent timing requirements. The timing diagram in Figure 9–14 shows READY causing one wait state (T_w), along with the required setup and hold times from the system clock. The timing requirement for this operation is met by the internal READY synchronization circuitry of the 8284A clock generator. When the 8284A is used for READY, the RDY (ready input to the 8284A) input occurs at the end of each T state.

RDY and the 8284A

RDY is the synchronized ready input to the 8284A clock generator. The timing diagram for this input is provided in Figure 9–15. Although it differs from the timing for the READY input to the
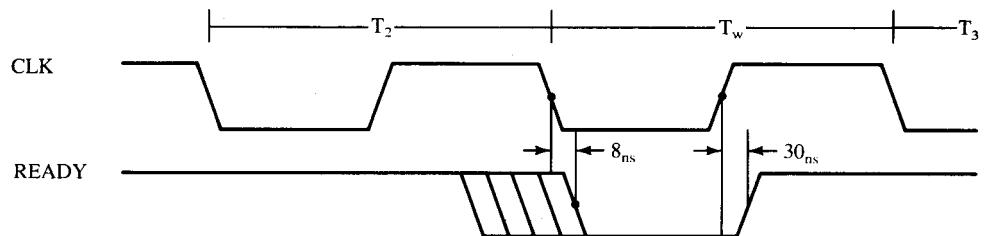
FIGURE 9–14   8086/8088 READY input timing.
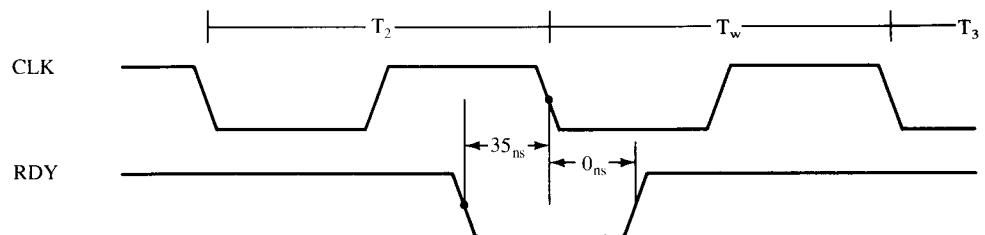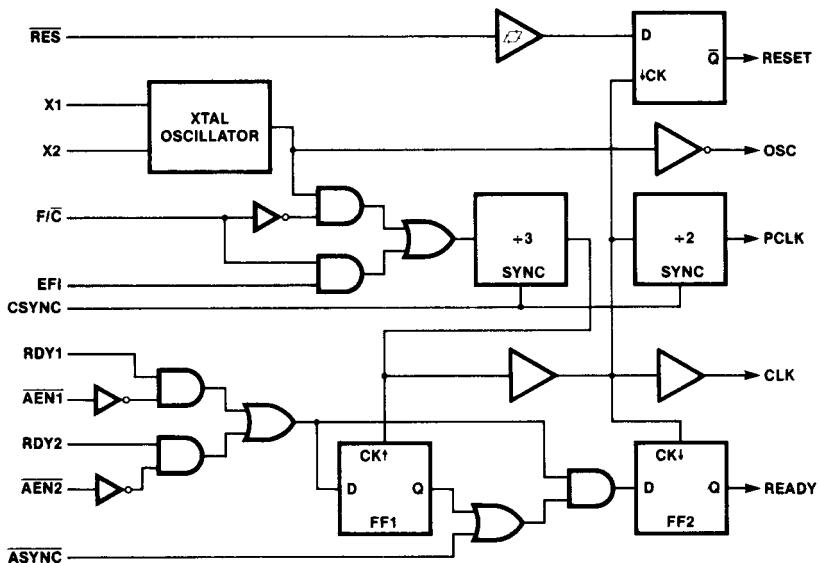
FIGURE 9–15   8284A RDY input timing.

FIGURE 9–16   The internal block diagram of the 8284A clock generator. (Courtesy of Intel Corporation.)

8086/8088, the internal 8284A circuitry guarantees the accuracy of the READY synchronization provided to the 8086/8088 microprocessors.

Figure 9–16 again depicts the internal structure of the 8284A. The bottom half of this diagram is the READY synchronization circuitry. At the leftmost side, the RDY_1 and $\overline{AEN1}$ inputs are ANDed, as are the RDY_2 and $\overline{AEN2}$ inputs. The outputs of the AND gates are then ORed to generate the input to the one or two stages of synchronization. In order to obtain a logic 1 at the inputs to the flip-flops, RDY1 ANDed with $\overline{AEN1}$ must be active or RDY2 ANDed with $\overline{AEN2}$ must be active.

The ASYNC input selects one stage of synchronization when it is a logic 1 and two stages when it is a logic 0. If one stage is selected, then the RDY signal is kept from reaching the 8086/8088 READY pin until the next negative edge of the clock. If two stages are selected, the first positive edge of the clock captures RDY in the first flip-flop. The output of this flip-flop is fed to the second flip-flop, so on the next negative edge of the clock, the second flip-flop captures RDY.

Figure 9–17 illustrates a circuit used to introduce almost any number of wait states for the 8086/8088 microprocessors. Here, an 8-bit serial shift register (74LS164) shifts a logic 0 for one or more clock periods from one of its Q outputs through to the RDY1 input of the 8284A. With appropriate strapping, this circuit can provide various numbers of wait states. Notice also how the shift register is cleared back to its starting point. The output of the register is forced high when the $\overline{RD}$, $\overline{WR}$, and $\overline{INTA}$ pins are all logic 1s. These three signals are high until state T_2, so the shift register shifts for the first time when the positive edge of the T_2 arrives. If one wait is desired, output Q_B is connected to the OR gate. If two waits are desired, output Q_C is connected, and so forth.

Notice in Figure 9–17 that this circuit does not always generate wait states. It is enabled from the memory only for memory devices that require the insertion of waits. If the selection signal from a memory device is a logic 0, the device is selected; then this circuit will generate a wait state.

Figure 9–18 illustrates the timing diagram for this shift register wait state generator when it is wired to insert one wait state. The timing diagram also illustrates the internal contents of the shift register's flip-flops to present a more detailed view of its operation. In this example, one wait state is generated.
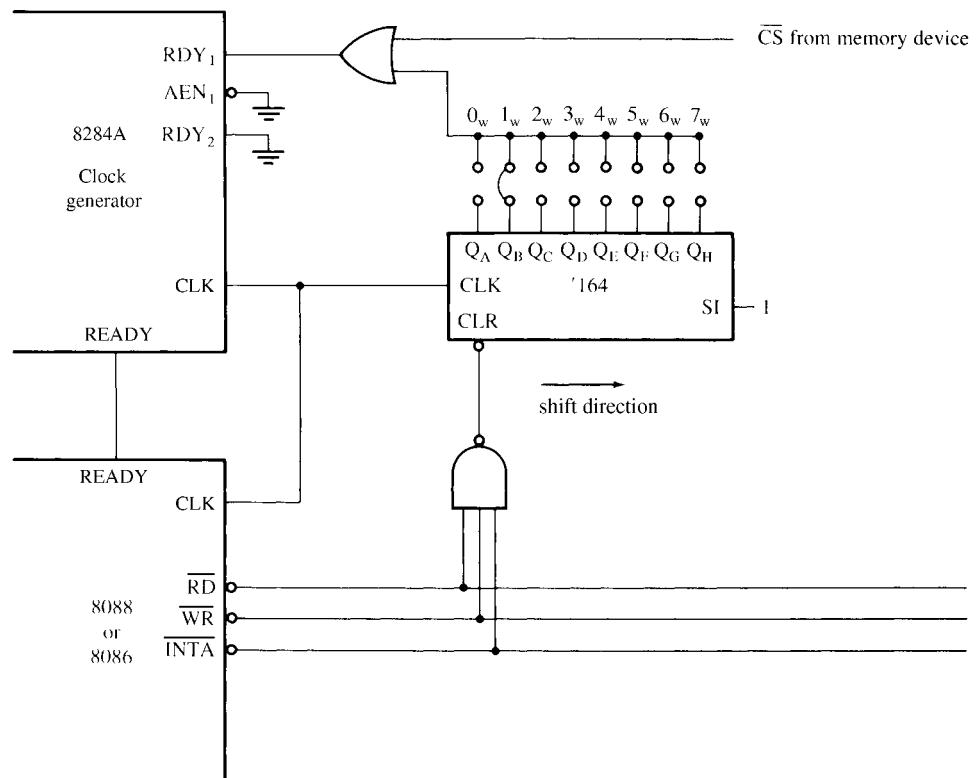
FIGURE 9–17   A circuit that will cause between 0 and 7 wait states.
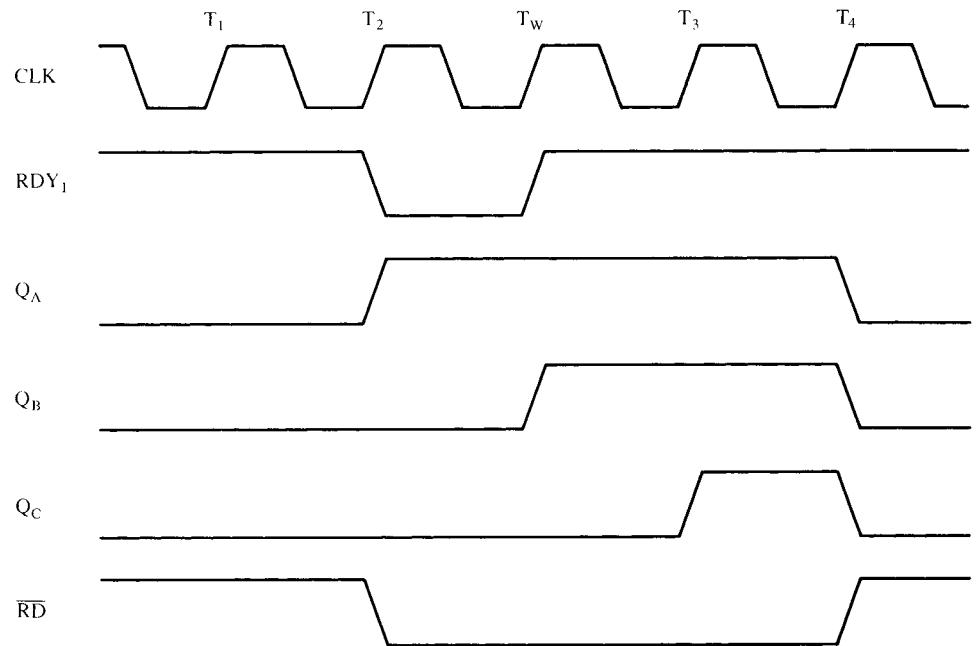
FIGURE 9–18   Wait state generation timing of the circuit of Figure 9–17.

9–6   ## MINIMUM MODE VERSUS MAXIMUM MODE

There are two available modes of operation for the 8086/8088 microprocessors: minimum mode and maximum mode. Minimum mode operation is obtained by connecting the mode selection pin MN/$\overline{\text{MX}}$ to +5.0 V, and maximum mode is selected by grounding this pin. Both modes enable different control structures for the 8086/8088 microprocessors. The mode of operation provided by minimum mode is similar to that of the 8085A, the most recent Intel 8-bit microprocessor. The maximum mode is unique and designed to be used whenever a coprocessor exists in a system. Note that the maximum mode was dropped from the Intel family beginning with the 80286 microprocessor.

Minimum Mode Operation

Minimum mode operation is the least expensive way to operate the 8086/8088 microprocessors (see Figure 9–19 for the minimum mode 8088 system). It costs less because all the control signals for the memory and I/O are generated by the microprocessor. These control signals are identical to those of the Intel 8085A, an earlier 8-bit microprocessor. The minimum mode allows the 8085A 8-bit peripherals to be used with the 8086/8088 without any special considerations.

Maximum Mode Operation

Maximum mode operation differs from minimum mode in that some of the control signals must be externally generated. This requires the addition of an external bus controller—the 8288 bus controller (see Figure 9–20 for the maximum mode 8088 system). There are not enough pins on the 8086/8088 for bus control during maximum mode because new pins and new features have replaced some of them. Maximum mode is used only when the system contains external coprocessors such as the 8087 arithmetic coprocessor.
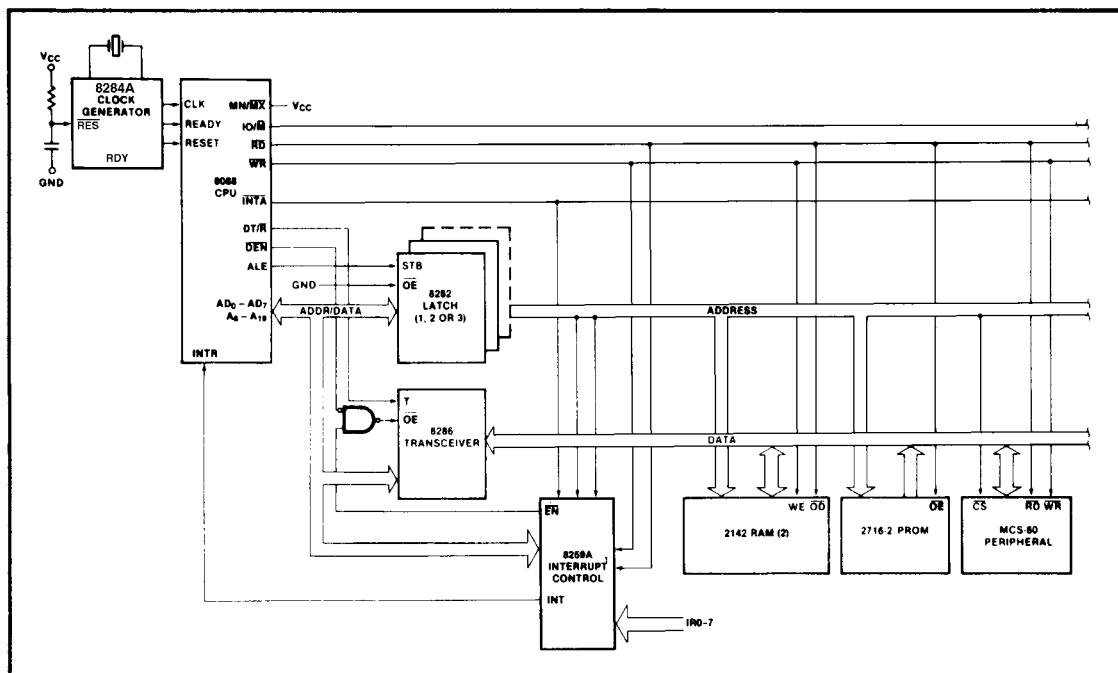
FIGURE 9–19   Minimum mode 8088 system.
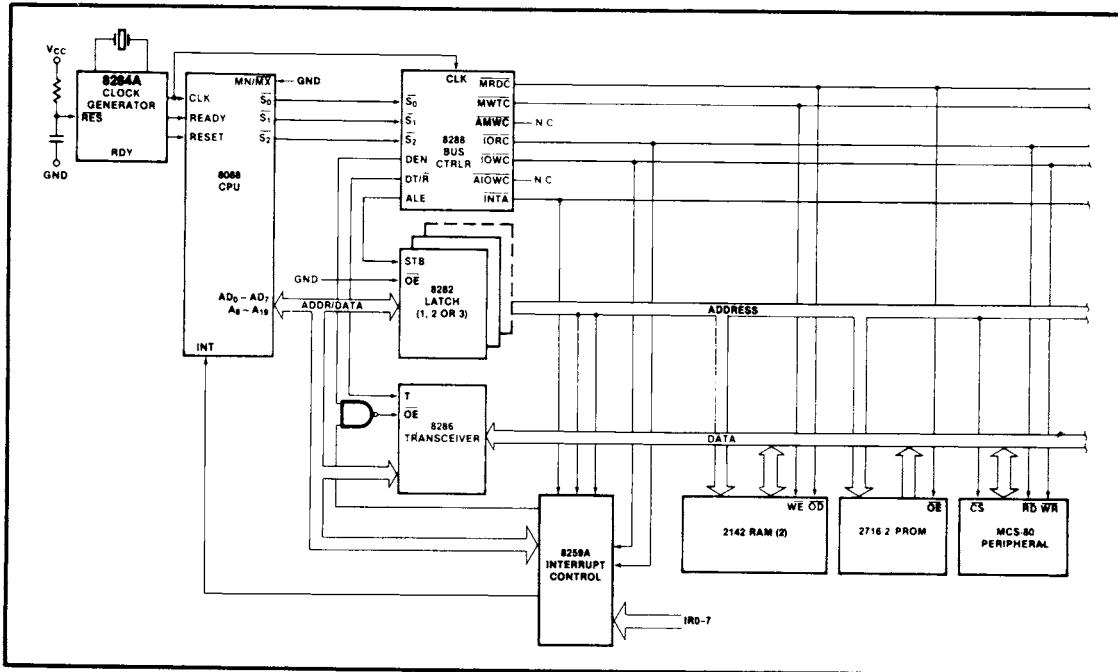
FIGURE 9–20   Maximum mode 8088 system.

The 8288 Bus Controller

An 8086/8088 system that is operated in maximum mode must have an 8288 bus controller to provide the signals eliminated from the 8086/8088 by the maximum mode operation. Figure 9–21 illustrates the block diagram and pin-out of the 8288 bus controller.

Notice that the control bus developed by the 8288 bus controller contains separate signals for I/O ($\overline{\text{IORC}}$ and $\overline{\text{IOWC}}$) and memory ($\overline{\text{MRDC}}$ and $\overline{\text{MWTC}}$). It also contains advanced memory ($\overline{\text{AMWC}}$) and I/O ($\overline{\text{AIOWC}}$) write strobes, and the $\overline{\text{INTA}}$ signal. These signals replace the minimum mode ALE, $\overline{\text{WR}}$, IO/$\overline{\text{M}}$, DT/$\overline{\text{R}}$, $\overline{\text{DEN}}$, and $\overline{\text{INTA}}$, which are lost when the 8086/8088 microprocessors are operated in the maximum mode.
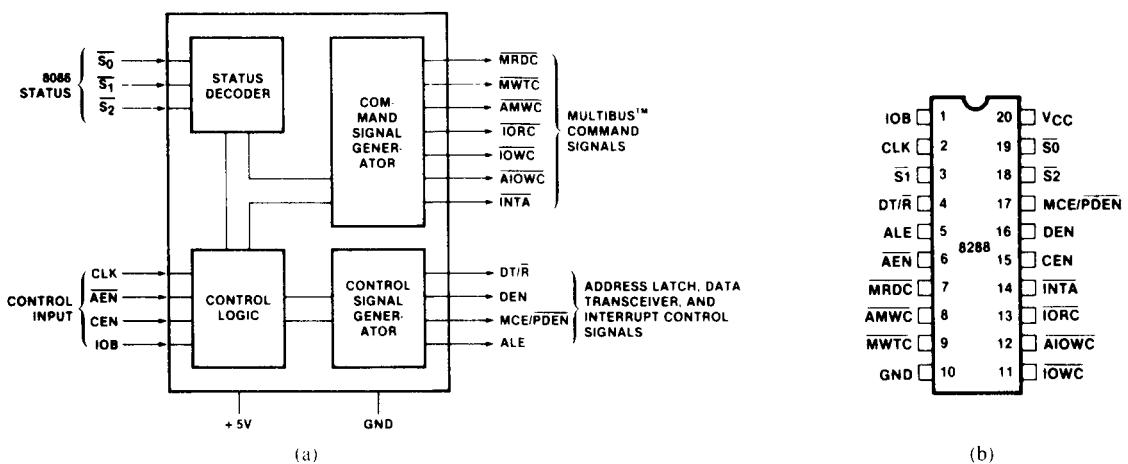
FIGURE 9–21   The 8288 bus controller; (a) block diagram and (b) pin-out.

Pin Functions

The following list provides a description of each pin of the 8288 bus controller.

S_2, S_1, and S_0	**Status inputs** are connected to the status output pins on the 8086/8088 microprocessor. These three signals are decoded to generate the timing signals for the system.
CLK	The **clock** input provides internal timing and must be connected to the CLK output pin of the 8284A clock generator.
ALE	The **address latch enable** output is used to demultiplex the address/data bus.
DEN	The **data bus enable** pin controls the bidirectional data bus buffers in the system. Note that this is an active high output pin that is the opposite polarity from the $\overline{\text{DEN}}$ signal found on the microprocessor when operated in the minimum mode.
DT/$\overline{\text{R}}$	The **data transmit/receive** signal is output by the 8288 to control the direction of the bidirectional data bus buffers.
$\overline{\text{AEN}}$	The **address enable** input causes the 8288 to enable the memory control signals.
CEN	The **control enable** input enables the command output pins on the 8288.
IOB	The **I/O bus mode** input selects either the I/O bus mode or system bus mode operation.
$\overline{\text{AIOWC}}$	The **advanced I/O write** is a command output used to provide I/O with an advanced I/O write control signal.
$\overline{\text{IORC}}$	The **I/O read command** output provides I/O with its read control signal.
$\overline{\text{IOWC}}$	The **I/O write command** output provides I/O with its main write signal.
$\overline{\text{AMWT}}$	The **advanced memory write** control pin provides memory with an early or advanced write signal.
$\overline{\text{MWTC}}$	The **memory write** control pin provides memory with its normal write control signal.
$\overline{\text{MRDC}}$	The **memory read** control pin provides memory with a read control signal.
$\overline{\text{INTA}}$	The **interrupt acknowledge** output acknowledges an interrupt request input applied to the INTR pin.
MCE/$\overline{\text{PDEN}}$	The **master cascade/peripheral data** output selects cascade operation for an interrupt controller if IOB is grounded, and enables the I/O bus transceivers if IOB is tied high.

9–7   SUMMARY

1. The main differences between the 8086 and 8088 are (1) an 8-bit data bus on the 8088 and a 16-bit data bus on the 8086, (2) an $\overline{\text{SS0}}$ pin on the 8088 in place of $\overline{\text{BHE}}$ /S7 on the 8086, and (3) an IO/$\overline{\text{M}}$ pin on the 8088 instead of an M/$\overline{\text{IO}}$ on the 8086.
2. Both the 8086 and 8088 require a single +5.0 V power supply with a tolerance of ±10%.
3. The 8086/8088 microprocessors are TTL-compatible if the noise immunity figure is derated to 350 mV from the customary 400 mV.
4. The 8086/8088 microprocessors can drive one 74XX, five 74LSXX, one 74SXX, ten 74ALSXX, and ten 74HCXX unit loads.

5. The 8284A clock generator provides the system clock (CLK), READY synchronization, and RESET synchronization.

6. The standard 5 MHz 8086/8088 operating frequency is obtained by attaching a 15 MHz crystal to the 8284A clock generator. The PCLK output contains a TTL-compatible signal at one half the CLK frequency.

7. Whenever the 8086/8088 microprocessors are reset, they begin executing software at memory location FFFF0H (FFFF:0000) with the interrupt request pin disabled.

8. Because the 8086/8088 buses are multiplexed and most memory and I/O devices aren't, the system must be demultiplexed before interfacing with memory or I/O. Demultiplexing is accomplished by an 8-bit latch whose clock pulse is obtained from the ALE signal.

9. In a large system, the buses must be buffered because the 8086/8088 microprocessors are capable of driving only 10 unit loads, and large systems often have many more.

10. Bus timing is very important to the remaining chapters in the text. A bus cycle that consists of four clocking periods acts as the basic system timing. Each bus cycle is able to read or write data between the microprocessor and the memory or I/O system.

11. A bus cycle is broken into four states, or T periods: T_1 is used by the microprocessor to send the address to the memory or I/O and the ALE signal to the demultiplexers; T_2 is used to send data to memory for a write and to test the READY pin and activate control signals $\overline{\text{RD}}$ or $\overline{\text{WR}}$; T_3 allows the memory time to access data and allows data to be transferred between the microprocessor and the memory or I/O; and T_4 is where data are written.

12. The 8086/8088 microprocessors allow the memory and I/O 460 ns to access data when they are operated with a 5 MHz clock.

13. Wait states (T_w) stretch the bus cycle by one or more clocking periods to allow the memory and I/O additional access time. Wait states are inserted by controlling the READY input to the 8086/8088. READY is sampled at the end of T_2 and during T_w.

14. Minimum mode operation is similar to that of the Intel 8085A microprocessor, whereas maximum mode operation is new and specifically designed for the operation of the 8087 arithmetic coprocessor.

15. The 8288 bus controller must be used in the maximum mode to provide the control bus signals to the memory and I/O. This is because the maximum mode operation of the 8086/8088 removes some of the system's control signal lines in favor of control signals for the coprocessors. The 8288 reconstructs these removed control signals.

9–8   QUESTIONS AND PROBLEMS

1. List the differences between the 8086 and the 8088 microprocessors.
2. Is the 8086/8088 TTL-compatible? Explain your answer.
3. What is the fan-out from the 8086/8088 to the following devices:
   (a) 74XXX TTL
   (b) 74ALSXXX TTL
   (c) 74HCXXX CMOS
4. What information appears on the address/data bus of the 8088 while ALE is active?
5. What are the purposes of status bits S_3 and S_4?
6. What condition does a logic 0 on the 8086/8088 $\overline{\text{RD}}$ pin indicate?
7. Explain the operation of the $\overline{\text{TEST}}$ pin and the WAIT instruction.
8. Describe the signal that is applied to the CLK input pin of the 8086/8088 microprocessors.
9. What mode of operation is selected when MN/$\overline{\text{MX}}$ is grounded?
10. What does the $\overline{\text{WR}}$ strobe signal from the 8086/8088 indicate about the operation of the 8086/8088?

11. When does ALE float to its high-impedance state?
12. When DT/$\overline{\text{R}}$ is a logic 1, what condition does it indicate about the operation of the 8086/8088?
13. What happens when the HOLD input to the 8086/8088 is placed at its logic 1 level?
14. What three minimum mode 8086/8088 pins are decoded to discover whether the processor is halted?
15. Explain the operation of the $\overline{\text{LOCK}}$ pin.
16. What conditions do the QS_1 and QS_0 pins indicate about the 8086/8088?
17. What three housekeeping chores are provided by the 8284A clock generator?
18. By what factor does the 8284A clock generator divide the crystal oscillator's output frequency?
19. If the F/$\overline{\text{C}}$ pin is placed at a logic 1 level, the crystal oscillator is disabled. Where is the timing input signal attached to the 8284A under this condition?
20. The PCLK output of the 8284A is _____ MHz if the crystal oscillator is operating at 14 MHz.
21. The $\overline{\text{RES}}$ input to the 8284A is placed at a logic _____ level in order to reset the 8086/8088.
22. Which bus connections on the 8086 microprocessor are typically demultiplexed?
23. Which bus connections on the 8088 microprocessor are typically demultiplexed?
24. Which TTL-integrated circuit is often used to demultiplex the buses on the 8086/8088?
25. What is the purpose of the demultiplexed $\overline{\text{BHE}}$ signal on the 8086 microprocessor?
26. Why are buffers often required in an 8086/8088-based system?
27. What 8086/8088 signal is used to select the direction of the data flows through the 74LS245 bidirectional bus buffer?
28. A bus cycle is equal to clocking _____ periods.
29. If the CLK input to the 8086/8088 is 4 MHz, how long is one bus cycle?
30. What two 8086/8088 operations occur during a bus cycle?
31. How many MIPS is the 8086/8088 capable of obtaining when operated with a 10 MHz clock?
32. Briefly describe the purpose of each T state listed:
    (a) T_1
    (b) T_2
    (c) T_3
    (d) T_4
    (e) T_w
33. How much time is allowed for memory access when the 8086/8088 is operated with a 5 MHz clock?
34. How wide is $\overline{\text{DEN}}$ if the 8088 is operated with a 5 MHz clock?
35. If the READY pin is grounded, it will introduce_____ states into the bus cycle of the 8086/8088.
36. What does the $\overline{\text{ASYNC}}$ input to the 8284A accomplish?
37. What logic levels must be applied to $\overline{\text{AEN1}}$ and RDY_1 to obtain a logic 1 at the READY pin? (Assume that $\overline{\text{AEN2}}$ is at a logic 1 level.)
38. Contrast minimum and maximum mode 8086/8088 operation.
39. What main function is provided by the 8288 bus controller when used with 8086/8088 maximum mode operation?

CHAPTER 10

Memory Interface

INTRODUCTION

Whether simple or complex, every microprocessor-based system has a memory system. The Intel family of microprocessors is no different from any other in this respect. Almost all systems contain two main types of memory: **read-only memory** (ROM) and **random access memory** (RAM) or read/write memory. Read-only memory contains system software and permanent system data, while RAM contains temporary data and application software. This chapter explains how to interface both memory types to the Intel family of microprocessors. We demonstrate memory interface to an 8-, 16-, 32-, and 64-bit data bus by using various memory address sizes. This allows virtually any microprocessor to be interfaced to any memory system.

CHAPTER OBJECTIVES

Upon completion of this chapter, you will be able to:

1. Decode the memory address and use the outputs of the decoder to select various memory components.
2. Use programmable logic devices (PLDs) to decode memory addresses.
3. Explain how to interface both RAM and ROM to a microprocessor.
4. Explain how error correction code (ECC) is used with memory.
5. Interface memory to an 8-, 16-, 32-, and 64-bit data bus.
6. Explain the operation of a dynamic RAM controller.
7. Interface dynamic RAM to the microprocessor.

10–1   **MEMORY DEVICES**

Before attempting to interface memory to the microprocessor, it is essential to completely understand the operation of memory components. In this section, we explain the functions of the four common types of memory: read-only memory (ROM), flash memory (EEPROM), static random access memory (SRAM), and dynamic random access memory (DRAM).

Memory Pin Connections

Pin connections common to all memory devices are the address inputs, data outputs or input/
outputs, some type of selection input, and at least one control input used to select a read or write
operation. See Figure 10–1 for ROM and RAM generic-memory devices.

Address Connections.   All memory devices have address inputs that select a memory location
within the memory device. Address inputs are almost always labeled from A_0, the least signifi-
cant address input, to A_n where subscript n can be any value but is always labeled as one less
than the total number of address pins. For example, a memory device with 10 address pins has its
address pins labeled from A_0 to A_9. The number of address pins found on a memory device is
determined by the number of memory locations found within it.
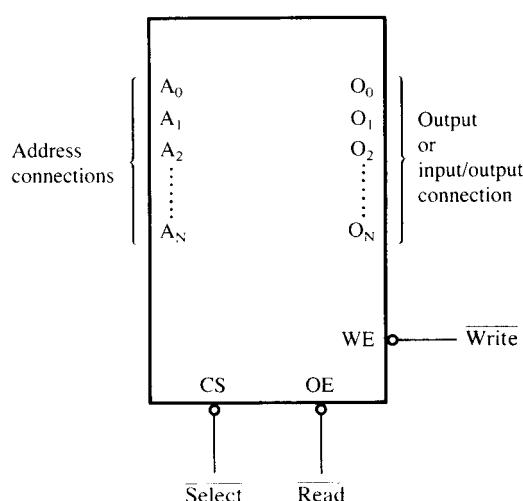
Today, the more common memory devices have between 1K (1024) to 1G (1,073,741,824)
memory locations, with 4G and larger memory location devices on the horizon. A 1K memory
device has 10 address pins (A_0–A_9); therefore, 10 address inputs are required to select any of its
1024 memory locations. It takes a 10-bit binary number (1024 different combinations) to select
any single location on a 1024-location device. If a memory device has 11 address connections
(A_0–A_{11}), it has 2048 (2K) internal memory locations. The number of memory locations can
thus be extrapolated from the number of address pins. For example, a 4K memory device has
12 address connections, an 8K device has 13, and so forth. A device that contains 1M locations
requires a 20-bit address (A_0–A_{19}).

The number 400H represents a 1K-byte section of the memory system. If a memory device is
decoded to begin at memory address 10000H and it is a 1K device, its last location is at address
103FFH—one location less than 400H. Another important hexadecimal number to remember is
1000H, because 1000H is 4K. A memory device that contains a starting address of 14000H that is 4K
bytes long ends at location 14FFFH—one location less than 1000H. A third number is 64K, or
10000H. A memory that starts at location 30000H and ends at location 3FFFFH is a 64K-byte mem-
ory. Finally, because 1M of memory is common, a 1M memory contains 100000H memory locations.

Data Connections.   All memory devices have a set of data outputs or input/outputs. The device
illustrated in Figure 10–1 has a common set of input/output (I/O) connections. Today, many
memory devices have bidirectional common I/O pins.

The data connections are the points at which data are entered for storage or extracted for
reading. Data pins on memory devices are labeled D_0 through D_7 for an 8-bit-wide memory

FIGURE 10–1   A pseudo-
memory component illustrat-
ing the address, data, and
control connections.

device. In this sample memory device there are 8 I/O connections, meaning that the memory device stores 8 bits of data in each of its memory locations. An 8-bit-wide memory device is often called a **byte-wide** memory. Although most devices are currently 8 bits wide, some devices are 16 bits, 4 bits, or just 1 bit wide.

Catalog listings of memory devices often refer to memory locations times bits per location. For example, a memory device with 1K memory locations and 8 bits in each location is often listed as a $1K \times 8$ by the manufacturer. A $16K \times 1$ is a memory device containing 16K 1-bit memory locations. Memory devices are often classified according to total bit capacity. For example, a $1K \times 8$-bit memory device is sometimes listed as an 8K memory device, or a $64K \times 4$ memory is listed as a 256K device. These variations occur from one manufacturer to another.

Selection Connections.   Each memory device has an input—sometimes more than one—that selects or enables the memory device. This type of input is most often called a **chip select** ($\overline{CS}$), **chip enable** ($\overline{CE}$), or simply **select** ($\overline{S}$) input. RAM memory generally has at least one $\overline{CS}$ or $\overline{S}$ input, and ROM has at least one $\overline{CE}$. If the $\overline{CE}$, $\overline{CS}$, or $\overline{S}$ input is active (a logic 0, in this case, because of the overbar), the memory device performs a read or write operation; if it is inactive (a logic 1, in this case), the memory device cannot do a read or a write because it is turned off or disabled. If more than one $\overline{CS}$ connection is present, all must be activated to read or write data.

Control Connections.   All memory devices have some form of control input or inputs. A ROM usually has only one control input, while a RAM often has one or two control inputs.

The control input most often found on a ROM is the **output enable** ($\overline{OE}$) or **gate** ($\overline{G}$) connection, which allows data to flow out of the output data pins of the ROM. If $\overline{OE}$ and the selection input ($\overline{CE}$) are both active, the output is enabled; if $\overline{OE}$ is inactive, the output is disabled at its high-impedance state. The $\overline{OE}$ connection enables and disables a set of three-state buffers located within the memory device and must be active to read data.

A RAM memory device has either one or two control inputs. If there is only one control input, it is often called $R/\overline{W}$. This pin selects a read operation or a write operation only if the device is selected by the selection input ($\overline{CS}$). If the RAM has two control inputs, they are usually labeled $\overline{WE}$ (or $\overline{W}$), and $\overline{OE}$ (or $\overline{G}$). Here, $\overline{WE}$ (**write enable**) must be active to perform a memory write, and $\overline{OE}$ must be active to perform a memory read operation. When these two controls ($\overline{WE}$ and $\overline{OE}$) are present, they must never both be active at the same time. If both control inputs are inactive (logic 1s), data are neither written nor read, and the data connections are at their high-impedance state.

ROM Memory

The **read-only memory** (ROM) permanently stores programs and data that are resident to the system and must not change when power supply is disconnected. The ROM is permanently programmed so that data are always present, even when power is disconnected. This type of memory is often called **nonvolatile memory**, because its contents *do not* change even if power is disconnected.

Today, the ROM is available in many forms. A device we call a ROM is purchased in mass quantities from a manufacturer and programmed during its fabrication at the factory. The EPROM (**erasable programmable read-only memory**), a type of ROM, is more commonly used when software must be changed often or when too few are in demand to make the ROM economical. For a ROM to be practical, we usually must purchase at least 10,000 devices to recoup the factory programming charge. An EPROM is programmed in the field on a device called an EPROM programmer. The EPROM is also erasable if exposed to high-intensity ultraviolet light for about 20 minutes or so, depending on the type of EPROM.

PROM memory devices are also available, although they are not as common today. The PROM (**programmable read-only memory**) is also programmed in the field by burning open tiny NI-chrome or silicon oxide fuses; but once it is programmed, it cannot be erased.

MODE SELECTION

PINS MODE	PD/PGM (18)	$\overline{CS}$ (20)	V_{PP} (21)	V_{CC} (24)	OUTPUTS (9-11, 13-17)
Read	V_{IL}	V_{IL}	+5	+5	D_{OUT}
Deselect	Don't care	V_{IH}	+5	+5	High Z
Power Down	V_{IH}	Don't care	+5	+5	High Z
Program	Pulsed V_{IL} to V_{IH}	V_{IH}	+25	+5	D_{IN}
Program Verify	V_{IL}	V_{IL}	+25	+5	D_{OUT}
Program Inhibit	V_{IL}	V_{IH}	+25	+5	High Z

PIN CONFIGURATION

PIN NAMES

A_0-A_{10}	ADDRESSES
PD/PGM	POWER DOWN/PROGRAM
$\overline{CS}$	CHIP SELECT
O_0-O_7	OUTPUTS

BLOCK DIAGRAM

FIGURE 10–2 The pin-out of the 2716, $2K \times 8$ EPROM. (Courtesy of Intel Corporation.)

Still another, newer type of **read-mostly memory** (RMM) is called the **flash memory**. The flash memory[1] is also often called an EEPROM (**electrically erasable programmable ROM**), EAROM (**electrically alterable ROM**), or a NOVRAM (**nonvolatile RAM**). These memory devices are electrically erasable in the system, but they require more time to erase than a normal RAM. The flash memory device is used to store setup information for systems such as the video card in the computer. It has all but replaced the EPROM in most computer systems for the BIOS memory. Some systems contain a password stored in the flash memory device. Flash memory has its biggest impact in memory cards for digital cameras and memory in MP3 audio players.

Figure 10–2 illustrates the 2716 EPROM, which is representative of most common EPROMs. This device contains 11 address inputs and eight data outputs. The 2716 is a $2K \times 8$ read-only memory device. The 27XXX series of the EPROMs includes the following part numbers: 2704 (512×8), 2708 ($1K \times 8$), 2716 ($2K \times 8$), 2732 ($4K \times 8$), 2764 ($8K \times 8$), 27128 ($16K \times 8$), 27256 ($32K \times 8$), 27512 ($64K \times 8$), and 271024 ($128K \times 8$). Each of these parts contains address pins, eight data connections, one or more chip selection inputs ($\overline{CE}$), and an output enable pin ($\overline{OE}$).

Figure 10–3 illustrates the timing diagram for the 2716 EPROM. Data appear on the output connections only after a logic 0 is placed on both $\overline{CE}$ and $\overline{OE}$ pin connections. If $\overline{CE}$ and $\overline{OE}$ are not both logic 0s, the data output connections remain at their high-impedance or off states. Note that the V_{PP} pin must be placed at a logic 1 level for data to be read from the EPROM. In some cases, the V_{PP} pin is in the same position as the $\overline{WE}$ pin on the SRAM. This will allow a single socket to hold either an EPROM or an SRAM. An example is the 27256 EPROM and the 62256 SRAM, both $32K \times 8$ devices that have the same pin-out, except for V_{PP} on the EPROM and $\overline{WE}$ on the SRAM.

---

[1]Flash memory is a trademark of Intel Corporation.

A.C. Characteristics

$T_A = 0°C$ to $70°C$, V_{CC}[1] $= +5V ±5\%$, V_{PP}[2] $= V_{CC} ±0.6V$[3]

Symbol	Parameter	Limits			Unit	Test Conditions
		Min.	Typ.[4]	Max.		
t_{ACC1}	Address to Output Delay		250	450	ns	PD/PGM = $\overline{CS}$ = V_{IL}
t_{ACC2}	PD/PGM to Output Delay		280	450	ns	$\overline{CS}$ = V_{IL}
t_{CO}	Chip Select to Output Delay			120	ns	PD/PGM = V_{IL}
t_{PF}	PD/PGM to Output Float	0		100	ns	$\overline{CS}$ = V_{IL}
t_{DF}	Chip Deselect to Output Float	0		100	ns	PD/PGM = V_{IL}
t_{OH}	Address to Output Hold	0			ns	PD/PGM = $\overline{CS}$ = V_{IL}

Capacitance[5] $T_A = 25°C$, f = 1 MHz

Symbol	Parameter	Typ.	Max.	Unit	Conditions
C_{IN}	Input Capacitance	4	6	pF	V_{IN} = 0V
C_{OUT}	Output Capacitance	8	12	pF	V_{OUT} = 0V

A.C. Test Conditions:

Output Load: 1 TTL gate and C_L = 100 pF
Input Rise and Fall Times: ≤20 ns
Input Pulse Levels: 0.8V to 2.2V
Timing Measurement Reference Level:
    Inputs   1V and 2V
    Outputs   0.8V and 2V
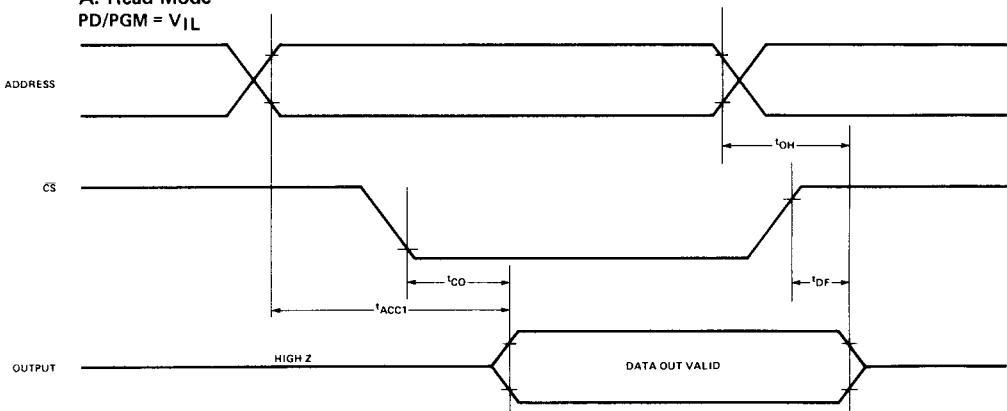
WAVEFORMS

A. Read Mode
PD/PGM = V_{IL}

FIGURE 10–3   The timing diagram of AC characteristics of the 2716 EPROM. (Courtesy of Intel Corporation.)
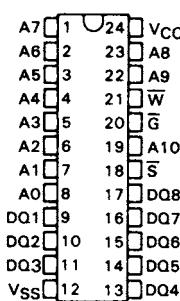
One important piece of information provided by the timing diagram and data sheet is the memory access time—the time that it takes the memory to read information. As Figure 10–3 illustrates, memory access time (T_{ACC}) is measured from the appearance of the address at the address inputs until the appearance of the data at the output connections. This is based on the assumption that the $\overline{CE}$ input goes low at the same time that the address inputs become stable. Also, $\overline{OE}$ must be a logic 0 for the output connections to become active. The basic speed of this EPROM is 450 ns. (Recall that the 8086/8088 operated with a 5 MHz clock allowed memory 460 ns to access data.) This type of memory component requires wait states to operate properly with the 8086/8088 microprocessors because of its rather long access time. If wait states are not desired, higher-speed versions of the EPROM are available at an additional cost. Today, EPROM memory is available with access times of as little as 100 ns. Obviously, wait states are required in modern microprocessors for any EPROM device.

Static RAM (SRAM) Devices

Static RAM memory devices retain data for as long as DC power is applied. Because no special action (except power) is required to retain stored data, these devices are called **static memory**. They are also called **volatile memory** because they will not retain data without power. The main

FIGURE 10–4   The pin-out
of the TMS4016, 2K × 8 static
RAM (SRAM). (Courtesy of
Texas Instruments
Incorporated.)

```
     TMS4016 . . . NL PACKAGE
            (TOP VIEW)

     A7 ▢ 1   ᴗ  24 ▢ Vcc
     A6 ▢ 2      23 ▢ A8
     A5 ▢ 3      22 ▢ A9
     A4 ▢ 4      21 ▢ W̄
     A3 ▢ 5      20 ▢ Ḡ
     A2 ▢ 6      19 ▢ A10
     A1 ▢ 7      18 ▢ S̄
     A0 ▢ 8      17 ▢ DQ8
    DQ1 ▢ 9      16 ▢ DQ7
    DQ2 ▢ 10     15 ▢ DQ6
    DQ3 ▢ 11     14 ▢ DQ5
    Vss ▢ 12     13 ▢ DQ4
```

PIN NOMENCLATURE	
A0 – A10	Addresses
DQ1 – DQ8	Data In/Data Out
Ḡ	Output Enable
S̄	Chip Select
Vcc	+5-V Supply
Vss	Ground
W̄	Write Enable

difference between a ROM and a RAM is that a RAM is written under normal operation, whereas a ROM is programmed outside the computer and normally is only read. The SRAM, which stores temporary data, is used when the size of the read/write memory is relatively small. Today, a small memory is one that is less than 1M byte.

Figure 10–4 illustrates the 4016 SRAM, which is a 2K × 8 read/write memory. This device has 11 address inputs and eight data input/output connections. This device is representative of all SRAM devices, except for the number of address and data connections.

The control inputs of this RAM are slightly different from those presented earlier. The $\overline{OE}$ pin is labeled $\overline{G}$, the $\overline{CS}$ pin is $\overline{S}$, and the $\overline{WE}$ pin is $\overline{W}$. Despite the altered designations, the control pins function exactly the same as those outlined previously. Other manufacturers make this popular SRAM under the part numbers 2016 and 6116.

Figure 10–5 depicts the timing diagram for the 4016 SRAM. As the read cycle timing reveals, the access time is $t_{a(A)}$. On the slowest version of the 4016, this time is 250 ns, which is fast enough to connect directly to an 8088 or an 8086 operated at 5 MHz without wait states. Again, it is important to remember that the access time must be checked to determine the compatibility of memory components with the microprocessor.

Figure 10–6 on p. 336 illustrates the pin-out of the 62256, 32K × 8 static RAM. This device is packaged in a 28-pin integrated circuit and is available with access times of 120 ns or 150 ns. Other common SRAM devices are available in 8K × 8, 128K × 8, 256K × 8, 512K × 8, and 1M × 8 sizes, with access times of as little as 1.0 ns for SRAM used in computer cache memory systems.

Dynamic RAM (DRAM) Memory

About the largest static RAM available today is a 1M × 8. Dynamic RAM, on the other hand, is available in much larger sizes: up to 256M × 8 (2G bits). In all other respects, DRAM is essentially the same as SRAM, except that it retains data for only 2 or 4 ms on an integrated capacitor. After 2 or 4 ms, the contents of the DRAM must be completely rewritten *(refreshed)* because the capacitors, which store a logic 1 or logic 0, lose their charges.

Instead of requiring the almost impossible task of reading the contents of each memory location with a program and then rewriting them, the manufacturer has internally constructed the DRAM differently from the SRAM. In the DRAM, the entire contents of the memory are

refreshed with 256 reads in a 2- or 4-ms interval. Refreshing also occurs during a write, a read, or during a special refresh cycle. Much more information about DRAM refreshing is provided in Section 10–6.

Another disadvantage of DRAM memory is that it requires so many address pins that the manufacturers have decided to multiplex the address inputs. Figure 10–7 illustrates a 64K × 4 DRAM, the TMS4464, which stores 256K bits of data. Notice that it contains only eight address inputs where it should contain 16—the number required to address 64K memory locations. The

electrical characteristics over recommended operating free-air temperature range (unless otherwise noted)

PARAMETER		TEST CONDITIONS		MIN	TYP†	MAX	UNIT
V_{OH}	High level voltage	$I_{OH} = -1$ mA,	$V_{CC} = 4.5$ V	2.4			V
V_{OL}	Low level voltage	$I_{OL} = 2.1$ mA,	$V_{CC} = 4.5$ V			0.4	V
I_I	Input current	$V_I = 0$ V to 5.5 V				10	μA
I_{OZ}	Off-state output current	$\overline{S}$ or $\overline{G}$ at 2 V or $\overline{W}$ at 0.8 V, $V_O = 0$ V to 5.5 V				10	μA
I_{CC}	Supply current from V_{CC}	$I_O = 0$ mA, $T_A = 0°C$ (worst case)	$V_{CC} = 5.5$ V,		40	70	mA
C_i	Input capacitance	$V_I = 0$ V,	$f = 1$ MHz			8	pF
C_O	Output capacitance	$V_O = 0$ V,	$f = 1$ MHz			12	pF

†All typical values are at $V_{CC} = 5$ V, $T_A = 25°C$.

timing requirements over recommended supply voltage range and operating free-air temperature range

PARAMETER		TMS4016-12		TMS4016-15		TMS4016-20		TMS4016-25		UNIT
		MIN	MAX	MIN	MAX	MIN	MAX	MIN	MAX	
$t_{c(rd)}$	Read cycle time	120		150		200		250		ns
$t_{c(wr)}$	Write cycle time	120		150		200		250		ns
$t_{w(W)}$	Write pulse width	60		80		100		120		ns
$t_{su(A)}$	Address setup time	20		20		20		20		ns
$t_{su(S)}$	Chip select setup time	60		80		100		120		ns
$t_{su(D)}$	Data setup time	50		60		80		100		ns
$t_{h(A)}$	Address hold time	0		0		0		0		ns
$t_{h(D)}$	Data hold time	5		10		10		10		ns

switching characteristics over recommended voltage range, $T_A = 0°C$ to $70°C$

PARAMETER		TMS4016-12		TMS4016-15		TMS4016-20		TMS4016-25		UNIT
		MIN	MAX	MIN	MAX	MIN	MAX	MIN	MAX	
$t_{a(A)}$	Access time from address		120		150		200		250	ns
$t_{a(S)}$	Access time from chip select low		60		75		100		120	ns
$t_{a(G)}$	Access time from output enable low		50		60		80		100	ns
$t_{v(A)}$	Output data valid after address change	10		15		15		15		ns
$t_{dis(S)}$	Output disable time after chip select high		40		50		60		80	ns
$t_{dis(G)}$	Output disable time after output enable high		40		50		60		80	ns
$t_{dis(W)}$	Output disable time after write enable low		50		60		60		80	ns
$t_{en(S)}$	Output enable time after chip select low	5		5		10		10		ns
$t_{en(G)}$	Output enable time after output enable low	5		5		10		10		ns
$t_{en(W)}$	Output enable time after write enable high	5		5		10		10		ns
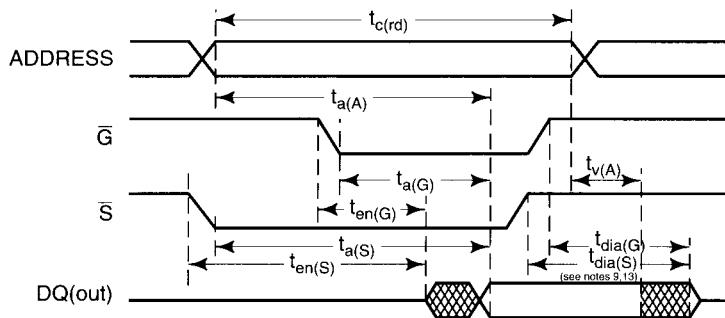
NOTES: 3. $C_L = 100$pF for all measurements except $t_{dis(W)}$ and $t_{en(W)}$.
         $C_L = 5$ pF for $t_{dis(W)}$ and $t_{en(W)}$.
      4. t_{dis} and t_{en} parameters are sampled and not 100% tested.
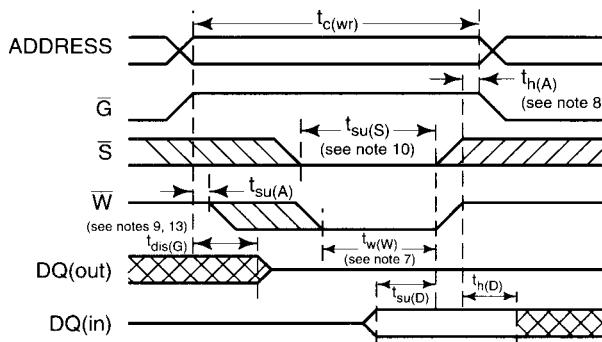
FIGURE 10–5   (a) The AC characteristics of the TMS4016 SRAM. (b) The timing diagrams of the TMS4016 SRAM. (Courtesy of Texas Instruments Incorporated.)
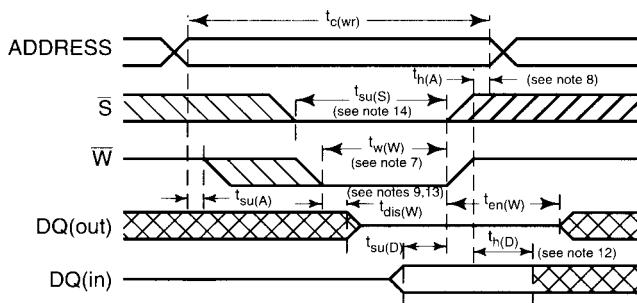
timing waveform of read cycle (see note 5)

timing waveform of write cycle no. 1 (see note 6)
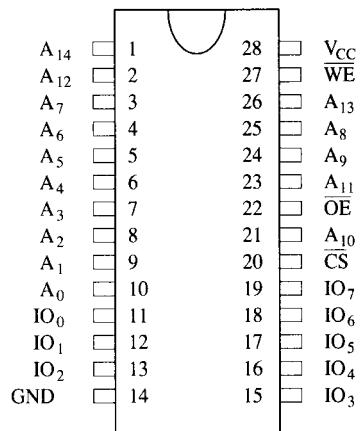
timing waveform of write cycle no. 2 (see notes 6 and 11)

NOTES: 5. $\overline{W}$ is high Read Cycle.
   6. $\overline{W}$ must be high during all address transitions.
   7. A write occurs during the overlap of a low $\overline{S}$ and a low $\overline{W}$.
   8. $t_{h(A)}$ is measured from the earlier of $\overline{S}$ or $\overline{W}$ going high to the end of the write cycle.
   9. During this period, I/O pins are in the output state so that the input signals of opposite phase to the outputs must not be applied.
   10. If the Slow transition occurs simultaneously with the $\overline{W}$ low transitions or after the $\overline{W}$ transition, output remains in a high impedance state.
   11. G is continuously low ($G = V_{IL}$).
   12. If $\overline{S}$ is low during this period, I/O pins are in the output state. Data input signals of opposite phase to the outputs must not be applied.
   13. Transition is measured $\pm$ 200 mV from steady-state voltage.
   14. If the $\overline{S}$ low transition occurs before the W low transition, then the data input signals of opposite phase to the outputs must not be applied for the duration of $t_{dis(W)}$ after the $\overline{W}$ low transition.

FIGURE 10–5   *(continued)*

FIGURE 10–6 Pin diagram of the 62256, 32K × 8 static RAM.
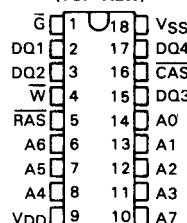
PIN FUNCTION

$A_0 - A_{14}$	Addresses
$IO_0 - IO_7$	Data connections
$\overline{CS}$	Chip select
$\overline{OE}$	Output enable
$\overline{WE}$	Write enable
V_{CC}	+5V Supply
GND	Ground

FIGURE 10–7 The pin-out of the TMS4464, 64K × 4 dynamic RAM (DRAM). (Courtesy of Texas Instruments Incorporated.)

TMS4464 . . . JL OR NL PACKAGE
(TOP VIEW)

(a)

PIN NOMENCLATURE	
A0-A7	Address Inputs
$\overline{CAS}$	Column Address Strobe
DQ1-DQ4	Data-In/Data-Out
$\overline{G}$	Output Enable
$\overline{RAS}$	Row Address Strobe
V_{DD}	+ 5-V Supply
V_{SS}	Ground
$\overline{W}$	Write Enable

(b)

only way that 16 address bits can be forced into eight address pins is in two 8-bit increments. This operation requires two special pins: the **column address strobe** ($\overline{CAS}$) and **row address strobe** ($\overline{RAS}$). First, A_0–A_7 are placed on the address pins and strobed into an internal row latch by $\overline{RAS}$ as the row address. Next, the address bits A_8-A_{15} are placed on the same eight address

inputs and strobed into an internal column latch by $\overline{\text{CAS}}$ as the column address (see Figure 10–8 for this timing). The 16-bit address held in these internal latches addresses the contents of one of the 4-bit memory locations. Note that $\overline{\text{CAS}}$ also performs the function of the chip selection input to the DRAM.

Figure 10–9 illustrates a set of multiplexers used to strobe the column and row addresses into the eight address pins on a pair of TMS4464 DRAMs. Here, the $\overline{\text{RAS}}$ signal not only strobes
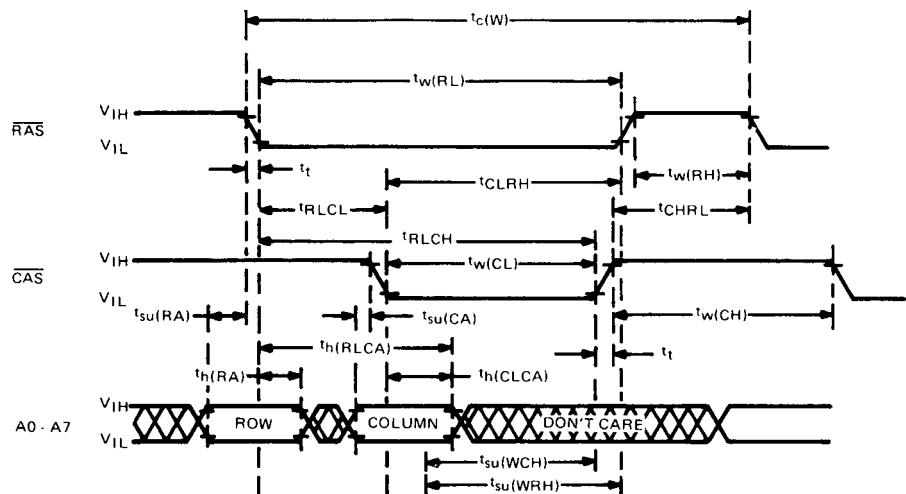
FIGURE 10–8 $\overline{\text{RAS}}$, $\overline{\text{CAS}}$, and address input timing for the TMS4464 DRAM. (Courtesy of Texas Instruments Incorporated.)
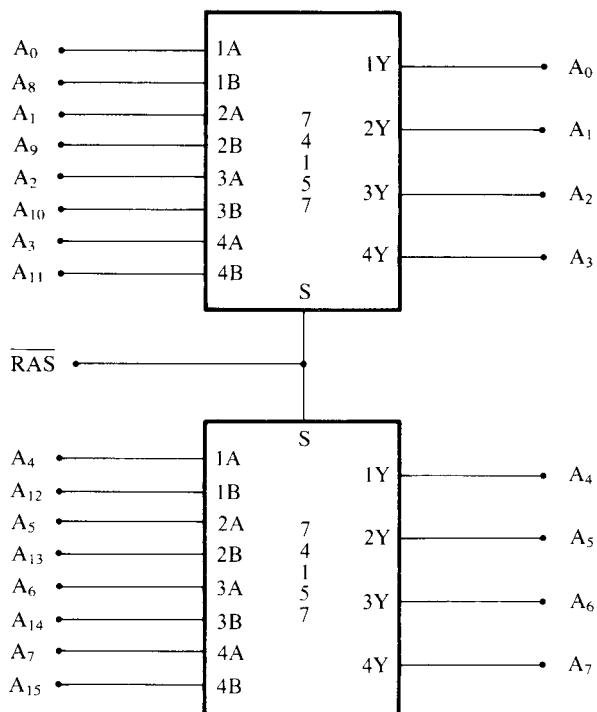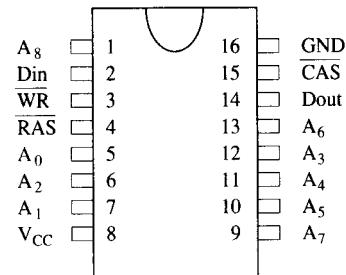
FIGURE 10–9 Address multiplexer for the TMS4464 DRAM.

FIGURE 10–10   The 41256 dynamic RAM organized as a 256K × 1 memory device.

A_8	1	16	GND
Din	2	15	$\overline{CAS}$
$\overline{WR}$	3	14	Dout
$\overline{RAS}$	4	13	A_6
A_0	5	12	A_3
A_2	6	11	A_4
A_1	7	10	A_5
V_{CC}	8	9	A_7

PIN FUNCTIONS

A_0 - A_8	Addresses
Din	Data in
Dout	Data out
$\overline{CAS}$	Column Address Strobe
$\overline{RAS}$	Row Address Strobe
$\overline{WR}$	Write enable
V_{CC}	+5V Supply
GND	Ground

the row address into the DRAMs, but it also selects which part of the address is applied to the address inputs. This is possible due to the long propagation-delay time of the multiplexers. When $\overline{RAS}$ is a logic 1, the B inputs are connected to the Y outputs of the multiplexers; when the $\overline{RAS}$ input goes to a logic 0, the A inputs connect to the Y outputs. Because the internal row address latch is edge-triggered, it captures the row address before the address at the inputs changes to the column address. More detail on DRAM and DRAM interfacing is provided in Section 10–6.

As with the SRAM, the R/$\overline{W}$ pin writes data to the DRAM when a logic 0, but there is no pin labeled G or enable. There also is no $\overline{S}$ (select) input to the DRAM. As mentioned, the $\overline{CAS}$ input selects the DRAM. If selected, the DRAM is written if R/$\overline{W}$ = 0 and read if R/$\overline{W}$ = 1.

Figure 10–10 shows the pin-out of the 41256 dynamic RAM. This device is organized as a 256K × 1 memory, requiring as little as 70 ns to access data.

More recently, larger DRAMs have become available that are organized as a 16M × 1, 256M × 1, 512M × 1, 1G × 1, and 2G × 1 memory. On the horizon is the 4G × 1 memory, which is in the planning stages. DRAM memory is often placed on small circuit boards called SIMMs (Single In-Line Memory Modules). Figure 10–11 shows the pin-outs of the two most common SIMMs. The 30-pin SIMM is organized most often as 1M × 8 or 1M × 9, and 4M × 8 or 4M × 9. (Illustrated in Figure 10–11 is a 4M × 9.) The ninth bit is the parity bit. Also shown is a newer 72 pin SIMM. The 72-pin SIMMs are often organized as 1M × 32 or 1M × 36 (with parity). Other sizes are 2M × 32, 4M × 32, 8M × 32, and 16M × 32. These are also available with parity. Figure 10–11 illustrates a 4M × 36 SIMM, which has 16M bytes of memory.

Lately, many systems are using the Pentium–Pentium 4 microprocessors. These microprocessors have a 64-bit wide data bus, which precludes the use of the 8-bit-wide SIMMs described here. Even the 72-pin SIMMs are cumbersome to use because they must be used in pairs to obtain a 64-bit-wide data connection. Today, the 64-bit-wide DIMMs (Dual In-line Memory Modules) have become the standard in most systems. The memory on these modules is organized as 64 bits wide. The common sizes available are 16M bytes (2M × 64), 32M bytes (4M × 64), 64M bytes (8M × 64), 128M bytes (16M × 64), 256M bytes (32M × 64), 512M bytes (64M × 64), and 1G bytes (128M × 64). The pin-out of the DIMM is illustrated in Figure 10–12. The DIMM module is available in DRAM, EDO, SDRAM, and DDR (double-data rate) forms,
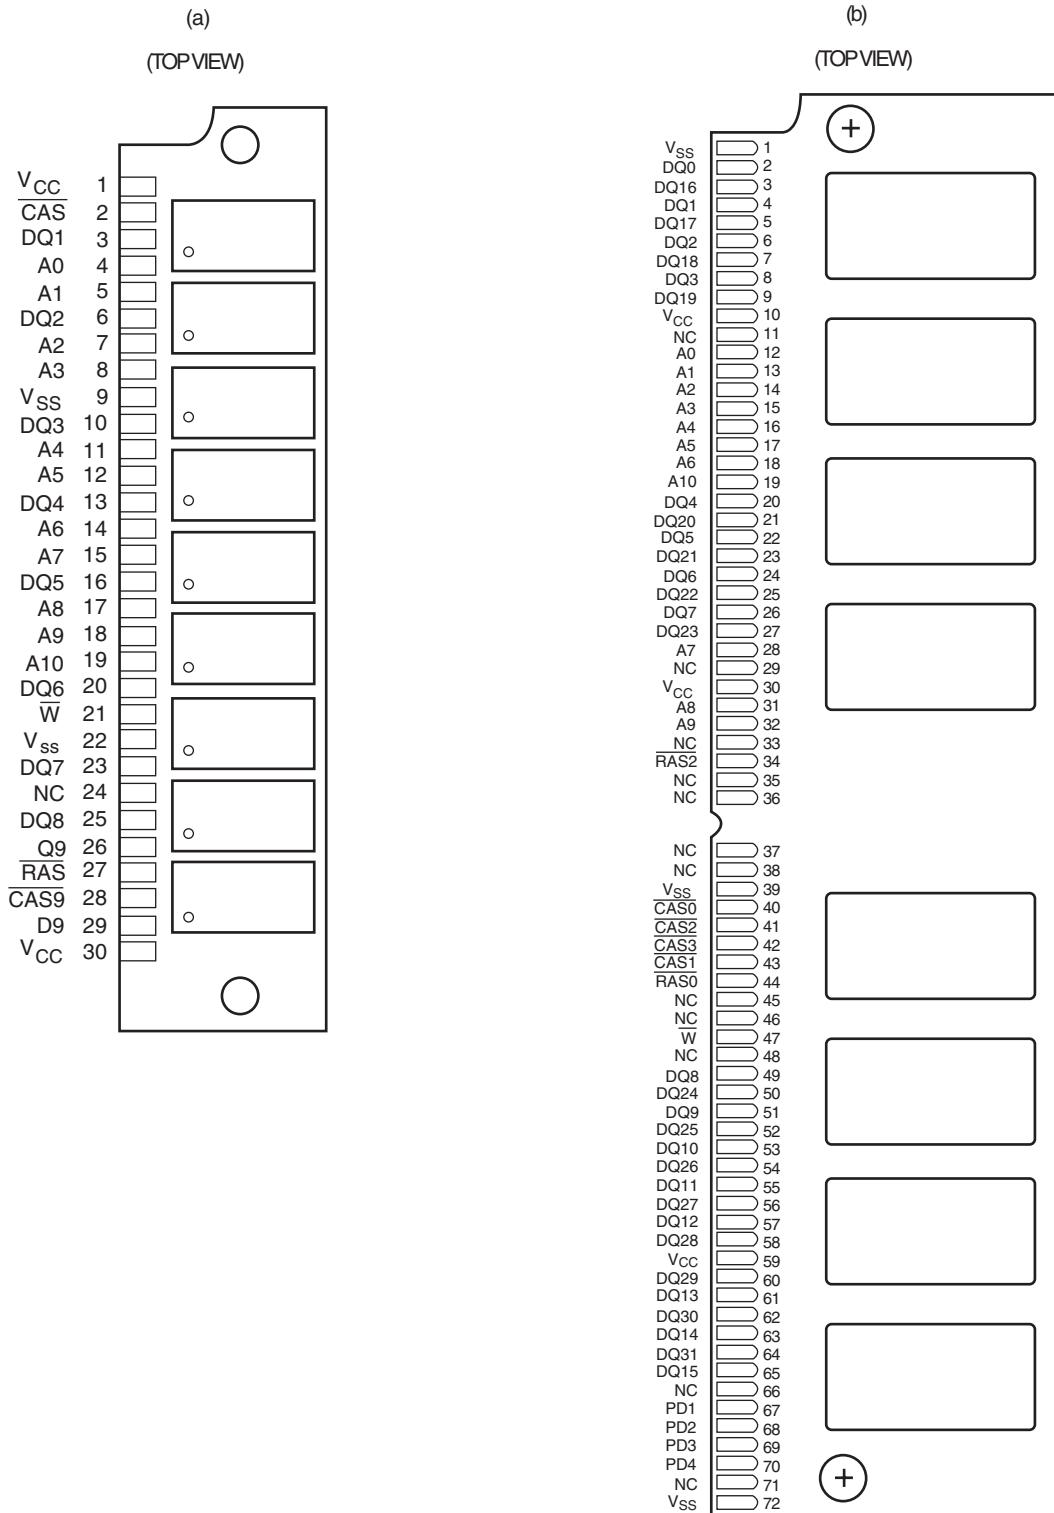
FIGURE 10–11 The pin-outs of the 30-pin and 72-pin SIMM. (a) A 30-pin SIMM organized as 4M × 9 and (b) a 72-pin SIMM organized as 4M × 36.
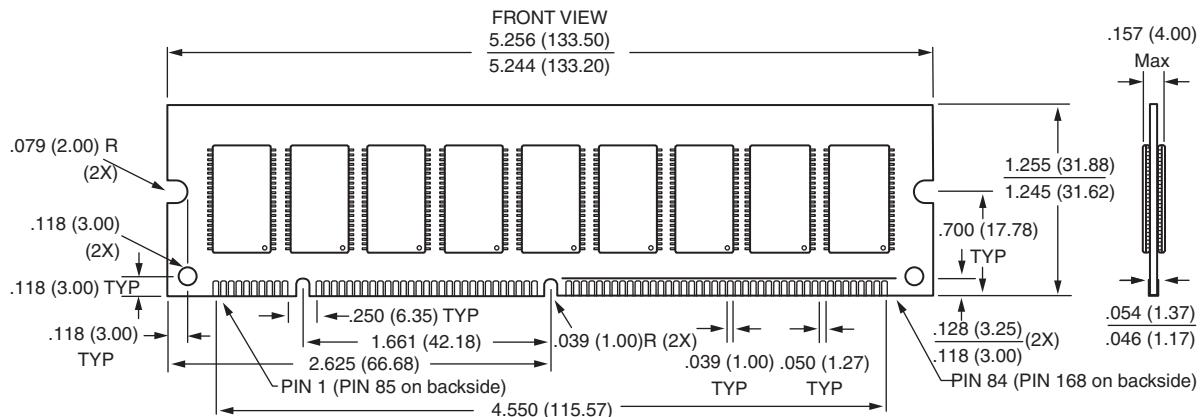
FIGURE 10–12   The pin-out of a 168-pin DIMM.

with or without an EPROM. The EPROM provides information to the system on the size and the speed of the memory device for plug-and-play applications.

Another addition to the memory market is the RIMM memory module from RAMBUS Corporation, although this memory type has faded from the market. Like the SDRAM, the RIMM contains 168 pins, but each pin is a two-level pin, bringing the total number of connections to 336. The fastest SDRAM currently available is the PC-4400 or 500 DDR, which operates at a rate of 4.4G bytes per second. By comparison, the 800 MHz RIMM operates at a rate of 3.2G bytes per second and is no longer supported in many systems. RDRAM had a fairly short life in the volatile computer market. The RIMM module is organized as a 32-bit-wide device. This means that to populate a Pentium 4 memory, RIMM memory is used in pairs. Intel claims that the Pentium 4 system using RIMM modules is 300% faster than a Pentium III using PC-100 memory. According to RAMBUS, the current 800 MHz RIMM has been increased to a speed of 1200 MHz, but it is still not fast enough to garner much of a market share.

Currently the latest DRAM is the DDR (**double-data rate**) memory device and DDR2. The DDR memory transfers data at each edge of the clock, making it operate at twice the speed of SDRAM. This does not affect the access time for the memory, so many wait states are still required to operate this type of memory, but it can be much faster than normal SDRAM memory and that includes RDRAM.

10–2   ADDRESS DECODING

In order to attach a memory device to the microprocessor, it is necessary to decode the address sent from the microprocessor. Decoding makes the memory function at a unique section or partition of the memory map. Without an address decoder, only one memory device can be connected to a microprocessor, which would make it virtually useless. In this section, we describe a few of the more common address-decoding techniques, as well as the decoders that are found in many systems.

Why Decode Memory?

When the 8088 microprocessor is compared to the 2716 EPROM, a difference in the number of address connections is apparent—the EPROM has 11 address connections and the microprocessor has 20. This means that the microprocessor sends out a 20-bit memory address whenever it

reads or writes data. Because the EPROM has only 11 address pins, there is a mismatch that must be corrected. If only 11 of the 8088's address pins are connected to the memory, the 8088 will see only 2K bytes of memory instead of the 1M bytes that it "expects" the memory to contain. The decoder corrects the mismatch by decoding the address pins that do not connect to the memory component.

Simple NAND Gate Decoder

When the 2K × 8 EPROM is used, address connections A_{10}–A_0 of the 8088 are connected to address inputs A_{10}–A_0 of the EPROM. The remaining nine address pins (A_{19}–A_{11}) are connected to the inputs of a NAND gate decoder (see Figure 10–13). The decoder selects the EPROM from one of the 2K-byte sections of the 1M-byte memory system in the 8088 microprocessor.

In this circuit, a single NAND gate decodes the memory address. The output of the NAND gate is a logic 0 whenever the 8088 address pins attached to its inputs (A_{19}–A_{11}) are all logic 1s. The active low, logic 0 output of the NAND gate decoder is connected to the $\overline{CE}$ input pin that selects (enables) the EPROM. Recall that whenever $\overline{CE}$ is a logic 0, data will be read from the EPROM only if $\overline{OE}$ is also a logic 0. The $\overline{OE}$ pin is activated by the 8088 $\overline{RD}$ signal or the $\overline{MRDC}$ (**memory read control**) signal of other family members.

If the 20-bit binary address, decoded by the NAND gate, is written so that the leftmost nine bits are 1s and the rightmost 11 bits are don't cares (X), the actual address range of the EPROM can be determined. (A *don't care* is a logic 1 or a logic 0, whichever is appropriate.)

Example 10–1 illustrates how the address range for this EPROM is determined by writing down the externally decoded address bits (A_{19}–A_{11}) and the address bits decoded by the EPROM (A_{10}–A_0) as don't cares. We really do not care about the address pins on the EPROM because they are internally decoded. As the example illustrates, the don't cares are first written as 0s to locate the lowest address and then as 1s to find the highest address. Example 10–1 also shows these binary boundaries as hexadecimal addresses. Here, the 2K EPROM is decoded at memory address locations FF800H–FFFFFH. Notice that this is a 2K-byte section of the memory and is also located at the reset location for the 8086/8088 (FFFF0H), the most likely place for an EPROM in a system.
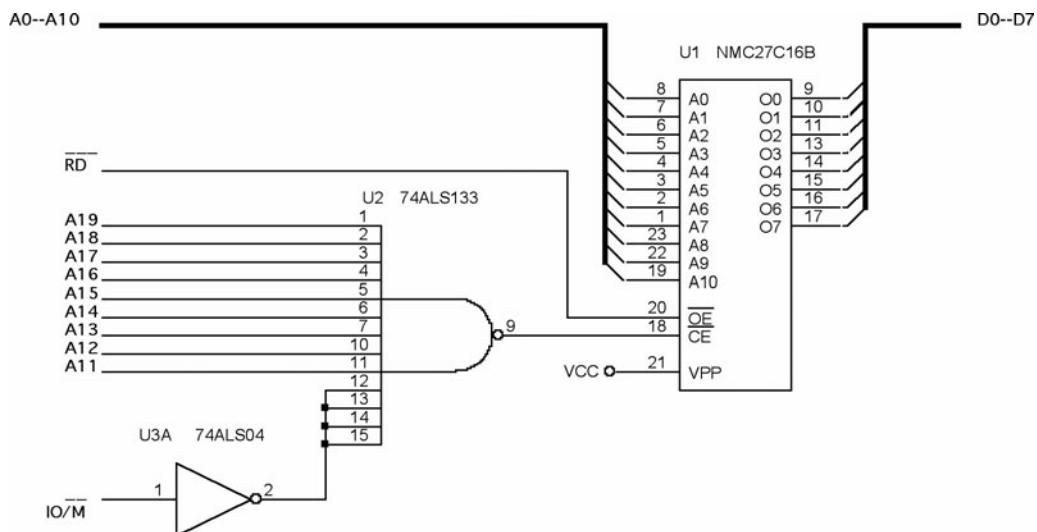
FIGURE 10–13 A simple NAND gate decoder that selects a 2716 EPROM for memory location FF800H–FFFFFH.

EXAMPLE 10–1

```
1111 1111 1XXX XXXX XXXX

          or

1111 1111 1000 0000 0000 = FF800H
          to
1111 1111 1111 1111 1111 = FFFFFH
```
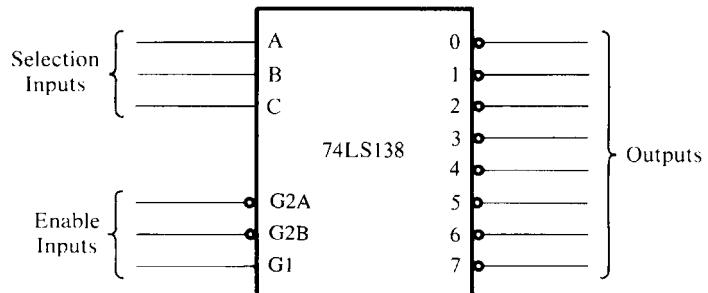
Although this example serves to illustrate decoding, NAND gates are rarely used to decode memory because each memory device requires its own NAND gate decoder. Because of the excessive cost of the NAND gate decoder and inverters that are often required, this option requires that an alternate be found.

The 3-to-8 Line Decoder (74LS138)

One of the more common, although not only, integrated circuit decoders found in many microprocessor-based systems is the 74LS138 3-to-8 line decoder. Figure 10–14 illustrates this decoder and its truth table.

The truth table shows that only one of the eight outputs ever goes low at any time. For any of the decoder's outputs to go low, the three enable inputs ($\overline{G2A}$, $\overline{G2B}$, and G1) must all be active. To be active, the $\overline{G2A}$ and $\overline{G2B}$ inputs must both be low (logic 0), and G1 must be high (logic 1). Once the 74LS138 is enabled, the address inputs (C, B, and A) select which output pin

FIGURE 10–14   The 74LS138 3-to-8 line decoder and function table.

Inputs						Outputs							
Enable			Select										
$\overline{G2A}$	$\overline{G2B}$	G1	C	B	A	$\overline{0}$	$\overline{1}$	$\overline{2}$	$\overline{3}$	$\overline{4}$	$\overline{5}$	$\overline{6}$	7
1	X	X	X	X	X	1	1	1	1	1	1	1	1
X	1	X	X	X	X	1	1	1	1	1	1	1	1
X	X	0	X	X	X	1	1	1	1	1	1	1	1
0	0	1	0	0	0	0	1	1	1	1	1	1	1
0	0	1	0	0	1	1	0	1	1	1	1	1	1
0	0	1	0	1	0	1	1	0	1	1	1	1	1
0	0	1	0	1	1	1	1	1	0	1	1	1	1
0	0	1	1	0	0	1	1	1	1	0	1	1	1
0	0	1	1	0	1	1	1	1	1	1	0	1	1
0	0	1	1	1	0	1	1	1	1	1	1	0	1
0	0	1	1	1	1	1	1	1	1	1	1	1	0
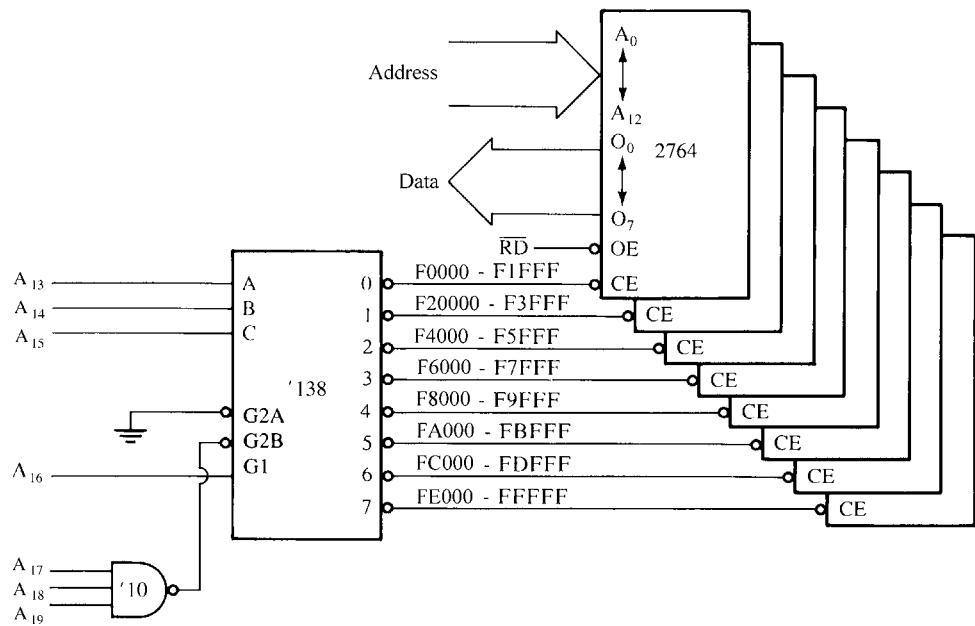
FIGURE 10–15 A circuit that uses eight 2764 EPROMs for a 64K × 8 section of memory in an 8088 microprocessor-based system. The addresses selected in this circuit are F0000H–FFFFFH.

goes low. Imagine eight EPROM $\overline{CE}$ inputs connected to the eight outputs of the decoder! This is a very powerful device because it selects eight different memory devices at the same time. Even today this device still finds wide application.

Sample Decoder Circuit. Notice that the outputs of the decoder, illustrated in Figure 10–15, are connected to eight different 2764 EPROM memory devices. Here, the decoder selects eight 8K-byte blocks of memory for a total memory capacity of 64K bytes. This figure also illustrates the address range of each memory device and the common connections to the memory devices. Notice that all of the address connections from the 8088 are connected to this circuit. Also, notice that the decoder's outputs are connected to the $\overline{CE}$ inputs of the EPROMs, and the $\overline{RD}$ signal from the 8088 is connected to the $\overline{OE}$ inputs of the EPROMs. This allows only the selected EPROM to be enabled and to send its data to the microprocessor through the data bus whenever $\overline{RD}$ becomes a logic 0.

In this circuit, a three-input NAND gate is connected to address bits A_{19}–A_{17}. When all three address inputs are high, the output of this NAND gate goes low and enables input $\overline{G2B}$ of the 74LS138. Input G1 is connected directly to A_{16}. In other words, in order to enable this decoder, the first four address connections (A_{19}–A_{16}) must all be high.

The address inputs C, B, and A connect to microprocessor address pins A_{15}–A_{13}. These three address inputs determine which output pin goes low and which EPROM is selected whenever the 8088 outputs a memory address within this range to the memory system.

Example 10–2 shows how the address range of the entire decoder is determined. Notice that the range is location F0000H–FFFFFH. This is a 64K-byte span of the memory.

EXAMPLE 10–2

```
1111 XXXX XXXX XXXX XXXX

        or

1111 0000 0000 0000 0000 = F0000H
        to
1111 1111 1111 1111 1111 = FFFFFH
```

How is it possible to determine the address range of each memory device attached to the decoder's outputs? Again, the binary bit pattern is written down; this time the C, B, and A address inputs are not don't cares. Example 10–3 shows how output 0 of the decoder is made to go low to select the EPROM attached to that pin. Here, C, B, and A are shown as logic 0s.

EXAMPLE 10–3

```
    CBA
1111 OOOX XXXX XXXX XXXX

        or

1111 0000 0000 0000 0000 = FOOOOH
        to
1111 0001 1111 1111 1111 = F1FFFH
```

If the address range of the EPROM connected to output 1 of the decoder is required, it is determined in exactly the same way as that of output 0. The only difference is that now the C, B, and A inputs contain a 001 instead of a 000 (see Example 10–4). The remaining output address ranges are determined in the same manner by substituting the binary address of the output pin into C, B, and A.

EXAMPLE 10–4

```
    CBA
1111 001X XXXX XXXX XXXX

        or

1111 0010 0000 0000 0000 = F2000H
        to
1111 0011 1111 1111 1111 = F3FFFH
```

The Dual 2-to-4 Line Decoder (74LS139)

Another decoder that finds some application is the 74LS139 dua1 2-to-4 line decoder. Figure 10–16 illustrates both the pin-out and the truth table for this decoder. The 74LS139 contains two separate 2-to-4 line decoders—each with its own address, enable, and output connections.
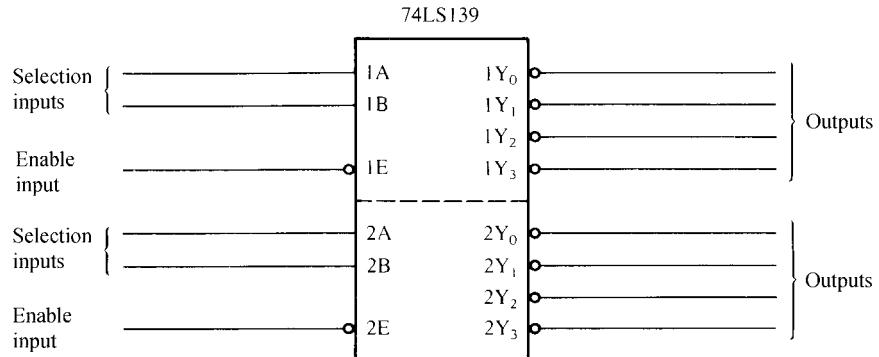
A more complicated decoder using the 74LS139 decoder appears in Figure 10–17. This circuit uses a 128K × 8 EPROM (271000) and a 128K × 8 SRAM (621000). The EPROM is decoded at memory locations E0000H–FFFFFH and the SRAM is decoded at addresses 00000H–1FFFFH. This is fairly typical of a small embedded system, where the EPROM is located at the top of the memory space and the SRAM at the bottom.

Output $\overline{Y0}$ of decoder U1A activates the SRAM whenever address bits A_{17} and A_{18} are both logic 0s if the $IO/\overline{M}$ signal is a logic 0 and address line A_{19} is a logic 0. This selects the SRAM for any address between 00000H and 1FFFFH. The second decoder (U1B) is slightly more complicated because the NAND gate (U4B) selects the decoder when $IO/\overline{M}$ is a logic 0 while A_{19} is a logic 1. This selects the EPROM for addresses E0000H through FFFFFH.

PLD Programmable Decoders

This section of the text explains the use of the programmable logic device, or PLD, as a decoder. There are three SPLD (**simple PLD**) devices that function in the same manner but have different names: PLA (**programmable logic array**), PAL (**programmable array logic**), and GAL (**gated array logic**). Although these devices have been in existence since the mid-1970s, they have only appeared in memory system and digital designs since the early 1990s. The PAL and the PLA are fuse-programmed, as is the PROM, and some PLD devices are erasable devices (as are EPROMs). In essence, all three devices are arrays of logic elements that are programmable.

FIGURE 10–16 The pin-out and truth table of the 74LS139, dual 2-to-4 line decoder.

Inputs			Outputs			
$\overline{E}$	A	B	$\overline{Y_0}$	$\overline{Y_1}$	$\overline{Y_2}$	$\overline{Y_3}$
0	0	0	0	1	1	1
0	0	1	1	0	1	1
0	1	0	1	1	0	1
0	1	1	1	1	1	0
1	X	X	1	1	1	1

Other PLDs are also available, such as CPLDs (**complex programmable logic devices**), FPGAs (**field programmable gate arrays**), and FPICs (**field programmable interconnect**). These types of PLDs are much more complex than the SPLDs that are used more commonly in designing a complete system. If the concentration is on decoding addresses, the SPLD is used and if the concentration is on a complete system, then the CPLD, FPLG, or FPIC is used to implement the design. These devices are also referred to as an ASIC (**application-specific integrated circuit**).

Combinatorial Programmable Logic Arrays. One of the two basic types of PALs is the combinatorial programmable logic array. This device is internally structured as a programmable array of combinational logic circuits. Figure 10–18 illustrates the internal structure of the PAL16L8 that is constructed with AND/OR gate logic. This device, which is representative of a PLD, has 10 fixed inputs, two fixed outputs, and six pins that are programmable as inputs of outputs. Each output signal is generated from a seven-input OR gate that has an AND gate attached to each input. The outputs of the OR gates pass through a three-state inverter that defines each output as an AND/NOR function. Initially, all of the fuses connect all of the vertical/horizontal connections illustrated in Figure 10–18. Programming is accomplished by blowing fuses to connect various inputs to the OR gate array. The wired-AND function is performed at each input connection, which allows a product term of up to 16 inputs. A logic expression using the PAL16L8 can have up to seven product terms with up to 16 inputs NORed together to generate the output expression. This device is ideal as a memory address decoder because of its structure. It is also ideal because the outputs are active low.

Fortunately, we don't have to choose the fuses by number for programming, as was customary when this device was first introduced. A PAL is programmed with a software package such as PALASM, the PAL assembler program. More recently, PLD design is accomplished using HDL (**hardware description language**) or VHDL (**verilog HDL**). The VHDL language and its syntax are currently the industry standard for programming PLD devices. Example 10–5 shows a program that decodes the same areas of memory as decoded in Figure 10–17. Note that this program
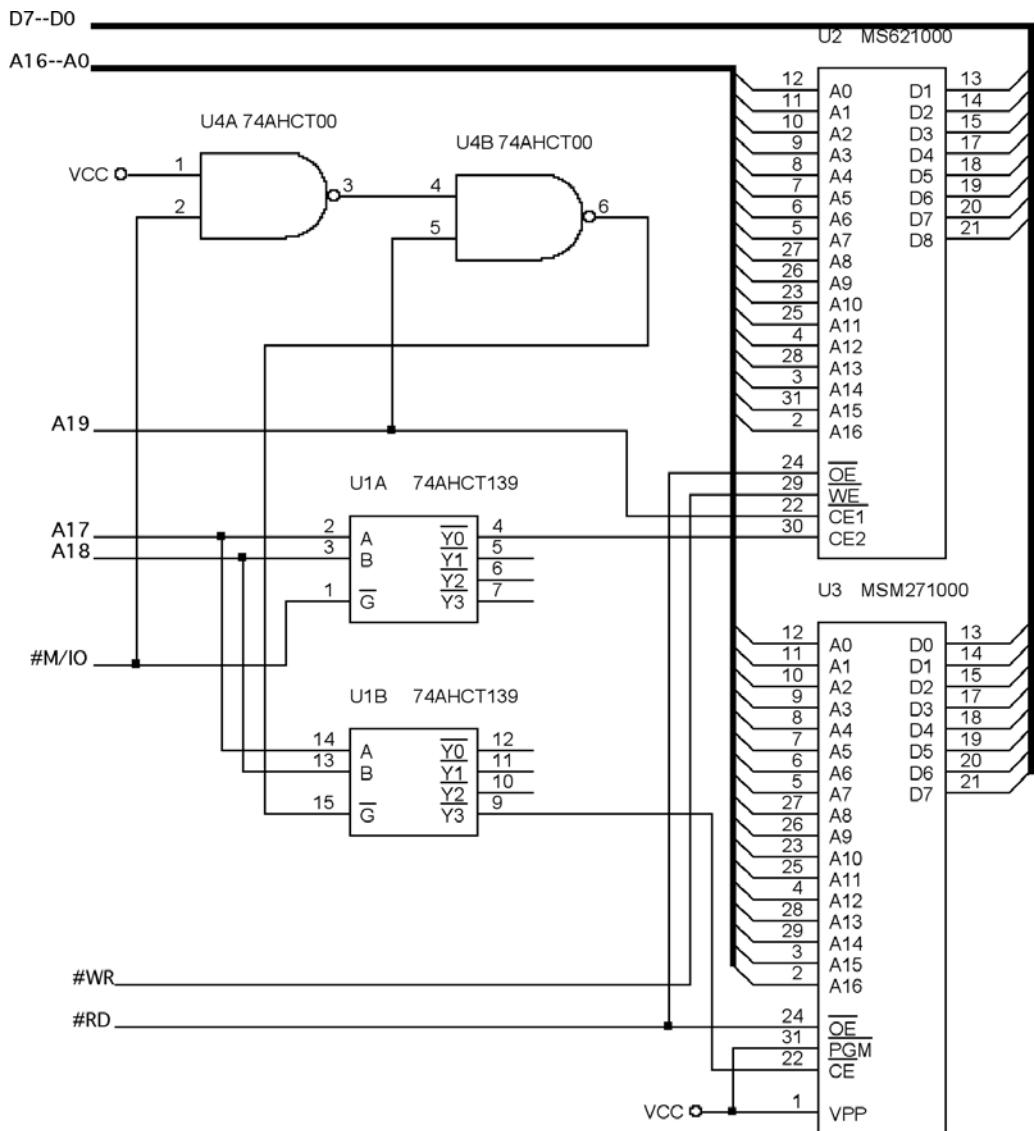
FIGURE 10–17   A sample memory system constructed with a 74HCT139.

was developed by using a text editor such as EDIT, available with Microsoft DOS version 7.1 with XP or NotePad in Windows. The program can also be developed by using an editor than comes with any of the many programming packages for PLDs. Various editors attempt to ease the task of defining the pins, but we believe it is easier to use NotePad and the listing as shown.

EXAMPLE 10–5

```
-- VHDL code for the decoder of Figure 10-17

library ieee;
use ieee.std_logic_1164.all;

entity DECODER_10_19 is

port (
      A19, A18, A17, MIO: in STD_LOGIC;
      ROM, RAM, AX19: out STD_LOGIC
```

Logic Diagram   16L8
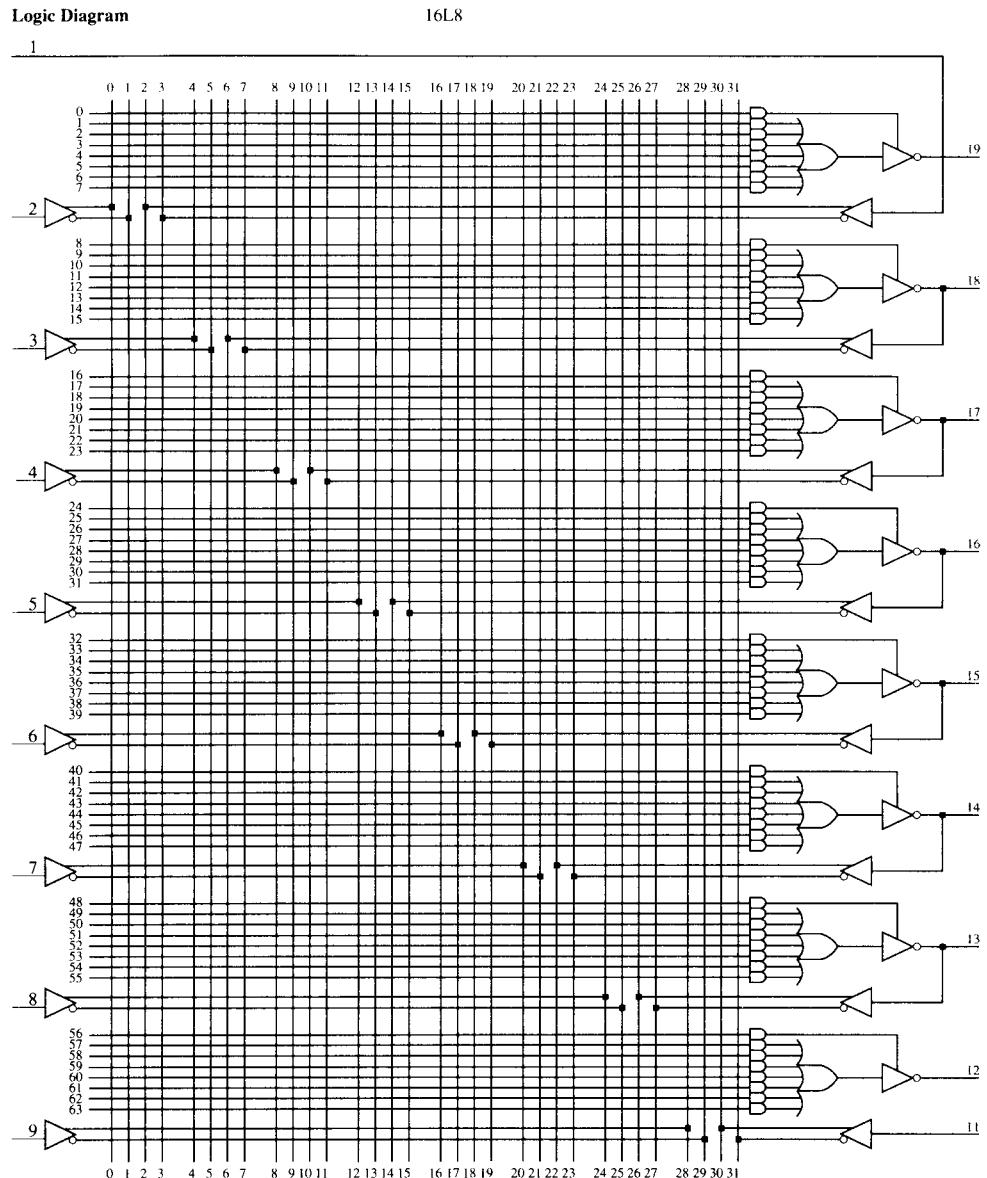
FIGURE 10–18   The PAL16L8. (Copyright Advanced Micro Devices, Inc., 1988. Reprinted with permission of copyright owner. All rights reserved.)

```
);

end;

architecture V1 of DECODER_10_19 is

begin

     ROM <= A19 or A18 or A17 or MIO;
     RAM <= not (A18 and A17 and (not MIO));
     AX19 <= not A19;

end V1;
```

Comments in VHDL programming begin with a pair of minus signs as illustrated in the first line of the VHDL code in Example 10–5. The library and use statements specify the standard IEEE library using standard logic. The entity statement names the VHDL module, in this case DECODER_10_17. The port statements define the in, out, and in-out pins used in the equations for the logic expression, which appears in the begin block. A_{19}, A_{18}, A_{17}, and MIO (this signal cannot be defined as $IO/\overline{M}$ so it was called MIO) are defined as input pins and ROM and RAM are the output pins for connection to the $\overline{CS}$ pins on the memory devices. The architecture statement merely refers to the version (V_1) of this design. Finally, the equations for the design are placed in the begin block. Each output pin has its own equation. The keyword not is used for logical inversion and the keyword and is used for the logical and operation. In this case the ROM equation causes the ROM pin to become a logic zero only when the A_{19}, A_{18}, A_{17}, and MIO are all logic zeros (00000H–1FFFFH). The RAM equation causes the RAM pin to become a logic zero when A_{18} and A_{17} are all ones at the same time that MIO is a logic zero. A_{19} is connected to the active high CE2 pin after being inverted by the PLD. The RAM is selected for addresses 60000H–7FFFFH. See Figure 10–19 for the PLD realization of Example 10–5.
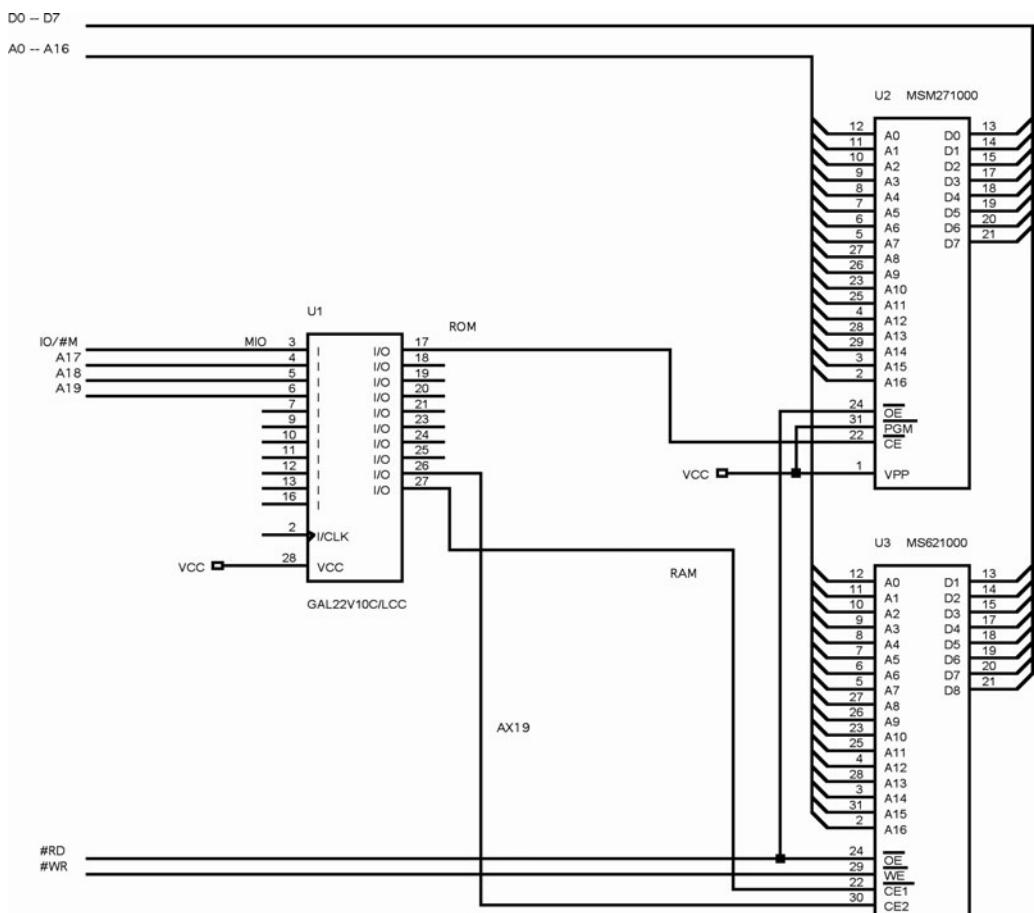
FIGURE 10–19   A RAM and ROM interface using a programmable logic device.

10–3   8088 AND 80188 (8-BIT) MEMORY INTERFACE

This text contains separate sections on memory interfacing for the 8088 and 80188 with their 8-bit data buses; the 8086, 80186, 80286, and 80386SX with their 16-bit data buses; the 80386DX and 80486 with their 32-bit data buses; the Pentium–Core2 with their 64-bit data buses. Separate sections are provided because the methods used to address the memory are slightly different in microprocessors that contain different data bus widths. Hardware engineers or technicians who wish to broaden their expertise in interfacing 16-bit, 32-bit, and 64-bit memory interface should cover all sections. This section is much more complete than the sections on the 16-, 32-, and 64-bit-wide memory interface, which cover material not explained in the 8088/80188 section.

In this section, we examine the memory interface to both RAM and ROM and explain the error-correction code (ECC), which is still is currently available to memory system designers. Many home computer systems do not use ECC because of the cost, but business machines often do use it.

Basic 8088/80188 Memory Interface

The 8088 and 80188 microprocessors have an 8-bit data bus, which makes them ideal to connect to the common 8-bit memory devices available today. The 8-bit memory size makes the 8088, and especially the 80188, ideal as a simple controller. For the 8088/80188 to function correctly with the memory, however, the memory system must decode the address to select a memory component. It must also use the $\overline{RD}$, $\overline{WR}$, and IO/$\overline{M}$ control signals provided by the 8088/80188 to control the memory system.

The minimum mode configuration is used in this section and is essentially the same as the maximum mode system for memory interface. The main difference is that, in maximum mode, the IO/$\overline{M}$ signal is combined with $\overline{RD}$ to generate the $\overline{MRDC}$ signal, and IO/$\overline{M}$ is combined with $\overline{WR}$ to generate the $\overline{MWTC}$ signal. The maximum mode control signals are developed inside the 8288 bus controller. In minimum mode, the memory sees the 8088 or the 80188 as a device with 20 address connections (A_{19}–A_0), eight data bus connections (AD_7–AD_0), and the control signals IO/$\overline{M}$, $\overline{RD}$, and $\overline{WR}$.

Interfacing EPROM to the 8088.   You will find this section very similar to Section 10–2 on decoders. The only difference is that, in this section, we discuss wait states and the use of the IO/$\overline{M}$ signal to enable the decoder.

Figure 10–20 illustrates an 8088/80188 microprocessor connected to three 27256 EPROMs, 32K × 8 memory devices. The 27256 has one more address input (A_{15}) than the 27128 and twice the memory. The 74HCT138 decoder in this illustration decodes three 32K × 8 blocks of memory for a total of 96K × 8 bits of the physical address space for the 8088/80188.

The decoder (74HCT138) is connected a little differently than might be expected because the slower version of this type of EPROM has a memory access time of 450 ns. Recall from Chapter 9 that when the 8088 is operated with a 5 MHz clock, it allows 460 ns for the memory to access data. Because of the decoder's added time delay (8 ns), it is impossible for this memory to function within 460 ns. In order to correct this problem, the output from the NAND gate can be used to generate a signal to enable the decoder and a signal for the wait state generator, covered in Chapter 9. (Note that the 80188 can internally insert from 0 to 15 wait states without any additional external hardware, so it does not require this NAND gate.) With a wait state inserted every time this section of the memory is accessed, the 8088 will allow 660 ns for the EPROM to access data. Recall that an extra wait state adds 200 ns (1 clock) to the access time. The 660 ns is ample time for a 450 ns memory component to access data, even with the delays introduced by the decoder and any buffers added to the data bus. The wait states are inserted in this system for memory locations C0000H–FFFFFH. If this creates a problem, a three-input OR gate can be added to the three outputs of the decoder to generate a wait signal only for the actual addresses for this system (E8000H–FFFFFH).
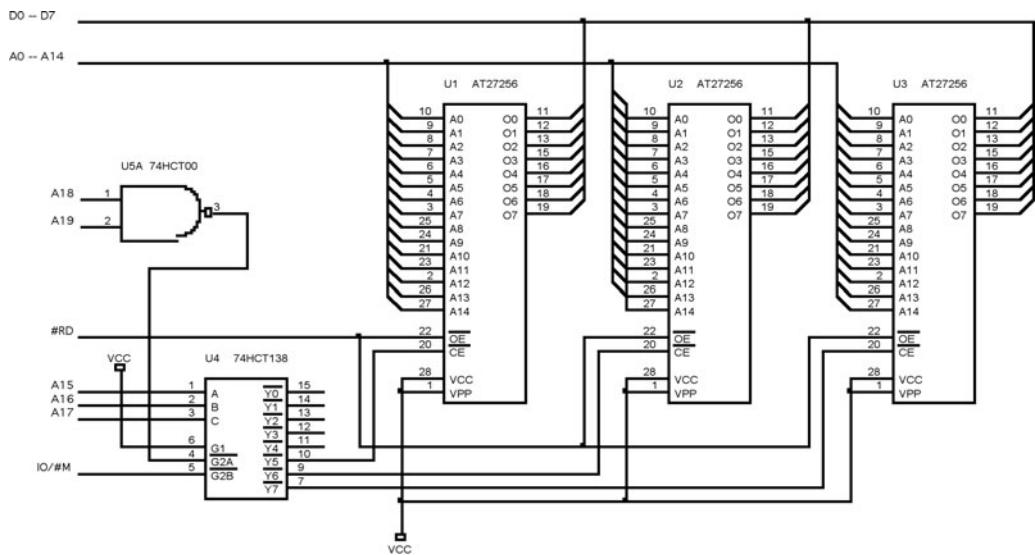
FIGURE 10–20   Three 27256 EPROMs interfaced to the 8088 microprocessor.

Notice that the decoder is selected for a memory address range that begins at location E8000H and continues through location FFFFFH—the upper 96K bytes of memory. This section of memory is an EPROM because FFFF0H is where the 8088 starts to execute instructions after a hardware reset. We often call location FFFF0H the **cold-start location**. The software stored in this section of memory would contain a JMP instruction at location FFFF0H that jumps to location E8000H so the remainder of the program can execute. In this circuit, U_1 is decoded at addresses E8000H–EFFFFH, U_2 is decoded at F0000H–F7FFFH, and U_3 is decoded at F8000H–FFFFFH.

Interfacing RAM to the 8088.   RAM is a little easier to interface than EPROM because most RAM memory components do not require wait states. An ideal section of the memory for the RAM is the very bottom, which contains vectors for interrupts. Interrupt vectors (discussed in more detail in Chapter 12) are often modified by software packages, so it is rather important to encode this section of the memory with RAM.

Figure 10–21 shows sixteen 62256, $32K \times 8$ static RAMs interfaced to the 8088, beginning at memory location 00000H. This circuit board uses two decoders to select the 16 different RAM memory components and a third to select the other decoders for the appropriate memory sections. Sixteen 32K RAMs fill memory from location 00000H through location 7FFFFH, for 512K bytes of memory.

The first decoder (U_4) in this circuit selects the other two decoders. An address beginning with 00 selects decoder U_3 and an address that begins with 01 selects decoder U_9. Notice that extra pins remain at the output of decoder U_4 for future expansion. These pins allow more $256K \times 8$ blocks of RAM for a total of $1M \times 8$, simply by adding the RAM and the additional secondary decoders.

Also notice from the circuit in Figure 10–21 that all the address inputs to this section of memory are buffered, as are the data bus connections and control signals $\overline{RD}$ and $\overline{WR}$. Buffering is important when many devices appear on a single board or in a single system. Suppose that three other boards like this are plugged into a system. Without the buffers on each board, the load on the system address, data, and control buses would be enough to prevent proper operation. (Excessive loading causes the logic 0 output to rise above the 0.8 V maximum allowed in a system.) Buffers are normally used if the memory will contain additions at some future date. If the memory will never grow, then buffers may not be needed.
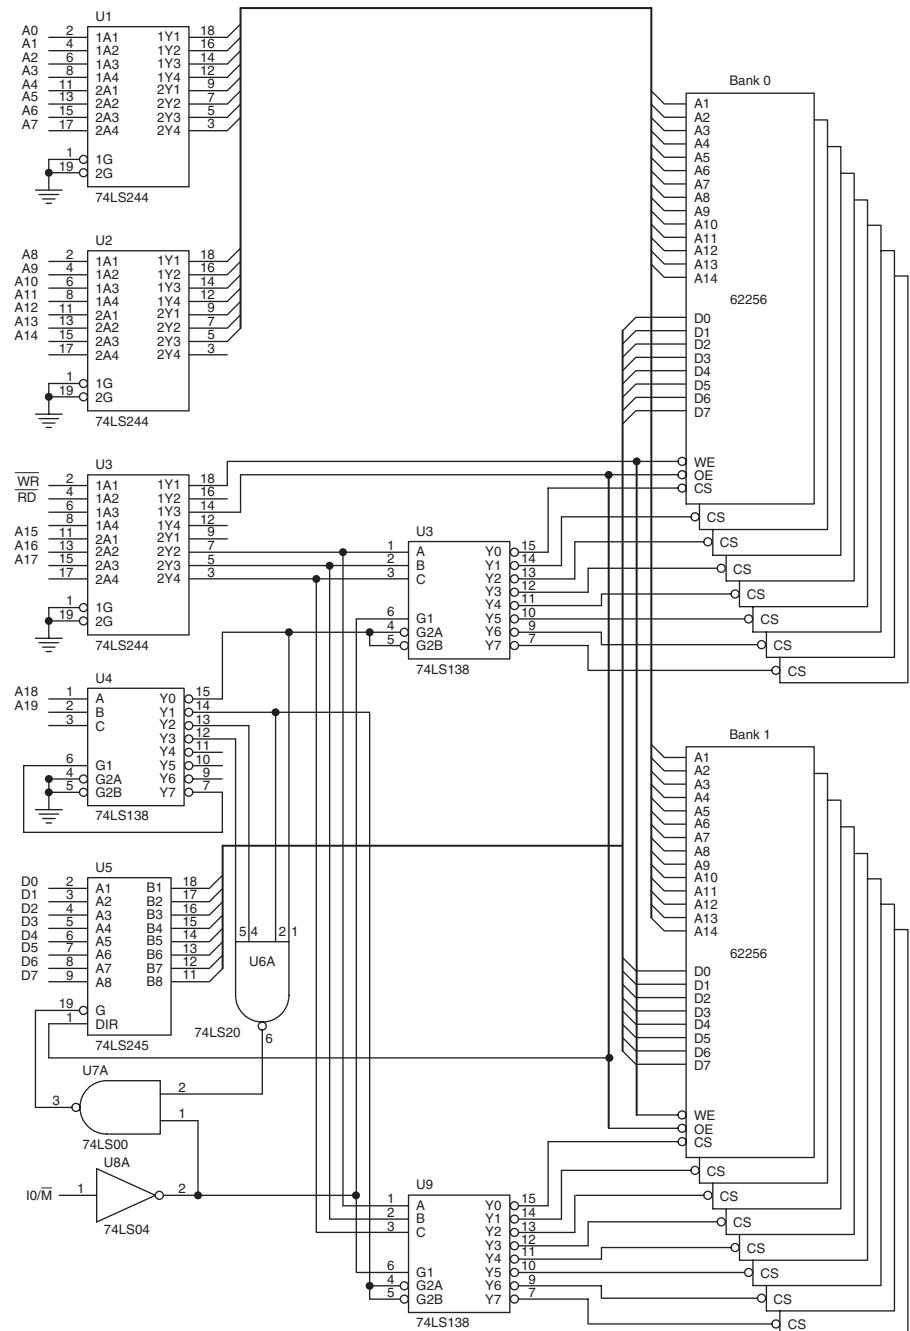
FIGURE 10–21   A 512K-byte static memory system using 16 62255 SRAMs.

Interfacing Flash Memory

Flash memory (EEPROM) is becoming commonplace for storing setup information on video cards, as well as for storing the system BIOS in the personal computer. It even finds application in MP3 audio players and USB pen drives. Flash memory is also found in many other applications to store information that is only changed occasionally.

The only difference between a flash memory device and SRAM is that the flash memory device requires a 12V programming voltage to erase and write new data. The 12V can be available either at the power supply or a 5V to 12V converter designed for use with flash memory can be obtained. The newest versions of flash memory are erased with a 5.0V or even a 3.3V signal so that a converter is not needed.

EEPROM is available as either a memory device with a parallel interface or as devices that are serial. The serial device is extremely small and is not suited for memory expansion, but as an I/O device it can store information such as in a flash drive. This section of the text details both memory types.

Figure 10–22 illustrates the 28F400 Intel flash memory device interfaced to the 8088 microprocessor using its parallel interface. The 28F400 can be used as either a $512K \times 8$ memory device or as a $256K \times 16$ memory device. Because it is interfaced to the 8088, its configuration is $512K \times 8$. Notice that the control connections on this device are identical to that of an SRAM: $\overline{CE}$, $\overline{OE}$, and $\overline{WE}$. The only new pins are V_{PP}, which is connected to 12V for erase and programming; $\overline{PWD}$, which selects the power-down mode when a logic 0 and is also used for programming; and $\overline{BYTE}$, which selects byte (0) or word (1) operation. Note that the pin DQ_{15} functions as the least significant address input when operated in the byte mode. Another difference is the amount of time required to accomplish a write operation. The SRAM can perform a write operation in as little as 10 ns, but the flash memory requires approximately 0.4 seconds to erase a byte. The topic of programming the flash memory device is covered in Chapter 11, along with I/O devices. The flash memory device has internal registers that are programmed by using I/O techniques not yet explained. This chapter concentrates on its interface to the microprocessor.

Notice in Figure 10–22 that the decoder chosen is the 74LS139 because only a simple decoder is needed for a flash memory device this large. The decoder uses address connection A_{19} and $IO/\overline{M}$ as inputs. The A_{15} signal selects the flash memory for locations 80000H through FFFFFH, and $IO/\overline{M}$ enables the decoder.
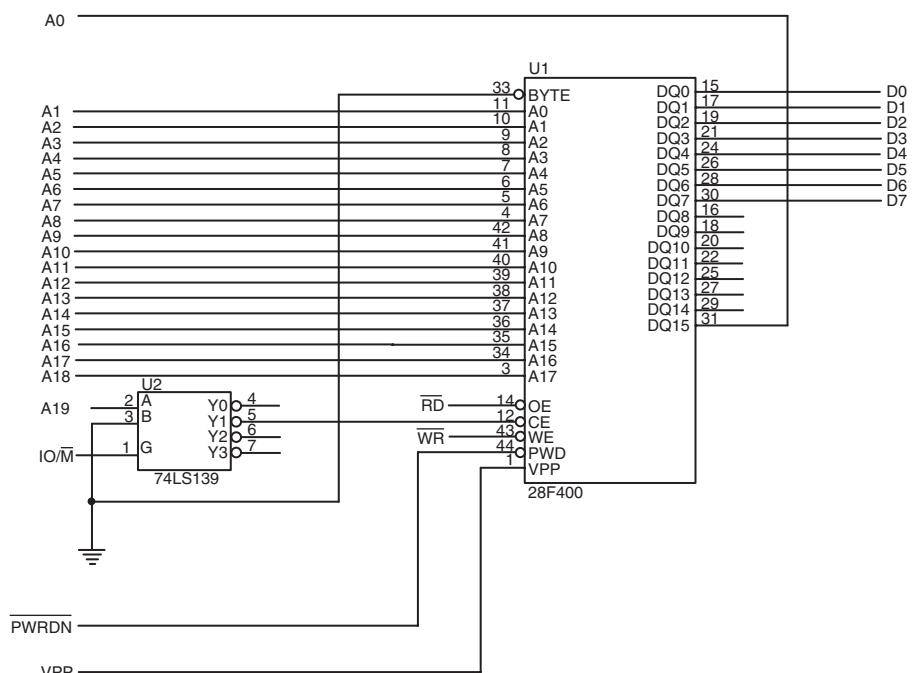
FIGURE 10–22   The 28F400 flash memory device interfaced to the 8088 microprocessor.
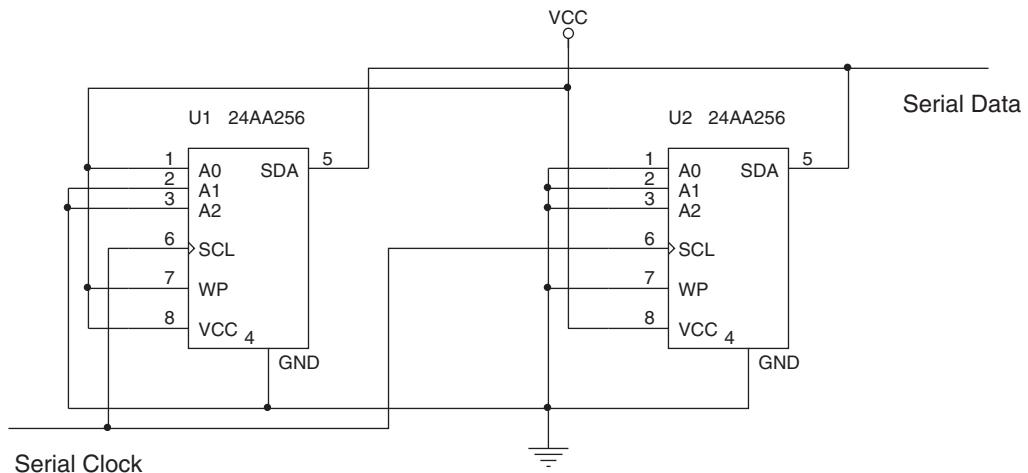
FIGURE 10–23   A serial EEPROM interface.

As mentioned, many newer flash memory devices use a serial interface to reduce the cost of the integrated circuit because of fewer pins and a smaller size. Serial flash memory is available in sizes to 4G bytes and has comparable speeds and erase times with the parallel flash devices. Most modern flash memory functions from 5V or 3.3V without the need for a higher programming voltage and has a life of 1,000,000 erases with a storage time of 200 years.

Figure 10–23 illustrates a small serial flash device (a 256K device, organized as a $32K \times 8$ memory). The three address pins are hardwired to allow more than one device to be placed on a serial bus. In the illustration U_1 is wired at address 001 and U_2 is wired at address 000. Not shown in the illustration is a pull-up resistor that is needed for the serial data connection. The pull-up may be located in the microprocessor or it may need to be connected externally, depending on the microprocessor and interface connected to the memory.
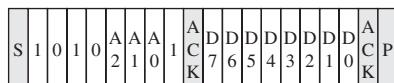
This memory interface has two signal lines. One is a serial clock (SCL) and the other is a bidirectional serial data line (SDA). The clock frequency can be anything up to 400 KHz, so this type of memory is not meant to replace the main memory in a system. It is fast enough for music or other low-speed data. The serial interface is explained in Chapter 11.

Figure 10–24 depicts the basic serial data format for the serial EEPROM. The serial data contains the address (the A_0, A_1, A_2 pins) in the first byte as well as a device code of 1010, which represents the EEPROM. Other serial devices have different device codes. This is followed by the memory location and the data in additional bytes.

Error Correction

Error-correction schemes have been around for a long time, but integrated circuit manufacturers have only recently started to produce error-correcting circuits. One such circuit is the 74LS636, an 8-bit error correction and detection circuit that corrects any single-bit memory read error and flags any 2-bit error called SECDED (**single error correction/double error correction**). This device is found in high-end computer systems because of the cost of implementing a system that uses error correction.

The newest computer systems are now using DDR memory with ECC (error-correction code). The scheme to correct the errors that might occur in these memory devices is identical to the scheme discussed in this text.

Write Address

S	1	0	1	0	A2	A1	A0	0	ACK	x	A14	A13	A12	A11	A10	A9	A8	ACK	A7	A6	A5	A4	A3	A2	A1	A0	ACK

Followed by for a Read Byte

S	1	0	1	0	A2	A1	A0	1	ACK	D7	D6	D5	D4	D3	D2	D1	D0	ACK	P

or Followed by for a Write Byte

D7	D6	D5	D4	D3	D2	D1	D0	ACK	P

S = Start
P = Stop
ACK = acknowledge

FIGURE 10–24 The data signals to the serial EEPROM for a read or a write.

The 74LS636 corrects errors by storing five parity bits with each byte of memory data. This does increase the amount of memory required, but it also provides automatic error correction for single-bit errors. If more than two bits are in error, this circuit may not detect it. Fortunately, this is rare, and the extra effort required to correct more than a single-bit error is very expensive and not worth the effort. Whenever a memory component fails completely, its bits are all high or all low. In this case, the circuit flags the processor with a multiple-bit error indication.

Figure 10–25 depicts the pin-out of the 74LS636. Notice that it has eight data I/O pins, five check bit I/O pins, two control inputs (SO and SI), and two error outputs: single-error flag (SEF) and double-error flag (DEF). The control inputs select the type of operation to be performed and are listed in the truth table of Table 10–1.

When a single error is detected, the 74LS636 goes through an error-correction cycle. It places a 01 on S_0 and S_1 by causing a wait and then a read following error correction.

Figure 10–26 illustrates a circuit used to correct single-bit errors with the 74LS636 and to interrupt the processor through the NMI pin for double-bit errors. To simplify the illustration, we depict only one 2K × 8 RAM and a second 2K × 8 RAM to store the 5-bit check code.

The connection of this memory component is different from that of the previous example. Notice that the S or $\overline{CS}$ pin is grounded, and data bus buffers control the flow to the system bus. This is necessary if the data are to be accessed from the memory before the $\overline{RD}$ strobe goes low.

On the next negative edge of the clock after the $\overline{RD}$ signal, the 74LS636 checks the single-error flag (SEF) to determine whether an error has occurred. If it has, a correction cycle causes the single-error defect to be corrected. If a double error occurs, an interrupt request is generated by the double-error flag (DEF) output, which is connected to the NMI pin of the microprocessor.

Modern DDR error-correction memory (ECC) does not actually have logic circuitry on board that detects and corrects errors. Since the Pentium, the microprocessor incorporates the logic circuitry to detect/correct errors provided the memory can store the extra 8 bits required for storing the ECC code. ECC memory is 72-bits wide using the additional 8 bits to store the ECC code. If an error occurs, the microprocessor runs the correction cycle to correct the error. Some memory devices such as Samsung memory also perform an internal error check. The Samsung ECC uses 3 bytes to check every 256 bytes of memory, which is far more efficient. Additional information on the Samsung ECC algorithm is available at the Samsung website.

pin assignments

J, N PACKAGES			
1	DEF	11	CB4
2	DB0	12	nc
3	DB1	13	CB3
4	DB2	14	CB2
5	DB3	15	CB1
6	DB4	16	CB0
7	DB5	17	S0
8	DB6	18	S1
9	DB7	19	SEF
10	GND	20	V_{CC}

(a)

functional block diagram

(b)
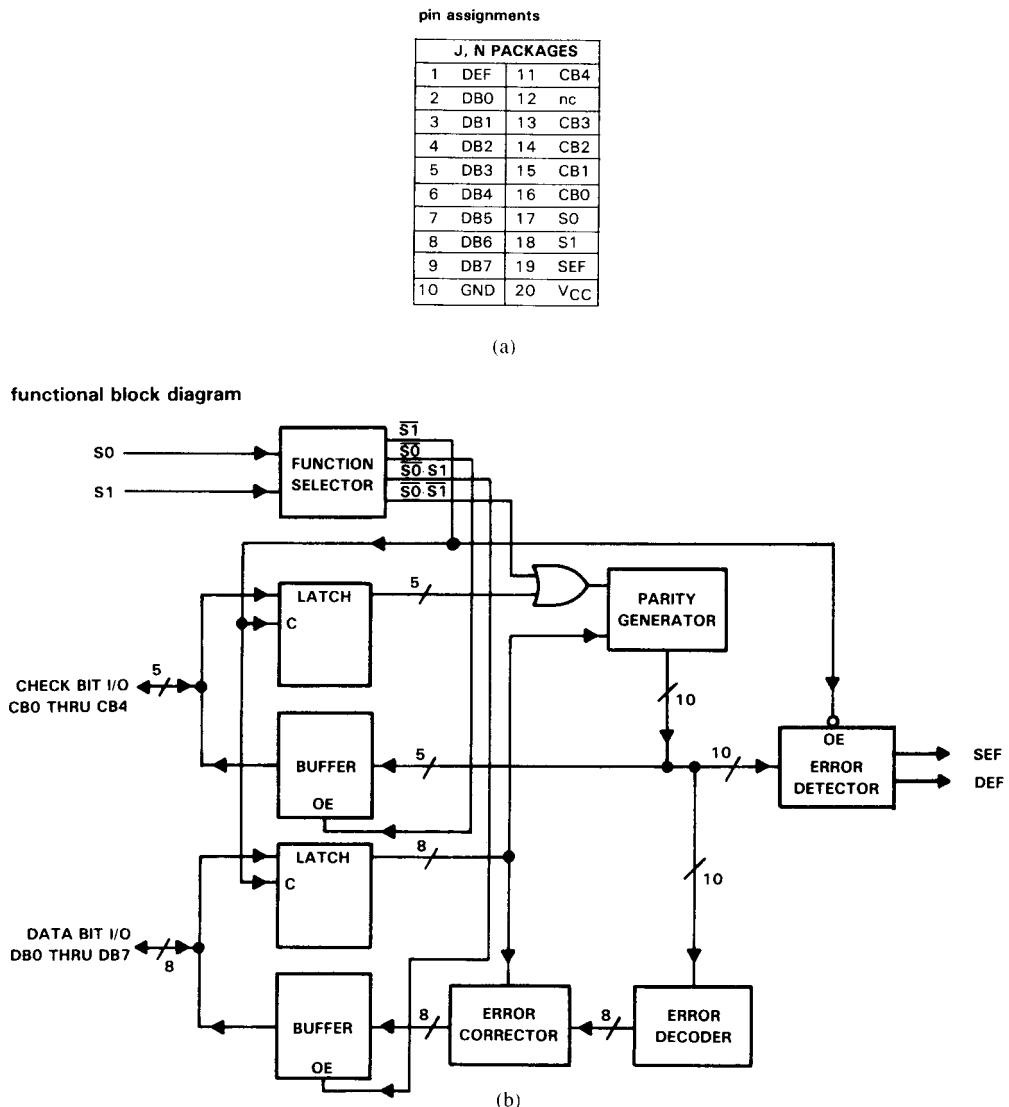
FIGURE 10–25 (a) The pin connections of the 74LS636. (b) The block diagram of the 74LS636. (Courtesy of Texas Instruments Incorporated.)

TABLE 10–1 Control bits S_0 and S_1 for the 74LS636.

S_1	S_0	Function	SEF	DEF
0	0	Write check word	0	0
0	1	Correct data word	*	*
1	0	Read data	0	0
1	1	Latch data	*	*

*Note: These levels are determined by the error type.
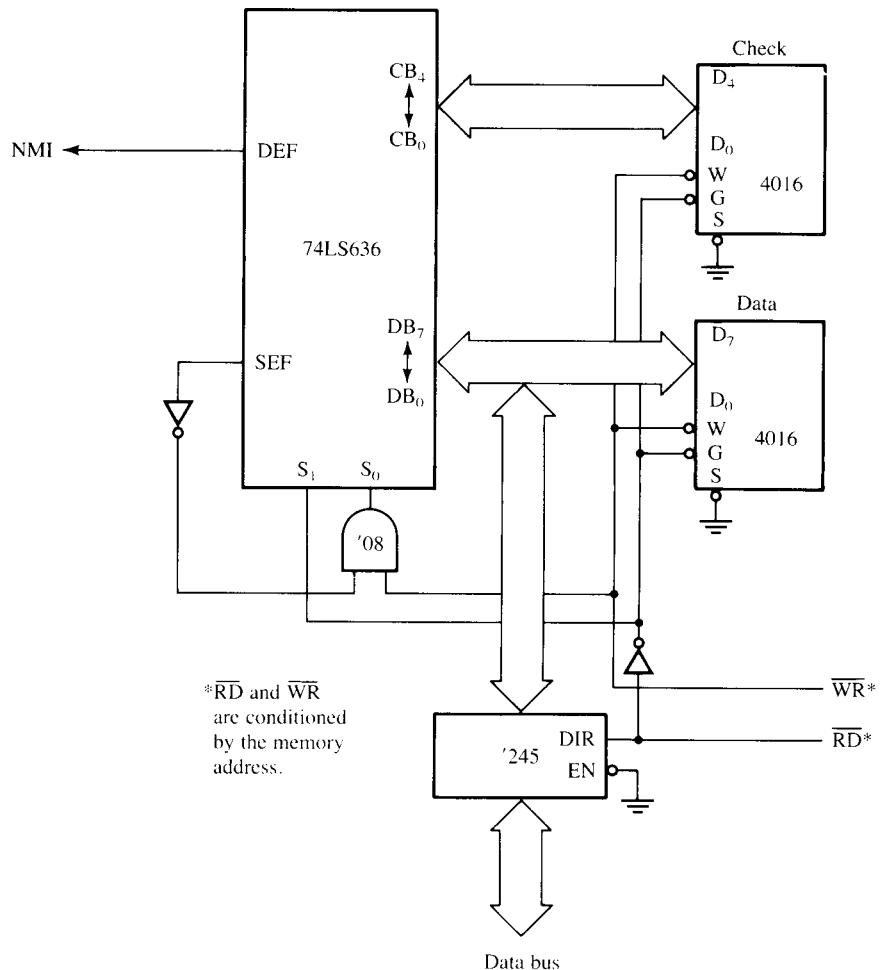
FIGURE 10–26   An error detection and correction circuit using the 74LS636.
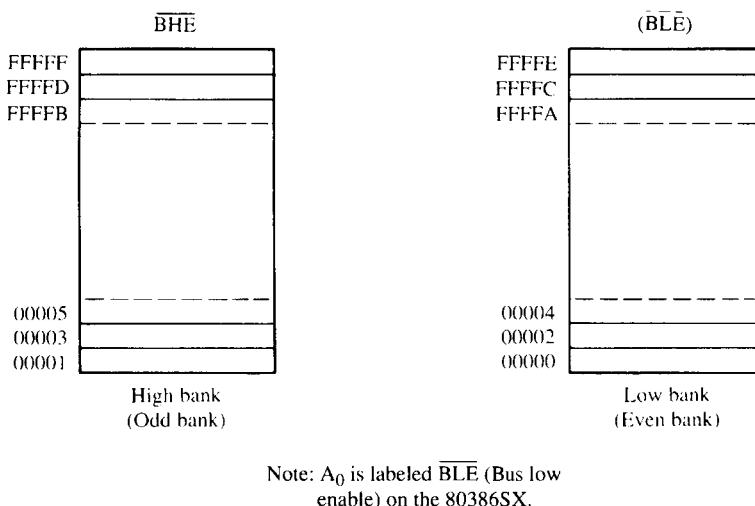
10–4   8086, 80186, 80286, AND 80386SX (16-BIT) MEMORY INTERFACE

The 8086, 80186, 80286, and 80386SX microprocessors differ from the 8088/80188 in three ways: (1) The data bus is 16 bits wide instead of 8 bits wide as on the 8088; (2) the IO/$\overline{\text{M}}$ pin of the 8088 is replaced with an M/$\overline{\text{IO}}$ pin; and (3) there is a new control signal called bus high enable ($\overline{\text{BHE}}$). The address bit A_0 or $\overline{\text{BLE}}$ is also used differently. (Because this section is based on information provided in Section 10–3, it is extremely important that you read the previous section first.) A few other differences exist between the 8086/80186 and the 80286/80386SX. The 80286/80386SX contains a 24-bit address bus (A_{23}–A_0) instead of the 20-bit address bus (A_{19}–A_0) of the 8086/80186. The 8086/80186 contain an M/$\overline{\text{IO}}$ signal, while the 80286 system and 80386SX microprocessor contain control signals $\overline{\text{MRDC}}$ and $\overline{\text{MWTC}}$ instead of $\overline{\text{RD}}$ and $\overline{\text{WR}}$.

16-Bit Bus Control

The data bus of the 8086, 80186, 80286, and 80386SX is twice as wide as the bus for the 8088/80188. This wider data bus presents us with a unique set of problems. The 8086, 80186, 80286, and 80386SX must be able to write data to any 16-bit location—or any 8-bit location.

FIGURE 10–27 The high (odd) and low (even) 8-bit memory banks of the 8086/80286/80386SX microprocessors.

$\overline{BHE}$

FFFFF
FFFFD
FFFFB

00005
00003
00001

High bank
(Odd bank)

$(\overline{BLE})$

FFFFE
FFFFC
FFFFA

00004
00002
00000

Low bank
(Even bank)

Note: A_0 is labeled $\overline{BLE}$ (Bus low enable) on the 80386SX.

This means that the 16-bit data bus must be divided into two separate sections (or **banks**) that are 8 bits wide so that the microprocessor can write to either half (8-bit) or both halves (16-bit). Figure 10–27 illustrates the two banks of the memory. One bank (**low bank**) holds all the even-numbered memory locations, and the other bank (**high bank**) holds all the odd-numbered memory locations.

The 8086, 80186, 80286, and 80386SX use the $\overline{BHE}$ signal (high bank) and the A_0 address bit or $\overline{BLE}$ (bus low enable) to select one or both banks of memory used for the data transfer. Table 10–2 depicts the logic levels on these two pins and the bank or banks selected.

Bank selection is accomplished in two ways: (1) A separate write signal is developed to select a write to each bank of the memory, or (2) separate decoders are used for each bank. As a careful comparison reveals, the first technique is by far the least costly approach to memory interface for the 8086, 80186, 80286, and 80386SX microprocessors. The second technique is only used in a system that must achieve the most efficient use of the power supply.

Separate Bank Decoders.   The use of separate bank decoders is often the least effective way to decode memory addresses for the 8086, 80186, 80286, and 80386SX microprocessors. This method is sometimes used, but it is difficult to understand why in most cases. One reason may be to conserve energy, because only the bank or banks selected are enabled. This is not always the case, as with the separate bank read and write signals that are discussed later.

Figure 10–28 illustrates two 74LS138 decoders used to select 64K RAM memory components for the 80386SX microprocessor (24-bit address). Here, decoder U_2 has the $\overline{BLE}$ (A_0) attached to $\overline{G2A}$, and decoder U_3 has the $\overline{BHE}$ signal attached to its $\overline{G2A}$ input. Because the decoder will not activate until all of its enable inputs are active, decoder U_2 activates only for a 16-bit operation or an 8-bit operation from the low bank. Decoder U_3 activates for a 16-bit operation

TABLE 10–2 Memory bank selection using $\overline{BHE}$ and $\overline{BLE}$ (A_0).

$\overline{BHE}$	$\overline{BLE}$	Function
0	0	Both banks enabled for a 16-bit transfer
0	1	High bank enabled for an 8-bit transfer
1	0	Low bank enabled for an 8-bit transfer
1	1	No bank enabled
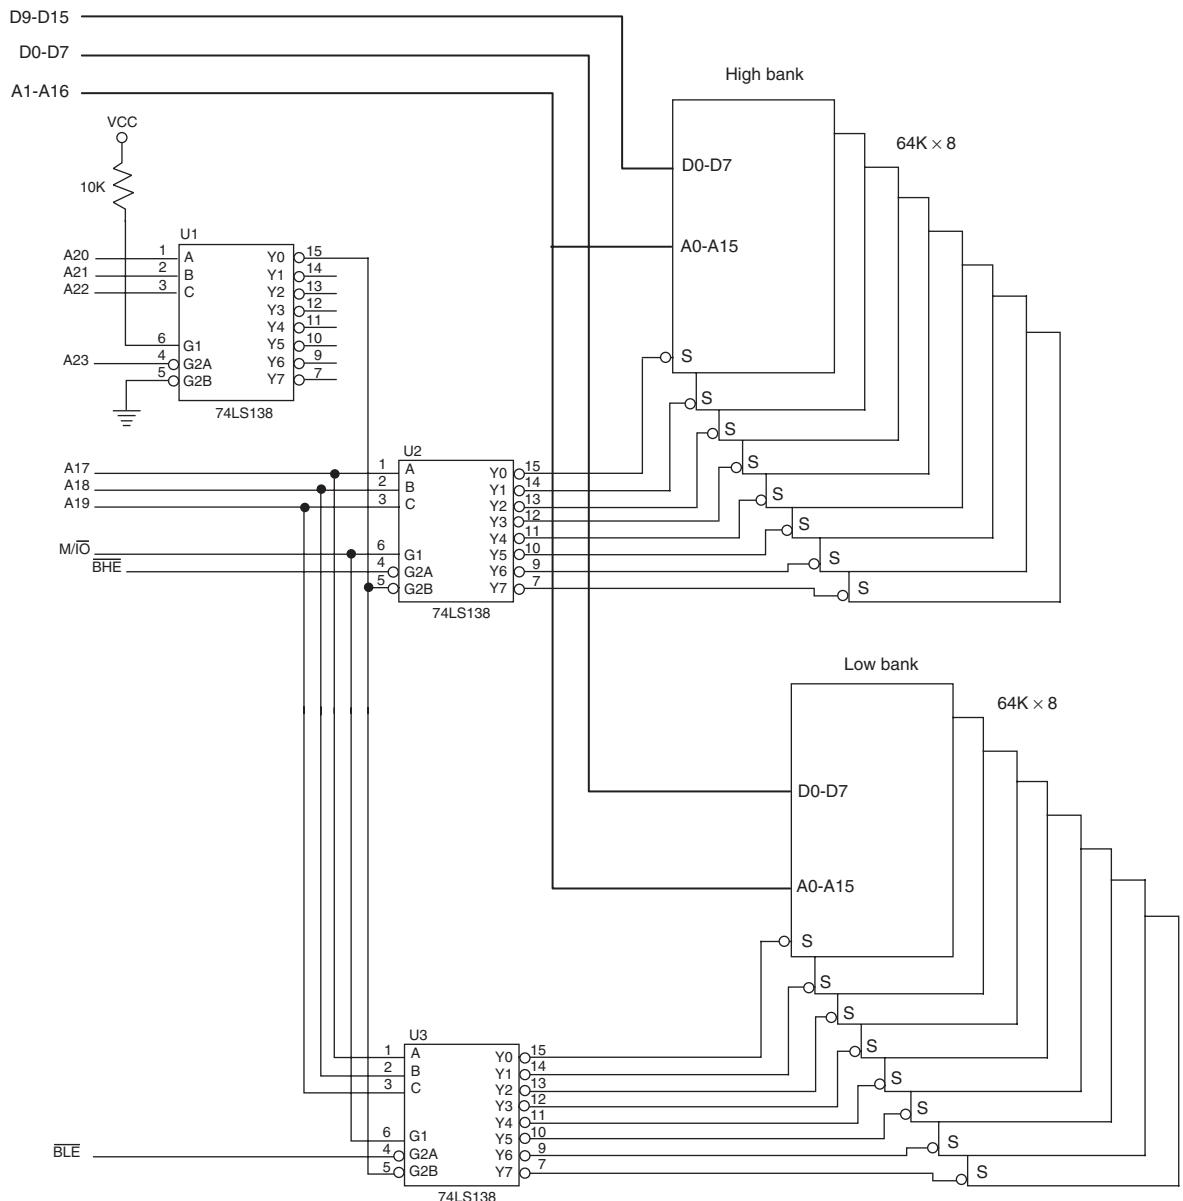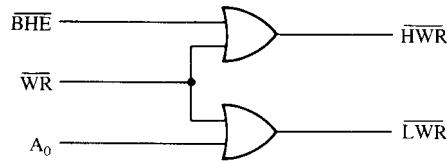
FIGURE 10–28 Separate bank decoders.

or an 8-bit operation to the high bank. These two decoders and the sixteen 64K-byte RAMs they control represent a 1M range of the 80386SX memory system. Decoder U_1 enables U_2 and U_3 for memory address range 000000H–0FFFFFH.

Notice in Figure 10–28 that the A_0 address pin does not connect to the memory because it does not exist on the 80386SX microprocessor. Also notice that address bus bit position A_1 is connected to memory address input A_0, A_2 is connected to A_1, and so forth. The reason is that A_0 from the 8086/80186 (or $\overline{\text{BLE}}$ from the 80286/80386SX) is already connected to decoder U_2 and does not need to be connected again to the memory. If A_0 or $\overline{\text{BLE}}$ is attached to the A_0 address pin of memory, every other memory location in each bank of memory would be used. This means that half of the memory is wasted if A_0 or $\overline{\text{BLE}}$ is connected to A_0.

FIGURE 10–29 The memory bank write selection input signals: $\overline{\text{HWR}}$ (high bank write) and $\overline{\text{LWR}}$ (low bank write).

Separate Bank Write Strobes. The most effective way to handle bank selection is to develop a separate write strobe for each memory bank. This technique requires only one decoder to select a 16-bit-wide memory, which often saves money and reduces the number of components in a system.

Why not also generate separate read strobes for each memory bank? This is usually unnecessary because the 8086, 80186, 80286, and 80386SX microprocessors read only the byte of data that they need at any given time from half of the data bus. If 16-bit sections of data are always presented to the data bus during a read, the microprocessor ignores the 8-bit section that it doesn't need, without any conflicts or special problems.

Figure 10–29 depicts the generation of separate 8086 write strobes for the memory. Here, a 74LS32 OR gate combines A_0 with $\overline{\text{WR}}$ for the low bank selection signal ($\overline{\text{LWR}}$), and $\overline{\text{BHE}}$ combines with $\overline{\text{WR}}$ for the high bank selection signal ($\overline{\text{HWR}}$). Write strobes, for the 80286/80386SX, are generated by using the $\overline{\text{MWTC}}$ signal instead of $\overline{\text{WR}}$.

A memory system that uses separate write strobes is constructed differently from either the 8-bit system (8088) or the system using separate memory banks. Memory in a system that uses separate write strobes is decoded as 16-bit-wide memory. For example, suppose that a memory system will contain 64K bytes of SRAM memory. This memory requires two 32K-byte memory devices (62256) so that a 16-bit-wide memory can be constructed. Because the memory is 16 bits wide and another circuit generates the bank write signals, address bit A_0 becomes a don't care. In fact, A_0 is not even a pin on the 80386SX microprocessor.

Example 10–6 shows how a 16-bit-wide memory stored at locations 060000H–06FFFFH is decoded for the 80286 or 80386 microprocessor. Memory in this example is decoded, so bit A_0 is a don't care for the decoder. Bit positions A_1–A_{15} are connected to memory component address pins A_0–A_{14}. The decoder (GAL22V10) enables both memory devices by using address connection A_{23}–A_{15} to select memory whenever address 06XXXXH appears on the address bus.

EXAMPLE 10–6

```
0000 0110 0000 0000 0000 0000 = 060000H
              to
0000 0110 1111 1111 1111 1111 = 06FFFFH

0000 0110 XXXX XXXX xxxx XXXX = 06XXXXH
```

Figure 10–30 illustrates this simple circuit by using a GAL22V10 to both decode memory and generate the separate write strobe. The program for the GAL22V10 decoder is illustrated in Example 10–7. Notice that not only is the memory selected, but both the lower and upper write strobes are also generated by the PLD.

EXAMPLE 10–7

```
-- VHDL code for the decoder of Figure 10-30

library ieee;
use ieee.std_logic_1164.all;

entity DECODER_10_30 is

port (
      A23, A22, A21, A20, A19, A18, A17, A16, A0, BHE, MWTC: in STD_LOGIC;
      SEL, LWR, HWR: out STD_LOGIC
);

end;
```

```
architecture V1 of DECODER_10_30 is

begin
        SEL <= A23 or A22 or A21 or A20 or A19 or (not A18) or (not A17) or A16;
        LWR <= A0 or MWTC;
        HWR <= BHE or MWTC;

end V1;
```

Figure 10–31 depicts a small memory system for the 8086 microprocessor that contains an EPROM section and a RAM section. Here, there are four 27128 EPROMs (16K × 8) that compose a 32K × 16-bit memory at locations F0000–FFFFFH and four 62256 (32K × 8) RAMs that compose 64K × 16-bit memory at locations 00000H–1FFFFH. (Remember that even though the memory is 16 bits wide, it is still numbered in bytes.)

This circuit uses a 74HC139 dual 2-to-4 line decoder that selects EPROM with one half and RAM with the other half. It decodes memory that is 16 bits wide, not 8 bits, as before. Notice that the $\overline{RD}$ strobe is connected to all the EPROM $\overline{OE}$ inputs and all RAM $\overline{OE}$ input pins. This is done because even if the 8086 is reading only 8 bits of data, the application of the remaining 8 bits to the data bus has no effect on the operation of the 8086.

The $\overline{LWR}$ and $\overline{HWR}$ strobes are connected to different banks of the RAM memory. Here, it does matter whether the microprocessor is doing a 16-bit or an 8-bit write. If the 8086 writes a 16-bit number to memory, both $\overline{LWR}$ and $\overline{HWR}$ go low and enable the $\overline{WE}$ pins in both memory banks. But if the 8086 does an 8-bit write, only one of the write strobes goes low, writing to only one memory bank. Again, the only time that the banks make a difference is for a memory write operation.

Notice that an EPROM decoder signal is sent to the 8086 wait state generator because EPROM memory usually requires a wait state. The signal comes from the NAND gate used to select the EPROM decoder section, so that if EPROM is selected, a wait state is requested.

Figure 10–32 illustrates a memory system connected to the 80386SX microprocessor by using a GAL22V10 as a decoder. This interface contains 256K bytes of EPROM in the form of
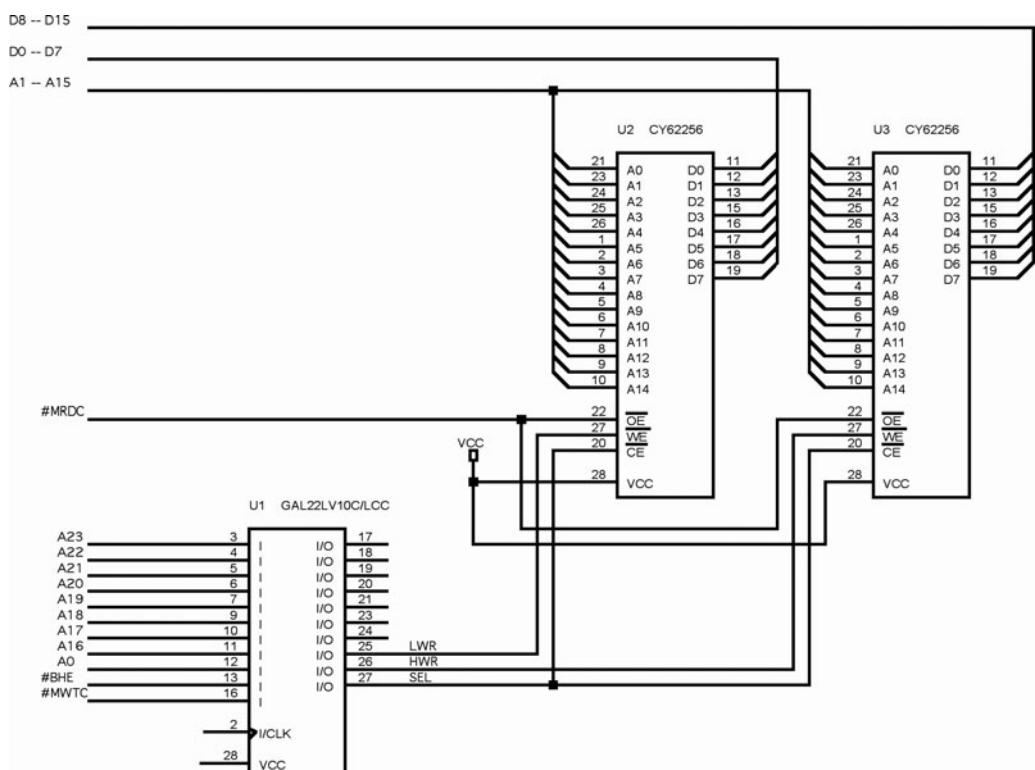
FIGURE 10–30   A 16-bit-wide memory interfaced at memory locations 06000H–06FFFH.
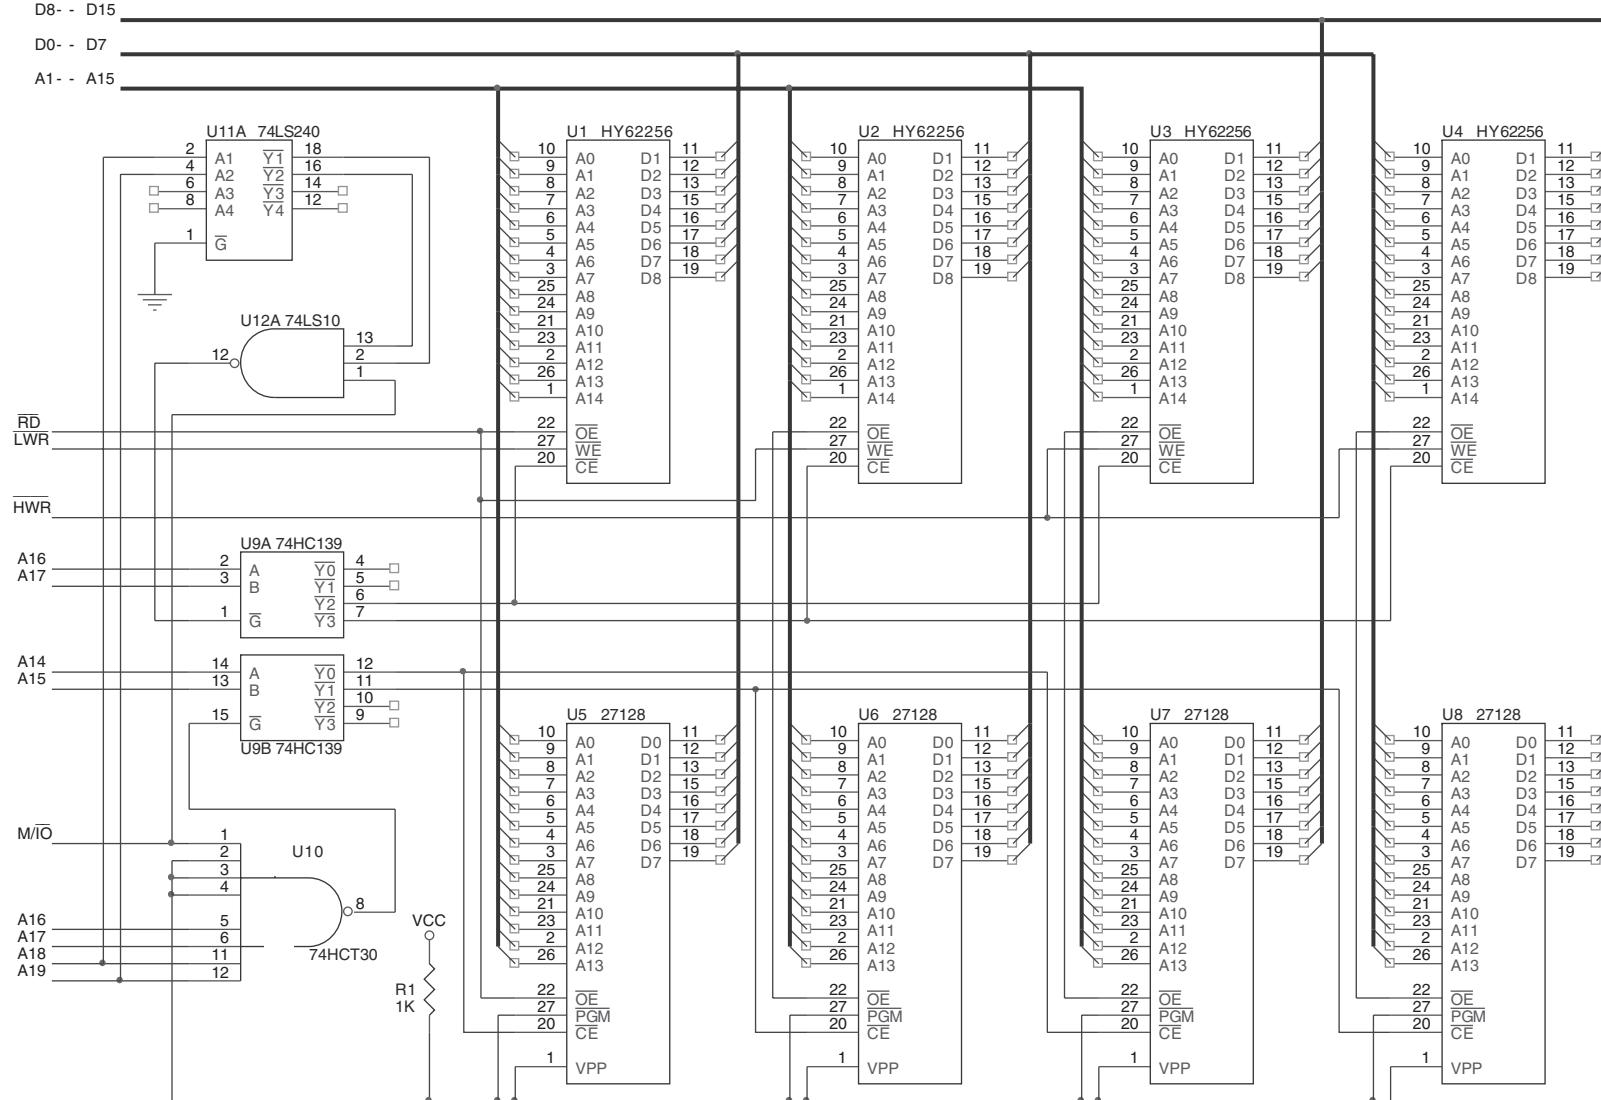
FIGURE 10–31 A memory system for the 8086 that contains a 64K-byte EPROM and a 128K-byte SRAM.

D8–D15

D0–D7

A1–A16

U2 M27512

U3 M27512

U4 M27512

U5 M27512

#MRDC

#MWTC

A0

A23

A22

A21

A20

A19

A18

A17

A16

#BHE

I/CLK

VCC

U1 GAL22LV10C/LCC

LWR
HWR
RB0
RB1
RB2
RB3

U6 HM62256

U7 HM62256

U8 HM62256

U9 HM62256
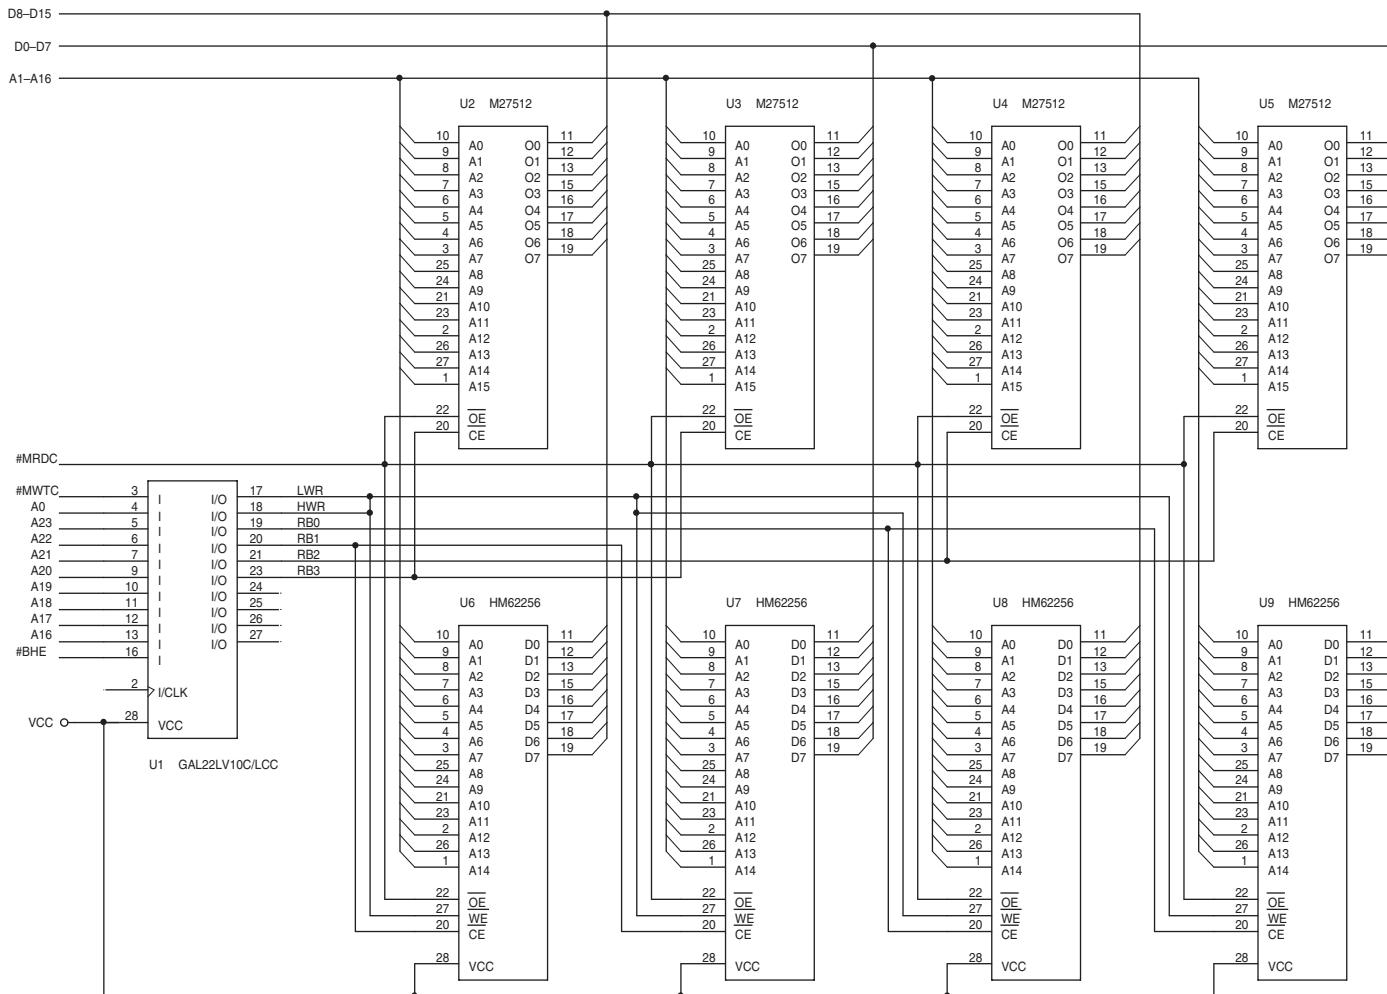
FIGURE 10–32   An 80386SX memory system containing 256K of EPROM and 128K of SRAM.

four 27512 (64K × 8) EPROMs and 128K bytes of SRAM memory found in four 62256 (32K × 8) SRAMs.

Notice in Figure 10–32 that the PLD also generates the memory bank write signals $\overline{\text{LWR}}$ and $\overline{\text{HWR}}$. As can be gleaned from this circuit, the number of components required to interface memory has been reduced to just one, in most cases (the PLD). The program listing for the PLD is located in Example 10–8. The PLD decodes the 16-bit-wide memory addresses at locations 000000H–01FFFFH for the SRAM and locations FC0000H–FFFFFFH for the EPROM.

EXAMPLE 10–8

```
-- VHDL code for the decoder of Figure 10–32

library ieee;
use ieee.std_logic_1164.all;

entity DECODER_10_32 is

port (
        A23, A22, A21, A20, A19, A18, A17, A16, A0, BHE, MWTC: in STD_LOGIC;
        LWR, HWR, RB0, RB1, RB2, RB3: out STD_LOGIC
);

end;

architecture V1 of DECODER_10_32 is

begin
        LWR <= A0 or MWTC;
        HWR <= BHE or MWTC;
        RB0 <= A23 or A22 or A21 or A20 or A19 or A18 or A17 or A16;
        RB1 <= A23 or A22 or A21 or A20 or A19 or A18 or A17 or not(A16));
        RB2 <= not(A23 and A22 and A21 and A20 and A19 and A18 and A17);
        RB3 <= not(A23 and A22 and A21 and A20 and A19 and A18 and not(A17));
end V1;
```

10–5   80386DX AND 80486 (32-BIT) MEMORY INTERFACE

As with 8- and 16-bit memory systems, the microprocessor interfaces to memory through its data bus and control signals that select separate memory banks. The only difference with a 32-bit memory system is that the microprocessor has a 32-bit data bus and four banks of memory, instead of one or two. Another difference is that both the 80386DX and 80486 (both SX and DX) contain a 32-bit address bus that usually requires PLD decoders instead of integrated decoders because of the sizable number of address bits.

Memory Banks

The memory banks for both the 80386DX and 80486 microprocessors are illustrated in Figure 10–33. Notice that these large memory systems contain four 8-bit-wide banks that each contain up to 1G bytes of memory. Bank selection is accomplished by the bank selection signals $\overline{\text{BE3}}$, $\overline{\text{BE2}}$, $\overline{\text{BE1}}$, and $\overline{\text{BE0}}$. If a 32-bit number is transferred, all four banks are selected; if a 16-bit number is transferred, two banks (usually $\overline{\text{BE3}}$ and $\overline{\text{BE2}}$ or $\overline{\text{BE1}}$ and $\overline{\text{BE0}}$) are selected; and if 8 bits are transferred, a single bank is selected.

BE3
FFFFFFFF
FFFFFFFB
FFFFFFF7
Bank 3
0000000B
00000007
00000003
D31   D24

BE2
FFFFFFFE
FFFFFFFA
FFFFFFF6
Bank 2
0000000A
00000006
00000002
D23   D16

BE1
FFFFFFFD
FFFFFFF9
FFFFFFF5
Bank 1
00000009
00000005
00000001
D15   D8

BE0
FFFFFFFC
FFFFFFF8
FFFFFFF4
Bank 0
00000008
00000004
00000000
D7   D0
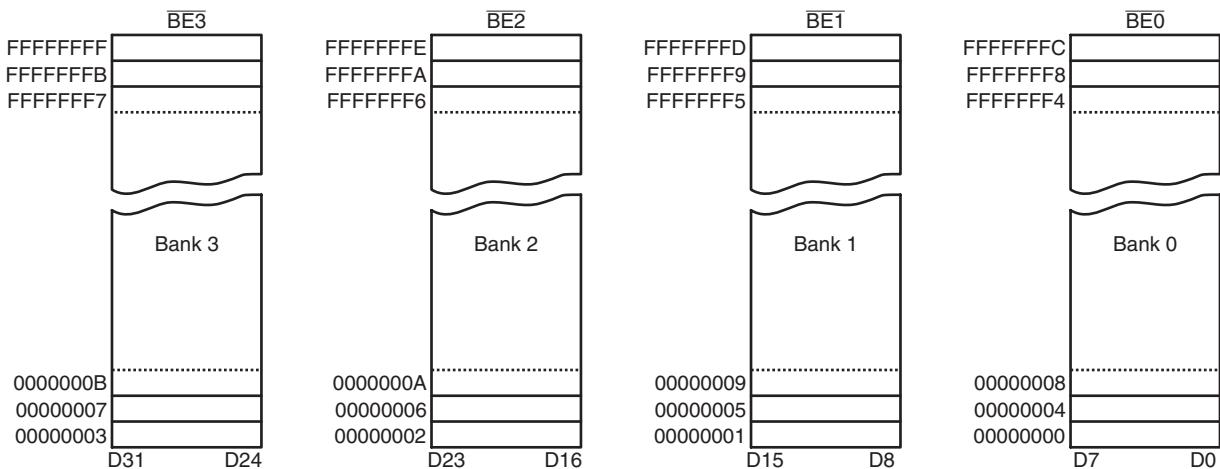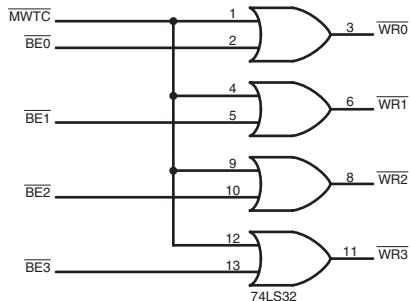
FIGURE 10–33   The memory organization for the 80386DX and 80486 microprocessors.

FIGURE 10–34   Bank write signals for the 80386DX and 80486 microprocessors.

MWTC ——— 1
BE0 ——— 2   3 — WR0

4
BE1 ——— 5   6 — WR1

9
BE2 ——— 10   8 — WR2

12
BE3 ——— 13   11 — WR3

74LS32

As with the 8086/80286/80386SX, the 80386DX and 80486 require separate write strobe signals for each memory bank. These separate write strobes are developed, as illustrated in Figure 10–34, by using a simple OR gate or other logic component.

32-Bit Memory Interface

As can be gathered from the prior discussion, a memory interface for the 80386DX or 80486 requires that we generate four bank write strobes and decode a 32-bit address. There are no integrated decoders, such as the 74LS138, that can easily accommodate a memory interface for the 80386DX or 80486 microprocessors. Note that address bits A_0 and A_1 are don't cares when 32-bit-wide memory is decoded. These address bits are used within the microprocessor to generate the bank enable signals. Notice that the address bus connection A_2 connects to memory address pin A_0. This occurs because there is no A_0 or A_1 pin on the 80486 microprocessor.

Figure 10–35 shows a 512K × 8 SRAM memory system for the 80486 microprocessor. This interface uses eight 64K × 8 SRAM memory devices, a PLD, and an OR gate. The OR gate is required because of the number of address connections found on the microprocessor. This system places the SRAM memory at locations 02000000H–0203FFFFH. The program for the PLD device is found in Example 10–9.
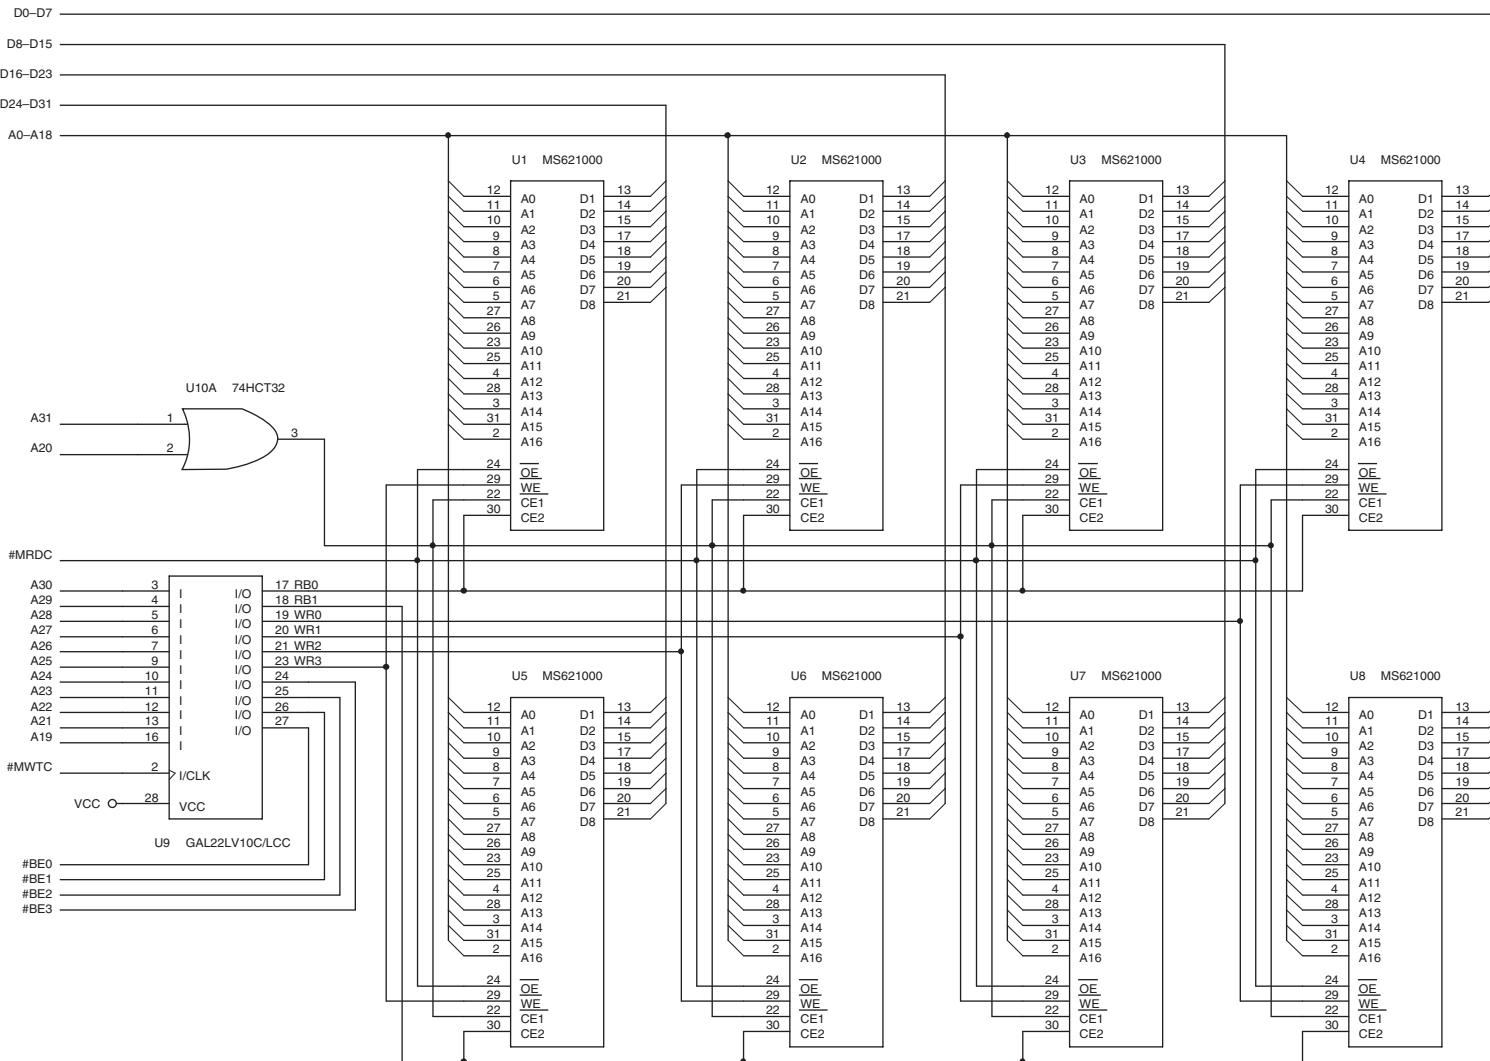
FIGURE 10–35   A small 512K-byte SRAM memory system for the 80486 microprocessor.

EXAMPLE 10–9

```
library ieee;
use ieee.std_logic_1164.all;

entity DECODER_10_35 is

port (
        A30, A29, A28, A27, A26, A25, A24, A23, A22, A21, A19, BE0, BE1, BE2,
            BE3, MWTC: in STD_LOGIC;
        RB0, RB1, WR0, WR1, WR2, WR3: out STD_LOGIC
);

end;

architecture V1 of DECODER_10_35 is

begin

        WR0 <= BE0 or MWTC;
        WR1 <= BE1 or MWTC;
        WR2 <= BE2 or MWTC;
        WR3 <= BE3 or MWTC;
        RB0 <= A30 or A29 or A28 or A27 or A26 or A25 or A24 or A23 or A22
            or A 21 or A19;
        RB1 <= A30 or A29 or A28 or A27 or A26 or A25 or A24 or A23 or A22
            or A 21 or not(A19);

end V1;
```

Although not mentioned in this section of the text, the 80386DX and 80486 microprocessors operate with very high clock rates that usually require wait states for memory access. Access time calculations for these microprocessors are discussed in Chapters 17 and 18. The interface provides a signal used with the wait state generator that is not illustrated in this section of the text. Other devices with these higher speed microprocessors are cache memory and interleaved memory systems. These are also presented in Chapter 17 with the 80386DX and 80486 microprocessors.

10–6 PENTIUM THROUGH CORE2 (64-BIT) MEMORY INTERFACE

The Pentium through Core2 microprocessors (except for the P24T version of the Pentium) contain a 64-bit data bus, which requires either eight decoders (one per bank) or eight separate write signals. In most systems, separate write signals are used with this microprocessor when interfacing memory. Figure 10–36 illustrates the Pentium's memory organization and its eight memory banks. Notice that this is almost identical to the 80486, except that it contains eight banks instead of four.

As with earlier versions of the Intel microprocessor, this organization is required for upward memory compatibility. The separate write strobe signals are obtained by combining the bank enable signals with the $\overline{\text{MWTC}}$ signal, which is generated by combining the M/$\overline{\text{IO}}$ with W/$\overline{\text{R}}$. The circuit employed for bank write signals appears in Figure 10–37. As can be imagined, we often find a PLD used for bank write signal generation.

64-Bit Memory Interface

Figure 10–38 illustrates a small Pentium–Core2 memory system. This system uses a PLD to decode the memory address. This system contains eight 27C4001 EPROM memory devices (512K × 8), interfaced to the Pentium–Core2 at locations FFC00000H through FFFFFFFFH.
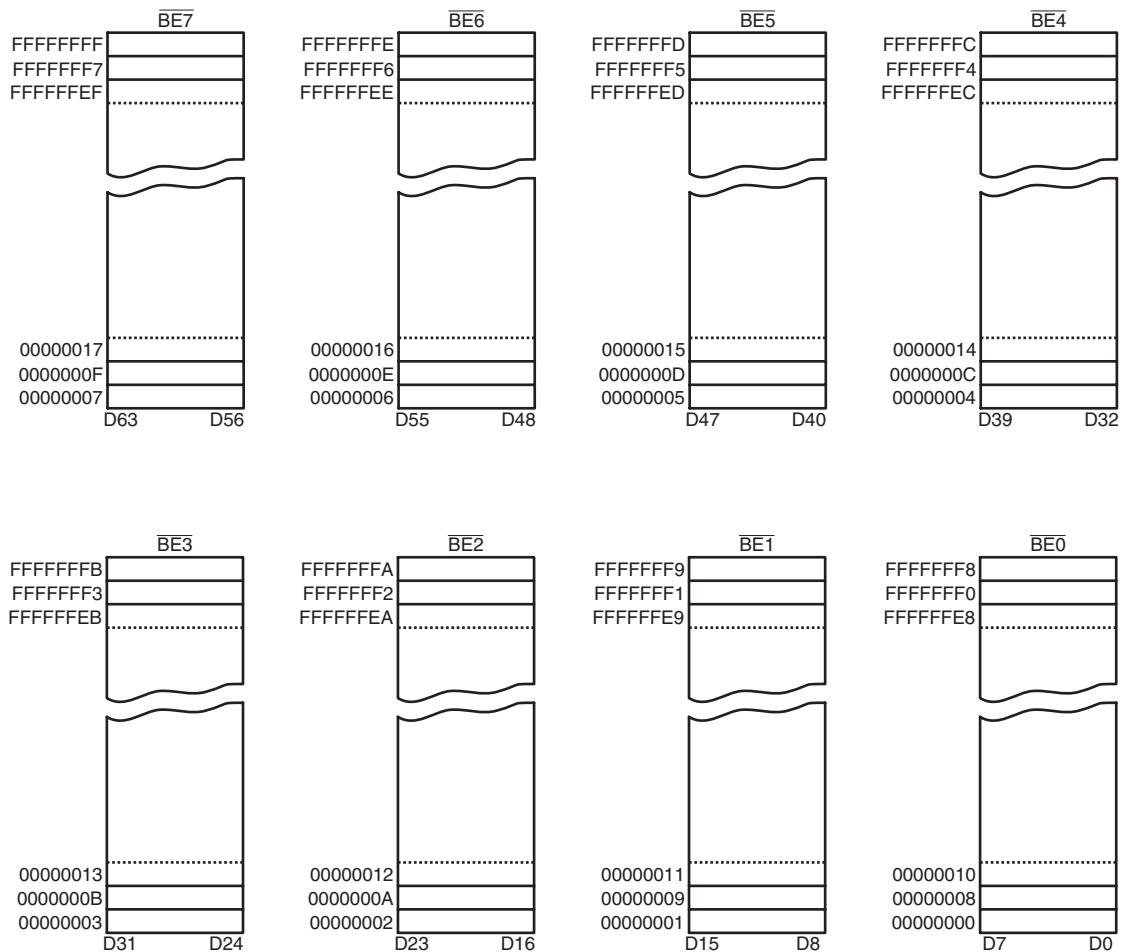
FIGURE 10–36   The memory organization of the Pentium–Core2 microprocessors.

This is a total memory size of 4M bytes organized so that each bank contains two memory components. Note that the Pentium Pro through the Core2 can be configured with 36 address connections, allowing up to 64G of memory. The Pentium 4 and the Core2 can also be configured in the flat mode and may contain up to 40 address connections. (The Core2 contains only 36.)

Memory decoding, as illustrated in Example 10–10, is similar to the earlier examples, except that with the Pentium–Core2 the rightmost three address bits (A_2–A_0) are ignored. In this case, the decoder selects sections of memory that are 64 bits wide and contain 4M bytes of EPROM memory.

The A_0 address input of each memory device connects to the A_3 address output of the Pentium and above. This A_1 address input of each memory device connects to the A_4 address output of the Pentium and above. This skewed address connection continues until the A_{18} address input to the memory is connected to the A_{22} address output of the Pentium. Address positions A_{22}–A_{31} are decoded by PLD. The program for the PLD device is listed in Example 10–10 for memory locations FFC00000H–FFFFFFFFH.

368
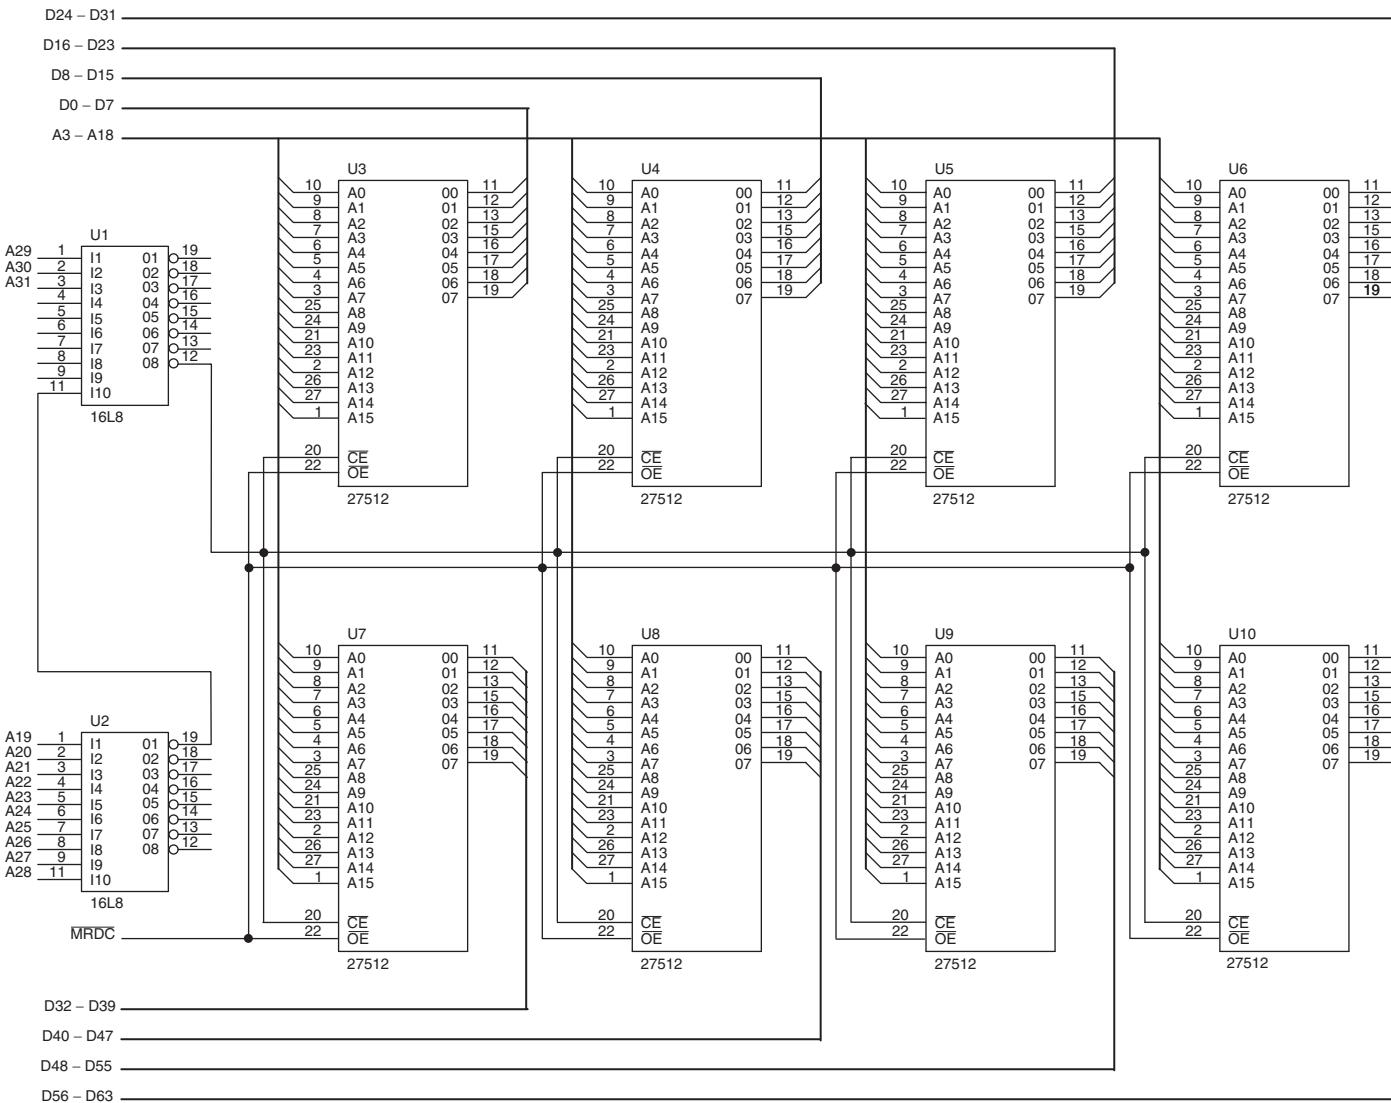
FIGURE 10–37   A small 512K-byte EPROM memory interfaced to the Pentium–Core2 microprocessors.
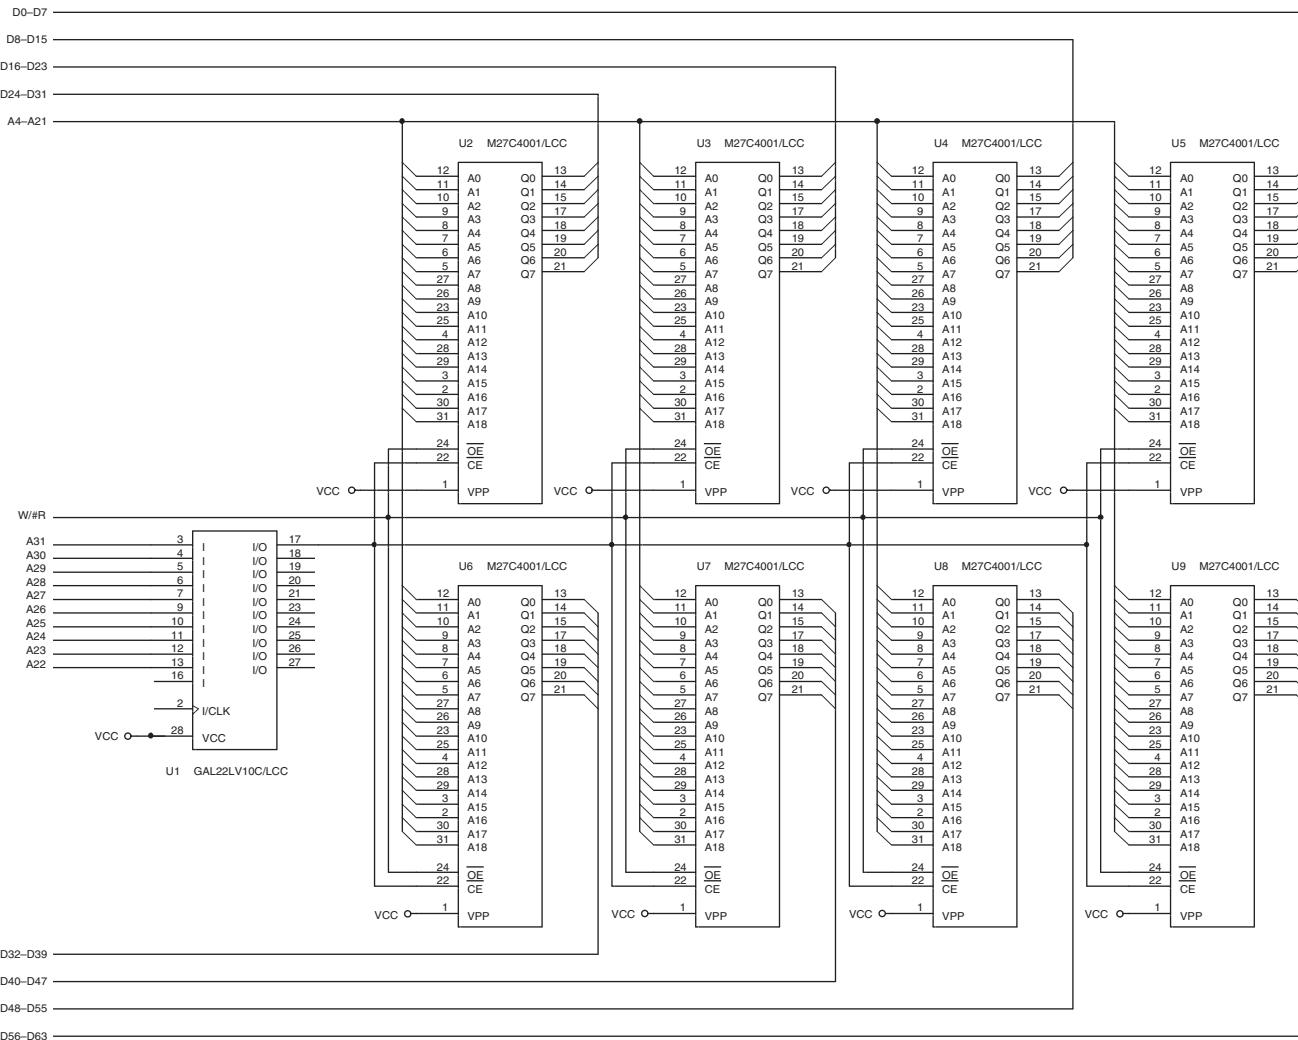
FIGURE 10–38 A small 4M-byte EPROM memory system for the Pentium–Core2 microprocessors.

EXAMPLE 10–10

```
library ieee;
use ieee.std_logic_1164.all;

entity DECODER_10_38 is

port (
        A31, A30, A29, A28, A27, A26, A25, A24, A23, A22: in STD_LOGIC;
        SEL: out STD_LOGIC
);

end;

architecture V1 of DECODER_10_38 is

begin
        SEL <= not(A31 and A30 and A29 and A28 and A27 and A26 and A25 and A24
                and A23 and A22);

end V1;
```

Not explained in this text is the memory interface for an Itanium and Itanium II from Intel, which contains a data bus width of 128 bits. From the information presented in this section of the chapter, it is a fairly easy task to create a memory with 16 banks for the Itanium.

10–7   DYNAMIC RAM

Because RAM memory is often very large, it requires many SRAM devices at a great cost or just a few DRAMs (**dynamic RAMs**) at a much reduced cost. The DRAM memory, as briefly discussed in Section 10–1, is fairly complex because it requires address multiplexing and refreshing. Luckily, the integrated circuit manufacturers have provided a dynamic RAM controller that includes the address multiplexers and all the timing circuitry necessary for refreshing.

This section of the text covers the DRAM memory device in much more detail than in Section 10–1 and provides information on the use of a dynamic controller in a memory system.

DRAM Revisited

As mentioned in Section 10–1, a DRAM retains data for only 2–4 ms and requires the multiplexing of address inputs. Although address multiplexers have already been covered in Section 10–1, the operation of the DRAM during refresh is explained in detail here.

As previously mentioned, a DRAM must be refreshed periodically because it stores data internally on capacitors that lose their charge in a short period of time. To refresh a DRAM, the contents of a section of the memory must periodically be read or written. Any read or write automatically refreshes an entire section of the DRAM. The number of bits that are refreshed depends on the size of the memory component and its internal organization.

Refresh cycles are accomplished by doing a read, a write, or a special refresh cycle that doesn't read or write data. The refresh cycle is internal to the DRAM and is often accomplished while other memory components in the system operate. This type of memory refresh is called either *hidden refresh, transparent refresh,* or sometimes *cycle stealing.*

In order to accomplish a hidden refresh while other memory components are functioning, an $\overline{RAS}$-only cycle strobes a row address into the DRAM to select a row of bits to be refreshed. The $\overline{RAS}$ input also causes the selected row to be read out internally and rewritten into the selected bits. This recharges the internal capacitors that store the data. This type of refresh is

hidden from the system because it occurs while the microprocessor is reading or writing to other sections of the memory.

The DRAM's internal organization contains a series of rows and columns. A 256K × 1 DRAM has 256 columns, each containing 256 bits, or rows organized into four sections of 64K bits each. Whenever a memory location is addressed, the column address selects a column (or internal memory word) of 1024 bits (one per section of the DRAM). Refer to Figure 10–39 for the internal structure of a 256K × 1 DRAM. Note that larger memory devices are structured similarly to the 256K × 1 device. The difference usually lies in either the size of each section or the number of sections in parallel.

Figure 10–40 illustrates the timing for an $\overline{RAS}$-only refresh cycle. The difference between the $\overline{RAS}$ and a read or write is that it applies only a refresh address, which is usually obtained from a 7- or 8-bit binary counter. The size of the counter is determined by the type of DRAM being refreshed. The refresh counter is incremented at the end of each refresh cycle so all the rows are refreshed in 2 or 4 ms, depending on the type of DRAM.

If there are 256 rows to be refreshed within 4 ms, as in a 256K × 1 DRAM, then the refresh cycle must be activated at least once every 15.6 μs in order to meet the refresh specification. For example, it takes the 8086/8088, running at a 5 MHz clock rate, 800 ns to do a read or a write. Because the DRAM must have a refresh cycle every 15.6 μs, for every 19 memory reads or writes, the memory system must run a refresh cycle or else memory data will be lost. This represents a loss of 5% of the computer's time, a small price to pay for the savings represented by using the dynamic RAM. In a modern system such as a 3.0 GHz Pentium 4, 15.6 μs is a great deal of time. Since the 3.0 GHz Pentium 4 executes an instruction in about one-third ns (many instructions execute in a single clock), it can execute about 46,000 instructions between refreshes. This means that in the new machines much less than 1% (≈ 0.002 %) is required for a refresh.
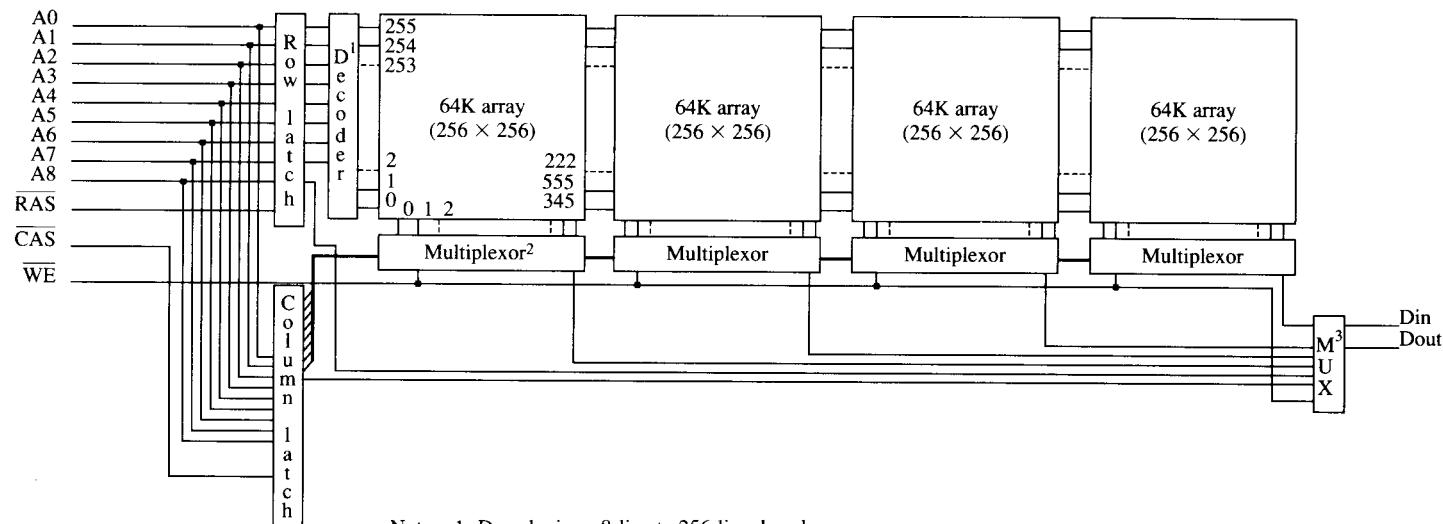
EDO Memory

A slight modification to the structure of the DRAM changes the device into an EDO (**extended data output**) DRAM device. In the EDO memory, any memory access, including a refresh, stores the 256 bits selected by $\overline{RAS}$ into latches. These latches hold the next 256 bits of information, so in most programs, which are sequentially executed, the data are available without any wait states. This slight modification to the internal structure of the DRAM increases system performance by about 15% to 25%. Although EDO memory is no longer available, this technique is still employed in all modern DRAM.

SDRAM

Synchronous dynamic RAM (**SDRAM**) is used with most newer systems in one form or another because of its speed. Versions are available with access times of 10 ns for use with a 66 MHz system bus; 8 ns for use with a 100 MHz system bus; and 7 ns for the 133 MHz bus. At first, the access time may lead one to think that these devices operate without wait states, but that is not true. After all, DRAM access time is 60 ns and SDRAM access time is 10 ns. The 10 ns access time is misleading because it only applies to the second, third, and fourth 64-bit reads from the device. The first read requires the same number of waits as a standard DRAM.

When a burst transfer occurs to the SDRAM from the microprocessor, it takes three or four bus clocks before the first 64-bit number is read. Each subsequent number is read without wait states and in one bus cycle each. Because SDRAM bursts read four 64-bit numbers, and the second through the fourth require no waits and can be read in one bus cycle each, SDRAM outperforms standard DRAM or even EDO memory. This means that if it takes three bus cycles for the first number and three more for the next three, it takes a total of seven bus clocks to read four 64-bit numbers. If this is compared to DRAM, which takes three clocks per number or 12 clocks,

A0
A1
A2
A3
A4
A5
A6
A7
A8
$\overline{RAS}$
$\overline{CAS}$
$\overline{WE}$

Row latch

D e c o d e r

255
254
253

64K array
(256 × 256)

2
1
0
0 1 2

222
555
345

64K array
(256 × 256)

64K array
(256 × 256)

64K array
(256 × 256)

Column latch

Multiplexor[2]

Multiplexor

Multiplexor

Multiplexor

M[3]
U
X

Din
Dout

Notes: 1. Decoder is an 8-line to 256-line decoder.
2. Multiplexor is 256 to 1 line.
3. Multiplexor is 4 to 1 line.

FIGURE 10–39 The internal structure of a 256K × 1 DRAM. Note that each of the internal 256 words are 1024 bits wide.
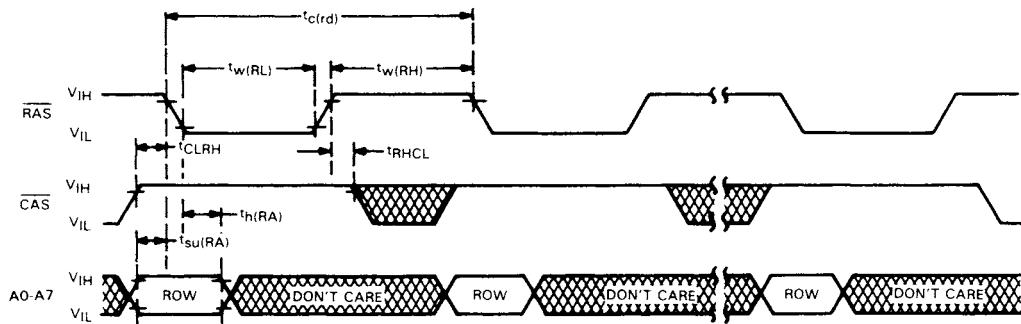
FIGURE 10–40   The timing diagram of the $\overline{RAS}$ refresh cycle for the TMS4464 DRAM. (Courtesy of Texas Instruments Corporation.)

you can see the increase in speed. Most estimates place SDRAM at about a 10% performance increase over EDO memory.

DDR

Double-data rate (DDR) memory is the latest improvement in a long string of modifications to DRAM. The DDR memory transfers data at double the rate of an SDRAM because it transfers data on each edge of the clock. The positive edge is used for a transfer and so is the negative edge. Even though this seems as if it doubles the speed of the memory, it really does not. The main reason is that the access time problem still exists, with even the most advanced memory requiring an access time of 40 ns. If you think about a microprocessor running at many GHz, this is a very long time to wait for the memory. Hence, the speed will not always be double as the name might suggest.

DRAM Controllers

In most systems, a DRAM controller-integrated circuit performs the task of address multiplexing and the generation of the DRAM control signals. Some newer embedded microprocessors, such as the 80186/80188, include the refresh circuitry as a part of the microprocessor. Most modern computers contain the DRAM controller in the chip set so a stand-alone DRAM controller is not available. The DRAM controller in the chip set for the microprocessor times refresh cycles and inserts refresh cycles into the timing. The memory refresh is transparent to the microprocessor, because it really does not control refreshing.

     For the Pentium II, III, and Pentium 4, the DRAM controller is built into the chip set provided by Intel or AMD. It has been many years since a separate DRAM controller has been used to build a computer system. In the future, even the chip set will undoubtedly be built into the microprocessor.

10–8   SUMMARY

1. All memory devices have address inputs; data inputs and outputs, or just outputs; a pin for selection; and one or more pins that control the operation of the memory.
2. Address connections on a memory component are used to select one of the memory locations within the device. Ten address pins have 1024 combinations and therefore are able to address 1024 different memory locations.
3. Data connections on a memory are used to enter information to be stored in a memory location and also to retrieve information read from a memory location. Manufacturers list their

memory as, for example, $4K \times 4$, which means that the device has 4K memory locations (4096) and that four bits are stored in each location.

4. Memory selection is accomplished via a chip selection pin ($\overline{CS}$) on many RAMs or a chip enable pin ($\overline{CE}$) on many EPROM or ROM memories.

5. Memory function is selected by an output enable pin ($\overline{CE}$) for reading data, which normally connects to the system read signal ($\overline{RD}$ or $\overline{MRDC}$). The write enable pin ($\overline{WE}$), for writing data, normally connects to the system write signal ($\overline{WR}$ or $\overline{MWTC}$).

6. An EPROM memory is programmed by an EPROM programmer and can be erased if exposed to ultraviolet light. Today, EPROMs are available in sizes from $1K \times 8$ all the way up to $512K \times 8$ and larger.

7. The flash memory (EEPROM) is programmed in the system by using a 12 V or 5.0 V programming pulse.

8. Static RAM (SRAM) retains data for as long as the system power supply is attached. These memory types are available in sizes up to $128K \times 8$.

9. Dynamic RAM (DRAM) retains data for only a short period, usually 2–4 ms. This creates problems for the memory system designer because the DRAM must be refreshed periodically. DRAMs also have multiplexed address inputs that require an external multiplexer to provide each half of the address at the appropriate time.

10. Memory address decoders select an EPROM or RAM at a particular area of the memory. Commonly found address decoders include the 74LS138 3-to-8 line decoder, the 74LS139 2-to-4 line decoder, and programmed selection logic in the form of a PLD.

11. The PLD address decoder for microprocessors like the 8088 through the Pentium 4 reduce the number of integrated circuits required to complete a functioning memory system.

12. The 8088 minimum mode memory interface contains 20 address lines, eight data lines, and three control lines: $\overline{RD}$, $\overline{WR}$, and IO/$\overline{M}$. The 8088 memory functions correctly only when all these lines are used for memory interface.

13. The access speed of the EPROM must be compatible with the microprocessor to which it is interfaced. Many EPROMs available today have an access time of 450 ns, which is too slow for the 5 MHz 8088. In order to circumvent this problem, a wait state is inserted to increase memory access time to 660 ns.

14. Error-correction features are also available for memory systems, but these require the storage of many more bits. If an 8-bit number is stored with an error-correction circuit, it actually takes 13 bits of memory: five for an error checking code and eight for the data. Most error-correction integrated circuits are able to correct only a single-bit error.

15. The 8086/80286/80386SX memory interface has a 16-bit data bus and contains an M/$\overline{IO}$ control pin, whereas the 8088 has an 8-bit data bus and contains an IO/$\overline{M}$ pin. In addition to these changes, there is an extra control signal, bus high enable ($\overline{BHE}$).

16. The 8086/80386/80386SX memory is organized in two 8-bit banks: high bank and low bank. The high bank of memory is enabled by the $\overline{BHE}$ control signal and the low bank is enabled by the A_0 address signal or by the $\overline{BLE}$ control signal.

17. Two common schemes for selecting the banks in an 8086/80286/80386SX-based system include (1) a separate decoder for each bank and (2) separate $\overline{WR}$ control signals for each bank with a common decoder.

18. Memory interfaced to the 80386DX and 80486 is 32 bits wide, as selected by a 32-bit address bus. Because of the width of this memory, it is organized in four memory banks that are each 8 bits wide. Bank selection signals are provided by the microprocessor as $\overline{BE3}$, $\overline{BE2}$, $\overline{BE1}$, and $\overline{BE0}$.

19. Memory interfaced to the Pentium–Core2 is 64 bits wide, as selected by a 32-bit address bus. Because of the width of the memory, it is organized in eight banks that are each 8 bits wide. Bank selection signals are provided by the microprocessor as $\overline{BE7}$–$\overline{BE0}$.

20. Dynamic RAM controllers are designed to control DRAM memory components. Many DRAM controllers today are built into the chip set and contain address multiplexers, refresh counters, and the circuitry required to do a periodic DRAM memory refresh.

10–9   QUESTIONS AND PROBLEMS

1. What types of connections are common to all memory devices?
2. List the number of words found in each memory device for the following numbers of address connections:
    (a) 8
    (b) 11
    (c) 12
    (d) 13
    (e) 20
3. List the number of data items stored in each of the following memory devices and the number of bits in each datum:
    (a) $2K \times 4$
    (b) $1K \times 1$
    (c) $4K \times 8$
    (d) $16K \times 1$
    (e) $64K \times 4$
4. What is the purpose of the $\overline{CS}$ or $\overline{CE}$ pin on a memory component?
5. What is the purpose of the $\overline{OE}$ pin on a memory device?
6. What is the purpose of the $\overline{WE}$ pin on a SRAM?
7. How many bytes of storage do the following EPROM memory devices contain?
    (a) 2708
    (b) 2716
    (c) 2732
    (d) 2764
    (e) 27512
8. Why won't a 450 ns EPROM work directly with a 5 MHz 8088?
9. What can be stated about the amount of time needed to erase and write a location in a flash memory device?
10. SRAM is an acronym for what type of device?
11. The 4016 memory has a $\overline{G}$ pin, an $\overline{S}$ pin, and a $\overline{W}$ pin. What are these pins used for in this RAM?
12. How much memory access time is required by the slowest 4016?
13. DRAM is an acronym for what type of device?
14. The 256M DIMM has 28 address inputs, yet it is a 256M DRAM. Explain how a 28-bit memory address is forced into 14 address inputs.
15. What are the purposes of the $\overline{CAS}$ and $\overline{RAS}$ inputs of a DRAM?
16. How much time is required to refresh the typical DRAM?
17. Why are memory address decoders important?
18. Modify the NAND gate decoder of Figure 10–13 to select the memory for address range DF800H–DFFFFH.
19. Modify the NAND gate decoder in Figure 10–13 to select the memory for address range 40000H–407FFH.
20. When the G1 input is high and both $\overline{G2A}$ and $\overline{G2B}$ are low, what happens to the outputs of the 74HCT138 3-to-8 line decoder?
21. Modify the circuit of Figure 10–15 to address memory range 70000H–7FFFFH.

22. Modify the circuit of Figure 10–15 to address memory range 40000H–4FFFFH.
23. Describe the 74LS139 decoder.
24. What is VHDL?
25. What are the five major keywords in VHDL for the five major logic functions (AND, OR, NAND, NOR, and invert)?
26. Equations are placed in what major block of a VHDL program?
27. Modify the circuit of Figure 10–19 by rewriting the PLD program to address memory at locations A0000H–BFFFFH for the ROM.
28. The $\overline{RD}$ and $\overline{WR}$ minimum mode control signals are replaced by what two control signals in the 8086 maximum mode?
29. Modify the circuit of Figure 10–20 to select memory at location 60000H–77FFFH.
30. Modify the circuit of Figure 10–20 to select eight 27256 32K × 8 EPROMs at memory locations 40000H–7FFFFH.
31. Add another decoder to the circuit of Figure 10–21 so that an additional eight 62256 SRAMs are added at locations C0000H–FFFFFH.
32. The 74LS636 error-correction and detection circuit stores a check code with each byte of data. How many bits are stored for the check code?
33. What is the purpose of the SEF pin on the 74LS636?
34. The 74LS636 will correct _____ bits that are in error.
35. Outline the major difference between the buses of the 8086 and 8088 microprocessors.
36. What is the purpose of the $\overline{BHE}$ and A_0 pins on the 8086 microprocessor?
37. What is the $\overline{BLE}$ pin and what other pin has it replaced?
38. What two methods are used to select the memory in the 8086 microprocessor?
39. If $\overline{BHE}$ is a logic 0, then the _____ memory bank is selected.
40. If A_0 is a logic 0, then the _____ memory bank is selected.
41. Why don't separate bank read ($\overline{RD}$) strobes need to be developed when interfacing memory to the 8086?
42. Modify the circuit of Figure 10–30 so that the RAM is located at memory range 30000H–4FFFFH.
43. Develop a 16-bit-wide memory interface that contains SRAM memory at locations 200000H–21FFFFH for the 80386SX microprocessor.
44. Develop a 32-bit-wide memory interface that contains EPROM memory at locations FFFF0000H–FFFFFFFFH.
45. Develop a 64-bit-wide memory for the Pentium–Core2 that contains EPROM at locations FFF00000H–FFFFFFFFH and SRAM at locations 00000000H–003FFFFFH.
46. On the Internet, search for the largest size EEPROM you can find. List its size and manufacturer.
47. What is an $\overline{RAS}$-only cycle?
48. Can a DRAM refresh be done while other sections of the memory operate?
49. If a 1M × 1 DRAM requires 4 ms for a refresh and has 256 rows to be refreshed, no more than _____ of time must pass before another row is refreshed.
50. How wide is the data bus in the Intel Itanium?
51. Scour the Internet to find the largest DRAM currently available.
52. Write a report on DDR memory. (Hint: Samsung makes them.)
53. Write a report that details RAMBUS RAM. Try to determine why this technology appears to have fallen by the wayside.

CHAPTER 11

Basic I/O Interface

INTRODUCTION

A microprocessor is great at solving problems, but if it can't communicate with the outside world, it is of little worth. This chapter outlines some of the basic methods of communications, both serial and parallel, between humans or machines and the microprocessor.

In this chapter, we first introduce the basic I/O interface and discuss decoding for I/O devices. Then, we provide detail on parallel and serial interfacing, both of which have a variety of applications. To study applications, we connect analog-to-digital and digital-to-analog converters, as well as both DC and stepper motors to the microprocessor.

CHAPTER OBJECTIVES

Upon completion of this chapter, you will be able to:

1. Explain the operation of the basic input and output interfaces.
2. Decode an 8-, 16-, and 32-bit I/O device so that they can be used at any I/O port address.
3. Define handshaking and explain how to use it with I/O devices.
4. Interface and program the 82C55 programmable parallel interface.
5. Interface LCD displays, LED displays, keyboards, ADC, DAC, and various other devices to the 82C55.
6. Interface and program the 16550 serial communications interface adapter.
7. Interface and program the 8254 programmable interval timer.
8. Interface an analog-to-digital converter and a digital-to-analog converter to the microprocessor.
9. Interface both DC and stepper motors to the microprocessor.

11–1 INTRODUCTION TO I/O INTERFACE

In this section of the text I/O instructions (IN, INS, OUT, and OUTS) are explained and used in example applications. Also explained here is the concept of isolated (sometimes called direct or I/O mapped I/O) and memory-mapped I/O, the basic input and output interfaces, and hand-shaking. A working knowledge of these topics makes it easier to understand the connection and

operation of the programmable interface components and I/O techniques presented in the remainder of this chapter and text.

The I/O Instructions

The instruction set contains one type of instruction that transfers information to an I/O device (OUT) and another to read information from an I/O device (IN). Instructions (INS and OUTS, found on all versions except the 8086/8088) are also provided to transfer strings of data between the memory and an I/O device. Table 11–1 lists all versions of each instruction found in the microprocessor's instruction set.

Instructions that transfer data between an I/O device and the microprocessor's accumulator (AL, AX, or EAX) are called **IN** and **OUT**. The I/O address is stored in register DX as a 16-bit I/O address or in the byte (p8) immediately following the opcode as an 8-bit I/O address. Intel calls the 8-bit form (p8) a **fixed address** because it is stored with the instruction, usually in a ROM. The 16-bit I/O address in DX is called a **variable address** because it is stored in a DX, and then used to address the I/O device. Other instructions that use DX to address I/O are the INS and OUTS instructions. I/O ports are 8 bits in width so whenever a 16-bit port is accessed two consecutive 8-bit ports are actually addressed. A 32-bit I/O port is actually four 8-bit ports. For example, port 100H is accessed as a word, then 100H and 101H are actually accessed. Port 100H contains the least significant part of the data and port 101H the most significant part.

TABLE 11–1   Input/Output instructions.

Instruction	Data Width	Function
IN AL, p8	8	A byte is input into AL from port p8
IN AX, p8	16	A word is input into AX from port p8
IN EAX, p8	32	A doubleword is input into EAX from port p8
IN AL, DX	8	A byte is input into AL from the port addressed by DX
IN AX, DX	16	A word is input into AX from the port addressed by DX
IN EAX, DX	32	A doubleword is input into EAX from the port addressed by DX
INSB	8	A byte is input from the port addressed by DI and stored into the extra segment memory location addressed by DI, then DI = DI ± 1
INSW	16	A word is input from the port addressed by DI and stored into the extra segment memory location addressed by DI, then DI = DI ± 2
INSD	32	A doubleword is input from the port addressed by DI and stored into the extra segment memory location addressed by DI, then DI = DI ± 4
OUT p8, AL	8	A byte is output from AL into port p8
OUT p8, AX	16	A word is output from AL into port p8
OUT p8, EAX	32	A doubleword is output from EAX into port p8
OUT DX, AL	8	A byte is output from AL into the port addressed by DX
OUT DX, AX	16	A word is output from AX into the port addressed by DX
OUT DX, EAX	32	A doubleword is output from EAX into the port addressed by DX
OUTSB	8	A byte is output from the data segment memory location addressed by SI into the port addressed by DX, then SI = SI ± 1
OUTSW	16	A word is output from the data segment memory location addressed by SI into the port addressed by DX, then SI = SI ± 2
OUTSD	32	A doubleword is output from the data segment memory location addressed by SI into the port addressed by DX, then SI = SI ± 4

Whenever data are transferred by using the IN or OUT instructions, the I/O address, often called a **port number** (or simply port), appears on the address bus. The external I/O interface decodes the port number in the same manner that it decodes a memory address. The 8-bit fixed port number (p8) appears on address bus connections A_7–A_0 with bits A_{15}–A_8 equal to 00000000_2. The address connections above A_{15} are undefined for an I/O instruction. The 16-bit variable port number (DX) appears on address connections A_{15}–A_0. This means that the first 256 I/O port addresses (00H–FFH) are accessed by both the fixed and variable I/O instructions, but any I/O address from 0100H to FFFFH is only accessed by the variable I/O address. In many dedicated systems, only the rightmost 8 bits of the address are decoded, thus reducing the amount of circuitry required for decoding. In a PC computer, all 16 address bus bits are decoded with locations 0000H–03FFH, which are the I/O addresses used for I/O inside the PC on the ISA (**industry standard architecture**) bus.

The INS and OUTS instructions address an I/O device by using the DX register, but do not transfer data between the accumulator and the I/O device as do the IN and OUT instructions. Instead, these instructions transfer data between memory and the I/O device. The memory address is located by ES:DI for the INS instruction and by DS:SI for the OUTS instruction. As with other string instructions, the contents of the pointers are incremented or decremented, as dictated by the state of the direction flag (DF). Both INS and OUTS can be prefixed with the REP prefix, allowing more than one byte, word, or doubleword to be transferred between I/O and memory.

The Pentium 4 and Core2 operating in the 64-bit mode have the same I/O instructions. There are no 64-bit I/O instructions in the 64-bit mode. The main reason is that most I/O is still 8 bits and likely will remain so for an indefinite time.
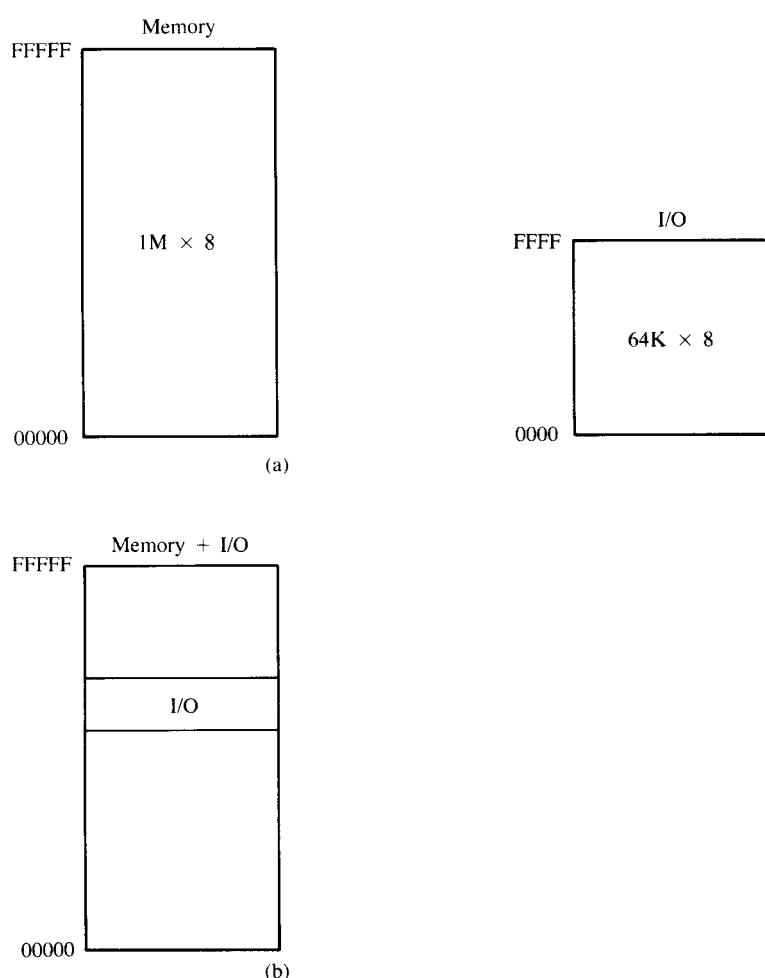
Isolated and Memory-Mapped I/O

There are two different methods of interfacing I/O to the microprocessor: **isolated I/O** and **memory-mapped I/O**. In the isolated I/O scheme, the IN, INS, OUT, and OUTS instructions transfer data between the microprocessor's accumulator or memory and the I/O device. In the memory-mapped I/O scheme, any instruction that references memory can accomplish the transfer. Both isolated and memory-mapped I/O are in use, so both are discussed in this text. The PC does not use memory-mapped I/O.

Isolated I/O.   The most common I/O transfer technique used in the Intel microprocessor-based system is isolated I/O. The term *isolated* describes how the I/O locations are isolated from the memory system in a separate I/O address space. (Figure 11–1 illustrates both the isolated and memory-mapped address spaces for any Intel 80X86 or Pentium–Core2 microprocessor.) The addresses for isolated I/O devices, called ports, are separate from the memory. Because the ports are separate, the user can expand the memory to its full size without using any of memory space for I/O devices. A disadvantage of isolated I/O is that the data transferred between I/O and the microprocessor must be accessed by the IN, INS, OUT, and OUTS instructions. Separate control signals for the I/O space are developed (using M/$\overline{\text{IO}}$ and W/$\overline{\text{R}}$), which indicate an I/O read ($\overline{\text{IORC}}$) or an I/O write ($\overline{\text{IOWC}}$) operation. These signals indicate that an I/O port address, which appears on the address bus, is used to select the I/O device. In the personal computer, isolated I/O ports are used for controlling peripheral devices. An 8-bit port address is used to access devices located on the system board, such as the timer and keyboard interface, while a 16-bit port is used to access serial and parallel ports as well as video and disk drive systems.

Memory-Mapped I/O.   Unlike isolated I/O, memory-mapped I/O does not use the IN, INS, OUT, or OUTS instructions. Instead, it uses any instruction that transfers data between the microprocessor and memory. A memory-mapped I/O device is treated as a memory location in the memory map. The main advantage of memory-mapped I/O is that any memory transfer instruction can be used to access the I/O device. The main disadvantage is that a portion of the memory system is used as the I/O map. This reduces the amount of memory available to applications. Another advantage is that the $\overline{\text{IORC}}$ and $\overline{\text{IOWC}}$ signals have no function in a memory-mapped I/O system and may reduce the amount of circuitry required for decoding.

FIGURE 11–1 The memory
and I/O maps for the 8086/
8088 microprocessors.
(a) Isolated I/O. (b) Memory-
mapped I/O.

Personal Computer I/O Map
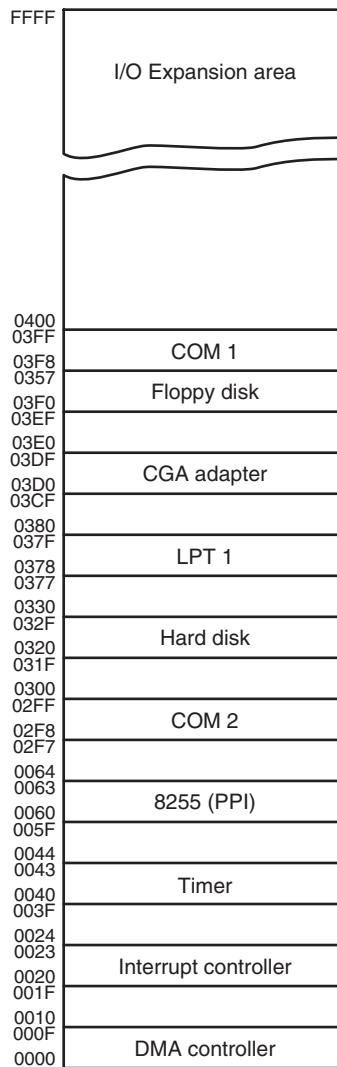
The personal computer uses part of the I/O map for dedicated functions. Figure 11–2 shows the I/O map for the PC. Note that I/O space between ports 0000H and 03FFH is normally reserved for the computer system and the ISA bus. The I/O ports located at 0400H–FFFFH are generally available for user applications, main-board functions, and the PCI bus. Note that the 80287 arithmetic coprocessor uses I/O address 00F8H–00FFH for communications. For this reason, Intel reserves I/O ports 00F0H–00FFH. The 80386–Core2 use I/O ports 800000F8–800000FFH for communications to their coprocessors. The I/O ports located between 0000H and 00FFH are accessed via the fixed port I/O instructions; the ports located above 00FFH are accessed via the variable I/O port instructions.

Basic Input and Output Interfaces

The basic input device is a set of three-state buffers. The basic output device is a set of data latches. The term IN refers to moving data from the I/O device into the microprocessor and the term OUT refers to moving data out of the microprocessor to the I/O device.

The Basic Input Interface.   Three-state buffers are used to construct the 8-bit input port depicted in Figure 11–3. The external TTL data (simple toggle switches in this example) are connected to the

FIGURE 11–2 The I/O map of a personal computer illustrating many of the fixed I/O areas.

FFFF	I/O Expansion area
0400 / 03FF	
03F8	COM 1
0357 / 03F0	Floppy disk
03EF / 03E0	
03DF / 03D0	CGA adapter
03CF / 0380	
037F / 0378	LPT 1
0377 / 0330	
032F / 0320	Hard disk
031F / 0300	
02FF / 02F8	COM 2
02F7 / 0064	
0063 / 0060	8255 (PPI)
005F / 0044	
0043 / 0040	Timer
003F / 0024	
0023 / 0020	Interrupt controller
001F / 0010	
000F / 0000	DMA controller

inputs of the buffers. The outputs of the buffers connect to the data bus. The exact data bus connections depend on the version of the microprocessor. For example, the 8088 has data bus connections D_7–D_0, the 80386/80486 has connections D_{31}–D_0, and the Pentium–Core2 have connections D_{63}–D_0. The circuit of Figure 11–3 allows the microprocessor to read the contents of the eight switches that connect to any 8-bit section of the data bus when the select signal $\overline{SEL}$ becomes a logic 0. Thus, whenever the IN instruction executes, the contents of the switches are copied into the AL register.

When the microprocessor executes an IN instruction, the I/O port address is decoded to generate the logic 0 on $\overline{SEL}$. A 0 placed on the output control inputs ($\overline{1G}$ and $\overline{2G}$) of the 74ALS244 buffer causes the data input connections (A) to be connected to the data output (Y) connections. If a logic 1 is placed on the output control inputs of the 74ALS244 buffer, the device enters the three-state high-impedance mode that effectively disconnects the switches from the data bus.

This basic input circuit is not optional and must appear any time that input data are interfaced to the microprocessor. Sometimes it appears as a discrete part of the circuit, as shown in Figure 11–3; many times it is built into a programmable I/O device.

Sixteen- or 32-bit data can also be interfaced to various versions of the microprocessor, but this is not nearly as common as using 8-bit data. To interface 16 bits of data, the circuit in

FIGURE 11–3 The basic input interface illustrating the connection of eight switches. Note that the 74ALS244 is a three-state buffer that controls the application of the switch data to the data bus.
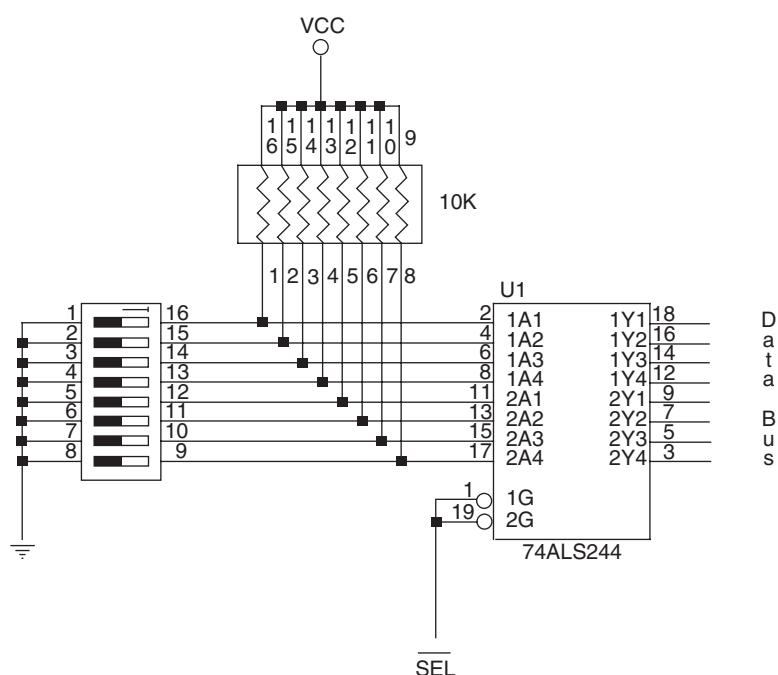
Figure 11–3 is doubled to include two 74ALS244 buffers that connect 16 bits of input data to the 16-bit data bus. To interface 32 bits of data, the circuit is expanded by a factor of 4.

The Basic Output Interface.   The basic output interface receives data from the microprocessor and usually must hold it for some external device. Its latches or flip-flops, like the buffers found in the input device, are often built into the I/O device.

Figure 11–4 shows how eight simple light-emitting diodes (LEDs) connect to the microprocessor through a set of eight data latches. The latch stores the number output by the microprocessor from the data bus so that the LEDs can be lit with any 8-bit binary number. Latches are needed to hold the data because when the microprocessor executes an OUT instruction, the data are only present on the data bus for less than 1.0 μs. Without a latch, the viewer would never see the LEDs illuminate.
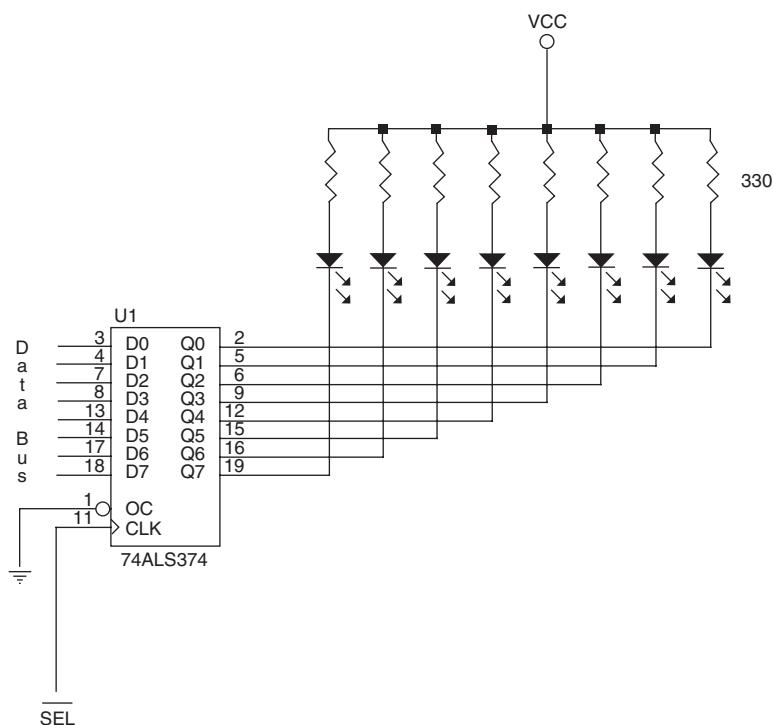
When the OUT instruction executes, the data from AL, AX, or EAX are transferred to the latch via the data bus. Here, the D inputs of a 74ALS374 octal latch are connected to the data bus to capture the output data, and the Q outputs of the latch are attached to the LEDs. When a Q output becomes a logic 0, the LED lights. Each time that the OUT instruction executes, the $\overline{SEL}$ signal to the latch activates, capturing the data output to the latch from any 8-bit section of the data bus. The data are held until the next OUT instruction executes. Thus, whenever the output instruction is executed in this circuit, the data from the AL register appear on the LEDs.

Handshaking

Many I/O devices accept or release information at a much slower rate than the microprocessor. Another method of I/O control, called **handshaking** or **polling**, synchronizes the I/O device with the microprocessor. An example of a device that requires handshaking is a parallel printer that prints a few hundred characters per second (CPS). It is obvious that the microprocessor can send more than a few hundred CPS to the printer, so a way to slow the microprocessor down to match speeds with the printer must be developed.

Figure 11–5 illustrates the typical input and output connections found on a printer. Here, data are transferred through a series of data connections (D_7–D_0). BUSY indicates that the printer is busy. $\overline{STB}$ is a clock pulse used to send data to the printer for printing.

FIGURE 11–4 The basic
output interface connected to
a set of LED displays.

The ASCII data to be printed by the printer are placed on D_7–D_0, and a pulse is then applied to the $\overline{\text{STB}}$ connection. The strobe signal sends or clocks the data into the printer so that they can be printed. As soon as the printer receives the data, it places a logic 1 on the BUSY pin, indicating that the printer is busy printing data. The microprocessor software polls or tests the BUSY pin to decide whether the printer is busy. If the printer is busy, the microprocessor waits; if it is not busy, the microprocessor sends the next ASCII character to the printer. This process of interrogating the printer, or any asynchronous device like a printer, is called handshaking or polling. Example 11–1 illustrates a simple procedure that tests the printer BUSY flag and then sends data to the printer if it is not busy. Here, the PRINT procedure prints the ASCII-coded contents of BL only if the BUSY flag is a logic 0, indicating that the printer is not busy. This procedure is called each time a character is to be printed.

EXAMPLE 11–1

```
;An assembly language procedure that prints the ASCII contents of BL.

PRINT   PROC    NEAR

        .REPEAT                     ;test the busy flag
                IN AL,BUSY
                TEST AL,BUSY_BIT
        .UNTIL ZERO
        MOV AL,BL                   ;position data in AL
        OUT PRINTER,AL              ;print data
        RET

PRINT   ENDP
```
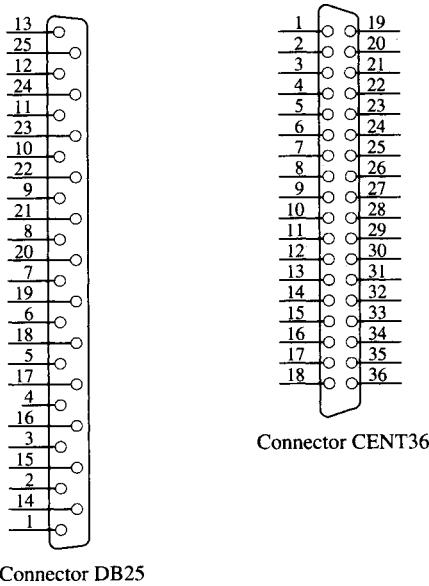
Notes about Interfacing Circuitry

A part of interfacing requires some knowledge about electronics. This portion of the introduction to interfacing examines some of the many facets of electronic interfacing. Before a circuit or

Connector DB25

Connector CENT36

DB25 Pin number	CENT36 Pin number	Function	DB25 Pin number	CENT36 Pin number	Function
1	1	$\overline{\text{Data Strobe}}$	12	12	Paper empty
2	2	Data 0 (D0)	13	13	Select
3	3	Data 1 (D1)	14	14	Afd
4	4	Data 2 (D2)	15	32	$\overline{\text{Error}}$
5	5	Data 3 (D3)	16	—	$\overline{\text{RESET}}$
6	6	Data 4 (D4)	17	31	Select in
7	7	Data 5 (D5)	18—25	19—30	Ground
8	8	Data 6 (D6)	—	17	Frame ground
9	9	Data 7 (D7)	—	16	Ground
10	10	$\overline{\text{Ack}}$	—	33	Ground
11	11	Busy			

FIGURE 11–5   The DB25 connector found on computers and the Centronics 36-pin connector found on printers for the Centronics parallel printer interface.

device can be interfaced to the microprocessor, the terminal characteristics of the microprocessor and its associated interfacing components must be known. (This subject was introduced at the start of Chapter 9.)

Input Devices.   Input devices are already TTL and compatible, and therefore can be connected to the microprocessor and its interfacing components, or they are switch-based. Most switch-based devices are either open or connected. These are not TTL levels—TTL levels are a logic 0 (0.0 V–0.8 V) or a logic 1 (2.0 V–5.0 V).

For a switch-based device to be used as a TTL-compatible input device, some conditioning must be applied. Figure 11–6 shows a simple toggle switch that is properly connected to function as an input device. Notice that a pull-up resistor is used to ensure that when the switch is open, the output signal is a logic 1; when the switch is closed, it connects to ground, producing a valid logic 0 level. The value of the pull-up resistor is not critical—it merely assures that the signal is

FIGURE 11–6 A single-pole, single-throw switch interfaced as a TTL device.
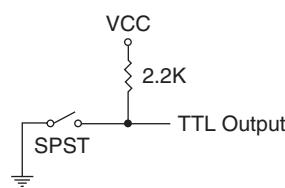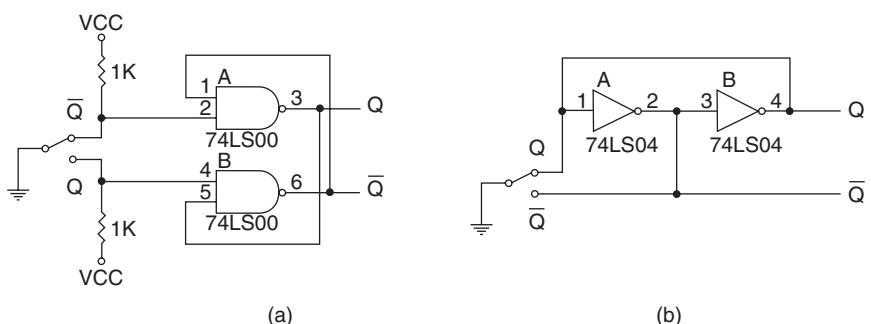
FIGURE 11–7 Debouncing switch contacts: (a) conventional debouncing and (b) practical debouncing.

at a logic 1 level. A standard range of values for pull-up resistors is usually anywhere between 1K Ω and 10K Ω.

Mechanical switch contacts physically bounce when they are closed, which can create a problem if a switch is used as a clocking signal for a digital circuit. To prevent problems with bounces, one of the two circuits depicted in Figure 11–7 can be constructed. The first circuit (a) is a classical textbook bounce eliminator; the second (b) is a more practical version of the same circuit. Because the first version costs more money to construct, in practice, the second would be used because it requires no pull-up resistors and only two inverters instead of two NAND gates.

You may notice that both circuits in Figure 11–7 are asynchronous flip-flops. The circuit of (b) functions in the following manner: Suppose that the switch is currently at position $\overline{Q}$. If it is moved toward Q but does not yet touch Q, the Q output of the circuit is a logic 0. The logic 0 state is remembered by the inverters. The output of inverter B connects to the input of inverter A. Because the output of inverter B is a logic 0, the output of inverter A is a logic 1. The logic 1 output of inverter A maintains the logic 0 output of inverter B. The flip-flop remains in this state until the moving switch-contact first touches the Q connection. As soon as the Q input from the switch becomes a logic 0, it changes the state of the flip-flop. If the contact bounces back away from the Q input, the flip-flop remembers and no change occurs, thus eliminating any bounce.

Output Devices.   Output devices are far more diverse than input devices, but many are interfaced in a uniform manner. Before any output device can be interfaced, we must understand what the voltages and currents are from the microprocessor or a TTL interface component. The voltages are TTL-compatible from the microprocessor of the interfacing element. (Logic 0 = 0.0 V to 0.4 V; logic 1 = 2.4 V to 5.0 V.) The currents for a microprocessor and many microprocessor-interfacing components are less than for standard TTL components. (Logic 0 = 0.0 to 2.0 mA; logic 1 = 0.0 to 400 μA.)

Once the output currents are known, a device can now be interfaced to one of the outputs. Figure 11–8 shows how to interface a simple LED to a microprocessor peripheral pin. Notice that a transistor driver is used in Figure 11–8(a) and a TTL inverter is used in Figure 11–8(b). The TTL inverter (standard version) provides up to 16 mA of current at a logic 0 level, which is more than enough to drive a standard LED. A standard LED requires 10 mA of forward bias current to light. In both circuits, we assume that the voltage drop across the LED is about 2.0 V.

FIGURE 11–8 Interfacing an LED: (a) using a transistor and (b) using an inverter.

(a)   (b)
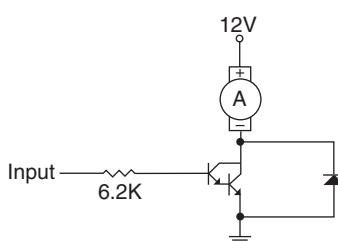
The data sheet for an LED states that the nominal drop is 1.65 V, but it is known from experience that the drop is anywhere between 1.5 V and 2.0 V. This means that the value of the current-limiting resistor is 3.0 V ÷ 10 mA or 300 Ω. Because 300 Ω is not a standard resistor value (the lowest cost), a 330 Ω resistor is chosen for this interface.

In the circuit of Figure 11–8(a), we elected to use a switching transistor in place of the TTL buffer. The 2N2222 is a good low-cost, general-purpose switching transistor that has a minimum gain of 100. In this circuit, the collector current is 10 mA, so the base current will be 1/100 of the collector current of 0.1 mA. To determine the value of the base current–limiting resistor, use the 0.1 mA base current and a voltage drop of 1.7 V across the base current–limiting resistor. The TTL input signal has a minimum value of 2.4 V and the drop across the emitter-base junction is 0.7 V. The difference is 1.7 V, which is the voltage drop across the resistor. The value of the resistor is 1.7 V ÷ 0.1 mA or 17K Ω. Because 17K Ω is not a standard value, an 18K Ω resistor is chosen.

Suppose that we need to interface a 12 V DC motor to the microprocessor and the motor current is 1A. Obviously, we cannot use a TTL inverter for two reasons: The 12 V signal would burn out the inverter and the amount of current far exceeds the 16 mA maximum current from the inverter. We cannot use a 2N2222 transistor either, because the maximum amount of current is 250 mA to 500 mA, depending on the package style chosen. The solution is to use a Darlington-pair, such as a TIP120. The TIP120 costs 25¢ and with the proper heat sink can handle 4A of current.

Figure 11–9 illustrates a motor connected to the Darlington-pair. The Darlington-pair has a minimum current gain of 7000 and a maximum current of 4A. The value of the bias resistor is calculated exactly the same as the one used in the LED driver. The current through the resistor is 1.0 A ÷ 7000, or about 0.143 mA. The voltage drop across the resistor is 0.9 V because of the two diode drops (base/emitter junctions) instead of one. The value of the bias resistor is 0.9 V ÷ 0.143 mA or 6.29K Ω. The standard value of 6.2 K Ω is used in the circuit. The Darlington-pair must use a heat sink because of the amount of current going through it. Typically any device that passes more than ½ A of current needs a heat sink. The diode must also be present to prevent the Darlington-pair from being destroyed by the inductive kickback from the motor. This circuit is also used to interface mechanical relays or just about any device that requires a large amount of current or a change in voltage.

FIGURE 11–9 A DC motor interfaced to a system by using a Darlington-pair.

11–2   **I/O PORT ADDRESS DECODING**

I/O port address decoding is very similar to memory address decoding, especially for memory-mapped I/O devices. In fact, we do not discuss memory-mapped I/O decoding because it is treated the same as memory (except that the $\overline{\text{IORC}}$ and $\overline{\text{IOWC}}$ are not used because there is no IN or OUT instruction). The decision to use memory-mapped I/O is often determined by the size of the memory system and the placement of the I/O devices in the system.

      The main difference between memory decoding and isolated I/O decoding is the number of address pins connected to the decoder. We decode A_{31}–A_0, A_{23}–A_0, or A_{19}–A_0 for memory, and A_{15}–A_0 for isolated I/O. Sometimes, if the I/O devices use only fixed I/O addressing, we decode only A_7–A_0. In the personal computer system, we always decode all 16 bits of the I/O port address. Another difference with isolated I/O is that $\overline{\text{IORC}}$ and $\overline{\text{IOWC}}$ activate I/O devices for a read or write operation. On earlier versions of the microprocessor, IO/$\overline{\text{M}}$ = 1 and $\overline{\text{RD}}$ or $\overline{\text{WR}}$ are used to activate I/O devices. On the newest versions of the microprocessor, the M/$\overline{\text{IO}}$ = 0 and W/$\overline{\text{R}}$ are combined and used to activate I/O devices.

Decoding 8-Bit I/O Port Addresses

As mentioned, the fixed I/O instruction uses an 8-bit I/O port address that appears on A_{15}–A_0 as 0000H–00FFH. If a system will never contain more than 256 I/O devices, we often decode only address connections A_7–A_0 for an 8-bit I/O port address. Thus, we ignore address connection A_{15}–A_8. Embedded systems often use 8-bit port addresses. Please note that the DX register can also address I/O ports 00H–FFH. If the address is decoded as an 8-bit address, we can never include I/O devices that use a 16-bit I/O address. The personal computer never uses or decodes an 8-bit address.

      Figure 11–10 illustrates a 74ALS138 decoder that decodes 8-bit I/O ports F0H through F7H. (We assume that this system will only use I/O ports 00H–FFH for this decoder example.) This decoder is identical to a memory address decoder except we only connect address bits A_7–A_0 to the inputs of the decoder. Figure 11–11 shows the PLD version, using a GAL22V10 (a low-cost device) for this decoder. The PLD is a better decoder circuit because the number of integrated circuits has been reduced to one device. The VHDL program for the PLD appears in Example 11–2.

EXAMPLE 11–2

```
-- VHDL code for the decoder of Figure 11–11

library ieee;
use ieee.std_logic_1164.all;

entity DECODER_11_11 is

port (
      A7, A6, A5, A4, A3, A2, A1, A0: in STD_LOGIC;
      D0, D1, D2, D3, D4, D5, D6, D7: out STD_LOGIC
);

end;

architecture V1 of DECODER_11_11 is

begin

      D0 <= not( A7 and A6 and A5 and A4 and not A3 and not A2 and not A1 and
            not A0 );
      D1 <= not( A7 and A6 and A5 and A4 and not A3 and not A2 and not A1 and
            A0 );
      D2 <= not( A7 and A6 and A5 and A4 and not A3 and not A2 and A1 and
            not A0 );
```

FIGURE 11–10 A port decoder that decodes 8-bit I/O ports. This decoder generates active low outputs for ports F0H–F7H.
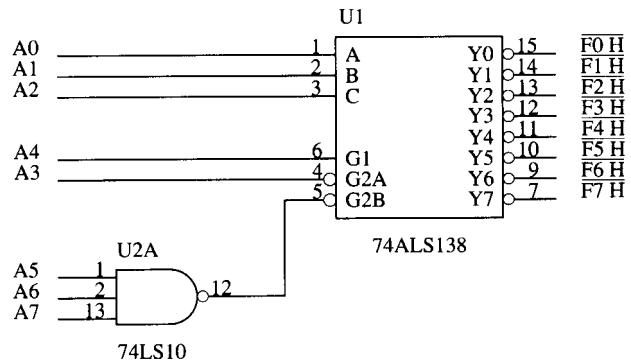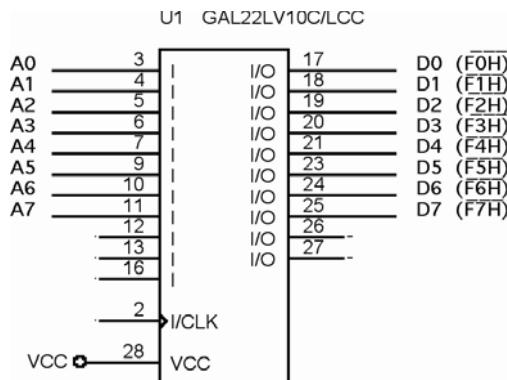
FIGURE 11–11 A PLD that generates part selection signals $\overline{\text{F0H}}$–$\overline{\text{F0H}}$.


```
D3 <= not( A7 and A6 and A5 and A4 and not A3 and not A2 and A1 and
      A0 );
D4 <= not( A7 and A6 and A5 and A4 and not A3 and A2 and not A1 and
      not A0 );
D5 <= not( A7 and A6 and A5 and A4 and not A3 and A2 and not A1 and
      A0 );
D6 <= not( A7 and A6 and A5 and A4 and not A3 and A2 and A1 and
      not A0 );
D0 <= not( A7 and A6 and A5 and A4 and not A3 and A2 and A1 and
      A0 );

end V1;
```

Decoding 16-Bit I/O Port Addresses

Personal computer systems typically use 16-bit I/O addresses. It is relatively rare to find 16-bit port addresses in embedded systems. The main difference between decoding an 8-bit I/O address and a 16-bit I/O address is that eight additional address lines (A_{15}–A_8) must be decoded. Figure 11–12 illustrates a circuit that contains a PLD and a 4-input NAND gate used to decode I/O ports EFF8H–EFFFH.

The NAND gate decodes part of the address (A_{15}, A_{14}, A_{13}, and A_{11}) because the PLD does not have enough address inputs. The output of the NAND gate connects to the Z input of the PLD and is decoded as a part of the I/O port address. The PLD generates address strobes for I/O ports $\overline{\text{EFF8H}}$–$\overline{\text{EFFFH}}$. The program for the PLD is listed in Example 11–3.

EXAMPLE 11–3
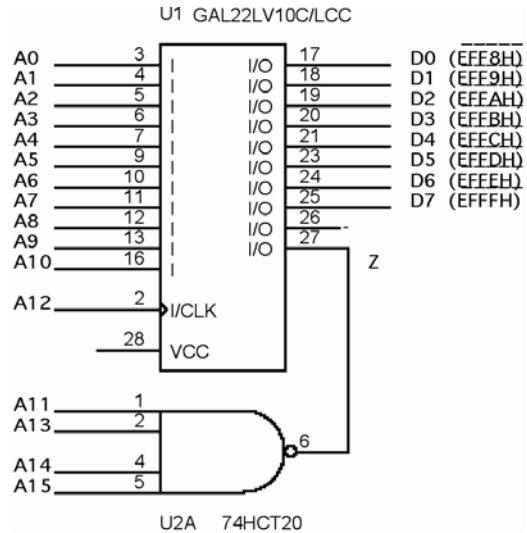
```
-- VHDL code for the decoder of Figure 11-12

library ieee;
use ieee.std_logic_1164.all;
```

FIGURE 11–12 A PLD that decodes 16-bit I/O ports EFF8H through EFFFH.


```
entity DECODER_11_12 is

port (
      Z, A12, A10, A9, A8, A7, A6, A5, A4, A3, A2, A1, A0: in STD_LOGIC;
      D0, D1, D2, D3, D4, D5, D6, D7: out STD_LOGIC
);

end;

architecture V1 of DECODER_11_12 is

begin

      D0 <= not ( not Z and not A12 and A10 and A9 and A8 and A7 and A6 and A5
            and A4 and A3 and not A2 and not A1 and not A0 );
      D1 <= not ( not Z and not A12 and A10 and A9 and A8 and A7 and A6 and A5
            and A4 and A3 and not A2 and not A1 and A0 );
      D2 <= not ( not Z and not A12 and A10 and A9 and A8 and A7 and A6 and A5
            and A4 and A3 and not A2 and A1 and not A0 );
      D3 <= not ( not Z and not A12 and A10 and A9 and A8 and A7 and A6 and A5
            and A4 and A3 and not A2 and A1 and A0 );
      D4 <= not ( not Z and not A12 and A10 and A9 and A8 and A7 and A6 and A5
            and A4 and A3 and A2 and not A1 and not A0 );
      D5 <= not ( not Z and not A12 and A10 and A9 and A8 and A7 and A6 and A5
            and A4 and A3 and A2 and not A1 and A0 );
      D6 <= not ( not Z and not A12 and A10 and A9 and A8 and A7 and A6 and A5
            and A4 and A3 and A2 and A1 and not A0 );
      D7 <= not ( not Z and not A12 and A10 and A9 and A8 and A7 and A6 and A5
            and A4 and A3 and A2 and A1 and A0 );

end V1;
```
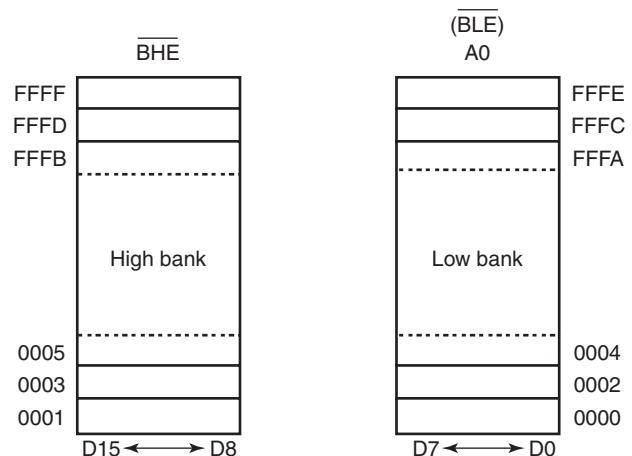
8- and 16-Bit Wide I/O Ports

Now that I/O port addresses are understood and we learned that an I/O port address is probably simpler to decode than a memory address (because of the number of bits), interfacing between the microprocessor and 8- or 16-bit-wide I/O devices is explained. Data transferred to an 8-bit I/O device exist in one of the I/O banks in a 16-bit microprocessor such as the 80386SX. There are 64K different 8-bit ports, but only 32K different 8-bit ports because a 16-bit port uses two 8-bit ports. The I/O system on such a microprocessor contains two 8-bit memory banks, just as memory does. This is illustrated in Figure 11–13, which shows the separate I/O banks for a 16-bit system such as the 80386SX.

FIGURE 11–13 The I/O banks found in the 8086, 80186, 80286, and 80386SX.

Because two I/O banks exist, any 8-bit I/O write requires a separate write strobe to function correctly. I/O reads do not require separate read strobes. As with memory, the microprocessor reads only the byte it expects and ignores the other byte. The only time that a read can cause problems is when the I/O device responds incorrectly to a read operation. In the case of an I/O device that responds to a read from the wrong bank, we may need to include separate read signals. This is discussed later in this chapter.

Figure 11–14 illustrates a system that contains two different 8-bit output devices, located at 8-bit I/O address 40H and 41H. Because these are 8-bit devices and because they appear in different I/O banks, separate I/O write signals are generated to clock a pair of latches that capture port data. Note that all I/O ports use 8-bit addresses. Thus, ports 40H and 41H can each be addressed as separate 8-bit ports, or together as one 16-bit port. The program for the PLD decoder used in Figure 11–14 is illustrated in Example 11–4.

EXAMPLE 11–4

```
-- VHDL code for the decoder of Figure 11–14

library ieee;

use ieee.std_logic_1164.all;
entity DECODER_11_14 is

port (
      BHE, IOWC, A7, A6, A5, A4, A3, A2, A1, A0: in STD_LOGIC;
      D0, D1: out STD_LOGIC
);

end;

architecture V1 of DECODER_11_14 is

begin

      D0 <= BHE or IOWC or A7 or not A6 or A5 or A4 or A3 or A2 or A1 or A0;
      D1 <= BHE or IOWC or A7 or not A6 or A5 or A4 or A3 or A2 or A1 or
            not A0;

end V1;
```
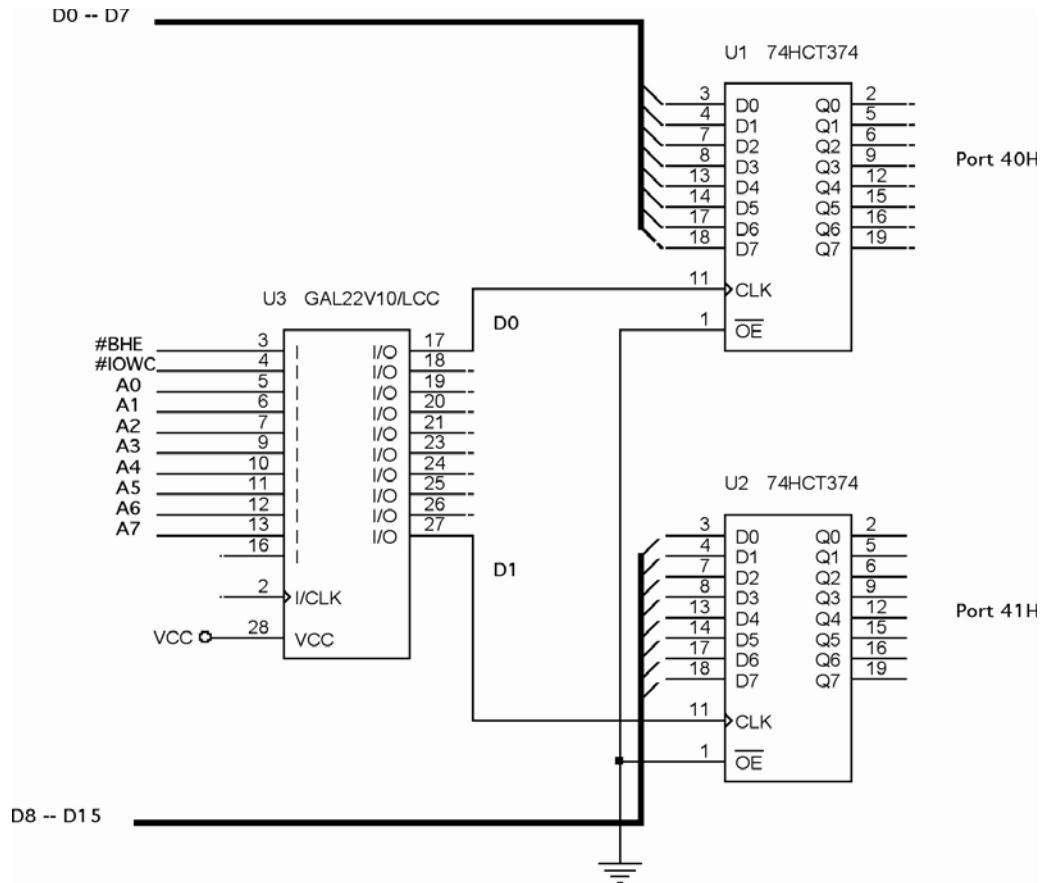
FIGURE 11–14 An I/O port decoder that selects ports 40H and 41H for output data.

When selecting 16-bit-wide I/O devices, the $\overline{BLE}$ (A_0) and $\overline{BHE}$ pins have no function because both I/O banks are selected together. Although 16-bit I/O devices are relatively rare, a few do exist for analog-to-digital and digit-to-analog converters, as well as for some video and disk interfaces.

Figure 11–15 illustrates a 16-bit input device connected to function at 8-bit I/O addresses 64H and 65H. Notice that the PLD decoder does not have a connection for address bits $\overline{BLE}$ (A_0) and $\overline{BHE}$ because these signals do not apply to 16-bit-wide I/O devices. The program for the PLD, illustrated in Example 11–5, shows how the enable signals are generated for the three-state buffers (74HCT244) used as input devices.

EXAMPLE 11–5

```
-- VHDL code for the decoder of Figure 11-15

library ieee;
use ieee.std_logic_1164.all;

entity DECODER_11_15 is

port (
      IORC, A7, A6, A5, A4, A3, A2, A1: in STD_LOGIC;
      D0: out STD_LOGIC
);
```
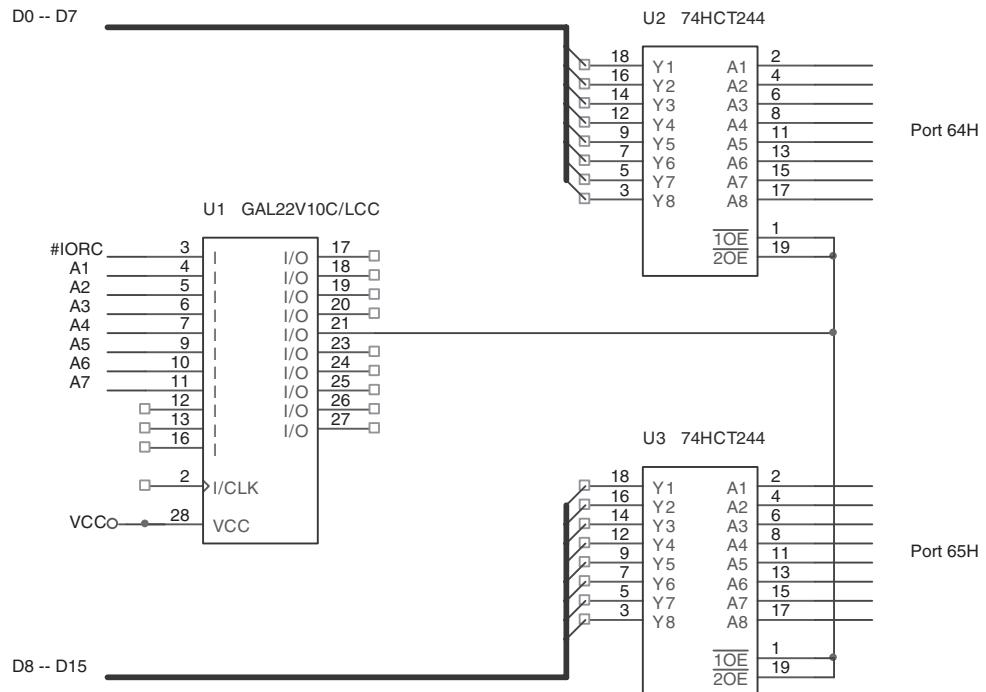
FIGURE 11–15   A 16-bit-wide port decoded at I/O addresses 64H and 65H.

```
end;

architecture V1 of DECODER_11_15 is

begin

     D0 <= IORC or A7 or not A6 or not A5 or A4 or A3 or not A2 or A1;

end V1;
```

32-Bit-Wide I/O Ports

Although 32-bit-wide I/O ports are not common, they may eventually become commonplace because of newer buses found in computer systems. The once-promising EISA system bus supports 32-bit I/O as well as the VESA local and current PCI bus, but not many I/O devices are 32 bits in width.

The circuit of Figure 11–16 illustrates a 32-bit input port for the 80386DX through the 80486DX microprocessor. As with prior interfaces, this circuit uses a single PLD to decode the I/O ports and four 74HCT244 buffers to connect the I/O data to the data bus. The I/O ports decoded by this interface are the 8-bit ports 70H–73H, as illustrated by the PLD program in Example 11–6. Again, we only decode an 8-bit I/O port address. When writing software to access this port, it is crucial to use the address 70H for the 32-bit input as in the instruction IN EAX, 70H.

EXAMPLE 11–6

```
-- VHDL code for the decoder of Figure 11–16

library ieee;
use ieee.std_logic_1164.all;

entity DECODER_11_16 is

port (
```
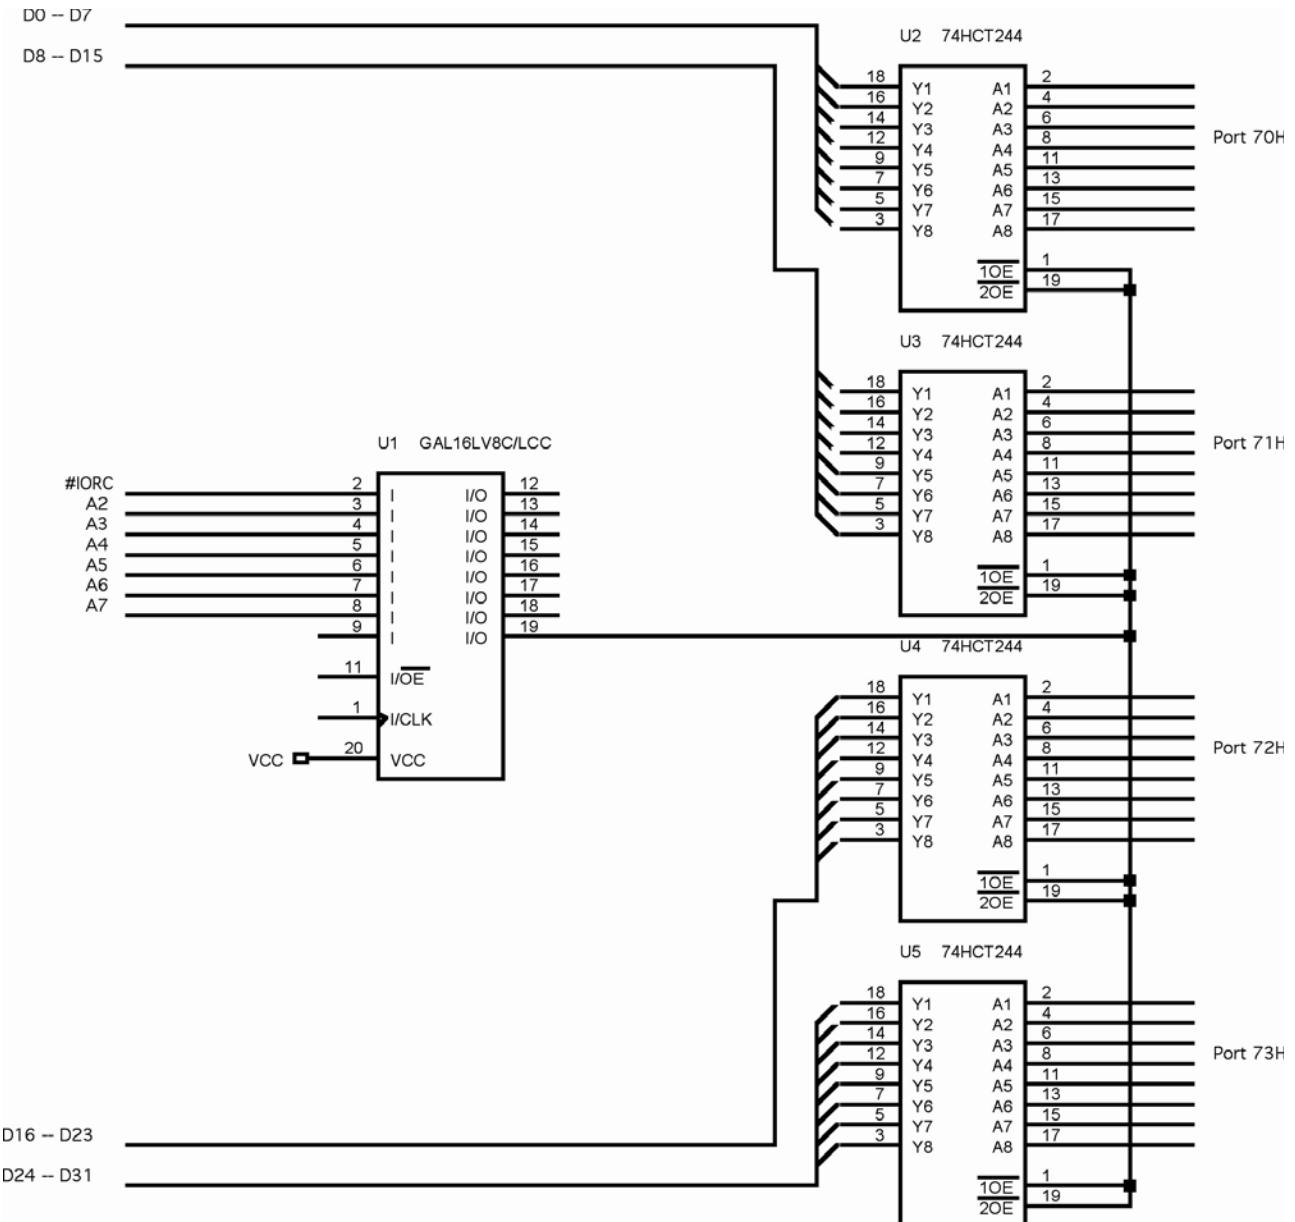
FIGURE 11–16 A 32-bit-wide port decoded at 70H through 73H for the 80486DX microprocessor.

```
            IORC, A7, A6, A5, A4, A3, A2: in STD_LOGIC;
            D0: out STD_LOGIC
);

end;

architecture V1 of DECODER_11_16 is

begin

      D0 <= IORC or A7 or not A6 or not A5 or not A4 or A3 or A2;

end V1;
```
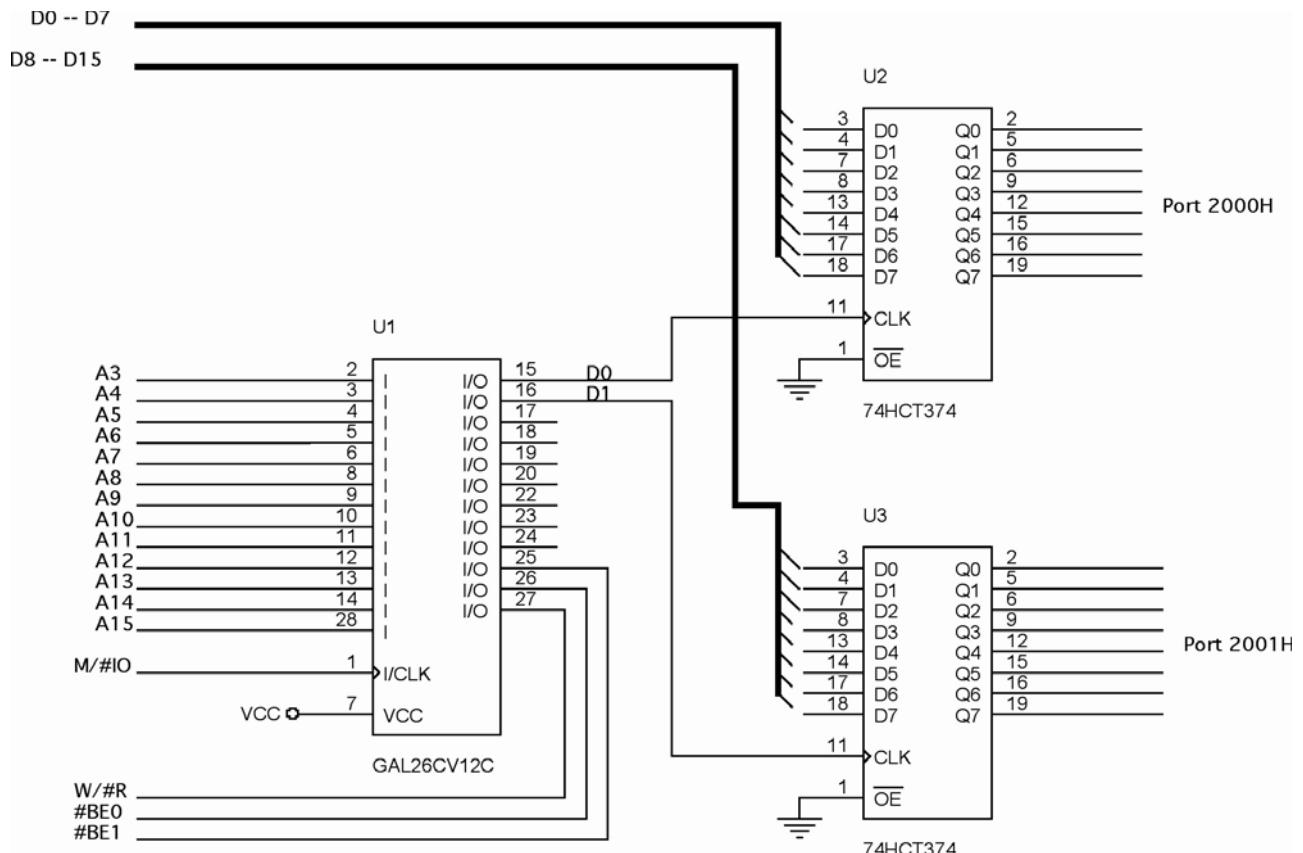
FIGURE 11–17   A Pentium 4 interfaced to a 16-bit-wide I/O port at port addresses 2000H and 2001H.

With the Pentium–Core2 microprocessors and their 64-bit data buses, I/O ports appear in various banks, as determined by the I/O port address. For example, 8-bit I/O port 0034H appears in Pentium I/O bank 4, while the l6-bit I/O ports 0034H–0035H appear in Pentium banks 4 and 5. A 32-bit I/O access in the Pentium system can appear in any four consecutive I/O banks. For example, 32-bit I/O ports 0100H–0103H appear in banks 0–3. The I/O address range must begin at a location where the rightmost two bits are zeros. Hence, 0100H–0103H is allowable but 0101H–0104H is not.

How is a 64-bit I/O device interfaced? The widest I/O transfers are 32 bits, and currently there are no 64-bit I/O instructions to support 64-bit transfers. This event is true for the Pentium 4 or Core2 operated in the 64-bit mode.

Suppose that we need to interface a simple 16-bit-wide output port at I/O port address 2000H and 2001H. The rightmost three bits of the lowest port address are 000 for port 2000H. This means that port 2000H is in memory bank 0. Likewise the rightmost three binary bits of I/O port 2001H are 001, which means that port 2001H is in bank 1. An interface is illustrated in Figure 11–17 and the PLD program is listed in Example 11–7.

The control signal M/$\overline{\text{IO}}$ and W/$\overline{\text{R}}$ must be combined to generate an I/O write signal for the latches and both $\overline{\text{BE0}}$ and $\overline{\text{BE1}}$ bank enable signals must be used to steer the write signal to the correct latch clock for address 2000H (bank 0) and 2001H (bank 1). The only problem that can arise in interfacing is when the I/O port spans across a 64-bit boundary, for example, a 16-bit-wide port located at 2007H and 2008H. In this case, port 2007H uses bank 7 and 2008H uses bank 0, but the address that is decoded is different for each location. A 0010 0000 0000 0XXX is decoded for 2007H and 0010 0000 0000 1XXX is decoded for 2008H. It is probably best to avoid situations such as this.

EXAMPLE 11–7

```
-- VHDL code for the decoder of Figure 11-17

library ieee;
use ieee.std_logic_1164.all;

entity DECODER_11_17 is

port (
      MIO, BE0, BE1, WR, A15, A14, A13, A12, A11, A10, A9, A8, A7, A6, A5, A4,
            A3: in STD_LOGIC;
      D0, D1: out STD_LOGIC
);

end;

architecture V1 of DECODER_11_17 is

begin

      D0 <= MIO or BE0 or not WR or A15 or A14 or not A13 or A12 or A11 or A10
            or A9 or A8 or A7 or A6 or A5 or A4 or A3;
      D1 <= MIO or BE1 or not WR or A15 or A14 or not A13 or A12 or A11 or A10
            or A9 or A8 or A7 or A6 or A5 or A4 or not A3;

end V1;
```

11–3 THE PROGRAMMABLE PERIPHERAL INTERFACE

The 82C55 **programmable peripheral interface** (PPI) is a very popular, low-cost interfacing component found in many applications. This is true even with all the programmable devices available for simple applications. The PPI, which has 24 pins for I/O that are programmable in groups of 12 pins, has groups that operate in three distinct modes of operation. The 82C55 can interface any TTL-compatible I/O device to the microprocessor. The 82C55 (CMOS version) requires the insertion of wait states if operated with a microprocessor using higher than an 8 MHz clock. It also provides at least 2.5 mA of sink (logic 0) current at each output, with a maximum of 4.0 mA. Because I/O devices are inherently slow, wait states used during I/O transfers do not impact significantly upon the speed of the system. The 82C55 still finds application (compatible for programming, although it may not appear in the system as a discrete 82C55), even in the latest Core2-based computer system. The modern computer uses a few 82C55s located inside the chip set for various features on the personal computer. The 82C55 is used for interface to the keyboard and the parallel printer port in many personal computers, but it is found as a function within a interfacing chip set. The chip set also controls the timer and reads data from the keyboard interface.

A low-cost experimentation board is available that plugs into the parallel port of a PC that allows access to an 8255 located on the board. The 8255 is programmed in either assembly language or Visual C++ through drivers available with the board. Visit the following Internet link for pricing and additional information: http://www.microdigitaled.com/hardware/mdelpt/MDELPT.htm.
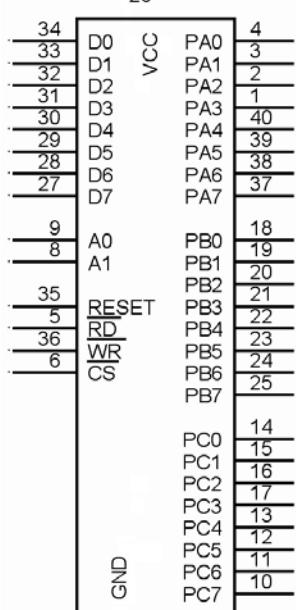
Basic Description of the 82C55

Figure 11–18 illustrates the pin-out diagram of the 82C55 in both the DIP format and the surface mount (flat pack). Its three I/O ports (labeled A, B, and C) are programmed as groups. Group A connections consist of port A (PA_7–PA_0) and the upper half of port C (PC_7–PC_4), and group B consists of port B (PB_7–PB_0) and the lower half of port C (PC_3–PC_0). The 82C55 is selected by its $\overline{CS}$ pin for programming and for reading or writing to a port. Register selection is accomplished through the A_1 and A_0 input pins that select an internal register for programming

FIGURE 11–18 The pin-out of the 82C55 peripheral interface adapter (PPI).
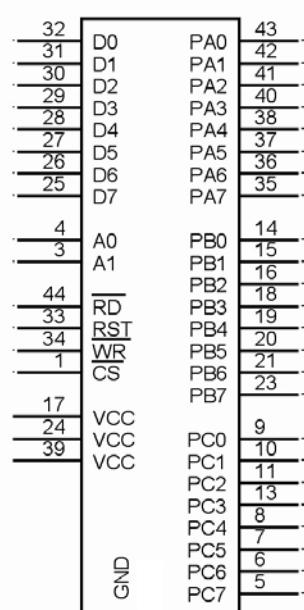
or operation. Table 11–2 shows the I/O port assignments used for programming and access to the I/O ports. In the personal computer a pair of 82C55s, or their equivalents, are decoded at I/O ports 60H–63H and also at ports 378H–37BH.

The 82C55 is a fairly simple device to interface to the microprocessor and program. For the 82C55 to be read or written, the $\overline{CS}$ input must be a logic 0 and the correct I/O address must be applied to the A_1 and A_0 pins. The remaining port address pins are don't cares as far as the 82C55 is concerned, and are externally decoded to select the 82C55.

Figure 11–19 shows an 82C55 connected to the 80386SX so that it functions at 8-bit I/O port addresses C0H (port A), C2H (port B), C4H (port C), and C6H (command register). This interface uses the low bank of the 80386SX I/O map. Notice from this interface that all the 82C55 pins are direct connections to the 80386SX, except for the $\overline{CS}$ pin. The $\overline{CS}$ pin is decoded and selected by a 74ALS138 decoder.

The RESET input to the 82C55 initializes the device whenever the microprocessor is reset. A RESET input to the 82C55 causes all ports to be set up as simple input ports using mode 0 operation. Because the port pins are internally programmed as input pins after a RESET, damage is prevented when the power is first applied to the system. After a RESET, no other commands are needed to program the 82C55, as long as it is used as an input device for all three ports. Note

TABLE 11–2 I/O port assignments for the 82C55.

A_1	A_0	Function
0	0	Port A
0	1	Port B
1	0	Port C
1	1	Command register
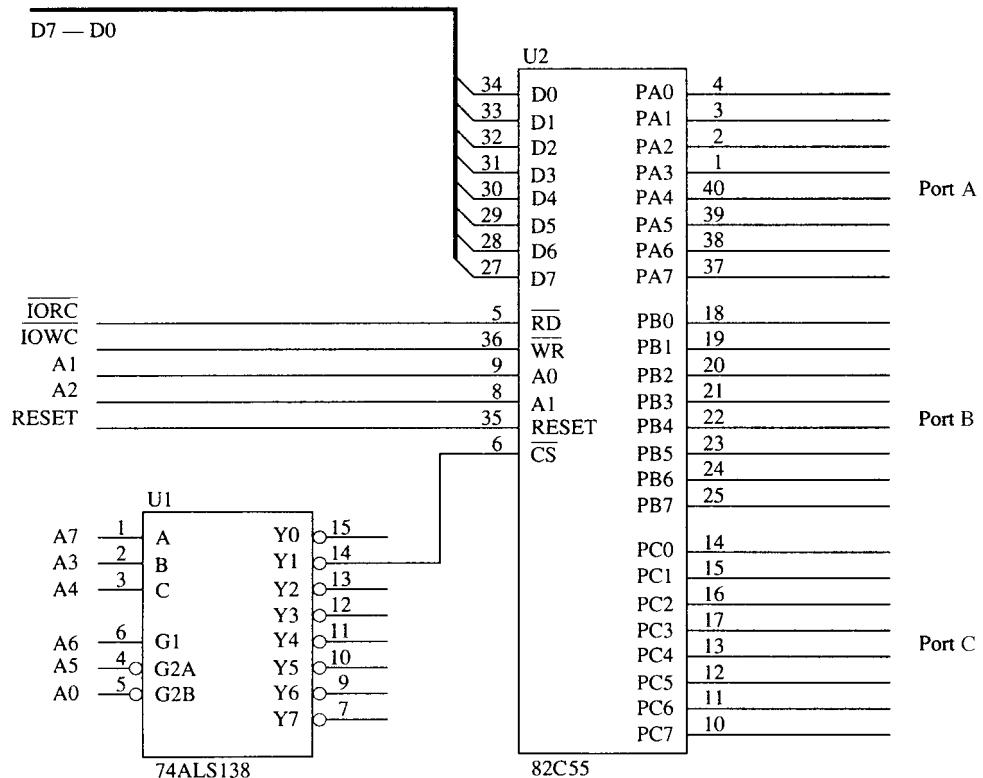
FIGURE 11–19   The 82C55 interfaced to the low bank of the 80386SX microprocessor.

that an 82C55 is interfaced to the personal computer at port addresses 60H–63H for keyboard control, and also for controlling the speaker, timer, and other internal devices such as memory expansion. It is also used for the parallel printer port at I/O ports 378H–37BH.

Programming the 82C55
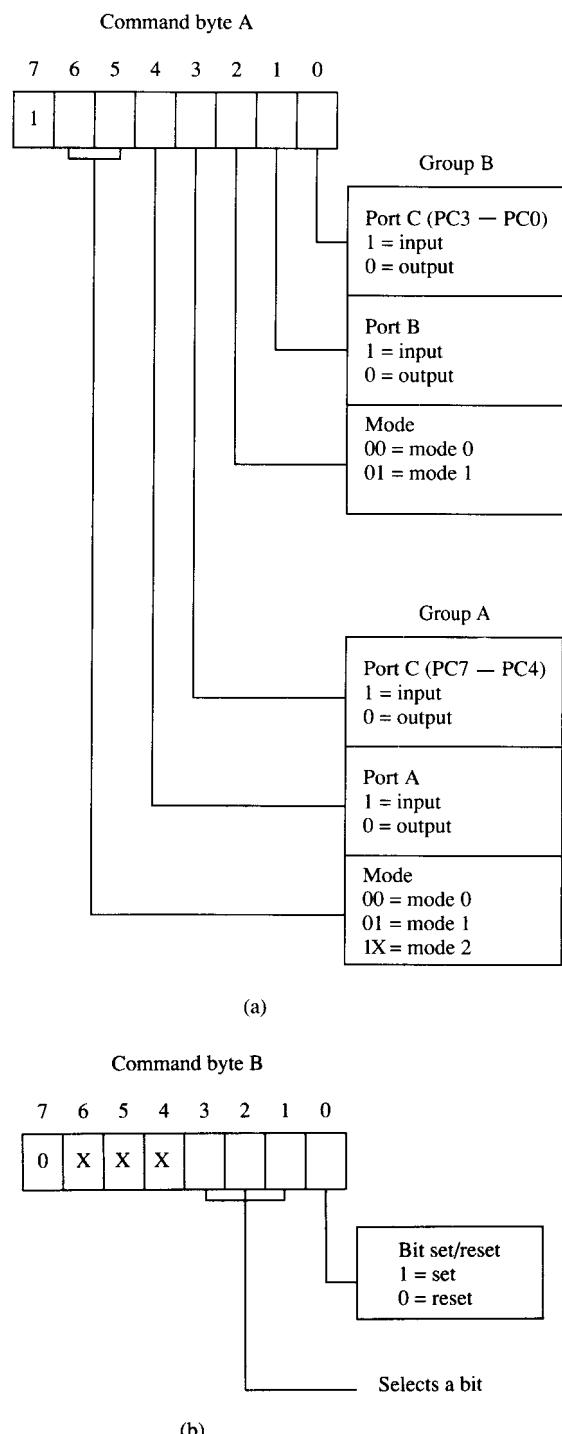
The 82C55 is programmed through the two internal command registers that are illustrated in Figure 11–20. Notice that bit position 7 selects either command byte A or command byte B. Command byte A programs the function of group A and B, whereas command byte B sets (1) or resets (0) bits of port C only if the 82C55 is programmed in mode 1 or 2.

Group B pins (port B and the lower part of port C) are programmed as either input or output pins. Group B operates in either mode 0 or mode 1. Mode 0 is the basic input/output mode that allows the pins of group B to be programmed as simple input and latched output connections. Mode 1 operation is the strobed operation for group B connections, where data are transferred through port B and handshaking signals are provided by port C.

Group A pins (port A and the upper part of port C) are programmed as either input or output pins. The difference is that group A can operate in modes 0, 1, and 2. Mode 2 operation is a bidirectional mode of operation for port A.

If a 0 is placed in bit position 7 of the command byte, command byte B is selected. This command allows any bit of port C to be set (1) or reset (0), if the 82C55 is operated in either mode 1 or 2. Otherwise, this command byte is not used for programming. The bit set/reset feature is often used in a control system to set or clear a control bit at port C. The bit set/reset function is glitch-free, which means that the other port C pins will not change during the bit set/reset command.

FIGURE 11–20 The command byte of the command register in the 82C55. (a) Programs ports A, B, and C. (b) Sets or resets the bit indicated in the select a bit field.

Command byte A

```
 7  6  5  4  3  2  1  0
┌──┬──┬──┬──┬──┬──┬──┬──┐
│1 │  │  │  │  │  │  │  │
└──┴──┴──┴──┴──┴──┴──┴──┘
```

Group B

Port C (PC3 — PC0)
1 = input
0 = output

Port B
1 = input
0 = output

Mode
00 = mode 0
01 = mode 1

Group A

Port C (PC7 — PC4)
1 = input
0 = output

Port A
1 = input
0 = output

Mode
00 = mode 0
01 = mode 1
1X = mode 2

(a)

Command byte B

```
 7  6  5  4  3  2  1  0
┌──┬──┬──┬──┬──┬──┬──┬──┐
│0 │X │X │X │  │  │  │  │
└──┴──┴──┴──┴──┴──┴──┴──┘
```

Bit set/reset
1 = set
0 = reset

Selects a bit

(b)

Mode 0 Operation

Mode 0 operation causes the 82C55 to function either as a buffered input device or as a latched output device. These are the same as the basic input and output circuits discussed in the first section of this chapter.

Figure 11–21 shows the 82C55 connected to a set of eight seven-segment LED displays. These are standard LEDs, but the interface can be modified with a change in resistor values for an organic LED (OLED) display or high-brightness LEDs. In this circuit, both ports A and B are programmed as (mode 0) simple latched output ports. Port A provides the segment data inputs to the display and port B provides a means of selecting one display position at a time for multiplexing the displays. The 82C55 is interfaced to an 8088 microprocessor through a PLD so that it functions at I/O port numbers 0700H–0703H. The program for the PLD is listed in Example 11–8. The PLD decodes the I/O address and develops the write strobe for the $\overline{WR}$ pin of the 82C55.

EXAMPLE 11–8

```
-- VHDL code for the decoder of Figure 11-21

library ieee;
use ieee.std_logic_1164.all;

entity DECODER_11_21 is

port (
      IOM, A15, A14, A13, A12, A11, A10, A9, A8, A7, A6, A5, A4, A3,
         A2: in STD_LOGIC;
      D0: out STD_LOGIC
);

end;

architecture V1 of DECODER_11_17 is

begin

      D0 <= not IOM or A15 or A14 or A13 or A12 or A11 or not A10
      or not A9 or not A8 or A7 or A6 or A5 or A4 or A3 or A2;

end V1;
```

The resistor values are chosen in Figure 11–21 so that the segment current is 80 mA. This current is required to produce an average current of 10 mA per segment as the displays are multiplexed. A six-digit display uses a segment current of 60 mA for an average of 10 mA per segment. In this type of display system, only one of the eight display positions is on at any given instant. The peak anode current in an eight-digit display is 560 mA (seven segments × 80 mA), but the average anode current is 80 mA. In a six-digit display, the peak current would be 420 mA (seven segments × 60 mA). Whenever displays are multiplexed, we increase the segment current from 10 mA (for a display that uses 10 mA per segment as the nominal current) to a value equal to the number of display positions times 10 mA. This means that a four-digit display uses 40 mA per segment, a five-digit display uses 50 mA, and so on.

In this display, the segment load resistor passes 80 mA of current and has a voltage of approximately 3.0 V across it. The LED (1.65 V nominally) and a few tenths are dropped across the anode switch and the segment switch, hence a voltage of 3.0 V appears across the segment load resistor. The value of the resistor is 3.0 V ÷ 180 mA = 37.5 Ω. The closest standard resistor value of 39 Ω is used in Figure 11–21 for the segment load.

The resistor in series with the base of the segment switch assumes that the minimum gain of the transistor is 100. The base current is therefore 80 mA ÷ 100 = 0.8 mA. The voltage across the base resistor is approximately 3.0 V (the minimum logic 1 voltage level of the 82C55), minus the drop across the emitter-base junction (0.7 V), or 2.3 V. The value of the base resistor is therefore 2.3 V ÷ 0.8mA = 2.875 KΩ. The closest standard resistor value is 2.7 KΩ, but 2.2 KΩ is chosen for this circuit.

The anode switch has a single resistor on its base. The current through the resistor is 560 mA ÷ 100 = 5.6 mA because the minimum gain of the transistor is 100. This exceeds the

39

D0 D1 D2 D3 D4 D5 D6 D7

14
13
12
11
10
9
8

1
2
3
4
5
6
7

2N2907

VCC

U1

D0 34 D0   PA0 4 1 14 2.2K
D1 33 D1   PA1 3 2 13
D2 32 D2   PA2 2 3 12
D3 31 D3   PA3 1 4 11
D4 30 D4   PA4 40 5 10
D5 29 D5   PA5 39 6 9
D6 28 D6   PA6 38 7 8
D7 27 D7   PA7 37

2N2222

RD 5 RD   PB0 18 1 16
WR 36 WR   PB1 19 2 15
A0 9 A0   PB2 20 3 14
A1 8 A1   PB3 21 4 13
RESET 35 RESET PB4 22 5 12
6 CS   PB5 23 6 11
               PB6 24 7 10
               PB7 25 8 9

690

PC0 14
PC1 15
PC2 16
PC3 17
PC4 13
PC5 12
PC6 11
PC7 10

8255

U2

A2 1 I1   O1 19
A3 2 I2   O2 18
A4 3 I3   O3 17
A5 4 I4   O4 16
A6 5 I5   O5 15
A7 6 I6   O6 14
A8 7 I7   O7 13
A9 8 I8   O8 12
A10 9 I9
A11 11 I10

16L8

IO/M
A12
A13
A14
A15

FIGURE 11–21 An 8-digit LED display interfaced to the 8088 microprocessor through an 82C55 PIA.

maximum current of 4.0 mA from the 82C55, but this is close enough so that it will work without problems. The maximum current assumes that you are using the port pin as a TIL input to another circuit. If the amount of current were over 8.0–10.0 mA, then appropriate circuitry (in the form of either a Darlington-pair or another transistor switch) would be required. Here, the voltage across the base resistor is 5.0 V, minus the drop across the emitter-base junction (0.7 V), minus the voltage at the port pin (0.4 V), for a logic 0 level. The value of the resistor is 3.9 V ÷ 5.66 mA = 68.9 Ω. The closest standard resistor value is 69 Ω, which is chosen for this example.

Before software to operate the display is examined, we must first program the 82C55. This is accomplished with the short sequence of instructions listed in Example 11–9. Here, ports A and B are both programmed as outputs.

EXAMPLE 11–9

```
;programming the 82C55 PIA

      MOV AL,10000000B        ;command
      MOV DX,703H             ;address port 703H
      OUT DX,AL               ;send command to port 703H
```

The procedure to multiplex the displays is listed in Example 11–10 in both assembly language and C++ with assembly language. For the display system to function correctly, we must call this procedure often. Notice that the procedure calls another procedure (DELAY) that causes a 1.0 ms time delay. The time delay is not illustrated in this example, but it is used to allow time for each display position to turn on. Manufacturers of LED displays recommend that the display flashes between 100 and 1500 times per second. Using a 1.0 ms time delay, each digit is displayed for 1.0 ms for a total display flash rate of 1000 Hz ÷ 8 or a flash rate of 125 Hz for all eight digits.

EXAMPLE 11–10

```
;An assembly language procedure that multiplexes the 8-digit display.
;This procedure must be called often for the display
;to appear correctly.

DISP  PROC  NEAR USES AX BX DX SI

      PUSHF
      MOV   BX,8                 ;load counter
      MOV   AH,7FH               ;load selection pattern
      MOV   SI,OFFSET MEM-1      ;address display data
      MOV   DX,701H              ;address Port B

;display all 8 digits

      .REPEAT
            MOV   AL,AH          ;send selection pattern to Port B
            OUT   DX,AL
            DEC   DX
            MOV   AL,[BX+SI]     ;send data to Port A
            OUT   DX,AL
            CALL  DELAY          ;wait 1.0 ms
            ROR   AH,1           ;adjust selection pattern
            INC   DX
            DEC   BX             ;decrement counter
      .UNTIL BX == 0

      POPF
      RET

DISP  ENDP

// A C/C++ function that multiplexes the 8-digit displays
// uses char sized array MEM
```

```
void Disp()
{
      unsigned int *Mem = &MEM[0];          //point to array element 0
      for ( int a = 0; a < 8; a++ )
      {
            unsigned char b = 0xff ^ ( 1 << a );        //form select pattern
            _asm
            {
                mov  al,b
                mov  dx,701H
                out  dx,al       ;send select pattern to Port B
                mov  al,Mem[a]
                dec  dx
                out  dx,al       ;send data to Port A
            }
            Sleep(1);                  ;wait 1.0 ms
      }
}
```

The display procedure (DISP) addresses an area of memory where the data, in seven-segment code, are stored for the eight display digits called MEM. The AH register is loaded with a code (7FH) that initially addresses the most significant display position. Once this position is selected, the contents of memory location MEM +7 is addressed and sent to the most significant digit. The selection code is then adjusted to select the next display digit. This process repeats eight times to display the contents of location MEM through MEM +7 on the eight display digits.

The time delay of 1.0 ms can be obtained by writing a procedure that uses the system clock frequency to determine how long each instruction requires to execute. The procedure listed in Example 11–11 causes a time delay of a duration determined by the number of times that the LOOP instruction executes. Here XXXX was used and will be filled in with a value after a few facts are discussed. The LOOP instruction requires a certain number of clocks to execute—how many can be located in Appendix B. Suppose that the interface is using the 80486 microprocessor running with a 20 MHz clock. Appendix B represents that the LOOP instruction requires 7/6 clocks. The first number is the number of clocks required when a jump to D_1 occurs and the second number is when the jump does not occur. With a 20 MHz clock, one clock requires $1 \div 20$ MHz = 50 ns. The LOOP instruction, in this case, requires 350 ns to execute in all but the very last iteration. To determine the count (XXXX) needed to accomplish a 1.0 ms time delay, divide 1.0 ms by 350 ns. In this case XXXX = 2,857 to accomplish a 1.0 ms time delay. If a larger count occurs, a LOOPD instruction can be used with the ECX register. The time required to execute the MOV CX, XXXX, and RET instructions can usually be ignored.

Suppose a Core2 with a 2.0 GHz clock is used for the delay. Here one clock is 0.5 ns and LOOP requires five clocks per iteration. This requires a count of 400,000, so LOOPD would be used with ECX.

EXAMPLE 11–11

```
;equation for the delay
;
;        Delay Time
;XXXX = -------------
;        time for LOOP
;

DELAY   PROC  NEAR USES CX

        MOV CX,XXXX
D1:
        LOOP D1
        RET

DELAY   ENDP
```
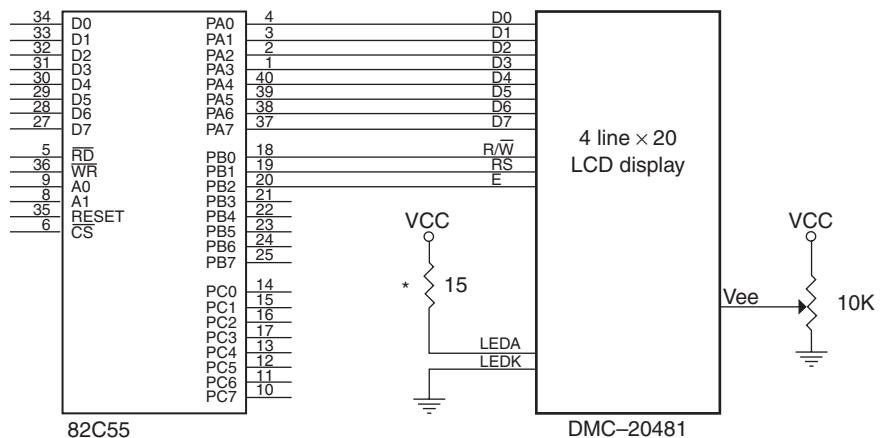
FIGURE 11–22 The DMC-20481 LCD display interfaced to the 82C55.

*Current max is 480 mA, nominal 260 mA.

If the program is written for the Windows environment, such as for use in an embedded system using embedded Windows, time delays can use a timer. The timer can operate with a precision of milliseconds, and in the embedded version of Windows, the delays are guaranteed.

An LCD Display Interfaced to the 82C55

LCDs (**liquid crystal displays**) have replaced LED displays in many applications. The only disadvantage of the LED display is that it is difficult to see in low-light situations in which the LED is still in limited use. An example is medical equipment for older people with poor eyesight. If the price of the OLED becomes low enough, LCD displays will disappear. A German company currently manufactures an OLED display panel that sells for under $10.

Figure 11–22 illustrates the connection of the Optrex DMC-20481 LCD display interfaced to an 82C55. The DMC-20481 is a 4-line by 20-characters-per-line display that accepts ASCII code as input data. It also accepts commands that initialize it and control its application. As you can see in Figure 11–22, the LCD display has few connections. The data connections, which are attached to the 82C55 port A, are used to input display data and to read information from the display. This illustrates an 8-bit interface. If a 4-bit interface is desired, D_4–D_7 pins are used for the data where the data must be formatted with the high nibble first, followed by the low nibble. A few newer OLED devices also contain a serial interface that uses a single pin for the data.

There are four control pins on the display. The V_{EE} connection is used to adjust the contrast of the LED display and is normally connected to a 10 KΩ potentiometer, as illustrated. The RS (register select) input selects data (RS = 1) or instructions (RS = 0). The E (enable) input must be a logic 1 for the DMC-20481 to read or write information and functions as a clock. Finally, the R/$\overline{W}$ pin selects a read or a write operation. Normally, the RS pin is placed at a 1 or 0, the R/$\overline{W}$ pin is set or cleared, data are placed on the data input pins, and then the E pin is pulsed to access the DMC-20481. This display also has two inputs (LEDA [anode] and LEDK [cathode]) for back-lighting LED diodes, which are not shown in the illustration.

In order to program the DMC-20481 we must first initialize it. This applies to any display that uses the HD44780 (Hitachi) display driver integrated circuit. The entire line of small display panels from Optrex and most other manufacturers is programmed in the same manner. Initialization is accomplished via the following steps:

1. Wait at least 15 ms after V_{CC} rises to 5.0 V.
2. Output the function set command (30H), and wait at least 4.1 ms.
3. Output the function set command (30H) a second time, and wait at least 100 μs.
4. Output the function set command (30H) a third time, and wait at least 40 μs.
5. Output the function set command (38H) a fourth time, and wait at least 40 μs.

403

6. Output 08H to disable the display, and wait at least 40 μs.
7. Output a 01H to home the cursor and clear the display, and wait at least 1.64 ms.
8. Output the enable display cursor off (0CH), and wait at least 40 μs.
9. Output 06H to select auto-increment, shift the cursor, and wait at least 40 μs.

The software to accomplish the initialization of the LCD display is listed in Example 11–12. It is long, but the display controller requires the long initialization dialog. Note that the software for the three time delays is not included in the listing. If you are interfacing to a PC, you can use the RDTSC instruction as discussed in the Pentium chapter for the time delay. If you are developing the interface for another application, then you must write separate time delays, which must provide the delay times indicated in the initialization dialog. The time delays can also be obtained by using a timer in C++.

EXAMPLE 11–12

```
PORTA_ADDRESS   EQU 700H          ;set port addresses
PORTB_ADDRESS   EQU 701H
COMMAND_ADDRESS EQU 703H

;macro to send a command or data to the LCD display
;
SEND  MACRO PORTA_DATA, PORTB_DATA, DELAY

      MOV  AL,PORTA_DATA       ;PORTA_DATA to Port A
      MOV  DX,PORTA_ADDRESS
      OUT  DX,AL
      MOV  AL,PORTB_DATA       ;PORTB_DATA to Port B
      MOV  DX,PORTB_ADDRESS
      OUT  DX,AL
      OR   AL,00000100B        ;Set E bit
      OUT  DX,AL               ;send to Port B
      AND  AL,11111011B        ;Clear E bit
      NOP                      ;a small delay
      NOP
      OUT  DX,AL               ;send to Port B
      MOV  BL,DELAY            ;BL = delay count
      CALL MS_DELAY            ;ms Time Delay
      ENDM

;Program to initialize the LCD display

START:
      MOV  AL,80H              ;Program the 82C55
      MOV  DX,COMMAND_ADDRESS
      OUT  DX,AL
      MOV  AL,0
      MOV  DX,PORTB_ADDRESS    ;Clear Port B
      SEND 30H, 2, 16          ;send 30H for 16 ms
      SEND 30H, 2, 5           ;send 30H for 5 ms
      SEND 30H, 2, 1           ;send 30H for 1 ms
      SEND 38H, 2, 1           ;send 38H for 1 ms
      SEND 8, 2, 1             ;send 8 for 1 ms
      SEND 1, 2, 2             ;send 1 for 2 ms
      SEND 0CH, 2, 1           ;send 0CH for 1 ms
      SEND 6, 2, 1             ;send 6 for 1 ms
```

The NOP instructions are added in the SEND macro to ensure that the E bit remains a logic 1 long enough to activate the LCD display. This process should work in most systems at most clock frequencies, but additional NOP instructions may be needed to lengthen this time in some cases. Also notice that equate statements are used to equate the port addresses to labels. This is done so that the software can be changed easily if the port numbers differ from those used in the program.

Before programming the display, the commands used in the initialization dialog must be explained. See Table 11–3 for a complete listing of the commands or instructions for the LCD display. Compare the commands sent to the LCD display in the initialization program to Table 11–3.

TABLE 11–3   Instructions for most LCD displays.

Instruction	Code	Description	Time
Clear display	0000 0001	Clears the display and homes the cursor	1.64 ms
Cursor home	0000 0010	Homes the cursor	1.64 ms
Entry mode set	0000 00AS	Sets cursor movement direction (A = 1, increment) and shift (S = 1, shift)	40 μs
Display on/off	0000 1DCB	Sets display on/off (D = 1, on) (C = 1, cursor on) (B = 1, cursor blink)	40 μs
Cursor/display shift	0001 SR00	Sets cursor movement and display shift (S = 1, shift display) (R = 1, right)	40 μs
Function set	001L NF00	Programs LCD circuit (L = 1, 8-bit interface) (N = 1, 2 lines) (F = 1, 5 × 10 characters) (F = 0, 5 × 7 characters)	40 μs
Set CGRAM address	01XX XXXX	Sets character generator RAM address	40 μs
Set DRAM address	10XX XXXX	Sets display RAM address	40 μs
Read busy flag	B000 0000	Reads busy flag (B = 1, busy)	0
Write data	Data	Writes data to the display RAM or the character generator RAM	40 μs
Read data	Data	Reads data from the display RAM or character generator RAM	40 μs

Once the LCD display is initialized, a few procedures are needed to display information and control the display. After initialization, time delays are no longer needed when sending data or many commands to the display. The clear display command still needs a time delay because the busy flag is not used with that command. Instead of a time delay, the busy flag is tested to see whether the display has completed an operation. A procedure to test the busy flag appears in Example 11–13. The BUSY procedure tests the LCD display and only returns when the display has completed a prior instruction.

EXAMPLE 11–13

```
PORTA_ADDRESS    EQU 700H    ;set port addresses
PORTB_ADDRESS    EQU 701H
COMMAND_ADDRESS  EQU 703H

BUSY   PROC   NEAR USES DX AX

       PUSHF
       MOV   DX,COMMAND_ADDRESS
       MOV   AL,90H           ;program Port A as IN
       OUT   DX,AL
       .REPEAT
            MOV    AL,5      ;select read from LCD
            MOV    DX,PORTB_ADDRESS
            OUT    DX,AL     ;and pulse E
            NOP
            NOP
            MOV    AL,1
            OUT    DX,AL
```

```
                    MOV  DX,PORTA_ADDRESS
                    MOV  AL,DX      ;read busy command
                    SHL  AL,1
              .UNTIL !CARRY?        ;until not busy
            NOV DX,COMMAND_ADDRESS
            MOV AL,80H
            OUT DX,AL               ;program Port A as OUT
            POPF
            RET

BUSY ENDP
```

Once the BUSY procedure is available, data can be sent to the display by writing another procedure called WRITE. The WRITE procedure uses BUSY to test before trying to write new data to the display. Example 11–14 shows the WRITE procedure, which transfers the ASCII character from the BL register to the current cursor position of the display. Note that the initialization dialog has sent the cursor for auto-increment, so if WRITE is called more than once, the characters written to the display will appear one next to the other, as they would on a video display.

EXAMPLE 11–14

```
WRITE PROC  NEAR
      MOV  AL,BL                    ;BL to Port A
      MOV  DX,PORTA_ADDRESS
      OUT  DX,AL
      MOV  AL,0                     ;write ASCII
      MOV  DX,PORTB_ADDRESS
      OUT  DX,AL
      OR   AL,00000100B             ;Set E bit
      OUT  DX,AL                    ;send to Port B
      AND  AL,11111011B             ;Clear E bit
      NOP                           ;a small delay
      NOP
      OUT  DX,AL                    ;send to Port B
      CALL BUSY                     ;wait for completion
      RET
WRITE ENDP
```

The only other procedure that is needed for a basic display is the clear and home cursor procedure, called CLS, shown in Example 11–15. This procedure uses the SEND macro from the initialization software to send the clear command to the display. With CLS and the procedures presented thus far, you can display any message on the display, clear it, display another message, and basically operate the display. As mentioned earlier, the clear command requires a time delay (at least 1.64 ms) instead of a call to BUSY for proper operation.

EXAMPLE 11–15

```
CLS     PROC    NEAR
        SEND    1, 2, 2
        RET
CLS     ENDP
```

Additional procedures that could be developed might select a display RAM position. The display RAM address starts at 0 and progresses across the display until the last character address on the first line is location 19, location 20 is the first display position of the second line, and so forth. Once you can move the display address, you can change individual characters on the display and even read data from the display. These procedures are for you to develop if they are needed.

A word about the display RAM inside of the LCD display. The LCD contains 128 bytes of memory, addressed from 00H to 7FH. Not all of this memory is always used. For example, the
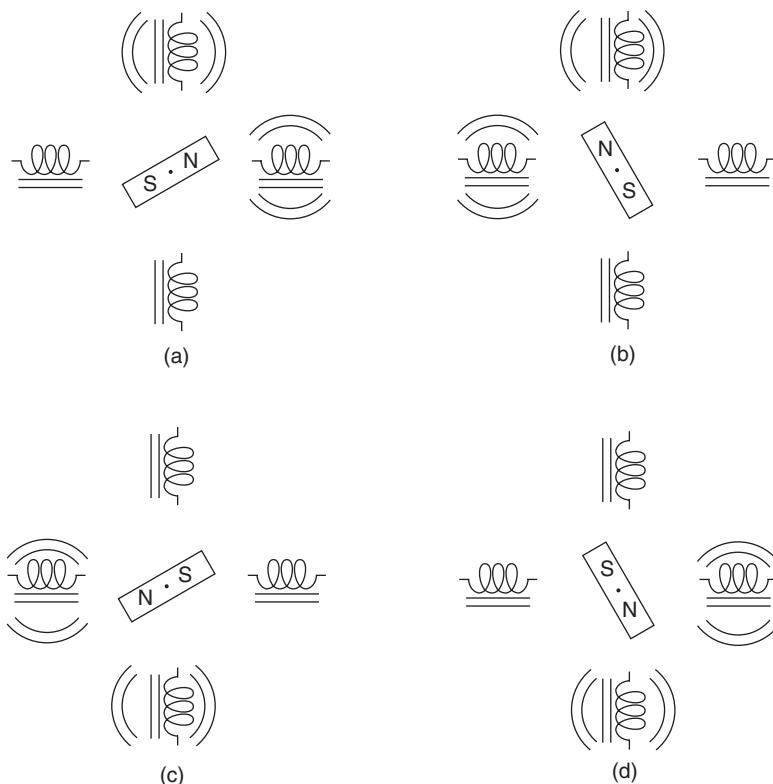
one-line × 20-character display uses only the first 20 bytes of memory (00–13H.) The first line of any of these displays always starts at address 00H. The second line of any display powered by the HD44780 always begins at address 40H. For example, a two-line × 40-character display uses addresses 00H–27H to store ASCII-coded data from the first line. The second line is stored at addresses 40H–67H for this display. In the four-line displays, the first line is at 00H, the second is at 40H, the third is at 14H, and the last line is at 54H. The largest display device that uses the HD44780 is a two-line × 40-character display. The four-line by 40-character display uses an M50530 or a pair of HD44780s. Because information on these devices can be readily found on the Internet, they are not covered in the text.

A Stepper Motor Interlaced to the 82C55. Another device often interfaced to a computer system is the **stepper motor**. A stepper motor is a digital motor because it is moved in discrete steps as it traverses through 360°. A common stepper motor is geared to move perhaps 15° per step in an inexpensive stepper motor, to 1° per step in a more costly, high-precision stepper motor. In all cases, these steps are gained through many magnetic poles and/or gearing. Notice that two coils are energized in Figure 11–23. If less power is required, one coil may be energized at a time, causing the motor to step at 45°, 135°, 225°, and 315°.

Figure 11–23 shows a four-coil stepper motor that uses an armature with a single pole. Notice that the stepper motor is shown four times with the armature (permanent magnetic) rotated to four discrete places. This is accomplished by energizing the coils, as shown. This is an illustration of full stepping. The stepper motor is driven by using NPN Darlington amplifier pairs to provide a large current to each coil.

A circuit that can drive this stepper motor is illustrated in Figure 11–24, with the four coils shown in place. This circuit uses the 82C55 to provide it with the drive signals that are used to rotate the armature of the motor in either the right-hand or left-hand direction.

FIGURE 11–23 The stepper motor showing full-step operation: (a) 45° (b) 135° (c) 225° (d) 315°.
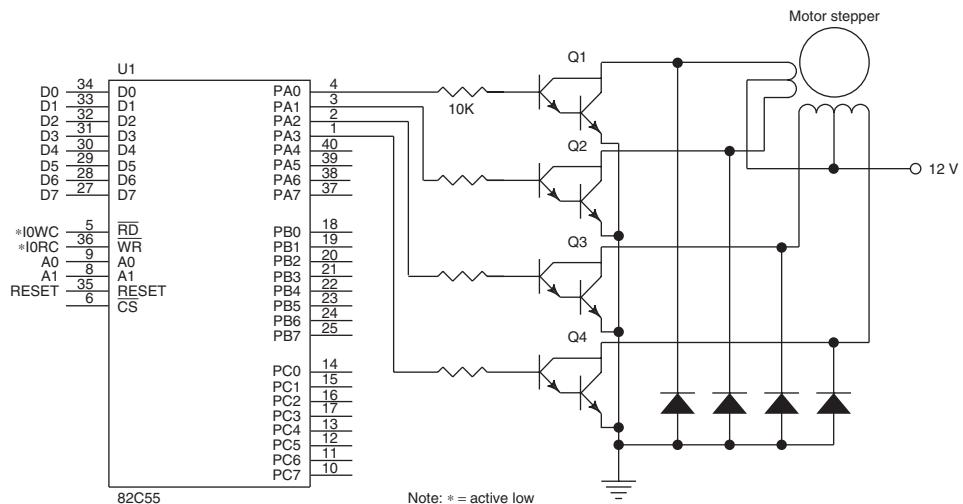
FIGURE 11–24   A stepper motor interfaced to the 82C55. This illustration does not show the decoder.

A simple procedure that drives the motor (assuming that port A is programmed in mode 0 as an output device) is listed in Example 11–16 in both assembly language and as a function in C++. This subroutine is called, with CX holding the number of steps and direction of the rotation. If CX is greater than 8000H, the motor spins in the right-hand direction; if CX is less than 8000H, it spins in the left-hand direction. For example, if the number of steps is 0003H, the motor moves in the left-hand direction three steps and if the number of steps is 8003H, it moves three steps in the right-hand direction. The leftmost bit of CX is removed and the remaining 15 bits contain the number of steps. Notice that the procedure uses a time delay (not illustrated) that causes a l ms time delay. This time delay is required to allow the stepper-motor armature time to move to its next position.

EXAMPLE 11–16

```
      PORT   EQU   40H

;An assembly language procedure that controls the stepper motor

STEP   PROC   NEAR USES CX AX

      MOV   AL,POS                  ;get position
      OR   CX,CX                    ;set flag bits
      IF !ZERO?
          .IF !SIGN?                ;if no sign
                .REPEAT
                   ROL   AL,1       ;rotate step left
                   OUT   PORT,AL
                   CALL DELAY       ;wait 1 ms
                .UNTILCXZ
          .ELSE
                   AND CX,7FFFH      ;make CX positive
                   .REPEAT
                   ROR   AL,1       ;rotate step right
                   OUT   PORT,AL
                   CALL DELAY       ;wait 1 ms
                .UNTILCXZ
          .ENDIF
      .ENDIF
      MOV POS,AL
      RET

STEP   ENDP
```

```
// A C++ function that controls the stepper motor

char Step(char Pos, short Step)
{
      char Direction = 0;
      if (Step < 0)
      {
            Direction = 1;
            Step =& 0x8000;
      }
      while (Step)
      {
            if (Direction)
            {
                  if ((Pos & 1) == 1)
                  {
                        Pos = (Pos >> 1) | 0x80;
                  }
                  else
                  {
                        Pos >>= 1;
                  }
            }
            else
            {
                  if ((Pos & 0x80) == 0x80)
                  {
                        Pos = (Pos << 1) | 1;
                  }
                  else
                  {
                        Pos <<= 1;
                  }
            }
            _asm
            {
                  mov al,Pos
                  out 40h, al
            }
      }
      return Pos;
}
```

The current position is stored in memory location POS, which must be initialized with 33H, 66H, 0EEH, or 99H. This allows a simple ROR (step right) or ROL (step left) instruction to rotate the binary bit pattern for the next step.

The C++ version has two parameters: Pos is the current position of the stepper motor and Step is the number of steps as described earlier. The new Pos is returned in the C++ version instead of being stored in a variable.

Stepper motors can also be operated in the half-step mode, which allows eight steps per sequence. This is accomplished by using the full-step sequence described with a half step obtained by energizing one coil interspersed between the full steps. Half-stepping allows the armature to be positioned at 0°, 90°, 180°, and 270°. The half-step position codes are 11H, 22H, 44H, and 88H. A complete sequence of eight steps would be as follows: 11H, 33H, 22H, 66H, 44H, 0CCH, 88H, and 99H. This sequence could be either output from a lookup table or generated with software.

Key Matrix Interface.   Keyboards come in a vast variety of sizes, from the standard 101-key QWERTY keyboards interfaced to the microprocessor to small specialized keyboards that may contain only four to 16 keys. This section of the text concentrates on the smaller keyboards that may be purchased preassembled or may be constructed from individual key switches.

Figure 11–25 illustrates a small key-matrix that contains 16 switches interfaced to ports A and B of an 82C55. In this example, the switches are formed into a 4×4 matrix, but any matrix could be used, such as a 2×8. Notice how the keys are organized into four rows (ROW_0–ROW_3)
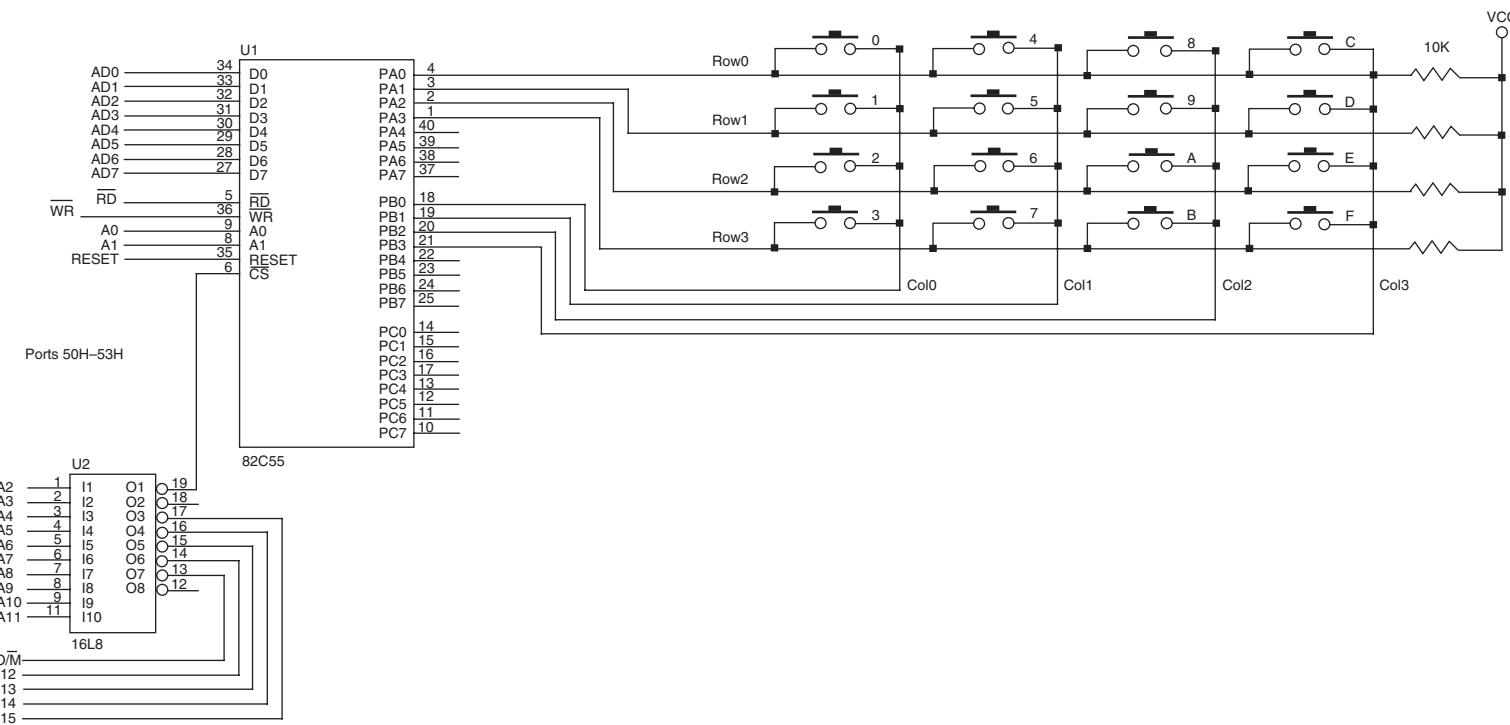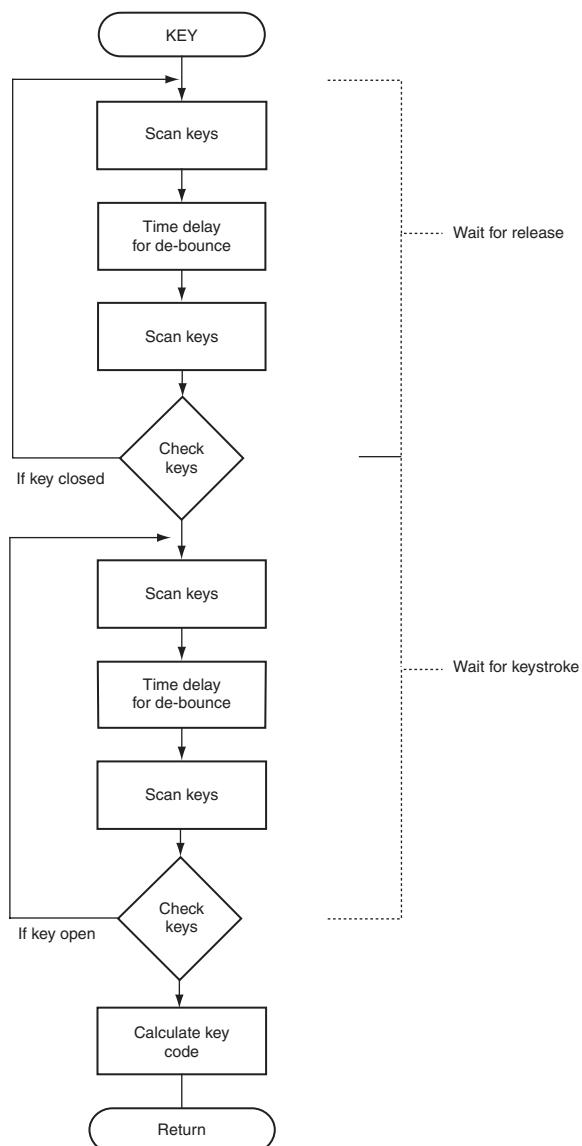
FIGURE 11–25   A 4 × 4 keyboard matrix connected to an 8088 microprocessor through the 82C55 PIA.

and four columns (COL_0–COL_3). Each row is connected to 5.0 V through a 10 KΩ pull-up resistor to ensure that the row is pulled high when no push-button switch is closed.

The 82C55 is decoded (the PLD program is not shown) at I/O ports 50H–53H for an 8088 microprocessor. Port A is programmed as an input port to read the rows and port B is programmed as an output port to select a column. For example, if 1110 is output to port B pins PB_3–PB_0, column 0 has a logic 1, so the four keys in column 0 are selected. Notice that with a logic 0 on PB0, the only switches that can place a logic 0 onto port A are switches 0–3. If switches 4–F are closed, the corresponding port A pins remain a logic 1. Likewise, if 1101 is output to port B, switches 4–7 are selected, and so forth.

A flowchart of the software required to read a key from the keyboard matrix and debounce the key is illustrated in Figure 11–26. Keys must be debounced, which is normally accomplished with a short time delay of 10–20 ms. The flowchart contains three main sections. The first waits for the release of a key. This seems awkward, but software executes very quickly in a microprocessor and there is a possibility that the program will return to the top of this program before

FIGURE 11–26 The flowchart of a keyboard-scanning procedure.

the key is released, so we must wait for a release first. Next, the flowchart shows that we wait for a keystroke. Once the keystroke is detected, the position of the key is calculated in the final part of the flowchart.

The software uses a procedure called SCAN to scan the keys and another called DELAY10 (not shown in this example) to waste 10 ms of time for debouncing. The main keyboard procedure is called KEY and it appears with the others in Example 11–17. Example 11–17 also lists a C++ function to accomplish a key read operation. Note that the KEY procedure is generic, so it can handle any keyboard configuration from a 1×1 matrix to an 8×8 matrix. Changing the two equates at the start of the program (ROWS and COLS) will change the configuration of the software for any size keyboard. Also note that the steps required to initialize the 82C55 so that port A is an input port and port B is an output port are not shown.

With certain keyboards that do not follow the way keys are scanned, a lookup table may be needed to convert the raw key codes returned by KEY into key codes that match the keys on the keyboard. The lookup software is placed just before returning from KEY. It is merely a MOV BX,OFFSET TABLE followed by the XLAT instruction.

EXAMPLE 11–17(a)

```
;assembly language version;

;KEY scans the keyboard and returns the key code in AL.

        COLS  EQU  4
        ROWS  EQU  4
        PORTA EQU  50H
        PORTB EQU  51H

KEY    PROC  NEAR USES CX BX

        MOV  BL,FFH              ;compute row mask
        SHL  BL,ROWS

        MOV  AL,0
        OUT  PORTB,AL            ;place zeros on Port B

        .REPEAT                  ;wait for release
              .REPEAT
                    CALL SCAN
              .UNTIL ZERO?
              CALL DELAY10
              CALL SCAN
        .UNTIL ZERO?
        .REPEAT                  ;wait for key
              .REPEAT
                      CALL SCAN
              .UNTIL !ZERO?
              CALL DELAY10
              CALL SCAN
        .UNTIL !ZERO?
        MOV CX,00FEH
        .WHILE 1                 ;find column
              MOV  AL,CL
              OUT  PORTB,AL
              CALL SHORTDELAY     ;see text
              CALL SCAN
              .BREAK !ZERO?
              ADD CH,COLS
              ROL CL,1
        .ENDW
        .WHILE 1                 ;find row
              SHR  AL,1
              .BREAK .IF !CARRY?
              INC  CH
        .ENDW
```

```
        MOV AL,CH                  ;get key code
        RET

KEY     ENDP

SCAN    PROC   NEAR

        IN   AL,PORTA              ;read rows
        OR   AL,BL
        CMP AL,0FFH                ;test for no keys
        RET

SCAN        ENDP
```

EXAMPLE 11–17(b)

```cpp
// C++ language version of keyboard scanning software

#define ROWS 4
#define COLS 4
#define PORTA 50h
#define PORTB 51h

char Key()
{
        char mask = 0xff << ROWS;
        _asm
        {
                mov al,0                 ;select all columns
                out PORTB,al
        }
        do
        {                                       //wait for release
                while (Scan(mask));
                Delay();
        }
        while (Scan(mask));
        do
        {                                       //wait for key press
                while (!Scan(mask));
                Delay();
        }
        while (!Scan(mask));
        unsigned char select = 0xfe;
        char key = 0;
        _asm
        {
                mov  al,select
                out  PortB,al
        }
        ShortDelay();
        while(!Scan(mask))
        {                                       //calculate key code
                _asm
                {
                        mov  al,select
                        rol  al,1
                        mov  select,al
                        out  PortB,al
                }
                ShortDelay();
                key += COLS;
        }
        _asm
        {
                in   al,PortA
                mov  select,al
        }
```

```
        while ((Select & 1) != 0)
        {
                Select <<= 1;
                key ++;
        }
        return key;
}

bool Scan(mask)
{
        bool flag;
        _asm
        {
                in   al,PORTA
                mov  flag,al
        }
        return (flag | mask);
}
```

The ShortDelay procedure is needed because the computer changes port B at a very high rate of speed. The short time delay allows time for the data sent to port B to settle to their final state. In most cases, this is not needed if the scan rate (time between output instructions) of this part of the software does not exceed 30 KHz. If the scanning frequency is higher, the device generates radio interference. If it does, the Federal Communications Commission (FCC) will not approve its application in any accepted system. Without FCC Type A or Type B certification the system cannot be sold.

Mode 1 Strobed Input
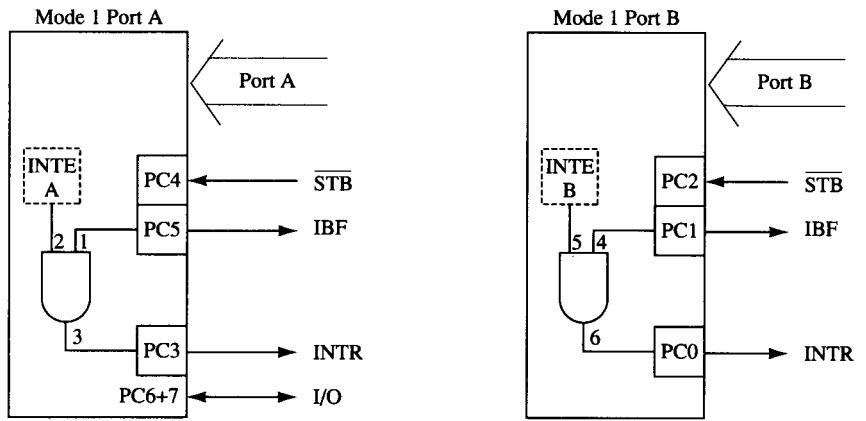
Mode 1 operation causes port A and/or port B to function as latching input devices. This allows external data to be stored into the port until the microprocessor is ready to retrieve it. Port C is also used in mode 1 operation—not for data, but for control or handshaking signals that help operate either or both port A and port B as strobed input ports. Figure 11–27 shows how both ports are structured for mode 1 strobed input operation and the timing diagram.

The strobed input port captures data from the port pins when the strobe ($\overline{STB}$) is activated. Note that the strobe captures the port data on the 0-to-1 transition. The $\overline{STB}$ signal causes data to be captured in the port, and it activates the IBF (**input buffer full**) and INTR (**interrupt request**) signals. Once the microprocessor, through software (IBF) or hardware (INTR), notices that data are strobed into the port, it executes an IN instruction to read the port $\overline{RD}$. The act of reading the port restores both IBF and INTR to their inactive states until the next datum is strobed into the port.
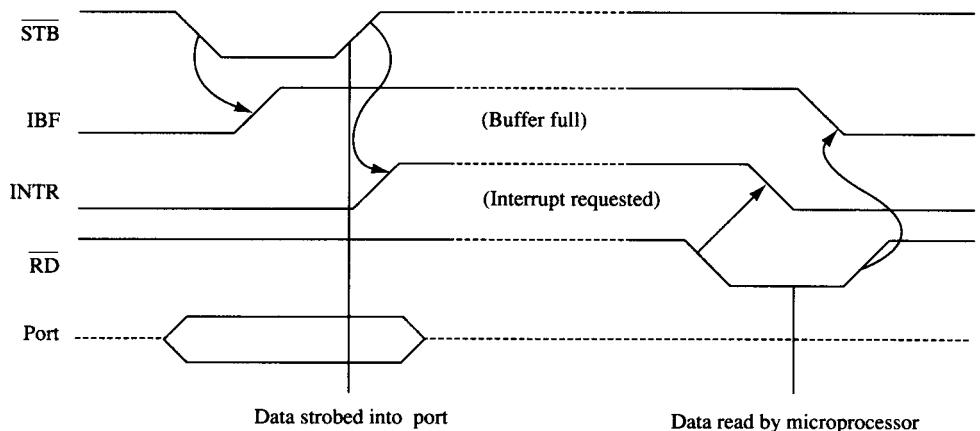
Signal Definitions for Mode 1 Strobed Input

$\overline{STB}$	The **strobe** input loads data into the port latch, which holds the information until it is input to the microprocessor via the IN instruction.
IBF	**Input buffer full** is an output indicating that the input latch contains information.
INTR	**Interrupt request** is an output that requests an interrupt. The INTR pin becomes a logic 1 when the $\overline{STB}$ input returns to a logic 1, and is cleared when the data are input from the port by the microprocessor.
INTE	The **interrupt** enable signal is neither an input nor an output; it is an internal bit programmed via the port PC_4 (port A) or PC_2 (port B) bit position.
PC_7, PC_6	The port C pins 7 and 6 are general-purpose I/O pins that are available for any purpose.

Strobed Input Example.   An excellent example of a strobed input device is a keyboard. The keyboard encoder debounces the key switches and provides a strobe signal whenever a key is

Mode 1 Port A

Mode 1 Port B

(a)

Data strobed into port

Data read by microprocessor
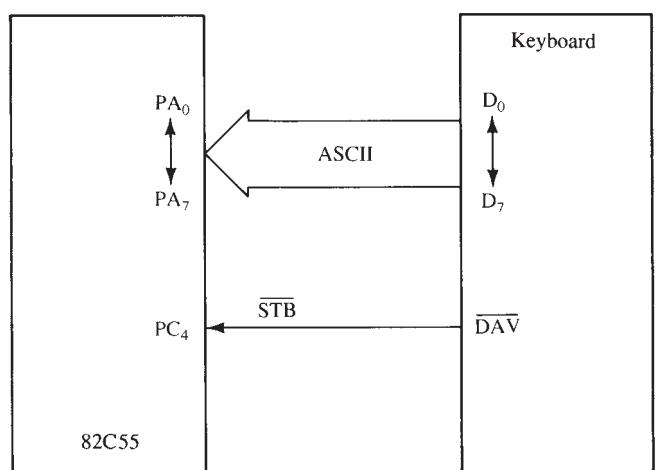
(b)

FIGURE 11–27 Strobed input operation (mode 1) of the 82C55. (a) Internal structure and (b) timing diagram.

FIGURE 11–28 Using the 82C55 for strobed input operation of a keyboard.

415

depressed and the data output contain the ASCII-coded key code. Figure 11–28 illustrates a keyboard connected to strobed input port A. Here $\overline{DAV}$ (data available) is activated for 1.0 μs each time that a key is typed on the keyboard. This causes data to be strobed into port A because $\overline{DAV}$ is connected to the $\overline{STB}$ input of port A. Each time a key is typed, therefore, it is stored into port A of the 82C55. The $\overline{STB}$ input also activates the IBF signal, indicating that data are in port A.

Example 11–17 shows a procedure that reads data from the keyboard each time a key is typed. This procedure reads the key from port A and returns with the ASCII code in AL. To detect a key, port C is read and the IBF bit (bit position PC_5) is tested to see whether the buffer is full. If the buffer is empty (IBF = 0), then the procedure keeps testing this bit, waiting for a character to be typed on the keyboard.

EXAMPLE 11–18

```
;A procedure that reads the keyboard encoder and
;returns the ASCII key code in AL

        BIT5  EQU   20H
        PORTC EQU   22H
        PORTA EQU   20H

        READ  PROC  NEAR

            .REPEAT                 ;poll IBF bit
                  IN    AL,PORTC
                  TEST AL,BIT5
            .UNTIL !ZERO?
            IN    AL.PORTA           ;get ASCII data
            RET

        READ  ENDP
```

Mode 1 Strobed Output

Figure 11–29 illustrates the internal configuration and timing diagram of the 82C55 when it is operated as a strobed output device under mode 1. Strobed output operation is similar to mode 0 output operation, except that control signals are included to provide handshaking.

Whenever data are written to a port programmed as a strobed output port, the $\overline{OBF}$ (**output buffer full**) signal becomes a logic 0 to indicate that data are present in the port latch. This signal indicates that data are available to an external I/O device that removes the data by strobing the $\overline{ACK}$ (acknowledge) input to the port. The $\overline{ACK}$ signal returns the $\overline{OBF}$ signal to a logic 1, indicating that the buffer is not full.

Signal Definitions for Mode 1 Strobed Output

$\overline{OBF}$	**Output buffer full** is an output that goes low whenever data are output (OUT) to the port A or port B latch. This signal is set to a logic 1 whenever the $\overline{ACK}$ pulse returns from the external device.
$\overline{ACK}$	The **acknowledge signal** causes the $\overline{OBF}$ pin to return to a logic 1 level. The $\overline{ACK}$ signal is a response from an external device, indicating that it has received the data from the 82C55 port.
INTR	**Interrupt request** is a signal that often interrupts the microprocessor when the external device receives the data via the $\overline{ACK}$ signal. This pin is qualified by the internal INTE (**interrupt enable**) bit.
INTE	**Interrupt enable** is neither an input nor an output; it is an internal bit programmed to enable or disable the INTR pin. The INTE A bit is programmed using the PC_6 bit and INTE B is programmed using the PC_2 bit.

(a)

(b)
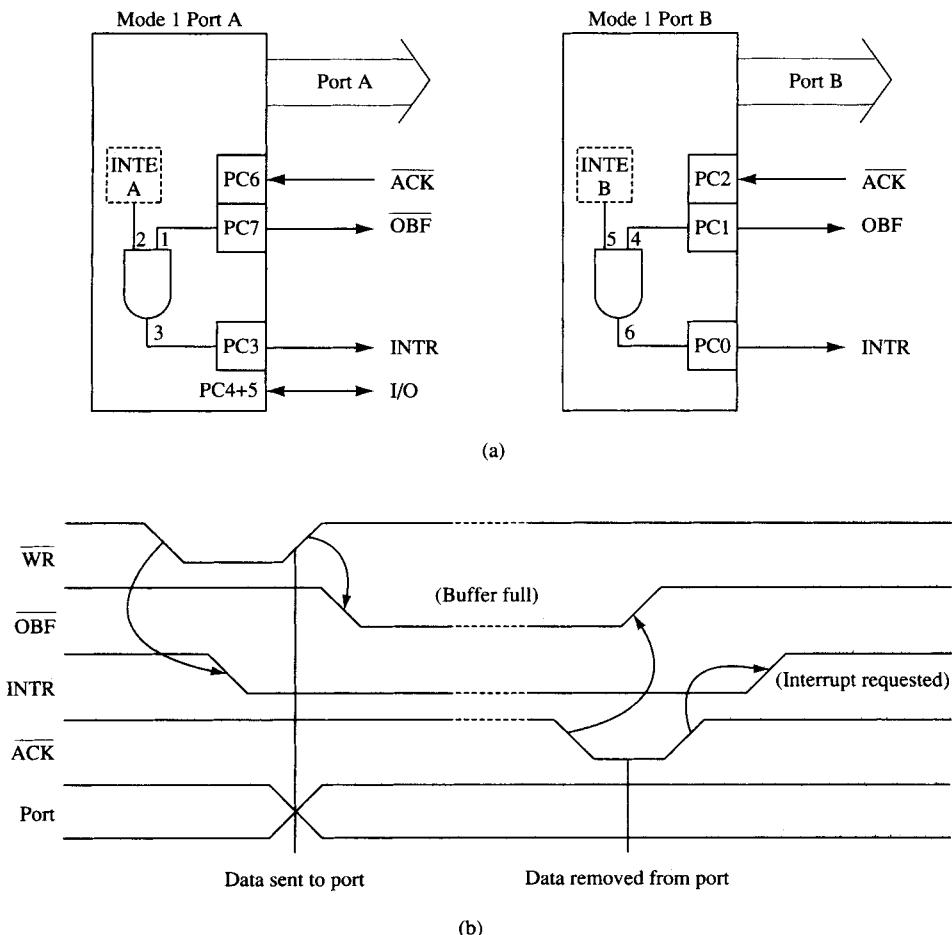
FIGURE 11–29 Strobed output operation (mode 1) of the 82C55. (a) Internal structure and (b) timing diagram.
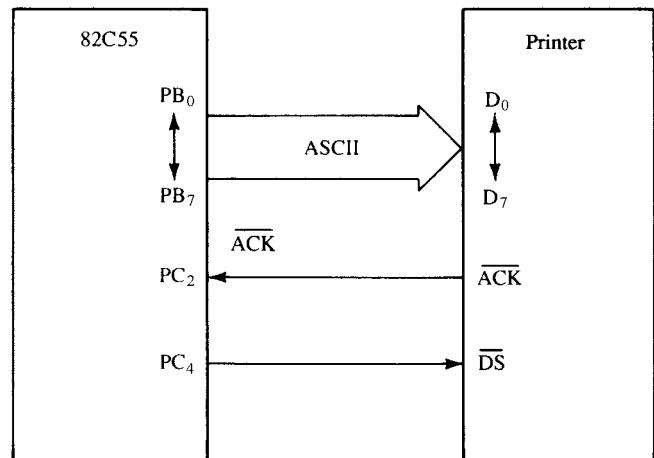
PC$_4$, PC$_5$ Port C pins PC$_4$ and PC$_5$ are general-purpose I/O pins. The bit set and reset command is used to set or reset these two pins.

Strobed Output Example. The printer interface discussed in Section 11–1 is used here to demonstrate how to achieve strobed output synchronization between the printer and the 82C55. Figure 11–30 illustrates port B connected to a parallel printer, with eight data inputs for receiving ASCII-coded data, a $\overline{DS}$ (**data strobe**) input to strobe data into the printer, and an $\overline{ACK}$ output to acknowledge the receipt of the ASCII character.

In this circuit, there is no signal to generate the $\overline{DS}$ signal to the printer, so PC$_4$ is used with software that generates the $\overline{DS}$ signal. The $\overline{ACK}$ signal that is returned from the printer acknowledges the receipt of the data and is connected to the $\overline{ACK}$ input of the 82C55.

Example 11–19 lists the software that sends the ASCII-coded character in AH to the printer. The procedure first tests $\overline{OBF}$ to decide whether the printer has removed the data from port B. If not, the procedure waits for the $\overline{ACK}$ signal to return from the printer. If $\overline{OBF} = 1$, then the procedure sends the contents of AH to the printer through port B and also sends the $\overline{DS}$ signal.

FIGURE 11–30 The 82C55 connected to a parallel printer interface that illustrates the strobed output mode of operation for the 82C55.

EXAMPLE 11–19

```
;A procedure that transfers an ASCII character from AH to the printer
;connected to port B

BIT1   EQU    2
PORTC  EQU    63H
PORTB  EQU    61H
CMD    EQU    63H

PRINT PROC   NEAR

      .REPEAT                  ;wait for printer ready
             IN   AL,PORTC
             TEST AL,BIT1
      .UNTIL !ZERO?
      MOV  AL,AH               ;send ASCII
      OUT  PORTB,AL
      MOV  AL,8                ;pulse data strobe
      OUT  CMD,AL
      MOV  AL,9
      OUT  CMD,AL
      RET

PRINT ENDP
```

Mode 2 Bidirectional Operation

In mode 2, which is allowed with group A only, port A becomes bidirectional, allowing data to be transmitted and received over the same eight wires. Bidirectional bused data are useful when interfacing two computers. It is also used for the IEEE-488 parallel high-speed GPIB (**general-purpose instrumentation bus**) interface standard. Figure 11–31 shows the internal structure and timing diagram for mode 2 bidirectional operation.

Signal Definitions for Bidirectional Mode 2

INTR   **Interrupt request** is an output used to interrupt the microprocessor for both input and output conditions.

$\overline{\text{OBF}}$ **Output buffer full** is an output indicating that the output buffer contains data for the bidirectional bus.

$\overline{\text{ACK}}$ **Acknowledge** is an input that enables the three-state buffers so that data can appear on port A. If $\overline{\text{ACK}}$ is a logic 1, the output buffers of port A are at their high-impedance state.

Bi-directional mode 2

(a)

(b)
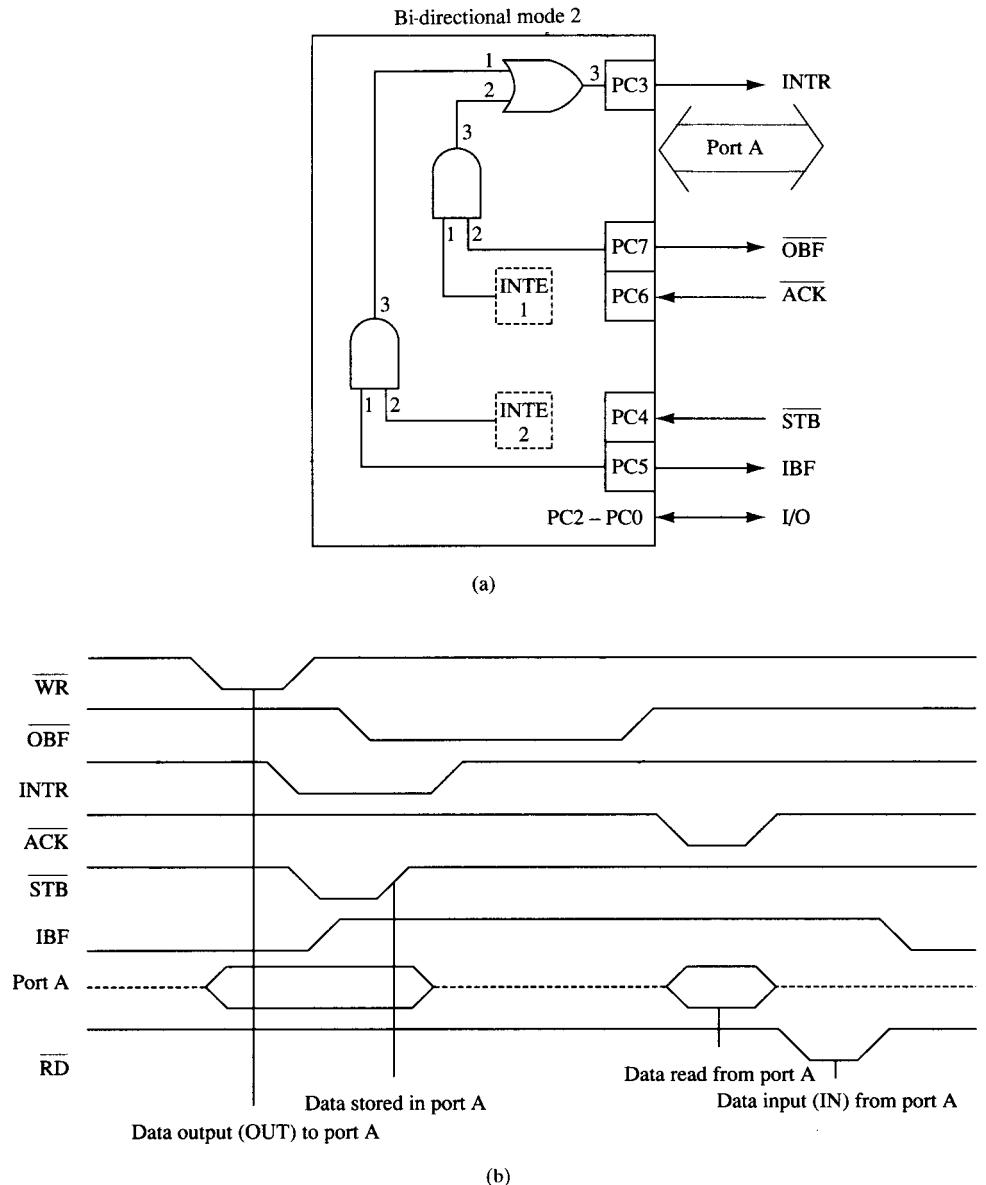
FIGURE 11–31   Mode 2 operation of the 82C55. (a) Internal structure and (b) timing diagram.

$\overline{\text{STB}}$ The **strobe** input loads the port A input latch with external data from the bidirectional port A bus.

IBF   **Input buffer full** is an output used to signal that the input buffer contains data for the external bidirectional bus.

INTE   **Interrupt enable** are internal bits (INTE1 and INTE2) that enable the INTR pin. The state of the INTR pin is controlled through port C bits PC_6 (INTE1) and PC_4 (INTE2).

PC_0, PC_1, and PC_2 These pins are general-purpose I/O pins in mode 2 controlled by the bit set and reset command.

The Bidirectional Bus.   The bidirectional bus is used by referencing port A with the IN and OUT instructions. To transmit data through the bidirectional bus, the program first tests the $\overline{\text{OBF}}$ signal

to determine whether the output buffer is empty. If it is, then data are sent to the output buffer via the OUT instruction. The external circuitry also monitors the $\overline{OBF}$ signal to decide whether the microprocessor has sent data to the bus. As soon as the output circuitry sees a logic 0 on $\overline{OBF}$, it sends back the $\overline{ACK}$ signal to remove it from the output buffer. The $\overline{ACK}$ signal sets the $\overline{OBF}$ bit and enables the three-state output buffers so that data may be read. Example 11–20 lists a procedure that transmits the contents of the AH register through bidirectional port A.

EXAMPLE 11–20

```
;A procedure transmits AH through the bidirectional bus

BIT7   EQU   80H
PORTC  EQU   62H
PORTA  EQU   60H

TRANS  PROC  NEAR

       .REPEAT                    ;test OBF
              IN   AL,PORTC
              TEST AL,BIT7
       .UNTIL !ZERO?
       MOV  AL,AH                  ;send data
       OUT  PORTA,AL
       RET

TRANS  ENDP
```

To receive data through the bidirectional port A bus, the IBF bit is tested with software to decide whether data have been strobed into the port. If IBF = 1, then data is input using the IN instruction. The external interface sends data into the port by using the $\overline{STB}$ signal. When $\overline{STB}$ is activated, the IBF signal becomes a logic 1 and the data at port A are held inside the port in a latch. When the IN instruction executes, the IBF bit is cleared and the data in the port are moved into AL. Example 11–21 lists a procedure that reads data from the port.

EXAMPLE 11–21

```
;A procedure that reads data from the bidirectional bus into AL

BIT5   EQU   20H
PORTC  EQU   62H
PORTA  EQU   60H

READ   PROC  NEAR

       .REPEAT                    ;test IBF
              IN   AL,PORTC
              TEST AL,BIT5
       .UNTIL !ZERO?
       IN   AL,PORTA
       RET

READ   ENDP
```

The INTR (**interrupt request**) pin can be activated from both directions of data flow through the bus. If INTR is enabled by both INTE bits, then the output and input buffers both cause interrupt requests. This occurs when data are strobed into the buffer using $\overline{STB}$ or when data are written using OUT.

82C55 Mode Summary

Figure 11–32 shows a graphical summary of the three modes of operation for the 82C55. Mode 0 provides simple I/O, mode 1 provides strobed I/O, and mode 2 provides bidirectional I/O. As mentioned, these modes are selected through the command register of the 82C55.

FIGURE 11–32 A summary of the port connections for the 82C55 PIA.

		Mode 0		Mode 1		Mode 2
Port A		IN	OUT	IN	OUT	I/O
Port B		IN	OUT	IN	OUT	Not used
Port C	0	IN	OUT	$INTR_B$	$INTR_B$	I/O
	1			IBF_B	$\overline{OBF_B}$	I/O
	2			$\overline{STB_B}$	$\overline{ACK_B}$	I/O
	3	IN	OUT	$INTR_A$	$INTR_A$	INTR
	4			$\overline{STB_A}$	I/O	$\overline{STB}$
	5			IBF_A	I/O	IBF
	6			I/O	$\overline{ACK_A}$	$\overline{ACK}$
	7			I/O	$\overline{OBF_A}$	$\overline{OBF}$

The Serial EEPROM Interface

In Chapter 10, Figure 10–23, a serial EEPROM is illustrated, but at that point in the text, no I/O existed for an interface. Suppose that port C of an 82C55 is used for connection to this interface and software is needed to drive the interface. It is assumed that pin PC_4 connects to the SCL input and pin PC_0 connects to the serial connection (SDA). PC_4 is programmed as an output pin to provide a clock signal. The PC_0 pin is programmed as an output to send data and as an input to receive data from the EEPROM.

Refer to Figure 10–24; the data format for the software for reading and writing data to the EEPROM is also illustrated in Example 11–22. This software is written in C with some assembly language, but it can also be developed in assembly language. The I/O port addresses are 0x1203 for the command register and 0x1202 for the port C register. The time delay should be 1.25 μs for a data rate of 400 KHz. Note that the time delay software is not illustrated here and the while loop is used to wait for an ACK signal after a write.

EXAMPLE 11–22

```
unsigned char void PC0in(unsigned char bit)
{
      _asm
      {
            mov dx,1203h
            mov al,81h
            out dx,al
            dec dx
            mov al,bit
            out dx,al
      }
      Delay();
      _asm
      {
            mov dx,1202h
            in al,dx               ;al is returned
      }
}

void PC0out(unsigned char bit)
{
      _asm
      {
            mov dx,1203h
            mov al,80h
            out dx,al
            dec dx
```

```
            mov al,bit
            out dx,al
        }
        Delay();
}

unsigned char void SendByte(unsigned char data)
{
        for (int a = 7; a >= 0; a—)
        {
                PC0out((data >> a) & 0xef);
                PC0out((data >> a) | 0x10);
        }
        PC0in(0xef);            //ack bit
        return PC0in(0x10);
}

unsigned char GetByte()
{
        unsigned char temp = 0;
        for (int a = 7; a >= 0; a—)
        {
                PC0in(0xef);
                temp |= PC0in(0x10) << a;
        }
        PC0in(0xef);            //ack bit
        PC0in(0x10);
        return temp;
}

void SendStart()
{       // start is one
        PC0out(0xef);
        PC0out(0x10);
}

void SendStop()
{       // stop is zero
        PC0out(0xee);
        PC0out(0x10);
}

void SendData(char device, short address, unsigned char data)
{
        Char c = 0;
        SendStart();
        SendByte(0xa0 | device << 1);
        SendByte(address >> 8);
        SendByte(address);
        SendByte(data);
        while (c == 0)
        {       // wait for ACK = 1;
                C = SendByte(0xa0 | device << 1);
        }
        SendStop();
}

unsigned char ReadData(char device, short address)
{
        SendStart();
        SendByte(0xa0 | device << 1);
        SendByte(address >> 8);
        SendByte(address);
        SendByte(0xa1 | device << 1);
        unsigned char temp = GetByte();
        SendStop();
        return temp;
}
```

11–4   8254 PROGRAMMABLE INTERVAL TIMER

The 8254 programmable interval timer consists of three independent 16-bit programmable counters (**timers**). Each counter is capable of counting in binary or binary-coded decimal (BCD). The maximum allowable input frequency to any counter is 10 MHz. This device is useful wherever the microprocessor must control real-time events. Some examples of usage include real-time clock and an events counter, and for motor speed and direction control.

This timer also appears in the personal computer decoded at ports 40H–43H to do the following:

1. Generate a basic timer interrupt that occurs at approximately 18.2 Hz.
2. Cause the DRAM memory system to be refreshed.
3. Provide a timing source to the internal speaker and other devices. The timer in the personal computer is an 8253 instead of an 8254.

8254 Functional Description

Figure 11–33 shows the pin-out of the 8254, which is a higher-speed version of the 8253, and a diagram of one of the three counters. Each timer contains a CLK input, a gate input, and an output (OUT) connection. The CLK input provides the basic operating frequency to the timer, the gate pin controls the timer in some modes, and the OUT pin is where we obtain the output of the timer.

The signals that connect to the microprocessor are the data bus pins (D_7–D_0), $\overline{RD}$, $\overline{WR}$, $\overline{CS}$, and address inputs A_1 and A_0. The address inputs are present to select any of the four internal registers used for programming, reading, or writing to a counter. The personal computer contains an 8253 timer or its equivalent, decoded at I/O ports 40H–43H. Timer zero is programmed to generate an 18.2 Hz signal that interrupts the microprocessor at interrupt vector 8 for a clock tick. The tick is often used to time programs and events in DOS. Timer 1 is programmed for 15 μs, which is used on the personal computer to request a DMA action used to refresh the dynamic RAM. Timer 2 is programmed to generate a tone on the personal computer speaker.
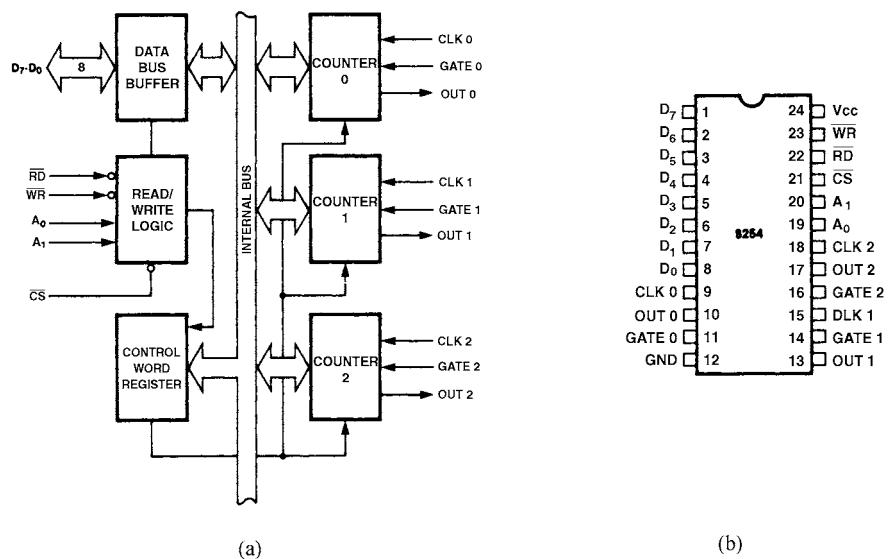
FIGURE 11–33 The 8254 programmable interval timer. (a) Internal structure and (b) pin-out. (Courtesy of Intel Corporation.)

TABLE 11–4   Address selection inputs to the 8254.

A_1	A_0	Function
0	0	Counter 0
0	1	Counter 1
1	0	Counter 2
1	1	Control word
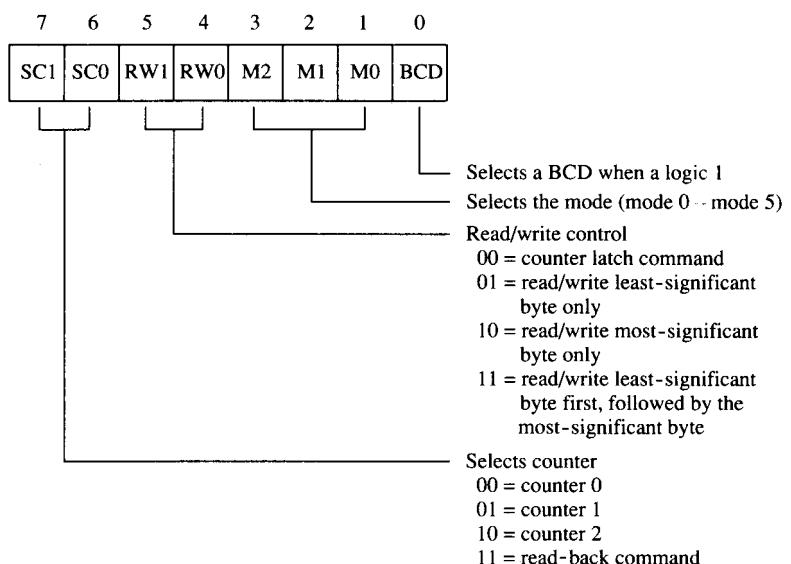
Pin Definitions

A_0, A_1	The **address inputs** select one of four internal registers within the 8254. See Table 11–4 for the function of the A_1 and A_0 address bits.
CLK	The **clock** input is the timing source for each of the internal counters. This input is often connected to the PCLK signal from the microprocessor system bus controller.
$\overline{CS}$	**Chip select** enables the 8254 for programming and reading or writing a counter.
G	The **gate input** controls the operation of the counter in some modes of operation.
GND	**Ground** connects to the system ground bus.
OUT	A **counter output** is where the waveform generated by the timer is available.
$\overline{RD}$	**Read** causes data to be read from the 8254 and often connects to the $\overline{IORC}$ signal.
Vcc	**Power** connects to the +5.0 V power supply.
$\overline{WR}$	**Write** causes data to be written to the 8254 and often connects to the write strobe ($\overline{IOWC}$).

Programming the 8254

Each counter is individually programmed by writing a control word, followed by the initial count. Figure 11–34 lists the program control word structure of the 8254. The control word allows the programmer to select the counter, mode of operation, and type of operation (read/write). The control word also selects either a binary or BCD count. Each counter may be programmed with a count of 1 to FFFFH. A count of 0 is equal to FFFFH+l (65,536) or 10,000

FIGURE 11–34   The control word for the 8254-2 timer.

7	6	5	4	3	2	1	0
SC1	SC0	RW1	RW0	M2	M1	M0	BCD

Selects a BCD when a logic 1

Selects the mode (mode 0 – mode 5)

Read/write control
00 = counter latch command
01 = read/write least-significant
     byte only
10 = read/write most-significant
     byte only
11 = read/write least-significant
     byte first, followed by the
     most-significant byte

Selects counter
00 = counter 0
01 = counter 1
10 = counter 2
11 = read-back command

in BCD. The minimum count of 1 applies to all modes of operation except modes 2 and 3, which have a minimum count of 2. Timer 0 is used in the personal computer with a divide-by count of 64K (FFFFH) to generate the 18.2 Hz (18.196 Hz) interrupt clock tick. Timer 0 has a clock input frequency of 4.77 MHz + 4 or 1.1925 MHz.

The control word uses the BCD bit to select a BCD count (BCD = 1) or a binary count (BCD = 0). The M_2, M_1, and M_0 bits select one of the six different modes of operation (000–101) for the counter. The RW_1 and RW_0 bits determine how the data are read from or written to the counter. The SC_1 and SC_0 bits select a counter or the special read-back mode of operation, discussed later in this section.

Each counter has a program control word used to select the way the counter operates. If two bytes are programmed into a counter, then the first byte (LSB) will stop the count, and the second byte (MSB) will start the counter with the new count. The order of programming is important for each counter, but programming of different counters may be interleaved for better control. For example, the control word may be sent to each counter before the counts for individual programming. Example 11–23 shows a few ways to program counters 1 and 2. The first method programs both control words, then the LSB of the count for each counter, which stops them from counting. Finally, the MSB portion of the count is programmed, starting both counters with the new count. The second example shows one counter programmed before the other.

EXAMPLE 11–23

```
PROGRAM CONTROL WORD 1 PROGRAM CONTROL WORD 2 PROGRAM LSB 1
PROGRAM LSB 2
PROGRAM MSB 1
PROGRAM MSB 2

;setup counter 1
;setup counter 2
;stop counter 1 and program LSB
;stop counter 2 and program LSB ;program MSB of counter 1 and start it
;program MSB of counter 2 and start it

       or

PROGRAM CONTROL WORD 1 PROGRAM LSB 1
PROGRAM MSB 1
PROGRAM CONTROL WORD 2 PROGRAM LSB 2
PROGRAM MSB 2

;setup counter 1
;stop counter 1 and program LSB ;program MSB of counter 1 and start it
;setup counter 2
;stop counter 2 and program LSB ;program MSB of counter 2 and start it
```

Modes of Operation.   Six modes (mode 0–mode 5) of operation are available to each of the 8254 counters. Figure 11–35 shows how each of these modes functions with the CLK input, the gate (G) control signal, and OUT signal. A description of each mode follows:

MODE 0	Allows the 8254 counter to be used as an events counter. In this mode, the output becomes a logic 0 when the control word is written and remains there until N plus the number of programmed counts. For example, if a count of 5 is programmed, the output will remain a logic 0 for 6 counts beginning with N. Note that the gate (G) input must be a logic 1 to allow the counter to count. If G becomes a logic 0 in the middle of the count, the counter will stop until G again becomes a logic 1.
MODE 1	Causes the counter to function as a retriggerable, monostable multivibrator (one-shot). In this mode the G input triggers the counter so that it develops a pulse at the OUT connection that becomes a logic 0 for the duration of the
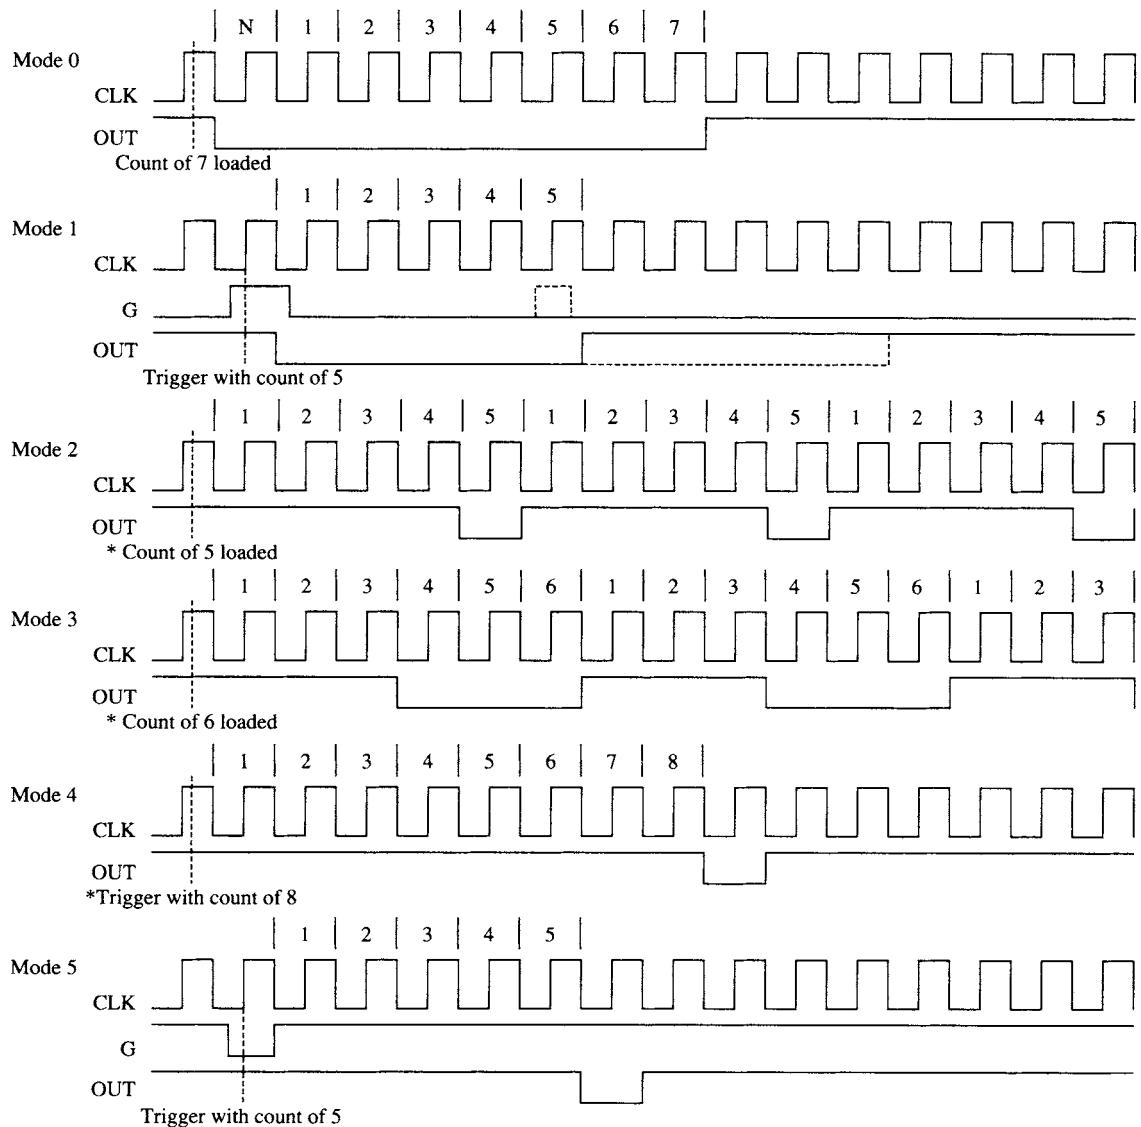
FIGURE 11–35 The six modes of operation for the 8254-2 programmable interval timer. The G input stops the count when 0 in modes 2, 3, and 4.

count. If the count is 10, then the OUT connection goes low for 10 clocking periods when triggered. If the G input occurs within the duration of the output pulse, the counter is again reloaded with the count and the OUT connection continues for the total length of the count.

MODE 2　Allows the counter to generate a series of continuous pulses that are one clock pulse wide. The separation between pulses is determined by the count. For example, for a count of 10, the output is a logic 1 for nine clock periods and low for one clock period. This cycle is repeated until the counter is programmed with a new count or until the G pin is placed at a logic 0 level. The G input must be a logic 1 for this mode to generate a continuous series of pulses.

MODE 3   Generates a continuous square wave at the OUT connection, provided that the G pin is a logic 1. If the count is even, the output is high for one half of the count and low for one half of the count. If the count is odd, the output is high for one clocking period longer than it is low. For example, if the counter is programmed for a count of 5, the output is high for three clocks and low for two clocks.

MODE 4   Allows the counter to produce a single pulse at the output. If the count is programmed as a 10, the output is high for 10 clocking periods and low for one clocking period. The cycle does not begin until the counter is loaded with its complete count. This mode operates as a software triggered one-shot. As with modes 2 and 3, this mode also uses the G input to enable the counter. The G input must be a logic 1 for the counter to operate for these three modes.

MODE 5   A hardware triggered one-shot that functions as mode 4, except that it is started by a trigger pulse on the G pin instead of by software. This mode is also similar to mode 1 because it is retriggerable.

Generating a Waveform with the 8254.   Figure 11–36 shows an 8254 connected to function at I/O ports 0700H, 0702H, 0704H, and 0706H of an 80386SX microprocessor. The addresses are decoded by using a PLD that also generates a write strobe signal for the 8254, which is connected to the low-order data bus connections. The PLD also generates a wait signal for the microprocessor that causes two wait states when the 8254 is accessed. The wait state generator connected to the microprocessor actually controls the number of wait states inserted into the timing. The program for the PLD is not illustrated here because it is the same as many of the prior examples.
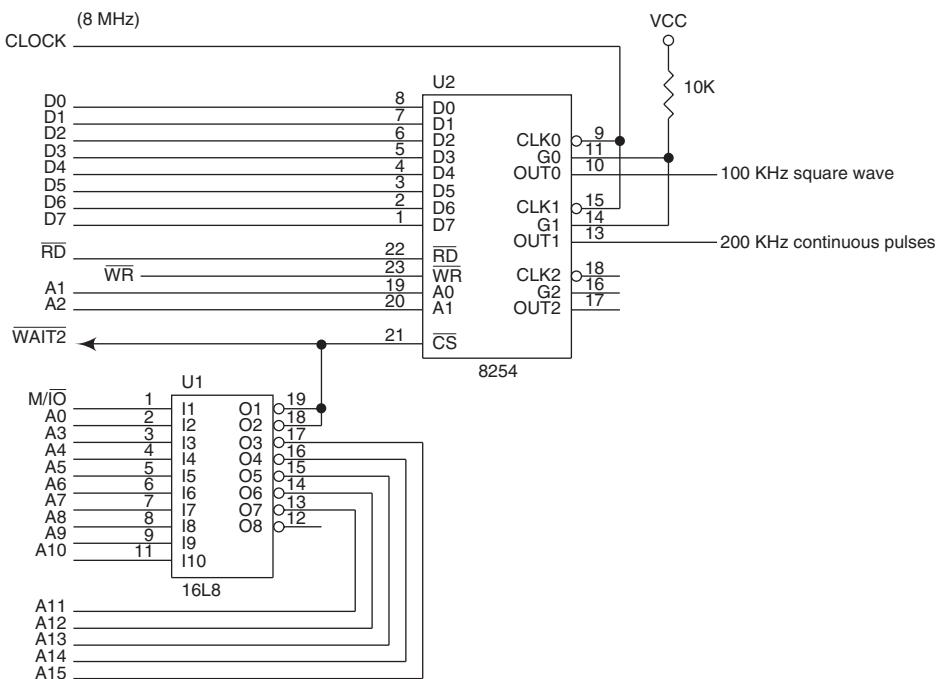
FIGURE 11–36   The 8254 interfaced to an 8 MHz 8086 so that it generates a 100 KHz square wave at OUT0 and a 200 KHz continuous pulse at OUT1.

Example 11–24 lists the program that generates a 100 KHz square-wave at OUT0 and a 200 KHz continuous pulse at OUT1. Counter 0 uses mode 3 and counter 1 uses mode 2. The count programmed into counter 0 is 80 and the count for counter 1 is 40. These counts generate the desired output frequencies with an 8 MHz input clock.

EXAMPLE 11–24

```
;A procedure that programs the 8254 timer to function
;as illustrated in Figure 11-36

TIME   PROC   NEAR USES AX DX

       MOV  DX,706H                    ;program counter 0 for mode 3
       MOV  AL,00110110B
       OUT  DX,AL
       MOV  AL,01110100B               ;program counter 1 for mode 2
       OUT  DX,AL

       MOV  DX,700H                    ;program counter 0 with 80
       MOV  AL,80
       OUT  DX,AL
       MOV  AL,0
       OUT  DX,AL

       MOV  DX,702H                    ;program counter 1 with 40
       MOV  AL,40
       OUT  DX,AL
       MOV  AL,0
       OUT  DX,AL

       RET

TIME   ENDP
```

Reading a Counter.   Each counter has an internal latch that is read with the read counter port operation. These latches will normally follow the count. If the contents of the counter are needed, then the latch can remember the count by programming the counter latch control word (see Figure 11–37), which causes the contents of the counter to be held in a latch until they is read. Whenever a read from the latch or the counter is programmed, the latch tracks the contents of the counter.

When it is necessary for the contents of more than one counter to be read at the same time, we use the read-back control word, illustrated in Figure 11–38. With the read-back control word, the $\overline{CNT}$ bit is a logic 0 to cause the counters selected by CNT0, CNT1, and CNT2 to be latched. If the status register is to be latched, then the $\overline{ST}$ bit is placed at a logic 0. Figure 11–39 shows the status register, which shows the state of the output pin, whether the counter is at its null state (0), and how the counter is programmed.

FIGURE 11–37 The 8254-2 counter latch control word.

Select counter
00 = counter 0
01 = counter 1
10 = counter 2
11 = read-back command
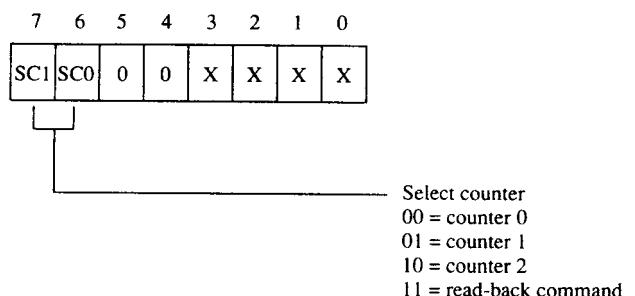
FIGURE 11–38   The 8254-2
read-back control word.

7	6	5	4	3	2	1	0
1	1	$\overline{CNT}$	$\overline{ST}$	CNT2	CNT1	CNT0	0

Select counter bits

Latch status of
selected counters

Latch count of
selected counters
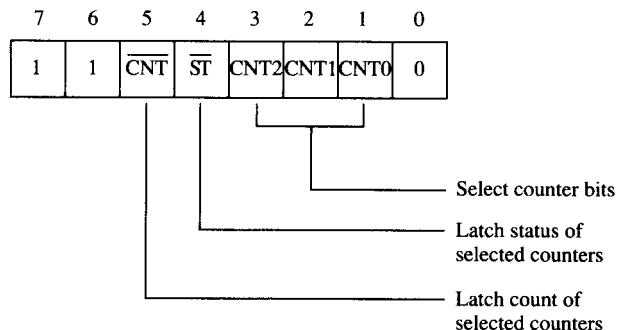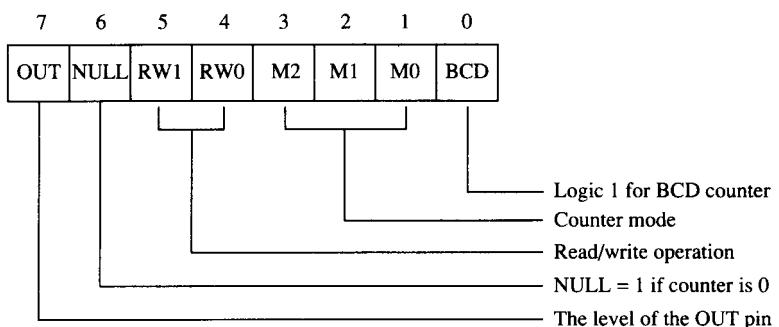
FIGURE 11–39   The 8254-2
status register.

7	6	5	4	3	2	1	0
OUT	NULL	RW1	RW0	M2	M1	M0	BCD

Logic 1 for BCD counter
Counter mode
Read/write operation
NULL = 1 if counter is 0
The level of the OUT pin

DC Motor Speed and Direction Control

One application of the 8254 timer is as a motor speed controller for a DC motor. Figure 11–40 shows the schematic diagram of the motor and its associated driver circuitry. It also illustrates the interconnection of the 8254, a flip-flop, and the motor and its driver.

The operation of the motor driver circuitry is straightforward. If the Q output of the 74ALS112 is a logic 1, the base Q_2 is pulled up to +12 V through the base pull-up resistor, and the base of Q_2 is open circuited. This means that Q_1 is off and Q_2 is on, with ground applied to the positive lead of the motor. The bases of both Q_3 and Q_4 are pulled low to ground through the inverters. This causes Q_3 to conduct or turn on and Q_4 to turn off, applying ground to the negative lead of the motor. The logic 1 at the Q output of the flip-flop therefore connects +12 V to the positive lead of the motor and ground to the negative lead. This connection causes the motor to spin in its forward direction. If the state of the Q output of the flip-flop becomes a logic 0, then the conditions of the transistors are reversed and +12 V is attached to the negative lead of the motor, with ground attached to the positive lead. This causes the motor to spin in the reverse direction.

If the output of the flip-flop is alternated between a logic 1 and 0, the motor spins in either direction at various speeds. If the duty cycle of the Q output is 50%, the motor will not spin at all and exhibits some holding torque because current flows through it. Figure 11–41 shows some timing diagrams and their effects on the speed and direction of the motor. Notice how each counter generates pulses at different positions to vary the duty cycle at the Q output of the flip-flop. This output is also called *pulse width modulation*

To generate these wave forms, counters 0 and 1 are both programmed to divide the input clock (PCLK) by 30,720. We change the duty cycle of Q by changing the point at which
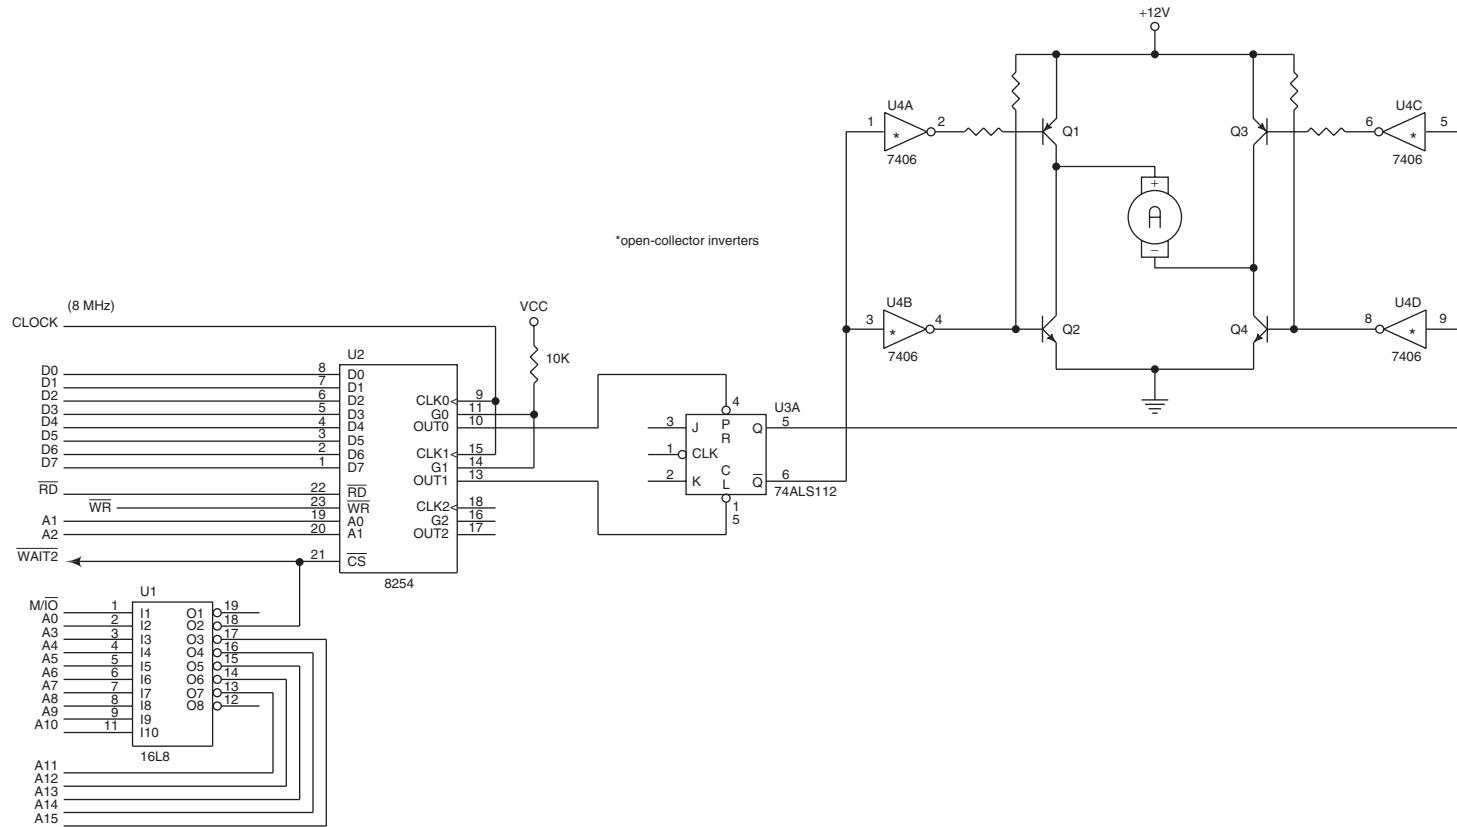
FIGURE 11–40 Motor speed and direction control using the 8254 timer.
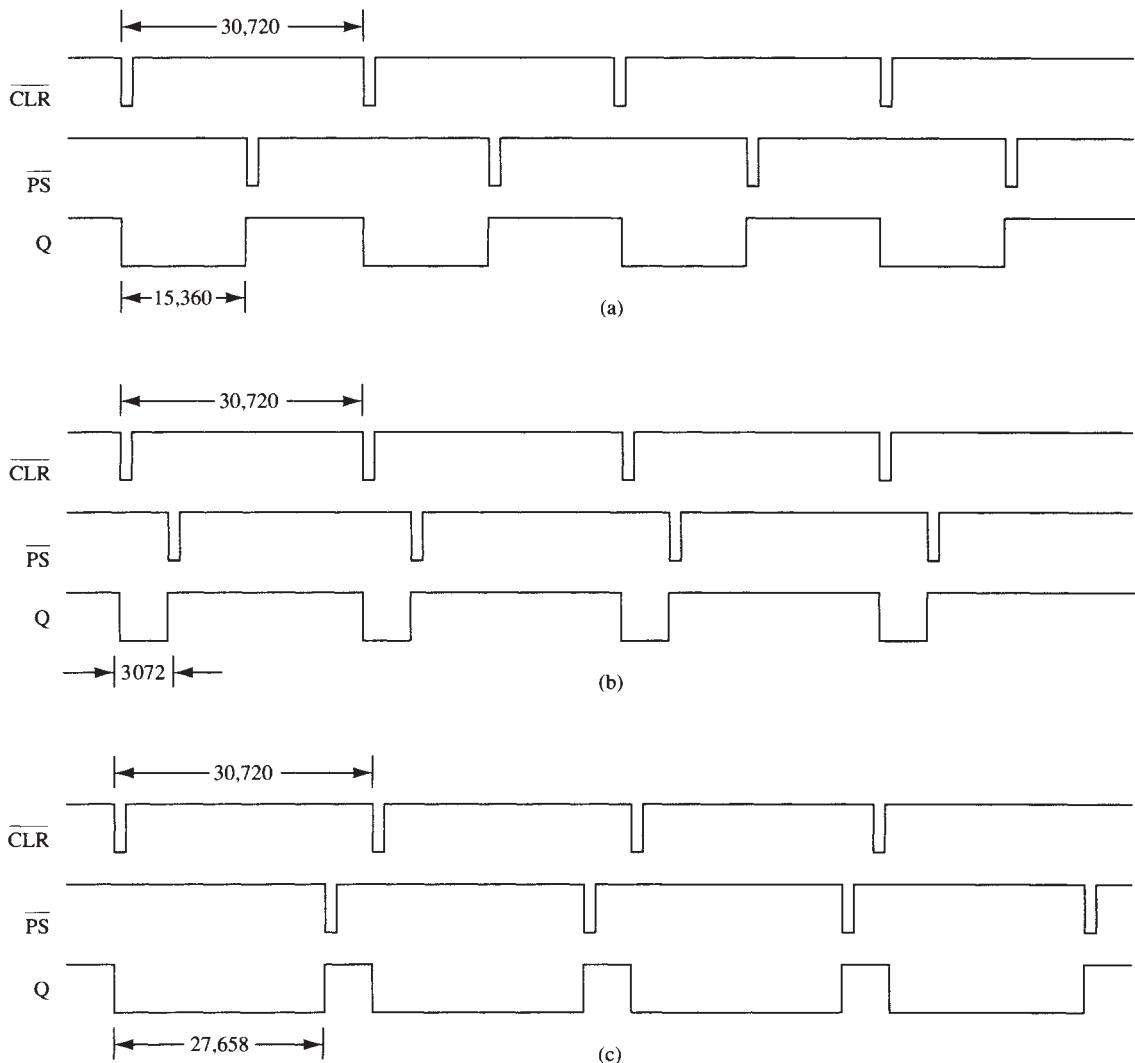
FIGURE 11–41   Timing for the motor speed and direction control circuit of Figure 11–40. (a) No rotation, (b) high-speed rotation in the reverse direction, and (c) high-speed rotation in the forward direction.

counter 1 is started in relationship to counter 0. This changes the direction and speed of the motor. But why divide the 8 MHz clock by 30,720? The divide rate of 30,720 is divisible by 256, so we can develop a short program that allows 256 different speeds. This also produces a basic operating frequency for the motor of about 260 Hz, which is low enough in frequency to power the motor. It is important to keep this operating frequency below 1000 Hz, but above 60 Hz.

   Example 11–25 lists a procedure that controls the speed and direction of the motor. The speed is controlled by the value of AH when this procedure is called. Because we have an 8-bit number to represent speed, a 50% duty cycle, for a stopped motor, is a count of 128. By changing the value in AH when the procedure is called, we can adjust the motor speed. The speed of the motor will increase in either direction by changing the number in AH when this procedure is called. As the value in AH approaches 00H, the motor begins to increase its speed in the reverse

direction. As the value of AH approaches FFH, the motor increases its speed in the forward direction.

EXAMPLE 11–25

```
;A procedure that controls the speed and direction of the motor
;in Figure 11-40.
;
;AH determines the speed and direction of the motor where
;AH is between 00H and FFH.

CNTR   EQU   706H
CNT0   EQU   700H
CNT1   EQU   702H
COUNT  EQU   30720

SPEED  PROC  NEAR USES BX DX AX

       MOV   BL,AH                ;calculate count
       MOV   AX,120
       MUL   BL
       MOV   BX,AX
       MOV   AX,COUNT
       SUB   AX,BX
       MOV   BX,AX

       MOV   DX,CNTR
       MOV   AL,00110100B         ;program control words
       OUT   DX,AL
       MOV   AL,01110100B
       OUT   DX,AL

       MOV   DX,CNT1              ;program counter 1
       MOV   AX,COUNT             ;to generate a clear
       OUT   DX,AL
       MOV   AL,AH
       OUT   DX,AL

       .REPEAT                    ;wait for counter 1
             IN    AL,DX
             XCHG AL,AH
             IN    AL,DX
             XCHG AL,AH
       .UNTIL BX == AX

       MOV   DX,CNT0              ;program counter 0
       MOV   AX,COUNT             ;to generate a set
       OUT   DX,AL
       MOV   AL,AH
       OUT   DX,AL

       RET

SPEED  ENDP
```

The procedure adjusts the wave form at Q by first calculating the count at which counter 0 is to start in relationship to counter 1. This is accomplished by multiplying AH by 120 and then subtracting it from 30,720. This is required because the counters are down-counters that count from the programmed count to 0 before restarting. Next, counter 1 is programmed with a count of 30,720 and started so it generates the clear-wave form for the flip-flop. After counter 1 is started, it is read and compared with the calculated count. Once it reaches this count, counter 0 is started with a count of 30,720. From this point forward, both counters continue generating the clear and set wave forms until the procedure is again called to adjust the speed and direction of the motor.

11–5   16550 PROGRAMMABLE COMMUNICATIONS INTERFACE

The National Semiconductor Corporation's PC16550D is a programmable communications interface designed to connect to virtually any type of serial interface. The 16550 is a universal asynchronous receiver/transmitter (UART) that is fully compatible with the Intel microprocessors. The 16550 is capable of operating at 0–1.5 M baud. Baud rate is the number of bits transferred per second (bps), including start, stop, data, and parity (Bps is bytes per second and bps are bits per second). The 16550 also includes a programmable baud rate generator and separate FIFOs for input and output data to ease the load on the microprocessor. Each FIFO contains 16 bytes of storage. This is the most common communications interface found in modem microprocessor-based equipment, including the personal computer and many modems.

Asynchronous Serial Data

Asynchronous serial data are transmitted and received without a clock or timing signal. Figure 11–42 illustrates two frames of asynchronous serial data. Each frame contains a start bit, seven data bits, parity, and one stop bit. The figure shows a frame that contains one ASCII character and 10 bits. Most dial-up communications systems of the past, such as CompuServe, Prodigy, and America Online, used 10 bits for asynchronous serial data with even parity. Most Internet and bulletin board services also use 10 bits, but they normally do not use parity. Instead, eight data bits are transferred, replacing parity with a data bit. This makes byte transfers of non-ASCII data much easier to accomplish.

16550 Functional Description

Figure 11–43 illustrates the pin-out of the 16550 UART. This device is available as a 40-pin DIP (**dual in-line package**) or as a 44-pin PLCC (**plastic leadless chip carrier**). Two completely separate sections are responsible for data communications: the receiver and the transmitter. Because each of these sections is independent, the 16550 is able to function in simplex, half-duplex, or full-duplex modes. One of the main features of the 16550 is its internal receiver and transmitter FIFO (first-in, first-out) memories. Because each is 16 bytes deep, the UART requires attention only from the microprocessor after receiving 16 bytes of data. It also holds 16 bytes before the microprocessor must wait for the transmitter. The FIFO makes this UART ideal when interfacing to high-speed systems because less time is required to service it.

An example **simplex** system is one in which the transmitter or receiver is used by itself such as in an FM (**frequency modulation**) radio station. An example **half-duplex** system is a CB (**citizens band**) radio, on which we transmit and receive, but not both at the same time. The **full-duplex** system allows transmission and reception in both directions simultaneously. An example of a full-duplex system is the telephone.

The 16550 can control a **modem** (**modulator/demodulator**), which is a device that converts TTL levels of serial data into audio tones that can pass through the telephone system. Six pins on

* = STOP
ST = START
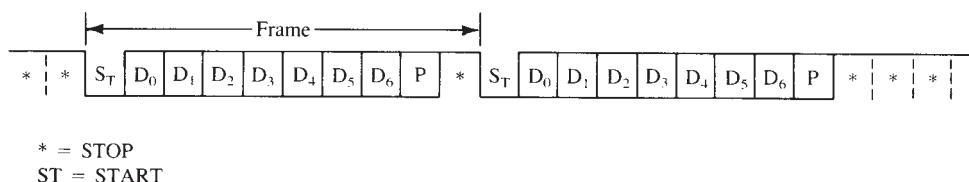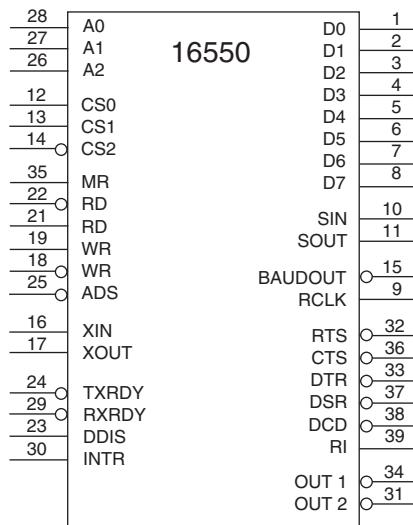
FIGURE 11–42   Asynchronous serial data.

FIGURE 11–43 The pin-out
of the 16550 UART.

the 16650 are devoted to modem control: $\overline{\text{DSR}}$ (**data set ready**), $\overline{\text{DTR}}$ (**data terminal ready**), $\overline{\text{CTS}}$ (**clear-to-send**), $\overline{\text{RTS}}$ (**request-to-send**), $\overline{\text{RI}}$ (**ring indicator**), and $\overline{\text{DCD}}$ (**data carrier detect**). The modem is referred to as the data set and the 16550 is referred to as the data terminal.

16550 Pin Functions

A_0, A_1, A_2	The **address inputs** are used to select an internal register for programming and also data transfer. See Table 11–5 for a list of each combination of the address inputs and the registers selected.
$\overline{\text{ADS}}$	The **address strobe** input is used to latch the address lines and chip select lines. If not needed (as in the Intel system), connect this pin to ground. The $\overline{\text{ADS}}$ pin is designed for use with Motorola microprocessors.
$\overline{\text{BAUDOUT}}$	The **baud out** pin is where the clock signal generated by the baud rate generator from the transmitter section is made available. It is most often connected to the RCLK input to generate a receiver clock that is equal to the transmitter clock.
$CS_0, CS_1, \overline{CS_2}$	The **chip select** inputs must all be active to enable the 16550 UART.
$\overline{\text{CTS}}$	The **clear-to-send** (if low) indicates that the modem or data set is ready to exchange information. This pin is often used in a half-duplex system to turn the line around.

TABLE 11–5 The register
selected by A_0, A_1, and A_2.

A_2	A_1	A_0	Function
0	0	0	Receiver buffer (read) and transmitter holding (write)
0	0	1	Interrupt enable
0	1	0	Interrupt identification (read) and FIFO control (write)
0	1	1	Line control
1	0	0	Modem control
1	0	1	Line status
1	1	0	Modem status
1	1	1	Scratch
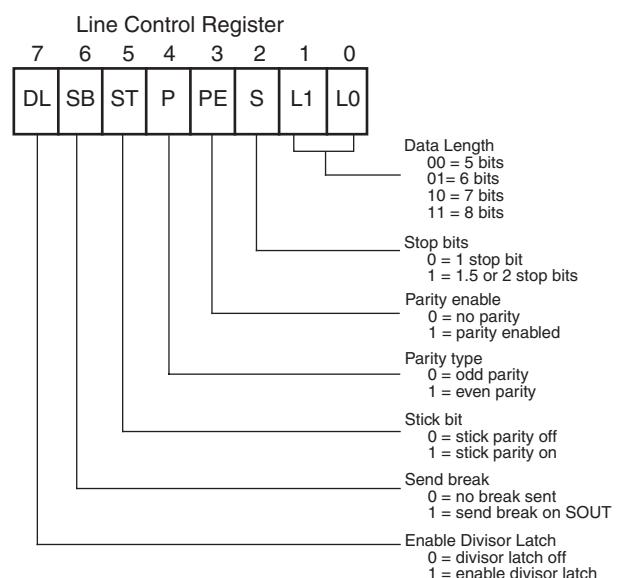
D_0–D_7	The **data bus** pins are connected to the microprocessor data bus.
$\overline{DCD}$	The **data carrier detect** input is used by the modem to signal the 16550 that a carrier is present.
DDIS	The **disable driver** output becomes a logic 0 to indicate that the microprocessor is reading data from the UART. DDIS can be used to change the direction of data flow through a buffer.
$\overline{DSR}$	**Data set ready** is an input to the 16550, indicating that the modem or data set is ready to operate.
$\overline{DTR}$	**Data terminal ready** is an output that indicates that the data terminal (16550) is ready to function.
INTR	**Interrupt request** is an output to the microprocessor used to request an interrupt (INTR = 1) whenever the 16550 has a receiver error, it has received data, and the transmitter is empty.
MR	**Master reset** initializes the 16550 and should be connected to the system RESET signal.
$\overline{OUT1}$, $\overline{OUT2}$	**User-defined output pins** that can provide signals to a modem or any other device as needed in a system.
RCLK	**Receiver clock** is the clock input to the receiver section of the UART. This input is always 16 times the desired receiver baud rate.
RD, $\overline{RD}$	**Read inputs** (either may be used) cause data to be read from the register specified by the address inputs to the UART.
$\overline{RI}$	The **ring indicator** input is placed at the logic 0 level by the modem to indicate that the telephone is ringing.
$\overline{RTS}$	**Request-to-send** is a signal to the modem indicating that the UART wishes to send data.
SIN, SOUT	These are the **serial data pins**. SIN accepts serial data and SOUT transmits serial data.
$\overline{RXRDY}$	**Receiver ready** is a signal used to transfer received data via DMA techniques (see text).
$\overline{TXRDY}$	**Transmitter ready** is a signal used to transfer transmitter data via DMA techniques (see text).
WR, $\overline{WR}$	**Write** (either may be used) connects to the microprocessor write signal to transfer commands and data to the 16550.
XIN, XOUT	These are the main **clock** connections. A crystal is connected across these pins to form a crystal oscillator, or XIN is connected to an external timing source.

Programming the 16550

Programming the 16550 is simple, although it may be slightly more involved when compared to some of the other programmable interfaces described in this chapter. Programming is a two-part process that includes the initialization dialog and operational dialog.

In the personal computer, which uses the 16550 or its programming equivalent, the I/O port addresses are decoded at 3F8H through 3FFH for COM port 0 and 2F8H through 2FFH for COM port 2. Although the examples in this section of the chapter are not written specifically for the personal computer, they can be adapted by changing the port numbers to control the COM ports on the PC.

FIGURE 11–44 The contents of the 16550 line control register.

Initializing the 16550.   Initialization dialog, which occurs after a hardware or software reset, consists of two parts: programming the line control register and the baud rate generator. The line control register selects the number of data bits, number of stop bits, and parity (whether it's even or odd, or if parity is sent as a 1 or a 0). The baud rate generator is programmed with a divisor that determines the baud rate of the transmitter section.

Figure 11–44 illustrates the line control register. The line control register is programmed by outputting information to I/O port 011 (A_2, A_1, A_0). The rightmost two bits of the line control register select the number of transmitted data bits (5, 6, 7, or 8). The number of stop bits is selected by S in the line control register. If S = 0, one stop bit is used; if S = 1, 1.5 stop bits are used for five data bits, and two stop bits are used with six, seven, or eight data bits.

The next three bits are used together to send even or odd parity, to send no parity, or to send a 1 or a 0 in the parity bit position. To send even or odd parity, the ST (**stick**) bit must be placed at a logic 0 level, and parity enable must be a logic 1. The value of the parity bit then determines even or odd parity. To send no parity (common in Internet connections), ST = 0 as well as the parity enable bit. This sends and receives data without parity. Finally, if a 1 or a 0 must be sent and received in the parity bit position for all data, ST = 1 with a 1 in parity enable. To send a 1 in the parity bit position, place a 0 in the parity bit; to send a 0, place a 1 in the parity bit. (See Table 11–6 for the operation of the parity and stick bits.)

TABLE 11–6 The operation of the ST and parity bits.

ST	P	PE	Function
0	0	0	No parity
0	0	1	Odd parity
0	1	0	No parity
0	1	1	Even parity
1	0	0	Undefined
1	0	1	Send/receive 1
1	1	0	Undefined
1	1	1	Send/receive 0

TABLE 11–7 The divisor used with the baud rate generator for an 18.432 MHz crystal illustrating common baud rates.

Baud Rate	Divisor Value
110	10,473
300	3840
1200	920
2400	480
4800	240
9600	120
19,200	60
38,400	30
57,600	20
115,200	10

The remaining bits in the line control register are used to send a break and to select programming for the baud rate divisor. If bit position 6 of the line control register is a logic 1, a break is transmitted. As long as this bit is a 1, the break is sent from the SOUT pin. A break, by definition, is at least two frames of logic 0 data. The software in the system is responsible for timing the transmission of the break. To end the break, bit position 6 or the line control register is returned to a logic 0 level. The baud rate divisor is only programmable when bit position 7 of the line control register is a logic 1.

Programming the Baud Rate. The baud rate generator is programmed at I/O addresses 000 and 001 (A_2, A_1, A_0). Port 000 is used to hold the least significant part of the 16-bit divisor and port 001 is used to hold the most significant part. The value used for the divisor depends on the external clock or crystal frequency. Table 11–7 illustrates common baud rates obtainable if an 18.432 MHz crystal is used as a timing source. It also shows the divisor values programmed into the baud rate generator to obtain these baud rates. The actual number programmed into the baud rate generator causes it to produce a clock that is 16 times the desired baud rate. For example, if 240 is programmed into the baud rate divisor, the baud rate is $(18.432 \text{ MHz} \div 16) \times 240 = 4800$ baud.

Sample Initialization. Suppose that an asynchronous system requires seven data bits, odd parity a baud rate of 9600, and one stop bit. Example 11–24 lists a procedure that initializes the 16550 to function in this manner. Figure 11–45 shows the interface to the 8088 microprocessor, using a PLD to decode the 8-bit port addresses F0H through F7H. (The PLD program is not shown.) Here port F3H accesses the line control register and F0H and F1H access the baud rate divisor registers. The last part of Example 11–26 is described with the function of the FIFO control register in the next few paragraphs.

EXAMPLE 11–26

```
;Initialization dialog for Figure 11-45
;Baud rate 9600, 7 data, odd parity, 1 stop

LINE   EQU   0F3H
LSB    EQU   0F0H
MSB    EQU   0F1H
FIFO   EQU   0F2H


INIT   PROC  NEAR

       MOV  AL,10001010B    ;enable baud rate divisor
       OUT  LINE,AL

       MOV  AL,120          ;program baud 9600
       OUT  LSB,AL
```
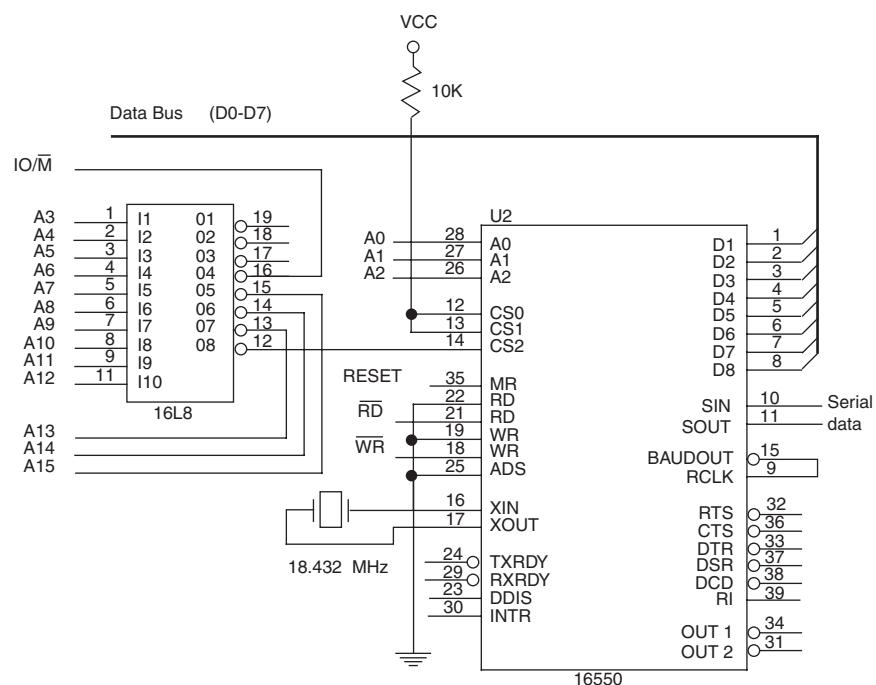
FIGURE 11–45 The 16550 interfaced to the 8088 micro-processor at ports 00F0H–00F7H.


```
        MOV   AL,0
        OUT   MSB,AL

        MOV   AL,00001010B      ;program 7 data, odd
        OUT   LINE,AL           ;parity, 1 stop

        MOV   AL,00000111B      ;enable transmitter and
        OUT   FIFO,AL           ;receiver

        RET

INIT    ENDP
```

After the line control register and baud rate divisor are programmed into the 16550, it is still not ready to function. After programming the line control register and baud rate, we still must program the FIFO control register, which is at port F2H in the circuit of Figure 11–45.

Figure 11–46 illustrates the FIFO control register for the 16550. This register enables the transmitter and receiver (bit 0 = 1), and clears the transmitter and receiver FIFOs. It also provides control for the 16550 interrupts, which are discussed in Chapter 12. Notice that the last section of Example 11–26 places a 7 into the FIFO control register. This enables the transmitter and receiver, and clears both FIFOs. The 16550 is now ready to operate, but without interrupts. Interrupts are automatically disabled when the MR (master reset) input is placed at a logic 1 by the system RESET signal.

Sending Serial Data. Before serial data can be sent or received through the 16550, we need to know the function of the line status register (see Figure 11–47). The line status register contains information about error conditions and the state of the transmitter and receiver. This register is tested before a byte is transmitted or can be received.

Suppose that a procedure (see Example 11–27) is written to transmit the contents of AH to the 16550 and out through its serial data pin (SOUT). The TH bit is polled by software to determine whether the transmitter is ready to receive data. This procedure uses the circuit of Figure 11–45.

FIGURE 11–46 The FIFO control register of the 16550 UART.

FIFO Control Register

7	6	5	4	3	2	1	0
RT1	RT0	0	0	DMA	XMIT RST	REVC RST	EN

FIFO Enable
 0 = disable the FIFO
 1 = enable the FIFO

Receiver Reset
 0 = no effect
 1 = reset receiver FIFO

Transmitter Reset
 0 = no effect
 1 = reset transmitter FIFO

DMA mode control
 0 = set to function as 16450 UART
 1 = FIFO mode

Receiver Trigger level
 00 = 1 bytes in FIFO
 01 = 4 bytes in FIFO
 10 = 8 bytes in FIFO
 11 = 14 bytes in FIFO
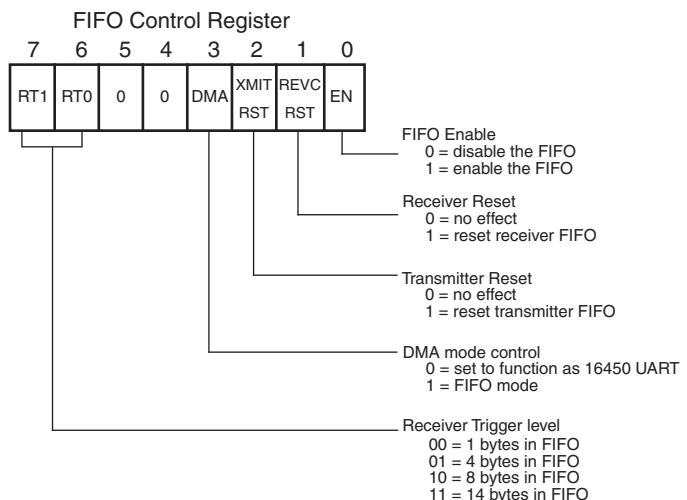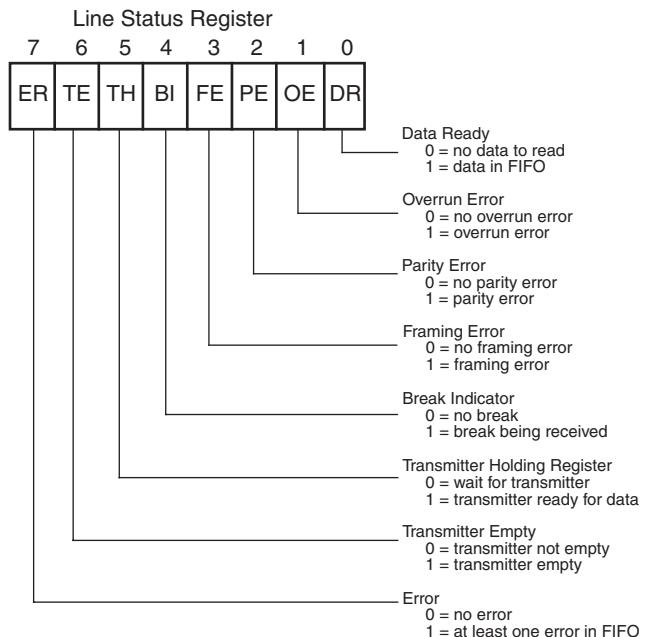
FIGURE 11–47 The contents of the line status register of the 16550 UART.

Line Status Register

7	6	5	4	3	2	1	0
ER	TE	TH	BI	FE	PE	OE	DR

Data Ready
 0 = no data to read
 1 = data in FIFO

Overrun Error
 0 = no overrun error
 1 = overrun error

Parity Error
 0 = no parity error
 1 = parity error

Framing Error
 0 = no framing error
 1 = framing error

Break Indicator
 0 = no break
 1 = break being received

Transmitter Holding Register
 0 = wait for transmitter
 1 = transmitter ready for data

Transmitter Empty
 0 = transmitter not empty
 1 = transmitter empty

Error
 0 = no error
 1 = at least one error in FIFO

EXAMPLE 11–27

```
;A procedure that transmits AH via the 16650 UART

LSTAT EQU    0F5H
DATA  EQU    0F0H

SEND  PROC   NEAR USES AX

      .REPEAT                 ;test the TH bit
            IN   AL,LSTAT
            TEST AL,20H
      .UNTIL !ZERO?

      MOV  AL,AH              ;send data
      OUT  DATA,AL
      RET

SEND  ENDP
```

Receiving Serial Data.   To read received information from the 16550, test the DR bit of the line status register. Example 11–28 lists a procedure that tests the DR bit to decide whether the 16550 has received any data. Upon the reception of data, the procedure tests for errors. If an error is detected, the procedure returns with AL equal to an ASCII '?'. If no error has occurred, then the procedure returns with AL equal to the received character.

EXAMPLE 11–28

```
;A procedure that receives data from the 16550 UART and
;returns it in AL.

LSTAT EQU   0F5H
DATA  EQU   0F0H

REVC  PROC  NEAR

      .REPEAT
            IN   AL,LSTAT       ;test DR bit
            TEST AL,1
      .UNTIL !ZERO?

      TEST AL,0EH               ;test for any error
      .IF ZERO?                 ;no error
            IN   AL,DATA
      .ELSE                     ;any error
            MOV  AL,'?'
      .ENDIF
      RET

RECV  ENDP
```

UART Errors.   The types of errors detected by the 16550 are parity error, framing error, and over-run error. A **parity error** indicates that the received data contain the wrong parity. A **framing error** indicates that the start and stop bits are not in their proper places. An **overrun error** indicates that data have overrun the internal receiver FIFO buffer. These errors should not occur during normal operation. If a parity error occurs, it indicates that noise was encountered during reception. A framing error occurs if the receiver is receiving data at an incorrect baud rate. An overrun error occurs only if the software fails to read the data from the UART before the receiver FIFO is full. This example does not test the BI (break indicator bit) for a break condition. Note that a break is two consecutive frames of logic 0s on the SIN pin of the UART. The remaining registers, which are used for interrupt control and modem control, are developed in Chapter 12.
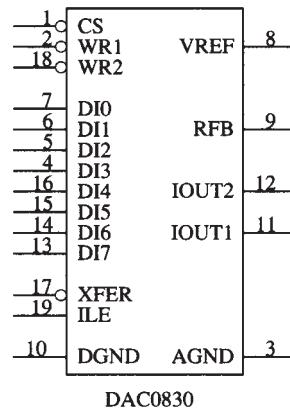
11–6   ANALOG-TO-DIGITAL (ADC) AND DIGITAL-TO-ANALOG (DAC) CONVERTERS

Analog-to-digital (ADC) and digital-to-analog (DAC) converters are used to interface the microprocessor to the analog world. Many events that are monitored and controlled by the microprocessor are analog events. These can range from monitoring all forms of events, even speech, to controlling motors and like devices. In order to interface the microprocessor to these events, we must have an understanding of the interface and control of the ADC and DAC, which convert between analog and digital data.

The DAC0830 Digital-to-Analog Converter

A fairly common and low-cost digital-to-analog converter is the DAC0830 (a product of National Semiconductor Corporation). This device is an 8-bit converter that transforms an 8-bit binary number into an analog voltage. Other converters are available that convert from 10-, 12-,

FIGURE 11–48 The pin-out of the DAC0830 digital-to-analog converter.
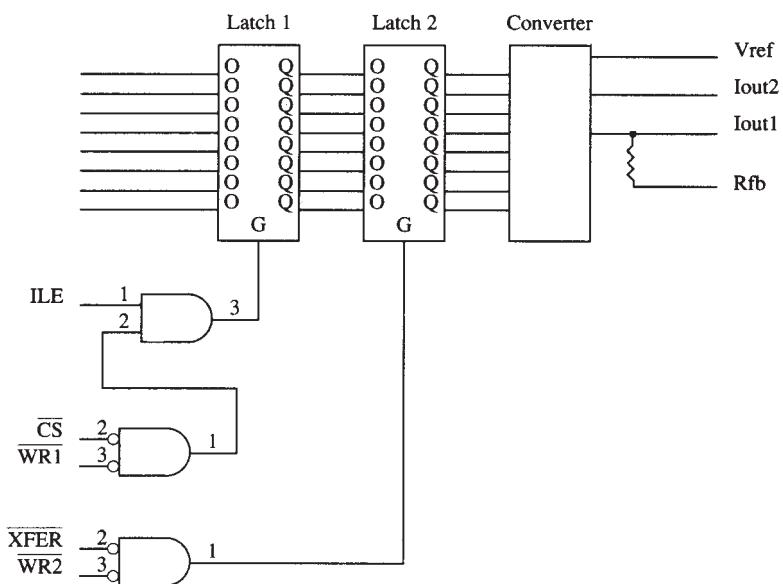
or 16-bit binary numbers into analog voltages. The number of voltage steps generated by the converter is equal to the number of binary input combinations. Therefore, an 8-bit converter generates 256 different voltage levels, a 10-bit converter generates 1024 levels, and so forth. The DAC0830 is a medium-speed converter that transforms a digital input to an analog output in approximately 1.0 μs.

Figure 11–48 illustrates the pin-out of the DAC0830. This device has a set of eight data bus connections for the application of the digital input code, and a pair of analog outputs labeled IOUT1 and IOUT2 that are designed as inputs to an external operational amplifier. Because this is an 8-bit converter, its output step voltage is defined as $-V_{REF}$ (reference voltage), divided by 255. For example, if the reference voltage is -5.0 V, its output step voltage is $+.0196$ V. Note that the output voltage is the opposite polarity of the reference voltage. If an input of $1001\ 0010_2$ is applied to the device, the output voltage will be the step voltage times $1001\ 0010_2$, or, in this case, $+2.862$ V. By changing the reference voltage to -5.1 V, the step voltage becomes $+.02$ V. The step voltage is also often called the resolution of the converter.

Internal Structure of the DAC0830. Figure 11–49 illustrates the internal structure of the DAC0830. Notice that this device contains two internal registers. The first is a holding register,

FIGURE 11–49 The internal structure of the DAC0830.
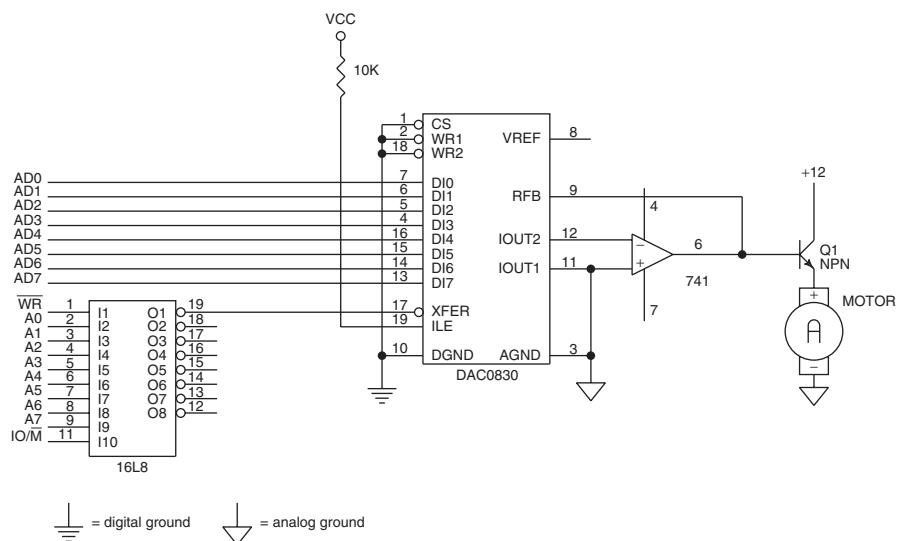
FIGURE 11–50   A DAC0830 interfaced to the 8086 microprocessor at 8-bit I/O location 20H.

and the second connects to the R–2R internal ladder converter. The two latches allow one byte to be held while another is converted. In many cases, we disable the first latch and only use the second for entering data into the converter. This is accomplished by connecting a logic 1 to ILE and a logic 0 to $\overline{\text{CS}}$ (chip select).

Both latches within the DAC0830 are transparent latches. That is, when the G input to the latch is a logic 1, data pass through the latch, but when the G input becomes a logic 0, data are latched or held. The converter has a reference input pin (V_{REF}) that establishes the full-scale output voltage. If −10 V is placed on V_{REF}, the full-scale (11111111_2) output voltage is + 10 V. The output of the R–2R ladder within the converter appears at IOUT1 and IOUT2. These outputs are designed to be applied to an operational amplifier such as a 741 or similar device.

Connecting the DAC0830 to the Microprocessor.   The DAC0830 is connected to the microprocessor as illustrated in Figure 11–50. Here, a PLD is used to decode the DAC0830 at 8-bit I/O port address 20H. Whenever an OUT 20H,AL instruction is executed, the contents of data bus connections AD_0–AD_7 are passed to the converter within the DAC0830. The 741 operational amplifier, along with the −12 V zener reference voltage, causes the full-scale output voltage to equal +12 V. The output of the operational amplifier feeds a driver that powers a 12 V DC motor. This driver is a Darlington amplifier for large motors. This example shows the converter driving a motor, but other devices could be used as outputs.

The ADC080X Analog-to-Digital Converter

A common, low-cost ADC is the ADC080X, which belongs to a family of converters that are all identical, except for accuracy. This device is compatible with a wide range of microprocessors such as the Intel family. Although there are faster ADCs available and some have more resolution than 8 bits, this device is ideal for many applications that do not require a high degree of accuracy. The ADC080X requires up to 100 μs to convert an analog input voltage into a digital output code.

Figure 11–51 shows the pin-out of the ADC0804 converter (a product of National Semiconductor Corporation). To operate the converter, the $\overline{\text{WR}}$ pin is pulsed with $\overline{\text{CS}}$ grounded to start the conversion process. Because this converter requires a considerable amount of time for the conversion, a pin labeled INTR signals the end of the conversion. Refer to Figure 11–52 for

FIGURE 11–51 The pin-out
of the ADC0804 analog-to-
digital converter.

ADC0804

Start conversion   Read data
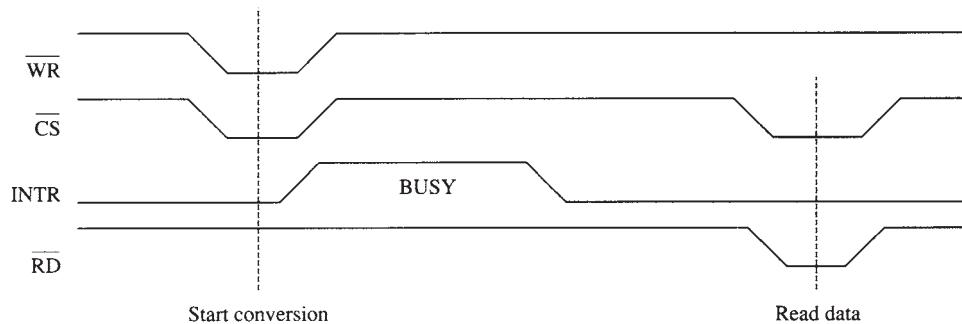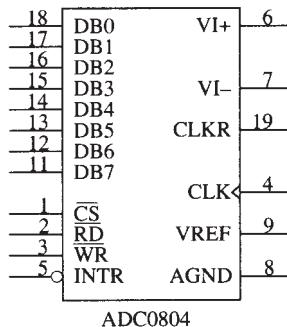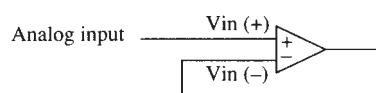
FIGURE 11–52 The timing diagram for the ADC0804 analog-to-digital converter.
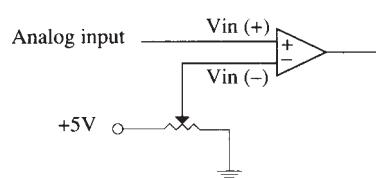
a timing diagram that shows the interaction of the control signals. As can be seen, we start the
converter with the $\overline{WR}$ pulse, we wait for INTR to return to a logic 0 level, and then we read the
data from the converter. If a time delay is used that allows at least 100 μs of time, then we don't
need to test the INTR pin. Another option is to connect the INTR pin to an interrupt input, so that
when the conversion is complete, an interrupt occurs.

The Analog Input Signal. Before the ADC0804 can be connected to the microprocessor, its ana-
log inputs must be understood. There are two analog inputs to the ADC0804: VIN(+) and VIN(−).
These inputs are connected to an internal operational amplifier and are differential inputs,
as shown in Figure 11–53. The differential inputs are summed by the operational amplifier to

FIGURE 11–53 The analog
inputs to the ADC0804 con-
verter. (a) To sense a 0- to
+5.0 V input. (b) To sense an
input offset from ground.
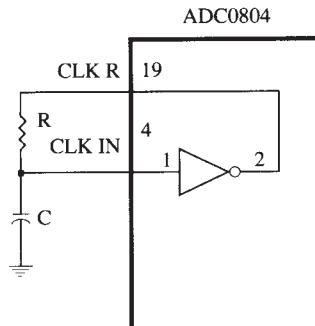
(a)

(b)

FIGURE 11–54 Connecting the RC circuit to the CLK IN and CLK R pins on the ADC0804.

produce a signal for the internal analog-to-digital converter. Figure 11–53 shows a few ways to use these differential inputs. The first way (see Figure 11–53a) uses a single input that can vary between 0 V and +5.0 V. The second way (see Figure 11–53b) shows a variable voltage applied to the VIN(−) pin, so the zero reference for VIN(+) can be adjusted.
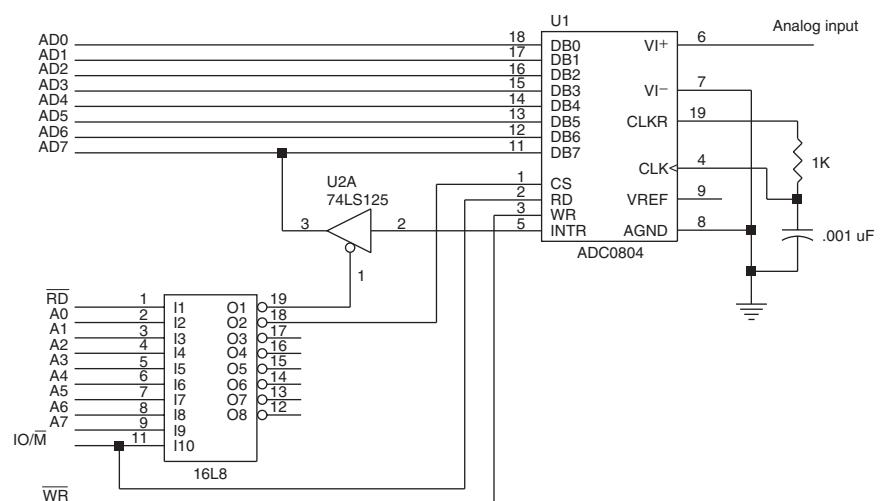
Generating the Clock Signal.   The ADC0804 requires a clock source for operation. The clock can be an external clock applied to the CLK IN pin or it can be generated with an RC circuit. The permissible range of clock frequencies is between 100 KHz and 1460 KHz. It is desirable to use a frequency that is as close as possible to 1460 KHz, so conversion time is kept to a minimum.

If the clock is generated with an RC circuit, we use the CLK IN and CLK R pins connected to an RC circuit, as illustrated in Figure 11–54. When this connection is in use, the clock frequency is calculated by the following equation:

$$\text{Fclk} = \frac{1}{1.1RC}$$

Connecting the ADC0804 to the Microprocessor.   The ADC0804 is interfaced to the 8086 microprocessor, as illustrated in Figure 11–55. Note that the V_{REF} signal is not attached to anything, which is normal. Suppose that the ADC0804 is decoded at 8-bit I/O port address 40H for the data and port address 42H for the INTR signal, and a procedure is required to start and read the data from the ADC. This procedure is listed in Example 11–29. Notice that the INTR bit is polled and if it becomes a logic 0, the procedure ends with AL, containing the converted digital code.

FIGURE 11–55 The ADC0804 interfaced to the microprocessor.

EXAMPLE 11–29

```
ADC    PROC    NEAR

       OUT 40H,AL
       .REPEAT                  ;test INTR
              IN   AL,42H
              TEST AL,80H
       .UNTIL ZERO?
       IN AL,40H
       RET

ADC    ENDP
```

Using the ADC0804 and the DAC0830

This section of the text illustrates an example that uses both the ADC0804 and the DAC0830 to capture and replay audio signals or speech. In the past, we often used a speech synthesizer to generate speech, but the quality of the speech was poor. For human quality speech, we can use the ADC0804 to capture an audio signal and store it in memory for later playback through the DAC0830.

Figure 11–56 illustrates the circuitry required to connect the ADC0804 at I/O ports 0700H and 0702H. The DAC0830 is interfaced at I/O port 704H. These I/O ports are in the low bank of
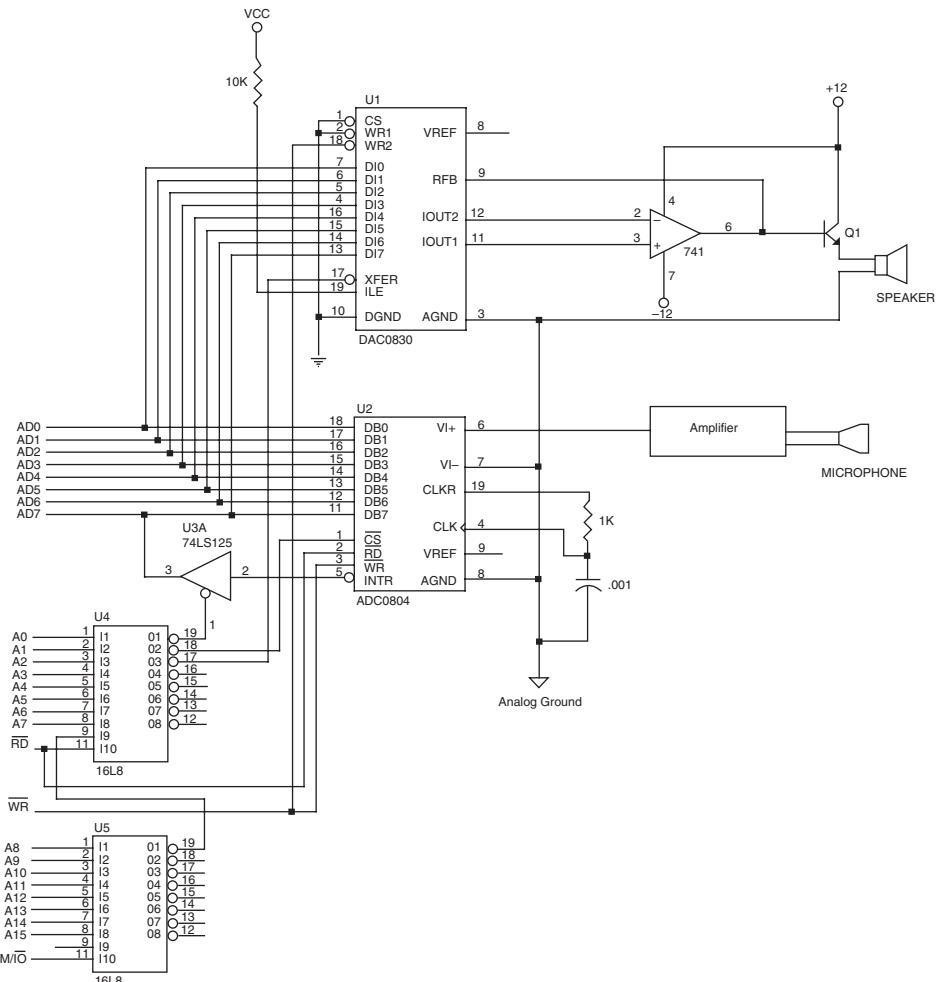
FIGURE 11–56 A circuit that stores speech and plays it back through the speaker.

a 16-bit microprocessor such as the 8086 or 80386SX. The software used to run these converters appears in Example 11–30. This software reads a 1-second burst of speech and then plays it back 10 times. One procedure reads the speech called READS and the other, called PLAYS, plays it back. The speech is sampled and stored in a section of memory called WORDS. The sample rate is chosen at 2048 samples per second, which renders acceptable-sounding speech.

EXAMPLE 11–30

```
;Software that records a second of speech and plays it back
;10 times.

;Assumes the clock frequency is 20 MHz on an 80386EX microprocessor

READS PROC   NEAR USES ECX DX

      MOV  ECX,2048             ;count = 2048
      MOV  DX,700H              ;address port 700H
      .REPEAT
            OUT  DX,AL          ;start conversion
            ADD  DX,2           ;address status port
            .REPEAT             ;wait for converter
                  IN   AL,DX
                  TEST AL,80H
            .UNTIL ZERO?
            SUB  DX,2           ;address data port
            IN   AL,DX          ;get data
            MOV  WORDS[ECX-1]
            CALL DELAY          ;wait for 1/2048 sec
      .UNTILCXZ
      RET

READS ENDP

PLAYS PROC   NEAR USES DX ECX

      MOV  ECX,2048             ;count = 2048
      MOV  DX,704H              ;address DAC
      .REPEAT
            MOV  AL,WORDS[EAX-1]
            OUT  DX,AL          ;send byte to DAC
            CALL DELAY          ;wait for 1/2048 sec
      .UNTILCXZ
      RET
PLAYS ENDP

DELAY PROC   NEAR USES CX

      MOV  CX,888
      .REPEAT                   ;waste 1/2048 sec
      .UNTILCXZ
      RET

DELAY     ENDP
```

11–7   SUMMARY

1. The 8086–Core2 microprocessors have two basic types of I/O instructions: IN and OUT. The IN instruction inputs data from an external I/O device into either the AL (8-bit) or AX (16-bit) register. The IN instruction is available as a fixed port instruction, a variable port instruction, or a string instruction (80286–Pentium 4) INSB or INSW. The OUT instruction

outputs data from AL or AX to an external I/O device and is available as a fixed, variable, or string instruction OUTSB or OUTSW. The fixed port instruction uses an 8-bit I/O port address, while the variable and string I/O instructions use a 16-bit port number found in the DX register.

2. Isolated I/O, sometimes called direct I/O, uses a separate map for the I/O space, freeing the entire memory for use by the program. Isolated I/O uses the IN and OUT instructions to transfer data between the I/O device and the microprocessor. The control structure of the isolated I/O map uses $\overline{IORC}$ (I/O read control) and $\overline{IOWC}$ (I/O write control), plus the bank selection signals $\overline{BHE}$ and $\overline{BLE}$ (A0 on the 8086 and 80286), to effect the I/O transfer. The early 8086/8088 used the M/$\overline{IO}$ (IO/$\overline{M}$) signal with $\overline{RD}$ and $\overline{WR}$ to generate the I/O control signals.

3. Memory-mapped I/O uses a portion of the memory space for I/O transfers. This reduces the amount of memory available, but it negates the need to use the $\overline{IORC}$ and $\overline{IOWC}$ signals for I/O transfers. In addition, any instruction that addresses a memory location using any addressing mode can be used to transfer data between the microprocessor and the I/O device using memory-mapped I/O.

4. All input devices are buffered so that the I/O data are connected only to the data bus during the execution of the IN instruction. The buffer is either built into a programmable peripheral or located separately.

5. All output devices use a latch to capture output data during the execution of the OUT instruction. This is necessary because data appear on the data bus for less than 100 ns for an OUT instruction, and most output devices require the data for a longer time. In many cases, the latch is built into the peripheral.

6. Handshaking or polling is the act of two independent devices synchronizing with a few control lines. For example, the computer asks a printer if it is busy by inputting the BUSY signal from the printer. If it isn't busy, the computer outputs data to the printer and informs the printer that data are available with a data strobe ($\overline{DS}$) signal. This communication between the computer and the printer is a handshake or a poll.

7. Interfaces are required for most switch-based input devices and for most output devices that are not TTL-compatible.

8. The I/O port number appears on address bus connections A_7–A_0 for a fixed port I/O instruction and on A_{15}–A_0 for a variable port I/O instruction (note that A_{15}–A_8 contains zeros for an 8-bit port). In both cases, address bits above A_{15} are undefined.

9. Because the 8086/80286/80386SX microprocessors contain a 16-bit data bus and the I/O addresses reference byte-sized I/O locations, the I/O space is also organized in banks, as is the memory system. In order to interface an 8-bit I/O device to the 16-bit data bus, we often require separate write strobes (an upper and a lower) for I/O write operations. Likewise, the 80486 and Pentium–Core2 also have I/O arranged in banks.

10. The I/O port decoder is much like the memory address decoder, except instead of decoding the entire address, the I/O port decoder decodes only a l6-bit address for variable port instructions and often an 8-bit port number for fixed I/O instructions.

11. The 82C55 is a programmable peripheral interface (PIA) that has 24 I/O pins that are programmable in two groups of 12 pins each (group A and group B). The 82C55 operates in three modes: simple I/O (mode 0), strobed I/O (mode 1), and bidirectional I/O (mode 2). When the 82C55 is interfaced to the 8086 operating at 8 MHz, we insert two wait states because the speed of the microprocessor is faster than the 82C55 can handle.

12. The LCD display device requires a fair amount of software, but it displays ASCII-coded information.

13. The 8254 is a programmable interval timer that contains three l6-bit counters that count in binary or binary-coded decimal (BCD). Each counter is independent and operates in six different modes: (1) events counter, (2) retriggerable, monostable multivibrator, (3) pulse

generator, (4) square-wave generator, (5) software-triggered pulse generator, and (6) hardware-triggered pulse generator.

14. The 16550 is a programmable communications interface, capable of receiving and transmitting asynchronous serial data.

15. The DAC0830 is an 8-bit digital-to-analog converter that converts a digital signal to an analog voltage within 1.0 μs.

16. The ADC0804 is an 8-bit analog-to-digital converter that converts an analog signal into a digital signal within 100 μs.

11–8   QUESTIONS AND PROBLEMS

1. Explain which way the data flow for an IN and an OUT instruction.
2. Where is the I/O port number stored for a fixed I/O instruction?
3. Where is the I/O port number stored for a variable I/O instruction?
4. Where is the I/O port number stored for a string I/O instruction?
5. To which register are data input by the 16-bit IN instruction?
6. Describe the operation of the OUTSB instruction.
7. Describe the operation of the INSW instruction.
8. Contrast a memory-mapped I/O system with an isolated I/O system.
9. What is the basic input interface?
10. What is the basic output interface?
11. Explain the term *handshaking* as it applies to computer I/O systems.
12. An even-number I/O port address is found in the _____ I/O bank in the 8086 microprocessor.
13. In the Pentium 4, what bank contains I/O port number 000AH?
14. How many I/O banks are found in the Pentium 4 or Core2 microprocessor?
15. Show the circuitry that generates the upper and lower I/O write strobes.
16. What is the purpose of a contact bounce eliminator?
17. Develop an interface to correctly drive a relay. The relay is 12 V and requires a coil current of 150 mA.
18. Develop a relay coil driver that can control a 5.0 V relay that requires 60 mA of coil current.
19. Develop an I/O port decoder, using a 74ALS138, that generates low-bank I/O strobes, for a 16-bit microprocessor, for the following 8-bit I/O port addresses: 10H, 12H, 14H, 16H, 18H, 1AH, 1CH, and 1EH.
20. Develop an I/O port decoder, using a 74ALS138, that generates high-bank I/O strobes, for a 16-bit microprocessor, for the following 8-bit I/O port addresses: 11H, 13H, 15H, 17H, 19H, 1BH, 1DH, and 1FH.
21. Develop an I/O port decoder, using a PLD, that generates 16-bit I/O strobes for the following 16-bit I/O port addresses: 1000H–1001H, 1002H–103H, 1004H–1005H, 1006H–1007H, 1008H–1009H, 100AH–100BH, 100CH–100DH, and 100EH–100FH.
22. Develop an I/O port decoder, using a PLD, that generates the following low-bank I/O strobes: 00A8H, 00B6H, and 00EEH.
23. Develop an I/O port decoder, using a PLD, that generates the following high-bank I/O strobes: 300DH, 300BH, 1005H, and 1007H.
24. Why are both $\overline{BHE}$ and $\overline{BLE}$ (A_0) ignored in a 16-bit port address decoder?
25. An 8-bit I/O device, located at I/O port address 0010H, is connected to which data bus connections in a Pentium 4?
26. An 8-bit I/O device, located at I/O port address 100DH, is connected to which data bus connections in a Core2 microprocessor?

27. The 82C55 has how many programmable I/O pin connections?
28. List the pins that belong to group A and to group B in the 82C55.
29. Which two 82C55 pins accomplish internal I/O port address selection?
30. The $\overline{RD}$ connection on the 82C55 is attached to which 8086 system control bus connection?
31. Using a PLD, interface an 82C55 to the 8086 microprocessor so that it functions at I/O locations 0380H, 0382H, 0384H, and 0386H.
32. When the 82C55 is reset, its I/O ports are all initialized as _____.
33. What three modes of operation are available to the 82C55?
34. What is the purpose of the $\overline{STB}$ signal in strobed input operation of the 82C55?
35. Develop a time delay procedure for the 2.0 GHz Pentium 4 that waits for 80 μs.
36. Develop a time delay procedure for the 3.0 GHz Pentium 4 that waits for 12 ms.
37. Explain the operation of a simple four-coil stepper motor.
38. What sets the IBF pin in strobed input operation of the 82C55?
39. Write the software required to place a logic 1 on the PC7 pin of the 82C55 during strobed input operation.
40. How is the interrupt request pin (INTR) enabled in the strobed input mode of operation of the 82C55?
41. In strobed output operation of the 82C55, what is the purpose of the $\overline{ACK}$ signal?
42. What clears the $\overline{OBF}$ signal in strobed output operation of the 82C55?
43. Write the software required to decide whether PC4 is a logic 1 when the 82C55 is operated in the strobed output mode.
44. Which group of pins is used during bidirectional operation of the 82C55?
45. Which pins are general-purpose I/O pins during mode 2 operation of the 82C55?
46. Describe how the display is cleared in the LCD display.
47. How is a display position selected in the LCD display?
48. Write a short procedure that places an ASCII null string in display position 6 on the LCD display.
49. How is the busy flag tested in the LCD display?
50. What changes must be made to Figure 11–25 so that it functions with a keyboard matrix that contains three rows and five columns?
51. What time is usually used to debounce a keyboard?
52. Develop the interface to a three- by four-key telephone-style keypad. You will need to use a lookup table to convert to the proper key code.
53. The 8254 interval timer functions from DC to _____ Hz.
54. Each counter in the 8254 functions in how many different modes?
55. Interface an 8254 to function at I/O port addresses XX10H, XX12H, XX14H, and XX16H.
56. Write the software that programs counter 2 to generate an 80 KHz square wave if the CLK input to counter 2 is 8 MHz.
57. What number is programmed in an 8254 counter to count 300 events?
58. If a 16-bit count is programmed into the 8254, which byte of the count is programmed first?
59. Explain how the read-back control word functions in the 8254.
60. Program counter 1 of the 8254 so that it generates a continuous series of pulses that have a high time of 100 μs and a low time of 1 μs. Make sure to indicate the CLK frequency required for this task.
61. Why does a 50% duty cycle cause the motor to stand still in the motor speed and direction control circuit presented in this chapter?
62. What is asynchronous serial data?
63. What is baud rate?
64. Program the 16550 for operation using six data bits, even parity, one stop bit, and a baud rate of 19,200 using a 18.432 MHz clock. (Assume that the I/O ports are numbered 20H and 22H.)

65. If the 16550 is to generate a serial signal at a baud rate of 2400 baud and the baud rate divisor is programmed for 16, what is the frequency of the signal?
66. Describe the following terms: *simplex, half-duplex,* and *full-duplex.*
67. How is the 16550 reset?
68. Write a procedure for the 16550 that transmits 16 bytes from a small buffer in the data segment address (DS is loaded externally) by SI (SI is loaded externally).
69. The DAC0830 converts an 8-bit digital input to an analog output in approximately_____.
70. What is the step voltage at the output of the DAC0830 if the reference voltage is −2.55 V?
71. Interface a DAC0830 to the 8086 so that it operates at I/O port 400H.
72. Develop a program for the interface of question 71 so the DAC0830 generates a triangular voltage wave-form. The frequency of this wave-form must be approximately 100 Hz.
73. The ADC080X requires approximately _____ to convert an analog voltage into a digital code.
74. What is the purpose of the INTR pin on the ADC080X?
75. The $\overline{WR}$ pin on the ADC080X is used for what purpose?
76. Interface an ADC080X at I/O port 0260H for data and 0270H to test the INTR pin.
77. Develop a program for the ADC080X in question 76 so that it reads an input voltage once per 100 ms and stores the results in a memory array that is 100H bytes long.
78. Rewrite Example 11–29 using C++ with inline assembly code.

CHAPTER 12

Interrupts

INTRODUCTION

In this chapter, the coverage of basic I/O and programmable peripheral interfaces is expanded by examining a technique called interrupt-processed I/O. An interrupt is a hardware-initiated procedure that interrupts whatever program is currently executing. This chapter provides examples and a detailed explanation of the interrupt structure of the entire Intel family of microprocessors.

CHAPTER OBJECTIVES

Upon completion of this chapter, you will be able to:

1. Explain the interrupt structure of the Intel family of microprocessors.
2. Explain the operation of software interrupt instructions INT, INTO, INT 3, and BOUND.
3. Explain how the interrupt enable flag bit (IF) modifies the interrupt structure.
4. Describe the function of the trap interrupt flag bit (TF) and the operation of trap-generated tracing.
5. Develop interrupt-service procedures that control lower-speed, external peripheral devices.
6. Expand the interrupt structure of the microprocessor by using the 82S9A programmable interrupt controller and other techniques.
7. Explain the purpose and operation of a real-time clock.
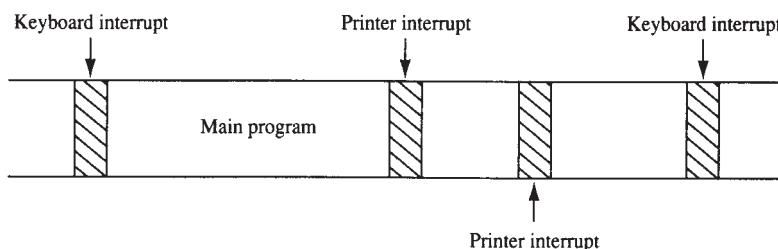
12–1   BASIC INTERRUPT PROCESSING

In this section, we discuss the function of an interrupt in a microprocessor-based system, and the structure and features of interrupts available to the Intel family of microprocessors.

The Purpose of Interrupts

Interrupts are particularly useful when interfacing I/O devices that provide or require data at relatively low data transfer rates. In Chapter 11, for instance, we showed a keyboard example using strobed input operation of the 82C55. In that example, software polled the 82C55 and its IBF bit to decide whether data were available from the keyboard. If the person using the keyboard typed

FIGURE 12–1 A time line that indicates interrupt usage in a typical system.

one character per second, the software for the 82C55 waited an entire second between each keystroke for the person to type another key. This process was such a tremendous waste of time that designers developed another process, *interrupt processing,* to handle this situation.

Unlike the polling technique, interrupt processing allows the microprocessor to execute other software while the keyboard operator is thinking about what key to type next. As soon as a key is pressed, the keyboard encoder debounces the switch and puts out one pulse that interrupts the microprocessor. The microprocessor executes other software until the key is actually pressed, when it reads a key and returns to the program that was interrupted. As a result, the microprocessor can print reports or complete any other task while the operator is typing a document and thinking about what to type next.

Figure 12–1 shows a time line that indicates a typist typing data on a keyboard, a printer removing data from the memory, and a program executing. The program is the main program that is interrupted for each keystroke and each character that is to print on the printer. Note that the keyboard interrupt service procedure, called by the keyboard interrupt, and the printer interrupt service procedure each take little time to execute.

Interrupts

The interrupts of the entire Intel family of microprocessors include two hardware pins that request interrupts (INTR and NMI), and one hardware pin ($\overline{\text{INTA}}$) that acknowledges the interrupt requested through INTR. In addition to the pins, the microprocessor also has software interrupts INT, INTO, INT 3, and BOUND. Two flag bits, IF (interrupt flag) and TF (trap flag), are also used with the interrupt structure and a special return instruction, IRET (or IRETD in the 80386, 80486, or Pentium–Pentium 4).
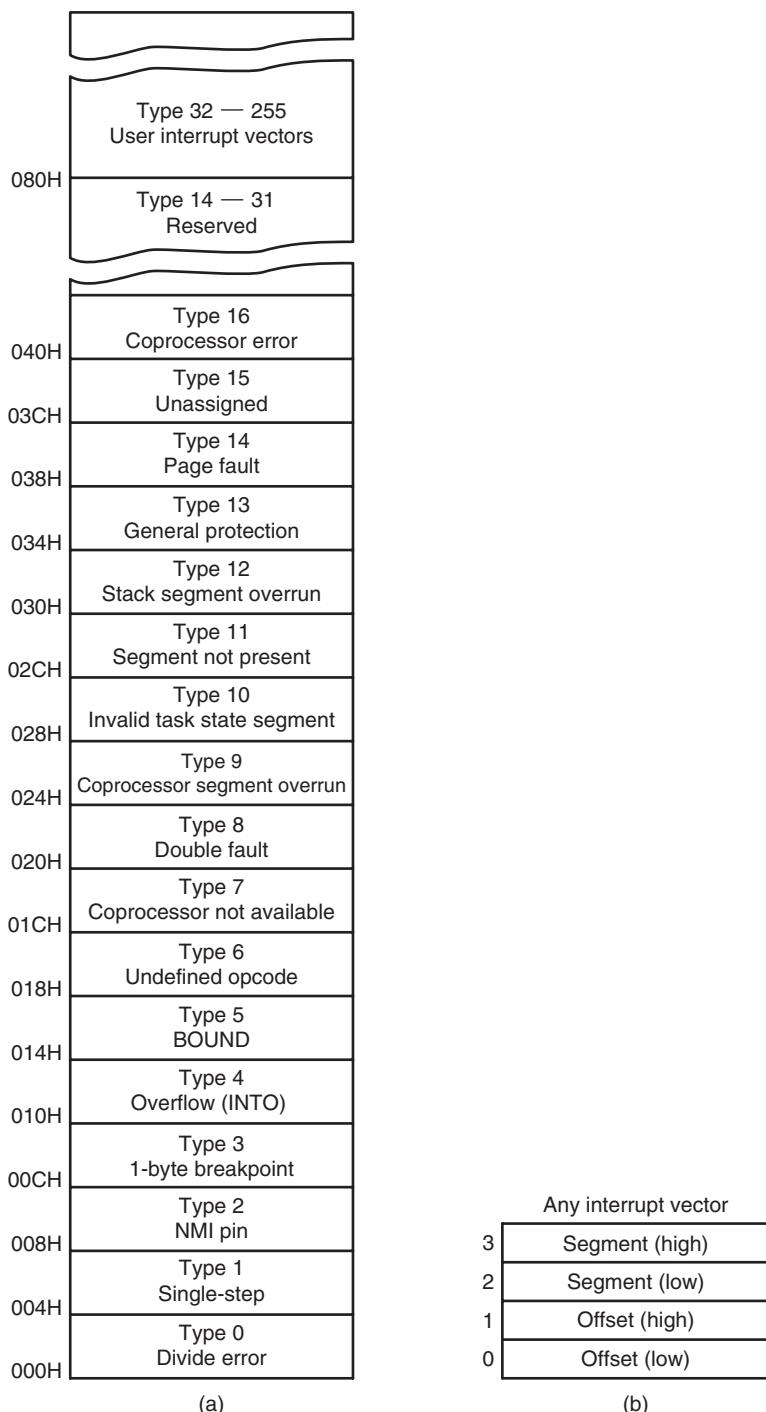
Interrupt Vectors.   The interrupt vectors and vector table are crucial to an understanding of hardware and software interrupts. The **interrupt vector table** is located in the first 1024 bytes of memory at addresses 000000H–0003FFH. It contains 256 different four-byte interrupt vectors. An interrupt vector contains the address (segment and offset) of the interrupt service procedure.

Figure 12–2 illustrates the interrupt vector table for the microprocessor. The first five interrupt vectors are identical in all Intel microprocessor family members, from the 8086 to the Pentium. Other interrupt vectors exist for the 80286 that are upward-compatible to the 80386, 80486, and Pentium–Core2, but not downward-compatible to the 8086 or 8088. Intel reserves the first 32 interrupt vectors for their use in various microprocessor family members. The last 224 vectors are available as user interrupt vectors. Each vector is four bytes long in the real mode and contains the starting address of the interrupt service procedure. The first two bytes of the vector contain the offset address and the last two bytes contain the segment address.

The following list describes the function of each dedicated interrupt in the microprocessor:

TYPE 0   The **divide error** whenever the result from a division overflows or an attempt is made to divide by zero.

TYPE 1   **Single-step** or **trap** occurs after the execution of each instruction if the trap (TF) flag bit is set. Upon accepting this interrupt, the TF bit is cleared so that the

FIGURE 12–2 (a) The interrupt vector table for the microprocessor and (b) the contents of an interrupt vector.


```
                        Type 32 — 255
                        User interrupt vectors
080H
                        Type 14 — 31
                        Reserved

                        Type 16
                        Coprocessor error
040H
                        Type 15
                        Unassigned
03CH
                        Type 14
                        Page fault
038H
                        Type 13
                        General protection
034H
                        Type 12
                        Stack segment overrun
030H
                        Type 11
                        Segment not present
02CH
                        Type 10
                        Invalid task state segment
028H
                        Type 9
                        Coprocessor segment overrun
024H
                        Type 8
                        Double fault
020H
                        Type 7
                        Coprocessor not available
01CH
                        Type 6
                        Undefined opcode
018H
                        Type 5
                        BOUND
014H
                        Type 4
                        Overflow (INTO)
010H
                        Type 3
                        1-byte breakpoint
00CH
                        Type 2
                        NMI pin
008H
                        Type 1
                        Single-step
004H
                        Type 0
                        Divide error
000H
                          (a)
```

```
        Any interrupt vector
    3   Segment (high)
    2   Segment (low)
    1   Offset (high)
    0   Offset (low)
              (b)
```

interrupt service procedure executes at full speed. (More detail is provided about this interrupt later in this section of the chapter.)

TYPE 2   The **non-maskable interrupt** occurs when a logic 1 is placed on the NMI input pin to the microprocessor. This input is non-maskable, which means that it cannot be disabled.

TYPE 3   A special one-byte instruction (INT 3) that uses this vector to access its interrupt-service procedure. The INT 3 instruction is often used to store a **breakpoint** in a program for debugging.

TYPE 4   **Overflow** is a special vector used with the INTO instruction. The INTO instruction interrupts the program if an overflow condition exists, as reflected by the overflow flag (OF).

TYPE 5   The **BOUND** instruction compares a register with boundaries stored in the memory. If the contents of the register are greater than or equal to the first word in memory and less than or equal to the second word, no interrupt occurs because the contents of the register are within bounds. If the contents of the register are out of bounds, a type 5 interrupt ensues.

TYPE 6   An **invalid opcode** interrupt occurs whenever an undefined opcode is encountered in a program.

TYPE 7   The **coprocessor not available** interrupt occurs when a coprocessor is not found in the system, as dictated by the machine status word (MSW or CR0) coprocessor control bits. If an ESC or WAIT instruction executes and the coprocessor is not found, a type 7 exception or interrupt occurs.

TYPE 8   A **double fault** interrupt is activated whenever two separate interrupts occur during the same instruction.

TYPE 9   The **coprocessor segment overrun** occurs if the ESC instruction (coprocessor opcode) memory operand extends beyond offset address FFFFH in real mode.

TYPE 10   An **invalid task state segment** interrupt occurs in the protected mode if the TSS is invalid because the segment limit field is not 002BH or higher. In most cases, this is caused because the TSS is not initialized.

TYPE 11   The **segment not present** interrupt occurs when the protected mode P bit (P = 0) in a descriptor indicates that the segment is not present or not valid.

TYPE 12   A **stack segment overrun** occurs if the stack segment is not present (P = 0) in the protected mode or if the limit of the stack segment is exceeded.

TYPE 13   The **general protection fault** occurs for most protection violations in the 80286–Core2 protected mode system. (These errors occur in Windows as general protection faults.) A list of these protection violations follows:
(a)   Descriptor table limit exceeded
(b)   Privilege rules violated
(c)   Invalid descriptor segment type loaded
(d)   Write to code segment that is protected
(e)   Read from execute-only code segment
(f)   Write to read-only data segment
(g)   Segment limit exceeded
(h)   CPL = IOPL when executing CTS, HLT, LGDT, LIDT, LLDT, LMSW, or LTR
(i)   CPL > IOPL when executing CLI, IN, INS, LOCK, OUT, OUTS, and STI

TYPE 14   **Page fault** interrupts occur for any page fault memory or code access in the 80386, 80486, and Pentium–Core2 microprocessors.

TYPE 16   **Coprocessor error** takes effect whenever a coprocessor error (ERROR = 0) occurs for the ESCape or WAIT instructions for the 80386, 80486, and Pentium–Core2 microprocessors only.

TYPE 17   **Alignment checks** indicate that word and doubleword data are addressed at an odd memory location (or an incorrect location, in the case of a doubleword). This interrupt is active in the 80486 and Pentium–Core2 microprocessors.

TYPE 18   A **machine check** activates a system memory management mode interrupt in the Pentium–Core2 microprocessors.

Interrupt Instructions: BOUND, INTO, INT, INT 3, and IRET

Of the five software interrupt instructions available to the microprocessor, INT and INT 3 are very similar, BOUND and INTO are conditional, and IRET is a special interrupt return instruction.

The BOUND instruction, which has two operands, compares a register with two words of memory data. For example, if the instruction BOUND AX,DATA is executed, AX is compared with the contents of DATA and DATA+1 and also with DATA+2 and DATA+3. If AX is less than the contents of DATA and DATA+1, a type 5 interrupt occurs. If AX is greater than DATA+2 and DATA+3, a type 5 interrupt occurs. If AX is within the bounds of these two memory words, no interrupt occurs.

The INTO instruction checks or tests the overflow flag (O). If $O = 1$, the INTO instruction calls the procedure whose address is stored in interrupt vector type number 4. If $O = 0$, then the INTO instruction performs no operation and the next sequential instruction in the program executes.

The INT n instruction calls the interrupt service procedure that begins at the address represented in vector number n. For example, an INT 80H or INT 128 calls the interrupt service procedure whose address is stored in vector type number 80H (000200H–00203H). To determine the vector address, just multiply the vector type number (n) by 4, which gives the beginning address of the four-byte long interrupt vector. For example, INT $5 = 4 \times 5$ or 20 (14H). The vector for INT 5 begins at address 0014H and continues to 0017H. Each INT instruction is stored in two bytes of memory: The first byte contains the opcode, and the second byte contains the interrupt type number. The only exception to this is the INT 3 instruction, a one-byte instruction. The INT 3 instruction is often used as a breakpoint-interrupt because it is easy to insert a one-byte instruction into a program. Breakpoints are often used to debug faulty software.

The IRET instruction is a special return instruction used to return for both software and hardware interrupts. The IRET instruction is much like a far RET, because it retrieves the return address from the stack. It is unlike the near return because it also retrieves a copy of the flag register from the stack. An IRET instruction removes six bytes from the stack: two for the IP, two for the CS, and two for the flags.

In the 80386–Core2, there is also an IRETD instruction because these microprocessors can push the EFLAG register (32 bits) on the stack, as well as the 32-bit EIP in the protected mode and 16-bit code segment register. If operated in the real mode, we use the IRET instruction with the 80386–Core2 microprocessors. If the Pentium 4 operates in 64-bit mode, an IRETQ instruction is used to return from an interrupt. The IRETQ instruction pops the EFLAG register into RFLAGS and also the 64-bit return address is placed into the RIP register.

The Operation of a Real Mode Interrupt

When the microprocessor completes executing the current instruction, it determines whether an interrupt is active by checking (1) instruction executions, (2) single-step, (3) NMI, (4) coprocessor segment overrun, (5) INTR, and (6) INT instructions in the order presented. If one or more of these interrupt conditions are present, the following sequence of events occurs:

1. The contents of the flag register are pushed onto the stack.
2. Both the interrupt (IF) and trap (TF) flags are cleared. This disables the INTR pin and the trap or single-step feature.
3. The contents of the code segment register (CS) are pushed onto the stack.

4. The contents of the instruction pointer (IP) are pushed onto the stack.
5. The interrupt vector contents are fetched, and then placed into both IP and CS so that the next instruction executes at the interrupt service procedure addressed by the vector.

Whenever an interrupt is accepted, the microprocessor stacks the contents of the flag register, CS and IP; clears both IF and TF; and jumps to the procedure addressed by the interrupt vector. After the flags are pushed onto the stack, IF and TF are cleared. These flags are returned to the state prior to the interrupt when the IRET instruction is encountered at the end of the interrupt service procedure. Therefore, if interrupts were enabled prior to the interrupt service procedure, they are automatically re-enabled by the IRET instruction at the end of the procedure.

The return address (in CS and IP) is pushed onto the stack during the interrupt. Sometimes the return address points to the next instruction in the program; sometimes it points to the instruction or point in the program where the interrupt occurred. Interrupt type numbers 0, 5, 6, 7, 8, 10, 11, 12, and 13 push a return address that points to the offending instruction, instead of to the next instruction in the program. This allows the interrupt service procedure to possibly retry the instruction in certain error cases.

Some of the protected mode interrupts (types 8, 10, 11, 12, and 13) place an error code on the stack following the return address. The error code identifies the selector that caused the interrupt. In cases where no selector is involved, the error code is a 0.

Operation of a Protected Mode Interrupt

In the protected mode, interrupts have exactly the same assignments as in the real mode, but the interrupt vector table is different. In place of interrupt vectors, protected mode uses a set of 256 interrupt descriptors that are stored in an interrupt descriptor table (IDT). The interrupt descriptor table is 256×8 (2K) bytes long, with each descriptor containing eight bytes. The interrupt descriptor table is located at any memory location in the system by the interrupt descriptor table address register (IDTR).

Each entry in the IDT contains the address of the interrupt service procedure in the form of a segment selector and a 32-bit offset address. It also contains the P bit (present) and DPL bits to describe the privilege level of the interrupt. Figure 12–3 shows the contents of the interrupt descriptor.

Real mode interrupt vectors can be converted into protected mode interrupts by copying the interrupt procedure addresses from the interrupt vector table and converting them to 32-bit offset addresses that are stored in the interrupt descriptors. A single selector and segment descriptor can be placed in the global descriptor table that identifies the first 1M byte of memory as the interrupt segment.

Other than the IDT and interrupt descriptors, the protected mode interrupt functions like the real mode interrupt. We return from both interrupts by using the IRET or IRETD instruction. The only difference is that in protected mode the microprocessor accesses the IDT instead of the

FIGURE 12–3 The protected mode interrupt descriptor.

```
FLAGS  [        |O|D|I|T|S|Z| |A| |P| |C|
        15       11 10 9 8 7 6 5 4 3 2 1 0
```
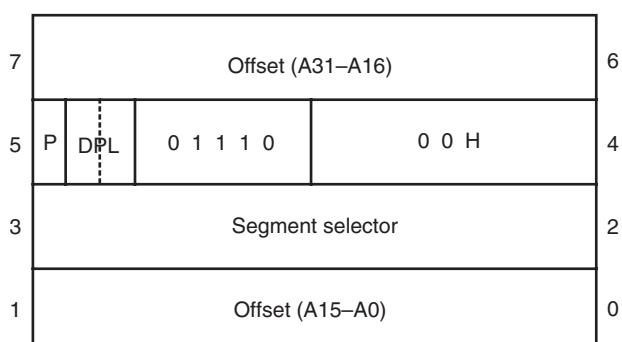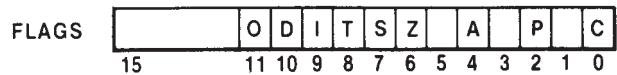
FIGURE 12–4 The flag register. (Courtesy of Intel Corporation.)

interrupt vector table. In the 64-bit mode of the Pentium 4 and Core2, an IRETQ must be used to return from an interrupt. This is one reason why there are different drivers and operating systems for the 64-bit mode.

Interrupt Flag Bits

The interrupt flag (IF) and the trap flag (TF) are both cleared after the contents of the flag register are stacked during an interrupt. Figure 12–4 illustrates the contents of the flag register and the location of IF and TF. When the IF bit is set, it allows the INTR pin to cause an interrupt; when the IF bit is cleared, it prevents the INTR pin from causing an interrupt. When TF = 1, it causes a trap interrupt (type number 1) to occur after each instruction executes. This is why we often call trap a *single-step.* When TF = 0, normal program execution occurs. This flag bit allows debugging, as explained in Chapters 17 through 19, which detail the 80386–Core2.

The interrupt flag is set and cleared by the STI and CLI instructions, respectively. There are no special instructions that set or clear the trap flag. Example 12–1 shows an interrupt service procedure that turns tracing on by setting the trap flag bit on the stack from inside the procedure. Example 12–2 shows an interrupt service procedure that turns tracing off by clearing the trap flag on the stack from within the procedure.

EXAMPLE 12–1

```
;A procedure that sets the TRAP flag bit to enable trapping

TRON  PROC   FAR USES AX BP

      MOV  BP,SP            ;get SP
      MOV  AX[BP+8]         ;retrieve flags from stack
      OR   AH,1             ;set trap flag
      MOV  [BP+8],AX
      IRET

TRON  ENDP
```

EXAMPLE 12–2

```
;A procedure that clears the TRAP flag to disable trapping

TROFF PROC   FAR USES AX BP

      MOV  BP,SP            ;get SP
      MOV  AX,[BP+8]        ;retrieve flags from stack
      AND  AH,0FEH          ;clear trap flag
      MOV  [BP+8],AX
      IRET

TROFF ENDP
```

In both examples, the flag register is retrieved from the stack by using the BP register, which, by default, addresses the stack segment. After the flags are retrieved, the TF bit is either set (TRON) or clears (TROFF) before returning from the interrupt service procedure. The IRET instruction restores the flag register with the new state of the trap flag.

Trace Procedure. Assuming that TRON is accessed by an INT 40H instruction and TROFF is accessed by an INT 41H instruction, Example 12–3 traces through a program immediately following the INT 40H instruction. The interrupt service procedure illustrated in Example 12–3

responds to interrupt type number 1 or a trap interrupt. Each time that a trap occurs—after each instruction executes following INT 40H—the TRACE procedure stores the contents of all the 32-bit microprocessor registers in an array called REGS. This provides a register trace of all the instructions between the INT 40H (TRON) and INT 41H (TROFF) if the contents of the registers stored in the array are saved.

EXAMPLE 12–3

```
REGS    DD   8 DUP(?)         ;space for registers

TRACE   PROC FAR USES EBX

        MOV  EBX,OFFSET REGS
        MOV  [EBX],EAX         ;save EAX
        POP  EAX
        PUSH EAX
        MOV  [EBX+4],EAX       ;save EBX
        MOV  [EBX+8],ECX       ;save ECX
        MOV  [EBX+12],EDX      ;save EDX
        MOV  [EBX+16],ESP      ;save ESP
        MOV  [EBX+20],EBP      ;save EBP
        MOV  [EBX+24],ESI      ;save ESI
        MOV  [EBX+28],EDI      ;save EDI
        IRET

TRACE   ENDP
```

Storing an Interrupt Vector in the Vector Table

In order to install an interrupt vector—sometimes called a **hook**—the assembler must address absolute memory. Example 12–4 shows how a new vector is added to the interrupt vector table by using the assembler and a DOS function call. Here, the vector for INT 40H, for interrupt procedure NEW40, is installed in memory at real mode vector location 100H–103H. The first thing accomplished by the procedure is that the old interrupt vector contents are saved in case we need to uninstall the vector. This step can be skipped if there is no need to uninstall the interrupt.

The function AX = 3100H for INT 21H, the DOS access function, installs the NEW40 procedure in memory until the computer is shut off. The number in DX is the length of the software in paragraphs (16-byte chunks). Refer to Appendix A for more detail about this DOS function.

Notice that the INT40 function has an IRET instruction before ENDP. This is required because the assembler has no way of determining if the FAR procedure is an interrupt procedure. Normal FAR procedures do not need a return instruction, but an interrupt procedure does need an IRET. Interrupts must always be defined as FAR.

EXAMPLE 12–4

```
.MODEL TINY
.CODE
.STARTUP
        JMP    START
OLD     DD     ?      ;space for old vector

NEW40   PROC   FAR    ;must be FAR

;
;Interrupt software for INT 40H
;

        IRET           ;must have an IRET

NEW40   ENDP
```

```
;start installation

START:
        MOV      AX,0                        ;address segment 0000H
        MOV      DS,AX
        MOV      AX,DS:[100H]                ;get INT 40H offset
        MOV      WORD PTR CS:OLD,AX          ;save it
        MOV      AX,DS:[102H]                ;get INT 40H segment
        MOV      WORD PTR CS:OLD+2,AX        ;save it
        MOV      DS:[100H],OFFSET NEW40      ;save offset
        MOV      DS:[102H],CS                ;save segment
        MOV      DX,OFFSET START
        SHR      DX,4
        INC      DX
        MOV      AX,3100H                    ;make NEW40 resident
        INT      21H

END
```
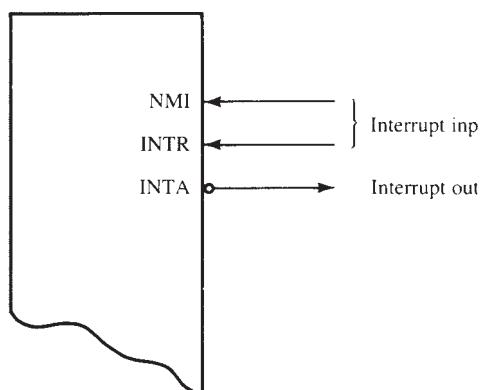
12–2   HARDWARE INTERRUPTS

The microprocessor has two hardware interrupt inputs: non-maskable interrupt (NMI) and interrupt request (INTR). Whenever the NMI input is activated, a type 2 interrupt occurs because NMI is internally decoded. The INTR input must be externally decoded to select a vector. Any interrupt vector can be chosen for the INTR pin, but we usually use an interrupt type number between 20H and FFH. Intel has reserved interrupts 00H through 1FH for internal and future expansion. The $\overline{\text{INTA}}$ signal is also an interrupt pin on the microprocessor, but it is an output that is used in response to the INTR input to apply a vector type number to the data bus connections D_7–D_0. Figure 12–5 shows the three user interrupt connections on the microprocessor.

The **non-maskable interrupt** (NMI) is an edge-triggered input that requests an interrupt on the positive edge (0-to-1 transition). After a positive edge, the NMI pin must remain a logic 1 until it is recognized by the microprocessor. Note that before the positive edge is recognized, the NMI pin must be a logic 0 for at least two clocking periods.

The NMI input is often used for parity errors and other major system faults, such as power failures. Power failures are easily detected by monitoring the AC power line and causing an NMI interrupt whenever AC power drops out. In response to this type of interrupt, the microprocessor stores all of the internal register in a battery-backed-up memory or an EEPROM. Figure 12–6 shows a power failure detection circuit that provides a logic 1 to the NMI input whenever AC power is interrupted.

FIGURE 12–5 The interrupt pins on all versions of the Intel microprocessor.
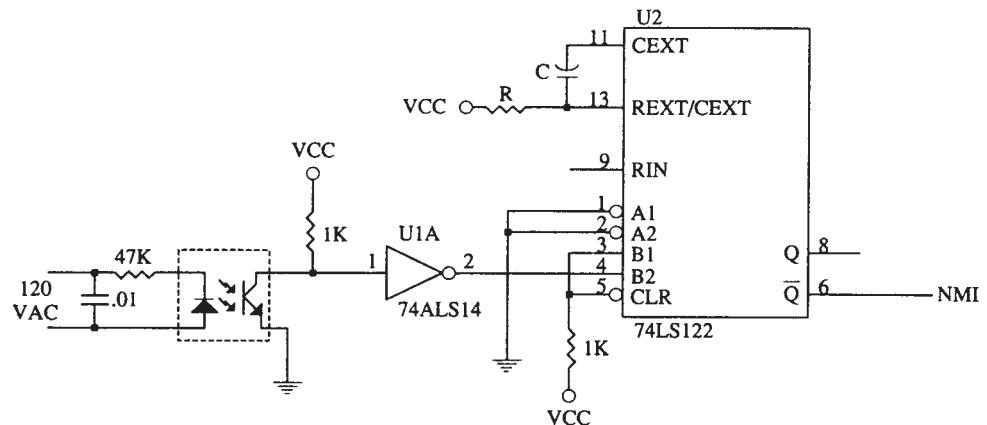
FIGURE 12–6   A power failure detection circuit.
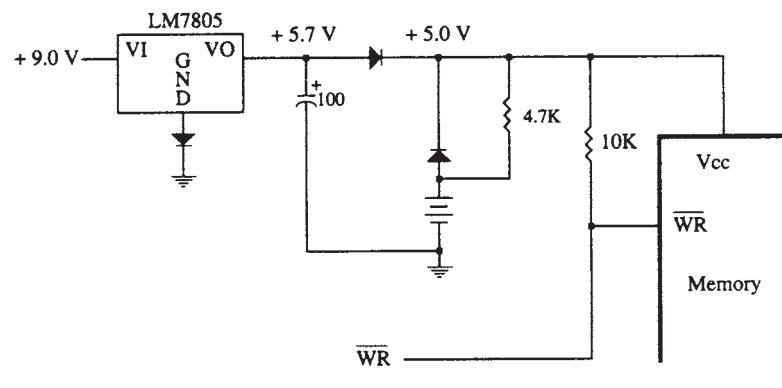
In this circuit, an optical isolator provides isolation from the AC power line. The output of the isolator is shaped by a Schmitt-trigger inverter that provides a 60 Hz pulse to the trigger input of the 74LS122 retriggerable, monostable multivibrator. The values of R and C are chosen so that the 74LS122 has an active pulse width of 33 ms or 2 AC input periods. Because the 74LS122 is retriggerable, as long as AC power is applied, the Q output remains triggered at a logic 1 and $\overline{Q}$ remains a logic 0.

If the AC power fails, the 74LS122 no longer receives trigger pulses from the 74ALS14, which means that Q becomes a logic 0 and $\overline{Q}$ becomes a logic 1, interrupting the microprocessor through the NMI pin. The interrupt service procedure, not shown here, stores the contents of all internal registers and other data into a battery-backed-up memory. This system assumes that the system power supply has a large enough filter capacitor to provide energy for at least 75 ms after the AC power ceases.

Figure 12–7 shows a circuit that supplies power to a memory after the DC power fails. Here, diodes are used to switch supply voltages from the DC power supply to the battery. The diodes used are standard silicon diodes because the power supply to this memory circuit is elevated above +5.0 V to +5.7 V. The resistor is used to trickle-charge the battery, which is either NiCAD, lithium, or a gel cell.

When DC power fails, the battery provides a reduced voltage to the V_{CC} connection on the memory device. Most memory devices will retain data with V_{CC} voltages as low as 1.5 V, so the battery voltage does not need to be +5.0 V. The $\overline{WR}$ pin is pulled to V_{CC} during a power outage, so no data will be written to the memory.

FIGURE 12–7   A battery-backed-up memory system using a NiCad, lithium, or gel cell.
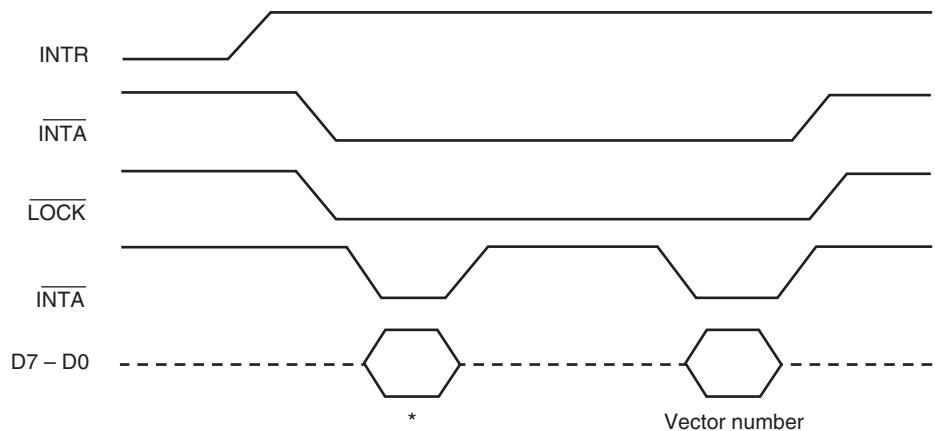
FIGURE 12–8   The timing of the INTR input and $\overline{\text{INTA}}$ output. *This portion of the data bus is ignored and usually contains the vector number.

INTR and $\overline{\text{INTA}}$

The interrupt request input (INTR) is level-sensitive, which means that it must be held at a logic 1 level until it is recognized. The INTR pin is set by an external event and cleared inside the interrupt service procedure. This input is automatically disabled once it is accepted by the microprocessor and re-enabled by the IRET instruction at the end of the interrupt service procedure. The 80386–Core2 use the IRETD instruction in the protected mode of operation. In the 64-bit mode, an IRETQ is used in protected mode.

The microprocessor responds to the INTR input by pulsing the $\overline{\text{INTA}}$ output in anticipation of receiving an interrupt vector type number on data bus connections D_7–D_0. Figure 12–8 shows the timing diagram for the INTR and $\overline{\text{INTA}}$ pins of the microprocessor. There are two $\overline{\text{INTA}}$ pulses generated by the system that are used to insert the vector type number on the data bus.

Figure 12–9 illustrates a simple circuit that applies interrupt vector type number FFH to the data bus in response to an INTR. Notice that the $\overline{\text{INTA}}$ pin is not connected in this circuit.

FIGURE 12–9   A simple method for generating interrupt vector type number FFH in response to INTR.
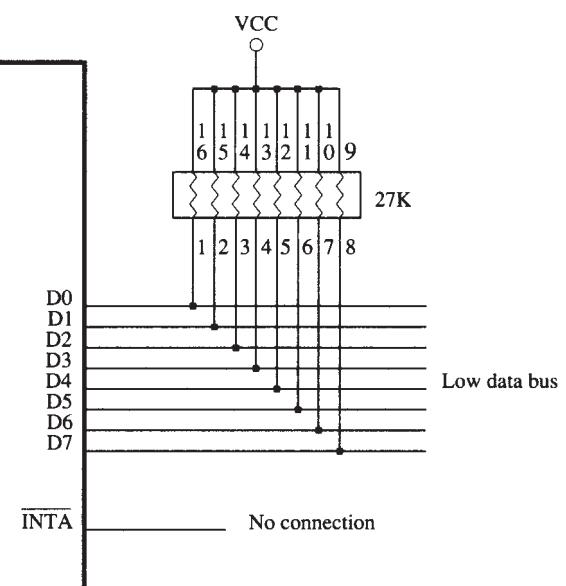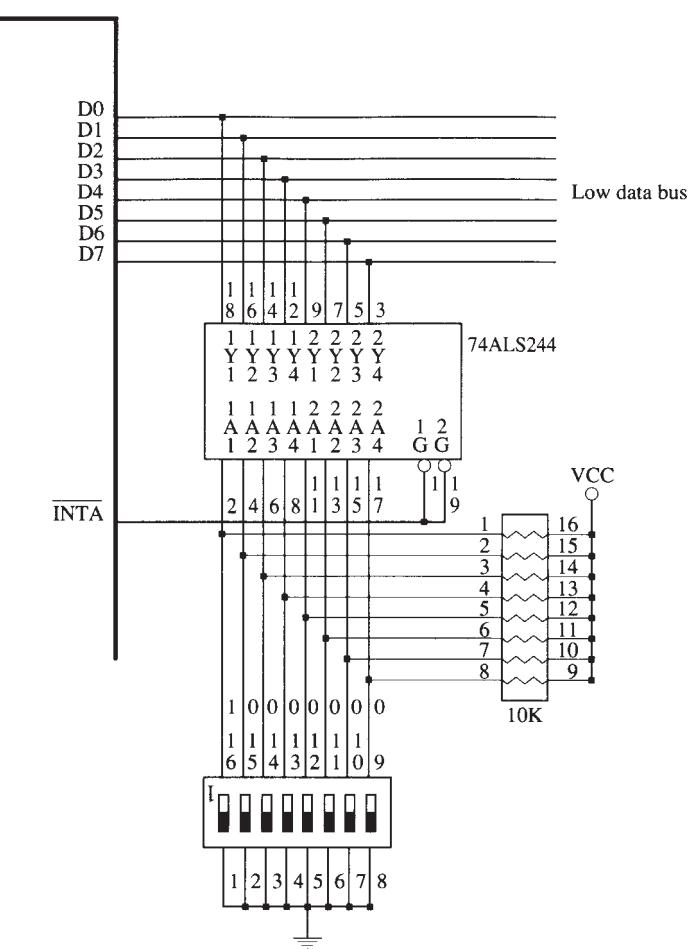
FIGURE 12–10 A circuit that applies any interrupt vector type number in response to $\overline{INTA}$. Here the circuit is applying type number 80H.

Because resistors are used to pull the data bus connections (D_0–D_7) high, the microprocessor automatically sees vector type number FFH in response to the INTR input. This is the least expensive way to implement the INTR pin on the microprocessor.

Using a Three-State Buffer for INTA.   Figure 12–10 shows how interrupt vector type number 80H is applied to the data bus (D_0–D_7) in response to an INTR. In response to the INTR, the microprocessor outputs the $\overline{INTA}$ that is used to enable a 74ALS244 three-state octal buffer. The octal buffer applies the interrupt vector type number to the data bus in response to the $\overline{INTA}$ pulse. The vector type number is easily changed with the DIP switches that are shown in this illustration.

Making the INTR Input Edge-Triggered.   Often, we need an edge-triggered input instead of a level-sensitive input. The INTR input can be converted to an edge-triggered input by using a D-type flip-flop, as illustrated in Figure 12–11. Here, the clock input becomes an edge-triggered interrupt request input, and the clear input is used to clear the request when the $\overline{INTA}$ signal is output by the microprocessor. The RESET signal initially clears the flip-flop so that no interrupt is requested when the system is first powered.

The 82C55 Keyboard Interrupt

The keyboard example presented in Chapter 11 provides a simple example of the operation of the INTR input and an interrupt. Figure 12–12 illustrates the interconnection of the 82C55 with the microprocessor and the keyboard. It also shows how a 74ALS244 octal buffer is used to provide

FIGURE 12–11 Converting INTR into an edge-triggered interrupt request input.
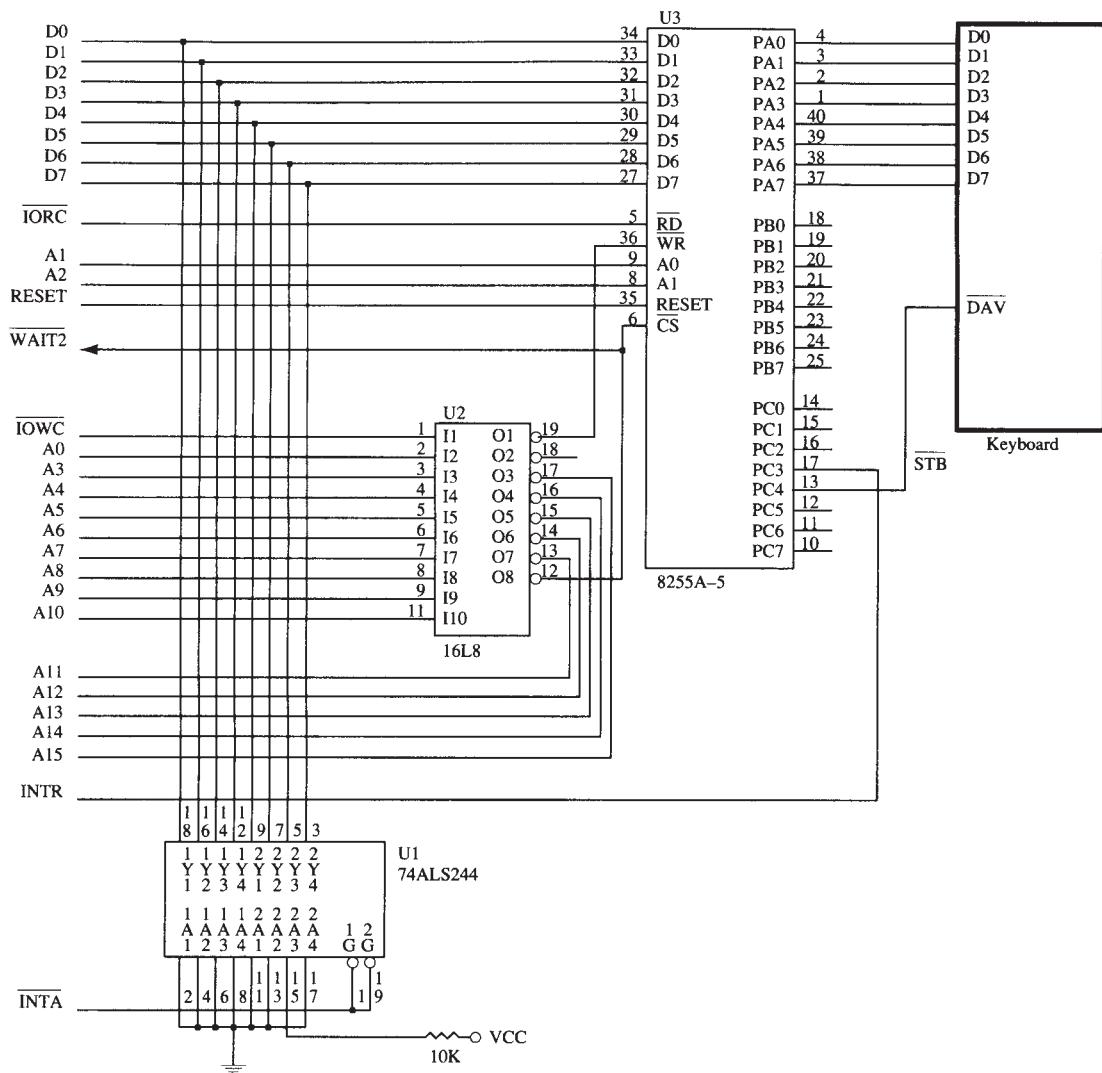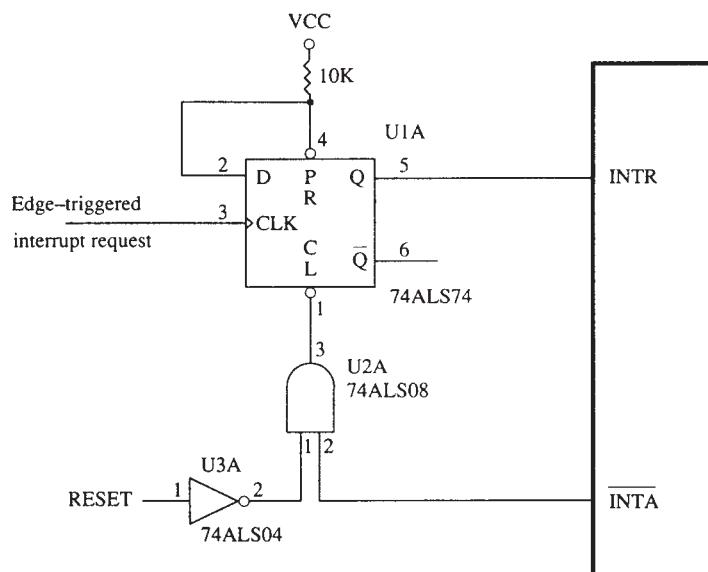
FIGURE 12–12 An 82C55 interfaced to a keyboard from the microprocessor system using interrupt vector 40H.

the microprocessor with interrupt vector type number 40H in response to the keyboard interrupt during the $\overline{\text{INTA}}$ pulse.

The 82C55 is decoded at 80386SX I/O port address 0500H, 0502H, 0504H, and 0506H by a PLD (the program is not illustrated). The 82C55 is operated in mode 1 (strobed input mode), so whenever a key is typed, the INTR output (PC3) becomes a logic 1 and requests an interrupt through the INTR pin on the microprocessor. The INTR pin remains high until the ASCII data are read from port A. In other words, every time a key is typed, the 82C55 requests a type 40H interrupt through the INTR pin. The $\overline{\text{DAV}}$ signal from the keyboard causes data to be latched into port A and causes INTR to become a logic 1.

Example 12–5 illustrates the interrupt service procedure for the keyboard. It is very important that all registers affected by an interrupt are saved before they are used. In the software required to initialize the 82C55 (not shown here), the FIFO is initialized so that both pointers are equal, the INTR request pin is enabled through the INTE bit inside the 82C55, and the mode of operation is programmed.

EXAMPLE 12–5

```
;An interrupt service procedure that reads a key from
;the keyboard depicted in Figure 12-12.

PORTA   EQU    500H
CNTR    EQU    506H

FIFO    DB     256 DUP(?)              ;queue

INP     DD     FIFO                    ;input pointer
OUTP    DD     FIFO                    ;output pointer

KEY     PROC   FAR USES EAX EBX EDX EDI

        MOV    EBX,CS:INP              ;get pointers
        MOV    EDI,CS:OUTP

        INC    BL
        .IF BX == DI                   ;if full
            MOV    AL,8
            MOV    DX,CNTR
            OUT    DX,AL               ;disable 82C55 interrupt
        .ELSE                          ;if not full
            DEC  BL
            MOV  DX,PORTA
            IN   AL,DX                 ;read key code
            MOV  CS:[BX]               ;save in queue
            INC  BYTE PTR CS:INP
        .ENDIF
        IRET

KEY     ENDP
```

The procedure is short because the microprocessor already knows that keyboard data are available when the procedure is called. Data are input from the keyboard and then stored in the FIFO (first-in, first-out) buffer or queue. Most keyboard interfaces contain an FIFO that is at least 16 bytes in depth. The FIFO in this example is 256 bytes, which is more than adequate for a keyboard interface. Note how the INC BYTE PTR CX:INP is used to add 1 to the input pointer and also make sure that it always addresses data in the queue.

This procedure first checks to see whether the FIFO is full. A full condition is indicated when the input pointer (INP) is one byte below the output pointer (OUTP). If the FIFO is full, the interrupt is disabled with a bit set/reset command to the 82C55, and a return from the interrupt occurs. If the FIFO is not full, the data are input from port A, and the input pointer is incremented before a return occurs.

Example 12–6 shows the procedure that removes data from the FIFO. This procedure first determines whether the FIFO is empty by comparing the two pointers. If the pointers are equal, the FIFO is empty, and the software waits at the EMPTY loop where it continuously tests the pointers. The EMPTY loop is interrupted by the keyboard interrupt, which stores data into the FIFO so that it is no longer empty. This procedure returns with the character in register AH.

EXAMPLE 12–6

```
;A procedure that reads data from the queue of Example 12-5
;and returns it in AH;

READQ   PROC   FAR USES EBX EDI EDX

        .REPEAT
            MOV  EBX,CS:INP          ;get pointers
            NOV  EDI,CS:OUTP
        .UNTIL EBX == EDI            ;while empty

        MOV   AH,CS:[EDI]            ;get data
        MOV   AL,9
        MOV   DX,CNTR
        OUT   DX,AL                  ;enable 52C55 interrupt
        INC   BYTE PTR CS:OUTP
        RET

READQ   ENDP
```

12–3   EXPANDING THE INTERRUPT STRUCTURE

This text covers three of the more common methods of expanding the interrupt structure of the microprocessor. In this section, we explain how, with software and some hardware modification of the circuit shown in Figure 12–10, it is possible to expand the INTR input so that it accepts seven interrupt inputs. We also explain how to "daisy-chain" interrupts by software polling. In the next section, we describe a third technique in which up to 63 interrupting inputs can be added by means of the 8259A programmable interrupt controller.

Using the 74ALS244 to Expand Interrupts

The modification shown in Figure 12–13 allows the circuit of Figure 12–10 to accommodate up to seven additional interrupt inputs. The only hardware change is the addition of an eight-input NAND gate, which provides the INTR signal to the microprocessor when any of the $\overline{IR}$ inputs becomes active.

Operation.   If any of the $\overline{IR}$ inputs becomes a logic 0, then the output of the NAND gate goes to a logic 1 and requests an interrupt through the INTR input. The interrupt vector that is fetched during the $\overline{INTA}$ pulse depends on which interrupt request line becomes active. Table 12–1 shows the interrupt vectors used by a single interrupt request input.

If two or more interrupt request inputs are simultaneously active, a new interrupt vector is generated. For example, if $\overline{IR1}$ and $\overline{IR0}$ are both active, the interrupt vector generated is FCH (252). Priority is resolved at this location. If the $\overline{IR0}$ input is to have the higher priority, the vector address for $\overline{IR0}$ is stored at vector location FCH. The entire top half of the vector table and its 128 interrupt vectors must be used to accommodate all possible conditions of these seven interrupt request inputs. This seems wasteful, but in many dedicated applications it is a cost-effective approach to interrupt expansion.
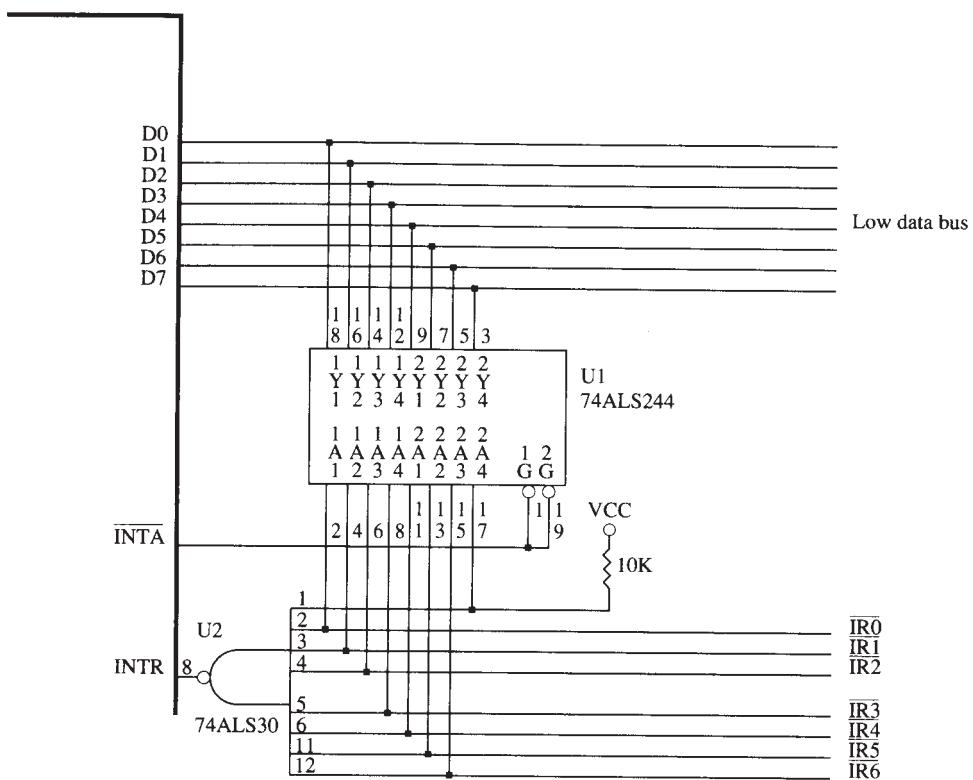
FIGURE 12–13   Expanding the INTR input from one to seven interrupt request lines.

TABLE 12–1   Single interrupt requests for Figure 12–13.

$\overline{IR6}$	$\overline{IR5}$	$\overline{IR4}$	$\overline{IR3}$	$\overline{IR2}$	$\overline{IR1}$	$\overline{IR0}$	*Vector*
1	1	1	1	1	1	0	FEH
1	1	1	1	1	0	1	FDH
1	1	1	1	0	1	1	FBH
1	1	1	0	1	1	1	F7H
1	1	0	1	1	1	1	EFH
1	0	1	1	1	1	1	DFH
0	1	1	1	1	1	1	BFH

Daisy-Chained Interrupt
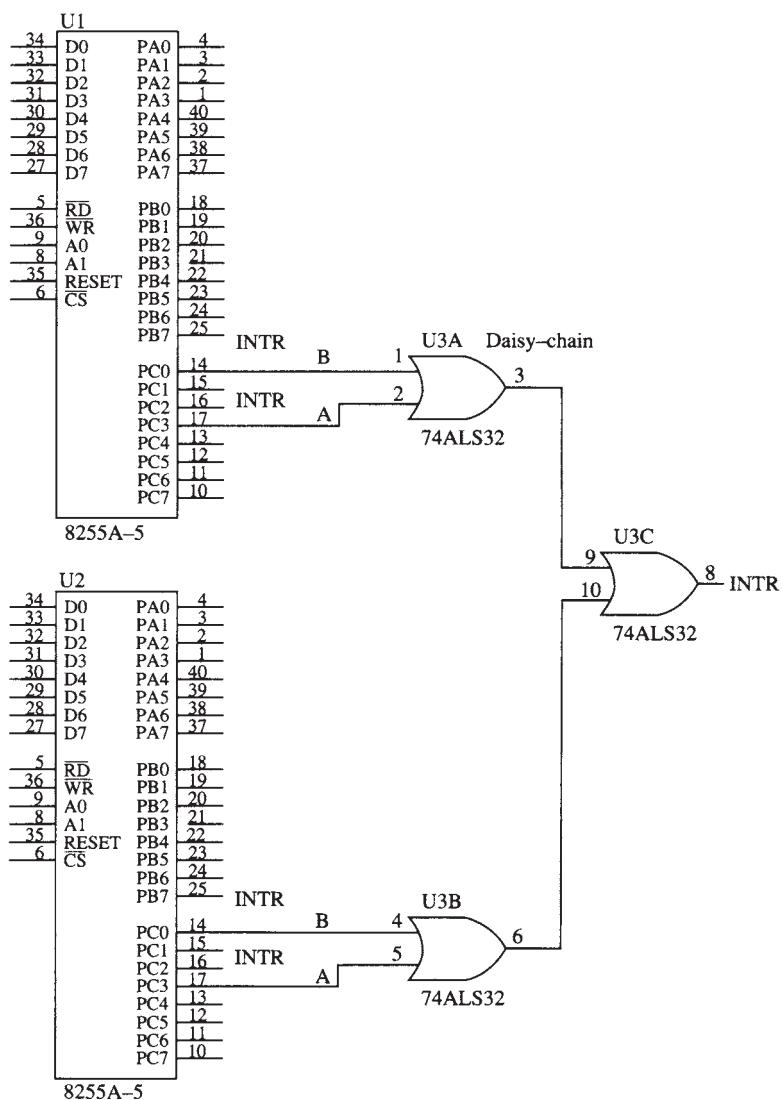
Expansion by means of a daisy-chained interrupt is in many ways better than using the 74ALS244 because it requires only one interrupt vector. The task of determining priority is left to the interrupt service procedure. Setting priority for a daisy-chain does require additional software execution time, but in general this is a much better approach to expanding the interrupt structure of the microprocessor.

Figure 12–14 illustrates a set of two 82C55 peripheral interfaces with their four INTR outputs daisy-chained and connected to the single INTR input of the microprocessor. If any interrupt output becomes a logic 1, so does the INTR input to the microprocessor causing an interrupt.

When a daisy-chain is used to request an interrupt, it is better to pull the data bus connections (D_0–D_7) high by using pull-up resistors so interrupt vector FFH is used for the chain. Any

FIGURE 12–14 Two 82C55 PIAs connected to the INTR outputs are daisy-chained to produce an INTR signal.

interrupt vector can be used to respond to a daisy-chain. In the circuit, any of the four INTR outputs from the two 82C55s will cause the INTR pin on the microprocessor to go high, requesting an interrupt.

When the INTR pin does go high with a daisy-chain, the hardware gives no direct indication as to which 82C55 or which INTR output caused the interrupt. The task of locating which INTR output became active is up to the interrupt service procedure, which must poll the 82C55s to determine which output caused the interrupt.

Example 12–7 illustrates the interrupt service procedure that responds to the daisy-chain interrupt request. The procedure polls each 82C55 and each INTR output to decide which interrupt service procedure to utilize.

EXAMPLE 12–7

```
;A procedure that services the daisy-chain interrupt
;of Figure 12-14.
```

```
C1      EQU    504H                    ;first 82C55
C2      EQU    604H                    ;second 82C55
MASK1   EQU    1                       ;INTRB
MASK2   EQU    8                       ;INTRA

POLL    PROC   FAR     USES EAX EDX

        MOV    DX,C1                   ;address first 82C55
        IN     AL,DX
        TEST   AL,MASK1                ;test INTRB
        .IF !ZERO?

                ;LEVEL 1 interrupt software here
        .ENDIF
        TEST   AL,MASK2                ;test INTRA
        .IF !ZERO?

                ;LEVEL 2 interrupt software here

        .ENDIF
        MOV    DX,C2                   ;address second 82C55
        TEST   AL,MASK1                ;test INTRB
        .IF !ZERO?

                ;LEVEL 3 interrupt software here

        .ENDIF

                ;LEVEL 4 interrupt software here

POLL    ENDP
```
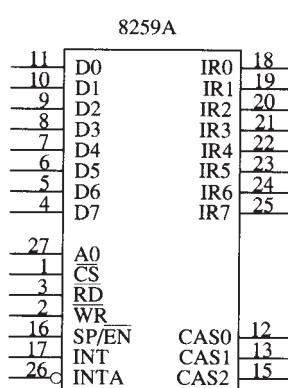
12–4   8259A PROGRAMMABLE INTERRUPT CONTROLLER

The 8259A programmable interrupt controller (PIC) adds eight vectored priority encoded inter-rupts to the microprocessor. This controller can be expanded, without additional hardware, to accept up to 64 interrupt requests. This expansion requires a master 8259A and eight 8259A slaves. A pair of these controllers still resides and is programmed as explained here in the latest chip sets from Intel and other manufacturers.

General Description of the 8259A

Figure 12–15 shows the pin-out of the 8259A. The 8259A is easy to connect to the microproces-sor because all of its pins are direct connections except the $\overline{CS}$ pin, which must be decoded, and

FIGURE 12–15   The pin-out of the 8259A programmable interrupt controller (PIC).

the $\overline{\text{WR}}$ pin, which must have an I/O bank write pulse. Following is a description of each pin on the 8259A:

$\mathbf{D_0}$–$\mathbf{D_7}$	The bidirectional **data connections** are normally connected to the data bus on the microprocessor.
$\mathbf{IR_0}$–$\mathbf{IR_7}$	**Interrupt request inputs** are used to request an interrupt and to connect to a slave in a system with multiple 8259As.
$\overline{\text{WR}}$	The **write input** connects to write strobe signal ($\overline{\text{IOWC}}$) on the microprocessor.
$\overline{\text{RD}}$	The **read input** connects to the $\overline{\text{IORC}}$ signal.
INT	The **interrupt output** connects to the INTR pin on the microprocessor from the master and is connected to a master IR pin on a slave.
$\overline{\text{INTA}}$	**Interrupt acknowledge** is an input that connects to the $\overline{\text{INTA}}$ signal on the system. In a system with a master and slaves, only the master $\overline{\text{INTA}}$ signal is connected.
$\mathbf{A_0}$	The $\mathbf{A_0}$ **address input** selects different command words within the 8259A.
$\overline{\text{CS}}$	**Chip select** enables the 8259A for programming and control.
$\text{SP}/\overline{\text{EN}}$	**Slave program/enable buffer** is a dual-function pin. When the 8259A is in buffered mode, this is an output that controls the data bus transceivers in a large microprocessor-based system. When the 8259A is not in the buffered mode, this pin programs the device as a master (1) or a slave (0).
$\mathbf{CAS_0}$–$\mathbf{CAS_2}$	The **cascade lines** are used as outputs from the master to the slaves for cascading multiple 8259As in a system.

Connecting a Single 8259A

Figure 12–16 shows a single 8259A connected to the microprocessor. Here the $\text{SP}/\overline{\text{EN}}$ pin is pulled high to indicate that it is a master. The 8259A is decoded at I/O ports 0400H and 0401H by the PLD (no program shown). Like other peripherals discussed in Chapter 11, the 8259A requires four wait states for it to function properly with a 16 MHz 80386SX and more for some other versions of the Intel microprocessor family.

Cascading Multiple 8259As
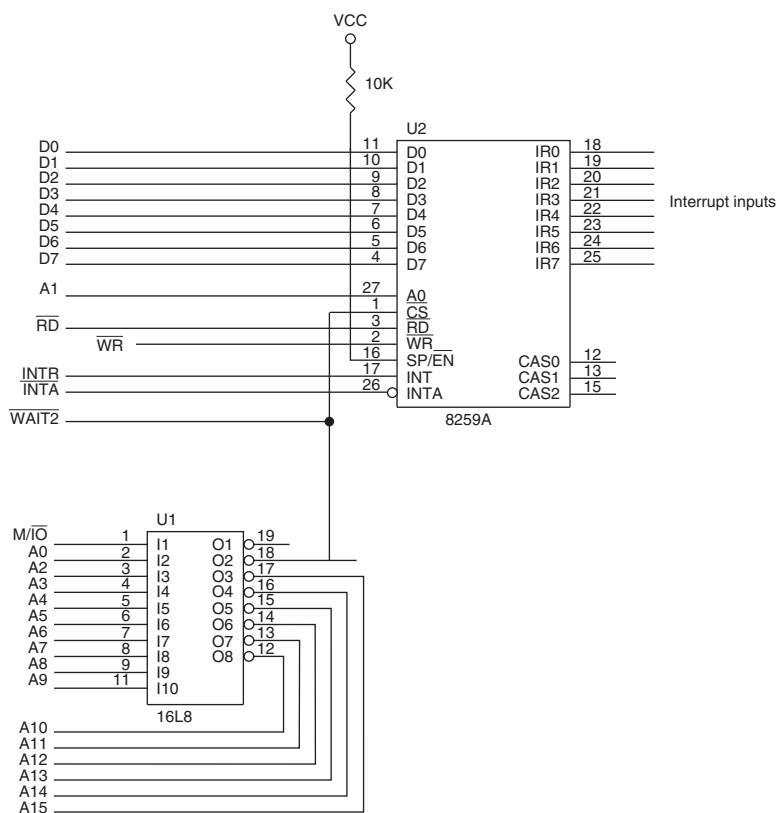
Figure 12–17 shows two 8259As connected to the microprocessor in a way that is often found in the ATX-style computer, which has two 8259As for interrupts. The XT- or PC-style computers use a single 8259A controller at interrupt vectors 08H–0FH. The ATX-style computer uses interrupt vector 0AH as a cascade input from a second 8259A located at vectors 70H through 77H. Appendix A contains a table that lists the functions of all the interrupt vectors used.

This circuit uses vectors 08H–0FH and I/O ports 0300H and 0302H for U1, the master; and vectors 70H–77H and I/O ports 0304H and 0306H for U2, the slave. Notice that we also include data bus buffers to illustrate the use of the $\text{SP}/\overline{\text{EN}}$ pin on the 8259A. These buffers are used only in very large systems that have many devices connected to their data bus connections. In practice, we seldom find these buffers.

Programming the 8259A

The 8259A is programmed by initialization and operation command words. **Initialization command words** (ICWs) are programmed before the 8259A is able to function in the system and dictate the basic operation of the 8259A. **Operation command words** (OCWs) are programmed during the normal course of operation. The OCWs control the operation of the 8259A.

FIGURE 12–16 An 8259A interfaced to the 8086 microprocessor.

Initialization Command Words. There are four initialization command words (ICWs) for the 8259A that are selected when the A_0 pin is a logic 1. When the 8259A is first powered up, it must be sent ICW_1, ICW_2, and ICW_4. If the 8259A is programmed in cascade mode by ICW_1, then we also must program ICW_3. So if a single 8259A is used in a system, ICW_1, ICW_2, and ICW_4 must be programmed. If cascade mode is used in a system, then all four ICWs must be programmed.

Refer to Figure 12–18 for the format of all four ICWs. The following is a description of each ICW:

ICW_1	Programs the basic operation of the 8259A. To program this ICW for 8086–Pentium 4 operation, place a logic 1 in bit IC_4. Bits AD_1, A_7, A_6, and A_5 are don't cares for microprocessor operation and only apply to the 8259A when used with an 8-bit 8085 microprocessor (not covered in this textbook). This ICW selects single or cascade operation by programming the SNGL bit. If cascade operation is selected, we must also program ICW_3. The LTIM bit determines whether the interrupt request inputs are positive edge-triggered or level-triggered.
ICW_2	Selects the vector number used with the interrupt request inputs. For example, if we decide to program the 8259A so it functions at vector locations 08H–0FH, we place 08H into this command word. Likewise, if we decide to program the 8259A for vectors 70H–77H, we place 70H in this ICW.
ICW_3	Only used when ICW_1 indicates that the system is operated in cascade mode. This ICW indicates where the slave is connected to the master. For example, in Figure 12–18 we connected a slave to IR_2. To program ICW_3
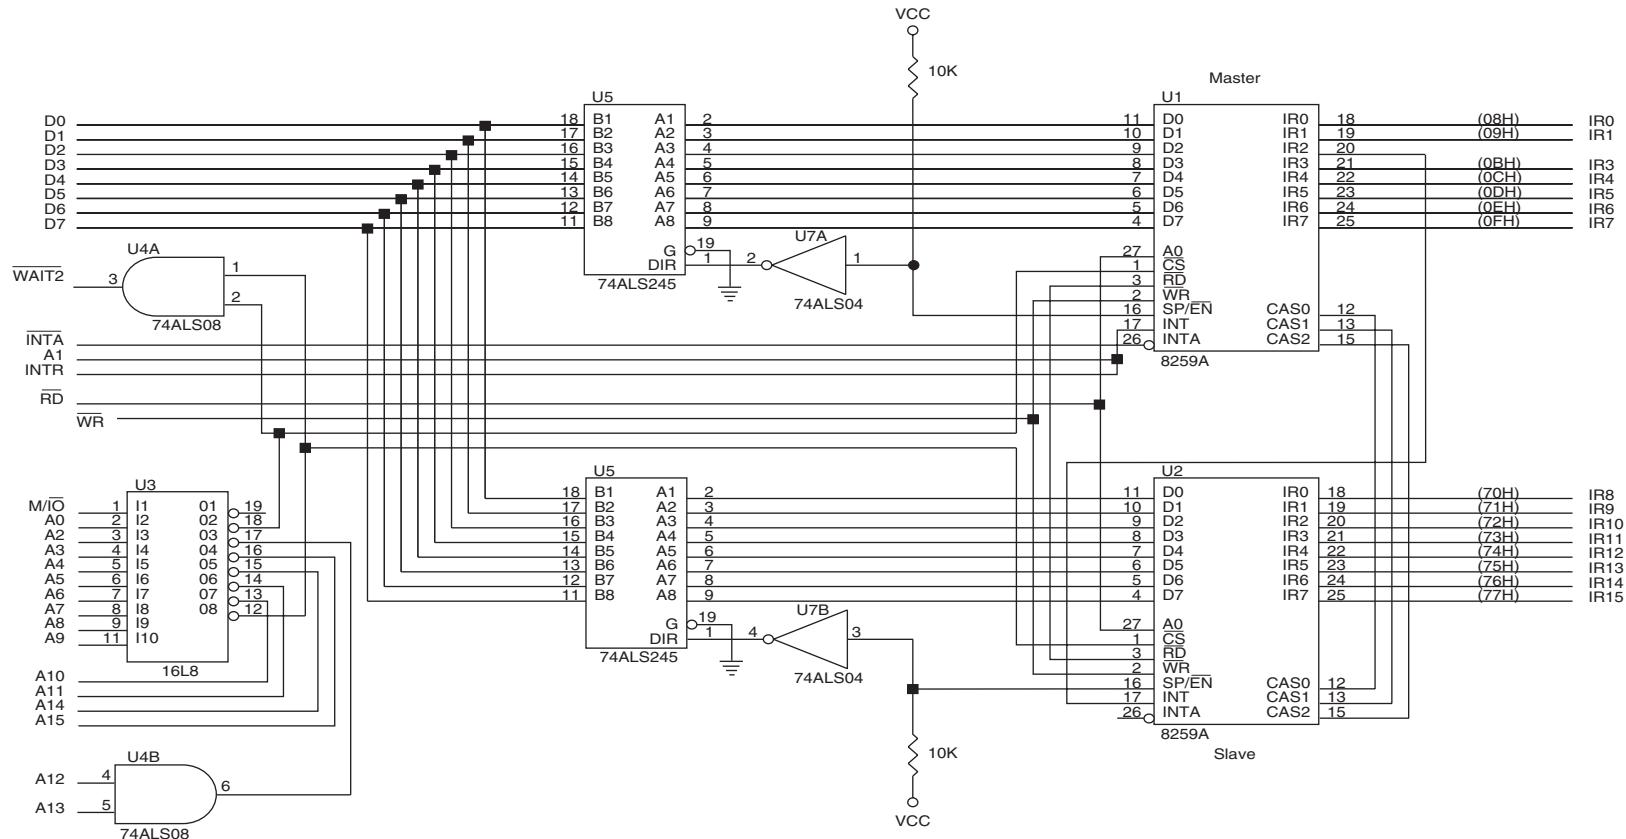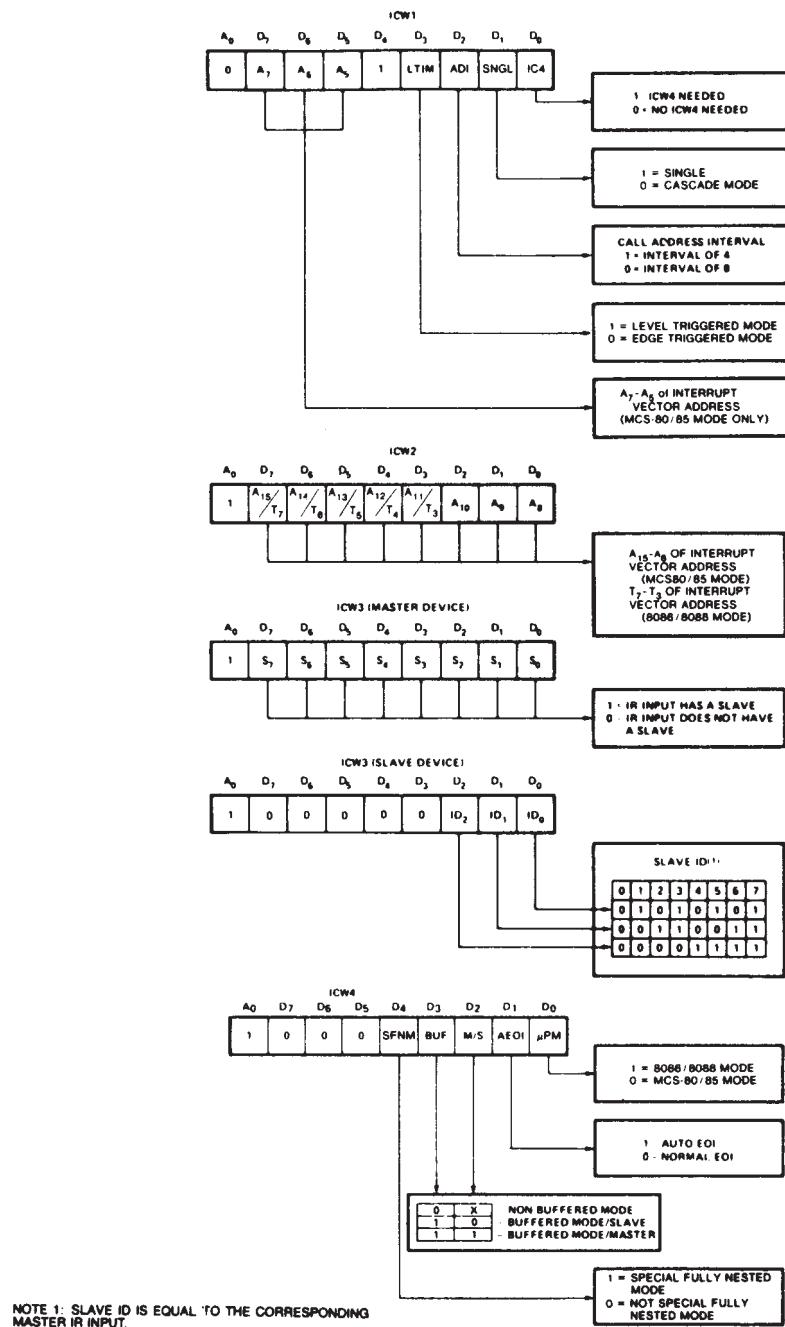
FIGURE 12–17   Two 8259As interfaced to the 8259A at I/O ports 0300H and 0302H for the master and 0304H and 0306H for the slave.

FIGURE 12–18 The 8259A initialization command words (ICWs). (Courtesy of Intel Corporation.)

for this connection, in both master and slave, we place 04H in ICW_3. Suppose we have two slaves connected to a master using IR_0 and IR_1. The master is programmed with an ICW_3 of 03H; one slave is programmed with an ICW_3 of 01H and the other with an ICW_3 of 02H.

ICW_4   Programmed for use with the 8086–Pentium 4 microprocessors, but is not programmed in a system that functions with the 8085 microprocessor. The

rightmost bit must be a logic 1 to select operation with the 8086–Pentium 4 microprocessors, and the remaining bits are programmed as follows:

SFNM—Selects the special fully nested mode of operation for the 8259A if a logic 1 is placed in this bit. This allows the highest priority interrupt request from a slave to be recognized by the master while it is processing another interrupt from a slave. Normally, only one interrupt request is processed at a time and others are ignored until the process is complete.

BUF and M/S—Buffered and master slave are used together to select buffered operation or nonbuffered operation for the 8259A as a master or a slave.

AEOI—Selects automatic or normal end of interrupt (discussed more fully under operation command words). The EOI commands of OCW2 are used only if the AEOI mode is not selected by ICW4. If AEOI is selected, the interrupt automatically resets the interrupt request bit and does not modify priority. This is the preferred mode of operation for the 8259A and reduces the length of the interrupt service procedure.

Operation Command Words. The operation command words (OCWs) are used to direct the operation of the 8259A once it is programmed with the ICW. The OCWs are selected when the A_0 pin is at a logic 0 level, except for OCW_1, which is selected when A_0 is a logic 1. Figure 12–19 lists the binary bit patterns for all three operation command words of the 8259A. Following is a list describing the function of each OCW:

OCW_1 Used to set and read the interrupt mask register. When a mask bit is set, it will turn off (mask) the corresponding interrupt input. The mask register is read when OCW_1 is read. Because the state of the mask bits is unknown when the 8259A is first initialized, OCW_1 must be programmed after programming the ICW upon initialization.

OCW_2 Programmed only when the AEOI mode is not selected for the 8259A. In this case, this OCW selects the way that the 8259A responds to an interrupt. The modes are listed as follows:

Nonspecific End-of-Interrupt—A command sent by the interrupt service procedure to signal the end of the interrupt. The 8259A automatically determines which interrupt level was active and resets the correct bit of the interrupt status register. Resetting the status bit allows the interrupt to take action again or a lower priority interrupt to take effect.

Specific End-of-Interrupt—A command that allows a specific interrupt request to be reset. The exact position is determined with bits L_2–L_0 of OCW_2.
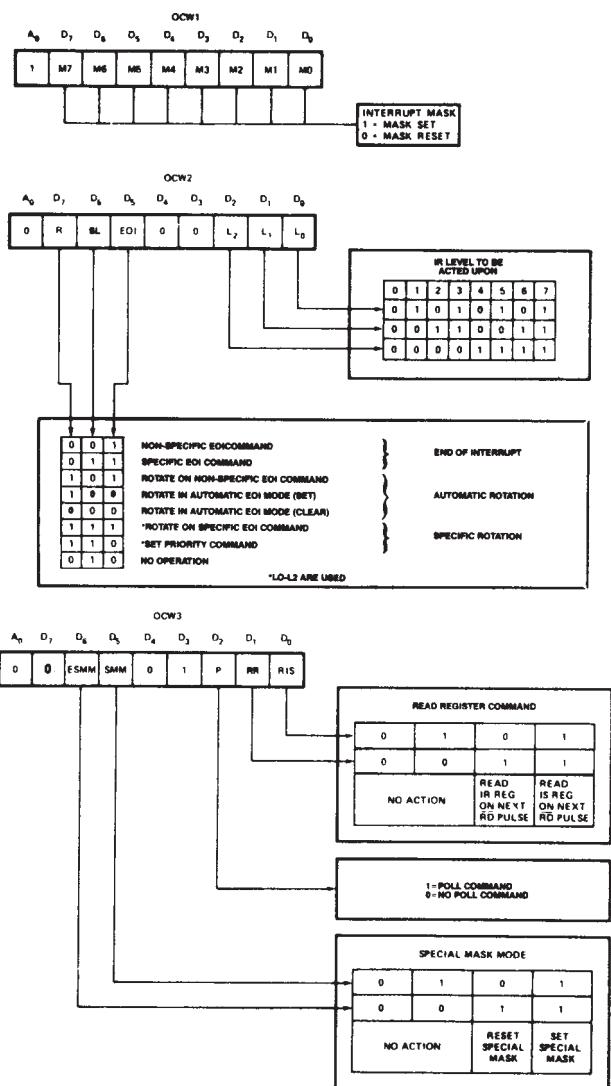
Rotate-on-Nonspecific EOI—A command that functions exactly like the Nonspecific End-of-Interrupt command, except that it rotates interrupt priorities after resetting the interrupt status register bit. The level reset by this command becomes the lowest priority interrupt. For example, if IR_4 was just serviced by this command, it becomes the lowest priority interrupt input and IR_5 becomes the highest priority.

Rotate-on-Automatic EOI—A command that selects automatic EOI with rotating priority. This command must only be sent to the 8259A once if this mode is desired. If this mode must be turned off, use the clear command.

Rotate-on-Specific EOI—Functions as the specific EOI, except that it selects rotating priority.

Set priority—Allows the programmer to set the lowest priority interrupt input using the L_2–L_0 bits.

FIGURE 12–19 The 8259A operation command words (OCWs). (Courtesy of Intel Corporation.)

FIGURE 12–19 The 8259A operation command words (OCWs). (Courtesy of Intel Corporation.)

OCW₃ Selects the register to be read, the operation of the special mask register, and the poll command. If polling is selected, the P bit must be set and then output to the 8259A. The next read operation will read the poll word. The rightmost three bits of the poll word indicate the active interrupt request with the highest priority. The leftmost bit indicates whether there is an interrupt and must be checked to determine whether the rightmost three bits contain valid information.

Status Register. Three status registers are readable in the 8259A: interrupt request register (IRR), in-service register (ISR), and interrupt mask register (IMR). (See Figure 12–20 for all
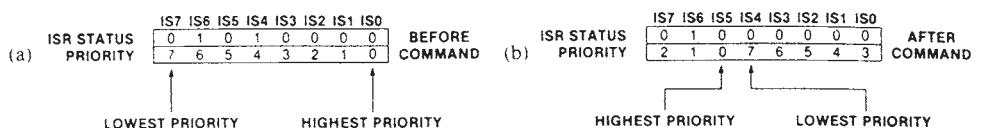
FIGURE 12–20 The 8259A in-service register (ISR). (a) Before IR₄ is accepted and (b) after IR₄ is accepted. (Courtesy of Intel Corporation.)

three status registers; they all have the same bit configuration.) The IRR is an 8-bit register that indicates which interrupt request inputs are active. The ISR is an 8-bit register that contains the level of the interrupt being serviced. The IMR is an 8-bit register that holds the interrupt mask bits and indicates which interrupts are masked off.

Both the IRR and ISR are read by programming OCW_3, and IMR is read through OCW_1. To read the IMR, $A_0 = 1$; to read either IRR or ISR, $A_0 = 0$. Bit positions D_0 and D_1 of OCW_3 select which register (IRR or ISR) is read when $A_0 = 0$.

8259A Programming Example

Figure 12–21 illustrates the 8259A programmable interrupt controller connected to a 16550 programmable communications controller. In this circuit, the INTR pin from the 16550 is connected to the programmable interrupt controller's interrupt request input IR_0. An IR_0 occurs whenever (1) the transmitter is ready to send another character, (2) the receiver has received a character, (3) an error is detected while receiving data, and (4) a modem interrupt occurs. Notice that the 16550
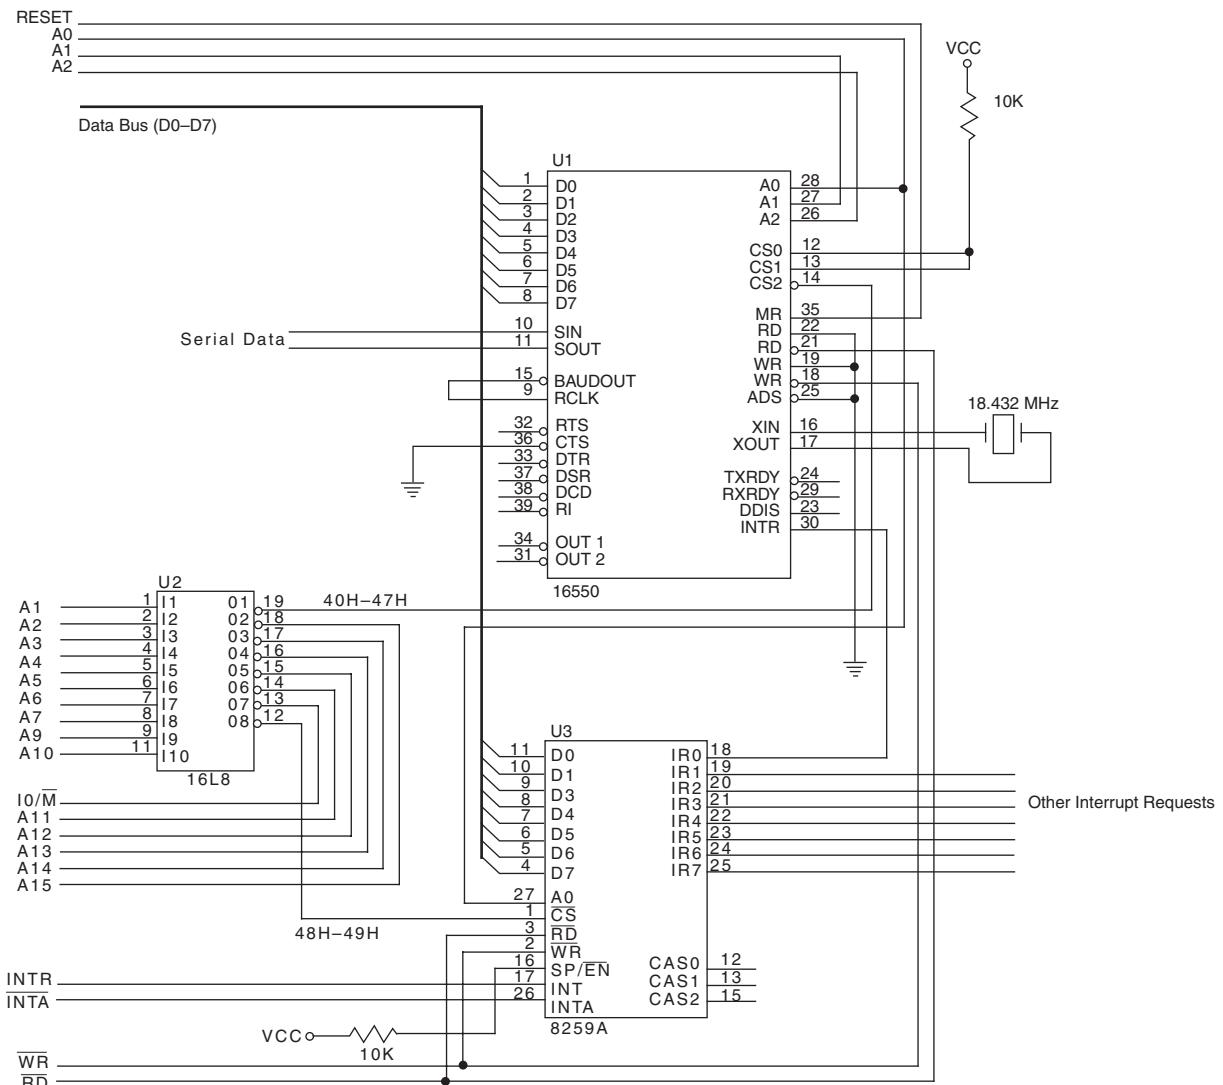
FIGURE 12–21 The 16550 UART interfaced to the 8088 microprocessor through the 8259A.

is decoded at I/O ports 40H and 47H, and the 8259A is decoded at 8-bit I/O ports 48H and 49H. Both devices are interfaced to the data bus of an 8088 microprocessor.

Initialization Software.   The first portion of the software for this system must program both the 16550 and the 8259A, and then enable the INTR pin on the 8088 so that interrupts can take effect. Example 12–8 lists the software required to program both devices and enable INTR. This software uses two memory FIFOs that hold data for the transmitter and for the receiver. Each memory FIFO is 16K bytes long and is addressed by a pair of pointers (input and output).

EXAMPLE 12–8

```
;Initialization software for the 16650 and 8259A
;of the circuit in Figure 12-21

PIC1    EQU  48H                    ;8259A control A0 = 0
PIC2    EQU  49H                    ;8259A control A0 = 1
ICW1    EQU  1BH                    ;8259A ICW1
ICW2    EQU  80H                    ;8259A ICW2
ICW4    EQU  3                      ;8259A ICW4
OCW1    EQU  0FEH                   ;8259A OCW1
LINE    EQU  43H                    ;16650 line register
LSB     EQU  40H                    ;16650 baud divisor LSB
MSB     EQU  41H                    ;16650 baud divisor MSB
FIFO    EQU  42H                    ;16650 FIFO register
ITR     EQU  41H                    ;16650 interrupt register

INIT  PROC  NEAR

;
;setup 16650
;
        MOV  AL,10001010B           ;enable baud rate divisor
        OUT  LINE,AL

        MOV  AL,120                 ;program baud 9600
        OUT  LSB,AL
        MOV  AL,0
        OUT  MSB,AL

        MOV  AL,00001010B           ;program 7 data, odd
        OUT  LINE,AL                ;parity, 1 stop

        MOV  AL,00000111B           ;enable transmitter and
        OUT  FIFO,AL                ;receiver
;
;program 8259A
;
        MOV  AL,ICW1                ;program ICW1
        OUT  PIC1,AL
        MOV  AL,ICW2                ;program ICW2
        OUT  PIC2,AL
        MOV  AL,ICW4                ;program ICW4
        OUT  PIC2,AL
        MOV  AL,OCW1                ;program OCW1
        OUT  PIC2,AL
        STI                         ;enable INTR pin
;
;enable 16650 interrupts
;
        MOV  AL,5
        OUT  ITR,AL                 ;enable interrupts
        RET

INIT    ENDP
```

The first portion of the procedure (INIT) programs the 16550 UART for operation with seven data bits, odd parity, one stop bit, and a baud rate clock of 9600. The FIFO control register also enables both the transmitter and receiver.

The second part of the procedure programs the 8259A, with its three ICWs and one OCW. The 8259A is set up so that it functions at interrupt vectors 80H–87H and operates with automatic EOI. The OCW enables the interrupt for the 16550 UART. The INTR pin of the microprocessor is also enabled by using the STI instruction.

The final part of the software enables the receiver and error interrupts of the 16550 UART through the interrupt control register. The transmitter interrupt is not enabled until data are available for transmission. See Figure 12–22 for the contents of the interrupt control register of the 16550 UART. Notice that the control register can enable or disable the receiver, transmitter, line status (error), and modem interrupts.

Handling the 16550 UART Interrupt Request. Because the 16550 generates only one interrupt request for various interrupts, the interrupt handler must poll the 16550 to determine what type of interrupt has occurred. This is accomplished by examining the interrupt identification register (see Figure 12–23). Note that the interrupt identification register (read-only) shares the same I/O port as the FIFO control register (write-only).

The interrupt identification register indicates whether an interrupt is pending, the type of interrupt, and whether the transmitter and receiver FIFO memories are enabled. See Table 12–2 for the contents of the interrupt control bits.

The interrupt service procedure must examine the contents of the interrupt identification register to determine what event caused the interrupt and pass control to the appropriate procedure for the event. Example 12–9 shows the first part of an interrupt handler that passes control to RECV for a receiver data interrupt, TRANS for a transmitter data interrupt, and ERR for a line status error interrupt. Note that the modem status is not tested in this example.

FIGURE 12–22 The 16550 interrupt control register.

Interrupt Control Register

7	6	5	4	3	2	1	0
0	0	0	0	EM	EL	ET	ER

Enable Receiver Interrupt
0 = disabled
1 = enabled

Enable Transmitter Interrupt
0 = disabled
1 = enabled

Enable Line Interrupt
0 = disabled
1 = enabled

Enable Modem Interrupt
0 = disabled
1 = enabled
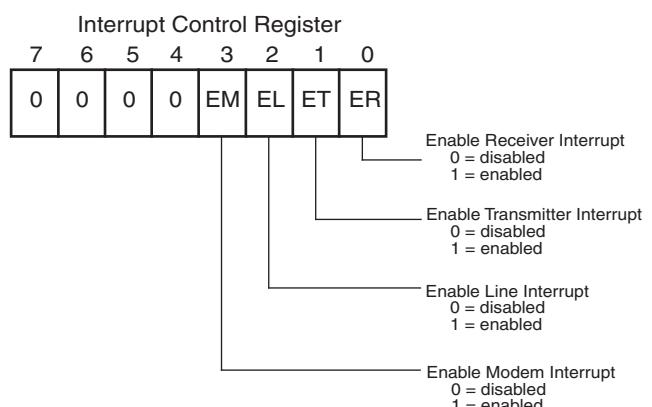
FIGURE 12–23 The 16550 interrupt identification register.

Interrupt Identification Register

7	6	5	4	3	2	1	0
0	0	0	0	ID	ID	ID	PN

Interrupt Pending
0 = interrupt pending
1 = no interrupt

Interrupt Identification Bits
(see Table 12–2)
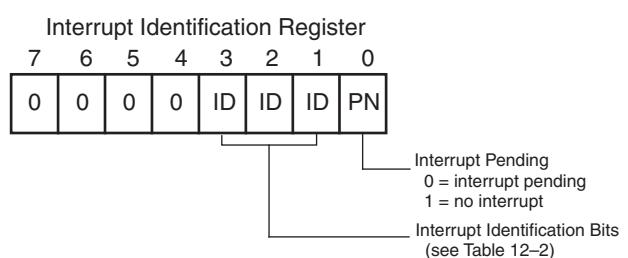
TABLE 12–2   The interrupt control bits of the 16650.

Bit 3	Bit 2	Bit 1	Bit 0	Priority	Type	Reset Control
0	0	0	1	—	No interrupt	—
0	1	1	0	1	Receiver error (parity, framing, overrun, or break)	Reset by reading the register
0	1	0	0	2	Receiver data available	Reset by reading the data
1	1	0	0	2	Character time-out, nothing has been removed from the receiver FIFO for at least four character times	Reset by reading the data
0	0	1	0	3	Transmitter empty	Reset by writing the transmitter
0	0	0	0	4	Modem status	Reset by reading the modem status

EXAMPLE 12–9

```
;Interrupt handler for the 16650 UART of Figure 12-21

INT80 PROC  FAR USES AX BX DI SI

        IN   AL,42H                  ;read interrupt ID
        .IF  AL == 6

                ;handle receiver error

        .ELSEIF AL == 2

                ;handle transmitter empty

                JMP    TRAN    ;example 12-13

        .ELSEIF AL == 4

                ;handle receiver ready

                JMP     RECV   ;example 12-11

            .ENDIF

            IRET

INT80   ENDP
```

Receiving data from the 16550 requires two procedures. One procedure reads the data register of the 16550 each time that the INTR pin requests an interrupt and stores it into the memory FIFO. The other procedure reads data from the memory FIFO from the main program.

Example 12–10 lists the procedure used to read data from the memory FIFO from the main program. This procedure assumes that the pointers (IIN and IOUT) are initialized in the initialization dialog for the system (not shown). The READ procedure returns with AL containing a character read from the memory FIFO. If the memory FIFO is empty, the procedure returns with the carry flag bit set to a logic 1. If AL contains a valid character, the carry flag bit is cleared upon return from READ.

Notice how the FIFO is reused by changing the address from the top of the FIFO to the bottom whenever it exceeds the start of the FIFO plus 16K. Notice that interrupts are enabled at the end of this procedure, in case they are disabled by a full memory FIFO condition by the RECV interrupt procedure.

EXAMPLE 12–10

```
;A procedure that reads one character from the FIFO
;and returns it on AL. If the FIFO is empty the return
;occurs with carry = 1.

READC   PROC    NEAR USES BX DX

        MOV   DI,IOUT              ;get pointer
        MOV   BX,IIN
        .IF   BX == DI             ;if empty
             STC                   ;set carry
          .ELSE                    ;if not empty
             MOV  AL,ES:[DI]       ;get data
             INC  DI               ;increment pointer
             .IF DI == OFFSET FIFO+16*1024
                  MOV  DI,OFFSET FIFO
             .ENDIF
             MOV IOUT,DI
             CLC
        .ENDIF
        PUSHF                      ;enable receiver
        IN    AL,41H
        OR    AL,5
        OUT   41H,AL
        POPF
        RET

READC      ENDP
```

Example 12–11 lists the RECV interrupt service procedure that is called each time the 16550 receives a character for the microprocessor. In this example, the interrupt uses vector type number 80H, which must address the interrupt handler of Example 12–9. Each time that this interrupt occurs, the REVC procedure is accessed by the interrupt handler reading a character from the 16550. The RECV procedure stores the character into the memory FIFO. If the memory FIFO is full, the receiver interrupt is disabled by the interrupt control register within the 16550. This may result in lost data, but at least it will not cause the interrupt to overrun valid data already stored in the memory FIFO. Any error conditions detected by the 8251A store a ? (3FH) in the memory FIFO. Note that errors are detected by the ERR portion of the interrupt handler (not shown).

EXAMPLE 12–11

```
;RECV portion of the interrupt handler of Example 12-9

RECV:
     MOV   BX,IOUT                ;get pointers
     MOV   DI,IIN
     MOV   SI,DI
     INC   SI
     .IF   SI == OFFSET FIFO+16*1024
          MOV   SI,OFFSET FIFO
     .ENDIF
     .IF SI == BX                 ;if FIFO full
          IN  AL,41H              ;disable receiver
          AND AL,0FAH
          OUT 41H,AL
     .ENDIF
     IN  AL,40H                   ;read data
     STOSB
     MOV  IIN,SI
     MOV  AL,20H                  ;8259A EOI command
     OUT  49H,AL
     IRET
```

Transmitting Data to the 16550. Data are transmitted to the 16550 in much the same manner as they are received, except that the interrupt service procedure removes transmit data from a second 16K-byte memory FIFO.

Example 12–12 lists the procedure that fills the output FIFO. It is similar to the procedure listed in Example 12–10, except it determines whether the FIFO is full instead of empty.

EXAMPLE 12–12

```
;A procedure that places data into the memory FIFO for
;transmission by the transmitter interrupt. AL = the
;character transmitted.

SAVEC  PROC  NEAR USES BX DI SI

      MOV  SI,OIN                ;load pointers
      MOV  BX,OOUT
      MOV  DI,SI
      INC  SI
      .IF SI == OFFSET OFIFO+16*1024
            MOV  SI,OFFSET OFIFO
      .ENDIF
      .IF BX == SI               ;if OFIFO full
            STC
      .ELSE
            STOSB
            MOV  OIN,SI
            CLC
      .ENDIF
      PUSHF
      IN   AL,41H                ;enable transmitter
      OR   AL,1
      OUT  41H,AL
      RET

      SAVEC  ENDP
```

Example 12–13 lists the interrupt service subroutine for the 16550 UART transmitter. This procedure is a continuation of the interrupt handler presented in Example 12–9 and is similar to the RECV procedure of Example 12–11, except that it determines whether the FIFO is empty rather than full. Note that we do not include an interrupt service procedure for the break interrupt or any errors.

EXAMPLE 12–13

```
;Interrupt service for the 16650 transmitter

TRAN:
      MOV  BX,OIN               ;load pointers
      MOV  DI,OOUT
      .IF BX == DI              ;if empty
            IN   AL,41H
            AND  AL,0FDH        ;disable transmitter
            OUT  41H,AL
      .ELSE                     ;if not empty
            MOV  AL,ES:[DI]
            OUT  40H,AL         ;send data
            INC  DI
            .IF DI == OFFSET OFIFO+16*1024
                  MOV  DI,OFFSET OFIFO
            .ENDIF
            MOV  OFIFO,DI
      .ENDIF
      MOV  AL,20H               ;send EOI to 8259A
      OUT  49H,AL
      IRET
```

The 16550 also contains a scratch register, which is a general-purpose register that can be used in any way deemed necessary by the programmer. Also contained in the 16550 are a modem control register and a modem status register. These registers allow the modem to cause interrupt and control the operation of the 16550 with a modem. See Figure 12–24 for the contents of both the modem status register and the modem control register.

The modem control register uses bit positions 0–3 to control various pins on the 16550. Bit position 4 enables the internal loop-back test for testing purposes. The modem status register allows the status of the modem pins to be tested; it also allows the modem pins to be checked for a change or, in the case of $\overline{RI}$, a trailing edge.

Figure 12–25 illustrates the 16550 UART, connected to an RS-232C interface that is often used to control a modem. Included in this interface are line driver and receiver circuits used to convert between TTL levels on the 16550 to RS-232C levels found on the interface. Note that RS-232C levels are usually +12 V for a logic 0 and −12 V for a logic 1 level.

In order to transmit or receive data through the modem, the $\overline{DTR}$ pin is activated (logic 0) and the UART then waits for the $\overline{DSR}$ pin to become a logic 0 from the modem, indicating that the modem is ready. Once this handshake is complete, the UART sends the modem a logic 0 on the $\overline{RTS}$ pin. When the modem is ready, it returns the $\overline{CTS}$ signal (logic 0) to the UART. Communications can now commence. The $\overline{DCD}$ signal from the modem is an indication that the modem has detected a carrier. This signal must also be tested before communications can begin.

12–5 INTERRUPT EXAMPLES

This section of the text presents a real-time clock and an interrupt-processed keyboard as examples of interrupt applications. A real-time (RTC) clock keeps time in real time—that is, in hours and minutes. It is also used for precision time delays. The example illustrated here keeps time in
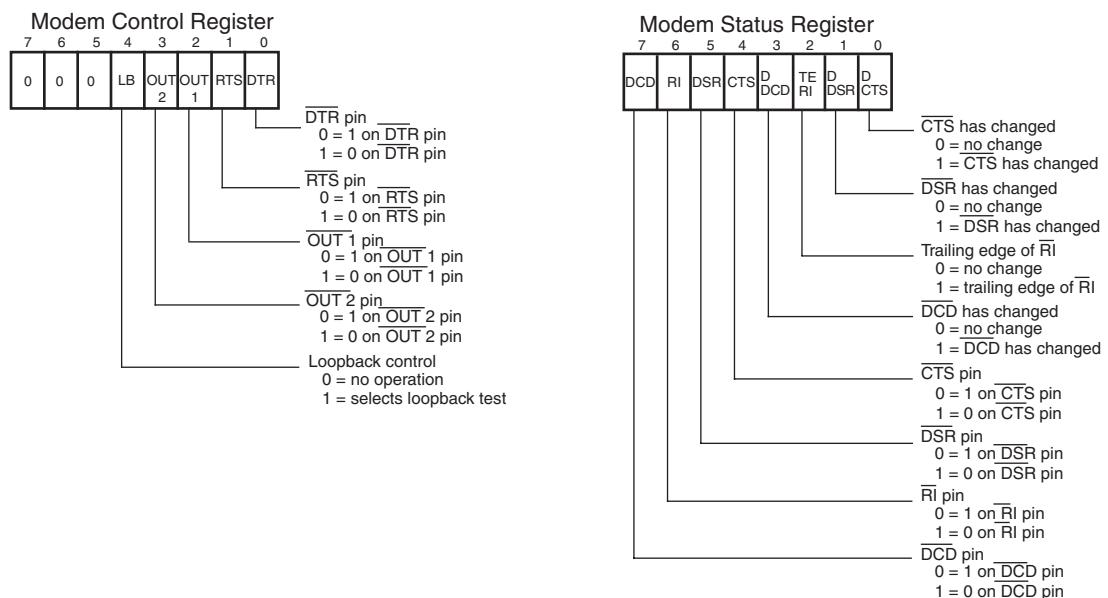
FIGURE 12–24 The 16550 modem control and modem status registers.
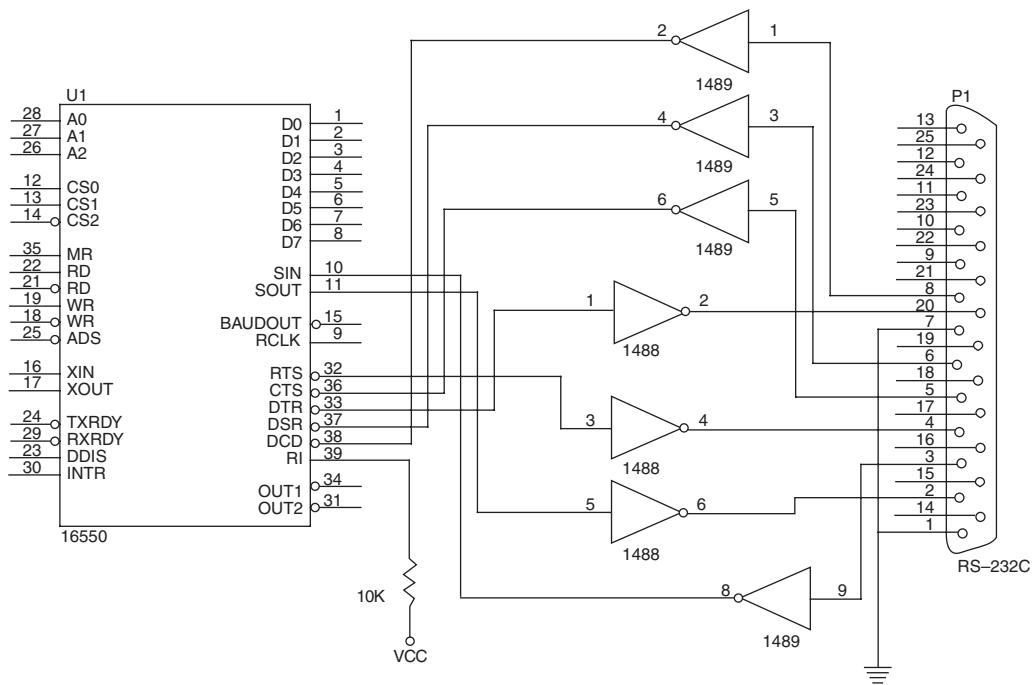
FIGURE 12–25   The 16550 interfaced to an RS-232C using 1488 line drivers and 1489 line receivers.
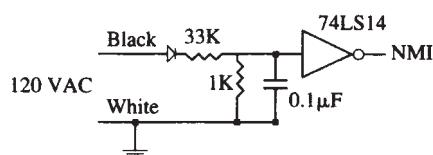
hours, minutes, seconds, and l/60 second, using four memory locations to hold the BCD time of day. The interrupt-processed keyboard uses a periodic interrupt to scan through the keys of the keyboard.

Real-Time Clock

Figure 12–26 illustrates a simple circuit that uses the 60 Hz AC power line to generate a periodic interrupt request signal for the NMI interrupt input pin. Although we are using a signal from the AC power line, which varies slightly in frequency from time to time, it is accurate over a period of time as mandated by the Federal Trade Commission (FTC).

The circuit uses a signal from the 120 V AC power line that is conditioned by a Schmitt trigger inverter before it is applied to the NMI interrupt input. Note that you must make certain that the power line ground is connected to the system ground in this schematic. The power line neutral (white wire) connection is the wide flat pin on the power line. The narrow flat pin is the hot (black wire) side or 120 V AC side of the line.

FIGURE 12–26   Converting the AC power line to a 60 Hz TTL signal for the NMI input.

The software for the real-time clock contains an interrupt service procedure that is called 60 times per second and a procedure that updates the count located in four memory locations. Example 12–14 lists both procedures, along with the four bytes of memory used to hold the BCD time of day. The memory locations for the TIME are stored somewhere in the system memory at the segment address (SEGMENT) and at the offset address TIME, which is first loaded in the TIMEP procedure. The lookup table (LOOK) for the modulus or each counter is stored in the code segment with the procedure.

EXAMPLE 12–14

```
TIME    DB    ?               ;1/60 sec counter (÷60)
        DB    ?               ;second counter (÷60)
        DB    ?               ;minute counter (÷60)
        DB    ?               ;hour counter (÷24)

LOOK    DB    60H, 60H, 60H, 24H

TIMEP   PROC  FAR USES AX BX DS

        MOV  AX,SEGMENT            ;load segment address of TIME
        MOV  DS,AX
        MOV  BX,0                  ;initialize pointer

        .REPEAT                    ;crank clock
            MOV  AL,DS:TIME[BX]
            ADD  AL,1              ;increment count
            DAA                    ;adjust for BCD
            .IF AL == BYTE PTR CS:LOOK[BX]
                  MOV  AL,0
            .ENDIF
            MOV DS:TIME[BX],AL
            INC BX
        .UNTIL AL != 0 || BX == 4
        IRET

TIMEP   ENDP
```

Another way to handle time is to use a single counter to store the time in memory and then determine the actual time with software. For example, the time can be stored in one single 32-bit counter (there are 5,184,000 1/60 sec in a day). In a counter such as this, a count of 0 is 12:00:00:00 AM and a count of 5,183,999 is 11:59:59:59 PM. Example 12–15 shows the interrupt procedure for this type of RTC, which requires the least of amount of time to execute.

EXAMPLE 12–15

```
TIME    DD    ?                   ;modulus 5,184,000 counter

TIMEP   PROC  FAR USES EAX

        MOV  AX,SEGMENT
        MOV  DS,AX

        INC  DS:TIME
        .IF DS:TIME == 5184000
              MOV DWORD PTR DS:TIME,0
        .ENDIF
        IRET

TIMEP   ENDP
```

Software to convert the count in the modulus 5,184,000 counter into hours, minutes, and seconds appears in Example 12–16. The procedure returns with the number of hours (0–23) in BL, number of minutes in BH, and number of seconds in AL. No attempt was made to retrieve the 1/60 second count.

EXAMPLE 12–16

```
;Time is returned as BL = hours, BH = minutes and AL = seconds

GETT    PROC    NEAR ECX EDX

        MOV  ECX,216000      ;divide by 216,000
        MOV  EAX,TIME
        SUB  EDX,EDX         ;clear EDX
        DIV  ECX             ;get hours
        MOV  BL,AL
        MOV  EAX,EDX
        MOV  ECX,3600        ;divide by 3600
        DIV  ECX             ;get minutes
        MOV  BH,AL
        SUB  EAX,EDX
        MOV  ECX,60          ;divide by 60
        DIV  ECX
        RET

GETT    ENDP
```

Suppose a time delay is needed. Time delays can be achieved using the RTC in Example 12–15 for any amount from 1/60 of a second to 24 hours. Example 12–17 shows a procedure that uses the RTC to perform time delays of the number of seconds passed to the procedure in the EAX register. This can be 1 second to an entire day's worth of seconds. It has an accuracy to within 1/60 second, the resolution of the RTC.

EXAMPLE 12–17

```
SEC    PROC    NEAR USES EAX EDX

       MOV   EDX,60
       MUL   EDX               ;get seconds as 1/60s count
       ADD   EAX,TIME          ;advance the TIME in EAX
       .IF   EAX >= 51840000
             SUB  EAX,5184000
       .ENDIF
       .REPEAT                 ;wait for TIME to catch up
       .UNTIL EAX == TIME
       RET

SEC    ENDP
```
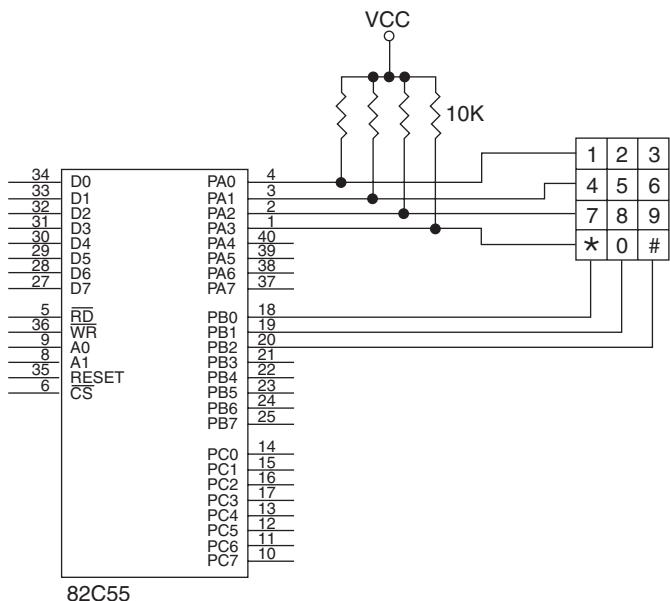
Interrupt-Processed Keyboard

The interrupt-processed keyboard scans through the keys on a keyboard through a periodic interrupt. Each time the interrupt occurs, the interrupt-service procedure tests for a key or debounces the key. Once a valid key is detected, the interrupt service procedure stores the key code into a keyboard queue for later reading by the system. The basis for this system is a periodic interrupt that can be caused by a timer, RTC, or other device in the system. Note that most systems already have a periodic interrupt for the real-time clock. In this example, we assume the interrupt calls the interrupt service procedure every 10 ms or, if the RTC is used with a 60 Hz clock, every 16.7 ms.

Figure 12–27 shows the keyboard interfaced to an 82C55. It does not show the timer or other circuitry required to call the interrupt once in every 10 ms or 16.7 ms. (Not shown in the software is programming of the 82C55.) The 82C55 must be programmed so that port A is an input port, port B is an output port, and the initialization software must store 00H at port B. This interfaces uses memory that is stored in the code segment for a queue and a few bytes that keep track of the keyboard scanning. Example 12–18 lists the interrupt service procedure for the keyboard.

FIGURE 12–27 A telephone-style keypad interfaced to the 82C55.

EXAMPLE 12–18

```
;Interrupt procedure for the keyboard of Figure 12-27

PORTA   EQU     1000H
PORTB   EQU     1001H

DBCNT   DB      0                       ;de-bounce counter
DBF     DB      0                       ;de-bounce flag
PNTR    DW      QUEUE                   ;input pointer to queue
OPNTR   DW      QUEUE                   ;output pointer to queue
QUEUE   DB      16 DUP(?)               ;16 byte queue

INTK    PROC    FAR USES AX BX DX

        MOV  DX,PORTA                 ;test for a key
        IN   AL,DX
        OR   AL,0F0H
        .IF AL != 0FFH                ;if key down
             INC   DBCNT              ;increment de-bounce count
             .IF DBCNT == 3           ;if key down for > 20 ms
                  DEC   DBCNT
                  .IF  DBF == 0
                       MOV  DBF,1
                       MOV  BX,00FEH
                       .WHILE 1       ;find key
                               MOV  AL,BL
                               MOV  DX,PORTB
                               OUT  DX,AL
                               ROL  BL,1
                               MOV  DX,PORTA
                               IN   AL,DX
                               OR   AL,0F0H
                               .BREAK .IF AL != 0
                               ADD  BH,4
                       .ENDW
                       MOV  BL,AL
                       MOV  AL,0
                       MOV  DX,PORTB
                       OUT  DX,AL
                       DEC  BH
                       .REPEAT
```

```
                                        SHR   BL,1
                                        INC   BH
                              .UNTIL !CARRY?
                              MOV  AL,BH
                              MOV  BX,PNTR
                              MOV  [BX],AL        ;key code to queue
                              INC    BX
                              .IF BX == OFFSET QUEUE+16
                                     MOV  DX,OFFSET QUEUE
                              .ENDIF
                              MOV  PNTR,BX
                        .ENDIF
                  .ENDIF
          .ELSE                          ;if no key down
                  DEC  DBCNT              ;decrement de-bounce count
                  .IF SIGN?               ;if below zero
                         MOV  DBCNT,0
                         MOV  DBF,0
                  .ENDIF
          .ENDIF
          IRET
INTK   ENDP
```

The keyboard-interrupt finds the key and stores the key code in the queue. The code stored in the queue is a raw code that does not indicate the key number. For example, the key code for the 1-key is 00H, the key code for the 4-key is 01H, and so on. There is no provision for a queue overflow in this software. It could be added, but in almost all cases it is difficult to out-type a 16-byte queue.

Example 12–19 illustrates a procedure that removes data from the keyboard queue. This procedure is not interrupt-driven and is called only when information from the keyboard is needed in a program. Example 12–20 shows the caller software for the key procedure.

EXAMPLE 12–19

```
LOOK   DB    1,4,7,10     ;lookup table
       DB    2,5,8,0
       DB    3,6,9,11

KEY    PROC  NEAR USES BX

   MOV     BX,OPNTR
   .IF BX == PNTR          ;if queue empty
          STC
   .ELSE
          MOV  AL,[BX]     ;get queue data
          INC  BX
          .IF BX == OFFSET QUEUE+16
                  MOV  BX,OFFSET QUEUE
          .ENDIF
          MOV  OPNTR,BX
          MOV  BX,LOOK
          XLAT
          CLC
   .ENDIF
   RET

KEY        ENDP
```

EXAMPLE 12–20

```
.REPEAT
       CALL KEY
.UNTIL !CARRY?
```

12–6   SUMMARY

1. An interrupt is a hardware- or software-initiated call that interrupts the currently executing program at any point and calls a procedure. The procedure is called by the interrupt handler or an interrupt service procedure.
2. Interrupts are useful when an I/O device needs to be serviced only occasionally at low data transfer rates.
3. The microprocessor has five instructions that apply to interrupts: BOUND, INT, INT 3, INTO, and IRET. The INT and INT 3 instructions call procedures with addresses stored in the interrupt vector whose type is indicated by the instruction. The BOUND instruction is a conditional interrupt that uses interrupt vector type number 5. The INTO instruction is a conditional interrupt that interrupts a program only if the overflow flag is set. Finally, the IRET, IRETD, or IRETQ instruction is used to return from interrupt service procedures.
4. The microprocessor has three pins that apply to its hardware interrupt structure: INTR, NMI, and $\overline{INTA}$. The interrupt inputs are INTR and NMI, which are used to request interrupts, and $\overline{INTA}$, an output used to acknowledge the INTR interrupt request.
5. Real mode interrupts are referenced through a vector table that occupies memory locations 0000H–03FFH. Each interrupt vector is four bytes long and contains the offset and segment addresses of the interrupt service procedure. In protected mode, the interrupts reference the interrupt descriptor table (IDT) that contains 256 interrupt descriptors. Each interrupt descriptor contains a segment selector and a 32-bit offset address.
6. Two flag bits are used with the interrupt structure of the microprocessor: trap (TF) and interrupt enable (IF). The IF flag bit enables the INTR interrupt input, and the TF flag bit causes interrupts to occur after the execution of each instruction, as long as TF is active.
7. The first 32 interrupt vector locations are reserved for Intel use, with many predefined in the microprocessor. The last 224 interrupt vectors are for the user's use and can perform any function desired.
8. Whenever an interrupt is detected, the following events occur: (1) the flags are pushed onto the stack, (2) the IF and TF flag bits are both cleared, (3) the IP and CS registers are both pushed onto the stack, and (4) the interrupt vector is fetched from the interrupt vector table and the interrupt service subroutine is accessed through the vector address.
9. Tracing or single-stepping is accomplished by setting the TF flag bit. This causes an interrupt to occur after the execution of each instruction for debugging.
10. The non-maskable interrupt input (NMI) calls the procedure whose address is stored at interrupt vector type number 2. This input is positive edge-triggered.
11. The INTR pin is not internally decoded, as is the NMI pin. Instead, $\overline{INTA}$ is used to apply the interrupt vector type number to data bus connections D_0–D_7 during the $\overline{INTA}$ pulse.
12. Methods of applying the interrupt vector type number to the data bus during $\overline{INTA}$ vary widely. One method uses resisters to apply interrupt type number FFH to the data bus, while another uses a three-state buffer to apply any vector type number.
13. The 8259A programmable interrupt controller (PIC) adds at least eight interrupt inputs to the microprocessor. If more interrupts are needed, this device can be cascaded to provide up to 64 interrupt inputs.
14. Programming the 8259A is a two-step process. First, a series of initialization command words (ICWs) are sent to the 8259A, then a series of operation command words (OCWs) are sent.
15. The 8259A contains three status registers: IMR (interrupt mask register), ISR (in-service register), and IRR (interrupt request register).
16. A real-time clock is used to keep time in real time. In most cases, time is stored in either binary or BCD form in several memory locations.

12–7   QUESTIONS AND PROBLEMS

1. What is interrupted by an interrupt?
2. Define the term *interrupt*.
3. What is called by an interrupt?
4. Why do interrupts free up time for the microprocessor?
5. List the interrupt pins found on the microprocessor.
6. List the five interrupt instructions for the microprocessor.
7. What is an interrupt vector?
8. Where are the interrupt vectors located in the microprocessor's memory?
9. How many different interrupt vectors are found in the interrupt vector table?
10. Which interrupt vectors are reserved by Intel?
11. Explain how a type 0 interrupt occurs.
12. Where is the interrupt descriptor table located for protected mode operation?
13. Each protected mode interrupt descriptor contains what information?
14. Describe the differences between a protected and real mode interrupt.
15. Describe the operation of the BOUND instruction.
16. Describe the operation of the INTO instruction.
17. What memory locations contain the vector for an INT 44H instruction?
18. Explain the operation of the IRET instruction.
19. Where is the IRETQ instruction used?
20. What is the purpose of interrupt vector type number 7?
21. List the events that occur when an interrupt becomes active.
22. Explain the purpose of the interrupt flag (IF).
23. Explain the purpose of the trap flag (TF).
24. How is IF cleared and set?
25. How is TF cleared and set?
26. The NMI interrupt input automatically vectors through which vector type number?
27. Does the $\overline{\text{INTA}}$ signal activate for the NMI pin?
28. The INTR input is _____-sensitive.
29. The NMI input is _____-sensitive.
30. When the $\overline{\text{INTA}}$ signal becomes a logic 0, it indicates that the microprocessor is waiting for an interrupt _____ number to be placed on the data bus (D_0–D_7).
31. What is an FIFO?
32. Develop a circuit that places interrupt type number CCH on the data bus in response to the INTR input.
33. Develop a circuit that places interrupt type number 86H on the data bus in response to the INTR input.
34. Explain why pull-up resistors on D_0–D_7 cause the microprocessor to respond with interrupt vector type number FFH for the $\overline{\text{INTA}}$ pulse.
35. What is a daisy-chain?
36. Why must interrupting devices be polled in a daisy-chained interrupt system?
37. What is the 8259A?
38. How many 8259As are required to have 64 interrupt inputs?
39. What is the purpose of the IR_0–IR_7 pins on the 8259A?
40. When are the CAS_2–CAS_0 pins used on the 8259A?
41. Where is a slave INT pin connected on the master 8259A in a cascaded system?
42. What is an OCW?
43. What is an ICW?
44. Where is the vector type number stored in the 8259A?

45. How many ICWs are needed to program the 8259A when operated as a single master in a system?
46. What is the purpose of ICW_1?
47. Where is the sensitivity of the IR pins programmed in the 8259A?
48. Explain priority rotation in the 8259A.
49. What is a nonspecific EOI?
50. At which interrupt vectors is the master 8259A found in the personal computer?
51. What is the purpose of IRR in the 8259A?
52. At which interrupt vectors is the slave 8259A found in the personal computer?

CHAPTER 13

Direct Memory Access and DMA-Controlled I/O

INTRODUCTION

In previous chapters, we discussed basic and interrupt-processed I/O. Now we turn to the final form of I/O called direct memory access (DMA). The DMA I/O technique provides direct access to the memory while the microprocessor is temporarily disabled. This allows data to be transferred between memory and the I/O device at a rate that is limited only by the speed of the memory components in the system or the DMA controller. The DMA transfer speed can approach 33 to 150 M-byte transfer rates with today's high-speed RAM memory components.

DMA transfers are used for many purposes, but more common are DRAM refresh, video displays for refreshing the screen, and disk memory system reads and writes. The DMA transfer is also used to do high-speed memory-to-memory transfers.

This chapter also explains the operation of disk memory systems and video systems that are often DMA-processed. Disk memory includes floppy, fixed, and optical disk storage. Video systems include digital and analog monitors.

CHAPTER OBJECTIVES

Upon completion of this chapter, you will be able to:
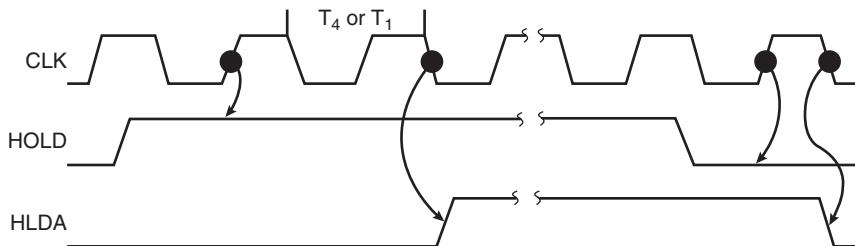
1. Describe a DMA transfer.
2. Explain the operation of the HOLD and HLDA direct memory access control signals.
3. Explain the function of the 8237 DMA controller when used for DMA transfers.
4. Program the 8237 to accomplish DMA transfers.
5. Describe the disk standards found in personal computer systems.
6. Describe the various video interface standards that are found in the personal computer.

13–1   BASIC DMA OPERATION

Two control signals are used to request and acknowledge a direct memory access (DMA) transfer in the microprocessor-based system. The HOLD pin is an input that is used to request a DMA action and the HLDA pin is an output that acknowledges the DMA action. Figure 13–1 shows the timing that is typically found on these two DMA control pins.

FIGURE 13–1 HOLD
and HLDA timing for the
microprocessor.

Whenever the HOLD input is placed at a logic 1 level, a DMA action (hold) is requested. The microprocessor responds, within a few clocks, by suspending the execution of the program and by placing its address, data, and control bus at their high-impedance states. The high-impedance state causes the microprocessor to appear as if it has been removed from its socket. This state allows external I/O devices or other microprocessors to gain access to the system buses so that memory can be accessed directly.

As the timing diagram indicates, HOLD is sampled in the middle of any clocking cycle. Thus, the hold can take effect any time during the operation of any instruction in the microprocessor's instruction set. As soon as the microprocessor recognizes the hold, it stops executing software and enters hold cycles. Note that the HOLD input has a higher priority than the INTR or NMI interrupt inputs. Interrupts take effect at the end of an instruction, whereas a HOLD takes effect in the middle of an instruction. The only microprocessor pin that has a higher priority than a HOLD is the RESET pin. Note that the HOLD input may not be active during a RESET or the reset is not guaranteed.
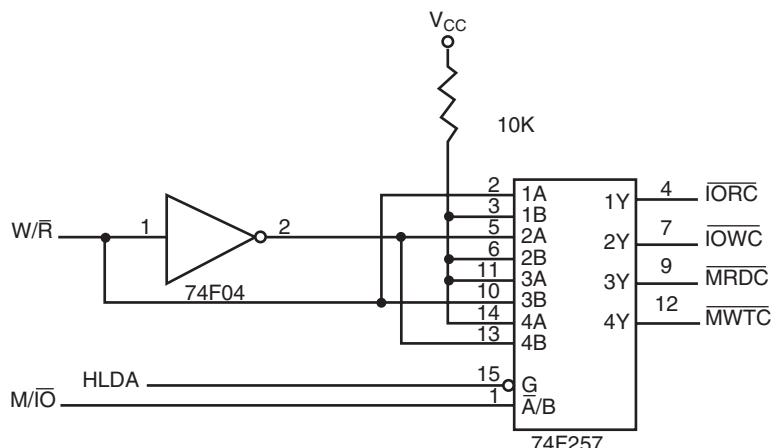
The HLDA signal becomes active to indicate that the microprocessor has indeed placed its buses at their high-impedance state, as can be seen in the timing diagram. Note that there are a few clock cycles between the time that HOLD changes and until HLDA changes. The HLDA output is a signal to the external requesting device that the microprocessor has relinquished control of its memory and I/O space. You could call the HOLD input a DMA request input and the HLDA output a DMA grant signal.

Basic DMA Definitions

Direct memory accesses normally occur between an I/O device and memory without the use of the microprocessor. A **DMA read** transfers data from the memory to the I/O device. A **DMA write** transfers data from an I/O device to memory. In both operations, the memory and I/O are controlled simultaneously, which is why the system contains separate memory and I/O control signals. This special control bus structure of the microprocessor allows DMA transfers. A DMA read causes both the $\overline{MRDC}$ and $\overline{IOWC}$ signals to activate simultaneously, transferring data from the memory to the I/O device. A DMA write causes the $\overline{MWTC}$ and $\overline{IORC}$ signals to both activate. These control bus signals are available to all microprocessors in the Intel family except the 8086/8088 system. The 8086/8088 require their generation with either a system controller or a circuit such as the one illustrated in Figure 13–2. The DMA controller provides the memory with its address and a signal from the controller ($\overline{DACK}$) selects the I/O device during the DMA transfer.

The data transfer speed is determined by the speed of the memory device or a DMA controller that often controls DMA transfers. If the memory speed is 50 ns, DMA transfers occur at rates of up to 1/50 ns or 20 M bytes per second. If the DMA controller in a system functions at a maximum rate of 15 MHz and we still use 50 ns memory, the maximum transfer rate is 15 MHz because the DMA controller is slower than the memory. In many cases, the DMA controller slows the speed of the system when DMA transfers occur.

FIGURE 13–2 A circuit that generates system control signals in a DMA environment.

Because of the switch to serial data transfers in modern computer systems, DMA is becoming less important. The PCI Express bus, which is serial, transfers data at rates that exceed DMA transfers. Even the SATA (serial ATA) interface for disk drives uses serial transfers at the rate of 300 Mbps, which has replaced DMA transfers for hard disk drives. Serial transfers on main-boards (motherboards) between components that use serial techniques can approach 20 Gbps for the PCI Express connection.

13–2   THE 8237 DMA CONTROLLER

The 8237 DMA controller supplies the memory and I/O with control signals and memory address information during the DMA transfer. The 8237 is actually a special-purpose microprocessor whose job is high-speed data transfer between memory and the I/O. Figure 13–3 shows the pin-out and block diagram of the 8237 programmable DMA controller. Although this device may not appear as a discrete component in modern microprocessor-based systems, it does appear within system controller chip sets found in most systems. Although not described because of its complexity, the modern chip set (ISP or integrated system peripheral controller) and its integral set of two DMA controllers are programmed almost exactly (*it does not support memory-to-memory transfers*) like the 8237. The ISP also provides a pair of 8259A programmable interrupt controllers for the system.

The 8237 is a four-channel device that is compatible with the 8086/8088 microprocessors. The 8237 can be expanded to include any number of DMA channel inputs, although four channels seem to be adequate for many small systems. The 8237 is capable of DMA transfers at rates of up to 1.6M bytes per second. Each channel is capable of addressing a full 64K-byte section of memory and can transfer up to 64K bytes with a single programming.

Pin Definitions

CLK	The **clock** input is connected to the system clock signal as long as that signal is 5 MHz or less. In the 8086/8088 system, the clock must be inverted for the proper operation of the 8237.
$\overline{\text{CS}}$	**Chip select** enables the 8237 for programming. The $\overline{\text{CS}}$ pin is normally connected to the output of a decoder. The decoder does not use the 8086/8088 control signal IO/$\overline{\text{M}}$(M/$\overline{\text{IO}}$) because it contains the new memory and I/O control signals ($\overline{\text{MEMR}}$, $\overline{\text{MEMW}}$, $\overline{\text{IOR}}$, and $\overline{\text{IOW}}$).
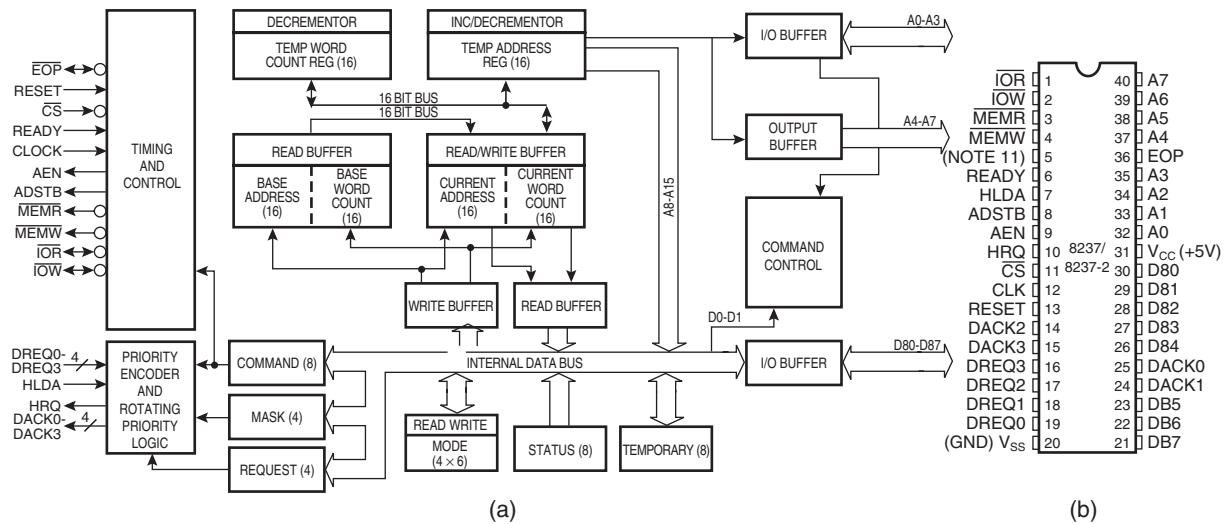
FIGURE 13–3 The 8237A-5 programmable DMA controller. (a) Block diagram and (b) pin-out. (Courtesy of Intel Corporation.)

RESET	The **reset** pin clears the command, status, request, and temporary registers. It also clears the first/last flip-flop and sets the mask register. This input primes the 8237 so it is disabled until programmed otherwise.
READY	A logic 0 on the **ready** input causes the 8237 to enter wait states for slower memory components.
HLDA	A **hold acknowledge** signals the 8237 that the microprocessor has relinquished control of the address, data, and control buses.
$DREQ_0$–$DREQ_3$	The **DMA request inputs** are used to request a DMA transfer for each of the four DMA channels. Because the polarity of these inputs is programmable, they are either active-high or active-low inputs.
DB_0–DB_7	The **data bus** pins are connected to the microprocessor data bus connections and are used during the programming of the DMA controller.
$\overline{IOR}$	**I/O read** is a bidirectional pin used during programming and during a DMA write cycle.
$\overline{IOW}$	**I/O write** is a bidirectional pin used during programming and during a DMA read cycle.
$\overline{EOP}$	**End-of-process** is a bidirectional signal that is used as an input to terminate a DMA process or as an output to signal the end of the DMA transfer. This input is often used to interrupt a DMA transfer at the end of a DMA cycle.
A_0–A_3	These **address pins** select an internal register during programming and also provide part of the DMA transfer address during a DMA action. The address pins are outputs that provide part of the DMA transfer address during a DMA action.
HRQ	**Hold request** is an output that connects to the HOLD input of the microprocessor in order to request a DMA transfer.
$DACK_0$–$DACK_3$	**DMA channel acknowledge** outputs acknowledge a channel DMA request. These outputs are programmable as either active-high or active-low signals. The DACK outputs are often used to select the DMA-controlled I/O device during the DMA transfer.
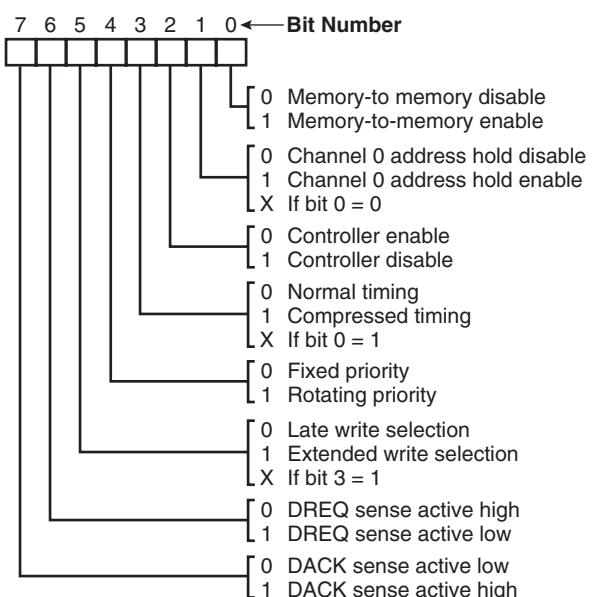
AEN	The **address enable** signal enables the DMA address latch connected to the DB_7–DB_0 pins on the 8237. It is also used to disable any buffers in the system connected to the microprocessor.
ADSTB	**Address strobe** functions as ALE, except that it is used by the DMA controller to latch address bits A_{15}–A_8 during the DMA transfer.
$\overline{MEMR}$	**Memory read** is an output that causes memory to read data during a DMA read cycle.
$\overline{MEMW}$	**Memory write** is an output that causes memory to write data during a DMA write cycle.

Internal Registers

CAR	The **current address register** is used to hold the 16-bit memory address used for the DMA transfer. Each channel has its own current address register for this purpose. When a byte of data is transferred during a DMA operation, the CAR is either incremented or decremented, depending on how it is programmed.
CWCR	The **current word count register** programs a channel for the number of bytes (up to 64K) transferred during a DMA action. The number loaded into this register is one less than the number of bytes transferred. For example, if a 10 is loaded into the CWCR, then 11 bytes are transferred during the DMA action.
BA and BWC	The **base address** (BA) and **base word count** (BWC) registers are used when auto-initialization is selected for a channel. In the auto-initialization mode, these registers are used to reload both the CAR and CWCR after the DMA action is completed. This allows the same count and address to be used to transfer data from the same memory area.
CR	The **command register** programs the operation of the 8237 DMA controller. Figure 13–4 depicts the function of the command register. The command register uses bit position 0 to select the memory-to-memory DMA transfer mode. Memory-to-memory DMA transfers use DMA

FIGURE 13–4 8237A-5 command register. (Courtesy of Intel Corporation.)

channel 0 to hold the source address and DMA channel 1 to hold the destination address. (This is similar to the operation of a MOVSB instruction.) A byte is read from the address accessed by channel 0 and saved within the 8237 in a temporary holding register. Next, the 8237 initiates a memory write cycle in which the contents of the temporary holding register are written into the address selected by DMA channel 1. The number of bytes transferred is determined by the channel 1 count register.

The channel 0 address hold enable bit (bit position 1) programs channel 0 for memory-to-memory transfers. For example, if you must fill an area of memory with data, channel 0 can be held at the same address while channel 1 changes for memory-to-memory transfer. This copies the contents of the address accessed by channel 0 into a block of memory accessed by channel 1.

The controller enable/disable bit (bit position 2) turns the entire controller on and off. The normal and compressed bit (bit position 3) determine whether a DMA cycle contains two (compressed) or four (normal) clocking periods. Bit position 5 is used in normal timing to extend the write pulse so it appears one clock earlier in the timing for I/O devices that require a wider write pulse.
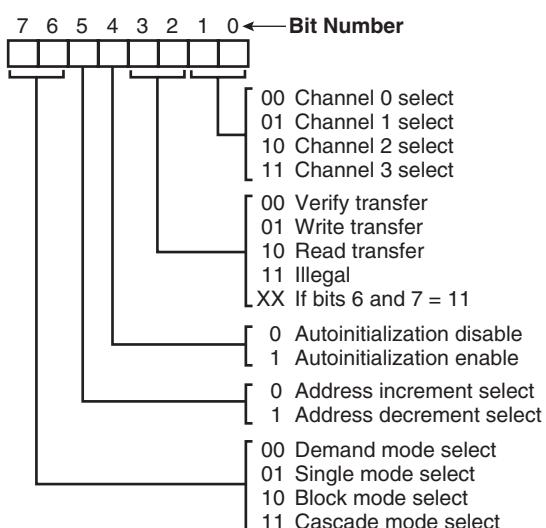
Bit position 4 selects priority for the four DMA channel DREQ inputs. In the fixed priority scheme, channel 0 has the highest priority and channel 3 has the lowest. In the rotating priority scheme, the most recently serviced channel assumes the lowest priority. For example, if channel 2 just had access to a DMA transfer, it assumes the lowest priority and channel 3 assumes the highest priority position. Rotating priority is an attempt to give all channels equal priority.

The remaining two bits (bit positions 6 and 7) program the polarities of the DREQ inputs and the DACK outputs.

MR   The **mode register** programs the mode of operation for a channel. Note that each channel has its own mode register (see Figure 13–5), as selected by bit positions 1 and 0. The remaining bits of the mode register select the operation, auto-initialization, increment/decrement, and

FIGURE 13–5   8237A-5 mode register. (Courtesy of Intel Corporation.)

7 6 5 4 3 2 1 0 ◄— Bit Number

00 Channel 0 select
01 Channel 1 select
10 Channel 2 select
11 Channel 3 select

00 Verify transfer
01 Write transfer
10 Read transfer
11 Illegal
XX If bits 6 and 7 = 11

0 Autoinitialization disable
1 Autoinitialization enable

0 Address increment select
1 Address decrement select

00 Demand mode select
01 Single mode select
10 Block mode select
11 Cascade mode select

mode for the channel. Verification operations generate the DMA addresses without generating the DMA memory and I/O control signals.

The modes of operation include demand mode, single mode, block mode, and cascade mode. Demand mode transfers data until an external EOP is input or until the DREQ input becomes inactive. Single mode releases the HOLD after each byte of data is transferred. If the DREQ pin is held active, the 8237 again requests a DMA transfer through the DRQ line to the microprocessor's HOLD input. Block mode automatically transfers the number of bytes indicated by the count register for the channel. DREQ need not be held active through the block mode transfer. Cascade mode is used when more than one 8237 is present in a system.

BR   The **bus request register** is used to request a DMA transfer via software (see Figure 13–6). This is very useful in memory-to-memory transfers, where an external signal is not available to begin the DMA transfer.

MRSR   The **mask register set/reset** sets or clears the channel mask, as illustrated in Figure 13–7. If the mask is set, the channel is disabled. Recall that the RESET signal sets all channel masks to disable them.

MSR   The **mask register** (see Figure 13–8) clears or sets all of the masks with one command instead of individual channels, as with the MRSR.

FIGURE 13–6   8237A-5 request register. (Courtesy of Intel Corporation.)
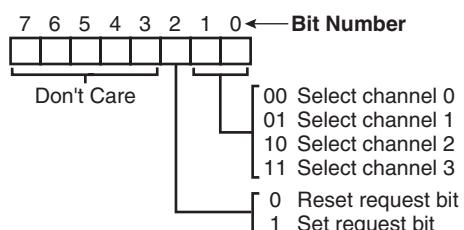
FIGURE 13–7   8237A-5 mask register set/reset mode. (Courtesy of Intel Corporation.)
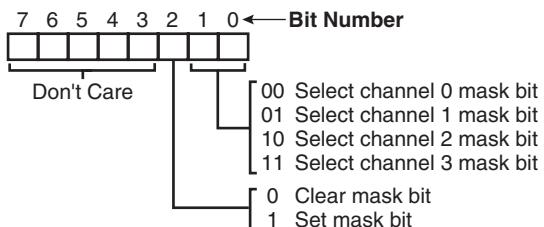
FIGURE 13–8   8237A-5 mask register. (Courtesy of Intel Corporation.)
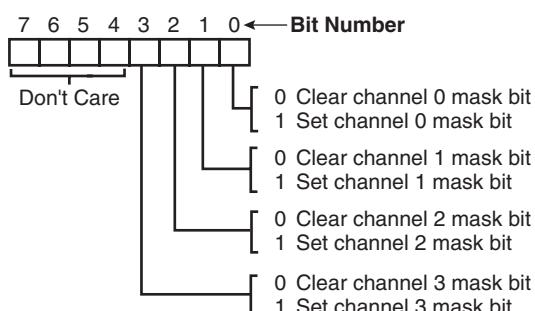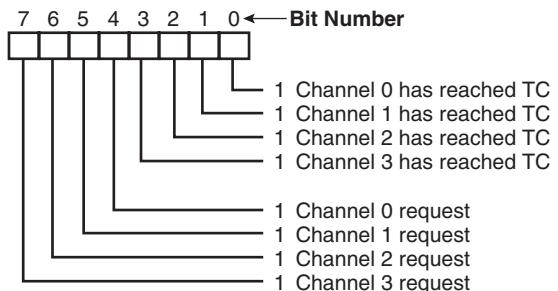
FIGURE 13–9   8237A-5 status register. (Courtesy of Intel Corporation.)

SR   The **status register** shows the status of each DMA channel (see Figure 13–9). The TC bits indicate whether the channel has reached its terminal count (transferred all its bytes). Whenever the terminal count is reached, the DMA transfer is terminated for most modes of operation. The request bits indicate whether the DREQ input for a given channel is active.

Software Commands

Three software commands are used to control the operation of the 8237. These commands do not have a binary bit pattern, as do the various control registers within the 8237. A simple output to the correct port number enables the software command. Figure 13–10 shows the I/O port assignments that access all registers and the software commands.

The functions of the software commands are explained in the following list:

1. **Clear the first/last flip-flop**—Clears the first/last (F/L) flip-flop within the 8237. The F/L flip-flop selects which byte (low or high order) is read/written in the current address and current count registers. If F/L = 0, the low-order byte is selected; if F/L = 1, the high-order byte is selected. Any read or write to the address or count register automatically toggles the F/L flip-flop.
2. **Master clear**—Acts exactly the same as the RESET signal to the 8237. As with the RESET signal, this command disables all channels.
3. **Clear mask register**—Enables all four DMA channels.

FIGURE 13–10   8237A-5 command and control port assignments. (Courtesy of Intel Corporation.)

Signals						Operation
A3	A2	A1	A0	$\overline{IOR}$	$\overline{IOW}$	
1	0	0	0	0	1	Read Status Register
1	0	0	0	1	0	Write Command Register
1	0	0	1	0	1	Illegal
1	0	0	1	1	0	Write Request Register
1	0	1	0	0	1	Illegal
1	0	1	0	1	0	Write Single Mask Register Bit
1	0	1	1	0	1	Illegal
1	0	1	1	1	0	Write Mode Register
1	1	0	0	0	1	Illegal
1	1	0	0	1	0	Clear Byte Pointer Flip/Flop
1	1	0	1	0	1	Read Temporary Register
1	1	0	1	1	0	Master Clear
1	1	1	0	0	1	Illegal
1	1	1	0	1	0	Clear Mask Register
1	1	1	1	0	1	Illegal
1	1	1	1	1	0	Write All Mask Register Bits

Programming the Address and Count Registers

Figure 13–11 illustrates the I/O port locations for programming the count and address registers for each channel. Notice that the state of the F/L flip-flop determines whether the LSB or MSB is programmed. If the state of the F/L flip-flop is unknown, the count and address could be programmed incorrectly. It is also important that the DMA channel be disabled before its address and count are programmed.

Four steps are required to program the 8237: (1) The F/L flip-flop is cleared using a clear F/L command; (2) the channel is disabled; (3) the LSB and then MSB of the address are programmed; and (4) the LSB and MSB of the count are programmed. Once these four operations are performed, the channel is programmed and ready to use. Additional programming is required to select the mode of operation before the channel is enabled and started.

The 8237 Connected to the 80X86 Microprocessor

Figure 13–12 shows an 80X86-based system that contains the 8237 DMA controller.

The address enable (AEN) output of the 8237 controls the output pins of the latches and the outputs of the 74LS257 (E). During normal 80X86 operation (AEN = 0), latches A and C and the multiplexer (E) provide address bus bits A_{19}–A_{16} and A_7–A_0. The multiplexer provides the

Channel	Register	Operation	Signals							Internal Flip-Flop	Data Bus DB0-DB7
			$\overline{CS}$	$\overline{IOR}$	$\overline{IOW}$	A3	A2	A1	A0		
0	Base and Current Address	Write	0	1	0	0	0	0	0	0	A0-A7
			0	1	0	0	0	0	0	1	A8-A15
	Current Address	Read	0	0	1	0	0	0	0	0	A0-A7
			0	0	1	0	0	0	0	1	A8-A15
	Base and Current Word Count	Write	0	1	0	0	0	0	1	0	W0-W7
			0	1	0	0	0	0	1	1	W8-W15
	Current Word Count	Read	0	0	1	0	0	0	1	0	W0-W7
			0	0	1	0	0	0	1	1	W8-W15
1	Base and Current Address	Write	0	1	0	0	0	1	0	0	A0-A7
			0	1	0	0	0	1	0	1	A8-A15
	Current Address	Read	0	0	1	0	0	1	0	0	A0-A7
			0	0	1	0	0	1	0	1	A8-A15
	Base and Current Word Count	Write	0	1	0	0	0	1	1	0	W0-W7
			0	1	0	0	0	1	1	1	W8-W15
	Current Word Count	Read	0	0	1	0	0	1	1	0	W0-W7
			0	0	1	0	0	1	1	1	W8-W15
2	Base and Current Address	Write	0	1	0	0	1	0	0	0	A0-A7
			0	1	0	0	1	0	0	1	A8-A15
	Current Address	Read	0	0	1	0	1	0	0	0	A0-A7
			0	0	1	0	1	0	0	1	A8-A15
	Base and Current Word Count	Write	0	1	0	0	1	0	1	0	W0-W7
			0	1	0	0	1	0	1	1	W8-W15
	Current Word Count	Read	0	0	1	0	1	0	1	0	W0-W7
			0	0	1	0	1	0	1	1	W8-W15
3	Base and Current Address	Write	0	1	0	0	1	1	0	0	A0-A7
			0	1	0	0	1	1	0	1	A8-A15
	Current Address	Read	0	0	1	0	1	1	0	0	A0-A7
			0	0	1	0	1	1	0	1	A8-A15
	Base and Current Word Count	Write	0	1	0	0	1	1	1	0	W0-W7
			0	1	0	0	1	1	1	1	W8-W15
	Current Word Count	Read	0	0	1	0	1	1	1	0	W0-W7
			0	0	1	0	1	1	1	1	W8-W15

FIGURE 13–11   8237A-5 DMA channel I/O port addresses. (Courtesy of Intel Corporation.)

TABLE 13–1 DMA page register ports.

Channel	Port for A_{16}–A_{23}	Port for A_{24}–A_{31}
0	87H	487H
1	83H	483H
2	81H	481H
3	82H	482H
4	8FH	48FH
5	8BH	48BH
6	89H	489H
7	8AH	486H

system control signals as long as the 80X86 is in control of the system. During a DMA action (AEN = 1), latches A and C are disabled along with the multiplexer (E). Latches D and B now provide address bits A_{19}–A_{16} and A_{15}–A_8. Address bus bits A_7–A_0 are provided directly by the 8237 and contain a part of the DMA transfer address. The control signals $\overline{MEMR}$, $\overline{MEMW}$, $\overline{IOR}$, and $\overline{IOW}$ are provided by the DMA controller.

The address strobe output (ADSTB) of the 8237 clocks the address (A_{15}–A_8) into latch D during the DMA action so that the entire DMA transfer address becomes available on the address bus. Address bus bits A_{19}–A_{16} are provided by latch B, which must be programmed with these four address bits before the controller is enabled for the DMA transfer. The DMA operation of the 8237 is limited to a transfer of not more than 64K bytes within the same 64K-byte section of the memory.

The decoder (F) selects the 8237 for programming and the 4-bit latch (B) for the uppermost four address bits. The latch in a PC is called the DMA page register (8 bits) that holds address bits A_{16}–A_{23} for a DMA transfer. A high page register also exists, but its address is chip-dependent. The port numbers for the DMA page registers are listed in Table 13–1 (these are for the Intel ISP). The decoder in this system enables the 8237 for I/O port addresses XX60H–XX7FH, and the I/O latch (B) for ports XX00H–XX1FH. Notice that the decoder output is combined with the $\overline{IOW}$ signal to generate an active-high clock for the latch (B).

During normal 80X86 operation, the DMA controller and integrated circuits B and D are disabled. During a DMA action, integrated circuits A, C, and E are disabled so that the 8237 can take control of the system through the address, data, and control buses.
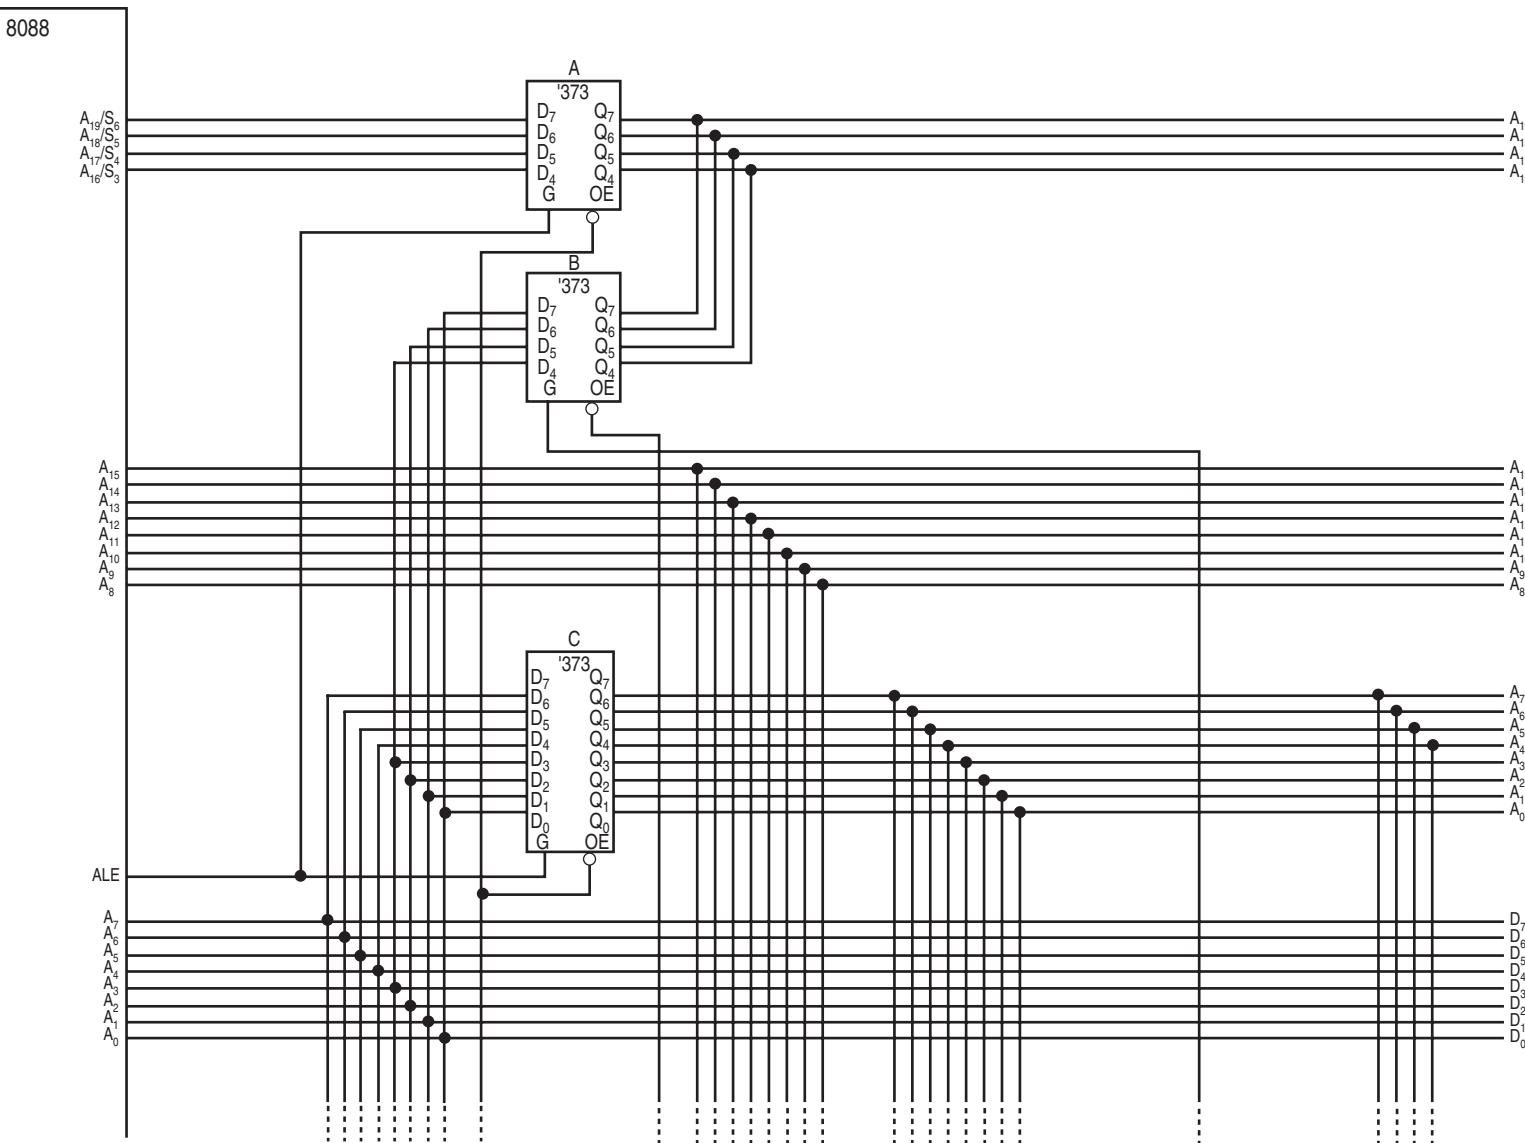
In the personal computer, the two DMA controllers are programmed at I/O ports 0000H–000FH for DMA channels 0–3, and at ports 00C0H–00DFH for DMA channels 4–7. Note that the second controller is programmed at even addresses only, so the channel 4 base and current address is programmed at I/O port 00C0H and the channel 4 base and current count is programmed at port 00C2H. The page register, which holds address bits A_{23}–A_{16} of the DMA address, is located at I/O ports 0087H (CH-0), 0083H (CH-1), 0081H (CH-2), 0082H (CH-3), (no channel 4), 008BH (CH-5), 0089H (CH-6), and 008AH (CH-7). The page register functions as the address latch described with the examples in this text.
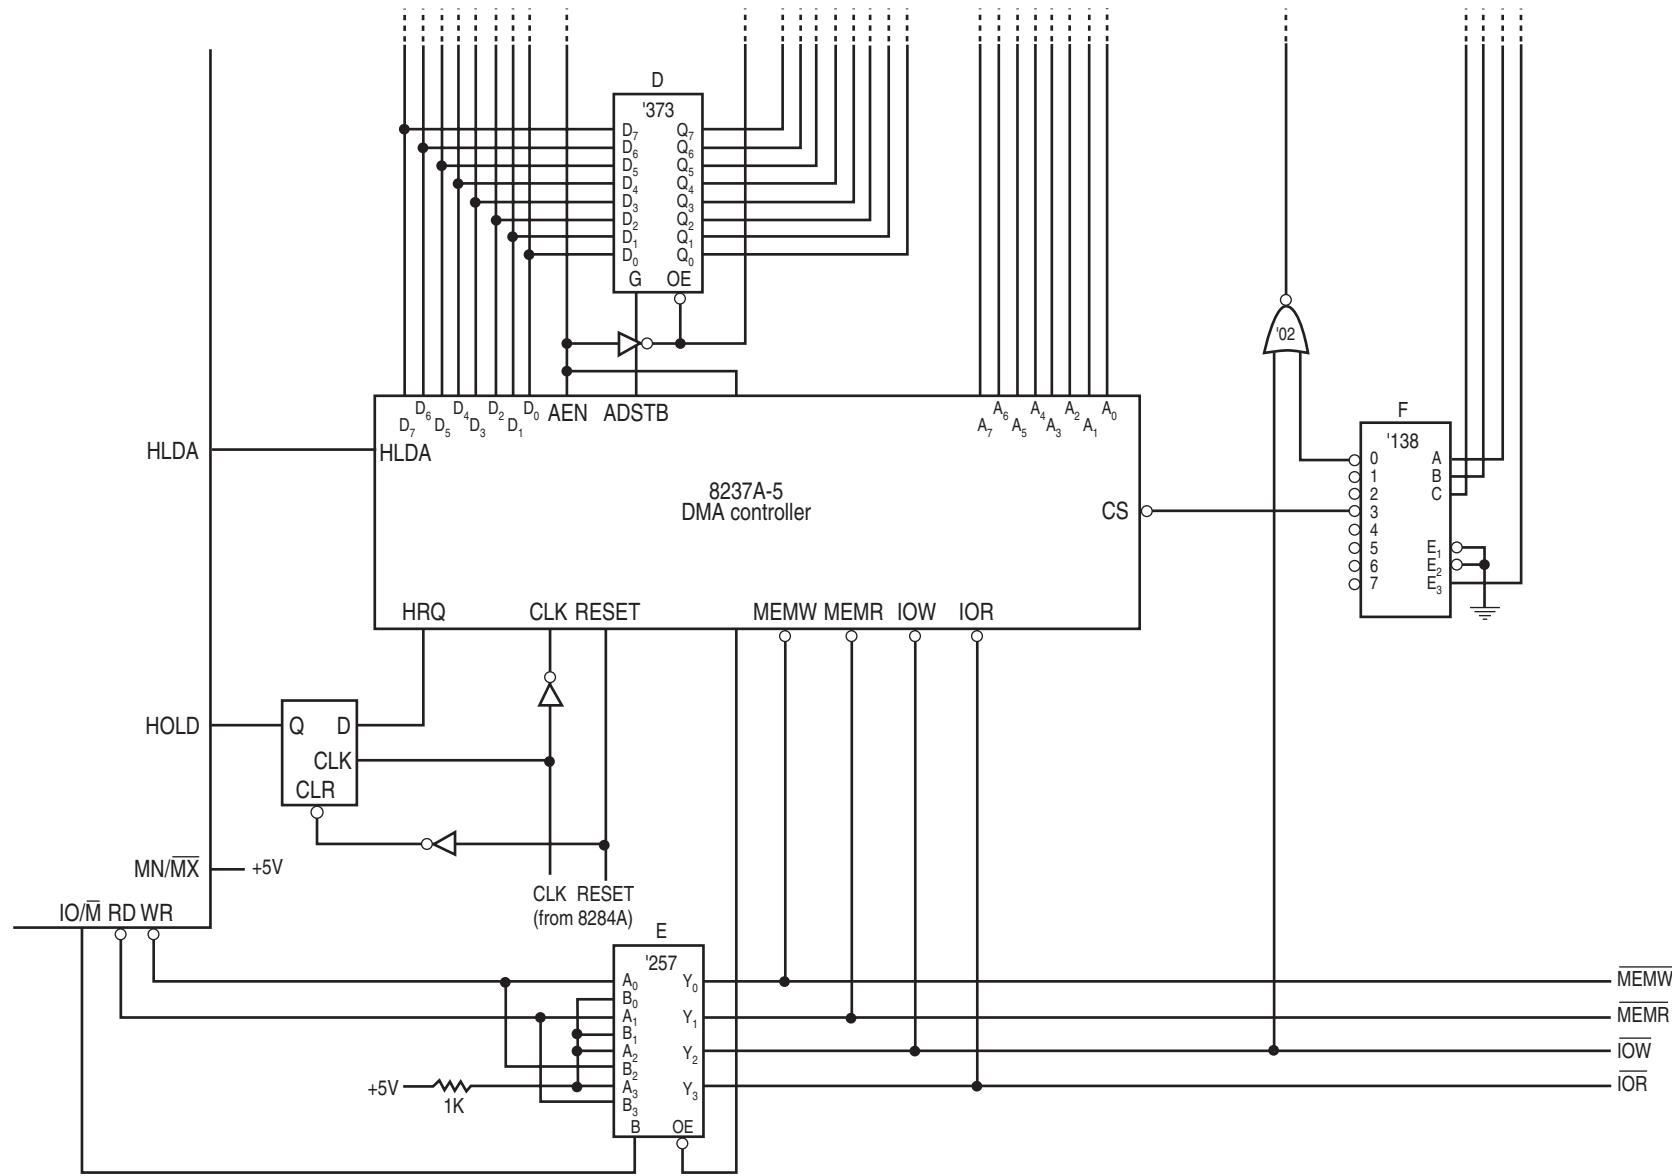
Memory-to-Memory Transfer with the 8237

The memory-to-memory transfer is much more powerful than even the automatically repeated MOVSB instruction. (Note: Most modern chip sets do not support the memory-to-memory feature.) Although the repeated MOVSB instruction tables show that the 8088 requires 4.2 μs per byte, the 8237 requires only 2.0 μs per byte, which is over twice as fast as a software data transfer. This is not true if an 80386, 80846, or Pentium through Pentium 4 is in use in the system.

Sample Memory-to-Memory DMA Transfer. Suppose that the contents of memory locations 10000H–13FFFH are to be transferred into memory locations 14000H–17FFFH. This is accomplished with a repeated string move instruction or, at a much faster rate, with the DMA controller.

FIGURE 13–12   Complete 8088 minimum mode DMA system.

Example 13–1 illustrates the software required to initialize the 8237 and program latch B in Figure 13–12 for this DMA transfer. This software is written for an embedded application. For it to function in the PC (if your chip set supports this feature), you must use the port addresses listed in Table 13–1 for the page registers.

EXAMPLE 13–1

```
;A procedure that transfers a block of data using the 8237A
;DMA controller in Figure 13-12. This is a memory-to-memory
;transfer.

;Calling parameters:
;    SI = source address
;    DI = destination address
;    CX = count
;    ES = segment of source and destination

LATCHB EQU    10H
CLEARF EQU    7CH
CHOA   EQU    70H
CH1A   EQU    72H
CH1C   EQU    73H
MODE   EQU    7BH
CMMD   EQU    78H
MASKS  EQU    7FH
REQ    EQU    79H
STATUS EQU    78H

TRANS PROC    NEAR USES AX

    MOV   AX,ES               ;program latch B
    MOV   AL,AH
    SHR   AL,4
    OUT   LATCHB,AL
    OUT   CLEARF,AL           ;clear F/L

    MOV   AX,ES               ;program source address
    SHL   AX,4
    ADD   AX,SI
    OUT   CH0A,AL
    MOV   AL,AH
    OUT   CH0A

    MOV   AX,ES               ;program destination address
    SHL   AX,4
    ADD   AX,DI
    OUT   CH1A,AL
    MOV   AL,AH
    OUT   CH1A,AL

    MOV   AX,CX               ;program count
    DEC   AX
    OUT   CH1C,AL
    MOV   AL,AH
    OUT   CH1C,AL

    MOV   AL,88H              ;program mode
    OUT   MODE,AL
    MOV   AL,85H
    OUT   MODE,AL

    MOV   AL,1                ;enable block transfer
    OUT   CMMD,AL

    MOV   AL,0EH              ;unmask channel 0
    OUT   MASKS,AL
```

```
      MOV   AL,4                 ;start DMA
      OUT   REQ,AL

      .REPEAT                    ;wait for completion
            IN  AL,STATUS
      .UNTIL AL & 1
      RET

TRANS ENDP
```

Programming the DMA controller requires a few steps, as illustrated in Example 13–1. The leftmost digit of the 5-digit address is sent to latch B. Next, the channels are programmed after the F/L flip-flop is cleared. Note that we use channel 0 as the source and channel 1 as the destination for a memory-to-memory transfer. The count is next programmed with a value that is one less than the number of bytes to be transferred. Next, the mode register of each channel is programmed, the command register selects a block move, channel 0 is enabled, and a software DMA request is initiated. Before return is made from the procedure, the status register is tested for a terminal count. Recall that the terminal count flag indicates that the DMA transfer is completed. The TC also disables the channel, preventing additional transfers.

Sample Memory Fill Using the 8237.   In order to fill an area of memory with the same data, the channel 0 source register is programmed to point to the same address throughout the transfer. This is accomplished with the channel 0 hold mode. The controller copies the contents of this single memory location to an entire block of memory addressed by channel 1. This has many useful applications.

For example, suppose that a DOS video display must be cleared. This operation can be performed using the DMA controller with the channel 0 hold mode and a memory-to-memory transfer. If the video display contains 80 columns and 25 lines, it has 2000 display positions that must be set to 20H (an ASCII space) to clear the screen.

Example 13–2 shows a procedure that clears an area of memory addressed by ES:DI. The CX register transfers the number of bytes to be cleared to the CLEAR procedure. Notice that this procedure is nearly identical to Example 13–1, except that the command register is programmed so the channel 0 address is held. The source address is programmed as the same address as ES:DI, and then the destination is programmed as one location beyond ES:DI. Also note that this program is designed to function with the hardware in Figure 13–12 and will not function in the personal computer unless you have the same hardware.

EXAMPLE 13–2

```
;A procedure that clears the DOS mode video screen using the DMA
;controller as depicted in Figure 13-12.

;Calling sequence:
;     DI = offset address of area cleared
;     ES = segment address of area cleared
;     CX = number of bytes cleared

LATCHB EQU   10H
CLEARF EQU   7CH
CHOA   EQU   70H
CH1A   EQU   72H
CH1C   EQU   73H
MODE   EQU   7BH
CMMD   EQU   78H
MASKS  EQU   7FH
REQ    EQU   79H
STATUS EQU   78H
ZERO   EQU   0

CLEAR  PROC  NEAR USES AX
```

```
        MOV    AX,ES              ;program latch B
        MOV    AL,AH
        SHR    AL,4
        OUT    LATCHB,AL
        OUT    CLEARF,AL          ;clear F/L

        MOV    AL,ZERO            ;save zero in first byte
        MOV    ES:[DI],AL

        MOV    AX,ES              ;program source address
        SHL    AX,4
        ADD    AX,SI
        OUT    CH0A,AL
        MOV    AL,AH
        OUT    CH0A

        MOV    AX,ES              ;program destination address
        SHL    AX,4
        ADD    AX,DI
        OUT    CH1A,AL
        MOV    AL,AH
        OUT    CH1A,AL

        MOV    AX,CX              ;program count
        DEC    AX
        OUT    CH1C,AL
        MOV    AL,AH
        OUT    CH1C,AL

        MOV    AL,88H             ;program mode
        OUT    MODE,AL
        MOV    AL,85H
        OUT    MODE,AL

        MOV    AL,03H             ;enable block hold transfer
        OUT    CMMD,AL

        MOV    AL,0EH             ;enable channel 0
        OUT    MASKS,AL

        MOV    AL,4               ;start DMA
        OUT    REQ,AL

        .REPEAT
            IN   AL,STATUS
        .UNTIL AL &  1
        RET

CLEAR   ENDP
```
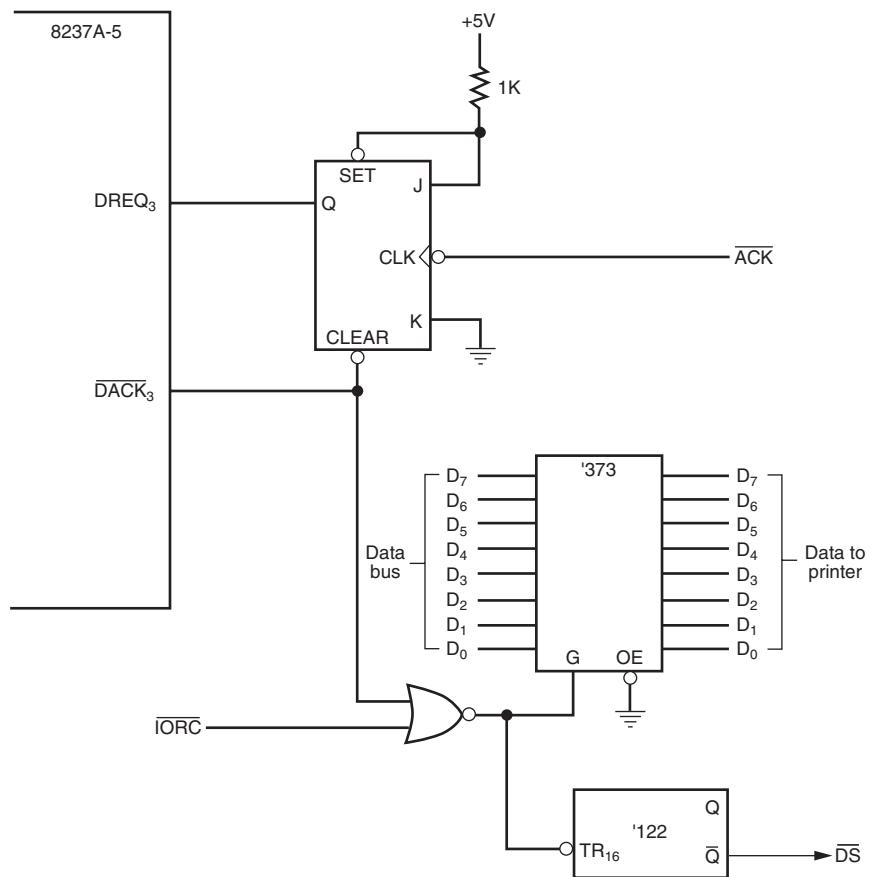
DMA-Processed Printer Interface

Figure 13–13 illustrates the hardware added to Figure 13–12 for a DMA-controlled printer inter-
face. Little additional circuitry is added for this interface to a Centronics-type parallel printer. The
latch is used to capture the data as it is sent to the printer during the DMA transfer. The write pulse
passed through to the latch during the DMA action also generates the data strobe ($\overline{DS}$) signal to
the printer through the single-shot. The $\overline{ACK}$ signal returns from the printer each time it is ready
for additional data. In this circuit, $\overline{ACK}$ is used to request a DMA action through a flip-flop.

Notice that the I/O device is not selected by decoding the address on the address bus.
During the DMA transfer, the address bus contains the memory address and cannot contain the
I/O port address. In place of the I/O port address, the $\overline{DACK3}$ output from the 8237 selects the
latch by gating the write pulse through an OR gate.

Software that controls this interface is simple because only the address of the data and
the number of characters to be printed are programmed. Once programmed, the channel is

FIGURE 13–13   DMA-processed printer interface.

enabled, and the DMA action transfers a byte at a time to the printer interface each time that the interface receives the $\overline{\text{ACK}}$ signal from the printer.

The procedure that prints data from the current data segment is illustrated in Example 13–3. This procedure programs the 8237, but doesn't actually print anything. Printing is accomplished by the DMA controller and the printer interface.

EXAMPLE 13–3

```
;A procedure that prints data via the printer interface in
;Figure 13–13

;Calling sequence:
;       BX = offset address of printer data
;       DS = segment address of printer data
;       CX = number of bytes to print

LATCHB EQU    10H
CLEARF EQU    7CH
CH3A   EQU    76H
CH1C   EQU    77H
MODE   EQU    7BH
CMMD   EQU    78H
MASKS  EQU    7FH
REQ    EQU    79H

PRINT  PROC   NEAR USES AX CX BX

    MOV  EAX,0
    MOV  AX,DS              ;program latch B
```

```
        SHR     EAX,4
        PUSH    AX
        SHR     EAX,16
        OUT     LATCHB,AL

        POP     AX              ;program address
        OUT     CH3A,AL
        MOV     AL,AH
        OUT     CH3A,AL

        MOV     AX,CX           ;program count
        DEC     AX
        OUT     CH3C,AL
        MOV     AL,AH
        OUT     CH3C,AL

        MOV     AL,0BH          ;program mode
        OUT     MODE,AL

        MOV     AL,00H          ;enable block mode transfer
        OUT     CMMD,AL

        MOV     AL,7            ;enable channel 3
        OUT     MASKS,AL
        RET

PRINT   ENDP
```

A secondary procedure is needed to determine whether the DMA action has been completed. Example 13–4 lists the secondary procedure that tests the DMA controller to see whether the DMA transfer is complete. The TESTP procedure is called before programming the DMA controller to see whether the prior transfer is complete.

EXAMPLE 13–4

```
;A procedure that tests for completion of the DMA action

STATUS  EQU   78H

TESTP   PROC  NEAR USES AX

        .REPEAT
                IN AL,STATUS
        .UNTIL AL & 8
        RET

TESTP   ENDP
```

Printed data can be double-buffered by first loading buffer 1 with data to be printed. Next, the PRINT procedure is called to begin printing buffer 1. Because it takes very little time to program the DMA controller, a second buffer (buffer 2) can be filled with new printer data while the first buffer (buffer 1) is printed by the printer interface and DMA controller. This process is repeated until all data are printed.

13–3   SHARED-BUS OPERATION

Complex present-day computer systems have so many tasks to perform that some systems are using more than one microprocessor to accomplish the work. This is called a multiprocessing system. We also sometimes call this a distributed system. A system that performs more than one task is called

a multitasking system. In systems that contain more than one microprocessor, some method of control must be developed and employed. In a distributed, multiprocessing, multitasking environment, each microprocessor accesses two buses: (1) the local bus and (2) the remote or shared bus.

This section of the text describes shared bus operation for the 8086 and 8088 microprocessors using the 8289 bus arbiter. The 80286 uses the 82289 bus arbiter and the 80386/80486 uses the 82389 bus arbiter. The Pentium–Pentium 4 directly support a multiuser environment, as described in Chapters 17, 18, and 19. These systems are much more complex and difficult to illustrate at this point in the text, but their terminology and operation is essentially the same as for the 8086/8088.

The local bus is connected to memory and I/O devices that are directly accessed by a single microprocessor without any special protocol or access rules. The remote (shared) bus contains memory and I/O that are accessed by any microprocessor in the system. Figure 13–14 illustrates this idea with a few microprocessors. Note that the personal computer is also configured in the same manner as the system in Figure 13–14. The bus master is the main microprocessor in the personal computer. What we call the local bus in the personal computer is the shared bus in this illustration. The ISA bus is operated as a slave to the personal computer's microprocessor as well as any other devices attached to the shared bus. The PCI bus can operate as a slave or a master.

Types of Buses Defined

The local bus is the bus that is resident to the microprocessor. The local bus contains the resident or local memory and I/O. All microprocessors studied thus far in this text are considered to be local bus systems. The local memory and local I/O are accessed by the microprocessor that is directly connected to them.

A shared bus is one that is connected to all microprocessors in the system. The shared bus is used to exchange data between microprocessors in the system. A shared bus may contain memory
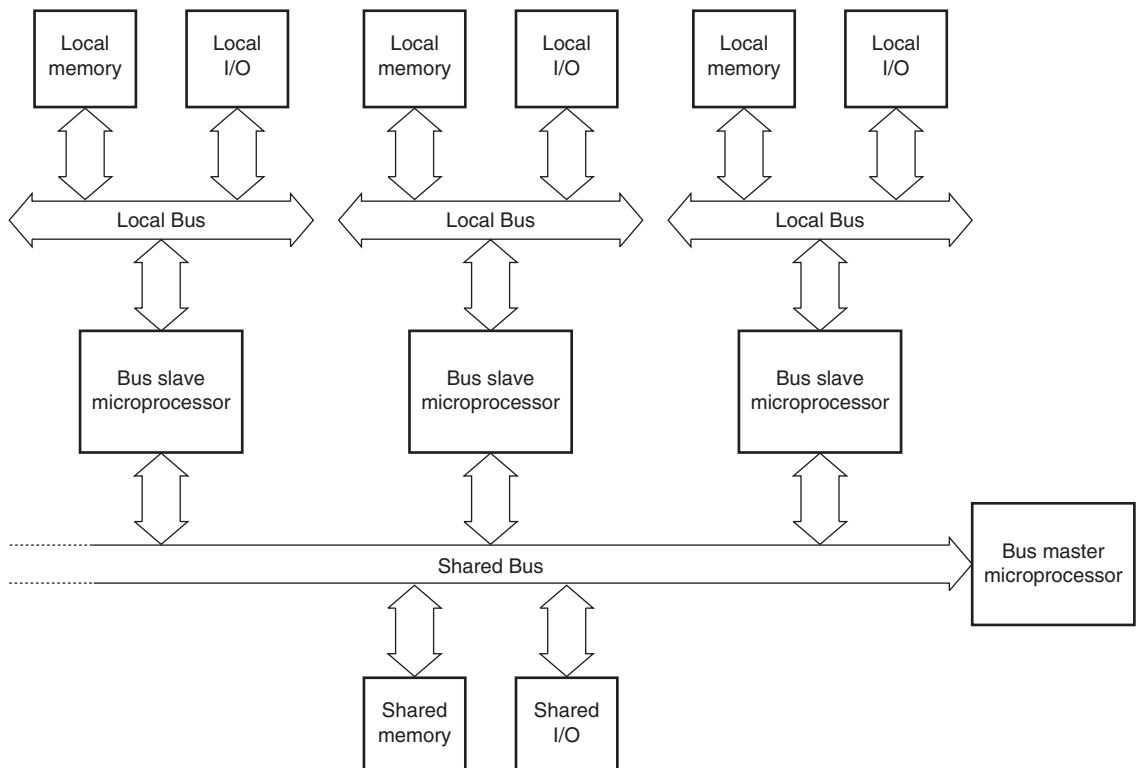
FIGURE 13–14   A block diagram illustrating the shared and local buses.

and I/O devices that are accessed by all microprocessors in the system. Access to the shared bus is controlled by some form or arbiter that allows only a single microprocessor to access the system's shared bus space. As mentioned, the shared bus in the personal computer is what we often call the local bus in the personal computer because it is local to the microprocessor in the personal computer.

Figure 13–15 shows an 8088 microprocessor that is connected as a remote bus master. The term *bus master* applies to any device (microprocessor or otherwise) that can control a bus
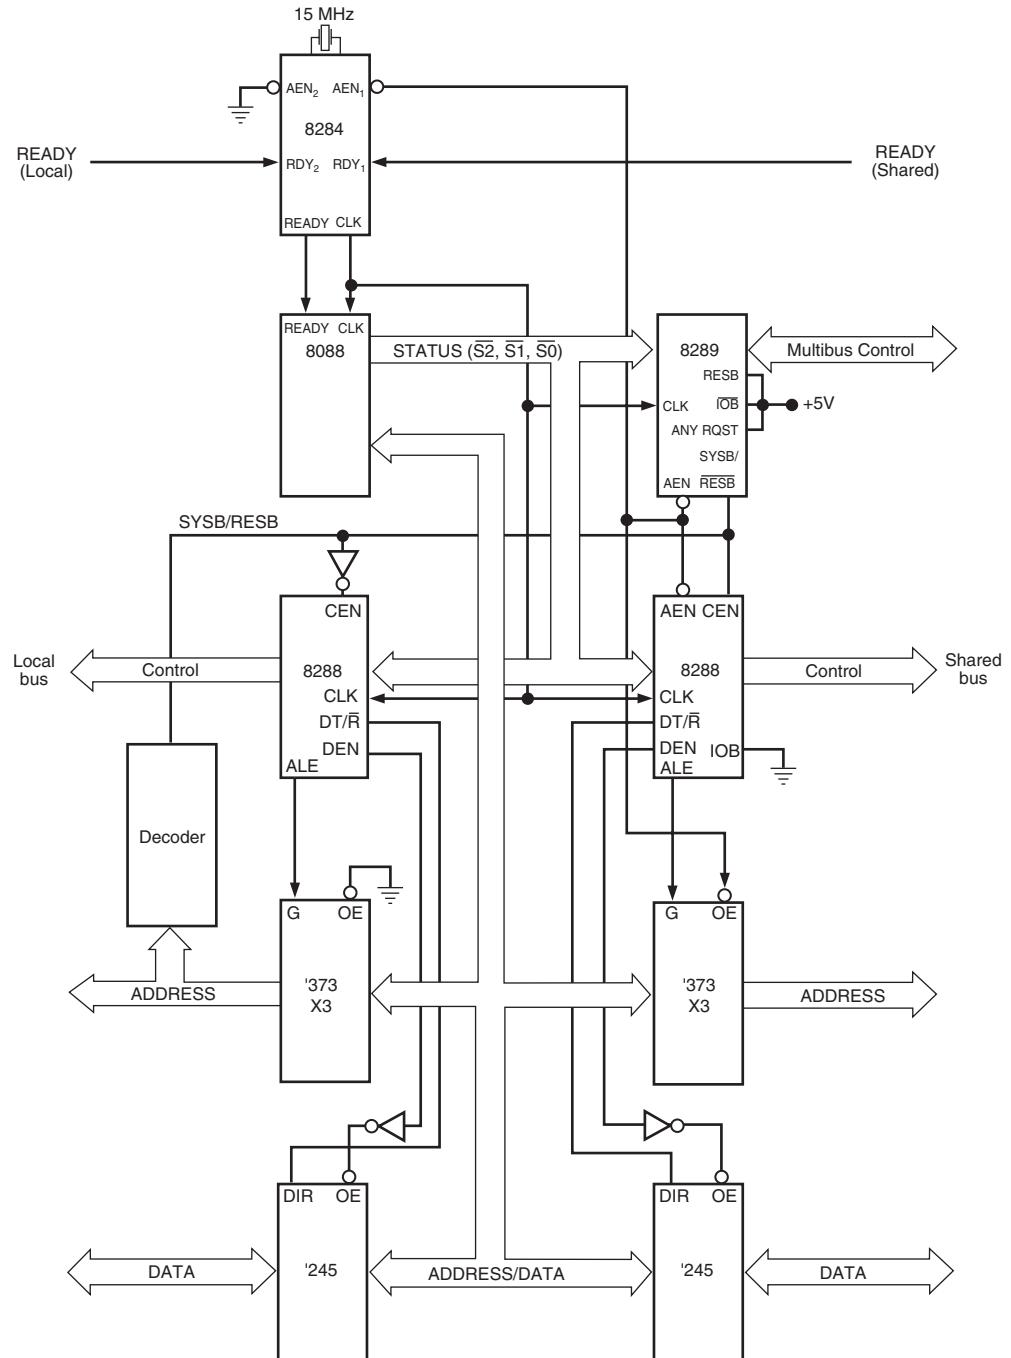
FIGURE 13–15   The 8088 operated in the remote mode, illustrating the local and shared bus connections.

containing memory and I/O. The 8237 DMA controller presented earlier in the chapter is an example of a remote bus master. The DMA controller gained access to the system memory and I/O space to cause a data transfer. Likewise, a remote bus master gains access to the shared bus for the same purpose. The difference is that the remote bus master microprocessor can execute variable software, whereas the DMA controller can only transfer data.

Access to the shared bus is accomplished by using the HOLD pin on the microprocessor for the DMA controller. Access to the shared bus for the remote bus master is accomplished via a bus arbiter, which functions to resolve priority between bus masters and allows only one device at a time to access the shared bus.

Notice in Figure 13–15 that the 8088 microprocessor has an interface to both a local, resident bus and the shared bus. This configuration allows the 8088 to access local memory and I/O or, through the bus arbiter and buffers, the shared bus. The task assigned to the microprocessor might be data communications. It may, after collecting a block of data from the communications interface, pass those data on to the shared bus and shared memory so that other microprocessors attached to the system can access the data. This allows many microprocessors to share common data. In the same manner, multiple microprocessors can be assigned various tasks in the system, drastically improving throughput.

The Bus Arbiter

Before Figure 13–15 can be fully understood, the operation of the bus arbiter must be grasped. The 8289 bus arbiter controls the interface of a bus master to a shared bus. Although the 8289 is not the only bus arbiter, it is designed to function with the 8086/8088 microprocessors, so it is presented here. Each bus master or microprocessor requires an arbiter for the interface to the shared bus, which Intel calls the Multibus and IBM calls the Micro Channel.

The shared bus is used only to pass information from one microprocessor to another; otherwise, the bus masters function in their own local bus modes by using their own local programs, memory, and I/O space. Microprocessors connected in this kind of system are often called parallel or distributed processors because they can execute software and perform tasks in parallel.
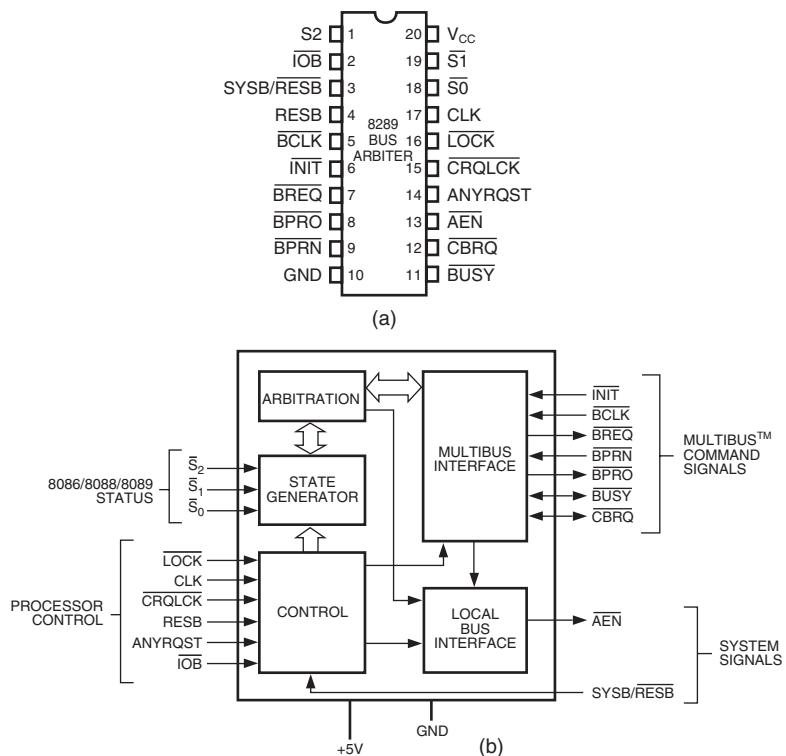
8289 Architecture.   Figure 13–16 illustrates the pin-out and block diagram of the 8289 bus arbiter. The left side of the block diagram depicts the connections to the microprocessor. The right side denotes the 8289 connection to the shared (remote) bus or Multibus.

The 8289 controls the shared bus by causing the READY input to the microprocessor to become a logic 0 (not ready) if access to the shared bus is denied. The blocking occurs whenever another microprocessor is accessing the shared bus. As a result, the microprocessor requesting access is blocked by the logic 0 applied to its READY input. When the READY pin is a logic 0, the microprocessor and its software wait until access to the shared bus is granted by the arbiter. In this manner, one microprocessor at a time gains access to the shared bus. No special instructions are required for bus arbitration with the 8289 bus arbiter because arbitration is accomplished strictly by the hardware.

Pin Definitions

$\overline{\text{AEN}}$	The **address enable** output causes the bus drivers in a system to switch to their three-state, high-impedance state.
ANYRQST	The **any request** input is a strapping option that prevents a lower-priority microprocessor from gaining access to the shared bus. If tied to a logic 0, normal arbitration occurs and a lower priority microprocessor can gain access to the shared bus if $\overline{\text{CBRQ}}$ is also a logic O.
$\overline{\text{BCLK}}$	The **bus clock** input synchronizes all shared-bus masters.
$\overline{\text{BPRN}}$	The **bus priority input** allows the 8289 to acquire the shared bus on the next falling edge of the $\overline{\text{BCLK}}$ signal.

FIGURE 13–16 The 8289
pin-out and block diagram.
(Courtesy of Intel
Corporation.)

(a)

(b)

$\overline{\text{BPRO}}$	The **bus priority output** is a signal that is used to resolve priority in a system that contains multiple bus masters.
$\overline{\text{BREQ}}$	The **bus request** output is used to request access to the shared bus.
BUSY	The **busy** input/output indicates, as an output, that an 8289 has acquired the shared bus. As an input, $\overline{\text{BUSY}}$ is used to detect that another 8289 has acquired the shared bus.
$\overline{\text{CBRQ}}$	The **common bus request** input/output is used when a lower priority microprocessor is asking for the use of the shared bus. As an output, $\overline{\text{CBRQ}}$ becomes a logic 0 whenever the 8289 requests the shared bus and remains low until the 8289 obtains access to the shared bus.
CLK	The **clock** input is generated by the 8284A clock generator and provides the internal timing source to the 8289.
$\overline{\text{CRQLCK}}$	The **common request lock** input prevents the 8289 from surrendering the shared bus to any of the 8289s in the system. This signal functions in conjunction with the $\overline{\text{CBRQ}}$ pin.
$\overline{\text{INIT}}$	The **initialization** input resets the 8289 and is normally connected to the system RESET signal.
$\overline{\text{IOB}}$	The **I/O bus** input selects whether the 8289 operates in a shared-bus system (if selected by RESB) with I/O ($\overline{\text{IOB}} = 0$) or with memory and I/O ($\overline{\text{IOB}} = 1$).
$\overline{\text{LOCK}}$	The **lock** input prevents the 8289 from allowing any other microprocessor from gaining access to the shared bus. An 8086/8088 instruction that contains a LOCK prefix will prevent other microprocessors from accessing the shared bus.
RESB	The **resident-bus** input is a strapping connection that allows the 8289 to operate in systems that have either a shared-bus or resident-bus system.

TABLE 13–2   8289 modes of operation.

Mode	Pin Connections
Single bus	$\overline{\text{IOB}}$ = 1 and RESB = 0
Resident bus	$\overline{\text{IOB}}$ = 1 and RESB = 1
I/O bus	$\overline{\text{IOB}}$ = 0 and RESB = 0
I/O bus and resident bus	$\overline{\text{IOB}}$ = 0 and RESB = 1

	If RESB is a logic 1, the 8289 is configured as a shared-bus master. If RESB is a logic 0, the 8289 is configured as a local-bus master. When configured as a shared-bus master, access is requested through the SYSB/$\overline{\text{RESB}}$ input pin.
S$_0$, S$_1$, and S$_2$	The **status** inputs initiate shared-bus requests and surrenders. These pins connect to the 8288 system bus controller status pins.
SYSB/$\overline{\text{RESB}}$	The **system bus/resident bus** input selects the shared-bus system when placed at a logic 1 or the resident local bus when placed at a logic 0.

General 8289 Operation.   As the pin descriptions demonstrate, the 8289 can be operated in three basic modes: (1) I/O peripheral-bus mode, (2) resident-bus mode, and (3) single-bus mode. See Table 13–2 for the connections required to operate the 8289 in these modes. In the I/O peripheral bus mode, all devices on the local bus are treated as I/O, including memory, and are accessed by all instructions. All memory references access the shared bus and all I/O access the resident-local bus. The resident-bus mode allows memory and I/O accesses on both the local and shared buses. Finally, the single-bus mode interfaces a microprocessor to a shared bus, but the microprocessor has no local memory or local I/O. In many systems, one microprocessor is set up as the shared-bus master (single-bus mode) to control the shared bus and become the shared-bus master. The shared-bus master controls the system through shared memory and I/O. Additional microprocessors are connected to the shared bus as resident- or I/O peripheral-bus masters. These additional bus masters usually perform independent tasks that are reported to the shared-bus master through the shared bus.

System Illustrating Single-Bus and Resident-Bus Connections.   Single-bus operation interfaces a microprocessor to a shared bus that contains both I/O and memory resources that are shared by other microprocessors. Figure 13–17 illustrates three 8088 microprocessors, each connected to a shared bus. Two of the three microprocessors operate in the resident-bus mode, while the third operates in the single-bus mode. Microprocessor A, in Figure 13–17, operates in the single-bus mode and has no local bus. This microprocessor accesses only the shared memory and I/O space. Microprocessor A is often referred to as the system-bus master because it is responsible for coordinating the main memory and I/O tasks. The remaining two microprocessors (B and C) are connected in the resident-bus mode, which allows them access to both the shared bus and their own local buses. These resident-bus microprocessors are used to perform tasks that are independent from the system-bus master. In fact, the only time that the system-bus master is interrupted from performing its tasks are when one of the two resident-bus microprocessors needs to transfer data between itself and the shared bus. This connection allows all three microprocessors to perform tasks simultaneously, yet data can be shared between microprocessors when needed.

In Figure 13–17, the bus master (A) allows the user to operate with a video terminal that allows the execution of programs and generally controls the system. Microprocessor B handles all telephone communications and passes this information to the shared memory in blocks. This means that microprocessor B waits for each character to be transmitted or received and controls the protocol used for the transfers. For example, suppose that a 1K-byte block of data is transmitted across the telephone interface at the rate of 100 characters per second. This means that the
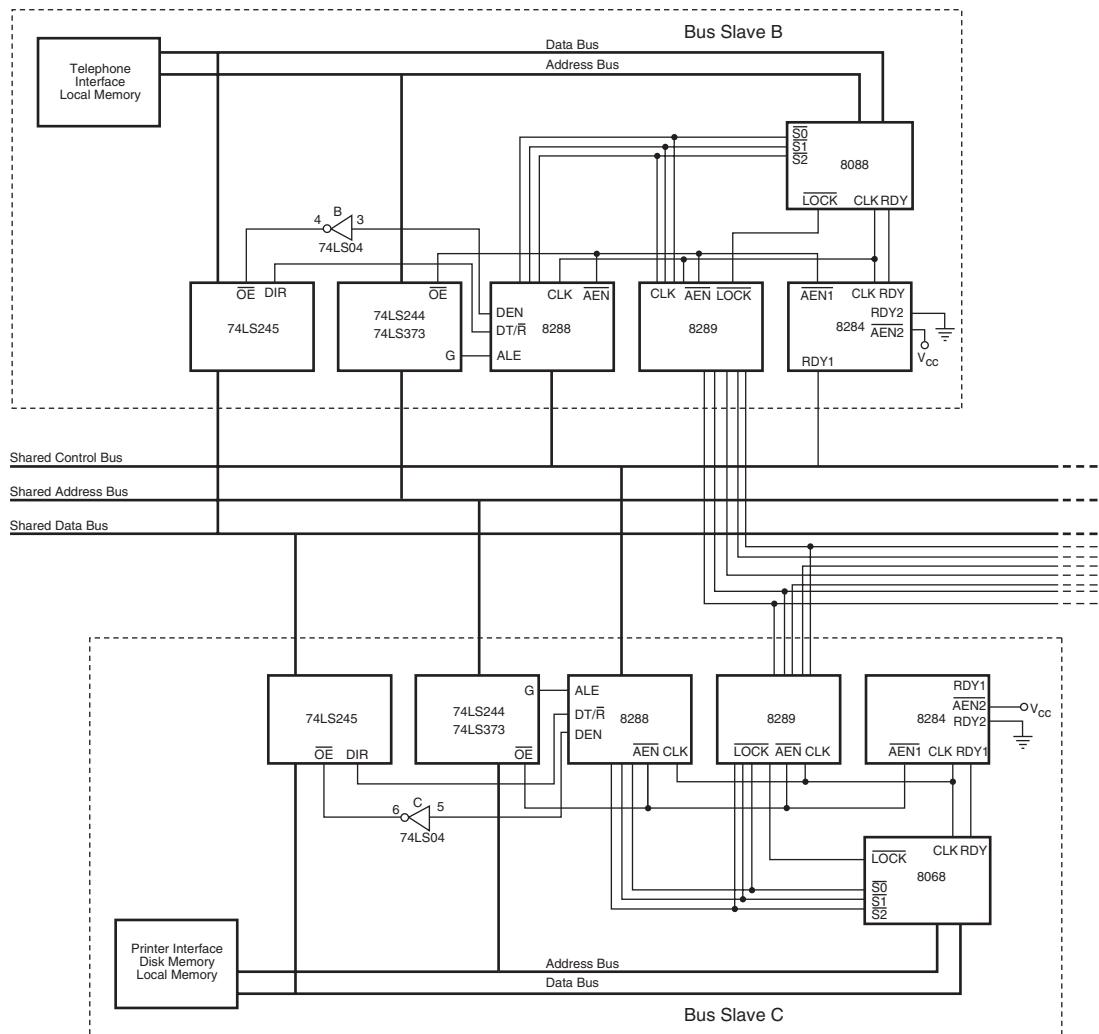
FIGURE 13–17   Three 8088 microprocessors that share a common bus system. Microprocessor A is the bus master in control of the shared memory and CRT terminal. Microprocessor B is a bus slave controlling its local telephone interface and memory. Microprocessor C is also a slave that controls a printer, disk memory system, and local memory.

transfer requires 10 seconds. Rather than tie up the bus master for 10 seconds, microprocessor B patiently performs the data transfer from its own local memory and the local communications interface. This frees the bus master for other tasks. The only time the microprocessor B interrupts the bus master is to transfer data between the shared memory and its local memory system. This data transfer between microprocessor B and the bus master requires only a few hundred microseconds.

Microprocessor C is used as a print spooler. Its only task is to print data on the printer. Whenever the bus master requires printed output, it transfers the task to microprocessor C. Microprocessor C then accesses the shared memory and captures the data to be printed and stores it in its own local memory. Data are then printed from the local memory, freeing the bus master to perform other tasks. This allows the system to execute a program with the bus master, transfer data through the communications interface with microprocessor B, and print information on the printer with microprocessor C. These tasks all execute simultaneously. There is no limit to the number of microprocessors connected to a system or the number of tasks performed simultaneously using this
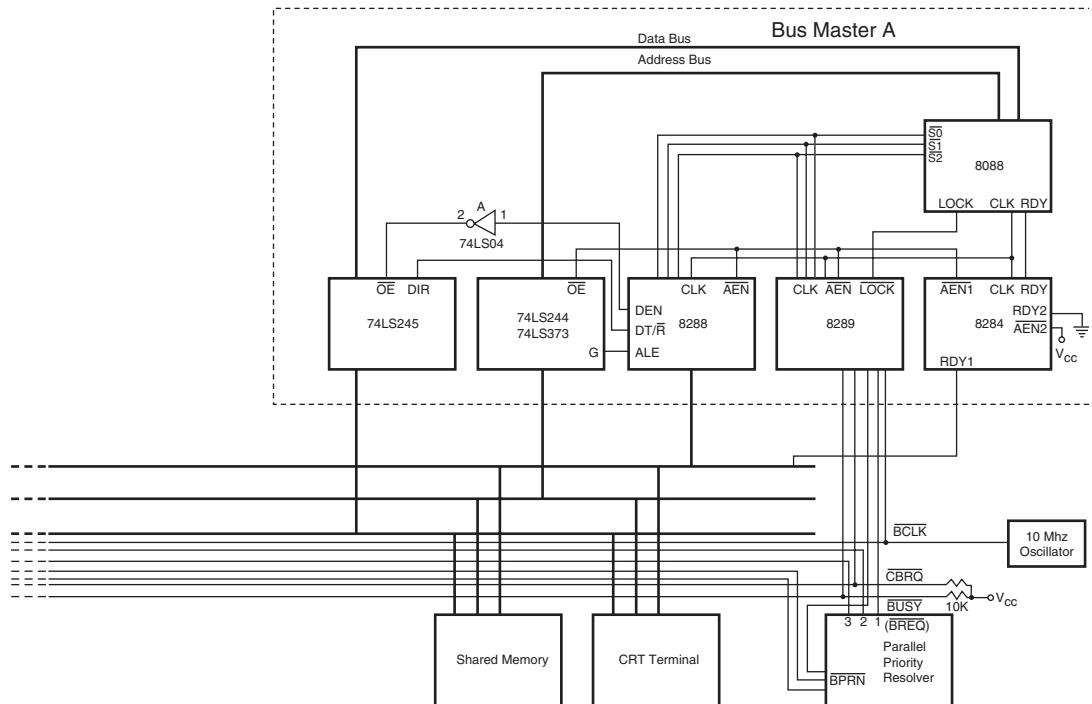
FIGURE 13–17 *(continued)*

technique. The only limit is that introduced by the system design and the designer's ingenuity. Lawrence Livermore Labs in California has a system that contains 4096 Pentium microprocessors.

13–4 DISK MEMORY SYSTEMS

Disk memory is used to store long-term data. Many types of disk storage systems are available today and they use magnetic media, except the optical disk memory that stores data on a plastic disk. Optical disk memory is either a CD-ROM (compact disk/read only memory) that is read, but never written, or a WORM (write once/read mostly) that is read most of the time, but can be written once by a laser beam. Also becoming available is optical disk memory that can be read and written many times, but there is still a limitation on the number of write operations allowed. The latest optical disk technology is called DVD (digital-versatile disk). The DVD (8.5G) is also available in high-resolution versions for video and data storage as Blu-ray (50G) or HD-DVD (30G). This section of the chapter provides an introduction to disk memory systems so that they may be used with computer systems. It also provides details of their operation.

Floppy Disk Memory

Once the most common and the most basic form of disk memory was the floppy, or flexible disk. Today the floppy is beginning to vanish and may completely disappear shortly in favor of the USB pen drive. The floppy disk magnetic recording media have been made available in three sizes: the 8″ standard, $5\frac{1}{4}$″ mini-floppy, and the $3\frac{1}{2}$″ micro-floppy. Today, the 8″ standard version and $5\frac{1}{4}$″ mini-floppy have all but disappeared, giving way to the micro-floppy disks and more

recently pen drives. The 8″ disk is too large and difficult to handle and stockpile. To solve this problem, industry developed the $5\frac{1}{4}$″ mini-floppy disk. Today, the micro-floppy disk has just about replaced the mini-floppy in newer systems because of its reduced size, ease of storage, and durability. Even so, systems are still marketed with the micro-floppy disk drives.

All disks and even the pen drives have several things in common. They are all organized so that data are stored in tracks. A track is a concentric ring of data that is stored on a surface of a disk. Figure 13–18 illustrates the surface of a $5\frac{1}{4}$″ mini-floppy disk, showing a track that is divided into sectors. A sector is a common subdivision of a track that is designed to hold a reasonable amount of data. In many systems, a sector holds either 512 or 1024 bytes of data. The size of a sector can vary from 128 bytes to the length of one entire track.

Notice in Figure 13–18 that there is a hole through the disk that is labeled an index hole. The index hole is designed so that the electronic system that reads the disk can find the beginning of a track and its first sector (00). Tracks are numbered from track 00, the outermost track, in increasing value toward the center or innermost track. Sectors are often numbered from sector 00 on the outermost track to whatever value is required to reach the innermost track and its last sector.

The 5 $\frac{1}{4}$″ Mini-floppy Disk.   Today, the $5\frac{1}{4}$″ floppy is very difficult to find and is used only with older microcomputer systems. Figure 13–19 illustrates this mini-floppy disk. The floppy disk is rotated at 300 RPM inside its semi-rigid plastic jacket. The head mechanism in a floppy disk drive makes physical contact with the surface of the disk, which eventually causes wear and damage to the disk.

Most mini-floppy disks are double-sided. This means that data are written on both the top and bottom surfaces of the disk. A set of tracks called a **cylinder** consists of one top and one bottom track. Cylinder 00, for example, consists of the outermost top and bottom tracks.

Floppy disk data are stored in the double-density format, which uses a recording technique called MFM (**modified frequency modulation**) to store the information. Double-density, double-sided (DSDD) disks are normally organized with 40 tracks of data on each side of the

FIGURE 13–18   The format of a $5\frac{1}{4}$″ mini-floppy disk.
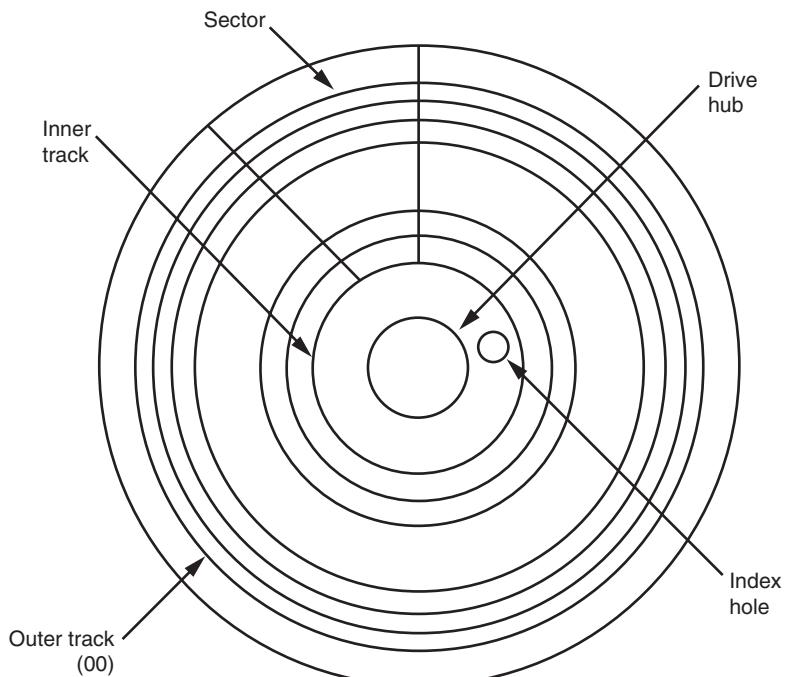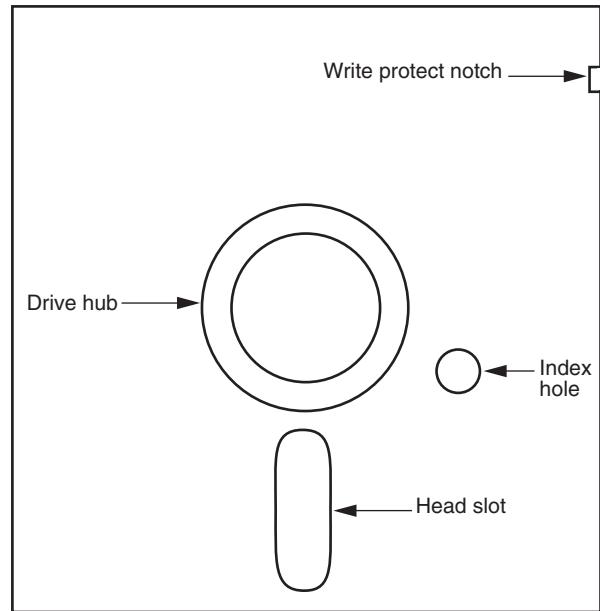
FIGURE 13–19 The $5\frac{1}{4}''$ mini-floppy disk.

disk. A double-density disk track is typically divided into nine sectors, with each sector containing 512 bytes of information. This means that the total capacity of a double-density, double-sided disk is 40 tracks per side $\times$ 2 sides $\times$ 9 sectors per track $\times$ 512 bytes per sector, or 368,640 (360K) bytes of information.

Also common are high-density (HD) mini-floppy disks. A high-density mini-floppy disk contains 80 tracks of information per side, with eight sectors per track. Each sector contains 1024 bytes of information. This gives the $5\frac{1}{4}''$ high-density, mini-floppy disk a total capacity of 80 tracks per side $\times$ 2 sides $\times$ 15 sectors per track $\times$ 512 bytes per sector, or 1,228,800 (approximately 1.2 M) bytes of information.

The magnetic recording technique used to store data on the surface of the disk is called non-return to zero (NRZ) recording. With NRZ recording, magnetic flux placed on the surface of the disk never returns to zero. Figure 13–20 illustrates the information stored in a portion of a track. It also shows how the magnetic field encodes the data. Note that arrows are used in this illustration to show the polarity of the magnetic field stored on the surface of the disk.

The main reason that this form of magnetic encoding was chosen is that it automatically erases old information when new information is recorded. If another technique were used, a separate erase head would be required. The mechanical alignment of a separate erase head and a
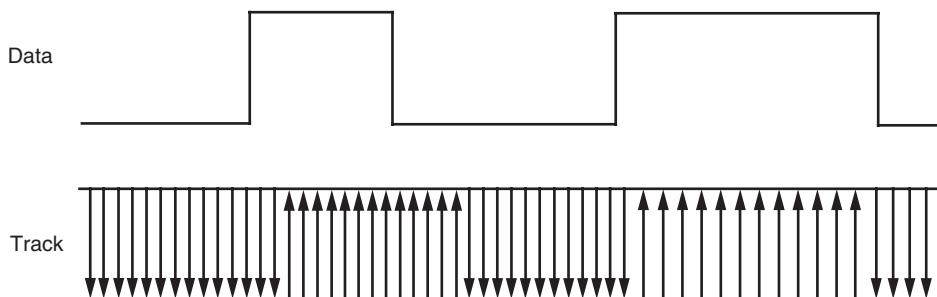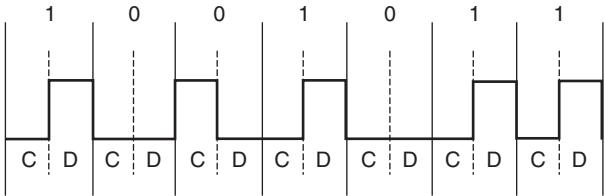
FIGURE 13–20 The non-return to zero (NRZ) recording technique.

FIGURE 13–21   Modified frequency modulation (MFM) used with disk memory.

separate read/write head is virtually impossible. The magnetic flux density of the NRZ signal is so intense that it completely saturates (magnetizes) the surface of the disk, erasing all prior data. It also ensures that information will not be affected by noise because the amplitude of the magnetic field contains no information. The information is stored in the placement of the changes of the magnetic field.

Data are stored in the form of MFM (modified frequency modulation) in modern floppy disk systems. The MFM recording technique stores data in the form illustrated in Figure 13–21. Notice that each bit time is 2.0 μs wide on a double-density disk. This means that data are recorded at the rate of 500,000 bits per second. Each 2.0 μs bit time is divided into two parts: One part is designated to hold a clock pulse and the other holds a data pulse. If a clock pulse is present, it is 1.0 μs wide, as is a data pulse. Clock and data pulses are never present at the same time in one bit period. (Note that high-density disk drives halve these times so that a bit time is 1.0 μs and a clock or data pulse is 0.5 μs wide. This also doubles the transfer rate to 1 million bits per second [1 Mbps]).

If a data pulse is present, the bit time represents a logic 1. If no data or no clock is present, the bit time represents a logic 0. If a clock pulse is present with no data pulse, the bit time also represents a logic 0. The rules followed when data are stored using MFM are as follows:

1. A data pulse is always stored for a logic l.
2. No data and no clock are stored for the first logic 0 in a string of logic 0s.
3. The second and subsequent logic 0s in a row contain a clock pulse, but no data pulse.
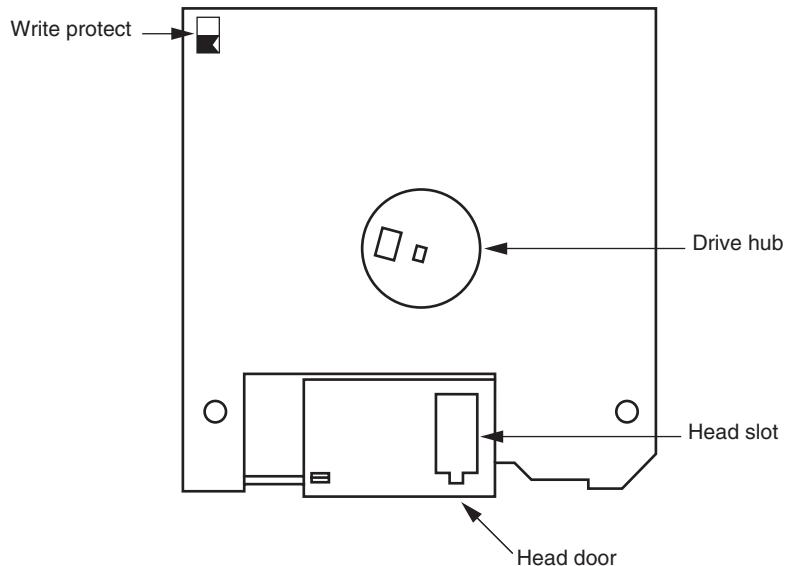
The reason that a clock is inserted as the second and subsequent zero in a row is to maintain synchronization as data are read from the disk. The electronics used to recapture the data from the disk drive use a phase-locked loop to generate a clock and a data window. The phase-locked loop needs a clock or data to maintain synchronized operation.

The 3 $\frac{1}{2}$″ Micro-Floppy Disk.   A popular disk size is the $3\frac{1}{2}$″ micro-floppy disk. Recently, this size floppy disk has begun to be replaced by the USB pen drive as the dominant transportable media. The micro-floppy disk is a much improved version of the mini-floppy disk described earlier. Figure 13–22 illustrates the $3\frac{1}{2}$″ micro-floppy disk.

Disk designers noticed several shortcomings of the mini-floppy, which is a scaled down version of the 8″ standard floppy, soon after it was released. Probably one of the biggest problems with the mini-floppy is that it is packaged in a semi-rigid plastic cover that bends easily. The micro-floppy is packaged in a rigid plastic jacket that will not bend easily. This provides a much greater degree of protection to the disk inside the jacket.

Another problem with the mini-floppy is the head slot that continually exposes the surface of the disk to contaminants. This problem is also corrected on the micro-floppy because it is constructed with a spring-loaded sliding head door. The head door remains closed until the disk is inserted into the drive. Once inside the drive, the drive mechanism slides open the door, exposing the surface of the disk to the read/write heads. This provides a great deal of protection to the surface of the micro-floppy disk.

FIGURE 13–22 The $3\frac{1}{2}''$ micro-floppy disk.

Write protect

Drive hub

Head slot

Head door

Yet another improvement is the sliding plastic write-protection mechanism on the micro-floppy disk. On the mini-floppy disk, a piece of tape was placed over a notch on the side of the jacket to prevent writing. This plastic tape easily became dislodged inside disk drives, causing problems. On the micro-floppy, an integrated plastic slide has replaced the tape write-protection mechanism. To write-protect (prevent writing) the micro-floppy disk, the plastic slide is moved to open the hole through the disk jacket. This allows light to strike a sensor that inhibits writing.

Still another improvement is the replacement of the index hole with a different drive mechanism. The drive mechanism on the mini-floppy allows the disk drive to grab the disk at any point. This requires an index hole so that the electronics can find the beginning of a track. The index hole is another trouble spot because it collects dirt and dust. The micro-floppy has a drive mechanism that is keyed so that it only fits one way inside the disk drive. The index hole is no longer required because of this keyed drive mechanism. Because of the sliding head mechanism and the fact that no index hole exists, the micro-floppy disk has no place to catch dust or dirt.

Two types of micro-floppy disks are widely available: the double-sided, double-density (DSDD) and the high-density (HD). The double-sided, double-density micro-floppy disk has 80 tracks per side, with each track containing nine sectors. Each sector contains 512 bytes of information. This allows 80 tracks per side × 2 sides × 9 sectors × 512 bytes per sector, or 737,280 (720K) bytes of data to be stored on a double-density, double-sided floppy disk.

The high-density, double-sided micro-floppy disk stores even more information. The high-density version has 80 tracks per side, but the number of sectors is doubled to 18 per track. This format still uses 512 bytes per sector, as did the double-density format. The total number of bytes on a high-density, double-sided micro-floppy disk is 80 tracks per side × 2 sides × 18 sectors per track × 512 bytes per sector, or 1,474,560 (1.44M) bytes of information.

Pen Drives

Pen drives, or flash drives, as they are often called, are replacements for floppy disk drives that use flash memory to store data. A driver, which is part of Windows (except for Windows 98), treats the pen drive as a floppy with tracks and sectors even though it really does not contain tracks and sectors. As with a floppy, the FAT system is used for the file structure. The memory in this type of drive is serial memory. When a pen drive is connected to the USB bus, the operating system recognizes it and allows data to be transferred between it and the computer.

Newer pen drives use the USB 2.0 bus specification to transfer data at a much higher rate of speed than the older USB 1.1 specification. Transfer speeds for USB 1.1 are a read speed of 750 KBps and a write speed of 450 KBps. The USB 2.0 pen drives have a transfer speed of about 48 MBps. The pen drive is currently available in sizes up to 4G bytes and has an erase cycle of up to 1,000,000 erases. The price is very reasonable when compared to the floppy disk.

Hard Disk Memory
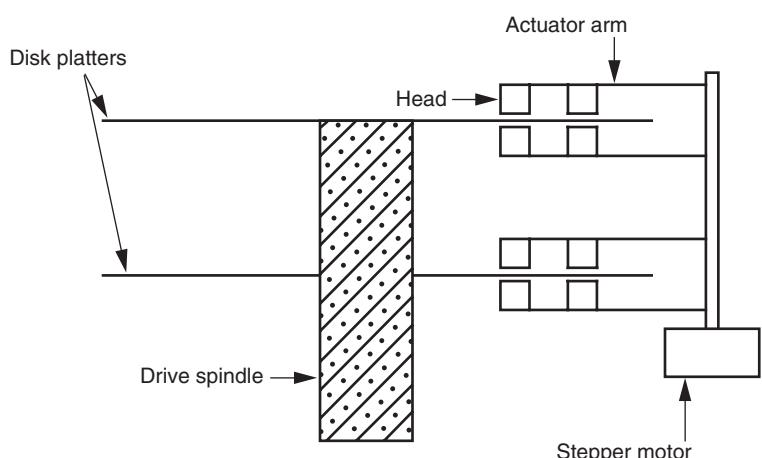
Larger disk memory is available in the form of the hard disk drive. The hard disk drive is often called a fixed disk because it is not removable like the floppy disk. A hard disk is also often called a rigid disk. The term *Winchester drive* is also used to describe a hard disk drive, but less commonly today. Hard disk memory has a much larger capacity than the floppy disk memory. Hard disk memory is available in sizes approaching 1 T (tera) bytes of data. Common, low-cost (less than $1 per gigabyte) sizes are presently 20G bytes to 500G bytes.

There are several differences between the floppy disk and the hard disk memory. The hard disk memory uses a flying head to store and read data from the surface of the disk. A flying head, which is very small and light, does not touch the surface of the disk. It flies above the surface on a film of air that is carried with the surface of the disk as it spins. The hard disk typically spins at 3000 to 15,000 RPM, which is many times faster than the floppy disk. This higher rotational speed allows the head to fly (just as an airplane flies) just over the top of the surface of the disk. This is an important feature because there is no wear on the hard disk's surface, as there is with the floppy disk.

Problems can arise because of flying heads. One problem is a head crash. If the power is abruptly interrupted or the hard disk drive is jarred, the head can crash onto the disk surface, which can damage the disk surface or the head. To help prevent crashes, some drive manufacturers have included a system that automatically parks the head when power is interrupted. This type of disk drive has auto-parking heads. When the heads are parked, they are moved to a safe landing zone (unused track) when the power is disconnected. Some drives are not auto-parking; they usually require a program that parks the heads on the innermost track before power is disconnected. The innermost track is a safe landing area because it is the very last track filled by the disk drive. Parking is the responsibility of the operator in this type of disk drive.

Another difference between a floppy disk drive and a hard disk drive is the number of heads and disk surfaces. A floppy disk drive has two heads, one for the upper surface and one for the lower surface. The hard disk drive may have up to eight disk surfaces (four platters), with up to two heads per surface. Each time that a new cylinder is obtained by moving the head assembly, 16 new tracks are available under the heads. See Figure 13–23, which illustrates a hard disk system.

FIGURE 13–23   A hard disk drive that uses four heads per platter.

Heads are moved from track to track by using either a stepper motor or a voice coil. The stepper motor is slow and noisy, while the voice coil mechanism is quiet and quick. Moving the head assembly requires one step per cylinder in a system that uses a stepper motor to position the heads. In a system that uses a voice coil, the heads can be moved many cylinders with one sweeping motion. This makes the disk drive faster when seeking new cylinders.

Another advantage of the voice coil system is that a servo mechanism can monitor the amplitude of the signal as it comes from the read head and make slight adjustments in the position of the heads. This is not possible with a stepper motor, which relies strictly on mechanics to position the head. Stepper-motor-type head positioning mechanisms can often become misaligned with use, while the voice coil mechanism corrects for any misalignment.

Hard disk drives often store information in sectors that are 512 bytes long. Data are addressed in clusters of eight or more sectors, which contain 4096 bytes (or more) on most hard disk drives. Hard disk drives use either MFM or RLL to store information. MFM is described with floppy disk drives. Run-length limited (RLL) is described here.

A typical older MFM hard disk drive uses 18 sectors per track so that 18 K bytes of data are stored per track. If a hard disk drive has a capacity of 40M bytes, it contains approximately 2280 tracks. If the disk drive has two heads, this means that it contains 1140 cylinders; if it contains four heads, then it has 570 cylinders. These specifications vary from disk drive to disk drive.

RLL Storage.   Run-length limited (RLL) disk drives use a different method for encoding the data than MFM. The term RLL means that the run of zeros (zeros in a row) is limited. A common RLL encoding scheme in use today is RLL 2,7. This means that the run of zeros is always between two and seven. Table 13–3 illustrates the coding used with standard RLL.

Data are first encoded by using Table 13–3 before being sent to the drive electronics for storage on the disk surface. Because of this encoding technique, it is possible to achieve a 50% increase in data storage on a disk drive when compared to MFM. The main difference is that the RLL drive often contains 27 tracks instead of the 18 found on the MFM drive. (Some RLL drives also use 35 sectors per track.)

In most cases, RLL encoding requires no change to the drive electronics or surface of the disk. The only difference is a slight decrease in the pulse width using RLL, which may require slightly finer oxide particles on the surface of the disk. Disk manufacturers test the surface of the disk and grade the disk drive as either an MFM-certified or an RLL-certified drive. Other than grading, there is no difference in the construction of the disk drive or the magnetic material that coats the surface of the disks.

Figure 13–24 shows a comparison of MFM data and RLL data. Notice that the amount of time (space) required to store RLL data is reduced when compared to MFM. Here 101001011 is coded in both MFM and RLL so that these two standards can be compared. Notice that the width

TABLE 13–3   Standard RLL 2,7 encoding.

Input Data Stream	RLL Output
000	000100
10	0100
010	100100
0010	00100100
11	1000
011	001000
0011	00001000

Data   RLL
101001011 = 010001001001001000
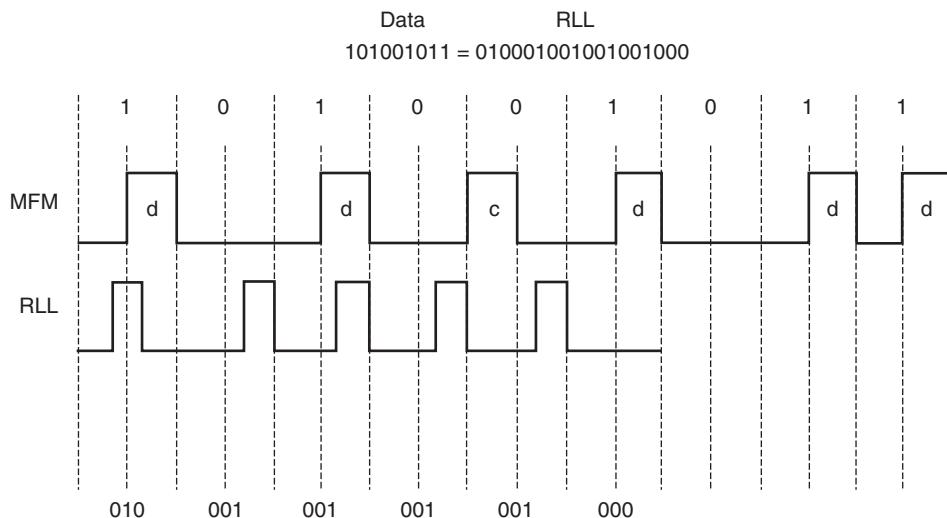
FIGURE 13–24   A comparison of MFM with RLL using data 101001011.

of the RLL signal has been reduced so that three pulses fit in the same space as a clock and a data pulse for MFM. A 40M-byte MFM disk can hold 60M bytes of RLL-encoded data. Besides holding more information, the RLL drive can be written and read at a higher rate.

All hard disk drives use today RLL encoding. There are a number of disk drive interfaces in use today. The oldest is the ST-506 interface, which uses either MFM or RLL data. A disk system using this interface is also called either MFM or RLL disk system. Newer standards are also found in use today, which include ESDI, SCSI, and IDE. All of these newer standards use RLL, even though they normally do not call attention to it. The main difference is the interface between the computer and the disk drive. The IDE system is becoming the standard hard disk memory interface.

The enhanced small disk interface (ESDI) system, which has disappeared, is capable of transferring data between itself and the computer at rates approaching 10M bytes per second. An ST-506 interface can approach a transfer rate of 860K bytes per second.

The small computer system interface (SCSI) system is also in use because it allows up to seven different disk or other interfaces to be connected to the computer through same interface controller. SCSI is found in some PC-type computers and also in the Apple Macintosh system. An improved version, SCSI-II, has started to appear in some systems. In the future, this interface may be replaced with IDE in most applications.

Today one of the most common systems is the **integrated drive electronics** (IDE) system, which incorporates the disk controller in the disk drive and attaches the disk drive to the host system through a small interface cable. This allows many disk drives to be connected to a system without worrying about bus conflicts or controller conflicts. IDE drives are found in newer IBM PS-2 systems and many clones. Even Apple computer systems are starting to be found with IDE drives in place of the SCSI drives found in older Apple computers. The IDE interface is also capable of driving other I/O devices besides the hard disk. This interface also usually contains at least a 256K- to 8M-byte cache memory for disk data. The cache speeds disk transfers. Common access times for an IDE drive are often less than 8 ms, whereas the access time for a floppy-disk is about 200 ms.

Sometimes IDE is also called ATA. ATA is an acronym for **AT attachment** where the AT means the Advanced Technology computer. The latest system is the serial ATA interface or SATA. This interface transfers serial data at rates of 150 MBps (or 300 MBps for SATA2), which

is faster than any IDE interface. Not yet released is SATA3, which transfers data at a rate of 600 MBps. The transfer rate is higher because the logic 1 level is no longer 5.0 V. In the SATA interface, the logic 1 level is 0.5 V, which allows data to be transferred at higher rates because it takes less time for the signal to rise to 0.5 V than it takes to rise to 5.0 V. Speeds of this interface should eventually reach 600 MBps with SATA3.

Optical Disk Memory

Optical disk memory (see Figure 13–25) is commonly available in two forms: the CD-ROM (compact disk/read only memory) and the WORM (write once/read mostly). The CD-ROM is the lowest cost optical disk, but it suffers from lack of speed. Access times for a CD-ROM are typically 300 ms or longer, about the same as a floppy disk. (Note that slower CD-ROM devices are on the market and should be avoided.) Hard disk magnetic memory can have access times as little as 11 ms. A CD-ROM stores 660M bytes of data, or a combination of data and musical passages. As systems develop and become more visually active, the use of the CD-ROM drive will become even more common.

The WORM drive sees far more commercial application than the CD-ROM. The problem is that its application is very specialized due to the nature of the WORM. Because data may be written only once, the main application is in the banking industry, insurance industry, and other massive data-storing organizations. The WORM is normally used to form an audit trail of transactions that are spooled onto the WORM and retrieved only during an audit. You might call the WORM an archiving device.

Many WORM and read/write optical disk memory systems are interfaced to the microprocessor by using the SCSI or ESDI interface standards used with hard disk memory. The difference is
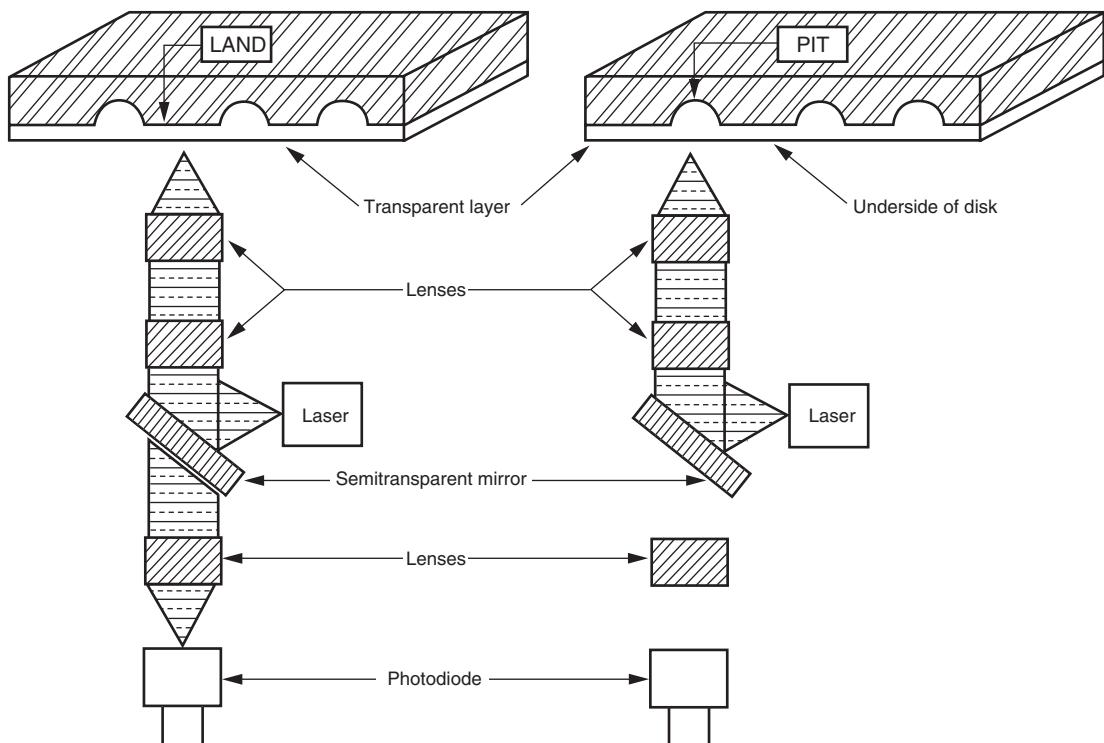
FIGURE 13–25   The optical CD-ROM memory system.

that the current optical disk drives are no faster than the most floppy drives. Some CD-ROM drives are interfaced to the microprocessor through proprietary interfaces that are not compatible with other disk drives.

The main advantage of the optical disk is its durability. Because a solid-state laser beam is used to read the data from the disk, and the focus point is below a protective plastic coating, the surface of the disk may contain small scratches and dirt particles and still be read correctly. This feature allows less care of the optical disk than a comparable floppy disk. About the only way to destroy data on an optical disk is to break it or deeply scar it.

The read/write CD-ROM drive is here and its cost is dropping rapidly. In the near future, we should start seeing the read/write CD-ROM replacing floppy disk drives. The main advantage is the vast storage available on the read/write CD-ROM. Soon, the format will change so that many G bytes of data will be available. The new versatile read/write CD-ROM, called a DVD, became available in late 1996 or early 1997. The DVD functions exactly like the CD-ROM except that the bit density is much higher. The CD-ROM stores 660M bytes of data, while the current-genre DVD stores 4.7G bytes or 9.4G bytes, depending on the current standard. Look for the DVD to eventually replace the CD-ROM format completely, at least for computer data storage, but maybe not for audio.

New to this technology are the Blu-ray DVD from Sony Corporation and the HD-DVD from Toshiba Corporation. The Blu-ray DVD has a capacity of 50 GB and the HD-DVD has a capacity of 30 GB. Which format will eventually become the standard is conjecture. The main advantage is to video, where high-resolution HD video (1080p) can be stored on either Blu-ray or HD-DVD. Because there are rumors of a higher resolution video standard in the future, even Blu-ray and HD-DVD may be replaced by some other technology. The big change from older DVDs and the newer technology is a switch from a red laser to a blue laser. A blue laser has a higher frequency, which means that it can read more information per second from the DVD, hence a high storage density.

13–5   VIDEO DISPLAYS

Modern video displays are OEM (original equipment manufacturer) devices that are usually purchased and incorporated into a system. Today, there are many different types of video displays available in either color or monochrome versions.

Monochrome versions usually display information using amber, green, or paper-white displays. The paper-white displays were once extremely popular for many applications. The most common of these applications are desktop publishing and computer-aided drafting (CAD).

The color displays are more diverse and have all but replaced the black-and-white display. Color display systems are available that accept information as a composite video signal, much like your home television, as TIL voltage level signals (0 or 5 V), and as analog signals (0–0.7 V). Composite video displays are disappearing because the available resolution is too low. Today, many applications require high-resolution graphics that cannot be displayed on a composite display such as a home television receiver. Early composite video displays were found with Commodore 64, Apple 2, and similar computer systems.

Video Signals

Figure 13–26 illustrates the signal sent to a composite video display. This signal is composed of several parts that are required for this type of display. The signals illustrated represent the signals sent to a color composite-video monitor. Notice that these signals include not only video, but also include sync pulses, sync pedestals, and a color burst. Notice that no audio signal is illustrated because one
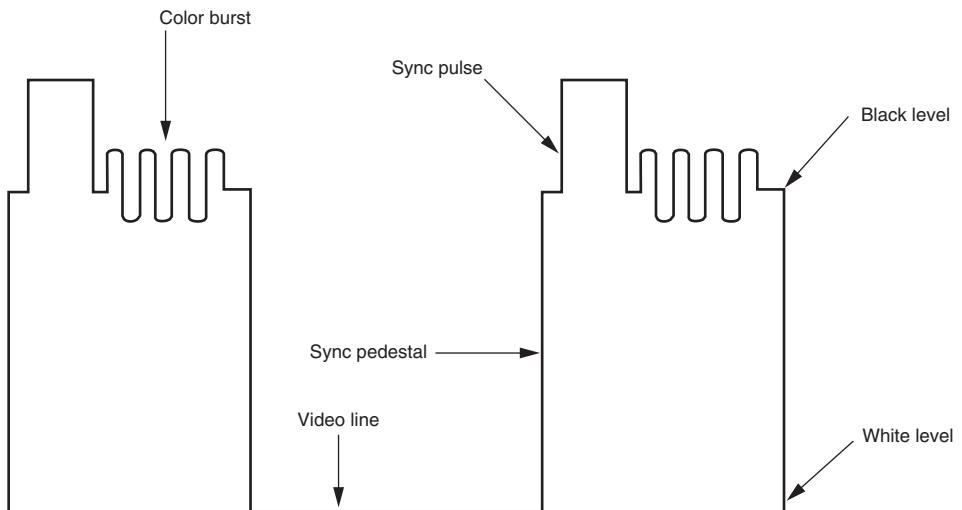
FIGURE 13–26   The composite video signal.

often does not exist. Rather than include audio with the composite video signal, audio is developed in the computer and output from a speaker inside the computer cabinet. It can also be developed by a sound system and output in stereo to external speakers. The major disadvantages of the composite video display are the resolution and color limitations. Composite video signals were designed to emulate television video signals so that a home television receiver could function as a video monitor.

Most modern video systems use direct video signals that are generated with separate sync signals. In a direct video system, video information is passed to the monitor through a cable that uses separate lines for video and also synchronization pulses. Recall that these signals were combined in a composite video signal.

A monochrome (one color) monitor uses one wire for video, one for horizontal sync, and one for vertical sync. Often, these are the only signal wires found. A color video monitor uses three video signals. One signal represents red, another green, and the third blue. These monitors are often called RGB monitors for the video primary colors of light: red (R), green (G), and blue (B).

The TTL RGB Monitor

The RGB monitor is available as either an analog or TTL monitor. The RGB monitor uses TTL level signals (0 or 5 V) as video inputs and a fourth line called intensity to allow a change in intensity. The RGB video TTL display can display a total of 16 different colors. The TTL RGB monitor is used in the CGA (color graphics adapter) system found in older computer systems.

Table 13–4 lists these 16 colors and also the TTL signals present to generate them. Eight of the 16 colors are generated at high intensity and the other eight at low intensity. The three video colors are red, green, and blue. These are primary colors of light. The secondary colors are cyan, magenta, and yellow. Cyan is a combination of blue and green video signals, and is blue-green in color. Magenta is a combination of blue and red video signals, and is a purple color.

Yellow (high intensity) and brown (low intensity) are both a combination of red and green video signals. If additional colors are desired, TTL video is not normally used. A scheme was developed by using low- and medium-color TTL video signals, which provided 32 colors, but it proved to have little application and never found widespread use in the field.

Figure 13–27 illustrates the connector most often found on the TTL RGB monitor or a TTL monochrome monitor. The connector illustrated is a 9-pin connector. Two of the connections are used for ground, three for video, two for synchronization or retrace signals, and one for

TABLE 13–4   The 16 colors found in a TTL display.

Intensity	Red	Green	Blue	Color
0	0	0	0	Black
0	0	0	1	Blue
0	0	1	0	Green
0	0	1	1	Cyan
0	1	0	0	Red
0	1	0	1	Magenta
0	1	1	0	Brown
0	1	1	1	White
1	0	0	0	Gray
1	0	0	1	Bright Blue
1	0	1	0	Bright Green
1	0	1	1	Bright Cyan
1	1	0	0	Bright Red
1	1	0	1	Bright Magenta
1	1	1	0	Yellow
1	1	1	1	Bright White

intensity. Notice that pin 7 is labeled normal video. This is the pin used on a monochrome monitor for the luminance or brightness signal. Monochrome TTL monitors use the same 9-pin connector as RGB TTL monitors.

The Analog RGB Monitor

In order to display more than 16 colors, an analog video display is required. These are often called analog RGB monitors. Analog RGB monitors still have three video input signals, but don't have the intensity input. Because the video signals are analog signals instead of two-level TTL signals, they are at any voltage level between 0.0 V and 0.7 V, which allows an infinite number of colors to be displayed. This is because an infinite number of voltage levels between the minimum and maximum could be generated. In practice, a finite number of levels are generated. This is usually either 256K, 16M, or 24M colors, depending on the standard.

Figure 13–28 illustrates the connector used for an analog RGB or analog monochrome monitor. Notice that the connector has 15 pins and supports both RGB and monochrome analog displays. The way data are displayed on an analog RGB monitor depends upon the interface standard used with the monitor. Pin 9 is a key, which means that no hole exists on the female connector for this pin.

Another type of connector for the analog RGB monitor that is becoming common is called the DVI-D (digital visual interface) connector. The -D is for digital and is the most common interface of

FIGURE 13–27   The 9-pin connector found on a TTL monitor.

Pin	Function
1	Ground
2	Ground
3	Red video
4	Green video
5	Blue video
6	Intensity
7	Normal video
8	Horizontal retrace
9	Vertical retrace
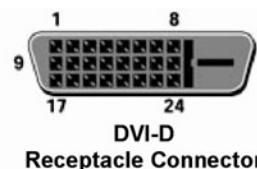
Pin	Function
1	Red video
2	Green video (monochrome video)
3	Blue video
4	Ground
5	Ground
6	Red ground
7	Green ground (monochrome ground)
8	Blue ground
9	Blocked as a key
10	Ground
11	Color detect (ground on a color monitor)
12	Monochrome detect (ground on a monochrome monitor)
13	Horizontal retrace
14	Vertical retrace
15	Ground

FIGURE 13–28   The 15-pin connector found on an analog monitor.

this type. Figure 13–29 illustrates the female connector found on newer monitors and video cards. Also found on television and video equipment is the HDMI (high-definition multimedia interface) connector. This has not made its way to digital video cards, but will probably appear in the future. Eventually all video equipment will use the HDMI connector for its connection.

Most analog displays use a digital-to-analog converter (DAC) to generate each color video voltage. A common standard uses a 8-bit DAC for each video signal to generate 256 different voltage levels between 0 V and 0.7 V. There are 256 different red video levels, 256 different green video levels, and 256 different blue video levels. This allows $256 \times 256 \times 256$, or 16,777,216 (16 M) colors to be displayed.

Figure 13–30 illustrates the video generation circuit employed in many common video standards such as the short-lived EGA (enhanced graphics adapter) and VGA (variable graphics array), as used with an IBM PC. This circuit is used to generate VGA video. Notice that each color is generated with an 18-bit digital code. Six of the 18 bits are used to generate each video color voltage when applied to the inputs of a 6-bit DAC.

**DVI-D
Receptacle Connector**

DIGITAL-ONLY CONNECTOR PIN ASSIGNMENTS					
Pin	Signal Assignment	Pin	Signal Assignment	Pin	Signal Assignment
1	Data2-	9	Data1-	17	Data0-
2	Data2+	10	Data1+	18	Data0+
3	Data2/4 Shield	11	Data1/3 Shield	19	Data0/5 Shield
4	Data4-	12	Data3-	20	Data5-
5	Data4+	13	Data3+	21	Data5+
6	DDC Clock	14	+5V Power	22	Clock Shield
7	DDC Data	15	Ground (for +5V)	23	Clock+
8	No Connect	16	Hot Plug Detect	24	Clock-

FIGURE 13–29   The DVI-D interface found on many newer monitors and video cards.
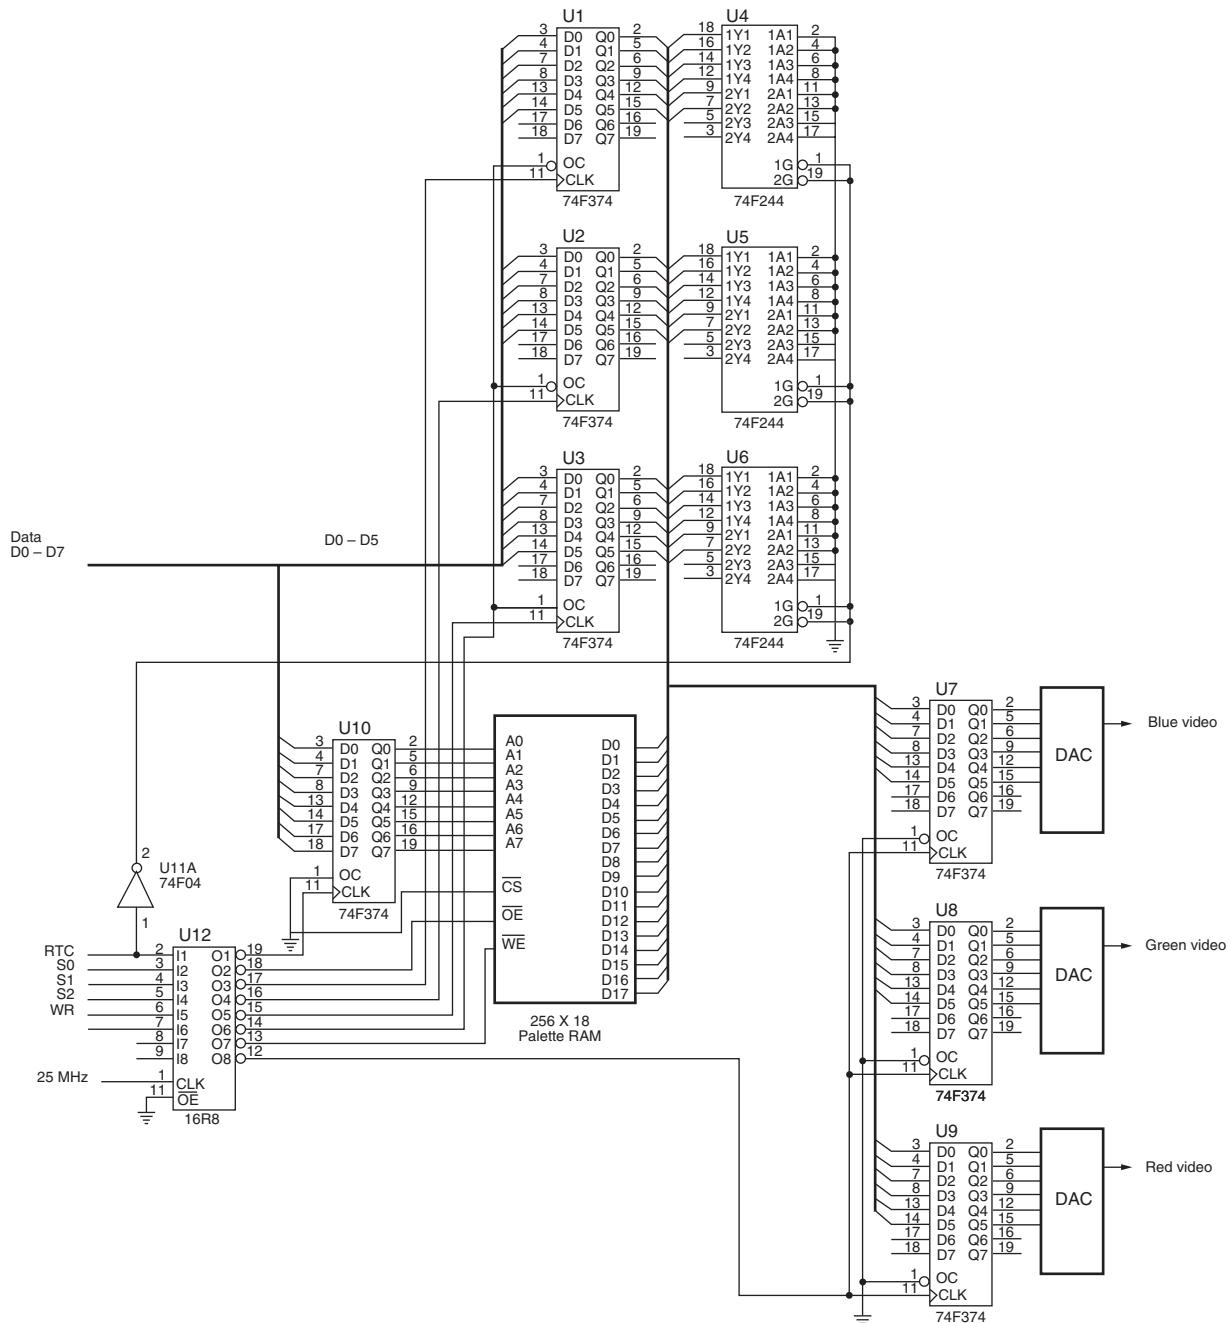
FIGURE 13–30   Generation of VGA video signals.

A high-speed palette SRAM (access time of less than 40 ns) is used to store 256 different 18-bit codes that represent 256 different hues. This 18-bit code is applied to the digital-to-analog converters. The address input to the SRAM selects one of the 256 colors stored as 18-bit binary codes. This system allows 256 colors out of a possible 256K colors to be displayed at one time. In order to select any of 256 colors, an 8-bit code that is stored in the computer's video display RAM is used to specify the color of a picture element. If more colors are used in a system, the

code must be wider. For example, a system that displays 1024 colors out of 256K colors requires a 10-bit code to address the SRAM that contains 1024 locations, each containing an 18-bit color code. Some newer systems use a larger palette SRAM to store up to 64K of different color codes.

Whenever a color is placed on the video display, provided that RTC is a logic 0, the system sends the 8-bit code that represents a color to the D_0–D_7 connections. The PLD then generates a clock pulse for U_{10}, which latches the color code. After 40 ns (one 25 MHz clock), the PLD generates a clock pulse for the DAC latches (U_7, U_8, and U_9). This amount of time is required for the palette SRAM to look up the 18-bit contents of the memory location selected by U_{10}. Once the color code (18-bit) is latched into U_7–U_9, the three DACs convert it to three video voltages for the monitor. This process is repeated for each 40-ns-wide picture element (pixel) that is displayed. The pixel is 40 ns wide because a 25 MHz clock is used in this system. Higher resolution is attainable if a higher clock frequency is used with the system.

If the color codes (18-bits) stored in the SRAM must be changed, this is always accomplished during retrace when RTC is a logic 1. This prevents any video noise from disrupting the image displayed on the monitor.

In order to change a color, the system uses the S_0, S_1, and S_2 inputs of the PLD to select U_1, U_2, U_3, and U_{10}. First, the address of the color to be changed is sent to latch U_{10}, which addresses a location in the palette SRAM. Next, each new video color is loaded into U_1, U_2, and U_3. Finally, the PLD generates a write pulse for the $\overline{WE}$ input to the SRAM to write the new color code into the palette SRAM.

Retrace occurs 70.1 times per second in the vertical direction and 31,500 times per second in the horizontal direction for a 640×480 display. During retrace, the video signal voltage sent to the display must be 0 V, which causes black to be displayed during the retrace. Retrace itself is used to move the electron beam to the upper left-hand corner for vertical retrace and to the left margin of the screen for horizontal retrace.

The circuit illustrated causes U_4–U_6 buffers to be enabled so that they apply 000000 each to the DAC latch for retrace. The DAC latches capture this code and generate 0 V for each video color signal to blank the screen. By definition, 0 V is considered to be the black level for video and 0.7 V is considered to be the full intensity on a video color signal.
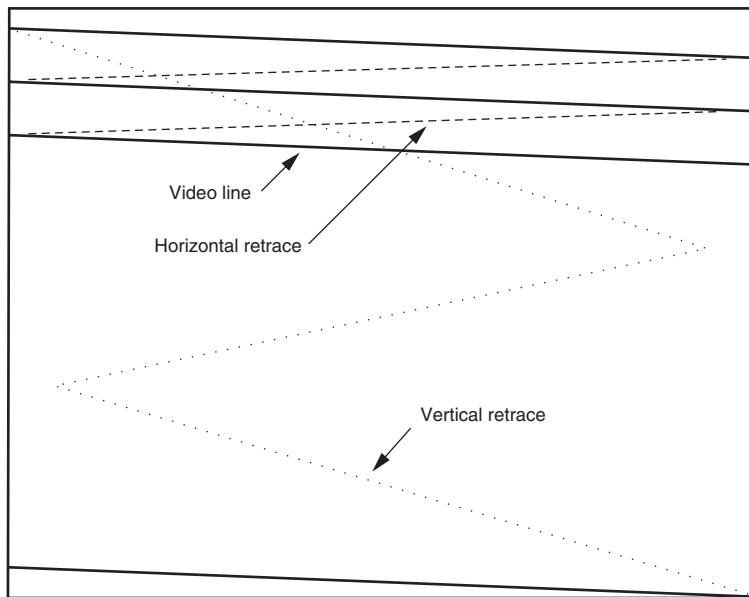
The resolution of the display, for example, 640×480, determines the amount of memory required for the video interface card. If this resolution is used with a 256-color display (8 bits per pixel), then 640×480 bytes of memory (307,200) are required to store all of the pixels for the display. Higher resolution displays are possible, but, as you can imagine, even more memory is required. A 640×480 display has 480 video raster lines and 640 pixels per line. A **raster line** is the horizontal line of video information that is displayed on the monitor. A pixel is the smallest subdivision of this horizontal line.

Figure 13–31 illustrates the video display, showing the video lines and retrace. The slant of each video line in this illustration is greatly exaggerated, as is the spacing between lines. This illustration shows retrace in both the vertical and horizontal directions. In the case of a VGA display, as described, the vertical retrace occurs exactly 70.1 times per second and the horizontal retrace occurs exactly 31,500 times per second.

In order to generate 640 pixels across one line, it takes 40 ns $\times$ 640, or 25.6 μs. A horizontal time of 31,500 Hz allows a horizontal line time of 1/31,500, or 31.746 μs. The difference between these two times is the retrace time allowed to the monitor. (The Apple Macintosh has a horizontal line time of 28.57 μs.)

Because the vertical retrace repetition rate is 70.1 Hz, the number of lines generated is determined by dividing the vertical time into the horizontal time. In the case of a VGA display (a 640×400 display), this is 449.358 lines. Only 400 of these lines are used to display information; the rest are lost during the retrace. Because 49.358 lines are lost during the retrace, the retrace time is 49.358 $\times$ 31.766 μs, or 1568 μs. It is during this relatively large amount of time that the color palette SRAM is changed or the display memory system is updated for a new video display.

FIGURE 13–31   A video
screen illustrating the raster
lines and retrace.

In the Apple Macintosh computer (640×480), the number of lines generated is 525 lines. Of the total number of lines, 45 are lost during vertical retrace.

Other display resolutions are 800×600 and 1024×768. The 800×600 SVGA (super VGA) display is ideal for a 14″ color monitor, while the 1024×768 EVGA or XVGA (extended VGA) is ideal for a 21″ or 25″ monitor used in CAD systems. These resolutions sound like just another set of numbers, but realize that an average home television receiver has a resolution approximately 400×300. The high-resolution display available on computer systems is much clearer than that available as home television. A resolution of 1024×768 approaches that found in 35 mm film. The only disadvantage of the video display on a computer screen is the number of colors displayed at a time, but as time passes, this will surely improve. Additional colors allow the image to appear more realistically because of subtle shadings that are required for a true high-quality, lifelike image.

If a display system operates with a 60 Hz vertical time and a 15,600 Hz horizontal time, the number of lines generated is 15,600/60, or 260 lines. The number of usable lines in this system is most likely 240, where 20 are lost during vertical retrace. It is clear that the number of scanning lines is adjustable by changing the vertical and horizontal scanning rates. The vertical scanning rate must be greater than or equal to 50 Hz or flickering will occur. The vertical rate must not be higher than about 75 Hz or problems with the vertical deflection coil may occur. The electron beam in a monitor is positioned by an electrical magnetic field generated by coils in a yoke that surrounds the neck of the picture tube. Because the magnetic field is generated by coils, the frequency of the signal applied to the coil is limited.

The horizontal scanning rate is also limited by the physical design of the coils in the yoke. Because of this, it is normal to find the frequency applied to the horizontal coils within a narrow range. This is usually 30,000 Hz–37,000 Hz or 15,000 Hz–17,000 Hz. Some newer monitors are called multisync monitors because the deflection coil is taped so that it can be driven with different deflection frequencies. Sometimes, both the vertical and horizontal coils are both taped for different vertical and horizontal scanning rates.

High-resolution displays use either interlaced or noninterlaced scanning. The noninterlaced scanning system is used in all standards except the highest. In the interlaced system, the video image is displayed by drawing half the image first with all of the odd scanning lines, then the other half is drawn using the even scanning lines. Obviously, this system is more complex and is only more efficient because the scanning frequencies are reduced by 50% in an interlaced

system. For example, a video system that uses 60 Hz for the vertical scanning frequency and 15,720 Hz for the horizontal frequency generates 262 (15,720/60) lines of video at the rate of 60 full frames per second. If the horizontal frequency is changed slightly to 15,750 Hz, 262.5 (15,750/60) lines are generated, so two full sweeps are required to draw one complete picture of 525 video lines. Notice how just a slight change in horizontal frequency doubled the number of raster lines.

13–6　SUMMARY

1. The HOLD input is used to request a DMA action, and the HLDA output signals that the hold is in effect. When a logic 1 is placed on the HOLD input, the microprocessor (1) stops executing the program; (2) places its address, data, and control bus at their high-impedance state; and (3) signals that the hold is in effect by placing a logic 1 on the HLDA pin.

2. A DMA read operation transfers data from a memory location to an external I/O device. A DMA write operation transfers data from an I/O device into the memory. Also available is a memory-to-memory transfer that allows data to be transferred between two memory locations by using DMA techniques.

3. The 8237 direct memory access (DMA) controller is a four-channel device that can be expanded to include an additional channel of DMA.

4. Disk memory comes in the form of floppy disk storage that is found as $3\frac{1}{2}$" micro-floppy disks. Disks are found as double-sided, double-density (DSDD), or as high-density (HD) storage devices. The DSDD $3\frac{1}{2}$" disk stores 720K bytes of data and the HD $3\frac{1}{2}$" disk stores 1.44M bytes of data.

5. Floppy disk memory data are stored using NRZ (non-return to zero) recording. This method saturates the disk with one polarity of magnetic energy for a logic 1 and the opposite polarity for a logic 0. In either case, the magnetic field never returns to 0. This technique eliminates the need for a separate erase head.

6. Data are recorded on disks by using either modified frequency modulation (MFM) or run-length limited (RLL) encoding schemes. The MFM scheme records a data pulse for a logic 1, no data or clock for the first logic 0 of a string of zeros, and a clock pulse for the second and subsequent logic 0 in a string of zeros. The RLL scheme encodes data so that 50% more information can be packed onto the same disk area. Most modern disk memory systems use the RLL encoding scheme.

7. Video monitors are either TTL or analog. The TTL monitor uses two discrete voltage levels of 0 V and 5.0 V. The analog monitor uses an infinite number of voltage levels between 0.0 V and 0.7 V. The analog monitor can display an infinite number of video levels, while the TTL monitor is limited to two video levels.

8. The color TTL monitor displays 16 different colors. This is accomplished through three video signals (red, green, and blue) and an intensity input. The analog color monitor can display an infinite number of colors through its three video inputs. In practice, the most common form of color analog display system (VGA) can display 16 M different colors.

9. The video standards found today include VGA (640×480), SVGA (800×600), and EVGA or XVGA (1024×768). In all three cases, the video information can be 16M colors.

13–7　QUESTIONS AND PROBLEMS

1. Which microprocessor pins are used to request and acknowledge a DMA transfer?
2. Explain what happens whenever a logic 1 is placed on the HOLD input pin.

3. A DMA read transfers data from _____ to _____.
4. A DMA write transfers data from _____ to _____.
5. The DMA controller selects the memory location used for a DMA transfer through what bus signals?
6. The DMA controller selects the I/O device used during a DMA transfer by which pin?
7. What is a memory-to-memory DMA transfer?
8. Describe the effect on the microprocessor and DMA controller when the HOLD and HLDA pins are at their logic 1 levels.
9. Describe the effect on the microprocessor and DMA controller when the HOLD and HLDA pins are at their logic 0 levels.
10. The 8237 DMA controller is a(n) _____ channel DMA controller.
11. If the 8237 DMA controller is decoded at I/O ports 2000H –200FH, what ports are used to program channel 1?
12. Which 8237 DMA controller register is programmed to initialize the controller?
13. How many bytes can be transferred by the 8237 DMA controller?
14. Write a sequence of instructions that transfer data from memory location 21000H –210FFH to 20000H –200FFH by using channel 2 of the 8237 DMA controller. You must initialize the 8237 and use the latch described in Section 12–1 to hold A_{19} –A_{16}.
15. Write a sequence of instructions that transfers data from memory to an external I/O device by using Channel 3 of the 8237. The memory area to be transferred is at location 20000H – 20FFFH.
16. What is a pen drive?
17. The 3 1/2″ disk is known as a(n) _____ floppy disk.
18. Data are recorded in concentric rings on the surface of a disk known as a(n) _____.
19. A track is divided into sections of data called _____.
20. On a double-sided disk, the upper and lower tracks together are called a(n) _____.
21. Why is NRZ recording used on a disk memory system?
22. Draw the timing diagram generated to write a 1001010000 using MFM encoding.
23. Draw the timing diagram generated to write a 1001010000 using RLL encoding.
24. What is a flying head?
25. Why must the heads on a hard disk be parked?
26. What is the difference between a voice coil head position mechanism and a stepper motor head positioning mechanism?
27. What is a WORM?
28. What is a CD-ROM?
29. How much data can be stored on a common DVD, an HD-DVD, and a Blu-ray DVD?
30. What is the difference between a TTL monitor and an analog monitor?
31. What are the three primary colors of light?
32. What are the three secondary colors of light?
33. What is a pixel?
34. A video display with a resolution of 1280×1024 contains _____ lines, with each line divided into _____ pixels.
35. Explain how a TTL RGB monitor can display 16 different colors.
36. What are the DVI-D and HDMI connectors?
37. Explain how an analog RGB monitor can display an infinite number of colors.
38. If an analog RGB video system uses 8-bit DACs, it can generate _____ different colors.
39. If a video system uses a vertical frequency of 60 Hz and a horizontal frequency of 32,400 Hz, how many raster lines are generated?

CHAPTER 14

The Arithmetic Coprocessor, MMX, and SIMD Technologies

INTRODUCTION

The Intel family of arithmetic coprocessors includes the 8087, 80287, 80387SX, 80387DX, and the 80487SX for use with the 80486SX microprocessor. The 80486DX–Core2 microprocessors contain their own built-in arithmetic coprocessors. Be aware that some of the cloned 80486 microprocessors (from IBM and Cyrix) did not contain arithmetic coprocessors. The instruction sets and programming for all devices are almost identical; the main difference is that each coprocessor is designed to function with a different Intel microprocessor. This chapter provides detail on the entire family of arithmetic coprocessors. Because the coprocessor is a part of the 80486DX–Core2, and because these microprocessors are commonplace, many programs now require or at least benefit from a coprocessor.

The family of coprocessors, which is labeled the 80X87, is able to multiply, divide, add, subtract, find the square root, and calculate the partial tangent, partial arctangent, and logarithms. Data types include 16-, 32-, and 64-bit signed integers; l8-digit BCD data; and 32-, 64-, and 80-bit floating-point numbers. The operations performed by the 80X87 generally execute many times faster than equivalent operations written with the most efficient programs that use the microprocessor's normal instruction set. With the improved Pentium coprocessor, operations execute about five times faster than those performed by the 80486 microprocessor with an equal clock frequency. Note that the Pentium can often execute a coprocessor instruction and two integer instructions simultaneously. The Pentium Pro through Core2 coprocessors are similar in performance to the Pentium coprocessor, except that a few new instructions have been added: FCMOV and FCOMI.

The multimedia extensions (MMX) to the Pentium–Core2 are instructions that share the arithmetic coprocessor register set. The MMX extension is a special internal processor designed to execute integer instructions at high-speed for external multimedia devices. For this reason, the MMX instruction set and specifications have been placed in this chapter. The SIMD (single-instruction, multiple data) extensions, which are called SSE (streaming SIMD extensions), are similar to the MMX instructions, but function with floating-point numbers instead of integers and do not use the coprocessor register space as do MMX instructions.

CHAPTER OBJECTIVES

Upon completion of this chapter, you will be able to:

1. Convert between decimal data and signed integer, BCD, and floating-point data for use by the arithmetic coprocessor, MMX, and SIMD technologies.

2. Explain the operation of the 80X87 arithmetic coprocessor and the MMX and SIMD units.
3. Explain the operation and addressing modes of each arithmetic coprocessor, MMX, and SSE instruction.
4. Develop programs that solve complex arithmetic problems using the arithmetic coprocessor, MMX, and SIMD instructions.

14–1   DATA FORMATS FOR THE ARITHMETIC COPROCESSOR

This section of the text presents the types of data used with all arithmetic coprocessor family members. (See Table 14–1 for a listing of all Intel microprocessors and their companion coprocessors.) These data types include signed integer, BCD, and floating-point. Each has a specific use in a system, and many systems require all three data types. Note that assembly language programming with the coprocessor is often limited to modifying the coding generated by a high-level language such as C/C++. In order to accomplish any such modification, the instruction set and some basic programming concepts are required, which are presented in this chapter.

Signed Integers

The signed integers used with the coprocessor are the same as those described in Chapter 1. When used with the arithmetic coprocessor, signed integers are 16- (word), 32- (doubleword integer), or 64-bits (quadword integer) wide. The long integer is new to the coprocessor and is not described in Chapter 1, but the principles are the same. Conversion between decimal and signed integer format is handled in exactly the same manner as for the signed integers described in Chapter 1. As you will recall, positive numbers are stored in true form with a leftmost sign-bit of 0, and negative numbers are stored in two's complement form with a leftmost sign-bit of 1.

The word integers range in value from $-32,768$ to $+32,767$, the doubleword integer range is $\pm 2 \times 10^9$, and the quadword integer range is $\pm 9 \times 10^{18}$. Integer data types are found in some applications that use the arithmetic coprocessor. See Figure 14–1, which shows these three forms of signed integer data.

Data are stored in memory using the same assembler directives described and used in earlier chapters. The DW directive defines words, DD defines doubleword integers, and DQ defines quadword integers. Example 14–1 shows how several different sizes of signed integers are defined for use by the assembler and arithmetic coprocessor.
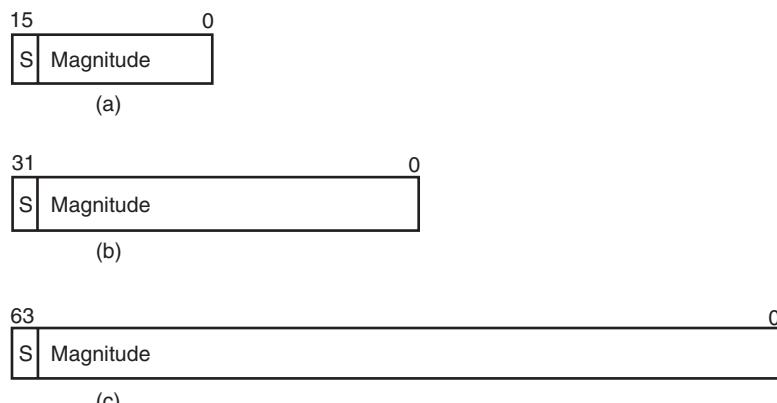
EXAMPLE 14–1

```
0000 0002                   DATA1   DW      2       ;16-bit integer
0002 FFDE                   DATA2   DW      -34     ;16-bit integer
0004 000004D2               DATA3   DD      1234    ;32-bit integer
0008 FFFFFF9C               DATA4   DD      -100    ;32-bit integer
000C 0000000000005BA0       DATA5   DQ      23456   ;64-bit integer
0014 FFFFFFFFFFFFFF86       DATA6   DQ      -122    ;64-bit integer
```

TABLE 14–1 Microprocessor and coprocessor compatibility.

Microprocessor	Coprocessor
8086/8088	8087
80186/80188	80187
80286	80287
80386	80387
80486SX	80487SX
80486DX–Core2	Built into the microprocessor

FIGURE 14–1 Integer formats for the 80×87 family of arithmetic coprocessors: (a) word, (b) short, and (c) long.

Note: S = sign-bit

Binary-Coded Decimal (BCD)

The binary-coded decimal (BCD) form requires 80 bits of memory. Each number is stored as an 18-digit packed integer in nine bytes of memory as two digits per byte. The tenth byte contains only a sign-bit for the 18-digit signed BCD number. Figure 14–2 shows the format of the BCD number used with the arithmetic coprocessor. Note that both positive and negative numbers are stored in true form and never in ten's complement form. The DT directive stores BCD data in the memory as illustrated in Example 14–2. This form is rarely used because it is unique to the Intel coprocessor.

EXAMPLE 14–2

```
0000  00000000000000000200    DATA1  DT      200     ;define 10 byte
000A  80000000000000000010    DATA2  DT      -10     ;define 10 byte
0014  00000000000000010020    DATA3  DT      10020   ;define 10 byte
```

Floating-Point

Floating-point numbers are often called *real numbers* because they hold signed integers, fractions, and mixed numbers. A floating-point number has three parts: a sign-bit, a biased exponent, and a significand. Floating-point numbers are written in scientific binary notation. The Intel family of arithmetic coprocessors supports three types of floating-point numbers: single (32 bits), double (64 bits), and temporary (80 bits). See Figure 14–3 for the three forms of the floating-point number. Please note that the single form is also called a single-precision number and the double form is called a double-precision number. Sometimes the 80-bit temporary form is called an extended-precision number. The floating-point numbers and the operations performed by the arithmetic coprocessor conform to the IEEE-754 standard, as adopted by all major personal computer software producers. This includes Microsoft, which in 1995 stopped supporting the Microsoft floating-point format and also the ANSI floating-point standard that is popular in some mainframe computer systems.
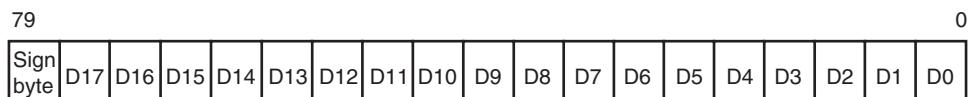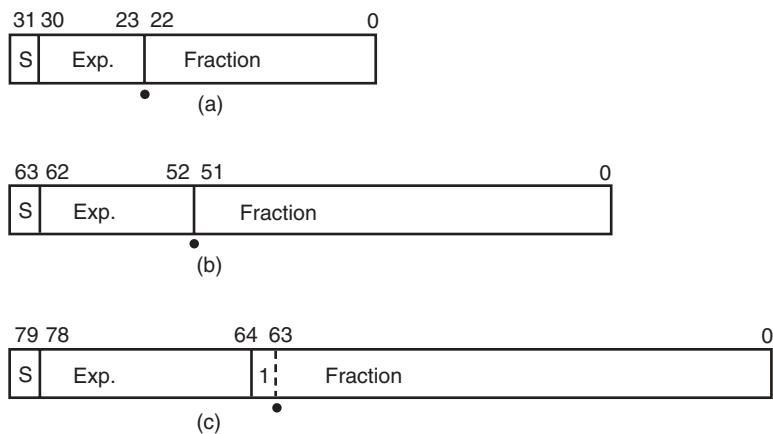
FIGURE 14–2 BCD data format for the 80X87 family of arithmetic coprocessors.

FIGURE 14–3 Floating-point (real) format for the 80X87 family of arithmetic coprocessors. (a) Short (single-precision) with a bias of 7FH, (b) long (double-precision) with a bias of 3FFH, and (c) temporary (extended-precision) with a bias of 3FFFH.

Note: S = sign-bit and Exp. = exponent

In Visual C++ 2008 or the Express edition, float, double, and decimal are used for the three data types. The float is a 32-bit version, double is the 64-bit version, and decimal is a special version developed for Visual studio that develops a very accurate floating-point number for use in banking transactions or anything else that requires a high degree of precision. The decimal variable form is new to Visual Studio 2005 and 2008.

Converting to Floating-Point Form. Converting from decimal to the floating-point form is a simple task that is accomplished through the following steps:

1. Convert the decimal number to binary.
2. Normalize the binary number.
3. Calculate the biased exponent.
4. Store the number in the floating-point format.

These four steps are illustrated for the decimal number 100.25_{10} in Example 14–3. Here, the decimal number is converted to a single-precision (32-bit) floating-point number.

EXAMPLE 14–3

```
Step          Result

1             100.25 => 1100100.01

2             1100100.01 => 1.10010001 × 2⁶

3             110 + 01111111 => 10000101

4             Sign => 0
              Exponent => 10000101
              Significand => 10010001000000000000000
```

In step 3 of Example 14–3, the biased exponent is the exponent, a 2^6 or 110, plus a bias of 01111111 (7FH) or 10000101 (85H). All single-precision numbers use a bias of 7FH, double-precision numbers use a bias of 3FFH, and extended-precision numbers use a bias of 3FFFH.

In step 4 of Example 14–3, the information found in the prior steps is combined to form the floating-point number. The leftmost bit is the sign-bit of the number. In this case, it is a 0 because the number is $+100.25_{10}$. The biased exponent follows the sign-bit. The significand is a 23-bit number with an implied one-bit. Note that the significand of a number l.XXXX is the XXXX

portion. The 1. is an implied one-bit that is only stored in the extended temporary-precision form of the floating-point number as an explicit one-bit.

Some special rules apply to a few numbers. The number 0, for example, is stored as all zeros except for the sign-bit, which can be a logic 1 to represent a negative zero. The plus and minus infinity is stored as logic 1s in the exponent with a significand of all zeros and the sign-bit that represents plus or minus. A NAN (not-a-number) is an invalid floating-point result that has all ones in the exponent with a significand that is *not* all zeros.

Converting from Floating-Point Form. Conversion to a decimal number from a floating-point number is summarized in the following steps:

1. Separate the sign-bit, biased exponent, and significand.
2. Convert the biased exponent into a true exponent by subtracting the bias.
3. Write the number as a normalized binary number.
4. Convert it to a denormalized binary number.
5. Convert the denormalized binary number to decimal.

These five steps convert a single-precision floating-point number to decimal, as shown in Example 14–4. Notice how the sign-bit of 1 makes the decimal result negative. Also notice that the implied one-bit is added to the normalized binary result in step 3.

EXAMPLE 14–4

```
Step            Result

1               Sign => 1
                Exponent => 10000011
                Significand => 10010010000000000000000

2               100 = 10000011 − 01111111

3               1.1001001 × 2⁴

4               11001.001

5               −25.125
```

Storing Floating-Point Data in Memory. Floating-point numbers are stored with the assembler using the DD directive for single-precision, DQ for double-precision, and DT for extended temporary-precision. Some examples of floating-point data storage are shown in Example 14–5. The author discovered that the Microsoft macro assembler contains an error that does not allow a plus sign to be used with positive floating-point numbers. A +92.45 must be defined as 92.45 for the assembler to function correctly. Microsoft has assured the author that this error has been corrected in version 6.11 of MASM if the REAL4, REAL8, or REAL10 directives are used in place of DD, DQ, and DT to specify floating-point data. The assembler provides access 8087 emulator if your system does not contain a microprocessor with a coprocessor. The emulator comes with all Microsoft high-level languages or as shareware programs such as EM87. Access the emulator by including the OPTION EMULATOR statement immediately following the .MODEL statement in a program. Be aware that the emulator does not emulate some of the coprocessor instructions. Do not use this option if your system contains a coprocessor. In all cases, you must include the .8087, .80187, .80287, .80387, .80487, .80587, or .80687 switch to enable the generation of coprocessor instructions.

EXAMPLE 14–5

```
0000 C377999A            DATA7  DD      -247.6      ;single-precision
0004 40000000            DATA8  DD      2.0         ;single precision
0008 486F4200            DATA9  REAL4   2,45E+5     ;single-precision
000C 4059100000000000    DATAA  DQ      100.25      ;double-precision
0014 3F543BF727136A40    DATAB  REAL8   0.001235    ;double-precision
001C 400487F34D6A161E4F76 DATAC REAL10  33.9876     ;temporary-precision
```

14–2   THE 80X87 ARCHITECTURE

The 80X87 is designed to operate concurrently with the microprocessor. Note that the 80486DX–Core2 microprocessors contain their own internal and fully compatible versions of the 80387. With other family members, the coprocessor is an external integrated circuit that parallels most of the connections on the microprocessor. The 80X87 executes 68 different instructions. The microprocessor executes all normal instructions and the 80X87 executes arithmetic coprocessor instructions. Both the microprocessor and coprocessor will execute their respective instructions simultaneously or concurrently. The numeric or arithmetic coprocessor is a special-purpose microprocessor that is especially designed to efficiently execute arithmetic and transcendental operations.

The microprocessor intercepts and executes the normal instruction set, and the coprocessor intercepts and executes only the coprocessor instructions. Recall that the coprocessor instructions are actually escape (ESC) instructions. These instructions are used by the microprocessor to generate a memory address for the coprocessor so that the coprocessor can execute a coprocessor instruction.

Internal Structure of the 80X87

Figure 14–4 shows the internal structure of the arithmetic coprocessor. Notice that this device is divided into two major sections: the control unit and the numeric execution unit.

The **control unit** interfaces the coprocessor to the microprocessor-system data bus. Both the devices monitor the instruction stream. If the instruction is an ESCape (coprocessor) instruction, the coprocessor executes it; if not, the microprocessor executes it.

The **numeric execution unit** (NEU) is responsible for executing all coprocessor instructions. The NEU has an eight-register stack that holds operands for arithmetic instructions and the results of arithmetic instructions. Instructions either address data in specific stack data registers or use a push-and-pop mechanism to store and retrieve data on the top of the stack. Other registers in the NEU are status, control, tag, and exception pointers. A few instructions transfer data between the coprocessor and the AX register in the microprocessor. The FSTSW AX instruction is the only instruction available to the coprocessor that allows direct communications to the microprocessor through the AX register. Note that the 8087 does not contain the FSTSW AX instruction, but all newer coprocessors do contain it.

The stack within the coprocessor contains eight registers that are each 80 bits wide. These stack registers always contain an 80-bit extended-precision floating-point number. The only time that data appear as any other form is when they reside in the memory system. The coprocessor converts from signed integer, BCD, single-precision, or double-precision form as the data are moved between the memory and the coprocessor register stack.

Status Register.   The status register (see Figure 14–5) reflects the overall operation of the coprocessor. The status register is accessed by executing the instruction (FSTSW), which stores the contents of the status register into a word of memory. The FSTSW AX instruction copies the status register directly into the microprocessor's AX register on the 80187 or above coprocessor. Once status is stored in memory or the AX register, the bit positions of the
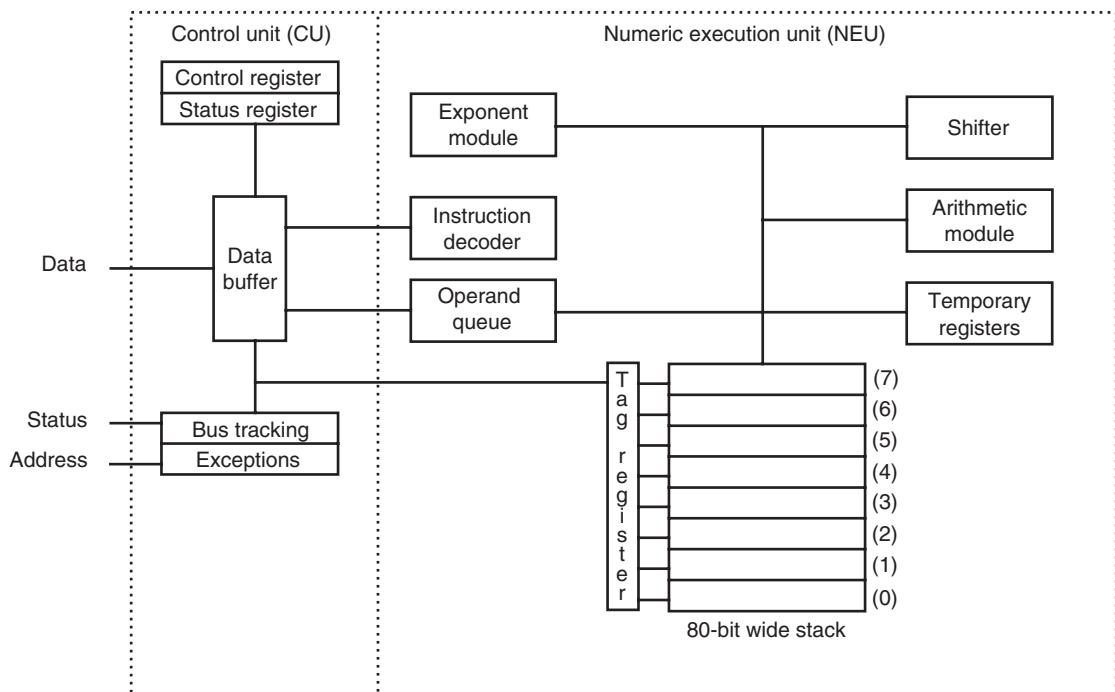
FIGURE 14–4   The internal structure of the 80X87 arithmetic coprocessor.
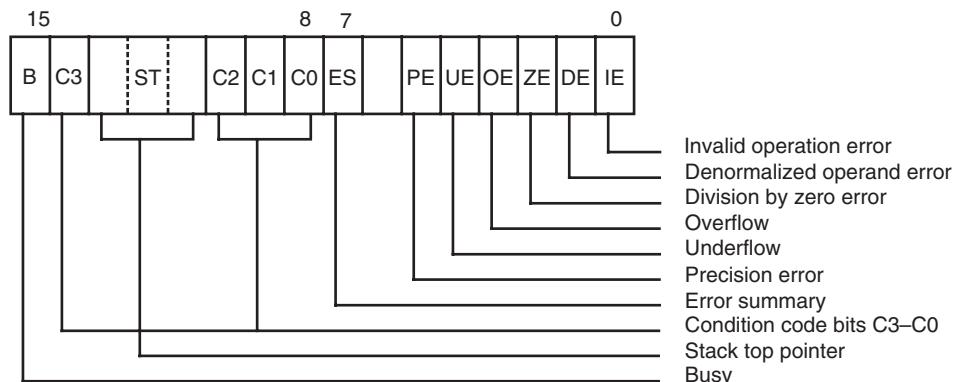
FIGURE 14–5   The 80X87 arithmetic coprocessor status register.

status register can be examined by normal software. The coprocessor/microprocessor communications are carried out through the I/O ports 00FAH–00FFH on the 80187 and 80287, and I/O ports 800000FAH–800000FFH on the 80386 through the Pentium 4. Never use these I/O ports for interfacing I/O devices to the microprocessor.

The newer coprocessors (80187 and above) use status bit position 6 (SF) to indicate a stack overflow or underflow error. Following is a list of the status bits, except for SF, and their applications:

B   The **busy bit** indicates that the coprocessor is busy executing a task. Busy is tested by examining the status register or by using the FWAIT instruction. Newer coprocessors automatically synchronize with the microprocessor, so the busy flag need not be tested before performing additional coprocessor tasks.

TABLE 14–2   The coprocessor status register condition code bits.

Instruction	C_3	C_2	C_1	C_0	Indication
FTST, FCOM	0	0	X	0	ST > Operand
	0	0	X	1	ST < Operand
	1	0	X	1	ST = Operand
	1	1	X	1	ST is not comparable
FPREM	Q1	0	Q0	Q2	Rightmost 3 bits of quotient
	?	1	?	?	Incomplete
FXAM	0	0	0	0	+ un-normal
	0	0	0	1	+ NAN
	0	0	1	0	− un-normal
	0	0	1	1	− NAN
	0	1	0	0	+ normal
	0	1	0	1	+ ∞
	0	1	1	0	− normal
	0	1	1	1	− ∞
	1	0	0	0	+ 0
	1	0	0	1	Empty
	1	0	1	0	− 0
	1	0	1	1	Empty
	1	1	0	0	+ denormal
	1	1	0	1	Empty
	1	1	1	0	− denormal
	1	1	1	1	Empty

Notes: Un-normal = leading bits of the significand are zero; denormal = exponent is at its most negative value; normal = standard floating-point form; NAN (not-a-number) = an exponent of all ones and a significand not equal to zero; and the operand for TST is zero.

C_0–C_3	The **condition code bits** indicate conditions about the coprocessor (see Table 14–2 for a complete listing of each combination of these bits and their functions). Note that these bits have different meanings for different instructions, as indicated in the table. The top of the stack is denoted as ST in this table.
TOP	The **top-of-stack** (ST) bit indicates the current register addressed as the top-of-the-stack (ST). This is normally register ST(0).
ES	The **error summary** bit is set if any unmasked error bit (PE, UE, OE, ZE, DE, or IE) is set. In the 8087 coprocessor, the error summary also caused a coprocessor interrupt. Since the 80187, the coprocessor interrupt has been absent from the family.
PE	The **precision error** indicates that the result or operands exceed the selected precision.
UE	An **underflow error** indicates a nonzero result that is too small to represent with the current precision selected by the control word.
OE	An **overflow error** indicates a result that is too large to be represented. If this error is masked, the coprocessor generates infinity for an overflow error.
ZE	A **zero error** indicates the divisor was zero while the dividend is a noninfinity or nonzero number.
DE	A **denormalized error** indicates that at least one of the operands is denormalized.
IE	An **invalid error** indicates a stack overflow or underflow, indeterminate form ($0 \div 0$, +∞, −∞, etc.), or the use of a NAN as an operand. This flag indicates errors such as those produced by taking the square root of a negative number, etc.

TABLE 14–3 Coprocessor conditions tested with conditional jumps as illustrated in Example 14–6.

C_3	C_2	C_0	Condition	Jump Instruction
0	0	0	ST > Operand	JA (jump if ST above)
0	0	1	ST < Operand	JB (jump if ST below)
1	0	0	ST = Operand	JE (jump if ST equal)

There are two ways to test the bits of the status register once they are moved into the AX register with the FSTSW AX instruction. One method uses the TEST instruction to test individual bits of the status register. The other uses the SAHF instruction to transfer the leftmost 8 bits of the status register into the microprocessor's flag register. Both methods are illustrated in Example 14–6. This example uses the DIV instruction to divide the top of the stack by the contents of DATA1 and the FSQRT instruction to find the square root of the top of the stack. The example also uses the FCOM instruction to compare the contents of the stack top with DATA1. Note that the conditional jump instructions are used with the SAHF instruction to test for the condition listed in Table 14–3. Although SAHF and conditional jumps cannot test all possible operating conditions of the coprocessor, they can help to reduce the complexity of certain tested conditions. Note that SAHF places C_0 into the carry flag, C_2 into the parity flag, and C_3 into the zero flag.

If the Pentium 4 or Core2 is operated in the 64-bit mode, the SAHF instruction does not function. In the 64-bit mode, another method of testing the coprocessor flags is needed, such as testing each bit of AX for C_0, C_2, and C_3. (See Example 14–6.)

EXAMPLE 14–6

```
;testing for a divide by zero error

      FDIV    DATA1
      FSTSW   AX              ;copy status register into AX
      TEST    AX,4            ;test ZE bit
      JNZ     DIVIDE_ERROR

;testing for an invalid operation after a FSQRT

      FSQRT
      FSTSW   AX
      TEST    AX,1            ;test IE
      JNZ     FSQRT_ERROR

;testing with SAHF so conditional jumps can be used

      FCOM    DATA1
      FSTSW   AX
      SAHF                    ;copy coprocessor flags to flags
      JE      ST_EQUAL
      JB      ST_BELOW
      JA      ST_ABOVE

;in the 64-bit mode of the Pentium 4 or Core2,
;the following code is needed to test for the prior conditions
;because the SAHF instruction does not function in 64-bit mode

;testing for conditions

      FCOM    DATA1
      FSTSW   AX              ;copy status register into AX
      TEST    AX,100H
      JNZ     ST_BELOW
      TEST    AX,4000H
      JNZ     ST_EQUAL
      JMP     ST_ABOVE
```
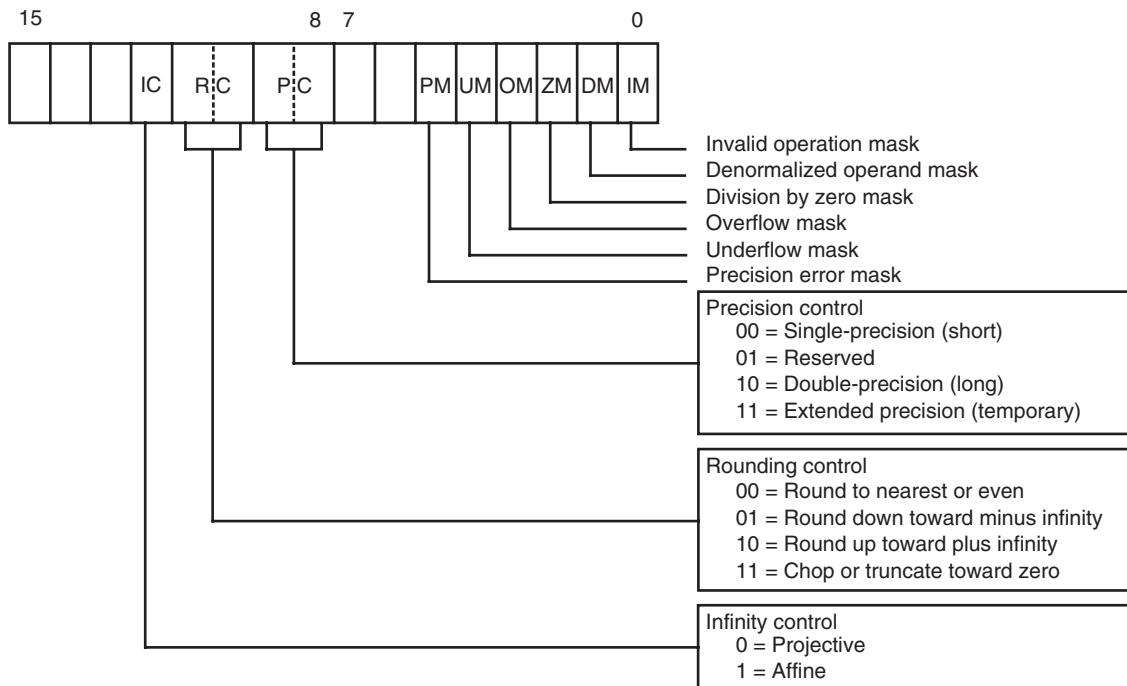
FIGURE 14–6   The 80X87 arithmetic coprocessor control register.

When the FXAM instruction and FSTSW AX are executed and followed by the SAHF instruction, the zero flag will contain C_3. The FXAM instruction could be used to test a divisor before a division for a zero value by using the JZ instruction following FXAM, FSTSW AX, and SAHF.
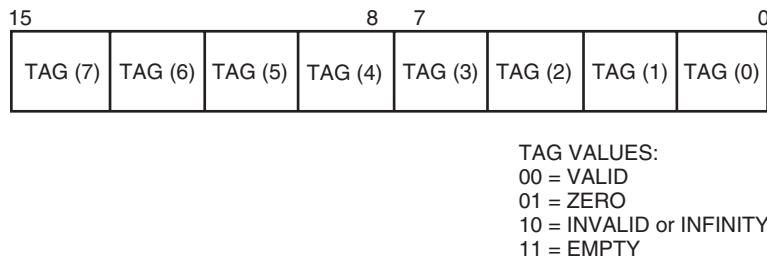
Control Register.   The control register is pictured in Figure 14–6. The control register selects the precision, rounding control, and infinity control. It also masks and unmasks the exception bits that correspond to the rightmost 6 bits of the status register. The FLDCW instruction is used to load a value into the control register.

Following is a description of each bit or grouping of bits found in the control register:

IC	**Infinity control** selects either affine or projective infinity. Affine allows positive and negative infinity; projective assumes infinity is unsigned.
RC	**Rounding control** determines the type of rounding, as defined in Figure 14–6.
PC	The **precision control** sets the precision of the result, as defined in Figure 14–6.
Exception masks	Determine whether the error indicated by the exception affects the error bit in the status register. If a logic 1 is placed in one of the exception control bits, the corresponding status register bit is masked off.

Tag Register.   The **tag register** indicates the contents of each location in the coprocessor stack. Figure 14–7 illustrates the tag register and the status indicated by each tag. The tag indicates whether a register is valid; zero; invalid or infinity; or empty. The only way that a program can view the tag register is by storing the coprocessor environment using the FSTENV, FSAVE, or FRSTOR instructions. Each of these instructions stores the tag register along with other coprocessor data.

TAG VALUES:
00 = VALID
01 = ZERO
10 = INVALID or INFINITY
11 = EMPTY

14–3   INSTRUCTION SET

The arithmetic coprocessor executes over 68 different instructions. Whenever a coprocessor instruction references memory, the microprocessor automatically generates the memory address for the instruction. The coprocessor uses the data bus for data transfers during coprocessor instructions and the microprocessor uses it during normal instructions. Also note that the 80287 uses the Intel-reserved I/O ports 00F8H–00FFH for communications between the coprocessor and the microprocessor (even though the coprocessor only uses ports 00FCH–00FFH). These ports are used mainly for the FSTSW AX instruction. The 80387–Core2 use I/O ports 800000F8H–800000FFH for these communications.

This section of the text describes the function of each instruction and lists its assembly language form. Because the coprocessor uses the microprocessor memory-addressing modes, not all forms of each instruction are illustrated. Each time that the assembler encounters a coprocessor mnemonic opcode, it converts it into a machine language ESC instruction. The ESC instruction represents an opcode to the coprocessor.

Data Transfer Instructions

There are three basic data transfers: floating-point, signed integer, and BCD. The only time that data ever appear in the signed integer or BCD form is in the memory. Inside the coprocessor, data are always stored as an 80-bit extended-precision floating-point number.

Floating-Point Data Transfers. There are four traditional floating-point data transfer instructions in the coprocessor instruction set: FLD (load real), FST (store real), FSTP (store real and pop), and FXCH (exchange). A new instruction is added to the Pentium Pro through Core2 called a conditional floating-point move instruction that uses the opcode FCMOV with a floating-point condition.

The FLD instruction loads floating-point memory data to the top of the internal stack, referred to as ST (stack top). This instruction stores the data on the top of the stack and then decrements the stack pointer by 1. Data loaded to the top of the stack are from any memory location or from another coprocessor register. For example, an FLD ST(2) instruction copies the contents of register 2 to the stack top, which is ST. The top of the stack is register 0 when the coprocessor is reset or initialized. Another example is the FLD DATA7 instruction, which copies the contents of memory location DATA 7 to the top of the stack. The size of the transfer is automatically determined by the assembler through the directives DD or REAL4 for single-precision, DQ or REAL 8 for double-precision, and DT or REAL10 for extended temporary-precision.

The FST instruction stores a copy of the top of the stack into the memory location or coprocessor register indicated by the operand. At the time of storage, the internal, extended temporary-precision floating-point number is rounded to the size of the floating-point number indicated by the control register.

TABLE 14–4   The FCMOV instructions and conditions tested by them.

Instruction	Condition
FCMOVB	Move if below
FCMOVE	Move if equal
FCMOVBE	Move if below or equal
FCMOVU	Move if unordered
FCMOVNB	Move if not below
FCMOVNE	Move if not equal
FCMOVNBE	Move if not below or equal
FCMOVNU	Move if not unordered

The FSTP (floating-point store and pop) instruction stores a copy of the top of the stack into memory or any coprocessor register, and then pops the data from the top of the stack. You might think of FST as a copy instruction and FSTP as a removal instruction.

The FXCH instruction exchanges the register indicated by the operand with the top of the stack. For example, the FXCH ST(2) instruction exchanges the top of the stack with register 2.

Integer Data Transfer Instructions.   The coprocessor supports three integer data transfer instructions: FILD (load integer), FIST (store integer), and FISTP (store integer and pop). These three instructions function as did FLD, FST, and FSTP, except that the data transferred are integer data. The coprocessor automatically converts the internal extended temporary-precision floating-point data to integer data. The size of the data is determined by the way that the label is defined with DW, DD, or DQ in the assembly language program.

BCD Data Transfer Instructions.   Two instructions load or store BCD signed-integer data. The FBLD instruction loads the top of the stack with BCD memory data, and the FBSTP stores the top of the stack and does a pop.

The Pentium Pro through Pentium 4 FCMOV Instruction.   The Pentium Pro–Pentium 4 microprocessors contain a new instruction called FCMOV, which also contains a condition. If the condition is true, the FCMOV instruction copies the source to the destination. The conditions tested by FCMOV and the opcodes used with FCMOV appear in Table 14–4. Notice that these conditions check for either an ordered or unordered condition. The testing for NAN and denormalized numbers are not checked with FCMOV.

Example 14–7 shows how the FCMOVB (move if below) instruction is used to copy the contents of ST(2) to the stack top (ST) if the contents of ST(2) is below ST. Notice that the FCOM instruction must be used to perform the compare and the contents of the status register must still be copied to the flags for this instruction to function. More about the FCMOV instruction appears with the FCOMI instruction, which is also new to the Pentium Pro through the Core2 microprocessors.

EXAMPLE 14–7

```
FCOM   ST(2)          ;compare ST and ST(2)
FSTSW  AX             ;floating flags to AX
SAHF                  ;floating flags to flags
FCMOVB ST(2)          ;copy ST(2) to ST if below
;
;OR
;
FCOMI  ST(2)
FCMOVB ST(2)
```

TABLE 14–5 Coprocessor addressing modes.

Mode	Form	Example
Stack	ST(1), ST	FADD
Register	ST, ST(n)	FADD ST, ST(1)
	ST(n), ST	FADD ST(4), ST
Register with pop	ST(n), ST	FADDP ST(3), ST
Memory	Operand	FADD DATA3

Note: Stack address is fixed as ST(1),ST and includes a pop, so only the result remains at the top of the stack; and n = register number 0–7.

Arithmetic Instructions

Arithmetic instructions for the coprocessor include addition, subtraction, multiplication, division, and calculating square roots. The arithmetic-related instructions are scaling, rounding, absolute value, and changing the sign.

Table 14–5 shows the basic addressing modes allowed for the arithmetic operations. Each addressing mode is shown with an example using the FADD (real addition) instruction. All arithmetic operations are floating-point, except some cases in which memory data are referenced as an operand.

The classic stack form of addressing operand data (stack addressing) uses the top of the stack as the source operand and the next to the top of the stack as the destination operand. Afterward, a pop removes the source datum from the stack and only the result in the destination register remains at the top of the stack. To use this addressing mode, the instruction is placed in the program without any operands such as FADD or FSUB. The FADD instruction adds ST to ST(1) and stores the answer at the top of the stack; it also removes the original two data from the stack by popping. Note carefully that FSUB subtracts ST from ST(1) and leaves the difference at ST. Therefore, a reverse subtraction (FSUBR) subtracts ST(1) from ST and leaves the difference at ST. (Note that an error exists in Intel documentation, including the Pentium data book, which describes the operation of some reverse instructions.) Another use for reverse operations is for finding a reciprocal (1/X). This is accomplished, if X is at the top of the stack, by loading a 1.0 to ST, followed by the FDIVR instruction. The FDIVR instruction divides ST(1) into ST or X into 1 and leaves the reciprocal (1/X) at ST.

The register-addressing mode uses ST for the top of the stack and ST(n) for another location, where n is the register number. With this form, one operand must be ST and the other is ST(n). Note that to double the top of the stack, the FADD ST,ST(0) instruction is used where ST(0) also addresses the top of the stack. One of the two operands in the register-addressing mode must be ST, while the other must be in the form ST(n), where n is a stack register 0–7. For many instructions, either ST or ST(n) can be the destination. It is fairly important that the top of the stack be ST(0). This is accomplished by resetting or initializing the coprocessor before using it in a program. Another example of register-addressing is FADD ST(1),ST where the contents of ST are added to ST(1) and the result is placed in ST(1).

The top of the stack is always used as the destination for the memory-addressing mode because the coprocessor is a stack-oriented machine. For example, the FADD DATA instruction adds the real number contents of memory location DATA to the top of the stack.

Arithmetic Operations. The letter P in an opcode specifies a register pop after the operation (FADDP compared to FADD). The letter R in an opcode (subtraction and division only) indicates reverse mode. The reverse mode is useful for memory data because memory data normally subtract from the top of the stack. A reversed subtract instruction subtracts the top of the stack from memory and stores the result in the top of the stack. For example, if the top of the stack contains a 10 and memory location DATA1 contains a 1, the FSUB DATA1

instruction results in a +9 on the stack top, and the FSUBR instruction results in a –9. Another example is FSUBR ST,ST(1), which will subtract ST from ST(1) and store the result on ST. A variant is FSUBR ST(1),ST, which will subtract ST(1) from ST and store the result on ST(1).

   The letter I as a second letter in an opcode indicates that the memory operand is an integer. For example, the FADD DATA instruction is a floating-point addition, while the FIADD DATA is an integer addition that adds the integer at memory location DATA to the floating-point number at the top of the stack. The same rules apply to FADD, FSUB, FMUL, and FDIV instructions.

Arithmetic-Related Operations.   Other operations that are arithmetic in nature include FSQRT (square root), FSCALE (scale a number), FPREM/FPREM1 (find partial remainder), FRNDINT (round to integer), FXTRACT (extract exponent and significand), FABS (find absolute value), and FCHG (change sign). These instructions and the functions that they perform follow:

FSQRT	Finds the **square root** of the top of the stack and leaves the resultant square root at the top of the stack. An invalid error occurs for the square root of a negative number. For this reason, the IE bit of the status register should be tested whenever an invalid result can occur. The IE bit can be tested by loading the status register to AX with the FSTSW AX instruction, followed by TEST AX,1 to test the IE status bit.
FSCALE	Adds the contents of ST(1) (interpreted as an integer) to the exponent at the top of the stack. **FSCALE** multiplies or divides rapidly by powers of two. The value in ST(1) must be between 2^{-15} and 2^{+15}.
FPREM/FPREM1	Performs **modulo division** of ST by ST(1). The resultant remainder is found in the top of the stack and has the same sign as the original dividend. Note that a modulo division results in a remainder without a quotient. Note also that FPREM is supported for the 8086 and 80287, and FPREM1 should be used in newer coprocessors.
FRNDINT	**Rounds** the top of the stack to an integer.
FXTRACT	**Decomposes** the number at the top of the stack into two separate parts that represent the value of the unbiased exponent and the value of the significand. The extracted significand is found at the top of the stack and the unbiased exponent at ST(1). This instruction is often used to convert a floating-point number into a form that can be printed as a mixed number.
FABS	**Changes the sign** of the top of the stack to positive.
FCHS	**Changes the sign** from positive to negative or negative to positive.

Comparison Instructions

The comparison instructions all examine data at the top of the stack in relation to another element and return the result of the comparison in the status register condition code bits C_3–C_0. Comparisons that are allowed by the coprocessor are FCOM (floating-point compare), FCOMP (floating-point compare with a pop), FCOMPP (floating-point compare with two pops), FICOM (integer compare), FICOMP (integer compare and pop), FSTS (test), and FXAM (examine). New with the introduction of the Pentium Pro is the floating compare and move

results to flags or FCOMI instruction. Following is a list of these instructions with a description of their functions:

FCOM	**Compares** the floating-point data at the top of the stack with an operand, which may be any register or any memory operand. If the operand is not coded with the instruction, the next stack element ST(1) is compared with the stack top ST.
FCOMP/FCOMPP	Both instructions perform as FCOM, but they also pop one or two data from the stack.
FICOM/FICOMP	The top of the stack is compared with the integer stored at a memory operand. In addition to the compare, FICOMP also pops the top of the stack.
FTST	**Tests** the contents of the top of the stack against a zero. The result of the comparison is coded in the status register condition code bits, as illustrated in Table 14–2 with the status register. Also, refer to Table 14–3 for a way of using SAHF and the conditional jump instruction with FTST.
FXAM	**Examines** the stack top and modifies the condition code bits to indicate whether the contents are positive, negative, normalized, and so on. Refer to the status register in Table 14–2.
FCOMI/FUCOMI	New to the Pentium Pro through the Pentium 4, this instruction compares in exactly the same manner as the FCOM instruction, with one additional feature: It moves the floating-point flags into the flag register, just as the FNSTSW AX and SAHF instructions do in Example 14–8. Intel has combined the FCOM, FNSTSW AX, and SAHF instructions to form FCOMI. Also available is the unordered compare or FUCOMI. Each is also available with a pop by appending the opcode with a P.

Transcendental Operations

The transcendental instructions include FPT AN (partial tangent), FPATAN (partial arctangent), FSIN (sine), FCOS (cosine), FSINCOS (sine and cosine), F2XM1 ($2^X - 1$), FYL2X (Y $\log_2$ X), and FYL2XP1 [Y $\log_2$ (X + 1)]. A list of these operations follows with a description of each transcendental operation:

FPTAN	Finds the **partial tangent** of Y/X = tan θ. The value of θ is at the top of the stack. It must be between 0 and n/4 radians for the 8087 and 80287, and must be less than 2^{63} for the 80387, 80486/7, and Pentium–Core2 microprocessors. The result is a ratio found as ST = X and ST(1) = Y. If the value is outside of the allowable range, an invalid error occurs, as indicated by the status register IE bit. Also note that ST(7) must be empty for this instruction to function properly.
FPATAN	Finds the **partial arctangent** as θ = ARCTAN X/Y. The value of X is at the top of the stack and Y is at ST(1). The values of X and Y must be as follows: $0 \leq Y < X < \infty$. The instruction pops the stack and leaves θ in radians at the top of the stack.
F2XM1	Finds the **function $2^X - 1$**. The value of X is taken from the top of the stack and the result is returned to the top of the stack. To obtain 2^X add one to the

TABLE 14–6 Exponential functions.

Function	Equation
10^{Y}	$2^{Y} \times \log_2 10$
ε^{Y}	$2^{Y} \times \log_2 \varepsilon$
X^{Y}	$2^{Y} \times \log_2 X$

result at the top of the stack. The value of X must be in the range of -1 and $+1$. The F2XM1 instruction is used to derive the functions listed in Table 14–6. Note that the constants $\log_2 10$ and $\log_2 \varepsilon$ are built in as standard values for the coprocessor.

FSIN/FCOS Finds the **sine** or **cosine** of the argument located in ST expressed in radians $(360° = 2\pi$ radians), with the result found in ST. The values of ST must be less than 2^{63}.

FSINCOS Finds the **sine** and **cosine** of ST, expressed in radians, and leaves the results as ST = sine and ST(1) = cosine. As with FSIN or FCOS, the initial value of ST must be less than 2^{63}.

FYL2X Finds **Y $\log_2$ X**. The value X is taken from the stack top, and Y is taken from ST(1). The result is found at the top of the stack after a pop. The value of X must range between 0 and ∞, and the value of Y must be between $-\infty$ and $+\infty$. A logarithm with any positive base (b) is found by the equation $LOG_b X = (LOG_2 b)^{-1} \times LOG_2 X$.

FYL2P1 Finds **Y $\log_2$ (X + 1)**. The value of X is taken from the stack top and Y is taken from ST(1). The result is found at the top of the stack after a pop. The value of X must range between 0 and $1 - \mathbf{1}\, \overline{2}/2$ and the value of Y must be between $-\infty$ and $+\infty$.

Constant Operations

The coprocessor instruction set includes opcodes that return constants to the top of the stack. A list of these instructions appears in Table 14–7.

Coprocessor Control Instructions

The coprocessor has control instructions for initialization, exception handling, and task switching. The control instructions have two forms. For example, FINIT initializes the coprocessor, as does FNINIT. The difference is that FNINIT does not cause any wait states, while FINIT does

TABLE 14–7 Constant operations.

Instruction	Constant Pushed in ST
FLDZ	$+ 0.0$
FLD1	$+ 1.0$
FLDPI	π
FLDL2T	$\log_2 10$
FLDL2E	$\log_2 \varepsilon$
FLDLG2	$\log_{10} 2$
FLDLN2	$\log_{\varepsilon} 2$

TABLE 14–8 Coprocessor state after the FINIT instruction.

Field	Value	Condition
Infinity	0	Projective
Rounding	00	Round to nearest
Precision	11	Extended-precision
Error masks	11111	Error bits disabled
Busy	0	Not busy
C_0–C_3	????	Unknown
TOP	000	Register 000 or ST(0)
ES	0	No errors
Error bits	00000	No errors
All tags	11	Empty
Registers	ST(0)–ST(7)	Not changed

cause waits. The microprocessor waits for the FINIT instruction by testing the BUSY pin on the coprocessor. All control instructions have these two forms. Following is a list of each control instruction with its function:

FINIT/FNINIT Performs a reset (**initialize**) operation on the arithmetic coprocessor (see Table 14–8 for the reset conditions). The coprocessor operates with a closure of projective (unsigned infinity), rounds to the nearest or even, and uses extended-precision when reset or initialized. It also sets register 0 as the top of the stack.

FSETPM Changes the addressing mode of the coprocessor to the **protected-addressing mode**. This mode is used when the microprocessor is also operated in the protected mode. As with the microprocessor, protected mode can only be exited by a hardware reset or, in the case of the 80386 through the Pentium 4, with a change to the control register.

FLDCW Loads the control register with the word addressed by the operand.

FSTCW Stores the control register into the word-sized memory operand.

FSTSW AX Copies the contents of the control register to the AX register. This instruction is not available to the 8087 coprocessor.

FCLEX Clears the error flags in the status register and also the busy flag.

FSAVE Writes the entire state of the machine to memory. Figure 14–8 shows the memory layout for this instruction.

FRSTOR Restores the state of the machine from memory. This instruction is used to restore the information saved by FSAVE.

FSTENV Stores the environment of the coprocessor, as shown in Figure 14–9.

FLDENV Reloads the environment saved by FSTENV.

FINCSP Increments the stack pointer.

FDECSP Decrements the stack pointer.

FFREE Frees a register by changing the destination register's tag to empty. It does not affect the contents of the register.

FNOP Floating-point coprocessor NOP.

FWAIT Causes the microprocessor to wait for the coprocessor to finish an operation. FWAIT should be used before the microprocessor accesses memory data that are affected by the coprocessor.

FIGURE 14–8 Memory format when the 80X87 registers are saved with the FSAVE instruction.

FIGURE 14–9 Memory format for the FSTENV instruction: (a) real mode and (b) protected mode.

Coprocessor Instructions

Although the microprocessor circuitry has not been discussed, the instruction sets of these coprocessors and their differences from the other versions of the coprocessor can be discussed. These newer coprocessors contain the same basic instructions provided by the earlier versions, with a few additional instructions.

The 80387, 80486, 80487SX, and Pentium through the Core2 contain the following additional instructions: FCOS (cosine), FPREM1 (partial remainder), FSIN (sine), FSINCOS (sine and

cosine), and FUCOM/FUCOMP/FUCOMPP (unordered compare). The sine and cosine instructions are the most significant addition to the instruction set. In the earlier versions of the coprocessor, the sine and cosine is calculated from the tangent. The Pentium Pro through the Core2 contain two new floating-point instructions: FCMOV (a conditional move) and FCOMI (a compare and move to flags).

Table 14–9 lists the instruction sets for all versions of the coprocessor. It also lists the number of clocking periods required to execute each instruction. Execution times are listed for the 8087, 80287, 80387, 80486, 80487, and Core2. (The timings for the Pentium through the Pentium 4 are the same because the coprocessor is identical in each of these microprocessors.) To determine the execution time of an instruction, the clock time is multiplied times the listed execution time. The FADD instruction requires 70–143 clocks for the 80287. Suppose that an 8 MHz clock is used with the 80287. The clocking period is 1/8 MHz, or 125 ns. The FADD instruction requires between 8.75 μs and 17.875 μs to execute. Using a 33 MHz (33 ns) 80486DX2, this instruction requires between 0.264 μs and 0.66 μs to execute. On the Pentium the FADD instruction requires from 1–7 clocks, so if operated at 133 MHz (7.52 ns), the FADD requires between 0.00752 μs and 0.05264 μs. The Pentium Pro through the Core2 are even faster than the Pentium. For example, in a 3 GHz Pentium 4, which has a clock period of 0.333 ns, the FADD instruction requires between 0.333 ns and 2.333 ns to execute.

Table 14–9 uses some shorthand notations to represent the displacement that may or may not be required for an instruction that uses a memory-addressing mode. It also uses the abbreviation *mmm* to represent a register/memory addressing mode and uses *rrr* to represent one of the floating-point coprocessor registers ST(0)–ST(7). The d (destination) bit that appears in some instruction opcodes defines the direction of the data flow, as in FADD ST,ST(2) or FADD ST(2),ST. The d bit is a logic 0 for flow toward ST, as in FADD ST,ST(2), where ST holds the sum after the addition; and a logic 1 for FADD ST(2),ST, where ST(2) holds the sum.

Also note that some instructions allow a choice of whether a wait is inserted. For example, the FSTSW AX instruction copies the status register into AX. The FNSTSW AX instruction also copies the status register to AX, but without a wait.

TABLE 14–9 The instruction set of the arithmetic coprocessor.

F2XM1	$2^{ST} - 1$		
11011001 11110000			
Example			Clocks
F2XM1		8087	310–630
		80287	310–630
		80387	211–476
		80486/7	140–279
		Pentium–Core2	13–57

FABS	Absolute value of ST		
11011001 11100001			
Example			Clocks
FABS		8087	10–17
		80287	10–17
		80387	22
		80486/7	3
		Pentium–Core2	1

FADD/FADDP/FIADD	Addition		
11011000 oo000mmm disp	32-bit memory (FADD)		
11011100 oo000mmm disp	64-bit memory (FADD)		
11011d00 11000rrr	FADD ST,ST(rrr)		
11011110 11000rrr	FADDP ST,ST(rrr)		
11011110 oo000mmm disp	16-bit memory (FIADD)		
11011010 oo000mmm disp	32-bit memory (FIADD)		
Format	Examples		Clocks
FADD	FADD DATA	8087	70–143
FADDP	FADD ST,ST(1)		
FIADD	FADDP	80287	70–143
	FIADD NUMBER	80387	23–72
	FADD ST,ST(3)		
	FADDP ST,ST(2)	80486/7	8–20
	FADD ST(2),ST	Pentium–Core2	1–7

FCLEX/FNCLEX   Clear errors

11011011 11100010

Example	Clocks	
FCLEX FNCLEX	8087	2–8
	80287	2–8
	80387	11
	80486/7	7
	Pentium–Core2	9

FCOM/FCOMP/FCOMPP/FICOM/FICOMP   Compare

11011000 oo010mmm disp	32-bit memory (FCOM)
11011100 oo010mmm disp	64-bit memory (FCOM)
11011000 11010rrr	FCOM ST(rrr)
11011000 oo011mmm disp	32-bit memory (FCOMP)
11011100 oo011mmm disp	64-bit memory (FCOMP)
11011000 11011rrr	FCOMP ST(rrr)
11011110 11011001	FCOMPP
11011110 oo010mmm disp	16-bit memory (FICOM)
11011010 oo010mmm disp	32-bit memory (FICOM)
11011110 oo011mmm disp	16-bit memory (FICOMP)
11011010 oo011mmm disp	32-bit memory (FICOMP)

Format	Examples	Clocks	
FCOM FCOMP FCOMPP FICOM FICOMP	FCOM ST(2) FCOMP DATA FCOMPP FICOM NUMBER FICOMP DATA3	8087	40–93
		80287	40–93
		80387	24–63
		80486/7	15–20
		Pentium–Core2	1–8

FCOMI/FUCOMI/COMIP/FUCOMIP   Compare and Load Flags

11011011 11110rrr	FCOMI ST(rrr)
11011011 11101rrr	FUCOMI ST(rrr)
11011111 11110rrr	FCOMIP ST(rrr)
11011111 11101rrr	FUCOMIP ST(rrr)

Format	Examples	Clocks	
FCOM FUCOMI FCOMIP FUCOMIP	FCOMI ST(2) FUCOMI ST(4) FCOMIP ST(0) FUCOMIP ST(1)	8087	—
		80287	—
		80387	—
		80486/7	—
		Pentium–Core2	—

FCMOVcc   Conditional Move

11011010 11000rrr	FCMOVB ST(rrr)
11011010 11001rrr	FCMOVE ST(rrr)
11011010 11010rrr	FCMOVBE ST(rrr)
11011010 11011rrr	FCMOVU ST(rrr)
11011011 11000rrr	FCMOVNB ST(rrr)
11011011 11001rrr	FCMOVNE ST(rrr)
11011011 11010rrr	FCMOVENBE ST(rrr)
11011011 11011rrr	FCMOVNU ST(rrr)

Format	Examples		Clocks
FCMOVB FCMOVB ST(2)		8087	—
FCMOVE FCMOVE ST(3)		80287	—
		80387	—
		80486/7	—
		Pentium–Core2	—

FCOS   Cosine of ST

11011001 11111111

Example		Clocks
FCOS	8087	—
	80287	—
	80387	123–772
	80486/7	193–279
	Pentium–Core2	18–124

FDECSTP   Decrement stack pointer

11011001 11110110

Example		Clocks
FDECSTP	8087	6–12
	80287	6–12
	80387	22
	80486/7	3
	Pentium–Core2	1

FDISI/FNDISI   Disable interrupts

11011011 11100001

(Ignored on the 80287, 80387, 80486/7, Pentium–Core 2)

Example		Clocks
FDISI FNDISI	8087	2–8
	80287	—
	80387	—
	80486/7	—
	Pentium–Core2	—

FDIV/FDIVP/FIDIV   Division

```
11011000  oo110mmm  disp        32-bit memory (FDIV)
11011100  oo100mmm  disp        64-bit memory (FDIV)
11011d00  11111rrr              FDIV ST,ST(rrr)
11011110  11111rrr              FDIVP ST,ST(rrr)
11011110  oo110mmm  disp        16-bit memory (FIDIV)
11011010  oo110mmm  disp        32-bit memory (FIDIV)
```

Format	Examples		Clocks
FDIV FDIVP FIDIV	FDIV DATA FDIV ST,ST(3) FDIVP FIDIV NUMBER FDIV ST,ST(5) FDIVP ST,ST(2) FDIV ST(2),ST	8087	191–243
		80287	191–243
		80387	88–140
		80486/7	8–89
		Pentium–Core2	39–42

FDIVR/FDIVRP/FIDIVR   Division reversed

```
11011000  oo111mmm  disp        32-bit memory (FDIVR)
11011100  oo111mmm  disp        64-bit memory (FDIVR)
11011d00  11110rrr              FDIVR ST,ST(rrr)
11011110  11110rrr              FDIVRP ST,ST(rrr)
11011110  oo111mmm  disp        16-bit memory (FIDIVR)
11011010  oo111mmm  disp        32-bit memory (FIDIVR)
```

Format	Examples		Clocks
FDIVR FDIVRP FIDIVR	FDIVR DATA FDIVR ST,ST(3) FDIVRP FIDIVR NUMBER FDIVR ST,ST(5) FDIVRP ST,ST(2) FDIVR ST(2),ST	8087	191–243
		80287	191–243
		80387	88–140
		80486/7	8–89
		Pentium–Core2	39–42

FENI/FNENI Disable interrupts		
11011011 11100000		
(Ignored on the 80287, 80387, 80486/7, Pentium–Core 2)		
Example		Clocks
FENI FNENI	8087	2–8
	80287	—
	80387	—
	80486/7	—
	Pentium–Core2	—

FFREE Free register			
11011101 11000rrr			
Format	Examples		Clocks
FFREE	FFREE FFREE ST(1) FFREE ST(2)	8087	9–16
		80287	9–16
		80387	18
		80486/7	3
		Pentium–Core2	1

FINCSTP Increment stack pointer		
11011001 11110111		
Example		Clocks
FINCSTP	8087	6–12
	80287	6–12
	80387	21
	80486/7	3
	Pentium–Core2	1

FINIT/FNINIT   Initialize coprocessor

11011001 11110110

Example		Clocks
FINIT FNINIT	8087	2–8
	80287	2–8
	80387	33
	80486/7	17
	Pentium–Core2	12–16

FLD/FILD/FBLD   Load data to ST(0)

11011001 oo000mmm disp	32-bit memory (FLD)	
11011101 oo000mmm disp	64-bit memory (FLD)	
11011011 oo101mmm disp	80-bit memory (FLD)	
11011111 oo000mmm disp	16-bit memory (FILD)	
11011011 oo000mmm disp	32-bit memory (FILD)	
11011111 oo101mmm disp	64-bit memory (FILD)	
11011111 oo100mmm disp	80-bit memory (FBLD)	

Format	Examples		Clocks
FLD FILD FBLD	FLD DATA FILD DATA1 FBLD DEC_DATA	8087	17–310
		80287	17–310
		80387	14–275
		80486/7	3–103
		Pentium–Core2	1–3

FLD1   Load +1.0 to ST(0)

11011001 11101000

Example		Clocks
FLD1	8087	15–21
	80287	15–21
	80387	24
	80486/7	4
	Pentium–Core2	2

FLDZ Load +0.0 to ST(0)		
11011001 11101110		
Example		Clocks
FLDZ	8087	11–17
	80287	11–17
	80387	20
	80486/7	4
	Pentium–Core2	2

FLDPI Load π to ST(0)		
11011001 11101011		
Example		Clocks
FLDPI	8087	16–22
	80287	16–22
	80387	40
	80486/7	8
	Pentium–Core2	3–5

FLDL2E Load $\log_2 e$ to ST(0)		
11011001 11101010		
Example		Clocks
FLDL2E	8087	15–21
	80287	15–21
	80387	40
	80486/7	8
	Pentium–Core2	3–5

FLDL2T Load $\log_2 10$ to ST(0)		
11011001 11101001		
Example		Clocks
FLDL2T	8087	16–22
	80287	16–22
	80387	40
	80486/7	8
	Pentium–Core2	3–5

FLDLG2   Load $\log_{10}2$ to ST(0)

11011001 11101000

Example		Clocks
FLDLG2	8087	18–24
	80287	18–24
	80387	41
	80486/7	8
	Pentium–Core2	3–5

FLDLN2   Load $\log_e 2$ to ST(0)

11011001 11101101

Example		Clocks
FLDLN2	8087	17–23
	80287	17–23
	80387	41
	80486/7	8
	Pentium–Core2	3–5

FLDCW   Load control register

11011001 oo101mmm disp

Format	Examples		Clocks
FLDCW	FLDCW DATA	8087	7–14
	FLDCW STATUS	80287	7–14
		80387	19
		80486/7	4
		Pentium–Core2	7

FLDENV   Load environment

11011001 oo100mmm disp

Format	Examples		Clocks
FLDENV	FLDENV ENVIRON	8087	35–45
	FLDENV DATA	80287	25–45
		80387	71
		80486/7	34–44
		Pentium–Core2	32–37

FMUL/FMULP/FIMUL   Multiplication

11011000 oo001mmm disp	32-bit memory (FMUL)	
11011100 oo001mmm disp	64-bit memory (FMUL)	
11011d00 11001rrr	FMUL ST,ST(rrr)	
11011110 11001rrr	FMULP ST,ST(rrr)	
11011110 oo001mmm disp	16-bit memory (FIMUL)	
11011010 oo001mmm disp	32-bit memory (FIMUL)	

Format	Examples		Clocks
FMUL	FMUL DATA	8087	110–168
FMULP	FMUL ST,ST(2)	80287	110–168
FIMUL	FMUL ST(2),ST	80387	29–82
	FMULP	80486/7	11–27
	FIMUL DATA3	Pentium–Core2	1–7

FNOP   No operation

11011001 11010000

Example		Clocks
FNOP	8087	10–16
	80287	10–16
	80387	12
	80486/7	3
	Pentium–Core2	1

FPATAN   Partial arctangent of ST(0)

11011001 11110011

Example		Clocks
FPATAN	8087	250–800
	80287	250–800
	80387	314–487
	80486/7	218–303
	Pentium–Core2	17–173

FPREM   Partial remainder

11011001 11111000

Example		Clocks
FPREM	8087	15–190
	80287	15–190
	80387	74–155
	80486/7	70–138
	Pentium–Core2	16–64

FPREM1   Partial remainder (IEEE)

11011001 11110101

Example		Clocks
FPREM1	8087	—
	80287	—
	80387	95–185
	80486/7	72–167
	Pentium–Core2	20–70

FPTAN   Partial tangent of ST(0)

11011001 11110010

Example		Clocks
FPTAN	8087	30–450
	80287	30–450
	80387	191–497
	80486/7	200–273
	Pentium–Core2	17–173

FRNDINT   Round ST(0) to an integer

11011001 11111100

Example		Clocks
FRNDINT	8087	16–50
	80287	16–50
	80387	66–80
	80486/7	21–30
	Pentium–Core2	9–20

FRSTOR	Restore state		
11011101 oo110mmm disp			
Format	Examples		Clocks
FRSTOR	FRSTOR DATA FRSTOR STATE FRSTOR MACHINE	8087	197–207
		80287	197–207
		80387	308
		80486/7	120–131
		Pentium–Core 2	70–95

FSAVE/FNSAVE	Save machine state		
11011101 oo110mmm disp			
Format	Examples		Clocks
FSAVE FNSAVE	FSAVE STATE FNSAVE STATUS FSAVE MACHINE	8087	197–207
		80287	197–207
		80387	375
		80486/7	143–154
		Pentium–Core2	124–151

FSCALE	Scale ST(0) by ST(1)		
11011001 11111101			
Example			Clocks
FSCALE		8087	32–38
		80287	32–38
		80387	67–86
		80486/7	30–32
		Pentium–Core2	20–31

FSETPM	Set protected mode		
11011011 11100100			
Example			Clocks
FSETPM		8087	—
		80287	2–18
		80387	12
		80486/7	—
		Pentium–Core2	—

FSIN   Sine of ST(0)

11011001 11111110

Example		Clocks
FSIN	8087	—
	80287	—
	80387	122–771
	80486/7	193–279
	Pentium–Core2	16–126

FSINCOS   Find sine and cosine of ST(0)

11011001 11111011

Example		Clocks
FSINCOS	8087	—
	80287	—
	80387	194–809
	80486/7	243–329
	Pentium–Core2	17–137

FSQRT   Square root of ST(0)

11011001 11111010

Example		Clocks
FSQRT	8087	180–186
	80287	180–186
	80387	122–129
	80486/7	83–87
	Pentium–Core2	70

FST/FSTP/FIST/FISTP/FBSTP   Store

11011001 oo010mmm disp	32-bit memory (FST)	
11011101 oo010mmm disp	64-bit memory (FST)	
11011101 11010rrr	FST ST(rrr)	
11011011 oo011mmm disp	32-bit memory (FSTP)	
11011101 oo011mmm disp	64-bit memory (FSTP)	
11011011 oo111mmm disp	80-bit memory (FSTP)	
11011101 11001rrr	FSTP ST(rrr)	
11011111 oo010mmm disp	16-bit memory (FIST)	
11011011 oo010mmm disp	32-bit memory (FIST)	
11011111 oo011mmm disp	16-bit memory (FISTP)	
11011011 oo011mmm disp	32-bit memory (FISTP)	
11011111 oo111mmm disp	64-bit memory (FISTP)	
11011111 oo110mmm disp	80-bit memory (FBSTP)	

Format	Examples		Clocks
FST	FST DATA	8087	15–540
FSTP	FST ST(3)	80287	15–540
FIST	FST	80387	11–534
FISTP	FSTP	80486/7	3–176
FBSTP	FIST DATA2	Pentium–Core2	1–3
	FBSTP DATA6		
	FISTP DATA9		

FSTCW/FNSTCW   Store control register

11011001 oo111mmm disp

Format	Examples		Clocks
FSTCW	FSTCW CONTROL	8087	12–18
FNSTCW	FNSTCW STATUS	80287	12–18
	FSTCW MACHINE	80387	15
		80486/7	3
		Pentium–Core2	2

FSTENV/FNSTENV   Store environment

11011001 oo110mmm disp

Format	Examples		Clocks
FSTENV	FSTENV CONTROL	8087	40–50
FNSTENV	FNSTENV STATUS	80287	40–50
	FSTENV MACHINE	80387	103–104
		80486/7	58–67
		Pentium–Core2	48–50

FSTSW/FNSTSW   Store status register

11011101 oo111mmm disp

Format	Examples		Clocks
FSTSW FNSTSW	FSTSW CONTROL FNSTSW STATUS FSTSW MACHINE FSTSW AX	8087	12–18
		80287	12–18
		80387	15
		80486/7	3
		Pentium–Core2	2–5

FSUB/FSUBP/FISUB   Subtraction

11011000	oo100mmm disp	32-bit memory (FSUB)
11011100	oo100mmm disp	64-bit memory (FSUB)
11011d00	11101rrr	FSUB ST,ST(rrr)
11011110	11101rrr	FSUBP ST,ST(rrr)
11011110	oo100mmm disp	16-bit memory (FISUB)
11011010	oo100mmm disp	32-bit memory (FISUB)

Format	Examples		Clocks
FSUB FSUBP FISUB	FSUB DATA FSUB ST,ST(2) FSUB ST(2),ST FSUBP FISUB DATA3	8087	70–143
		80287	70–143
		80387	29–82
		80486/7	8–35
		Pentium–Core2	1–7

FSUBR/FSUBRP/FISUBR   Reverse subtraction

11011000	oo101mmm disp	32-bit memory (FSUBR)
11011100	oo101mmm disp	64-bit memory (FSUBR)
11011d00	11100rrr	FSUBR ST,ST(rrr)
11011110	11100rrr	FSUBRP ST,ST(rrr)
11011110	oo101mmm disp	16-bit memory (FISUBR)
11011010	oo101mmm disp	32-bit memory (FISUBR)

Format	Examples		Clocks
FSUBR FSUBRP FISUBR	FSUBR DATA FSUBR ST,ST(2) FSUBR ST(2),ST FSUBRP FISUBR DATA3	8087	70–143
		80287	70–143
		80387	29–82
		80486/7	8–35
		Pentium–Core2	1–7

FTST   Compare ST(0) with + 0.0

11011001 11100100

Example	Clocks	
FTST	8087	38–48
	80287	38–48
	80387	28
	80486/7	4
	Pentium–Core2	1–4

FUCOM/FUCOMP/FUCOMPP   Unordered compare

11011101 11100rrr   FUCOM ST,ST(rrr)
11011101 11101rrr   FUCOMP ST,ST(rrr)
11011101 11101001   FUCOMPP

Format	Examples	Clocks	
FUCOM	FUCOM ST,ST(2)	8087	—
FUCOMP	FUCOM	80287	—
FUCOMPP	FUCOMP ST,ST(3)	80387	24–26
	FUCOMP	80486/7	4–5
	FUCOMPP	Pentium–Core2	1–4

FWAIT   Wait

10011011

Example	Clocks	
FWAIT	8087	4
	80287	3
	80387	6
	80486/7	1–3
	Pentium–Core2	1–3

FXAM   Examine ST(0)

11011001 11100101

Example	Clocks	
FXAM	8087	12–23
	80287	12–23
	80387	30–38
	80486/7	8
	Pentium–Core2	21

14–4   PROGRAMMING WITH THE ARITHMETIC COPROCESSOR

This section of the chapter provides programming examples for the arithmetic coprocessor. Each example is chosen to illustrate a programming technique for the coprocessor.

Calculating the Area of a Circle

This first programming example illustrates a simple method of addressing the coprocessor stack. First, recall that the equation for calculating the area of a circle is $A = \pi R^2$. A program that performs this calculation is listed in Example 14–8. Note that this program takes test data from array RAD that contains five sample radii. The five areas are stored in a second array called AREA. No attempt is made in this program to use the data from the AREA array.

EXAMPLE 14–8

```
;A short procedure that finds the areas of 5 circles whose radii are stored
;in array RAD.

RAD     DD     2.34                     ;array of radii
        DD     5.66
        DD     9.33
        DD     234.5
        DD     23.4

AREA    DD     5 DUP(?)                 ;array for areas

FINDA   PROC   NEAR

        FLDPI                           ;load pi
        MOV ECX,0                       ;initialize pointer
        .REPEAT
             FLD    RAD[ECX*4]          ;get radius
             FMUL   ST,ST(0)            ;square radius
              FMUL   ST,ST(1)           ;multiply radius squared times pi
              FSTP   AREA[ECX*4]        ;store area
              INC    ECX                ;index next radius
        .UNTIL ECX = 5                  ;repeat 5 times
        FCOMP                           ;clear pi from coprocessor stack
        RET

FINDA   ENDP
```

Although this is a simple program, it does illustrate the operation of the stack. To provide a better understanding of the operation of the stack, Figure 14–10 shows that the contents of the stack after each instruction of Example 14–8 executes. Note only one pass through the loop is illustrated because the program calculates five areas and each pass is identical.

FIGURE 14–10 Operation of the coprocessor stack for one iteration of the loop in Example 14–8.

Instruction	ST(0)	ST(1)
FLDPI	π	
FLD RAD[ECX*4]	2.34	π
FMUL ST,ST(0)	5.4756	π
FMUL ST,ST(1)	17.202	π
FSTP AREA[ECX*4]	π	

The first instruction loads π to the top of the stack. Next, the contents of memory location RAD [ECX*4], one of the elements of the array, is loaded to the top of the stack. This pushes π to ST(1). Next, the FMUL ST,ST(0) instruction squares the radius on the top of the stack. The FMUL ST,ST(1) instruction forms the area. Finally, the top of the stack is stored in the AREA array and also pops it from the stack in preparation for the next iteration.

Notice how care is taken to always remove all stack data. The last instruction before the RET pops π from the stack. This is important because if data remain on the stack at the end of the procedure, the stack top will no longer be register 0. This could cause problems because software assumes that the top of the stack is register 0. Another way of ensuring that the coprocessor is initialized is to place the FINIT (initialization) instruction at the start of the program.

Finding the Resonant Frequency

An equation commonly used in electronics is the formula for determining the resonant frequency of an LC circuit. The equation solved by the program illustrated in Example 14–9 is

$$Fr = \frac{1}{2\pi\sqrt{LC}}$$

This example uses L_1 for the inductance L, C_1 for the capacitor C, and RES for the resultant resonant frequency.

EXAMPLE 14–9

```
RES    DD     ?                  ;resonant frequency
L1     DD     0.0001             ;1 mH inductor
C1     DD     47E-6              ;47 µF capacitor

FR     PROC   NEAR

       FLD    L1                 ;get L
       FMUL   C1                 ;form LC
       FSQRT                     ;form square root of LC
       FLDPI                     ;get pi
       FADD   ST,ST(0)           ;form 2 pi
       FMUL                      ;form 2 pi square root LC
       FLD1
       FDIVR                     ;form reciprocal
       FSTP   RES
       RET

FR     ENDP
```

Notice the straightforward manner in which the program solves this equation. Very little extra data manipulation is required because of the stack inside the coprocessor. Notice how FDIVR, using classic stack addressing, is used to form the reciprocal. If you own a reverse Polish notation calculator, such as those produced by Hewlett-Packard, you are familiar with stack addressing. If not, using the coprocessor will increase your experience with this type of entry.

Finding the Roots Using the Quadratic Equation

This example illustrates how to find the roots of a polynomial expression ($ax^2 + bx + c = 0$) by using the quadratic equation. The quadratic equation is:

$$\frac{b \pm \sqrt{b^2 - 4ac}}{2a}$$

Example 14–10 illustrates a program that finds the roots (R_1 and R_2) for the quadratic equation. The constants are stored in memory locations A_1, B_1, and C_1. Note that no attempt is made to determine the roots if they are imaginary. This example tests for imaginary roots and then exits to DOS with a zero in the roots (R_1 and R_2), if it finds them. In practice, imaginary roots could be solved for and stored in a separate set of result memory locations.

EXAMPLE 14–10

```
;A procedure that finds the roots of a polynomial equation using
;the quadratic equation. Imaginary roots are indicated if both
;R1 and R2 are returned as zero.


FOUR    DW     4             ;integer of 4
A1      DD     ?             ;value for a
B1      DD     ?             ;value for b
C1      DD     ?             ;value for c
R1      DD     ?             ;root 1
R2      DD     ?             ;root 2


ROOTS   PROC   NEAR

        FLDZ                 ;get 0.0
        FST    R1            ;clear roots
        FSTP   R2
        FLD    A1            ;form 2a
        FADD   ST,ST(0)
        FILD   FOUR          ;get 4
        FMUL   A1            ;form 4ac
        FMUL   C1
        FLD    B1            ;from b²
        FMUL   ST,ST(0)
        FSUBR                ;form b² – 4ac
        FTST                 ;test result against zero
        SAHF
        .IF !ZERO?
             FSQRT                 ;find square root of b² – 4ac
             FSTSW AX              ;test for invalid error
             TEST AX,1
             .IF !ZERO?
                  FCOMPP      ;clear stack
                  RET
             .ENDIF
        .ENDIF
        FLD    B1
        FSUB   ST,ST(1)
        FDIV   ST,ST(2)
        FSTP   R1            ;save root 1
        FLD    B1
        FADD
        FDIVR
        FSTP   R2            ;save root 2
        RET

ROOTS   ENDP
```

Using a Memory Array to Store Results

The next programming example illustrates the use of a memory array and the scaled-indexed addressing mode to access the array. Example 14–11 shows a program that calculates 100 values of inductive reactance. The equation for inductive reactance is $XL = 2\pi FL$. In this example, the frequency range is from 10 Hz to 1000 Hz for F and an inductance of 4 mH. Notice how the instruction FSTP XL[ECX*4 + 4] is used to store the reactance for each frequency, beginning with the last at 1000 Hz and ending with the first at 10 Hz. Also notice how the FCOMP instruction is used to clear the stack just before the RET instruction.

EXAMPLE 14–11

```
;A procedure that calculates the inductive reactance of L at a
;frequencies from 10 Hz to 1000 Hz in steps of 10 Hz and stores
;them in array called XL.

XL     DD     100 DUP(?)              ;array for XL
L      DD     4E-3                    ;L = 4 mH
F      DW     10                      ;integer of 10 for F

XLS    PROC   NEAR

       MOV    ECX,100                 ;count = 100
       FLDPI                          ;get pi
       FADD   ST,ST(0)                ;form 2 pi
       FMUL   L                       ;form 2 pi L
       .REPEAT
              FILD   F                ;get F
              FMUL   ST,ST(1)         ;find XL
              FSTP   XL[ECX*4+4]      ;save result
              MOV    AX,F             ;add 10 to F
              ADD    AX,10
              MOV    F,AX
       .UNTILCXZ
       FCOMP                          ;clear stack
       RET

XLS    ENDP
```

Converting a Single-Precision Floating-Point Number to a String

This section of the text shows how to take the floating-point contents of a 32-bit single-precision floating-point number and store it as an ASCII character string. The procedure converts the floating-point number as a mixed number with an integer part and a fractional part, separated by a decimal point. In order to simplify the procedure, a limit is placed on the size of the mixed number so the integer portion is a 32-bit binary number (± 2 G) and the fraction is a 24-bit binary number (1/16M). The procedure will not function properly for larger or smaller numbers.

Example 14–12 lists a procedure that converts the contents of memory location NUMB to a string stored in the STR array. The procedure first tests the sign of the number and stores a minus sign for a negative number. After storing a minus sign, if needed, the number is made positive by the FABS instruction. Next, it is divided into an integer and fractional part and stored at WHOLE and FRACT. Notice how the FRNDINT instruction is used to round (using the chop mode) the top of the stack to form the whole number part of NUMB. The whole number part is then subtracted from the original number to generate the fractional part. This is accomplished with the FSUB instruction that subtracts the contents of ST(1) from ST.

EXAMPLE 14–12

```
;A procedure that converts a floating-point number into an ASCII
;character string.

STR    DB     40 DUP(?)               ;storage for string
NUMB   DD     -2224.125               ;test number
WHOLE  DD     ?
FRACT  DD     ?
TEMP   DW     ?                       ;place for CW
TEN    DW     10                      ;integer of 10
```

```
FTOA    PROC    NEAR USES EBX ECX EDX


        MOV     ECX,0                   ;initialize pointer
        FSTCW   TEMP                    ;save current control word
        MOV     AX,TEMP                 ;change rounding to chop
        PUSH    AX
        OR      AX,0C00H
        MOV     TEMP,AX
        FLDCW   TEMP
        FTST    NUMB                    ;test NUMB
        FSTSW   AX
        AND     AX,4500H                ;get C0, C2, and C3
        .IF AX == 100H                  ;if negative
                MOV STR[ECX],'-'
                INC ECX
                FABS                    ;make positive
        .ENDIF
        FRNDINT                         ;round to integer
        FIST    WHOLE                   ;save whole number part
        FLD     NUMB                    ;form and store fraction
        FABS
        FSUBR
        FSTP    FRACT
        MOV     EAX,WHOLE               ;convert whole part
        MOV     EBX,10
        PUSH    EBX
        .REPEAT
                MOV     EDX,0
                DIV     EBX
                ADD     DL,30H          ;convert to ASCII
                PUSH    EDX
        .UNTIL EAX == 0
        POP     EDX
        MOV     AH,3                    ;comma counter
        .WHILE EDX !== 10               ;whole part to ASCII
                POP     EBX
                DEC     AH
                .IF AH == 0 && EBX != 10
                        MOV     STR[ECX],'-'
                        INC     ECX
                        MOV     AH,3
                .ENDIF
                MOV     STR[ECX],DL
                INC     ECX
                MOV     EDX,EBX
        .ENDW
        MOV     STR[ECX],'-'            ;store decimal point
        INC     ECX
        POP     TEMP                    ;restore original CW
        FLDCW   TEMP
        FLD     FRACT                   ;convert fractional part
        .REPEAT
                FIMUL TEN
                FIST  TEMP
                MOV   AX,TEMP
                ADD   AL,30H
                MOV   STR[ECX],AL
                INC   ECX
                FISUB TEMP
                FXAM
                SAHF
        .UNTIL ZERO?
        FCOMP                           ;clear stack
        MOV     STR[ECX],0              ;store null
        RET


FTOA    ENDP
```

14–5   INTRODUCTION TO MMX TECHNOLOGY

The MMX[1] (multimedia extensions) technology adds 57 new instructions to the instruction set of the Pentium–Pentium 4 microprocessors. The MMX technology also introduces new general-purpose instructions. The new MMX instructions are designed for applications such as motion video, combined graphics with video, image processing, audio synthesis, speech synthesis and compression, telephony, video conferencing, 2D graphics, and 3D graphics. These instructions (new beginning with the Pentium in 1995) operate in parallel with other operations as the instructions for the arithmetic coprocessor.

Data Types

The MMX architecture introduces new packed data types. The data types are eight packed, consecutive 8-bit bytes; four packed, consecutive 16-bit words; and two packed, consecutive 32-bit doublewords. Bytes in this multibyte format have consecutive memory addresses and use the little endian form, as with other Intel data. See Figure 14–11 for the format for these new data types.

The MMX technology registers have the same format as a 64-bit quantity in memory and have two data access modes: 64-bit access mode and 32-bit access mode. The 64-bit access mode is used for 64-bit memory and registers transfers for most instructions. The 32-bit access mode is used for 32-bit memory and also register transfers for most instructions. The 32-bit transfers occur between microprocessor registers, and the 64-bit transfers occur between floating-point coprocessor registers.

Figure 14–12 illustrates the internal register set of the MMX technology extension and how it uses the floating-point coprocessor register set. This technique is called aliasing because the floating-point registers are shared as the MMX registers. That is, the MMX registers (MM_0–MM_7) are the same as the floating-point registers. Note that the MMX register set is 64 bits wide and uses the rightmost 64 bits of the floating-point register set.
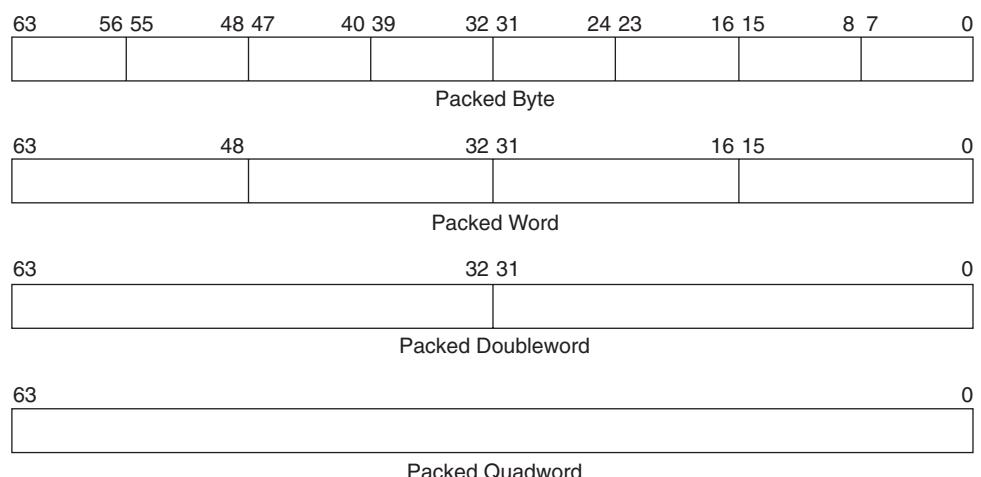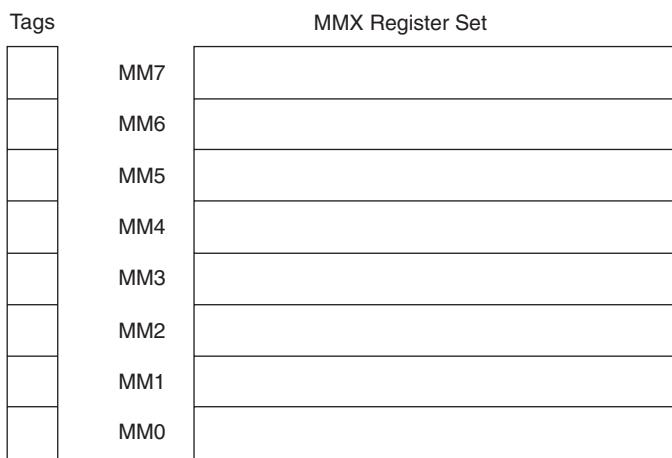

FIGURE 14–11   The structure of data stored in the MMX registers.

---

[1]MMX is a registered trademark of Intel Corporation.

FIGURE 14–12 The structure of the MMX register set. Note that MM_0 and FF_0 through MM_7 and FP_7 interchange with each other.

Tags   MMX Register Set

MM7

MM6

MM5

MM4

MM3

MM2

MM1

MM0

Instruction Set

The instruction for MMX technology includes arithmetic, comparison, conversion, logical, shift, and data transfer instructions. Although the instruction types are similar to the microprocessor's instruction set, the main difference is that the MMX instructions use the data types shown in Figure 14–11 instead of the normal data types used with the microprocessor.

Arithmetic Instructions. The set of arithmetic instructions includes addition, subtraction, multiplication, a special multiplication with an addition, and so on. Three additions exist. The PADD and PSUB instructions add or subtract packed signed or unsigned packed bytes, packed words, or packed doubleword data. The add instructions are appended with a B, W, or D to select the size, as in PADDB for a byte, PADDW for a word, and PADDD for a doubleword. The same is true for the PSUB instruction. The PMULHW and the PMULLW instructions perform multiplication on four pairs of l6-bit operands, producing 32-bit results. The PMULHW instruction multiplies the high-order l6 bits, and the PMULLW instruction multiplies the low-order 16 bits. The PMADDWD instruction multiplies and adds. After multiplying, the four 32-bit results are added to produce two 32-bit doubleword results.

The MMX instructions use operands just as the integer or floating-point instructions do. The difference is the register names (MM_0–MM_7). For example, the PADDB MM_1, MM_2 instruction adds the entire 64-bit contents of MM_2 to MM_1, byte by byte. The result is steered into MM_1. When each 8-bit section is added, any carries generated are dropped. For example, the byte A0H added to the byte 70H produces the byte sum of 10H. The true sum is 110H, but the carry is dropped. Note that the second operand or source can be a memory location containing the 64-bit packed source or an MMX register. You might say that this instruction performs the same function as eight separate byte-sized ADD instructions! If used in an application, this certainly speeds execution of the application. Like PADD, PSUB also does not carry or borrow. The difference is that if an overflow or underflow occurs, the difference becomes 7FH (+127) for an overflow and 80H (–128) for an underflow. Intel calls this **saturation**, because these values represent the largest and smallest signed bytes.

Comparison Instructions. There are two comparison instructions: PCMPEQ (equal) and PCMPGT (greater than). As with PADD and PSUB, there are three versions of each compare instruction: for example, PCMPEQUB (compares bytes), PCMPEQUW (compares words), and PCMPEQUD (compares doublewords). These instructions do not change the microprocessor flag bits; instead, the result is all ones for a true condition and all zeros for a false condition. For example, if the PCMPEQB MM_2, MM_3 instruction is executed and the least significant bytes of MM_2 and MM_3 = 10H and 11H, respectively, the result found in the least significant byte of

MM_2 is 00H. This indicates that the least significant bytes were not equal. If the least significant byte contained an FFH, it indicates that the two bytes were equal.

Conversion Instructions.   There are two basic conversion instructions: PACK and PUNPCK. PACK is available as PACKSS (signed saturation) and PACKUS (unsigned saturation). PUNPCK is available as PUNPCKH (unpack high data) and PUNPCKL (unpack low data). Similar to the prior instructions, these can be appended with B, W, or D for byte, word, and doubleword pack and unpack, but they must be used in combinations WB (word to byte) or DW (doubleword to word). For example, the PACKUSWB MM_3, MM_6 instruction packs the words from MM_6 into bytes in MM_3. If the unsigned word does not fit (too large) into a byte, the destination byte becomes an FFH. For signed saturation, we use the same values explained under addition.

Logic Instructions.   The logic instructions are PAND (AND), PANDN (NAND), POR (OR), and PXOR (Exclusive-OR). These instructions do not have size extensions, and perform these bit-wise operations on all 64 bits of the data. For example, the POR MM_2, MM_3 instruction ORs all 64 bits of MM_3 with MM_2. The logical sum is placed into MM_2 after the OR operation.

Shift Instruction.   This instruction contains logical shifts and an arithmetic shift right instruction. The logic shifts are PSLL (left) and PSRL (right). Variations are word (W), doubleword (D), and quadword (Q). For example, the PSLLQ $MM_{3,2}$ instruction shifts all 64 bits in MM_3 left two places. Another example is the PSLLD $MM_{3,2}$ instruction that shifts the two 32-bit doublewords in MM_3 left two places each.

The PSRA (arithmetic right shift) instruction functions in the same manner as the logical shifts, except that the sign-bit is preserved.

Data Transfer Instructions.   There are two data transfer instructions: MOVED and MOVEQ. These instructions allow transfers between registers and between a register and memory. The MOVED instruction transfers 32 bits of data between an integer register or memory location and an MMX register. For example, the MOVED ECX, MM_2 instruction copies the rightmost 32 bits of MM_2 into ECX. There is no instruction to transfer the leftmost 32 bits of an MMX register. You could use a shift right before a MOVED to do the transfer.

The MOVEQ instruction copies all 64 bits of an MMX register between memory or another MMX register. The MOVEQ MM_2, MM_3 instruction transfers 64 bits of MM_3 into MM_2.

EMMS Instruction.   The EMMS (empty MMX-state) instruction sets (11) all the tags in the floating-point unit, so the floating-point registers are listed as empty. The EMMS instruction must be executed before the return instruction at the end of any MMX procedure, or a subsequent floating-point operation will cause a floating-point interrupt error, crashing Windows or any other application. If you plan to use floating-point instructions within an MMX procedure, you must use the EMMS instruction before executing the floating-point instruction. All other MMX instructions clear the tags, which indicate that all floating-point registers are in use.

Instruction Listing.   Table 14–10 lists all the MMX instructions with the machine code so these instructions can be used with the assembler. At present, MASM does not support these new instructions unless you have upgraded to the latest version (6.15). The latest version can be found in the Windows Driver Development Kit (Windows DDK), which is available for a small shipping charge from Microsoft Corporation. It is also available in Visual Studio Express (search for ML.EXE). Any MMX instruction can be used inside Visual C++ using the inline assembler.

Programming Example.   Example 14–13 on p. 581 illustrates a simple programming example that uses the MMAX instructions to perform a task that takes eight times longer using normal microprocessor instruction. In this example an array of 1000 bytes of data (BLOCKA) is added to a second array of 1000 bytes (BLOCKB). The result is stored in a third array called BLOCKC. Example 14–13(a) lists a procedure that uses traditional assembly language to perform the addition and Example 14–13(b) shows the same process using MMX instructions.

TABLE 14–10 The MMX instruction set extension.

EMMS Empty MMX state	
0000 1111 0111 1111	
Example	
EMMS	
MOVED Move doubleword	
0000 1111 0110 1110 11 xxx rrr	reg → xreg
Examples	
MOVED MM3, EDX MOVED MM4, EAX	
0000 1111 0111 1110 11 xxx rrr	xreg → reg
Examples	
MOVED EAX, MM3 MOVED EBP, MM7	
0000 1111 0110 1110 oo xxx mmm	mem → xreg
Examples	
MOVED MM3, DATA1 MOVED MM5, BIG_ONE	
0000 1111 0111 1110 oo xxx mmm	xreg → mem
Examples	
MOVED DATA2, MM3 MOVED SMALL_POTS, MM7	
MOVEQ Move quadword	
0000 1111 0110 1111 11 xxx1 xxx2	xreg2 → xreg1
Examples	
MOVEQ MM3, MM2 ;copies MM2 to MM3 MOVEQ MM7, MM3	
0000 1111 0111 1111 11 xxx1 xxx2	xreg1 → reg2
Examples	
MOVEQ MM3, MM2 ;copies MM3 to MM2 MOVEQ MM7, MM3	
0000 1111 0110 1111 oo xxx mmm	mem → xreg
Examples	
MOVEQ MM3, DATA1 MOVEQ MM5, DATA3	
0000 1111 0111 1111 oo xxx mmm	xreg → mem
Examples	
MOVEQ DATA2, MM0 MOVEQ SMALL_POTS, MM3	

PACKSSDW	Pack signed doubleword to word	
0000 1111 0110 1011 11 xxx1 xxx2		xreg2 → xreg1
Examples		
PACKSSDW MM1, MM2 PACKSSDW MM7, MM3		
0000 1111 0111 1011 oo xxx mmm		mem → xreg
Examples		
PACKSSDW MM3, BUTTON PACKSSDW MM7, SOUND		

PACKSSWB	Pack signed word to byte	
0000 1111 0110 0011 11 xxx1 xxx2		xreg2 → xreg1
Examples		
PACKSSWB MM1, MM2 PACKSSWB MM7, MM3		
0000 1111 0111 0011 oo xxx mmm		mem → xreg
Examples		
PACKSSWB MM3, BUTTON PACKSSWB MM7, SOUND		

PACKUSWB	Pack unsigned word to byte	
0000 1111 0110 0111 11 xxx1 xxx2		xreg2 → xreg1
Examples		
PACKUSWB MM1, MM2 PACKUSWB MM7, MM3		
0000 1111 0111 0111 oo xxx mmm		mem → xreg
Examples		
PACKUSWB MM3, BUTTON PACKUSWB MM7, SOUND		

PADD	Add with truncation	Byte, word, and doubleword
0000 1111 1111 11gg 11 xxx1 xxx2		xreg2 → xreg1
Examples		
PADDB MM1, MM2 PADDW MM7, MM3 PADDD MM3, MM4		
0000 1111 1111 11gg oo xxx mmm		mem → xreg
Examples		
PADDB MM3, BUTTON PADDW MM7, SOUND PADDD MM3, BUTTER		

PADDS	Add with signed saturation	Byte and word
0000 1111 1110 11gg 11 xxx1 xxx2 Examples		xreg2 → xreg1
PADDSB MM1, MM2 PADDSW MM7, MM3		
0000 1111 1110 11gg oo xxx mmm Examples		mem → xreg
PADDSB MM3, BUTTON PADDSW MM7, SOUND		

PADDUS	Add with unsigned saturation	Byte and word
0000 1111 1101 11gg 11 xxx1 xxx2 Examples		xreg2 → xreg1
PADDUSB MM1, MM2 PADDUSW MM7, MM3		
0000 1111 1101 11gg oo xxx mmm Examples		mem → xreg
PADDUSB MM3, BUTTON PADDUSW MM7, SOUND		

PAND	And	
0000 1111 1101 1011 11 xxx1 xxx2 Examples		xreg2 → xreg1
PAND MM1, MM2 PAND MM7, MM3		
0000 1111 1101 1011 oo xxx mmm Examples		mem → xreg
PAND MM3, BUTTON PAND MM7, SOUND		

PANDN	Nand	
0000 1111 1101 1111 11 xxx1 xxx2 Examples		xreg2 → xreg1
PANDN MM1, MM2 PANDN MM7, MM3		
0000 1111 1101 1111 oo xxx mmm Examples		mem → xreg
PANDN MM3, BUTTON PANDN MM7, SOUND		

PCMPEQU	Compare for equality	Byte, word, and doubleword
0000 1111 0111 01gg 11 xxx1 xxx2 Examples		xreg2 → xreg1
PCMPEQUB MM1, MM2 PCMPEQUW MM7, MM3 PCMPEQUD MM0, MM5		
0000 1111 0111 01gg oo xxx mmm Examples		mem → xreg
PCMPEQUB MM3, BUTTON PCMPEQUW MM7, SOUND PCMPEQUD MM0, FROG		

PCMPGT	Compare for greater than	Byte, word, and doubleword
0000 1111 0110 01gg 11 xxx1 xxx2 Examples		xreg2 → xreg1
PCMPGTB MM1, MM2 PCMPGTW MM7, MM3 PCMPGTD MM0, MM5		
0000 1111 0110 01gg oo xxx mmm Examples		mem → xreg
PCMPGTB MM3, BUTTON PCMPGTW MM7, SOUND PCMPGTD MM0, FROG		

PMADD	Multiply and add	
0000 1111 1111 0101 11 xxx1 xxx2 Examples		xreg2 → xreg1
PMADD MM1, MM2 PMADD MM7, MM3		
0000 1111 1111 0101 oo xxx mmm Examples		mem → xreg
PMADD MM3, BUTTON PMADD MM7, SOUND		

PMULH	Multiplication–high	
0000 1111 1110 0101 11 xxx1 xxx2 Examples		xreg2 → xreg1
PMULH MM1, MM2 PMULH MM7, MM3		
0000 1111 1110 0101 oo xxx mmm Examples		mem → xreg
PMULH MM3, BUTTON PMULH MM7, SOUND		

PMULL	Multiplication–low	
0000 1111 1101 0101 11 xxx1 xxx2		xreg2 → xreg1
Examples		
PMULL MM1, MM2 PMULL MM7, MM3		
0000 1111 1101 0101 oo xxx mmm		mem → xreg
Examples		
PMULL MM3, BUTTON PMULL MM7, SOUND		

POR	Or	
0000 1111 1110 1011 11 xxx1 xxx2		xreg2 → xreg1
Examples		
POR MM1, MM2 POR MM7, MM3		
0000 1111 1110 1011 oo xxx mmm		mem → xreg
Examples		
POR MM3, BUTTON POR MM7, SOUND		

PSLL	Shift left	Word, doubleword, and quadword
0000 1111 1111 00gg 11 xxx1 xxx2		xreg2 → xreg1
Examples		
PSLLW MM1, MM2 PSLLD MM7, MM3 PSLLQ MM6, MM5		
0000 1111 1111 00gg oo xxx mmm		mem → xreg shift count in memory
Examples		
PSLLW MM3, BUTTON PSLLD MM7, SOUND PSLLQ MM2, COUNT1		
0000 1111 0111 00gg 11 110 mmm data8		xreg by count shift count is data8
Examples		
PSLLW MM3, 2 PSLLD MM0, 6 PSLLQ MM7, 1		

PSRA	Shift arithmetic right	Word, doubleword, and quadword
0000 1111 1110 00gg 11 xxx1 xxx2		xreg2 → xreg1
Examples		
PSRAW MM1, MM2 PSRAD MM7, MM3 PSRAQ MM6, MM5		
0000 1111 1110 00gg oo xxx mmm		mem → xreg shift count in memory
Examples		
PSRAW MM3, BUTTON MM7, SOUND PSRAQ MM2, COUNT1		
0000 1111 0111 00gg 11 100 mmm data8		xreg by count shift count is data8
Examples		
PSRAW MM3, 2 PSRAD MM0, 6 PSRAQ MM7, 1		

PSRL	Shift right	Word, doubleword, and quadword
0000 1111 1101 00gg 11 xxx1 xxx2		xreg2 → xreg1
Examples		
PSRLW MM1, MM2 PSRLD MM7, MM3 PSRLQ MM6, MM5		
0000 1111 1101 00gg oo xxx mmm		mem → xreg shift count in memory
Examples		
PSRLW MM3, BUTTON PSRLD MM7, SOUND PSRLQ MM2, COUNT1		
0000 1111 0111 00gg 11 010 mmm data8		xreg by count shift count is data8
Examples		
PSRLW MM3, 2 PSRLD MM0, 6 PSRLQ MM7, 1		

PSUB	Subtract with truncation	Byte, word, and doubleword
0000 1111 1111 10gg 11 xxx1 xxx2 Examples		xreg2 → xreg1
PSUBB MM1, MM2 PSUBW MM7, MM3 PSUBD MM3, MM4		
0000 1111 1111 10gg oo xxx mmm Examples		mem → xreg
PSUBB MM3, BUTTON PSUBW MM7, SOUND PSUBD MM3, BUTTER		
PSUBS	Subtract with signed saturation	Byte, word, and doubleword
0000 1111 1110 10gg 11 xxx1 xxx2 Examples		xreg2 → xreg1
PSUBSB MM1, MM2 PSUBSW MM7, MM3 PSUBSD MM3, MM4		
0000 1111 1110 10gg oo xxx mmm Examples		mem → xreg
PSUBSB MM3, BUTTON PSUBSW MM7, SOUND PSUBSD MM3, BUTTER		
PSUBUS	Subtract with unsigned saturation	Byte, word, and doubleword
0000 1111 1101 10gg 11 xxx1 xxx2 Examples		xreg2 → xreg1
PSUBUSB MM1, MM2 PSUBUSW MM7, MM3 PSUBUSD MM3, MM4		
0000 1111 1101 10gg oo xxx mmm Examples		mem → xreg
PSUBUSB MM3, BUTTON PSUBUSW MM7, SOUND PSUBUSD MM3, BUTTER		

PUNPCKH	Unpack high data to next larger	Byte, word, doubleword
0000 1111　0110 10gg　11 xxx1 xxx2 Examples		xreg2 → xreg1
PUNPCKH　MM1, MM2 PUNPCKH　MM3, MM4		
0000 1111　0110 10gg　oo xxx mmm Examples		mem → xreg
PUNPCKH　MM7, WATER PUNPCKH　MM2, DOGGY		
PUNPCKL	Unpack low data to next larger	Byte, word, doubleword
0000 1111　0110 00gg　11 xxx1 xxx2 Examples		xreg2 → xreg1
PUNPCKL　MM1, MM2 PUNPCKL　MM3, MM4		
0000 1111　0110 00gg　oo xxx mmm Examples		mem → xreg
PUNPCKL　MM7, WATER PUNPCKL　MM2, DOGGY		
PXOR	Bitwise Exclusive-OR	Byte, word, doubleword
0000 1111　1110 1111　11 xxx1 xxx2 Examples		xreg2 → xreg1
PXOR　MM2, MM3 PXOR　MM4, MM7 PXOR　MM0, MM1		
0000 1111　1110 1111　oo xxx mmm Examples		mem → xreg
PXOR　MM2, FROGS PXOR　MM4, WALTER		

EXAMPLE 14–13(a)

```
;Procedure that adds BLOCKA0 to BLOCKB and stores the sums in BLOCKC

BLOCKA DB     1000 DUP(?)
BLOCKB DB     1000 DUP(?)
BLOCKC DB     1000 DUP(?)

SUM    PROC   NEAR

       MOV    ECX,1000
       .REPEAT
              MOV    AL,BLOCKA[ECX-1]
              ADD    AL.BLOCKB[ECX-1]
              MOV    BLOCKC[ECX-1]
       .UNTILCXZ
       RET

SUM    ENDP
```

EXAMPLE 14–13(b)

```
;Procedure that adds BLOCKA0 to BLOCKB and stores the sums in BLOCKC

BLOCKA DB     1000 DUP(?)
BLOCKB DB     1000 DUP(?)
BLOCKC DB     1000 DUP(?)

SUMM   PROC   NEAR

       MOV    ECX,125
       .REPEAT
              MOVEQ  MM0,QWORD PTR BLOCKA[ECX-8]
              PADDB  MM0,QWORD PTR BLOCKB[ECX-8]
              MOVEQ  QWORD PTR BLOCKC[ECX-8],MM0
       .UNTILCXZ
       RET

SUMM   ENDP
```

If you closely compare the programs, notice that the MMX version goes through its loop of three instructions 125 times, while the traditional software goes through its loop 1000 times. The MMX version executes eight times faster than the traditional version. This occurs because eight bytes (QWORD) are added at a time.

14–6   INTRODUCTION TO SSE TECHNOLOGY

The latest type of instruction added to the instruction set of the Pentium 4 is SIMD (**single-instruction multiple data**). As the name implies, a single instruction operates on multiple data in much the same way as do the MMX instructions, which are SIMD instructions that operate on multiple data. The MMX instruction set functions with integers; the SIMD instruction set functions with floating-point numbers as well as integers. The SIMD extension instructions first appeared in the Pentium III as SSE (**streaming SIMD extensions**) instructions. Later, SSE 2 instructions were added to the Pentium 4, and new to the Pentium 4 (beginning with the 90-nanometer E model) are SSE 3 instructions. The SSE 3 extensions are also found in the Core2 microprocessor.

Recall that the MMX instructions shared registers with the arithmetic coprocessor. The SSE instructions use a new and separate register array to operate on data. Figure 14–13 illustrates an array of eight 128-bit-wide registers that function with the SSE instructions. These new registers

**FIGURE 14–13**   The XMM registers used by the SSE instructions.
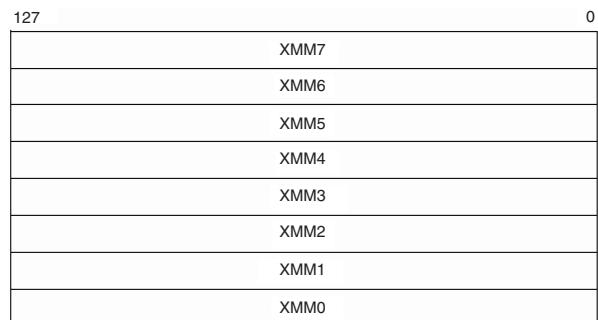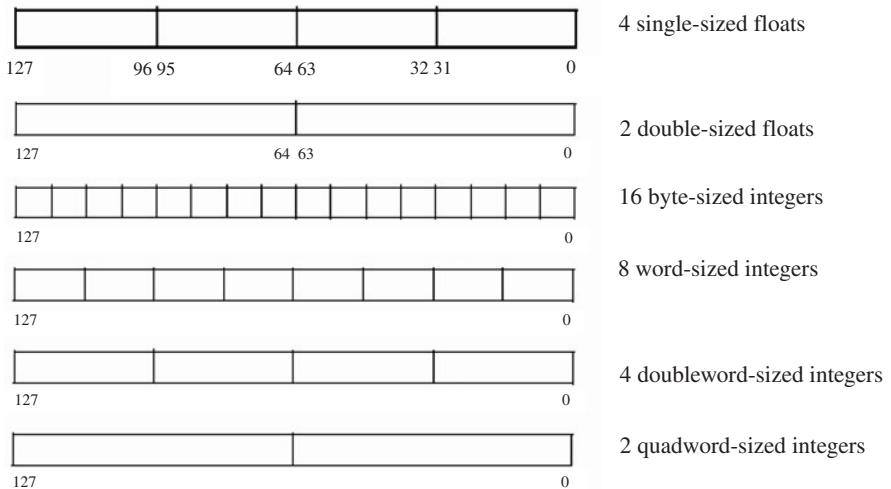
FIGURE 14–14   Data formats for the SSE 2 and SSE 3 instructions.
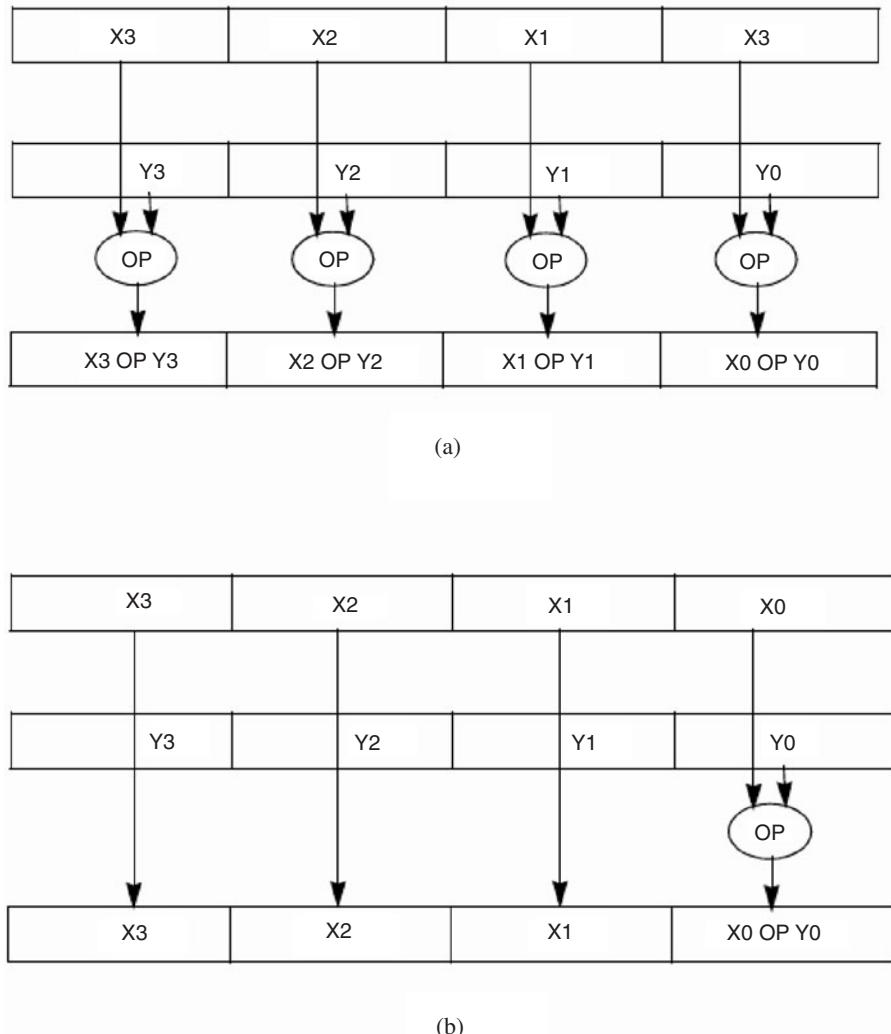

are called **XMM registers (XMM$_0$–XMM7)**, which denote extended multimedia registers. To accommodate this new 128-bit-wide data size, a new keyword is added called **OWORD**. An OWORD (**octalword**) designates a 128-bit variable, as in OWORD PTR for the SSE instruction set. A double quadword is also used at times to specify a 128-bit number.

   Just as the MMX registers can contain multiple data types, so can the XMM registers of the SSE unit. Figure 14–14 illustrates the data types that can appear in any XMM register for various SSE instructions. An XMM register can hold four single-precision floating-point numbers or two double-precision floating-point numbers. XMM registers can also hold sixteen 8-bit integers, eight 16-bit integers, four 32-bit integers, or two 64-bit integers. This is a twofold increase in the capacity of the system when compared to the integers contained in MMX registers and hence a twofold increase in execution speeds of integers operations that use the XMM registers and SSE instructions. For new applications that are designed to execute on a Pentium 4 or newer microprocessor, the SSE instructions are used in place of the MMX instructions. Because not all machines are yet Pentium 4 class machines, there still is a need to include MMX technology instructions in a program for compatibility to these older systems.

## Floating-Point Data

Floating point data are operated upon as either packed or scalar, and either single-precision or double-precision. The packed operation is performed on all sections at a time; the scalar form is only operated on the rightmost section of the register contents. Figure 14–15 shows both the packed and scalar operations on SSE data in XMM registers. The scalar form is comparable to

FIGURE 14–15 Packed (a) and scalar (b) operations for single-precision floating-point numbers.


(a)

(b)
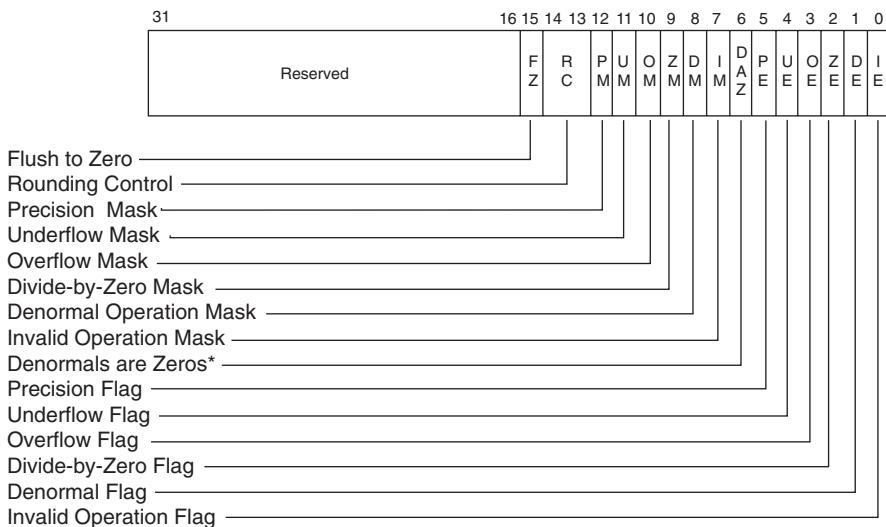
the operation performed by the arithmetic coprocessor. Opcodes are appended with PS (packed single), SS (scalar single), PD (packed double), or SD (scaled double) to form the desired instruction. For example, the opcode for a multiply is MUL, but the opcode for a packed double is MULPD and MULSD for a scalar double multiplication. The single-precision multiplies are MULPS and MULSS. In other words, once the two-letter extension and its meaning are understood, it is relatively easy to master the new SSE instructions.

The Instruction Set

The SSE instructions have a few new types added to the instruction set. The floating-point unit does not have a reciprocal instruction, which is used quite often to solve complex equations. The reciprocal instruction ($\frac{1}{n}$) now appears in the SSE extensions as the RCP instruction, which generates reciprocals and is written as RCPPS, RCPSS, RCPPD, and RCPSD. There is also a reciprocal of a square root ($\frac{1}{\sqrt{n}}$) instruction, called RSQRT, which is written as RSQRTPS, RSQRTSS, RSQRTPD, and RSQRTSD.

The remainder of the instructions for the SSE unit are basically the same as for the microprocessor and MMX unit except for a few cases. The instruction table in Appendix B lists the instructions, but does not list the extensions (PS, SS, PD, and DS) to the instructions. Again note

FIGURE 14–16   The MXCSR (control/status) register of the SSE unit.

*The denormals-are-zeros flag was introduced in the Pentium 4 and Intel Xeon processors.

that SSE 2 and SSE 3 contain double-precision operations and SSE does not. Instructions that start with the letter P operate on integer data that is byte, word, doubleword, or quadword sized. For example, the PADDB XMM_0, XMM_1 instruction adds the 16 byte-sized integers in the XMM_1 register to the 16 byte-sized integers in the XMM_0 register. PADDW adds 16-bit integers, PADDD adds doublewords, and PADDQ adds quadwords. The execution times are not provided by Intel so they do not appear in the appendix for these instructions.

The Control/Status Register

The SSE unit also contains a control/status register accessed as MXCSR. Figure 14–16 illustrates the MXCSR for the SSE unit. Notice that this register is very similar to the control/status register of the arithmetic coprocessor presented earlier in this chapter. This register sets the precision and rounding modes for the coprocessor, as does the control register for the arithmetic coprocessor, and it provides information about the operation the SSE unit.

The SSE control/status register is loaded from memory using the LDMXCSR and FXRSTOR instructions or stored into the memory using the STMXCSR and FXSAVE instructions. Suppose the rounding control (see Figure 14–6 for the state of the rounding control bits) needs to be changed to round toward positive infinity (RC = 10). Example 14–14 shows the software that changes only the rounding control bits of the control/status register.

EXAMPLE 14–14

```
;change the rounding control to 10.

      STMXCSR  CONTROL      ;save the control/status register
      BTS  CONTROL,14       ;set bit 14
      BTR  CONTROL,13       ;clear bit 13
      LDMXCSR CONTROL       ;reload control/status register
```

Programming Examples

A few programming examples are needed to show how to use the SSE unit. As mentioned, the SSE unit allows floating-point and integer operations on multiple data. Suppose that the capacitive

reactance is needed for a circuit that contains a 1.0 μF capacitor at various frequencies from 100 Hz to 10,000 Hz in 100 Hz steps. The equation used to calculate capacitive reactance is:

$$XC = \frac{1}{2\pi FC}$$

Example 14–15 illustrates a procedure that generates the 100 outcomes for this equation using the SSE unit and single-precision floating-point data. The program listed in Example 14–15(a) uses the SSE unit to perform four calculations per iteration, while the program in Example 14–15(b) uses the floating-point coprocessor to calculate XC one at a time. Example 14–15(c) is yet another example in C++. Examine the loop to see that the first example goes through the loop 25 times and the second goes through the loop 100 times. Each time the loop executes in Example 14–15(a) it executes seven instructions ($25 \times 7 = 175$), which takes 175 instruction times. Example 14–15(b) executes eight instructions per iteration of its loop ($100 \times 8 = 800$), which requires 800 instruction times. By using this parallelism, the SSE unit allows the calculations to be accomplished in much less time than any other method. The C++ version in Example 14–15(c) uses the directive `__declspec(align(16))` before each variable to make certain that they are aligned properly in the memory. If these are missing, the program will not function because the SSE memory variables must be aligned on at least quadword boundaries (16). This final version executes at about $4\frac{1}{2}$ times faster than Example 14–15(b);

EXAMPLE 14–15(a)

```
;using the SSE unit

XC      DD      100 DUP(?)
CAP     DD      1.0E-6, 1.0E-6, 1.0E-6, 1.0E-6
F       DD      100.0, 200.0, 300.0, 400.0
INCR    DD      400.0, 400.0, 400.0, 400.0
PI      DD      4 DUP(?)

FXC     PROC    NEAR

        MOV     ECX,0
        FLDPI                                   ;get π
        FADD    ST,ST(0)                        ;double π
        FST     PI                              ;store four 2π
        FST     PI+4
        FST     PI+8
        FSTP    PI+12
        MOVUPS  XMM0,OWORD PTR PI               ;get four 2πs
        MOVUPS  XMM1,OWORD PTR INCR             ;get increment
        .REPEAT
                MOVUPS  XMM2,OWORD PTR F         ;load frequencies
                MULPS   XMM2,XMM0                ;generate 2πFs
                MULPS   XMM2,CAPS                ;2πFC
                RCPPS   XMM3,XMM2                ;find reciprocal
                MOVUPS  OWORD PTR XC[ECX],XMM3   ;save four XCs
                ADD     ECX,16                   ;move pointer
                ADDPS   OWORD PTR F,XMM1         ;increment Fs
        .UNTIL ECX == 100
        RET

FXC     ENDP
```

EXAMPLE 14–15(b)

```
;using the coprocessor

XC      DD      100 DUP(?)
CAP     DD      1.0E-6
```

```
F       DD      0
INCR    DD      100.0

FXC1    PROC    NEAR

        FLDPI                           ;get π
        FADD    ST,ST(0)                ;form 2π
        FMUL    ST,CAP                  ;form 2πC
        MOV     ECX,0
        .REPEAT
                FLD     F               ;get frequency
                FADD    INCR            ;add increment
                FST     F               ;save it for next time
                FMUL    ST,ST(1)        ;2πFC
                FLD1                    ;form reciprocal
                FDIVR
                FSTP    XC[ECX*4]       ;save XC
                INC     ECX
        .UNTIL  ECX == 100
        FCOMP                           ;clear coprocessor stack
        RET

FXC1    ENDP
```

EXAMPLE 14–15(c)

```
void FindXC()
{
     //floating-point example using C++ with the inline assembler

     __declspec(align(16)) float f[4] = {-300,-200,-100,0};
     __declspec(align(16)) float pi[4];
     __declspec(align(16)) float caps[4] = {1.0E-6, 1.0E-6, 1.0E-6, 1.0E-6};
     __declspec(align(16)) float incr[4] = {400, 400, 400, 400};
     __declspec(align(16)) float Xc[400];
     _asm
     {

             fldpi                                   ;form 2π
             fadd    st,st(0)
             fst     pi
             fst     pi+4
             fst     pi+8
             fstp    pi+12
             movaps  xmm0,oword ptr pi
             movaps  xmm1,oword ptr incr
             movaps  xmm3,oword ptr f
             mulps   xmm0,oword ptr caps             ;2πC
             mov     ecx,0
LOOP1:
             movaps  xmm2,xmm3
             addps   xmm2,xmm1
             movaps  xmm3,xmm2
             mulps   xmm2,xmm0
             rcpps   xmm2,xmm2                       ;recipocal
             movaps  oword ptr Xc[ecx],xmm2
             add     ecx,16
             cmp     ecx,400
             jnz     LOOP1
     }
}
```

The first example in this section (Example 14–15) used floating-point number to perform multiple calculations, but the SSE unit can also operate on integers. The example illustrated in Example 14–16 uses integer operation to add BlockA to BlockB and store the sum in BlockC. Each block contains 4000 eight-bit numbers. Example 14–16(a) lists an assembly language

procedure that forms the sums using the standard integer unit of the microprocessor, which requires 4000 iterations to accomplish.

EXAMPLE 14–16(a)

```
;A procedure that forms 4000 eight-bit sums

SUMS   PROC   NEAR

       MOV    ECX,0
       .REPEAT
               MOV  AL,BLOCKA[ECX]
               ADD  AL,BLOCKB[ECX]
               MOV  BLOCKC,[ECX]
               INC  ECX
       .UNTIL ECX == 4000
       RET

SUMS   EMDP
```

EXAMPLE 14–16(b)

```
;A procedure that uses SSE to form 4000 eight-bit sums

SUMS1  PROC   NEAR

       MOV    ECX,0
       .REPEAT
               MOVDQA   XMM0,OWORD PTR BLOCKA[ECX]
               PADDB    XMM0,OWORD PTR BLOCKB[ECX]
               MOVDQA   OWORD PTR BLOCKC[ECX]
               ADD      ECX,16
       .UNTIL ECX == 4000
       RET

SUMS1  ENDP
```

Both example programs generate 4000 sums, but the second example using the SSE unit does it by passing through its loop 250 times, while the first example requires 4000 passes. Hence, the second example functions 16 times faster because of the SSE unit. Notice how the PADDB (an instruction presented with the MMX unit) is used with the SSE unit. The SSE unit uses the same commands as the MMX except the registers are different. The MMX unit uses 64-bit-wide MM registers and the SSE unit uses 128-bit-wide XMM registers.

Optimization

The compiler in Visual C++ does have optimization for the SSE unit, but it does not optimize the examples presented in this chapter. It will attempt to optimize a single equation in a statement if the SSE unit can be utilized for the equation. It does not look at a program for blocks of operations that can be optimized as in the examples presented here. Until a compiler and extensions are developed so parallel operations such as these can be included, programs that require high speeds will require hand-coded assembly language for optimization. This is especially true of the SSE unit.

14–7   SUMMARY

1. The arithmetic coprocessor functions in parallel with the microprocessor. This means that the microprocessor and coprocessor can execute their respective instructions simultaneously.

2. The data types manipulated by the coprocessor include signed integer, floating-point, and binary-coded decimal (BCD).

3. Three forms of integers are used with the coprocessor: word (16 bits), short (32 bits), and long (64 bits). Each integer contains a signed number in true magnitude for positive numbers and two's complement form for negative numbers.

4. A BCD number is stored as an 18-digit number in 10 bytes of memory. The most significant byte contains the sign-bit, and the remaining nine bytes contain an 18-digit packed BCD number.

5. The coprocessor supports three types of floating-point numbers: single-precision (32 bits), double-precision (64 bits), and temporary extended-precision (80 bits). A floating-point number has three parts: the sign, biased exponent, and significant. In the coprocessor, the exponent is biased with a constant and the integer bit of the normalized number is not stored in the significant, except in the temporary extended-precision form.

6. Decimal numbers are converted to floating-point numbers by (a) converting the number to binary, (b) normalizing the binary number, (c) adding the bias to the exponent, and (d) storing the number in floating-point form.

7. Floating-point numbers are converted to decimal by (a) subtracting the bias from the exponent, (b) un-normalizing the number, and (c) converting it to decimal.

8. The 80287 uses I/O space for the execution of some of its instructions. This space is invisible to the program and is used internally by the 80286/80287 system. These 16-bit I/O addresses (00F8H–00FFH) must not be used for I/O data transfers in a system that contains an 80287. The 80387, 80486/7, and Pentium through Core2 use I/O addresses 800000F8H–800000FFH.

9. The coprocessor contains a status register that indicates busy, the conditions that follow a compare or test, the location of the top of the stack, and the state of the error bits. The FSTSW AX instruction, followed by SAHF, is often used with conditional jump instructions to test for some coprocessor conditions.

10. The control register of the coprocessor contains control bits that select infinity, rounding, precision, and error masks.

11. The following directives are often used with the coprocessor for storing data: DW (defineword), DD (define doubleword), DQ (define quadword), and DT (define 10 bytes).

12. The coprocessor uses a stack to transfer data between itself and the memory system. Generally, data are loaded to the top of the stack or removed from the top of the stack for storage.

13. All internal coprocessor data are always in the 80-bit extended-precision form. The only time that data are in any other form is when they are stored or loaded from the memory.

14. The coprocessor addressing modes include the classic stack mode, register, register with apop, and memory. Stack addressing is implied. The data at ST become the source, at ST(1) the destination, and the result is found in ST after a pop.

15. The coprocessor's arithmetic operations include addition, subtraction, multiplication, division, and square root calculation.

16. There are transcendental functions in the coprocessor's instruction set. These functions find the partial tangent or arctangent, $2X - 1$, $Y \log_2 X$, and $Y \log_2 (X + 1)$. The 80387, 80486/7, and Pentium–Core2 also include sine and cosine functions.

17. Constants are stored inside the coprocessor that provide $+0.0$, $+1.0$, π, $\log_2 10$, $\log_2 \varepsilon$, $\log_2 2$, and $\log_\varepsilon 2$.

18. The 80387 functions with the 80386 microprocessor and the 80487SX functions with the 80486SX microprocessor, but the 80486DX and Pentium–Core2 contain their own internal arithmetic coprocessor. The instructions performed by the earlier versions are available on these coprocessors. In addition to these instructions, the 80387, 80486/7, and Pentium–Core2 also can find the sine and cosine.

19. The Pentium Pro through Core2 contain two new floating-point instructions: FCMOV and FCOMI. The FCMOV instruction is a conditional move and the FCOMI performs the same task as FCOM, but it also places the floating-point flags into the system flag register.
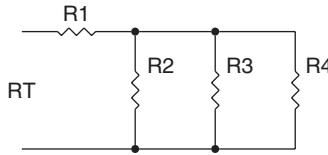
20. The MMX extension uses the arithmetic coprocessor registers for MM_0–MM_7. Therefore, it is important that coprocessor software and MMX software do not try to use them at the same time.

21. The instructions for the MMX extensions perform arithmetic and logic operations on bytes (eight at a time), words (four at a time), doublewords (two at a time), and quadwords. The operations performed are addition, subtraction, multiplication, division, AND, OR, Exclusive-OR, and NAND.

22. Both the MMX unit and the SSE unit employ SIMD techniques to perform parallel operations on multiple data with a single instruction. The SSE unit performs operations on integers and floating-point numbers. The registers in the SSE unit are 128 bits in width and can hold (SSE 2 or newer) 16 bytes at a time or four single-precision floating-point numbers. The SSE unit contains registers XMM_0–XMM_7.

23. New applications written for the Pentium 4 should contain SSE instructions in place of MMX instructions.

24. The OWORD pointer has been added to address 128-bit-wide numbers, which are referred to as octal words or double quadwords.

14–8   QUESTIONS AND PROBLEMS

1. List the three types of data that are loaded or stored in memory by the coprocessor.
2. List the three integer data types, the range of the integers stored in them, and the number of bits allotted to each.
3. Explain how a BCD number is stored in memory by the coprocessor.
4. List the three types of floating-point numbers used with the coprocessor and the number of binary bits assigned to each.
5. Convert the following decimal numbers into single-precision floating-point numbers:
    (a) 28.75
    (b) 624
    (c) – 0.615
    (d) + 0.0
    (e) – 1000.5
6. Convert the following single-precision floating-point numbers into decimal:
    (a) 11000000 11110000 00000000 00000000
    (b) 00111111 00010000 00000000 00000000
    (c) 01000011 10011001 00000000 00000000
    (d) 01000000 00000000 00000000 00000000
    (e) 01000001 00100000. 00000000 00000000
    (f) 00000000 00000000 00000000 00000000
7. Explain what the coprocessor does when a normal microprocessor instruction executes.
8. Explain what the microprocessor does when a coprocessor instruction executes.
9. What is the purpose of the C_3–C_0 bits in the status register?
10. What operation is accomplished with the FSTSW AX instruction?
11. What is the purpose of the IE bit in the status register?
12. How can SAHF and a conditional jump instruction be used to determine whether the top of the stack (ST) is equal to register ST(2)?
13. How is the rounding mode selected in the 80X87?
14. What coprocessor instruction uses the microprocessor's AX register?
15. What I/O ports are reserved for coprocessor use with the 80287?
16. How are data stored inside the coprocessor?
17. What is a NAN?

18. Whenever the coprocessor is reset, the top of the stack register is register number _____.

19. What does the term *chop* mean in the rounding control bits of the control register?

20. What is the difference between affine and projective infinity control?

21. What microprocessor instruction forms the opcodes for the coprocessor?

22. The FINIT instruction selects _____-precision for all coprocessor operations.

23. Using assembler pseudo-opcodes, form statements that accomplish the following:
    (a) Store a 23.44 into a double-precision floating-point memory location FROG.
    (b) Store a –123 into a 32-bit signed integer location DATA3.
    (c) Store a –23.8 into a single-precision floating-point memory location DATAL.
    (d) Reserve double-precision memory location DATA2.

24. Describe how the FST DATA instruction functions. Assume that DATA is defined as a 64-bit memory location.

25. What does the FILD DATA instruction accomplish?

26. Form an instruction that adds the contents of register 3 to the top of the stack.

27. Describe the operation of the FADD instruction.

28. Choose an instruction that subtracts the contents of register 2 from the top of the stack and stores the result in register 2.

29. What is the function of the FBSTP DATA instruction?

30. What is the difference between a forward and a reverse division?

31. What is the purpose of the Pentium Pro FCOMI instruction?

32. What does a Pentium Pro FCMOVB instruction accomplish?

33. What must occur before executing any FCMOV instruction?

34. Develop a procedure that finds the reciprocal of the single-precision floating-point number. The number is passed to the procedure in EAX and must be returned as a reciprocal in EAX.

35. What is the difference between the FTST instruction and FXAM?

36. Explain what the F2XM1 instruction calculates.

37. Which coprocessor status register bit should be tested after the FSQRT instruction executes?

38. Which coprocessor instruction pushes π onto the top of the stack?

39. Which coprocessor instruction places 1.0 at the top of the stack?

40. What will FFREE ST(2) accomplish when executed?

41. Which instruction stores the environment?

42. What does the FSAVE instruction save?

43. Develop a procedure that finds the area of a rectangle ($A = L \times W$). Memory locations for this procedure are single-precision floating-point locations *A, L,* and *W*.

44. Write a procedure that finds the capacitive reactance ($XC = \frac{1}{2\pi FC}$). Memory locations for this procedure are single-precision floating-point locations *XC, F,* and *C1* for *C*.

45. Develop a procedure that generates a table of square roots for the integers 2 through 10. The results must be stored as single-precision floating-point numbers in an array called ROOTS.

46. When is the FWAIT instruction used in a program?

47. What is the difference between the FSTSW and FNSTSW instructions?

48. Given the series/parallel circuit and equation illustrated in Figure 14–17, develop a program using single-precision values for R_1, R_2, R_3, and R_4 that finds the total resistance and stores the result at single-precision location RT.

49. Develop a procedure that finds the cosine of a single-precision floating-point number. The angle, in degrees, is passed to the procedure in EAX and the cosine is returned in EAX. Recall that FCOS finds the cosine of an angle expressed in radians.

50. Given two arrays of double-precision floating-point data ($ARRAY_1$ and $ARRAY_2$) that each contain 100 elements, develop a procedure that finds the product of $ARRAY_1$ times $ARRAY_2$, and then stores the double-precision floating-point result in a third array ($ARRAY_3$).

FIGURE 14–17   The series/parallel circuit (question 48).

$$RT = R1 + \cfrac{1}{\cfrac{1}{R2} + \cfrac{1}{R3} + \cfrac{1}{R4}}$$

51. Develop a procedure that takes the single-precision contents of register EBX times π and stores the result in register EBX as a single-precision floating-point number. You must use memory to accomplish this task.

52. Write a procedure that raises a single-precision floating-point number X to the power Y. Parameters are passed to the procedure with *EAX* = X and *EBX* = Y. The result is passed back to the calling sequence in ECX.

53. Given that the $LOG_{10} X = (LOG_2 10)^{-1} \times LOG_2 X$, write a procedure called LOG_{10} that finds the LOG_{10} of the value (X) at the stack top. Return the LOG_{10} at the stack top at the end of the procedure.

54. Use the procedure developed in question 53 to solve the equation

$$\text{Gain in decibels} = 20\log_{10}\frac{V_{out}}{V_{in}}$$

54. The program should take arrays of single-precision values for V_{out} and V_{in} and store the decibel gains in a third array called DBG. These are 100 values V_{out} and V_{in}.

55. What is the MMX extension to the Pentium–Core2 microprocessors?

56. What is the purpose of the EMMS instruction?

57. Where are the MM_0–MM_7 registers found in the microprocessor?

58. What is signed saturation?

59. What is unsigned saturation?

60. How could all of the MMX registers be stored in the memory with one instruction?

61. Write a short program that uses MMX instruction to multiply the word-size numbers in arrays and store the 32-bit results in a third array. The source arrays are 256 words long.

62. What are SIMD instructions?

63. What are SSE instructions?

64. The XMM registers are _____ bits wide.

65. A single XMM register can hold _____ single-precision floating-point numbers.

66. A single XMM register can hold _____ byte-sized integers.

67. What is an OWORD?

68. Can floating-point instructions for the arithmetic coprocessor execute at the same time as SSE instructions?

69. Develop a C++ function (using inline assembly code) that computes (using scalar SSE instructions and floating-point instructions) and returns a single-precision number that represents the resonant frequency from parameters (L and C) passed to it to solve the following equation:

$$Fr = \frac{1}{2\pi\mathbf{2}\,\overline{LC}}$$

CHAPTER 15

Bus Interface

INTRODUCTION

Many applications require some knowledge of the bus systems located within the personal computer. At times, main boards from personal computers are used as core systems in industrial applications. These systems often require custom interfaces that are attached to one of the buses on the main board. This chapter presents the ISA (industry standard architecture) bus, the PCI (peripheral component interconnect) and PCI Express buses, the USB (universal serial bus), and the AGP (advanced graphics port). Also provided are some simple interfaces to many of these bus systems as design guides.

Although it is likely that they will not be on personal computers of the future, the parallel port and serial communications ports are discussed. These were the first I/O ports on the personal computer and they have stood the test of time, but the universal serial bus seems to have all but replaced their utility.

CHAPTER OBJECTIVES

Upon completion of this chapter, you will be able to:

1. Detail the pin connections and signal bus connections on the parallel and serial ports as well as on ISA, AGP, PCI, and PCI Express buses.
2. Develop simple interfaces that connect to the parallel and serial ports and the ISA and PCI buses.
3. Program interfaces located on boards that connect to the ISA and PCI buses.
4. Describe the operation of the USB and develop some short programs that transfer data.
5. Explain how the AGP increases the efficiency of the graphics subsystem.

15–1   THE ISA BUS

The ISA, or industry standard architecture, bus has been around since the very start of the IBM-compatible personal computer system (circa 1982). In fact, any card from the very first personal computer will plug into and function in any of the modern Pentium 4-based computers provided they have an ISA slot. This is all made possible by the ISA bus interface found in some of these

machines, which is still compatible with the early personal computers. The ISA bus has all but disappeared on the home PC, but is still found in many industrial applications and is presented here for this reason. The main reason it is still used in industrial application is the low cost of the interface and the number of existing interface cards. This will eventually change.

Evolution of the ISA Bus

The ISA bus has changed from its early days. Over the years, the ISA bus has evolved from its original 8-bit standard to the 16-bit standard found in some systems today. The last computer system that contained the ISA bus en masse was the Pentium III. When the Pentium 4 started to appear, the ISA bus started to disappear. Along the way, there was even a 32-bit version called the EISA bus (extended ISA), but that seems to have all but disappeared. What remains today in some personal computers is an ISA slot (connection) on the main board that can accept either an 8-bit ISA card or a 16-bit ISA printed circuit card. The 32-bit printed circuit cards are the PCI bus or, in some older 80486-based machines, the VESA cards. The ISA bus has all but vanished recently in home computers, but it is available as a special order in most main boards. The ISA bus is still found in many industrial applications, but its days now seem limited.

The 8-Bit ISA Bus Output Interface

Figure 15–1 illustrates the 8-bit ISA connector found on the main board of all personal computer systems (again, this may be combined with a 16-bit connector). The ISA bus connector contains the entire demultiplexed address bus (A_{19}–A_0) for the 1M-byte 8088 system, the 8-bit data bus (D_7–D_0), and the four control signals $\overline{MEMR}$, $\overline{MEMW}$, $\overline{IOR}$, and $\overline{IOW}$ for controlling I/O and any memory that might be placed on the printed circuit card. Memory is seldom added to any

FIGURE 15–1 The 8-bit ISA bus.

ISA bus card today because the ISA card only operates at an 8 MHz rate. There might be an EPROM or flash memory used for setup information on some ISA cards, but never any RAM.

Other signals, which are useful for I/O interface, are the **interrupt request lines** IRQ_2–IRQ_7. Note that IRQ_2 is redirected to IRQ_9 on modern systems and is so labeled on the connector in Figure 15–1. The DMA channels 0–3 control signals are also present on the connector. The **DMA request inputs** are labeled DRQ_1–DRQ_3 and the DMA acknowledge outputs are labeled $\overline{DACK0}$–$\overline{DACK3}$. Notice that the DRQ_0 input pin is missing because the early personal computers used it and the $DACK_0$ output as a refresh signal to refresh any DRAM that might be located on the ISA card. Today, this output pin contains a 15.2 μs clock signal that was used for refreshing DRAM. The remaining pins are for power and RESET.

Suppose that a series of four 8-bit latches must be interfaced to the personal computer for 32 bits of parallel data. This is accomplished by purchasing an ISA interface card (part number 4713-1) from a company like Vector Electronics or other companies. In addition to the edge connector for the ISA bus, the card also contains room at the back for interface connectors. A 37-pin subminiature D-type connector can be placed on the back of the card to transfer the 32 bits of data to the external source.

Figure 15–2 shows a simple interface for the ISA bus, which provides 32 bits of parallel TTL data. This example system illustrates some important points about any system interface. First, it is extremely important that the loading to the ISA bus be kept to one low-power (LS) TTL load. In this circuit, a 74LS244 buffer is used to reduce the loading on the data bus. If the


FIGURE 15–2   A 32-bit parallel port interfaced to the 8-bit ISA bus.
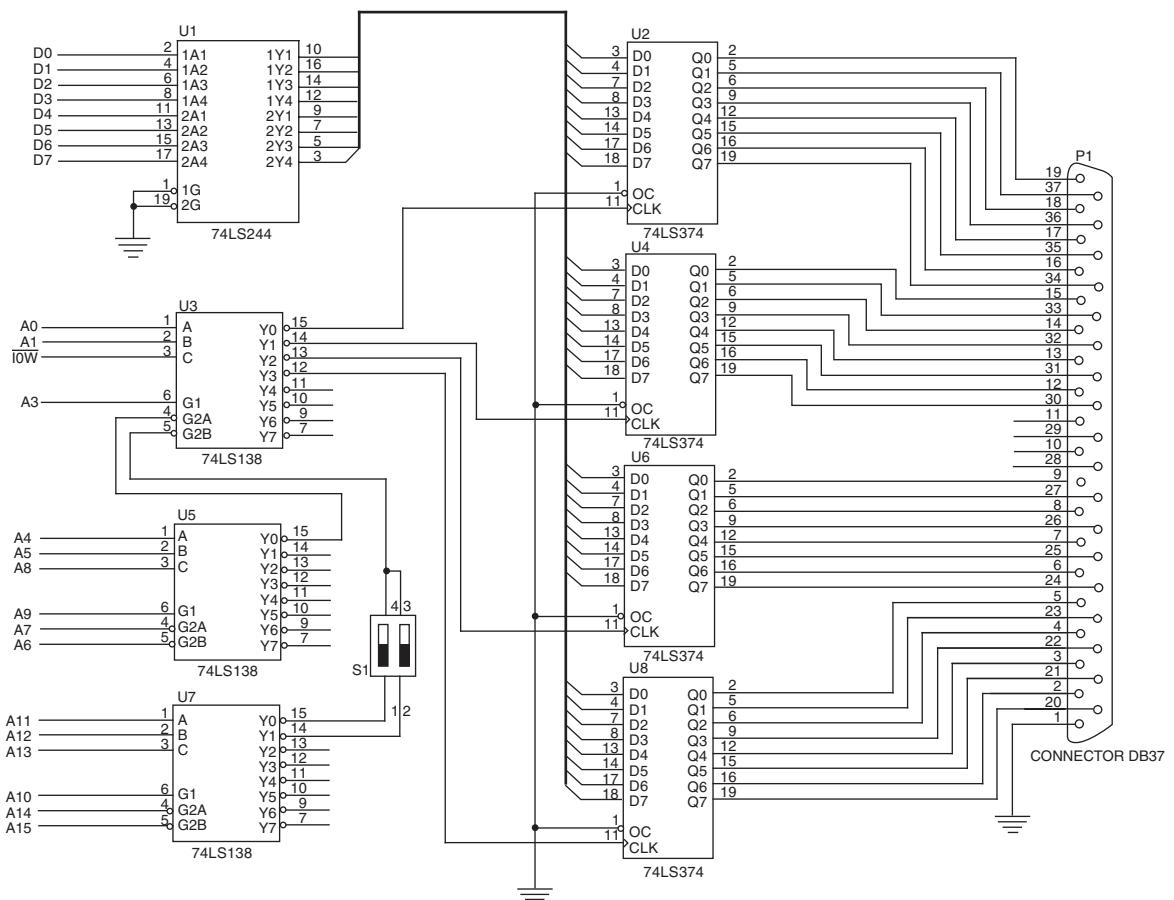
TABLE 15–1   The I/O port assignments of Figure 15–2.

DIP Switch	Latch U_2	Latch U_4	Latch U_6	Latch U_8
1–4 On	0608H or 060CH	0609H or 060DH	060AH or 060EH	060BH or 060FH
2–3 On	0E08H or 0E0CH	0E09H or 0E0DH	0E0AH or 0E0EH	0E0BH or 0E0FH

74LS244 were not there, this system would present the data bus with four unit loads. If all bus cards were to present heavy loads, the system would not operate properly (or perhaps not at all).

Output from the ISA card is provided in this circuit by a 37-pin connector labeled P_1. The output pins from the circuit connect to P_1, and a ground wire is attached. You must provide ground to the outside world, or else the TTL data on the parallel ports are useless. If needed, the output control pins ($\overline{OC}$) on each of the 74LS374 latch chips can also be removed from ground and connected to the four remaining pins on P_1. This allows an external circuit to control the outputs from the latches.

A small DIP switch is placed on two of the outputs of D_7, so the address can be changed if an address conflict occurs with another card. This is unlikely, unless you plan to use two of these cards in the same system. Address connection A_2 is not decoded in this system so it becomes a don't care (x). See Table 15–1 for the addresses of each latch and each position of the S_1. Note that only one of the two switches may be on at a time and that each port has two possible addresses for each switch setting because A_2 is not connected.

In the personal computer, the ISA bus is designed to operate at I/O address 0000H through 03FFH. Depending on the version and manufacturer of the main board, ISA cards may or may not function above these locations. Some newer systems often allow ISA ports at locations above 03FFH, but older systems do not. The ports in this example may need to be changed for some systems. Some older cards only decode I/O addresses 0000H–03FFH and may have address conflicts if the port addresses above 03FFH conflict. The ports are decoded in this example by three 74LS138 decoders. It would be more efficient and cost-effective to decode the ports with a programmable logic device.

Figure 15–3 shows the circuit of Figure 15–2 reworked using a PLD to decode the addresses for the system. Notice that address bits A_{15}–A_4 are decoded by the PLD and the switch is connected to two of the PLD inputs. This change allows four different I/O port addresses for each latch, making the circuit more flexible. Table 15–2 shows the port number selected by switch 1–4 and switch 2–3. Example 15–1 shows the program for the PLD that causes the port assignments of Table 15–2.

EXAMPLE 15–1

```
-- VHDL code for the decoder of Figure 15-3

library ieee;
use ieee.std_logic_1164.all;

entity DECODER_15_3 is

port (
        IOW, A14, A13, A12, A11, A10, A9, A8, A7, A6
            A5, A4, A3, A2, A1, A0, S1, S2: in STD_LOGIC;
        U3, U4, U5, U6: out STD_LOGIC
);

end;

architecture V1 of DECODER_15_3 is
```
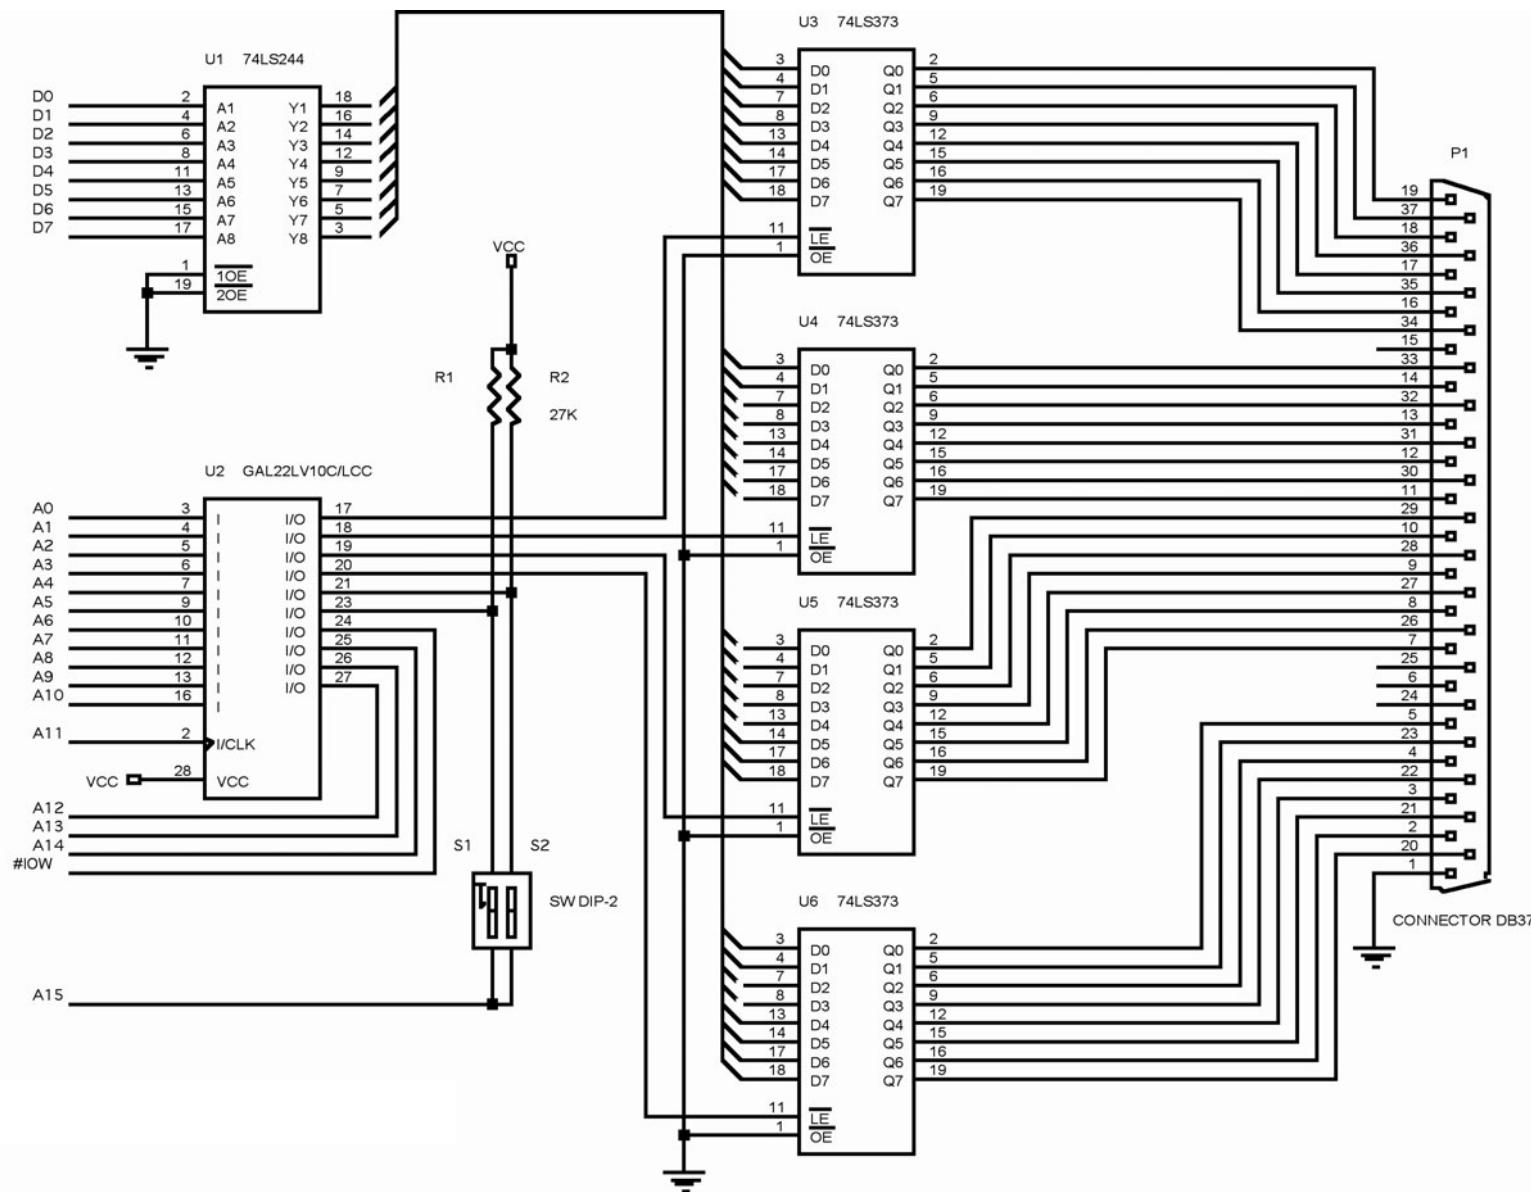
FIGURE 15–3 A 32-bit parallel interface for the ISA bus.

TABLE 15–2 Port assignments of Figure 15–3.

S_2	S_1	U_3	U_4	U_5	U_6
On	On	0300H	0301H	0302H	0303H
On	Off	0304H	0305H	0306H	0307H
Off	On	0308H	0309H	030AH	030BH
Off	Off	030CH	030DH	030EH	030FH

Note: On is a closed switch (0) and off is open (1).

```
begin

        U3 <= IOW or A14 or A13 or A12 or A11 or A10 or not A9 or not A8 or A7
              or A6 or A5 or A4 or A1 or A0 or (S2 or S1 or A3 or A2) and (S2 or
              not S1 or A3 or not A2) and (not S2 or S1 or not A3 or A2) and
              (not S2 or not S1 or not A3 or not A2);
        U4 <= IOW or A14 or A13 or A12 or A11 or A10 or not A9 or not A8 or A7
              or A6 or A5 or A4 or A1 or not A0 or (S2 or S1 or A3 or A2) and
              (S2 or not S1 or A3 or not A2) and (not S2 or S1 or not A3 or A2)
              and (not S2 or not S1 or not A3 or not A2);
        U5 <= IOW or A14 or A13 or A12 or A11 or A10 or not A9 or not A8 or A7
              or A6 or A5 or A4 or not A1 or A0 or (S2 or S1 or A3 or A2) and
              (S2 or not S1 or A3 or not A2) and (not S2 or S1 or not A3 or A2)
              and (not S2 or not S1 or not A3 or not A2);
        U6 <= IOW or A14 or A13 or A12 or A11 or A10 or not A9 or not A8 or A7
              or A6 or A5 or A4 or not A1 or not A0 or (S2 or S1 or A3 or A2)
              and (S2 or not S1 or A3 or not A2) and (not S2 or S1 or not A3 or
              A2) and (not S2 or not S1 or not A3 or not A2);
end V1;
```

Notice in Example 15–1 how the first term (U_3) generates a logic 0 on the output to the decoder only when both switches are in their off positions for I/O port 0300H. It also generates a clock for U_3 for I/O ports 304H, 308H, or 30CH, depending on the switch settings. The second term (U_4) is active for ports 301H, 305H, 309H, or 30DH, depending on the switch settings. Again, refer to Table 15–2 for the complete set of port assignments for various switch settings. Since A_{15} is connected to the bottom of the switches, this circuit will also activate the latches for other I/O locations, because it is not decoded. I/O addresses 830XH will also generate clock signals to the latch because A_{15} is not decoded.

Example 15–2 shows two C++ functions that transfer an integer to the 32-bit port. Either of these functions sends data to the port; the first is more efficient, but the second may be more readable. (Example 15–2(c) shows Example 15–2(b) in disassembled form.) Two parameters are passed to the function: One is the data to be sent to the port, and the other is the base port address. The base address is 0300H, 0304H, 0308H, or 030CH and must match the switch settings of Figure 15–3.

EXAMPLE 15–2(a)

```
void OutPort(int address, int data)
{
        _asm
        {
            mov edx,address
            mov eax,data
            mov ecx,4
OutPort1:
            out dx,al               ;output 8-bits
            shr eax,8               ;get next 8-bit section
```

```
                    inc dx                  ;address next port
                    loop OutPort1           ;repeat 4 times
            }
}
```

EXAMPLE 15–2(b)

```
void OutPrt(int address, int data)
{
        for ( int a = address; a < address + 4; a++ )
        {
                _asm
                {
                        mov edx,a
                        mov eax,data
                        out dx,al
                }
                data >>= 8;              //get next 8-bit section
        }
}
```

EXAMPLE 15–2(c)

```
//Example 15-2(b) disassembled

            for ( int a = address; a < address + 4; a++ )
00413823  mov        eax,dword ptr [address]
00413826  mov        dword ptr [a],eax
00413829  jmp        CSSEDlg::OutPrt+54h (413834h)
0041382B  mov        eax,dword ptr [a]
0041382E  add        eax,1
00413831  mov        dword ptr [a],eax
00413834  mov        eax,dword ptr [address]
00413837  add        eax,4
0041383A  cmp        dword ptr [a],eax
0041383D  jge        CSSEDlg::OutPrt+71h (413851h)
      {
            _asm
            {
                    mov edx,a
0041383F  mov        edx,dword ptr [a]
                    mov eax,data
00413842  mov        eax,dword ptr [data]
                    out dx,al
00413845  out        dx,al
            }
            data >>= 8;              //get next 8-bit section
00413846  mov        eax,dword ptr [data]
00413849  sar        eax,8
0041384C  mov        dword ptr [data],eax
      }
    0041384F  jmp        CSSEDlg::OutPrt+4Bh (41382Bh)
```

The 8-Bit ISA Bus Input Interface

To illustrate the input interface to the ISA bus, a pair of ADC804 analog-to-digital converters are interfaced to the ISA bus in Figure 15–4. The connections to the converters are made through a nine-pin DB_9 connector. The task of decoding the I/O port addresses is more complex, because each converter needs a write pulse to start a conversion, a read pulse to read the digital data once they have been converted from the analog input data, and a pulse to enable the selection of the $\overline{INTR}$ output. Notice that the $\overline{INTR}$ output is connected to data bus bit position D_0. When $\overline{INTR}$ is
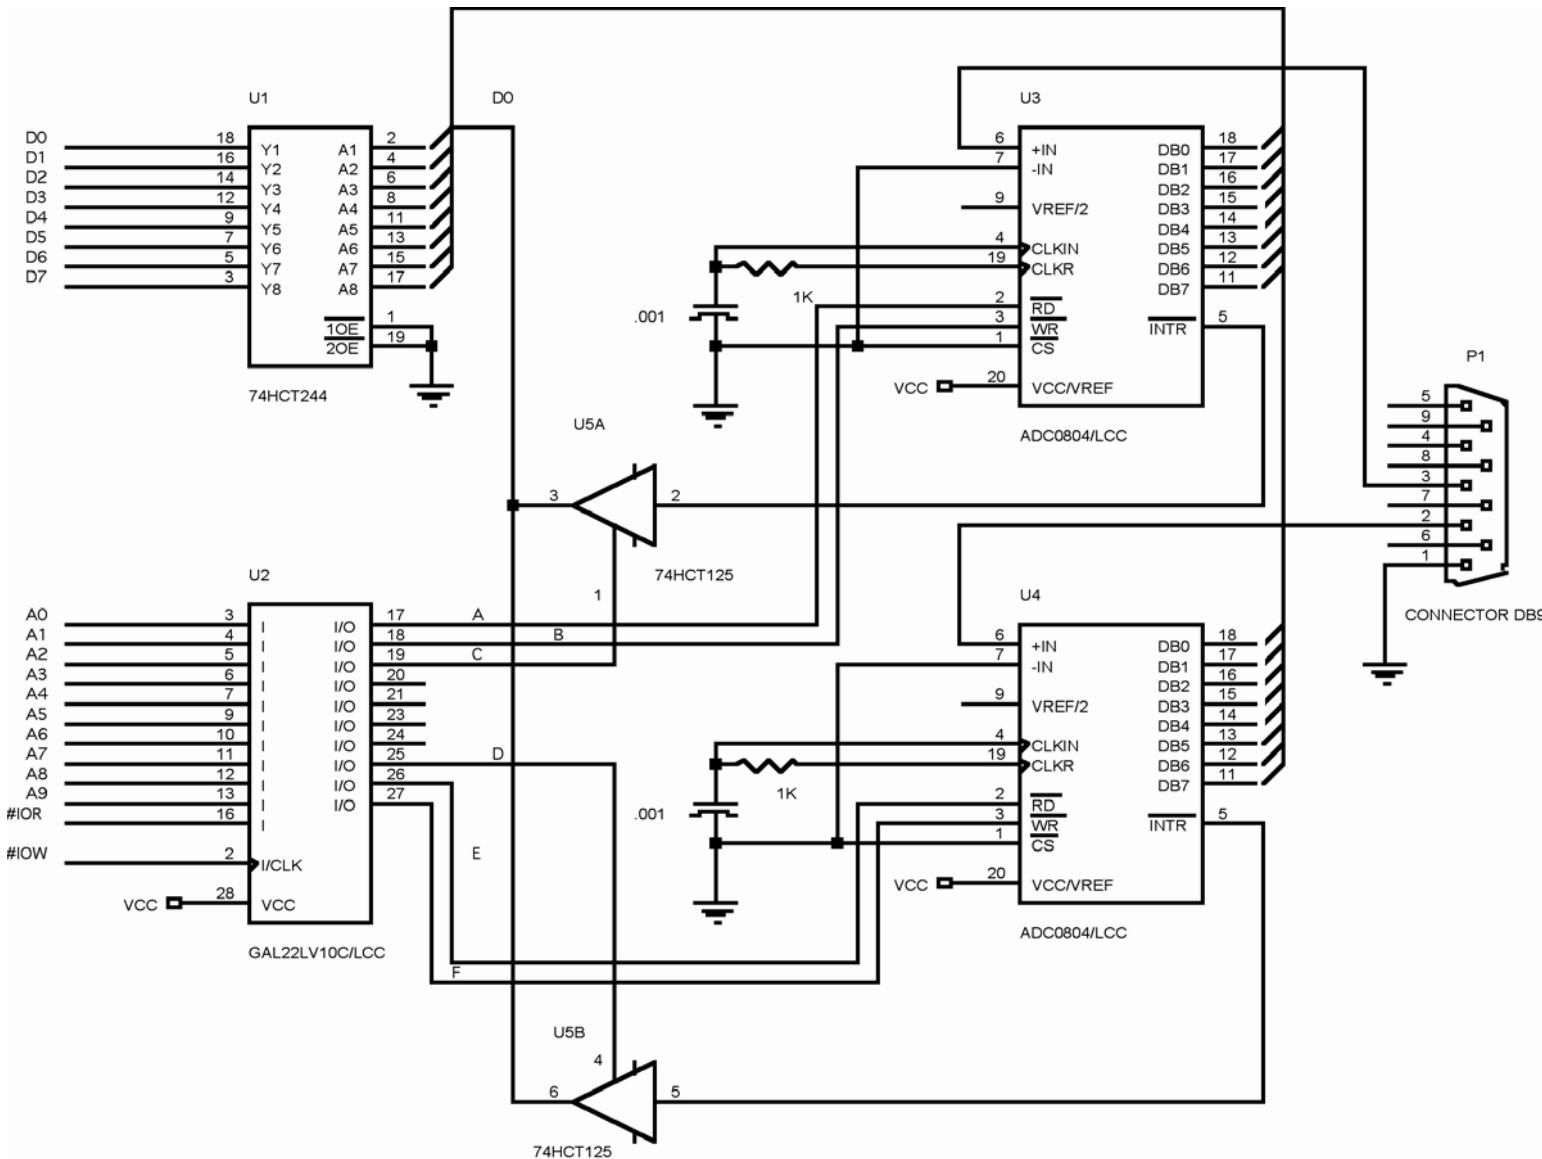
**FIGURE 15–4**   A pair of analog-to-digital converters interfaced to the ISA bus.

TABLE 15–3 Port assignments for Figure 15–4.

Device	Port
Start ADC (U_3)	0300H
Read ADC (U_3)	0300H
Read INTR (U_3)	0301H
Start ADC (U_4)	0302H
Read SDC (U_4)	0302H
Read INTR (U_4)	0303H

input to the microprocessor, the rightmost bit of AL is tested to determine whether the converter is busy.

As before, great care is taken so that the connections to the ISA bus present one unit load to the system. Table 15–3 illustrates the I/O port assignment decoded by the PLD (see Example 15–3 for the program). In this example we assumed that the standard ISA bus is used, which only contains address connection A_0 through A_9.

EXAMPLE 15–3

```
-- VHDL code for the decoder of Figure 15-4

library ieee;
use ieee.std_logic_1164.all;

entity DECODER_15_4 is

port (
      IOW, IOR, A9, A8, A7, A6  A5, A4, A3, A2, A1, A0: in STD_LOGIC;
      A, B, C, D, E, F: out STD_LOGIC
);

end;

architecture V1 of DECODER_15_4 is

begin

      A <= not A9 or not A8 or A7 or A6 or A5 or A4 or A3 or A2 or A1 or A0 or
            IOR;
      B <= not A9 or not A8 or A7 or A6 or A5 or A4 or A3 or A2 or A1 or A0 or
            IOW;
      C <= not A9 or not A8 or A7 or A6 or A5 or A4 or A3 or A2 or A1 or not A0
            or IOR;
      D <= not A9 or not A8 or A7 or A6 or A5 or A4 or A3 or A2 or not A1 or
            not A0 or IOR;
      E <= not A9 or not A8 or A7 or A6 or A5 or A4 or A3 or A2 or not A1 or A0
            or IOR;
      F <= not A9 or not A8 or A7 or A6 or A5 or A4 or A3 or A2 or not A1 or A0
            or IOW;

end V1;
```

Example l5–4 lists a function that can read either ADC U_3 or U_4. The address is generated by passing either a 0 for U_3 or a 1 for U_4 to the address parameter of the function. The function starts the converter by writing to it, and then waits until the $\overline{INTR}$ pin returns to a logic 0, indicating that the conversion is complete before the data are read and returned by the function as a char.

EXAMPLE 15–4
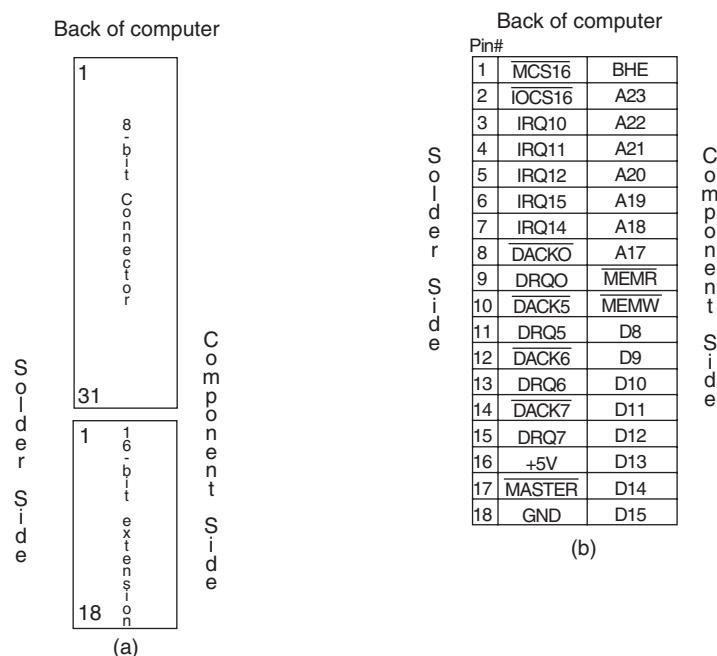
```
char ADC(int address)
{
        char temp = 1;
        if ( address )
              address = 2;
        address += 0x300;
        _asm
        {                          ;start converter
              mov edx,address
              out dx,al
        }
        while ( temp )      //wait if busy
        {
              _asm
              {
                     mov edx,address
                     inc edx
                     in al,dx
                     mov temp,al
                     and al,1
              }
        }
        _asm
        {                          ;get data
              mov edx,address
              in al,dx
              mov temp,al
        }
        return temp;
}
```

The 16-Bit ISA Bus

The only difference between the 8- and 16-bit ISA bus is that an additional connector is attached behind the 8-bit connector. A 16-bit ISA card contains two edge connectors: One plugs into the original 8-bit connector and the other plugs into the new 16-bit connector. Figure 15–5 shows the pin-out

FIGURE 15–5 The 16-bit ISA bus. (a) Both 8- and 16-bit connectors and (b) the pinout of the 16-bit connector.

of the additional connector and its placement in the computer in relation to the 8-bit connector. Unless additional memory is added on the ISA card, the extra address connections A_{23}–A_{20} do not serve any function for I/O operations. The added features that are most often used are the additional interrupt request inputs and the DMA request signals. In some systems, 16-bit I/O uses the additional eight data bus connections (D_8–D_{15}), but more often today the PCI bus is used for peripherals that are wider than 8 bits. About the only recent interfaces found for the ISA bus are a few modems and sound cards.

15–2   THE PERIPHERAL COMPONENT INTERCONNECT (PCI) BUS

The PCI (**peripheral component interconnect**) bus is virtually the only bus found in the newest Pentium 4 systems and just about all the Pentium systems. In all of the newer systems, the ISA bus still exists by special order, but as an interface for older 8-bit and 16-bit interface cards. Many new systems contain only two ISA bus slots or no ISA slots. In time, the ISA bus may disappear, but it is still an important interface for many industrial applications. The PCI bus has replaced the VESA local bus. One reason is that the PCI bus has plug-and-play characteristics and the ability to function with a 64-bit data bus. A PCI interface contains a series of registers, located in a small memory device on the PCI interface, that contain information about the board. This same memory can provide plug-and-play characteristics to the ISA bus or any other bus. The information in these registers allows the computer to automatically configure the PCI card. This feature, called plug-and-play (PnP), is probably the main reason that the PCI bus has become so popular in the most systems.

Figure 15–6 shows the system structure for the PCI bus in a personal computer system. Notice that the microprocessor bus is separate and independent of the PCI bus. The microprocessor
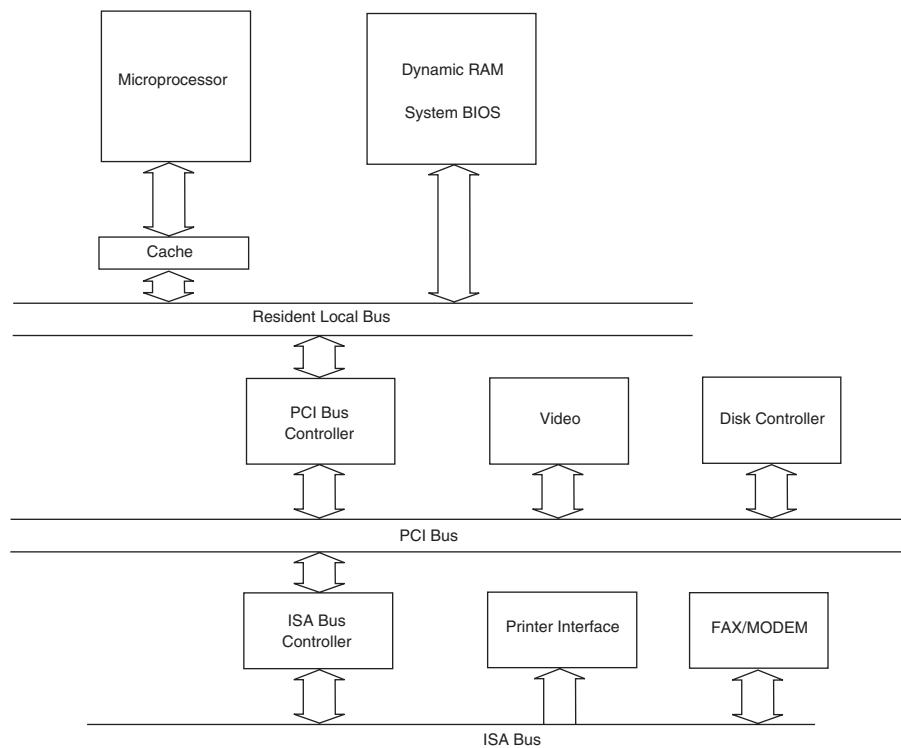

FIGURE 15–6   The system block diagram for the personal computer that contains a PCI bus.

connects to the PCI bus through an integrated circuit called a PCI bridge. This means that virtually any microprocessor can be interfaced to the PCI bus, as long as a PCI controller or bridge is designed for the system. In the future, all computer systems may use the same bus. Even the Apple Macintosh system is switching to the PCI bus. The resident local bus is often called a front side bus.

The PCI Bus Pin-Out

As with the other buses described in this chapter, the PCI bus contains all of the system control signals. Unlike the other buses, the PCI bus functions with either a 32-bit or a 64-bit data bus and a full 32-bit address bus. Another difference is that the address and data buses are multiplexed to reduce the size of the edge connector. These multiplexed pins are labeled AD_0–AD_{63} on the connector. The 32-bit card (which is found in most computers) has only connections 1 through 62, while the 64-bit card has all 94 connections. The 64-bit card can accommodate a 64-bit address if it is required at some point in the future. Figure 15–7 on the next page illustrates the PCI bus pin-out.

  As with the other bus systems, the PCI bus is most often used for interfacing I/O components to the microprocessor. Memory could be interfaced, but it would operate only at a 33 MHz rate with the Pentium, which is half the speed of the 66 MHz resident local bus of the Pentium system. A more recent version of PCI (2.1-compliant) operates at 66 MHz and at 33 MHz for older interface cards. Pentium 4 systems use a 200 MHz system bus speed (although it is often listed as 800 MHz), but there is no planned modification to the PCI bus speed yet.

The PCI Address/Data Connections

The PCI address appears on AD_0–AD_{31} and it is multiplexed with data. In some systems, there is a 64-bit data bus that uses AD_{32}–AD_{63} for data transfer only. In the future, these pins can be used for extending the address to 64 bits. Figure 15–8 illustrates the timing diagram for the PCI bus, which shows the way that the address is multiplexed with data and also the control signals used for multiplexing.

  During the first clocking period, the address of the memory or I/O location appears on the AD connections, and the command to a PCI peripheral appears on the C/$\overline{\text{BE}}$ pins. Table 15–4 illustrates the bus commands found on the PCI bus.

INTA Sequence   During the interrupt acknowledge sequence, an interrupt controller (the controller that caused the interrupt) is addressed and interrogated for the interrupt vector. The byte-sized interrupt vector is returned during a byte read operation.

TABLE 15–4 PC bus commands.

C/$\overline{\text{BE}}$3–C/$\overline{\text{BE}}$0	*Command*
0000	INTA sequence
0001	Special cycle
0010	I/O read cycle
0011	I/O write cycle
0100–1001	Reserved
1010	Configuration read
1011	Configuration write
1100	Memory multiple access
1101	Dual addressing cycle
1110	Line memory access
1111	Memory write with invalidation

FIGURE 15–7 The pin-out of the PCI bus.

Back of computer

Pin #	Solder Side	Component Side
1	−12V	$\overline{TRST}$
2	TCK	+12V
3	GND	TM5
4	TD0	TD1
5	+5V	+5V
6	+5V	$\overline{INTA}$
7	$\overline{INTB}$	$\overline{INTC}$
8	$\overline{INTD}$	+5V
9	$\overline{PRSNT\ 1}$	
10		+VI/O
11	$\overline{PRSNT\ 2}$	
12	KEY	KEY
13	KEY	KEY
14		
15	GND	$\overline{RST}$
16	CLK	VI/O
17	GND	$\overline{VNT}$
18	$\overline{REQ}$	GND
19	+V I0	
20	AD31	AD30
21	AD29	+3.3V
22	GND	AD28
23	AD27	AD26
24	AD25	GND
25	+3.3V	AD24
26	$C/\overline{BE3}$	IDSEL
27	AD23	+3.3V
28	GND	AD22
29	AD21	AD20
30	AD19	GND
31	+3.3V	AD18
32	AD17	AD16
33	$C/\overline{BE2}$	+3.3V
34	GND	$\overline{FRAME}$
35	$\overline{IRDY}$	GND
36	+3.3V	$\overline{TRDY}$
37	$\overline{DEVSEL}$	GND
38	GND	$\overline{STOP}$
39	$\overline{LOCK}$	+3.3V
40	$\overline{PERR}$	SDONE

Pin #	Solder Side	Component Side
41	+3.3V	$\overline{SBO}$
42	$\overline{SERR}$	GND
43	+3.3V	PAR
44	$C/\overline{BE1}$	AD15
45	AD14	+3.3V
46	GND	AD13
47	AD12	AD11
48	AD10	GND
49	GND	AD9
50	KEY	KEY
51	KEY	KEY
52	AD8	$C/\overline{BE0}$
53	AD7	+3.3V
54	+3.3V	AD6
55	AD5	AD4
56	AD3	GND
57	GND	AD2
58	AD1	AD0
59	+V IO	+V IO
60	$\overline{ACK64F}$	$\overline{REQ64}$
61	+5V	+5V
62	+5V	+5V
63		GND
64	GND	$C/\overline{BE7}$
65	$C/\overline{BE6}$	$C/\overline{BE5}$
66	$C/\overline{BE4}$	+V IO
67	GND	PAR64
68	AD63	AD62
69	AD61	GND
70	+V IO	AD60
71	AD59	AD58
72	AD57	GND
73	GND	AD56
74	AD55	AD54
75	AD53	+V IO
76	GND	AD52
77	AD51	AD50
78	AD49	GND
79	+V IO	AD48
80	AD47	AD46
81	AD45	GND
82	GND	AD44
83	AD43	AD42
84	AD41	+VI/O
85	GND	AD40
86	AD39	AD38
87	AD37	GND
88	+VI/O	AD36
89	AD35	AD34
90	AD33	GND
91	GND	AD32
92		
93		GND
94	GND	

Notes: (1) pins 63–94 exist only on the 64-bit PCI card
(2) + VI/O is 3.3V on a 3.3V board and +5V on a 5V board
(3) blank pins are reserved

|	
Special Cycle	The special cycle is used to transfer data to all PCI components. During this cycle, the rightmost 16 bits of the data bus contain a 0000H, indicating a processor shutdown, 0001H for a processor halt, or 0002H for 80X86 specific code or data.
I/O Read Cycle	Data are read from an I/O device using the I/O address that appears on AD_0–AD_{15}. Burst reads are not supported for I/O devices.
I/O Write Cycle	As with I/O read, this cycle accesses an I/O device, but writes data.
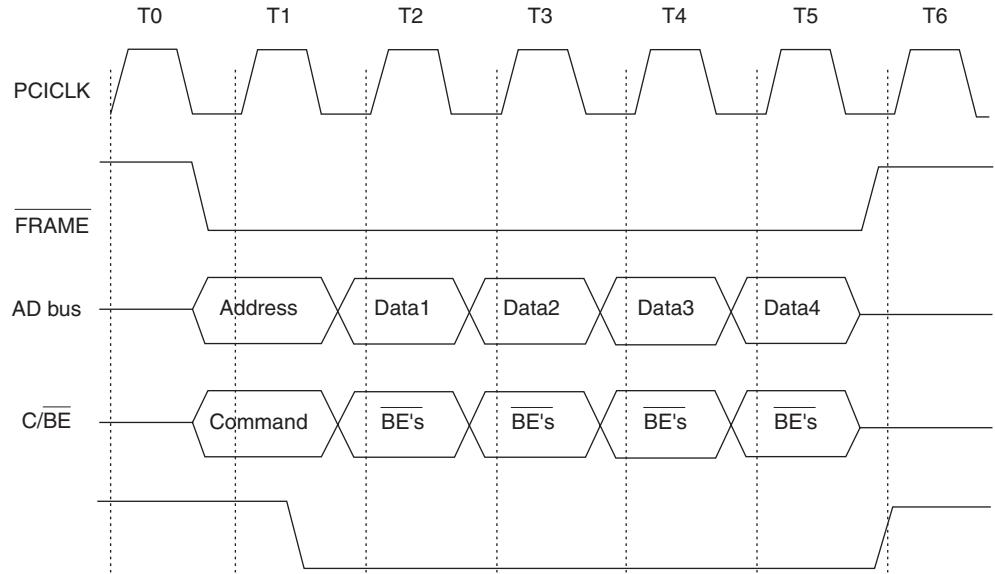
FIGURE 15–8   The basic burst mode timing for the PCI bus system. Note that this transfers either four 32-bit numbers (32-bit PCI) or four 64-bit numbers (64-bit PCI).

Memory Read Cycle	Data are read from a memory device located on the PCI bus.
Memory Write Cycle	As with memory read, data are accessed in a device located on the PCI bus. The location is written.
Configuration Read	Configuration information is read from the PCI device using the configuration read cycle.
Configuration Write	The configuration write allows data to be written to the configuration area in a PCI device. Note that the address is specified by the configuration read.
Memory Multiple Access	This is similar to the memory read access, except that it is usually used to access many data instead of one.
Dual Addressing Cycle	Used for transferring address information to a 64-bit PCI device, which only contains a 32-bit data path.
Line Memory Addressing	Used to read more than two 32-bit numbers from the PCI bus.
Memory Write with Invalidation	This is the same as line memory access, but it is used with a write. This write bypasses the write-back function of the cache.

Configuration Space

The PCI interface contains a 256-byte configuration memory that allows the computer to interrogate the PCI interface. This feature allows the system to automatically configure itself for the PCI plug-board. Microsoft Corporation calls this plug-and-play (PnP). Figure 15–9 illustrates the configuration memory and its contents.

The first 64 bytes of the configuration memory contain the header that holds information about the PCI interface. The first 32-bit doubleword contains the unit ID code and the vendor ID code. The unit ID code is a 16-bit number (D_{31}–D_{16}) that is an FFFFH if the unit is not installed, and a number between 0000H and FFFEH that identifies the unit if it is installed. The class codes identify the class of the PCI interface. The class code is found in bits D_{31}–D_{16} of configuration memory at location 08H. Note that bits D_{15}–D_0 are defined by the manufacturer. The current

**FIGURE 15–9** The contents of the configuration memory on a PCI expansion board.
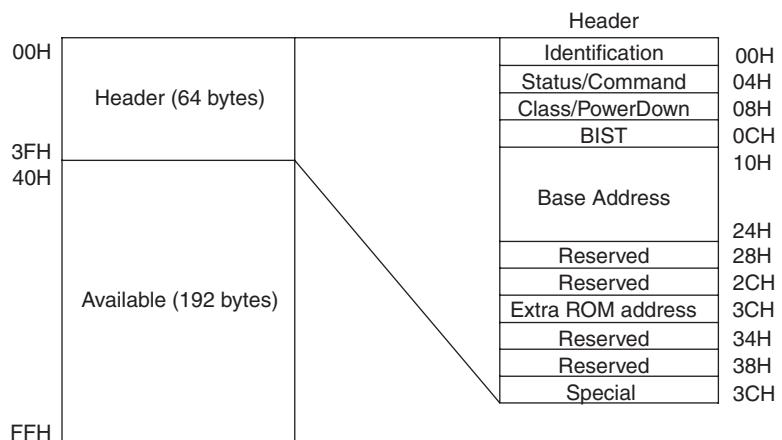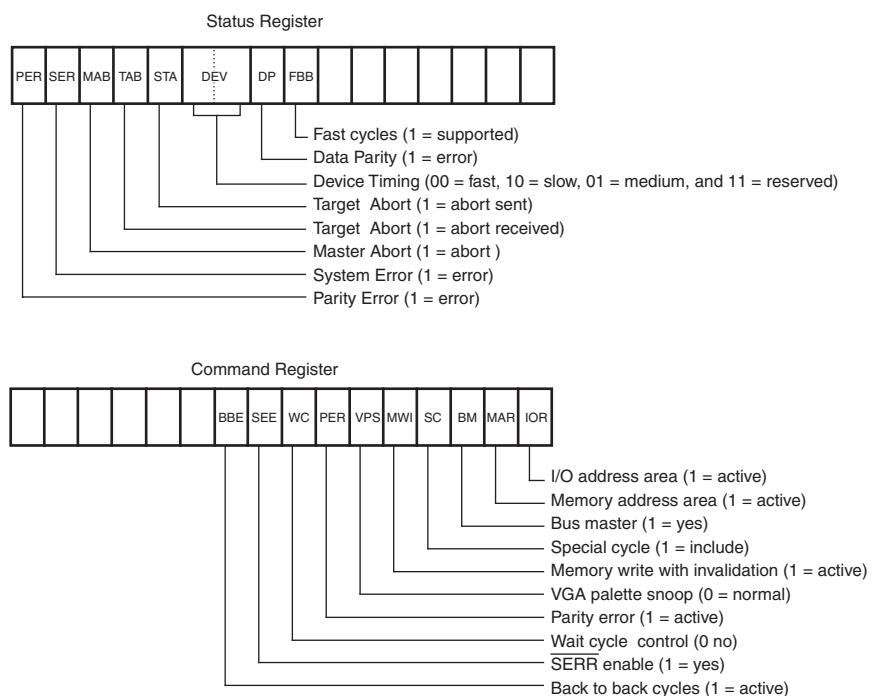
TABLE 15–5 The class codes.

Class Code	Function
0000H	Older non-VGA device (not PnP)
0001H	Older VGA device (not PnP)
0100H	SCSI controller
0101H	IDE controller
0102H	Floppy disk controller
0103H	IPI controller
0180H	Other hard/floppy controller
0200H	Ethernet controller
0201H	Token ring controller
0202H	FDDI
0280H	Other network controller
0300H	VGA controller
0301H	XGA controller
0380H	Other video controller
0400H	Video multimedia
0480H	Other multimedia controller
0500H	RAM controller
0580H	Other memory bridge controller
0600H	Host bridge
0601H	ISA bridge
0602H	EISA bridge
0603H	MCA bridge
0604H	PCI–PCI bridge
0605H	PCMIA bridge
0680H	Other bridge
0700H–FFFEH	Reserved
FFFFH	Not installed

FIGURE 15–10 The con-
tents of the status and control
words in the configuration
memory.

Status Register

Fast cycles (1 = supported)
Data Parity (1 = error)
Device Timing (00 = fast, 10 = slow, 01 = medium, and 11 = reserved)
Target Abort (1 = abort sent)
Target Abort (1 = abort received)
Master Abort (1 = abort)
System Error (1 = error)
Parity Error (1 = error)

Command Register

I/O address area (1 = active)
Memory address area (1 = active)
Bus master (1 = yes)
Special cycle (1 = include)
Memory write with invalidation (1 = active)
VGA palette snoop (0 = normal)
Parity error (1 = active)
Wait cycle control (0 no)
SERR enable (1 = yes)
Back to back cycles (1 = active)

class codes are listed in Table 15–5 and are assigned by the PCI SIG, which is the governing body for the PCI bus interface standard. The vendor ID (D_{15}–D_0) is also allocated by the PCI SIG.

The status word is loaded in bits D_{31}–D_{16} of configuration memory location 04H and the command is at bits D_{15}–D_0 of location 04H. Figure 15–10 illustrates the format of both the status and command registers.

The base address space consists of a base address for the memory, a second for the I/O space, and a third for the expansion ROM. The first two doublewords of the base address space contain either the 32- or 64-bit base address for the memory present on the PCI interface. The next doubleword contains the base address of the I/O space. Note that even though the Intel microprocessors only use a 16-bit I/O address, there is room for expanding the I/O address to 32 bits. This allows systems that use the 680X0 family and PowerPC access to the PCI bus because they do have I/O space that is accessed via a 32-bit address. The 600X0 and PowerPC use memory-mapped I/O, discussed at the beginning of Chapter 11.

BIOS for PCI

Most modem personal computers contain the PCI bus and an extension to the normal system BIOS that supports the PCI bus. These newer systems contain access to the PCI bus at interrupt vector 1AH. Table 15–6 lists the functions currently available through the DOS INT 1AH instruction with AH = 0B1H for the PCI bus.

Example 15–5 shows how the BIOS is used to determine whether the PCI bus extension available. Once the presence of the BIOS is established, the contents of the configuration memory can be read using the BIOS functions. Note that the BIOS does not support data transfers between the computer and the PCI interface. Data transfers are handled by drivers that are provided with the interface. These drivers control the flow of data between the microprocessor and the component found on the PCI interface.

TABLE 15–6   BIOS INT1AH functions for the PCI bus.

01H	BIOS Available?
Entry	AH = 0B1H AL = 01H
Exit	AH = 00H if PCI BIOS extension is available BX = version number EDX = ASCII string 'PCI' CARRY = 1 if no PCI extension present

02H	PCI Unit Search
Entry	AH = 0B1H AL = 02H CX = Unit DX = Manufacturer SI = index
Exit	AH = result code (see notes) BX = bus and unit number Carry = 1 for error
Notes	The result codes are: 00H = successful search 81H = function not supported 83H = invalid manufacturer ID code 86H = unit not found 87H = invalid register number

03H	PCI Class Code Search
Entry	AH = 0B1H AL = 03H ECX = class code SI = index
Exit	AH = result code (see notes for function 02H) BX = bus and unit number Carry = 1 for an error

06H	Start Special Cycle
Entry	AH = 0B1H AL = 06H BX = bus and unit number EDX = data
Exit	AH = result code (see notes for function 02H) Carry = 1 for error
Notes	The value passed in EDX is sent to the PCI bus during the address phase.

08H	Configuration Byte-Sized Read
Entry	AH = 0B1H AL = 08H BX = bus and unit number DI = register number
Exit	AH = result code (see notes for function 02H) CL = data from configuration register Carry = 1 for error

TABLE 15–6 (continued).

09H	Configuration Word-Sized Read
Entry	AH = 0B1H AL = 08H BX = bus and unit number DI = register number
Exit	AH = result code (see notes for function 02H) CX = data from configuration register Carry = 1 for error

0AH	Configuration Doubleword-Sized Read
Entry	AH = 0B1H AL = 08H BX = bus and unit number DI = register number
Exit	AH = result code (see notes for function 02H) ECX = data from configuration register Carry = 1 for error

0BH	Configuration Byte-Sized Write
Entry	AH = 0B1H AL = 08H BX = bus and unit number CL = data to be written to configuration register DI = register number
Exit	AH = result code (see notes for function 02H) Carry = 1 for error

0CH	Configuration Word-Sized Write
Entry	AH = 0B1H AL = 08H BX = bus and unit number CX = data to be written to configuration register DI = register number
Exit	AH = result code (see notes for function 02H) Carry = 1 for error

0DH	Configuration Doubleword-Sized Write
Entry	AH = 0B1H AL = 08H BX = bus and unit number ECX = data to be written to configuration register DI = register number
Exit	AH = result code (see notes for function 02H) Carry = 1 for error

EXAMPLE 15–5

```
;DOS program that determines whether PCI exists

.MODEL SMALL

.DATA
      MES1   DB      "PCI BUS IS PRESENT$"
      MES2   DB      "PCI BUS IS NOT FOUND$"

.CODE
.STARTUP
      MOV AH,0B1H               ;access PCI BIOS
      MOV AL,1
      INT 1AH
      MOV DX,OFFSET MES2
      .IF CARRY?        ;if PCI is present
            MOV  DX,OFFSET MES1
      .ENDIF
      MOV  AH,9         ;display MES1 or MES2
      INT  21H
      .EXIT
END
```

PCI Interface

The PCI interface is complex, and normally an integrated PCI bus controller is used for interfacing to the PCI bus. It requires memory (EPROM) to store vendor information and other information, as explained earlier in this section of the chapter. The basic structure of the PCI interface is illustrated in Figure 15–11. The contents of this block diagram illustrate the required components for a functioning PCI interface; it does not illustrate the interface itself. The Registers, Parity Block, Initiator, Target, and Vendor ID EPROM are required components of any PCI interface. If a PCI interface is constructed, a PCI controller is often used because of the complexity of this interface. The PCI controller provides the structures shown in Figure 15–11.

PCI Express Bus

The PCI Express transfers data in serial at the rate of 2.5 GHz to legacy PCI applications, increasing the data link speed to 250 MBps to 8 GBps for PCI Express interfaces. The standard PCI bus delivers data at a speed of about 133 MBps, in comparison. The big improvement is on the motherboard, where the interconnections are in serial and at 2.5 GHz. Each serial connection on the PCI Express bus is called a **lane**. The slots on the main board are single lane slots with a total transfer speed of 1 GBps. The PCI Express video card connector currently has 16 lanes with a

**FIGURE 15–11** The block diagram of the PCI interface.
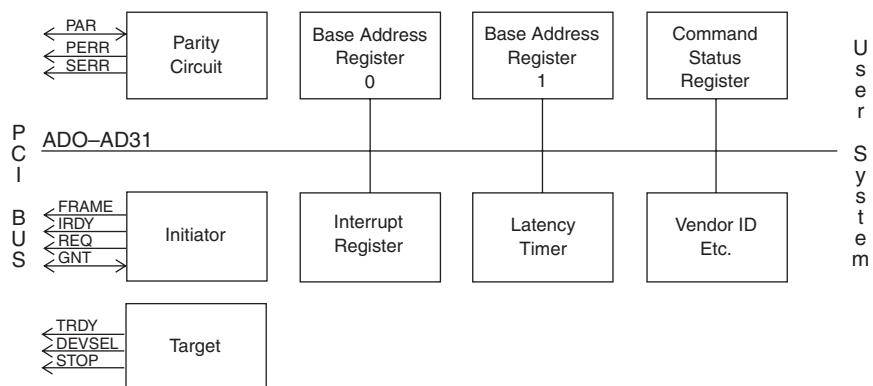
**TABLE 15–7**   The PCI Express single lane pin-out (PCI X1).
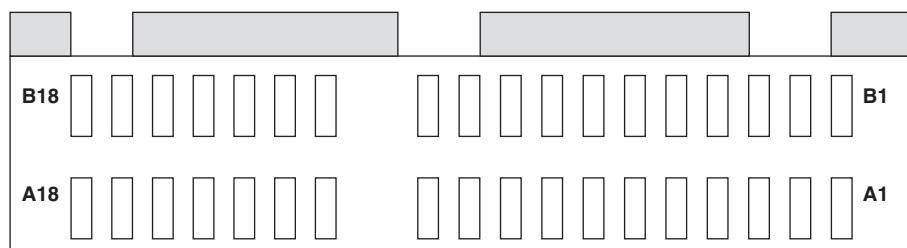
Pin #	Side A Name	Side A Description	Side B Name	Side B Description
1	PRSNT1	Present	+12V	+12 Volt power
2	+12V	+12 Volt power	+12V	+12 Volt power
3	+12V	+12 Volt power	Reserved	Not used
4	GND	Ground	GND	Ground
5	JTAG2	TCK	SMCLK	SMBus clock
6	JTAG3	TDI	SMDAT	SMBus data
7	JTAG4	TDO	GND	Ground
8	JTAG5	TMS	+3.3V	+3.3 Volt power
9	+3.3V	+3.3 Volt power	JTAG1	+TRST#
10	+3.3V	+3.3 Volt power	+3.3V	+3.3 Volt power
11	PWRGD	Power good	WAKE#	Link reactivation
12	GND	Ground	Reserved	Not used
13	REFCLK+	Reference clock	GND	Ground
14	REFCLK−	Reference clock	HSOp(0)	Lane 0 output data+
15	GND	Ground	HSOm(0)	Lane 0 output data−
16	HSIp(0)	Lane 0 input data+	GND	Ground
17	HSIm(0)	Lane 0 input data−	PRSNT2	Present
18	GND	Ground	GND	Ground

transfer speed of 4 GBps. The standard allows up to 32 lanes, but at present the widest slot is the 16 lanes on the video card. Most current main boards contain four single lane slots for peripherals and one 16 lane slot for the video card. A few newer main boards contain two 16 lane slots. In the future the standard PCI slots will all be replaced with the lower cost PCI Express connectors.

The PCI Express 2 bus was released in late 2007 with a transfer speed that is twice that of the PCI Express bus. This means that the speed per lane increased from 250 MBps to 500 MBps.

This new version of the PCI bus is replacing most current video cards on the AGP port with a yet higher speed version of the PCI Express bus. This technology (serial) allows main board manufacturers to use less space on the main board for interconnection and thus reduce the cost of manufacturing a main board. The connectors are smaller, which also reduces connector cost. The software used with the PCI Express bus remains the same as that used with the PCI bus so new programs are not needed to develop drivers for the PCI Express bus.

The PCI Express pin-out for the most commonly interfaced connector, the single lane connector, appears in Table 15–7. The connector is a 36-pin connector as illustrated in Figure 15–12. Signaling on the PCI Express bus uses 3.3 V with differential signals that are 180 degrees

FIGURE 15–12   The single lane PCI Express connector.

out of phase. The lane is constructed from a pair of data pipes, one for input data and one for output data.

15–3   THE PARALLEL PRINTER INTERFACE (LPT)

The parallel printer interface (LPT) is located on the rear of the personal computer, and as long as it is a part of the PC, it can be used as an interface to the PC. LPT stands for line printer. The printer interface gives the user access to eight lines that can be programmed to receive or send parallel data.

Port Details

The parallel port (LPT$_1$) is normally at I/O port addresses 378H, 379H, and 37AH from DOS or using a driver in Windows. The secondary (LPT$_2$) port, if present, is located at I/O port addresses 278H, 279H, and 27AH. The following information applies to both ports, but LPT$_1$ port addresses are used throughout.

The Centronics interface implemented by the parallel port uses two connectors, a 25-pin D-type on the back of the PC and a 36-pin Centronics on the back of the printer. The pin-outs of these connectors are listed in Table 15–8, and the connectors are shown in Figure 15–13.

The parallel port can work as both a receiver and a transmitter at its data pins (D$_0$–D$_7$). This allows devices other than printers, such as CD-ROMs, to be connected to and used by the PC through the parallel port. Anything that can receive and/or send data through an 8-bit interface can and often does connect to the parallel port (LPT$_1$) of a PC.

Figure 15–14 illustrates the contents of the data port (378H), the status register (379H), and an additional status port (37AH). Some of the status bits are true when they are a logic zero.

TABLE 15–8   The pin-outs of the parallel port.

Signal	Description	25-pin	36-pin
#STR	Strobe to printer	1	1
D$_0$	Data bit 0	2	2
D$_1$	Data bit 1	3	3
D$_2$	Data bit 2	4	4
D$_3$	Data bit 3	5	5
D$_4$	Data bit 4	6	6
D$_5$	Data bit 5	7	7
D$_6$	Data bit 6	8	8
D$_7$	Data bit 7	9	9
#ACK	Acknowledge from printer	10	10
BUSY	Busy from printer	11	11
PAPER	Out of paper	12	12
ONLINE	Printer is online	13	13
#ALF	Low if printer issues LF after CR	14	14
#ERROR	Printer error	15	32
#RESET	Resets the printer	16	31
#SEL	Selects the printer	17	36
+5V	5V from printer	—	18
Protective Ground	Earth ground	—	17
Signal Ground	Signal Ground	All other pins	All other pins

Note: # indicates an active low signal.
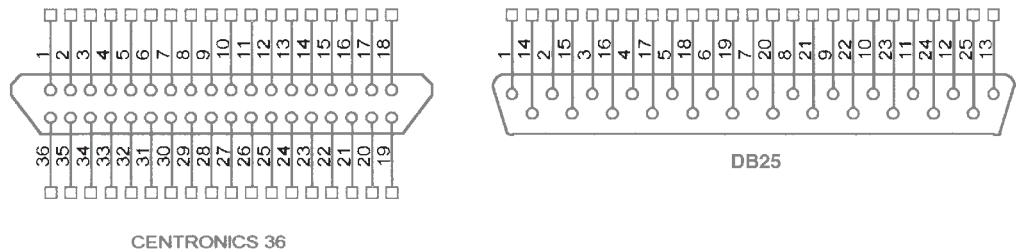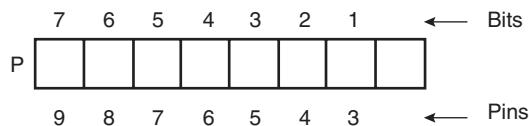
**FIGURE 15–13** The connectors used for the parallel port.
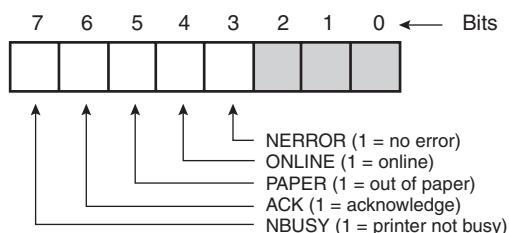
FIGURE 15–14 Ports 378H, 379H, and 37AH as used by the parallel port.

Port 378H

The data port that connects to bits D_0–D_7 (pins 2–9)
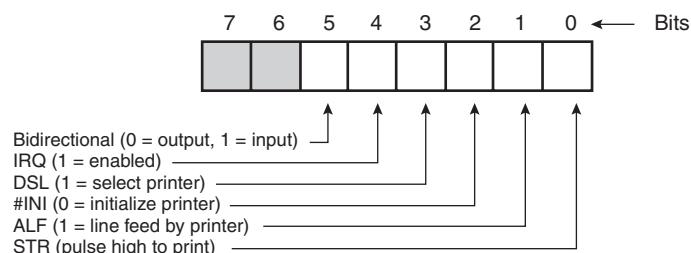
Port 379H

This is a read-only port that returns the information from the printer through signals such as BUSY, #ERROR, and so forth. (Careful! Some of the bits are inverted.)

NERROR (1 = no error)
ONLINE (1 = online)
PAPER (1 = out of paper)
ACK (1 = acknowledge)
NBUSY (1 = printer not busy)

Port 37AH

Bidirectional (0 = output, 1 = input)
IRQ (1 = enabled)
DSL (1 = select printer)
#INI (0 = initialize printer)
ALF (1 = line feed by printer)
STR (pulse high to print)

613

Using the Parallel Port Without ECP Support

For most systems since the PS/2 was released by IBM, you can basically follow the information presented in Figure 15–14 to use the parallel port without ECP. To read the port, it must first be initialized by sending 20H to register 37AH as illustrated in Example 15–6. As indicated in Figure 15–14, this sets the bidirectional that selects input operation for the parallel port. If the bit is cleared, output operation is selected.

EXAMPLE 15–6

```
MOV  AL,20H
MOV  DX,37AH
OUT  DX,AL
```

Once the parallel port is programmed as an input, it is read as depicted in Example 15–7. Once the parallel port is programmed to function as an input port, reading is accomplished by accessing the data port at address 378H.

EXAMPLE 15–7

```
MOV  DX,378H
IN   AL,DX
```

To write data to the parallel port, reprogram the command register at address 37A by writing 00H to program the bidirectional bit with a zero. Once the bidirectional bit is programmed, data are sent to the parallel port through the data port at address 378H. Example 15–8 illustrates how data are sent to the parallel port.

EXAMPLE 15–8

```
MOV  DX,378H
MOV  AL,WRITE_DATA
OUT  DX,AL
```

On older (80286-based) machines the bidirectional bit is missing from the interface. In order to read information from the parallel port, write 0FFH to the port (378H), so that it can be read. These older systems do not have a register at location 37AH.

Accessing the printer port from Windows is difficult because a driver must be written to do so if Windows 2000 or Windows XP is in use. In Windows 98 or Windows ME, access to the port is accomplished as explained in this section.

There is a way to access the parallel port through Windows 2000 and Windows XP without writing a driver. A driver called UserPort (readily available on the Internet) opens up the protected I/O ports in Windows and allows direct access to the parallel port through assembly blocks in Visual C++ using port 378H. It also allows access to any I/O ports between 0000H and 03FFH. Another useful tool is available for a 30-day trial at www.jungo.com. The Jungo tool is a driver development tool, with many example drivers for most subsystems.

15–4   THE SERIAL COM PORTS

The serial communications ports are COM_1–COM_8 in older systems and may contain any number of ports in modern systems, but most computers only have COM_1 and COM_2 installed. Some have a single communication port (COM_1). These ports are controlled and accessed in the DOS environment as described in Chapter 11 with the 16550 serial interface component and will not be discussed again. Instead, we will discuss the Windows API functions for operating the COM ports for the 16550 communications interface. USB devices are often interfaced using the HID (**human interface device**) as a COM port. This allows standard serial software to access USB devices.

TABLE 15–9   Allowable
baud rates for the COM ports.

Keyword	Speed in Bits per Second
CBR_110	110
CBR_300	300
CBR_600	600
CBR_1200	1200
CBR_2400	2400
CBR_4800	4800
CBR_9600	9600
CBR_14400	14400
CBR_19200	19200
CBR_38400	38400
CBR_56000	56000
CBR_57600	57600
CBR_115200	115200
CBR_128000	128000
CBR_256000	256000

Communication Control

The serial ports are accessed through any version of Windows and Visual C++ by using a few system
application interface (API) functions. An example of a short C++ function that accesses the serial
ports is listed in Example 15–9 for Visual Studio.net 2003. The function is called WriteComPort and
it contains two parameters. The first parameter is the port, as in COM_1, COM_2, and so on, and the
second parameter is the character to be sent through the port. A return true indicates that the charac-
ter was sent and a return false indicates that a problem exists. To use the function to send the letter A
through the COM_1 port call it with a WriteComPort ("COM_1", "A"). This function is written to send
only a single byte through the serial COM port, but it could be modified to send strings. To send 00H
(no other number can be sent this way) through COM_2 use WriteComPort ("COM_2", 0x00). Notice
that the COM port is set to 9600 baud, but this is easily changed by changing the CBR_9600 to
another acceptable value. See Table 15–9 for the allowed baud rates.

EXAMPLE 15–9

```
bool WriteComPort(CString PortSpecifier, CString data)
{
     DCB dcb;
     DWORD byteswritten;

     HANDLE hPort = CreateFile(PortSpecifier,
     GENERIC_WRITE,
     0,
     NULL,
     OPEN_EXISTING,
     0,
     NULL);

     if (!GetCommState(hPort,&dcb)){
          return false;
     }

     dcb.BaudRate = CBR_9600;        //9600 baud
     dcb.ByteSize = 8;               //8 data bits
     dcb.Parity = NOPARITY;          //no parity
     dcb.StopBits = ONESTOPBIT;      //1 stop

     if (!SetCommState(hPort,&dcb))
     return false;
```

```
      bool retVal = WriteFile(hPort,data,1,&byteswritten,NULL);
      CloseHandle(hPort);                  //close the handle
      return retVal;}
```

The CreateFile structure creates a handle to the COM ports that can be used to write data to the port. After getting and changing the state of the port to meet the baud rate requirements, the WriteFile function sends data to the port. The parameters used with the WriteFile function are the file handle (hPort), the data to be written as a string, the number of bytes to write (1 in this example), and a place to store the number of bytes actually written to the port.

Receiving data through the COM port is a little more challenging because errors occur more frequently than with transmission. There are also many types of errors that can be detected that often should be reported to the user. Example 15–10 illustrates a C++ function that is used to read a character from the serial port called ReadByte. The ReadByte function returns either the character read from the port or an error code of 0×100 if the port could not be opened, or 0×101 if the receiver detected an error. If data are not received, this function will hang because no timeouts were set.

EXAMPLE 15–10

```
int ReadByte(CString PortSpecifier)
{
      DCB dcb;
      int retVal;
      BYTE Byte;
      DWORD dwBytesTransferred;
      DWORD dwCommModemStatus;

      HANDLE hPort = CreateFile(PortSpecifier,
      GENERIC_READ,
      0,
      NULL,
      OPEN_EXISTING,
      0,
      NULL);

      if (!GetCommState(hPort,&dcb))
            return 0x100;

      dcb.BaudRate = CBR_9600;             //9600 baud
      dcb.ByteSize = 8;                    //8 data bits
      dcb.Parity = NOPARITY;               //no parity
      dcb.StopBits = ONESTOPBIT;           //1 stop

      if (!SetCommState(hPort,&dcb))
            return 0x100;

      SetCommMask (hPort, EV_RXCHAR | EV_ERR);      //receive character event
      WaitCommEvent (hPort, &dwCommModemStatus, 0);  //wait for character

      if (dwCommModemStatus & EV_RXCHAR)
            ReadFile (hPort, &Byte, 1, &dwBytesTransferred, 0); //read 1
      else if (dwCommModemStatus & EV_ERR)
            retVal = 0x101;
      retVal = Byte;
      CloseHandle(hPort);
      return retVal;
}
```

If Visual Studio Express is in use, the toolbox contains the serial port control that allows access to any COM port. For some reason, this was available in Visual Studio 5, it then vanished in Visual Studio 5 and Visual Studio.net and was re-added to the 2005 Express edition. Many USB devices appear as COM ports and are accessed through the serial port control, as well as classic COM ports. The HID USB device is the main reason that Microsoft added the serial port control to Visual Studio.

Once the serial port control is added to a program, it is typically set up for communication in its properties and then an event handler is used when data are received. Sending data occurs as illustrated in the function listed in Example 15–11.

EXAMPLE 15–11

```
private: System::Void SendPort(String^ portDataString)
{
        serialPort1->WriteLine(portDataString);
}
```

To receive data, install the handler for data received. Each time that information is received on the serial port, the data received event is called where the information is processed. Example 15–12 shows the data received function. What does not appear here is that the port must be open to send or receive information using the Open function in the serial port class.

EXAMPLE 15–12

```
private: System::Void serialPort1_DataReceived(System::Object^ sender,
   ystem::IO::Ports::SerialDataReceivedEventArgs^ e)'
{
        String^ receivedString = serialPort1->ReadLine();

        // process the line read from the port
}
```

15–5

THE UNIVERSAL SERIAL BUS (USB)

The universal serial bus (USB) has solved a problem with the personal computer system. The current PCI sound cards use the internal PC power supply, which generates a tremendous amount of noise. Because the USB allows the sound card to have its own power supply, the noise associated with the PC power supply can be eliminated, allowing for high-fidelity sound without 60 Hz hum. Other benefits are ease of user connection and access to up to 127 different connections through a four-connection serial cable. This interface is ideal for keyboards, sound cards, simple video-retrieval devices, and modems. Data transfer speeds are 480 Mbps for full-speed USB 2.0 operation, 11 Mbps for USB 1.1 compliant transfers, and 1.5 Mbps for slow-speed operation.

Cable lengths are limited to five meters maximum for the full-speed interface and three meters maximum for the low-speed interface. The maximum power available through these cables is rated at 100 mA and maximum current at 5.0 V. If the amount of current exceeds 100 mA, Windows will display a yellow exclamation point next to the device, indicating an overload condition.

The Connector

Figure 15–15 illustrates the pin-out of the USB connector. There are two types of connectors specified and both are in use. In either case, there are four pins on each connector, which contain the signals indicated in Table 15–10. As mentioned, the +5.0 V and ground signals can be used to power devices connected to the bus as long as the amount of current does not exceed 100 mA per device. The data signals are biphase signals. When +data are at 5.0 V, −data are at zero volts and vice versa.

USB Data

The data signals are biphase signals that are generated using a circuit such as the one illustrated in Figure 15–16. The line receiver is also illustrated in Figure 15–16. Placed on the transmission pair is a noise-suppression circuit that is available from Texas Instruments (SN75240). Once the

FIGURE 15–15   The front view of the two common types of USB connectors.

Front View

Front View
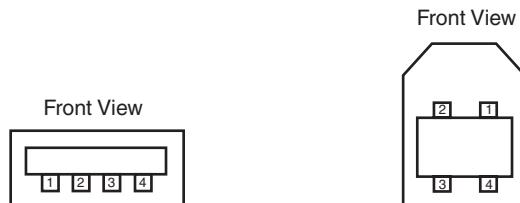
TABLE 15–10   USB pin configuration.

Pin Number	Signal
1	+ 5.0 V
2	− Data
3	+ Data
4	Ground

transceiver is in place, interfacing to the USB is complete. The 75773 integrated circuit from Texas Instruments functions as both the differential line driver and receiver for this schematic.

The next phase is learning how the signals interact on the USB. These signals allow data to be sent and received from the host computer system. The USB uses NRZI (non-return to zero, inverted) data encoding for transmitting packets. This encoding method does not change the signal level for the transmission of a logic 1, but the signal level is inverted for each change to a logic 0. Figure 15–17 illustrates a digital data stream and the USB signal produced using this encoding method.

The actual data transmitted includes sync bits using a method called *bit stuffing*. If a logic 1 is transmitted for more than 6 bits in a row, the bit stuffing technique adds an extra bit (logic 0) after six continuous 1s in a row. Because this lengthens the data stream, it is called bit stuffing. Figure 15–18 shows a bit-stuffed serial data stream and the algorithm used to create it from raw digital serial data. Bit stuffing ensures that the receiver can maintain synchronization for long strings of 1s. Data are always transmitted beginning with the least-significant bit first, followed by subsequent bits.

USB Commands

Now that the USB data format is understood, we will discuss the commands used to transfer data and select the receptor. To begin communications, the sync byte (80H) is transmitted first, followed by the packet identification byte (PID). The PID contains 8 bits, but only the rightmost 4 bits contain the type of packet that follows, if any. The leftmost 4 bits of the PID are the ones complementing the rightmost 4 bits. For example, if a command of 1000 is sent, the actual byte sent for the PID is 0111 1000. Table 15–11 shows the available 4-bit PIDs and their 8-bit codes. Notice that PIDs are used as token indicators, as data indicators, and for handshaking.

**FIGURE 15–16**   The interface to the USB using a pair of CMOS buffers.
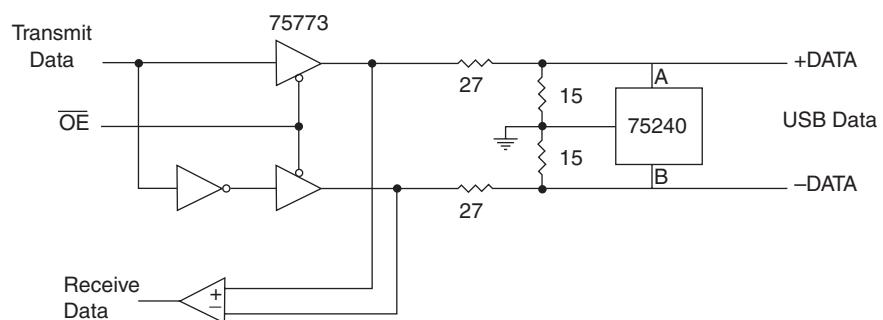
**FIGURE 15–17**   NRZI
encoding used with the USB.
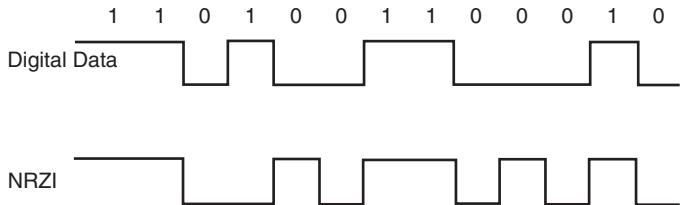
Figure 15–19 lists the formats of the data, token, handshaking, and start-of-frame packets found on the USB. In the token packet, the ADDR (address field) contains the 7-bit address of the USB device. As mentioned earlier, there are up to 127 devices present on the USB at a time. The ENDP (endpoint) is a 4-bit number used by the USB. Endpoint 0000 is used for initialization; other endpoint numbers are unique to each USB device.

There are two types of CRC (**cyclic redundancy checks**) used on the USB: One is a 5-bit CRC and the other (used for data packets) is a 16-bit CRC. The 5-bit CRC is generated with the $X_5 + X_2 + 1$ polynomial; the 16-bit CRC is generated with the $X_{16} + X_{15} + X_2 + 1$ polynomial. When constructing circuitry to generate or detect the CRC, the plus signs represent Exclusive-OR circuits. The CRC circuit or program is a serial checking mechanism. When using the 5-bit CRC, a residual of 01100 is received for no error in all five bits of the CRC and the data bits. With the 16-bit CRC, the residual is 1000000000001101 for no error.

The USB uses the ACK and NAK tokens to coordinate the transfer of data packets between the host system and the USB device. Once a data packet is transferred from the host to the USB device, the USB device either transmits an ACK (acknowledge) or a NAK (not acknowledge)
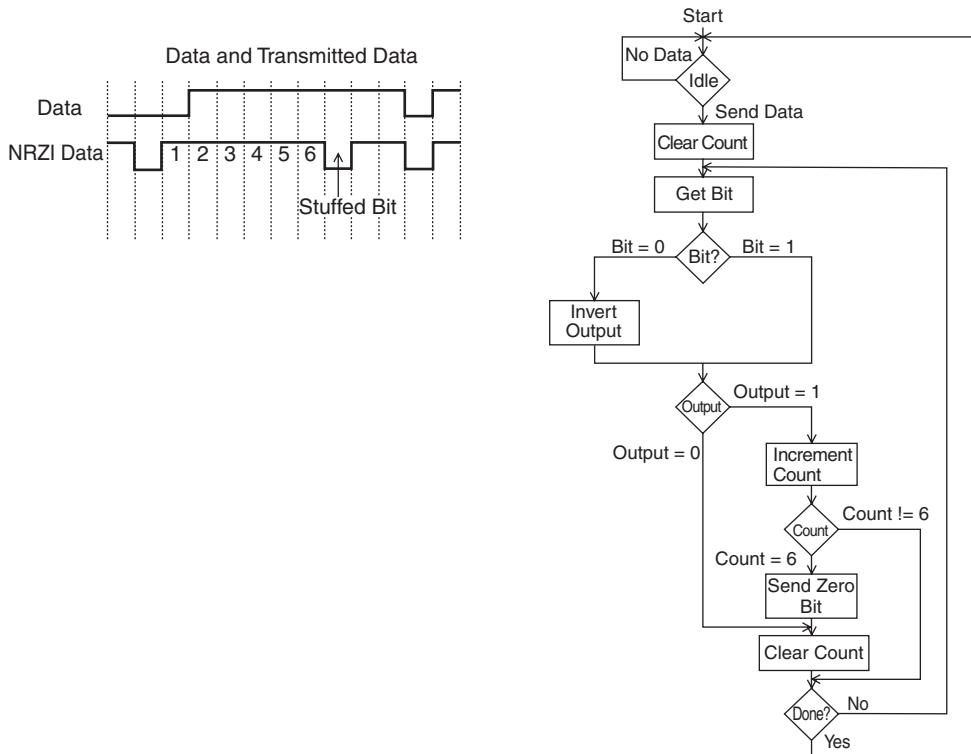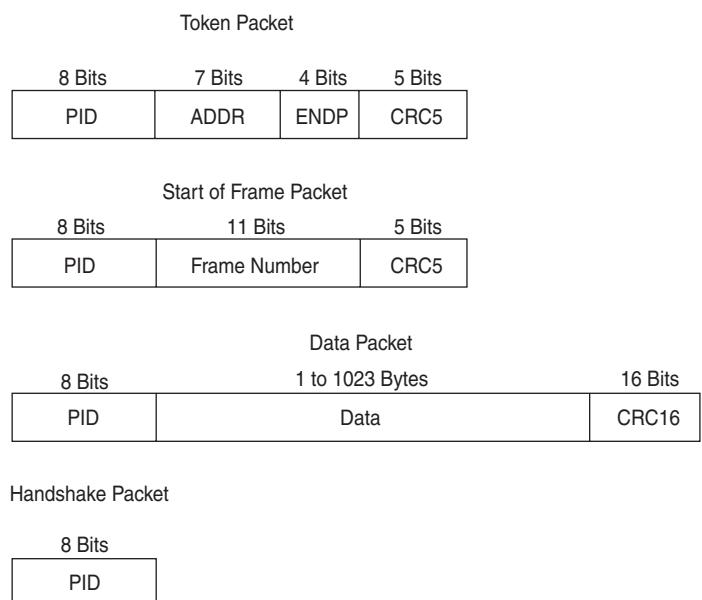

FIGURE 15–18   The data stream and the flowchart used to generate USB data.

TABLE 15–11 PID codes.

PID	Name	Type	Description
E_1	OUT	Token	Host → function transaction
D_2	ACK	Handshake	Receiver accepts packet
C_3	Data0	Data	Data packet (PID even)
A_5	SOF	Token	Start of frame
69	IN	Token	Function → host transaction
5A	NAK	Handshake	Receiver does not accept packet
4B	Data1	Data	Data packet (PID odd)
3C	PRE	Special	Host preamble
2D	Setup	Token	Setup command
1E	Stall	Token	Stalled

FIGURE 15–19 The types of packets and contents found on the USB.

Token Packet

8 Bits	7 Bits	4 Bits	5 Bits
PID	ADDR	ENDP	CRC5

Start of Frame Packet

8 Bits	11 Bits	5 Bits
PID	Frame Number	CRC5

Data Packet

8 Bits	1 to 1023 Bytes	16 Bits
PID	Data	CRC16

Handshake Packet

8 Bits
PID
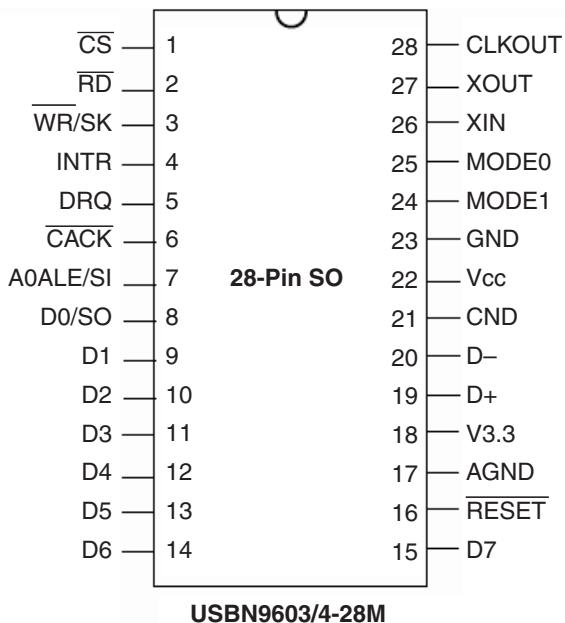
token back to the host. If the data and CRC are received correctly, the ACK is sent; if not, the NAK is sent. If the host receives a NAK token, it retransmits the data packet until the receiver finally receives it correctly. This method of data transfer is often called **stop and wait flow control**. The host must wait for the client to send an ACK or NAK before transferring additional data packets.

The USB Bus Node

National Semiconductor produces a USB bus interface that is fairly easy to interface to the micro-processor. Figure 15–20 illustrates the USBN9604 USB node. Connecting this device to a system using non-DMA access is accomplished by connecting the data bus to D_0–D_7, the control inputs $\overline{RD}$, $\overline{WR}$, and $\overline{CS}$, and a 24 MHz fundamental crystal across the X_{In} and X_{Out} pins. The USB bus connection is located on the D− and D+ pins. The simplest interface is achieved by connecting the two mode inputs to ground. This places the device into a nonmultiplexed parallel mode. In this mode the A_0 pin is used to select address (1) or data (0). Figure 15–21 shows this connection to the microprocessor decodes at I/O port addresses 0300H (data) and 0301H (address).

The USBN9604 is a USB bus transceiver that can receive USB data and transmit USB data. This provides an interface point to the USB bus for a minimal cost of about two dollars.

FIGURE 15–20 The USB bus node from National Semiconductor.

USBN9603/4-28M

Software for the USBN9604/3

The software presented here functions with the interface in Figure 15–21. Not provided is the driver software for the host system. Example 15–13 illustrates the code required to initialize the USB controller. The USBINT procedure sets the USB controller to use end point zero for data transfers.

EXAMPLE 15–13

```
SEND    MACRO   ADDR, DATA
        MOV     DX,301H
        MOV     AL,ADDR
        OUT     DX,AL
        MOV     DX,300H
        MOV     AL,DATA
        OUT     DX,AL
        ENDM

USBINT PROC NEAR

        SEND    0,5             ;interrupts off, software reset USB
        SEND    0,4             ;clear reset
        CALL    DELAY1          ;wait 1ms
        SEND    9,40H           ;enable reset check
        SEND    0DH,3           ;enable EP0 for receive data
        SEND    0BH,3           ;enable EP0 for transmit data
        SEND    20H,0           ;EPO control to no default address
        SEND    4,80H           ;set FAR to accept default address
        SEND    0,8CH           ;USB is ready to send or receive data

USBINT ENDP
```

Once the USB controller is initialized, data can be sent or received to the host system through the USB. To accomplish data transmission, the procedure illustrated in Example 15–14 is called to send a 1-byte packet using the TXD0 FIFO. This procedure uses the SEND macro listed in Example 15–13 to transfer the byte in BL through the USB to the host system.
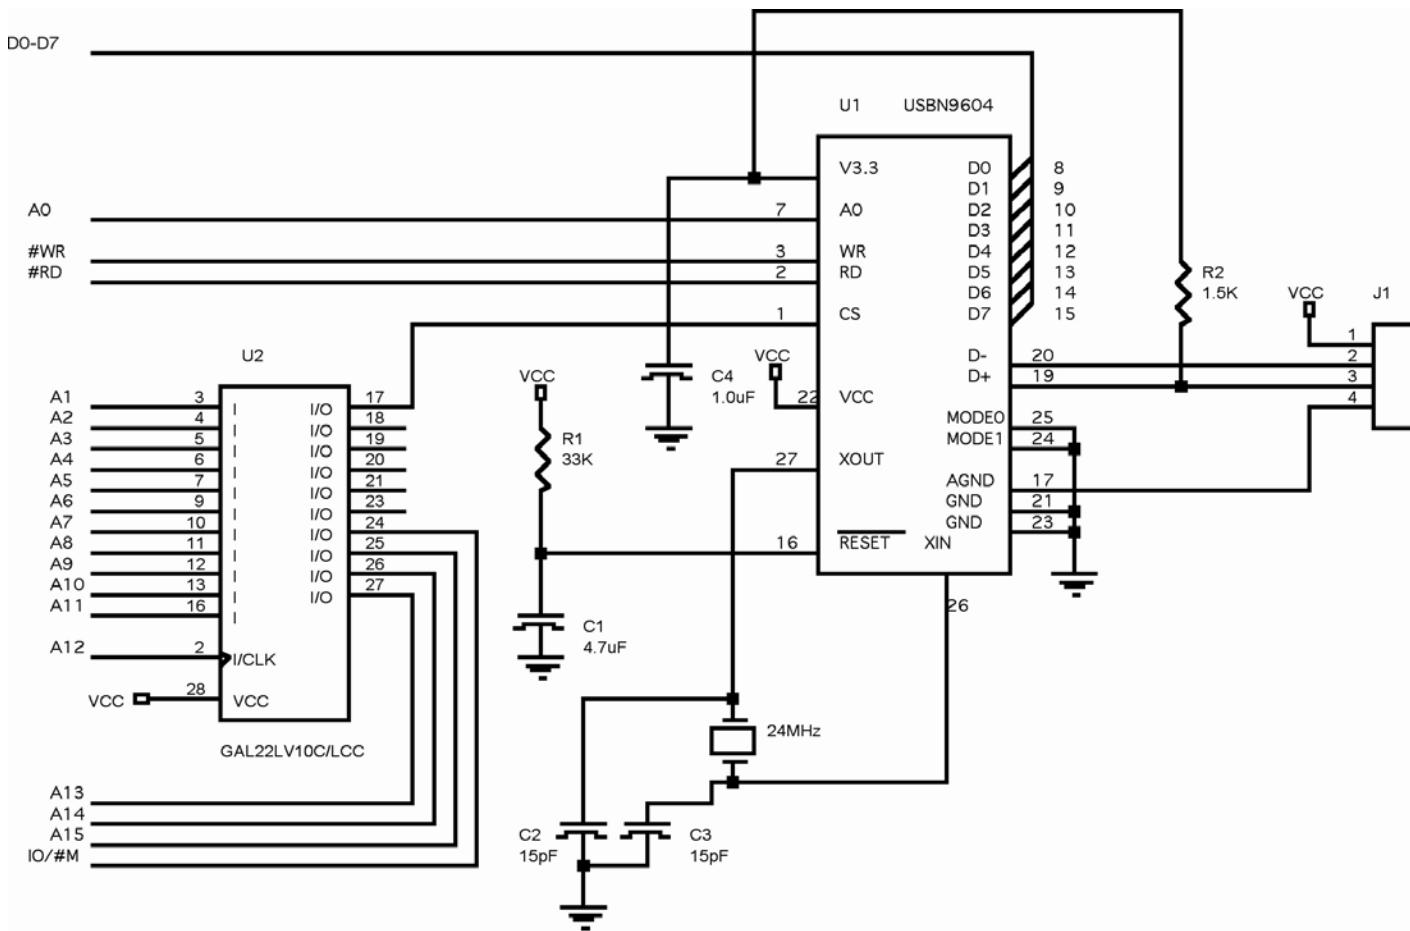
FIGURE 15–21 The USBN9604 interfaced to a microprocessor at I/O addresses 300H and 301H.

EXAMPLE 15–14

```
TRANS  PROC   NEAR

       SEND 21H,BL              ;send BL to FIFO
       SEND 23H,3               ;send the byte

TRANS ENDP
```

To receive data from the USB, two functions are required. One tests to see if data are available and the other reads a byte from the USB and places it into the BL register. Both procedures are listed in Example 15–15. The STATUS procedure checks to see if data are in the receiver FIFO. If data are present, carry is set upon return and if no data are received, carry is cleared. The READS procedure retrieves a byte for the receiver FIFO and returns it in BL.

EXAMPLE 15–15

```
READ   MACRO  ADDR

       MOV DX,301H
       MOV AL,ADDR
       OUT DX,AL
       MOV DX,300H
       IN  AL,DX
       ENDM

STATUS PROC NEAR

       SEND 6
       SEND 6
       SHL AL,2
       RET

STATUS ENDP

READS PROC NEAR

       READ 25H
       RET

READS ENDP
```
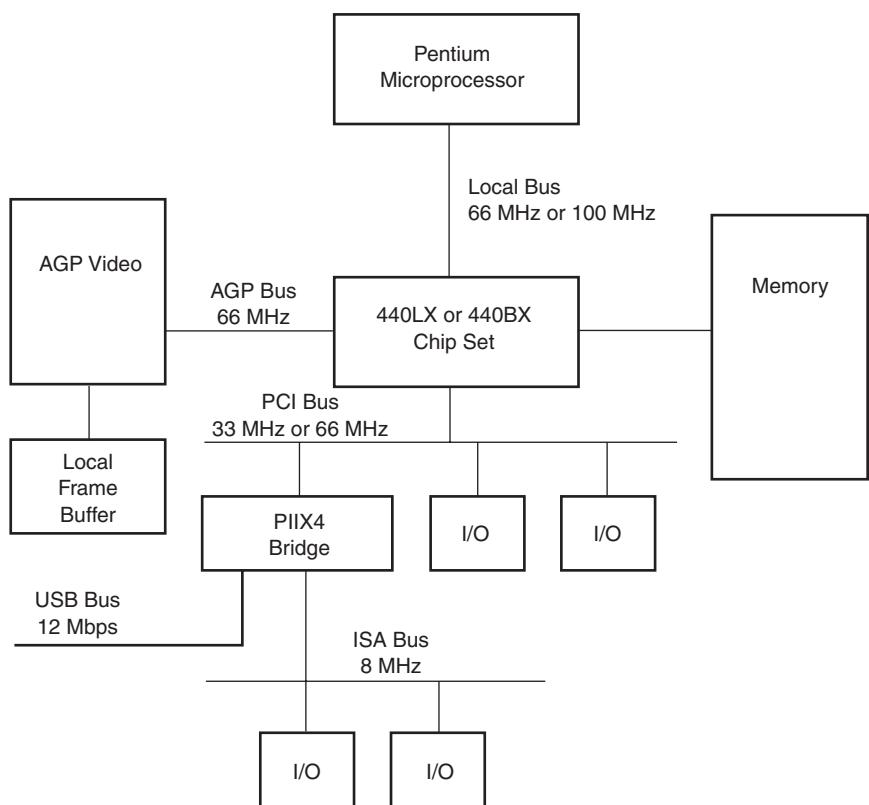
15–6   ACCELERATED GRAPHICS PORT (AGP)

The latest addition to most computer systems was the inclusion of the accelerated graphics port (AGP), until the PCI Express interface became available for video. The AGP operates at the bus clock frequency of the microprocessor. It is designed so that a transfer between the video card and the system memory can progress at a maximum speed. The AGP can transfer data at a maximum rate of 2G bytes per second. This port probably will never be used for any devices other than the video card, so we do not devote much space to its coverage. Because PCI Express video cards use 8 lanes, data transfer occurs at a rate of 4 GBps for the x16 PCI Express video card.

Figure 15–22 illustrates the interface of the AGP to a Pentium 4 system and the placement of other buses in the system. The main advantage of the AGP bus over the PCI bus is that the AGP can sustain transfers (using the 8X compliant system) at speeds up to 2G bytes per second. The 4X system transfers data at rates of over 1G byte per second. The PCI bus has a maximum transfer speed of about 133M bytes per second. The AGP is designed specifically to allow high-speed transfers between the video card frame buffer and the system memory through the chip set.

FIGURE 15–22 Structure of a modern computer, illustrating all the buses.

15–7   SUMMARY

1. The bus systems (ISA, PCI, and USB) allow I/O and memory systems to be interfaced to the personal computer.
2. The ISA bus is either 8 or 16 bits, and supports either memory or I/O transfers at rates of 8 MHz.
3. The PCI (peripheral component interconnect) supports 32- or 64-bit transfers between the personal computer and memory or I/O at rates of 33 MHz. This bus also allows virtually any microprocessor to be interfaced to the PCI bus via the use of a bridge interface.
4. The PCI Express bus found on most computers is in the form of single lane or 16-lane ports. The single lane port is interfaced to I/O devices, whereas the 16-lane port is interfaced to the video card replacing AGP.
5. A plug-and-play (PnP) interface is one that contains a memory that holds configuration information for the system.
6. The parallel port called LPT_1 is used to transfer 8-bit data in parallel to printers and other devices.
7. The serial COM ports are used for serial data transfer. The Windows API is used in a Windows Visual C++ application to effect serial data transfer through the COM ports.
8. The universal serial bus (USB) has all but replaced the ISA bus in the most advanced systems. The USB has three data transfer rates: 1.5 Mbps, 12 Mbps, and 480 Mbps.
9. The USB uses the NRZI system to encode data, and uses bit stuffing for logic 1 transmission more than 6 bits long.
10. The accelerated graphics port (AGP) is a high-speed connection between the memory system and the video graphics card.

15–8 QUESTIONS AND PROBLEMS

1. The letters ISA are an acronym for what phrase?
2. The ISA bus system supports what size data transfers?
3. Is the ISA bus interface often used for memory expansion?
4. Develop an ISA bus interface that is decoded at addresses 310H–313H. This interface must contain an 8255 accessed via these port addresses. (Don't forget to buffer all inputs to the ISA bus card.)
5. Develop an ISA bus interface that decodes ports 0340H–0343H to control a single 8254 timer.
6. Develop a 32-bit PCI bus interface that adds a 27C256 EPROM at memory addresses FFFF0000H–FFFF7FFFH.
7. Given a 74LS244 buffer and a 74LS374 latch, develop an ISA bus interface that contains an 8-bit input port at I/O address 308H and an 8-bit output port at I/O address 30AH.
8. Create an ISA bus interface that allows four channels of analog output signals from 0 to 5.0 V each. These four channels must be decoded at I/O addresses 300H, 310H, 320H, and 330H. Also develop software that supports the four channels.
9. Redo question 8, but instead of four output channels, use four ADCs to create four analog input channels at the same addresses.
10. Using an 8254 timer or timers, develop a darkroom timer on an ISA bus card. Your timer must generate a logic 0 for 1/100-second intervals from 1/100 second to five minutes. Use the system clock of 8 MHz as a timing source. The software you develop must allow the user to select the time from the keyboard. The output signal from the timer must be a logic 0 for the duration of the selected time and must be passed through an inverter to enable a solid-state relay that controls the photographic enlarger.
11. Interface a 16550 UART to the personal computer through the PCI bus interface. Develop software that transmits and receives data at baud rates of 300, 1200, 9600, and 19,200. The UART must respond to I/O ports 1E3XH.
12. The ISA bus can transfer data that are _____ wide at the rate of 8 MHz.
13. Describe how the address can be captured from the PCI bus.
14. What is the purpose of the configuration memory found on the PCI bus interface?
15. Define the term *plug-and-play*.
16. What is the purpose of the C/$\overline{\text{BE}}$ connection on the PCI bus system?
17. How is the BIOS tested for the PCI BIOS extension?
18. Develop a short program that interrogates the PCI bus, using the extension to the BIOS, and reads the 32-bit contents of configuration register 08H. For this problem, consider that the bus and unit numbers are 0000H.
19. What advantage does the PCI bus exhibit over the ISA bus?
20. How fast does the PCI Express bus transfer serial data?
21. What is a lane in a PCI Express interface?
22. The parallel port is decoded at which I/O addresses in a personal computer?
23. Can data be read from the parallel port?
24. The parallel port connecter found on the back of the computer has _____ pins.
25. Most computers contain at least one serial communication port. What is this port called?
26. Develop a C++ function that sends the letters ABC through the serial port and continues to do so until the letters ABC are returned through the serial port. Show all functions needed to accomplish this, including any initialization.
27. Modify Example 15–9 so it sends a character string of any length.
28. Search the Internet and detail, in a short report, variants as used in the Visual Programming environment.

29. Can a USB device appear as a COM device?
30. What data rates are available for use on the USB?
31. How are data encoded on the USB?
32. What is the maximum cable length for use with the USB?
33. Will the USB ever replace the ISA bus?
34. How many device addresses are available on the USB?
35. What is NRZI encoding?
36. What is a stuffed bit?
37. If the following raw data are sent on the USB, draw the waveform of the signal found on the USB: (1100110000110011011010).
38. How long can a data packet be on the USB?
39. What is the purpose of the NAK and ACK tokens on the USB?
40. Describe the difference in data transfer rates on the PCI bus when compared with the AGP.
41. What is the transfer rate in a system using an 8X AGP video card?
42. What is the transfer rate of a PCI Express 16X video card?
43. On the Internet, locate a few video card manufacturers and find how much memory is available on AGP video cards. List the manufacturers and the amount of memory on the cards.
44. Using the Internet, write a report that details any USB controller.

CHAPTER 16

The 80186, 80188, and 80286 Microprocessors

INTRODUCTION

The Intel 80186/80188 and the 80286 are enhanced versions of the earlier versions of the 80X86 family of microprocessors. The 80186/80188 and 80286 are all 16-bit microprocessors that are upward-compatible to the 8086/8088. Even the hardware of these microprocessors is similar to the earlier versions. This chapter presents an overview of each microprocessor and points out the differences or enhancements that are present in each version. The first part of the chapter describes the 80186/80188 microprocessors, and the last part shows the 80286 microprocessor.

New to recent editions is expanded coverage of the 80186/80188 family. Intel has added four new versions of each of these embedded controllers to its lineup of microprocessors. Each is a CMOS version and is designated with a two-letter suffix: XL, EA, EB, and EC. The 80C186XL and 80C188XL models are most similar to the earlier 80186/80188 models.

CHAPTER OBJECTIVES

Upon completion of this chapter, you will be able to:

1. Describe the hardware and software enhancements of the 80186/80188 and the 80286 microprocessors as compared to the 8086/8088.
2. Detail the differences between the various versions of the 80186 and 80188 embedded controllers.
3. Interface the 80186/80188 and the 80286 to memory and I/O.
4. Develop software using the enhancements provided in these microprocessors.
5. Describe the operation of the memory management unit (MMU) within the 80286 microprocessor.
6. Define and detail the operation of a real-time operating system (RTOS).

16–1   80186/80188 ARCHITECTURE

The 80186 and 80188, like the 8086 and 8088, are nearly identical. The only difference between the 80186 and 80188 is the width of their data buses. The 80186 (like the 8086) contains a 16-bit data bus, while the 80188 (like the 8088) contains an 8-bit data bus. The internal register structure

of the 80186/80188 is virtually identical to that of the 8086/8088. About the only difference is that the 80186/80188 contain additional reserved interrupt vectors and some very powerful built-in I/O features. The 80186 and 80188 are often called embedded controllers because of their application as a controller, not as a microprocessor-based computer.

Versions of the 80186/80188

As mentioned, the 80186 and 80188 are available in four different versions, which are all CMOS microprocessors. Table 16–1 lists each version and the major features provided. The 80C186XL and 80C188XL are the most basic versions of the 80186/80188; the 80C186EC and 80C188EC are the most advanced. This text details the 80C186XL/80C188XL, and then describes the additional features and enhancements provided in the other versions.

80186 Basic Block Diagram

Figure 16–1 provides the block diagram of the 80188 microprocessor that generically represents all versions except for the enhancements and additional features outlined in Table 16–1. Notice that this microprocessor has a great deal more internal circuitry than the 8088. The block diagrams of the 80186 and 80188 are identical except for the prefetch queue, which is four bytes in the 80188 and six bytes in the 80186. Like the 8088, the 80188 contains a bus interface unit (BIU) and an execution unit (ED).

In addition to the BIU and ED, the 80186/80188 family contains a clock generator, a programmable interrupt controller, programmable timers, a programmable DMA controller, and a programmable chip selection unit. These enhancements greatly increase the utility of the 80186/80188 and reduce the number of peripheral components required to implement a system. Many popular subsystems for the personal computer use the 80186/80188 microprocessors as

TABLE 16–1 The four versions of the 80186/80188 embedded controller.

| Feature	80C186XL 80C188XL	80C186EA 80C188EA	80C186EB 80C188EB	80C186EC 80C188EC
80286-like instruction set	✓	✓	✓	✓
Power-save (green mode)	✓	✓		✓
Power down mode		✓	✓	✓
80C187 interface	✓	✓	✓	✓
ONCE mode	✓	✓	✓	✓
Interrupt controller	✓	✓	✓	✓ 8259-like
Timer unit	✓	✓	✓	✓
Chip selection unit	✓	✓	✓ enhanced	✓ enhanced
DMA controller	✓ 2-channel	✓ 2-channel		✓ 4-channel
Serial communications unit			✓	✓
Refresh controller	✓	✓	✓ enhanced	✓ enhanced
Watchdog timer				✓
I/O ports			✓ 16 bits	✓ 22 bits
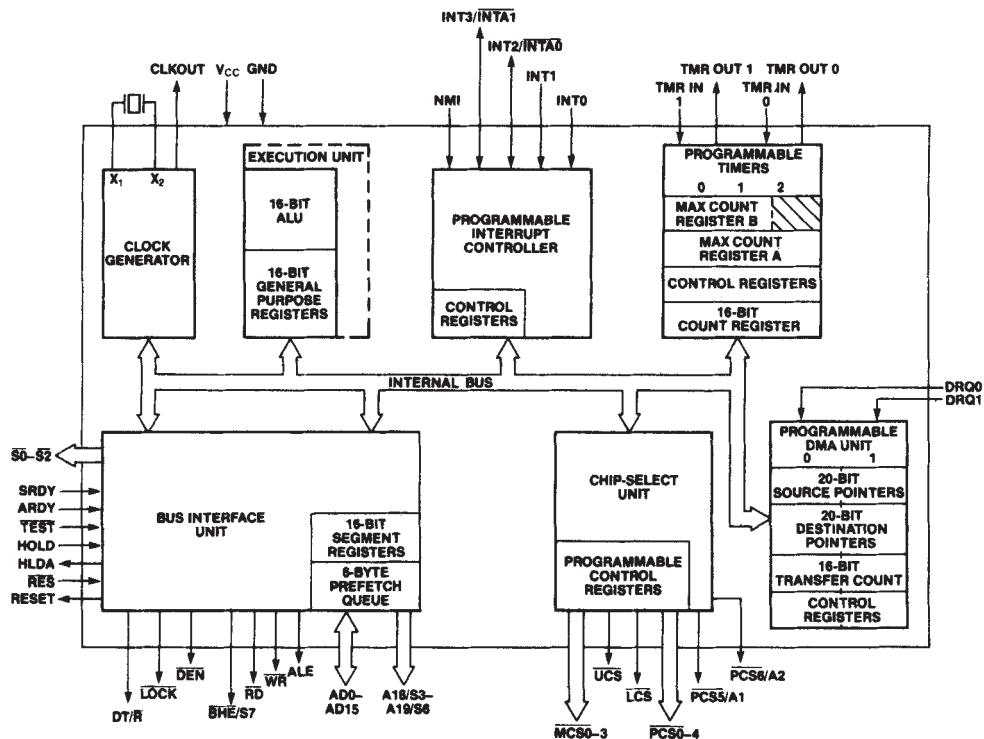
FIGURE 16–1 The block diagram of the 80186 microprocessor. Note that the block diagram of the 80188 is identical, except that $\overline{BHE}$/S7 is missing and AD_{15}–AD_8 are relabeled A_{15}–A_8. (Courtesy of Intel Corporation.)

caching disk controllers, local area network (LAN) controllers, and so forth. The 80186/80188 also finds application in the cellular telephone network as a switcher.

Software for the 80186/80188 is identical to that for the 80286 microprocessor, without the memory management instructions. This means that the 80286-like instructions for immediate multiplication, immediate shift counts, string I/O, PUSHA, POPA, BOUND, ENTER, and LEAVE all function on the 80186/80188 microprocessors.

80186/80188 Basic Features

In this segment of the text, we introduce the enhancements of the 80186/80188 microprocessors or embedded controllers that apply to all versions except where noted, but we do not provide exclusive coverage. More details on the operation of each enhancement and details of each advanced version are provided later in the chapter.

Clock Generator. The internal clock generator replaces the external 8284A clock generator used with the 8086/8088 microprocessors. This reduces the component count in a system.

The internal clock generator has three pin connections: X_1, X_2, and CLKOUT (or on some versions: CLKIN, OSCOUT, and CLKOUT). The X_1 (CLKIN) and X_2 (OSCOUT) pins are connected to a crystal that resonates at twice the operating frequency of the microprocessor. In the 8 MHz version of the 80186/80188, a 16 MHz crystal is attached to X_1 (CLKIN) and X_2 (OSCOUT). The 80186/80188 is available in 6 MHz, 8 MHz, 12 MHz, 16 MHz, or 25 MHz versions.

The CLKOUT pin provides a system clock signal that is one half the crystal frequency, with a 50% duty cycle. The CLKOUT pin drives other devices in a system and provides a timing source to additional microprocessors in the system.

In addition to these external pins, the clock generator provides the internal timing for synchronizing the READY input pin, whereas in the 8086/8088 system, READY synchronization is provided by the 8284A clock generator.

Programmable Interrupt Controller.   The programmable interrupt controller (PIC) arbitrates the internal and external interrupts and controls up to two external 8259A PICs. When an external 8259 is attached, the 80186/80188 microprocessors function as the master and the 8259 functions as the slave. The 80C186EC and 80C188EC models contain an 8259A-compatible interrupt controller in place of the one described here for the other versions (XL, EA, and EB).

If the PIC is operated without the external 8259, it has five interrupt inputs: INT0–INT3 and NMI. Note that the number of available interrupts depends on the version: The EB version has six interrupt inputs and the EC version has 16. This is an expansion from the two interrupt inputs available on the 8086/8088 microprocessors. In many systems, the five interrupt inputs are adequate.

Timers.   The timer section contains three fully programmable l6-bit timers. Timers 0 and 1 generate waveforms for external use and are driven by either the master clock of the 80186/80188 or by an external clock. They are also used to count external events. The third timer, timer 2, is internal and clocked by the master clock. The output of timer 2 generates an interrupt after a specified number of clocks and can provide a clock to the other timers. Timer 2 can also be used as a watchdog timer because it can be programmed to interrupt the microprocessor after a certain length of time.

The 80C186EC and 80C188EC models have an additional timer called a *watchdog*. The watchdog timer is a 32-bit counter that is clocked internally by the CLKOUT signal (one half the crystal frequency). Each time the counter hits zero, it reloads and generates a pulse on the WDTOUT pin that is four CLKOUT periods wide. This output can be used for any purpose: It can be wired to the reset input to cause a reset or to the NMI input to cause an interrupt. Note that if it is connected to the reset or NMI inputs, it is periodically reprogrammed so that it never counts down to zero. The purpose of a watchdog timer is to reset or interrupt the system if the software goes awry.

Programmable DMA Unit.   The programmable DMA unit contains two DMA channels or four DMA channels in the 80C186EC/80C188EC models. Each channel can transfer data between memory locations, between memory and I/O, or between I/O devices. This DMA controller is similar to the 8237 DMA controller discussed in Chapter 13. The main difference is that the 8237 DMA controller has four DMA channels, as does the EC model.

Programmable Chip Selection Unit.   The chip selection is a built-in programmable memory and I/O decoder. It has six output lines to select memory, seven lines to select I/O on the XL and EA models, and 10 lines that select either memory or I/O on the EB and EC models.

On the XL and EA models, the memory selection lines are divided into three groups that select memory for the major sections of the 80186/80188 memory map. The lower memory select signal enables memory for the interrupt vectors, the upper memory select signal enables memory for reset, and the middle memory select signals enable up to four middle memory devices. The boundary of the lower memory begins at location 00000H and the boundary of the upper memory ends at location FFFFFH. The sizes of the memory areas are programmable, and wait states (0–3 waits) can be automatically inserted with the selection of an area of memory.

On the XL and EA models, each programmable I/O selection signal addresses a 128-byte block of I/O space. The programmable I/O area starts at a base I/O address programmed by the user, and all seven 128-byte blocks are contiguous.

On the EB and EC models, there is an upper and lower memory chip selection pin and eight general-purpose memory or I/O chip selection pins. Another difference is that from 0 to 15 wait states can be programmed in these two versions of the 80186/80188 embedded controllers.

Power Save/Power Down Feature. The power save feature allows the system clock to be divided by 4, 8, or 16 to reduce power consumption. The power-saving feature is started by software and exited by a hardware event such as an interrupt. The power down feature stops the clock completely, but it is not available on the XL version. The power down mode is entered by execution of an HLT instruction and is exited by any interrupt.

Refresh Control Unit. The refresh control unit generates the refresh row address at the interval programmed. The refresh control unit does not multiplex the address for the DRAM—this is still the responsibility of the system designer. The refresh address is provided to the memory system at the end of the programmed refresh interval, along with the $\overline{\text{RFSH}}$ control signal. The memory system must run a refresh cycle during the active time of the $\overline{\text{RFSH}}$ control signal. More on memory and refreshing is provided in the section that explains the chip selection unit.

Pin-Out

Figure 16–2 illustrates the pin-out of the 80C186XL microprocessor. Note that the 80C186XL is packaged in either a 68-pin leadless chip carrier (LCC) or in a pin grid array (PGA). The LCC package and PGA packages are illustrated in Figure 16–3.

Pin Definitions. The following list defines each 80C186XL pin and notes any differences between the 80C186XL and 80C188XL microprocessors. The enhanced versions are described later in this chapter.

|	
V_{CC}	This is the system **power** supply connection for ±10%, +5.0 V.
V_{SS}	This is the system **ground** connection.
X_1 **and** X_2	The **clock** pins are generally connected to a fundamental-mode parallel resonant crystal that operates an internal crystal oscillator. An external clock signal may be connected to the X_1 pin. The internal master clock operates at one half the external crystal or clock input signal. Note that these pins are labeled CLKIN (X_1) and OSCOUT (X_2) on some versions of the 80186/80188.
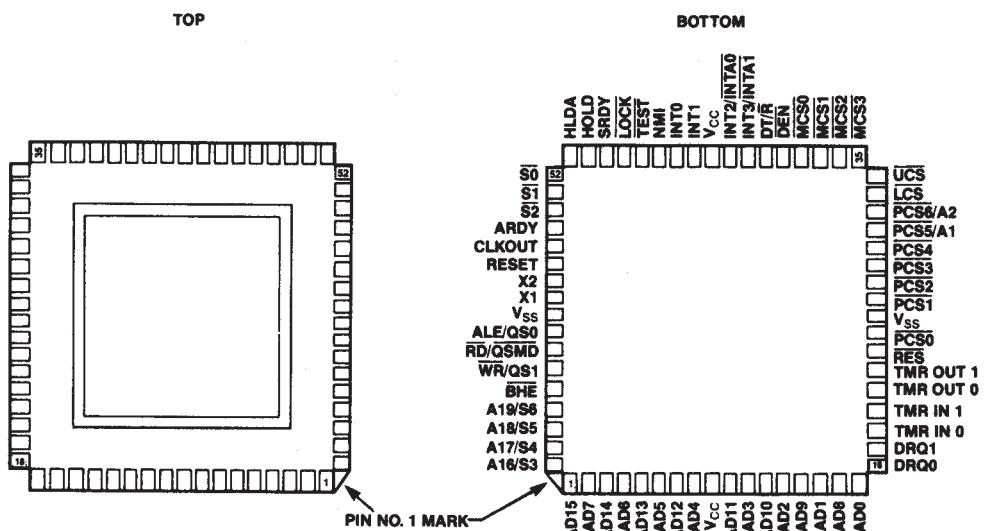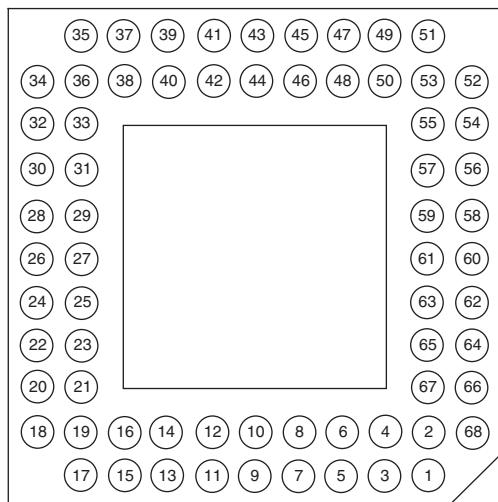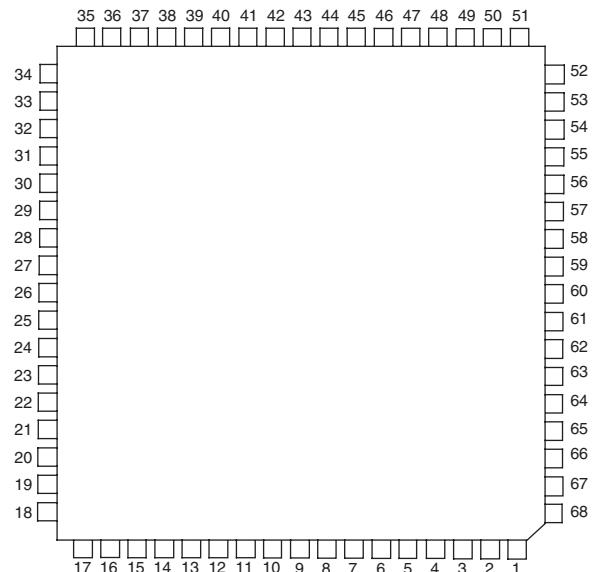
**FIGURE 16–2** Pin-out of the 80186 microprocessor. (Courtesy of Intel Corporation.)

PGA Bottom View

LCC Bottom View

FIGURE 16–3 The bottom views of the PGA and LCC style versions of the 80C188XL microprocessor.

CLKOUT	Clock out provides a **timing signal** to system peripherals at one half the clock input frequency with a 50% duty cycle.
$\overline{\text{RES}}$	The **reset input** pin resets the 80186/80188. For a proper reset, the $\overline{\text{RES}}$ must be held low for at least 50 ms after power is applied. This pin is often connected to an RC circuit that generates a reset signal after power is applied. The reset location is identical to that of the 8086/8088 microprocessor—FFFF0H.
RESET	The companion **reset output** pin (goes high for a reset) connects to system peripherals to initialize them whenever the $\overline{\text{RES}}$ input goes low.
$\overline{\text{TEST}}$	This test pin connects to the BUSY output of the 80187 numeric coprocessor. The $\overline{\text{TEST}}$ pin is interrogated with the FWAIT or WAIT instruction.
$T_{in}0$ and $T_{in}1$	These pins are used as **external clocking sources** to timers 0 and 1.
$T_{out}0$ and $T_{out}1$	These pins provide the **output signals** from timers 0 and 1, which can be programmed to provide square waves or pulses.
DRQ0 and DRQ1	These pins are active-high-level triggered **DMA request** lines for DMA channels 0 and 1.
NMI	This is a **non-maskable interrupt** input. It is positive edge-triggered and always active. When NMI is activated, it uses interrupt vector 2.
INT_0, INT_1, $INT_2/\overline{\text{INTA0}}$, and $INT_3/\overline{\text{INTA1}}$	These are **maskable interrupt inputs.** They are active-high and are programmed as either level or edge-triggered. These pins are configureed as four interrupt inputs if no external 8259 is present, or as two interrupt inputs if the 8259A interrupt controller is present.

$A_{19}/\overline{ONCE}$, A_{18}, A_{17}, and A_{16}	These are **multiplexed address/status** connections that provide the address (A_{19}–A_{16}) and status (S_6–S_3). Status bits found on address pins A_{18}–A_{16} have no system function and are used during manufacturing for testing. The A_{19} pin is an input for the $\overline{ONCE}$ function on a reset. If $\overline{ONCE}$ is held low on a reset, the microprocessor enters a testing mode.
AD_{15}–AD_0	These are **multiplexed address/data** bus connections. During T_1, the 80186 places A_{15}–A_0 on these pins; during T_2, T_3, and T_4, the 80186 uses these pins as the data bus for signals D_{15}–D_0. Note that the 80188 has pins AD_7–AD_0 and A_{15}–A_8.
$\overline{BHE}$	The **bus high enable** pin indicates (when a logic 0) that valid data are transferred through data bus connections D_{15}–D_8.
ALE	**Address latch enable** is an output pin that contains ALE one-half clock cycle earlier than in the 8086. It is used to demultiplex the address/data and address/status buses. (Even though the status bits on A_{19}–A_{16} are not used in the system, they must still be demultiplexed.)
$\overline{WR}$	The **write** pin causes data to be written to memory or I/O.
$\overline{RD}$	The **read** pin causes data to be read from memory or I/O.
ARDY	The **asynchronous READY** input informs the 80186/80188 that the memory or I/O is ready for the 80186/80188 to read or write data. If this pin is tied to +5.0 V, the microprocessor functions normally; if it is grounded, the microprocessor enters wait states.
SRDY	The **synchronous READY** input is synchronized with the system clock to provide a relaxed timing for the ready input. Like ARDY, SRDY is tied to +5.0 V for no wait states.
$\overline{LOCK}$	The **lock** pin is an output controlled by the LOCK prefix. If an instruction is prefixed with LOCK, the $\overline{LOCK}$ pin becomes a logic 0 for the duration of the locked instruction.
S_2, S_1, and S_0	These are status bits that provide the system with the type of bus transfer in effect. See Table 16–2 for the states of the status bits. The **upper-memory chip select** pin selects memory on the upper portion of the memory map.
$\overline{UCS}$	The **upper-memory chip select** output is programmable to enable memory sizes of 1K to 256K bytes ending at location FFFFFH. Note that this pin is programmed differently on the EB and EC versions and enables memory between 1K and 1M long.

TABLE 16–2 The S_2, S_1, and S_0 status bits.

S_2	S_1	S_0	Function
0	0	0	Interrupt acknowledge
0	0	1	I/O read
0	1	0	I/O write
0	1	1	Halt
1	0	0	Opcode fetch
1	0	1	Memory read
1	1	0	Memory write
1	1	1	Passive

$\overline{\text{LCS}}$	The **lower-memory chip select** pin enables memory beginning at location 00000H. This pin is programmed to select memory sizes from 1K to 256K bytes. Note that this pin functions differently for the EB and EC versions and enables memory between 1K and 1M bytes long.
$\overline{\text{MCS0}}-\overline{\text{MCS3}}$	The **middle-memory chip select** pins enable four middle memory devices. These pins are programmable to select an 8K to 512K byte block of memory, containing four devices. Note that these pins are not present on the EB and EC versions.
$\overline{\text{PCS0}}-\overline{\text{PCS4}}$	These are five different **peripheral selection** lines. Note that the lines are not present on the EB and EC versions.
$\overline{\text{PCS5}}/\text{A}_1$ **and** $\overline{\text{PCS6}}/\text{A}_2$	These are programmed as peripheral selection lines or as internally latched address bits A_2 and A_1. These lines are not present on the EB and EC versions.
$\text{DT}/\overline{\text{R}}$	The **data transmit/receive** pin controls the direction of data bus buffers if attached to the system.
$\overline{\text{DEN}}$	The **data bus enable** pin enables the external data bus buffers.

DC Operating Characteristics

It is necessary to know the DC operating characteristics before attempting to interface or operate the microprocessor. The 80C186/801C88 microprocessors require between 42 mA and 63 mA of power-supply current. Each output pin provides 3.0 mA of logic 0 current and −2 mA of logic 1 current.

80186/80188 Timing
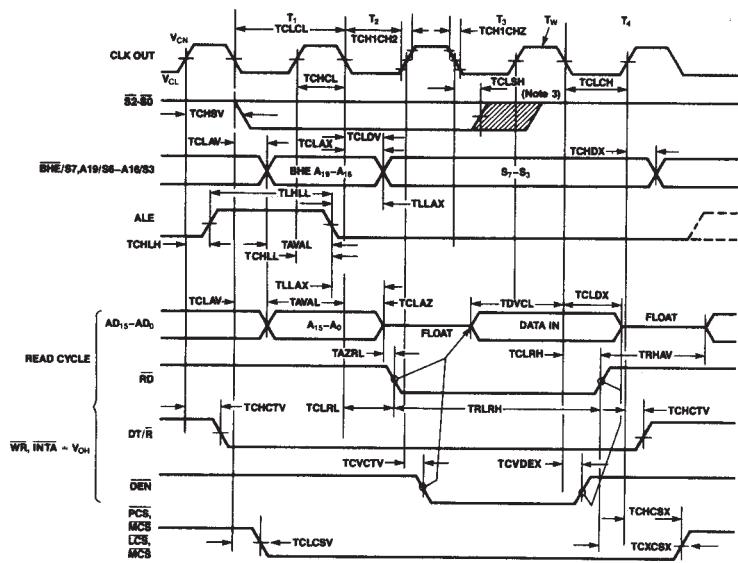
The timing diagram for the 80186 is provided in Figure 16–4. Timing for the 80188 is identical except for the multiplexed address connection, which are AD_7-AD_0 instead of $\text{AD}_{15}-\text{AD}_0$, and the $\overline{\text{BHE}}$, which does not exist on the 80188.

The basic timing for the 80186/80188 is composed of four clocking periods just as in the 8086/8088. A bus cycle for the 8 MHz version requires 500 ns, while the 16 MHz version requires 250 ns.
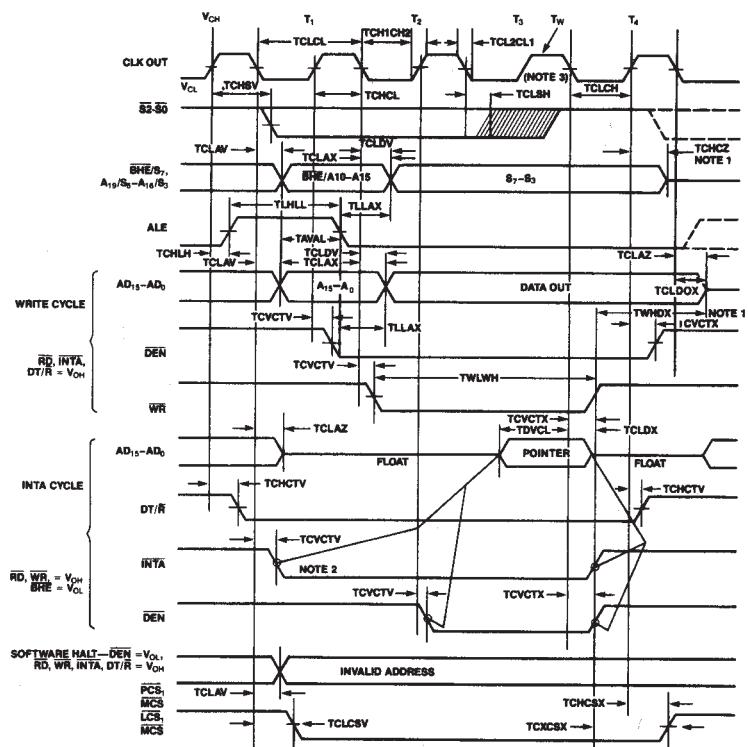
There are very few differences between the timing for the 80186/80188 and the 8086/8088. The most noticeable difference is that ALE appears one-half clock cycle earlier in the 80186/80188.

Memory Access Time.   One of the more important points in any microprocessor's timing diagram is the memory access time. Access time calculations for the 80186/80188 are identical to that of the 8086/8088. Recall that the access time is the time allotted to the memory and I/O to provide data to the microprocessor after the microprocessor sends the memory or I/O its address.

A close examination of the timing diagram reveals that the address appears on the address bus T_{CLAV} time after the start of T_1. T_{CLAV} is listed as 44 ns for the 8 MHz version. (See Figure 16–5.) Data are sampled from the data bus at the end of T_3, but a setup time is required before the clock defined as T_{DVCL}. The times listed for T_{DVCL} are 20 ns for both versions of the microprocessor. Access time is therefore equal to three clocking periods minus both T_{CLAV} and T_{DVCL}. Access time for the 8 MHz microprocessor is 375 ns − 44 ns − 20 ns, or 311 ns. The access time for the 16 MHz version is calculated in the same manner, except that T_{CLAV} is 25 ns and T_{DVCL} is 15 ns.

FIGURE 16–4 80186/80188 timing. (a) Read cycle timing and (b) write cycle timing. (Courtesy of Intel Corporation.)

80186 Master Interface Timing Responses

Symbol	Parameters	80188 (8 MHz) Min.	Max.	80188-6 (6 MHz) Min.	Max.	Units	Test Conditions
T_{CLAV}	Address Valid Delay	5	44	5	63	ns	C_L = 20-200 pF all outputs
T_{CLAX}	Address Hold	10		10		ns	
T_{CLAZ}	Address Float Delay	T_{CLAX}	35	T_{CLAX}	44	ns	
T_{CHCZ}	Command Lines Float Delay		45		56	ns	
T_{CHCV}	Command Lines Valid Delay (after float)		55		76	ns	
T_{LHLL}	ALE Width	$T_{CLCL-35}$		$T_{CLCL-35}$		ns	
T_{CHLH}	ALE Active Delay		35		44	ns	
T_{CHLL}	ALE Inactive Delay		35		44	ns	
T_{LLAX}	Address Hold to ALE Inactive	$T_{CHCL-25}$		$T_{CHCL-30}$		ns	
T_{CLDV}	Data Valid Delay	10	44	10	55	ns	
T_{CLDOX}	Data Hold Time	10		10		ns	
T_{WHDX}	Data Hold after WR	$T_{CLCL-40}$		$T_{CLCL-50}$		ns	
T_{CVCTV}	Control Active Delay 1	5	70	5	87	ns	
T_{CHCTV}	Control Active Delay 2	10	55	10	76	ns	
T_{CVCTX}	Control Inactive Delay	5	55	5	76	ns	
T_{CVDEX}	DEN Inactive Delay (Non-Write Cycle)		70		87	ns	
T_{AZRL}	Address Float to RD Active	0		0		ns	
T_{CLRL}	RD Active Delay	10	70	10	87	ns	
T_{CLRH}	RD Inactive Delay	10	55	10	76	ns	
T_{RHAV}	RD Inactive to Address Active	$T_{CLCL-40}$		$T_{CLCL-50}$		ns	
T_{CLHAV}	HLDA Valid Delay	10	50	10	67	ns	
T_{RLRH}	RD Width	$2T_{CLCL-50}$		$2T_{CLCL-50}$		ns	
T_{WLWH}	WR Width	$2T_{CLCL-40}$		$2T_{CLCL-40}$		ns	
T_{AVAL}	Address Valid to ALE Low	$T_{CLCH-25}$		$T_{CLCH-45}$		ns	
T_{CHSV}	Status Active Delay	10	55	10	76	ns	
T_{CLSH}	Status Inactive Delay	10	55	10	76	ns	
T_{CLTMV}	Timer Output Delay		60		75	ns	100 pF max
T_{CLRO}	Reset Delay		60		75	ns	
T_{CHQSV}	Queue Status Delay		35		44	ns	

80186 Chip-Select Timing Responses

Symbol	Parameter	Min.	Max.	Min.	Max.	Units	Test Conditions
T_{CLCSV}	Chip-Select Active Delay		66		80	ns	
T_{CXCSX}	Chip-Selct Hold from Command Inactive	35		35		ns	
T_{CHCSX}	Chip-Select Inactive Delay	5	35	5	47	ns	

Symbol	Parameter	Min.	Max.	Units	Test Conditions
TDVCL	Data in Setup (A/D)	20		ns	
TCLDX	Data in Hold (A/D)	10		ns	
TARYHCH	Asynchronous Ready (AREADY) active setup time*	20		ns	
TARYLCL	AREADY inactive setup time	35		ns	
TCHARYX	AREADY hold time	15		ns	
TSRYCL	Synchronous Ready (SREADY) transition setup time	35		ns	
TCLSRY	SREADY transition hold time	15		ns	
THVCL	HOLD Setup*	25		ns	
TINVCH	INTR, NMI, TEST, TIMERIN, Setup*	25		ns	
TINVCL	DRQ0, DRQ1, Setup*	25		ns	

*To guarantee recognition at next clock.

FIGURE 16–5 80186 AC characteristics. (Courtesy of Intel Corporation.)

16–2   ## PROGRAMMING THE 80186/80188 ENHANCEMENTS

This section provides detail on the programming and operation of the 80186/80188 enhancements of all versions (XL, EA, EB, and EC). The next section details the use of the 80C188EB in a system that uses many of the enhancements discussed here. The only new feature not discussed here is the clock generator, which is described in the previous section on architecture.

Peripheral Control Block

All internal peripherals are controlled by a set of 16-bit-wide registers located in the peripheral control block (PCB). The PCB (see Figure 16–6) is a set of 256 registers located in the I/O or memory space. Note that this set applies to the XL and EA versions. Later in this section, the EB and EC versions of the PCB are defined and described.

Whenever the 80186/80188 is reset, the peripheral control block is automatically located at the top of the I/O map (I/O addresses FF00H–FFFFH). In most cases, it stays in this area of I/O space, but the PCB may be relocated at any time to any other area of memory or I/O. Relocation is accomplished by changing the contents of the relocation register (see Figure 16–7) located at offset addresses FEH and FFH.

The relocation register is set to 20FFH when the 80186/80188 is reset. This locates the PCB at I/O addresses FF00H–FFFFH afterwards. To relocate the PCB, the user need only send a word OUT to I/O address FFFEH with a new bit pattern. For example, to relocate the PCB to

FIGURE 16–6   Peripheral control block (PCB) of the 80186/80188. (Courtesy of Intel Corporation.)
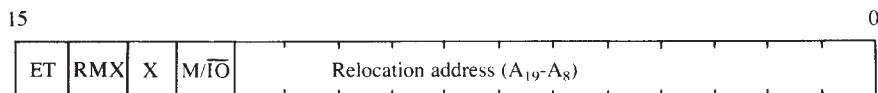
	OFFSET
Relocation Register	FEH
DMA Descriptors Channel 1	DAH / D0H
DMA Descriptors Channel 0	CAH / C0H
Chip-Select Control Registers	A8H / A0H
Timer 2 Control Registers	66H / 60H
Timer 1 Control Registers	5EH / 58H
Timer 0 Control Registers	56H / 50H
Interrupt Controller Registers	3EH / 20H

ET = ESC/NO ESC TRAP
RMX = iRM × 86 mode/master mode
M/$\overline{\text{IO}}$ = Memory/IO space
X = Unused

FIGURE 16–7   Peripheral control register.

memory locations 20000H–200FFH, send 1200H to I/O address FFFEH. Notice that M/$\overline{\text{IO}}$ is a logic 1 to select memory, and that 200H selects memory address 20000H as the base address of the PCB. Note that all accesses to the PCB must be word accesses because it is organized as 16-bit-wide registers. Example 16–1 shows the software required to relocate the PCB to memory locations 20000H–200FFH. Note that either an 8- or 16-bit output can be used to program the 80186; in the 80188, never use the OUT DX,AX instruction because it takes additional clocking periods to execute.

EXAMPLE 16–1

```
MOV    DX,0FFFEH                ;address relocation register
MOV    AX,1200H                 ;new PCB location
OUT    DX,AL                    ;this can also be an OUT DX,AX
```

The EB and EC versions use a different address for programming the PCB location. Both versions have the PCB relocation register stored at offset XXA8H, instead of at offset XXFEH for the XL and EA versions. The bit pattern of these versions is the same as for the XL and EA versions, except that the RMX bit is missing.

Interrupts in the 80186/80188

The interrupts in the 80186/80188 are identical to the 8086/8088, except that additional interrupt vectors are defined for some of the internal devices. A complete listing of the reserved interrupt vectors appears in Table 16–3. The first five are identical to the 8086/8088.

The array BOUND instruction interrupt is requested if the boundary of an index register is outside the values set up in the memory. The unused opcode interrupt occurs whenever the 80186/80188 executes any undefined opcode. This is important if a program begins to run awry. Note that the unused opcode interrupt can be accessed by an instruction, but the assembler does not include it in the instruction set. On the Pentium Pro–Pentium 4 and some earlier Intel microprocessors, the 0F0BH or 0FB9H instruction will cause the program to call the procedure whose address is stored at the unused opcode interrupt vector.

The ESC opcode interrupt occurs if ESC opcodes D8H–DFH are executed. This occurs only if the ET (escape trap) bit of the relocation register is set. If an ESC interrupt occurs, the address stored on the stack by the interrupt points to the ESC instruction or to its segment override prefix, if one is used.

The internal hardware interrupts must be enabled by the I flag bit and must be unmasked to function. The I flag bit is set (enabled) with STI and cleared (disabled) with CLI. The remaining internally decoded interrupts are discussed with the timers and DMA controller, later in this section.

Interrupt Controller
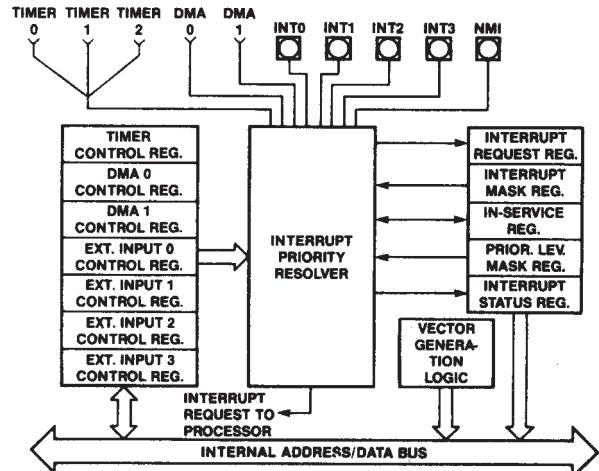
The interrupt controller inside the 80186/80188 is a sophisticated device. It has many interrupt inputs that arrive from the five external interrupt inputs, the DMA controller, and the three timers. Figure 16–8 provides a block diagram of the interrupt structure of the 80186/80188 interrupt

TABLE 16–3 80186/80188
interrupt vectors.

Name	Type	Address	Priority
Divide error	0	00000–00003	1
Single-step	1	00004–00007	1A
NMI pin	2	00008–0000B	1
Breakpoint	3	0000C–0000F	1
Overflow	4	00010–00013	1
BOUND instruction	5	00014–00017	1
Unused opcode	6	00018–0001B	1
ESCape opcode	7	0001C–0001F	1
Timer 0	8	00020–00023	2A
Reserved	9	00024–00027	—
DMA 0	A	00028–0002B	4
DMA 1	B	0002C–0002F	5
INT_0	C	00030–00033	6
INT_1	D	00034–00037	7
INT_2	E	00038–0003B	8
INT_3	F	0003C–0003F	9
80187 coprocessor	10	00040–00043	1
Reserved	11	00044–00047	—
Timer 1	12	00048–0004B	2B
Timer 2	13	0004C–0004F	2C
Serial receiver	14	00050–00053	3A
Serial transmitter	15	00054–00057	3B

Notes: Priority level 1 is the highest and 9 is the lowest. Some interrupts
have the same priority. Only the EB and EC models contain the serial port.

FIGURE 16–8 80186/80188
programmable interrupt
controller. (Courtesy of Intel
Corporation.)

controller. This controller appears in the XL, EA, and EB versions, but the EC version contains the exact equivalent to a pair of 8259As, as found in Chapter 12. In the EB version, the DMA inputs are replaced with inputs from the serial unit for receive and transmit.

The interrupt controller operates in two modes: master and slave mode. The mode is selected by a bit in the interrupt control register (EB and EC versions) called the *CAS bit*. If the CAS bit is a logic 1, the interrupt controller connects to external 8259A programmable interrupt controllers (see Figure 16–9); if CAS is a logic 0, the internal interrupt controller is selected. In many cases, there are enough interrupts within the 80186/80188, so the slave mode is not
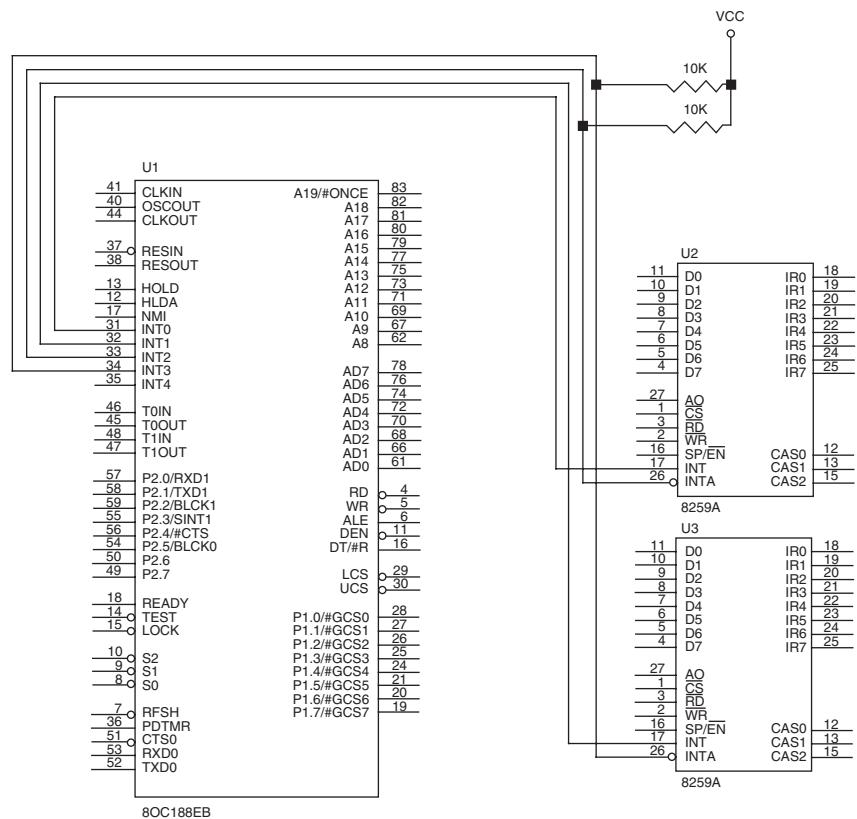
FIGURE 16–9 The interconnection between the 80C188EB and two 8259A programmable interrupt controllers. *Note:* Only the connections vital for this interface are shown.

normally used. In the XL and EA versions, the master and slave modes are selected in the peripheral control register at offset address FEH.

This portion of the text does not detail the programming of the interrupt controller. Instead, it is limited to a discussion of the internal structure of the interrupt controller. The programming and application of the interrupt controller is discussed in the sections that describe the timer and DMA controller.

Interrupt Controller Registers.   Figure 16–10 illustrates the interrupt controller's registers. These registers are located in the peripheral control block beginning at offset address 22H. For the EC version, which is compatible with the 8259A, the interrupt controller ports are at offset addresses 00H and 02H for the master and ports 04H and 06H for the slave. In the EB version, the interrupt controller is programmed at offset address 02H. Note that the EB version has an additional interrupt input (INT4).

Slave Mode.   When the interrupt controller operates in the slave mode, it uses up to two external 8259A programmable interrupt controllers for interrupt input expansion. Figure 16–9 shows how the external interrupt controllers connect to the 80186/80188 interrupt input pins for slave operation. Here, the INT_0 and INT_1 inputs are used as external connections to the interrupt request outputs of the 8259s, and $\overline{INTA0}(INT_2)$ and $\overline{INTA1}(INT_3)$ are used as interrupt acknowledge signals to the external controllers.

Interrupt Control Registers.   There are interrupt control registers in both modes of operation, which each control a single interrupt source. Figure 16–11 depicts the binary bit pattern of each of these interrupt control registers. The mask bit enables (0) or disables (1) the interrupt input

XL and EA Versions

3EH	INT3 Control Register
3CH	INT2 Control Register
3AH	INT1 Control Register
38H	INT0 Control Register
36H	DMA1 Control Register
34H	DMA0 Control Register
32H	Timer Control Register
30H	Interrupt Status
2EH	Request
2CH	In Service
2AH	PRIMSK
28H	Interrupt Masks
26H	POLL Status
24H	POLL
22H	EOI

EB Version

1EH	INT3 Control Register
1CH	INT2 Control Register
1AH	INT1 Control Register
18H	INT0 Control Register
16H	INT4 Control Register
14H	Serial Control Register
12H	Timer Control Register
10H	Interrupt Status
0EH	Request
0CH	In Service
0AH	PRIMSK
08H	Interrupt Masks
06H	POLL Status
04H	POLL
02H	EOI

FIGURE 16–10   The I/O offset port assignment for the interrupt control unit.

represented by the control word, and the priority bits set the priority level of the interrupt source. The highest priority level is 000, and the lowest is 111. The CAS bit is used to enable slave or cascade mode (0 enables slave mode), and the SFNM bit selects the **special fully nested mode**. The SFNM allows the priority structure of the 8259A to be maintained.

Interrupt Request Register.   The interrupt request register contains an image of the interrupt sources in each mode of operation. Whenever an interrupt is requested, the corresponding interrupt request bit becomes a logic 1, even if the interrupt is masked. The request is cleared whenever the 80186/80188 acknowledges the interrupt. Figure 16–12 illustrates the binary bit pattern of the interrupt request register for both the master and slave modes.

Mask and Priority Mask Registers.   The interrupt mask register has the same format as the interrupt register illustrated in Figure 16–12. If a source is masked (disabled), the corresponding bit of the interrupt mask register contains a logic 1; if enabled, it contains a logic 0. The interrupt mask register is read to determine which interrupt sources are masked and which are enabled. A source is masked by setting the source's mask bit in its interrupt control register.

The priority mask register, illustrated in Figure 16–13, shows the priority of the interrupt currently being serviced by the 80186/80188. The level of the interrupt is indicated by priority bits P_2–P_0. Internally, these bits prevent an interrupt by a lower priority source. These bits are automatically set to the next lower level at the end of an interrupt, as issued by the 80186/80188.
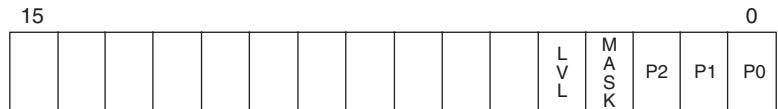
In-Service Register.   The in-service register has the same binary bit pattern as the request register of Figure 16–12. The bit that corresponds to the interrupt source is set if the 80186/80188 is currently acknowledging the interrupt. The bit is reset at the end of an interrupt.

The Poll and Poll Status Registers.   Both the interrupt poll and interrupt poll status registers share the same binary bit patterns as those illustrated in Figure 16–14. These registers have a bit

Timer and Serial Control Registers

INT2, INT3, and INT4 Control Registers

INT0 and INT1 Control Registers

P2–P0 = Priority Level
Mask = 0 enables interrupt
LVL = 0 = edge and 1 = level triggering
CAS = 1 selects slave mode
SFNM = 1 selects special fully nested mode
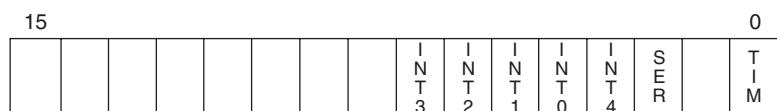
FIGURE 16–11 The interrupt control registers.

FIGURE 16–12 The interrupt request register.

Interrupt Request Register (EB version)

Interrupt Request Register (XL and EA versions)
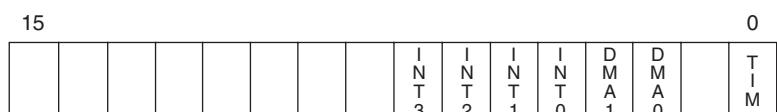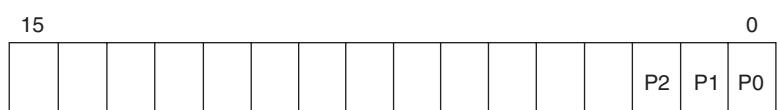
FIGURE 16–13 The priority mask register.

Priority Mask Register

P2–P0 = Priority Level

642

FIGURE 16–14 The poll and poll status registers.

Poll and Poll Status Registers

IREQ = 1 = Interrupt pending
VT4–VT0 = Interrupt type number of highest priority pending interrupt

FIGURE 16–15 The end-of-interrupt (EOI) register.

End-of-Interrupt Register

(INT REQ) that indicates an interrupt is pending. This bit is set if an interrupt is received with sufficient priority, and cleared when an interrupt is acknowledged. The S bits indicate the interrupt vector type number of the highest priority pending interrupt.

The poll and poll status registers may appear to be identical because they contain the same information. However, they differ in function. When the interrupt poll register is read, the interrupt is acknowledged. When the interrupt poll status register is read, no acknowledge is sent. These registers are used only in the master mode, not in the slave mode.

End-of-Interrupt Register.   The end-of-interrupt (EOI) register causes the termination of an interrupt when written by a program. Figure 16–15 shows the contents of the EOI register for both the master and slave mode.

In the master mode, writing to the EOI register ends either a specific interrupt level (vector number) or whatever level is currently active (nonspecific). In the nonspecific mode, the NSPEC bit must be set before the EOI register is written to end a nonspecific interrupt. The nonspecific EOI clears the highest level interrupt bit in the in-service register. The specific EOI clears the selected bit in the in-service register, which informs the microprocessor that the interrupt has been serviced and another interrupt of the same type can be accepted. The nonspecific mode is used unless there is a special circumstance that requires a different order for interrupt acknowledges. If a specific EOI is required, the vector number is placed in the EOI command. For example, to clear the timer 2 interrupt the EOI command is 13H (vector for timer 2).

In the slave mode, the level of the interrupt to be terminated is written to the EOI register. The slave mode does not allow a nonspecific EOI.

Interrupt Status Register.   The format of interrupt status register is depicted in Figure 16–16. In the master mode, T_2–T_0 indicates which timer (timer 0, timer 1, or timer 2) is causing an interrupt. This is necessary because all three timers have the same interrupt priority level. These bits are set when the timer requests an interrupt and are cleared when the interrupt is acknowledged. The DHLT (DMA halt) bit is only used in the master mode; when set, it stops a DMA action. Note that the interrupt status register is different for the EB version.

Interrupt Vector Register.   The interrupt vector register is present only in the slave mode, and only in the XL and EA versions at offset address 20H. It is used to specify the most significant five bits of the interrupt vector type number. Figure 16–17 illustrates the format of this register.

Timers

The 80186/80188 microprocessors contain three fully programmable 16-bit timers and each is totally independent of the others. Two of the timers (timer 0 and timer 1) have input and output

FIGURE 16–16 The interrupt status register.

Interrupt Status Register (XL and EA)

15   0

D H L T													T2	T1	T0

Interrupt Status Register (EB)

15   0

NMI											STX	SRX	T2	T1	T0

SRX = serial receiver
STX = serial transmitter
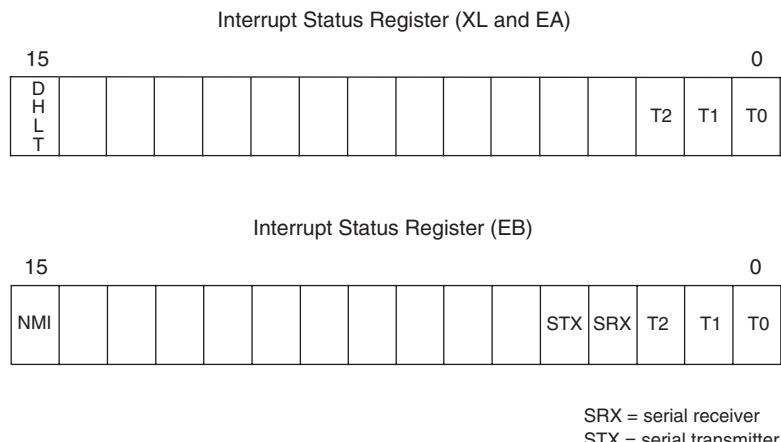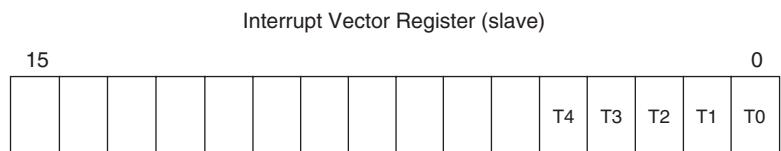
FIGURE 16–17 The interrupt vector register.

Interrupt Vector Register (slave)

15   0

										T4	T3	T2	T1	T0

pins that allow them to count external events or generate wave-forms. The third timer (timer 2) connects to the 80186/80188 clock. Timer 2 is used as a DMA request source, as a prescalar for other timers, or as a watchdog timer.

Figure 16–18 shows the internal structure of the timer unit. Notice that the timer unit contains one counting element that is responsible for updating all three counters. Each timer is actually a register that is rewritten from the counting element (a circuit that reads a value from a timer register and increments it before returning it). The counter element is also responsible for generating the outputs on the pins $T0_{OUT}$ and $T1_{OUT}$, reading the $T0_{IN}$ and $T1_{IN}$ pins, and causing a DMA request from the terminal count (TC) of timer 2 if timer 2 is programmed to request a DMA action.

Timer Register Operation.   The timers are controlled by a block of registers in the peripheral control block (see Figure 16–19). Each timer has a count register, maximum-count register or

FIGURE 16–18 Internal structure of the 80186/80188 timers. (Courtesy of Intel Corporation.)
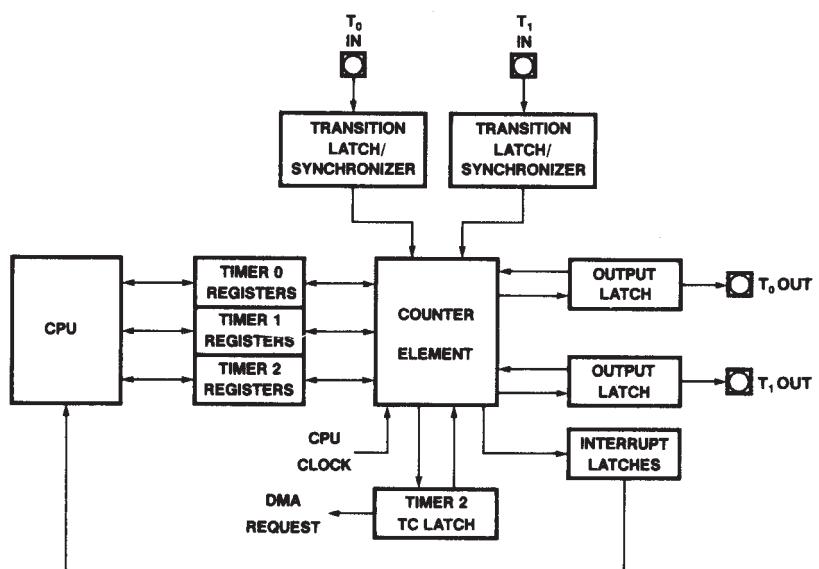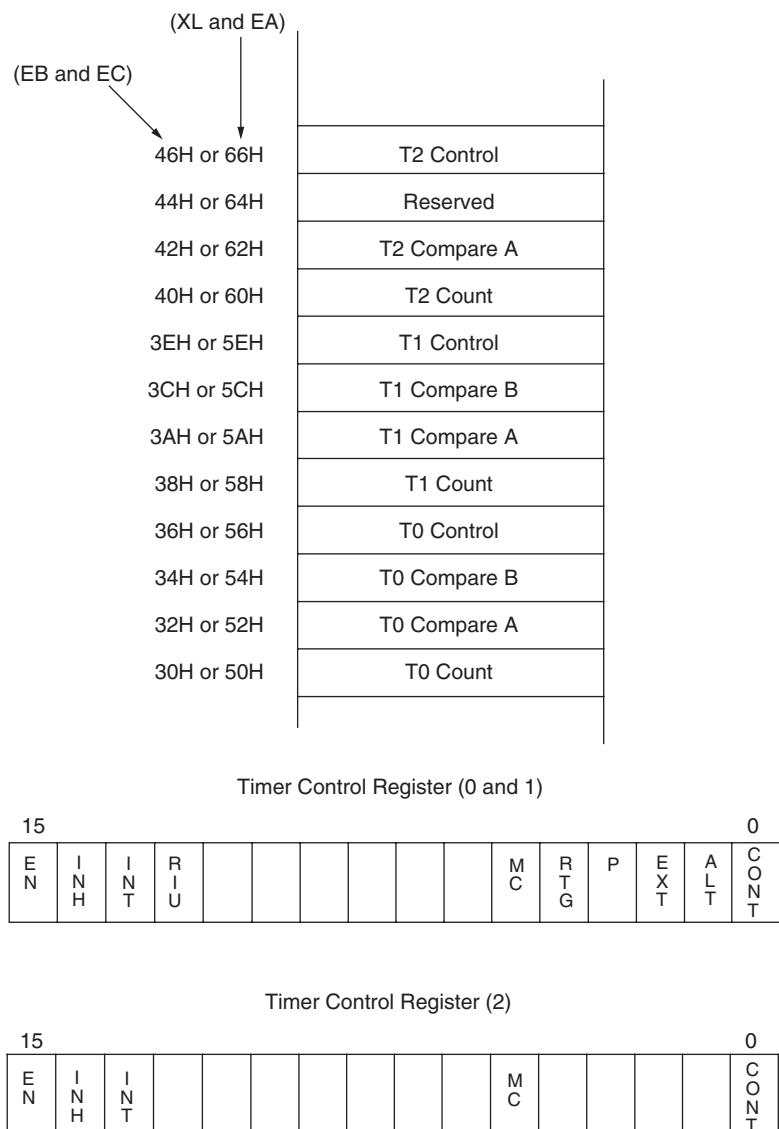
FIGURE 16–19 The offset locations and contents of the registers used to control the timers.

(XL and EA)

(EB and EC)

Offset	Contents
46H or 66H	T2 Control
44H or 64H	Reserved
42H or 62H	T2 Compare A
40H or 60H	T2 Count
3EH or 5EH	T1 Control
3CH or 5CH	T1 Compare B
3AH or 5AH	T1 Compare A
38H or 58H	T1 Count
36H or 56H	T0 Control
34H or 54H	T0 Compare B
32H or 52H	T0 Compare A
30H or 50H	T0 Count

Timer Control Register (0 and 1)

15 ... 0

| EN | INH | INT | RIU | | | | | MC | RTG | P | EXT | ALT | CONT |

Timer Control Register (2)

15 ... 0

| EN | INH | INT | | | | | | MC | | | | CONT |

registers, and a control register. These registers may all be read or written at any time because the 80186/80188 microprocessors ensure that the contents never change during a read or write.

The timer count register contains a 16-bit number that is incremented whenever an input to the timer occurs. Timers 0 and 1 are incremented at the positive edge on an external input pin, every fourth 80186/80188 clock, or by the output of timer 2. Timer 2 is clocked on every fourth 80186/80188 clock pulse and has no other timing source. This means that in the 8 MHz version of the 80186/80188, timer 2 operates at 2 MHz, and the maximum counting frequency of timers 0 and 1 is 2 MHz. Figure 16–20 depicts these four clocking periods, which are not related to the bus timing.

Each timer has at least one maximum-count register, called a **compare** register (compare register A for timers 0 and 1), which is loaded with the maximum count of the count register to generate an output. Note that a timer is an up counter. Whenever the count register is equal to the maximum-count compare register, it is cleared to 0. With a maximum count of 0000H, the counter counts 65,536 times. For any other value, the timer counts the true value of the count. For example, if the maximum count is 0002H, the counter will count from 0 to 1 and then be cleared to 0—a modulus 2 counter has two states.

1. Timer in 0 resolution time
2. Timer in 1 resolution time
3. Modified count value written into 80186 timer 0 count register
4. Modified count value written into 80186 timer
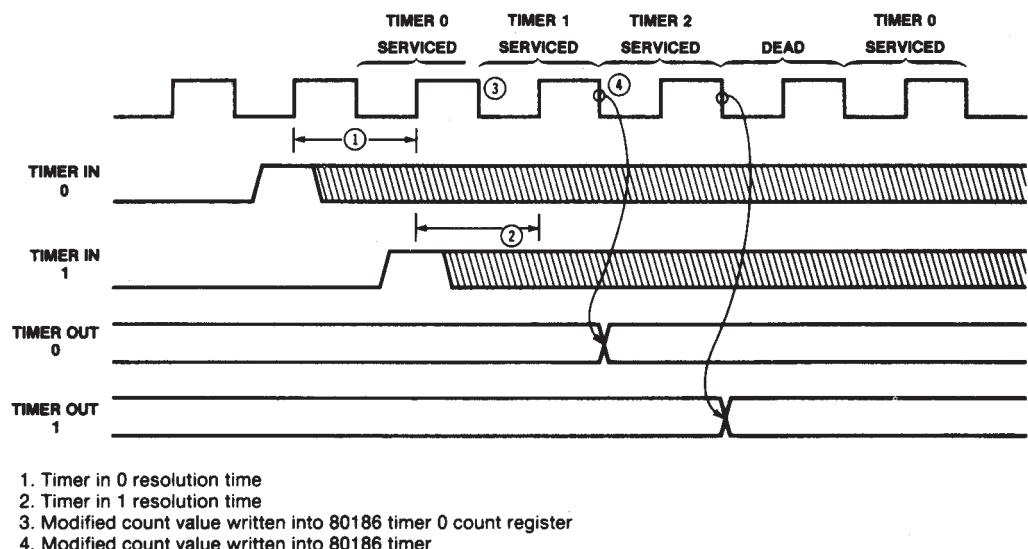
FIGURE 16–20   Timing for the 80186/80188 timers. (Courtesy of Intel Corporation.)

Timers 0 and 1 each have a second maximum-count compare register (compare register B) that is selected by the control register for the timer. Either maximum-count compare register A or both maximum-count compare registers A and B are used with these timers, as programmed by the ALT bit in the control register for the timer. When both maximum-count compare registers are used, the timer counts up to the value in maximum-count compare register A, clears to 0, and then counts up to the count in maximum-count compare register B. This process is then repeated. Using both maximum-count registers allows the timer to count up to 131,072.

The control register (refer to Figure 16–19) of each timer is 16 bits wide and specifies the operation of the timer. A definition of each control bit follows:

EN   The **enable** bit allows the timer to start counting. If EN is cleared, the timer does not count; if it is set, the timer counts.

INH   The **inhibit** bit allows a write to the timer control register to affect the enable bit (EN). If INH is set, then the EN bit can be set or cleared to control the counting. If INH is cleared, EN is not affected by a write to the timer control register. This allows other features of the timer to be modified without enabling or disabling the timer.

INT   The **interrupt** bit allows an interrupt to be generated by the timer. If INT is set, an interrupt will occur each time that the maximum count is reached in either maximum-count compare register. If this bit is cleared, no interrupt is generated. When the interrupt request is generated, it remains in force, even if the EN bit is cleared after the interrupt request.

RIU   The **register in use** bit indicates which maximum-count, compare register is currently in use by the timer. If RIU is a logic 0, then maximum-count compare register A is in use. This bit is a read-only bit, and writes do not affect it.

MC   The **maximum count** bit indicates that the timer has reached its maximum count. This bit becomes a logic 1 when the timer reaches its maximum count and remains a logic 1 until the MC bit is cleared by writing a logic 0. This allows the maximum count to be detected by software.

RTG   The **retrigger** bit is active only for external clocking (EXT = 0). The RTG bit is used only with timers 0 and 1 to select the operation of the timer input pins ($T0_{IN}$ and $T1_{IN}$).

If RTG is a logic 0, the external input will cause the timer to count if it is a logic 1; the timer will hold its count (stop counting) if it is a logic 0. If RTG is a logic 1, the external input pin clears the timer count to 0000H each time a positive-edge occurs.

P   The **prescalar** bit selects the clocking source for timers 0 and 1. If EXT = 0 and P = 0, the source is one fourth the system clock frequency. If P = 1, the source is timer 2.

EXT   The **external** bit selects internal timing (EXT = 0) or external timing (EXT = 1). If EXT = 1, the timing source is applied to the $T0_{IN}$ or $T1_{IN}$ pins. In this mode, the timer increments after each positive-edge on the timer input pin. If EXT = 0, the clocking source is from one of the internal sources.

ALT   The **alternate** bit selects single maximum-count mode (maximum-count compare register A) if a logic 0, or alternate maximum-count mode (maximum-count compare registers A and B) if a logic 1.

CONT   The **continuous** bit selects continuous operation if a logic 1. In continuous operation, the counter automatically continues counting after it reaches its maximum count. If CONT is a logic 0, the timer will automatically stop counting and clear the EN bit. Note that whenever the 80186/80188 are reset, the timers are automatically disabled.

Timer Output Pin.   Timers 0 and 1 have an output pin used to generate either square waves or pulses. To produce pulses, the timer is operated in single maximum-count mode (ALT = 0). In this mode, the output pin goes low for one clock period when the counter reaches its maximum count. By controlling the CONT bit in the control register, either a single pulse or continuous pulses can be generated.

To produce square waves or varying duty cycles, the alternate mode (ALT = 1) is selected. In this mode, the output pin is a logic 1 while maximum-count compare register A controls the timer; it is a logic 0 while maximum-count compare register B controls the timer. As with the single maximum-count mode, the timer can generate either a single square wave or continuous square waves. See Table 16–4 for the function of the ALT and CONT control bits.

Almost any duty cycle can be generated in the alternate mode. For example, suppose that a 10% duty cycle is required at a timer output pin. Maximum-count register A is loaded with a 10 and maximum-count register B is loaded with a 90 to produce an output that is a logic 1 for 10 clocks and a logic 0 for 90 clocks. This also divides the frequency of the timing source by a factor of 100.

Real-Time Clock Example.   Many systems require the time of day. This is often called a real-time clock (RTC). A timer within the 80186/80188 can provide the timing source for software that maintains the time of day.

The hardware required for this application is not illustrated because all that is required is that the $T1_{IN}$ pin must be connected to +5.0 V through a pull-up resistor to enable timer 1. In the example, timers 1 and 2 are used to generate a one-second interrupt that provides the software with a timing source.

The software required to implement a real-time clock is listed in Examples 16–2 and 16–3. Example 16–2 illustrates the software required to initialize the timers. Example 16–3 shows an interrupt service procedure, which keeps time. There is another procedure in Example 16–3 that increments a BCD modulus counter. None of the software required to install the interrupt vector and set or display time of day is illustrated here.

TABLE 16–4   The ALT and CONT bits in the timer control register.

ALT	CONT	Mode
0	0	Single pulse
0	1	Continuous pulses
1	0	Single square wave
1	1	Continuous square waves

EXAMPLE 16–2

```
;software is written for the 80186/80188 EB version that
;initializes and starts both timer 1 and 2.

;address equates

T2_CA   EQU   0FF42H              ;timer 2 compare A register
T2_CON  EQU   0FF46H              ;timer 2 control register
T2_CNT  EQU   0FF40H              ;timer 2 count register
T1_CA   EQU   0FF3AH              ;timer 1 compare A register
T1_CON  EQU   0FF38H              ;timer 1 control register
T1_CNT  EQU   0FF3EH              ;timer 1 count register

        MOV  AX,20000             ;program timer 2 for 10 msec
        MOV  DX,T2_CA
        OUT  DX,AX

        MOV  AX,100               ;program timer 1 for 1 sec
        MOV  DX,T1_CA
        OUT  DX,AX

        MOV  AX,0                 ;clear count registers
        MOV  DX,T2_CNT
        OUT  DX,AX
        MOV  DX,T1_CNT
        OUT  DX,AX

        MOV  AX,0C001H            ;enable timer 2 and start it
        MOV  DX,T2_CON
        OUT  DX,AX

        MOV  AX,0E009H            ;enable timer 1 with interrupt
        MOV  DX,T1_CON            ;and start it
        OUT  DX,AX
```

Timer 2 is programmed to divide by a factor of 20,000. This causes the clock (assuming a 2 MHz on the 8 MHz version of the 80186/80188) to be divided down to one pulse every 10 ms. The clock for timer 1 is derived internally from the timer 2 output. Timer 1 is programmed to divide the Timer 2 clock by 100 and generate a pulse once per second. The control register of timer 1 is programmed so that the one-second pulse internally generates an interrupt.

The interrupt service procedure is called once per second to keep time. The interrupt service procedure adds a one to the content of memory location SECONDS on each interrupt. Once every 60 seconds, the content of the next memory location (SECONDS + 1) is incremented. Finally, once per hour, the content of memory location SECONDS + 2 is incremented. The time is stored in these three consecutive memory locations in BCD, so the system software can easily access the time.

EXAMPLE 16–3

```
SECONDS      DB    ?
MINUTES      DB    ?
HOURS        DB    ?

INTRS PROC   FAR    USES DS AX SI

        MOV    AX,SEGMENT_ADDRESS
        MOV    DS,AX
        MOV    AH,60H              ;load modulus 60
        MOV    SI,OFFSET SECONDS   ;address clock
        CALL   UPS                 ;increment seconds
        .IF ZERO?                  ;if seconds became 0
             CALL UPS              ;increment minutes
             MOV  AH,24H           ;load modulus 24
             .IF ZERO?             ;if minutes became 0
                  CALL UPS         ;increment hours
```

```
                    .ENDIF
            .ENDIF
            MOV   DX,0FF02H            ;clear interrupt
            MOV   AX,8000H
            OUT   DX,AX

            RET

INTRS ENDP

UPS   PROC  NEAR

            MOV   AL,[SI]
            ADD   AL,1                 ;increment counter
            DAA                        ;make it BCD
            INC   SI
            .IF   AL == AH             ;test for modulus
                   MOV  AL,0
            .ENDIF
            MOV   [SI-1],AL
            RET

UPS   ENDP
```

DMA Controller

The DMA controller within the 80186/80188 has two fully independent DMA channels. Each has its own set of 20-bit address registers, so any memory or I/O location is accessible for a DMA transfer. In addition, each channel is programmable for auto-increment or auto-decrement to either source or destination registers. This controller is not available in the EB or EC versions. The EC version contains a modified four-channel DMA controller; the EB version contains no DMA controller. This text does not describe the DMA controller within the EC version.

Figure 16–21 illustrates the internal register structure of the DMA controller. These registers are located in the peripheral control block at offset addresses C0H–DFH.
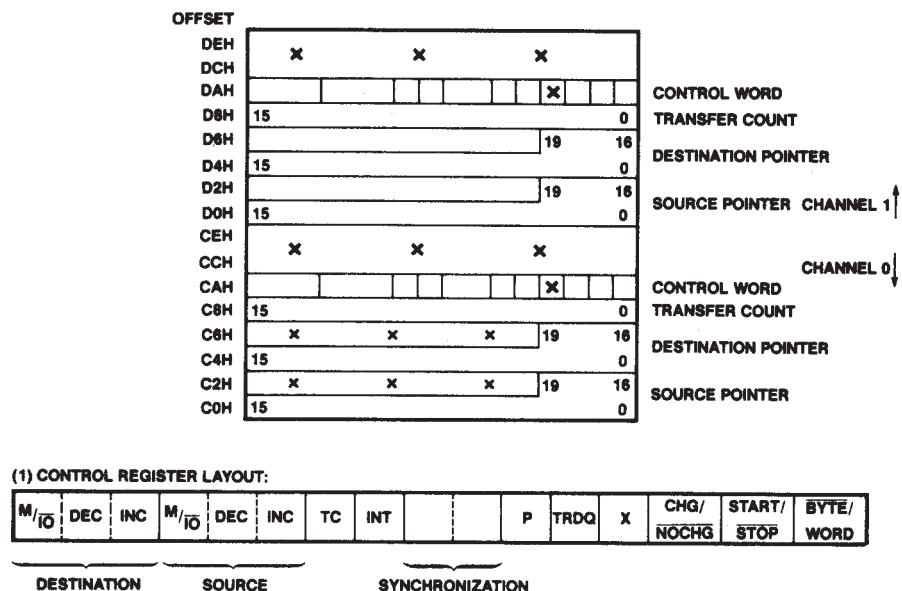
FIGURE 16–21 Register structure of the 80186/80188 DMA controller. (Courtesy of Intel Corporation.)

Notice that both DMA channel register sets are identical; each channel contains a control word, a source and destination pointer, and a transfer count. The transfer count is 16 bits wide and allows unattended DMA transfers of bytes (80188/80186) and words (80186 only). Each time that a byte or word is transferred, the count is decremented by one until it reaches 0000H—the terminal count.

The source and destination pointers are each 20 bits wide, so DMA transfers can occur to any memory location or I/O address without concern for segment and offset addresses. If the source or destination address is an I/O port, bits A_{19}–A_{16} must be 0000 or a malfunction may occur.

Channel Control Register.   Each DMA channel contains its own channel control register (refer to Figure 16–21), which defines its operation. The leftmost six bits specify the operation of the source and destination registers. The M/$\overline{\text{IO}}$ bit indicates a memory or I/O location, DEC causes the pointer to be decremented, and INC causes the pointer to be incremented. If both the INC and DEC bits are 1, then the pointer is unchanged after each DMA transfer. Notice that memory-to-memory transfers are possible with this DMA controller.

The TC (terminal count) bit causes the DMA channel to stop transfers when the channel count register is decremented to 0000H. If this bit is a logic 1, the DMA controller continues to transfer data, even after the terminal count is reached.

The INT bit enables interrupts to the interrupt controller. If set, the INT bit causes an interrupt to be issued when the terminal count of the channel is reached.

The SYN bit selects the type of synchronization for the channel: 00 = no synchronization, 01 = source synchronization, and 10 = destination synchronization. When either unsynchronized or source synchronization is selected, data are transferred at the rate of 2M bytes per second. These two types of synchronization allow transfers to occur without interruption. If destination synchronization is selected, the transfer rate is slower (1.3M bytes per second), and the controller relinquishes control to the 80186/80188 after each DMA transfer.

The P bit selects the channel priority. If P = 1, the channel has the highest priority. If both channels have the same priority, the controller alternates transfers between channels.

The TRDQ bit enables DMA transfers from timer 2. If this bit is a logic 1, the DMA request originates from timer 2. This can prevent the DMA transfers from using all of the microprocessor's time for the transfer.

The CHG/$\overline{\text{NOCHG}}$ bit determines whether START/$\overline{\text{STOP}}$ changes for a write to the control register. The START/$\overline{\text{STOP}}$ bit starts or stops the DMA transfer. To start a DMA transfer, both CHG/$\overline{\text{NOCHG}}$ and START/$\overline{\text{STOP}}$ are placed at a logic 1 level.

The BYTE/$\overline{\text{WORD}}$ selects whether the transfer is byte- or word-sized.

Sample Memory-to-Memory Transfer.   The built-in DMA controller is capable of performing memory-to-memory transfers. The procedure used to program the controller and start the transfer is listed in Example 16–4.

EXAMPLE 16–4

```
.MODEL SMALL
.186
.CODE
;
;Memory-to-Memory Transfer using DMA
;
;Calling Sequence for MOVES
;
;     DS:DI = source address
;     ED:DI = destination address
;     CX = Number of bytes
;
GETA  MACRO SEGA, OFFA, DMAA
      MOV   AX,SEGA              ;;get segment address
      SHL   AX,4                 ;;shift left 4 places
```

```
        ADD     AX,OFFA             ;;add in offset
        MOV     DX,DMAA             ;;address DMA controller
        OUT     DX,AX               ;;program address
        PUSHF
        MOV     AX,SEGA
        SHR     AX,12
        POPF
        ADC     AX,0
        ADD     DX,2
        OUT     DX,AX
        ENDM


MOVES   PROC    NEAR

        GETA    DS,SI,0FFC0H        ;program source address
        GETA    ES,DI,0FFC4H        ;program destination address

        MOV     DX,0FFC8H           ;program count
        MOV     AX,CX
        OUT     DX,AX
        MOV     DX,0FFCAH           ;program DMA control
        MOV     AX,0B606H
        OUT     DX,AX               ;start transfer

        RET

MOVES   ENDP
```

The procedure in Example 16–4 transfers data from the data segment location addressed by SI into the extra segment location addressed by DI. The number of bytes transferred is held in register CX. This operation is identical to the REP MOVSB instruction, but execution occurs at a much higher speed through the use of the DMA controller.

Chip Selection Unit

The chip selection unit simplifies the interface of memory and I/O to the 80186/80188. This unit contains programmable chip selection logic. In small- and medium-sized systems, no external decoder is required to select memory and I/O. Large systems, however, may still require external decoders. There are two forms of the chip selection unit; one form found in the XL and EA versions differs from the unit found in the EB and EC versions.

Memory Chip Selects.   Six pins (XL and EA versions) or 10 pins (EB and EC versions) are used to select different external memory components in a small- or medium-sized 80186/80188-based system. The $\overline{\text{UCS}}$ (**upper chip select**) pin enables the memory device located in the upper portion of the memory map that is most often populated with ROM. This programmable pin allows the size of the ROM to be specified and the number of wait states required. Note that the ending address of the ROM is FFFFFH. The $\overline{\text{LCS}}$ (**lower chip select**) pin selects the memory device (usually a RAM) that begins at memory location 00000H. As with the $\overline{\text{UCS}}$ pin, the memory size and number of wait states are programmable. The remaining four or eight chip select pins select middle memory devices. The four pins in the XL and EA version ($\overline{\text{MCS3}}-\overline{\text{MCS0}}$) are programmed for both the starting (base) address and memory size. Note that all devices must be of the same size. The eight pins ($\overline{\text{GCS7}}-\overline{\text{GCS0}}$) in the EB and EC versions are programmed by size and also by starting address. These selection signals represent a memory device or an I/O device.

Peripheral Chip Selects.   The 80186/80188 addresses up to seven external peripheral devices with pins $\overline{\text{PCS6}}-\overline{\text{PCS0}}$ (in the XL and EA versions). The $\overline{\text{GCS}}$ pins are used in the EB and EC versions to select up to eight memory or I/O devices. The base I/O address is programmed at any 1K-byte interval with port address block sizes of 128 bytes (64 bytes on the EB and EC versions).

TABLE 16–5 Wait state control bits R_2, R_1, and R_0 (XL and EA versions).

R_2	R_1	R_0	Number of Waits	READY Required
0	X	X	—	Yes
1	0	0	0	No
1	0	1	1	No
1	1	0	2	No
1	1	1	3	No

Programming the Chip Selection Unit for XL and EA Versions. The number of wait states in each section of the memory and the I/O are programmable. The 80186/80188 microprocessors have a built-in wait state generator that can introduce between 0 and 3 wait states (XL and EA version). Table 16–5 lists the logic levels required on bits R_2–R_0 in each programmable register to select various numbers of wait states. These three lines also select if an external READY signal is required to generate wait states. If READY is selected, the external READY signal is in parallel with the internal wait state generator. For example, if READY is a logic 0 for three clocking periods but the internal wait state generator is programmed to insert two wait states, three wait states are inserted.

Suppose that a 64K-byte EPROM is located at the top of the memory system and requires two wait states for proper operation. To select this device for this section of memory, the $\overline{UCS}$ pin is programmed for a memory range of F0000H–FFFFFH with two wait states. Figure 16–22

FIGURE 16–22 The chip selection registers for the XL and EA versions of the 80186/80188.

Upper Memory Control Register (offset = A0H)

15															0
0	0	A17	A16	A15	A14	A13	A12	A11	A10	0	0	0	R2	R1	R0

Lower Memory Control Register (offset = A2H)

15															0
0	0	A17	A16	A15	A14	A13	A12	A11	A10	0	0	0	R2	R1	R0

Peripheral Chip Select Base Address (offset = A4H)

15															0
A19	A18	A17	A16	A15	A14	A13	0	0	0	0	0	0	R2	R1	R0

Mid-range Memory Base Address (offset = A6H)

15															0
A19	A18	A17	A16	A15	A14	A13	0	0	0	0	0	0	R2	R1	R0

Mid-range Memory Size (offset = A8H)

15															0
0	M6	M5	M4	M3	M2	M1	M0	EX	MS	0	0	0	R2	R1	R0

TABLE 16–6 Upper memory programming for the A$_0$H register (XL and EA versions).

Start Address	Block Size	Value for No Waits, No READY
FFC00H	1K	3FC4H
FF800H	2K	3F84H
FF000H	4K	3F04H
FE000H	8K	3E04H
FC000H	16K	3C04H
F8000H	32K	3804H
F0000H	64K	3004H
E0000H	128K	1004H
C0000H	256K	0004H

lists the control registers for all memory and I/O selections in the peripheral control block at offset addresses A$_0$–A$_9$H. Notice that the rightmost three bits of these control registers are from Table 16–5. The control register for the upper memory area is at location PCB offset address A$_0$H. This 16-bit register is programmed with the starting address of the memory area (F0000H, in this case) and the number of wait states. Please note that the upper two bits of the address must be programmed as 00, and that only address bits A$_{17}$–A$_{10}$ are programmed into the control register. See Table 16–6 for examples illustrating the codes for various memory sizes. Because our example requires two wait states, the basic address is the same as in the table for a 64K device, except that the rightmost three bits are 110 instead of 100. The datum sent to the upper memory control register is 3006H.

Suppose that a 32K-byte SRAM that requires no waits and no READY input is located at the bottom of the memory system. To program the $\overline{LCS}$ pin to select this device, register A$_2$ is loaded in exactly the same manner as register A$_0$H. In this example, a 07FCH is sent to register A$_2$H. Table 16–7 lists the programming values for the lower chip-selection output.

The central part of the memory is programmed via two registers: A$_6$H and A$_8$H. Register A$_6$H programs the beginning or base address of the middle memory select lines ($\overline{MCS3}$–$\overline{MCS0}$) and number of waits. Register A$_8$H defines the size of the block of memory and the individual memory device size (see Table 16–8). In addition to block size, the number of peripheral wait states are programmed as with other areas of memory. The EX (bit 7) and MS (bit 6) specify the peripheral selection lines, and will be discussed shortly.

For example, suppose that four 32K-byte SRAMs are added to the middle memory area, beginning at location 80000H and ending at location 9FFFFH with no wait states. To program the middle memory selection lines for this area of memory, we place the leftmost seven address bits in

TABLE 16–7 Lower memory programming for the A$_2$H register (XL and EA versions).

Ending Address	Block Size	Value for No Waits, No READY
003FFH	1K	0004H
007FFH	2K	0044H
00FFFH	4K	00C4H
01FFFH	8K	01C4H
03FFFH	16K	03C4H
07FFFH	32K	07C4H
0FFFFH	64K	0FC4H
1FFFFH	128K	1FC4H
3FFFFH	256K	3FC4H

TABLE 16–8 Middle memory programming for the A_8H register (XL and EA versions).

Block Size	Chip Size	Value for No Waits, No READY, and EX = 0, MS = 1
8K	2K	0144H
16K	4K	0344H
32K	8K	0744H
64K	16K	0F44H
128K	32K	1F44H
256K	64K	3F44H
512K	128K	7F44H

register A_6H, with bits 8–3 containing logic 0s, and the rightmost three bits containing the ready control bits. For this example, register A_6H is loaded with 8004H. Register A_8H is programmed with 1F44H, assuming that EX = 0 and MS = 1 and no wait states and no READY are required for the peripherals.

Register A_4H programs the peripheral chip selection pins ($\overline{PCS6}$–$\overline{PCS0}$) along with the EX and MS bits of register A_8H. Register A_4H holds the beginning or base address of the peripheral selection lines. The peripherals may be placed in memory or in the I/O map. If they are placed in the I/O map, A_{19}–A_{16} of the port number must be 0000. Once the starting address is programmed on any 1K-byte I/O address boundary, the $\overline{PCS}$ pins are spaced at 128-byte intervals.

For example, if register A_4H is programmed with a 0204H, with no waits and no READY synchronization, the memory address begins at 02000H or the I/O port begins at 2000H. In this case, the I/O ports are: $\overline{PCS0}$ = 2000H, $\overline{PCS1}$ = 2080H, $\overline{PCS2}$ = 2100H, $\overline{PCS3}$ = 2180H, $\overline{PCS4}$ = 2200H, $\overline{PCS5}$ = 2280H, and $\overline{PCS5}$ = 2300H.

The MS bit of register A_8H selects memory mapping or I/O mapping for the peripheral select pins. If MS is a logic 0, then the $\overline{PCS}$ lines are decoded in the memory map; if it is a logic 1, then the $\overline{PCS}$ lines are in the I/O map.

The EX bit selects the function of the $\overline{PCS5}$ and $\overline{PCS6}$ pins. If EX = 1, these $\overline{PCS}$ pins select I/O devices; if EX = 0, these pins provide the system with latched address lines A_1 and A_2. The A_1 and A_2 pins are used by some I/O devices to select internal registers and are provided for this purpose.

Programming the Chip Selection Unit for EB and EC Versions.   As mentioned earlier, the EB and EC versions have a different chip selection unit. These newer versions of the 80186/80188 contain an upper and lower memory chip selection pin as do earlier versions, but they do not contain middle selection and peripheral selection pins. In place of the middle and peripheral chip selection pins, the EB and EC versions contain eight general chip selection pins ($\overline{GCS7}$–$\overline{GCS0}$) that select either a memory device or an I/O device.

Programming is also different because each of the chip selection pins contains a starting address register and an ending address register. See Figure 16–23 for the offset address of each pin and the contents of the start and end registers.

Notice that programming for the EB and EC versions of the 80186/80188 is much easier than for the earlier XL and EA versions. For example, to program the $\overline{UCS}$ pin for an address that begins at location F0000H and ends at location FFFFFH (64K bytes), the starting address register (offset = A4H) is programmed with F002H for a starting address of F0000H with two wait states. The ending address register (offset = A6H) is programmed with 000EH for an ending address of FFFFFH for memory with no external ready synchronization. The other chip selection pins are programmed in a similar fashion.

Address	Function
A6H	$\overline{\text{UCS}}$ stop
A4H	$\overline{\text{UCS}}$ start
A2H	$\overline{\text{LCS}}$ stop
A0H	$\overline{\text{LCS}}$ start
9EH	$\overline{\text{CGS7}}$ stop
9CH	$\overline{\text{CGS7}}$ start
9AH	$\overline{\text{GCS6}}$ stop
98H	$\overline{\text{GCS6}}$ start
96H	$\overline{\text{GCS5}}$ stop
94H	$\overline{\text{GCS5}}$ start
92H	$\overline{\text{GCS4}}$ stop
90H	$\overline{\text{GCS4}}$ start
8EH	$\overline{\text{GCS3}}$ stop
8CH	$\overline{\text{GCS3}}$ start
8AH	$\overline{\text{GCS2}}$ stop
88H	$\overline{\text{GCS2}}$ start
86H	$\overline{\text{GCS1}}$ stop
84H	$\overline{\text{GCS1}}$ start
82H	$\overline{\text{GCS0}}$ stop
80H	$\overline{\text{GCS0}}$ start

Start register

15															0
A19	A18	A17	A16	A15	A14	A13	A12	A11	A10	0	0	WS3	WS2	WS1	WS0

Stop register

15															0
A19	A18	A17	A16	A15	A14	A13	A12	A11	A10	0	0	CSEN	ISTOP	MEM	RDY

Notes: A19–A10 are memory address A19–A10 or I/O address bits A15–A6.
WS3–WS0 select between 0 and 15 wait states.
CSEN enables the pin if CSEN = 1.
ISTOP = if ISTOP = 1 the memory address is OFFFFFH or the I/O address is OFFFFH.
MEM = MEM = 1 selects memory and MEM = 0 selects I/O.
RDY = enables external ready if RDY = 1 for more than 15 wait states.

FIGURE 16–23 The chip selection unit in the EB and EC versions of the 80186/80188.
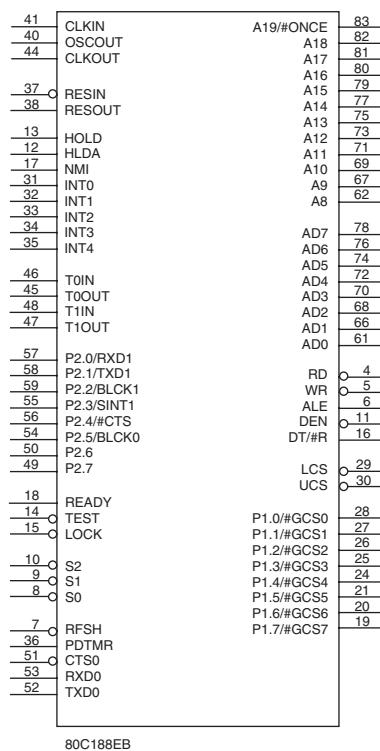
16–3   80C188EB EXAMPLE INTERFACE

Because the 80186/80188 microprocessors are designed as embedded controllers, this section of the text provides an example of such an application. The example illustrates simple memory and I/O attached to the 80C188EB microprocessor. It also lists the software required to program the 80C188EB and its internal registers after a system reset. Figure 16–24 illustrates the pin-out of the 80C188EB version of the 80188 microprocessor. Notice the differences between this version and the XL version presented earlier in the text.

The 80C188EB version contains some new features that were not present on earlier versions. These features include two I/O ports (P_1 and P_2) that are shared with other functions and two serial communications interfaces that are built into the processor. This version does not contain a DMA controller, as did the XL version.

The 80188EB can be interfaced with a small system designed to be used as a microprocessor trainer. The trainer illustrated in this text uses a 27256 EPROM for program storage, a 62256 SRAM for data storage, and an 8255 for a keyboard and LCD display interface. Figure 16–25 illustrates a small microprocessor trainer that is based on the 80C188EB microprocessor.

FIGURE 16–24 The pin-out of the 80C188EB version of the 80188 microprocessor.

Memory is selected by the $\overline{\text{UCS}}$ pin for the 27C256 EPROM and the $\overline{\text{LCS}}$ pin for the 62256 SRAM; the $\overline{\text{GCS0}}$ pin selects the 8255. The system software places the EPROM at memory addresses F8000H–FFFFFH; the SRAM at 00000H–07FFFH; and the 8255 at I/O ports 0000H–003FH (software uses ports 0, 1, 2, and 3). In this system, as is normally the case, we do not modify the address of the peripheral control block, which resides at I/O ports FF00H–FFFFH.

Example 16–5 lists the software required to initialize the 80C188EB microprocessor. This example completely programs the 80C188EB and also the entire system. The software is discussed in the next section of this chapter.

EXAMPLE 16–5

```
;Simple real-time operating test system for an 80188EB.

; INT 40H delays for BL milliseconds (range 1 to 99)
; INT 41H delays for BL seconds (range 1 to 59)
; Note* delay times must be written in hexadecimal!
; ie. 15 milliseconds is 15H.
; INT 42H displays character string on LCD
;       ES:BX addresses the NULL string
;       AL = where (80H line 1, C0H line2)
; INT 43H clears the LCD
; INT 44H reads a key from the keypad; AL = keycode

.MODEL TINY
.186              ;switch to the 80186/80188 instruction set
.CODE
.STARTUP

;program for the 80188EB microprocessor trainer.
;USE MASM 6.11
;command line = ML /AT FILENAME.ASM
```
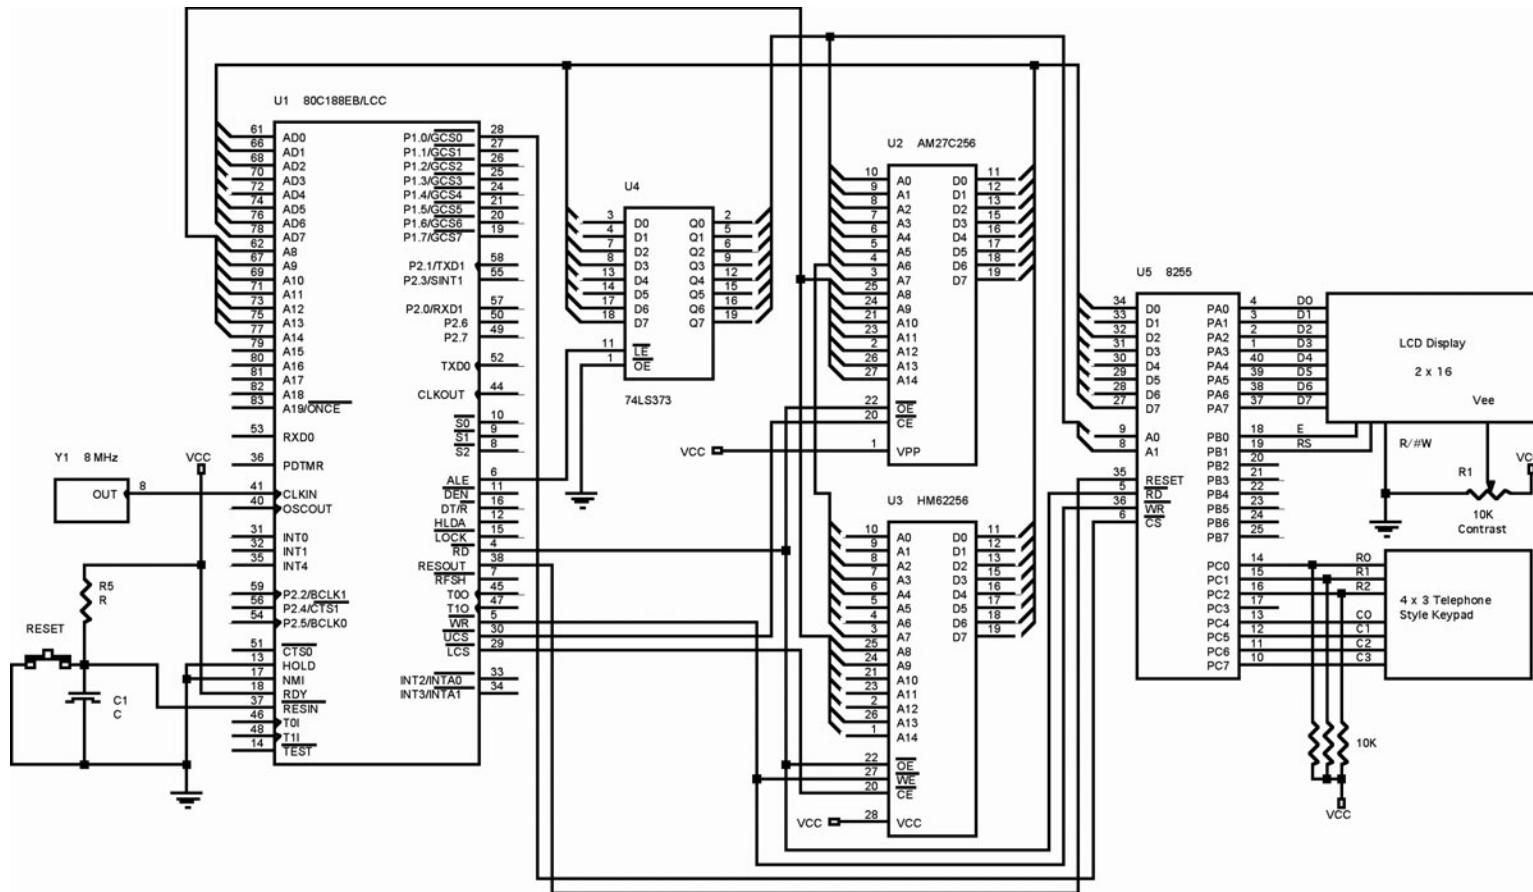
FIGURE 16–25 A sample 80C188EB-based microprocessor system.

```
;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;
;MACROs placed here
;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;

IO      MACRO   PORT,DATA
        MOV     DX,PORT
        MOV     AX,DATA
        OUT     DX,AL               ;;AL is more efficient
        ENDM

CS_IO   MACRO   PORT,START,STOP
        IO      PORT,START
        IO      PORT+2,STOP
        ENDM

SEND    MACRO   VALUE,COMMAND,DELAY
        MOV     AL,VALUE
        OUT     0,AL
        MOV     AL,COMMAND
        OUT     1,AL
        OR      AL,1
        OUT     1,AL
        AND     AL,2
        OUT     1,AL
        PUSH    BX
        MOV     BL,DELAY
        INT     40H
        POP     BX
        ENDM

BUT     MACRO
        IN  AL,2                     ;;test for key
        OR  AL,0F8H
        CMP AL,0FFH
        ENDM

;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;
;Initialization placed here
;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;

        IO      0FFA6H,000EH        ;UCS Stop address
        CS_IO   0FFA0H,0,80AH       ;program LCS
        CS_IO   0FF80H,0,48H        ;program GCS0
        IO      0FF54H,1            ;Port 1 control
        IO      0FF5CH,0            ;Port 2 control
        IO      0FF58H,00FFH        ;port 2 direction
        IO      3,81H               ;program 8255
        MOV     AX,0                ;address segment 0000
        MOV     DS,AX               ;for DS, ES, and SS
        MOV     ES,AX
        MOV     SS,AX
        MOV     SP,8000H ;setup stack pointer (0000:8000)

        MOV     BX,OFFSET INTT-100H        ;install interrupt vectors
        .WHILE WORD PTR CS:[BX] != 0
                MOV     AX,CS:[BX]
                MOV     DI,CS:[BX+2]
                MOV     DS:[DI],AX
                MOV     DS:[DI+2],CS
                ADD     BX,4
        .ENDW

        MOV     BYTE PTR DS:[40FH],0    ;don't display time
        IO      0FF40H,0            ;Timer 2 count
        IO      0FF42H,1000         ;Timer 2 compare
        IO      0FF46H,0E001H       ;Timer 2 control
        IO      0FF08H,00FCH        ;interrupt mask
        MOV     AL,0
```

```
        OUT     1,AL                ;place E at 0 for LCD
        STI                         ;enable interrupts
        CALL    INIT                ;initialize LCD

;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;
                        ;SYSTEM SOFTWARE HERE
;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;

;The following instructions are temporary to test the system and
;are replaced when a new system is created.

        MOV     BYTE PTR DS:[40FH],0FFH      ;set up to display time
        MOV     WORD PTR DS:[40CH],0         ;clear clock to 00:00:00 AM
        MOV     BYTE PTR DS:[40EH],0
        MOV     AX,CS                        ;line 1 message
        MOV     ES,AX
        MOV     AL,80H
        MOV     BX,OFFSET MES1 - 100H
        INT     42h

;System software placed here

        .WHILE 1                ;end of system loop
        .ENDW


;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;
;procedures and data follow the system software
;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;

MES1            DB          'The 80188 rules!',0

;Interrupt vector table

INTT    DW      TIM2-100H               ;interrupt procedure
        DW      13H * 4                 ;vector address
        DW      DELAYM-100H
        DW      40H * 4
        DW      DELAYS-100H
        DW      41H * 4
        DW      STRING-100H
        DW      42H * 4
        DW      CLEAR-100H
        DW      43H * 4
        DW      KEY-100H
        DW      44H * 4
        DW      0                       ;end of table

;Interrupt service procedure for TIMER 2 interrupt. (Once per millisecond)

TIM2    PROC    FAR USES ES DS AX BX SI DX

        MOV     AX,0
        MOV     DS,AX
        MOV     ES,AX
        MOV     BX,409H                 ;address real-time clock - 1
        MOV     SI,OFFSET MODU-101H     ;address Modulus Table - 1
        .REPEAT
                INC     SI              ;point to modulus
                INC     BX              ;point to counter
                MOV     AL,[BX]         ;get counter
                ADD     AL,1            ;add 1
                DAA                     ;make it BCD
                .IF AL == BYTE PTR CS:[SI] ;check modulus
                        MOV     AL,0
                .ENDIF
                MOV     [BX],AL    ;save new count
        .UNTIL !ZERO? || BX == 40FH
        IO      0FF02H,8000H                 ;end of interrupt
        .IF     BYTE PTR DS:[40AH] == 0 && BYTE PTR DS:[40BH] == 0
                CALL    DISPLAY    ;start Display Thread
```

```
        .ENDIF
        IRET

TIM2    ENDP

MODU    DB      0                           ;Mod 100
        DB      10H                         ;Mod 10
        DB      60H                         ;Mod 60
        DB      60H                         ;Mod 60
        DB      24H                         ;Mod 24

DISPLAY PROC NEAR     ;display time of day (once a second)

        .IF     BYTE PTR DS:[40FH] != 0         ;if display time is not zero
                STI                             ;enable future interrupts
                MOV     BX,3F0H
                MOV     SI,40EH                 ;address clock
                MOV     AL,[SI]                 ;get hours
                .IF AL > 12H                    ;for AM / PM
                        SUB     AL,12H
                        DAS
                .ELSEIF AL == 0
                        MOV     AL,12H
                .ENDIF
                CALL    STORE
                MOV     BYTE PTR [BX],':'
                INC     BX
                MOV     AL,[SI]                 ;get minutes
                CALL    STORE
                MOV     BYTE PTR [BX],':'
                INC     BX
                MOV     AL,[SI]                 ;get seconds
                CALL    STORE
                MOV     DL,'A'                  ;for AM / PM
                .IF     BYTE PTR DS:[40EH] > 11H
                        MOV     DL,'P'
                .ENDIF
                MOV     BYTE PTR [BX],' '
                MOV     [BX+1],DL
                MOV     BYTE PTR [BX+2],'M'
                MOV     BYTE PTR [BX+3],0       ;end of string
                MOV     BX,3F0H                 ;display buffer
                MOV     AL,0C2H                 ;LCD starting position
                INT     42H
        .ENDIF
        RET

DISPLAY ENDP

STORE   PROC    NEAR

        PUSH    AX
        SHR     AL,4
        MOV     DL,AL
        POP     AX
        AND     AL,15
        MOV     DH,AL
        ADD     DX,3030H
        MOV     [BX],DX
        ADD     BX,2
        DEC     SI
        RET

STORE   ENDP

DELAYM  PROC    FAR USES DS BX AX

        STI                     ;enable future interrupts
        MOV     AX,0
        MOV     DS,AX
```

```
        MOV     AL,DS:[40AH]        ;get milli counter
        ADD     AL,BL               ;BL = no. of milliseconds
        DAA
        .REPEAT
        .UNTIL AL == DS:[40AH]
        IRET

DELAYM  ENDP

DELAYS  PROC    FAR USES DS BX AX

        STI                         ;enable future interrupts
        MOV     AX,0
        MOV     DS,AX
        MOV     AL,DS:[40CH]        ;get seconds
        ADD     AL,BL
        DAA
        .IF AL >= 60H
                SUB   AL,60H
                DAS
        .ENDIF
        .REPEAT
        .UNTIL AL == DS:[40CH]
        IRET

DELAYS  ENDP

INIT    PROC    NEAR

        MOV     BL,30H              ;wait 30 milliseconds
        INT     40H
        MOV     CX,4
        .REPEAT
                SEND  38H,0,6
        .UNTILCXZ
        SEND    8,0,2
        SEND    1,0,2
        SEND    12,0,2
        SEND    6,0,2
        RET

INIT_LCD ENDP

STRING  PROC    FAR USES BX AX    ;display string

        STI                              ;enable future interrupts
        SEND    AL,0,1                   ;send start position
        .REPEAT
                SEND  BYTE PTR ES:[BX],2,1
                INC   BX
        .UNTIL BYTE PTR ES:[BX] == 0
        IRET

STRING  ENDP

CLEAR   PROC    FAR USES AX BX    ;clear LCD

        STI                              ;enable future interrupts
        SEND    1,0,2
        IRET
CLEAR   ENDP

KEY     PROC    FAR USES BX ;read key

        STI
        MOV  AL,0     ;clear C0 through C3
        OUT  2,AL
        .REPEAT       ;wait for key release
                .REPEAT
                        BUT
                .UNTIL ZERO?
```

```
                        MOV     BL,12H      ;time delay
                        INT 40H
                        BUT
                .UNTIL ZERO?
                .REPEAT                    ;wait for key press
                        .REPEAT
                                BUT
                        .UNTIL !ZERO?
                        MOV     BL,12H      ;time delay
                        INT 40H
                        BUT
                .UNTIL !ZERO?
                MOV     BX,0FDEFH
                .REPEAT
                        MOV  AL,BL
                        OUT  2,AL
                        ADD  BH,3
                        ROL  BL,1
                        BUT
                .UNTIL !ZERO?
                .WHILE 1
                        SHR AL,1
                        .BREAK .IF !CARRY?
                        INC BH
                .ENDW
                MOV     BX,OFFSET LOOK-100H
                XLAT CS:LOOK      ;specify code segment (EPROM)

                IRET

KEY     ENDP

LOOK    DB      3,2,1
        DB      6,5,4
        DB      9,8,7
        DB      10,0,11


;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;
;ANY OTHER SUBROUTINES OR DATA THAT ARE NEEDED SHOULD GO HERE
;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;

        ORG     080F0H                  ;get to reset location
RESET:
        IO      0FFA4H,0F800H           ;UCS start address
        DB      0EAH                    ;JMP   F800:0000    (F8000H)
        DW      0000H,0F800H
END
```
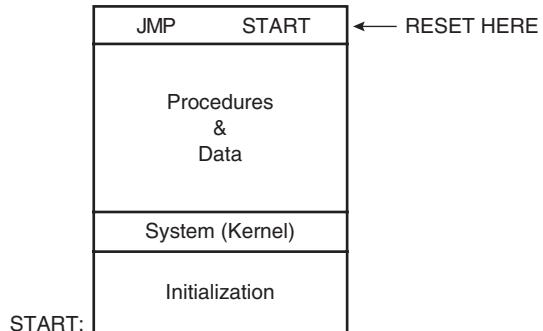
16–4 REAL-TIME OPERATING SYSTEMS (RTOS)

This section of the text describes the real-time operating system (RTOS). Interrupts are used to develop RTOSs because they are used in embedded applications of the microprocessor. All systems, from the simplest embedded application the most sophisticated system, must have an operating system.

What Is a Real-Time Operating System (RTOS)?

The RTOS is an operating system used in embedded applications that performs tasks in a predictable amount of time. Operating systems, like Windows, defer many tasks and do not guarantee their execution in a predictable time. The RTOS is much like any other operating system in that it contains the same basic sections. Figure 16–26 illustrates the basic structure of an operating system as it might be placed on an EPROM or flash memory device.

FIGURE 16–26 The structure of an RTOS operating system.

```
                  ┌──────────────────────┐
                  │ JMP        START     │  ←─── RESET HERE
                  ├──────────────────────┤
                  │                      │
                  │     Procedures       │
                  │         &            │
                  │       Data           │
                  │                      │
                  ├──────────────────────┤
                  │   System (Kernel)    │
                  ├──────────────────────┤
                  │   Initialization     │
          START:  └──────────────────────┘
```

There are three components to all operating systems: (1) initialization, (2) the kernel, (3) data and procedures. If Example 16–5 (last section) is compared to the Figure 16–26, all three sections will be seen. The initialization section is used to program all hardware components in the system, load drivers specific to a system, and program the contents of the microprocessor's registers. The kernel performs the basic system task, provides system calls or functions, and comprises the embedded system. The data and procedure section holds all procedures and any static data used by the operating system.

The RESET Section. The last part of the software in Example 16–5 shows the reset block of the RTOS. The ORG statement places the reset instructions at a location that is 16 bytes from the end of the memory device. In this case the EPROM is 32K bytes, which means it begins at 0000H and ends at 7FFFH. Recall that a 32K device has 15 address pins. The $\overline{CS}$ input selects the EPROM for location F8000H through FFFFFH in the system. The ORG statement in the program places the origin of the reset section at location 80F0H because all tiny model (.COM) programs are assembled from offset address 100H even though the first byte of the program is the first byte stored in the file. Because of this bias, all the addresses on the EPROM must be adjusted by 100H as is the ORG statement.

Only 16 bytes of memory exist for the reset instruction because the reset location is FFFF0H in the system. In this example there is only enough room to program the $\overline{UCS}$ starting address as F8000H before a jump to the start of the EPROM. Far jumps are not allowed in the tiny model, so it was forced by storing the actual hexadecimal opcode for a far jump (EAH).

Initialization Section. The initialization section of Example 16–5 begins in the reset block and continues at the start of the EPROM. If the initialization section is viewed, all of the programmable devices in the system are programmed and the segment registers are loaded. The initialization section also programs timer 2 so it causes an interrupt to the TIM_2 procedure each millisecond. The TIM_2 interrupt service procedure updates the clock once per second and is also the basis of precision time delays in the software.

The Kernel. The kernel in Example 16–5 is very short, because the system is incomplete and serves only as a test system. In this example all that the system does is display a sign-on message and display the time of day on the second line of the LCD. Once this is accomplished, the system ends at an infinite WHILE loop. All system programs are infinite loops unless they crash.

An Example System

Figure 16–27 illustrates a simple embedded system based on the 80188EB embedded microprocessor. This schematic depicts only the parts added to Figure 16–25 in order to read a temperature from the LM-70. This system contains a 2-line × 16 character-per-line LCD display that shows the time of day and the temperature. The system itself is stored on a small 32K × 8 EPROM. A 32K × 8 SRAM is included to act as a stack and store the time. A database holds the most recent temperatures and the times at which the temperatures were obtained.

FIGURE 16–27 Additional circuitry for Figure 16–25 so it can read a temperature.
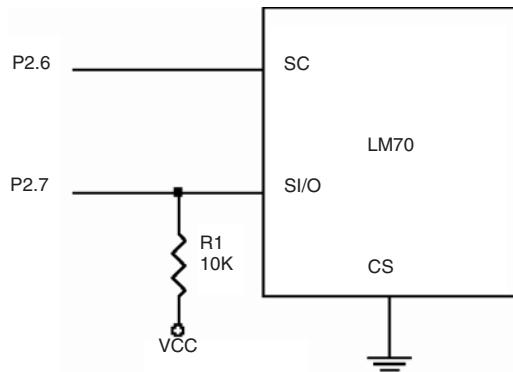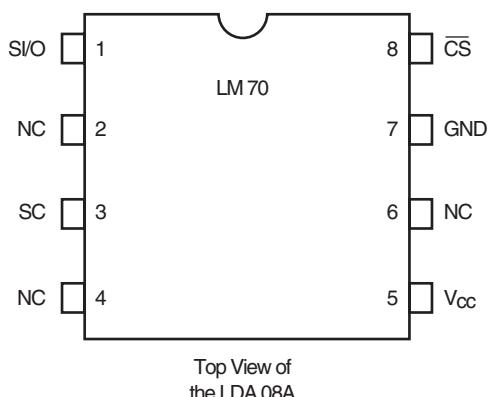
FIGURE 16–28 The LM_{70} temperature sensor.

Top View of
the LDA 08A

The temperature sensor is located inside the LM_{70} digital temperature sensor manufactured by National Semiconductor Corporation for less than \$1.00. The interface to the microprocessor is in serial format, and the converter has a resolution of 10 bits plus a sign bit. Figure 16–28 illustrates the pin-out of the LM_{70} temperature sensor.

The LM_{70} transfers data to and from the microprocessor through the SIO pin, which is a bidirectional serial data pin. Information is clocked through the SIO pin by the SC (clock) pin. The LM_{70} contains three 16-bit registers: the configuration register, the temperature sensor register, and the identification register. The configuration register selects either the shutdown mode (XXFFH) or continuous conversion mode (XX00). The temperature register contains the signed temperature in the leftmost 11 bits of the 16-bit data word. If the temperature is negative, it is in their respective complement forms. The identification register presents an 8100 when it is read.

When the temperature is read from the LM_{70}, it is read in Celsius and each step is equal to 0.25°C. For example, if the temperature register is 0000 1100 100X XXXX or a value of 100 decimal, the temperature is 25.0°C.

Example 16–6 illustrates the software added to the operating system listed in Example 16–5. The system samples the temperatures once per minute and stores them in a circular queue along with the day and the time in hours and minutes. The day is a number that starts at zero when the system is initialized. The size of the queue has been set to 16K bytes, so the most recent 4,096 measurements can be stored. In this example the keyboard is not used, but some of the system calls are used to display the temperature on line 1 of the display. The real-time clock is also interrogated to determine the start of each minute so a sample can occur. The software in the listing replaces the software section in Example 16–5 where it states, ";System software placed here". This software replaces the infinite WHILE loop in the example.

The LM$_{70}$ is initialized by output 16 bits of 0s to it and then read by reading all 16 bits of the temperature. The reading of the LM$_{70}$ is accomplished in the software by the TEMP procedure and initialization is by the INITT procedure.

EXAMPLE 16–6

```
;System software that reads the temperature once per minute and
;logs it into a queue at 0500H – 44FFH.

        MOV    WORD PTR DS:[4FCH],500H     ;in queue pointer = 500H
        MOV    WORD PTR DS:[4FEH],500H     ;out queue pointer = 500H
        CALL   INITT                       ;initialize LM70
        .while 1
            .IF DS:[40CH] == 0 && DS:[40BH] == 0 && DS:[40AH] == 0
                    CALL TEMP               ;once per minute
                    CALL ENQUE              ;queue temperature
                    CALL DTEMP              ;display temperature
            .ENDIF
        .endw

INITT PROC   NEAR                      ;send 0000H to LM70 to reset it

        IO   0FF58H,003FH                  ;p2.6 and p2.7 set to outputs
        MOV  CX,16                          ;bit count to 16
        MOV  DX,0FF5EH                      ;address port 2 latch
        .REPEAT
                MOV  AL,40H
                OUT  DX,AL
                MOV  AL,0
                OUT  DX,AL
        .UNTILCXZ
        RET

INITT ENDP

TEMP  PROC   NEAR            ;read temperature

        IO    0FF58K,00BFH        ;p2.7 as input
        MOV   CX,16
        MOV   BX,0
        .REPEAT                    ;get all 16-bits
                IO   0FF5EH,0C0H
                MOV  DX,0FF5AH
                IN   AL,DX          ;read bit
                SHR  AL,1
                RCR  BX,1           ;into BX
                IO   0FF5EH,40H
        .UNTILCXZ
        MOV   AX,BX
        SAR   AX,6                  ;convert to integer temperature
        RET

TEMP  EMDP

ENQUE PROC   NEAR USES AX        ;queue temperature, no check for full

        MOV    BX,DS:[4FCH]
        MOV    [BX],AX              ;save temp
        ADD    BX,2
        MOV    AX,DS:[40DH]         ;get time HH:MM
        MOV    [BX],AX
        ADD    BX,2
        .IF BX == 4500H
            MOV  BX,500H
        .ENDIF
        RET
```

```
ENQUE  ENDP

DTEMP  PROC

       MOV    BX,410H              ;address string buffer
       OR     AX,AX
       .IF    SIGN?                ;if negative
              MOV BYTE PTR [BX],'-'
              NEG AX
              INC BX
       .ENDIF
       AAM                         ;convert to BCD
       .IF AH != 0
              ADD  H,30H
              MOV  [BX],AH
              INC  BX
       .ENDIF
       ADD    AL,30H
       MOV    [BX],AL
       INC    BX
       MOV    BYTE PTR [BX],'°'
       MOV    BYTE PTR [BX+1],'C'
       MOV    AX,DS
       MOV    ES,AX
       MOV    AL,86H
       MOV    BX,410H
       INT    42H                  ;display temp on line 1

DTEMP  ENDP
```

A Threaded System

At times an operating system is needed that can process multiple threads. Multiple threads are handled by the kernel using a real-time clock interrupt. One method for scheduling processes in a small RTOS is to use a time slice to switch between various processes. The basic time slice can be any duration and is somewhat dependent on the execution speed of the microprocessor. For example, in a system using a 100 MHz clock, many instructions will execute in one or two clocks on a modem microprocessor. Assuming the machine executes one instruction every two clocks and a time slice of 1 ms is chosen, the machine can execute about 50,000 instructions one time slice, which should be adequate for most systems. If a lower clock frequency is employed, then a time slice of 10 ms or even 100 ms is selected.

Each time slice is activated by a timer interrupt. The interrupt service procedure must look to the queue to determine if a task is available to execute, and if it is, it must start execution of the new task. If no new task is present, it must continue executing old tasks or enter an idle state and wait for a new task to be queued. The queue is circular and may contain any number of tasks for the system up to some finite limit. For example, it might be a small queue in a small system with 10 entries. The size is determined by the intended overall system needs and could be made larger or smaller.

Each scheduling queue entry must contain a pointer to the process (CS:IP) and the entire context state of the machine. Scheduling queue entries may also contain some form of a time-to-live entry in case of a deadlock, a priority entry, and an entry that can lengthen the slice activation time. In the following example, a priority entry or an entry to lengthen the amount of consecutive time slices allowed a program will not be used. The kernel will service processes strictly on a linear basis or on a round-robin fashion as they come from the queue.

To implement a scheduler for the embedded system, procedures, or macros, are implemented to start a new application, kill an application when it completes, and pause an application if it needs time to access I/O. Each of these macros accesses a scheduling queue located in the memory system at an available address such as 0500H. The scheduling queue uses the data structure in Example 16–7 to make creating the queue fairly easy, and it will have room for 10 entries. This scheduling queue allows us to start up to 10 processes at a time.

EXAMPLE 16–7

```
PRESENT     DB    0       ;0 = not present
DUMMY1      DB    ?
RAX         DW    ?
RBX         DW    ?
RCX         DW    ?
RDX         DW    ?
RSP         DW    ?
RBP         DW    ?
RSI         DW    ?
RDI         DW    ?
RFLAG       DW    ?
RIP         DW    ?
RCS         DW    ?
RDS         DW    ?
RES         DW    ?
RSS         DW    ?
DUMMY2      DW    ?       ;padding to 32 bytes
```

The data structure of Example 16–7 is copied into memory 10 times to complete the queue structure during system initialization; hence, it contains no active process at initialization. We also need a queue pointer initialized to 500H. The queue pointer is stored at location 4FEH in this example. Example 16–8 provides one possible initialization. This stores the data structure in the RAM with 10 copies beginning at 500H. This software assumes that a system clock of 32 MHz operates timer 2 used as a prescalar to divide the clock input (system clock divided by 8) of 4 MHz by 40,000. This causes the output of timer 2 to be 1 KHz (1.0 ms). Timer 1 is programmed to divide the timer 2 clock signal by 10 to generate an interrupt every 10 ms.

EXAMPLE 16–8

```
;initialization of thread queue

PUSH   DS
MOV    AX,0
MOV    DS,AX
MOV    SI,500H
MOV    BX,OFFSET PRESENT-100H    ;empty all queue entries
MOV    CX,10
.REPEAT
   MOV  BYTE PTR DS:[SI],0
   ADD  SI,32
.UNTILCXZ
MOV    DS:[4FEH],500H            ;set queue pointer
POP    DS
MOV    DX,0FF42H                 ;Timer 2 CMPA = 40000
MOV    AX,40000
OUT    DX,AL
MOV    DX,0FF32H                 ;Timer 1 CMPA = 10
MOV    AX,10
OUT    DX,AL
MOV    AX,0                      ;Clear Timer count registers
MOV    DX,0FF30H
OUT    DX,AL
MOV    DX,0FF40H
OUT    DX,AL
MOV    DX,0FF46H
MOV    AX,0C001H                 ;Start Timer 2
OUT    DX,AL
MOV    DX,0FF36H                 ;Start Timer 1
MOV    AX,0E009H
OUT    DX,AL
```

The NEW procedure (installed at INT 60H in Example 16–9) adds a process to the queue. It searches through the 10 entries until it finds a zero in the first byte (PRESENT), which indicates

that the entry is empty. If it finds an empty entry, it places the starting address of the process into RCS and RIP and a 0200H into the RFLAG location. A 200H in RFLAG makes sure that the interrupt is enabled when the process begins, which prevents the system from crashing. The NEW procedure waits, if 10 processes are already scheduled, until a process ends. Each process is also assigned stack space in 256-byte sections beginning at offset address 7600H, so the lowest process has stack space 7500H–75FFH, the next has stack space 7600H–76FFH, and so on. The assignment of a stack area could be allocated by a memory manager algorithm.

EXAMPLE 16–9

```
INT60   PROC    FAR USES DS AX DX SI

        MOV     AX,0
        MOV     DS,AX                   ;address segment 0000
        STI                             ;interrupts on

        .REPEAT                         ;sticks here is full
                MOV    SI,500H
                HLT                     ;synchronize with RTC interrupt
                .WHILE BYTE PTR DS:[SI] != 0 && SI != 660H
                        ADD SI,32
                .ENDW
        .UNTIL BYTE PTR DS:[SI] == 0

        MOV     BYTE PTR DS:[SI],0FFH ;activate process
        MOV     WORD PTR DS:[SI+18],200H      ;flags
        MOV     DS:[SI+20],BX           ;save IP
        MOV     DS:[SI+22,DX            ;save CS
        MOV     DS:[SI+28],SS           ;save SS
        MOV     AX,SI
        AND     AX,3FFH
        SHL     AX,3
        ADD     AX,7500H
        MOV     DS:[SI+10],AX           ;save SP
        IRET

INT60   ENDP
```

The final control procedure (KILL), located at interrupt vector 61H as illustrated in Example 16–10, kills an application by placing 00H into PRESENT of the queue data structure, which removes it from the scheduling queue.

EXAMPLE 16–10

```
INT61 PROC   FAR USES DS AX DX SI

      MOV     AX,0
      MOV     DS,AX
      MOV     SI,DS:[4FEH]              ;get queue pointer
      MOV     BYTE PTR DS:[SI],0        ;kill thread
      JMP     INT12A

INT61 ENDP
```

The PAUSE procedure is merely a call to the time slice procedure (INT 12H) that bails out of the process and returns control to the time slice procedure, prematurely ending the time slice for the process. This *early out* allows other processes to continue before returning to the current process.

The time slice interrupt service procedure for an 80188EB using a 10 ms time slice appears in Example 16–11. Because this is an interrupt service procedure, care has been taken to make it as efficient as possible. Example 16–11 illustrates the time slice procedure located at interrupt vector 12H for operation with timer 1 in the 80188EB microprocessor. Although not shown, this software assumes that timer 2 is used as a prescalar and timer 1 uses the signal from timer 2 to generate the 10 ms interrupt. The software also assumes that no other interrupt is in use in the system.

EXAMPLE 16–11

```
INT12 PROC   FAR USES DS AX DX SI

      MOV   AX,0
      MOV   DS,AX                ;address segment 0000
      MOV   SI,DS:[4FEH]         ;get queue pointer
      CALL  SAVES                ;save state
INT12A:                          ;from Kill
      ADD   SI,32                ;get next process
      .IF SI == 660H             ;make queue circular
          MOV  SI,500H
      .ENDIF
      .WHILE BYTE PTR DS:[SI] != 0FFH   ;find next process
          ADD SI,32
          .IF SI == 660H
              MOV SI,500H
          .ENDIF
      .ENDW
      MOV   DX,0FF30H            ;get a complete 10ms time slice
      MOV   AX,0
      OUT   DX,AL
      MOV   DX,0FF02H            ;clear interrupt
      MOV   AX,8000H
      OUT   DX,AX
      JMP   LOADS                ;load state of next process

INT12 ENDP
```

The system startup (placed after system initialization) must be a single process inside an infinite wait loop. The startup is illustrated in Example 16–12.

EXAMPLE 16–12

```
SYSTEM_STARTUP:

;fork WAITS thread

      MOV   BX,OFFSET WAITS-100H       ;offset address of thread
      MOV   DX,0F800H                  ;segment address of thread
      INT   60H                        ;start wait thread
      STI

;other system processes started here

WAITS:                          ;system idle loop
      .WHILE 1
          INT 12H               ;bail out
      .ENDW
```

Finally, the SAVES and LOADS procedures are used to load and save the machine context as a switch from one process to another occurs. These procedures are called by the time slice interrupt (INT 12H) and start a process (INT 60H) and are listed in Example 16–13.

EXAMPLE 16–13

```
SAVES   PROC   NEAR

      MOV  DS:[SI+4],BX        ;save BX
      MOV  DS:[SI+6],CX        ;save CX
      MOV  DS:[SI+10],SP       ;save SP
      MOV  DS:[SI+12],BP       ;save BP
      MOV  DS:[SI+16],DI       ;save DI
      MOV  DS:[SI+26],ES       ;save ES
      MOV  DS:[SI+28],SS       ;save SS
```

```
                MOV  BP,SP                 ;get SP
                MOV  AX,[SP+2]             ;get SI
                MOV  DS:[SI+14],AX         ;save SI
                MOV  AX,[BP+4]             ;get DX
                MOV  DS:[SI+8],AX          ;save DX
                MOV  AX,[BP+6]             ;get AX
                MOV  DS:[SI+2],AX          ;save AX
                MOV  AX,[BP+8]             ;get DS
                MOV  DS:[SI+24],AX         ;save DS
                MOV  AX,[BP+10]            ;get flags
                MOV  DS:[SI+18],AX         ;save flags
                MOV  AX,[BP+12]            ;get CS
                MOV  DS:[SI+22],AX         ;save CS
                MOV  AX,[BP+14]            ;get IP
                MOV  DS:[SI+20],AX         ;save IP
                RET

SAVES   ENDP

LOADS   PROC    FAR

                MOV  SS,DS:[SI+28]         ;get SS
                MOV  SP,DS:[SI+10]         ;get SP
                PUSH WORD PTR DS:[SI+20]   ;PUSH IP
                PUSH WORD PTR DS:[SI+22]   ;PUSH CS
                PUSH WORD PTR DS:[SI+18]   ;PUSH Flags
                PUSH WORD PTR DS:[SI+24]   ;PUSH DS
                PUSH WORD PTR DS:[SI+2]    ;PUSH AX
                PSUH WORD PTR DS:[SI+8]    ;PUSH DX
                PUSH WORD PTR DS:[SI+14]   ;PUSH SI
                MOV  BX,DS:[SI+4]          ;get BX
                MOV  CX,DS:[SI+6]          ;get CX
                MOV  BP,DS:[SI+12]         ;get BP
                MOV  DI,DS:[SI+16]         ;get DI
                MOV  ES,DS:[SI+26]         ;get ES
                POP  SI
                POP  DX
                POP  AX
                POP  DS
                IRET

LOADS   ENDP
```

16–5   INTRODUCTION TO THE 80286

The 80286 microprocessor is an advanced version of the 8086 microprocessor that was designed for multiuser and multitasking environments. The 80286 addresses 16M bytes of physical memory and 1G bytes of virtual memory by using its memory-management system. This section of the text introduces the 80286 microprocessor, which finds use in earlier AT-style personal computers that once pervaded the computer market and still find some applications. The 80286 is basically an 8086 that is optimized to execute instructions in fewer clocking periods than the 8086. The 80286 is also an enhanced version of the 8086 because it contains a memory manager. At this time, the 80286 no longer has a place in the personal computer system, but it does find applications in control systems as an embedded controller.

Hardware Features

Figure 16–29 shows the internal block diagram of the 80286 microprocessor. Note that like the 80186/80188, the 80286 does not incorporate internal peripherals; instead, it contains a memory-management unit (MMU) that is called the address unit in the block diagram.
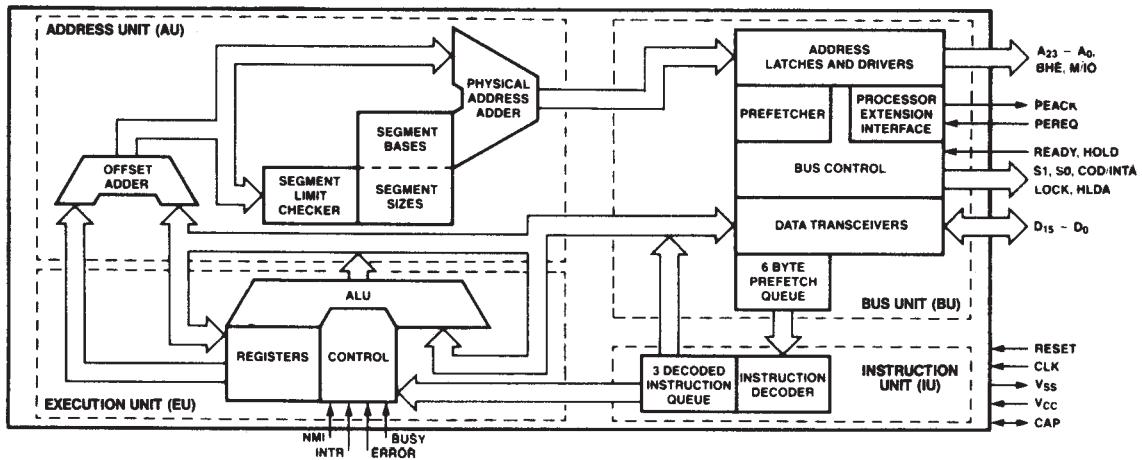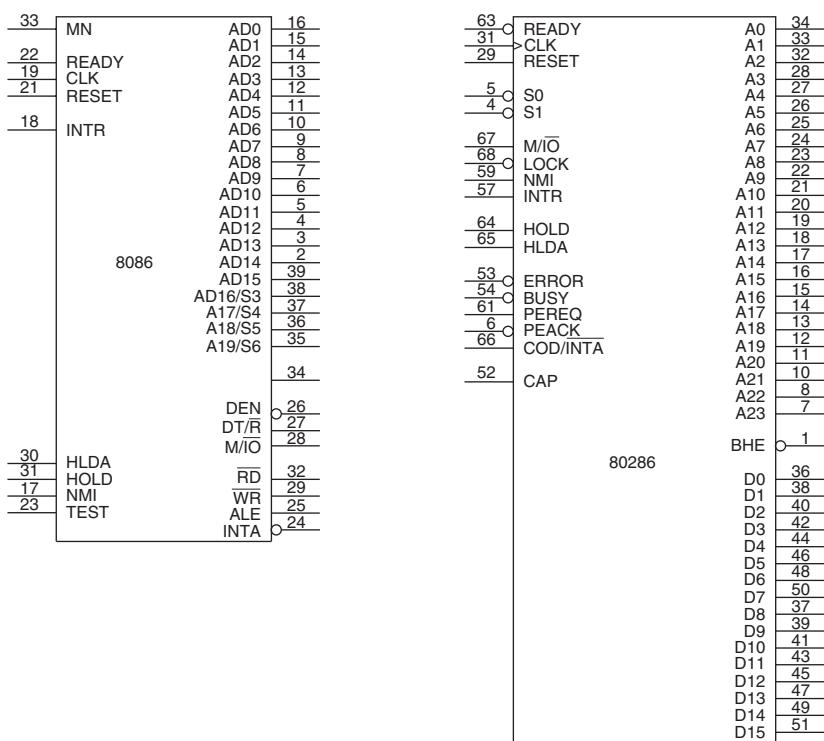
FIGURE 16–29   The block diagram of the 80286 microprocessor. (Courtesy of Intel Corporation.)

As a careful examination of the block diagram reveals, address pins A_{23}–A_0, $\overline{BUSY}$, CAP, $\overline{ERROR}$, $\overline{PEREQ}$, and $\overline{PEACK}$ are new or additional pins that do not appear on the 8086 microprocessor. The $\overline{BUSY}$, $\overline{ERROR}$, $\overline{PEREQ}$, and $\overline{PEACK}$ signals are used with the microprocessor extension or coprocessor, of which the 80287 is an example. (Note that the $\overline{TEST}$ pin is now referred to as the $\overline{BUSY}$ pin.) The address bus is now 24 bits wide to accommodate the 16M bytes of physical memory. The CAP pin is connected to a 0.047 μF, ±20% capacitor that acts as a 12 V filter and connects to ground. The pin-outs of the 8086 and 80286 are illustrated in Figure 16–30 for comparative purposes. Note that the 80286 does not contain a multiplexed address/data bus.

FIGURE 16–30   The 8086 and 80286 microprocessor pin-outs. Notice that the 80286 does not have a multiplexed address/data bus.
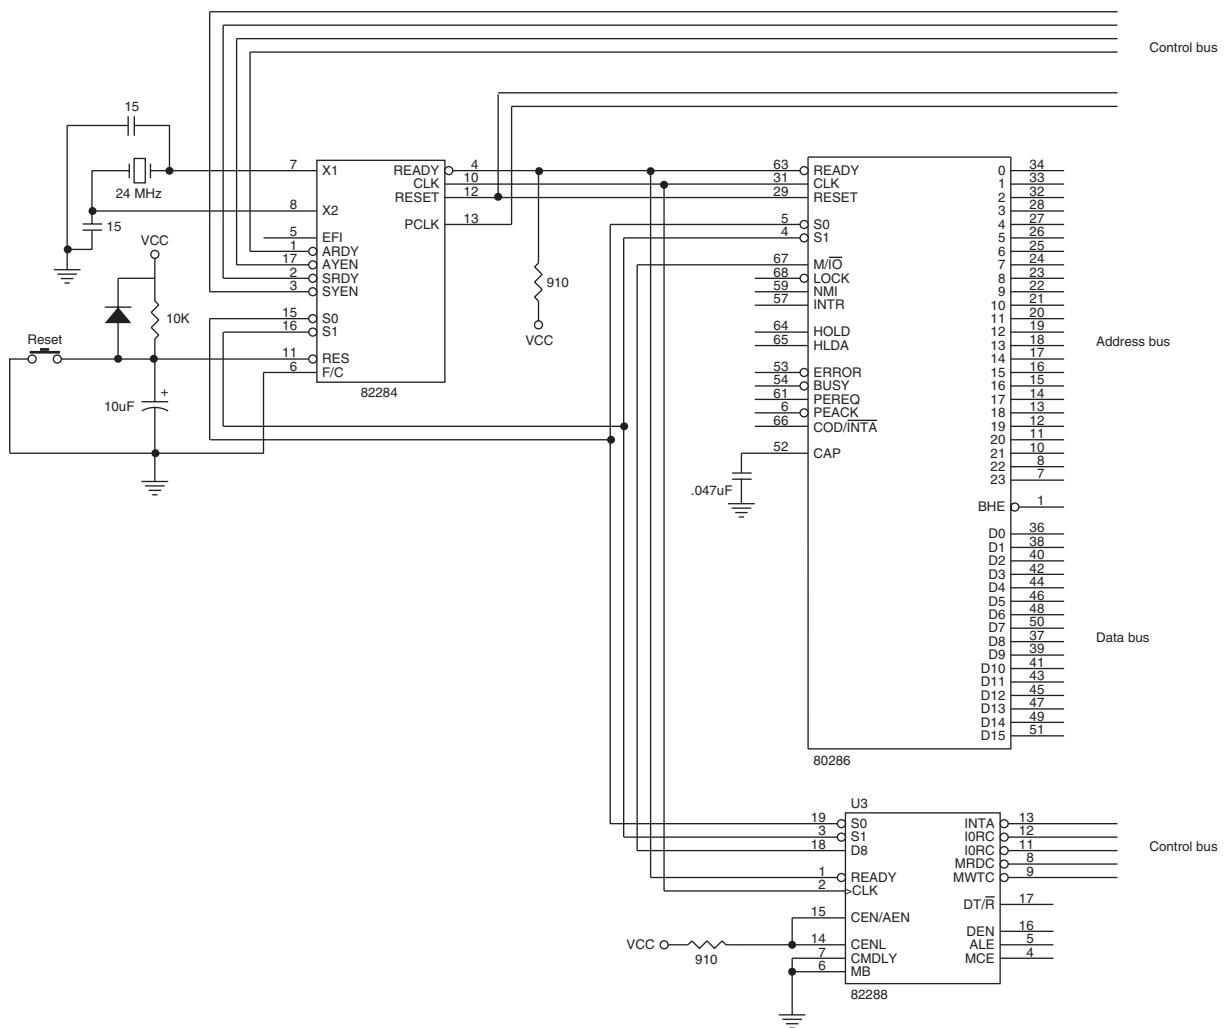
FIGURE 16–31   The interconnection of the 80286 microprocessor, 82284 clock generator, and 8288 system bus controller.

As mentioned in Chapter 1, the 80286 operates in both the real and protected modes. In the real mode, the 80286 addresses a 1M-byte memory address space and is virtually identical to the 8086. In the protected mode, the 80286 addresses a 16M-byte memory space.

Figure 16–31 illustrates the basic 80286 microprocessor-based system. Notice that the clock is provided by the 82284 clock generator (similar to the 8284A) and the system control signals are provided by the 82288 system bus controller (similar to the 8288). Also, note the absence of the latch circuits used to demultiplex the 8086 address/data bus.

Additional Instructions

The 80286 has even more instructions than its predecessors. These extra instructions control the virtual memory system through the memory manager of the 80286. Table 16–9 lists the additional 80286 instructions with a comment about the purpose of each instruction. These instructions are the only new instructions added to the 80286. Note that the 80286 contains the new instructions added to the 80186/80188 such as INS, OUTS, BOUND, ENTER, LEAVE, PUSHA, POPA, and the immediate multiplication and immediate shift and rotate counts.

TABLE 16–9 Additional 80286 instructions.

Instruction	Purpose
CLTS	Clear the task-switched bit
LDGT	Load global descriptor table register
SGDT	Store global descriptor table register
LIDT	Load interrupt descriptor table register
SIDT	Store interrupt descriptor table register
LLDT	Load local descriptor table register
SLDT	Store local descriptor table register
LMSW	Load machine status register
SMSW	Store machine status register
LAR	Load access rights
LSL	Load segment limit
SAR	Store access rights
ARPL	Adjust requested privilege level
VERR	Verify a read access
VERW	Verify a write access

Following are descriptions of instructions not explained in the memory-management section. The instructions described here are special and only used for the conditions indicated.

CLTS   The **clear task-switched flag** (CLTS) instruction clears the TS (task-switched) flag bit to a logic 0. If the TS flag bit is a logic 1 and the 80287 numeric coprocessor is used by the task, an interrupt occurs (vector type 7). This allows the function of the coprocessor to be emulated with software. The CLTS instruction is used in a system and is considered a privileged instruction because it can be executed only in the protected mode at privilege level 0. There is no set TS flag instruction; this is accomplished by writing a logic 1 to bit position 3 (TS) of the machine status word (MSW) by using the LMSW instruction.

LAR   The **load access rights** (LAR) instruction reads the segment descriptor and places a copy of the access rights byte into a 16-bit register. An example is the LAR AX,BX instruction that loads AX with the access rights byte from the descriptor selected by the selector value found in BX. This instruction is used to get the access rights so that it can be checked before a program uses the segment of memory described by the descriptor.

LSL   The **load segment limit** (LSL) instruction loads a user-specified register with the segment limit. For example, the LSL AX,BX instruction loads AX with the limit of the segment described by the descriptor selected by the selector in BX. This instruction is used to test the limit of a segment.

ARPL   The **adjust requested privilege level** (ARPL) instruction is used to test a selector so that the privilege level of the requested selector is not violated. An example is ARPL AX,CX: AX contains the requested privilege level and CX contains the selector value to be used to access a descriptor. If the requested privilege level is of a lower priority than the descriptor under test, the zero flag is set. This may require that a program adjust the requested privilege level or indicate a privilege violation.

VERR   The **verify for read access** (VERR) instruction verifies that a segment can be read. Recall from Chapter 1 that a code segment can be read-protected. If the code segment can be read, the zero flag bit is set. The VERR AX instruction tests the descriptor selected by the AX register.

VERW   The **verify for write access** (VERW) instruction is used to verify that a segment can be written. Recall from Chapter 1 that a data segment can be write-protected. If the data segment can be written, the zero flag bit is set.

The Virtual Memory Machine

A virtual memory machine is a machine that maps a larger memory space (1G bytes for the 80286) into a much smaller physical memory space (l6M bytes for the 80286), which allows a very large system to execute in smaller physical memory systems. This is accomplished by spooling the data and programs between the fixed disk memory system and the physical memory. Addressing a 1G-byte memory system is accomplished by the descriptors in the 80286 micro-processor. Each 80286 descriptor describes a 64K-byte memory segment and the 80286 allows 16K descriptors. This (64K × 16K) allows a maximum of 1G bytes of memory to be described for the system.

As mentioned in Chapter 1, descriptors describe the memory segment in the protected mode. The 80286 has descriptors that define codes, data, stack segments, interrupts, procedures, and tasks. Descriptor accesses are performed by loading a segment register with a selector in the protected mode. The selector accesses a descriptor that describes an area of the memory. Additional details on descriptors and their applications are defined in Chapter 1, and also Chapters 17, 18, and 19. Please refer to these chapters for a detailed view of the protected mode memory-management system.

16–6   SUMMARY

1. The 80186/80188 microprocessors contain the same basic instruction set as the 8086/8088 microprocessors, except that a few additional instructions are added. The 80186/80188 are thus enhanced versions of the 8086/8088 microprocessors. The new instructions include PUSHA, POPA, INS, OUTS, BOUND, ENTER, LEAVE, and immediate multiplication and shift/rotate counts.
2. Hardware enhancements to the 80186/80188 include a clock generator, a programmable interrupt controller, three programmable timers, a programmable DMA controller, a pro-grammable chip selection logic unit, a watchdog timer, a dynamic RAM refresh logic cir-cuit, and additional features on various versions.
3. The clock generator allows the 80186/80188 to operate from an external TTL-level clock source, or from a crystal attached to the X_1 (CLK_{IN}) and X_2 (OSCOUT) pins. The frequency of the crystal is twice the operating frequency of the microprocessor. The 80186/80188 microprocessors are available in speeds of 6 to 20 MHz.
4. The programmable interrupt controller arbitrates all internal and external interrupt requests. It is also capable of operating with two external 8259A interrupt controllers.
5. Three programmable timers are located within the 80186/80188. Each timer is a fully program-mable, 16-bit counter used to generate wave-forms or count events. Two of the timers, timers 0 and 1, have external inputs and outputs. The third timer, timer 2, is clocked from the system clock and is used either to provide a clock for another timer or to request a DMA action.
6. The programmable DMA controller is a fully programmable, two-channel controller. DMA transfers are made between memory and I/O, I/O and I/O, or between memory locations. DMA requests occur from software, hardware, or the output of timer 2.
7. The programmable chip selection unit is an internal decoder that provides up to 13 output pins to select memory (6 pins) and I/O (7 pins). It also inserts 0 to 3 wait states, with or with-

out external READY synchronization. On the EB and EC versions, the number of waits can be programmed from 0 to 15 and 10 chip selection pins.

8. The only difference between the timing of the 80186/80188 and the 8086/8088 is that ALE appears one-half clock pulse earlier. Otherwise, the timing is identical.

9. The 6 MHz version of the 80186/80188 allows 417 ns of access time for the memory; the 8 MHz version allows 309 ns of access time.

10. The interna180186/80188 peripherals are programmed via a peripheral control block (PCB), initialized at I/O ports FF00H–FFFFH. The PCB may be moved to any area of memory or I/O by changing the contents of the PCB relocation register at initial I/O location FFFEH and FFFFH.

11. The 80286 is an 8086 that has been enhanced to include a memory-management unit (MMU). The 80286 is capable of addressing a 16M-byte physical memory space because of the management unit.

12. The 80286 contains the same instructions as the 80186/80188, except for a handful of additional instructions that control the memory-management unit.

13. Through the memory-management unit, the 80286 microprocessor addresses a virtual memory space of 1G bytes, as specified by the 16K descriptors stored in two descriptor tables.

16–7   QUESTIONS AND PROBLEMS

1. List the differences between the 8086/8088 and the 80186/80188 microprocessors.
2. What hardware enhancements are added to the 80186/80188 that are not present in the 8086/8088?
3. The 80186/80188 is packaged in what types of integrated circuits?
4. If the 20 MHz crystal is connected to X_1 and X_2, what frequency signal is found at CLKOUT?
5. Describe the differences between the 80C188XL and the 80C188EB versions of the 80188 embedded controller.
6. The fan-out from any 80186/80188 pin is _____ for a logic 0.
7. How many clocking periods are found in an 80186/80188 bus cycle?
8. What is the main difference between the 8086/8088 and 80186/80188 timing?
9. What is the importance of memory access time?
10. How much memory access time is allowed by the 80186/80188 if operated with a 10 MHz clock?
11. Where is the peripheral control block located after the 80186/80188 is reset?
12. Write the software required to move the peripheral control block to memory locations 10000H–100FFH.
13. Which interrupt vector is used by the INT_0 pin on the 80186/80188 microprocessors?
14. How many interrupt vectors are available to the interrupt controller located within the 80186/80188 microprocessors?
15. Which two modes of operation are available to the interrupt controller?
16. What is the purpose of the interrupt control register?
17. Whenever an interrupt source is masked, the mask bit in the interrupt mask register is a logic _____.
18. What is the difference between the interrupt poll and interrupt poll status registers?
19. What is the purpose of the end-of-interrupt (EOI) register?
20. How many 16-bit timers are found within the 80186/80188?
21. Which timers have input and output pin connections?
22. Which timer connects to the system clock?

23. If two maximum-count compare registers are used with a timer, explain the operation of the timer.
24. What is the purpose of the INH timer control register bit?
25. What is the purpose of the P timer control register bit?
26. The timer control register bit ALT selects what type of operation for timers 0 and 1?
27. Explain how the timer output pins are used.
28. Develop a program that causes timer 1 to generate a continuous signal that is a logic 1 for 123 counts and a logic 0 for 23 counts.
29. Develop a program that causes timer 0 to generate a single pulse after 105 clock pulses on its input pin have occurred.
30. How many DMA channels are controlled by the DMA controller in the 80C186XL?
31. The DMA controller's source and destination registers are each _____ - bits wide.
32. How is the DMA channel started with software?
33. The chip selection unit (XL and EA) has _____ pins to select memory devices.
34. The chip selection unit (XL and EA) has _____ pins to select peripheral devices.
35. The last location of the upper memory block, as selected by the $\overline{\text{UCS}}$ pin, is location _____.
36. The middle memory chip selection pins (XL and EA) are programmed for a(n) _____ size and a block size. ____
37. The lower memory area, as selected by $\overline{\text{LCS}}$, begins at address _____.
38. The internal wait state generator (EB and EC versions) is capable of inserting between zero and _____ wait states.
39. Program register A_8H (XL and EA) so that the mid-range memory block size is 128K bytes and a chip size of 32K.
40. What is the purpose of the EX bit in register A_8H?
41. Develop the software required to program the $\overline{\text{GCS3}}$ pin so that it selects memory from locations 20000H–2FFFFH and inserts two wait states.
42. Develop the software required to program the $\overline{\text{GCS4}}$ pin so that it selects an I/O device for ports 1000H–103FH and inserts one wait state.
43. The 80286 microprocessor addresses _____ bytes of physical memory.
44. When the memory manager is in use, the 80286 addresses _____ bytes of virtual memory.
45. The instruction set of the 80286 is identical to the _____, except for the memory-management instructions.
46. What is the purpose of the VERR instruction?
47. What is the purpose of the LSL instruction?
48. What is an RTOS?
49. How are multiple threads handled with the RTOS?
50. Search the Internet for at least two different RTOSs and write a short report comparing them.

CHAPTER 17

The 80386 and 80486 Microprocessors

INTRODUCTION

The 80386 microprocessor is a full 32-bit version of the earlier 8086/80286 16-bit microprocessors, and represents a major advancement in the architecture—a switch from a 16-bit architecture to a 32-bit architecture. Along with this larger word size are many improvements and additional features. The 80386 microprocessor features multitasking, memory management, virtual memory (with or without paging), software protection, and a large memory system. All software written for the early 8086/8088 and the 80286 are upward-compatible to the 80386 microprocessor. The amount of memory addressable by the 80386 is increased from the 1M byte found in the 8086/8088 and the 16M bytes found in the 80286, to 4G bytes in the 80386. The 80386 can switch between protected mode and real mode without resetting the microprocessor. Switching from protected mode to real mode was a problem on the 80286 microprocessor because it required a hardware reset.

The 80486 microprocessor is an enhanced version of the 80386 microprocessor that executes many of its instructions in one clocking period. The 80486 microprocessor also contains an 8K-byte cache memory and an improved 80387 numeric coprocessor. (Note that the 80486DX4 contains a 16K-byte cache.) When the 80486 is operated at the same clock frequency as an 80386, it performs with about a 50% speed improvement. In Chapter 18, the Pentium and Pentium Pro are detailed. These microprocessors both contain a 16K cache memory, and perform at better than twice the speed of the 80486 microprocessor. The Pentium and Pentium Pro also contain improved numeric coprocessors that operate five times faster than the 80486 numeric coprocessor. Chapter 19 deals with additional improvements in the Pentium II–Core2 microprocessors.

CHAPTER OBJECTIVES

Upon completion of this chapter, you will be able to:

1. Contrast the 80386 and 80486 microprocessors with earlier Intel microprocessors.
2. Describe the operation of the 80386 and 80486 memory-management unit and paging unit.
3. Switch between protected mode and real mode.
4. Define the operation of additional 80386/80486 instructions and addressing modes.
5. Explain the operation of a cache memory system.
6. Detail the interrupt structure and direct memory access structure of the 80386/80486.

7. Contrast the 80486 with the 80386 microprocessor.

8. Explain the operation of the 80486 cache memory.

17–1   INTRODUCTION TO THE 80386 MICROPROCESSOR

Before the 80386 or any other microprocessor can be used in a system, the function of each pin must be understood. This section of the chapter details the operation of each pin, along with the external memory system and I/O structures of the 80386.

Figure 17–1 illustrates the pin-out of the 80386DX microprocessor. The 80386DX is packaged in a 132-pin PGA (pin grid array). Two versions of the 80386 are commonly available: the 80386DX, which is illustrated and described in this chapter and the 80386SX, which is a reduced bus version of the 80386. A new version of the 80386—the 80386EX—incorporates the AT bus system, dynamic RAM controller, programmable chip selection logic, 26 address pins, 16 data pins, and 24 I/O pins. Figure 17–2 illustrates the 80386EX embedded PC.

The 80386DX addresses 4G bytes of memory through its 32-bit data bus and 32-bit address. The 80386SX, more like the 80286, addresses 16M bytes of memory with its 24-bit address bus via its 16-bit data bus. The 80386SX was developed after the 80386DX for applications that didn't require the full 32-bit bus version. The 80386SX was found in many early personal computers that used the same basic motherboard design as the 80286. At the time that the 80386SX was popular, most applications, including Windows 3.11, required fewer than 16M bytes of memory, so the 80386SX is a popular and a less costly version of the 80386 microprocessor. Even though the 80486 has become a less expensive upgrade path for newer systems, the 80386 still can be used for many applications. For example, the 80386EX does not appear in computer systems, but it is becoming very popular in embedded applications.

FIGURE 17–1 The pin-outs of the 80386DX and 80386SX microprocessors.
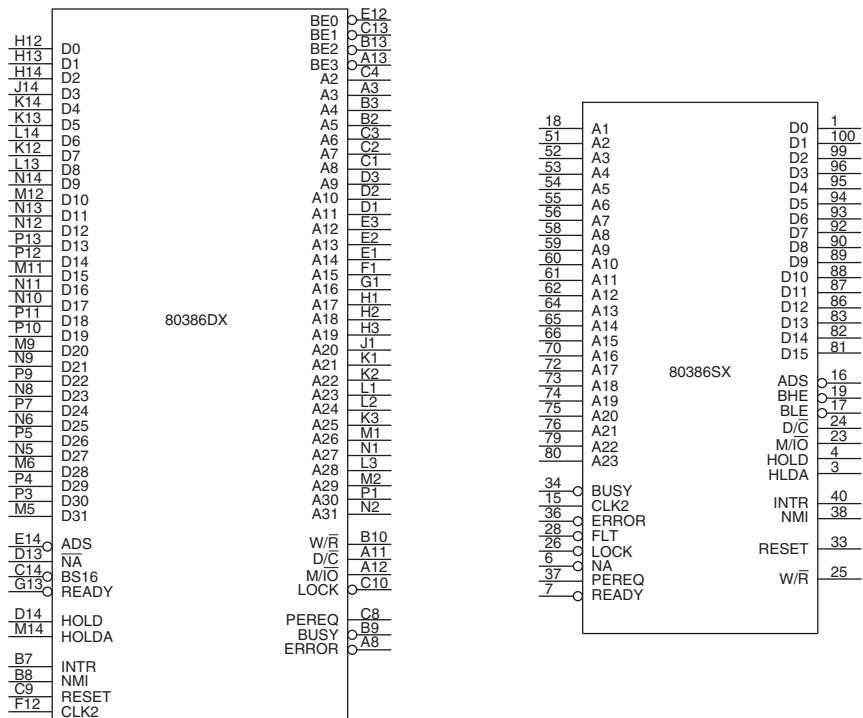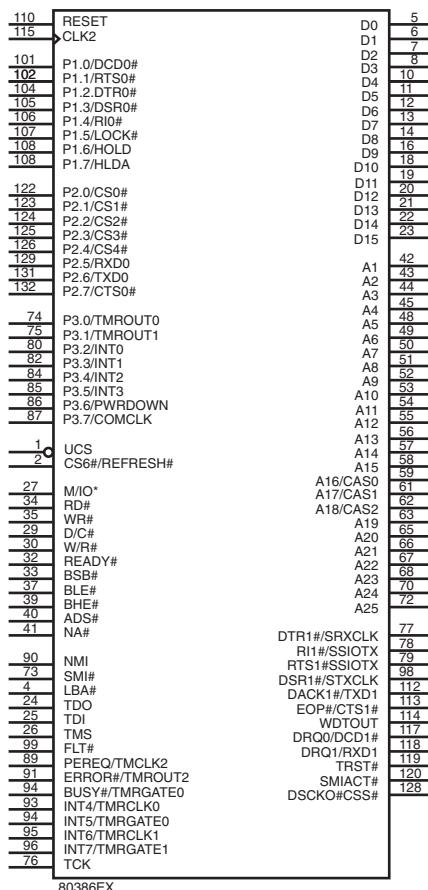
FIGURE 17–2 The 80386EX embedded PC.

110	RESET		D0	5
115	CLK2		D1	6
			D2	7
101	P1.0/DCD0#		D3	8
102	P1.1/RTS0#		D4	10
104	P1.2.DTR0#		D5	11
105	P1.3/DSR0#		D6	12
106	P1.4/RI0#		D7	13
107	P1.5/LOCK#		D8	14
108	P1.6/HOLD		D9	16
108	P1.7/HLDA		D10	18
			D11	19
122	P2.0/CS0#		D12	20
123	P2.1/CS1#		D13	21
124	P2.2/CS2#		D14	22
125	P2.3/CS3#		D15	23
126	P2.4/CS4#			
129	P2.5/RXD0		A1	42
131	P2.6/TXD0		A2	43
132	P2.7/CTS0#		A3	44
			A4	45
74	P3.0/TMROUT0		A5	48
75	P3.1/TMROUT1		A6	49
80	P3.2/INT0		A7	50
82	P3.3/INT1		A8	51
84	P3.4/INT2		A9	52
85	P3.5/INT3		A10	53
86	P3.6/PWRDOWN		A11	54
87	P3.7/COMCLK		A12	55
			A13	56
1	UCS		A14	57
2	CS6#/REFRESH#		A15	58
			A16/CAS0	59
27	M/IO*		A17/CAS1	61
34	RD#		A18/CAS2	62
35	WR#		A19	63
29	D/C#		A20	65
30	W/R#		A21	66
32	READY#		A22	67
33	BSB#		A23	68
37	BLE#		A24	70
39	BHE#		A25	72
40	ADS#			
41	NA#		DTR1#/SRXCLK	77
			RI1#/SSIOTX	78
90	NMI		RTS1#/SSIOTX	79
73	SMI#		DSR1#/STXCLK	98
4	LBA#		DACK1#/TXD1	112
24	TDO		EOP#/CTS1#	113
25	TDI		WDTOUT	114
26	TMS		DRQ0/DCD1#	117
99	FLT#		DRQ1/RXD1	118
89	PEREQ/TMCLK2		TRST#	119
91	ERROR#/TMROUT2		SMIACT#	120
94	BUSY#/TMRGATE0		DSCKO#CSS#	128
93	INT4/TMRCLK0			
94	INT5/TMRGATE0			
95	INT6/TMRCLK1			
96	INT7/TMRGATE1			
76	TCK			

80386EX

As with earlier versions of the Intel family of microprocessors, the 80386 requires a single +5.0 V power supply for operation. The power supply current averages 550 mA for the 25 MHz version of the 80386, 500 mA for the 20 MHz version, and 450 mA for the 16 MHz version. Also available is a 33 MHz version that requires 600 mA of power supply current. The power supply current for the 80386EX is 320 mA when operated at 33 MHz. Note that during some modes of normal operation, power supply current can surge to over 1.0 A. This means that the power supply and power distribution network must be capable of supplying these current surges. This device contains multiple V_{CC} and V_{SS} connections that must all be connected to +5.0 V and grounded for proper operation. Some of the pins are labeled N/C (no connection) and must not be connected. Additional versions of the 80386SX and 80386EX are available with a +3.3 V power supply. They are often found in portable notebook or laptop computers and are usually packaged in a surface mount device.

Each 80386 output pin is capable of providing 4.0 mA (address and data connections) or 5.0 mA (other connections). This represents an increase in drive current compared to the 2.0 mA available on earlier 8086, 8088, and 80286 output pins. The output current available on most 80386EX output pins is 8.0 mA. Each input pin represents a small load, requiring only ±10 μA of current. In some systems, except the smallest, these current levels require bus buffers.

The function of each 80386DX group of pins follows:

A_{31}–A_2 **Address bus** connections address any of the 1G × 32 (4G bytes) memory locations found in the 80386 memory system. Note that A_0 and A_1 are encoded in the bus enable ($\overline{BE3}$–$\overline{BE0}$) to select any or all of the four bytes in a

32-bit-wide memory location. Also note that because the 80386SX contains a 16-bit data bus in place of the 32-bit data bus found on the 80386DX, A_1 is present on the 80386SX, and the bank selection signals are replaced with $\overline{BHE}$ and $\overline{BLE}$. The $\overline{BHE}$ signal enables the upper data bus half; the $\overline{BLE}$ signal enables the lower data bus half.

D_{31}–D_0	**Data bus** connections transfer data between the microprocessor and its memory and I/O system. Note that the 80386SX contains D_{15}–D_0.
$\overline{BE3}$–$\overline{BE0}$	**Bank enable** signals select the access of a byte, word, or doubleword of data. These signals are generated internally by the microprocessor from address bits A_1 and A_0. On the 80386SX, these pins are replaced by $\overline{BHE}$, $\overline{BLE}$, and A_1.
$M/\overline{IO}$	**Memory/IO** selects a memory device when a logic 1 or an I/O device when a logic 0. During the I/O operation, the address bus contains a 16-bit I/O address on address connections A_{15}–A_2.
$W/\overline{R}$	**Write/$\overline{Read}$** indicates that the current bus cycle is a write when a logic 1 or a read when a logic 0.
$\overline{ADS}$	The **address data strobe** becomes active whenever the 80386 has issued a valid memory or I/O address. This signal is combined with the $W/\overline{R}$ signal to generate the separate read and write signals present in the earlier 8086–80286 microprocessor-based systems.
RESET	**Reset** initializes the 80386, causing it to begin executing software at memory location FFFFFFF0H. The 80386 is reset to the real mode, and the leftmost 12 address connections remain logic 1s (FFFH) until a far jump or far call is executed. This allows compatibility with earlier microprocessors.
CLK_2	**Clock times 2** is driven by a clock signal that is twice the operating frequency of the 80386. For example, to operate the 80386 at 16 MHz, apply a 32 MHz clock to this pin.
$\overline{READY}$	**Ready** controls the number of wait states inserted into the timing to lengthen memory accesses.
$\overline{LOCK}$	**Lock** becomes a logic 0 whenever an instruction is prefixed with the LOCK: prefix. This is used most often during DMA accesses.
$D/\overline{C}$	**Data/control** indicates that the data bus contains data for or from memory or I/O when a logic 1. If $D/\overline{C}$ is a logic 0, the microprocessor is halted or executes an interrupt acknowledge.
$\overline{BS16}$	**Bus size 16** selects either a 32-bit data bus ($\overline{BS16}$ = 1) or a 16-bit data bus ($\overline{BS16}$ = 0). In most cases, if an 80386DX is operated on a 16-bit data bus, we use the 80386SX that has a 16-bit data bus. On the 80386EX, the $\overline{BS8}$ pin selects an 8-bit data bus.
$\overline{NA}$	**Next address** causes the 80386 to output the address of the next instruction or data in the current bus cycle. This pin is often used for pipelining the address.
HOLD	**Hold** requests a DMA action.
HLDA	**Hold acknowledge** indicates that the 80386 is currently in a hold condition.
$\overline{PEREQ}$	The **coprocessor request** asks the 80386 to relinquish control and is a direct connection to the 80387 arithmetic coprocessor.
$\overline{BUSY}$	**Busy** is an input used by the WAIT or FWAIT instruction that waits for the coprocessor to become not busy. This is also a direct connection to the 80387 from the 80386.

$\overline{\text{ERROR}}$	**Error** indicates to the microprocessor that an error is detected by the coprocessor.
INTR	An **interrupt request** is used by external circuitry to request an interrupt.
NMI	A **non-maskable interrupt** requests a non-maskable interrupt as it did on the earlier versions of the microprocessor.

The Memory System
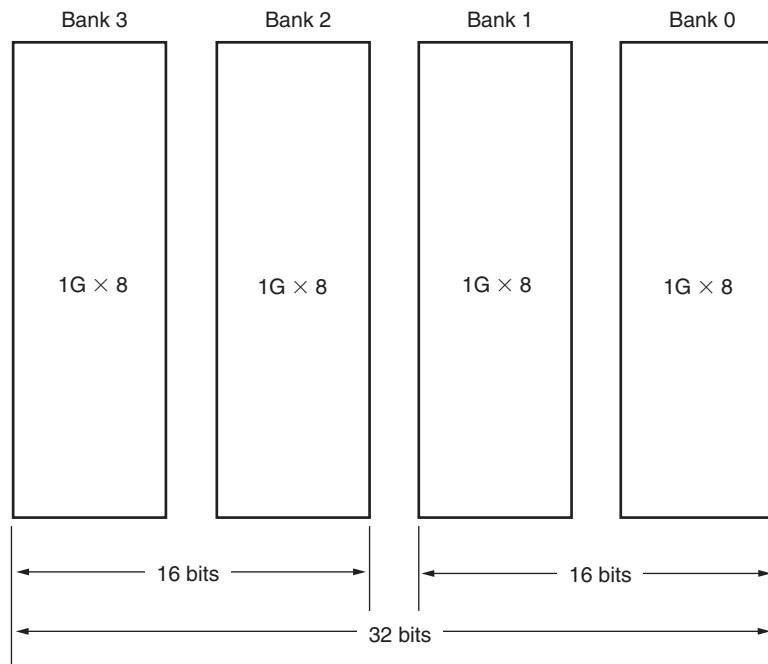
The physical memory system of the 80386DX is 4G bytes in size and is addressed as such. If virtual addressing is used, 64T bytes are mapped into the 4G bytes of physical space by the memory management unit and descriptors. (Note that virtual addressing allows a program to be larger than 4G bytes if a method of swapping with a large hard disk drive exists.) Figure 17–3 shows the organization of the 80386DX physical memory system.

The memory is divided into four 8-bit wide memory banks, each containing up to 1G bytes of memory. This 32-bit-wide memory organization allows bytes, words, or doublewords of memory data to be accessed directly. The 80386DX transfers up to a 32-bit-wide number in a single memory cycle, whereas the early 8088 requires four cycles to accomplish the same transfer, and the 80286 and 80386SX require two cycles. Today, the data width is important, especially with single-precision floating-point numbers that are 32 bits wide. High-level software normally uses floating-point numbers for data storage, so 32-bit memory locations speed the execution of high-level software when it is written to take advantage of this wider memory.

Each memory byte is numbered in hexadecimal as they were in prior versions of the family. The difference is that the 80386DX uses a 32-bit-wide memory address, with memory bytes numbered from location 00000000H to FFFFFFFFH.

The two memory banks in the 8086, 80286, and 80386SX system are accessed via $\overline{\text{BLE}}$ (A_0 on the 8086 and 80286) and $\overline{\text{BHE}}$. In the 80386DX, the memory banks are accessed via four bank enable signals, $\overline{\text{BE3}} - \overline{\text{BE0}}$. This arrangement allows a single byte to be accessed when one bank enable signal is activated by the microprocessor. It also allows a word to be addressed when two

FIGURE 17–3 The memory system for the 80386 microprocessor. Notice that the memory is organized as four banks, each containing 1G byte. Memory is accessed as 8-, 16-, or 32-bit data.

bank enable signals are activated. In most cases, a word is addressed in banks 0 and 1, or in banks 2 and 3. Memory location 00000000H is in bank 0, location 00000001H is in bank 1, location 00000002H is in bank 2, and location 00000003H is in bank 3. The 80386DX does not contain address connections A_0 and A_1 because these have been encoded as the bank enable signals. Likewise, the 80386SX does not contain the A_0 address pin because it is encoded in the $\overline{BLE}$ and $\overline{BHE}$ signals. The 80386EX addresses data either in two banks for a 16-bit-wide memory system if $\overline{BS8}$ = 1 or as an 8-bit system if $\overline{BS8}$ = 0.

Buffered System.   Figure 17–4 shows the 80386DX connected to buffers that increase fan-out from its address, data, and control connections. This microprocessor is operated at 25 MHz using a 50 MHz clock input signal that is generated by an integrated oscillator module. Oscillator modules are usually used to provide a clock in modern microprocessor-based equipment. The HLDA signal is used to enable all buffers in a system that uses direct memory access. Otherwise, the buffer enable pins are connected to ground in a non-DMA system.

Pipelines and Caches.   The cache memory is a buffer that allows the 80386 to function more efficiently with lower DRAM speeds. A pipeline is a special way of handling memory accesses so the memory has additional time to access data. A 16 MHz 80386 allows memory devices with access times of 50 ns or less to operate at full speed. Obviously, there are few DRAMs currently available with these access times. In fact, the fastest DRAMs currently in use have an access time of 40 ns or longer. This means that some technique must be found to interface these memory devices, which are slower than required by the microprocessor. Three techniques are available: interleaved memory, caching, and a pipeline.

The pipeline is the preferred means of interfacing memory because the 80386 microprocessor supports pipelined memory accesses. Pipelining allows memory an extra clocking period to access data. The extra clock extends the access time from 50 ns to 81 ns on an 80386 operating with a 16 MHz clock. The pipe, as it is often called, is set up by the microprocessor. When an instruction is fetched from memory, the microprocessor often has extra time before the next instruction is fetched. During this extra time, the address of the next instruction is sent out from the address bus ahead of time. This extra time (one clock period) is used to allow additional access time to slower memory components.

Not all memory references can take advantage of the pipe, which means that some memory cycles are not pipelined. These nonpipelined memory cycles request one wait state if the normal pipeline cycle requires no wait states. Overall, a pipe is a cost-saving feature that reduces the access time required by the memory system in low-speed systems.

Not all systems can take advantage of the pipe. Those systems typically operate at 20, 25, or 33 MHz. In these higher-speed systems, another technique must be used to increase the memory system speed. The cache memory system improves overall performance of the memory systems for data that are accessed more than once. Note that the 80486 contains an internal cache called a level 1 cache and the 80386 can only contain an external cache called a level 2 cache.

A cache is a high-speed memory (SRAM) system that is placed between the microprocessor and the DRAM memory system. Cache memory devices are usually static RAM memory components with access times of less than 10 ns. In many cases, we see level 2 cache memory systems with sizes between 32K and 1M byte. The size of the cache memory is determined more by the application than by the microprocessor. If a program is small and refers to little memory data, a small cache is beneficial. If a program is large and references large blocks of memory, the largest cache size possible is recommended. In many cases, a 64K-byte cache improves speed sufficiently, but the maximum benefit is often derived from a 256K-byte cache. It has been found that increasing the cache size much beyond 256K provides little benefit to the operating speed of the system that contains an 80386 microprocessor.
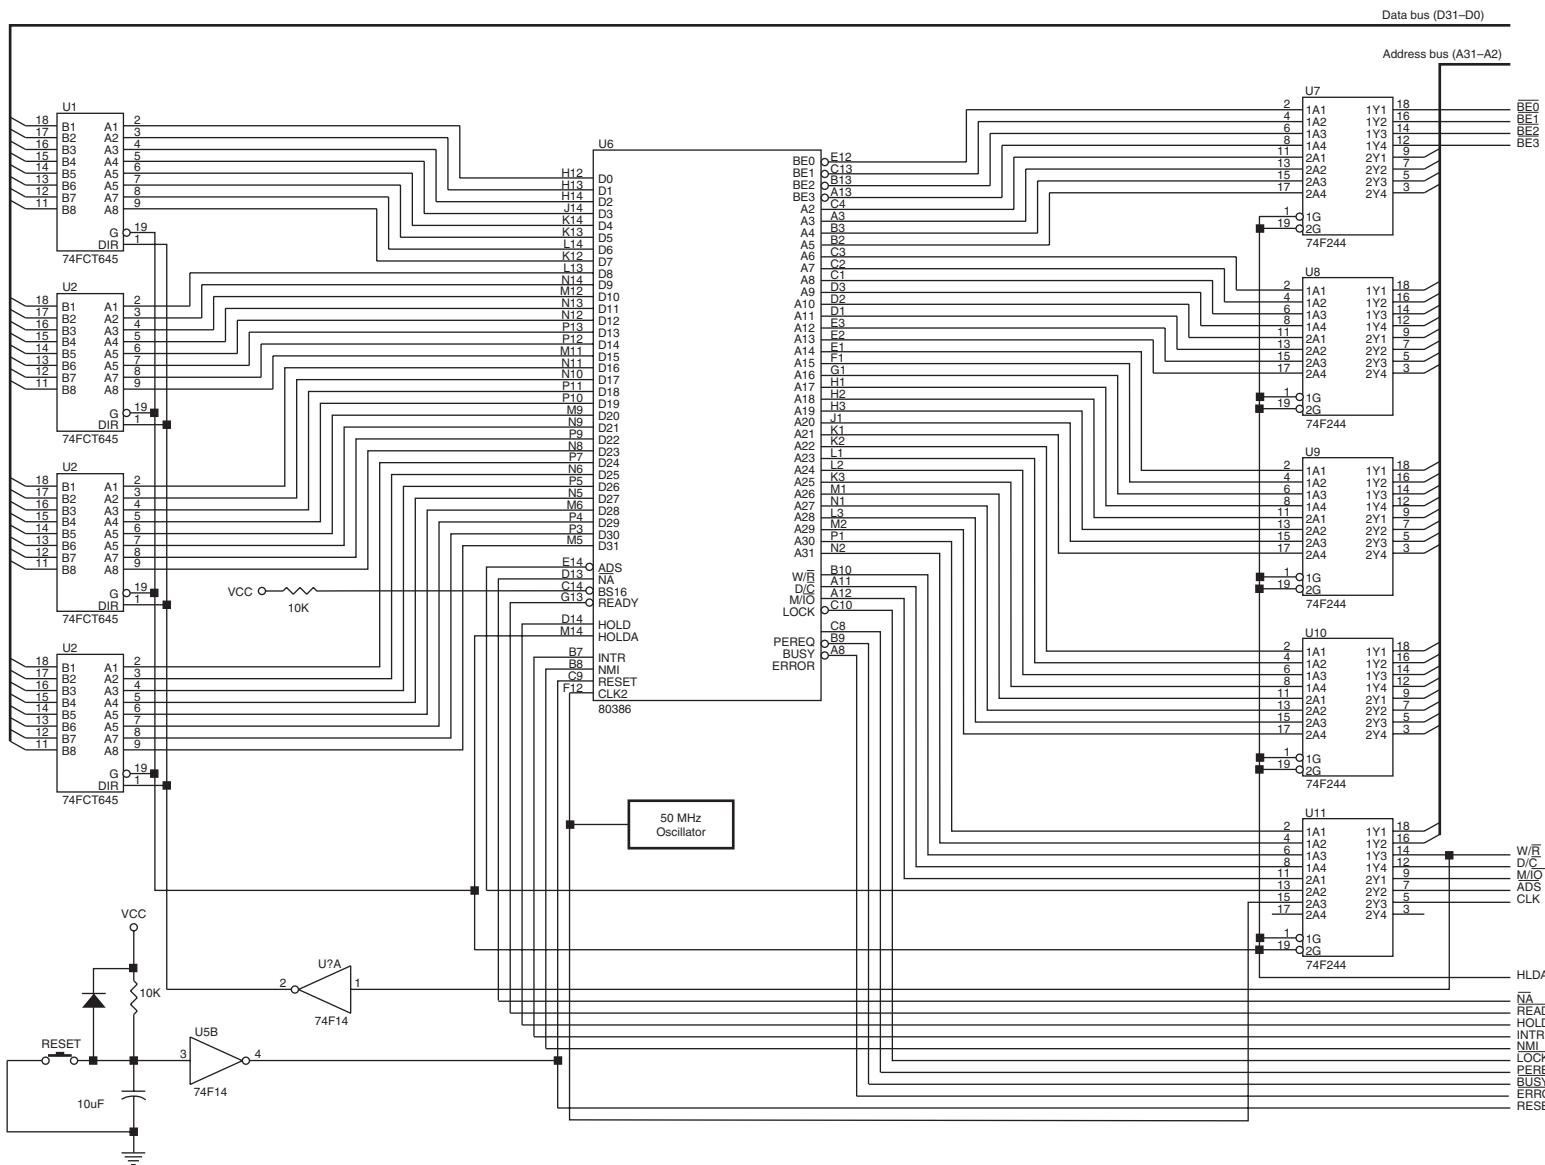
FIGURE 17–4 A fully buffered 25 MHz 80386DX.

Interleaved Memory Systems.   An interleaved memory system is another method of improving the speed of a system. Its only disadvantage is that it costs considerably more memory because of its structure. Interleaved memory systems are present in some systems, so memory access times can be lengthened without the need for wait states. In some systems, an interleaved memory may still require wait states, but may reduce their number. An interleaved memory system requires two or more complete sets of address buses and a controller that provides addresses for each bus. Systems that employ two complete buses are called a two-way interleave; systems that use four complete buses are called a four-way interleave.

An interleaved memory is divided into two or four parts. For example, if an interleaved memory system is developed for the 80386SX microprocessor, one section contains the 16-bit addresses 000000H–000001H, 000004H–000005H, and so on; the other section contains addresses 000002–000003, 000006H–000007H, and so forth. While the microprocessor accesses locations 000000H–000001H, the interleave control logic generates the address strobe signal for locations 000002H–000003H. This selects and accesses the word at location 000002H–000003H, while the microprocessor processes the word at location 000000H–000001H. This process alternates memory sections, thus increasing the performance of the memory system.

Interleaving increases the amount of access time provided to the memory because the address is generated to select the memory before the microprocessor accesses it. This is because the microprocessor pipelines memory addresses, sending the next address out before the data are read from the last address.

The problem with interleaving, although not major, is that the memory addresses must be accessed so that each section is alternately addressed. This does not always happen as a program executes. Under normal program execution, the microprocessor alternately addresses memory approximately 93% of the time. For the remaining 7%, the microprocessor addresses data in the same memory section, which means that in these 7% of the memory accesses, the memory system must cause wait states because of the reduced access time. The access time is reduced because the memory must wait until the previous data are transferred before it can obtain its address. This leaves the memory with less access time; therefore, a wait state is required for accesses in the same memory bank.

See Figure 17–5 for the timing diagram of the address as it appears at the microprocessor address pins. This timing diagram shows how the next address is output before the current data are accessed. It also shows how access time is increased by using interleaved memory addresses for each section of memory compared to a non-interleaved access, which requires a wait state.

Figure 17–6 pictures the interleave controller. Admittedly, this is a complex logic circuit, which needs some explanation. First, if the SEL input (used to select this section of the memory) is inactive (logic 0), then the $\overline{\text{WAIT}}$ signal is a logic 1. Also, both ALE_0 and ALE_1, used to strobe the address to the memory sections, are both logic 1s, causing the latches connected to them to become transparent.

As soon as the SEL input becomes a logic 1, this circuit begins to function. The A_1 input is used to determine which latch (U2B or U5A) becomes a logic 0, selecting a section of the memory. Also the ALE pin that becomes a logic 0 is compared with the previous state of the ALE pins. If the same section of memory is accessed a second time, the $\overline{\text{WAIT}}$ signal becomes a logic 0, requesting a wait state.

Figure 17–7 illustrates an interleaved memory system that uses the circuit of Figure 17–6. Notice how the ALE_0 and ALE_1 signals are used to capture the address for either section of memory. The memory in each bank is 16 bits wide. If accesses to memory require 8-bit data, the system causes wait states, in most cases. As a program executes, the 80386SX fetches instructions 16 bits at a time from normally sequential memory locations. Program execution uses interleaving in most cases. If a system is going to access mostly 8-bit data, it is doubtful that memory interleaving will reduce the number of wait states.
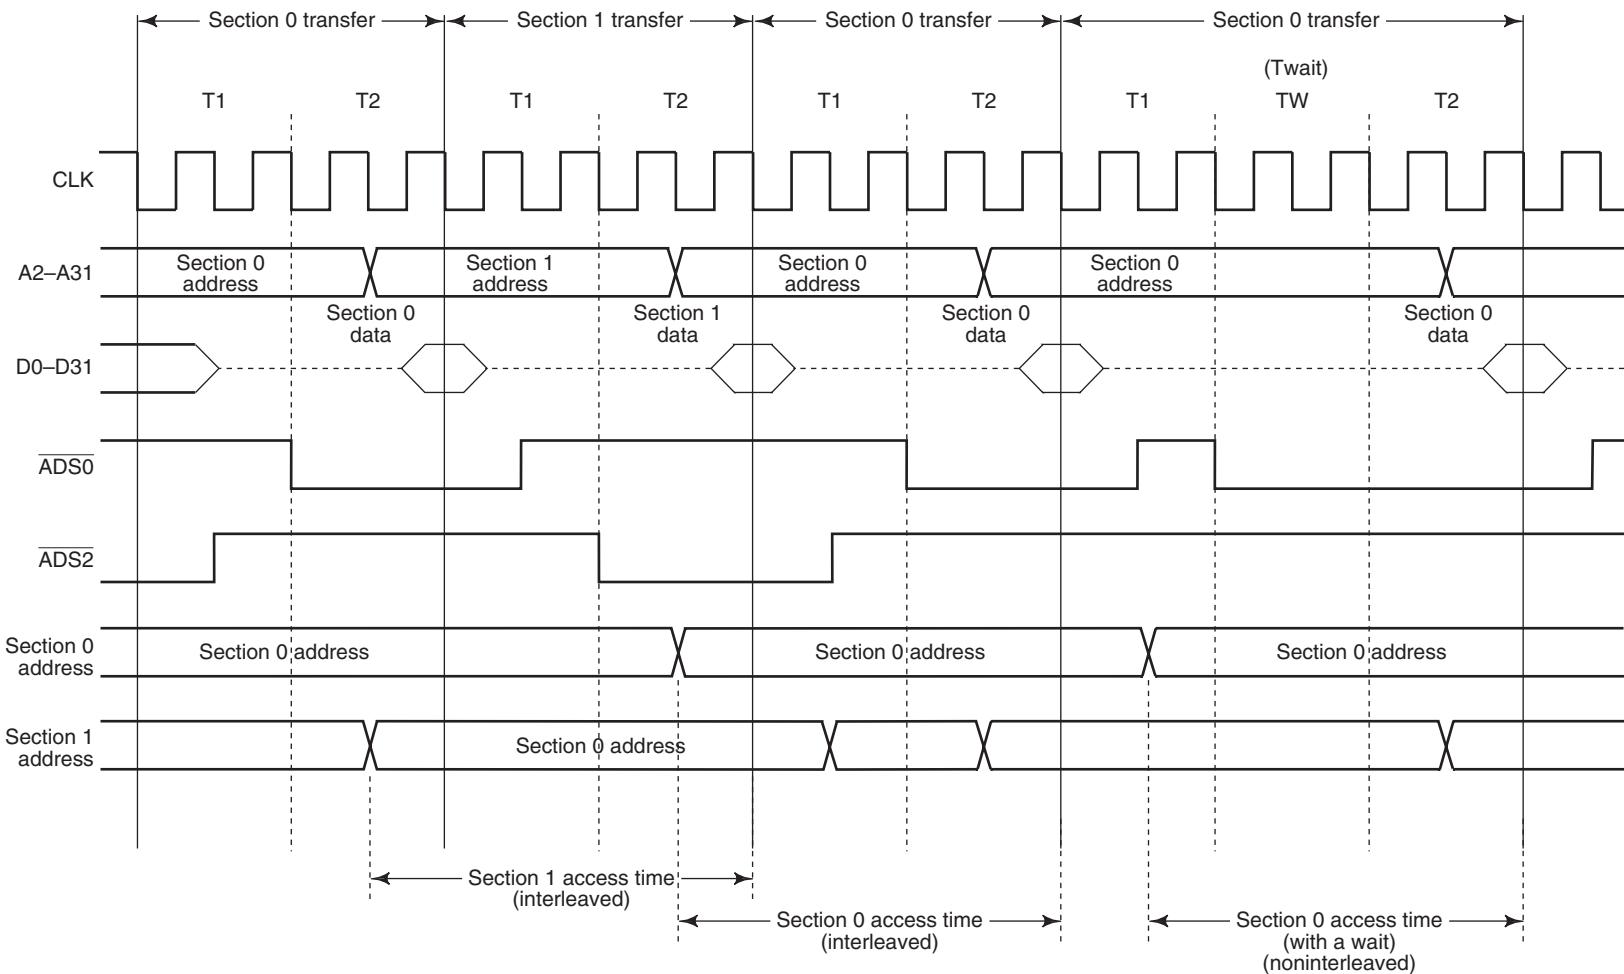
FIGURE 17–5 The timing diagram of an interleaved memory system showing the access times and address signals for both sections of memory.
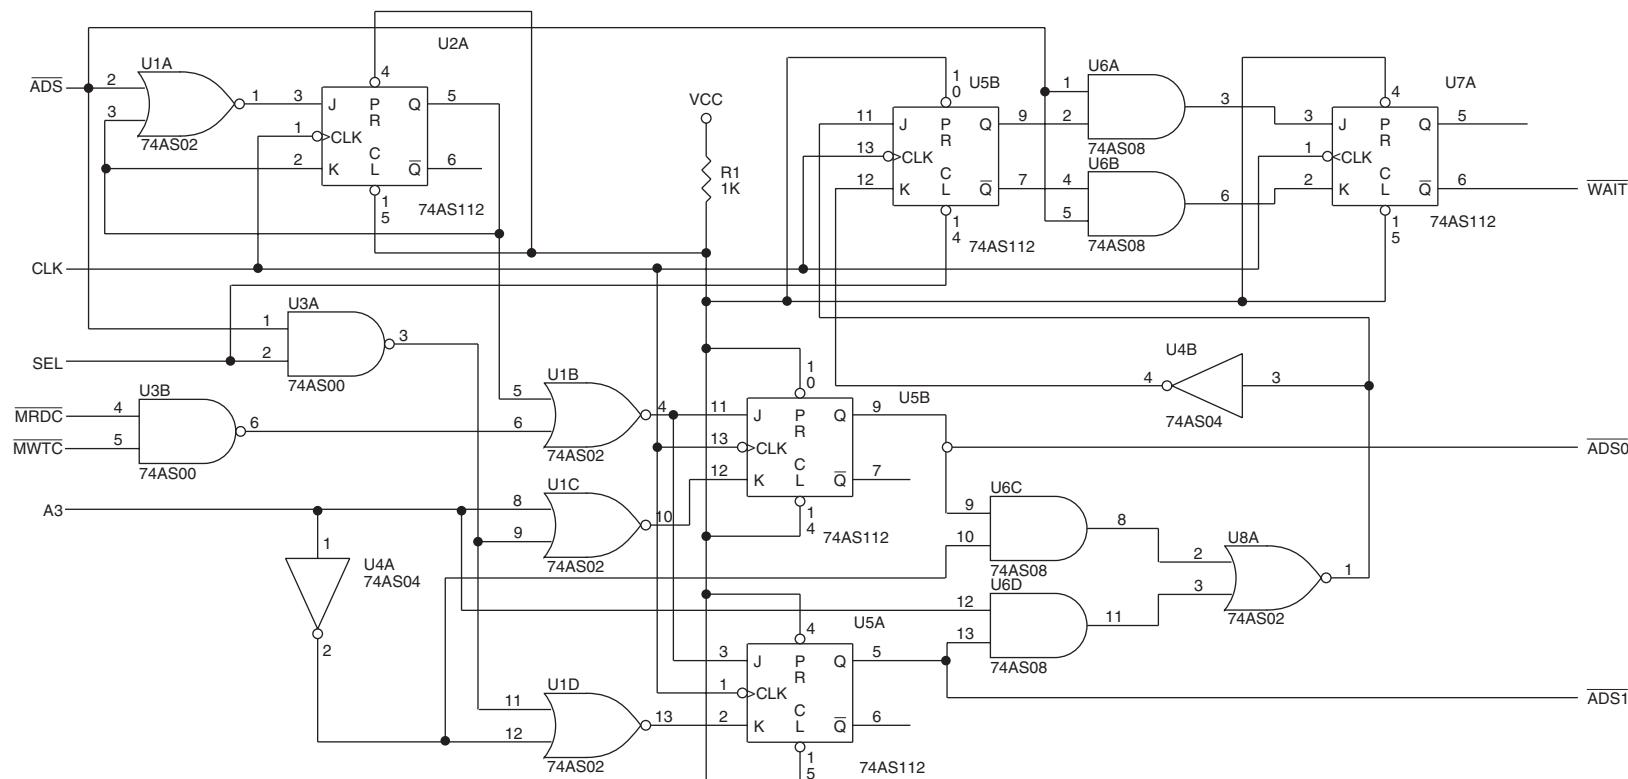
FIGURE 17–6 The interleaved control logic, which generates separate ADS signals and a WAIT signal used to control interleaved memory.
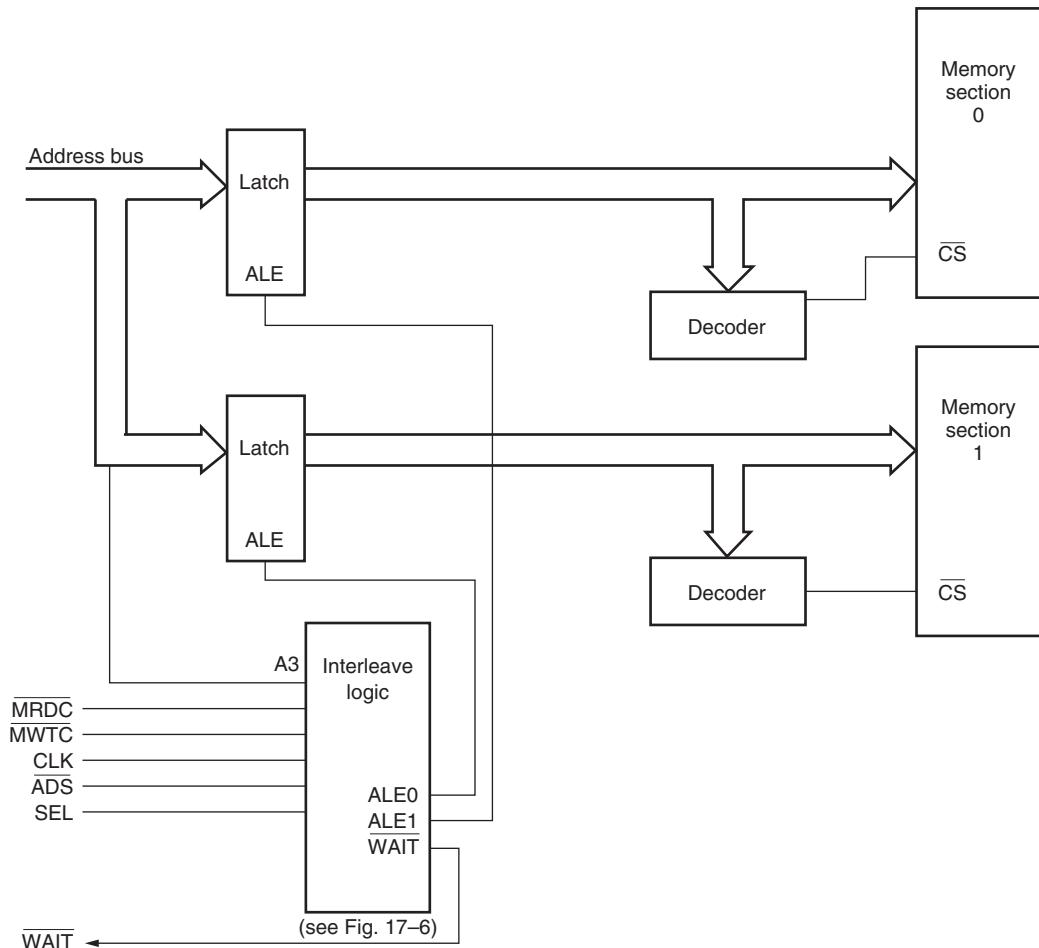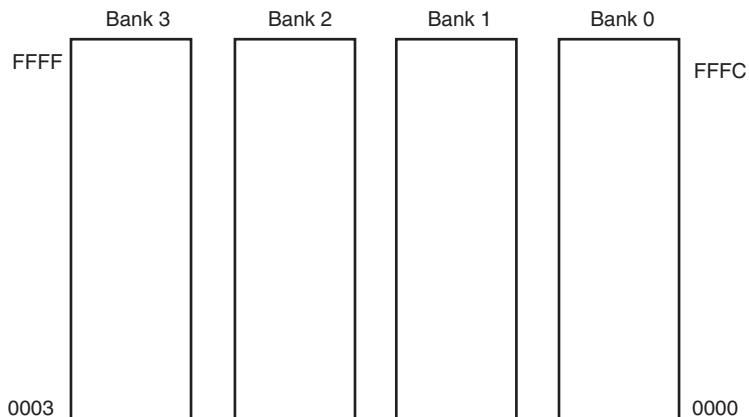
FIGURE 17–7   An interleaved memory system showing the address latches and the interleaved logic circuit.

The access time allowed by an interleaved system, such as the one shown in Figure 17–7, is increased to 112 ns from 69 ns by using a 16 MHz system clock. (If a wait state is inserted, access time with a 16 MHz clock is 136 ns, which means that an interleaved system performs at about the same rate as a system with one wait state.) If the clock is increased to 20 MHz, the interleaved memory requires 89.6 ns, where standard, noninterleaved memory interfaces allow 48 ns for memory access. At this higher clock rate, 80 ns DRAMs function properly without wait states when the memory addresses are interleaved. If an access to the same section occurs, a wait state is inserted.

The Input/Output System

The 80386 input/output system is the same as that found in any Intel 8086 family microprocessor-based system. There are 64K different bytes of I/O space available if isolated I/O is implemented. With isolated I/O, the IN and OUT instructions are used to transfer I/O data between the microprocessor and I/O devices. The I/O port address appears on address bus connections A_{15}–A_2, with $\overline{BE3}$–$\overline{BE0}$ used to select a byte, word, or doubleword of I/O data. If memory-mapped I/O is implemented, then the number of I/O locations can be any amount up to 4G bytes. With memory-mapped I/O, any instruction that transfers data between the microprocessor and memory system can be used for I/O transfers because the I/O device is treated as a memory

FIGURE 17–8   The isolated I/O map for the 80386 micro- processor. Here four banks of 8 bits each are used to address 64K different I/O locations. I/O is numbered from location 0000H to FFFFH.

device. Almost all 80386 systems use isolated I/O because of the I/O protection scheme provided by the 80386 in protected mode operation.

Figure 17–8 shows the I/O map for the 80386 microprocessor. Unlike the I/O map of ear- lier Intel microprocessors, which were 16 bits wide, the 80386 uses a full 32-bit-wide I/O system divided into four banks. This is identical to the memory system, which is also divided into four banks. Most I/O transfers are 8 bits wide because we often use ASCII code (a 7-bit code) for transferring alphanumeric data between the microprocessor and printers and keyboards. This may change if Unicode, a 16-bit alphanumeric code, becomes common and replaces ASCII code. Recently, I/O devices that are 16 and even 32 bits wide have appeared for systems such as disk memory and video display interfaces. These wider I/O paths increase the data transfer rate between the microprocessor and the I/O device when compared to 8-bit transfers.

The I/O locations are numbered from 0000H to FFFFH. A portion of the I/O map is desig- nated for the 80387 arithmetic coprocessor. Although the port numbers for the coprocessor are well above the normal I/O map, it is important that they be taken into account when decoding I/O space (overlaps). The coprocessor uses I/O locations 800000F8H–800000FFH for communica- tions between the 80387 and 80386. The 80287 numeric coprocessor designed for use with the 80286 uses the I/O addresses 00F8H–00FFH for coprocessor communications. Because we often decode only address connections A_{15}–A_2 to select an I/O device, be aware that the coprocessor will activate devices 00F8H–00FFH unless address line A_{31} is also decoded. This should present no problem because you really should not be using I/O ports 00F8H–00FFH for any purpose.
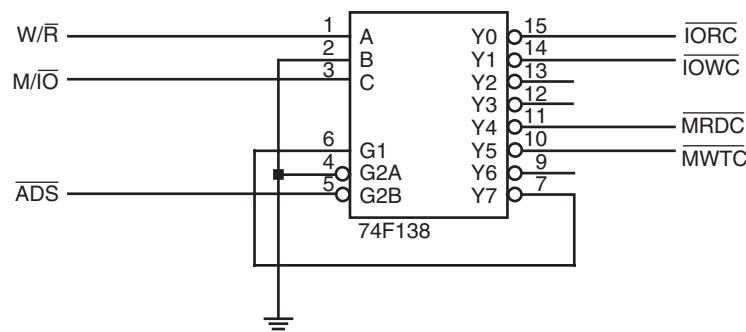
The only new feature that was added to the 80386 with respect to I/O is the I/O privilege information added to the tail end of the TSS when the 80386 is operated in protected mode. As described in the section on memory management, an I/O location can be blocked or inhibited in the protected mode. If the blocked I/O location is addressed, an interrupt (type 13, general fault) is generated. This scheme is added so that I/O access can be prohibited in a multiuser environ- ment. Blocking is an extension of the protected mode operation, as are privilege levels.

Memory and I/O Control Signals

The memory and I/O are controlled with separate signals. The $M/\overline{IO}$ signal indicates whether the data transfer is between the microprocessor and the memory ($M/\overline{IO} = 1$) or I/O ($M/\overline{IO} = 0$). In addition to $M/\overline{IO}$, the memory and I/O systems must read or write data. The $W/\overline{R}$ signal is a logic 0 for a read operation and a logic 1 for a write operation. The $\overline{ADS}$ signal is used to qualify the $M/\overline{IO}$ and $W/\overline{R}$ control signals. This is a slight deviation from earlier Intel microprocessors, which didn't use $\overline{ADS}$ for qualification.

See Figure 17–9 for a simple circuit that generates four control signals for the memory and I/O devices in the system. Notice that two control signals are developed for memory control ($\overline{MRDC}$

FIGURE 17–9 Generation of memory and I/O control signals for the 80386, 80486, and Pentium.

and $\overline{\text{MWTC}}$) and two for I/O control ($\overline{\text{IORC}}$ and $\overline{\text{IOWC}}$). These signals are consistent with the memory and I/O control signals generated for use in earlier versions of the Intel microprocessor.
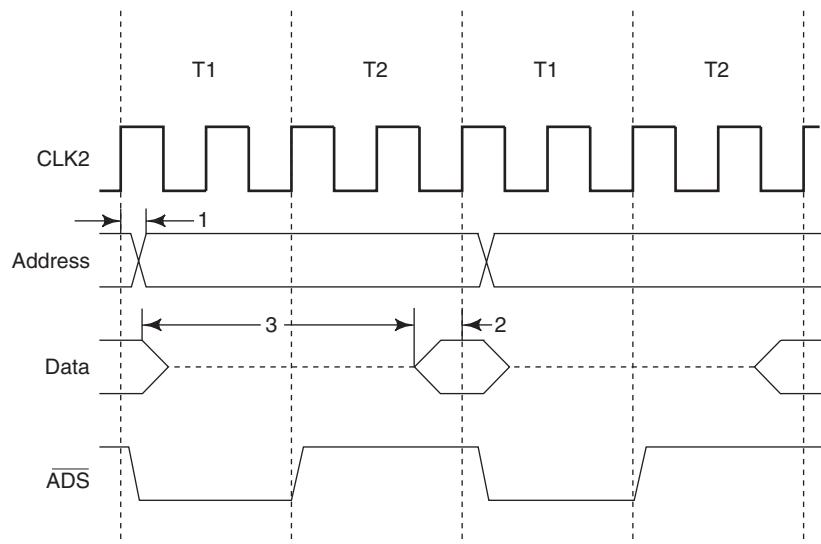
Timing

Timing is important for understanding how to interface memory and I/O to the 80386 microprocessor. Figure 17–10 shows the timing diagram of a nonpipelined memory read cycle. Note that the timing is referenced to the CLK_2 input signal and that a bus cycle consists of four clocking periods.

Each bus cycle contains two clocking states with each state (T_1 and T_2) containing two clocking periods. Note in Figure 17–10 that the access time is listed as time number 3. The 16 MHz version allows memory an access time of 78 ns before wait states are inserted in this nonpipelined mode of operation. To select the nonpipelined mode, we place a logic 1 on the $\overline{\text{NA}}$ pin.
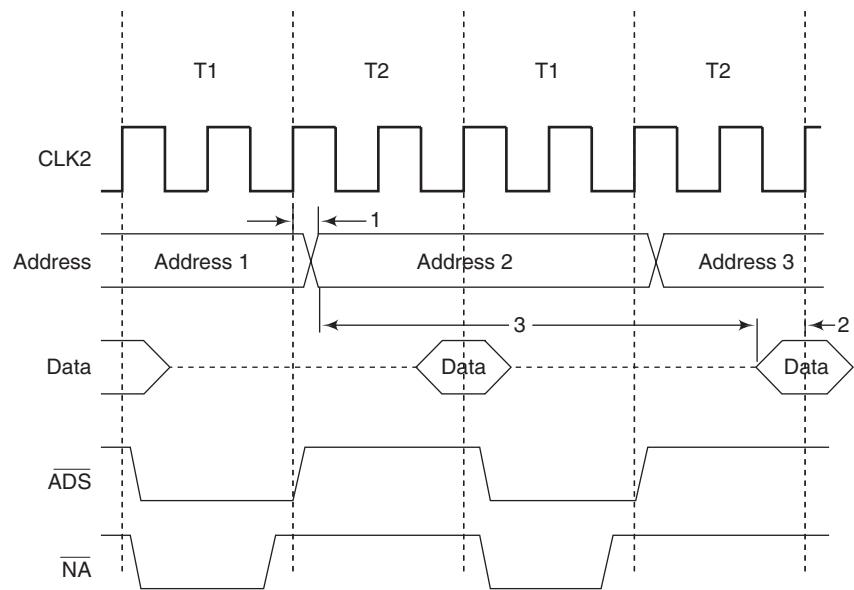
Figure 17–11 illustrates the read timing when the 80386 is operated in the pipelined mode. Notice that additional time is allowed to the memory for accessing data because the address is sent out early. Pipelined mode is selected by placing a logic 0 on the $\overline{\text{NA}}$ pin and by using address latches to capture the pipelined address. The clock pulse that is applied to the address

FIGURE 17–10 The non-pipelined read timing for the 80386 microprocessor.

	33 MHz	25 MHz	20 MHz	16 MHz
Time 1:	4–15 ns	4–21 ns	4–30 ns	4–36 ns
Time 2:	5 ns	7 ns	11 ns	11 ns
Time 3:	46 ns	52 ns	59 ns	78 ns

FIGURE 17–11 The pipelined read timing for the 80386 microprocessor.

latches comes from the $\overline{\text{ADS}}$ signal. Address latches must be used with a pipelined system, as well as with interleaved memory banks. The minimum number of interleaved banks of two and four have been successfully used in some applications.

Notice that the pipelined address appears one complete clocking state before it normally appears with nonpipelined addressing. In the 16 MHz version of the 80386, this allows an additional 62.5 ns for memory access. In a nonpipelined system, a memory access time of 78 ns is allowed to the memory system; in a pipelined system, 140.5 ns is allowed. The advantages of the pipelined system are that no wait states are required (in many, but not all bus cycles) and much lower-speed memory devices may be connected to the microprocessor. The disadvantage is that we need to interleave memory to use a pipe, which requires additional circuitry and occasional wait states.
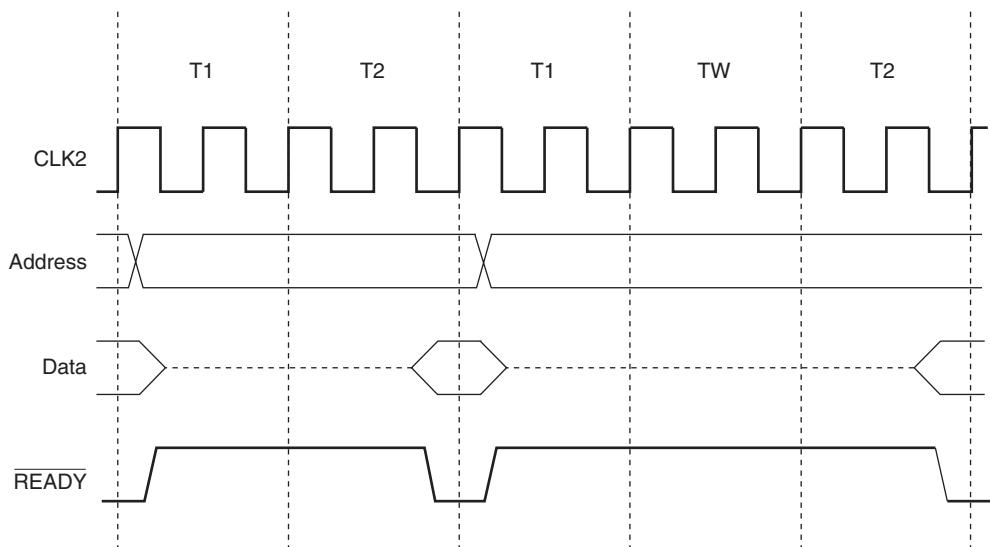
FIGURE 17–12 A nonpipelined 80386 with 0 and 1 wait states.

Wait States

Wait states are needed if memory access times are long compared with the time allowed by the 80386 for memory access. In a nonpipelined 33 MHz system, memory access time is only 46 ns. Currently, only a few DRAM memories exist that have an access time of 46 ns. This means that often wait states must be introduced to access the DRAM (one wait for 60 ns DRAM) or an EPROM that has an access time of 100 ns (two waits). Note that this wait state is built into a motherboard and cannot be removed.

The $\overline{\text{READY}}$ input controls whether or not wait states are inserted into the timing. The $\overline{\text{READY}}$ input on the 80386 is a dynamic input that must be activated during each bus cycle. Figure 17–12 on the previous page shows a few bus cycles with one normal (no wait) cycle and one that contains a single wait state. Notice how the $\overline{\text{READY}}$ is controlled to cause 0 or 1 wait.

(a)

(b)
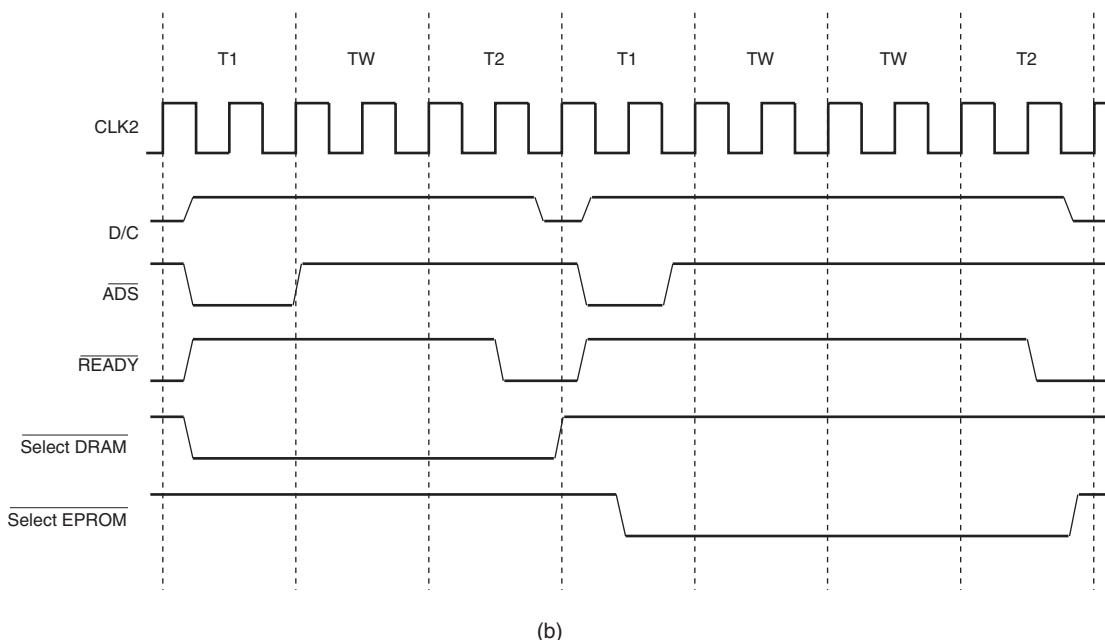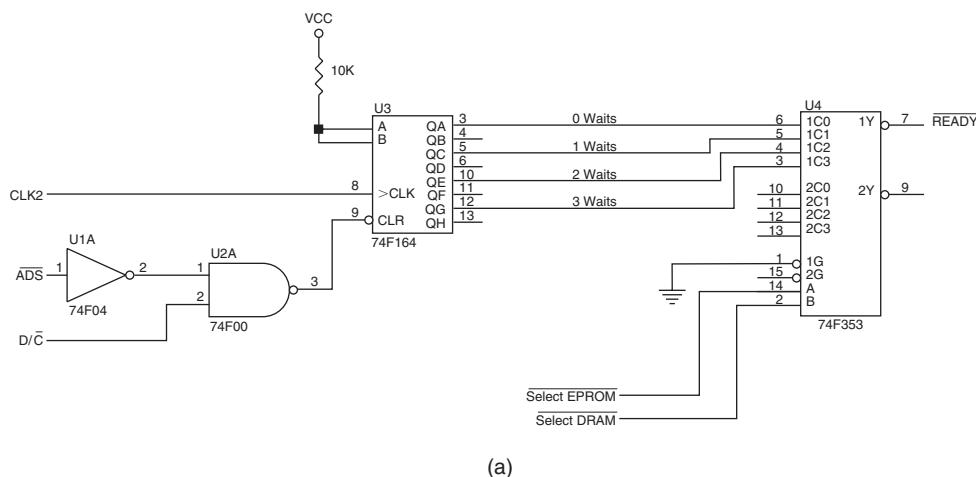
FIGURE 17–13   (a) Circuit and (b) timing that selects 1 wait state for DRAM and 2 waits for EPROM.

The $\overline{\text{READY}}$ signal is sampled at the end of a bus cycle to determine whether the clock cycle is T_2 or TW. If $\overline{\text{READY}} = 0$ at this time, it is the end of the bus cycle or T_2. If $\overline{\text{READY}}$ is 1 at the end of a clock cycle, the cycle is a TW and the microprocessor continues to test $\overline{\text{READY}}$, searching for a logic 0 and the end of the bus cycle.

In the nonpipelined system, whenever $\overline{\text{ADS}}$ becomes a logic 0, a wait state is inserted if $\overline{\text{READY}} = 1$. After $\overline{\text{ADS}}$ returns to a logic 1, the positive edges of the clock are counted to generate the $\overline{\text{READY}}$ signal. The $\overline{\text{READY}}$ signal becomes a logic 0 after the first clock to insert 0 wait states. If one wait state is inserted, the $\overline{\text{READY}}$ line must remain a logic 1 until at least two clocks have elapsed. If additional wait states are desired, then additional time must elapse before $\overline{\text{READY}}$ is cleared. This essentially allows any number of wait states to be inserted into the timing.

Figure 17–13 on the previous page shows a circuit that inserts 0 through 3 wait states for various memory addresses. In the example, one wait state is produced for a DRAM access and two wait states for an EPROM access. The 74F164 clears whenever $\overline{\text{ADS}}$ is low and $D/\overline{C}$ is high. It begins to shift after $\overline{\text{ADS}}$ returns to a logic 1 level. As it shifts, the 00000000 in the shift register begins to fill with logic 1s from the QA connection toward the QH connection. The four different outputs are connected to an inverting multiplexer that generates the active low $\overline{\text{READY}}$ signal.

17–2   SPECIAL 80386 REGISTERS
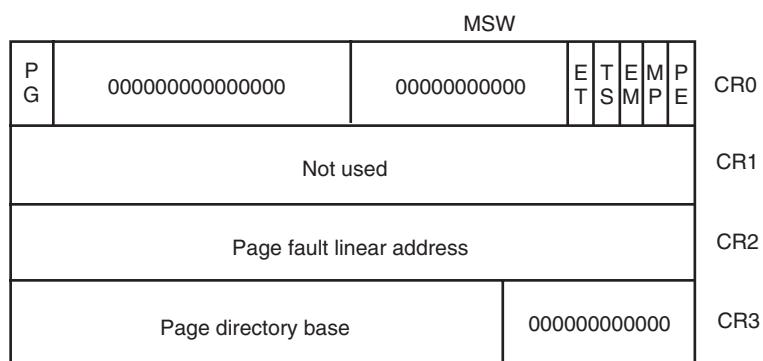
A new series of registers, not found in earlier Intel microprocessors, appears in the 80386 as control, debug, and test registers. Control registers CR_0–CR_3 control various features, DR_0–DR_7 facilitate debugging, and registers TR_6 and TR_7 are used to test paging and caching.

Control Registers

In addition to the EFLAGS and EIP as described earlier, there are other control registers found in the 80386. Control register 0 (CR_0) is identical to the MSW (machine status word) found in the 80286 microprocessor, except that it is 32 bits wide instead of 16 bits wide. Additional control registers are CR_1, CR_2, and CR_3.

Figure 17–14 illustrates the control register structure of the 80386. Control register CR_1 is not used in the 80386, but is reserved for future products. Control register CR_2 holds the linear page address of the last page accessed before a page fault interrupt. Finally, control register CR_3 holds the base address of the page directory. The rightmost 12 bits of the 32-bit page table address contain zeros and combine with the remainder of the register to locate the start of the 4K-long page table.

FIGURE 17–14   The control register structure of the 80386 microprocessor.

Register CR_0 contains a number of special control bits that are defined as follows in the 80386:

PG   Selects **page table translation** of linear addresses into physical addresses when PG = 1. Page table translation allows any linear address to be assigned any physical memory location.

ET   Selects the **80287 coprocessor** when ET = 0 or the **80387 coprocessor** when ET = l. This bit was installed because there was no 80387 available when the 80386 first appeared. In most systems, ET is set to indicate that an 80387 is present in the system.

TS   Indicates that the 80386 has **switched tasks** (in protected mode, changing the contents of TR places a 1 into TS). If TS = 1, a numeric coprocessor instruction causes a type 7 (coprocessor not available) interrupt.

EM   The **emulate** bit is set to cause a type 7 interrupt for each ESC instruction. (ESCape instructions are used to encode instructions for the 80387 coprocessor.) Once this feature was used to emulate interrupts with software, the function of the coprocessor. Emulation reduces the system cost, but it often requires at least 100 times longer to execute the emulated coprocessor instructions.

MP   Is set to indicate that the arithmetic **coprocessor is present** in the system.

PE   Is set to select the **protected mode** of operation for the 80386. It may also be cleared to reenter the real mode. This bit can only be set in the 80286. The 80286 could not return to real mode without a hardware reset, which precludes its use in most systems that use protected mode.

Debug and Test Registers

Figure 17–15 shows the sets of debug and test registers. The first four debug registers contain 32-bit linear breakpoint addresses. (A linear address is a 32-bit address generated by a microprocessor instruction that may or may not be the same as the physical address.) The breakpoint addresses, which may locate an instruction or datum, are constantly compared with the addresses generated by the program. If a match occurs, the 80386 will cause a type 1 interrupt (TRAP or

31	16 15	0	
BREAKPOINT 0 LINEAR ADDRESS			DR0
BREAKPOINT 1 LINEAR ADDRESS			DR1
BREAKPOINT 2 LINEAR ADDRESS			DR2
BREAKPOINT 3 LINEAR ADDRESS			DR3
Intel reserved. Do not define.			DR4
Intel reserved. Do not define.			DR5
0 — BT BS BD 0 0 0 0 0 0 0 0 0 0 B3 B2 B1 B0			DR6
LEN3 R3 W3 LEN2 R2 W2 LEN1 R1 W1 LEN0 R0 W0 0 0 GD 0 0 0 GE LE G3 L3 G2 L2 G1 L1 G0 L0			DR7
31	16 15	0	

31	12 11	0	
LINEAR ADDRESS	V D D# U U# W W# 0 0 0 0 C		TR6
PHYSICAL ADDRESS	0 0 0 0 0 0 0 PL REP 0 0		TR7

FIGURE 17–15   The debug and test registers of the 80386. (Courtesy of Intel Corporation.)

debug interrupt) to occur, if directed by debug registers DR_6 and DR_7. This feature is a much-expanded version of the basic trapping or tracing allowed with the earlier Intel microprocessors through the type 1 interrupt. The breakpoint addresses are very useful in debugging faulty software. The control bits in DR_6 and DR_7 are defined as follows:

BT   If set (1), the debug interrupt was caused by a task switch.

BS   If set, the debug interrupt was caused by the TF bit in the flag register.

BD   If set, the debug interrupt was caused by an attempt to read the debug register with the GD bit set. The GD bit protects access to the debug registers.

B_3–B_0   Indicate which of the four debug breakpoint addresses caused the debug interrupt.

LEN   Each of the four length fields pertains to each of the four breakpoint addresses stored in DR_0–DR_3. These bits further define the size of access at the breakpoint address as 00 (byte), 01 (word), or 11 (doubleword).

RW   Each of the four read/write fields pertains to each of the four breakpoint addresses stored in DR_0–DR_3. The RW field selects the cause of action that enabled a breakpoint address as 00 (instruction access), 01 (data write), and 11 (data read and write).

GD   If set, GD prevents any read or write of a debug register by generating the debug interrupt. This bit is automatically cleared during the debug interrupt so that the debug registers can be read or changed, if needed.

GE   If set, selects a global breakpoint address for any of the four breakpoint address registers.

LE   If set, selects a local breakpoint address for any of the four breakpoint address registers.

The test registers, TR_6 and TR_7, are used to test the translation look-aside buffer (TLB). The TLB is used with the paging unit within the 80386. The TLB, which holds the most commonly used page table address translations, reduces the number of memory reads required for looking up page translation addresses in the page translation tables. The TLB holds the most common 32 entries from the page table, and it is tested with the TR_6 and TR_7 test registers.

Test register TR_6 holds the tag field (linear address) of the TLB, and TR_7 holds the physical address of the TLB. To write a TLB entry, perform the following steps:

1. Write TR_7 for the desired physical address, PL, and REP values.
2. Write TR_6 with the linear address, making sure that C = 0.

    To read a TLB entry:

1. Write TR_6 with the linear address, making sure that C = 1.
2. Read both TR_6 and TR_7. If the PL bit indicates a hit, then the desired values of TR_6 and TR_7 indicate the contents of the TLB.

The bits found in TR_6 and TR_7 indicate the following conditions:

V   Shows that the entry in the TLB is valid.

D   Indicates that the entry in the TLB is invalid or dirty.

U   A bit for the TLB.

W   Indicates that the area addressed by the TLB entry is writable.

C   Selects a write (0) or immediate lookup (1) for the TLB.

PL   Indicates a hit if a logic 1.

REP   Selects which block of the TLB is written.

Refer to the section on memory management and the paging unit for more detail on the function of the TLB.

17–3   80386 MEMORY MANAGEMENT

The memory-management unit (MMU) within the 80386 is similar to the MMU inside the 80286, except that the 80386 contains a paging unit not found in the 80286. The MMU performs the task of converting linear addresses, as they appear as outputs from a program, into physical addresses that access a physical memory location located anywhere within the memory system. The 80386 uses the paging mechanism to allocate any physical address to any logical address. Therefore, even though the program is accessing memory location A0000H with an instruction, the actual physical address could be memory location 100000H, or any other location if paging is enabled. This feature allows virtually any software, written to operate at any memory location, to function in an 80386 because any linear location can become any physical location. Earlier Intel microprocessors did not have this flexibility. Paging is used with DOS to relocate 80386 and 80486 memory at addresses above FFFFFH and into spaces between ROMs at locations D0000–DFFFFH and other areas as they are available. The area between ROMs is often referred to as *upper memory;* the area above FFFFFH is referred to as *extended memory*.

Descriptors and Selectors

Before the memory paging unit is discussed, we examine the descriptor and selector for the 80386 microprocessor. The 80386 uses descriptors in much the same fashion as the 80286. In both microprocessors, a descriptor is a series of eight bytes that describes and locates a memory segment. A selector (segment register) is used to index a descriptor from a table of descriptors. The main difference between the 80286 and 80386 is that the latter has two additional selectors (FS and GS) and the most significant two bytes of the descriptor are defined for the 80386. Another difference is that 80386 descriptors use a 32-bit base address and a 20-bit limit, instead of the 24-bit base address and a 16-bit limit found on the 80286.
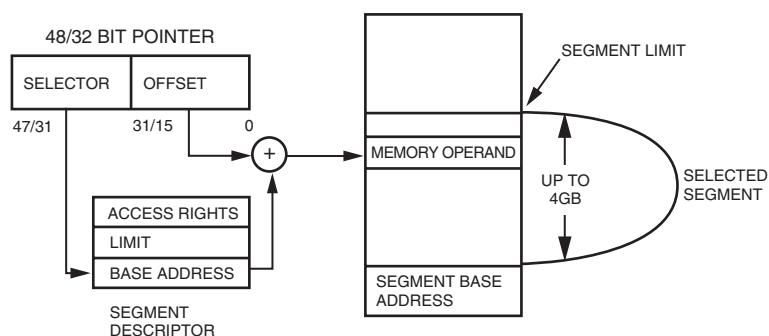
The 80286 addresses a 16M-byte memory space with its 24-bit base address and has a segment length limit of 64K bytes, due to the 16-bit limit. The 80386 addresses a 4G-byte memory space with its 32-bit base address and has a segment length limit of 1M byte or 4G bytes, due to a 20-bit limit that is used in two different ways. The 20-bit limit can access a segment with a length of 1M byte if the granularity bit (G) = 0. If G = 1, the 20-bit limit allows a segment length of 4G bytes.

The granularity bit is found in the 80386 descriptor. If G = 0, the number stored in the limit is interpreted directly as a limit, allowing it to contain any limit between 00000H and FFFFFH for a segment size up to 1M byte. If G = 1, the number stored in the limit is interpreted as 00000XXXH–FFFFFXXXH, where the XXX is any value between 000H and FFFH. This allows the limit of the segment to range between 0 bytes to 4G bytes in steps of 4K bytes. A limit of 00001H indicates that the limit is 4K bytes when G = 1 and 1 byte when G = 0. An example is a segment that begins at physical address 10000000H. If the limit is 00001H and G = 0, this segment begins at 10000000H and ends at 10000001H. If G = 1 with the same limit (00001H), the segment begins at location 10000000H and ends at location 10001FFFH.

Figure 17–16 shows how the 80386 addresses a memory segment in the protected mode using a selector and a descriptor. Note that this is identical to the way that a segment is addressed by the 80286. The difference is the size of the segment accessed by the 80386. The selector uses its leftmost 13 bits to select a descriptor from a descriptor table. The TI bit indicates either the local (TI = 1) or global (TI = 0) descriptor table. The rightmost two bits of the selector define the requested privilege level of the access.

Because the selector uses a 13-bit code to access a descriptor, there are at most 8192 descriptors in each table—local or global. Because each segment (in an 80386) can be 4G bytes in length, 16,384 segments can be accessed at a time with the two descriptor tables. This allows the 80386 to access a virtual memory size of 64T bytes. Of course, only 4G bytes of memory actually exist in the memory system (1T byte = 1024G bytes). If a program requires more than

FIGURE 17–16 Protected mode addressing using a segment register as a selector. (Courtesy of Intel Corporation.)

4G bytes of memory at a time, it can be swapped between the memory system and a disk drive or other form of large volume storage.
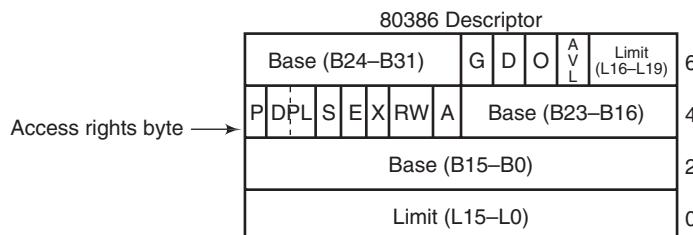
The 80386 uses descriptor tables for both global (GDT) and local (LDT) descriptors. A third descriptor table appears for interrupt (IDT) descriptors or gates. The first six bytes of the descriptor are the same as in the 80286, which allows 80286 software to be upward compatible with the 80386. (An 80286 descriptor used 00H for its most significant two bytes.) See Figure 17–17 for the 80286 and 80386 descriptor. The base address is 32 bits in the 80386, the limit is 20 bits, and a G bit selects the limit multiplier (1 or 4K times). The fields in the descriptor for the 80386 are defined as follows:

Limit (L_{19}–L_0)	Defines the starting 32-bit address of the segment within the 4G-byte physical address space of the 80386 microprocessor. Defines the limit of the segment in units of bytes if the G bit = 0, or in units of 4K bytes if G = 1. This allows a segment to be of any length from 1 byte to 1M byte if G = 0, and from 4K bytes to 4G bytes if G = 1. Recall that the limit indicates the last byte in a segment.
Access Rights	Determines privilege level and other information about the segment. This byte varies with different types of descriptors and is elaborated with each descriptor type.
G	The granularity bit selects a multiplier of 1 or 4K times for the limit field. If G = 0, the multiplier is 1; if G = 1, the multiplier is 4K.
D	Selects the default instruction mode. If D = 0, the registers and memory pointers are 16 bits wide, as in the 80286; if D = 1, they are 32 bits wide, as in the 80386. This bit determines whether prefixes are required for 32-bit data and index registers. If D = 0, then a prefix is required to access 32-bit registers and to use 32-bit pointers. If D = 1, then a prefix is required to access 16-bit registers and 16-bit pointers. The USE_{16} and USE_{32} directives appended to the SEGMENT statement in assembly language control the setting of the D bit. In the real mode, it is always

FIGURE 17–17 The descriptors for the 80286 and 80386 microprocessors.

FIGURE 17–18 The format
of the 80386 segment
descriptor.

assumed that the registers are 16 bits wide, so any instruction that refer-
ences a 32-bit register or pointer must be prefixed. The current version of
DOS assumes D = 0, and most Windows programs assume D = 1.

AVL   This bit is available to the operating system to use in any way that it sees
fit. It often indicates that the segment described by the descriptor is avail-
able.

Descriptors appear in two forms in the 80386 microprocessor: the segment descriptor and
the system descriptor. The segment descriptor defines data, stack, and code segments; the system
descriptor defines information about the system's tables, tasks, and gates.

Segment Descriptors.   Figure 17–18 shows the segment descriptor. This descriptor fits the gen-
eral form, as dictated in Figure 17–17, but the access rights bits are defined to indicate how the
data, stack, or code segment described by the descriptor functions. Bit position 4 of the access
rights byte determines whether the descriptor is a data or code segment descriptor (S = 1) or a
system segment descriptor (S = 0). Note that the labels used for these bits may vary in different
versions of Intel literature, but they perform the same tasks.

Following is a description of the access rights bits and their function in the segment
descriptor:

P   **Present** is a logic 1 to indicate that the segment is present. If P = 0 and the segment
is accessed through the descriptor, a type 11 interrupt occurs. This interrupt indi-
cates that a segment was accessed that is not present in the system.

DPL   **Descriptor privilege level** sets the privilege level of the descriptor; 00 has the high-
est privilege and 11 has the lowest. This is used to protect access to segments. If a
segment is accessed with a privilege level that is lower (higher in number) than the
DPL, a privilege violation interrupt occurs. Privilege levels are used in multiuser
systems to prevent access to an area of the system memory.

S   **Segment** indicates a data or code segment descriptor (S = 1), or a system segment
descriptor (S = 0).

E   **Executable** selects a data (stack) segment (E = 0) or a code segment (E = 1). E also
defines the function of the next two bits (X and RW).

X   If E = 0, then X indicates the direction of **expansion** for the data segment. If X = 0,
the segment expands upward, as in a data segment; if X = 1, the segment expands
downward as in a stack segment. If E = 1, then X indicates whether the privilege
level of the code segment is ignored (X = 0) or observed (X = 1).

RW   If E = 0, then the **read/write** bit (RW) indicates that the data segment may be writ-
ten (RW = 1) or not written (RW = 0). If E = 1, then RW indicates that the code
segment may be read (RW = 1) or not read (RW = 0).

A   **Accessed** is set each time that the microprocessor accesses the segment. It is some-
times used by the operating system to keep track of which segments have been
accessed.

FIGURE 17–19 The general format of an 80386 system descriptor.

80386 Descriptor

Access rights byte →
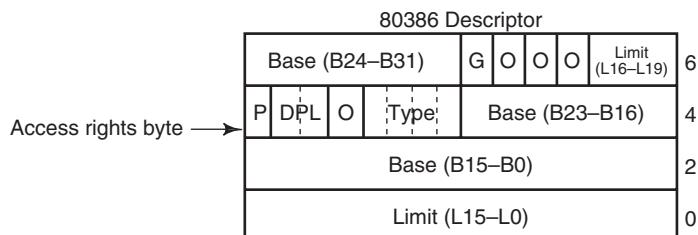
TABLE 17–1 System descriptor types.

Type	Purpose
0000	Invalid
0001	Available 80286 TSS
0010	LDT
0011	Busy 80286 TSS
0100	80286 task gate
0101	Task gate (80386 and above)
0110	80286 interrupt gate
0111	80286 trap gate
1000	Invalid
1001	Available 80386 and above TSS
1010	Reserved
1011	Busy 80386 and above TSS
1100	80386 and above CALL gate
1101	Reserved
1110	80386 and above interrupt gate
1111	80386 and above trap gate

System Descriptor. The system descriptor is illustrated in Figure 17–19. There are 16 possible system descriptor types (see Table 17–1 for the different descriptor types), but not all are used in the 80386 microprocessor. Some of these types are defined for the 80286 so that the 80286 software is compatible with the 80386. Some of the types are new and unique to the 80386; some have yet to be defined and are reserved for future Intel products.
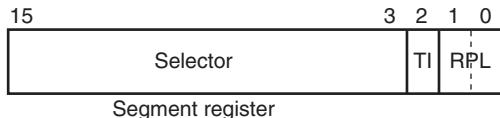
Descriptor Tables

The descriptor tables define all the segments used in the 80386 when it operates in the protected mode. There are three types of descriptor tables: the global descriptor table (GDT), the local descriptor table (LDT), and the interrupt descriptor table (IDT). The registers used by the 80386 to address these three tables are called the global descriptor table register (GDTR), the local descriptor table register (LDTR), and the interrupt descriptor table register (IDTR). These registers are loaded with the LGDT, LLDT, and LIDT instructions, respectively.

The descriptor table is a variable-length array of data, with each entry holding an 8-byte long descriptor. The local and global descriptor tables hold up to 8192 entries each, and the interrupt descriptor table holds up to 256 entries. A descriptor is indexed from either the local or global descriptor table by the selector that appears in a segment register. Figure 17–20 shows a segment register and the selector that it holds in the protected mode. The leftmost 13 bits index a descriptor; the TI bit selects either the local (TI = 1) or global (TI = 1) descriptor table; and the RPL bits indicate the requested privilege level.

Whenever a new selector is placed into one of the segment registers, the 80386 accesses one of the descriptor tables and automatically loads the descriptor into a program-invisible cache portion of the segment register. As long as the selector remains the same in the segment register, no additional accesses are required to the descriptor table. The operation of fetching a new descriptor

FIGURE 17–20 A segment register showing the selector, T_1 bit, and requested privilege level (RPL) bits.

from the descriptor table is program-invisible because the microprocessor automatically accomplishes this each time that the segment register contents are changed in the protected mode.

Figure 17–21 shows how a sample global descriptor table (GDT), which is stored at memory address 00010000H, is accessed through the segment register and its selector. This table contains four entries. The first is a null (0) descriptor. Descriptor 0 must always be a null descriptor. The other entries address various segments in the 80386 protected mode memory system. In this illustration, the data segment register contains 0008H. This means that the selector is indexing descriptor location 1 in the global descriptor table (TI = 0), with a requested privilege level of 00. Descriptor 1 is located eight bytes above the base descriptor table address, beginning at location 00010008H. The descriptor located in this memory location accesses a base address of 00200000H and a limit of 100H. This means that this descriptor addresses memory locations 00200000H–00200100H. Because this is the DS (data segment) register, the data segment is located at these locations in the memory system. If data are accessed outside of these boundaries, an interrupt occurs.

The local descriptor table (LDT) is accessed in the same manner as the global descriptor table (GDT). The only difference in access is that the TI bit is cleared for a global access and set for a local access. Another difference exists if the local and global descriptor table registers are examined. The global descriptor table register (GDTR) contains the base address of the global
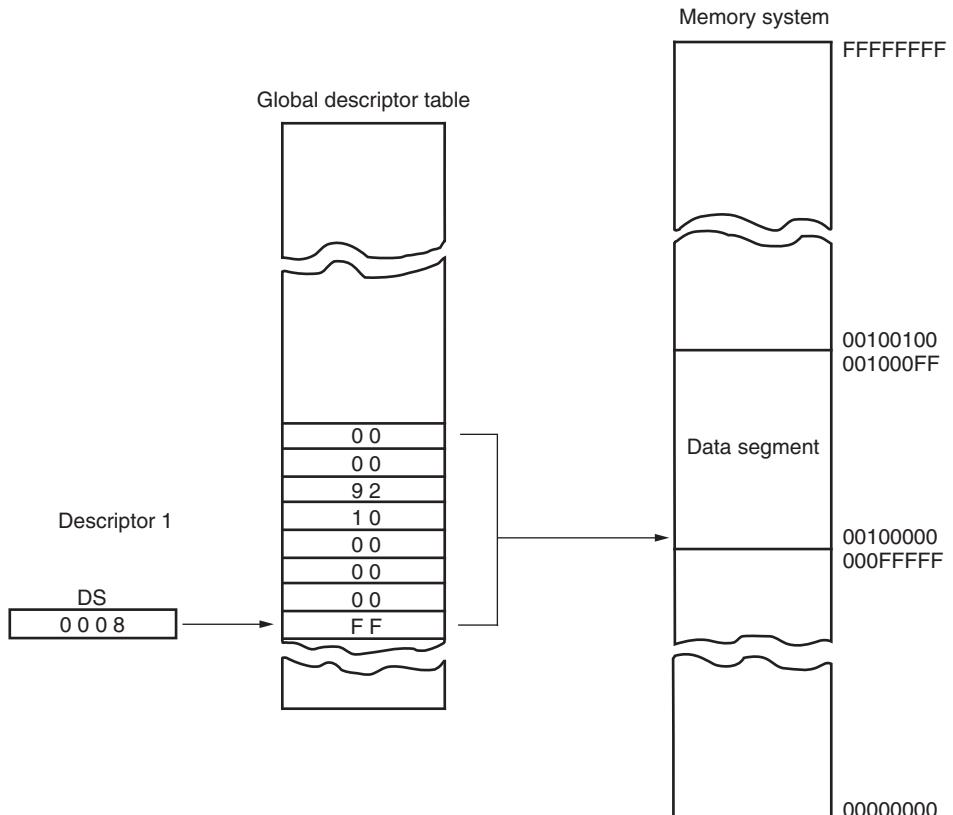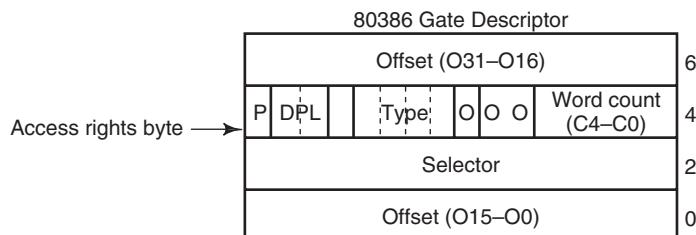
FIGURE 17–21 Using the DS register to select a descriptor from the global descriptor table. In this example, the DS register accesses memory locations 00100000H–001000FFH as a data segment.

FIGURE 17–22   The gate descriptor for the 80386 microprocessor.

descriptor table and the limit. The local descriptor table register (LDTR) contains only a selector, and it is 16 bits wide. The contents of the LDTR addresses a type 0010 system descriptor that contains the base address and limit of the LDT. This scheme allows one global table for all tasks, but allows many local tables, one or more for each task, if necessary. Global descriptors describe memory for the system, while local descriptors describe memory for applications or tasks.

Like the GDT, the interrupt descriptor table (IDT) is addressed by storing the base address and limit in the interrupt descriptor table register (IDTR). The main difference between the GDT and IDT is that the IDT contains only interrupt gates. The GDT and LDT contain segment and system descriptors, but never contain interrupt gates.

Figure 17–22 shows the gate descriptor, a special form of the system descriptor described earlier. (Refer to Table 17–1 for the different gate descriptor types.) Notice that the gate descriptor contains a 32-bit offset address, a word count, and a selector. The 32-bit offset address points to the location of the interrupt service procedure or other procedure. The word count indicates how many words are transferred from the caller's stack to the stack of the procedure accessed by a call gate. This feature of transferring data from the caller's stack is useful for implementing high-level languages such as C/C++. Note that the word count field is not used with an interrupt gate. The selector is used to indicate the location of task state segment (TSS) in the GDT or LDT if it is a local procedure.

When a gate is accessed, the contents of the selector are loaded into the task register (TR), causing a task switch. The acceptance of the gate depends on the privilege and priority levels. A return instruction (RET) ends a call gate procedure and a return from interrupt instruction (IRET) ends an interrupt gate procedure. Tasks are usually accessed with a CALL or an INT instruction, where the call instruction addresses a call gate in the descriptor table and the interrupt addresses an interrupt descriptor.
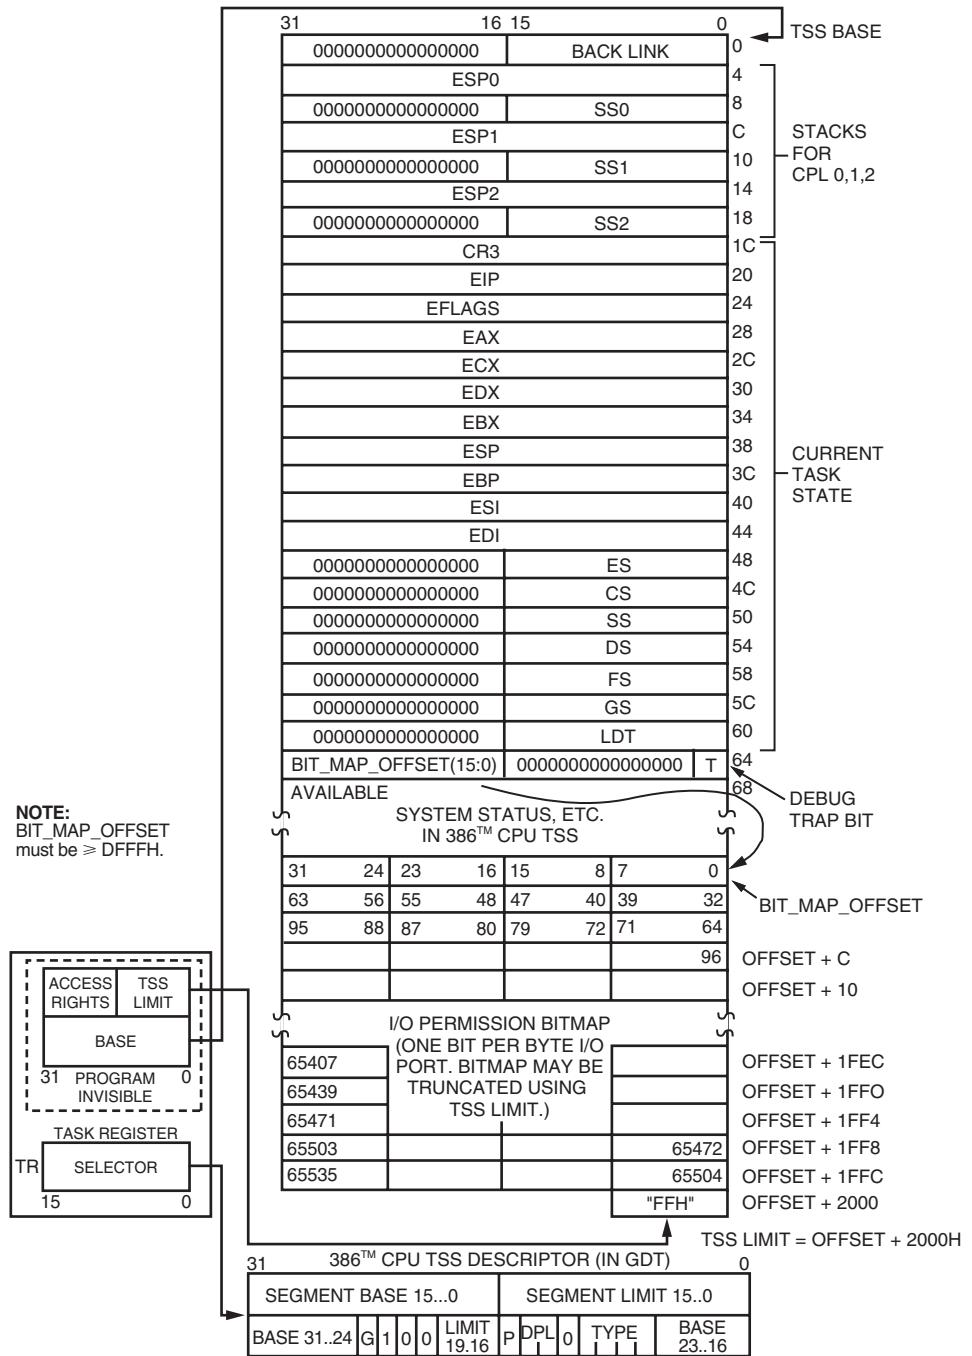
The difference between real mode interrupts and protected mode interrupts is that the interrupt vector table is an IDT in the protected mode. The IDT still contains up to 256 interrupt levels, but each level is accessed through an interrupt gate instead of an interrupt vector. Thus, interrupt type number 2 (INT 2) is located at IDT descriptor number 2 at 16 locations above the base address of the IDT. This also means that the first IK byte of memory no longer contains interrupt vectors, as it did in the real mode. The IDT can be located at any location in the memory system.

The Task State Segment (TSS)

The task state segment (TSS) descriptor contains information about the location, size, and privilege level of the task state segment, just like any other descriptor. The difference is that the TSS described by the TSS descriptor does not contain data or code. It contains the state of the task and linkage so tasks can be nested (one task can call a second, which can call a third, and so forth). The TSS descriptor is addressed by the task register (TR). The contents of the TR are changed by the LTR instruction. Whenever the protected mode program executes a JMP or CALL instruction, the contents of TR are also changed. The LTR instruction is used to initially access a task during system initialization. After initialization, the CALL or JUMP instructions normally switch tasks. In most cases, we use the CALL instruction to initiate a new task.

The TSS is illustrated in Figure 17–23. As can be seen, the TSS is quite a formidable data structure, containing many different types of information. The first word of the TSS is

labeled **back-link.** This is the selector that is used, on a return (RET or IRET), to link back to the prior TSS by loading the back-link selector into the TR. The following word must contain a 0. The second through the seventh doublewords contain the ESP and ESS values for privilege levels 0–2. These are required in case the current task is interrupted so these privilege level

FIGURE 17–23   The task state segment (TSS) descriptor. (Courtesy of Intel Corporation.)

(PL) stacks can be addressed. The eighth word (offset 1CH) contains the contents of CR_3, which stores the base address of the prior state's page directory register. This must be restored if paging is in effect. The contents of the next 17 doublewords are loaded into the registers indicated. Whenever a task is accessed, the entire state of the machine (all of the registers) is stored in these memory locations and then reloaded from the same locations in the new TSS. The last word (offset 66H) contains the I/O permission bit map base address.

The I/O permission bit map allows the TSS to block I/O operations to inhibited I/O port addresses via an I/O permission denial interrupt. The permission denial interrupt is type number 13, the general protection fault interrupt. The I/O permission bit map base address is the offset address from the start of the TSS. This allows the same permission map to be used by many TSSs.

Each I/O permission bit map is 64K bits long (8K bytes), beginning at the offset address indicated by the I/O permission bit map base address. The first byte of the I/O permission bit map contains I/O permission for I/O ports 0000H–0007H. The rightmost bit contains the permission for port number 0000H. The leftmost bit contains the permission for port number 0007H. This sequence continues for the very last port address (FFFFH) stored in the leftmost bit of the last byte of the I/O permission bit map. A logic 0 placed in an I/O permission bit map bit enables the I/O port address, while a logic 1 inhibits or blocks the I/O port address. At present, only Windows NT, Windows 2000, and Windows XP uses the I/O permission scheme to disable I/O ports dependent on the application or the user.

To review the operation of a task switch, which requires only 17 μs to execute on an 80386 microprocessor, we list the following steps:

1. The gate contains the address of the procedure or location jumped to by the task switch. It also contains the selector number of the TSS descriptor and the number of words transferred from the caller to the user stack area for parameter passing.
2. The selector is loaded into TR from the gate. (This step is accomplished by a CALL or JMP that refers to a valid TSS descriptor.)
3. The TR selects the TSS.
4. The current state is saved in the current TSS and the new TSS is accessed with the state of the new task (all the registers) loaded into the microprocessor. The current state is saved at the TSS selector currently found in the TR. Once the current state is saved, a new value (by the JMP or CALL) for the TSS selector is loaded into TR and the new state is loaded from the new TSS.

The return from a task is accomplished by the following steps:

1. The current state of the microprocessor is saved in the current TSS.
2. The back-link selector is loaded to the TR to access the prior TSS so that the prior state of the machine can be returned to and be restored to the microprocessor. The return for a called TSS is accomplished by the IRET instruction.

17–4   MOVING TO PROTECTED MODE

In order to change the operation of the 80386 from the real mode to the protected mode, several steps must be followed. Real mode operation is accessed after a hardware reset or by changing the PE bit to a logic 0 in CR_0. Protected mode is accessed by placing a logic 1 into the PE bit of CR_0; before this is done, however, some other things must be initialized. The following steps accomplish the switch from the real mode to the protected mode:
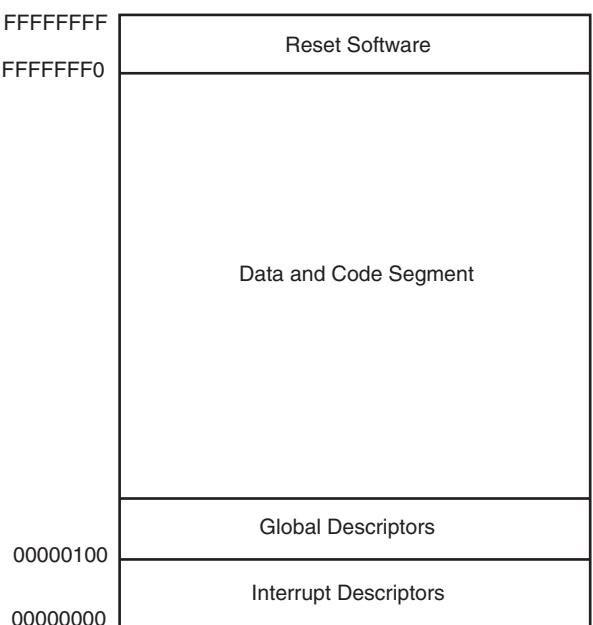
1. Initialize the interrupt descriptor table so that it contains valid interrupt gates for at least the first 32 interrupt type numbers. The IDT may (and often does) contain up to 256 eight-byte interrupt gates defining all 256 interrupt types.

2. Initialize the global descriptor table (GDT) so that it contains a null descriptor at descriptor 0 and valid descriptors for at least one code, one stack, and one data segment.
3. Switch to protected mode by setting the PE bit in CR_0.
4. Perform an intersegment (far) JMP to flush the internal instruction queue.
5. Load all the data selectors (segment registers) with their initial selector values.
6. The 80386 is now operating in the protected mode, using the segment descriptors that are defined in GDT and IDT.

Figure 17–24 shows the protected system memory map set up by following steps 1–6. The software for this task is listed in Example 17–1. This system contains one data segment descriptor and one code segment descriptor with each segment set to 4G bytes in length. This is the simplest protected mode system possible (called the **flat model**): loading all the segment registers, except code, with the same data segment descriptor from the GDT. The privilege level is initialized to 00, the highest level. This system is most often used where one user has access to the microprocessor and requires the entire memory space. This program is designed for use in a system that does not use DOS or does not shell from Windows to DOS. Later in this section, we show how to go to protected mode in a DOS environment. (Please note that the software in Example 17–1 is designed for a stand-alone system such as the 80386EX embedded microprocessor, and not for use in the PC.)

Example 17–1 does not store any interrupt vectors into the interrupt descriptor table, because none are used in the example. If interrupt vectors are used, then software must be included to load the addresses of the interrupt service procedures into the IDT. The software must be generated as two separate parts that are then connected together and burned on a ROM. The first part is written as shown in the real mode and the second part (see comment in listing) with the assembler set to generate protected mode code using the 32-bit flat model. This software will not function on a personal computer because it is written to function on an embedded system. The code must be converted to a binary file using EXE2BIN after it is assembled and before burning to a ROM.

FIGURE 17–24   The memory map for Example 17–1.

EXAMPLE 17–1

```
.MODEL SMALL
.386P
ADR     STRUC
        DW      ?               ;address structure
        DD      ?
ADR     ENDS

.DATA
IDT     DQ      32 DUP(?)       ;interrupt descriptor table
GDT     DQ      8               ;global descriptor table
DESC1   DW      0FFFFH          ;code segment descriptor
        DW      0
        DW      0
        DW      9EH
        DW      8FH
        DW      0
DESC2   DW      0FFFFH          ;data segment descriptor
        DW      0
        DW      0
        DW      92H
        DW      8FH
        DW      0
IDTR    ADR     <0FFH,IDT>      ;IDTR data
GDTR    ADR     <17H,GDT>       ;GDTR data
JADR    ADR     <8,PM>          ;far JMP data
.CODE
.STARTUP
        MOV  AX,0               ;initialize DS
        MOV  DS,AX

        LIDT IDTR               ;initialize IDTR
        LGDT GDTR               ;initialize GDTR

        MOV  EAX,CR0            ;set PE
        OR   EAX,1
        MOV  CR0,EAX

        JMP  JADR               ;far jump to PM
PM::                            ;force a far label

;the software that follows must be developed separately
;with the assembler set to generate 32-bit protected mode code
;ie:    .MODEL FLAT

        MOV  AX,10H             ;load segment registers
        MOV  DS,AX              ;now in protected mode
        MOV  ES,AX
        MOV  SS,AX
        MOV  FS,AX
        MOV  GS,AX
        MOV  SP,0FFFF000H

;other initialization appears here

end
```

In more complex systems (very unlikely to appear in embedded systems), the steps required to initialize the system in the protected mode are more involved. For complex systems that are often multiuser systems, the registers are loaded by using the task state segment (TSS). The steps required to place the 80386 into protected mode operation for a more complex system using a task switch follow:

1. Initialize the interrupt descriptor table so that it refers to valid interrupt descriptors with at least 32 descriptors in the IDT.

2. Initialize the global descriptor table so that it contains a task state segment (TSS) descriptor, and the initial code and data segments required for the initial task.
3. Initialize the task register (TR) so that it points to a TSS.
4. Switch to protected mode by using an intersegment (far) jump to flush the internal instruction queue. This loads the TR with the current TSS selector and initial task.
5. The 80386 is now operating in the protected mode under control of the first task.

Example 17–2 illustrates the software required to initialize the system and switch to protected mode by using a task switch. The initial system task operates at the highest level of protection (00) and controls the entire operating environment for the 80386. In many cases, it is used to boot (load) software that allows many users to access the system in a multiuser environment. As with Example 17–2 this software will not function on a personal computer and is designed to function only on an embedded system.

EXAMPLE 17–2

```
.MODEL SMALL
.386P
.DATA
ADR     STRUC                   ;structure for 48 bit address
        DW      ?               ;selector
        DD      ?               ;offset
ADR     ENDS

DESC    STRUC                   ;structure of a descriptor
        DW      ?
        DW      ?
        DB      ?
        DB      ?
        DB      ?
        DB      ?
DESC    ENDS

TSS     STRUC                   ;structure of TSS
        DD      18 DUP(?)
        DD      18H             ;ES
        DD      10H             ;CS
        DD      4 DUP(18H)
        DD      28H             ;LDT
        DD      IOBP            ;IO privilege map
TSS     ENDS

GDT     DESC    <>                      ;null
        DESC    <2067H,TS1,0,89H,90H,0>  ;TSS descriptor
        DESC    <-1,0,0,9AH,0CFH,0>      ;code segment
        DESC    <-1,0,0,92H,0CFH,0>      ;data segment
        DESC    <0,0,0,0,0,0>            ;LDT for TSS

LDT     DESC    <>              ;null

IOBP    DB      2000H DUP(0)    ;all I/O on

IDT     DQ      32 DUP(?)       ;IDT

TS1     TSS     <>              ;make TSS

IDTA    ADR     <0FFH,IDT>      ;IDTR
GDTA    ADR     <27H,GDT>       ;GDTR
JADR    ADR     <10H,PM>        ;jump address

.CODE
.STARTUP

        MOV  AX,0
        MOV  DS,AX
```

```
                    LGDT  GDTA
                    LIDT  IDTA

                    MOV   EAX,CR0
                    OR    EAX,1
                    MOV   CR0,EAX

                    MOV   AX,8
                    LTR   AX
                    JMP   JADR
          PM:

                    ;protected mode

          END
```

Neither Example 17–1 nor Example 17–2 is written to function in the personal computer environment. The personal computer environment requires the use of either the VCPI (virtual control program interface) driver provided by the HIMEM.SYS driver in DOS or the DPMI (**DOS protected mode interface**) driver provided by Windows when shelling to DOS. Example 17–3 shows how to switch to protected mode using DPMI and then display the contents of any area of memory. This includes memory in the extended memory area or anywhere else. This DOS application allows the contents of any memory location to be displayed in hexadecimal format on the monitor, including locations above the first 1M byte of the memory system.

EXAMPLE 17–3

```
                    ;A program that displays the contents of any area of memory
                    ;including extended memory.
                    ;***command line syntax***
                    ;EDUMP XXXX,YYYY where XXXX is the start address and YYYY is
                    ;the end address.
                    ;Note: this program must be executed from WINDOWS.
                    ;
                          .MODEL SMALL
                          .386
                          .STACK 1024           ;stack area of 1,024 bytes
0000                      .DATA
0000  00000000       ENTRY DD   ?              ;DPMI entry point
0004  00000000       EXIT  DD   ?              ;DPMI exit point
0008  00000000       FIRST DD   ?              ;first address
000C  00000000       LAST1 DD   ?              ;last address
0010  0000           MSIZE DW   ?              ;memory needed for DPMI
0012  0D 0A 0A 50 61 ERR1  DB   13,10,10,'Parameter error.$'
      72 61 6D 65 74
      65 72 20 65 72
      72 6F 72 2E 24
0026  0D 0A 0A 44 50 ERR2  DB   13,10,10,'DPMI not present.$'
      4D 49 20 6E 6F
      74 20 70 72 65
      73 65 6E 74 2E
      24
003B  0D 0A 0A 4E 6F ERR3  DB   13,10,10,'Not enough real memory.$'
      74 20 65 6E 6F
      75 67 68 20 72
      65 61 6C 20 6D
      65 6D 6F 72 79
      2E 24
0056  0D 0A 0A 43 6F ERR4  DB   13,10,10,'Could not move to protected mode.$'
      75 6C 64 20 6E
      6F 74 20 6D 6F
      76 65 20 74 6F
      20 70 72 6F 74
      65 63 74 65 64
      20 6D 6F 64 65
      2E 24
```

```
007B  0D 0A 0A 43 61   ERR5  DB   13,10,10,'Cannot allocate selector.$'
      6E 6E 6F 74 20
      61 6C 6C 6F 63
      61 74 65 20 73
      65 6C 65 63 74
      6F 72 2E 24
0098  0D 0A 0A 43 61   ERR6  DB   13,10,10,'Cannot use base address.$'
      6E 6E 6F 74 20
      75 73 65 20 62
      61 73 65 20 61
      64 64 72 65 73
      73 2E 24
00B4  0D 0A 0A 43 61   ERR7  DB   13,10,10,'Cannot allocate 64K to limit.$'
      6E 6E 6F 74 20
      61 6C 6C 6F 63
      61 74 65 20 36
      34 4B 20 74 6F
      20 6C 69 6D 69
      74 2E 24
00D5  0D 0A 24         CRLF  DB   13,10,'$'
00D8  50 72 65 73 73   MES1  DB   'Press any key...$'
      20 61 6E 79 20
      6B 65 79 2E 2E
      2E 24
                       ;
                       ;register array storage for DPMI function 0300H
                       ;
00E9  = 00E9           ARRAY EQU  THIS BYTE
00E9  00000000         REDI  DD   0                    ;EDI
00ED  00000000         RESI  DD   0                    ;ESI
00F1  00000000         REBP  DD   0                    ;EBP
00F5  00000000               DD   0                    ;reserved
00F9  00000000         REBX  DD   0                    ;EBX
00FD  00000000         REDX  DD   0                    ;EDX
0101  00000000         RECX  DD   0                    ;ECX
0105  00000000         REAX  DD   0                    ;EAX
0109  0000             RFLAG DW   0                    ;flags
010B  0000             RES   DW   0                    ;ES
010D  0000             RDS   DW   0                    ;DS
010F  0000             RFS   DW   0                    ;FS
0111  0000             RGS   DW   0                    ;GS
0113  0000             RIP   DW   0                    ;IP
0115  0000             RCS   DW   0                    ;CS
0117  0000             RSP   DW   0                    ;SP
0119  0000             RSS   DW   0                    ;SS
0000                         .CODE
                             .STARTUP
0010  8C C0                  MOV  AX,ES
0012  8C DB                  MOV  BX,DS                ;find size of program and data
0014  2B D8                  SUB  BX,AX
0016  8B C4                  MOV  AX,SP                ;find stack size
0018  C1 E8 04               SHR  AX,4
001B  40 INC                 AX
001C  03 D8                  ADD  BX,AX                ;BX = length in paragraphs
001E  B4 4A                  MOV  AH,4AH
0020  CD 21                  INT  21H                  ;modify memory allocation
0022  E8 00D1                CALL GETDA                ;get command line information
0025  73 0A                  JNC  MAIN1               ;if parameters are good
0027  B4 09                  MOV  AH,9                 ;parameter error
0029  BA 012 R               MOV  DX,OFFSET ERR1
002C  CD 21                  INT  21H
002E  E9 00AA                JMP  MAINE                ;exit to DOS
0031             MAIN1:
0031  E8 00AB                CALL ISDPMI               ;is DPMI loaded?
0034  72 0A                  JC   MAIN2                ;if DPMI present
0036  B4 09                  MOV  AH,9
0038  BA 0026 R              MOV  DX,OFFSET ERR2
003B  CD 21                  INT  21H                  ;display DPMI not present
003D  E9 009B                JMP  MAINE                ;exit to DOS
```

```
0040                    MAIN2:
0040  B8 0000           MOV  AX,0                 ;indicate 0 memory needed
0043  83 3E 0010 R 00   CMP  MSIZE,0
0048  74 F6             JE   MAIN2                ;if DPMI needs no memory
004A  8B 1E 0010 R      MOV  BX,MSIZE             ;get amount
004E  B4 48             MOV  AH,48H
0050  CD 21             INT  21H                  ;allocate memory for DPMI
0052  73 09             JNC  MAIN3
0054  B4 09             MOV  AH,9                 ;if not enough real memory
0056  BA 003B R         MOV  DX,OFFSET ERR3
0059  CD 21             INT  21H
005B  EB 7E             JMP  MAINE                ;exit to DOS
005D                    MAIN3:
005D  8E C0             MOV  ES,AX
005F  B8 0000           MOV  AX,0                 ;16-bit application
0062  FF 1E 0000 R      CALL DS:ENTRY             ;switch to protected mode
0066  73 09             JNC  MAIN4
0068  B4 09             MOV  AH,9                 ;if switch failed
006A  BA 0056 R         MOV  DX,OFFSET ERR4
006D  CD 21             INT  21H
006F  EB 6A             JMP  MAINE                ;exit to DOS
                        ;
                        ;PROTECTED MODE
                        ;
0071                    MAIN4:
0071  B8 0000           MOV  AX,0000H             ;get local selector
0074  B9 0001           MOV  CX,1                 ;only one is needed
0077  CD 31             INT  31H
0079  72 48             JC   MAIN7                ;if error
007B  8B D8             MOV  BX,AX                ;save selector
007D  8E C0             MOV  ES,AX                ;load ES with selector
007F  B8 0007           MOV  AX,0007H             ;set base address
0082  8B 0E 000A R      MOV  CX,WORD PTR FIRST+2
0086  8B 16 0008 R      MOV  DX,WORD PTR FIRST
008A  CD 31             INT  31H
008C  72 3D             JC   MAIN8                ;if error
008E  B8 0008           MOV  AX,0008H
0091  B9 0000           MOV  CX,0
0094  BA FFFF           MOV  DX,0FFFFH            ;set limit to 64K
0097  CD 31             INT  31H
0099  72 38             JC   MAIN9                ;if error
009B  B9 0018           MOV  CX,24                ;load line count
009E  BE 0000           MOV  SI,0                 ;load offset
00A1                    MAIN5:
00A1  E8 00F4           CALL DADDR                ;display address, if needed
00A4  E8 00CE           CALL DDATA                ;display data
00A7  46                INC  SI                   ;point to next data
00A8  66| A1 0008 R     MOV  EAX,FIRST            ;test for end
00AC  66| 3B 06 000C R  CMP  EAX,LAST1
00B1  74 07             JE   MAIN6                ;if done
00B3  66| FF 06 0008 R  INC  FIRST
00B8  EB E7             JMP  MAIN5
00BA                    MAIN6:
00BA  B8 0001           MOV  AX,0001H             ;release descriptor
00BD  8C C3             MOV  BX,ES
00BF  CD 31             INT  31H
00C1  EB 18             JMP  MAINE                ;exit to DOS
00C3                    MAIN7:
00C3  BA 007B R         MOV  DX,OFFSET ERR5
00C6  E8 0096           CALL DISPS                ;display cannot allocate selector
00C9  EB 10             JMP  MAINE                ;exit to DOS
00CB                    MAIN8:
00CB  BA 0098 R         MOV  DX,OFFSET ERR6
00CE  E8 008E           CALL DISPS                ;display cannot use base address
00D1  EB E7             JMP  MAIN6                ;release descriptor
00D3                    MAIN9:
00D3  BA 00B4 R         MOV DX,OFFSET ERR7
00D6  E8 0086           CALL DISPS                ;display cannot allocate 64K limit
```

```
00D9  EB DF                 JMP  MAIN6               ;release descriptor
00DB             MAINE:
                        .EXIT
                 ;
                 ;The ISDPMI procedure tests for the presence of DPMI.
                 ;***exit parameters***
                 ;carry = 1; if DPMI is present
                 ;carry = 0; if DPMI is not present
                 ;
00DF             ISDPMI PROC NEAR

00DF  B8 1687            MOV  AX,1687H            ;get DPMI status
00E2  CD 2F              INT  2FH                 ;DOS multiplex
00E4  0B C0              OR   AX,AX
00E6  75 0D              JNZ  ISDPMI1             ;if no DPMI
00E8  89 36 0010 R       MOV  MSIZE,SI            ;save amount of memory needed
00EC  89 3E 0000 R       MOV  WORD PTR ENTRY,DI
00F0  8C 06 0002 R       MOV  WORD PTR ENTRY+2,ES
00F4  F9                 STC
00F5             ISDPMI1:
00F5  C3                 RET

00F6             ISDPMI ENDP
                 ;
                 ;The GETDA procedure retrieves the command line parameters
                 ;for memory display in hexadecimal.
                 ;FIRST = the first address from the command line
                 ;LAST1 = the last address from the command line
                 ;***return parameters***
                 ;carry = 1; if error
                 ;carry = 0; for no error
                 ;
00F6             GETDA PROC NEAR

00F6  1E                 PUSH DS
00F7  06                 PUSH ES
00F8  1F                 POP  DS
00F9  07                 POP  ES                  ;exchange ES with DS
00FA  BE 0081            MOV  SI,81H              ;address command line
00FD             GETDA1:
00FD  AC                 LODSB                    ;skip spaces
00FE  3C 20              CMP  AL,' '
0100  74 FB              JE   GETDA1              ;if space
0102  3C 0D              CMP  AL,13
0104  74 1E              JE   GETDA3              ;if enter = error
0106  4E                 DEC  SI                  ;adjust SI
0107             GETDA2:
0107  E8 0020            CALL GETNU               ;get first number
010A  3C 2C              CMP AL,','
010C  75 16              JNE  GETDA3              ;if no comma = error
010E  66| 26: 89 16 0008 R MOV  ES:FIRST,EDX
0114  E8 0013            CALL GETNU               ;get second number
0117  3C 0D              CMP  AL,13
0119  75 09              JNE  GETDA3              ;if error
011B  66| 26: 89 16 000C R MOV  ES:LAST1,EDX
0121  F8                 CLC                      ;indicate no error
0122  EB 01              JMP GETDA4               ;return no error
0124             GETDA3:
0124  F9                 STC                      ;indicate error
0125             GETDA4:
0125  1E                 PUSH DS                  ;exchange ES with DS
0126  06                 PUSH ES
0127  1F                 POP  DS
0128  07                 POP  ES
0129  C3                 RET
012A             GETDA ENDP
                 ;
                 ;The GETNU procedure extracts a number from the command line
```

```
                         ;and returns with it in EDX and last command line character in
                         ;AL as a delimiter.
                         ;
012A                     GETNU PROC NEAR

012A  66| BA 00000000       MOV  EDX,0               ;clear result
0130                     GETNU1:
0130  AC                    LODSB                    ;get digit from command line
                            .IF  AL >= 'a' && AL <= 'z'
0139  2C 20                    SUB AL,20H            ;make uppercase
                            .ENDIF
013B  2C 30                 SUB  AL,'0'              ;convert from ASCII
013D  72 12                 JB   GETNU2              ;if not a number
                            .IF  AL > 9              ;convert A-F from ASCII
0143  2C 07                    SUB AL,7
                            .ENDIF
0145  3C 0F                 CMP  AL,0FH
0147  77 08                 JA   GETNU2              ;if not 0-F
0149  66| C1 E2 04          SHL  EDX,4
014D  02 D0                 ADD  DL,AL               ;add digit to EDX
014F  EB DF                 JMP  GETNU1              ;get next digit
0151                     GETNU2:
0151  8A 44 FF              MOV  AL,[SI-1]           ;get delimiter
0154  C3                    RET

0155                     GETNU ENDP
                         ;
                         ;The DISPC procedure displays the ASCII character found
                         ;in register AL.
                         ;***uses***
                         ;INT21H
                         ;
0155                     DISPC PROC NEAR

0155  52                   PUSH DX
0156  8A D0                MOV  DL,AL
0158  B4 06                MOV  AH,6
015A  E8 0084              CALL INT21H               ;do real INT 21H
015D  5A                   POP  DX
015E  C3                   RET

015F                     DISPC ENDP
                         ;
                         ;The DISPS procedure displays a character string from
                         ;protected mode addressed by DS:EDX.
                         ;***uses***
                         ;DISPC
                         ;
015F                     DISPS PROC NEAR

015F  66| 81 E2 0000FFFF    AND  EDX,0FFFFH
0166  67& 8A 02            MOV  AL,[EDX]             ;get character
0169  3C 24                CMP  AL,'$'               ;test for end
016B  74 07                JE   DISP1                ;if end
016D  66| 42               INC  EDX                  ;address next character
016F  E8 FFE3              CALL DISPC                ;display character
0172  EB EB                JMP  DISPS                ;repeat until $
0174                     DISP1:
0174  C3                    RET

0175                     DISPS ENDP
                         ;
                         ;The DDATA procedure displays a byte of data at the location
                         ;addressed by ES:SI. The byte is followed by one space.
                         ;***uses***
                         ;DIP and DISPC
                         ;
0175                     DDATA PROC NEAR
```

```
0175  26: 8A 04              MOV  AL,ES:[SI]        ;get byte
0178  C0 E8 04               SHR  AL,4
017B  E8 000C                CALL DIP               ;display first digit
017E  26: 8A 04              MOV  AL,ES:[SI]        ;get byte
0181  E8 0006                CALL DIP               ;display second digit
0184  B0 20                  MOV  AL,' '            ;display space
0186  E8 FFCC                CALL DISPC
0189  C3                     RET

018A              DDATA ENDP
                  ;
                  ;The DIP procedure displays the right nibble found in AL as a
                  ;hexadecimal digit.
                  ;***uses***
                  ;DISPC
                  ;
018A              DIP   PROC NEAR
018A  24 0F               AND  AL,0FH          ;get right nibble
018C  04 30               ADD  AL,30H          ;convert to ASCII
                          .IF  AL > 39H        ;if A-F
0192  04 07                 ADD    AL,7
                          .ENDIF
0194  E8 FFBE             CALL DISPC           ;display digit
0197  C3                  RET

0198              DIP ENDP
                  ;
                  ;The DADDR procedure displays the hexadecimal address found
                  ;in DS:FIRST if it is a paragraph boundary.
                  ;***uses***
                  ;DIP, DISPS, DISPC, and INT21H
                  ;
0198              DADDR PROC NEAR

0198  66| A1 0008 R        MOV  EAX,FIRST       ;get address
019C  A8 0F               TEST AL,0FH          ;test for XXXXXXX0
019E  75 40               JNZ  DADDR4          ;if not, don't display address
01A0  BA 00D5 R           MOV  DX,OFFSET CRLF
01A3  E8 FFB9             CALL DISPS           ;display CR and LF
01A6  49                  DEC  CX              ;decrement line count
01A7  75 18               JNZ  DADDR2          ;if not end of page
01A9  BA 00D8 R           MOV  DX,OFFSET MES1  ;if end of page
01AC  E8 FFB0             CALL DISPS           ;display press any key
01AF              DADDR1:
01AF  B4 06               MOV  AH,6            ;get any key, no echo
01B1  B2 FF               MOV  DL,0FFH
01B3  E8 002B             CALL INT21H          ;do real INT 21H
01B6  74 F7               JZ   DADDR1          ;if nothing typed
01B8  BA 00D5 R           MOV  DX,OFFSET CRLF
01BB  E8 FFA1             CALL DISPS           ;display CRLF
01BE  B9 0018             MOV  CX,24            ;reset line count
01C1              DADDR2:
01C1  51                  PUSH CX              ;save line count
01C2  B9 0008             MOV  CX,8            ;load digit count
01C5  66| 8B 16 0008 R    MOV  EDX,FIRST       ;get address
01CA              DADDR3:
01CA  66| C1 C2 04        ROL  EDX,4
01CE  8A C2               MOV  AL,DL
01D0  E8 FFB7             CALL DIP             ;display digit
01D3  E2 F5               LOOP DADDR3          ;repeat 8 times
01D5  59                  POP  CX              ;retrieve line count
01D6  B0 3A               MOV  AL,':'
01D8  E8 FF7A             CALL DISPC           ;display colon
01DB  B0 20               MOV  AL,' '
01DD  E8 FF75             CALL DISPC           ;display space
01E0              DADDR4:
01E0  C3                  RET
```

```
01E1                    DADDR ENDP
                        ;
                        ;The INT21H procedure gains access to the real mode DOS
                        ;INT 21H instruction with the parameters intact.
                        ;
01E1                    INT21H PROC NEAR

01E1  66| A3 0105 R        MOV  REAX,EAX           ;save registers
01E5  66| 89 1E 00F9 R     MOV  REBX,EBX
01EA  66| 89 0E 0101 R     MOV  RECX,ECX
01EF  66| 89 16 00FD R     MOV  REDX,EDX
01F4  66| 89 36 00ED R     MOV  RESI,ESI
01F9  66| 89 3E 00E9 R     MOV  REDI,EDI
01FE  66| 89 2E 00F1 R     MOV  REBP,EBP
0203  9C                   PUSHF
0204  58                   POP  AX
0205  A3 0109 R            MOV  RFLAG,AX
0208  06                   PUSH ES                 ;do DOS interrupt
0209  B8 0300              MOV  AX,0300H
020C  BB 0021              MOV  BX,21H
020F  B9 0000              MOV  CX,0
0212  1E                   PUSH DS
0213  07                   POP  ES
0214  BF 00E9 R            MOV  DI,OFFSET ARRAY
0217  CD 31                INT  31H
0219  07                   POP  ES
021A  A1 0109 R            MOV  AX,RFLAG            ;restore registers
021D  50                   PUSH AX
021E  9D                   POPF
021F  66| 8B 3E 00E9 R     MOV  EDI,REDI
0224  66| 8B 36 00ED R     MOV  ESI,RESI
0229  66| 8B 2E 00F1 R     MOV  EBP,REBP
022E  66| A1 0105 R        MOV  EAX,REAX
0232  66| 8B 1E 00F9 R     MOV  EBX,REBX
0237  66| 8B 0E 0101 R     MOV  ECX,RECX
023C  66| 8B 16 00FD R     MOV  EDX,REDX
0241  C3                   RET

0242                    INT21H ENDP
                        END
```
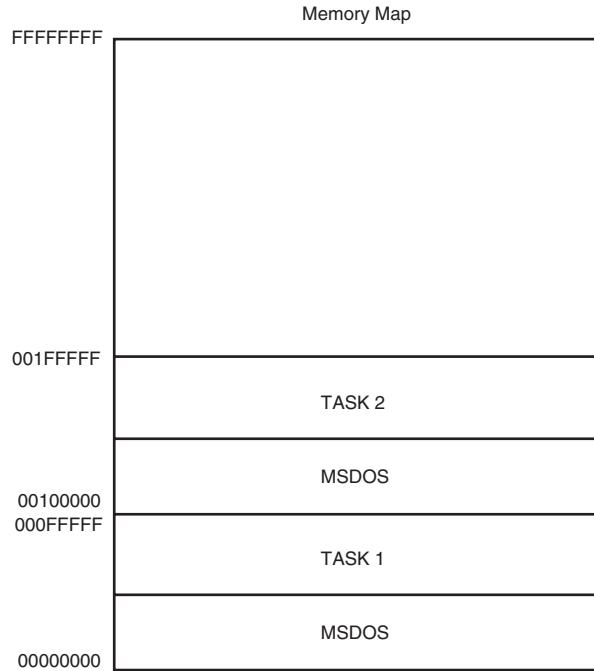
You might notice that the DOS INT 21H function call must be treated differently when operating in the protected mode. The procedure that calls a DOS INT 21H is at the end of Example 17–3. Because this is extremely long and time consuming, we have tended to move away from using the DOS interrupts from a Windows application. The best way to develop software for Windows is through the use of C/C++ with the inclusion of assembly language procedures for arduous tasks.

17–5   VIRTUAL 8086 MODE

One special mode of operation not discussed thus far is the virtual 8086 mode. This special mode is designed so that multiple 8086 real-mode software applications can execute at one time. The PC operates in this mode for DOS applications using the DOS emulator cmd.exe (the command prompt). Figure 17–25 illustrates two 8086 applications mapped into the 80386 using the virtual mode. The operating system allows multiple applications to execute, usually done through a technique called time-slicing. The operating system allocates a set amount of time to each task. For example, if three tasks are executing, the operating system can allocate 1 ms to each task. This means that after each millisecond, a task switch occurs to the next task. In this manner, all tasks receive a portion of the microprocessor's execution time, resulting in a system that appears to execute more than one task at a time. The task times can be adjusted to give any task any percentage of the microprocessor's execution time.

FIGURE 17–25 Two tasks resident in an 80386 operated in the virtual 8086 mode.

A system that can use this technique is a print spooler. The print spooler can function in one DOS partition and be accessed 10% of the time. This allows the system to print using the print spooler, but it doesn't detract from the system because it uses only 10% of the system time.

The main difference between 80386 protected mode operation and the virtual 8086 mode is the way the segment registers are interpreted by the microprocessor. In the virtual 8086 mode, the segment registers are used as they are in the real mode: as a segment address and an offset address capable of accessing a 1M-byte memory space from locations 00000H–FFFFFH. Access to many virtual 8086 mode systems is made possible by the paging unit that is explained in the next section. Through paging, the program still accesses memory below the 1M-byte boundary, yet the microprocessor can access a physical memory space at any location in the 4G-byte range of the memory system.

Virtual 8086 mode is entered by changing the VM bit in the EFLAG register to a logic 1. This mode is entered via an IRET instruction if the privilege level is 00. This bit cannot be set in any other manner. An attempt to access a memory address above the 1M-byte boundary will cause a type 13 interrupt to occur.

The virtual 8086 mode can be used to share one microprocessor with many users by partitioning the memory so that each user has its own DOS partition. User 1 can be allocated memory locations 00100000H–01FFFFFH, user 2 can be allocated locations 0020000H–01FFFFFFH, and so forth. The system software located at memory locations 00000000H–000FFFFFH can then share the microprocessor between users by switching from one to another to execute software. In this manner, one microprocessor is shared by many users.

17–6   THE MEMORY PAGING MECHANISM

The paging mechanism allows any linear (logical) address, as it is generated by a program, to be placed into any physical memory page, as generated by the paging mechanism. A linear memory page is a page that is addressed with a selector and an offset in either the real or protected mode.

A physical memory page is a page that exists at some actual physical memory location. For example, linear memory location 20000H could be mapped into physical memory location 30000H, or any other location, with the paging unit. This means that an instruction that accesses location 20000H actually accesses location 30000H.

Each 80386 memory page is 4K bytes long. Paging allows the system software to be placed at any physical address with the paging mechanism. Three components are used in page address translation: the page directory, the page table, and the actual physical memory page. Note that EEM386.EXE, the extended memory manager, uses the paging mechanism to simulate expanded memory in extended memory and to generate upper memory blocks between system ROMs.

The Page Directory

The page directory contains the location of up to 1024 page translation tables. Each page translation table translates a logic address into a physical address. The page directory is stored in the memory and accessed by the page descriptor address register (CR_3) (see Figure 17–14). Control register CR_3 holds the base address of the page directory, which starts at any 4K-byte boundary in the memory system. The MOV CR_3,reg instruction is used to initialize CR_3 for paging. In a virtual 8086 mode system, each 8086 DOS partition would have its own page directory.

The page directory contains up to 1024 entries, which are each four bytes long. The page directory itself occupies one 4K-byte memory page. Each entry in the page directory (see Figure 17–26) translates the leftmost 10 bits of the memory address. This 10-bit portion of the linear address is used to locate different page tables for different page table entries. The page table address (A_{32}–A_{12}), stored in a page directory entry, accesses a 4K-byte-long page translation table. To completely translate any linear address into any physical address requires 1024 page tables that are each 4K bytes long, plus the page table directory, which is also 4K bytes long. This translation scheme requires up to 4M plus 4K bytes of memory for a full address translation. Only the largest operating systems support this size address translation. Many commonly found operating systems translate only the first 16M bytes of the memory system if paging is enabled. This includes programs such as Windows. This translation requires four entries in the page directory (16 bytes) and four complete page tables (16K bytes).

The page table directory entry control bits, as illustrated in Figure 17–26, each perform the following functions:

D	**Dirty** is undefined for page table directory entries by the 80386 microprocessor and is provided for use by the operating system.
A	**Accessed** is set to a logic 1 whenever the microprocessor accesses the page directory entry.
R/W and U/S	**Read/write** and **user/supervisor** are both used in the protection scheme, as listed in Table 17–2. Both bits combine to develop paging priority level protection for level 3, the lowest user level.
P	**Present**, if a logic 1, indicates that the entry can be used in address translation. If P = 0, the entry cannot be used for translation. A not present entry can be used for other purposes, such as indicating that the page is currently stored on the disk. If P = 0, the remaining bits of the entry can be used to indicate the location of the page on the disk memory system.

FIGURE 17–26 The page table directory entry.
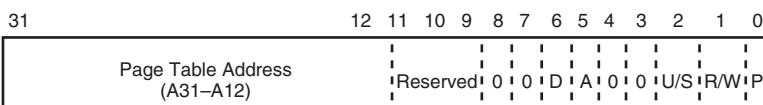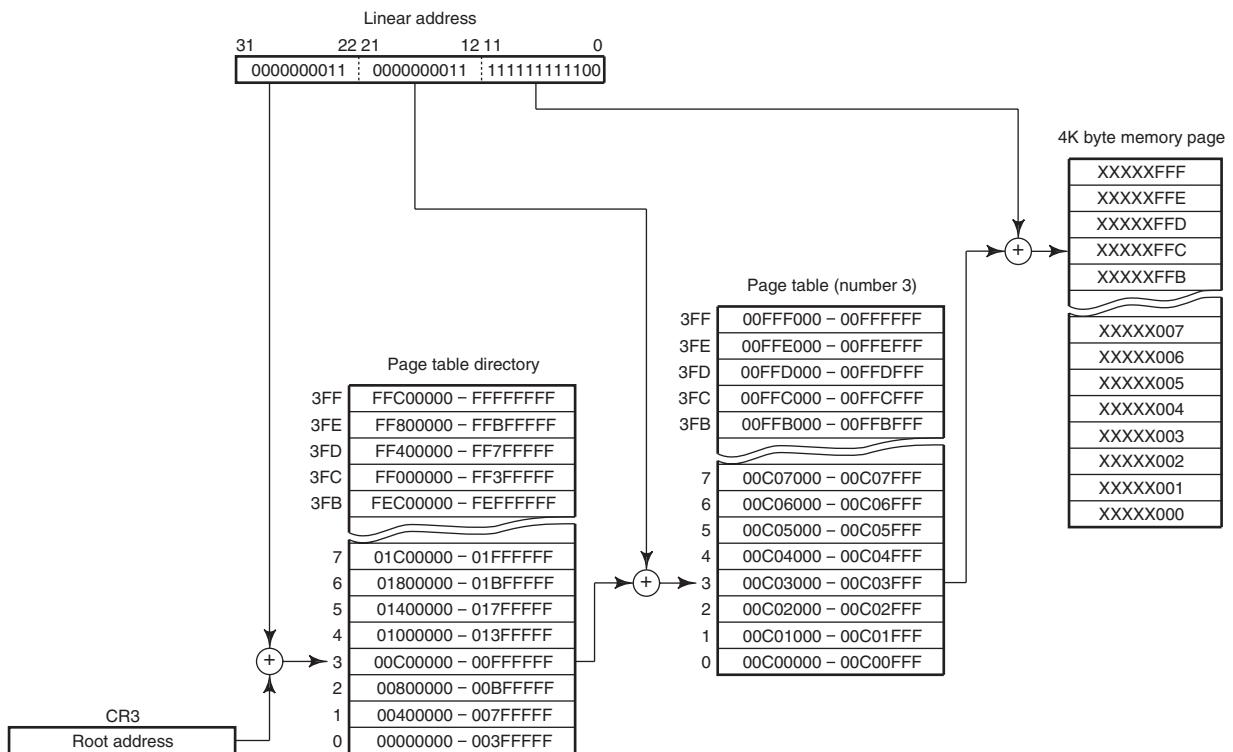
TABLE 17–2 Protection for level 3 using U/S and R/W.

U/S	R/W	Access Level 3
0	0	None
0	1	None
1	0	Read-only
1	1	Write-only

The Page Table

The page table contains 1024 physical page addresses, accessed to translate a linear address into a physical address. Each page table translates a 4M section of the linear memory into 4M of physical memory. The format for the page table entry is the same as for the page directory entry (refer to Figure 17–26). The main difference is that the page directory entry contains the physical address of a page table, while the page table entry contains the physical address of a 4K-byte physical page of memory. The other difference is the D (dirty bit), which has no function in the page directory entry, but indicates that a page has been written to in a page table entry.

Figure 17–27 illustrates the paging mechanism in the 80386 microprocessor. Here, the linear address 00C03FFCH, as generated by a program, is converted to physical address XXXXXF-

Note: 1. The address ranges illustrated in the page directory and page table represent the linear address ranges selected and not the contents of these tables.

2. The addresses (XXXXX) listed in the memory page are selected by the page table entry.

FIGURE 17–27 The translation of linear address 00C03FFCH to physical memory address XXXXXFFCH. The value of XXXXX is determined by the page table entry (not shown here).

FCH, as translated by the paging mechanism. (Note: XXXXX is any 4K-byte physical page address.) The paging mechanism functions in the following manner:

1. The 4K-byte long page directory is stored as the physical address located by CR_3. This address is often called the root address. One page directory exists in a system at a time. In the 8086 virtual mode, each task has its own page directory, allowing different areas of physical memory to be assigned to different 8086 virtual tasks.
2. The upper 10 bits of the linear address (bits 31–22), as determined by the descriptors described earlier in this chapter or by a real address, are applied to the paging mechanism to select an entry in the page directory. This maps the page directory entry to the leftmost 10 bits of the linear address.
3. The page table is addressed by the entry stored in the page directory. This allows up to 4K page tables in a fully populated and translated system.
4. An entry in the page table is addressed by the next 10 bits of the linear address (bits 21–12).
5. The page table entry contains the actual physical address of the 4K-byte memory page.
6. The rightmost 12 bits of the linear address (bits 11–0) select a location in the memory page.

The paging mechanism allows the physical memory to be assigned to any linear address through the paging mechanism. For example, suppose that linear address 20000000H is selected by a program, but this memory location does not exist in the physical memory system. The 4K-byte linear page is referenced as locations 20000000H–20000FFFH by the program. Because this section of physical memory does not exist, the operating system might assign an existing physical memory page such as 12000000H–12000FFFH to this linear address range.

In the address translation process, the leftmost 10 bits of the linear address select page directory entry 200H located at offset address 800H in the page directory. This page directory entry contains the address of the page table for linear addresses 20000000H–203FFFFFH. Linear address bits (21–12) select an entry in this page table that corresponds to a 4K-byte memory page. For linear addresses 20000000H–20000FFFH, the first entry (entry 0) in the page table is selected. This first entry contains the physical address of the actual memory page, or 12000000H–12000FFFH in this example.
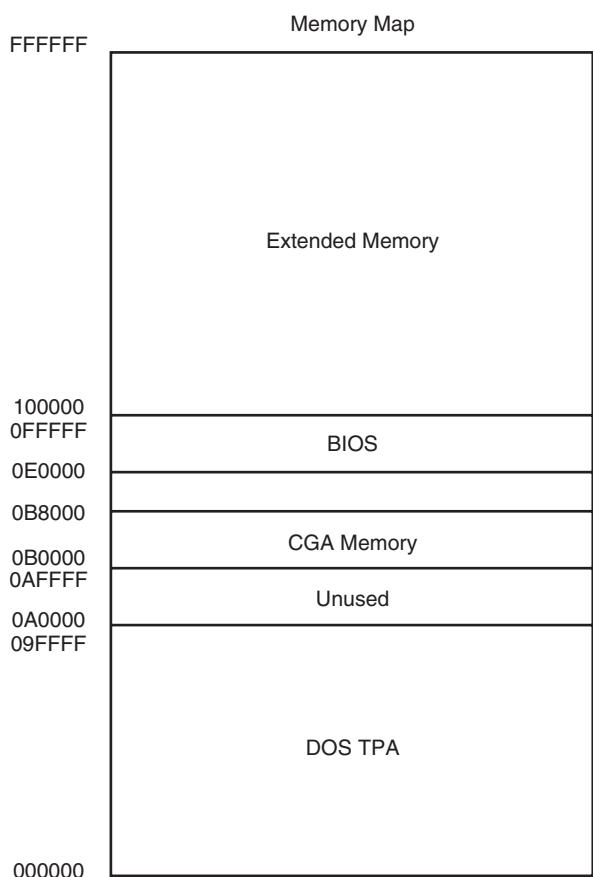
Take, for example, a typical DOS-based computer system. The memory map for the system appears in Figure 17–28. Note from the map that there are unused areas of memory, which can be paged to a different location, giving a DOS real mode application program more memory. The normal DOS memory system begins at location 00000H and extends to location 9FFFFH, which is 640K bytes of memory. Above location 9FFFFH, we find sections devoted to video cards, disk cards, and the system BIOS ROM. In this example, an area of memory just above 9FFFFH is unused (A0000–AFFFFH). This section of the memory could be used by DOS, so that the total application-memory area is 704K instead of 640K. Be careful when using A0000H–AFFFFH for additional RAM because the video card uses this area for bit-mapped graphics in mode 12H and 13H.

This section of memory can be used by mapping it into extended memory at locations 1002000H–11FFFFFH. Software to accomplish this translation and initialize the page table directory, and page tables required to set up memory, are illustrated in Example 17–4. Note that this procedure initializes the page table directory, a page table, and loads CR3. It does not switch to protected mode and it does enable paging. Note that paging functions in real mode memory operation.

EXAMPLE 17–4

```
.MODEL SMALL
.386P
.DATA

;page directory
```

FIGURE 17–28 Memory
map for an AT-style clone.


```
PDIR   DD     4

;page table 0

TAB0   DD     1024 dup(?)

.CODE
.STARTUP
       MOV  EAX,0
       MOV  AX,CS
       SHL  EAX,4
       ADD  EAX,OFFSET TAB0
       AND  EAX,0FFFFF000H
       ADD  EAX,7
       MOV  PDIR,EAX             ;address page directory
       MOV  ECX,256
       MOV  EDI,OFFSET TAB0
       MOV  AX,DS
       MOV  ES,AX
       MOV  EAX,7
       .REPEAT                   ;remap 00000H-9FFFFH
             STOSD               ;to 00000H-9FFFFH
             ADD  EAX,4096
       .UNTILCXZ
       MOV  EAX,102007H
       MOV  ECX,16
       .REPEAT                   ;remap A0000H-AFFFFH
             STOSD               ;to 102000H-11FFFFH
             ADD  EAX,4096
```

```
        .UNTILCXZ
        MOV  EAX,0
        MOV  AX,DS
        SHL  EAX,4
        ADD  EAX,OFFSET PDIR      ;load CR3
        MOV  CR3,EAX

;additional software to remap other areas of memory

        END
```

17–7   INTRODUCTION TO THE 80486 MICROPROCESSOR

The 80486 microprocessor is a highly integrated device, containing well over 1.2 million transistors. Located within this device circuit are a memory-management unit (MMU), a complete numeric coprocessor that is compatible with the 80387, a high-speed level 1 cache memory that contains 8K bytes of space, and a full 32-bit microprocessor that is upward-compatible with the 80386 microprocessor. The 80486 is currently available as a 25 MHz, 33 MHz, 50 MHz, 66 MHz, or 100 MHz device. Note that the 66 MHz version is double-clocked and the 100 MHz version is triple-clocked. In 1990, Intel demonstrated a 100 MHz version (not double-clocked) of the 80486 for *Computer Design* magazine, but it has yet to be released. Advanced Micro Devices (AMD) has produced a 40 MHz version that is also available in an 80 MHz (double-clocked) and a 120 MHz (triple-clocked) form. The 80486 is available as an 80486DX or an 80486SX. The only difference between these devices is that the 80486SX does not contain the numeric coprocessor, which reduces its price. The 80487SX numeric coprocessor is available as a separate component for the 80486SX microprocessor.

This section details the differences between the 80486 and 80386 microprocessors. These differences are few, as shall be seen. The most notable differences apply to the cache memory system and parity generator.
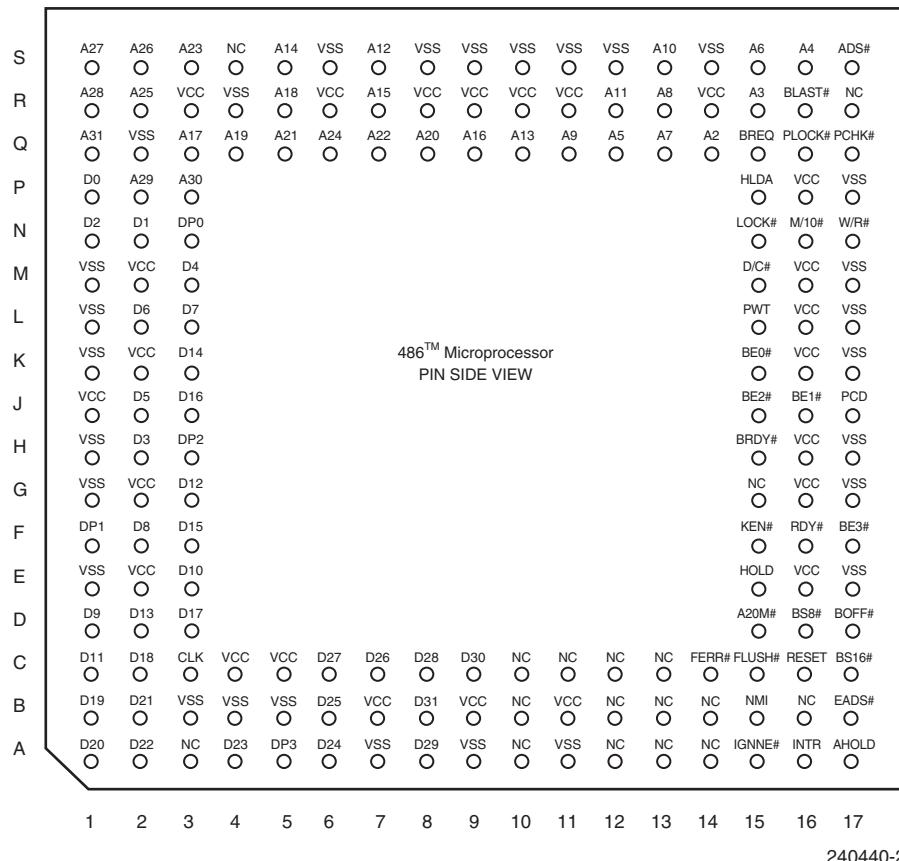
Pin-Out of the 80486DX and 80486SX Microprocessors

Figure 17–29 illustrates the pin-out of the 80486DX microprocessor, a 168-pin PGA. The 80486SX, also packaged in a 168-pin PGA, is not illustrated because only a few differences exist. Note that pin B_{15} is NMI on the 80486DX and pin A_{15} is NMI on the 80486SX. The only other differences are that pin A_{15} is $\overline{\text{IGNNE}}$ (ignore numeric error) on the 80486DX (not present on the 80486SX), pin C_{14} is $\overline{\text{FERR}}$ (floating-point error) on the 80486DX, and pins B_{15} and C_{14} on the 80486SX are not connected.

When connecting the 80486 microprocessor, all V_{CC} and V_{SS} pins must be connected to the power supply for proper operation. The power supply must be capable of supplying 5.0 V $\pm 10\%$, with up to 1.2 A of surge current for the 33 MHz version. The average supply current is 650 mA for the 33 MHz version. Intel has also produced a 3.3 V version that requires an average of 500 mA at a triple-clock speed of 100 MHz. Logic 0 outputs allow up to 4.0 mA of current, and logic 1 outputs allow up to 1.0 mA. If larger currents are required, as they often are, then the 80486 must be buffered. Figure 17–30 shows a buffered 80486DX system. In the circuit shown, only the address, data, and parity signals are buffered.

Pin Definitions

A_{31}–A_2 **Address outputs** A_{31}–A_2 provide the memory and I/O with the address during normal operation; during a cache line invalidation, A_{31}–A_4 are used to drive the microprocessor.

	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17
S	A27	A26	A23	NC	A14	VSS	A12	VSS	VSS	VSS	VSS	VSS	A10	VSS	A6	A4	ADS#
R	A28	A25	VCC	VSS	A18	VCC	A15	VCC	VCC	VCC	VCC	A11	A8	VCC	A3	BLAST#	NC
Q	A31	VSS	A17	A19	A21	A24	A22	A20	A16	A13	A9	A5	A7	A2	BREQ	PLOCK#	PCHK#
P	D0	A29	A30												HLDA	VCC	VSS
N	D2	D1	DP0												LOCK#	M/IO#	W/R#
M	VSS	VCC	D4												D/C#	VCC	VSS
L	VSS	D6	D7												PWT	VCC	VSS
K	VSS	VCC	D14												BE0#	VCC	VSS
J	VCC	D5	D16												BE2#	BE1#	PCD
H	VSS	D3	DP2												BRDY#	VCC	VSS
G	VSS	VCC	D12												NC	VCC	VSS
F	DP1	D8	D15												KEN#	RDY#	BE3#
E	VSS	VCC	D10												HOLD	VCC	VSS
D	D9	D13	D17												A20M#	BS8#	BOFF#
C	D11	D18	CLK	VCC	VCC	D27	D26	D28	D30	NC	NC	NC	NC	FERR#	FLUSH#	RESET	BS16#
B	D19	D21	VSS	VSS	VSS	D25	VCC	D31	VCC	NC	VCC	NC	NC	NC	NMI	NC	EADS#
A	D20	D22	NC	D23	DP3	D24	VSS	D29	VSS	NC	VSS	NC	NC	NC	IGNNE#	INTR	AHOLD

486™ Microprocessor
PIN SIDE VIEW

240440-2

FIGURE 17–29 The pin-out of the 80486. (Courtesy of Intel Corporation.)

$\overline{\text{A20M}}$ **Address bit 20 mask** causes the 80486 to wrap its address around from location 000FFFFFH to 00000000H, as does the 8086 microprocessor. This provides a memory system that functions like the 1M-byte real memory system in the 8086 microprocessor.

$\overline{\text{ADS}}$ **Address data strobe** becomes a logic 0 to indicate that the address bus contains a valid memory address.

AHOLD **Address hold** input causes the microprocessor to place its address bus connections at their high-impedance state, with the remainder of the buses staying active. It is often used by another bus master to gain access for a cache invalidation cycle.

$\overline{\text{BE3}}$–$\overline{\text{BE0}}$ **Byte enable** outputs select a bank of the memory system when information is transferred between the microprocessors and its memory and I/O space. The $\overline{\text{BE3}}$ signal enables D_{31}–D_{24}, $\overline{\text{BE2}}$ enables D_{23}–D_{16}, $\overline{\text{BE1}}$ enables D_{15}–D_8, and $\overline{\text{BE0}}$ enables D_7–D_0.

$\overline{\text{BLAST}}$ The **burst last** output shows that the burst bus cycle is complete on the next activation of the $\overline{\text{BRDY}}$ signal.

$\overline{\text{BOFF}}$ The **back-off** input causes the microprocessor to place its buses at their high-impedance state during the next clock cycle. The microprocessor remains in the bus hold state until the $\overline{\text{BOFF}}$ pin is placed at a logic 1 level.
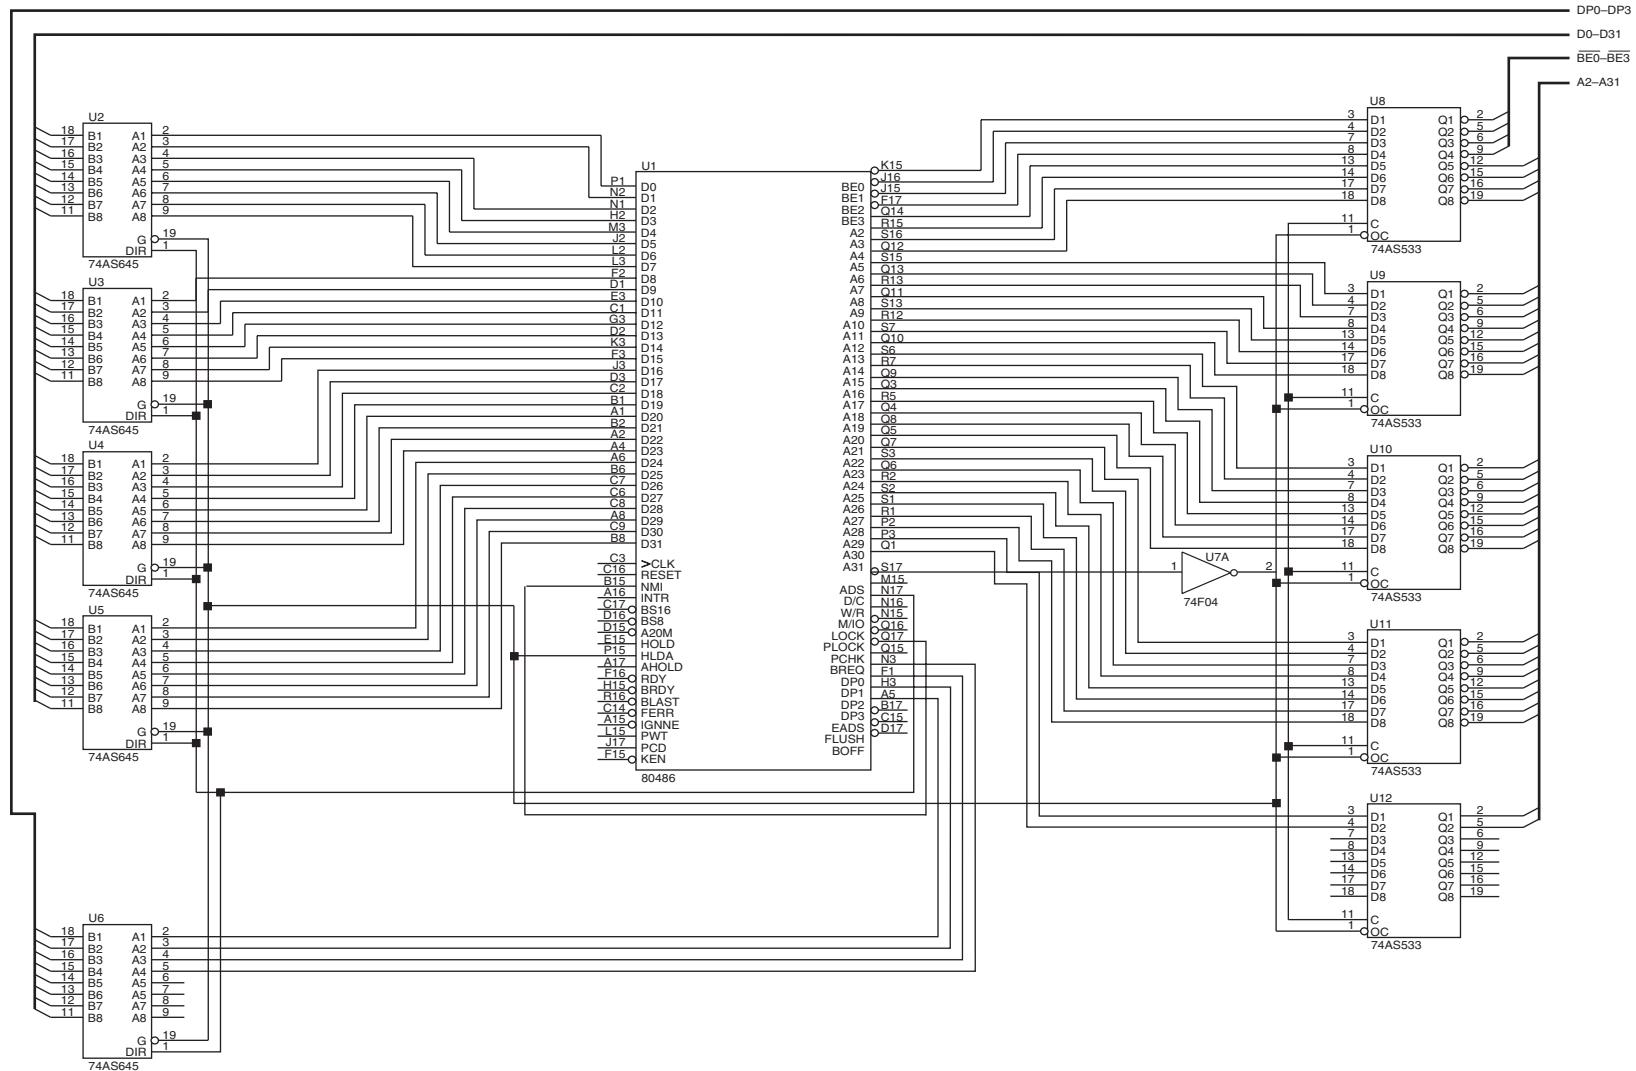
FIGURE 17–30 An 80486 microprocessor showing the buffered address, data, and parity buses.

$\overline{\text{BRDY}}$ — The **burst ready** input is used to signal the microprocessor that a burst cycle is complete.

BREQ — The **bus request** output indicates that the 80486 has generated an internal bus request.

$\overline{\text{BS8}}$ — The **bus size 8** input causes the 80486 to structure itself with an 8-bit data bus to access byte-wide memory and I/O components.

$\overline{\text{BS16}}$ — The **bus size 16** input causes the 80486 to structure itself with a 16-bit data bus to access word-wide memory and I/O components.

CLK — The **clock** input provides the 80486 with its basic timing signal. The clock input is a TTL-compatible input that is 25 MHz to operate the 80486 at 25 MHz.

D_{31}–D_0 — The **data bus** transfers data between the microprocessor and its memory and I/O system. Data bus connections D_7–D_0 are also used to accept the interrupt vector type number during an interrupt acknowledge cycle.

$D/\overline{C}$ — The **data/control** output indicates whether the current operation is a data transfer or control cycle. See Table 17–3 for the function of $D/\overline{C}$, $M/\overline{IO}$, and $W/\overline{R}$.

DP_3–DP_0 — **Data parity** I/O provides even parity for a write operation and check parity for a read operation. If a parity error is detected during a read, the $\overline{\text{PCHK}}$ output becomes a logic 0 to indicate a parity error. If parity is not used in a system, these lines must be pulled high to +5.0 V or to 3.3 V in a system that uses a 3.3 V supply.

$\overline{\text{EADS}}$ — The **external address strobe** input is used with AHOLD to signal that an external address is used to perform a cache-invalidation cycle.

$\overline{\text{FERR}}$ — The floating-point error output indicates that the floating-point coprocessor has detected an error condition. It is used to maintain compatibility with DOS software.

$\overline{\text{FLUSH}}$ — The **cache flush** input forces the microprocessor to erase the contents of its 8K-byte internal cache.

HLDA — The **hold acknowledge** output indicates that the HOLD input is active and that the microprocessor has placed its buses at their high-impedance state.

HOLD — The **hold** input requests a DMA action. It causes the address, data, and control buses to be placed at their high-impedance state and also, once recognized, causes HLDA to become a logic 0.

$\overline{\text{IGNNE}}$ — The **ignore numeric error** input causes the coprocessor to ignore floating-point errors and to continue processing data. This signal does not affect the state of the $\overline{\text{FERR}}$ pin.

TABLE 17–3 Bus cycle identification.

M/IO	D/C	W/R	Bus Cycle Type
0	0	0	Interrupt acknowledge
0	0	1	Halt/special
0	1	0	I/O read
0	1	1	I/O write
1	0	0	Opcode fetch
1	0	1	Reserved
1	1	0	Memory read
1	1	1	Memory write

INTR	The **interrupt request** input requests a maskable interrupt, as it does in all other family members.
$\overline{\text{KEN}}$	The **cache enable** input causes the current bus to be stored in the internal cache.
$\overline{\text{LOCK}}$	The **lock** output becomes a logic 0 for any instruction that is prefixed with the lock prefix.
M/$\overline{\text{IO}}$	**Memory/IO** defines whether the address bus contains a memory address or an I/O port number. It is also combined with the W/$\overline{\text{R}}$ signal to generate memory and I/O read and write control signals.
NMI	The **non-maskable interrupt** input requests a type 2 interrupt.
PCD	The **page cache disable** output reflects the state of the PCD attribute bit in the page table entry or the page directory entry.
$\overline{\text{PCHK}}$	The **parity check** output indicates that a parity error was detected during a read operation on the DP_3–DP_0 pins.
$\overline{\text{PLOCK}}$	The **pseudo-lock** output indicates that the current operation requires more than one bus cycle to perform. This signal becomes a logic 0 for arithmetic coprocessor operations that access 64- or 80-bit memory data.
PWT	The **page write through** output indicates the state of the PWT attribute bit in the page table entry or the page directory entry.
$\overline{\text{RDY}}$	The **ready** input indicates that a non-burst bus cycle is complete. The $\overline{\text{RDY}}$ signal must be returned, or the microprocessor places wait states into its timing until $\overline{\text{RDY}}$ is asserted.
RESET	The **reset** input initializes the 80486, as it does in other family members. Table 17–4 shows the effect of the RESET input on the 80486 microprocessor.
W/$\overline{\text{R}}$	**Write/read** signals that the current bus cycle is either a read or a write.
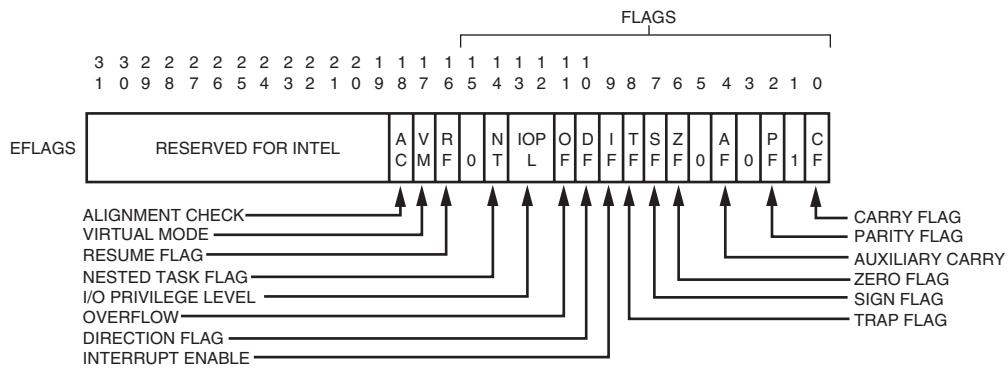
Basic 80486 Architecture

The architecture of the 80486DX is almost identical to the 80386. Added to the 80386 architecture inside the 80486DX is a math coprocessor and an 8K-byte level 1 cache memory. The 80486SX is almost identical to an 80386 with an 8K-byte cache, but no numeric coprocessor.

TABLE 17–4 State of the microprocessor after a RESET.

Register	Initial Value with Self Test	Initial Value without Self Test
EAX	00000000H	?
EDX	00000400H + ID*	00000400H + ID*
EFLAGS	00000002H	00000002H
EIP	0000FFF0H	0000FFF0H
ES	0000H	0000H
CS	F000H	F000H
DS	0000H	0000H
SS	0000H	0000H
GS	0000H	0000H
FS	0000H	0000H
IDTR	Base = 0, limit = 3FFH	Base = 0, limit = 3FFH
CR_0	60000010H	60000010H
DR_7	00000000H	00000000H

*Note: Revision ID number is supplied by Intel for revisions to the microprocessor.

FIGURE 17–31 The EFLAG register of the 80486. (Courtesy of Intel Corporation.)

The extended flag register (EFLAG) is illustrated in Figure 17–31. As with other family members, the rightmost flag bits perform the same functions for compatibility. The only new flag bit is the AC (alignment check), used to indicate that the microprocessor has accessed a word at an odd address or a doubleword stored at a non-doubleword boundary. Efficient software and execution require that data be stored at word or doubleword boundaries.

80486 Memory System
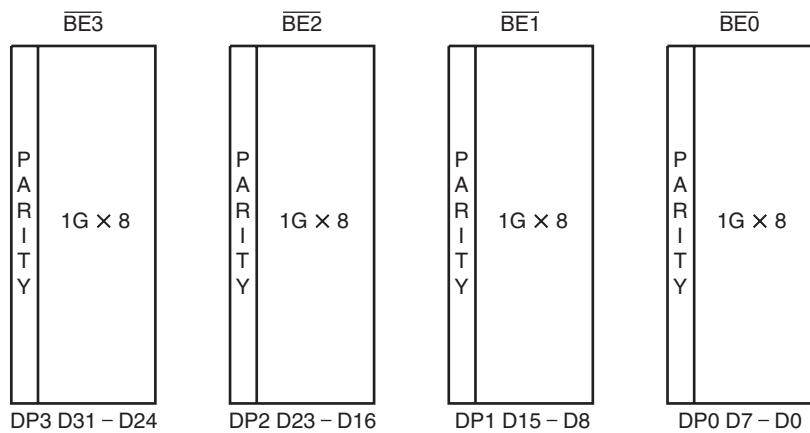
The memory system for the 80486 is identical to the 80386 microprocessor. The 80486 contains 4G bytes of memory, beginning at location 00000000H and ending at location FFFFFFFFH. The major change to the memory system is internal to the 80486 in the form of an 8K-byte cache memory, which speeds the execution of instructions and the acquisition of data. Another addition is the parity checker/generator built into the 80486 microprocessor.

Parity Checker/Generator. Parity is often used to determine if data are correctly read from a memory location. To facilitate this, Intel has incorporated an internal parity generator/detector. Parity is generated by the 80486 during each write cycle. Parity is generated as even parity, and a parity bit is provided for each byte of memory. The parity check bits appear on pins DP_0–DP_3, which are also parity inputs as well as outputs. These are typically stored in memory during each write cycle and read from memory during each read cycle.

On a read, the microprocessor checks parity and generates a parity check error, if it occurs, on the $\overline{PCHK}$ pin. A parity error causes no change in processing unless the user applies the $\overline{PCHK}$ signal to an interrupt input. Interrupts are often used to signal a parity error in DOS-based computer systems. Figure 17–32 shows the organization of the 80486 memory system that includes parity storage. Note that this is the same as for the 80386, except for the parity bit storage. If parity is not used, Intel recommends that the DP_0–DP_3 pins be pulled up to +5.0 V.

Cache Memory. The cache memory system caches (stores) data used by a program and also the instructions of the program. The cache is organized as a four-way set associative cache, with each location (line) containing 16 bytes or four doublewords of data. The cache operates as a write-through cache. Note that the cache changes only if a miss occurs. This means that data written to a memory location not already cached are not written to the cache. In many cases, much of the active portion of a program is found completely inside the cache memory. This causes execution to occur at the rate of one clock cycle for many of the instructions that are

FIGURE 17–32 The organization of the 80486 memory, showing parity.

commonly used in a program. About the only way that these efficient instructions are slowed is when the microprocessor must fill a line in the cache. Data are also stored in the cache, but it has less of an impact on the execution speed of a program because data are not referenced repeatedly as many portions of a program are.

Control register 0 (CR_0) is used to control the cache with two new control bits not present in the 80386 microprocessor. (See Figure 17–33 for CR_0 in the 80486 microprocessor.) The CD (cache disable) and NW (noncache write-through) bits are new to the 80486 and are used to control the 8K-byte cache. If the CD bit is a logic 1, all cache operations are inhibited. This setting is used only for debugging software and normally remains cleared. The NW bit is used to inhibit cache write-through operations. As with CD, cache write-through is inhibited only for testing. For normal program operation, CD = 0 and NW = 0.

Because the cache is new to the 80486 microprocessor and the cache is filled by using burst cycles not present on the 80386, some detail is required to understand bus-filling cycles. When a bus line is filled, the 80486 must acquire four 32-bit numbers from the memory system to fill a line in the cache. Filling is accomplished with a burst cycle. The burst cycle is a special memory where four 32-bit numbers are fetched from the memory system in five clocking periods. This assumes that the speed of the memory is sufficient and that no wait states are required. If the clock frequency of the 80486 is 33 MHz, we can fill a cache line in 167 ns, which is very efficient considering that a normal, nonburst 32-bit memory read operation requires two clocking periods.

Memory Read Timing. Figure 17–34 illustrates the read timing for the 80486 for a nonburst memory operation. Note that two clocking periods are used to transfer data. Clocking period T_1 provides the memory address and control signals, and clocking period T_2 is where the data are transferred between the memory and the microprocessor. Note that the $\overline{RDY}$ must become a logic 0 to cause data to be transferred and to terminate the bus cycle. Access time for a nonburst access is determined by taking two clocking periods, minus the time required for the address to appear on the address bus connection, minus a setup time for the data bus connections. For the 20 MHz version of the 80486, two clocking periods require 100 ns minus 28 ns for address setup time and 6 ns for data setup time. This yields a nonburst access time of 100 ns − 34 ns, or 76 ns. Of course, if decoder time and delay times are included, the access time allowed the memory is
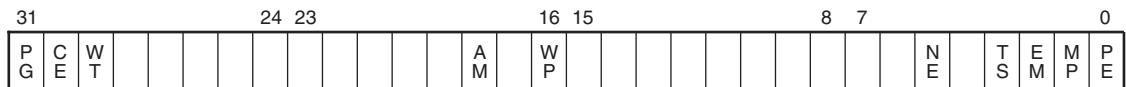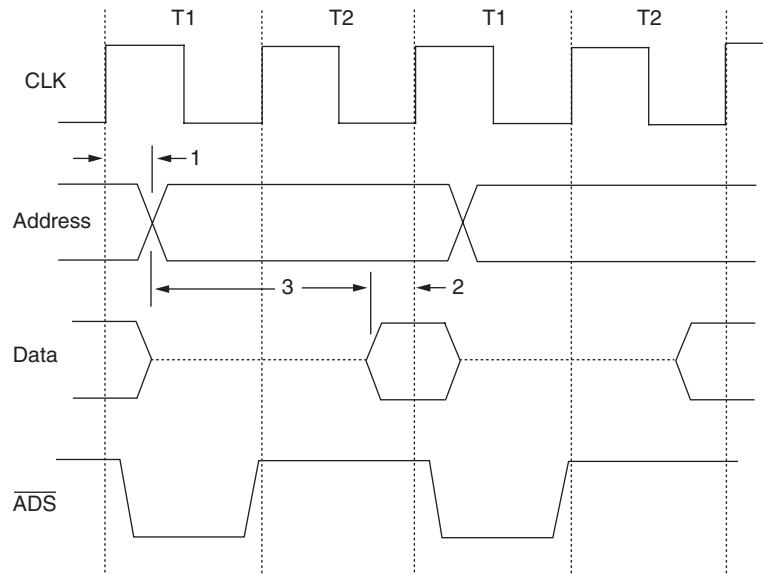
FIGURE 17–33   Control register zero (CR0) for the 80486 microprocessor.

FIGURE 17–34 The non-burst read timing for the 80486 microprocessor.

even less for no wait-state operation. If a higher frequency version of the 80486 is used in a system, memory access time is still less.

The 80486 33 MHz, 66 MHz, and 100 MHz processors all access bus data at a 33 MHz rate. In other words, the microprocessor may operate at 100 MHz, but the system bus operates at 33 MHz. Notice that the nonburst access timing for the 33 MHz system bus allows 60 ns − 24 ns = 36 ns. It is obvious that wait states are required for operation with standard DRAM memory devices.

Figure 17–35 illustrates the timing diagram for filling a cache line with four 32-bit numbers using a burst. Note that the addresses (A_{31}–A_4) appear during T_1 and remain constant throughout the burst cycle. Also, note that A_2 and A_3 change during each T_1 after the first to address four consecutive 32-bit numbers in the memory system. As mentioned, cache fills using bursts require only five clocking periods (one T_1 and four T_2s) to fill a cache line with four doublewords of data.
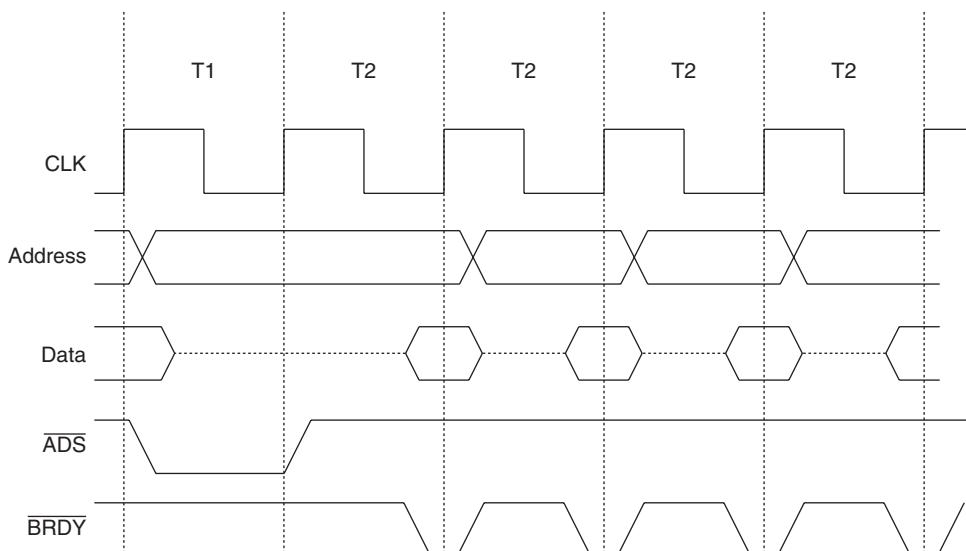
FIGURE 17–35 A burst cycle reads four doublewords in five clocking periods.

Access time using a 20 MHz version of the 80486 for the second and subsequent doublewords is 50 ns − 28 ns − 5 ns, or 17 ns, assuming no delays in the system. To use burst mode transfers, we need high-speed memory. Because DRAM memory access times are 40 ns at best, we are forced to use SRAM for burst cycle transfers. The 33 MHz system allows an access time of 30 ns − 19 ns − 5 ns, or 6 ns for the second and subsequent bytes. If an external counter is used in place of address bits A_2 and A_3, the 19 ns can be eliminated and the access time becomes 30 ns − 5 ns, or 25 ns, which is enough time for even the slowest SRAM connected to the system as a cache. This circuit is often called a *synchronous burst mode cache* if a SRAM cache is used with the system. Note that the $\overline{BRDY}$ pin acknowledges a burst transfer rather than the $\overline{RDY}$ pin, which acknowledges a normal memory transfer.

The PWT controls how the cache functions for a write operation of the external cache memory; it does not control writing to the internal cache. The logic level of this bit is found on the PWT pin of the 80486 microprocessor. Externally, it can be used to dictate the write-through policy of the external cache.

The PCD bit controls the on-chip cache. If the PCD = 0, the on-chip cache is enabled for the current page of memory. Note that 80386 page table entries place a logic 0 in the PCD bit position, enabling caching. If PCD = 1, the on-chip cache is disabled. Caching is disabled, regardless of the condition of $\overline{KEN}$, CD, and NW.

17–8   SUMMARY

1. The 80386 microprocessor is an enhanced version of the 80286 microprocessor and includes a memory-management unit that is enhanced to provide memory paging. The 80386 also includes 32-bit extended registers, and a 32-bit address and data bus. A scaled-down version of the 80386DX with a 16-bit data and 24-bit address bus is available as the 80386SX microprocessor. The 80386EX is a complete AT-style personal computer on a chip.
2. The 80386 has a physical memory size of 40 bytes that can be addressed as a virtual memory with up to 64T bytes. The 80386 memory is 32 bits wide, and it is addressed as bytes, words, or doublewords.
3. When the 80386 is operated in the pipelined mode, it sends the address of the next instruction or memory data to the memory system prior to completing the execution of the current instruction. This allows the memory system to begin fetching the next instruction or data before the current is completed. This increases access time, thus reducing the speed of the memory.
4. A cache memory system allows data that are frequently read to be accessed in less time because they are stored in high-speed semiconductor memory. If data are written to memory, they are also written to the cache, so the most current data are always present in the cache.
5. The I/O structure of the 80386 is almost identical to the 80286, except that I/O can be inhibited when the 80386 is operated in the protected mode through the I/O bit protection map stored with the TSS.
6. In the 80386 microprocessor, interrupts have been expanded to include additional predefined interrupts in the interrupt vector table. These additional interrupts are used with the memory-management system.
7. The 80386 memory manager is similar to the 80286, except that the physical addresses generated by the MMU are 32 bits wide instead of 24 bits wide. The 80386 MMU is also capable of paging.
8. The 80386 is operated in the real mode (8086 mode) when it is reset. The real mode allows the microprocessor to address data in the first 1M byte of memory. In the protected mode, the 80386 addresses any location in its 4G bytes of physical address space.

9. A descriptor is a series of eight bytes that specifies how a code or data segment is used by the 80386. The descriptor is selected by a selector that is stored in one of the segment registers. Descriptors are used only in the protected mode.

10. Memory management is accomplished through a series of descriptors, stored in descriptor tables. To facilitate memory management, the 80386 uses three descriptor tables: the global descriptor table (GDT), the local descriptor table (LDT), and the interrupt descriptor table (IDT). The GDT and LDT each hold up to 8192 descriptors; the IDT holds up to 256 descriptors. The GDT and LDT describe code and data segments as well as tasks. The IDT describes the 256 different interrupt levels through interrupt gate descriptors.

11. The TSS (task state segment) contains information about the current task and the previous task. Appended to the end of the TSS is an I/O bit protection map that inhibits selected I/O port addresses.

12. The memory paging mechanism allows any 4K-byte physical memory page to be mapped to any 4K-byte linear memory page. For example, memory location 00A00000H can be assigned memory location A0000000H through the paging mechanism. A page directory and page tables are used to assign any physical address to any linear address. The paging mechanism can be used in the protected mode or the virtual mode.

13. The 80486 microprocessor is an improved version of the 80386 microprocessor that contains an 8K-byte cache and an 80387 arithmetic coprocessor; it executes many instructions in one clocking period.

14. The 80486 microprocessor executes a few new instructions that control the internal cache memory and allow addition (XADD) and comparison (CMPXCHG) with an exchange and a byte swap (BSWAP) operation. Other than these few additional instructions, the 80486 is 100% upward-compatible with the 80386 and 80387.

15. A new feature found in the 80486 is the BIST (built-in self-test) that tests the microprocessor, coprocessor, and cache at reset time. If the 80486 passes the test, EAX contains a zero.

17–9   QUESTIONS AND PROBLEMS

1. The 80386 microprocessor addresses _____ bytes of memory in the protected mode.
2. The 80386 microprocessor addresses _____ bytes of virtual memory through the memory-management unit.
3. Describe the differences between the 80386DX and the 80386SX.
4. Draw the memory map of the 80386 when operated in the
   (a) protected mode
   (b) real mode
5. How much current is available on various 80386 output pin connections? Compare these currents with the currents available at the output pin connection of an 8086 microprocessor.
6. Describe the 80386 memory system, and explain the purpose and operation of the bank selection signals.
7. Explain the action of a hardware reset on the address bus connections of the 80386.
8. Explain how pipelining lengthens the access time for many memory references in the 80386 microprocessor-based system.
9. Briefly describe how the cache memory system functions.
10. I/O ports in the 80386 start at I/O address _____ and extend to I/O address _____.
11. What I/O ports communicate data between the 80386 and its companion 80387 coprocessor?
12. Compare and contrast the memory and I/O connections found on the 80386 with those found in earlier microprocessors.

13. If the 80386 operates at 20 MHz, what clocking frequency is applied to the CLK_2 pin?

14. What is the purpose of the $\overline{BS16}$ pin on the 80386 microprocessor?

15. Define the purpose of each of the control registers (CR_0, CR_1, CR_2, and CR_3) found within the 80386.

16. Define the purpose of each 80386 debug register.

17. The debug registers cause which level of interrupt?

18. What are the test registers?

19. Select an instruction that copies control register 0 into EAX.

20. Describe the purpose of PE in CR_0.

21. Form an instruction that accesses data in the FS segment at the location indirectly addressed by the DI register. The instruction should store the contents of EAX into this memory location.

22. What is scaled index addressing?

23. Is the following instruction legal? MOV AX,[EBX+ECX]

24. Explain how the following instructions calculate the memory address:
    (a) ADD [EBX+8*ECX],AL
    (b) MOV DATA[EAX+EBX],CX
    (c) SUB EAX,DATA
    (d) MOV ECX,[EBX]

25. What is the purpose of interrupt type number 7?

26. Which interrupt vector type number is activated for a protection privilege violation?

27. What is a double interrupt fault?

28. If an interrupt occurs in the protected mode, what defines the interrupt vectors?

29. What is a descriptor?

30. What is a selector?

31. How does the selector choose the local descriptor table?

32. What register is used to address the global descriptor table?

33. How many global descriptors can be stored in the GDT?

34. Explain how the 80386 can address a virtual memory space of 64T bytes when the physical memory contains only 4G bytes of memory.

35. What is the difference between a segment descriptor and a system descriptor?

36. What is the task state segment (TSS)?

37. How is the TSS addressed?

38. Describe how the 80386 switches from the real mode to the protected mode.

39. Describe how the 80386 switches from the protected mode to the real mode.

40. What is virtual 8086 mode operation of the 80386 microprocessor?

41. How is the paging directory located by the 80386?

42. How many bytes are found in a page of memory?

43. Explain how linear memory address D0000000H can be assigned to physical memory address C0000000H with the paging unit of the 80386.

44. What are the differences between an 80386 and 80486 microprocessor?

45. What is the purpose of the $\overline{FLUSH}$ input pin on the 80486 microprocessor?

46. Compare the register set of the 80386 with the 80486 microprocessor.

47. What differences exist in the flags of the 80486 when compared to the 80386 microprocessor?

48. Which pins are used for parity checking on the 80486 microprocessor?

49. The 80486 microprocessor uses _____ parity.

50. The cache inside the 80486 microprocessor is _____ -K bytes.

51. A cache line is filled by reading _____ -bytes from the memory system.

52. What is an 80486 burst?

53. Define the term *cache write-through*.

54. What is a BIST?

CHAPTER 18

The Pentium and Pentium Pro Microprocessors

INTRODUCTION

The Pentium microprocessor signals an improvement to the architecture found in the 80486 microprocessor. The changes include an improved cache structure, a wider data bus width, a faster numeric coprocessor, a dual integer processor, and branch prediction logic. The cache has been reorganized to form two caches that are each 8K bytes in size, one for caching data, and the other for instructions. The data bus width has been increased from 32 bits to 64 bits. The numeric coprocessor operates at approximately five times faster than the 80486 numeric coprocessor. A dual integer processor often allows two instructions per clock. Finally, the branch prediction logic allows programs that branch to execute more efficiently. Notice that these changes are internal to the Pentium, which makes software upward-compatible from earlier Intel 80X86 microprocessors. A later improvement to the Pentium was the addition of the MMX instructions.

The Pentium Pro is a still faster version of the Pentium. It contains a modified internal architecture that can schedule up to five instructions for execution and an even faster floating-point unit. The Pentium Pro also contains a 256K-byte or 512K-byte level 2 cache in addition to the 16K-byte (8K for data and 8K for instruction) level 1 cache. The Pentium Pro includes error correction circuitry (ECC), as described in Chapter 10, to correct a one-bit error and indicate a two-bit error. Also added are four additional address lines, giving the Pentium Pro access to an astounding 64G bytes of directly addressable memory space.

CHAPTER OBJECTIVES

Upon completion of this chapter, you will be able to:

1. Contrast the Pentium and Pentium Pro with the 80386 and 80486 microprocessors.
2. Describe the organization and interface of the 64-bit-wide Pentium memory system and its variations.
3. Contrast the changes in the memory-management unit and paging unit when compared to the 80386 and 80486 microprocessors.
4. Detail the new instructions found with the Pentium microprocessor.
5. Explain how the superscalar dual integer units improve performance of the Pentium microprocessor.
6. Describe the operation of the branch prediction logic.
7. Detail the improvements in the Pentium Pro when compared with the Pentium.
8. Explain how the dynamic execution architecture of the Pentium Pro functions.

18–1   INTRODUCTION TO THE PENTIUM MICROPROCESSOR

Before the Pentium or any other microprocessor can be used in a system, the function of each pin must be understood. This section of the chapter details the operation of each pin, along with the external memory system and I/O structures of the Pentium microprocessor.

Figure 18–1 illustrates the pin-out of the Pentium microprocessor, which is packaged in a huge 237-pin PGA (pin grid array). The Pentium was made available in two versions: the full-blown Pentium and the P24T version called the Pentium OverDrive. The P24T version contains a 32-bit data bus, compatible for insertion into 80486 machines, which contains the P24T socket. The P24T version also comes with a fan built into the unit. The most notable difference in the

FIGURE 18–1 The pin-out of the Pentium microprocessor.

pin-out of the Pentium, when compared to earlier 80486 microprocessors, is that there are 64 data bus connections instead of 32, which require a larger physical footprint.

As with earlier versions of the Intel family of microprocessors, the early versions of the Pentium require a single +5.0 V power supply for operation. The power supply current averages 3.3 A for the 66 MHz version of the Pentium, and 2.91 A for the 60 MHz version. Because these currents are significant, so are the power dissipations of these microprocessors: 13 W for the 66 MHz version and 11.9 W for the 60 MHz version. The current versions of the Pentium, 90 MHz and above, use a 3.3 V power supply with reduced current consumption. At present, a good heat sink with considerable airflow is required to keep the Pentium cool. The Pentium contains multiple V_{CC} and V_{SS} connections that must all be connected to +5.0 V or +3.3 V and ground for proper operation. Some of the pins are labeled N/C (no connection) and must not be connected. The latest versions of the Pentium have been improved to reduce the power dissipation. For example, the 233 MHz Pentium requires 3.4 A or current, which is only slightly more than the 3.3 A required by the early 66 MHz version.

Each Pentium output pin is capable of providing 4.0 mA of current at a logic 0 level and 2.0 mA at a logic 1 level. This represents an increase in drive current, compared to the 2.0 mA available on earlier 8086, 8088, and 80286 output pins. Each input pin represents a small load requiring only 15 μA of current. In some systems, except the smallest, these current levels require bus buffers.

The function of each Pentium group of pins follows:

$\overline{A20}$	The **address A20 mask** is an input that is asserted in the real mode to signal the Pentium to perform address wraparound, as in the 8086 microprocessor, for use of the HIMEM.SYS driver.
A_{31}–A_3	**Address bus** connections address any of the 5l2K × 64 memory locations found in the Pentium memory system. Note that A_0, A_1, and A_2 are encoded in the bus enable ($\overline{BE7}$–$\overline{BE0}$), described elsewhere, to select any or all of the eight bytes in a 64-bit-wide memory location.
$\overline{ADS}$	The **address data strobe** becomes active whenever the Pentium has issued a valid memory or I/O address. This signal is combined with the $W/\overline{R}$ and $M/\overline{IO}$ signals to generate the separate read and write signals present in the earlier 8086–80286 microprocessor-based systems.
AHOLD	**Address hold** is an input that causes the Pentium to hold the address and AP signals for the next clock.
$\overline{APCHK}$	**Address parity** provides even parity for the memory address on all Pentium-initiated memory and I/O transfers. The AP pin must also be driven with even parity information on all inquire cycles in the same clocking period as the $\overline{EADS}$ signal. Address parity check becomes a logic 0 whenever the Pentium detects an address parity error.
$\overline{BE7}$–$\overline{BE0}$	**Bank enable** signals select the access of a byte, word, doubleword, or quadword of data. These signals are generated internally by the microprocessor from address bits A_0, A_1, and A_2.
$\overline{BOFF}$	The **back-off** input aborts all outstanding bus cycles and floats the Pentium buses until $\overline{BOFF}$ is negated. After $\overline{BOFF}$ is negated, the Pentium restarts all aborted bus cycles in their entirety.
BP_3–BP_0	The **breakpoint** pins BP_3–BP_0 indicate a breakpoint match when the debug registers are programmed to monitor for matches.
PM_1–PM_0	The **performance monitoring** pins PM_1 and PM_0 indicate the settings of the performance monitoring bits in the debug mode control register.

$\overline{\text{BRDY}}$	The **burst ready** input signals the Pentium that the external system has applied or extracted data from the data bus connections. This signal is used to insert wait states into the Pentium timing.
BREQ	The **bus request** output indicates that the Pentium has generated a bus request.
$BT_3–BT_0$	The **branch trace** outputs provide bits 2–0 of the branch target linear address and the default operand size on BT_3. These outputs become valid during a branch trace special message cycle.
$\overline{\text{BUSCHK}}$	The **bus check** input allows the system to signal the Pentium that the bus transfer has been unsuccessful.
$\overline{\text{CACHE}}$	The **cache output** indicates that the current Pentium cycle can cache data.
CLK	The **clock** is driven by a clock signal that is at the operating frequency of the Pentium. For example, to operate the Pentium at 66 MHz, apply a 66 MHz clock to this pin.
$D_{63}–D_0$	**Data bus** connections transfer byte, word, doubleword, and quadword data between the microprocessor and its memory and I/O system.
$D/\overline{C}$	**Data/control** indicates that the data bus contains data for or from memory or I/O when a logic 1. If $D/\overline{C}$ is a logic 0, the microprocessor is either halted or executing an interrupt acknowledge.
$DP_7–DP_0$	**Data parity** is generated by the Pentium and detects its eight memory banks through these connections.
$\overline{\text{EADS}}$	The **external address strobe** input signals that the address bus contains an address for an inquire cycle.
$\overline{\text{EWBE}}$	The **external write buffer empty** input indicates that a write cycle is pending in the external system.
$\overline{\text{FERR}}$	The **floating-point error** is comparable to the $\overline{\text{ERROR}}$ line in the 80386 and shows that the internal coprocessor has erred.
$\overline{\text{FLUSH}}$	The **flush cache** input causes the cache to flush all write-back lines and invalidate its internal caches. If the $\overline{\text{FLUSH}}$ input is a logic 0 during a reset operation, the Pentium enters its test mode.
$\overline{\text{FRCMC}}$	The **functional redundancy check** is sampled during a reset to configure the Pentium in the master (1) or checker mode (0).
$\overline{\text{HIT}}$	**Hit** shows that the internal cache contains valid data in the inquire mode.
$\overline{\text{HITM}}$	**Hit modified** shows that the inquire cycle found a modified cache line. This output is used to inhibit other master units from accessing data until the cache line is written to memory.
HOLD	**Hold** requests a DMA action.
HLDA	**Hold acknowledge** indicates that the Pentium is currently in a hold condition.
IBT	**Instruction branch taken** indicates that the Pentium has taken an instruction branch.
$\overline{\text{IERR}}$	The **internal error** output shows that the Pentium has detected an internal parity error or functional redundancy error.
$\overline{\text{IGNNE}}$	The **ignore numeric error** input causes the Pentium to ignore a numeric coprocessor error.
INIT	The **initialization** input performs a reset without initializing the caches, write-back buffers, and floating-point registers. This may not be used to reset the microprocessor in lieu of RESET after power-up.

INTR	The **interrupt request** is used by external circuitry to request an interrupt.
INV	The **invalidation** input determines the cache line state after an inquiry.
IU	The **U-pipe instruction complete** output shows that the instruction in the U-pipe is complete.
IV	The **V-pipe instruction complete** output shows that the instruction in the V-pipe is complete.
$\overline{\text{KEN}}$	The **cache enable** input enables internal caching.
$\overline{\text{LOCK}}$	**LOCK** becomes a logic 0 whenever an instruction is prefixed with the LOCK: prefix. This is most often used during DMA accesses.
M/$\overline{\text{IO}}$	**Memory/IO** selects a memory device when a logic 1 or an I/O device when a logic 0. During the I/O operation, the address bus contains a 16-bit I/O address on address connections A_{15}–A_3.
$\overline{\text{NA}}$	**Next address** indicates that the external memory system is ready to accept a new bus cycle.
NMI	The **non-maskable interrupt** requests a non-maskable interrupt, just as on the earlier versions of the microprocessor.
PCD	The **page cache disable** output shows that the internal page caching is disabled by reflecting the state of the CR_3 PCD bit.
$\overline{\text{PCHK}}$	The **parity check** output signals a parity check error for data read from memory or I/O.
$\overline{\text{PEN}}$	The **parity enable** input enables the machine check interrupt or exception.
PRDY	The **probe ready** output indicates that the probe mode has been entered for debugging.
PWT	The **page write-through** output shows the state of the PWT bit in CR_3. This pin is provided for use with the Intel Debugging Port and causes an interrupt.
RESET	**Reset** initializes the Pentium, causing it to begin executing software at memory location FFFFFFF0H. The Pentium is reset to the real mode and the leftmost 12 address connections remain logic 1s (FFFH) until a far jump or far call is executed. This allows compatibility with earlier microprocessors. See Table 18–1 for the state of the Pentium after a hardware reset.

TABLE 18–1   State of the Pentium after a RESET.

Register	RESET Value	RESET + BIST Value
EAX	0	0 (it test passes)
EDX	0500XXXXH	0500XXXXH
EBX, ECX, ESP, EBP, ESI, and EDI	0	0
EFLAGS	2	2
EIP	0000FFF0H	0000FFF0H
CS	F000H	F000H
DS, ES, FS, GS, and SS	0	0
GDTR and TSS	0	0
CR_0	60000010H	60000010H
CR_2, CR_3, and CR_4	0	0
DR_0–DR_3	0	0
DR_6	FFFF0FF0H	FFFF0FF0H
DR_7	00000400H	00000400H

SCYC   The **split cycle** output signals a misaligned LOCKed bus cycle.

$\overline{\textbf{SMI}}$ The **system management interrupt** input causes the Pentium to enter the system management mode of operation.

$\overline{\textbf{SMIACT}}$ The **system management interrupt active** output shows that the Pentium is operating in the system management mode.

TCK   The **testability clock** input selects the clocking function in accordance to the IEEE 1149.1 Boundary Scan interface.

TDI   The **test data input** is used to test data clocked into the Pentium with the TCK signal.

TDO   The **test data output** is used to gather test data and instructions shifted out of the Pentium with TCK.

TMS   The **test mode select** input controls the operation of the Pentium in test mode. The test reset input allows the test mode to be reset.

$\textbf{W}/\overline{\textbf{R}}$ **Write/read** indicates that the current bus cycle is a write when a logic 1 or a read when a logic 0.

$\textbf{WB}/\overline{\textbf{WT}}$ **Write-back/write-through** selects the operation for the Pentium data cache.

The Memory System

The memory system for the Pentium microprocessor is 4G bytes in size, just as in the 80386DX and 80486 microprocessors. The difference lies in the width of the memory data bus. The Pentium uses a 64-bit data bus to address memory organized in eight banks that each contain 512M bytes of data. See Figure 18–2 for the organization of the Pentium physical memory system.

The Pentium memory system is divided into eight banks where each bank stores byte-wide data with a parity bit. The Pentium, like the 80486, employs internal parity generation and checking logic for the memory system's data bus information. (Note that most Pentium systems do not use parity checks, because ECC is available.) The 64-bit-wide memory is important to double-precision floating-point data. Recall that a double-precision floating-point number is 64 bits wide. Because of the change to a 64-bit-wide data bus, the Pentium is able to retrieve floating-point data with one read cycle, instead of two as in the 80486. This causes the Pentium to function at a higher throughput than an 80486. As with earlier 32-bit Intel microprocessors, the memory system is numbered in bytes, from byte 00000000H to byte FFFFFFFFH.

Memory selection is accomplished with the bank enable signals ($\overline{BE7}$–$\overline{BE0}$). These separate memory banks allow the Pentium to access any single byte, word, doubleword, or quadword with one memory transfer cycle. As with earlier memory selection logic, eight separate write strobes are generated for writing to the memory system.

A new feature added to the Pentium is its capability to check and generate parity for the address bus (A_{31}–A_5) during certain operations. The AP pin provides the system with parity
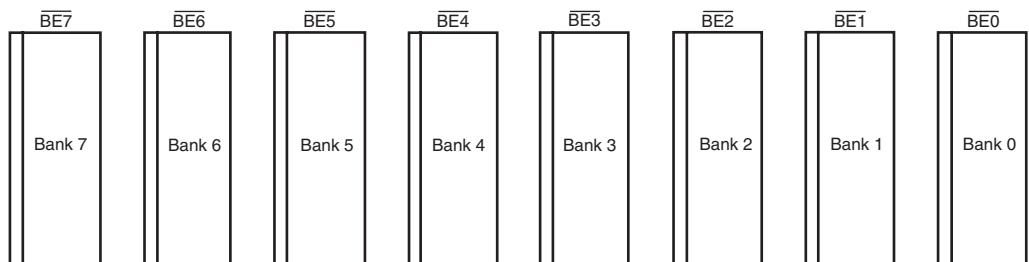
FIGURE 18–2   The 8-byte-wide memory banks of the Pentium microprocessor.

information and the $\overline{\text{APCHK}}$ indicates a bad parity check for the address bus. The Pentium takes no action when an address parity error is detected. The error must be assessed by the system and the system must take appropriate action (an interrupt), if so desired.

Input/Output System

The input/output system of the Pentium is completely compatible with earlier Intel microprocessors. The I/O port number appears on address lines A_{15}–A_3 with the bank enable signals used to select the actual memory banks used for the I/O transfer.

Beginning with the 80386 microprocessor, I/O privilege information is added to the TSS segment when the Pentium is operated in the protected mode. Recall that this allows I/O ports to be selectively inhibited. If the blocked I/O location is accessed, the Pentium generates a type 13 interrupt to signal an I/O privilege violation.

System Timing

As with any microprocessor, the system timing signals must be understood in order to interface the microprocessor. This portion of the text details the operation of the Pentium through its timing diagrams and shows how to determine memory access times.

The basic Pentium nonpipelined memory cycle consists of two clocking periods: T_1 and T_2. See Figure 18–3 for the basic nonpipelined read cycle. Notice from the timing diagram that the 66 MHz Pentium is capable of 33 million memory transfers per second. This assumes that the memory can operate at that speed.

Also notice from the timing diagram that the $W/\overline{R}$ signal becomes valid if $\overline{\text{ADS}}$ is a logic 0 at the positive edge of the clock (end of T_1). This clock must be used to qualify the cycle as a read or a write.

During T_1, the microprocessor issues the $\overline{\text{ADS}}$, $W/\overline{R}$, address, and $M/\overline{\text{IO}}$ signals. In order to qualify the $W/\overline{R}$ signal and generate appropriate $\overline{\text{MRDC}}$ and $\overline{\text{MWTC}}$ signals, we use a flip-flop
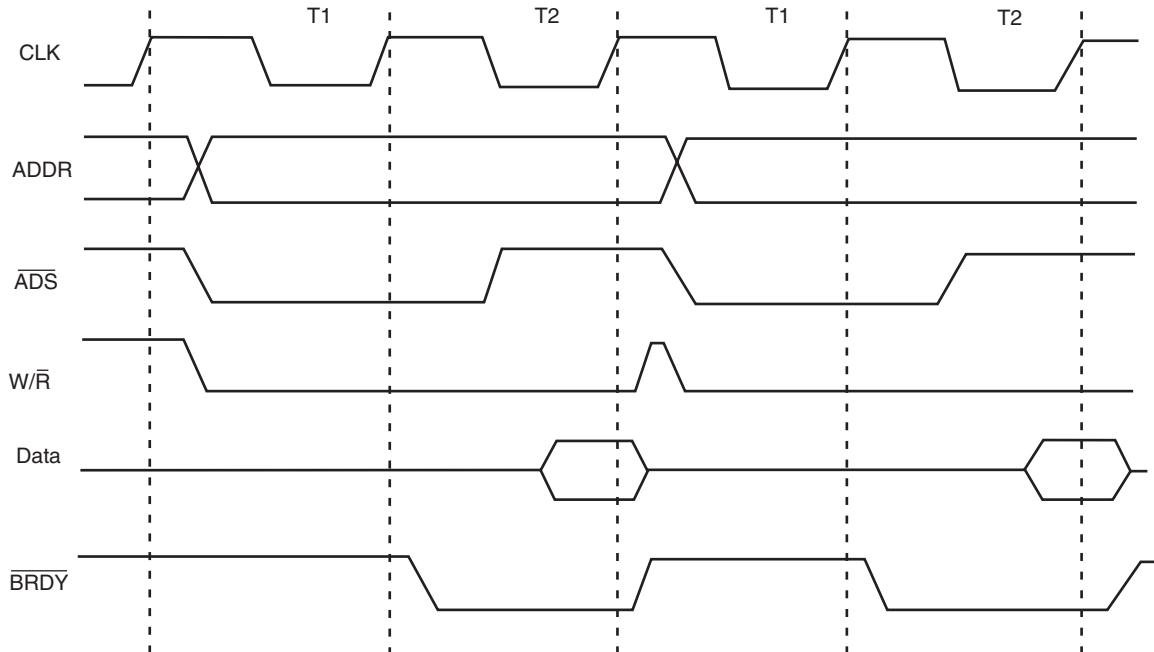
FIGURE 18–3   The nonpipelined read cycle for the Pentium microprocessor.
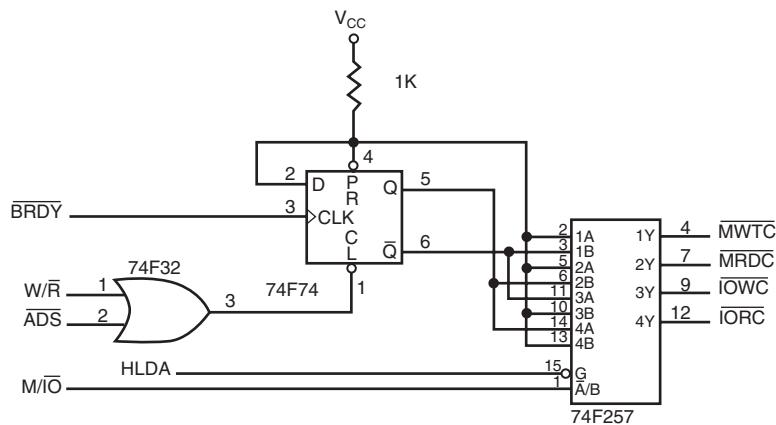
FIGURE 18–4   A circuit that generates the memory and I/O control signals.

to generate the W/$\overline{\text{R}}$ signal. Then a two-line-to-one-line multiplexer generates the memory and I/O control signals. See Figure 18–4 for a circuit that generates the memory and I/O control signals for the Pentium microprocessor.

During T_2, the data bus is sampled in synchronization with the end of T_2 at the positive transition of the clock pulse. The setup time before the clock is given as 3.8 ns, and the hold time after the clock is given as 2.0 ns. This means that the data window around the clock is 5.8 ns. The address appears on the 8.0 ns maximum after the start of T_1. This means that the Pentium microprocessor operating at 66 MHz allows 30.3 ns (two clocking periods), minus the address delay time of 8.0 ns and minus the data setup time of 3.8 ns. Memory access time without any wait states is $30.3 - 8.0 - 3.8$, or 18.5 ns. This is enough time to allow access to a SRAM, but not to any DRAM without inserting wait states into the timing. The SRAM is normally found in the form of an external level 2 cache.

Wait states are inserted into the timing by controlling the $\overline{\text{BRDY}}$ input to the Pentium. The $\overline{\text{BRDY}}$ signal must become a logic 0 by the end of T_2 or additional T_2 states are inserted into the timing. See Figure 18–5 for a read cycle timing diagram that contains wait states for slower
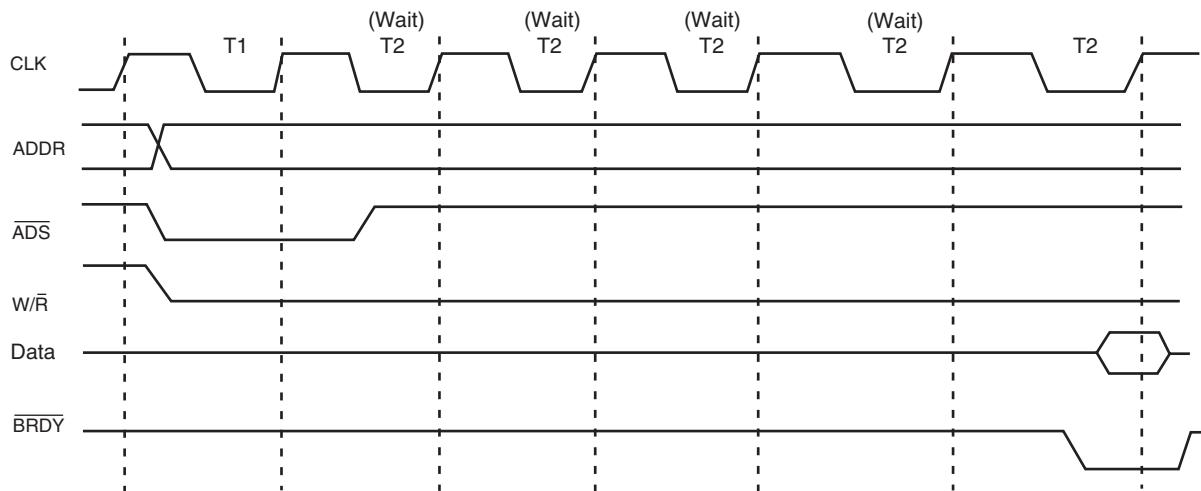
FIGURE 18–5   The Pentium timing diagram with four wait states inserted for an access time of 79.5 ns.
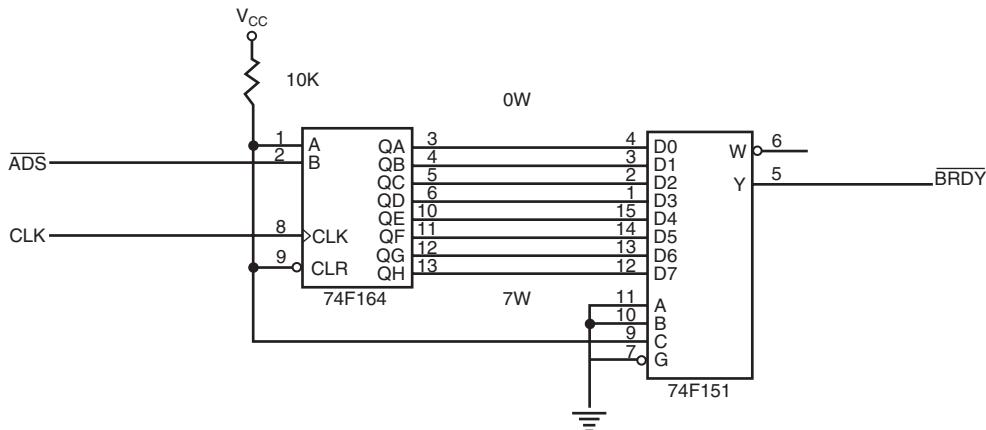
FIGURE 18–6 A circuit that generates wait states by delaying ADS. This circuit is wired to generate four wait states.

memory. The effect of inserting wait states into the timing is to lengthen the timing, allowing additional time for the memory to access data. In the timing shown, the access time has been lengthened so that standard 60 ns DRAM can be used in a system. Note that this requires the insertion of four wait states of 15.2 ns (one clocking period) each to lengthen the access time to 79.5 ns. This is enough time for the DRAM and any decoder in the system to function.

The $\overline{\text{BRDY}}$ signal is a synchronous signal generated by using the system clock. Figure 18–6 illustrates a circuit that can be used to generate $\overline{\text{BRDY}}$ for inserting any number of wait states into the Pentium timing diagram. You may recall a similar circuit inserting wait states into the timing diagram of the 80386 microprocessor. The $\overline{\text{ADS}}$ signal is delayed between 0 and 7 clocking periods by the 74Fl61 shift register to generate the $\overline{\text{BRDY}}$ signal. The exact number of wait states is selected by the 74F151 eight-line-to-one-line multiplexer. In this example, the multiplexer selects the four-wait output from the shift register.

A more efficient method of reading memory data is via the burst cycle. The burst cycle in the Pentium transfers four 64-bit numbers per burst cycle in five clocking periods. A burst without wait states requires that the memory system transfers data every 15.2 ns. If a level 2 cache is in place, this speed is no problem as long as the data are read from the cache. If the cache does not contain the data, then wait states must be inserted, which will reduce the data throughput. See Figure 18–7 for the Pentium burst cycle transfer without wait states. As before, wait states can be inserted to allow more time to the memory system for accesses.
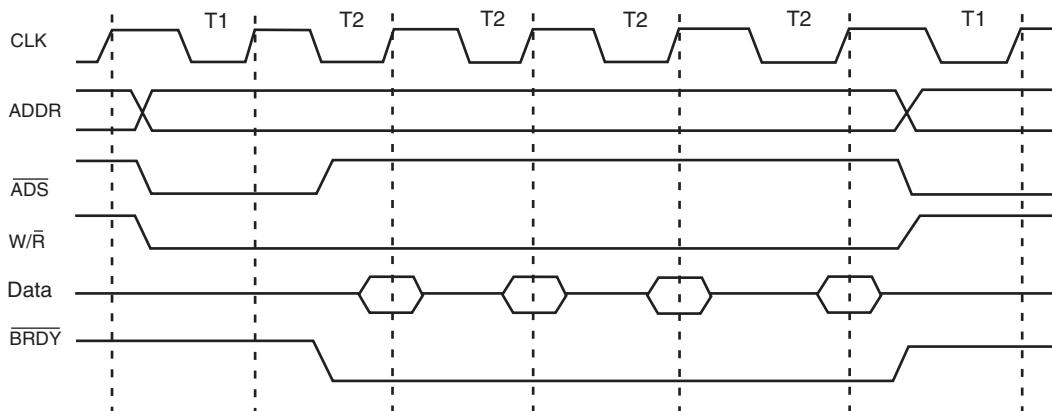
FIGURE 18–7 The Pentium burst cycle operation that transfers four 64-bit data between the microprocessor and memory.

Branch Prediction Logic

The Pentium microprocessor uses branch prediction logic to reduce the time required for a branch caused by internal delays. These delays are minimized because when a branch instruction (short or near only) is encountered, the microprocessor begins prefetch instruction at the branch address. The instructions are loaded into the instruction cache, so when the branch occurs, the instructions are present and allow the branch to execute in one clocking period. If for any reason the branch prediction logic errs, the branch requires an extra three clocking periods to execute. In most cases, the branch prediction is correct and no delay ensues.

Cache Structure

The cache in the Pentium has been changed from the one found in the 80486 microprocessor. The Pentium contains two 8K-byte cache memories instead of one as in the 80486. There is an 8K-byte data cache and an 8K-byte instruction cache. The instruction cache stores only instructions, while the data cache stores data used by instructions.

    In the 80486 with its unified cache, a program that was data-intensive quickly filled the cache, allowing little room for instructions. This slowed the execution speed of the 80486 microprocessor. In the Pentium, this cannot occur because of the separate instruction cache.

Superscalar Architecture

The Pentium microprocessor is organized with three execution units. One executes floating-point instructions, and the other two (U-pipe and V-pipe) execute integer instructions. This means that it is possible to execute three instructions simultaneously. For example, the FADD ST,ST(2) instruction, MOV EAX,10H instruction, and MOV EBX,12H instruction can all execute simultaneously because none of these instructions depend on each other. The FADD ST,ST(2) instruction is executed by the coprocessor; the MOV EAX,10H is executed by the U-pipe; and the MOV EBX,12H instruction is executed by the V-pipe. Because the floating-point unit is also used for MMX instructions, if available, the Pentium can execute two integers and one MMX instruction simultaneously.

    Software should be written to take advantage of this feature by looking at the instructions in a program, and then modifying them when cases are discovered in which dependent instructions can be separated by nondependent instructions. These changes can result in up to a 40% execution speed improvement in some software. Make sure that any new compiler or other application package takes advantage of this new superscalar feature of the Pentium.

18–2   SPECIAL PENTIUM REGISTERS

The Pentium is essentially the same microprocessor as the 80386 and 80486, except that some additional features and changes to the control register set have occurred. This section highlights the differences between the 80386 control register structure and the flag register.
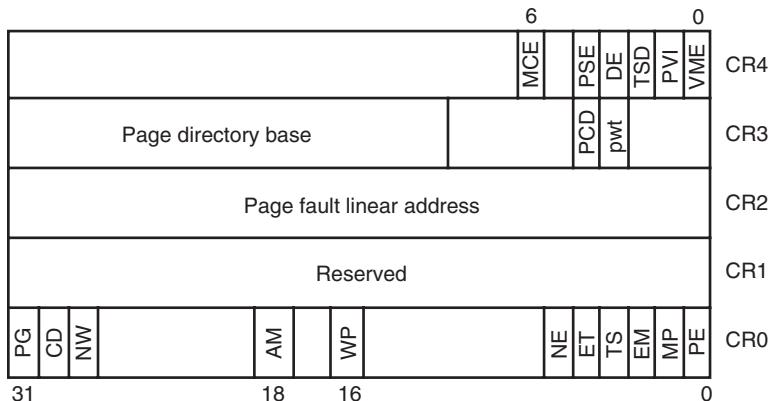
Control Registers

Figure 18–8 shows the control register structure for the Pentium microprocessor. Note that a new control register, CR_4, has been added to the control register array.

    This section of the text only explains the new Pentium components in the control registers. See Figure 17-14 for a description and illustration of the 80386 control registers. Following is a description of the new control bits and new control register CR_4:

CD   **Cache disable** controls the internal cache. If CD = 1, the cache will not fill with new data for cache misses, but it will continue to function for cache hits. If CD = 0, misses will cause the cache to fill with new data.

FIGURE 18–8 The structure of the Pentium control registers.

NW   **Not write-through** selects the mode of operation for the data cache. If NW = 1, the data cache is inhibited from cache write-through.

AM   **Alignment mask** enables alignment checking when set. Note that alignment checking only occurs for protected mode operation when the user is at privilege level 3.

WP   **Write protect** protects user-level pages against supervisor-level write operations. When WP = 1, the supervisor can write to user-level segments.

NE   **Numeric error** enables standard numeric coprocessor error detection. If NE = 1, the $\overline{\text{FERR}}$ pin becomes active for a numeric coprocessor error. If NE = 0, any coprocessor error is ignored.

VME   **Virtual mode extension** enables support for the virtual interrupt flag in protected mode. If VME = 0, virtual interrupt support is disabled.

PVI   **Protected mode virtual interrupt** enables support for the virtual interrupt flag in protected mode.

TSD   **Time stamp disable** controls the RDTSC instruction.

DE   **Debugging extension** enables I/O breakpoint debugging extensions when set.

PSE   **Page size extension** enables 4M-byte memory pages when set.

MCE   **Machine check enable** enables the machine checking interrupt.

The Pentium contains new features that are controlled by CR_4 and a few bits in CR_0. These new features are explained in later sections of the text.

EFLAG Register

The extended flag (EFLAG) register has been changed in the Pentium microprocessor. Figure 18–9 pictures the contents of the EFLAG register. Note that four new flag bits have been added to this register to control or indicate conditions about some of the new features in the Pentium. Following is a list of the four new flags and the function of each:

ID   The **identification flag** is used to test for the CPUID instruction. If a program can set and clear the ID flag, the processor supports the CPUID instruction.

Note: The blank bits in the flag register are reserved for future use and must not be defined.

FIGURE 18–9 The structure of the Pentium EFLAG register.

VIP **Virtual interrupt pending** indicates that a virtual interrupt is pending.

VIF **Virtual interrupt** is the image of the interrupt flag IF used with VIP.

AC **Alignment check** indicates the state of the AM bit in control register 0.

Built-In Self-Test (BIST)

The built-in self-test (BIST) is accessed on power-up by placing a logic 1 on INIT while the RESET pin changes from 1 to 0. The BIST tests 70% of the internal structure of the Pentium in approximately 150 μs. Upon completion of the BIST, the Pentium reports the outcome in register EAX. If EAX = 0, the BIST has passed and the Pentium is ready for operation. If EAX contains any other value, the Pentium has malfunctioned and is faulty.

18–3   PENTIUM MEMORY MANAGEMENT

The memory-management unit within the Pentium is upward-compatible with the 80386 and 80486 microprocessors. Many of the features of these earlier microprocessors are basically unchanged in the Pentium. The main change is in the paging unit and a new system memory-management mode.

Paging Unit

The paging mechanism functions with 4K-byte memory pages or with a new extension available to the Pentium with 4M-byte memory pages. As detailed in Chapters 1 and 17, the size of the paging table structure can become large in a system that contains a large memory. Recall that to fully repage 4G bytes of memory, the microprocessor requires slightly over 4M bytes of memory just for the page tables. In the Pentium, with the new 4M-byte paging feature, this is dramatically reduced to just a single page directory and no page tables. The new 4M-byte page sizes are selected by the PSE bit in control register 0.

The main difference between 4K paging and 4M paging is that in the 4M paging scheme there is no page table entry in the linear address. See Figure 18–10 for the 4M paging system in the Pentium microprocessor. Pay close attention to the way the linear address is used with this scheme. Notice that the leftmost 10 bits of the linear address select an entry in the page directory (just as with 4K pages). Unlike 4K pages, there are no page tables; instead, the page directory addresses a 4M-byte memory page.

Memory-Management Mode

The system memory-management mode (SMM) is on the same level as protected mode, real mode, and virtual mode, but it is provided to function as a manager. The SMM is not intended to be used as an application or a system-level feature. It is intended for high-level system functions such as power management and security, which most Pentiums use during operation, but that are controlled by the operating system.

Access to the SMM is accomplished via a new external hardware interrupt applied to the SMI pin on the Pentium. When the SMM interrupt is activated, the processor begins executing system-level software in an area of memory called the system management RAM, or SMMRAM, called the SMM state dump record. The SMI interrupt disables all other interrupts that are normally handled by user applications and the operating system. A return from the SMM interrupt is accomplished with a new instruction called RSM. RSM returns from the memory-management mode interrupt and returns to the interrupted program at the point of the interruption.
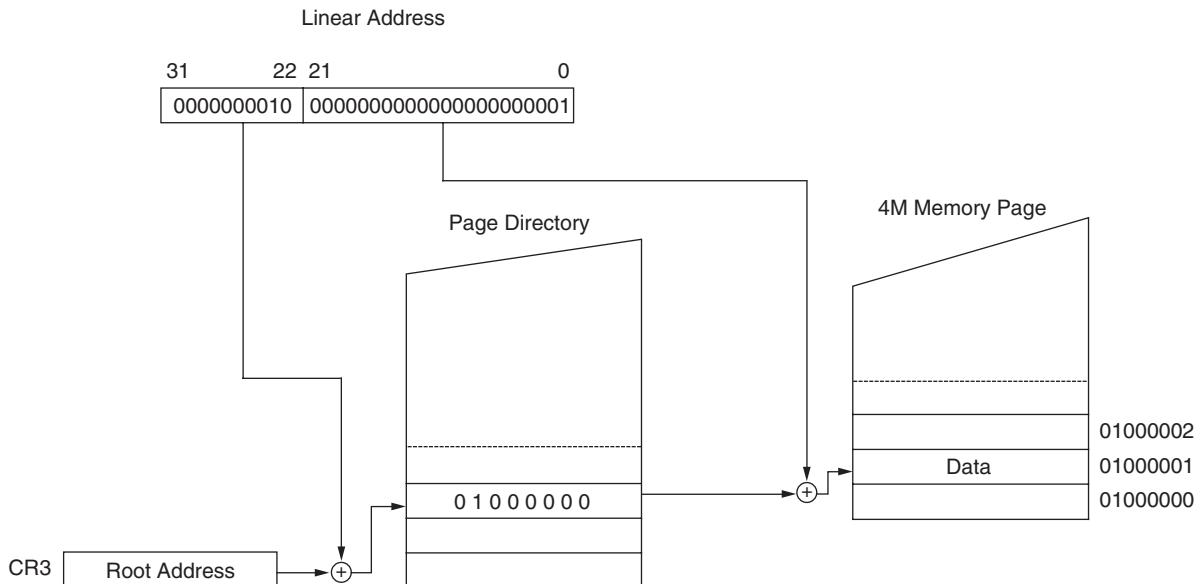
FIGURE 18–10 The linear address 00200001H repaged to memory location 01000002H in 4M-byte pages. Note that there are no page tables.

The SMM interrupt calls the software, initially stored at memory location 38000H, using CS = 3000H and EIP = 8000H. This initial state can be changed using a jump to any location within the first 1M byte of the memory. An environment similar to real mode memory addressing is entered by the management mode interrupt, but it is different because, instead of being able to address the first 1M of memory, SMM mode allows the Pentium to treat the memory system as a flat, 4G-byte system.

In addition to executing software that begins at location 38000H, the SMM interrupt also stores the state of the Pentium in what is called a **dump record**. The dump record is stored at memory locations 3FFA8H through 3FFFFH, with an area at locations 3FE00H through 3FEF7H that is reserved by Intel. The dump record allows a Pentium-based system to enter a sleep mode and reactivate at the point of program interruption. This requires that the SMMRAM be powered during the sleep period. Many laptop computers have a separate battery to power the SMMRAM for many hours during sleep mode. Table 18–2 lists the contents of the dump record.

The Halt auto restart and I/O trap restarts are used when the SMM mode is exited by the RSM instruction. These data allow the RSM instruction to return to the halt-state or return to the interrupt I/O instruction. If neither a halt nor an I/O operation is in effect upon entering the SMM mode, the RSM instruction reloads the state of the machine from the state dump and returns to the point of interruption.

The SMM mode can be used by the system before the normal operating system is placed in the memory and executed. It can also be used periodically to manage the system, provided that normal software doesn't exist at locations 38000H–3FFFFH. If the system relocates the SMRAM before booting the normal operating system, it becomes available for use in addition to the normal system.

The base address of the SMM mode SMRAM is changed by modifying the value in the state dump base address register (locations 3FEF8H through 3F3FBH) after the first memory-management mode interrupt. When the first RSM instruction is executed, returning control back to the interrupted system, the new value from these locations changes the base address of the SMM interrupt for all future uses. For example, if the state dump base address is changed to

TABLE 18–2   Pentium SMM
state dump record.

Offset Address	Register
FFFCH	CR_0
FFF8H	CR_3
FFF4H	EFLAGS
FFF0H	EIP
FFECH	EDI
FFE8H	ESI
FFE4H	EBP
FFE0H	ESP
FFDCH	EBX
FFD8H	EDX
FFD4H	ECX
FFD0H	EAX
FFCCH	DR_6
FFC8H	DR_7
FFC4H	TR
FFC0H	LDTR
FFBCH	GS
FFB8H	FS
FFB4H	DS
FFB0H	SS
FFACH	CS
FFA8H	ES
FF04H–FFA7H	Reserved
FF02H	Halt auto start
FF00H	I/O trap restart
FEFCH	SMM revision identifier
FED8H	State dump base
FE00H–FEF7H	Reserved

Note: The offset addresses are initially located at
base address 00003000H.

000E8000H, all subsequent SMM interrupts use locations E8000H–EFFFFH for the Pentium state dump. These locations are compatible with DOS and Windows.

18–4   NEW PENTIUM INSTRUCTIONS

The Pentium contains only one new instruction that functions with normal system software; the remainder of the new instructions are added to control the memory-management mode feature and serializing instructions. Table 18–3 lists the new instructions added to the Pentium instruction set.

The CMPXCHG8B instruction is an extension of the CMPXCHG instruction added to the 80486 instruction set. The CMPXCHG8B instruction compares the 64-bit number stored in EDX and EAX with the contents of a 64-bit memory location or register pair. For example, the CMPXCHG8B $DATA_2$ instruction compared the eight bytes stored in memory location $DATA_2$ with the 64-bit number in EDX and EAX. If $DATA_2$ equals EDX:EAX, the 64-bit number stored in ECX:EBX is stored in memory location $DATA_2$. If they are not equal, the contents of $DATA_2$ are stored into EDX:EAX. Note that the zero flag bit indicates that the contents of EDX:EAX were equal or not equal to $DATA_2$.

The CPUID instruction reads the CPU identification code and other information from the Pentium. Table 18–4 shows different information returned from the CPUID instruction for various input values for EAX. To use the CPUID instruction, first load EAX with the input value and then execute CPUID. The information is returned in the registers indicated in the table.

TABLE 18–3 New Pentium instructions.

Instruction	Function
CMPXCHG8B	Compare and exchange eight bytes
CPUID	Return CPU identification code
RDTSC	Read time-stamp counter
RDMSR	Read model-specific register
WRMSR	Write model-specific register
RSM	Return from system management interrupt

TABLE 18–4 CPUID instruction information.

Input Value (EAX)	Result after CPUID
0	EAX = 1 for all microprocessor
	EBX–EDX–ECX = Vendor information
1	EAX (bits 3–0) = Stepping ID
	EAX (bits 7–4) = Model
	EAX (bits 11–8) = Family
	EAX (bits 13–12) = Type
	EAX (bits 31–14) = Reserved
	EDX (bit 0) = CPU contains FPU
	EDX (bit 1) = Enhanced 8086 virtual mode supported
	EDX (bit 2) = I/O breakpoints supported
	EDX (bit 3) = Page size extensions supported
	EDX (bit 4) = Time-stamp counter supported
	EDX (bit 5) = Pentium-style MSR supported
	EDX (bit 6) = Reserved
	EDX (bit 7) = Machine check exception supported
	EDX (bit 8) = CMPXCHG8B supported
	EDX (bit 9) = 3.3V microprocessor
	EDX (bits 10–31) = Reserved

If a 0 is placed in EAX before executing the CPUID instruction, the microprocessor returns the vendor identification in EBX, EDX, and EBX. For example, the Intel Pentium returns "GenuineIntel" in ASCII code with the "Genu" in the EBX, "ineI" in EDX, and "ntel" in ECX. The EDX register returns information if EAX is loaded with a 1 before executing the CPUID instruction.

Example 18-1 illustrates a short program that reads the vendor information with the CPUID instruction. This software was placed into the TODO: section of the OnInitDialog function of a simple dialog application. It then displays it on the video screen in an ActiveX label as illustrated in Figure 18–11. The CPUID instruction functions in both the real and protected mode and can be used in any Windows application.

EXAMPLE 18–1

```
CString temp;
int a, b, c;
_asm
{
      mov eax,0
      cpuid
      mov a,ebx
      mov b,edx
      mov c,ecx
}
for (int d = 0; d < 4; d++ )
{
      temp += (char)a;
      a >>= 8;
}
```
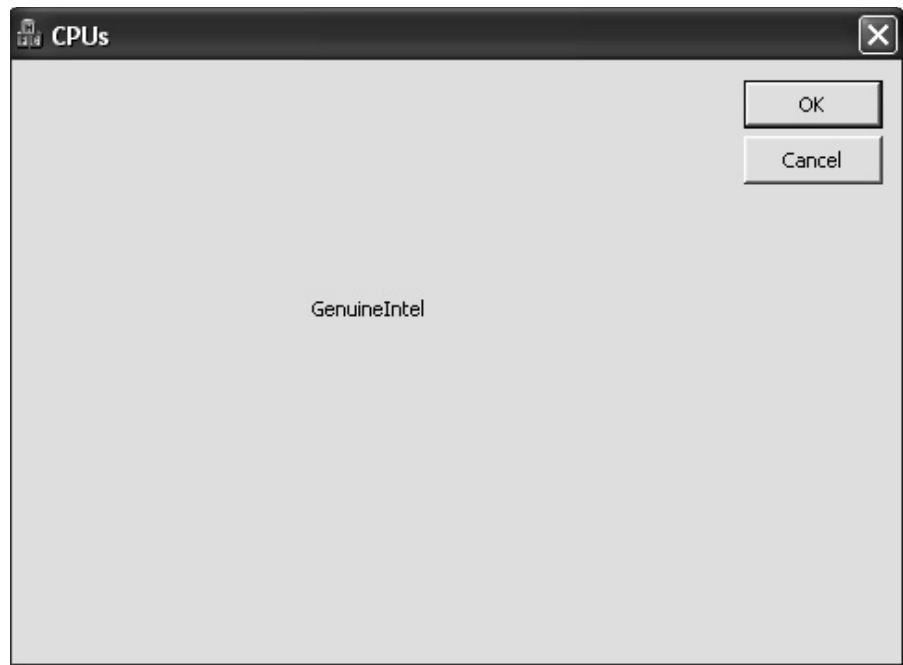
FIGURE 18–11 Screen shot of the program of Example 18–1 using the CPUID instruction.

```
for (d = 0; d < 4; d++ )
{
        temp += (char)b;
        b >>= 8;
}
for (d = 0; d < 4; d++ )
{
        temp += (char)c;
        c >>= 8;
}
Label1.put_Caption(temp);
```

The RDTSC instruction reads the time-stamp counter into EDX:EAX. The time-stamp counter counts CPU clocks from the time the microprocessor is reset, where the time-stamp counter is initialized to an unknown count. Because this is a 64-bit count, a 1GHz microprocessor can accumulate a count of over 580 years before the time-stamp counter rolls over. This instruction functions only in real mode or privilege level 0 in protected mode.

Example 18-2 shows a class written for Windows that provides member functions for accurate time delays and also member functions to measure software execution times. This class is added by right-clicking on the project name and inserting an MFC generic class named TimeD. It contains three member functions called Start, Stop, and Delay.The Start( ) function is used to start a measurement and Stop( ) is used to end a time measurement. The Stop( ) function returns a double floating-point value that is the amount of time in microseconds between Start( ) and Stop( ).

The Delay function causes a precision time delay based on the time-stamp counter. The parameter transferred to the Delay function is in milliseconds. This means that a Delay(1000) causes exactly 1000 ms of delay.

When TimeD is initialized in a program, it reads the microprocessor frequency in MHz from the Windows registry file using the RegQueryValueEx function after opening it with the RegOpenKeyEx function. The microprocessor clock frequency is returned in the MicroFrequency class variable.

EXAMPLE 18–2

```
#include "StdAfx.h"
#include ".\timed.h"

int MicroFrequency;                    //frequncy in MHz
_int64 Count;

TimeD::TimeD(void)
{
      HKEY hKey;
      DWORD dataSize;
      // Get the processor frequncy

      if (RegOpenKeyEx   (HKEY_LOCAL_MACHINE,
            "Hardware\\Description\\System\\CentralProcessor\\0",
            0, KEY_QUERY_VALUE, &hKey) == ERROR_SUCCESS )
      {
            RegQueryValueEx (hKey, _T("~MHz"), NULL, NULL,
                  (LPBYTE)&MicroFrequency, &dataSize);
            RegCloseKey (hKey);
      }
}

TimeD::~TimeD(void)
{
}

void TimeD::Start(void)
{
      _asm
      {
      rdtsc                           ;get and store TSC
      mov       dword ptr Count,eax
      mov       dword ptr Count+4,edx
      }
}

double TimeD::Stop(void)
{
      _asm
      {
      rdtsc
      sub       eax,dword ptr Count
      mov       dword ptr Count,eax
      sbb       edx,dword ptr Count+4
      mov       dword ptr Count+4,edx
      }
   return (double)Count/MicroFrequency;
}

void TimeD::Delay(__int64 milliseconds)
{
      milliseconds *= 1000;           //convert to microseconds
      milliseconds *= MicroFrequency; //convert to raw count
      _asm {
            mov   ebx, dword ptr milliseconds       ;64-bit delay in ms
            mov   ecx, dword ptr milliseconds+4
            rdtsc                                   ;get count
            add   ebx, eax
            adc   ecx, edx                          ;advance count by delay

      Delay_LOOP1:                                  ;wait for count to catch up

            rdtsc
            cmp   edx, ecx
            jb    Delay_LOOP1
            cmp   eax, ebx
            jb    Delay_LOOP1
      }
}
```

If an additional Delay is needed, it could be added to the class to cause delays in microseconds, but a restriction should be made so it is no less than about 2 or 3 microseconds, because of the time that it takes to add the time to the count from the time-stamp counter.

Example 18-3 shows a sample dialog application that used Delay( ) to wait for a second after clicking the button before changing the foreground color of an ActiveX Label. What does not appear in the example is that at the beginning of the dialog class an #include "TimeD.h" statement appears. The software itself is in the TODO: section of the OnInitDialog function.

EXAMPLE 18–3

```
void CRDTSCDlg::OnBnClickedButton1()
{
      TimeD timer;
      timer.Delay(1000);
      Label1.put_ForeColor(0xff0000);
}
```

The RDMSR and WRMSR instructions allow the model-specific registers to be read or written. The model-specific registers are unique to the Pentium and are used to trace, check performance, test, and check for machine errors. Both instructions use ECX to convey the register number to the microprocessor and use EDX:EAX for the 64-bit-wide read or write. Note that the register addresses are 0H–13H. See Table 18–5 for a list of the Pentium model-specific registers and their contents. As with the RDTSC instruction, these model-specific registers operate in the real or privilege level 0 of protected mode.

Never use an undefined value in ECX before using the RDMSR or WRMSR instructions. If ECX = 0 before the read or write machine-specific register instruction, the value returned, EDX:EAX, is the machine check exception address. (EDX:EAX is where all data reside when written or read from the model-specific registers.) If ECX = 1, the value is the machine check exception type; if ECX = 0EH, the test register 12 (TR_{12}) is accessed. Note that these are internal registers designed for in-house testing. The contents of these registers are proprietary to Intel and should not be used during normal programming.

TABLE 18–5   The Pentium model-specific registers.

Address (ECX)	Size	Function
00H	64 bits	Machine check exception address
01H	5 bits	Machine check exception type
02H	14 bits	TR_1 parity reversal test register
03H	—	—
04H	4 bits	TR_2 instruction cache end bits
05H	32 bits	TR_3 cache data
06H	32 bits	TR_4 cache tag
07H	15 bits	TR_4 cache control
08H	32 bits	TR_6 TLB command
09H	32 bits	TR_7 TLB data
0AH	—	—
0BH	32 bits	TR_9 BTB tag
0CH	32 bits	TR_{10} BTB target
0DH	12 bits	TR_{11} BTB control
0EH	10 bits	TR_{12} new feature control
0FH	—	—
10H	64 bits	Time-stamp counter (can be written)
11H	26 bits	Events counter selection and control
12H	40 bits	Events counter 0
13H	40 bits	Events counter 1

18–5 INTRODUCTION TO THE PENTIUM PRO MICROPROCESSOR

Before this or any other microprocessor can be used in a system, the function of each pin must be understood. This section of the chapter details the operation of each pin, along with the external memory system and I/O structures of the Pentium Pro microprocessor.

Figure 18–12 illustrates the pin-out of the Pentium Pro microprocessor, which is packaged in an immense 387-pin PGA (pin grid array). The Pentium Pro is available in two

FIGURE 18–12 The pin-out of the Pentium Pro microprocessor.

versions: One version contains a 256K level 2 cache; the other contains a 512K level 2 cache. The most notable difference in the pin-out of the Pentium Pro, when compared to the Pentium, is that there are provisions for a 36-bit address bus, which allows access to 64G bytes of memory. This is meant for future use because no system today contains anywhere near that amount of memory.

As with most recent versions of the Pentium microprocessor, the Pentium Pro requires a single +3.3 V or +2.7 V power supply for operation. The power supply current is a maximum of 9.9 A for the 150 MHz version of the Pentium Pro, which also has a maximum power dissipation of 26.7 W. A good heat sink with considerable airflow is required to keep the Pentium Pro cool. As with the Pentium, the Pentium Pro contains multiple V_{CC} and V_{SS} connections that must all be connected for proper operation. The Pentium Pro contains $V_{CC}P$ pins (primary V_{CC}) that connect to +3.1 V, $V_{CC}S$ (secondary V_{CC}) pins that connect to +3.3 V, and $V_{CC}5$ (standard V_{CC}) pins that connect to +5.0 V. There are some pins that are labeled N/C (no connection) and must not be connected.

Each Pentium Pro output pin is capable of providing an ample 48.0 mA of current at a logic 0 level. This represents a considerable increase in drive current, compared to the 2.0 mA available on earlier microprocessor output pins. Each input pin represents a small load, requiring only 15 μA of current. Because of the 48.0 mA of drive current available on each output, only an extremely large system requires bus buffers.

Internal Structure of the Pentium Pro

The Pentium Pro is structured differently than earlier microprocessors. Early microprocessors contained an execution unit and a bus interface unit with a small cache buffering the execution unit for the bus interface unit. This structure was modified in later microprocessors, but the modifications were just additional stages within the microprocessors. The Pentium architecture is also a modification, but more significant than earlier microprocessors. Figure 18–13 shows a block diagram of the internal structure of the Pentium Pro microprocessor.

The system buses, which communicate to the memory and I/O, connect to an internal level 2 cache that is often on the main board in most other microprocessor systems. The level 2 cache in the Pentium Pro is either 256K bytes or 512K bytes. The integration of the level 2 cache speeds processing and reduces the number of components in a system.
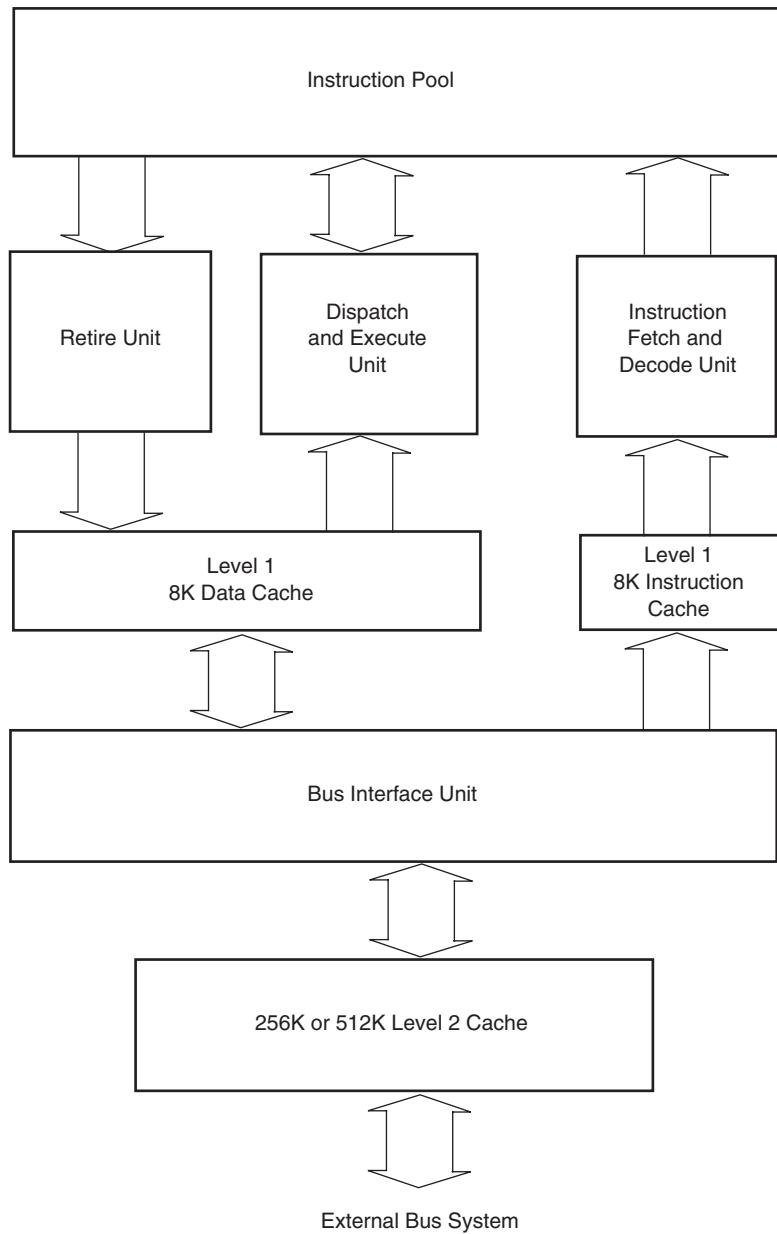
The bus interface unit (BIU) controls the access to the system buses through the level 2 cache, as it does in most other microprocessors. Again, the difference is that the level 2 cache is integrated. The BIU generates the memory address and control signals, and passes and fetches data or instructions to either a level 1 data cache or a level 1 instruction cache. Each cache is 8K bytes in size at present and may be made larger in future versions of the microprocessor. Earlier versions of the Intel microprocessor contained a unified cache that held both instructions and data. The implementation of separate caches improves performance because data-intensive programs no longer fill the cache with data.

The instruction cache is connected to the instruction fetch and decode unit (IFDU). Although not shown, the IFDU contains three separate instruction decoders that decode three instructions simultaneously. Once decoded, the outputs of the three decoders are passed to the instruction pool, where they remain until the dispatch and execution unit or retire unit obtains them. Also included within the IFDU is a branch prediction logic section that looks ahead in code sequences that contain conditional jump instructions. If a conditional jump is located, the branch prediction logic tries to determine the next instruction in the flow of a program.

Once decoded instructions are passed to the instruction pool, they are held for processing. The instruction pool is a content-addressable memory, but Intel never states its size in the literature.

The dispatch and execute unit (DEU) retrieves decoded instructions from the instruction pool when they are complete, and then executes them. The internal structure of the DEU is
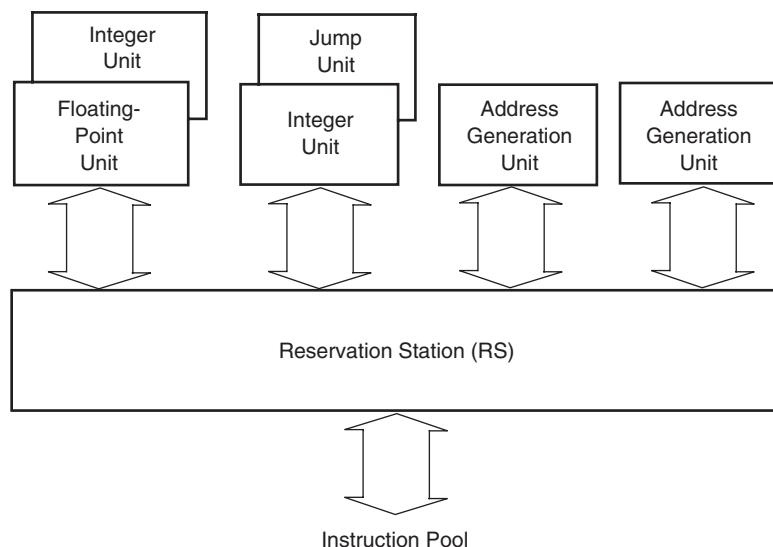
FIGURE 18–13 The internal structure of the Pentium Pro microprocessor.

illustrated in Figure 18–14. Notice that the DEU contains three instruction execution units: two for processing integer instructions and one for floating-point instructions. This means that the Pentium Pro can process two integer instructions and one floating-point instruction simultaneously. The Pentium also contains three execution units, but the architecture is different because the Pentium does not contain a jump execution unit or address generation units, as does the Pentium Pro. The reservation station (RS) can schedule up to five events for execution and process four simultaneously. Note that there are two station components connected to one of the address generation units that does not appear in the illustration of Figure 18–14.

The last internal structure of the Pentium Pro is the retire unit (RU). The RU checks the instruction pool and removes decoded instructions that have been executed. The RU can remove three decoded instructions per clock pulse.

FIGURE 18–14 The Pentium Pro dispatch and execution unit (DEU).

Pin Connections

The number of pins on the Pentium Pro has increased from the 237 pins on the Pentium to 387 pins on the Pentium Pro. Following is a description of each pin or grouping of pins:

$\overline{\text{A20M}}$	The **address A20 mask** is an input that is asserted in the real mode to signal the Pentium Pro to perform address wraparound, as in the 8086 microprocessor, for use of the HIMEM.SYS driver.
$\overline{\text{A35}}$–$\overline{\text{A3}}$	**Address bus** connections address any of the $8G \times 64$ memory locations found in the Pentium Pro memory system.
$\overline{\text{ADS}}$	The **address data strobe** becomes active whenever the Pentium Pro has issued a valid memory or I/O address.
$\overline{\text{AP1}}, \overline{\text{AP0}}$	**Address parity** provides even parity for the memory address on all Pentium Pro–initiated memory and I/O transfers. The $\overline{\text{AP0}}$ output provides parity for address connections A_{23}–A_3, and the $\overline{\text{AP1}}$ output provides parity for address connections A_{35}–A_{24}.
$\overline{\text{ASZ1}}, \overline{\text{ASZ0}}$	**Address size** inputs are driven to select the size of the memory access. Table 18–6 illustrates the size of the memory access for the binary bit patterns on these two inputs to the Pentium Pro.
BCLK	The **bus clock** input determines the operating frequency of the Pentium Pro microprocessor. For example, if BCLK is 66 MHz, various internal clocking speeds are selected by the logic levels applied to the pins in Table 18–7. A BCLK frequency of 66 MHz runs the system bus at 66 MHz.
$\overline{\text{BERR}}$	The **bus error** input/output either signals a bus error along or is asserted by an external device to cause a machine check interrupt or a non-maskable interrupt.

TABLE 18–6 Memory size dictated by the ASZ pins.

$\overline{ASZ1}$	$\overline{ASZ0}$	*Memory Size*
0	0	0–4G
0	1	4G–64G
1	X	Reserved

TABLE 18–7 The BCLK signal and its effect on the Pentium clock speed.

LINT1/NMI	LINT0/INTR	$\overline{IGNNE}$	$\overline{A20M}$	Ratio	BCLK = 50 MHz	BCLK = 66 MHz
0	0	0	0	2	100 MHz	133 MHz
0	0	0	1	4	200 MHz	266 MHz
0	0	1	0	3	150 MHz	200 MHz
0	0	1	1	5	250 MHz	333 MHz
0	1	0	0	5/2	125 MHz	166 MHz
0	1	0	1	9/2	225 MHz	300 MHz
0	1	1	0	7/2	175 MHz	233 MHz
0	1	1	1	11/2	275 MHz	366 MHz
1	1	1	1	2	100 MHz	133 MHz

$\overline{\text{BINIT}}$	**Bus initialization** is active on power-up to initialize the bus system.
$\overline{\text{BNR}}$	**Block next request** is used to halt the system in a multiple microprocessor system.
$\overline{\text{BP3}}, \overline{\text{BP2}}$	The **breakpoint status** outputs indicate the status of the Pentium Pro break-points.
$\overline{\text{BPM1}}, \overline{\text{BPM0}}$	The **break point monitor** outputs indicate the status of the breakpoints and programmable counters.
$\overline{\text{BPRI}}$	The **priority agent bus request** is an input that causes the microprocessor to cease bus requests.
$\overline{\text{BR3}}\text{–}\overline{\text{BR0}}$	The **bus request** inputs allow up to four Pentium Pro microprocessors to coexist on the same bus system.
$\overline{\text{BREQ3}}\text{–}\overline{\text{BREQ0}}$	**Bus request signals** are used for multiple microprocessors on the same system bus.
$\overline{\text{D63}}\text{–}\overline{\text{D0}}$	**Data bus** connections transfer byte, word, doubleword, and quadword data between the microprocessor and its memory and I/O system.
$\overline{\text{DBSY}}$	**Data bus busy** is asserted to indicate that the data bus is busy transferring data.
$\overline{\text{DEFER}}$	The **defer** input is asserted during the snoop phase to indicate that the transaction cannot be guaranteed in-order completion.
$\overline{\text{DEN}}$	The **defer enable** signal is driven to the bus on the second phase of a request phase.
$\overline{\text{DEP7}}\text{–}\overline{\text{DEP0}}$	**Data bus ECC protection** signals provide error-correction codes for correcting a single-bit error and detecting a double-bit error.
$\overline{\text{FERR}}$	The **floating-point error**, comparable to the ERROR line in the 80386, shows that the internal coprocessor has erred.
$\overline{\text{FLUSH}}$	The **flush** cache input causes the cache to flush all write-back lines and invalidate its internal caches. If the $\overline{\text{FLUSH}}$ input is a logic 0 during a reset operation, the Pentium Pro enters its test mode.
$\overline{\text{FRCERR}}$	**Functional redundancy check error** is used if two Pentium Pro microprocessors are configured in a pair.
$\overline{\text{HIT}}$	**Hit** shows that the internal cache contains valid data in the inquire mode.
$\overline{\text{HITM}}$	**Hit modified** shows that the inquire cycle found a modified cache line. This output is used to inhibit other master units from accessing data until the cache line is written to memory.

TABLE 18–8 The $\overline{\text{LEN}}$ bits and data size.

LEN1	LEN0	Data Transfer Size
0	0	0–8 bytes
0	1	16 bytes
1	0	32 bytes
1	1	Reserved

$\overline{\text{IERR}}$	**Internal error** output shows that the Pentium Pro has detected an internal parity error or functional redundancy error.
$\overline{\text{IGNNE}}$	The **ignore numeric error** input causes the Pentium Pro to ignore a numeric coprocessor error.
INIT	The **initialization** input performs a reset without initializing the caches, write-back buffers, and floating-point registers. This input may not be used to reset the microprocessor in lieu of RESET after power-up.
INTR	The **interrupt request** is used by external circuitry to request an interrupt.
$\overline{\text{LEN1}}, \overline{\text{LEN0}}$	Length signals (bit 0 and 1) indicate the size of the data transfer, as illustrated in Table 18–8.
$\overline{\text{LINT1}}, \overline{\text{LINT0}}$	The **local interrupt** inputs function as NMI and INTR, and also set the clock divider frequency on reset.
$\overline{\text{LOCK}}$	$\overline{\text{LOCK}}$ becomes a logic 0 whenever an instruction is prefixed with the LOCK: prefix. This is most often used during DMA accesses.
NMI	The non-maskable interrupt requests a non-maskable interrupt, as it did on the earlier versions of the microprocessor.
PICCLK	The **clock** signal input is used for synchronous data transfers.
PICD	The **processor interface serial data** is used to transfer bidirectional serial messages between Pentium Pro microprocessors.
PWRGOOD	**Power good** is an input that is placed at a logic 1 level when the power supply and clock have stabilized.
$\overline{\text{REQ4}}$–$\overline{\text{REQ0}}$	**Request signals** (bits 0–4) define the type of data-transfer operation, as illustrated in Tables 18–9 and 18–10.
$\overline{\text{RESET}}$	**Reset** initializes the Pentium Pro, causing it to begin executing software at memory location FFFFFFF0H. The Pentium Pro is reset to the real mode and

TABLE 18–9 Function of the request inputs on the first clock pulse.

REQ4	REQ3	REQ2	REQ1	REQ0	Function
0	0	0	0	0	Deferred reply
0	0	0	0	1	Reserved
0	1	0	0	0	Case 1*
0	1	0	0	1	Case 2*
1	0	0	0	0	I/O read
1	0	0	0	1	I/O write
X	X	0	1	0	Memory read
X	X	0	1	1	Memory write
X	X	1	0	0	Memory code read
X	X	1	1	0	Memory data read
X	X	1	X	1	Memory write

*See Table 18–10 for the second clock pulse for Case 1 and Case 2.

TABLE 18–10 Function of the request inputs for Case 1 and Case 2.

Case	$\overline{REQ4}$	$\overline{REQ3}$	$\overline{REQ2}$	$\overline{REQ1}$	$\overline{REQ0}$	Function
1	X	X	X	0	0	Interrupt acknowledge
1	X	X	X	0	1	Special transaction
1	X	X	X	1	X	Reserved
2	X	X	X	0	0	Branch trace message
2	X	X	X	0	1	Reserved
2	X	X	X	1	X	Reserved

the leftmost 12 address connections remain logic 1s (FFFH) until a far jump or far call is executed. This allows compatibility with earlier microprocessors.

$\overline{RP}$ **Request parity** provides a means of requesting that the Pentium Pro checks parity.

$\overline{RS2}$–$\overline{RS0}$ The **response status** inputs cause the Pentium Pro to perform the functions listed in Table 18–11.

$\overline{RSP}$ The **response parity** input applies a parity error signal from an external parity checker.

$\overline{SMI}$ The **system management interrupt** input causes the Pentium Pro to enter the system management mode of operation.

$\overline{SMMEM}$ The **system memory-management mode** signal becomes a logic 0 whenever the Pentium Pro is executing in the system memory-management mode interrupt and address space.

$\overline{SPCLK}$ The **split lock** signal is placed at a logic 0 level to indicate that the transfer will contain four locked transactions.

$\overline{SPCLK}$ **Stop clock** causes the Pentium Pro to enter the power-down state when placed at a logic 0 level.

TCK   The **testability clock** input selects the clocking function in accordance with the IEEE 1149.1 Boundary Scan interface.

TDI   The **test data input** is used to test data clocked into the Pentium Pro with the TCK signal.

TDO   The **test data output** is used to gather test data and instructions shifted out of the Pentium with TCK.

TMS   The **test mode select** input controls the operation of the Pentium Pro in test mode.

$\overline{TRDY}$ The **target ready** input is asserted when the target is ready for a data transfer operation.

TABLE 18–11 Operation of the response status inputs.

$\overline{RS2}$	$\overline{RS1}$	$\overline{RS0}$	Function	$\overline{HITM}$	$\overline{DEFER}$
0	0	0	Idle state	X	X
0	0	1	Retry	0	1
0	1	0	Defer	0	1
0	1	1	Reserved	0	1
1	0	0	Hard failure	X	X
1	0	1	Normal, no data	0	0
1	1	0	Implicit write-back	1	X
1	1	1	Normal, with data	0	0

The Memory System

The memory system for the Pentium Pro microprocessor is 4G bytes in size, just as in the 80386DX–Pentium microprocessors, but access to an area between 4G and 64G is made possible by additional address signals A_{32}–A_{35}. The Pentium Pro uses a 64-bit data bus to address memory organized in eight banks that each contain 8G bytes of data. Note that the additional memory is enabled with bit position 5 of CR_4 and is accessible only when 2M paging is enabled. Note also that 2M paging is new to the Pentium Pro to allow memory above 4G to be accessed. More information is presented on Pentium Pro paging later in this chapter. Refer to Figure 18–15 for the organization of the Pentium Pro physical memory system.

The Pentium Pro memory system is divided into eight banks where each bank stores a byte-wide data with a parity bit. Note that most Pentium and Pentium Pro microprocessor-based systems forgo the use of the parity bit. The Pentium Pro, like the 80486 and Pentium, employs internal parity generation and checking logic for the memory system data bus information. The 64-bit-wide memory is important to double-precision floating-point data. Recall that a double-precision floating-point number is 64 bits wide. As with earlier Intel microprocessors, the memory system is numbered in bytes from byte 000000000H to byte FFFFFFFFFH. This nine-digit hexadecimal address is employed in a system that addresses 64G of memory.

Memory selection is accomplished with the bank enable signals ($\overline{BE7}$–$\overline{BE0}$). In the Pentium Pro microprocessor, the bank enable signals are presented on the address bus (A_{15}–A_8) during the second clock cycle of a memory or I/O access. These must be extracted from the address bus to access memory banks. The separate memory banks allow the Pentium Pro to access any single byte, word, doubleword, or quadword with one memory transfer cycle. As with earlier memory selection logic, we often generate eight separate write strobes for writing to the memory system. Note that the memory write information is provided on the request lines from the microprocessor during the second clock phase of a memory or I/O access.

A new feature added to the Pentium and Pentium Pro is the capability to check and generate parity for the address bus during certain operations. The $\overline{AP}$ pin (Pentium) or pins (Pentium Pro) provide the system with parity information, and the $\overline{APCHK}$ (Pentium) or $\overline{AP}$ pins (Pentium Pro) indicate a bad parity check for the address bus. The Pentium Pro takes no action when an address-parity error is detected. The error must be assessed by the system, and the system must take appropriate action (an interrupt) if so desired.

New to the Pentium Pro is a built-in error-correction circuit (ECC) that allows the correction of a one-bit error and the detection of a two-bit error. To accomplish the detection and
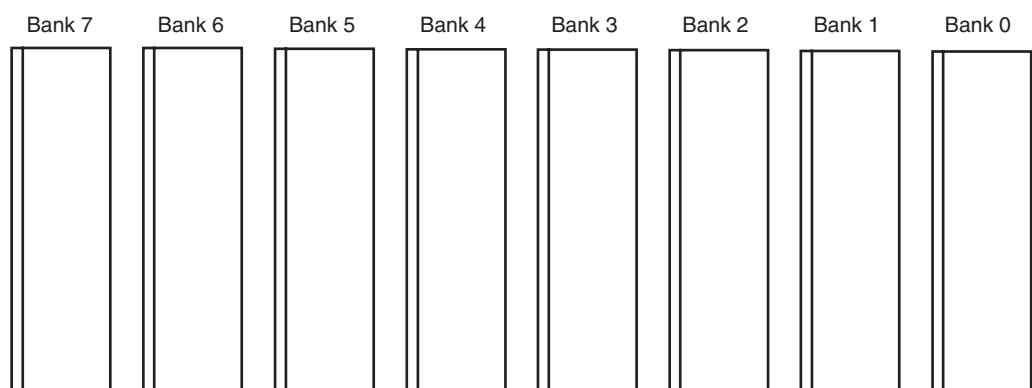
FIGURE 18–15 The eight memory banks in the Pentium Pro system. Note that each bank is 8 bits wide and 8G long if 36-bit addressing is enabled.

correction of errors, the memory system must have room for an extra 8-bit number that is stored with each 64-bit number. The extra 8 bits are used to store an error-correction code that allows the Pentium Pro to automatically correct any single-bit error. A $1M \times 64$ is a 64M SDRAM without ECC, and a $1M \times 72$ is an SDRAM with EEC support. The ECC code is much more reliable than the old parity scheme, which is rarely used in modern systems. The only drawback of the ECC scheme is the additional cost of SDRAM that is 72 bits wide.

Input/Output System

The input/output system of the Pentium Pro is completely compatible with earlier Intel microprocessors. The I/O port number appears on address lines A_{15}–A_3 with the bank enable signals used to select the actual memory banks used for the I/O transfer.
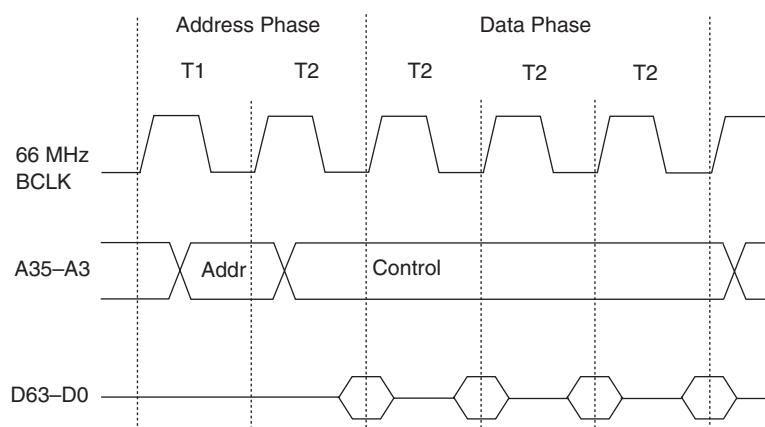
System Timing

As with any microprocessor, the system timing signals must be understood in order to interface the microprocessor. This portion of the text details the operation of the Pentium Pro through its timing diagrams and shows how to determine memory access times.

The basic Pentium Pro memory cycle consists of two sections: the address phase and the data phase. During the address phase, the Pentium Pro sends the address (T_1) to the memory and I/O system, and also the control signals (T_2). The control signals include the ATTR lines (A_{31}–A_{24}), the DID lines (A_{23}–A_{16}), the bank enable signals (A_{15}–A_8), and the EXF lines (A_7–A_3). See Figure 18–16 for the basic timing cycle. The type of memory cycle appears on the request pins. During the data phase, four 64-bit-wide numbers are fetched or written to the memory. This operation is most common because data from the main memory are transferred between the internal 256K or 512K write-back cache and the memory system. Operations that write a byte, word, or doubleword, such as I/O transfers, use the bank selection signals and have only one clock in the data transfer phase. Notice from the timing diagram that the 66 MHz Pentium Pro is capable of 33 million memory transfers per second. (This assumes that the memory can operate at that speed.)

The setup time before the clock is given as 5.0 ns and the hold time after the clock is given as 1.5 ns. This means that the data window around the clock is 6.5 ns. The address appears on the 8.0 ns maximum after the start of T_1. This means that the Pentium Pro microprocessor operating at 66 MHz allows 30 ns (two clocking periods), minus the address delay time of 8.0 ns and also

FIGURE 18–16 The basic Pentium Pro timing.

minus the data setup time of 5.0 ns. Memory access time without any wait states is $30 - 8.0 - 5.0$, or 17.0 ns. This is enough time to allow access to an SRAM, but not to any DRAM without inserting wait states into the timing.

Wait states are inserted into the timing by controlling the $\overline{\text{TRDY}}$ input to the Pentium Pro. The $\overline{\text{TRDY}}$ signal must become a logic 0 by the end of T_2; otherwise, additional T_2 states are inserted into the timing. Note that 60 ns DRAM requires the insertion of four wait states of 15 ns (one clocking period) each to lengthen the access time to 77 ns. This is enough time for the DRAM and any decoder in the system to function. Because many EPROM or flash memory devices require an access time of 100 ns, EPROM or flash requires the addition of seven wait states to lengthen the access time to 122 ns.

18–6   SPECIAL PENTIUM PRO FEATURES

The Pentium Pro is essentially the same microprocessor as the 80386, 80486, and Pentium, except that some additional features and changes to the control register set have occurred. This section highlights the differences between the 80386 control register structure and the Pentium Pro control register.

Control Register 4

Figure 18–17 shows control register 4 of the Pentium Pro microprocessor. Notice that CR_4 has two new control bits that are added to the control register array.

This section of the text explains only the two new Pentium Pro components in the control register 4. (Refer to Figure 18–8 for a description and illustration of the Pentium control registers.) Following is a description of the Pentium CR_4 bits and the new Pentium Pro control bits in control register CRM_4:

VME	**Virtual mode extension** enables support for the virtual interrupt flag in protected mode. If VME = 0, virtual interrupt support is disabled.
PVI	**Protected mode virtual interrupt** enables support for the virtual interrupt flag in protected mode.
TSD	**Time stamp disable** controls the RDTSC instruction.
DE	**Debugging extension** enables I/O breakpoint debugging extensions when set.
PSE	**Page size extension** enables 4M-byte memory pages when set in the Pentium, or 2M-byte pages when set in the Pentium Pro whenever PSE is also set.
PAE	**Page address extension** enables address lines A_{35}–A_{32} whenever a special new addressing mode, controlled by PGE, is enabled for the Pentium Pro.
MCE	**Machine check** enable enables the machine checking interrupt.
PGE	**Page extension** controls the new, larger 64G addressing mode whenever it is set along with PAE and PSE.

FIGURE 18–17 The new control register 4 (CR_4) in the Pentium Pro microprocessor.

31								7	6	5	4	3	2	1	0
								PGE	MCE	PAE	PSE	DE	TSD	PVI	VME

18–7 SUMMARY

1. The Pentium microprocessor is almost identical to the earlier 80386 and 80486 micro-processors. The main difference is that the Pentium has been modified internally to contain a dual cache (instruction and data) and a dual integer unit. The Pentium also operates at a higher clock speed of 66 MHz.

2. The 66 MHz Pentium requires 3.3 A of current, and the 60 MHz version requires 2.91 A. The power supply must be a +5.0 V supply with a regulation of ±5%. Newer versions of the Pentium require a 3.3 V or 2.7 V power supply.

3. The data bus on the Pentium is 64 bits wide and contains eight byte-wide memory banks selected with bank enable signals ($\overline{BE7}$–$\overline{BE0}$).

4. Memory access time, without wait states, is only about 18 ns in the 66 MHz Pentium. In many cases, this short access time requires wait states that are introduced by controlling the $\overline{BRDY}$ input to the Pentium.

5. The superscalar structure of the Pentium contains three independent processing units: a floating-point processor and two integer processing units labeled U and V by Intel.

6. The cache structure of the Pentium is modified to include two caches. One 8K × 8 cache is designed as an instruction cache; the other 8K × 8 cache is a data cache. The data cache can be operated as either a write-through or a write-back cache.

7. A new mode of operation called the system memory-management (SMM) mode has been added to the Pentium. The SMM mode is accessed via the system memory-management interrupt applied to the $\overline{SMI}$ input pin. In response to $\overline{SMI}$, the Pentium begins executing software at memory location 38000H.

8. New instructions include the CMPXCHG8B, RSM, RDMSR, WRMSR, and CPUID. The CMPXCHG8B instruction is similar to the 80486 CMPXCHG instruction. The RSM instruction returns from the system memory-management interrupt. The RDMSR and WRMSR instructions read or write to the machine-specific registers. The CPUID instruction reads the CPU identification code from the Pentium.

9. The built-in self-test (BIST) allows the Pentium to be tested when power is first applied to the system. A normal power-up reset activates the RESET input to the Pentium. A BIST power-up reset activates INIT and then deactivates the RESET pin. EAX is equal to a 00000000H in the BIST passes.

10. A new proprietary Intel modification to the paging unit allows 4M-byte memory pages instead of the 4K-byte pages. This is accomplished by using the page directory to address 1024 page tables that each contains 4M of memory.

11. The Pentium Pro is an enhanced version of the Pentium microprocessor that contains not only the level 1 caches found inside the Pentium, but also the level 2 cache of 256K or 512K found on most main boards.

12. The Pentium Pro operates by using the same 66 MHz bus speed as the Pentium and the 80486. It uses an internal clock generator to multiply the bus speed by various factors to obtain higher internal execution speeds.

13. The only significant software difference between the Pentium Pro and earlier microproces-sors is the addition of the FCMOV and CMOV instructions.

14. The only hardware difference between the Pentium Pro and earlier microprocessors is the addition of 2M paging and four extra address lines that allow access to a memory address space of 64G bytes.

15. Error correction code has been added to the Pentium Pro, which corrects any single-bit error and detects any two-bit error.

18–8   QUESTIONS AND PROBLEMS

1. How much memory is accessible to the Pentium microprocessor?
2. How much memory is accessible to the Pentium Pro microprocessor?
3. The memory data bus width is _____ in the Pentium.
4. What is the purpose of the DP_0–DP_7 pins on the Pentium?
5. If the Pentium operates at 66 MHz, what frequency clock signal is applied to the CLK pin?
6. What is the purpose of the $\overline{BRDY}$ pin on the Pentium?
7. What is the purpose of the AP pin on the Pentium?
8. How much memory access time is allowed by the Pentium, without wait states, when it is operated at 66 MHz?
9. What Pentium pin is used to insert wait states into the timing?
10. A wait state is an extra _____ clocking period.
11. Explain how two integer units allow the Pentium to execute two nondependent instructions simultaneously.
12. How many caches are found in the Pentium and what are their sizes?
13. How wide is the Pentium memory data sample window for a memory read operation?
14. Can the Pentium execute three instructions simultaneously?
15. What is the purpose of the $\overline{SMI}$ pin?
16. What is the system memory-management mode of operation for the Pentium?
17. How is the system memory-management mode exited?
18. Where does the Pentium begin to execute software for an $\overline{SMI}$ interrupt input?
19. How can the system memory-management unit dump address be modified?
20. Explain the operation of the CMPXCHG8B instruction.
21. What information is returned in register EAX after the CPUID instruction executes with an initial value of 0 in EAX?
22. What new flag bits are added to the Pentium microprocessor?
23. What new control register is added to the Pentium microprocessor?
24. Describe how the Pentium accesses 4M pages.
25. Explain how the time-stamp counter functions and how it can be used to time events.
26. Contrast the Pentium with the Pentium Pro microprocessor.
27. Where are the bank enable signals found in the Pentium Pro microprocessor?
28. How many address lines are found in the Pentium Pro system?
29. What changes have been made to CR_4 in the Pentium Pro and for what purpose?
30. Compare access times in the Pentium system with the Pentium Pro system.
31. What is ECC?
32. What type of SDRAM must be purchased to use ECC?

CHAPTER 19

The Pentium II, Pentium III, Pentium 4, and Core2 Microprocessors

INTRODUCTION

The Pentium II, Pentium III, Pentium 4, and Core2 microprocessors may well signal the end to the evolution of the 32-bit architecture with the advent of the Itanium[1] and Itanium II microprocessors from Intel. The Itanium is a 64-bit architecture microprocessor. The Pentium II, Pentium III, Pentium 4, and Core2 architectures are extensions of the Pentium Pro architecture, with some differences. The most notable difference is that the internal cache from the Pentium Pro architecture has been moved out of the microprocessor in the Pentium II. Another major change is that the Pentium II is not available in integrated circuit form. Instead, the Pentium II is found on a small plug-in circuit board called a cartridge along with a separate level 2 cache chip. Various versions of the Pentium II are available. The Celeron[2] is a version of the Pentium II that does not contain the level 2 cache on the Pentium II circuit board. The Xeon[3] is an enhanced version of the Pentium II that contains up to a 2M-byte cache on the circuit board.

Similar to the Pentium II, early Pentium III microprocessors were packaged in a cartridge instead of an integrated circuit. More recent versions, such as the Coppermine, are again packaged in an integrated circuit (370 pins). The Pentium III Coppermine, like the Pentium Pro, contains an internal cache. The Pentium 4 is packaged in a larger integrated circuit, with 423 or 478 pins and most recently the Pentium 4 and Core 2 are manufactured in a 775-pin LGA package (leadless grid array). The Pentium 4 also uses physically smaller transistors, which makes it much smaller and faster than the Pentium III. To date Intel has released versions of the Pentium 4 and Core2 that operate at frequencies over 3 GHz with a limit of possibly 10 GHz at some future date. Also available to the Pentium 4 and Core2 are the extreme model with a 2M-byte cache and the extreme edition model with a 4M-byte cache. These versions are available in the 65 nm (0.065 micron) form as compared to earlier P_4 microprocessors that use the 0.13 micron form. The latest versions are the Core2 Duo and Core2 Quad versions that use 45 nm technology and either two or four cores.

CHAPTER OBJECTIVES

Upon completion of this chapter, you will be able to:

1. Detail the differences between the Pentium II, Pentium III, Pentium 4, and Core2 and prior Intel microprocessors.

---

[1]Itanium is a registered trademark of Intel Corporation.

[2]Celeron is a registered trademark of Intel Corporation.

[3]Xeon is a registered trademark of Intel Corporation.

2. Explain how the architectures of the Pentium II, Pentium III, Pentium 4, and Core2 improve system speed.
3. Explain how the basic architecture of the computer system has changed by using the Pentium II, Pentium III, Pentium 4, and Core2 microprocessors.
4. Detail the changes to the CPUID instruction and model-specific registers.
5. Describe the operation of the SYSENTER and SYSEXIT instructions.
6. Describe the operation of the FXSAVE and FXRSTOR instructions.

19–1   INTRODUCTION TO THE PENTIUM II MICROPROCESSOR

Before the Pentium II or any other microprocessor can be used in a system, the function of each pin must be understood. This section of the chapter details the operation of each pin, along with the external memory system and I/O structures of the Pentium II microprocessor.

Figure 19–1 illustrates the basic outline of the Pentium II microprocessor slot 1 connector and the signals used to interface to the chip set. Figure 19–2 shows a simplified diagram of the components on the cartridge, and the placement of the Pentium II cartridge and bus components in the typical Pentium II system. There are 242 pins on the slot 1 connector for the microprocessor. (These connections are a reduction in the number of pins found on the Pentium and the Pentium II microprocessors.) The Pentium II is packaged on a printed circuit board instead of the integrated circuits of the past Intel microprocessors. The level 1 cache is 32K bytes as it was in the Pentium Pro, but the level 2 cache is no longer inside the integrated circuit. Intel changed the architecture so that a level 2 cache could be placed very close to the microprocessor. This change makes the microprocessor less expensive and still allows the level 2 cache to operate efficiently. The Pentium level 2 cache operates at one half the microprocessor clock frequency, instead of the 66 MHz of the Pentium microprocessor. A 400 MHz Pentium II has a cache speed of 200 MHz. The Pentium II is available in three versions. The first is the full-blown Pentium II, which is the Pentium II for the slot 1 connector. The second is the Celeron, which is like the Pentium II, except that the slot 1 circuit board does not contain a level 2 cache; the level 2 cache in the Celeron system is located on the main board and operates at 66 MHz. The most recent version is the Xeon, which, because it uses a level 2 cache of 512K, 1M, or 2M, represents a significant speed improvement over the Pentium II. The Xeon's level 2 cache operates at the clock frequency of the microprocessor. A 400 MHz Xeon has a level 2 cache speed of 400 MHz, which is twice the speed of the regular Pentium II.

The early versions of the Pentium II require a 5.0 V, 3.3 V, and variable voltage power supply for operation. The main variable power supply voltages vary from 3.5 V to as low as 1.8 V at the microprocessor. This is known as the core microprocessor voltage. The power-supply current averages 14.2 A to 8.4 A, depending on the operating frequency and voltage of the Pentium II. Because these currents are significant, so is the power dissipation of these microprocessors. At present, a good heat sink with considerable airflow is required to keep the Pentium II cool. Luckily, the heat sink and fan are built into the Pentium II cartridge. The latest versions of the Pentium II have been improved to reduce the power dissipation.

Each Pentium II cartridge output pin is capable of providing at least 36 mA of current at a logic 0 level on the signal connections. Some of the output control signals provide only 14 mA of current. Another change to the Pentium II is that the outputs are open-drain and require an external pull-up resister for proper operation.

The function of each Pentium II group of pins follows:

$\overline{A20}$ **Address A$_{20}$ mask** is an input that is asserted in the real mode to signal the Pentium II to perform address wraparound, as in the 8086 microprocessor, for use of the HIMEM.SYS driver.

FIGURE 19–1   The pin-out of the slot 1 connector showing the connections to the system.
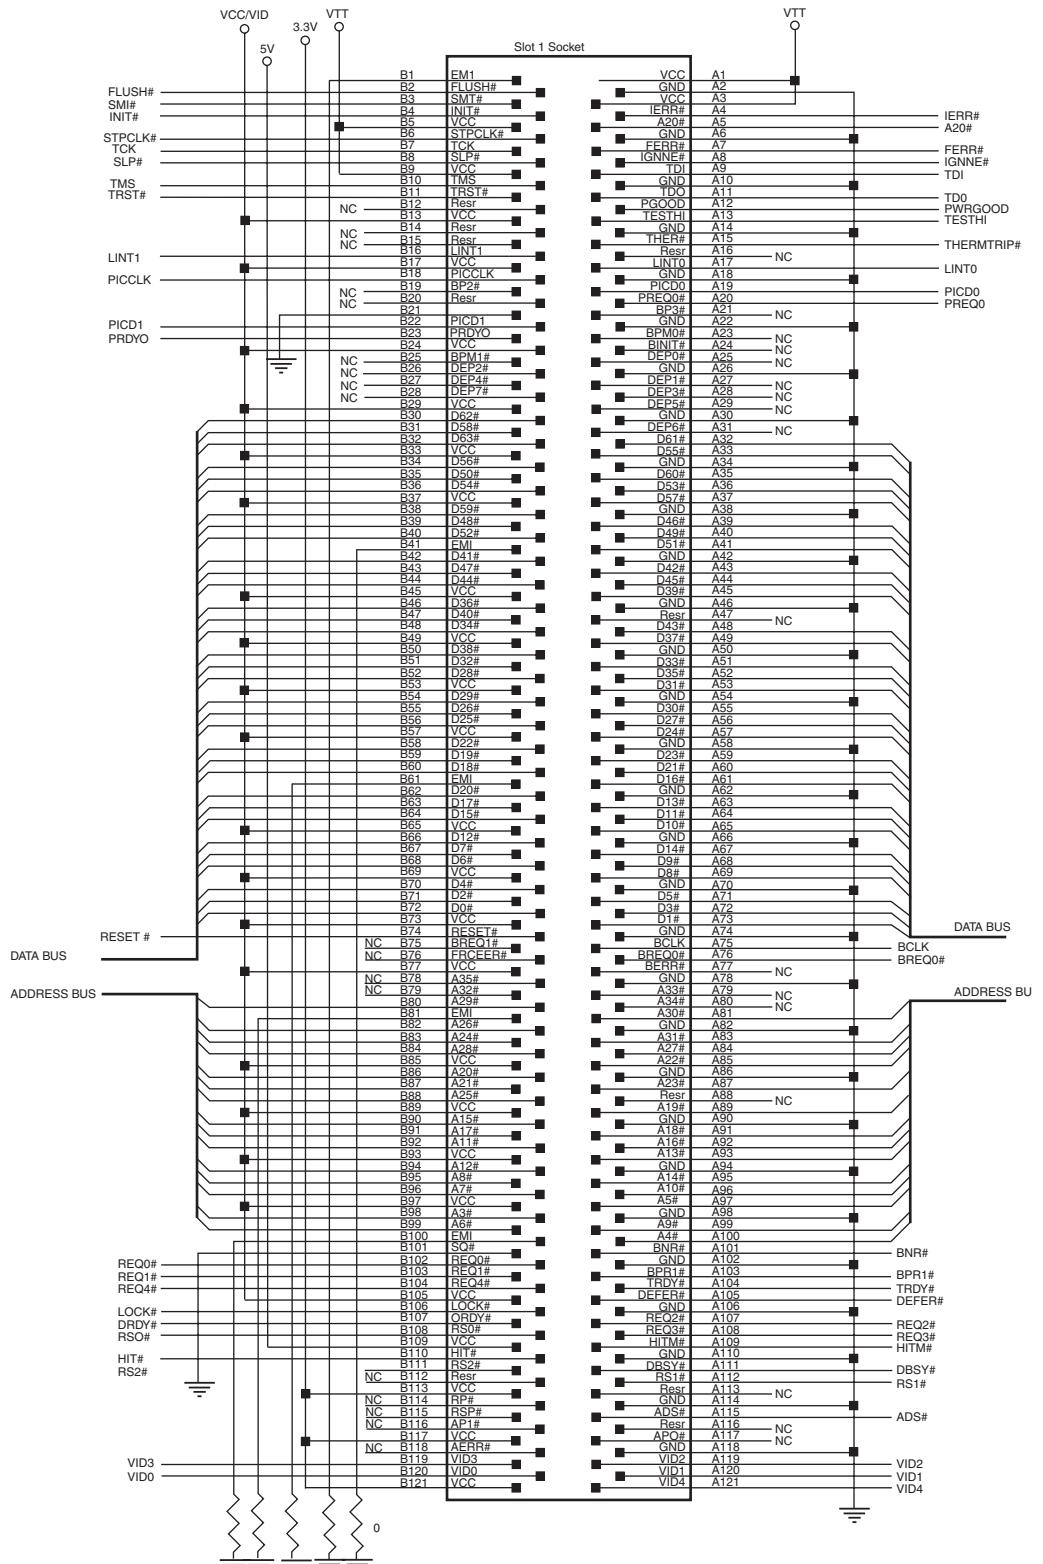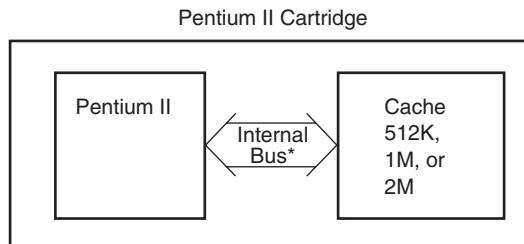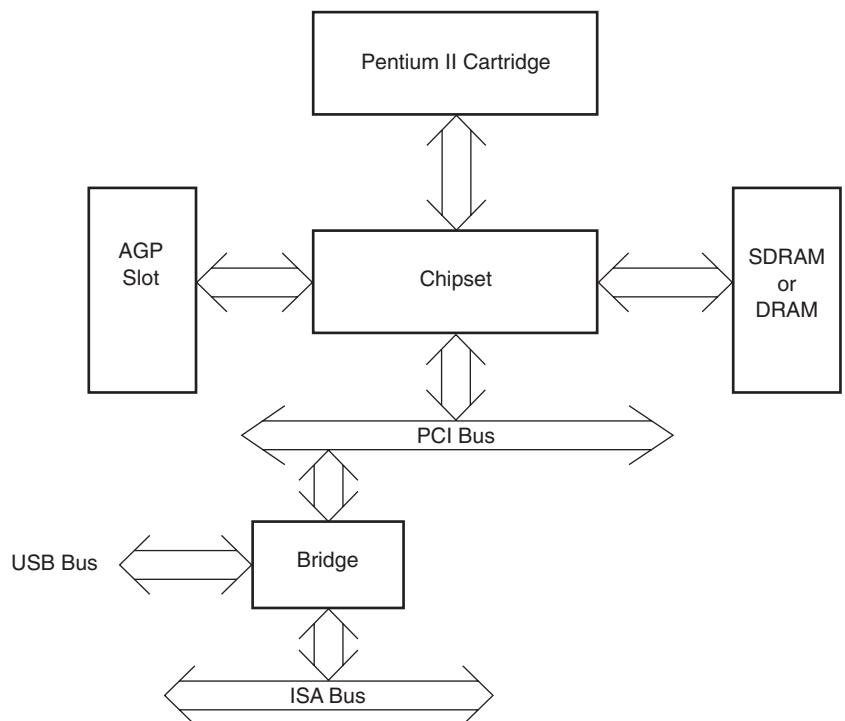
FIGURE 19–2   The structure of the Pentium II cartridge and the structure of the Pentium II system.

Pentium II Cartridge

* The bus speed is 1/2 Pentium speed or the same as the Pentium speed in the Xeon.

$\overline{A35}$–$\overline{A3}$	**Address buses**, which are active low connections, address any of the memory locations found in the Pentium II memory system. Note that A_0, A_1, and A_2 are encoded in the bus enable ($\overline{BE7}$–$\overline{BE0}$), which are generated by the chip set, to select any or all of the eight bytes in a 64-bit-wide memory location.
$\overline{ADS}$	**Address data strobe** is an input that is activated to indicate to the Pentium II that the system is ready to perform a memory or I/O operation. This signal causes the microprocessor to provide the address to the system.
$\overline{AERR}$	**Address error** is an input used to cause the Pentium II to check for an address parity error if it is activated.
$\overline{AP1}$, $\overline{AP0}$	**Address parity** inputs indicate an address parity error.
$\overline{BCLK}$	**Bus clock** is an input that sets the bus clock frequency. This is either 66 MHz or 100 MHz in the Pentium II.

$\overline{\text{BERR}}$	**Bus error** is asserted to indicate that an error has occurred on the bus system.
$\overline{\text{BINIT}}$	**Bus initialization** is a logic 0 during system reset or initialization. It is an input to indicate that a bus error has occurred and the system needs to be reinitialized.
$\overline{\text{BNR}}$	**Bus not ready** is an input used to insert wait states into the timing for the Pentium II. Placing a logic 0 on this pin causes the Pentium II to enter stall states or wait states.
$\overline{\text{BP3}}$, $\overline{\text{BP2}}$, $\text{PM1}/\overline{\text{BP1}}$, and $\text{PM0}/\overline{\text{BP0}}$	The **breakpoint** pins $\overline{\text{BP3}}$–$\overline{\text{BP0}}$ indicate a breakpoint match when the debug registers are programmed to monitor for matches. The performance monitoring pins $\overline{\text{PM1}}$ and $\overline{\text{PM0}}$ indicate the settings of the performance monitoring bits in the debug mode control register.
$\overline{\text{BPRI}}$	The **bus priority request input** is used to request the system bus from the Pentium II.
$\overline{\text{BR1}}$ and $\overline{\text{BR0}}$	**Bus requests** indicate that the Pentium II has generated a bus request. During initialization, the $\overline{\text{BR0}}$ pin must be activated.
BSEL	**Bus select** is currently not used by the Pentium II and must be connected to ground for proper operation.
$\overline{\text{D63}}$–$\overline{\text{D0}}$	**Data bus** connections transfer byte, word, doubleword, and quadword data between the microprocessor and its memory and I/O system.
$\overline{\text{DEFER}}$	The **defer** signal indicates that the external system cannot complete the bus cycle.
$\overline{\text{EP7}}$–$\overline{\text{EP0}}$	**Data EEC** pins are used in the error-correction scheme of the Pentium II and normally connect to an extra 8-bit memory section. This means that ECC memory modules are 72 bits wide instead of 64 bits wide.
$\overline{\text{DRDY}}$	**Data ready** is activated to indicate that the system is presenting valid data to the Pentium II.
EMI	**Electromagnetic interference** must be grounded to prevent the Pentium II from generating or receiving noise.
$\overline{\text{FERR}}$	**Floating-point error**, comparable to the ERROR line in the 80386, shows that the internal coprocessor has erred.
$\overline{\text{FLUSH}}$	The **flush cache** input causes the cache to flush all write-back lines and invalidate its internal caches. If the $\overline{\text{FLUSH}}$ input is a logic 0 during a reset operation, the Pentium enters its test mode.
FRCERR	**Functional redundancy check** is sampled during a reset to configure the Pentium II in the master (1) or checker (0) mode.
$\overline{\text{HIT}}$	**Hit** shows that the internal cache contains valid data in the inquire mode.
$\overline{\text{HITM}}$	**Hit modified** shows that the inquire cycle found a modified cache line. This output is used to inhibit other master units from accessing data until the cache line is written to memory.
$\overline{\text{IERR}}$	The **internal error** output shows that the Pentium II has detected an internal error or functional redundancy error.
$\overline{\text{IGNNE}}$	The **ignore numeric error** input causes the Pentium II to ignore a numeric coprocessor error.

$\overline{\text{INIT}}$	The **initialization** input performs a reset without initializing the caches, write-back buffers, and floating-point registers. This input may not be used to reset the microprocessor in lieu of RESET after power-up.
INTR	**Interrupt request** is used by external circuitry to request an interrupt.
LINT$_1$, LINT$_0$	**Local APIC interrupt** signals must connect the appropriate pins of all APIC bus agents. When the APIC is disabled, the LINT$_0$ signal becomes INTR, a maskable interrupt request signal; LINT$_1$ becomes NMI, a non-maskable interrupt.
$\overline{\text{LOCK}}$	$\overline{\text{LOCK}}$ becomes a logic 0 whenever an instruction is prefixed with the LOCK: prefix. This is most often used during DMA accesses.
NMI	**Non-maskable interrupt** requests a non-maskable interrupt as it did on the earlier versions of the microprocessor.
PICCLK	This **clock** signal must be ¼ the frequency of $\overline{\text{BCLK}}$.
PICD$_1$, PICD$_0$	Used for serial messages between the Pentium II and APIC.
PRDY	The **probe ready** output indicates that the probe mode has been entered for debugging.
$\overline{\text{PREQ}}$	The **probe request** is used to request debugging.
PWRGOOD	The **power good** input that indicates that the system power supply is operational.
$\overline{\text{REQ4}} - \overline{\text{REQ0}}$	**Request signals** communicate commands between bus controllers and the Pentium II.
$\overline{\text{RESET}}$	**Reset** initializes the Pentium II, causing it to begin executing software at memory location FFFFFFF0H or 000FFFF0H. The A_{35}–A_{32} address bits are set as logic 0s during the reset operation. The Pentium II is reset to the real mode and the leftmost 12 address connections remain logic 1s (FFFH) until a far jump or far call is executed. This allows compatibility with earlier microprocessors. See Table 19–1 for the state of the Pentium II after a hardware reset.

TABLE 19–1   State of the Pentium II after a reset.

Register	Reset	Reset + BIST
EAX	0	0 (if test passes)
EDX	0500XXXXH	0500XXXXH
EBX, ECX, ESP, EBP, ESI, and EDI	0	0
EFLAGS	2	2
EIP	0000FFF0H	0000FFF0H
CS	F000H	F000H
DS, ES, FS, GS, and SS	0	0
GDTR and TSS	0	0
CR$_0$	60000010H	60000010H
CR$_2$, CR$_3$, and CR$_4$	0	0
DR$_3$–DR$_0$	0	0
DR$_6$	FFFF0FF0H	FFFF0FF0H
DR$_7$	00000400H	00000400H

Notes: BIST = built-in self-test, XXXX = Pentium II version number.

$\overline{\text{RP}}$	**Request parity** is used to request parity.
$\overline{\text{RS2}}$–$\overline{\text{RS0}}$	**Request status** inputs are used to request the current status of the Pentium II.
$\overline{\text{RSP}}$	The **response parity** input is activated to request parity.
$\overline{\text{SLOTOCC}}$	The **slot occupied** output is a logic 0 if slot zero contains either a Pentium II or a dummy terminator.
$\overline{\text{SLP}}$	**Sleep** is an input that, when inserted in the stop-grant state, causes the Pentium II to enter the sleep state.
$\overline{\text{SMI}}$	The **system management interrupt** input causes the Pentium II to enter the system management mode of operation.
$\overline{\text{STPCLK}}$	The **stop clock** input causes the Pentium II to enter the low-power stop-grant state.
TCK	The **testability clock** input selects the clocking function in accordance with the IEEE 1149.1 Boundary Scan interface.
TDI	The **test data input** is used to test data clocked into the Pentium II with the TCK signal.
TDO	The **test data output** is used to gather test data and instruction shifted out of the Pentium II with TCK.
TESTHI	**Test high** is an input that must be connected to +2.5 V through a 1K–10K Ω resister for proper Pentium II operation.
$\overline{\text{THERMTRIP}}$	**Thermal sensor trip** is an output that becomes a zero when the temperature of the Pentium II exceeds 130°C.
TMS	The **test mode select** input controls the operation of the Pentium in test mode.
$\overline{\text{TRDY}}$	**Target ready** is an input that is used to cause the Pentium II to perform a write-back operation.
$\overline{\text{VID4}}$–$\overline{\text{VID0}}$	**Voltage data** output pins are either open or grounded signals that indicate what supply voltage is currently required by the Pentium II. The power supply must apply the request voltage to the Pentium II, as listed in Table 19–2.

The Memory System

The memory system for the Pentium II microprocessor is 64G bytes in size, just like the Pentium Pro microprocessor. Both microprocessors address a memory system that is 64 bits wide with an address bus that is 36 bits wide. Most systems use SDRAM operating at 66 MHz or 100 MHz for the Pentium II. The SDRAM for the 66 MHz system has an access time of 10 ns and the SDRAM for the 100 MHz system has an access time of 8 ns. The memory system, which connects to the chip set, is not illustrated in this chapter. Refer to earlier chapters to see the organization of a 64-bit-wide memory system without ECC.

The Pentium II memory system is divided into eight or nine banks that each store a byte of data. If the ninth byte is present, it stores an error-checking code (ECC). The Pentium II, like the 80486–Pentium Pro, employs internal parity generation and checking logic for the memory system's data bus information. (Note that most Pentium II systems do not use parity checks, but it is available.) If parity checks are employed, each memory bank contains a ninth bit. The 64-bit-wide memory is important to double-precision floating-point data. Recall that a double-precision floating-point number is 64 bits wide. As with the Pentium Pro, the memory system is numbered in bytes from byte 000000000H to byte FFFFFFFFFH. Please note that none of the current chip sets support more than 1G byte of system memory, so the additional address connections are for

TABLE 19–2 Power supply voltages for the Pentium II as requested by the $\overline{VID}$ pins.

$\overline{VID4}$	$\overline{VID3}$	$\overline{VID2}$	$\overline{VID1}$	$\overline{VID0}$	Vcc
0	0	0	0	0	2.05 V
0	0	0	0	1	2.00 V
0	0	0	1	0	1.95 V
0	0	0	1	1	1.90 V
0	0	1	0	0	1.85 V
0	0	1	0	1	1.80 V
0	0	1	1	0	—
0	0	1	1	1	—
0	1	0	0	0	—
0	1	0	0	1	—
0	1	0	1	0	—
0	1	0	1	1	—
0	1	1	0	0	—
0	1	1	0	1	—
0	1	1	1	0	—
0	1	1	1	1	—
1	0	0	0	0	3.5 V
1	0	0	0	1	3.4 V
1	0	0	1	0	3.3 V
1	0	0	1	1	3.2 V
1	0	1	0	0	3.1 V
1	0	1	0	1	3.0 V
1	0	1	1	0	2.9 V
1	0	1	1	1	2.8 V
1	1	0	0	0	2.7 V
1	1	0	0	1	2.6 V
1	1	0	1	0	2.5 V
1	1	0	1	1	2.4 V
1	1	1	0	0	2.3 V
1	1	1	0	1	2.2 V
1	1	1	1	0	2.1 V
1	1	1	1	1	—
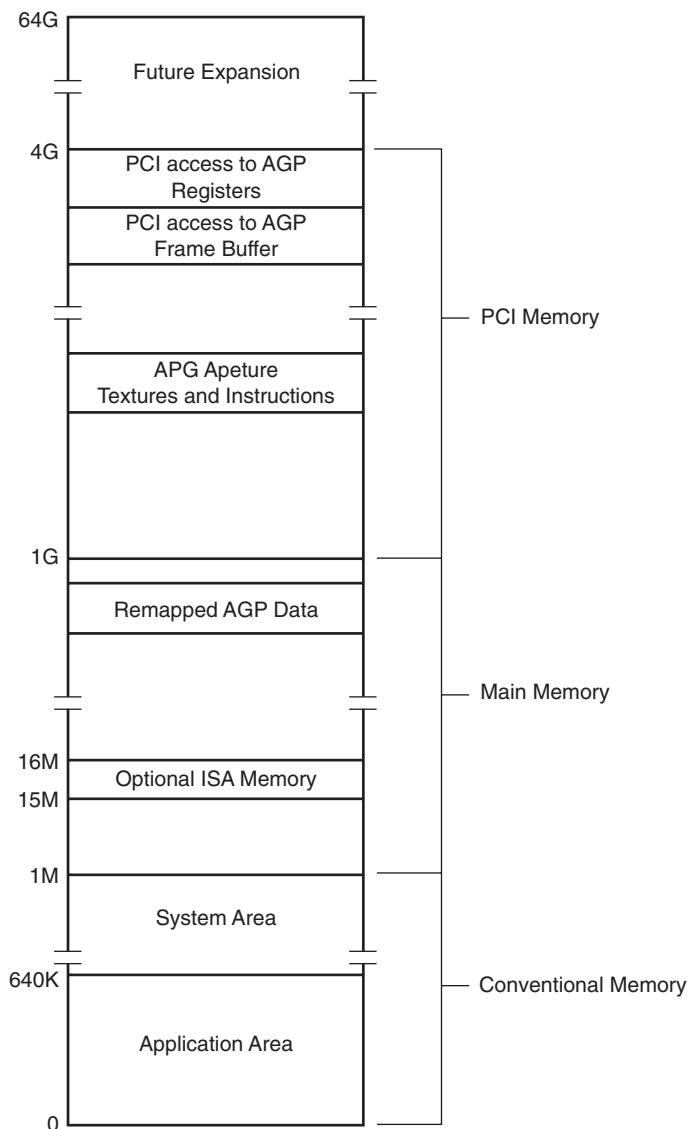
future expansion. Figure 19–3 illustrates the basic memory map of the Pentium II system, using the AGP for the video card.

The memory map for the Pentium II system is similar to the map illustrated in earlier chapters, except that an area of the memory is used for the AGP area. The AGP area allows the video card and Windows to access the video information in a linear address space. This is unlike the 128K-byte window in the DOS area for a standard VGA video card. The benefit is much faster video updates because the video card does not need to page through the 128K-byte DOS video memory.

Transfers between the Pentium II and the memory system are controlled by the 440 LX or 440 BX chip set. Data transfers between the Pentium II and the chip set are eight bytes wide. The chip set communicates to the microprocessor through the five $\overline{REQ}$ signals, as listed in Table 19–3. In essence, the chip set controls the Pentium II, which is a departure from the traditional method of connecting a microprocessor to the system directly to the memory.

The Pentium II connects only directly to the cache, which is on the Pentium II cartridge. As mentioned, the Pentium II cache operates at one half the clock frequency of the micro-processor. Therefore, a 400 MHz Pentium II cache operates at 200 MHz. The Pentium II Xeon

FIGURE 19–3 The memory map of a Pentium II–based computer system.

cache operates at the same frequency as the microprocessor, which means that the Xeon, with its 5I2K, 1M, or 2M cache, outperforms the standard Pentium II.

Input/Output System

The input/output system of the Pentium II is completely compatible with earlier Intel microprocessors. The I/O port number appears on address lines A_{15}–A_3 with the bank-enable signals used to select the actual memory banks used for the I/O transfer. Transfers are controlled by the chip set, which is a departure from the standard microprocessor architecture before the Pentium II.

Beginning with the 80386 microprocessor, I/O privilege information is added to the TSS segment when the Pentium II is operated in the protected mode. Recall that this allows I/O ports to be selectively inhibited. If the blocked I/O location is accessed, the Pentium II generates a type 13 interrupt to signal an I/O privilege violation.

TABLE 19–3   The $\overline{REQ}$ signals to the Pentium II.

$\overline{REQ4}-\overline{REQ0}$	Name	Comment
00000	Deferred reply	Deferred replies are issued for previously deferred transactions
00001	Reserved	—
00010	Memory read & invalidate	Memory read from DRAM or PCI write to DRAM from PCI
00011	Reserved	—
00100	Memory code read	Memory read cycle
00101	Memory write-back	Memory write-back cycle
00110	Memory data read	Memory read cycle
00111	Memory write	Normal memory write cycle
01000	Interrupt acknowledge or special cycle	Interrupt acknowledge for PCI bus
01001	Reserved	—
10000	I/O read	I/O read operation
10001	I/O write	I/O write operation

System Timing

As with any microprocessor, the system timing signals must be understood in order to interface the microprocessor, or so it was at one time. Because the Pentium II is designed to be controlled by the chip set, the timing signals between the microprocessor and the chip set have become proprietary to Intel.

19–2   PENTIUM II SOFTWARE CHANGES

The Pentium II microprocessor core is a Pentium Pro. This means that the Pentium II and the Pentium Pro are essentially the same device for software. This section of the text lists the changes to the CPUID instruction and the SYSENTER, SYSEXIT, FXSAVE, and FXRSTORE instructions (the only modifications to the software).

CPUID Instruction

Table 19–4 lists the values passed between the Pentium II and the CPUID instruction. These are changed from earlier versions of the Pentium microprocessor.

The version information returned after executing the CPUID instruction with a logic 0 in EAX is returned in EAX. The family ID is returned in bits 8 to 11; the model ID is returned in bits 4 to 7. The stepping ID is returned in bits 0 to 3. For the Pentium II, the model number is 6 and the family ID is a 3. The stepping number refers to an update number—the higher the stepping number, the newer the version.

The features are indicated in the EDX register after executing the CPUID instruction with a zero in EAX. Only two new features are returned in EDX for the Pentium II. Bit position 11 indicates whether the microprocessor supports the two new fast call instructions, SYSENTER and SYSEXIT. Bit position 23 indicates whether the microprocessor supports the MMX instruction set introduced in Chapter 14. The remaining bits are identical to earlier versions of the microprocessor and are not described. Bit 16 indicates whether the microprocessor supports the page attribute table or PAT. Bit 17 indicates whether the microprocessor supports the page size

TABLE 19–4 CPUID instruction for the Pentium II.

Input EAX	Output Register	Contents
0	EAX	Maximum allowed input to EAX for CPUID
0	EBX	"uneG"
0	ECX	"Inei"
0	EDX	"letn"
1	EAX	Version number
1	EDX	Feature information
2	EAX	Cache data
2	EBX	Cache data
2	ECX	Cache data
2	EDX	Cache data

extension found with the Pentium Pro and Pentium II microprocessors. The page size extension allows memory above 4G through 64G to be addressed. Finally, bit 24 indicates whether the fast floating-point save (FXSAVE) and restore (FXRSTOR) instructions are implemented.

SYSENTER and SYSEXIT Instructions

The SYSENTER and SYSEXIT instructions use the fast call facility introduced in the Pentium II microprocessor. Please note that these instructions function only in ring 0 (privilege level 0) in protected mode. Windows operates in ring 0, but does not allow applications access to ring 0. These new instructions are meant for operating system software because they will not function at any other privilege level.

The SYSENTER instruction uses some of the model-specific registers to store CS, EIP, and ESP to execute a fast call to a procedure defined by the model-specific register. The fast call is different from a regular call because it does not push the return address onto the stack as a regular call. Table 19–5 illustrates the model-specific register used with SYSENTER and SYSEXIT. Note that the model-specific registers are read with the RDMSR instruction and written with the WRMSR instruction.

To use the RDMSR or WRMSR instructions, place the register number in the ECX register. If the WRMSR is used, place the new data for the register in EDS:EAX. For the SYSENTER instruction, you need use only the EAX register, but place a zero into EDX. If the RDMSR register instruction is used in a program the data is returned in the EDX:EAX registers.

To use the SYSENTER instruction, you must first load the model-specific registers with the address of the system entrance point into the SYSENTER_CS, SYSENTER_ESP, and SYSENTER_EIP registers. This would normally be the entrance address and stack area of the operating system such as Windows 2000 or Windows XP. Note that this instruction is meant as a system instruction to access code or software in ring 0. The stack segment register is loaded with the value placed into SYSENTER_CS plus 8. In other words, the selector pair addressed by SYSENTER_CS selector value is loaded into CS and SS. The value of the stack offset is loaded into SYSENTER_ESP.

TABLE 19–5 The model-specific registers used with SYSENTER and SYSEXIT.

Name	Number	Function
SYSENTER_CS	174H	SYSENTER target code segment
SYSENTER_ESP	175H	SYSENTER target stack pointer
SYSENTER_EIP	176H	SYSENTER target instruction pointer

TABLE 19–6 Selectors addressed by the SYSENTER_CS value.

SYSENTER_CS (MSR 174H)	Function
SYSENTER_CS value	SYSENTER code segment selector
SYSENTER_CS value +8	SYSENTER stack segment selector
SYSENTER_CS value + 16	SYSEXIT code segment selector
SYSENTER_CS value + 24	SYSEXIT stack segment selector

The SYSEXIT instruction loads CS and SS with the selector pair addressed by SYSENTER_CS plus 16 and 24. Table 19–6 illustrates the selectors from the global selector table, as addressed by SYSENTER_CS. In addition to the code and stack segment selector and the memory segments that they represent, the SYSEXIT instruction passes the value in EDX to the EIP register and the value in ECX to the ESP register. The SYSEXIT instruction returns control back to application ring 3. As mentioned, these instructions appear to have been designed for quick entrance and return from the Windows or Windows NT operating systems on the personal computer.

To use SYSENTER and SYSEXIT, the SYSENTER instruction must pass the return address to the system. This is accomplished by loading the EDX register with the return offset and by placing the segment address into the global descriptor table at location SYSENTER_CS+16. The stack segment is transferred by loading the stack segment selector into SYSENTER_CS+24 and the ESP into the ECX.

FXSAVE and FXRSTOR Instructions

The last two new instructions added to the Pentium II microprocessor are the FXSAVE and FXRSTOR instructions, which are almost identical to the FSAVE and FRSTOR instructions detailed in Chapter 14. The main difference is that the FXSAVE instruction is designed to properly store the state of the MMX machine, while the FSAVE properly stores the state of the floating-point coprocessor. The FSAVE instruction stores the entire tag field, whereas the FXSAVE instruction only stores the valid bits of the tag field. The valid tag field is used to reconstruct the restore tag field when the FXRSTOR instruction executes. This means that if the MMX state of the machine is saved, use the FXSAVE instruction; if the floating-point state of the machine is saved, use the FSAVE instruction. For new applications, it is recommended that the FXSAVE and FXRSTOR instructions should be used to save the MMX state and floating-point state of the machine. Do not use the FSAVE and FRSTOR instructions in new applications.

19–3   THE PENTIUM III

The Pentium III microprocessor is an improved version of the Pentium II microprocessor. Even though it is newer than the Pentium II, it is still based on the Pentium Pro architecture.

There are two versions of the Pentium III. One version is available with a nonblocking 512K-byte cache and packaged in the slot 1 cartridge, and the other version is available with a 256K-byte advanced transfer cache and packaged in an integrated circuit. The slot-1 version cache runs at half the processor speed, and the integrated-cache version runs at the processor clock frequency. As shown in most benchmarks of cache performance, increasing the cache size from 256K bytes to 512K bytes only improves performance by a few percent.

Chip Sets

The chip set for the Pentium III is different from the Pentium II. The Pentium III uses an Intel 810, 815, or 820 chip set. The 815 is most commonly found in newer systems that use the Pentium III. A few other vendors' chip sets are available, but problems with drivers for new

peripherals, such as the video cards, have been reported. An 840 chip set also was developed for the Pentium III, but Intel did not make it available.

Bus

The Coppermine version of the Pentium III increases the bus speed to either 100 MHz or 133 MHz. The faster version allows transfers between the microprocessor and the memory at higher speeds. The last released version of the Pentium III was a 1 GHz microprocessor with a 133 MHz bus.

Suppose that you have a 1 GHz microprocessor that uses a 133 MHz memory bus. You might think that the memory bus speed could be faster to improve performance, and we agree. However, the connections between the microprocessor and the memory preclude using a higher speed for the memory. If we decided to use a 200 MHz bus speed, we must recognize that a wavelength at 200 MHz is 300,000,000/200,000,000 or 3/2 meter. An antenna is 1/4 of a wavelength. At 200 MHz, an antenna is 14.8 inches. We do not want to radiate energy at 200 MHz, so we need to keep the printed circuit board connections shorter than 1/4-wavelength. In practice, we would keep the connections to no more than 1/10 of 1/4-wavelength. This means that the connections in a 200 MHz system should be no longer than 1.48 inches. This size presents the main board manufacturer with a problem when placing the sockets for a 200 MHz memory system. A 200 MHz bus system may be the limit for the technology. If bus is tuned, there may be a way to go higher in frequency; only time will determine if it is possible. At present all that can be done is a play on words in advertisements such as 800M bytes per second to rate a bus. (Since 64 bits [8 bytes] are transferred at a time, 800M bytes per second is really 100 MHz.)

Will it be possible to exceed the 200 MHz memory system? Yes, if we develop a new technology for interconnecting the microprocessor, chip set, and memory. At present the memory functions in bursts of four 64-bit numbers each time we read the main memory. This burst of 32 bytes is read into the cache. The main memory requires three wait states at 100 MHz to access the first 64-bit number and then zero wait states for each of the three remaining 64-bit wide numbers for a total of seven 100 MHz bus clocks. This means we are reading data at 70 ns / 32 = 2.1875 ns per byte, which is a bus speed of 457M bytes per second. This is slower than the clock on a 1 GHz microprocessor, but because most programs are cyclic and the instructions are stored in an internal cache, we can and often do approach the operating frequency of the microprocessor.

Pin-Out

Figure 19–4 shows the pin-out of the socket 370 version of the Pentium III microprocessor. This integrated circuit is packaged in a 370-pin, pin grid array (PGA) socket. It is designed to function with one of the chip sets available from Intel. In addition to the full version of the Pentium III, the Celeron, which uses a 66 MHz memory bus speed, is available. The Pentium III Xeon, also manufactured by Intel, allows larger cache sizes for server applications.

19–4   THE PENTIUM 4 AND CORE2

The most recent version of the Pentium Pro architecture microprocessor is the Pentium 4 microprocessor and recently the Core2 from Intel. The Pentium II, Pentium III, Pentium 4, and Core2 are all versions of the Pentium Pro architecture. The Pentium 4 was released initially in November 2000 with a speed of 1.3 GHz. It is currently available in speeds up to 3.8 GHz. Two packages are available for early versions of this integrated microprocessor, the 423-pin PGA and the 478-pin FC-PGA2. Both versions of the original issue of the Pentium 4 used the 0.18 micron technology for fabrication. The most recent versions use either the 0.13 micron technology or the 90 nm (0.09 micron) technology. Newer versions of the Pentium 4 use the LGA (leadless grid array) 775 package, which has 775 pins. Intel is currently developing a 45 nm technology for
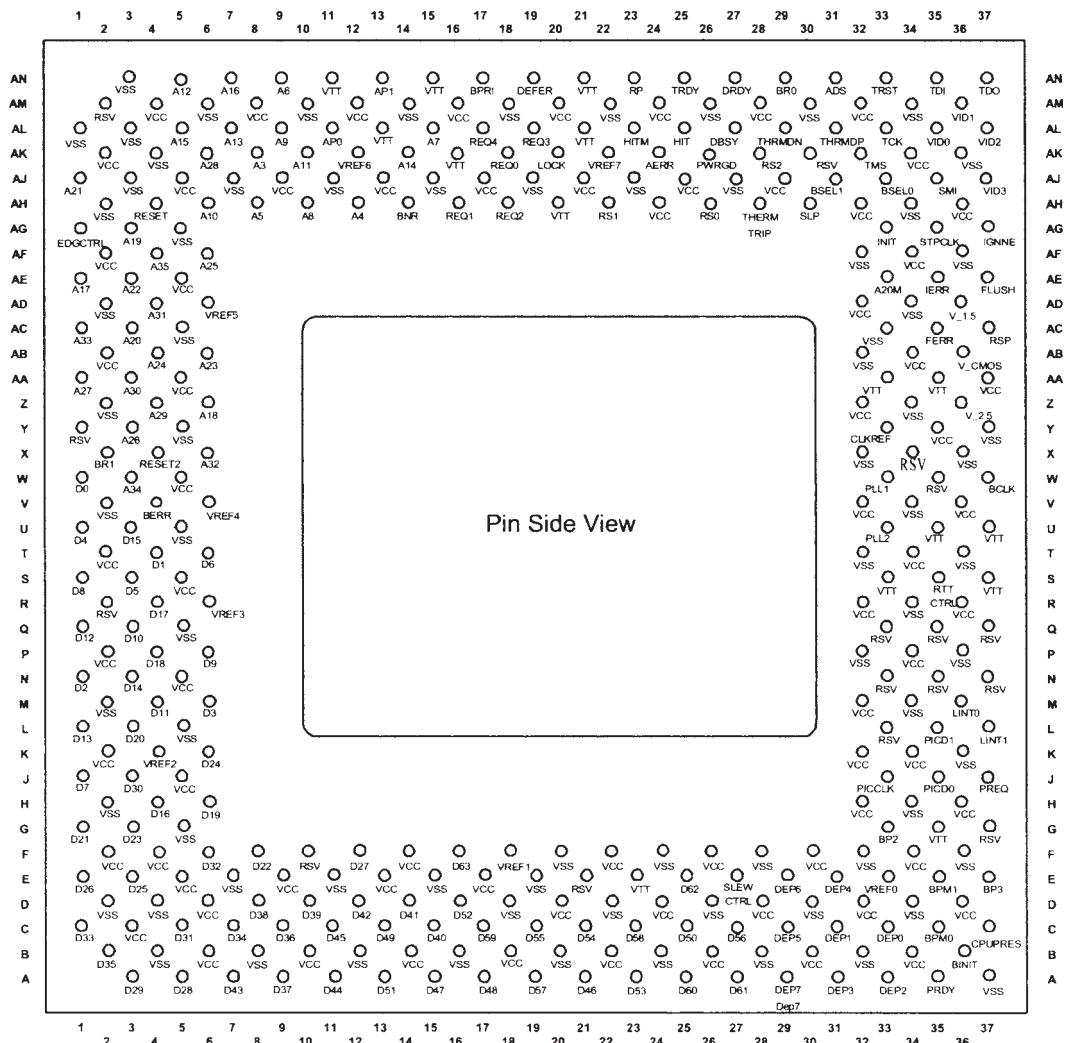
FIGURE 19–4 The pin-out of the socket 370 version of the Pentium III microprocessor. (Courtesy of Intel Corporation.)

future products. As with earlier versions of the Pentium III, the Pentium 4 uses a 100 MHz memory bus speed, but because it is quad pumped, the bus speed can approach 400 MHz. More recent versions use the 133 MHz bus listed as 533 MHz because of quad pumping or 200 MHz listed as 800 MHz. Some newer versions use a 1033 MHz or 1333 MHz front side bus; another package called the LGA 771 has appeared in newer versions of the Xeon. Figure 19–5 illustrates the pin-out of the 423-pin PGA of the Pentium 4 microprocessor.

Memory Interface

The memory interface to the Pentium 4 typically uses the Intel 945, 965, or 975 chip set. These chip sets provide a dual-pipe memory bus to the microprocessor with each pipe interfaced to a 32-bit-wide section of the memory. The two pipes function together to comprise the 64-bit-wide data path to the microprocessor. Because of the dual-pipe arrangement, the memory must be populated with pairs of DDR2 memory devices operating at 600 MHz, 800 MHz, or 1033 MHz. According to Intel the DDR2 arrangement provides a 300% increase in speed over a memory populated with PC-100 memory.
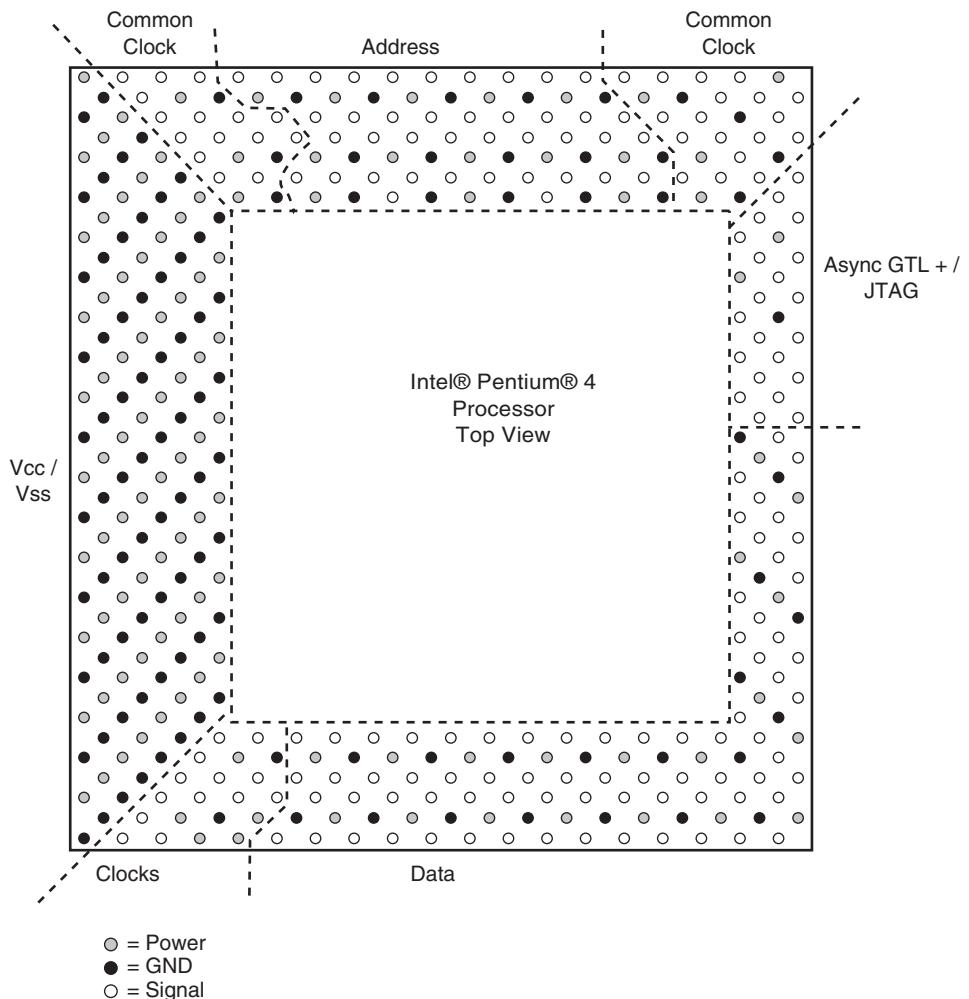
FIGURE 19–5   The pin-out of the Pentium 4, 423 PGA. (Courtesy of Intel Corporation.)

Intel has abandoned RDRAM in favor of DDR2 (double data rate) memory beginning with the 965 and 975 chip sets. Apparently the claim of a 300% increase in RDRAM speed failed to prove factual. In addition to the inclusion of support for DDR2, memory support for the serial ATA disk interface has also been added.

Newer chip sets such as the 945 and 965 contain the PCI Express interface and do not contain the AGP interface. The AGP interface is replaced by PCI Express for video support. IDE support remains for interface to legacy devices such as older HDD, CD-ROM, and DVD drives.

Register Set

The Pentium 4 and Core2 register set is nearly identical to all other versions of the Pentium except that the MMX registers are separate entities from the floating-point registers. In addition, eight 128-bit-wide XMM registers are added for use with the SIMD (single-instruction, multiple data) instructions as explained in Chapter 14 and the extended 128-bit packed doubled floating-point numbers.

You might think of the XMM registers as double-wide MMX registers that can hold a pair of 64-bit double-precision floating-point numbers or four single-precision floating-point numbers. Likewise they can also hold 16-byte-wide numbers as the MMX registers hold 8-byte-wide numbers. The XMM registers are double-width MMX registers.

If the new patch for MASM 6.15 is downloaded from Microsoft, programs can be assembled using both the MMX and XMM instructions. The ML.EXE program is also found in Microsoft Visual Studio.net 2003. To assemble programs that include MMX instructions, use the .MMX switch. For programs that include the SIMD instructions, use the .XMM switch. Example 19–1 illustrates a very simple program that uses the MMX instructions to add two 8-byte-wide numbers together. Notice how the .MMX switch is used to select the MMX instruction set. The MOVQ instructions transfer numbers between memory and the MMX registers. The MMX registers are numbered from MM_0 to MM_7. You can also use the MMX and SIMD instructions in Microsoft Visual C++ using the inline assembler if you download the latest patch from Microsoft for Visual Studio version 6.0 or use a newer version of Visual Studio. It is recommended that Visual Studio Express, which contains the patch, is used for software development.

EXAMPLE 19–1

```
.MMX
.DATA
      DATA1 DQ   1FFH
      DATA2 DQ   101H
      DATA3 DQ   ?
.CODE
      MOVQ  MM0,DATA1
      MOVQ  MM1,DATA2
      PADDB MM0,MM1
      MOVQ  DATA3,MM0
```

Similarly, the XMM software can be used in a program with the .XMM switch. Most modern programs use the XMM registers and the XMM instruction set to accomplish multimedia and other high-speed operations. Example 19–2 shows a short program that illustrates the use of a few XMM instructions. This program multiplies two sets of four single-precision floating-point numbers and stores the four products into the four doublewords at ANS. In order to enable access to octal words (128-byte-wide numbers), we use the OWORD PTR directive. Also notice that the FLAT model is used with the C profile. The SIMD instructions only function in protected mode so the program uses the FLAT model format. This means that the .686 and .XMM switches are both placed before the model statement.

EXAMPLE 19–2

```
.686
.XMM
.MODEL FLAT,C
.DATA
      DATA1 DD  1.0              ;define four floats for DATA1
            DD  2.0
            DD  3.0
            DD  4.0
      DATA2 DD  6.3              ;define four floats for DATA2
            DD  4.6
            DD  4.5
            DD  -2.3
      ANS   DD  4 DUP(?)
.CODE
      MOVAPS XMM0,OWORD PTR DATA1
      MOVAPS XMM1,OWORD PTR DATA2
      MULPS  XMM0,XMM1
      MOVAPS OWORD PTR ANS,XMM0

;additional code here

END
```

Hyper-Threading Technology
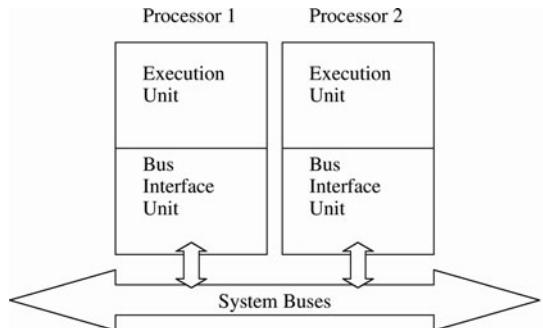
The most recent innovation and new to the Pentium is called hyper-threading technology. This significant advancement combines two microprocessors into a single package. To understand this new technology, refer to Figure 19–6, which shows a traditional dual processor system and a hyper-threaded system.

The hyper-threaded processor contains two execution units that each contain a complete set of the registers capable of running software independently or concurrently. These two separate machine contexts share a common bus interface unit. During machine operation each processor is capable of running a thread (process) independently, increasing the execution speed of an application that is written using multiple threads. The bus interface unit contains the level 2 and level 3 caches and the interface to the memory and I/O structure of the machine. When either microprocessor needs to access memory or I/O, it must share the bus interface unit.
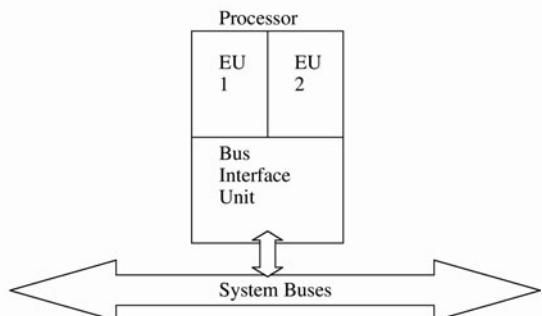
The bus interface unit is in use to access memory, but since memory is accessed in bursts that fill caches, it is often idle. Because of this, a second processor can use this idle time to access memory while the other processor is busy executing instructions. Does the speed of the system double? Yes and no. Some threads can run independently of each other as long as they do not access the same area of memory. If each thread accesses the same area of memory, the machine can actually run slower with hyper-threaded technology. This does not occur very often, so in most cases the system performance increases with hyper-threading achieving nearly the same performance as with a dual processor system.

Eventually most machines will use hyper-threading technology, which means that more attention should be given to developing software that is multi-threaded. Each thread runs on a different processor in a system that has either dual processors or hyper-threaded processors, increasing performance. In the future the architecture may include even more processors to handle additional threads.

FIGURE 19–6 Systems illustrating dual processors and a hyper-threaded processor.

(a) Dual processor system.

(b) A hyper-threaded system.

Multiple Core Technology

Most new versions of the Pentium 4 and Core2 contain either dual or quad cores. Each core is a separate version of the microprocessor that independently executes a separate task. Three versions are currently available: the Pentium D, which contains two cores; with separate caches; a Core2 Duo version that contains a shared cache, but two cores and a quad core version, which contains four cores. Intel seems to have migrated to a shared cache for multiple core microprocessors. A recent article from Intel stated that in the future the Pentium or whatever it will be called may contain up to 80 cores. The Core2 Duo contains either a 2M or 4M byte cache and operates at frequencies to 3 GHz. It certainly appears that the speed race is over and the clock frequency has stabilized at between 3 and 4GHz. Does this mean that in the future a 5 GHz version will never become available? It is possible, but at this time a much higher clock frequency appears to be impossible, so multiple cores using threaded application seem to be the prospect for some time to come. It appears silicon technology has reached its apex. What this means is that efficient programming will become the avenue for increasing the speed of computer systems.

CPUID

As in earlier versions of the Pentium, the CPUID instruction accesses information that indicates the type of microprocessor as well as the features supported by the microprocessor. In the ever-evolving series of microprocessors it is important to be able to access this information so that efficient software can be written to operate on many different versions of the microprocessor.

Table 19–7 lists the latest features available to the CPUID instruction. To access these features, EAX is loaded with the input number listed in the table, then the CPUID instruction is executed. The CPUID instruction usually returns information in the EAX, EBX, ECX, and EDX registers in the real or protected mode. As can be gleaned from the table, additional features have been added to the CPUID instruction when compared to previous versions.

TABLE 19–7   Pentium 4 CPUID instruction.

EXA Input Value	Output Registers	Notes
0	EAX = Maximum input value EBX = "uneG" ECX = "Iene" EDX = "Ietn"	"GenuineIntel" is returned in little endian format
1	EAX = Version information EBX = Feature information ECX = Extended feature information EDX = Feature information	Feature information
2	EAX, EBX, ECX, and EDX	Cache and TLB information
3	ECX and EDX	Serial number in the Pentium III only
4	EAX, EBX, ECX, and EDX	Deterministic cache parameters
5	EAX, EBX, ECX, and EDX	Monitor/Mwait information
80000000H	EAX	Extended function information
80000001H	EAX	Reserved
80000002H, 80000003H, and 80000004H	EAX, EBX, ECX, and EDX	Processor brand string
80000006H	ECX	Cache information

In Chapter 18 software was developed to read and display the data available after the CPUID instruction was invoked with EAX = 1. Here we deal with reading the processor brand string and prepare it for display in a Visual C++ function. The brand string, if supported, contains the frequency that the microprocessor is certified to operate and also the genuine Intel keyword. The BrandString function (see Example 19–3) returns a CString that contains the information stored in the CPUID members 80000002H–80000004H. This software requires a Pentium 4 system for proper operation as tested for in BrandString function. The Convert function reads the contents of EAX, EBX, ECX, and EDX from the register specified as the parameter and converts them to a CString that is returned. The author's system shows that the brand string is

"Intel(R) Pentium(R) 4 CPU 3.06GHz"

EXAMPLE 19–3

```
int getCPU(int EAXvalue)
{
      int temp;
      _asm
      {
            mov  eax,EAXvalue
            cpuid
            mov  temp1,eax
      }
      return temp;
}


private: System::String^ BrandString(void)
{
      String^ temp;
      int temp1 = getCpu(0x8000000);
      if ( temp1 >= 0x80000004 )    //if brand string present
      {
            temp += Convert(0x80000002);      //read register 80000002H
            temp += Convert(0x80000003);      //read register 80000003H
            temp += Convert(0x80000004);      //read register 80000004H
      }
      return temp;
}


private: System::String^ Convert(int EAXvalue)
{
      CString temp ="                   ";  //must be 16 spaces
      int temp1, temp2, temp3, temp4;
      _asm
      {
            mov  eax,EAXvalue
            cpuid
            mov  temp1,eax
            mov  temp2,ebx
            mov  temp3,ecx
            mov  temp4,edx
      }
      for ( int a = 0; a <4; a++ )
      {
            temp.SetAt(a, temp1);
            temp.SetAt(a + 4, temp2);
            temp.SetAt(a + 8, temp3);
            temp.SetAt(a + 12, temp4);
            temp1 >>= 8;
            temp2 >>= 8;
            temp3 >>= 8;
            temp4 >> =8;
      }
      return temp;
}
```
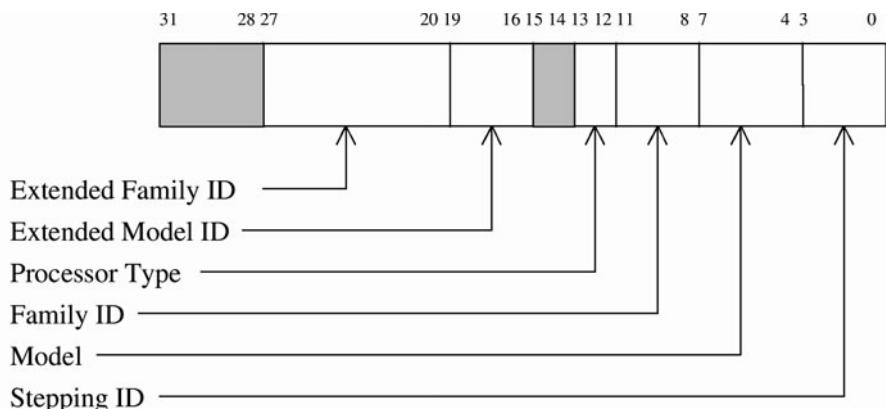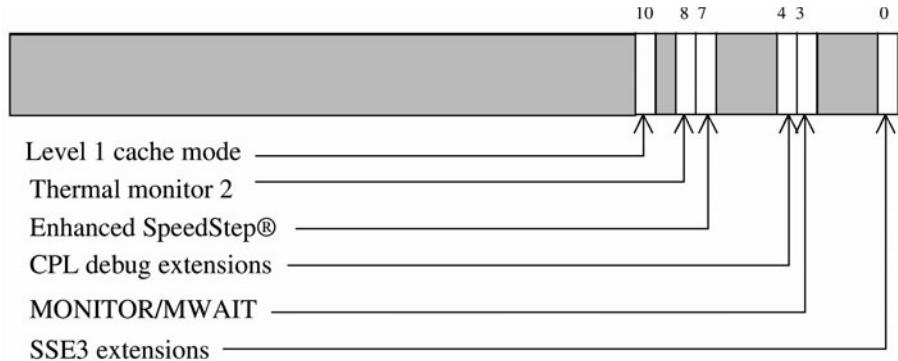
FIGURE 19–7   EAX after a CPUID instruction showing version information.

The other information available about the system is returned in EAX, EBX, ECX, and EDX after executing CPUID after loading EAX with a 1. The EAX register contains the version information as the model, family, and stepping information, as illustrated in Figure 19–7. The EBX register contains information about the cache, such as the size of the cache line flushed by the CFLUSH instruction in bits 15–8 and the ID assigned the local APIC on reset in bits 31–24. Bits 23–16 indicate how many internal processors are available to hyper-threading (two for the current Pentium 4 microprocessor). Example 19–4 shows a function that identifies the number of processors in a hyper-threaded CPU and returns it as a character string. If more than nine processors are eventually added to the microprocessor, then the software in Example 19–4 would need to be modified.

EXAMPLE 19–4
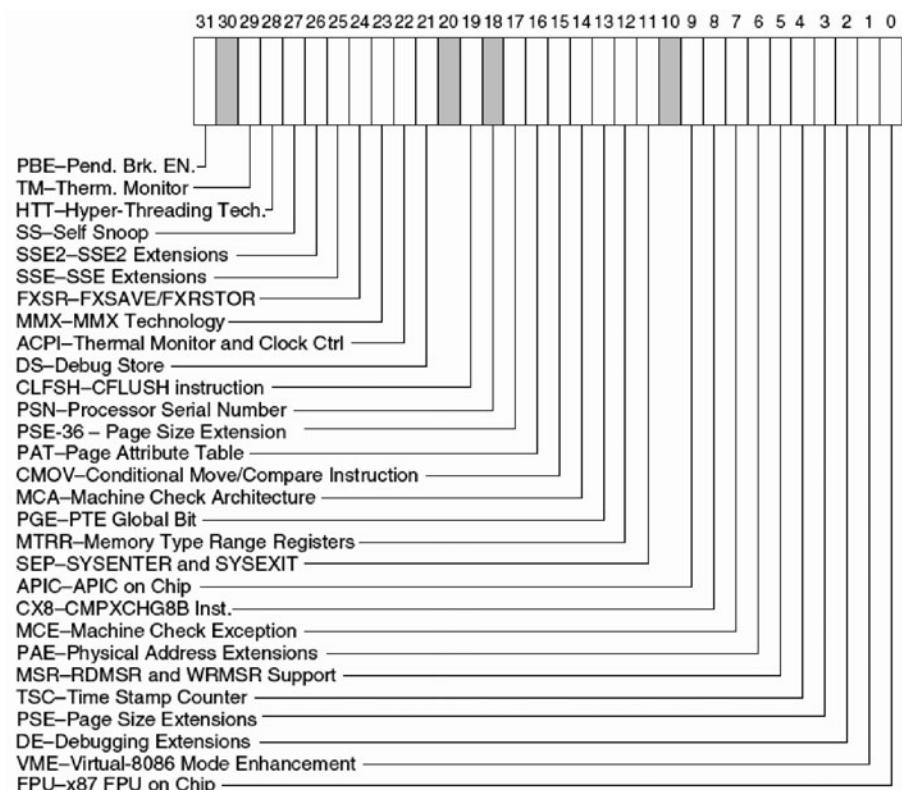
```
CString CCPUIDDlg::GetProcessorCount(void)
{
      CString temp = "This CPU has ";
      char temp1;
      _asm
      {
            mov  eax,1
            cpuid
            mov  temp1,31h
            bt   edx,28                    ;check for hyper-threading
            jnc  GetPro1                   ;if no hyper-threading, temp1 = 1
            bswap ebx
            add  bh,30h
            mov  temp1,bh
GetPro1:    }
      return temp + temp1 + "processors.";
}
```

Feature information for the microprocessor is returned in ECX and EDX as indicated in Figures 19–8 and 19–9. Each bit is a logic 1 if the feature is present. For example, if hyper-threading is needed in an application bit, position 28 is tested in EDX to see if hyper-threading is supported. This appears in Example 19–4 along with reading the number of processors found in a hyper-threaded microprocessor. The BT instruction tests the bit indicated and places it into the carry flag. If the bit under test is a 1, then the resultant carry is one and if the bit under test is a 0, the resultant carry is zero.

FIGURE 19–8   ECX after a CPUID instruction showing the version extensions.

FIGURE 19–9   EDX after a CPUID instruction showing the version extensions.

Model-Specific Registers

As with earlier versions of the Pentium, the Pentium 4 and Core2 also contain model-specific registers that are read with the RDMSR instruction and written with the WRMSR instruction. The Pentium 4 and Core2 each have 1743 model-specific registers numbered from 0H to 6CFH.

Intel does not provide information on all of them. The registers not identified are either reserved by Intel or used for some undocumented feature or function.

Both the read and write model-specific register instructions function in the same manner. Register ECX is loaded with the register number to be accessed, and the data are transferred through the EDX:EAX register pair as a 64-bit number where EDX is the most significant 32 bits and EAX is the least significant bits. These registers must be accessed in either the real mode (DOS) or in ring 0 of protected mode. These registers are normally accessed by the operating system and cannot be accessed in normal Visual C++ programming.

Performance-Monitoring Registers

Another feature in the Pentium 4 is a set of performance-monitoring registers (PMR) that, like the model-specific registers, can only be used in real mode or at ring 0 of protected mode. The only register that can be accessed via user software is the time-stamp counter, which is a performance-monitoring register. The remaining PMRs are accessed with the RDPMR. This instruction is similar to the RDMSR instruction in that it uses ECX to specify the register number and the result appears in EDX:EAX. There is no write instruction for the PMRs.

64-Bit Extension Technology

Intel has released its 64-bit extension technology for most members of the Intel 32-bit architecture family. The instruction set and architecture is backwards compatible to the 8086, which means that the instructions and register set have remained compatible. (The only things that are not compatible are a few of the legacy instructions and some instructions that deal with AH, BH, CH, and DH.) What is changed is that the register set is stretched to 64 bits in width in place of the current 32-bit-wide registers. Refer to Figure 19–10 for the programming model of the Pentium 4 and Core2 in 64-bit mode.

Notice that the register set now contains sixteen 64-bit-wide general-purpose registers, RAX, RBX, RCX, RDX, RSP, RBP, RDI, RSI, R_8–R_{15}. The instruction pointer is also stretched to a width of 64 bits, allowing the microprocessor to address memory using a 64-bit memory address. This allows the microprocessor to address as much memory as the specific implementation of the microprocessor has address pins.

The registers are addressed as 64-bit, 32-bit, 16-bit, or 8-bit registers. An example is R8 (64 bits), R8D (32 bits), R8W (16 bits), and R8L (8 bits). There is no way to address the high byte (as in BH) for a numbered register; only the low byte of a numbered register can be addressed. Legacy addressing such as MOV AH,AL functions correctly, but addressing a legacy high-byte register and a numbered low-byte register is not allowed. In other words, MOV AH,R9L is not allowed, but MOV AL,R9L is allowed. If the MOV AH,R9L instruction is included in a program no error will occur; instead the instruction will be changed to MOV BPL, R9L. AH, BH, CH, and DH are changed to the low-order 8 bits (the L is for low order) of BPL, SPL, DIL, and SIL, respectively. Otherwise the legacy registers can be mixed with the new numbered registers R_8–R_{15} as in MOV R11, RAX, MOV R11D, ECX, or MOV BX, R14W.

Another addition to the architecture is a set of additional SSE registers numbered XMM_8–XMM_{15}. These registers are accessed by the SSE, SSE_2, or SSE_3 instructions. Otherwise, the SSE unit has not been changed. The control and debug registers are expanded to 64 bits in width. A new model-specific register is added to control the extended features at address C0000080H. Figure 19–11 depicts the extended feature control register.

SCE   The **system CALL enable** bit is set to enable the SYSCALL and SYSRET instructions in the 64-bit mode.

LME   The **mode enable** bit is set to allow the microprocessor to use the 64-bit extended mode.
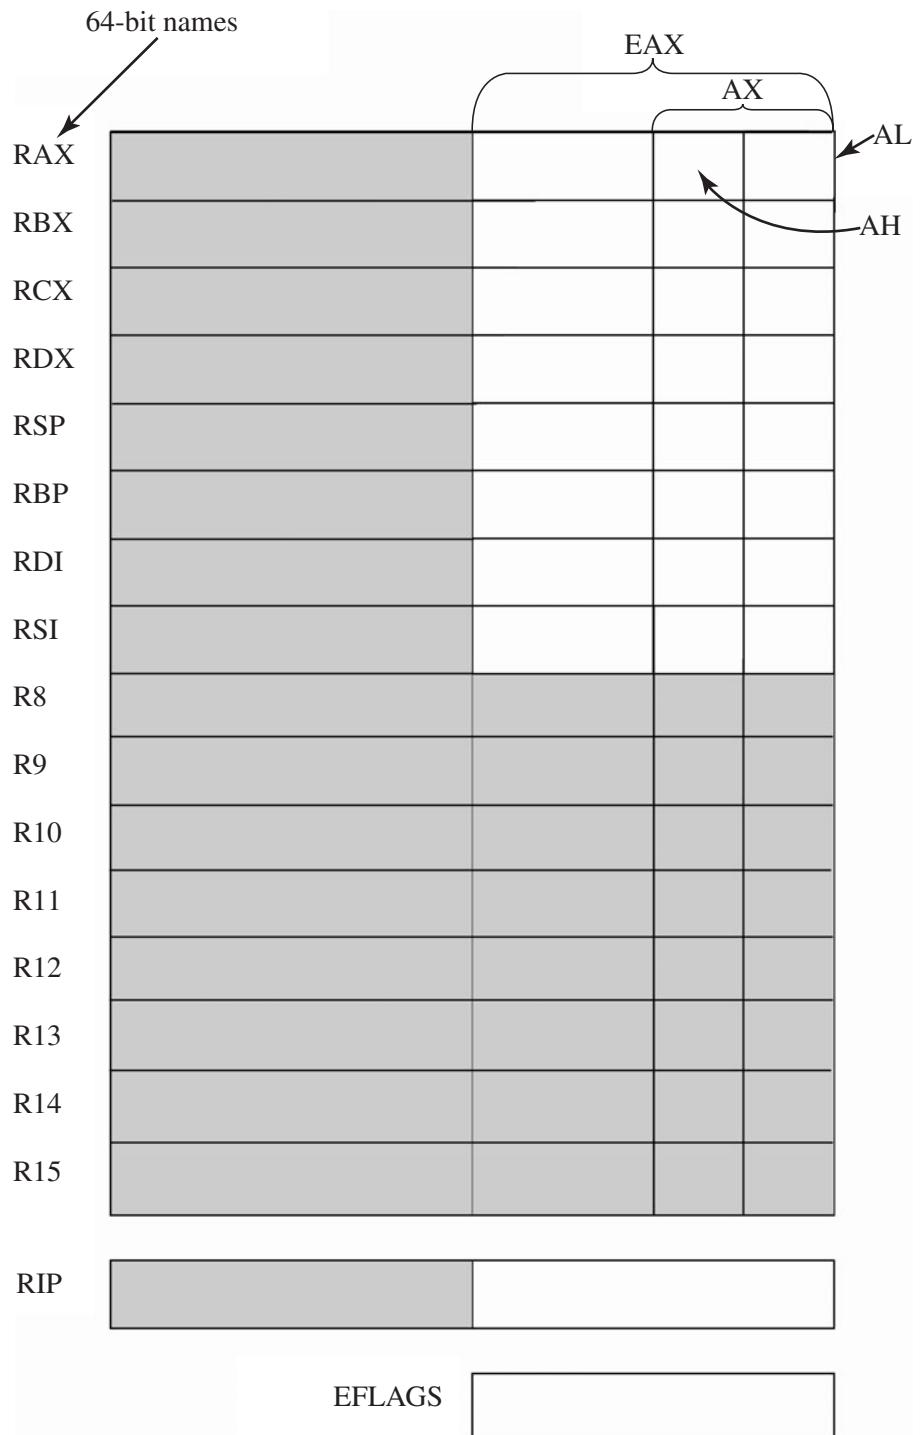
FIGURE 19–10 The integer register set in the 64-bit mode of the Pentium 4. The shaded areas are new to the Pentium 4 operated in the 64-bit mode.

FIGURE 19–11   The contents of the extended feature model-specific register.

LMA   The **mode active** bit shows that the microprocessor is operating in the 64-bit
               extended mode.

     The protected mode descriptor table registers are expanded in the extended 64-bit mode so
that each descriptor table register, GDTR, LDTR, IDTR, and the task register (TR) hold a 64-bit
base address instead of a 32-bit base address. The biggest change is that the base address and lim-
its of the segment descriptors are ignored. The system uses a base address of 0000000000000000H
for the code segment and the DS, ES, and SS segments are ignored.
     Paging is also modified to include a paging unit that supports the translation of a 64-bit lin-
ear address into a 52-bit physical address. Intel states that in the first version of this 64-bit
Pentium the linear address will be 48 bits and the physical address will be 40 bits. This means
that there will be a 40-bit address to support 1T (terra) byte of physical memory translated from
a linear address space of 256T bytes. The 52-bit address accesses 4P (peta) bytes of memory and
a 64-bit linear address accesses 16E (exa) bytes or memory. The translation is accomplished with
additional tables in the paging unit. In place of two tables (a page directory and a page table), the
64-bit extended paging unit uses four levels of page tables.

19–5   SUMMARY

1. The Pentium II differs from earlier microprocessors because instead of being offered as an
   integrated circuit, the Pentium II is available on a plug-in cartridge or printed circuit board.
2. The level 2 cache for the Pentium II is mounted inside of the cartridge, except for the
   Celeron, which has no level 2 cache. The cache speed is one half the Pentium II clock speed,
   except in the Xeon, where it is at the same speed as the Pentium II. All versions of the
   Pentium II contain an internal level 1 cache that stores 32K bytes of data.
3. The Pentium II is the first Intel microprocessor that is controlled from an external bus con-
   troller. Unlike earlier versions of the microprocessor, which issued read and write signals,
   the Pentium II is ordered to read or write information by an external bus controller.
4. The Pentium II operates at clock frequencies from 233 MHz to 450 MHz with bus speeds of
   66 MHz or 100 MHz. The level 2 cache can be 512K, 1M, or 2M bytes in size. The Pentium
   II contains a 64-bit data bus and a 36-bit address bus that allow up to 64G bytes of memory
   to be accessed.
5. The new instructions added to the Pentium II are SYSENTER, SYSEXIT, FXSAVE, and
   FXRSTOR.
6. The SYSENTER and SYSEXIT commands are optimized to access the operating system in
   privilege level 0 from a privilege level 3 access. These instructions operate at a much higher
   speed than a task switch or even a call and return combination.
7. The FXSAVE and FXRSTOR instructions are optimized to properly store the state of both
   the MMX technology unit and the floating-point coprocessor.
8. The Pentium III microprocessor is an extension of the Pentium Pro architecture with the
   addition of the SIMD instruction set that uses the XMM registers.

9. The Pentium 4 and Core2 microprocessors are extensions of the Pentium Pro architecture, which includes enhancements that allow it to operate at higher clock frequencies than previously possible because of the 0.13 micron and the latest 45 nm fabrication technologies.

10. The Pentium 4 and Core2 microprocessors require a modified ATX power supply and case to function properly in a system.

11. Version 6.15 of the MASM program and Visual Studio version 6 now support the new MMX and SIMD instructions using the .686 switch with the .MMX and .XMM switches.

12. The Pentium II, Pentium III, Pentium 4, and Core2 microprocessors are all variations of the Pentium Pro microprocessor.

13. Future Pentium 4 and Core2 microprocessors will all use the 64-bit extension to the 32-bit architecture. This will become important in systems with more than 4G bytes of memory.

19–6 QUESTIONS AND PROBLEMS

1. What is the size of the level 1 cache in the Pentium II microprocessor?
2. What sizes are available for the level 2 cache in the Pentium II microprocessor? (List all versions.)
3. What is the difference between the level 2 cache on the Pentium-based system and the Pentium II-based system?
4. What is the difference between the level 2 cache in the Pentium Pro and the Pentium II?
5. The speed of the Pentium II Xeon level 2 cache is _____ times faster than the cache in the Pentium II (excluding the Celeron).
6. How much memory can be addressed by the Pentium II?
7. Is the Pentium II available in integrated circuit form?
8. How many pin connections are found on the Pentium II cartridge?
9. What is the purpose of the PICD control signals?
10. What happened to the read and write pins on the Pentium II?
11. At what bus speeds does the Pentium II operate?
12. How fast is the SDRAM connected to the Pentium II system for a 100 MHz bus speed version?
13. How wide is the Pentium II memory if ECC is employed?
14. What new model-specific registers (MSR) have been added to the Pentium II microprocessor?
15. What new CPUID identification information has been added to the Pentium II microprocessor?
16. How is a model-specific register addressed and what instruction is used to read it?
17. Write software that stores 12H into model-specific register 175H.
18. Write a short procedure that determines whether the microprocessor contains the SYSENTER and SYSEXIT instructions. Your procedure must return carry set if the instructions are present, and return carry cleared if not present.
19. How is the return address transferred to the system when using the SYSENTER instruction?
20. How is the return address retrieved when using the SYSEXIT instruction to return to the application?
21. The SYSENTER instruction transfers control to software at what privilege level?
22. The SYSEXIT instruction transfers control to software at what privilege level?
23. What is the difference between the FSAVE and the FXSAVE instructions?
24. The Pentium III is an extension of the _____ architecture.
25. What new instructions appear in the Pentium III microprocessor that do not appear in the Pentium Pro microprocessor?

26. What changes to the power supply does the Pentium 4 or Core2 microprocessor require?
27. Write a short program that reads and displays the serial number of the Pentium III micro-processor on the video screen.
28. Develop a short C++ function that returns a bool value of true if the Pentium 4 supports hyper-threaded technology and false if it does not support it.
29. Develop a short C++ function that returns a bool value of true if the Pentium 4 or Core2 support SSE, SSE_2, and SSE_3 extensions.
30. Compare, in your own words, hyper-threading to dual processing. Postulate on the possibility of including additional processors beyond four.
31. What is a Core2 processor?

Appendix A

The Assembler, Visual C++, and DOS

This appendix is provided so that the assembler can be used for program development in the DOS environment and also the Visual C++ environment. The DOS environment is essentially gone (unless Windows 98 is still in use), but lives on through the emulation program called CMD.EXE in the accessory folder of Microsoft Windows. Some may shed a tear at the departure of DOS, but realize that the DOS environment was a vast headache to many of use who spent years programming in it. It had only a 1M memory system and drivers were a problem, especially in recent years. Microsoft never really provided a decent protected mode DOS. DOS displayed text information well, but graphics were another story because of the DOS architecture of the video memory and the lack of drivers.

Windows solved many of the problems that plagued DOS and ushered in the GUI age, which is a great improvement over DOS text-based applications. Windows is just so much easier for a human to use and control. The author remembers the old days where he had to write batch files so his wife could use his computer. Now she is a real pro because of Windows. Why—she can surf the net, and surf the net, and well, she is a real professional now. Oh, she can send an e-mail as long as it doesn't need a picture or anything attached. Teasing aside, Windows is a tremendous system—bar none.

THE ASSEMBLER

Although the assembler program is not often used as a stand-alone programming medium, it still finds some application in developing modules that are linked to Visual C++ programs (see Chapter 7). The program itself is provided with Visual C++ in the C:\Program Files\Microsoft Visual Studio .NET 2003\Vc7\bin directory as ML.EXE. Also found in the same directory is the LIB.EXE (library) program for creating library collections and the LINK program used for linking object modules.

Example A–1 illustrates how to assemble a program written in assembly language. The example uses a file called WOW.TXT (it does not need the .ASM extension even though it is often used for assembly language modules). The file WOW.TXT is compiled for errors (/c = lowercase c) and generates a listing file (/Fl) called WOW.LST. If other switches are needed, just type ML /? at the command prompt to display a list of the switches. The /coff (c object file format) switch might also be included to generate an object file that can be linked to a Visual C++ program as in `ML /c /coff WOW.TXT`.

EXAMPLE A–1

```
ML /c /FlWOW.LST WOW.TXT
```

If the LINK program is used from Visual C++, you cannot generate a DOS-compatible execution file because it is a 32-bit linker. The 16-bit linker for DOS is not in the Visual Studio package. If DOS software must be developed, obtain the Windows Driver Development Kit (Windows DDK) from Microsoft. The DDK contains the 16-bit linker needed to develop DOS applications. The linker is located in the C:\WINDDK\2600.1106\bin\win_me\bin16 folder of the DDK. In addition to the linker program, a 16-bit version of the C++ language for DOS appears as CL.EXE. These are provided for legacy applications.

Example A–2 shows how to link a program generated by the assembler. This assumes that you are using the 16-bit DOS real mode linker program. The 32-bit linker is normally used from Visual C++ for Windows applications. Here the object file generated by Example A–1 is linked to produce an executable program called WOW.EXE or, if the tiny model is in effect, a WOW.COM.

EXAMPLE A–2

```
LINK WOW.OBJ
```

ASSEMBLER MEMORY MODELS

Although the flat model is most often used with Visual C++, there are other memory models that are used with DOS applications and embedded program development. Table A–1 lists the most commonly used models for these applications. The origin is set by the .STARTUP directive in a DOS program and automatically in a flat program.

Table A–2 lists the default information for each of the models listed in Table A–1. If additional information on models is required please visit the Microsoft website and search for assembler models.

TABLE A–1   Commonly used assembler models.

Model	Description
.TINY	All data and code must fit into a single 64K-byte memory segment. Tiny programs assemble as DOS.COM files and must use an origin at 0100H for the code.
.SMALL	A two-segment model with a single code segment and a single data segment. Small programs generate DOS.EXE files and have an origin of 0000H.
.FLAT	The flat model uses a single segment of up to 4G bytes in length. Flat programs are programs that will only function in Windows with an origin of 00000000H.

TABLE A–2 Defaults for the more common assembly language models.

Model	Directives	Name	Align	Combine	Class	Group
.TINY	.CODE	_TEXT	Word	PUBLIC	'CODE'	DGROUP
	.FARDATA	FAR_DATA	Para	Private	'FAR_DATA'	
	.FARDATA?	FAR_BSS	Para	Private	'FAR_BSS'	
	.DATA	_DATA	Word	PUBLIC	'DATA'	DGROUP
	.CONST	CONST	Word	PUBLIC	'CONST'	DGROUP
	.DATA?	_BSS	Word	PUBLIC	'BSS'	DGROUP
.SMALL	.CODE	_TEXT	Word	PUBLIC	'CODE'	
	.FARDATA	FAR_DATA	Para	Private	'FAR_DATA'	
	.FARDATA?	FAR_BSS	Para	Private	'FAR_BSS'	
	.DATA	_DATA	Word	PUBLIC	'DATA'	DGROUP
	.CONST	CONST	Word	PUBLIC	'CONST'	DGROUP
	.DATA?	_BSS	Word	PUBLIC	'BSS'	DGROUP
	.STACK	STACK	Para	STACK	'STACK'	DGROUP
.FLAT	.CODE	_TEXT	Dword	PUBLIC	'CODE'	
	.FARDATA	_DATA	Dword	PUBLIC	'DATA'	
	.FARDATA?	_BSS	Dword	PUBLIC	'FBSS'	
	.DATA	_DATA	Dword	PUBLIC	'DATA'	DGROUP
	.CONST	CONST	Dword	PUBLIC	'CONST'	DGROUP
	.DATA?	_BSS	Dword	PUBLIC	'BSS'	DGROUP
	.STACK	STACK	Dword	STACK	'STACK'	DGROUP

SELECTED DOS FUNCTION CALLS

Not all DOS function calls are included because it is doubtful that they will all be used. The most recent version of DOS has function calls from function 00H to function 6CH. This text only lists the function calls that are used for simple applications. Many of the function calls were from DOS version 1.0 and have been obsolete for many years and others are used to access the disk system, which is accessed in Visual C++.

To use a DOS function call in a DOS program, place the function number in AH and other data that might be necessary in other registers, as indicated in Table A–3. Example A–3 shows an example of DOS function number 01H. This function reads the DOS keyboard and returns an ASCII character in AL. Once everything is loaded, execute the INT 21H instruction to perform the task.

EXAMPLE A–3

```
MOV AH,01H          ;load DOS function number

INT 21H             ;access DOS

;returns with AL = ASCII key code
```

TABLE A–3 DOS function calls.

00H	TERMINATE A PROGRAM
Entry	AH = 00H CS = program segment prefix address
Exit	DOS is entered
01H	**READ THE KEYBOARD**
Entry	AH = 01H
Exit	AL = ASCII character
Notes	If AL = 00H, the function call must be invoked again to read an extended ASCII character. Refer to Chapter 1, Table 1–9 for a listing of the extended ASCII keyboard codes. This function call automatically echoes whatever is typed to the video screen.
02H	**WRITE TO STANDARD OUTPUT DEVICE**
Entry	AH = 02H DL = ASCII character to be displayed
Notes	This function call normally displays data on the video display.
03H	**READ CHARACTER FROM COM1**
Entry	AH = 03H
Exit	AL = ASCII character read from the communications port
Notes	This function call reads data from the serial communications port.
04H	**WRITE TO COM1**
Entry	AH = 04H DL = character to be sent out of COM1
Notes	This function transmits data through the serial communications port. The COM port assignment can be changed to use other COM ports with functions 03H and 04H by using the DOS MODE command to reassign COM1 to another COM port.

05H	WRITE TO LPT1
Entry	AH = 05H DL = ASCII character to be printed
Notes	Prints DL on the line printer attached to LPT1. Note that the line printer port can be changed with the DOS MODE command.

06H	DIRECT CONSOLE READ/WRITE
Entry	AH = 06H DL = 0FFH or DL = ASCII character
Exit	AL = ASCII character
Notes	If DL = 0FFH on entry, then this function reads the console. If DL = ASCII character, then this function displays the ASCII character on the console (CON) video screen. If a character is read from the console keyboard, the zero flag (ZF) indicates whether a character was typed. A zero condition indicates that no key was typed, and a not-zero condition indicates that AL contains the ASCII code of the key or a 00H. If AL = 00H, the function must again be invoked to read an extended ASCII character from the keyboard. Note that the key does not echo to the video screen.

07H	DIRECT CONSOLE INPUT WITHOUT ECHO
Entry	AH = 07H
Exit	AL = ASCII character
Notes	This functions exactly as function number 06H with DL = 0FFH, but it will not return from the function until the key is typed.

08H	READ STANDARD INPUT WITHOUT ECHO
Entry	AH = 08H
Exit	AL = ASCII character
Notes	Performs as function 07H, except that it reads the standard input device. The standard input device can be assigned as either the keyboard or the COM port. This function also responds to a control-break, where function 06H and 07H do not. A control-break causes INT 23H to execute. By default, this functions as does function 07H.

09H	DISPLAY A CHARACTER STRING
Entry	AH = 09H DS:DX = address of the character string
Notes	The character string must end with an ASCII $ (24H). The character string can be of any length and may contain control characters such as carriage return (0DH) and line feed (0AH).
0AH	BUFFERED KEYBOARD INPUT
Entry	AH = 0AH DS:DX = address of keyboard input buffer
Notes	The first byte of the buffer contains the size of the buffer (up to 255). The second byte is filled with the number of characters typed upon return. The third byte through the end of the buffer contains the character string typed, followed by a carriage return (0DH). This function continues to read the keyboard (displaying data as typed) until either the specified number of characters are typed or until the enter key is typed.
0BH	TEST STATUS OF THE STANDARD INPUT DEVICE
Entry	AH = 0BH
Exit	AL = status of the input device
Notes	This function tests the standard input device to determine if data are available. If AL = 00, no data are available. If AL = 0FFH, then data are available that must be input using function number 08H.
0CH	CLEAR KEYBOARD BUFFER AND INVOKE KEYBOARD FUNCTION
Entry	AH = 0CH AL = 01H, 06H, 07H, or 0AH
Exit	See exit for functions 01H, 06H, 07H, or 0AH
Notes	The keyboard buffer holds keystrokes while programs execute other tasks. This function empties or clears the buffer and then invokes the keyboard function located in register AL.

USING VISUAL C++

Many of the new examples in the text use Visual C++ Express. Very few, if any, programs are written in assembly language. If assembly language is used, it normally appears in a C++ program to accomplish a special task or to increase the performance of a section of a program.

Not everyone is familiar with the C++ environment so this section has been added to act as a guide in setting up programs that use assembly language within Visual C++. The easiest application type for this is a Forms-based application using the Microsoft CLR (common language runtime).

Create a Dialog Application

Start Visual C++ Express and the screen in Figure A–1 should appear. Click on the word "Project" to the right of "Create" to start a new C++ project. The screen in Figure A–2 should appear. Select a Windows Form Application listed with CLR and give it a unique, but applicable name in the name box, then click OK.

At this point you should see the screen illustrated in Figure A–3, which is the application showing the blank form.

To summarize the process for creating a dialog-based application:

1. Start Visual C++ Express.
2. Click on "Project" to the right of "Create".
3. Select a Windows Form application from CLR and name it, then click on OK.

The screen that appears in Figure A–3 is the forms-based application in the resource editor. At this point you are ready to start placing objects onto the dialog form. The screen you see and the one displayed in the figure may vary somewhat. How it appears is determined by the Tools menu and can be changed by clicking Customize.

The window in the upper left corner of Figure A–3 has tabs above it. These tabs are used to select different pages of the program, such as start page. In Figure A–3 the plus next to Test1 (upper left in the Class View window) is clicked to expand the class. In many of the programs in the text, the Form_Load function is used for setting up the form.
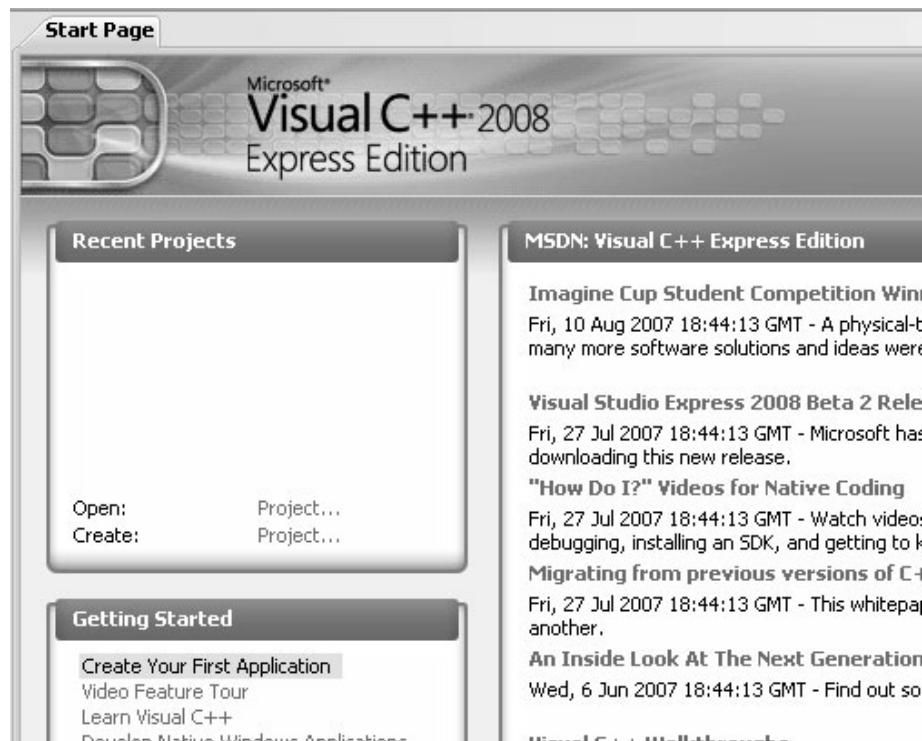
FIGURE A–1   The startup screen of Visual C++ Express.

FIGURE A–2 The new project screen.

FIGURE A–3 The new project design screen.

FIGURE A–4 The Form_Load function.

Figure A–4 illustrates the software in the Form_Load function. To install the Form_Load function, double-click on the form and the Design window will switch to the Code View and display the program with the Form_Load function in place. To view this, double-click on Form_Load in the Class View. The CLR framework calls the Form_Load function before displaying the form on the screen. Any software added for initialization is placed at the bottom of Figure A–4 in the Form_Load function. Once you arrive at this point, you are ready to enter and execute any of the Visual C++ with assembly programs in the textbook.

APPENDIX B

Instruction Set Summary

The instruction set summary contains a complete alphabetical listing of the entire 8086–Pentium 4 instruction set. The coprocessor and MMX instructions are listed in Chapter 14 and are not repeated in this appendix. The SIMD instructions appear at the end of this appendix after the main instruction set summary.

Each instruction entry lists the mnemonic opcode plus a brief description of the purpose of the instruction. Also listed is the binary machine language coding of each instruction and any other data required to form the instruction, such as the displacement or immediate data. Listed to the right of each binary machine language version of the instruction are the flag bits and any change that might occur for the instruction. The flags are described in the following manner: A blank indicates no effect or change; a ? indicates a change with an unpredictable outcome; a * indicates a change with a predictable outcome; a 1 indicates the flag is set; and a 0 indicates that the flag is cleared. If the flag bits ODITSZAPC are not illustrated with an instruction, the instruction does not modify any of these flags.

Before the instruction listing begins, some information about the bit settings in binary machine language versions of the instructions is presented. Table B–1 lists the modifier bits, coded as OO in the instruction listing.

Table B–2 lists the memory-addressing modes available using a register field coding of mmm. This table applies to all versions of the microprocessor, as long as the operating mode is 16 bits.

Table B–3 lists the register selections provided by the rrr field in an instruction. This table includes the register selections for 8-, 16-, and 32-bit registers.

Table B–4 lists the segment register bit assignment (rrr) found with the MOV, PUSH, and POP instructions.

TABLE B–1 The modifier bits, coded as oo in the instruction listing.

oo	Function
00	If mmm = 110, a displacement follows the opcode; otherwise no displacement is used.
01	An 8-bit signed displacement flows the opcode.
10	A 16- or 32-bit signed displacement follows the opcode.
11	mmm specifies a register instead of an addressing mode.

TABLE B–2   The 16-bit register/memory (mmm) field description.

mmm	16-Bit
000	DS:[BX+SI]
001	DS:[BX+DI]
010	SS:[BP+SI]
011	SS:[BP+DI]
100	DS:[SI]
101	DS:[DI]
110	SS:[BP]
111	DS:[BX]

TABLE B–3   The register (rrr) field.

rrr	W=0	W=1 (16-Bit)	W=1 (32-Bit)
000	AL	AX	EAX
001	CL	CX	ECX
010	DL	DX	EDX
011	BL	BX	EBX
100	AH	SP	ESP
101	CH	BP	EBP
110	DH	SI	ESI
111	BH	DI	EDI

TABLE B–4   Register field assignments (rrr) for the segment registers.

rrr	Segment Register
000	ES
001	CS
010	SS
011	DS
100	FS
101	GS

When the 80386–Core2 are used, some of the definitions provided in Table B–1 through B–3 change. See Tables B–5 and B–6 for these changes as they apply to the 80386–Core2 microprocessors.

TABLE B–5   Index registers specified by rrr when the 80386–Core2 are operated in 32-bit mode.

rrr	Index Register
000	DS:[EAX]
001	DS:[ECX]
010	DS:[EDX]
011	DS:[EBX]
100	(see Table B–6)
101	SS:[EBP]
110	DS:[ESI]
111	DS:[EDI]

TABLE B–6 Possible combinations of oo, mmm, and rrr for the 80386–Core2 microprocessors using 32-bit addressing.

oo	mmm	rrr (Base in Scaled-Index Byte)	Addressing Mode
00	000	—	DS:[EAX]
00	001	—	DS:[ECX]
00	010	—	DS:[EDX]
00	011	—	DS:[EBX]
00	100	000	DS:[EAX+scaled index]
00	100	001	DS:[ECX+scaled index]
00	100	010	DS:[EDX+scaled index]
00	100	011	DS:[EBX+scaled index]
00	100	100	SS:[ESP+scaled index]
00	100	101	DS:[disp32+scaled index]
00	100	110	DS:[ESI+scaled index]
00	100	111	DS:[EDI+scaled index]
00	101	—	DS:disp32
00	110	—	DS:[ESI]
00	111	—	DS:[EDI]
01	000	—	DS:[EAX+disp8]
01	001	—	DS:[ECX+disp8]
01	010	—	DS:[EDX+disp8]
01	011	—	DS:[EBX+disp8]
01	100	000	DS:[EAX+scaled index+disp8]
01	100	001	DS:[ECX+scaled index+disp8]
01	100	010	DS:[EDX+scaled index+disp8]
01	100	011	DS:[EBX+scaled index+disp8]
01	100	100	SS:[ESP+scaled index+disp8]
01	100	101	SS:[EBP+scaled index+disp8]
01	100	110	DS:[ESI+scaled index+disp8]
01	100	111	DS:[EDI+scaled index+disp8]
01	101	—	SS:[EBP+disp8]
01	110	—	DS:[ESI+disp8]
01	111	—	DS:[EDI+disp8]
10	000	—	DS:[EAX+disp32]
10	001	—	DS:[ECX+disp32]
10	010	—	DS:[EDX+disp32]
10	011	—	DS:[EBX+disp32]
10	100	000	DS:[EAX+scaled index+disp32]
10	100	001	DS:[ECX+scaled index+disp32]
10	100	010	DS:[EDX+scaled index+disp32]
10	100	011	DS:[EBX+scaled index+disp32]
10	100	100	SS:[ESP+scaled index+disp32]
10	100	101	SS:[EBP+scaled index+disp32]
10	100	110	DS:[ESI+scaled index+disp32]
10	100	111	DS:[EDI+scaled index+disp32]
10	101	—	SS:[EBP+disp32]
10	110	—	DS:[ESI+disp32]
10	111	—	DS:[EDI+disp32]

Note: disp8 = 8-bit displacement and disp32 = 32-bit displacement.

In order to use the scaled-index addressing modes listed in Table B–6, code oo and mmm in the second byte of the opcode. The scaled-index byte is usually the third byte and contains three fields. The leftmost two bits determine the scaling factor ($00 = \times1$, $01 = \times2$, $10 = \times4$, or $11 = \times8$). The next three bits toward the right contain the scaled-index register number (this is obtained from Table B–5). The rightmost three bits are from the mmm field listed in Table B–6. For example, the MOV AL,[EBX+2*ECX] instruction has a scaled-index byte of 01 001 011, where $01 = X_2$, $001 = ECX$, and $011 = EBX$.

Some instructions are prefixed to change the default segment or to override the instruction mode. Table B–7 lists the segment and instruction mode override prefixes with append at the beginning of an instruction if they are used to form the instruction. For example, the MOV AL,ES:[BX] instruction used the extra segment because of the override prefix ES:.

In the 8086 and 8088 microprocessors, the effective address calculation required additional clocks that are added to the times in the instruction set summary. These additional times are listed in Table B–8. No such times are added to the 80286–Core2. Note that the instruction set summary does not include clock times for the Pentium Pro through the Core2. Intel has not released these times and has decided that the RDTSC instruction can be used to have the microprocessor count the number of clocks required for a given application. Even though the timings do not appear for these new microprocessors, they are very similar to the Pentium, which can be used as a guide.

TABLE B–7 Override prefixes.

Prefix Byte	Purpose
26H	ES: segment override
2EH	CS: segment override
36H	SS: segment override
3EH	DS: segment override
64H	FS: segment override
65H	GS: segment override
66H	Memory operand instruction mode override
67H	Register operand instruction mode override

TABLE B–8 Effective address calculations for the 8086 and 8088 microprocessors.

Type	Clocks	Example Instruction
Base or index	5	MOV CL,[DI]
Displacement	3	MOV AL,DATA1
Base plus index	7	MOV AL,[BP+SI]
Displacement plus base or index	9	MOV DH,[DI+20H]
Base plus index plus displacement	11	MOV CL,[BX+DI+2]
Segment override	ea + 2	MOV AL,ED:[DI]

INSTRUCTION SET SUMMARY

AAA	ASCII adjust AL after addition		

| 00110111 | | O D I T   S Z A P C | |
| | | ?   ? ? * ? * | |
Example		Microprocessor	Clocks
AAA		8086	8
		8088	8
		80286	3
		80386	4
		80486	3
		Pentium–Core2	

AAD	ASCII adjust AX before division		

| 11010101 00001010 | | O D I T   S Z A P C | |
| | | ?   * * ? * ? | |
Example		Microprocessor	Clocks
AAD		8086	60
		8088	60
		80286	14
		80386	19
		80486	14
		Pentium–Core2	10

AAM	ASCII adjust AX after multiplication		

| 11010100 00001010 | | O D I T   S Z A P C | |
| | | ?   * * ? * ? | |
Example		Microprocessor	Clocks
AAM		8086	83
		8088	83
		80286	16
		80386	17
		80486	15
		Pentium–Core2	18

AAS	ASCII adjust AL after subtraction		

| 00111111 | | O D I T S Z A P C | |
| | | ? ? ? * ? * | |
Example		Microprocessor	Clocks
AAS		8086	8
		8088	8
		80286	3
		80386	4
		80486	3
		Pentium–Core2	3

ADC	Addition with carry		

| 000100dw oorrrmmm disp | | O D I T S Z A P C | |
| | | * * * * * * | |
Format	Examples	Microprocessor	Clocks
ADC reg,reg	ADC AX,BX	8086	3
	ADC AL,BL	8088	3
	ADC EAX,EBX	80286	3
	ADC CX,SI	80386	3
	ADC ESI,EDI	80486	1
		Pentium–Core2	1 or 3
ADC mem,reg	ADC DATAY,AL	8086	16 + ea
	ADC LIST,SI	8088	24 + ea
	ADC DATA2[DI],CL	80286	7
	ADC [EAX],BL	80386	7
	ADC [EBX+2*ECX],EDX	80486	3
		Pentium–Core2	1 or 3

ADC reg,mem	ADC BL,DATA1 ADC SI,LIST1 ADC CL,DATA2[SI] ADC CX,[ESI] ADC ESI,[2*ECX]	8086	9 + ea
		8088	13 + ea
		80286	7
		80386	6
		80486	2
		Pentium–Core2	1 or 2

100000sw oo010mmm disp data			
Format	Examples	Microprocessor	Clocks
ADC reg,imm	ADC CX,3 ADC DI,1AH ADC DL,34H ADC EAX,12345 ADC CX,1234H	8086	4
		8088	4
		80286	3
		80386	2
		80486	1
		Pentium–Core2	1 or 3
ADC mem,imm	ADC DATA4,33 ADC LIST,'A' ADC DATA3[DI],2 ADC BYTE PTR[EBX],3 ADC WORD PTR[DI],669H	8086	17 + ea
		8088	23 + ea
		80286	7
		80386	7
		80486	3
		Pentium–Core2	1 or 3
ADC acc,imm	ADC AX,3 ADC AL,1AH ADC AH,34H ADC EAX,2 ADC AL,'Z'	8086	4
		8088	4
		80286	3
		80386	2
		80486	1
		Pentium–Core2	1

ADD	Addition		

000000dw oorrrmmm disp		O D I T S Z A P C * * * * * *	
Format	Examples	Microprocessor	Clocks

ADD reg,reg	ADD AX,BX
	ADD AL,BL
	ADD EAX,EBX
	ADD CX,SI
	ADD ESI,EDI
		8086	3

		8088	3
		80286	2
		80386	2
		80486	1
		Pentium–Core2	1 or 3

ADD mem,reg	ADD DATAY,AL
	ADD LIST,SI
	ADD DATA6[DI],CL
	ADD [EAX],CL
	ADD [EDX+4*ECX],EBX
		8086	16 + ea

		8088	24 + ea
		80286	7
		80386	7
		80486	3
		Pentium–Core2	1 or 3

ADD reg,mem	ADD BL,DATA2
	ADD SI,LIST3
	ADD CL,DATA2[DI]
	ADD CX,[EDI]
	ADD ESI,[ECX+200H]
		8086	9 + ea

		8088	13 + ea
		80286	7
		80386	6
		80486	2
		Pentium–Core2	1 or 2

100000sw oo000mmm disp data			
Format	Examples	Microprocessor	Clocks

ADD reg,imm	ADD CX,3
	ADD DI,1AH
	ADD DL,34H
	ADD EDX,1345H
	ADD CX,1834H
		8086	4

		8088	4
		80286	3
		80386	2
		80486	1
		Pentium–Core2	1 or 3

ADD mem,imm	ADD DATA4,33
	ADD LIST,'A'
	ADD DATA3[DI],2
	ADD BYTE PTR[EBX],3
	ADD WORD PTR[DI],669H
		8086	17 + ea

		8088	23 + ea
		80286	7
		80386	7
		80486	3
		Pentium–Core2	1 or 3

ADD acc,imm	ADD AX,3
	ADD AL,1AH
	ADD AH,34H
	ADD EAX,2
	ADD AL,'Z'
		8086	4

		8088	4
		80286	3
		80386	2
		80486	1
		Pentium–Core2	1

AND   Logical AND

001000dw oorrrmmm disp		O D I T S Z A P C 0 * * ? * 0	
Format	Examples	Microprocessor	Clocks
AND reg,reg	AND CX,BX AND DL,BL AND ECX,EBX AND BP,SI AND EDX,EDI	8086	3
		8088	3
		80286	2
		80386	2
		80486	1
		Pentium–Core2	1 or 3
AND mem,reg	AND BIT,AL AND LIST,DI AND DATAZ[BX],CL AND [EAX],BL AND [ESI+4*ECX],EDX	8086	16 + ea
		8088	24 + ea
		80286	7
		80386	7
		80486	3
		Pentium–Core2	1 or 3

AND reg,mem	AND BL,DATAW AND SI,LIST AND CL,DATAQ[SI] AND CX,[EAX] AND ESI,[ECX+43H]	8086	9 + ea
		8088	13 + ea
		80286	7
		80386	6
		80486	2
		Pentium–Core2	1 or 2

100000sw oo100mmm disp data			
Format	Examples	Microprocessor	Clocks
AND reg,imm	AND BP,1 AND DI,10H AND DL,34H AND EBP,1345H AND SP,1834H	8086	4
		8088	4
		80286	3
		80386	2
		80486	1
		Pentium–Core2	1 or 3
AND mem,imm	AND DATA4,33 AND LIST,'A' AND DATA3[DI],2 AND BYTE PTR[EBX],3 AND DWORD PTR[DI],66H	8086	17 + ea
		8088	23 + ea
		80286	7
		80386	7
		80486	3
		Pentium–Core2	1 or 3
AND acc,imm	AND AX,3 AND AL,1AH AND AH,34H AND EAX,2 AND AL,'r'	8086	4
		8088	4
		80286	3
		80386	2
		80486	1
		Pentium–Core2	1

| **ARPL** | Adjust requested privilege level | | |

01100011 oorrrmmm disp		O D I T S Z A P C *	
Format	Examples	Microprocessor	Clocks
ARPL reg,reg	ARPL AX,BX ARPL BX,SI ARPL AX,DX ARPL BX,AX ARPL SI,DI	8086	—
		8088	—
		80286	10
		80386	20
		80486	9
		Pentium–Core2	7
ARPL mem,reg	ARPL DATAY,AX ARPL LIST,DI ARPL DATA3[DI],CX ARPL [EBX],AX ARPL [EDX+4*ECX],BP	8086	—
		8088	—
		80286	11
		80386	21
		80486	9
		Pentium–Core2	7

| **BOUND** | Check array against boundary | | |

01100010 oorrrmmm disp			
Format	Examples	Microprocessor	Clocks
BOUND reg,mem	BOUND AX,BETS BOUND BP,LISTG BOUND CX,DATAX BOUND BX,[DI] BOUND SI,[BX+2]	8086	—
		8088	—
		80286	13
		80386	10
		80486	7
		Pentium–Core2	8

BSF	Bit scan forward			

00001111 10111100 oorrrmmm disp		O D I T S Z A P C ? ? * ? ? ?		
Format	Examples	Microprocessor		Clocks
BSF reg,reg	BSF AX,BX BSF BX,SI BSF EAX,EDX BSF EBX,EAX BSF SI,DI	8086		—
		8088		—
		80286		—
		80386		10 + 3n
		80486		6–42
		Pentium–Core2		6–42
BSF reg,mem	BSF AX,DATAY BSF SI,LIST BSF CX,DATA3[DI] BSF EAX,[EBX] BSF EBP,[EDX+4*ECX]	8086		—
		8088		—
		80286		—
		80386		10 + 3n
		80486		7–43
		Pentium–Core2		6–43

BSR	Bit scan reverse			

00001111 10111101 oorrrmmm disp		O D I T S Z A P C ? ? * ? ? ?		
Format	Examples	Microprocessor		Clocks
BSR reg,reg	BSR AX,BX BSR BX,SI BSR EAX,EDX BSR EBX,EAX BSR SI,DI	8086		—
		8088		—
		80286		—
		80386		10 + 3n
		80486		6–103
		Pentium–Core2		7–71

BSR reg,mem	BSR AX,DATAY BSR SI,LIST BSR CX,DATA3[DI] BSR EAX,[EBX] BSR EBP,[EDX+4*ECX]	8086	—
		8088	—
		80286	—
		80386	10 + 3n
		80486	7–104
		Pentium–Core2	7–72

BSWAP   Byte swap

00001111 11001rrr

Format	Examples	Microprocessor	Clocks
BSWAP reg32	BSWAP EAX BSWAP EBX BSWAP EDX BSWAP ECX BSWAP ESI	8086	—
		8088	—
		80286	—
		80386	—
		80486	1
		Pentium–Core2	1

BT   Bit test

00001111 10111010 oo100mmm disp data

O D I T   S Z A P C
                            *

Format	Examples	Microprocessor	Clocks
BT reg,imm8	BT AX,2 BT CX,4 BT BP,10H BT CX,8 BT BX,2	8086	—
		8088	—
		80286	—
		80386	3
		80486	3
		Pentium–Core2	4

BT mem,imm8	BT DATA1,2 BT LIST,2 BT DATA2[DI],3 BT [EAX],1 BT FROG,6	8086	—
		8088	—
		80286	—
		80386	6
		80486	3
		Pentium–Core2	4

00001111 10100011 disp			
Format	Examples	Microprocessor	Clocks
BT reg,reg	BT AX,CX BT CX,DX BT BP,AX BT SI,CX BT EAX,EBX	8086	—
		8088	—
		80286	—
		80386	3
		80486	3
		Pentium–Core2	4 or 9
BT mem,reg	BT DATA4,AX BT LIST,BX BT DATA3[DI],CX BT [EBX],DX BT [DI],DI	8086	—
		8088	—
		80286	—
		80386	12
		80486	8
		Pentium–Core2	4 or 9

BTC	Bit test and complement		

00001111 10111010 oo111mmm disp data		O D I T S Z A P C *	
Format	Examples	Microprocessor	Clocks
BTC reg,imm8	BTC AX,2 BTC CX,4 BTC BP,10H BTC CX,8 BTC BX,2	8086	—
		8088	—
		80286	—
		80386	6
		80486	6
		Pentium–Core2	7 or 8
BTC mem,imm8	BTC DATA1,2 BTC LIST,2 BTC DATA2[DI],3 BTC [EAX],1 BTC FROG,6	8086	—
		8088	—
		80286	—
		80386	7 or 8
		80486	8
		Pentium–Core2	8

00001111 10111011 disp			
Format	Examples	Microprocessor	Clocks
BTC reg,reg	BTC AX,CX BTC CX,DX BTC BP,AX BTC SI,CX BTC EAX,EBX	8086	—
		8088	—
		80286	—
		80386	6
		80486	6
		Pentium–Core2	7 or 13
BTC mem,reg	BTC DATA4,AX BTC LIST,BX BTC DATA3[DI],CX BTC [EBX],DX BTC [DI],DI	8086	—
		8088	—
		80286	—
		80386	13
		80486	13
		Pentium–Core2	7 or 13

BTR	Bit test and reset		

00001111 10111010 oo110mmm disp data		O D I T S Z A P C *	
Format	Examples	Microprocessor	Clocks
BTR reg,imm8	BTR AX,2 BTR CX,4 BTR BP,10H BTR CX,8 BTR BX,2	8086	—
		8088	—
		80286	—
		80386	6
		80486	6
		Pentium–Core2	7 or 8
BTR mem,imm8	BTR DATA1,2 BTR LIST,2 BTR DATA2[DI],3 BTR [EAX],1 BTR FROG,6	8086	—
		8088	—
		80286	—
		80386	8
		80486	8
		Pentium–Core2	7 or 8

00001111 10110011 disp			
Format	Examples	Microprocessor	Clocks
BTR reg,reg	BTR AX,CX BTR CX,DX BTR BP,AX BTR SI,CX BTR EAX,EBX	8086	—
		8088	—
		80286	—
		80386	6
		80486	6
		Pentium–Core2	7 or 13
BTR mem,reg	BTR DATA4,AX BTR LIST,BX BTR DATA3[DI],CX BTR [EBX],DX BTR [DI],DI BTC [DI],DI	8086	—
		8088	—
		80286	—
		80386	13
		80486	13
		Pentium–Core2	7 or 13

BTS	Bit test and set		
00001111 10111010 oo101mmm disp data		O D I T S Z A P C *	
Format	Examples	Microprocessor	Clocks
BTS reg,imm8	BTS AX,2 BTS CX,4 BTS BP,10H BTS CX,8 BTS BX,2	8086	—
		8088	—
		80286	—
		80386	6
		80486	6
		Pentium–Core2	7 or 8
BTS mem,imm8	BTS DATA1,2 BTS LIST,2 BTS DATA2[DI],3 BTS [EAX],1 BTS FROG,6	8086	—
		8088	—
		80286	—
		80386	8
		80486	8
		Pentium–Core2	7 or 8
00001111 10101011 disp			
Format	Examples	Microprocessor	Clocks
BTS reg,reg	BTS AX,CX BTS CX,DX BTS BP,AX BTS SI,CX BTS EAX,EBX	8086	—
		8088	—
		80286	—
		80386	6
		80486	6
		Pentium–Core2	7 or 13
BTS mem,reg	BTS DATA4,AX BTS LIST,BX BTS DATA3[DI],CX BTS [EBX],DX BTS [DI],DI	8086	—
		8088	—
		80286	—
		80386	13
		80486	13
		Pentium–Core2	7 or 13

CALL	Call procedure (subroutine)		

11101000 disp Format	Examples	Microprocessor	Clocks
CALL label (near)	CALL FOR_FUN CALL HOME CALL ET CALL WAITING CALL SOMEONE	8086	19
		8088	23
		80286	7
		80386	3
		80486	3
		Pentium–Core2	1

10011010 disp Format	Examples	Microprocessor	Clocks
CALL label (far)	CALL FAR PTR DATES CALL WHAT CALL WHERE CALL FARCE CALL WHOM	8086	28
		8088	36
		80286	13
		80386	17
		80486	18
		Pentium–Core2	4

11111111 oo010mmm Format	Examples	Microprocessor	Clocks
CALL reg (near)	CALL AX CALL BX CALL CX CALL DI CALL SI	8086	16
		8088	20
		80286	7
		80386	7
		80486	5
		Pentium–Core2	2

CALL mem (near)	CALL ADDRESS CALL NEAR PTR [DI] CALL DATA1 CALL FROG CALL ME_NOW	8086	21 + ea
		8088	29 + ea
		80286	11
		80386	10
		80486	5
		Pentium–Core2	2

11111111 oo011mmm Format	Examples	Microprocessor	Clocks
CALL mem (far)	CALL FAR_LIST[SI] CALL FROM_HERE CALL TO_THERE CALL SIXX CALL OCT	8086	16
		8088	20
		80286	7
		80386	7
		80486	5
		Pentium–Core2	2

CBW   Convert byte to word (AL $\Rightarrow$ AX)

10011000 Example	Microprocessor	Clocks
CBW	8086	2
	8088	2
	80286	2
	80386	3
	80486	3
	Pentium–Core2	3

CDQ	Convert doubleword to quadword (EAX ⇒ EDX:EAX)	

11010100 00001010		
Example	Microprocessor	Clocks
CDQ	8086	—
	8088	—
	80286	—
	80386	2
	80486	2
	Pentium–Core2	2

CLC	Clear carry flag	

11111000	O D I T S Z A P C 0	
Example	Microprocessor	Clocks
CLC	8086	2
	8088	2
	80286	2
	80386	2
	80486	2
	Pentium–Core2	2

CLD	Clear direction flag	

11111100	O D I T S Z A P C 0	
Example	Microprocessor	Clocks
CLD	8086	2
	8088	2
	80286	2
	80386	2
	80486	2
	Pentium–Core2	2

APPENDIX B — page 814

CLI — Clear interrupt flag

11111010   O D I T S Z A P C (I=0)

Example CLI	Microprocessor	Clocks
	8086	2
	8088	2
	80286	3
	80386	3
	80486	5
	Pentium–Core2	7

CLTS — Clear task switched flag (CR0)

00001111 00000110

Example CLTS	Microprocessor	Clocks
	8086	—
	8088	—
	80286	2
	80386	5
	80486	7
	Pentium–Core2	10

CMC — Complement carry flag

10011000   O D I T S Z A P C (C=*)

Example CMC	Microprocessor	Clocks
	8086	2
	8088	2
	80286	2
	80386	2
	80486	2
	Pentium–Core2	2

CMOVcondition	Conditional move		

00001111 0100cccc oorrrmmm			
Format	Examples	Microprocessor	Clocks
CMOVcc reg,mem	CMOVNZ AX,FROG CMOVC EAX,[EDI] CMOVNC BX,DATA1 CMOVP EBX,WAITING CMOVNE DI,[SI]	8086	—
		8088	—
		80286	—
		80386	—
		80486	—
		Pentium–Core2	—

Condition Codes	Mnemonic	Flag	Description
0000	CMOVO	O = 1	Move if overflow
0001	CMOVNO	O = 0	Move if no overflow
0010	CMOVB	C = 1	Move if below
0011	CMOVAE	C = 0	Move if above or equal
0100	CMOVE	Z = 1	Move if equal/zero
0101	CMOVNE	Z = 0	Move if not equal/zero
0110	CMOVBE	C = 1 + Z = 1	Move if below or equal
0111	CMOVA	C = 0 • Z = 0	Move if above
1000	CMOVS	S = 1	Move if sign
1001	CMOVNS	S = 0	Move if no sign
1010	CMOVP	P = 1	Move if parity
1011	CMOVNP	P = 0	Move if no parity
1100	CMOVL	S • O	Move if less than
1101	CMOVGE	S = 0	Move if greater than or equal
1110	CMOVLE	Z = 1 + S • O	Move if less than or equal
1111	CMOVG	Z = 0 + S = O	Move if greater than

CMP	Compare		

001110dw oorrrmmm disp		O D I T S Z A P C * * * * * *	
Format	Examples	Microprocessor	Clocks

CMP reg,reg	CMP AX,BX
	CMP AL,BL
	CMP EAX,EBX
	CMP CX,SI
	CMP ESI,EDI
		8086	3

		8088	3
		80286	2
		80386	2
		80486	1
		Pentium–Core2	1 or 2

CMP mem,reg	CMP DATAY,AL
	CMP LIST,SI
	CMP DATA6[DI],CL
	CMP [EAX],CL
	CMP [EDX+4*ECX],EBX
		8086	9 + ea

		8088	13 + ea
		80286	7
		80386	5
		80486	2
		Pentium–Core2	1 or 2

CMP reg,mem	CMP BL,DATA2
	CMP SI,LIST3
	CMP CL,DATA2[DI]
	CMP CX,[EDI]
	CMP ESI,[ECX+200H]
		8086	9 + ea

		8088	13 + ea
		80286	6
		80386	6
		80486	2
		Pentium–Core2	1 or 2

100000sw oo111mmm disp data

Format	Examples	Microprocessor	Clocks

CMP reg,imm	CMP CX,3
	CMP DI,1AH
	CMP DL,34H
	CMP EDX,1345H
	CMP CX,1834H
		8086	4

		8088	4
		80286	3
		80386	2
		80486	1
		Pentium–Core2	1 or 2

CMP mem,imm	CMP DATAS,3
	CMP BYTE PTR[EDI],1AH
	CMP DADDY,34H
	CMP LIST,'A'
	CMP TOAD,1834H
		8086	10 + ea

		8088	14 + ea
		80286	6
		80386	5
		80486	2
		Pentium–Core2	1 or 2

0001111w data			
Format	Examples	Microprocessor	Clocks

CMP acc,imm	CMP AX,3
	CMP AL,1AH
	CMP AH,34H
	CMP EAX,1345H
	CMP AL,'Y'
		8086	4

		8088	4
		80286	3
		80386	2
		80486	1
		Pentium–Core2	1

CMPS   Compare strings

1010011w		O D I T S Z A P C * * * * * *	
Format	Examples	Microprocessor	Clocks
CMPSB CMPSW CMPSD	CMPSB CMPSW CMPSD CMPSB DATA1,DATA2 REPE CMPSB REPNE CMPSW	8086	32
		8088	30
		80286	8
		80386	10
		80486	8
		Pentium–Core2	5

CMPXCHG   Compare and exchange

00001111 1011000w 11rrrrrr		O D I T S Z A P C * * * * * *	
Format	Examples	Microprocessor	Clocks
CMPXCHG reg,reg	CMPXCHG EAX,EBX CMPXCHG ECX,EDX	8086	—
		8088	—
		80286	—
		80386	—
		80486	6
		Pentium–Core2	6

| 0001111w data | | | |
Format	Examples	Microprocessor	Clocks
CMPXCHG mem,reg	CMPXCHG DATAD,EAX CMPXCHG DATA2,EDI	8086	—
		8088	—
		80286	—
		80386	—
		80486	7
		Pentium–Core2	6

CMPXCHG8B   Compare and exchange 8 bytes

| 00001111 11000111 oorrrmmm | | O D I T   S Z A P C
                        * | |
Format	Examples	Microprocessor	Clocks
CMPXCHG8B mem64	CMPXCHG8B DATA3	8086	—
		8088	—
		80286	—
		80386	—
		80486	—
		Pentium–Core2	10

CPUID   CPU identification code

| 00001111 10100010 | | | |
Example		Microprocessor	Clocks
CPUID		8086	—
		8088	—
		80286	—
		80386	—
		80486	—
		Pentium–Core2	14

CWD	Convert word to doubleword (AX $\Rightarrow$ DX:AX)		
10011000 Example		Microprocessor	Clocks
CWD		8086	5
		8088	5
		80286	2
		80386	2
		80486	3
		Pentium–Core2	2

CWDE	Convert word to extended doubleword (AX $\Rightarrow$ EAX)		
10011000 Example		Microprocessor	Clocks
CWDE		8086	—
		8088	—
		80286	—
		80386	3
		80486	3
		Pentium–Core2	3

DAA	Decimal adjust AL after addition		
00100111 Example		O D I T S Z A P C ? * * * * * Microprocessor	Clocks
DAA		8086	4
		8088	4
		80286	3
		80386	4
		80486	2
		Pentium–Core2	3

DAS	Decimal adjust AL after subtraction		

| 00101111 | | O D I T S Z A P C
?   * * * * * | |
Example		Microprocessor	Clocks
DAS		8086	4
		8088	4
		80286	3
		80386	4
		80486	2
		Pentium–Core2	3

DEC	Decrement		

| 1111111w oo001mmm disp | | O D I T S Z A P C
*   * * * * | |
Format	Examples	Microprocessor	Clocks
DEC reg8	DEC BL	8086	3
	DEC BH	8088	3
	DEC CL		
	DEC DH	80286	2
	DEC AH	80386	2
		80486	1
		Pentium–Core2	1 or 3
DEC mem	DEC DATAY	8086	15 + ea
	DEC LIST		
	DEC DATA6[DI]	8088	23 + ea
	DEC BYTE PTR [BX]	80286	7
	DEC WORD PTR [EBX]		
		80386	6
		80486	3
		Pentium–Core2	1 or 3

01001rrr			
Format	Examples	Microprocessor	Clocks
DEC reg16 DEC reg32	DEC CX DEC DI DEC EDX DEC ECX DEC BP	8086	3
		8088	3
		80286	2
		80386	2
		80486	1
		Pentium–Core2	1

DIV   Divide

1111011w oo110mmm disp		O D I T S Z A P C ? ? ? ? ? ?	
Format	Examples	Microprocessor	Clocks
DIV reg	DIV BL DIV BH DIV ECX DIV DH DIV CX	8086	162
		8088	162
		80286	22
		80386	38
		80486	40
		Pentium–Core2	17–41
DIV mem	DIV DATAY DIV LIST DIV DATA6[DI] DIV BYTE PTR [BX] DIV WORD PTR [EBX]	8086	168
		8088	176
		80286	25
		80386	41
		80486	40
		Pentium–Core2	17–41

ENTER	Create a stack frame		

11001000 data Format	Examples	Microprocessor	Clocks
ENTER imm,0	ENTER 4,0 ENTER 8,0 ENTER 100,0 ENTER 200,0 ENTER 1024,0	8086	—
		8088	—
		80286	11
		80386	10
		80486	14
		Pentium–Core2	11
ENTER imm,1	ENTER 4,1 ENTER 10,1	8086	—
		8088	—
		80286	12
		80386	15
		80486	17
		Pentium–Core2	15
ENTER imm,imm	ENTER 3,6 ENTER 100,3	8086	—
		8088	—
		80286	12
		80386	15
		80486	17
		Pentium–Core2	15 + 2n

ESC	Escape (obsolete–see coprocessor)

HLT	Halt		

11110100 Example		Microprocessor	Clocks
HLT		8086	2
		8088	2
		80286	2
		80386	5
		80486	4
		Pentium–Core2	varies

IDIV	Integer (signed) division		

1111011w oo111mmm disp		O D I T S Z A P C ? ? ? ? ? ?	
Format	Examples	Microprocessor	Clocks
IDIV reg	IDIV BL IDIV BH IDIV ECX IDIV DH IDIV CX	8086	184
		8088	184
		80286	25
		80386	43
		80486	43
		Pentium–Core2	22–46
IDIV mem	IDIV DATAY IDIV LIST IDIV DATA6[DI] IDIV BYTE PTR [BX] IDIV WORD PTR [EBX]	8086	190
		8088	194
		80286	28
		80386	46
		80486	44
		Pentium–Core2	22–46

IMUL	Integer (signed) multiplication

1111011w oo101mmm disp		O D I T S Z A P C * ? ? ? ? *

Format	Examples	Microprocessor	Clocks
IMUL reg	IMUL BL IMUL CX IMUL ECX IMUL DH IMUL AL	8086	154
		8088	154
		80286	21
		80386	38
		80486	42
		Pentium–Core2	10–11
IMUL mem	IMUL DATAY IMUL LIST IMUL DATA6[DI] IMUL BYTE PTR [BX] IMUL WORD PTR [EBX]	8086	160
		8088	164
		80286	24
		80386	41
		80486	42
		Pentium–Core2	10–11

011010s1 oorrmmm disp data			
Format	Examples	Microprocessor	Clocks
IMUL reg,imm	IMUL CX,16 IMUL DI,100 IMUL EDX,20	8086	—
		8088	—
		80286	21
		80386	38
		80486	42
		Pentium–Core2	10
IMUL reg,reg,imm	IMUL DX,AX,2 IMUL CX,DX,3 IMUL BX,AX,33	8086	—
		8088	—
		80286	21
		80386	38
		80486	42
		Pentium–Core2	10

IMUL reg,mem,imm	IMUL CX,DATAY,99	8086	—
		8088	—
		80286	24
		80386	38
		80486	42
		Pentium–Core2	10

00001111 10101111 oorrmmm disp Format	Examples	Microprocessor	Clocks
IMUL reg,reg	IMUL CX,DX IMUL DI,BX IMUL EDX,EBX	8086	—
		8088	—
		80286	—
		80386	38
		80486	42
		Pentium–Core2	10
IMUL reg,mem	IMUL DX,DATAY IMUL CX,LIST IMUL ECX,DATA6[DI]	8086	—
		8088	—
		80286	—
		80386	41
		80486	42
		Pentium–Core2	10

IN   Input data from port

1110010w port# Format	Examples	Microprocessor	Clocks
IN acc,pt	IN AL,12H IN AX,12H IN AL,0FFH IN AX,0A0H IN EAX,10H	8086	10
		8088	14
		80286	5
		80386	12
		80486	14
		Pentium–Core2	7

1110110w			
Format	Examples	Microprocessor	Clocks
IN acc,DX	IN AL,DX IN AX,DX IN EAX,DX	8086	8
		8088	12
		80286	5
		80386	13
		80486	14
		Pentium–Core2	7

INC   Increment

1111111w oo000mmm disp		O D I T S Z A P C * * * * *	
Format	Examples	Microprocessor	Clocks
INC reg8	INC BL INC BH INC AL INC AH INC DH	8086	3
		8088	3
		80286	2
		80386	2
		80486	1
		Pentium–Core2	1 or 3
INC mem	INC DATA3 INC LIST INC COUNT INC BYTE PTR [DI] INC WORD PTR [ECX]	8086	15 + ea
		8088	23 + ea
		80286	7
		80386	6
		80486	3
		Pentium–Core2	1 or 3
INC reg16 INC reg32	INC CX INC DX INC BP INC ECX INC ESP	8086	3
		8088	3
		80286	2
		80386	2
		80486	1
		Pentium–Core2	1

INS	Input string from port		

0110110w			
Format	Examples	Microprocessor	Clocks
INSB	INSB	8086	—
INSW	INSW		
INSD	INSD	8088	—
	INS DATA2	80286	5
	REP INSB	80386	15
		80486	17
		Pentium–Core2	9

INT	Interrupt		

11001101 type			
Format	Examples	Microprocessor	Clocks
INT type	INT12H	8086	51
	INT15H	8088	71
	INT 21H		
	INT 2FH	80286	23
	INT 10H	80386	37
		80486	30
		Pentium–Core2	16–82

INT 3	Interrupt 3		

11001100			
Example		Microprocessor	Clocks
INT 3		8086	52
		8088	72
		80286	23
		80386	33
		80486	26
		Pentium–Core2	13–56

INTO	Interrupt on overflow		
11001110 Example		Microprocessor	Clocks
INTO		8086	53
		8088	73
		80286	24
		80386	35
		80486	28
		Pentium–Core2	13–56

INVD	Invalidate data cache		
00001111 00001000 Example		Microprocessor	Clocks
INTVD		8086	—
		8088	—
		80286	—
		80386	—
		80486	4
		Pentium–Core2	15

IRET/IRETD	Return from interrupt		
11001101 data		O D I T S Z A P C * * * * * * * * *	
Format	Examples	Microprocessor	Clocks
IRET IRETD	IRET IRETD IRET 100	8086	32
		8088	44
		80286	17
		80386	22
		80486	15
		Pentium–Core2	8–27

Jcondition   Conditional jump

0111cccc disp

Format	Examples		Microprocessor	Clocks
Jcnd label (8-bit disp)	JA ABOVE JB BELOW JG GREATER JE EQUAL JZ ZERO		8086	16/4
			8088	16/4
			80286	7/3
			80386	7/3
			80486	3/1
			Pentium–Core2	1

00001111 1000cccc disp

Format	Examples	Microprocessor	Clocks
Jcnd label (16-bit disp)	JNE NOT_MORE JLE LESS_OR_SO	8086	—
		8088	—
		80286	—
		80386	7/3
		80486	3/1
		Pentium–Core2	1

Condition Codes	Mnemonic	Flag	Description
0000	JO	$O = 1$	Jump if overflow
0001	JNO	$O = 0$	Jump if no overflow
0010	JB/NAE	$C = 1$	Jump if below
0011	JAE/JNB	$C = 0$	Jump if above or equal
0100	JE/JZ	$Z = 1$	Jump if equal/zero
0101	JNE/JNZ	$Z = 0$	Jump if not equal/zero
0110	JBE/JNA	$C = 1 + Z = 1$	Jump if below or equal
0111	JA/JNBE	$C = 0 \bullet Z = 0$	Jump if above
1000	JS	$S = 1$	Jump if sign
1001	JNS	$S = 0$	Jump if no sign
1010	JP/JPE	$P = 1$	Jump if parity
1011	JNP/JPO	$P = 0$	Jump if no parity
1100	JL/JNGE	$S \bullet O$	Jump if less than
1101	JGE/JNL	$S = 0$	Jump if greater than or equal
1110	JLE/JNG	$Z = 1 + S \bullet O$	Jump if less than or equal
1111	JG/JNLE	$Z = 0 + S = O$	Jump if greater than

JCXZ/JECXZ Jump if CX (ECX) equals zero

11100011 Format	Examples	Microprocessor	Clocks
JCXZ label JECXZ label	JCXZ ABOVE JCXZ BELOW JECXZ GREATER JECXZ EQUAL JCXZ NEXT	8086	18/6
		8088	18/6
		80286	8/4
		80386	9/5
		80486	8/5
		Pentium–Core2	6/5

JMP   Jump

11101011 disp Format	Examples	Microprocessor	Clocks
JMP label (short)	JMP SHORT UP JMP SHORT DOWN JMP SHORT OVER JMP SHORT CIRCUIT JMP SHORT JOKE	8086	15
		8088	15
		80286	7
		80386	7
		80486	3
		Pentium–Core2	1

11101001 disp Format	Examples	Microprocessor	Clocks
JMP label (near)	JMP VERS JMP FROG JMP UNDER JMP NEAR PTR OVER	8086	15
		8088	15
		80286	7
		80386	7
		80486	3
		Pentium–Core2	1

| 11101010 disp | | | |
Format	Examples	Microprocessor	Clocks
JMP label (far)	JMP NOT_MORE JMP UNDER JMP AGAIN JMP FAR PTR THERE	8086	15
		8088	15
		80286	11
		80386	12
		80486	17
		Pentium–Core2	3

| 11111111 oo100mmm | | | |
Format	Examples	Microprocessor	Clocks
JMP reg (near)	JMP AX JMP EAX JMP CX JMP DX	8086	11
		8088	11
		80286	7
		80386	7
		80486	3
		Pentium–Core2	2
JMP mem (near)	JMP VERS JMP FROG JMP CS:UNDER JMP DATA1[DI+2]	8086	18 + ea
		8088	18 + ea
		80286	11
		80386	10
		80486	5
		Pentium–Core2	4

| 11111111 oo101mmm | | | |
Format	Examples	Microprocessor	Clocks
JMP mem (far)	JMP WAY_OFF JMP TABLE JMP UP JMP OUT_OF_HERE	8086	24 + ea
		8088	24 + ea
		80286	15
		80386	12
		80486	13
		Pentium–Core2	4

LAHF	Load AH from flags		
10011111 Example		Microprocessor	Clocks
LAHF		8086	4
		8088	4
		80286	2
		80386	2
		80486	3
		Pentium–Core2	2

LAR	Load access rights byte		
00001111 00000010 oorrrmmm disp		O D I T S Z A P C *	
Format	Examples	Microprocessor	Clocks
LAR reg,reg	LAR AX,BX LAR CX,DX LAR ECX,EDX	8086	—
		8088	—
		80286	14
		80386	15
		80486	11
		Pentium–Core2	8
LAR reg,mem	LAR CX,DATA1 LAR AX,LIST3 LAR ECX,TOAD	8086	—
		8088	—
		80286	16
		80386	16
		80486	11
		Pentium–Core2	8

LDS — Load far pointer to DS and register

11000101 oorrrmmm

Format	Examples	Microprocessor	Clocks
LDS reg,mem	LDS DI,DATA3	8086	16 + ea
	LDS SI,LIST2	8088	24 + ea
	LDS BX,ARRAY_PTR	80286	7
	LDS CX,PNTR	80386	7
		80486	6
		Pentium–Core2	4

LEA — Load effective address

10001101 oorrrmmm disp

Format	Examples	Microprocessor	Clocks
LEA reg,mem	LEA DI,DATA3	8086	2 + ea
	LEA SI,LIST2	8088	2 + ea
	LEA BX,ARRAY_PTR	80286	3
	LEA CX,PNTR	80386	2
		80486	2
		Pentium–Core2	1

LEAVE — Leave high-level procedure

11001001

Example	Microprocessor	Clocks
LEAVE	8086	—
	8088	—
	80286	5
	80386	4
	80486	5
	Pentium–Core2	3

LES	Load far pointer to ES and register		
11000100 oorrrmmm Format	Examples	Microprocessor	Clocks
LES reg,mem	LES DI,DATA3 LES SI,LIST2 LES BX,ARRAY_PTR LES CX,PNTR	8086	16 + ea
		8088	24 + ea
		80286	7
		80386	7
		80486	6
		Pentium–Core2	4

LFS	Load far pointer to FS and register		
00001111 10110100 oorrrmmm disp Format	Examples	Microprocessor	Clocks
LFS reg,mem	LFS DI,DATA3 LFS SI,LIST2 LFS BX,ARRAY_PTR LFS CX,PNTR	8086	—
		8088	—
		80286	—
		80386	7
		80486	6
		Pentium–Core2	4

LGDT	Load global descriptor table		
00001111 00000001 oo010mmm disp Format	Examples	Microprocessor	Clocks
LGDT mem64	LGDT DESCRIP LGDT TABLED	8086	—
		8088	—
		80286	11
		80386	11
		80486	11
		Pentium–Core2	6

LGS	Load far pointer to GS and register		
00001111 10110101 oorrrmmm disp			
Format	Examples	Microprocessor	Clocks
LGS reg,mem	LGS DI,DATA3 LGS SI,LIST2 LGS BX,ARRAY_PTR LGS CX,PNTR	8086	—
		8088	—
		80286	—
		80386	7
		80486	6
		Pentium–Core2	4

LIDT	Load interrupt descriptor table		
00001111 00000001 oo011mmm disp			
Format	Examples	Microprocessor	Clocks
LIDT mem64	LIDT DATA3 LIDT LIST2	8086	—
		8088	—
		80286	12
		80386	11
		80486	11
		Pentium–Core2	6

LLDT	Load local descriptor table		
00001111 00000000 oo010mmm disp			
Format	Examples	Microprocessor	Clocks
LLDT reg	LLDT BX LLDT DX LLDT CX	8086	—
		8088	—
		80286	17
		80386	20
		80486	11
		Pentium–Core2	9

LLDT mem	LLDT DATA1 LLDT LIST3 LLDT TOAD	8086	—
		8088	—
		80286	19
		80386	24
		80486	11
		Pentium–Core2	9

LMSW   Load machine status word (80286 only)

00001111 00000001 oo110mmm disp

Format	Examples	Microprocessor	Clocks
LMSW reg	LMSW BX LMSW DX LMSW CX	8086	—
		8088	—
		80286	3
		80386	10
		80486	2
		Pentium–Core2	8
LMSW mem	LMSW DATA1 LMSW LIST3 LMSW TOAD	8086	—
		8088	—
		80286	6
		80386	13
		80486	3
		Pentium–Core2	8

LOCK	Lock the bus		

11110000 Format	Examples	Microprocessor	Clocks
LOCK:inst	LOCK:XCHG AX,BX LOCK:ADD AL,3	8086	2
		8088	3
		80286	0
		80386	0
		80486	1
		Pentium–Core2	1

LODS	Load string operand		

1010110w Format	Examples	Microprocessor	Clocks
LODSB LODSW LODSD	LODSB LODSW LODSD LODS DATA3	8086	12
		8088	15
		80286	5
		80386	5
		80486	5
		Pentium–Core2	2

LOOP/LOOPD	Loop until CX = 0 or ECX = 0		

11100010 disp Format	Examples	Microprocessor	Clocks
LOOP label LOOPD label	LOOP NEXT LOOP BACK LOOPD LOOPS	8086	17/5
		8088	17/5
		80286	8/4
		80386	11
		80486	7/6
		Pentium–Core2	5/6

LOOPE/LOOPED   Loop while equal

| 11100001 disp | | | |
Format	Examples	Microprocessor	Clocks
LOOPE label	LOOPE AGAIN	8086	18/6
LOOPED label	LOOPED UNTIL	8088	18/6
LOOPZ label	LOOPZ ZORRO	80286	8/4
LOOPZD label	LOOPZD WOW	80386	11
		80486	9/6
		Pentium–Core2	7/8

LOOPNE/LOOPNED Loop while not equal

| 11100000 disp | | | |
Format	Examples	Microprocessor	Clocks
LOOPNE label	LOOPNE FORWARD	8086	19/5
LOOPNED label	LOOPNED UPS	8088	19/5
LOOPNZ label	LOOPNZ TRY_AGAIN	80286	8/4
LOOPNZD label	LOOPNZD WOO	80386	11
		80486	9/6
		Pentium–Core2	7/8

LSL   Load segment limit

| 00001111 00000011 oorrrmmm disp | | O D I T   S Z A P C * | |
Format	Examples	Microprocessor	Clocks
LSL reg,reg	LSL AX,BX	8086	—
	LSL CX,BX	8088	—
	LSL EDX,EAX	80286	14
		80386	25
		80486	10
		Pentium–Core2	8

LSL reg,mem	LSL AX,LIMIT LSL EAX,NUM	8086	—
		8088	—
		80286	16
		80386	26
		80486	10
		Pentium–Core2	8

LSS — Load far pointer to SS and register

00001111 10110010 oorrrmmm disp

Format	Examples	Microprocessor	Clocks
LSS reg,mem	LSS DI,DATA1 LSS SP,STACK_TOP LSS CX,ARRAY	8086	—
		8088	—
		80286	—
		80386	7
		80486	6
		Pentium–Core2	4

LTR — Load task register

00001111 00000000 oo001mmm disp

Format	Examples	Microprocessor	Clocks
LTR reg	LTR AX LTR CX LTR DX	8086	—
		8088	—
		80286	17
		80386	23
		80486	20
		Pentium–Core2	10

LTR mem16	LTR TASK LTR NUM	8086	—
		8088	—
		80286	19
		80386	27
		80486	20
		Pentium–Core2	10

MOV   Move data

100010dw oorrrmmm disp

Format	Examples	Microprocessor	Clocks

MOV reg,reg	MOV CL,CH
	MOV BH,CL
	MOV CX,DX
	MOV EAX,EBP
	MOV ESP,ESI
		8086	2

		8088	2
		80286	2
		80386	2
		80486	1
		Pentium–Core2	1

MOV mem,reg	MOV DATA7,DL
	MOV NUMB,CX
	MOV TEMP,EBX
	MOV [ECX],BL
	MOV [DI],DH
		8086	9 + ea

		8088	13 + ea
		80286	3
		80386	2
		80486	1
		Pentium–Core2	1

MOV reg,mem	MOV DL,DATA8
	MOV DX,NUMB
	MOV EBX,TEMP+3
	MOV CH,TEMP[EDI]
	MOV CL,DATA2
		8086	10 + ea

		8088	12 + ea
		80286	5
		80386	4
		80486	1
		Pentium–Core2	1

| 1100011w oo000mmm disp data | | | |
Format	Examples	Microprocessor	Clocks

MOV mem,imm	MOV DATAF,23H
	MOV LIST,12H
	MOV BYTE PTR [DI],2
	MOV NUMB,234H
	MOV DWORD PTR[ECX],1
		8086	10 + ea

		8088	14 + ea
		80286	3
		80386	2
		80486	1
		Pentium–Core2	1

| 1011wrrr data | | | |
Format	Examples	Microprocessor	Clocks

MOV reg,imm	MOV BX,22H
	MOV CX,12H
	MOV CL,2
	MOV ECX,123456H
	MOV DI,100
		8086	4

		8088	4
		80286	3
		80386	2
		80486	1
		Pentium–Core2	1

| 101000dw disp | | | |
Format	Examples	Microprocessor	Clocks

MOV mem,acc	MOV DATAF,AL
	MOV LIST,AX
	MOV NUMB,EAX
		8086	10

		8088	14
		80286	3
		80386	2
		80486	1
		Pentium–Core2	1

MOV acc,mem	MOV AL,DATAE
	MOV AX,LIST
	MOV EAX,LUTE
		8086	10

		8088	14
		80286	5
		80386	4
		80486	1
		Pentium–Core2	1

100011d0 oosssmmm disp			
Format	Examples	Microprocessor	Clocks

MOV seg,reg	MOV SS,AX
	MOV DS,DX
	MOV ES,CX
	MOV FS,BX
	MOV GS,AX
		8086	2

		8088	2
		80286	2
		80386	2
		80486	1
		Pentium–Core2	1

MOV seg,mem	MOV SS,STACK_TOP
	MOV DS,DATAS
	MOV ES,TEMP1
		8086	8 + ea

		8088	12 + ea
		80286	2
		80386	2
		80486	1
		Pentium–Core2	2 or 3

MOV reg,seg	MOV BX,DS
	MOV CX,FS
	MOV CX,ES
		8086	2

		8088	2
		80286	2
		80386	2
		80486	1
		Pentium–Core2	1

MOV mem,seg	MOV DATA2,CS
	MOV TEMP,DS
	MOV NUMB1,SS
	MOV TEMP2,GS
		8086	9 + ea

		8088	13 + ea
		80286	3
		80386	2
		80486	1
		Pentium–Core2	1

| 00001111 001000d0 11rrmmm | | | |
Format	Examples	Microprocessor	Clocks

MOV reg,cr	MOV EBX,CR0
	MOV ECX,CR2
	MOV EBX,CR3
		8086	—

		8088	—
		80286	—
		80386	6
		80486	4
		Pentium–Core2	4

MOV cr,reg	MOV CR0,EAX
	MOV CR1,EBX
	MOV CR3,EDX
		8086	—

		8088	—
		80286	—
		80386	10
		80486	4
		Pentium–Core2	12–46

| 00001111 001000d1 11rrmmm | | | |
Format	Examples	Microprocessor	Clocks

MOV reg,dr	MOV EBX,DR6
	MOV ECX,DR7
	MOV EBX,DR1
		8086	—

		8088	—
		80286	—
		80386	22
		80486	10
		Pentium–Core2	11

MOV dr,reg	MOV DR0,EAX
	MOV DR1,EBX
	MOV DR3,EDX
		8086	—

		8088	—
		80286	—
		80386	22
		80486	11
		Pentium–Core2	11

00001111 001001d0 11rrrmmm			
Format	Examples	Microprocessor	Clocks
MOV reg,tr	MOV EBX,TR6 MOV ECX,TR7	8086	—
		8088	—
		80286	—
		80386	12
		80486	4
		Pentium–Core2	11
MOV tr,reg	MOV TR6,EAX MOV TR7,EBX	8086	—
		8088	—
		80286	—
		80386	12
		80486	6
		Pentium–Core2	11

MOVS   Move string data

1010010w			
Format	Examples	Microprocessor	Clocks
MOVSB MOVSW MOVSD	MOVSB MOVSW MOVSD MOVS DATA1,DATA2	8086	18
		8088	26
		80286	5
		80386	7
		80486	7
		Pentium–Core2	4

MOVSX	Move with sign extend		

00001111 1011111w oorrrmmm disp

Format	Examples	Microprocessor	Clocks
MOVSX reg,reg	MOVSX BX,AL MOVSX EAX,DX	8086	—
		8088	—
		80286	—
		80386	3
		80486	3
		Pentium–Core2	3
MOVSX reg,mem	MOVSX AX,DATA34 MOVSX EAX,NUMB	8086	—
		8088	—
		80286	—
		80386	6
		80486	3
		Pentium–Core2	3

MOVZX	Move with zero extend		

00001111 1011011w oorrrmmm disp

Format	Examples	Microprocessor	Clocks
MOVZX reg,reg	MOVZX BX,AL MOVZX EAX,DX	8086	—
		8088	—
		80286	—
		80386	3
		80486	3
		Pentium–Core2	3
MOVZX reg,mem	MOVZX AX,DATA34 MOVZX EAX,NUMB	8086	—
		8088	—
		80286	—
		80386	6
		80486	3
		Pentium–Core2	3

MUL	Multiply			

1111011w oo100mmm disp		O D I T S Z A P C * ? ? ? ? *		
Format	Examples	Microprocessor		Clocks
MUL reg	MUL BL MUL CX MUL EDX	8086		118
		8088		143
		80286		21
		80386		38
		80486		42
		Pentium–Core2		10 or 11
MUL mem	MUL DATA9 MUL WORD PTR [ESI]	8086		139
		8088		143
		80286		24
		80386		41
		80486		42
		Pentium–Core2		11

NEG	Negate			

1111011w oo011mmm disp		O D I T S Z A P C * * * * * *		
Format	Examples	Microprocessor		Clocks
NEG reg	NEG BL NEG CX NEG EDI	8086		3
		8088		3
		80286		2
		80386		2
		80486		1
		Pentium–Core2		1 or 3

NEG mem	NEG DATA9 NEG WORD PTR [ESI]	8086	16 + ea
		8088	24 + ea
		80286	7
		80386	6
		80486	3
		Pentium–Core2	1 or 3

NOP   No operation

10010000 Example		Microprocessor	Clocks
NOP		8086	3
		8088	3
		80286	3
		80386	3
		80486	3
		Pentium–Core2	1

NOT   One's complement

1111011w oo010mmm disp Format	Examples	Microprocessor	Clocks
NOT reg	NOT BL NOT CX NOT EDI	8086	3
		8088	3
		80286	2
		80386	2
		80486	1
		Pentium–Core2	1 or 3
NOT mem	NOT DATA9 NOT WORD PTR [ESI]	8086	16 + ea
		8088	24 + ea
		80286	7
		80386	6
		80486	3
		Pentium–Core2	1 or 3

OR	Inclusive-OR			

| 000010dw oorrrmmm disp | | | O D I T   S Z A P C | |
| | | | 0   * * ? * 0 | |
Format	Examples		Microprocessor	Clocks
OR reg,reg	OR AX,BX OR AL,BL OR EAX,EBX OR CX,SI OR ESI,EDI		8086	3
			8088	3
			80286	2
			80386	2
			80486	1
			Pentium–Core2	1 or 2
OR mem,reg	OR DATAY,AL OR LIST,SI OR DATA2[DI],CL OR [EAX],BL OR [EBX+2*ECX],EDX		8086	16 + ea
			8088	24 + ea
			80286	7
			80386	7
			80486	3
			Pentium–Core2	1 or 3
OR reg,mem	OR BL,DATA1 OR SI,LIST1 OR CL,DATA2[SI] OR CX,[ESI] OR ESI,[2*ECX]		8086	9 + ea
			8088	13 + ea
			80286	7
			80386	6
			80486	2
			Pentium–Core2	1 or 3

| 100000sw oo001mmm disp data | | | | |
Format	Examples		Microprocessor	Clocks
OR reg,imm	OR CX,3 OR DI,1AH OR DL,34H OR EDX,1345H OR CX,1834H		8086	4
			8088	4
			80286	3
			80386	2
			80486	1
			Pentium–Core2	1 or 3

OR mem,imm	OR DATAS,3 OR BYTE PTR[EDI],1AH OR DADDY,34H OR LIST,'A' OR TOAD,1834H	8086	17 + ea
		8088	25 + ea
		80286	7
		80386	7
		80486	3
		Pentium–Core2	1 or 3

0000110w data Format	Examples	Microprocessor	Clocks
OR acc,imm	OR AX,3 OR AL,1AH OR AH,34H OR EAX,1345H OR AL,'Y'	8086	4
		8088	4
		80286	3
		80386	2
		80486	1
		Pentium–Core2	1

OUT   Output data to port

1110011w port# Format	Examples	Microprocessor	Clocks
OUT pt,acc	OUT 12H,AL OUT 12H,AX OUT 0FFH,AL OUT 0A0H,AX OUT 10H,EAX	8086	10
		8088	14
		80286	3
		80386	10
		80486	10
		Pentium–Core2	12–26

1110111w Format	Examples	Microprocessor	Clocks
OUT DX,acc	OUT DX,AL OUT DX,AX OUT DX,EAX	8086	8
		8088	12
		80286	3
		80386	11
		80486	10
		Pentium–Core2	12–26

850   APPENDIX B

OUTS	Output string to port		
0110111w Format	Examples	Microprocessor	Clocks
OUTSB OUTSW OUTSD	OUTSB OUTSW OUTSD OUTS DATA2 REP OUTSB	8086	—
		8088	—
		80286	5
		80386	14
		80486	10
		Pentium–Core2	13–27

POP	Pop data from stack		
01011rrr Format	Examples	Microprocessor	Clocks
POP reg	POP CX POP AX POP EDI	8086	8
		8088	12
		80286	5
		80386	4
		80486	1
		Pentium–Core2	1
10001111 oo000mmm disp Format	Examples	Microprocessor	Clocks
POP mem	POP DATA1 POP LISTS POP NUMBS	8086	17 + ea
		8088	25 + ea
		80286	5
		80386	5
		80486	4
		Pentium–Core2	3

| 00sss111 | | | |
Format	Examples	Microprocessor	Clocks
POP seg	POP DS POP ES POP SS	8086	8
		8088	12
		80286	5
		80386	7
		80486	3
		Pentium–Core2	3

| 00001111 10sss001 | | | |
Format	Examples	Microprocessor	Clocks
POP seg	POP FS POP GS	8086	—
		8088	—
		80286	—
		80386	7
		80486	3
		Pentium–Core2	3

POPA/POPAD Pop all registers from stack

| 01100001 | | | |
Example		Microprocessor	Clocks
POPA POPAD		8086	—
		8088	—
		80286	19
		80386	24
		80486	9
		Pentium–Core2	5

POPF/POPFD	Pop flags from stack		
10010000		O D I T S Z A P C * * * * * * * *	
Example		Microprocessor	Clocks
POPF POPFD		8086	8
		8088	12
		80286	5
		80386	5
		80486	6
		Pentium–Core2	4 or 6

PUSH	Push data onto stack		
01010rrr Format	Examples	Microprocessor	Clocks
PUSH reg	PUSH CX PUSH AX PUSH EDI	8086	11
		8088	15
		80286	3
		80386	2
		80486	1
		Pentium–Core2	1
11111111 oo110mmm disp Format	Examples	Microprocessor	Clocks
PUSH mem	PUSH DATA1 PUSH LISTS PUSH NUMBS	8086	16 + ea
		8088	24 + ea
		80286	5
		80386	5
		80486	4
		Pentium–Core2	1 or 2

| 00ss110 | | | |
Format	Examples	Microprocessor	Clocks
PUSH seg	PUSH ES PUSH CS PUSH DS	8086	10
		8088	14
		80286	3
		80386	2
		80486	3
		Pentium–Core2	1

| 00001111 10sss000 | | | |
Format	Examples	Microprocessor	Clocks
PUSH seg	PUSH FS PUSH GS	8086	—
		8088	—
		80286	—
		80386	2
		80486	3
		Pentium–Core2	1

| 011010s0 data | | | |
Format	Examples	Microprocessor	Clocks
PUSH imm	PUSH 2000H PUSH 53220 PUSHW 10H PUSH ';' PUSHD 100000H	8086	—
		8088	—
		80286	3
		80386	2
		80486	1
		Pentium–Core2	1

PUSHA/PUSHAD   Push all registers onto stack

01100000 Example		Microprocessor	Clocks
PUSHA PUSHAD		8086	—
		8088	—
		80286	17
		80386	18
		80486	11
		Pentium–Core2	5

PUSHF/PUSHFD   Push flags onto stack

10011100 Example		Microprocessor	Clocks
PUSHF PUSHFD		8086	10
		8088	14
		80286	3
		80386	4
		80486	3
		Pentium–Core2	3 or 4

RCL/RCR/ROL/ROR Rotate

1101000w ooTTTmmm disp

		O D I T	S Z A P C
		*	*

TTT = 000 = ROL, TTT = 001 = ROR, TTT = 010 = RCL, and TTT = 011 = RCR

Format	Examples	Microprocessor	Clocks
ROL reg,1 ROR reg,1	ROL CL,1 ROL DX,1 ROR CH,1 ROR SI,1	8086	2
		8088	2
		80286	2
		80386	3
		80486	3
		Pentium–Core2	1 or 3

RCL reg,1 RCR reg,1	RCL CL,1 RCL SI,1 RCR AH,1 RCR EBX,1	8086	2
		8088	2
		80286	2
		80386	9
		80486	3
		Pentium–Core2	1 or 3
ROL mem,1 ROR mem,1	ROL DATAY,1 ROL LIST,1 ROR DATA2[DI],1 ROR BYTE PTR [EAX],1	8086	15 + ea
		8088	23 + ea
		80286	7
		80386	7
		80486	4
		Pentium–Core2	1 or 3
RCL mem,1 RCR mem,1	RCL DATA1,1 RCL LIST,1 RCR DATA2[SI],1 RCR WORD PTR [ESI],1	8086	15 + ea
		8088	23 + ea
		80286	7
		80386	10
		80486	4
		Pentium–Core2	1 or 3
1101001w ooTTTmmm disp			
Format	Examples	Microprocessor	Clocks
ROL reg,CL ROR reg,CL	ROL CH,CL ROL DX,CL ROR AL,CL ROR ESI,CL	8086	8 + 4n
		8088	8 + 4n
		80286	5 + n
		80386	3
		80486	3
		Pentium–Core2	4

RCL reg,CL RCR reg,CL	RCL CH,CL RCL SI,CL RCR AH,CL RCR EBX,CL	8086	8 + 4n
		8088	8 + 4n
		80286	5 + n
		80386	9
		80486	3
		Pentium–Core2	7–27
ROL mem,CL ROR mem,CL	ROL DATAY,CL ROL LIST,CL ROR DATA2[DI],CL ROR BYTE PTR [EAX],CL	8086	20 + 4n
		8088	28 + 4n
		80286	8 + n
		80386	7
		80486	4
		Pentium–Core2	4
RCL mem,CL RCR mem,CL	RCL DATA1,CL RCL LIST,CL RCR DATA2[SI],CL RCR WORD PTR [ESI],CL	8086	20 + 4n
		8088	28 + 4n
		80286	8 + n
		80386	10
		80486	9
		Pentium–Core2	9–26

1100000w ooTTTmmm disp data

Format	Examples	Microprocessor	Clocks
ROL reg,imm ROR reg,imm	ROL CH,4 ROL DX,5 ROR AL,2 ROR ESI,14	8086	—
		8088	—
		80286	5 + n
		80386	3
		80486	2
		Pentium–Core2	1 or 3

RCL reg,imm RCR reg,imm	RCL CL,2 RCL SI,12 RCR AH,5 RCR EBX,18	8086	—
		8088	—
		80286	5 + n
		80386	9
		80486	8
		Pentium–Core2	8–27
ROL mem,imm ROR mem,imm	ROL DATAY,4 ROL LIST,3 ROR DATA2[DI],7 ROR BYTE PTR [EAX],11	8086	—
		8088	—
		80286	8 + n
		80386	7
		80486	4
		Pentium–Core2	1 or 3
RCL mem,imm RCR mem,imm	RCL DATA1,5 RCL LIST,3 RCR DATA2[SI],9 RCR WORD PTR [ESI],8	8086	—
		8088	—
		80286	8 + n
		80386	10
		80486	9
		Pentium–Core2	8–27

RDMSR   Read model specific register

00001111 00110010 Example	Microprocessor	Clocks
RDMSR	8086	—
	8088	—
	80286	—
	80386	—
	80486	—
	Pentium–Core2	20–24

REP	Repeat prefix			

11110011 1010010w

Format	Examples	Microprocessor	Clocks
REP MOVS	REP MOVSB REP MOVSW REP MOVSD REP MOVS DATA1,DATA2	8086	$9 + 17n$
		8088	$9 + 25n$
		80286	$5 + 4n$
		80386	$8 + 4n$
		80486	$12 + 3n$
		Pentium–Core2	$13 + n$

11110011 1010101w

Format	Examples	Microprocessor	Clocks
REP STOS	REP STOSB REP STOSW REP STOSD REP STOS ARRAY	8086	$9 + 10n$
		8088	$9 + 14n$
		80286	$4 + 3n$
		80386	$5 + 5n$
		80486	$7 + 4n$
		Pentium–Core2	$9 + n$

11110011 0110110w

Format	Examples	Microprocessor	Clocks
REP INS	REP INSB REP INSW REP INSD REP INS ARRAY	8086	—
		8088	—
		80286	$5 + 4n$
		80386	$12 + 5n$
		80486	$17 + 5n$
		Pentium–Core2	$25 + 3n$

| 11110011 0110111w | | | |
Format	Examples	Microprocessor	Clocks
REP OUTS	REP OUTSB REP OUTSW REP OUTSD REP OUTS ARRAY	8086	—
		8088	—
		80286	$5 + 4n$
		80386	$12 + 5n$
		80486	$17 + 5n$
		Pentium–Core2	$25 + 4n$

REPE/REPNE   Repeat conditional

| 11110011 1010011w | | | |
Format	Examples	Microprocessor	Clocks
REPE CMPS	REPE CMPSB REPE CMPSW REPE CMPSD REPE CMPS DATA1,DATA2	8086	$9 + 22n$
		8088	$9 + 30n$
		80286	$5 + 9n$
		80386	$5 + 9n$
		80486	$7 + 7n$
		Pentium–Core2	$9 + 4n$

| 11110011 1010111w | | | |
Format	Examples	Microprocessor	Clocks
REPE SCAS	REPE SCASB REPE SCASW REPE SCASD REPE SCAS ARRAY	8086	$9 + 15n$
		8088	$9 + 19n$
		80286	$5 + 8n$
		80386	$5 + 8n$
		80486	$7 + 5n$
		Pentium–Core2	$9 + 4n$

11110010 1010011w Format	Examples	Microprocessor	Clocks
REPNE CMPS	REPNE CMPSB REPNE CMPSW REPNE CMPSD REPNE CMPS ARRAY,LIST	8086	9 + 22n
		8088	9 + 30n
		80286	5 + 9n
		80386	5 + 9n
		80486	7 + 7n
		Pentium–Core2	8 + 4n

11110010 101011w Format	Examples	Microprocessor	Clocks
REPNE SCAS	REPNE SCASB REPNE SCASW REPNE SCASD REPNE SCAS ARRAY	8086	9 + 15n
		8088	9 + 19N
		80286	5 + 8n
		80386	5 + 8n
		80486	7 + 5n
		Pentium–Core2	9 + 4n

RET	Return from procedure		

11000011 Example		Microprocessor	Clocks
RET (near)		8086	16
		8088	20
		80286	11
		80386	10
		80486	5
		Pentium–Core2	2

11000010 data			
Format	Examples	Microprocessor	Clocks
RET imm (near)	RET 4 RET 100H	8086	20
		8088	24
		80286	11
		80386	10
		80486	5
		Pentium–Core2	3

11001011		
Example	Microprocessor	Clocks
RET (far)	8086	26
	8088	34
	80286	15
	80386	18
	80486	13
	Pentium–Core2	4–23

11001010 data			
Format	Examples	Microprocessor	Clocks
RET imm (far)	RET 4 RET 100H	8086	25
		8088	33
		80286	11
		80386	10
		80486	5
		Pentium–Core2	4–23

RSM	Resume from system management mode		

00001111 10101010		O D I T S Z A P C	
		* * * * * * * * *	
Example		Microprocessor	Clocks
RSM		8086	—
		8088	—
		80286	—
		80386	—
		80486	—
		Pentium–Core2	83

SAHF	Store AH into flags		

10011110		O D I T S Z A P C	
		* * * * *	
Example		Microprocessor	Clocks
SAHF		8086	4
		8088	4
		80286	2
		80386	3
		80486	2
		Pentium–Core2	2

SAL/SAR/SHL/SHR Shift			

1101000w ooTTTmmm disp		O D I T S Z A P C	
		* * * ? * *	
TTT = 100 = SHL/SAL , TTT = 101 = SHR, and TTT = 111 = SAR			
Format	Examples	Microprocessor	Clocks
SAL reg,1	SAL CL,1	8086	2
SHL reg,1	SHL DX,1		
SHR reg,1	SAR CH,1	8088	2
SAR reg,1	SHR SI,1		
		80286	2
		80386	3
		80486	3
		Pentium–Core2	1 or 3

SAL mem,1 SHL mem,1 SHR mem,1 SAR mem,1	SAL DATA1,1 SHL BYTE PTR [DI],1 SAR NUMB,1 SHR WORD PTR[EDI],1	8086	15 + ea
		8088	23 + ea
		80286	7
		80386	7
		80486	4
		Pentium–Core2	1 or 3

1101001w ooTTTmmm disp

Format	Examples	Microprocessor	Clocks
SAL reg,CL SHL reg,CL SAR reg,CL SHR reg,CL	SAL CH,CL SHL DX,CL SAR AL,CL SHR ESI,CL	8086	8 + 4n
		8088	8 + 4n
		80286	5 + n
		80386	3
		80486	3
		Pentium–Core2	4
SAL mem,CL SHL mem,CL SAR mem,CL SHR mem,CL	SAL DATAU,CL SHL BYTE PTR [ESI],CL SAR NUMB,CL SHR TEMP,CL	8086	20 + 4n
		8088	28 + 4n
		80286	8 + n
		80386	7
		80486	4
		Pentium–Core2	4

1100000w ooTTTmmm disp data

Format	Examples	Microprocessor	Clocks
SAL reg,imm SHL reg,imm SAR reg,imm SHR reg,imm	SAL CH,4 SHL DX,10 SAR AL,2 SHR ESI,23	8086	—
		8088	—
		80286	5 + n
		80386	3
		80486	2
		Pentium–Core2	1 or 3

SAL mem,imm SHL mem,imm SAR mem,imm SHR mem,imm	SAL DATAU,3 SHL BYTE PTR [ESI],15 SAR NUMB,3 SHR TEMP,5	8086	—
		8088	—
		80286	8 + n
		80386	7
		80486	4
		Pentium–Core2	1 or 3

SBB   Subtract with borrow

000110dw oorrrmmm disp		O D I T S Z A P C\n* * * * * *	
Format	Examples	Microprocessor	Clocks
SBB reg,reg	SBB CL,DL SBB AX,DX SBB CH,CL SBB EAX,EBX SBB ESI,EDI	8086	3
		8088	3
		80286	2
		80386	2
		80486	1
		Pentium–Core2	1 or 2
SBB mem,reg	SBB DATAJ,CL SBB BYTES,CX SBB NUMBS,ECX SBB [EAX],CX	8086	16 + ea
		8088	24 + ea
		80286	7
		80386	6
		80486	3
		Pentium–Core2	1 or 3
SBB reg,mem	SBB CL,DATAL SBB CX,BYTES SBB ECX,NUMBS SBB DX,[EBX+EDI]	8086	9 + ea
		8088	13 + ea
		80286	7
		80386	7
		80486	2
		Pentium–Core2	1 or 2

| 100000sw oo011mmm disp data | | | |
Format	Examples	Microprocessor	Clocks
SBB reg,imm	SBB CX,3 SBB DI,1AH SBB DL,34H SBB EDX,1345H SBB CX,1834H	8086	4
		8088	4
		80286	3
		80386	2
		80486	1
		Pentium–Core2	1 or 3
SBB mem,imm	SBB DATAS,3 SBB BYTE PTR[EDI],1AH SBB DADDY,34H SBB LIST,'A' SBB TOAD,1834H	8086	17 + ea
		8088	25 + ea
		80286	7
		80386	7
		80486	3
		Pentium–Core2	1 or 3

| 0001110w data | | | |
Format	Examples	Microprocessor	Clocks
SBB acc,imm	SBB AX,3 SBB AL,1AH SBB AH,34H SBB EAX,1345H SBB AL,'Y'	8086	4
		8088	4
		80286	3
		80386	2
		80486	1
		Pentium–Core2	1

SCAS   Scan string

| 1010111w | | O D I T   S Z A P C
*   * * * * * | |
Format	Examples	Microprocessor	Clocks
SCASB SCASW SCASD	SCASB SCASW SCASD SCAS DATAF REP SCASB	8086	15
		8088	19
		80286	7
		80386	7
		80486	6
		Pentium–Core2	4

SETcondition	Conditional set		

00001111 1001cccc oo000mmm

Format	Examples	Microprocessor	Clocks
SETcnd reg8	SETA BL SETB CH SETG DL SETE BH SETZ AL	8086	—
		8088	—
		80286	—
		80386	4
		80486	3
		Pentium–Core2	1 or 2
SETcnd mem8	SETE DATAK SETAE LESS_OR_SO	8086	—
		8088	—
		80286	—
		80386	5
		80486	3
		Pentium–Core2	1 or 2

Condition Codes	Mnemonic	Flag	Description
0000	SETO	O = 1	Set if overflow
0001	SETNO	O = 0	Set if no overflow
0010	SETB/SETAE	C = 1	Set if below
0011	SETAE/SETNB	C = 0	Set if above or equal
0100	SETE/SETZ	Z = 1	Set if equal/zero
0101	SETNE/SETNZ	Z = 0	Set if not equal/zero
0110	SETBE/SETNA	C = 1 + Z = 1	Set if below or equal
0111	SETA/SETNBE	C = 0 • Z = 0	Set if above
1000	SETS	S = 1	Set if sign
1001	SETNS	S = 0	Set if no sign
1010	SETP/SETPE	P = 1	Set if parity
1011	SETNP/SETPO	P = 0	Set if no parity
1100	SETL/SETNGE	S • O	Set if less than
1101	SETGE/SETNL	S = 0	Set if greater than or equal
1110	SETLE/SETNG	Z = 1 + S • O	Set if less than or equal
1111	SETG/SETNLE	Z = 0 + S = O	Set if greater than

SGDT/SIDT/SLDT   Store descriptor table registers

00001111 00000001 oo000mmm disp			
Format	Examples	Microprocessor	Clocks
SGDT mem	SGDT MEMORY SGDT GLOBAL	8086	—
		8088	—
		80286	11
		80386	9
		80486	10
		Pentium–Core2	4

00001111 00000001 oo001mmm disp			
Format	Examples	Microprocessor	Clocks
SIDT mem	SIDT DATAS SIDT INTERRUPT	8086	—
		8088	—
		80286	12
		80386	9
		80486	10
		Pentium–Core2	4

00001111 00000000 oo000mmm disp			
Format	Examples	Microprocessor	Clocks
SLDT reg	SLDT CX SLDT DX	8086	—
		8088	—
		80286	2
		80386	2
		80486	2
		Pentium–Core2	2
SLDT mem	SLDT NUMBS SLDT LOCALS	8086	—
		8088	—
		80286	3
		80386	2
		80486	3
		Pentium–Core2	2

SHLD/SHRD Double precision shift

00001111 10100100 oorrrmmm disp data		O D I T S Z A P C	
		? * * ? * *	
Format	Examples	Microprocessor	Clocks
SHLD	SHLD AX,CX,10	8086	—
reg,reg,imm	SHLD DX,BX,8	8088	—
	SHLD CX,DX,2	80286	—
		80386	3
		80486	2
		Pentium–Core2	4
SHLD	SHLD DATAQ,CX,8	8086	—
mem,reg,imm		8088	—
		80286	—
		80386	7
		80486	3
		Pentium–Core2	4

00001111 10101100 oorrrmmm disp data			
Format	Examples	Microprocessor	Clocks
SHRD	SHRD CX,DX,2	8086	—
reg,reg,imm		8088	—
		80286	—
		80386	3
		80486	2
		Pentium–Core2	4
SHRD	SHRD DATAZ,DX,4	8086	—
mem,reg,imm		8088	—
		80286	—
		80386	7
		80486	2
		Pentium–Core2	4

| 00001111 10100101 oorrrmmm disp | | | |
Format	Examples	Microprocessor	Clocks
SHLD reg,reg,CL	SHLD BX,DX,CL	8086	—
		8088	—
		80286	—
		80386	3
		80486	3
		Pentium–Core2	4 or 5
SHLD mem,reg,CL	SHLD DATAZ,DX,CL	8086	—
		8088	—
		80286	—
		80386	7
		80486	3
		Pentium–Core2	4 or 5

| 00001111 10101101 oorrrmmm disp | | | |
Format	Examples	Microprocessor	Clocks
SHRD reg,reg,CL	SHRD AX,DX,CL	8086	—
		8088	—
		80286	—
		80386	3
		80486	3
		Pentium–Core2	4 or 5
SHRD mem,reg,CL	SHRD DATAZ,DX,CL	8086	—
		8088	—
		80286	—
		80386	7
		80486	3
		Pentium–Core2	4 or 5

SMSW	Store machine status word (80286)		
00001111 00000001 oo100mmm disp			
Format	Examples	Microprocessor	Clocks
SMSW reg	SMSW AX SMSW DX SMSW BP	8086	—
		8088	—
		80286	2
		80386	10
		80486	2
		Pentium–Core2	4
SMSW mem	SMSW DATAQ	8086	—
		8088	—
		80286	3
		80386	3
		80486	3
		Pentium–Core2	4

STC	Set carry flag		
11111001		O D I T S Z A P C 1	
Example		Microprocessor	Clocks
STC		8086	2
		8088	2
		80286	2
		80386	2
		80486	2
		Pentium–Core2	2

STD	Set direction flag		

11111101		O D I T S Z A P C 1	
Example		Microprocessor	Clocks
STD		8086	2
		8088	2
		80286	2
		80386	2
		80486	2
		Pentium–Core2	2

STI	Set interrupt flag		

11111011		O D I T S Z A P C 1	
Example		Microprocessor	Clocks
STI		8086	2
		8088	2
		80286	2
		80386	3
		80486	5
		Pentium–Core2	7

STOS	Store string data		

1010101w Format	Examples	Microprocessor	Clocks
STOSB STOSW STOSD	STOSB STOSW STOSD STOS DATA_LIST REP STOSB	8086	11
		8088	15
		80286	3
		80386	40
		80486	5
		Pentium–Core2	3

STR	Store task register		

00001111 00000000 oo001mmm disp

Format	Examples	Microprocessor	Clocks
STR reg	STR AX STR DX STR BP	8086	—
		8088	—
		80286	2
		80386	2
		80486	2
		Pentium–Core2	2
STR mem	STR DATA3	8086	—
		8088	—
		80286	2
		80386	2
		80486	2
		Pentium–Core2	2

SUB	Subtract		

000101dw oorrrmmm disp

	O	D	I	T		S	Z	A	P	C
	*					*	*	*	*	*

Format	Examples	Microprocessor	Clocks

SUB reg,reg	SUB CL,DL
	SUB AX,DX
	SUB CH,CL
	SUB EAX,EBX
	SUB ESI,EDI
		8086	3

		8088	3
		80286	2
		80386	2
		80486	1
		Pentium–Core2	1 or 2

SUB mem,reg	SUB DATAJ,CL
	SUB BYTES,CX
	SUB NUMBS,ECX
	SUB [EAX],CX
		8086	16 + ea

		8088	24 + ea
		80286	7
		80386	6
		80486	3
		Pentium–Core2	1 or 3

SUB reg,mem	SUB CL,DATAL
	SUB CX,BYTES
	SUB ECX,NUMBS
	SUB DX,[EBX+EDI]
		8086	9 + ea

		8088	13 + ea
		80286	7
		80386	7
		80486	2
		Pentium–Core2	1 or 2

100000sw oo101mmm disp data			
Format	Examples	Microprocessor	Clocks

SUB reg,imm	SUB CX,3
	SUB DI,1AH
	SUB DL,34H
	SUB EDX,1345H
	SUB CX,1834H
		8086	4

		8088	4
		80286	3
		80386	2
		80486	1
		Pentium–Core2	1 or 3

SUB mem,imm	SUB DATAS,3
	SUB BYTE PTR[EDI],1AH
	SUB DADDY,34H
	SUB LIST,'A'
	SUB TOAD,1834H
		8086	17 + ea

		8088	25 + ea
		80286	7
		80386	7
		80486	3
		Pentium–Core2	1 or 3

0010110w data			
Format	Examples	Microprocessor	Clocks

SUB acc,imm	SUB AL,3
	SUB AX,1AH
	SUB EAX,34H
		8086	4

		8088	4
		80286	3
		80386	2
		80486	1
		Pentium–Core2	1

TEST   Test operands (logical compare)

1000001w oorrrmmm disp		O D I T S Z A P C 0 * * ? * 0	
Format	Examples	Microprocessor	Clocks
TEST reg,reg	TEST CL,DL TEST BX,DX TEST DH,CL TEST EBP,EBX TEST EAX,EDI	8086	5
		8088	5
		80286	2
		80386	2
		80486	1
		Pentium–Core2	1 or 2
TEST mem,reg reg,mem	TEST DATAJ,CL TEST BYTES,CX TEST NUMBS,ECX TEST [EAX],CX TEST CL,POPS	8086	9 + ea
		8088	13 + ea
		80286	6
		80386	5
		80486	2
		Pentium–Core2	1 or 2

| 1111011sw oo000mmm disp data | | | |
Format	Examples	Microprocessor	Clocks
TEST reg,imm	TEST BX,3 TEST DI,1AH TEST DH,44H TEST EDX,1AB345H TEST SI,1834H	8086	4
		8088	4
		80286	3
		80386	2
		80486	1
		Pentium–Core2	1 or 2
TEST mem,imm	TEST DATAS,3 TEST BYTE PTR[EDI],1AH TEST DADDY,34H TEST LIST,'A' TEST TOAD,1834H	8086	11 + ea
		8088	11 + ea
		80286	6
		80386	5
		80486	2
		Pentium–Core2	1 or 2

| 1010100w data | | | |
Format	Examples	Microprocessor	Clocks
TEST acc,imm	TEST AL,3 TEST AX,1AH TEST EAX,34H	8086	4
		8088	4
		80286	3
		80386	2
		80486	1
		Pentium–Core2	1

VERR/VERW Verify read/write			

00001111 00000000 oo100mmm disp		O D I T S Z A P C *	
Format	Examples	Microprocessor	Clocks
VERR reg	VERR CX VERR DX VERR DI	8086	—
		8088	—
		80286	14
		80386	10
		80486	11
		Pentium–Core2	7
VERR mem	VERR DATAJ VERR TESTB	8086	—
		8088	—
		80286	16
		80386	11
		80486	11
		Pentium–Core2	7

00001111 00000000 oo101mmm disp			
Format	Examples	Microprocessor	Clocks
VERW reg	VERW CX VERW DX VERW DI	8086	—
		8088	—
		80286	14
		80386	15
		80486	11
		Pentium–Core2	7
VERW mem	VERW DATAJ VERW TESTB	8086	—
		8088	—
		80286	16
		80386	16
		80486	11
		Pentium–Core2	7

WAIT	Wait for coprocessor		
10011011 Example		Microprocessor	Clocks
WAIT FWAIT		8086	4
		8088	4
		80286	3
		80386	6
		80486	6
		Pentium–Core2	1

WBINVD	Write-back cache invalidate data cache		
00001111 00001001 Example		Microprocessor	Clocks
WBINVD		8086	—
		8088	—
		80286	—
		80386	—
		80486	5
		Pentium–Core2	2000+

WRMSR	Write to model specific register		
00001111 00110000 Example		Microprocessor	Clocks
WRMSR		8086	—
		8088	—
		80286	—
		80386	—
		80486	—
		Pentium–Core2	30–45

XADD	Exchange and add		

00001111 1100000w 11rrrrrr		O D I T S Z A P C	
		* * * * * *	
Format	Examples	Microprocessor	Clocks
XADD reg,reg	XADD EBX,ECX	8086	—
	XADD EDX,EAX	8088	—
	XADD EDI,EBP	80286	—
		80386	—
		80486	3
		Pentium–Core2	3 or 4

00001111 1100000w oorrrmmm disp			
Format	Examples	Microprocessor	Clocks
XADD mem,reg	XADD DATA5,ECX	8086	—
	XADD [EBX],EAX	8088	—
	XADD [ECX+4],EBP	80286	—
		80386	—
		80486	4
		Pentium–Core2	3 or 4

XCHG	Exchange		

1000011w oorrrmmm			
Format	Examples	Microprocessor	Clocks
XCHG reg,reg	XCHG CL,DL	8086	4
	XCHG BX,DX	8088	4
	XCHG DH,CL	80286	3
	XCHG EBP,EBX	80386	3
	XCHG EAX,EDI	80486	3
		Pentium–Core2	3

XCHG mem,reg reg,mem	XCHG DATAJ,CL XCHG BYTES,CX XCHG NUMBS,ECX XCHG [EAX],CX XCHG CL,POPS	8086	17 + ea
		8088	25 + ea
		80286	5
		80386	5
		80486	5
		Pentium–Core2	3

10010reg Format	Examples	Microprocessor	Clocks
XCHG acc,reg reg,acc	XCHG BX,AX XCHG AX,DI XCHG DH,AL XCHG EDX,EAX XCHG SI,AX	8086	3
		8088	3
		80286	3
		80386	3
		80486	3
		Pentium–Core2	2

XLAT   Translate

11010111 Example	Microprocessor	Clocks
XLAT	8086	11
	8088	11
	80286	5
	80386	3
	80486	4
	Pentium–Core2	4

XOR	Exclusive-OR		

| 000110dw oorrrmmm disp | | O D I T   S Z A P C | |
| | | 0   * * ? * 0 | |
Format	Examples	Microprocessor	Clocks
XOR reg,reg	XOR CL,DL	8086	3
	XOR AX,DX	8088	3
	XOR CH,CL	80286	2
	XOR EAX,EBX	80386	2
	XOR ESI,EDI	80486	1
		Pentium–Core2	1 or 2
XOR mem,reg	XOR DATAJ,CL	8086	16 + ea
	XOR BYTES,CX	8088	24 + ea
	XOR NUMBS,ECX	80286	7
	XOR [EAX],CX	80386	6
		80486	3
		Pentium–Core2	1 or 3
XOR reg,mem	XOR CL,DATAL	8086	9 + ea
	XOR CX,BYTES	8088	13 + ea
	XOR ECX,NUMBS	80286	7
	XOR DX,[EBX+EDI]	80386	7
		80486	2
		Pentium–Core2	1 or 2

| 100000sw oo110mmm disp data | | | |
Format	Examples	Microprocessor	Clocks
XOR reg,imm	XOR CX,3	8086	4
	XOR DI,1AH	8088	4
	XOR DL,34H	80286	3
	XOR EDX,1345H	80386	2
	XOR CX,1834H	80486	1
		Pentium–Core2	1 or 3

XOR mem,imm	XOR DATAS,3 XOR BYTE PTR[EDI],1AH XOR DADDY,34H XOR LIST,'A' XOR TOAD,1834H	8086	17 + ea
		8088	25 + ea
		80286	7
		80386	7
		80486	3
		Pentium–Core2	1 or 3

0010101w data Format	Examples	Microprocessor	Clocks
XOR acc,imm	XOR AL,3 XOR AX,1AH XOR EAX,34H	8086	4
		8088	4
		80286	3
		80386	2
		80486	1
		Pentium–Core2	1

SIMD INSTRUCTION SET SUMMARY

The SIMD (single-instruction, multiple data) instructions add a new dimension to the use of the microprocessor for performing multimedia and other operations. The XMM registers are numbered from XMM_0 to XMM_7 and are each 128 bits in width. Data formats stored in the XMM registers and used by the SIMD instructions appear in Figure B–1.

Packed 64-bit double-precision floating-point data

Packed byte integer data

Packed word integer data

Packed 32-bit integer data

Packed 64-bit integer data
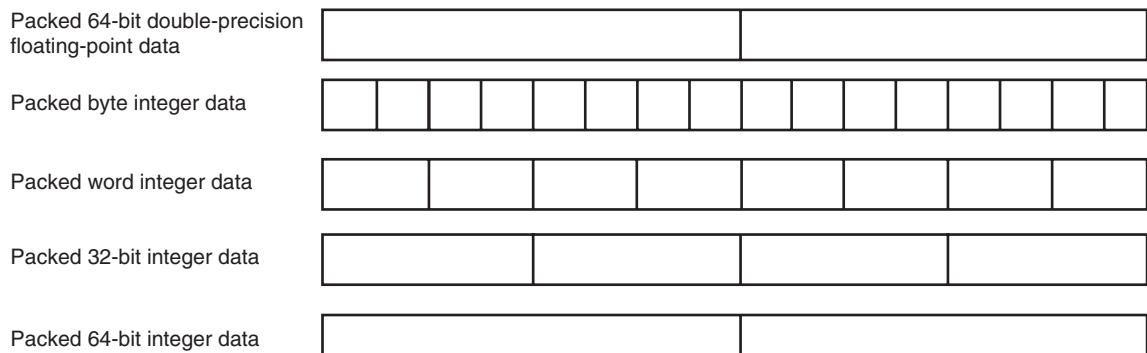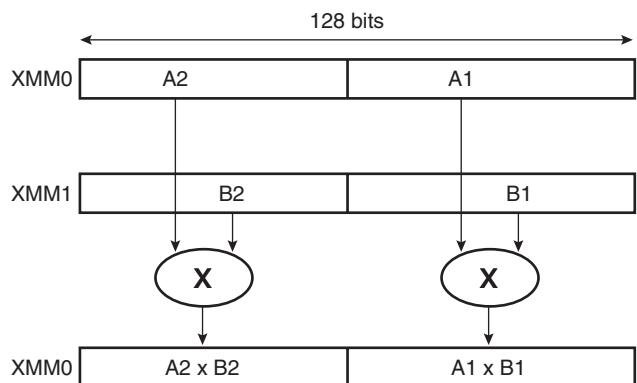
FIGURE B–1   Data formats for the 128-bit-wide XMM registers in the Pentium III and Pentium 4 microprocessors.

Data stored in the memory must be stored as 16-byte-long data in a series of memory locations accessed by using the OWORD PTR override when addressed by an instruction. The OWORD PTR override is used to address an octalword of data or 16 bytes. The SIMD instructions allow operations on packed and scalar double-precision floating-point numbers. The operation of both forms is illustrated in Figure B–2, which shows both packed and scalar multiplication. Notice that scalar only copies the leftmost double-precision number into the destination register and does not use the leftmost number in the source. The scalar instructions are meant to be compatible with the floating-point coprocessor instructions.
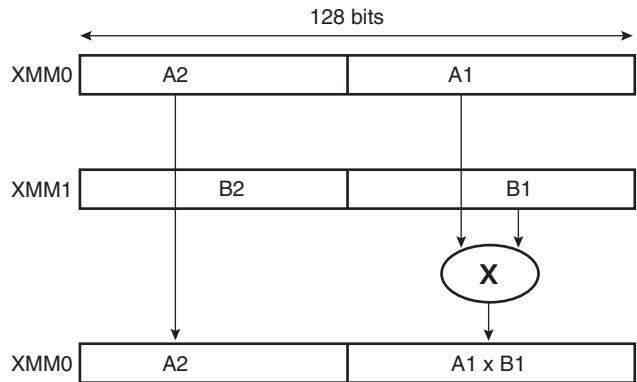
This section of the appendix details many of the SIMD instructions and provides examples of their usage.

FIGURE B–2 Packed and scalar double-precision floating-point operation.

Packed double-precision multiplication MULPD XMM0, XMM1

Packed double-precision multiplication MULSD XMM0, XMM1

DATA MOVEMENT INSTRUCTIONS

MOVAPD	Move aligned packed double-precision data, data must be aligned on 16-byte boundaries

Examples

```
MOVAPD  XMM0, OWORD DATA3        ;copies DATA3 to XMM0
MOVAPD  OWORD PTR DATA4, XMM2 ;copies XMM4 to DATA4
```

MOVUPD	Move unaligned packed double-precision data

Examples

```
MOVUPD  XMM0, OWORD DATA3        ;copies DATA3 to XMM0
MOVUPD  OWORD PTR DATA4, XMM2 ;copies XMM4 to DATA4
```

MOVSD	Move scalar packed double-precision data to low quadword

Examples

```
MOVSD  XMM0, DWORD DATA3         ;copies DATA3 to XMM0
MOVSD  DWORD PTR DATA4, XMM2  ;copies XMM4 to DATA4
```

MOVHPD	Move packed double-precision data to high quadword

Examples

```
MOVHPD  XMM0, DWORD DATA3         ;copies DATA3 to XMM0
MOVHPD  DWORD PTR DATA4, XMM2 ;copies XMM4 to DATA4
```

MOVLPD	Move packed double-precision data into low quadword

Examples

```
MOVLPD  XMM0, DWORD DATA3         ;copies DATA3 to XMM0
MOVLPD  DWORD PTR DATA4, XMM2 ;copies XMM4 to DATA4
```

MOVMSKPD	Move packed double-precision mask

Examples

```
MOVMSKPD   EAX, XMM1    ;copies 2 sign bits to general-purpose register
```

MOVAPS	Move 4 aligned packed single-precision data, data must be aligned on 16-byte boundaries
Examples	MOVAPS XMM0, OWORD DATA3 ;copies DATA3 to XMM0 MOVAPS OWORD PTR DATA4, XMM2 ;copies XMM4 to DATA4
MOVUPS	Move 4 unaligned packed single-precision data
Examples	MOVUPS XMM0, OWORD DATA3 ;copies DATA3 to XMM0 MOVUPS OWORD PTR DATA4, XMM2 ;copies XMM4 to DATA4
MOVLPS	Move 2 packed single-precision numbers to low-order quadword
Examples	MOVLPS XMM0, OWORD DATA3 ;copies DATA3 to XMM0 MOVLPS OWORD PTR DATA4, XMM2 ;copies XMM4 to DATA4
MOVHPS	Move packed single-precision numbers to high-order quadword
Examples	MOVHPS XMM0, OWORD DATA3 ;copies DATA3 to XMM0 MOVHPS OWORD PTR DATA4, XMM2 ;copies XMM4 to DATA4
MOVAPD	Move aligned packed double-precision data, data must be aligned on 16-byte boundaries
Examples	MOVAPD XMM0, OWORD DATA3 ;copies DATA3 to XMM0 MOVAPD OWORD PTR DATA4, XMM2 ;copies XMM4 to DATA4
MOVLHPS	Move 2 packed single-precision numbers from the low-order quadword to the high-order quadword
Examples	MOVLHPS XMM0, XMM1 ;copies XMM1 low to XMM0 high MOVLHPS XMM3, XMM2 ;copies XMM2 low to XMM3 high

MOVHLPS	Move 2 packed single-precision numbers from high-order quadword to low-order quadword

Examples

   MOVHLPS XMM0, XMM2   ;copies high XMM2 to low XMM0

   MOVHLPS XMM4, XMM5   ;copies high XMM5 to low XMM4

MOVMSKPS	Move 4-sign bits of 4 packed single-precision numbers to general-purpose register

Examples

   MOVMSKPS EBX, XMM0   ;copies sign bits of XMM0 to EBX

   MOVMSKPS EDX, XMM2   ;copies sign bits of XMM2 to EDX

ARITHMETIC INSTRUCTIONS

ADDPD	Adds packed double-precision data

Examples

   ADDPD   XMM0, OWORD DATA3   ;adds DATA3 to XMM0

   ADDPD   XMM2, XMM3   ;adds XMM3 to XMM2

ADDSD	Adds scalar double-precision data

Examples

   ADDSD   XMM0, OWORD DATA3   ;adds DATA3 to XMM0

   ADDSD   XMM4, XMM2   ;adds XMM2 to XMM4

ADDPS	Adds 2 packed single-precision numbers

Examples

   ADDPS   XMM0, QWORD DATA3   ;adds DATA3 to XMM0

   ADDPS   XMM3, XMM2   ;adds XMM2 to XMM3

ADDLS	Adds scalar single-precision data

Examples

```
ADDLS    XMM0, DWORD DATA3       ;adds DATA3 to XMM0
ADDLS    XMM7, XMM2              ;adds XMM2 to XMM7
```

SUBPD	Subtracts packed double-precision data

Examples

```
SUBPD    XMM0, OWORD DATA3       ;subtracts DATA3 from XMM0
SUBPD    XMM2, XMM3              ;subtracts XMM3 from XMM2
```

SUBSD	Subtracts scalar double-precision data

Examples

```
SUBSD    XMM0, OWORD DATA3       ;subtracts DATA3 from XMM0
SUBSD    XMM4, XMM2              ;subtracts XMM2 from XMM4
```

SUBPS	Subtracts 2 packed single-precision numbers

Examples

```
SUBPS    XMM0, QWORD DATA3       ;subtracts DATA3 from XMM0
SUBPS    XMM3, XMM2              ;subtracts XMM2 from XMM3
```

SUBLS	Subtracts scalar single-precision data

Examples

```
SUBLS    XMM0, DWORD DATA3       ;subtracts DATA3 from XMM0
SUBLS    XMM7, XMM2              ;subtracts XMM2 from XMM7
```

MULPD	Multiplies packed double-precision data

Examples

```
MULPD    XMM0, OWORD DATA3       ;multiplies DATA3 times XMM0
MULPD    XMM3, XMM2              ;multiplies XMM2 times XMM3
```

MULSD	Multiplies scalar double-precision data

Examples

```
MULSD    XMM0, OWORD DATA3        ;multiplies DATA3 times XMM0
MULSD    XMM3, XMM6               ;multiplies XMM6 times XMM3
```

MULPS	Multiplies 2 packed single-precision numbers

Examples

```
MULPS    XMM0, QWORD DATA3        ;multiplies DATA3 times XMM0
MULPS    XMM0, XMM2               ;multiplies XMM2 times XMM0
```

MULSS	Multiplies a single-precision number

Examples

```
MULSS    XMM0, DWORD DATA3        ;multiplies DATA3 times XMM0
MULSS    XMM1, XMM2               ;multiplies XMM2 times XMM1
```

DIVPD	Divides packed double-precision data

Examples

```
DIVPD    XMM0, OWORD DATA3     ;divides XMM0 by DATA3
DIVPD    XMM3, XMM2            ;divides XMM3 by XMM2
```

DIVSD	Divides scalar double-precision data

Examples

```
DIVSD    XMM0, OWORD DATA3     ;divides XMM0 by DATA3
DIVSD    XMM3, XMM6            ;divides XMM3 by XMM6
```

DIVPS	Divides 2 packed single-precision numbers

Examples

```
DIVPS    XMM0, QWORD DATA3     ;divides XMM0 by DATA3
DIVPS    XMM0, XMM2            ;divides XMM0 by XMM2
```

DIVSS	Divides a single-precision number

Examples		
DIVSS	XMM0, DWORD DATA3	;divides XMM0 by DATA3
DIVSS	XMM1, XMM2	;divides XMM1 by XMM2

SQRTPD	Finds the square root of packed double-precision data

Examples		
SQRTPD	XMM0, OWORD DATA3	;finds square root of DATA3, result to XMM0
SQRTPD	XMM3, XMM2	;finds square root of XMM2, result to XMM3

SQRTSD	Finds the square root of scalar double-precision data

Examples		
SQRTSD	XMM0, OWORD DATA3	;finds square root of DATA3, result to XMM0
SQRTSD	XMM3, XMM6	;finds square root of XMM6, result to XMM3

SQRTPS	Finds the square root of 2 packed single-precision numbers

Examples		
SQRTPS	XMM0, QWORD DATA3	;finds square root of DATA3, result to XMM0
SQRTPS	XMM0, XMM2	;finds square root of XMM2, result to XMM0

SQRTSS	Finds the square root of a single-precision number

Examples		
SQRTSS	XMM0, DWORD DATA3	;finds the square root of DATA3, result to XMM0
SQRTSS	XMM1, XMM2	;finds the square root of XMM2, result to XMM1

RCPPS	Finds the reciprocal of a packed single-precision number

Examples		
RCPPS	XMM0, OWORD DATA3	;finds the reciprocal of DATA3, result to XMM0
RCPPS	XMM3, XMM2	;finds the reciprocal of XMM2, result to XMM3

RCPSS	Finds the reciprocal of a single-precision number

Examples

| RCPSS | XMM0, OWORD DATA3 | ;finds the reciprocal of DATA3, result to XMM0 |
| RCPSS | XMM3, XMM6 | ;finds the reciprocal of XMM6, result to XMM3 |

RSQRTPS	Finds reciprocals of packed single-precision data

Examples

| RSQRTPS | XMM0, OWORD DATA3 | ;finds reciprocal of square root of DATA3 |
| RSQRTPS | XMM3, XMM2 | ;finds reciprocal of square root of XMM2 |

RSQRTSS	Finds the reciprocal of square root of a scalar single-precision number

Examples

| RSQRTSS | XMM0, OWORD DATA3 | ;finds reciprocal of square root of DATA3 |
| RSQRTSS | XMM3, XMM6 | ;finds reciprocal of square root of XMM6 |

MAXPD	Compares and returns the maximum packed double-precision floating-point number

Examples

| MAXPD | XMM0, OWORD DATA3 | ;compares numbers in DATA3, largest to XMM0 |
| MAXPD | XMM3, XMM2 | ;compares numbers in XMM2, largest to XMM3 |

MAXSD	Compares scalar double-precision data and returns the largest

Examples

| MAXSD | XMM0, OWORD DATA3 | ;compares numbers in DATA3, largest to XMM0 |
| MAXSD | XMM3, XMM6 | ;compares numbers in XMM6, largest to XMM3 |

MAXPS	Compares and returns the largest packed single-precision number

Examples

| MAXPS | XMM0, QWORD DATA3 | ;compares numbers in DATA3, largest to XMM0 |
| MAXPS | XMM0, XMM2 | ;compares numbers in XMM2, largest to XMM0 |

MAXSS	Compares scalar single-precision numbers and returns the largest
Examples	

```
    MAXSS    XMM0, DWORD DATA3  ;compares numbers in DATA3, largest to XMM0
    MAXSS    XMM1, XMM2              ;compares numbers in XMM2, largest to XMM1
```

MINPD	Compares and returns the minimum packed double-precision floating-point number
Examples	

```
    MINPD    XMM0, OWORD DATA3  ;compares numbers in DATA3, least to XMM0
    MINPD    XMM3, XMM2              ;compares numbers in XMM2, least to XMM3
```

MINSD	Compares scalar double-precision data and returns the smallest
Examples	

```
    MINSD    XMM0, OWORD DATA3  ;compares numbers in DATA3, least to XMM0
    MINSD    XMM3, XMM6              ;compares numbers in XMM6, least to XMM3
```

MINPS	Compares and returns the smallest packed single-precision number
Examples	

```
    MINPS    XMM0, QWORD DATA3  ;compares numbers in DATA3, least to XMM0
    MINPS    XMM0, XMM2              ;compares numbers in XMM2, least to XMM0
```

MINSS	Compares scalar single-precision numbers and returns the smallest
Examples	

```
    MINSS    XMM0, DWORD DATA3  ;compares numbers in DATA3, least to XMM0
    MINSS    XMM1, XMM2              ;compares numbers in XMM2, least to XMM1
```

LOGIC INSTRUCTIONS

ANDPD	ANDs packed double-precision data

Examples
 	
    ANDPD   XMM0, OWORD DATA3   ;ands DATA3 to XMM0
    ANDPD XMM2, XMM3 ;ands XMM3 to XMM2	

ANDNPD	NANDs packed double-precision data

Examples
 	
    ANDNPD   XMM0, OWORD DATA3   ;Nands DATA3 to XMM0
    ANDNPD XMM4, XMM2 ;Nands XMM2 to XMM4	

ANDPS	ANDs 2 packed single-precision data

Examples
 	
    ANDPS   XMM0, QWORD DATA3   ;ands DATA3 to XMM0
    ANDPS XMM3, XMM2 ;ands XMM2 to XMM3	

ANDNPS	NANDs 2 packed single-precision data

Examples
 	
    ANDNPS   XMM0, DWORD DATA3   ;Nands DATA3 to XMM0
    ANDNPS XMM7, XMM2 ;Nands XMM2 to XMM7	

ORPD	ORs packed double-precision data

Examples
 	
    ORPD   XMM0, OWORD DATA3   ;ors DATA3 to XMM0
    ORPD XMM2, XMM3 ;ors XMM3 to XMM2	

ORPS	ORs 2 packed single-precision numbers

Examples
 	
    ORPS   XMM0, OWORD DATA3   ;ors DATA3 to XMM0
    ORPS XMM3, XMM2 ;ors XMM2 to XMM3	

APPENDIX B

XORPD	Exclusive-ORs packed double-precision data

Examples		
XORPD	XMM0, OWORD DATA3	;exclusive-ors DATA3 to XMM0
XORPD	XMM2, XMM3	;exclusive-ors XMM3 to XMM2

XORPS	Exclusive-ORs packed double-precision data

Examples		
XORPS	XMM0, OWORD DATA3	;exclusive-ors DATA3 to XMM0
XORPS	XMM2, XMM3	;exclusive-ors XMM3 to XMM2

COMPARISION INSTRUCTIONS

CMPPD	Compares packed double-precision numbers

Examples		
CMPPD	XMM0, OWORD DATA3	;compares DATA3 with XMM0
CMPPD	XMM2, XMM3	;compares XMM3 with XMM2

CMPSD	Compares scalar double-precision data

Examples		
CMPSD	XMM0, QWORD DATA3	;compares DATA3 with XMM0
CMPSD	XMM3, XMM2	;compares XMM2 with XMM3

CMPISD	Compares scalar double-precision data and sets EFAGS

Examples		
CMPISD	XMM0, OWORD DATA3	;compares DATA3 with XMM0
CMPISD	XMM2, XMM3	;compares XMM3 with XMM2

UCOMISD	Compares scalar unordered double-precision numbers and changes EFLAGS

Examples

    UCOMISD   XMM0, QWORD DATA3   ;compares DATA3 with XMM0
    UCOMISD   XMM3, XMM2   ;compares XMM2 with XMM3

CMPPS	Compares packed single-precision data

Examples

    CMPPS   XMM0, OWORD DATA3   ;compares DATA3 with XMM0
    CMPPS   XMM2, XMM3   ;compares XMM3 with XMM2

CMPSS	Compares 2 packed single-precision numbers

Examples

    CMPSS   XMM0, QWORD DATA3   ;compares DATA3 with XMM0
    CMPSS   XMM3, XMM2   ;compares XMM2 with XMM3

COMISS	Compares scalar single-precision data and changes EFLAGS

Examples

    COMISS   XMM0, OWORD DATA3   ;compares DATA3 with XMM0
    COMISS   XMM2, XMM3   ;compares XMM3 with XMM2

UCOMISS	Compares unordered single-precision numbers and changes EFLAGS

Examples

    UCOMISS   XMM0, QWORD DATA3   ;compares DATA3 with XMM0
    UCOMISS   XMM3, XMM2   ;compares XMM2 with XMM3

DATA CONVERSION INSTRUCTIONS

SHUFPD	Shuffles packed double-precision numbers

Examples

    SHUFPD   XMM0, OWORD DATA3   ;shuffles DATA3 with XMM0
    SHUFPD   XMM2, XMM2   ;swaps upper and lower quadword in XMM2

UNPCKHPD	Unpacks the upper double-precision number

Examples

    UNPCKHPD XMM0, OWORD DATA3   ;unpacks DATA3 into XMM0
    UNPCKHPD XXM3, XMM2   ;unpacks XMM2 into XMM3

UNPCKLPD	Unpacks the lower double-precision number

Examples

    UNPCKLPD XMM0, OWORD DATA3   ;unpacks DATA3 into XMM0
    UNPCKLPD XMM3, XMM2   ;unpacks XMM2 into XMM3

SHUFPS	Shuffles packed single-precision numbers

Examples

    SHUFPS   XMM0, QWORD DATA3   ;shuffles DATA3 with XMM0
    SHUFPS   XMM2, XMM2   ;swaps upper and lower quadword in XMM2

UNPCKHPS	Unpacks the lower double-precision number

Examples

    UNPCKHPS XMM0, QWORD DATA3   ;unpacks DATA3 into XMM0
    UNPCKHPS XMM3, XMM2   ;unpacks XMM2 into XMM3

UNPCKLPSD	Unpacks the lower double-precision number

Examples

    UNPCKLPSD XMM0, QWORD DATA3   ;unpacks DATA3 into XMM0
    UNPCKLPSD XXM3, XMM2   ;unpacks XMM2 into XMM3

APPENDIX C

Flag-Bit Changes

This appendix shows only the instructions that actually change the flag bits. Any instruction not listed does not affect any of the flag bits.

Instruction	O	D	I	T	S	Z	A	P	C
AAA	?				?	?	*	?	*
AAD	?				*	*	?	*	?
AAM	?				*	*	?	*	?
AAS	?				?	?	*	?	*
ADC	*				*	*	*	*	*
ADD	*				*	*	*	*	*
AND	0				*	*	?	*	0
ARPL						*			
BSF						*			
BSR						*			
BT									*
BTC									*
BTR									*
BTS									*
CLC									0
CLD		0							
CLI			0						
CMC									*
CMP	*				*	*	*	*	*
CMPS	*				*	*	*	*	*
CMPXCHG	*				*	*	*	*	*
CMPXCHG8B						*			
DAA	?				*	*	*	*	*
DAS	?				*	*	*	*	*
DEC	*				*	*	*	*	
DIV	?				?	?	?	?	?
IDIV	?				?	?	?	?	?
IMUL	*				?	?	?	?	*
INC	*				*	*	*	*	

895

Instruction	O	D	I	T	S	Z	A	P	C
IRET	*	*	*	*	*	*	*	*	*
LAR						*			
LSL						*			
MUL	*				?	?	?	?	*
NEG	*				*	*	*	*	*
OR	0				*	*	?	*	0
POPF	*	*	*	*	*	*	*	*	*
RCL/RCR	*								*
REPE/REPNE						*			
ROL/ROR	*								*
SAHF					*	*	*	*	*
SAL/SAR	*				*	*	?	*	*
SHL/SHR	*				*	*	?	*	*
SBB	*				*	*	*	*	*
SCAS	*				*	*	*	*	*
SHLD/SHRD	?				*	*	?	*	*
STC									1
STD		1							
STI			1						
SUB	*				*	*	*	*	*
TEST	0				*	*	?	*	0
VERR/VERW						*			
XADD	*				*	*	*	*	*
XOR	0				*	*	?	*	0

APPENDIX D

Answers to Selected Even-Numbered Questions and Problems

CHAPTER 1

2. Herman Hollerith
4. Konrad Zuse
6. ENIAC
8. Augusta Ada Byron
10. A machine that stores the instructions of a program in the memory system.
12. 200 million
14. 16M bytes
16. 1993
18. 2000
20. Millions of instructions per second
22. A binary bit stores a 1 or a 0.
24. 1024K
26. 1024
28. System area and transient program area
30. 640K
32. 1M
34. 80386, 80486, Pentium, Pentium Pro, PII, PIII, P4, and Core2
36. The basic I/O system
38. The XT was used with the 8088 and 8086 and beginning with the 80286, the AT became the name of the system.
40. 8-bit and 16-bit
42. The advanced graphics port is designed to support video cards.
44. The serial ATA interface is designed to support disk drive memory.
46. 64K
48. See Figure 1–6.
50. Address, data, and control buses.
52. $\overline{\text{MRDC}}$
54. Memory read operation
56. (a) 8-bit signed number (b) 16-bit signed number

(c) 32-bit signed number (d) 32-bit floating-point number (e) 64-bit floating-point number
58. (a) 156.625 (b) 18.375 (c) 4087.109375 (d) 83.578125 (e) 58.90625
60. (a) 10111_2, 27_8, and 17_{16} (b) 1101011_2, 153_8, and 6B (c) 10011010110_2, 2326_8, and $4D6_{16}$ (d) 1011100_2, 134_8, and $5C_{16}$ (e) 10101101_2, 255_8, and AD
62. (a) 0010 0011 (b) 1010 1101 0100 (c) 0011 0100 . 1010 1101 (d) 1011 1101 0011 0010 (e) 0010 0011 0100 . 0011
64. (a) 0111 0111 (b) 1010 0101 (c) 1000 1000 (d) 0111 1111
66. Byte is an 8-bit binary number, word is a 16-bit binary number, doubleword is a 32-bit binary number.
68. Enter is a 0DH and it is used to return the cursor/print head to the left margin of the screen or page of paper.
70. LINE1 DB 'What time is it?'
72. (a) 0000 0011 1110 1000 (b) 1111 1111 1000 1000 (c) 0000 0011 0010 0000 (d) 1111 0011 0111 0100
74. char Fred1 = –34
76. Little endian numbers are stored so the least significant portion is in the lowest numbered memory location and big endian numbers are stored so the most significant part is stored in the lowest numbered memory location.
78. (a) packed = 00000001 00000010 and unpacked 00000001 00000000 00000010 (b) packed = 01000100 and unpacked 00000100 00000100 (c) packed = 00000011 00000001 and unpacked 00000011 00000000 00000001 (d) packed = 00010000 00000000 and unpacked 00000001 00000000 00000000 00000000
80. (a) 89 (b) 9 (c) 32 (d) 1
82. (a) +3.5 (b) −1.0 (c) +12.5

CHAPTER 2

2. 16
4. EBX
6. Holds the offset address of the next step in the program.
8. No, if you add +1 and –1 you have zero, which is a valid number.
10. The I-flag.
12. The segment register addresses the lowest address in a 64K memory segment.
14. (a) 12000H (b) 21000H (c) 24A00H (d) 25000H (e) 3F12DH
16. DI
18. SS plus either SP or ESP
20. (a) 12000H (b) 21002H (c) 26200H (d) A1000H (e) 2CA00H
22. All 16M bytes
24. The segment register is a selector that selects the descriptor from a descriptor table. It also sets privilege level of the request and chooses either the global or local table.
26. A00000H–A01000H
28. 00280000H–00290FFFH
30. 3
32. 64K
34.

0000 0011	1101 0000
1001 0010	0000 0000
0000 0000	0000 0000
0010 1111	1111 1111

36. Through a descriptor stored in the global table
38. The program invisible registers are the cache portions of the segment registers and also the GDTR, LDTR, and IDTR registers.
40. 4K
42. 1024
44. Entry zero or the first entry
46. The TLB caches the most recent memory accesses through the paging mechanism.
50. 1T

CHAPTER 3

2. AL, AH, BL, BH, CL, CH, DL, and DH
4. EAX, EBX, ECX, EDX, ESP, EBP, EDI, and ESI
6. CS, DS, ES, SS, FS, and GS
8. You may not specify mixed register sizes.
10. (a) MOV AL,12H (b) MOV AX,123AH (c) MOV CL,0CDH (d) MOV RAX,1000H (e) MOV EBX,1200A2H

12. Selects an assembly language programming model that contains a single segment that compiles as a .COM program.
14. A label is a symbolic memory address.
16. A label may begin with a letter and some special characters, but not with a number.
18. The .TINY model creates a .COM program.
20. A displacement is a distance and in MOV DS:[2000H],AL the displacement of 2000H is added to the contents of DS times 10H to form the memory address.
22. (a) 3234H (b) 2300H (c) 2400H
24. MOV BYTE PTR [2000H],6
26. MOV DWORD PTR DATA1, 5
28. The MOV BX,DATA instruction copies the word from memory location data into the BX register where the MOV BX,OFFSET DATA instruction copies the offset address of DATA into BX.
30. Nothing is wrong with the instruction; it just uses an alternative addressing style.
32. (a) 11750H (b) 11950H (c) 11700H
34. BP or as an extended version EBP
36.
```
FIELDS      STRUC
F1          DW    ?
F2          DW    ?
F3          DW    ?
F4          DW    ?
F5          DW    ?
FIELDS      ENDS
```
38. Direct, relative, and indirect
40. The intersegment jump allows jumps between segments or to anywhere in the memory system while the intrasegment jump allows a jump to any location within the current code segment.
42. 32
44. Short
46. JMP BX
48. 2
50. AX, CX, DX, BX, SP, BP, DI, and SI in the same order as listed
52. PUSHFD

CHAPTER 4

2. The D-bit indicates the direction of flow for the data (REG to R/M or R/M to REG) and the W-bit indicates the size of the data (byte or word/doubleword).
4. DL
6. DS:[BX+DI]
8. MOV AL,[BX]
10. 8B 77 02
12. The REX prefix, which is used in the 64-bit flat mode, is the register extension that allows the 64-bit registers to be addressed in an instruction.

14.
```
MOV   AX,1000H
MOV   DS,AX
```
16. PUSH RAX
18. AX, CX, DX, BX, SP, BP, SI, and DI
20. (a) AX is copied to the stack. (b) A 32-bit number is retrieved from the stack and placed into ESI. (c) The word contents of the data segment memory location addressed by BX is pushed onto the stack. (d) EFLAGS are pushed onto the stack. (e) A word is retrieved from the stack and placed into DS. (f) A 32-bit number 4 is pushed onto the stack.
22. Bits 24–31 of EAX are stored in location 020FFH, bits 16–23 of EAX are stored into location 020FEH, bits 8–15 of EAX are stored into location 020FDH, and bits 0–7 of EAX are stored into location 020FCH. SP is then decremented by 4 to a value of 00FCH.
24. There are many possible locations, but SP = 0200H and SS = 0200H is one of them.
26. Both instruction load the address of NUMB into DI. The difference is that the MOV DI,OFFSET NUMB assembles as a move immediate and the LEA DI,NUMB assembles as an LEA instruction.
28. The LDS BX,NUMB instruction loads BX with the word stored at data segment memory location NUMB and DS is loaded from the data segment memory location addressed by NUMB+2.
30.
```
MOV   BX,NUMB
MOV   DX,BX
MOV   SI,DX
```
32. CLD clears the direction flag and STD sets the direction flag.
34. The LODSB instruction copies a byte of data from the data segment memory location addressed by SI into the AL register and then increments SI by one if the direction flag is cleared.
36. The OUTSB instruction sends the contents of the data segment memory location addressed by SI to the I/O port address by DX, then SI is incremented by one if the direction flag is cleared.
38.
```
MOV   SI,OFFSET SOURCE
MOV   DI,OFFSET DEST
MOV   CX,12
REP   MOVSB
```
40. XCHG EBX,ESI
42. The LAHF and SAHF instructions in non-64-bit application with the arithmetic coprocessor.
44. The XLAT instruction passes the contents of AL to BX to form an offset address that accesses a memory location whose content is then copied into AL.
46. The OUT DX,AX instruction copies the 16-bit contents of AX into the data segment memory location addressed by the DX register.
48. MOV AH,ES:[BX]

50. An assembly language directive is a special command to the assembler that may or may not generate code or data for the memory.
52. The directives, DB, DW, and DD, are used to define memory as a byte (DB), a word (DW), and a double-word (DD).
54. The EQU (equate) directive allows a memory location to be equated to another memory location.
56. The .MODEL directive specifies the type of memory model used for a program.
58. Full segment definitions
60. PROC indicates the start of a procedure and ENDP indicates its end.
62.
```
STORE  PROC    NEAR
       MOV    [DI],AL
       MOV    [DI+1],AL
       MOV    [DI+2],AL
       MOV    [DI+3],AL
       RET
STORE  ENDP
```
64.
```
COPY   PROC    FAR
       MOV    AX,CS:DATA4
       MOV    BX,AX
       MOV    CX,AX
       MOV    DX,AX
       MOV    SI,AX
       RET
COPY   ENDP
```

CHAPTER 5

2. You cannot use mixed-size registers.
4. AX = 3100H, C = 0, A = 1, S = 0, Z = 0, and O = 0.
6.
```
ADD   AX,BX
ADD   AX,CX
ADD   AX,DX
ADD   AX,SP
```
8.
```
MOV   DI,AX
MOV   R12,RCX
ADD   R12,RDX
ADD   R12,RSI
```
10. INC SP
12. (a) SUB CX,BX (b) SUB DH,0EEH (c) SUB SI,DI (d) SUB EBP,3322H (e) SUB CH,[SI] (f) SUB DX,[SI+10] (g) SUB FROG,AL (h) SUB R10,R9
14.
```
MOV   BX,AX
SUB   BX,DI
SUB   BX,SI
SUB   BX,BP
```
16. The contents of DX and the carry flag are subtracted from the 16-bit contents of the data segment memory addressed by DI – 4 and the result is placed into DX.
18. AH (most significant) and AL (least significant)
20. The O and C flags contain the state of the most significant portion of the product. If the most significant part of the product is zero, then C and O are zero.
22.
```
MOV   DL,5
MOV   AL,DL
MUL   DL
MUL   DL
```
24. BX = DX times 100H

26. AX

28. The errors detected during a division are a divide overflow and a divide by zero.

30. AH

32.
```
MOV  AH,0
MOV  AL,BL
DIV  CL
ADD  AL,AL
MOV  DL,AL
MOV  DH,0
ADC  DH,0
```

34. It divides by AL by 10. This causes numbers between 0 and 99 decimal to be converted to unpacked BCD in AH (quotient) and AL (remainder).

36.
```
PUSH DX
PUSH CX
MOV  CX,1000
DIV  CX
MOV  [BX],AL
MOV  AX,DX
POP  CX
POP  DX
PUSH AX
AAM
MOV  [BX+1],AH
MOV  [BX+2],AL
POP  AX
MOV  AL,AH
AAM
MOV  [BX+3],AH
MOV  [BX+4],AL
```

38. Neither the BCD or the ASCII instructions function in the 64-bit mode.

40.
```
MOV  BH,DH
AND  BH,1FH
```

42.
```
MOV  SI,DI
OR   SI,1FH
```

44.
```
OR   AX,0FH
AND  AX,1FFFH
XOR  AX,0380H
```

46. TEST CH,4

48. (a) SHR DI,3 (b) SHL AL,1 (c) ROL AL,3 (d) RCR EDX,1 (e) SAR DH,1

50. Extra

52. The SCASB instruction is repeated while the condition is equal as long as CX is not zero.

54. CMPSB compares the byte contents of the byte in the data segment addressed by SI with the byte in the extra segment addressed by DI.

56. In DOS the letter C is displayed.

CHAPTER 6

2. A near JMP instruction

4. A far jump

6. (a) near (b) short (c) far

8. The IP or EIP register

10. The JMP AX instruction jumps to the offset address stored in AX. This can only be a near jump.

12. The JMP [DI] instruction jumps to the memory location addressed by the offset address stored in the data segment memory location addressed by DI. The JMP FAR PTR[DI] instruction jumps to the new offset address stored in the data segment memory location addressed by DI and the new segment addressed by the data segment memory location address by DI+2. JMP [DI] is a near jump and JMP FAR PTR [DI] is a far jump.

14. JA tests the condition of an arithmetic or logic instruction to determine if the outcome is above. If the outcome is above a jump occurs, otherwise no jump occurs.

16. JNE, JE, JG, JGE, JL, or JLE

18. JA and JBE

20. SETZ or SETE

22. ECX

24.
```
        MOV  DI,OFFSET DATAZ
        MOV  CX,150H
        CLD
        MOV  AL,00H
L1:     STOSB
        LOOP L1
```

26.
```
        CMP  AL,3
        JNE  @C0001
        ADD  AL,2
@C0001:
```

28.
```
        MOV SI,OFFSET BLOCKA
        MOV DI,OFFSET BLOCKB
        CLD
        .REPEAT
             LODSB
             STOSB
        .UNTIL AL == 0
```

30.
```
        MOV AL,0
        MOV SI,OFFSET BLOCKA
        MOV DI,OFFSET BLOCKB
        CLD
        .WHILE AL != 12H
             LODSB
             ADD AL,[DI]
             MOV [DI],AL
             INC DI
        .ENDW
```

32. A procedure is a reusable group of instructions that ends with a RET.

34. RET

36. By using NEAR or FAR to the right of the PROC directive

38.
```
CUBE  PROC  NEAR USES AX DX
      MOV   AX,CX
      MUL   CX
      MUL   CX
      RET
CUBE  ENDP
```

40.
```
SUMS  PROC  NEAR
      MOV   EDI,0
      ADD   EAX,EBX
      ADD   EAX,ECX
      ADD   EAX,EDX
      ADC   EDI,0
      RET
SUMS  ENDP
```

42. An interrupt is a hardware-initiated function call.

44. INT 0 through INT 255

46. The interrupt vector is used to detect and respond to divide errors.

48. The IRETD instruction is a 32-bit return that pops the return address into EIP.
50. When overflow is a 1
52. CLI and STI
54. If the value in the register or memory location under test in the destination operand is below or above the boundaries stored in the memory address by the source operand.
56. BP

CHAPTER 7

2. No, bytes must be defined in C++ using char.
4. EAX, EBX, ECX, EDX, and ES
6. Floating-point coprocessor stack
8. Data are accessed in array string using register SI to index the string element.
10. If no headers are used for a C++ program, it will be much smaller.
12. No. INT 21H is a 16-bit DOS call that cannot be used in the Windows 32-bit environment.
14.
```
#include "stdafx.h"
#include <conio.h>

int _tmain(int argc, _TCHAR* argv[])
{
        char a = 0;
        while ( a != ?@?)
        {
                a = _getche();
                _putch(a);
        }
        return 0;
}
```
16. The _putch(10) instruction displays the new line function and the _putch(13) returns the cursor to the left margin of the display.
18. Separate assembly modules are the most flexible.
20. The flat model must be used with the C prototype as in .MODEL FLAT,C and the function that is linked to C++ must be made public.
22. A 16-bit word is defined with the short directive.
24. Examples of events are mouse move, key down, etc. Event handlers catch these events so they can be used in a program.
26. Yes. The C++ editor can be used to edit an assembly language module, but the module must use the .TXT extension instead of .ASM.
28. `#define RDTSC _asm _emit 0x0f _asm _emit 0x31`
30.
```
;
;External function rotates a byte 3 places
 left
;
.586                    ;select Pentium and 32-
                        bit model
.model flat, C          ;select flat model with
                        C/C++ linkage
.stack 1024             ;allocate stack space
```

```
.code                   ;start code segment

public RotateLeft3  ;define RotateLeft3 as a
                        public function

RotateLeft3 proc    ;define procedure
Rotatedata:byte     ;define byte

        mov  al,Rotatedata
        rol  al,3
        ret

RotateLeft3 endp
```
32.
```
;Function that converts
;
.model flat,c
.stack 1024
.code

Public Upper

Upper proc
Char:byte

        mov  al,Char
        .if al >= 'a' && a; <= 'z'
                sub al,30h
        .endif
        Ret
Upper endp
```
34. Properties contains information about an object such as the foreground and background colors, etc.
36. `_asm inc ptr;`

CHAPTER 8

2. The TEST.ASM file, when assembled, generates the TEST.OBJ file and the TEST.EXE file if no switches appear on the command line.
4. PUBLIC indicates that a label is available to other program modules.
6. EXTRN
8. MACRO and ENDM
10. Parameters are passed to a macro through a parameter list that follows the MACRO keyword (on the same line).
12. The LOCAL directive defines local labels and must be on the line immediately following the MACRO line.
14.
```
ADDM  MACRO  LIST,LENGTH
      PUSH CX
      PUSH SI
      MOV CX,LENGTH
      MOV SI,LIST
      MOV AX,0
      .REPEAT
          ADD AX,[SI]
          INC SI
      .UNTILCXZ
      POP SI
      POP CX
      ENDM
```
16.
```
private: System::Void textBox1_KeyDown

(System::Object^ sender,System::

Windows::Forms::KeyEventArgs^ e)
```

```
{
        // this is called first
        keyHandled = true;
        if (e->KeyCode >= Keys::D0 &&
            e->KeyCode <= Keys::D9 &&
            e->Shift == false)
        {
            keyHandled = false;
        }
}
```

18.
```
private: System::Void textBox1_KeyDown
(System::Object^ sender,System::Windows::
Forms::KeyEventArgs^ e)
{
    {
        RandomNumber++;  //a global variable
        if ( RandomNumber == 63 )
            RandomNumber = 9;
    }
}
```

20.
```
bool direction;
private: System::Void button1_Click(System
::Object^ sender,System::EventArgs^ e)
{
        label1->Text = "Shift Left = ";
        shift = true;
        data1 = 1;
        label2->Text = "00000001";
        timer1->Start();
}
private: System::Void button1_Click(System::
Object^ sender,System::EventArgs^ e)
{
        label1->Text = "Rotate Left = ";
        shift = false;
        data1 = 1;
        label2->Text = "00000001";
        timer1->Start();
}
private: System::Void radiobutton1_Click
(System::Object^ sender,System::EventArgs^ e)
{
        // left button
        direction = true;   //new bool variable
        if ( shift )
            label1->Text = "Shift Left = ";
        else
            label1->Text = "Rotate Left = ";
}
private: System::Void radiobutton2_Click
(System::Object^ sender,System::EventArgs^ e)
{
        // left button
        direction = false;  //new bool variable
        if ( shift )
            label1->Text = "Shift Right = ";
        else
            label1->Text = "Rotate Left = ";
}
private: System::Void timer1_Tick(System::
Object^ sender,System::EventArgs^ e)
{
        String^ temp = "";
        char temp1 = data1;
        if ( shift )
                if ( direction )
                    _asm shl temp1,1;
                else
                    _asm shr temp1,1;
        else
```

```
            if ( direction )
                    _asm rol temp1,1;
            else
                    _asm ror temp1,1;
        data1 = temp1;
        for (int a = 128; a > 0; a>>=1)
        {
                if ( ( temp1 & a ) == a )
                    temp += "1";
                else
                    temp += "0";
        }
        label2->Text = temp;
}
```

22. The MouseEventArgs in MouseDown contains the Button state that is tested against ::mouses::Mouse Buttons::Right to intercept the right mouse button in a program.

24.
```
private: System::Void Form1_MouseDown
(System::Object^ sender,System::Windows::
Forms::MouseEventArgs^ e)
{
        if (e->Button ==
         ::mouses::MouseButtons::Left &&
         e->Button == ::mouses::MouseButtons::
         Right)
        {
                //left and right
        }
}
```

26. ForeColor sets the color of the text or characters in a control.

28. A large number is converted by repeated divisions by the number 10. After each digit the remainder is saved as a significant digit of the BCD result.

30. 30H

32.
```
int GetOct(void)
{
        String^ temp;
        int result = 0;
        char temp1;
        temp = textBox1-≥Text;
        for ( int a = 0; a < temp->Length; a++ )
        {
                temp1 = temp.[a];
                _asm
                {
                        shl result,3
                        mov eax,0
                        mov al,temp1
                        sub al,30h
                        or result,eax
                }
        }
        return result;
}
```

34.
```
char Up(char temp)
{
        if ( temp >= 'a' && temp <= 'z' )
            _asm sub temp,20h;
        Return temp;
}
```

36. The boot sector is where a bootstrap loader program is located that loads the operating system. The FAT is a table that contains numbers that indicate whether a cluster is free, bad, or occupied. If occupied, an

FFFFH indicates the end of a file chain or the next cluster number in the file chain. The director holds information about a file or a folder.

38. Sectors.

40. A cluster is a grouping of sectors.

42. 4G bytes

44. 8

46. 256

48. `File::Replace(?TEST.LST?, ?TEST.LIS?, ?TEST.BAK?);`
`// Creates TEST.LIS and TEST.BAK, Deletes TEXT.LST`

50. A control is a common object that can be used in any visual programming language.

52. See Figure D–1 for the output (the stock ListBox contains the output).

```
String CPowersDlg::GetNumb(int temp)
{
      char numb[10];
      _asm
      {
            mov   eax,temp
            mov   ebx,10
            push  ebx
            mov   ecx,0
loop1:
            mov   edx,0
            div   ebx
            push  edx
            cmp   eax,0
            jnz   loop1
loop2:
            pop   edx
            cmp   edx,ebx
            je    loop3
            add   dl,30h
            mov   numb[ecx],dl
            inc   ecx
            jmp   loop2
loop3:
            mov   byte ptr numb[ecx],0
      }
```
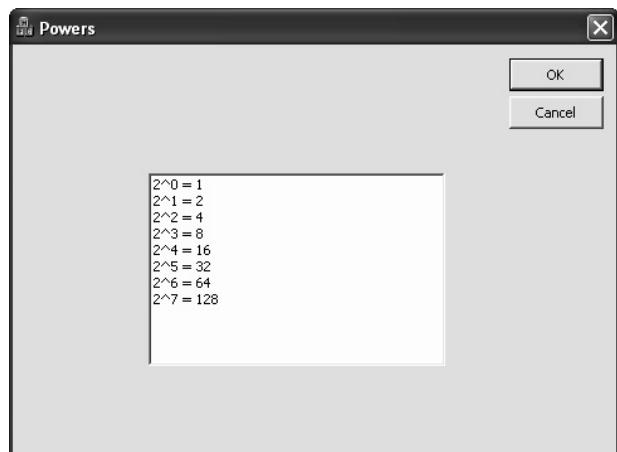

FIGURE D–1

```
      return numb;
}

//code placed in the OnInitDlg function
int tempval = 1;
for ( int a = 0; a < 8; a++ )
{
      CString temp = "2^ = ";
      temp.SetAt(2, a + 0x30);
      temp += GetNumb(tempval);
      List.InsertString(a, temp);
      tempval <<= 1;
}
```

54.
```
private: System::Void Clear()
{
      panel1->Visible = false;
      panel2->Visible = false;
      panel3->Visible = false;
      panel4->Visible = false;
      panel5->Visible = false;
      panel6->Visible = false;
      panel7->Visible = false;
}
private: System::Void Form1_KeyDown(System::
Object^  sender,System::Windows::Forms::
KeyEventArgs^ e)
{
      char lookup[] = {0x3f, 6, 0x5b, 0x4f,
            0x66, 0x6d, 0x7d, 7, 0x7f,0x6f,
            0x77, 0x7c, 0x39, 0x5e, 0x79,
            0x71};
      if (e->KeyCode >= Keys::D0 &&
            e->KeyCode <= Keys::D9)
      {
            ShowDigit(lookup
                  [e->KeyValue -
                  0x30]);  //display
                        the digit
      }
      else if (e->KeyCode >= Keys::A &&
            e->KeyCode <= Keys::F)
      {
            ShowDigit(lookup
                  [e->KeyValue -
                  0x37]);  //display
                        letter
      )
}
private: System::Void ShowDigit(unsigned
char code)
{
      Clear();
      if (( code & 1 ) == 1)  //test a
                        segment
            panel1->Visible = true;
      if (( code & 2 ) == 2)  //test b
                        segment
            panel4->Visible = true;
      if (( code & 4 ) == 4)  //test c
                        segment
            panel5->Visible = true;
      if (( code & 8 ) == 8)  //test d
                        segment
            panel3->Visible = true;
      if (( code & 16 ) == 16)  //test e
                        segment
            panel6->Visible = true;
      if (( code & 32 ) == 32)  //test f
                        segment
            panel7->Visible = true;
      if (( code & 64 ) == 64)  //test g
                        segment
            panel2->Visible = true;
}
```

```
private: System::Void Form1_Load(System::
Object^ sender,System::EventArgs^ e)
{
      Clear();
}
```

CHAPTER 9

2. Yes and no. The current drive of a logic zero is reduced to 2.0 mA and the noise immunity is reduced to 350 mV.
4. Address bits A_0–A_7
6. A read operation
8. The duty cycle must be 33%.
10. A write is occurring.
12. The data bus is sending data to the memory or I/O.
14. IO/$\overline{\text{M}}$, DT/$\overline{\text{R}}$, and $\overline{\text{SS0}}$
16. Signals to the coprocessor, which indicate what the microprocessor queue is doing.
18. 3
20. 14 MHz/6 = 2.33 MHz
22. Address bus connection A_0–A_{15}
24. 74LS373 transparent latch
26. If too many memory and/or I/O devices are attached to a system the buses must be buffered.
28. 4
30. Fetch and execute
32. (a) The address is output along with ALE (b) Time is allowed for memory access and the READY input is sampled. (c) The read or write signal is issued. (d) Data are transferred and read or write is deactivated. (e) Wait allows additional time for memory access.
36. Selects one or two stages of synchronization for READY
38. Minimum mode operation is most often used in embedded applications and maximum mode operation was most often used in early personal computers.

CHAPTER 10

2. (a) 256 (b) 2K (c) 4K (d) 8K (e) 1M
4. Select the memory device
6. Cause a write to occur
8. The microprocessor allows 460 ns for memory at 5 MHz, but because there is a small delay in the connections to the memory, it would be best not to use a 450 ns memory device in such a system without one wait state.
10. Static random access memory
12. 250 ns
14. The address inputs to many DRAMs are multiplexed so one address input accepts two different address bits, reducing the number of pins required to address memory in a DRAM.
16. Generally the amount of time is equal to a read cycle and represents only a small amount of time in a modern memory system.
18. See Figure D–2.
20. One of the eight outputs becomes a logic zero as dictated by the address inputs.
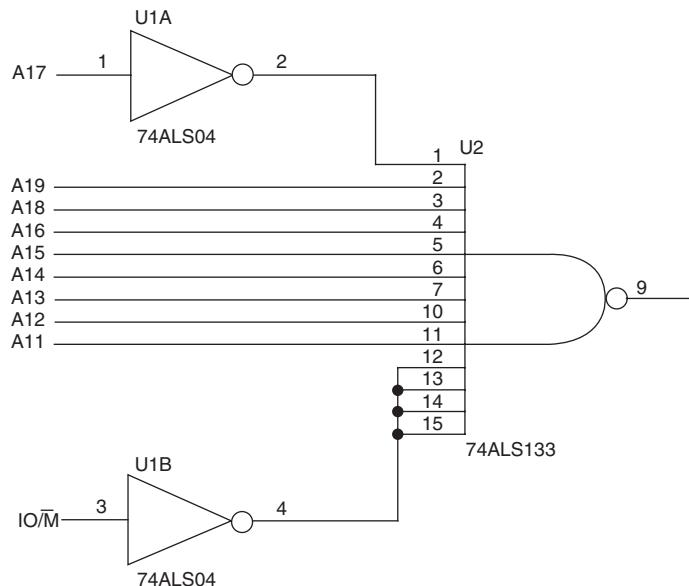
FIGURE D–2
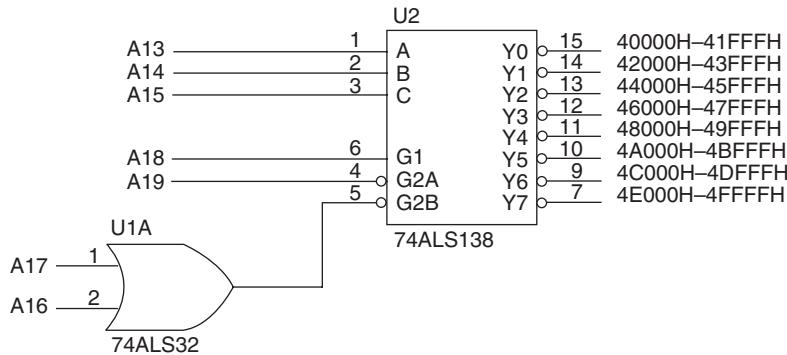
FIGURE D–3
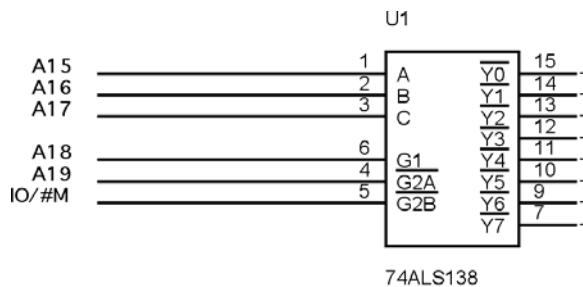
FIGURE D–4

22. See Figure D–3.
24. Verilog hardware description language
26. The architecture block between begin and end
28. $\overline{MRDC}$ and $\overline{MWTC}$
30. See Figure D–4.
32. 5
34. 1
36. $\overline{BHE}$ selects the upper memory bank and A_0 selects the lower memory bank.
38. Separate decoders and separate write signals
40. Lower memory bank
42.
```
library ieee;
use ieee.std_logic_1164.all;
entity DECODER_10_28 is
port (
      A23, A22, A21, A20, A19, A18, A17,
      A16, A0, BHE, MWTC: in STD_LOGIC;
      SEL, LWR, HWR: out STD_LOGIC
);

end;

architecture V1 of DECODER_10_28 is

begin

     SEL <= A23 or A22 or A21 or A20 or A19
     or A18 or (not A17) or (not A16);
     LWR <= A0 or MWTC;
     HWR <= BHE or MWTC;

end V1;
```
44. See Figure D–5.

48. Yes, as long as a memory location on the DRAM is not accessed.
50. 128 bits wide

CHAPTER 11

2. The I/O address is stored in the second byte of the instruction.
4. DX
6. The OUTSB instruction transfers the data segment byte addressed by SI to the I/O port addressed by DX, then SI is incremented by one.
8. Memory mapped I/O uses any instruction that transfers data to or from the memory for I/O, while isolated I/O requires the use of the IN or OUT instruction.
10. The basic output interface is a latch that captures output data and holds it for the output device.
12. Lower
14. 4
16. It removes mechanical bounces from a switch.
18. See Figure D–6.
20. See Figure D–7.
22. See Figure D–8.
24. If the port is 16 bits wide, there is no need to enable either the low or high half.
26. D_{47}–D_{40}
28. Group A is port A and PC_4–PC_7, while group B is port B and PC_3–PC_0.
30. $\overline{RD}$
32. Inputs
34. The strobe input latches the input data and sets the buffer full flag and interrupt request.
36.
```
DELAY  PROC  NEAR USES ECX

       MOV ECX, 7272727
D1:
       LOOPD D1
       RET

DELAY  ENDP
```
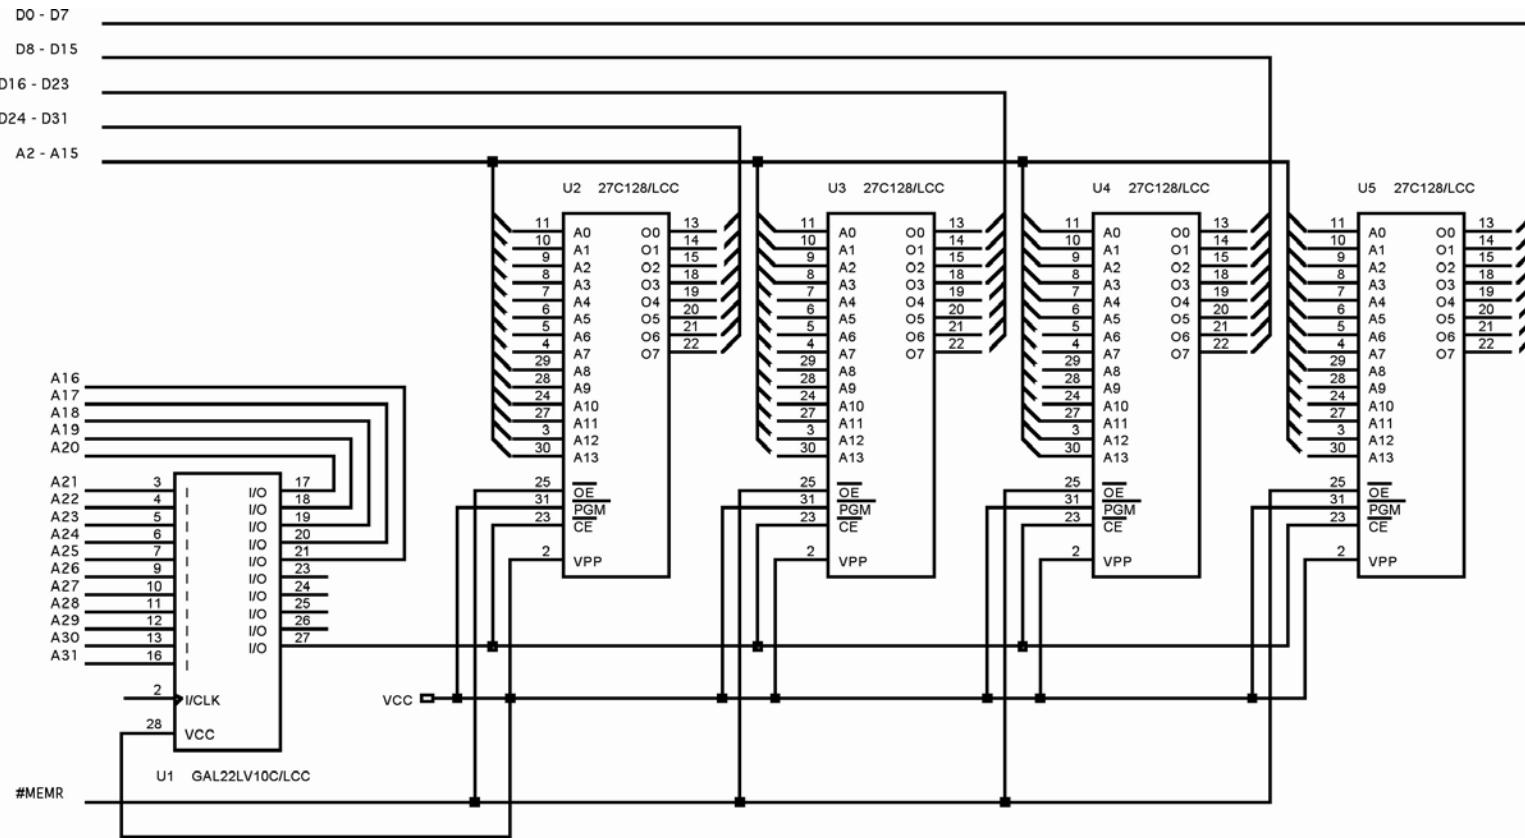
FIGURE D–5
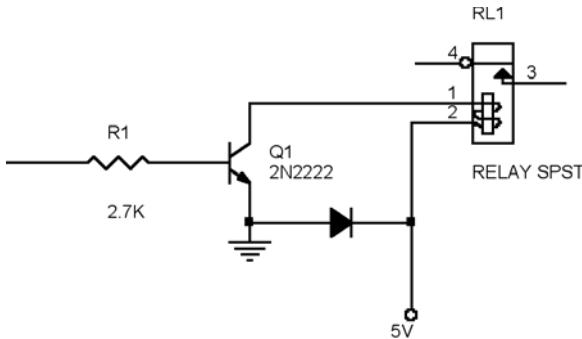
FIGURE D–6
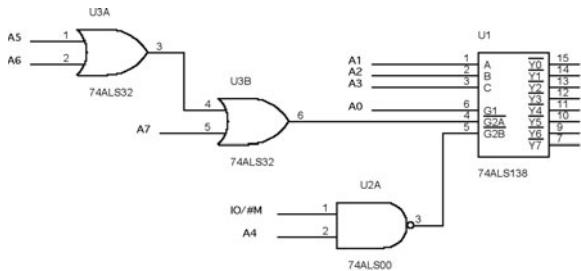
FIGURE D–7
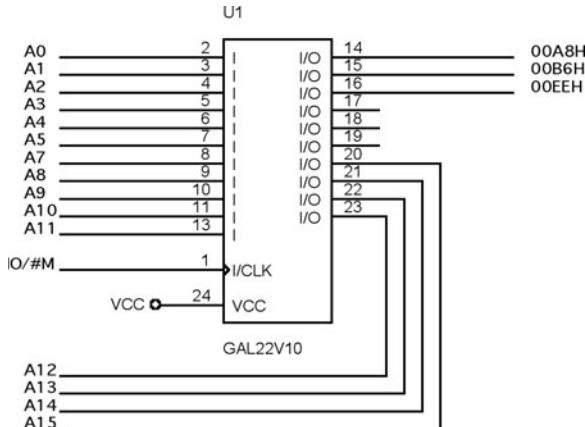
FIGURE D–8

38. The strobe signal ($\overline{STB}$)
40. The INTR pin is enabled by setting the INTE bit in PC_4 (port A) or PC_2 (port B).
42. When data are output to the port $\overline{OBF}$ becomes a 0 and when $\overline{ACK}$ is sent to the port $\overline{OBF}$ becomes a 1.
44. Group or port A contains the bidirectional data.
46. The 01H command is sent to the LCD display.
48.
```
;Displays the null terminated string addressed
 by DS:BX
```

```
;uses a macro called SEND to send data to
the display
;
DISP     PROC  NEAR  USES BX
         SEND  86H,2,1      ;move cursor to
                             position 6
         .WHILE BYTE PTR [BX] != 0
             SEND [BX],0,1
             INC  BX
         .ENDW
         RET
DISP     ENDP
```

50. The only changes that need to be made are that instead of four rows there are three rows and three pull-up resisters connected to port A and five columns to connect to port B. Of course, the software also needs some minor changes.
54. 6
58. Least significant
62. Data that are sent a bit at a time without any clocking pulses
64.
```
LINE   EQU   023H
LSB    EQU   020H
MSB    EQU   021H
FIFO   EQU   022H

       MOV   AL,10001010B ;enable baud
divisor
       OUT   LINE,AL

       MOV   AL,60         ;program baud rate
       OUT   LSB,AL
       MOV   AL,0
       OUT   MSB,AL

       MOV   AL,00011001B  ;program 7-data, odd
       OUT   LINE,AL       ;parity, one stop

       MOV   AL,00000111B  ;enable transmitter
                            and
       OUT   FIFO,AL       ;and receiver
```

66. Simplex = receiving or sending data; half-duplex = receiving and sending data; but only one direction at a time; and full-duplex = receiving and sending data at the same time.
68.
```
SENDS  PROC  NEAR

MOV    CX,16
       .REPEAT
         .REPEAT
           IN    AL,LSTAT ;get line
                           status register
           TEST  AL,20H    ;test TH bit
         .UNTIL !ZERO?
         LODSB             ;get data
         OUT  DATA,AL      ;transmit data
       .UNTILCXZ

   RET

SENDS ENDP
```

70. 0.01V
72.
```
       .MODEL  TINY
       .CODE
       .STARTUP
           MOV   DX,400H
           .WHILE  1
```

```
         MOV    CX,256
         MOV    AL,0
         .REPEAT
             OUT    DX,AL
             INC    AL
             CALL   DELAY
         .UNTILCXZ
         MOV    CX,256
         .REPEAT
             OUT    DX,AL
             DEC    AL
             CALL   DELAY
         .UNTILCXZ
         .ENDW
DELAY          PROC    NEAR

; 39 microsecond time delay

DELAY    ENDP
         END
```

74. INTR indicates that the converter has completed a conversion.

76. See Figure D–9.

CHAPTER 12

2. An interrupt is a hardware- or software-initiated subroutine call.

4. Interrupts only use computer time when the interrupt is activated.

6. INT, INT3, INTO, CLI, and STI

8. The first 1K byte of the memory system in real mode and anywhere in protected mode.

10. 00H through 1FH

12. Anywhere in the memory system

14. A real mode interrupt pushes CS, IP, and the FLAGS onto the stack, while a protected mode interrupt pushes CS, EIP, and the EFLAGS onto the stack.

16. The INTO occurs if overflow is set.

18. The IRET instruction pops the flags and the return address from the stack.

20. The state of the interrupt structure is stored on the stack, so when the return occurs, it is restored. Both the interrupt and trace flags are cleared.

22. The IF flag controls whether the INTR pin is enabled or disabled.

24. The IF flag is enabled or disabled by using the STI or CLI instructions.

26. 2

28. Level sensitive

30. Vector

32. See Figure D–10.

34. The pull-up resistors guarantee that vector number fetched from the data bus during an interrupt acknowledge is an FFH.

36. Because the interrupt request signal (INTR) is generated by ORing all the requests together, the software must ask or poll each device to determine which device caused the request.

38. 9

40. The CAS pins are cascade pins used to cascade 8259s for more than eight interrupt inputs.

42. An OCW is an operational control word for the 8259.

44. ICW_2
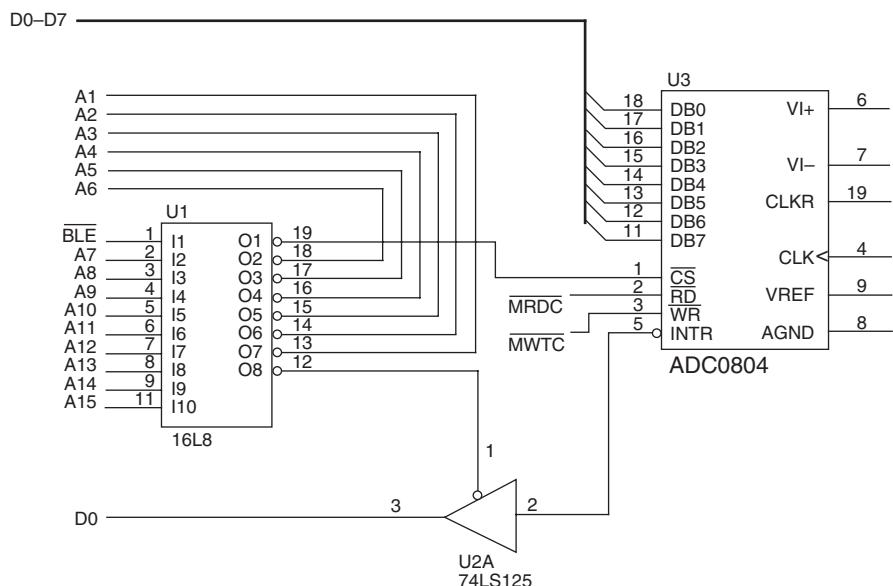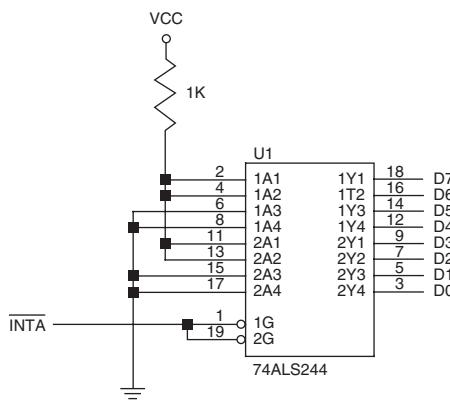
FIGURE D–9

FIGURE D–10

46. Program sensitivity and single or multiple 8259s
48. The most recent interrupt request level becomes the lowest level interrupt after being serviced.
50. INT 8 through INT 0FH

CHAPTER 13

2. When a 1 is placed on HOLD, the program stops executing and the address, data, and control buses go to their high-impedance state.
4. I/O to memory
6. DACK
8. The microprocessor is in its hold state and the DMA controller has control of the buses.
10. 4
12. The command register
16. A pen drive is a USB device that acts as a storage device using a flash memory.
18. Tracks
20. Cylinder
22. See Figure D–11.
24. The heads in a hard disk drive are aerodynamically designed to ride on a cushion of air as the disk spins and are therefore called flying heads.
26. The stepper motor positioning mechanism is noisy and not very precise, while the voice coil positioning mechanism is silent and very accurate because its placement can be continuously adjusted.

28. A CD-ROM is an optical device for storing music or digital data and has a capacity of about 660M or 700M (80 minute) bytes.
30. A TTL monitor uses TTL signals to generate a display and an analog monitor uses analog signals.
32. Cyan, magenta, and yellow
34. 1024 lines with 1280 horizontal elements per line
36. The DVI-D and HDMI connectors are the latest style of digital video input connectors for all types of video equipment.
38. 16 million colors

CHAPTER 14

2. Word (16-bits, ±32K), doubleword (32-bits, ±2G), and quadword (64-bits, $\pm 9 \times 10^{18}$)
4. Single-precision (32 bits), double-precision (64 bits), and temporary-precision (80 bits)
6. (a) −7.75 (b) .5625 (c) 76.5 (d) 2.0 (e) 10.0 (f) 0.0
8. The microprocessor continues executing microprocessor (integer) instructions while the coprocessor executes a floating-point instructions.
10. It copies the coprocessor status register to AX.
12. By comparing the two registers and then by transferring the status word to the AX register. If the SAHF instruction is next executed, a JZ instruction can be used to test the outcome of the coprocessor compare instruction.
14. FSTSW AX
16. Data are always stored as an 80-bit temporary precision number.
18. 0
20. Affine allows positive and negative infinity, while projective assumes infinity is unsigned.
22. Extended (temporary) precision
24. The contents of the top of the stack are copied into memory location DATA as a floating-point number.
26. FADD ST,ST(3)
28. FSUB ST(2),ST
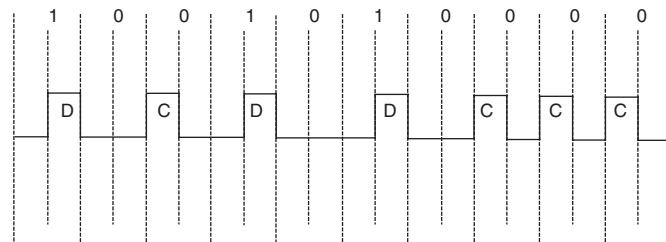
FIGURE D–11

30. Forward division divides the top of the stack by the contents of a memory location and returns the quotient to the top of the stack. Reverse division divides the top of the stack into the contents of the memory location and returns the result to the top of the stack. If no operand exists, then forward division divides ST(1) by ST and reverse division divides ST by ST(1).

32. It performs a MOV to ST if the condition is below.

34.
```
RECIP   PROC    NEAR
        MOV     TEMP,EAX
        FLD     TEMP
        FLD1
        FDIVR
        FSTP    TEMP
        MOV     EAX,TEMP
RECIP   ENDP
TEMP    DD  ?
```

36. Finds the function $2^X - 1$.

38. FLDPI

40. It indicates that register ST(2) is free.

42. The state of the machine

44.
```
CAPR    PROC    NEAR
        FLDPI
        FADD    ST,ST(1)
        FMUL    F
        FMUL    C1
        FLD1
        FDIVR
        FTSP    XC
        RET
CAPR    ENDP
```

46. In modern software it is never used.

48.
```
TOT     PROC    NEAR
        FLD     R2
        FLD1
        FDIVR
        FLD     R3
        FLD1
        FDIVR
        FLD     R4
        FLD1
        FDIVR
        FADD
        FADD
        FLD1
        FDIV
        FADD    R1
        FSTP    RT
        RET
TOT     ENDP
```

50.
```
PROD    PROC    NEAR
        MOV     ECX,100
        .REPEAT
            FLD     ARRAY1[ECX*8-8]
            FUML    ARRAY2[ECX*8-8]
            FSTP    ARRAY3[ECX+8-8]
        .UNTILCXZ
        RET
PROD    ENDP
```

52.
```
POW     PROC     NEAR
        MOV     TEMP,EBX
        FLD     TEMP
        F2XM1
        FLD1
        FADD
        MOV     TEMP,EAX
```

```
        FLD     TEMP
        FYL2X
        FSTP    TEMP
        MOV     ECX,TEMP
        RET
POW     ENDP
```

54.
```
GAIN    PROC     NEAR
        MOV     ECX,100
        .REPEAT
            FLD     DWORD PTR VOUT[ECX*4-4]
            FDIV    DWORD PTR VIN[ECX*4-4]
            CALL    LOG10
            FIMUL   TWENTY
            FSTP    DWORD PTR DBG[ECX*4-4]
        .UNTILCXZ
        RET
TWENTY  DW      20
GAIN    ENDP
```

56. The EMMS instruction clears the coprocessor stack to indicate that the MMX unit has completed using the stack.

58. Signed saturation occurs when byte-sized numbers are added and have values of 7FH for an overflow and 80H for an underflow.

60. The FSAVE instruction stores all the MMX registers in memory.

62. Single-instruction, multiple-data instructions

64. 128 bits

66. 16

68. Yes

CHAPTER 15

2. 8- or 16-bit depending on the socket configuration.

4. See Figure D–12.

6. See Figure D–13.

8. See Figure D–14.

12. 16 bits

14. The configuration memory identifies the vendor and also information about the interrupts.

16. This is the command/bus enable signal that is high to indicate the PCI bus contains a command and low for data.

18.
```
MOV   AX,0B108H
MOV   BX,0
MOV   DI,8
INT   1AH
```

20. 2.5 GHz

22. Yes

24. COM_1

30. 1.5 Mbps, 12 Mbps, and 480 Mbps

32. 5 meters

34. 127

36. An extra bit that is thrown in the data stream if more than six ones are sent in a row.

38. 1 to 1023 bytes

40. The PCI transfers data at 33 MBs, while AGP transfers data at 2 GBps (8 ×).

CHAPTER 16

2. The hardware enhancements include internal timers, additional interrupt inputs, chip selection logic, serial communications ports, parallel pins, DMA controller, and an interrupt controller.

4. 10 MHz

6. 3 mA

8. The point at which the address appears

10. 260 ns for the 16 MHz version operated at 10 MHz

12.
```
MOV  AX,1000H
MOV  DX,0FFFEH
OUT  DX,AX
```

14. 10 on most versions of the 80186/80188 including the internal interrupts.
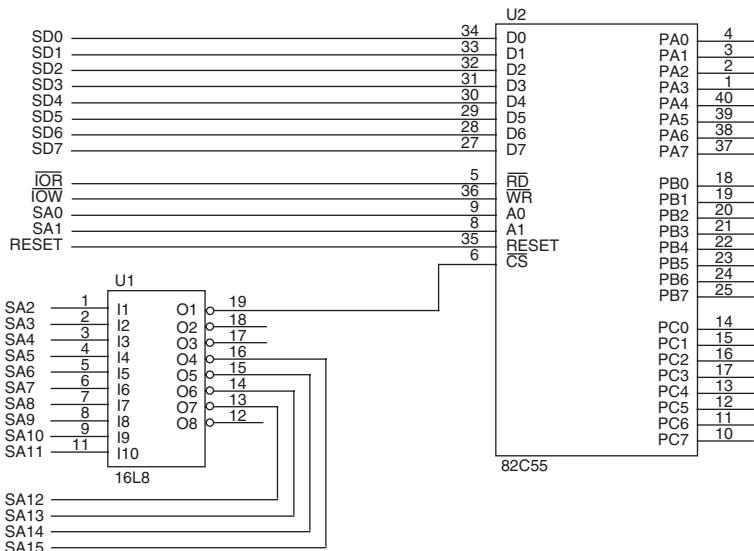
FIGURE D–12
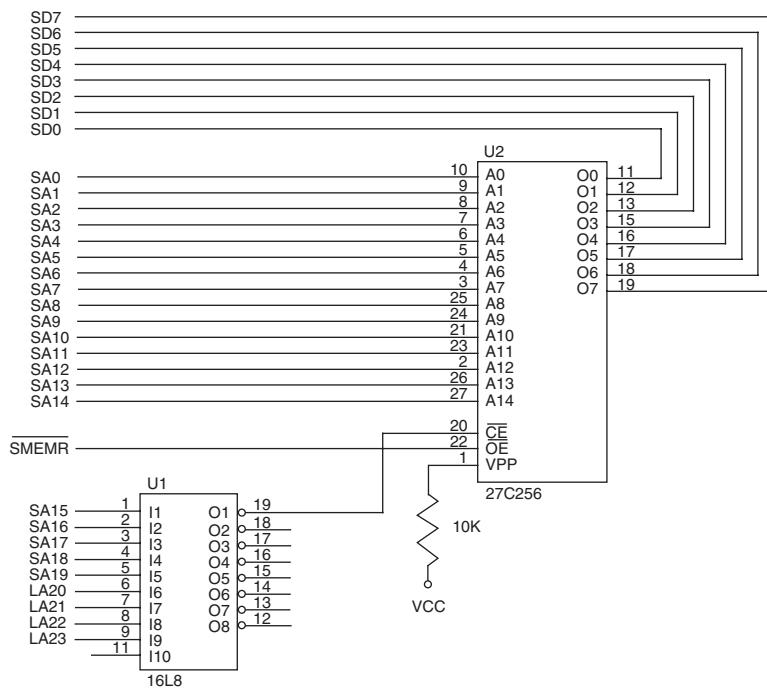
FIGURE D–13
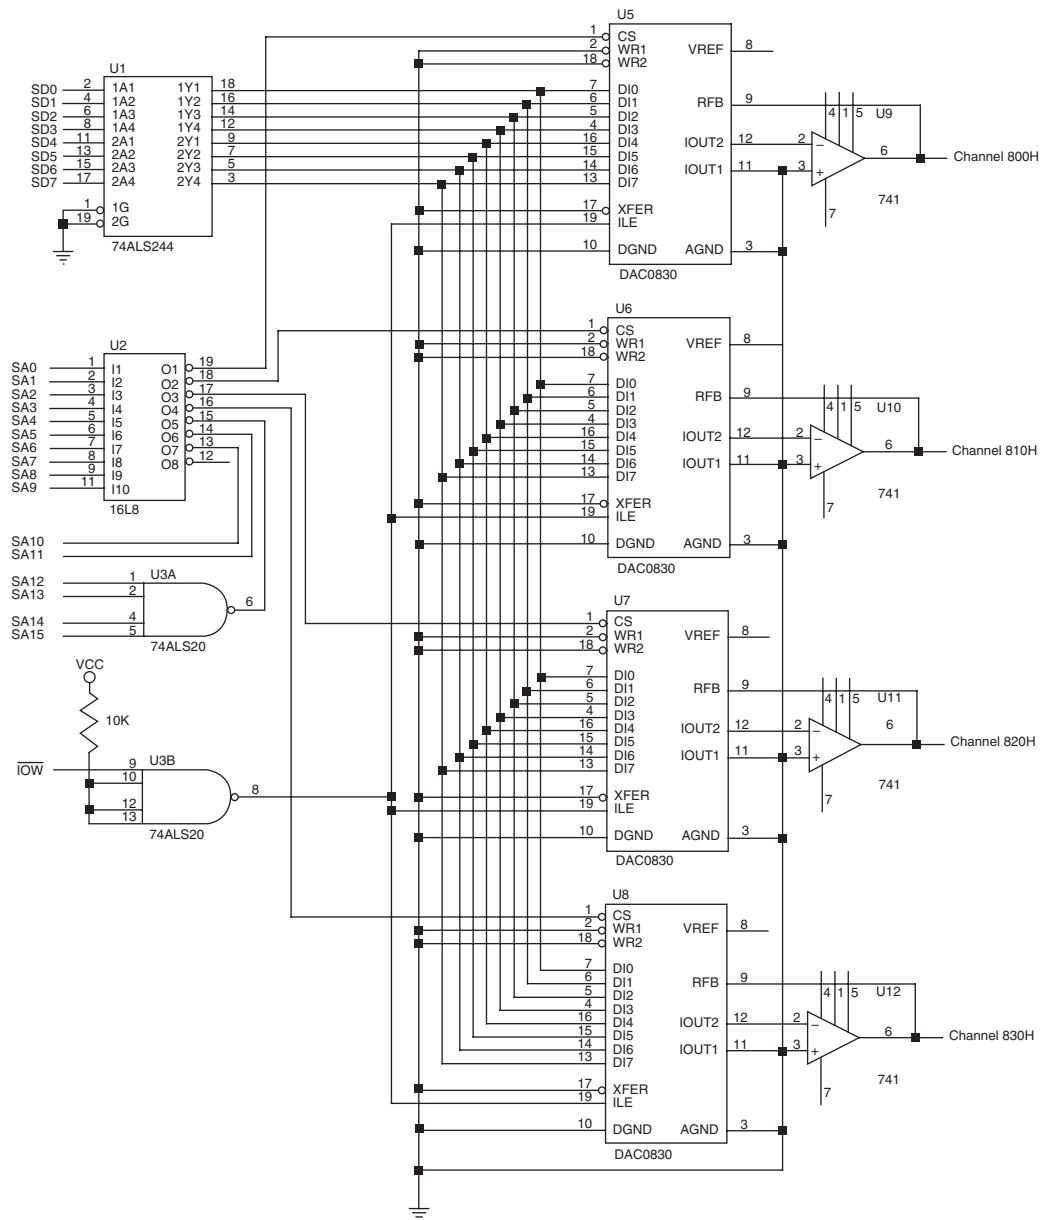
FIGURE D–14

16. The interrupt control registers control a single interrupt.
18. The interrupt poll register acknowledges the interrupt, while the interrupt poll status register does not acknowledge the interrupt.
20. 3
22. Timer 2
24. It determines whether the enable counter bit functions.
26. The ALT bit selects both compare registers so the duration of the logic 1 and logic 0 output times can be programmed.

28.
```
MOV  AX,123
MOV  DX,0FF5AH
OUT  DX,AX
MOV  AX,23
ADD  DX,2
OUT  DX,AX
MOV  AX,0C007H
MOV  DX,0FF58H
OUT  DX,AX
```

30. 2
32. Place a logic 1 in both the CHG/$\overline{\text{NOCHG}}$ and START/$\overline{\text{STOP}}$ bits of the control register.
34. 7

36. Chip
38. 15
40. It determines the operation of the $\overline{PCS5}$ and $\overline{PCS6}$ pins.
42.
```
MOV  AX,1001H
MOV  DX,0FF90H
OUT  DX,AX
MOV  AX,1048H
OUT  DX,AX
```
44. 1G
46. Verify for read access.
48. An RTOS is a real-time operating system that has a predictable and guaranteed time for threads access.

CHAPTER 17

2. 64T
4. See Figure D–15.
6. The memory system has up to 4G bytes and the bank enable signals select one or more of the 8-bit-wide banks of memory.
8. The pipeline allows the microprocessor to send the address of the next memory location, while it fetches the data from the prior memory operation. This allows the memory additional time to access the data.
10. 0000H–FFFFH
12. I/O has the same address as earlier models of the microprocessor. The difference is that the I/O is arranged as a 32-bit-wide space with four 8-bit banks that are selected by the bank enable signals.
14. The $\overline{BS16}$ pin causes the microprocessor to function with an 8-bit-wide data bus.
16. The first four debug registers (DR_0–DR_3) contain breakpoint addresses; registers DR_4 and DR_5 are reserved for Intel's use; DR_6 and DR_7 are used to control debugging.
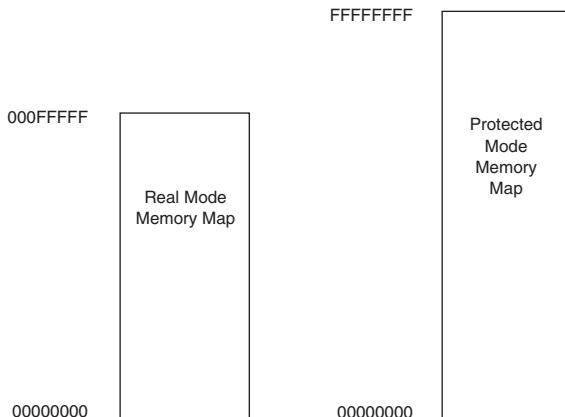
FIGURE D–15

18. The test registers are used to test the translation look-aside buffer.
20. The PE bit switches the microprocessor into protected mode if set and real mode if cleared.
22. Scaled-index addressing used a scaling factor of 1, 2, 4, or 8 times to scale addressing from byte, word, doubleword, or quadword.
24. (a) the address in the data segment at the location pointed to by EBX times 8 plus ECX (b) the address in the data segment array DATA pointed to by the sum of EAX plus EBX (c) the address at data segment location DATA (d) the address in the data segment pointed to by EBX
26. Type 13 (0DH)
28. The interrupt descriptor table and its interrupt descriptors
30. A selector appears in a segment register and it selects a descriptor from a descriptor table. It also contains the requested privilege level of the request.
32. The global descriptor table register
34. Because a descriptor addresses up to 4G of memory and there are 8K local and 8K global descriptor available at a time, 4G times 16K = 64T.
36. The TSS holds linkages and registers of a task so tasks can be switched efficiently.
38. The switch occurs when a logic 1 is placed into the PE bit of CR_0.
40. Virtual mode, which simulates DOS in protected mode, sets up 1M memory spans that can operate in the real mode.
42. 4K
44. The 80486 has an internal 8K cache and also contains a coprocessor.
46. The register sets are virtually identical.
48. $\overline{PCHK}$ and DP_0–DP_3
50. 8K
52. A burst is when four 32-bit numbers are read or written between the cache and memory.
54. Built-in self test

CHAPTER 18

2. 64G bytes
4. These pins generate and check parity.
6. The burst ready pin is used to insert wait state into the bus cycle.
8. 18.5 ns
10. T_2
12. An 8K byte data cache and an 8K-byte instruction cache.
14. Yes, if one is a coprocessor instruction and the integer instructions are not dependent.

16. The SSM mode is used for power management in most systems.
18. 38000H
20. The CMPXCH8B instruction compares the 64-bit number in EDX:EAX with a 64-bit number stored in memory. If they are equal, ECX:EBX is stored in memory. If not equal, the contents of memory are moved into EDX:EAX.
22. ID, VIP, VIF, and AC
24. To access 4M pages, the page tables are dropped and only the page directory is used with a 22-bit offset address.
26. The Pentium Pro is an improved version of the Pentium that contains three integer units, an MMX unit, and a 36-bit address bus.
28. 36 address bits on A_3 through A_{35} (A_0–A_2 are encoded in the bank selection signals)
30. The access time in a 66 MHz Pentium is 18.5 ns and in the Pentium Pro at 66 MHz access time is 17 ns.
32. SDRAM that is 72 bits wide is purchased for ECC memory applications instead of 64-bit-wide memory.

CHAPTER 19

2. 512K, 1M, or 2M
4. The Pentium Pro cache is on the main board and the Pentium 2 cache is in the cartridge and operates at a high speed.
6. 64G bytes
8. 242

10. The read and write signals are developed by the chip set instead of the microprocessor.
12. 8 ns after the first quadword is accessed. The first quadword still requires 60 ns for access.
14. Model-specific registers have been added for SYSENTER_CS, STSENTER_SS, and SYSENTER_ESP.
16. The ECX register address the MSR number when the RDMSR instruction executes. After execution, EDX:EAX contains the contents of the register
18.
```
TESTS   PROC    NEAR
        CPUID
        BT     EDX,800H
        RET
TESTS   ENDP
```
20. EDX to the EIP register and the value in ECX to the ESP register.
22. Ring 3
24. Pentium Pro
26. The Pentium 4 or Core2 requires a power supply with an additional 12 V connector for the main board. A Pentium 4–compliant supply must be used.
28.
```
bool Hyper()
{
    _asm
    {
        bool State = true;
        mov  eax,1
        cpuid
        mov  temp1,31h
        bt   edx,28      ;check for hyper-
                         threading
        jc   Hyper1
        mov  State, 0
Hyper1:
    }
    return State;
}
```

INDEX